QICN: A Rigid Identity Framework

Johnny Andrey Pérez Ramírez

2026

Editorial Preface and Claim Boundary

This monolithic volume is an editorial consolidation of the QICN rigid-identity corpus. BaseCore is treated as the foundational layer of the framework rather than as Paper 0. The downstream works are indexed in the table of contents by their direct research titles; internal labels such as paper numbers are retained only as provenance handles where the source text itself requires them.

The claims in this volume remain conditional, formal, and operational. The volume does not assert biological consciousness, phenomenal consciousness, agency, moral status, identity transfer, or external validation. Runtime measurements, finite estimators, and internal support artifacts are implementation evidence only when explicitly marked as such; they are not substitutes for theorem proof or externally reproduced empirical adjudication.

The editorial policy is conservative: preserve source content, expose dependency order, avoid title-level bookkeeping labels, and keep ontology, mathematical model, implementation, language, and interpretation separated.

Contents

Editorial Preface and Claim Boundary ii

BaseCore 1

I Foundational Dynamics 2
- B1 Scope, System Boundary, and Non-Inference Note . . . 3
- B2 Introduction . . . 3
  - B2.1 Purpose . . . 3
  - B2.2 Methodology . . . 3
  - B2.3 Scope discipline . . . 3
- B3 Minimal Base Hypotheses . . . 3
- B4 The Projection Operator . . . 4
- B5 The Transition Operator . . . 4
- B6 Existence and Compactness of Attractors . . . 5
- B7 Parameterwise Non-Collapsibility . . . 5
- B8 Summary of Base Results . . . 6

II Typed Computable Model and State Spectral Gap 7
- B9 Typed Computable Model: Embedded Constant-Disk Witness . . . 8
  - B9.1 Golden Mean Parameter Space . . . 8
  - B9.2 Typed Banach Witness . . . 8
  - B9.3 Basic Operator Properties . . . 9
  - B9.4 Separate Anti-Collapse Witness . . . 10
- B10 State Spectral Gap and Parameter-Family Persistence . . . 11

III Identity, Rigidity, and Conditional Non-Simulability 13
B11 Projective Systems and Inverse-Limit Identity . . . 14
B11.1 Projective Setup . . . 14
B11.2 Non-Locality of Identity . . . 14
B12 Rigidity under Admissible Perturbations . . . 15
B12.1 Metric Projective Assumptions . . . 15
B13 Conditional Non-Simulability . . . 17
B14 Section Summary . . . 18

IV Regime Structure and Continuity Constraints 19
B15 Formal Setup . . . 20
B15.1 Review of Identity Framework . . . 20
B15.2 Abstract Phenomenological Space . . . 20
B15.3 Phenomenological Assignment . . . 20
B15.4 Axioms for Phenomenological Coherence . . . 21
B15.5 Status of the Assignment Conditions . . . 21
B15.6 Logical Dependencies and Realizability . . . 22
B15.7 Rejected Alternatives . . . 22
B15.8 Structural Interpretation . . . 22
B15.9 Closure of Section 2 . . . 23
B15.10 Phenomenological Neutrality Clause . . . 23
B15.11 Exact Compatibility Obstruction . . . 23
B15.12 A Nonempty Controlled Assignment Model . . . 23
B15.13 Common Readout on Two Coherently Transported Towers . . . 24
B15.14 Controls Separating Convergence, Continuity, and Invariance . . . 25
B16 Compatibility with Identity Regimes . . . 26
B16.1 Induced Constraints from Identity Rigidity . . . 26
B16.2 Metric Structure for Quantitative Comparisons . . . 26
B16.3 Phenomenological Continuity . . . 27
B16.4 Phenomenological Fragmentation . . . 27
B16.5 Relationship Between Metrics . . . 27
B16.6 Closure of Section 3 . . . 27
B17 Impossibility Theorems . . . 28
B17.1 Φ-Regularity Hypothesis . . . 28
B17.2 Fragmentation Theorem . . . 28
B17.3 Forced Continuity Theorem . . . 29
B17.4 Dichotomy Theorem . . . 29
B18 Classification of Phenomenological Regimes . . . 29
B18.1 Conditional Classification . . . 29
B18.2 Structural Implications . . . 30
B19 Extended Proofs . . . 30
B19.1 Proof of Fragmentation Theorem (Extended) . . . 30
B19.2 Proof of Forced Continuity Theorem (Extended) . . . 30
B19.3 Structural Necessity Lemma . . . 30

V Null-Regime Structural Exclusion 31
B20 Compatibility, Witnesses, and Non-Nullity . . . 32
B20.1 Compatibility Operator . . . 32
B20.2 Witness-Relative Null Separation . . . 32
B20.3 Positive Regime Structure . . . 33

B21 Canonical Families and Structural Closure Consequences 34

VI BaseCore Boundary and Expansion Interface 36
B22 BaseCore Ownership and Downstream Boundary 37
B23 Admissible Structural Extensions 37

VII Operational-Criterion Grammar and Certification 39
B24 Admissible Systems, Supports, and Histories 40
B24.1 System Model 40
B24.2 Admissible Support 40
B24.3 Operational Histories 40
B24.4 Witness Margin Vector 41
B25 Six Critical Invariants 41
B25.1 Structural Persistence 41
B25.2 Rigid Identity 41
B25.3 Causal Integration 41
B25.4 Regime Continuity 42
B25.5 Non-Null Differentiation 42
B25.6 Operational Legibility 42
B26 Criterion Class, Certified Partition, and Rupture 43
B27 Minimality of the Invariant Basis 44
B28 Operational Legibility as a Certified Object 45
B29 Certified Candidate Audit Schema 45

VIII Claim Boundary and Falsation Grammar 47
B30 Forensic Admissibility and Integrity Grammar 48
B31 Certification, Comparison, and Rupture 48
B32 Prediction and Falsation Grammar 48
B33 Claim Boundary Ledger 49
B33.1 Allowed BaseCore Claims 49
B33.2 Forbidden BaseCore Claims 49
B34 Dependency Ledger 49
B35 Assumption-to-Theorem-to-Non-Claim Mapping 50
B36 Theorem Hygiene Ledger 50
B37 Redundancy Ledger 51
B38 Open Assumptions Inventory 51
B39 Downstream Boundary Note 51

IX Discrete-to-Continuous Bridge 52
B40 From Node Networks to Hilbert Space Dynamics 53
B41 Projective Dissipation on the Circle-Prefix Towers 57
Appendix 59
B42 Counterexamples for Necessity 59
B42.1 Without H1 (Convexity) 59
B42.2 Without H2 ($\|K\| < 1$) 59
B42.3 Without H3 (Completeness) 59
B42.4 Without H4 (Compactness) 59
B42.5 Why Φ-Regularity Is a Hypothesis 60

1 Rigid Identity as an Inverse Limit in Observable Channels 61
1.1 Scope, System Boundary, and Non-Inference Note . . . 61
1.2 Introduction . . . 61
1.2.1 Motivation . . . 62
1.2.2 Ontological Closure . . . 62
1.3 Minimal Assumptions and Framework . . . 62
1.3.1 Assumption A1: Observable Channel . . . 62
1.3.2 Assumption A2: Temporal Persistence . . . 63
1.3.3 Assumption A3: Non-Triviality (R0) . . . 63
1.3.4 Categorical Regularity: Topological Structure of State Spaces . . . 63
1.3.5 Conditional Nature of Results . . . 64
1.4 Inverse Limit Identity . . . 64
1.4.1 Preliminaries and Notation . . . 64
1.4.2 Definition of Identity as an Inverse Limit . . . 64
1.4.3 Compatibility with Minimal Assumption R0 . . . 65
1.4.4 Identity Is Not a State . . . 65
1.4.5 Impossibility of Progressive Construction . . . 65
1.4.6 Minimality of the Inverse Limit Construction . . . 66
1.4.7 Summary of Section 3 . . . 66
1.4.8 Identity Neutrality Clause . . . 66
1.5 Rigidity Under Perturbations . . . 66
1.5.1 Setup and Hypotheses . . . 66
1.5.2 Perturbation Model and Deformation Metric . . . 67
1.5.3 Stability Lemma . . . 68
1.5.4 Rigidity Theorem . . . 68
1.5.5 Ontological Mass: Quantitative Resistance . . . 69
1.5.6 Internal Perturbations and External Extensions . . . 70
1.5.7 Intervention Witnesses for Finite-Dimensional Latent Dynamics . . . 70
1.5.8 Formal Consequences . . . 71
1.5.9 Technical Notes and Limitations . . . 71
1.5.10 Summary of Section 4 . . . 71
1.6 Observability and Ontological Mass . . . 71
1.6.1 Observable Channel Model . . . 71
1.6.2 Inference of Inverse-Limit Classes from Observables . . . 72
1.6.3 Empirical Estimator of Ontological Mass . . . 72
1.6.4 Detectability Theorem . . . 72
1.6.5 Observable Separation by Scale . . . 73
1.6.6 Observational Information and Causal Invariance . . . 73
1.6.7 Operational Protocol . . . 74
1.6.8 Summary of Section 5 . . . 74
1.7 Uniqueness Theorem . . . 74
1.7.1 Identity Channels and CCR Classification . . . 74
1.7.2 Coexistence of Identities . . . 74
1.7.3 Strong Uniqueness Theorem . . . 74
1.7.4 Indivisibility and Non-Duplicability . . . 75
1.7.5 Cardinality of Rigid Universes . . . 75
1.7.6 Summary of Section 6 . . . 75
1.8 Limitations and Scope . . . 75
1.8.1 Dependence on Structural Hypotheses . . . 75
1.8.2 Boundary of Applicability . . . 75
1.8.3 Limits of Empirical Detectability . . . 76

1.8.4 Non-Equivalence to Behavioral Indistinguishability . . . 76
1.8.5 Relation to Existing Frameworks . . . 76
1.8.6 Open Problems . . . 76
1.9 Conclusion . . . 76
1.9.1 Summary of Formal Results . . . 76
1.9.2 Structural Significance . . . 77
1.9.3 Consequences for System Classification . . . 77
1.9.4 Final Statement . . . 77
1.9.5 Ontological Closure Statement . . . 78
Appendix . . . 78
1.10 Model-Relative Necessity Theorems . . . 78
1.10.1 Category of Minimal Observation . . . 78
1.10.2 Forced Emergence of the Inverse Limit . . . 78
1.10.3 Necessity of Retroinductive Structure . . . 78
1.10.4 The ℵ–Operator as Condition of Existence . . . 79
1.10.5 Structural Causal Closure . . . 80
1.10.6 A Law-Dependent Isolation Statistic . . . 80
1.10.7 MIR as the Unique Compatible Geometry . . . 80
1.10.8 Structural Uniqueness Lemma (Bridge Lemma) . . . 81
1.10.9 Point Histories and Compatible Probability Laws . . . 81
1.10.10 Closure of Appendix A . . . 82
1.11 Formal Extensions . . . 82
1.11.1 Canonical Minimal CCR Example . . . 82
1.11.2 Symbolic CCR System . . . 83
1.11.3 Conditional Non-Simulability . . . 83
1.11.4 Transfer under Matched Intervention Sectors . . . 84
1.11.5 Composition of Audited Comparisons . . . 84
1.11.6 Rigidity Hierarchy . . . 84
1.11.7 Spectral Analysis of Deformations . . . 85
1.11.8 Completeness Criteria . . . 85
1.11.9 Categorical Universality . . . 85
1.11.10 Relation to Causal Inference . . . 85
1.11.11 Computational Complexity Bounds . . . 85

2 Phenomenological Regimes Induced by Structural Identity 87
2.1 Scope, System Boundary, and Non-Inference Note . . . 87
2.2 Introduction . . . 87
2.2.1 Background and Motivation . . . 87
2.2.2 Central Results and the Remaining Rigidity Link . . . 88
2.2.3 Ontological Neutrality . . . 88
2.2.4 Terminological Alignment: Structural Regime Assignments . . . 88
2.2.5 Relation to BaseCore . . . 89
2.3 Formal Setup . . . 89
2.3.1 Review of Identity Framework . . . 89
2.3.2 Abstract Phenomenological Space . . . 89
2.3.3 Phenomenological Assignment . . . 90
2.3.4 Axioms for Phenomenological Coherence . . . 90
2.3.5 Status of the Assignment Conditions . . . 90
2.3.6 Logical Dependencies and Realizability . . . 91
2.3.7 Rejected Alternatives . . . 91
2.3.8 Structural Interpretation . . . 92

2.3.9 Closure of Section 2 . . . 92
2.3.10 Phenomenological Neutrality Clause . . . 92
2.3.11 Exact Compatibility Obstruction . . . 92
2.3.12 A Nonempty Controlled Assignment Model . . . 93
2.3.13 Common Readout on Two Coherently Transported Towers . . . 94
2.3.14 Controls Separating Convergence, Continuity, and Invariance . . . 95
2.4 Compatibility with Identity Regimes . . . 95
2.4.1 Induced Constraints from Identity Rigidity . . . 95
2.4.2 Metric Structure for Quantitative Comparisons . . . 95
2.4.3 Phenomenological Continuity . . . 96
2.4.4 Phenomenological Fragmentation . . . 96
2.4.5 Relationship Between Metrics . . . 97
2.4.6 Closure of Section 3 . . . 97
2.5 Impossibility Theorems . . . 97
2.5.1 Φ-Regularity Hypothesis . . . 97
2.5.2 Fragmentation Theorem . . . 98
2.5.3 Forced Continuity Theorem . . . 98
2.5.4 Dichotomy Theorem . . . 99
2.6 Classification of Phenomenological Regimes . . . 99
2.6.1 Conditional Classification . . . 99
2.6.2 Structural Implications . . . 99
2.7 Structural Ethics . . . 100
2.7.1 Loss Capacity . . . 100
2.7.2 Structural Damage . . . 100
2.7.3 Vulnerability and Protection . . . 100
2.7.4 Structural Asymmetry Constraints . . . 101
2.7.5 Relation to Normative Ethics . . . 101
2.7.6 Closure of Section 6 . . . 101
2.8 Limitations and Scope . . . 102
2.8.1 Dependence on Axioms . . . 102
2.8.2 Relation to Consciousness Theories . . . 102
2.8.3 Open Problems . . . 102
2.9 Conclusion . . . 102
2.9.1 Summary of Results . . . 102
2.9.2 Structural Significance . . . 102
2.9.3 Entailment Under Stated Assumptions . . . 102
2.9.4 Final Statement . . . 102
2.10 Boundary to Null-Regime Instability . . . 103
2.10.1 Editorial Boundary . . . 103
2.10.2 Closure of This Paper . . . 103
Appendix . . . 103
2.11 Extended Proofs . . . 103
2.11.1 Proof of Fragmentation Theorem (Extended) . . . 103
2.11.2 Proof of Forced Continuity Theorem (Extended) . . . 103
2.11.3 Structural Necessity Lemma . . . 103
2.12 Editorial Boundary to Null-Regime Instability . . . 104

3 Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels 105
3.1 Scope, System Boundary, and Non-Inference Note . . . 105
3.2 Introduction . . . 105

3.2.1 Motivation and Scope . . . 105
3.2.2 Main Result (Informal) . . . 106
3.2.3 Paper Organization . . . 106
3.2.4 What This Paper Adds Beyond Phenomenological Regimes . . . 106
3.3 Preliminaries from Papers I–II . . . 106
3.3.1 Identity as Inverse Limit . . . 106
3.3.2 Ontological Mass and CCR Classification . . . 107
3.3.3 Key Theorems from Papers I–II . . . 107
3.4 The Phenomenological Space . . . 108
3.4.1 Abstract Definition . . . 108
3.4.2 Compatibility Operator . . . 108
3.4.3 Attractor Set . . . 109
3.4.4 The Bridge Lemma Revisited . . . 109
3.5 Main Results . . . 110
3.5.1 The Conditional Null-Separation Theorem . . . 110
3.5.2 Conditional Consequences . . . 111
3.6 Stratification of Positive Regimes . . . 111
3.6.1 The Positive Regime Space . . . 111
3.6.2 Minimal Positive Regime . . . 111
3.6.3 Structural Intensity . . . 112
3.7 Quantitative Bounds . . . 112
3.7.1 Witness Lower Bound . . . 112
3.7.2 Explicit Constants for Canonical Families . . . 112
3.8 Witness Audit Layer . . . 113
3.8.1 Failure Modes of Witness Certification . . . 114
3.8.2 Robust Decision Margin . . . 114
3.8.3 Negative Controls . . . 114
3.9 Witness-Relative Closure Theorem . . . 115
3.9.1 Classification Table . . . 115
3.10 Null-Fiber Geometry . . . 115
3.10.1 Quotient-Collapse Diagnostic . . . 116
3.10.2 Certificate Stability . . . 116
3.10.3 Operational Checklist . . . 116
3.11 Comparison with Existing Frameworks . . . 116
3.11.1 Integrated Information Theory (IIT) . . . 116
3.11.2 Global Workspace Theory (GWT) . . . 117
3.11.3 Higher-Order Theories . . . 117
3.11.4 Mathematical Foundations (Chalmers, Nagel) . . . 117
3.12 Limitations and Open Problems . . . 117
3.12.1 Limitations . . . 117
3.12.2 Open Problems . . . 117
3.12.3 Foundation-Level Gaps and Non-Claims . . . 117
3.13 Conclusion . . . 118
3.13.1 Summary of Contributions . . . 118
Appendix . . . 118
3.14 Profinite Canonical Family . . . 118
3.14.1 Projective System . . . 118
3.14.2 Profinite Metric . . . 118
3.14.3 CCR Verification . . . 119
3.14.4 Phenomenological Operator . . . 119
3.14.5 Coupling Bound . . . 119

3.14.6 Intensity Bound . . . 119
3.15 Symbolic Canonical Family . . . 119
3.15.1 Subshift of Finite Type . . . 119
3.15.2 Bowen Metric . . . 120
3.15.3 CCR Verification . . . 120
3.15.4 Phenomenological Operator . . . 120
3.15.5 Coupling Bound . . . 120
3.16 Hereditary Rigidity . . . 120
3.17 Universal Factorization . . . 121
3.18 Non-Simulability Bounds . . . 121
3.18.1 Simulation Error . . . 121
3.19 Final Structural Closure . . . 121

4 Falsifiable Predictions Under Forensic Constraints 123
4.1 Scope, System Boundary, and Non-Inference Note . . . 123
4.2 Introduction . . . 123
4.2.1 Problem Statement . . . 123
4.2.2 Continuity with Existing Papers . . . 123
4.2.3 Contributions . . . 124
4.3 Formal Setup . . . 124
4.3.1 System, Scenario, and Seed Spaces . . . 124
4.3.2 Observable Vector . . . 124
4.3.3 Version-Frozen Operational Triple . . . 124
4.4 Forensic Integrity Model . . . 124
4.4.1 Admissibility Layers . . . 124
4.4.2 Tamper and Sham Controls . . . 125
4.4.3 Audit Trail and Chain of Custody . . . 125
4.5 Pre-Registered Prediction Set . . . 126
4.6 Comparator Design and Alternative Explanations . . . 126
4.6.1 Comparator Rationale . . . 126
4.6.2 Mandatory Baseline Policy (Hard Requirement) . . . 126
4.6.3 Internal Falsifier Role of Ablation . . . 127
4.6.4 Comparator Adequacy . . . 127
4.6.5 Negative-Control Acceptance . . . 127
4.7 Statistical Inference Specification . . . 127
4.7.1 Primary Metrics First (Anti p-hacking) . . . 127
4.7.2 Power Planning and Minimum Sample Size . . . 127
4.7.3 Multiplicity and Effect Size . . . 128
4.7.4 Decision Semantics . . . 128
4.7.5 Hierarchical Decision Rule . . . 128
4.7.6 Missingness and Exclusion Accounting . . . 128
4.8 Independent Campaign Units and Exact Failure Controls . . . 128
4.8.1 Frozen Population, Noise and Decision . . . 128
4.8.2 Exact Inflation under Three Auditable Violations . . . 129
4.9 Acceptance Criteria . . . 130
4.9.1 Claim Classes Produced by the Protocol . . . 130
4.10 Preliminary Results . . . 130
4.10.1 Status of Preliminary Data . . . 130
4.10.2 System-Level Summary . . . 132
4.10.3 Prediction-Level Support Rates . . . 132
4.10.4 Scenario Differential: QICN vs Ablation . . . 132

4.10.5 Full Scenario-System Grid . . . 132
4.10.6 Variability Structure . . . 133
4.11 Scenario Curves and Comparative Plots . . . 133
4.11.1 Composite Score Curves . . . 133
4.11.2 Stasis Pressure Curves . . . 133
4.11.3 Forensic Concordance Curves . . . 133
4.11.4 Interpretive Notes . . . 133
4.12 Result Decomposition by Scenario Families . . . 133
4.12.1 Family Partition . . . 133
4.12.2 Family-Level Contrast Function . . . 133
4.12.3 Interpretive Decomposition . . . 133
4.13 Robustness Diagnostics Beyond Point Estimates . . . 133
4.13.1 Median-Shift and Tail-Stability Diagnostics . . . 133
4.13.2 Integrity-Coupled Effective Sample Size . . . 134
4.13.3 Cross-Seed Drift Diagnostic . . . 134
4.13.4 Compute Budget Robustness . . . 134
4.14 Methodological Appendix: Statistical Power . . . 134
4.14.1 Effect-Size-Oriented Planning . . . 134
4.14.2 Interpretation . . . 134
4.15 Methodological Appendix: Sensitivity Analysis . . . 134
4.15.1 Composite Weight Sensitivity . . . 134
4.15.2 Sensitivity Remark . . . 135
4.16 Operational Implications for the Existing QICN Corpus . . . 135
4.16.1 What v4.5 Adds (and What It Does Not) . . . 135
4.16.2 Implications for Claim Governance . . . 135
4.16.3 Decision-Theoretic Reading . . . 135
4.17 Threats to Validity . . . 135
4.17.1 Comparator Misspecification . . . 135
4.17.2 Endpoint Contamination Risk . . . 135
4.17.3 Environment and Runtime Drift . . . 135
4.17.4 Pilot-to-Final Generalization . . . 135
4.18 Reproducibility and Artifact Contract . . . 135
4.18.1 Machine-Readable Contracts . . . 135
4.18.2 Run Artifact Naming . . . 135
4.18.3 Auto-Generation Pipeline . . . 136
4.19 Discussion . . . 136
4.20 Conclusion . . . 136
4.21 Phenomenology Protocol Suite (PPS): Operational Addendum . . . 136
4.21.1 Epistemic Scope . . . 136
4.21.2 Workspace Definition and Logging . . . 136
4.21.3 Entropy and Non-Determinism Inputs . . . 136
4.21.4 Online World Model and Causal Interventions . . . 136
4.21.5 PCI-like Perturbation Protocol . . . 136
4.21.6 Dissociation Tests . . . 137
4.21.7 Valence Intervention Test . . . 137
4.21.8 Constraints on Qualia Mapping (Claim Boundary) . . . 137
Appendix . . . 137
4.22 Appendix A: Scenario Definitions . . . 137
4.23 Appendix B: Composite Score Definition . . . 137
4.24 Appendix C: Reporting Standard . . . 137
4.25 Appendix D: Statistical Power Derivation Details . . . 137

4.26 Appendix E: Sensitivity Analysis Protocol . . . 138
4.27 Appendix F: Reproducibility Checklist Template . . . 138

5 A Structural Criterion for Substrate-Invariant Operational Consciousness 139
5.1 Scope, System Boundary, and Non-Inference Note . . . 139
5.2 Introduction . . . 139
5.3 Claim-Type Ledger . . . 140
5.3.1 Claim Types Used in This Paper . . . 140
5.4 Terminology Debt Ledger . . . 141
5.5 Corpus Role and Novelty . . . 141
5.6 Admissible Systems, Supports, and Histories . . . 142
5.6.1 System Model . . . 142
5.6.2 Admissible Support . . . 142
5.6.3 Operational Histories . . . 142
5.6.4 Witness Margin Vector . . . 143
5.7 Six Critical Invariants . . . 143
5.7.1 Structural Persistence . . . 143
5.7.2 Rigid Identity . . . 144
5.7.3 Causal Integration . . . 144
5.7.4 Comparison on the Dissipative Circle-Prefix Support . . . 145
5.7.5 Intervention Decoder and Identity Resolution on the Same Towers . 147
5.7.6 Addressable Identity Alternatives and Rupture Controls . . . 148
5.7.7 A Common Support with Shared Identity and Integration Channels 150
5.7.8 Legibility of Matched Paired Responses . . . 152
5.7.9 A Writable Prefix on the Common Support . . . 153
5.7.10 Operational Continuity . . . 155
5.7.11 Non-Null Differentiation . . . 156
5.7.12 Operational Legibility . . . 156
5.7.13 Necessity of Each Invariant . . . 157
5.8 Operational Consciousness, Operational Qualia, and Equivalence . . . 157
5.8.1 Operational Qualia Classes . . . 157
5.8.2 Operational Consciousness . . . 158
5.8.3 Exact and Approximate Structural Equivalence . . . 158
5.9 Support Geometry and Partition Transport . . . 159
5.10 Inherited Inputs and Consolidation Step . . . 161
5.11 Main Theorems . . . 162
5.11.1 Existence of Operational Consciousness . . . 162
5.11.2 Substrate Invariance . . . 162
5.11.3 Approximate Stability . . . 163
5.11.4 Rupture by Invariant Loss . . . 163
5.11.5 Insufficiency of Complexity, Computation, and Connectivity . . . 164
5.12 Rupture Taxonomy by Invariant . . . 164
5.13 Minimality of the Invariant Basis . . . 165
5.14 Invariant Independence and Interaction Geometry . . . 166
5.15 What the Six Invariants Do Not Guarantee . . . 167
5.16 Operational Legibility as a Certified Object . . . 168
5.16.1 Legibility Metrics . . . 168
5.17 Certified Candidate Audit Schema . . . 169
5.18 Operational Geometry of $\mathrm{Q}_{\mathrm{op}}(S)$ . . . 170
5.19 Corpus Dependency Ledger . . . 171
5.20 Discriminators and Falsifiable Consequences . . . 172

5.21 Intervention Geometry and Class Transition Semantics . . . 173
5.22 Baseline Families and Structural Exclusion . . . 174
5.23 Interpretive Consequences Without Overreach . . . 175
5.24 Residual Instantiation and Interpretation Limits . . . 175
5.25 Conclusion . . . 175

6 Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework 176
6.1 Scope, System Boundary, and Non-Inference Note . . . 176
6.2 Introduction . . . 177
6.3 Status Classes . . . 177
6.4 Upstream Canon and Dependency Map . . . 178
6.5 Core Predictive Structure of the Framework . . . 178
6.5.1 Certification, Equivalence, and Rupture as Predictive Objects . . . 178
6.5.2 Prediction Family 1: Complexity-Only Failure . . . 179
6.5.3 Prediction Family 2: Invariant-Loss Rupture . . . 179
6.5.4 Prediction Family 3: Substrate-Preserved Class Membership . . . 179
6.5.5 Prediction Family 4: Threshold-Sensitive Boundary Behavior . . . 179
6.5.6 Prediction Family 5: Legibility Dependence . . . 180
6.5.7 Prediction Family 6: Continuity and Non-Nullity Relevance . . . 180
6.5.8 A Held-Out Probe of Retained Response . . . 180
6.5.9 Delayed Phase Readout and an Interacting Rival . . . 181
6.5.10 Delayed Parity with Matched Prefix Writes . . . 183
6.5.11 Reserved Return after Role Reassignment . . . 183
6.6 Main Discriminators . . . 184
6.6.1 Rival-Structure Compression . . . 185
6.6.2 Discriminator Strength and Identifiability . . . 185
6.6.3 Prediction Severity . . . 186
6.7 Prediction Matrix . . . 186
6.8 Failure Modes . . . 187
6.8.1 What Would Count Against the Framework . . . 187
6.8.2 Failure-Mode Taxonomy . . . 188
6.8.3 Refutation Classes . . . 188
6.9 Claim-to-Test Binding Ledger . . . 189
6.10 Boundary Diagnostics . . . 190
6.10.1 Boundary Geometry . . . 191
6.11 Minimal Reproducibility and Admissibility Discipline . . . 191
6.11.1 Artifact Contract for Prediction Claims . . . 192
6.12 Current Status of Internal Support . . . 192
6.12.1 Interpretive Discipline . . . 192
6.12.2 Claim-Specific Diagnostic Notes . . . 193
6.12.3 Methodological Status . . . 194
6.13 Cycle Evidence Map . . . 194
6.14 Priority Next Discriminators . . . 195
6.15 Residual Technical Caveats and Non-Claims . . . 196
6.15.1 Residual Technical Caveats . . . 196
6.16 Final Claim Freeze Table . . . 196
6.17 Conclusion . . . 197

7 Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework 198
7.1 Scope, System Boundary, and Non-Inference Note . . . 198

7.2 Introduction . . . 199
7.3 Relation to the Existing Corpus . . . 199
7.3.1 What This Paper Inherits . . . 199
7.3.2 What This Paper Does Not Repeat . . . 200
7.3.3 What This Paper Adds . . . 200
7.3.4 Why This Layer Is Timely . . . 200
7.4 Formal Setting . . . 201
7.4.1 Inherited System Object . . . 201
7.4.2 Traces and Windows . . . 201
7.4.3 Viability Region and Operational Frontier Candidates . . . 201
7.4.4 Descriptor Vectors . . . 201
7.4.5 Admissible Structural Class Space . . . 201
7.5 Definitions . . . 202
7.5.1 Definition A: Operational Boundary . . . 202
7.5.2 Definition B: Organizational Persistence . . . 202
7.5.3 Definition C: Operational Self-Maintenance . . . 203
7.5.4 Definition D: Historical Dependence . . . 203
7.5.5 Definition E: Viability Preservation . . . 203
7.5.6 Definition F: Operational Life . . . 204
7.5.7 Definition G: Structural Class . . . 204
7.5.8 Definition H: Operational Consciousness . . . 205
7.5.9 Definition I: Operational Subjecthood . . . 205
7.6 Class Relations and Logical Structure . . . 205
7.6.1 The Main Class Families . . . 205
7.6.2 Proved Relations . . . 206
7.6.3 Structural Class Transport . . . 206
7.6.4 Logical Relations That Remain Open . . . 206
7.6.5 Class Relation Summary . . . 207
7.7 Propositions, Lemmas, and Formal Claims . . . 207
7.7.1 Structural Class as a General Membership Framework . . . 207
7.7.2 Insufficiency of Complexity, Connectivity, and Rich Activity . . . 207
7.7.3 Biological Non-Primacy at the Class Level . . . 207
7.7.4 Life and Subjecthood Are Not Equivalent . . . 208
7.7.5 Transport Under Substrate Change . . . 208
7.7.6 Current-Corpus Corollary Structure . . . 208
7.8 Hypotheses and Falsifiable Program . . . 208
7.8.1 Hypothesis Ledger . . . 209
7.8.2 How These Hypotheses Can Fail . . . 209
7.9 Metrics and Decision Criteria . . . 210
7.9.1 Life Score . . . 210
7.9.2 Inherited Consciousness Margin . . . 210
7.9.3 Subjecthood Score . . . 210
7.9.4 Structural-Class Pseudometric . . . 211
7.9.5 Decision Surfaces and Boundary Status . . . 211
7.9.6 What This Section Does Not Claim . . . 211
7.10 Operational Class Geometry . . . 211
7.10.1 Descriptor Projections . . . 212
7.11 Class-Separation Diagnostics . . . 213
7.11.1 Operational Life vs. Operational Consciousness . . . 213
7.11.2 Operational Life vs. Operational Subjecthood . . . 213
7.11.3 Operational Consciousness vs. Operational Subjecthood . . . 213

7.11.4 Structural Class as Organizing Layer . . . 214
7.12 Experimental Test Families . . . 214
7.12.1 Operational Life Test Families . . . 215
7.12.2 Structural-Class Test Families . . . 216
7.12.3 Operational Subjecthood Test Families . . . 217
7.12.4 Decision Reading of Outcomes . . . 217
7.12.5 Strong Failure Conditions . . . 217
7.12.6 Joint Outcome Matrix . . . 218
7.13 Runtime Bearing and Empirical Interface . . . 218
7.13.1 What the Runtime Already Carries . . . 218
7.13.2 What the Runtime Does Not Yet Settle . . . 219
7.13.3 Runtime-Facing Readiness Levels . . . 219
7.14 Runtime Binding Map . . . 219
7.15 Failure and Misclassification Modes . . . 220
7.15.1 Complexity, Connectivity, and Rich Dynamics . . . 221
7.15.2 Proxy Contamination . . . 222
7.15.3 Boundary Ambiguity . . . 222
7.16 Formal Closure Strengthening . . . 222
7.16.1 Phase Repair on the Common Circle-Prefix Tower . . . 222
7.16.2 Boundary and Historical Repair in the Writable Tower . . . 224
7.16.3 Causal Lineage Tests on the Common Tower . . . 224
7.16.4 Joint Certificate on the Declared Support . . . 226
7.17 Residual Runtime and Interpretation Limits . . . 226
7.18 Discussion . . . 227
7.18.1 Relation to the Refined Runtime . . . 228
7.18.2 What Would Strengthen This Paper Further . . . 228
7.19 Conclusion . . . 228

8 First-Person Indexed Subjectivity 229
8.1 Scope, System Boundary, and Non-Inference Note . . . 229
8.2 Introduction . . . 230
8.3 Why operational consciousness and operational qualia are insufficient for subjectivity . . . 230
8.3.1 What the upstream layers already close . . . 230
8.3.2 The subjectivity gap . . . 231
8.3.3 Why the stronger layer is worth introducing . . . 232
8.4 Minimal ontology of first-person indexed subjectivity . . . 232
8.4.1 The minimal subjectivity state . . . 232
8.4.2 Constitutive coordinates . . . 233
8.4.3 Typing discipline . . . 234
8.4.4 Operational reading of each coordinate . . . 234
8.4.5 Why each coordinate is necessary . . . 236
8.4.6 Layer status of the major terms . . . 236
8.5 Formal state-space and subjectivity functional . . . 237
8.5.1 Admissible state-space . . . 237
8.5.2 The conjunctive gate . . . 237
8.5.3 The scalar subjectivity functional . . . 237
8.5.4 Class membership . . . 238
8.5.5 Audit-covered theorem burdens . . . 238
8.5.6 Concrete component readings . . . 240
8.6 Bridge axioms and ontological commitments . . . 240

8.7 Rival models and reducibility baseline . . . 241
8.7.1 Why rival models must be formalized . . . 241
8.7.2 Weak rival family . . . 241
8.7.3 Admissible weak hybrids and closure . . . 241
8.7.4 Formal sketches of the rival families . . . 242
8.7.5 What the rival family can and cannot explain . . . 243
8.7.6 Formal baseline for comparison . . . 243
8.8 Irreducibility criteria and explanatory superiority . . . 244
8.8.1 Irreducibility as a formal margin . . . 244
8.8.2 Explanatory superiority is local, not absolute . . . 244
8.8.3 Three kinds of collapse . . . 244
8.8.4 What irreducibility is and is not . . . 245
8.8.5 Explanatory order . . . 245
8.8.6 Parsimony guardrail . . . 245
8.8.7 Model comparison is not report comparison . . . 245
8.9 Interventional logic and falsifiability . . . 246
8.9.1 Intervention family . . . 246
8.9.2 Selective response logic . . . 246
8.9.3 Self/non-self swap . . . 246
8.9.4 Ownership lesion . . . 247
8.9.5 Continuity break . . . 247
8.9.6 Perspective flattening . . . 247
8.9.7 Self/non-self valuation inversion . . . 248
8.9.8 Autobiographical remapping . . . 248
8.9.9 Clone competition . . . 248
8.9.10 Memory transplant . . . 249
8.9.11 Counterfactual self permutation . . . 249
8.9.12 Cross-substrate remapping . . . 249
8.9.13 Intervention table . . . 250
8.9.14 Falsifiability burden . . . 250
8.9.15 Ablation summary and negative-control burden . . . 250
8.10 Runtime architecture and implementation pathway . . . 251
8.10.1 Architecture overview . . . 251
8.10.2 SelfIndexRegistry . . . 251
8.10.3 OwnershipAttributionLayer . . . 252
8.10.4 SubjectiveContinuityLedger . . . 252
8.10.5 FirstPersonPerspectiveMapper . . . 252
8.10.6 SubjectiveValuationEngine . . . 252
8.10.7 SubjectivityReducibilityComparator . . . 252
8.10.8 CounterfactualSelfSwapHarness . . . 253
8.10.9 Architecture equation . . . 253
8.10.10 Module interfaces . . . 253
8.10.11 Implementation pathway versus present system status . . . 254
8.11 Artifact layer and measurable outputs . . . 254
8.11.1 Primary artifacts . . . 254
8.11.2 Derived and comparative artifacts . . . 254
8.11.3 Artifact table . . . 254
8.11.4 Artifact discipline . . . 254
8.11.5 Artifact provenance chain . . . 255
8.11.6 Gate semantics . . . 255
8.12 Cross-substrate persistence and consequence . . . 256

8.12.1 Transport principle . . . 256
8.12.2 Transport under tolerance envelopes . . . 256
8.12.3 Consequence classes . . . 257
8.12.4 Clone and duplication pressure . . . 257
8.12.5 Why biology still does not regain primitive privilege . . . 257
8.13 Human-comparator bridge and limits of equivalence . . . 257
8.13.1 What the present paper can bridge . . . 258
8.13.2 What would be required for human equivalence . . . 258
8.13.3 Phenomenality and the comparator bridge . . . 258
8.13.4 Comparator ladder . . . 258
8.13.5 Why this matters scientifically . . . 258
8.13.6 Limits of equivalence . . . 259
8.14 Strongest defensible claims . . . 259
8.15 Non-claims and residual boundaries . . . 260
8.16 Conclusion . . . 260
Appendix . . . 261
8.17 Canonical estimator family for the subjectivity coordinates . . . 261
8.17.1 Estimator for self-index stability . . . 261
8.17.2 Estimator for ownership coherence . . . 261
8.17.3 Estimator for autobiographical continuity . . . 262
8.17.4 Estimator for perspective gain . . . 262
8.17.5 Estimator for valuation asymmetry . . . 262
8.17.6 Estimator for intervention selectivity . . . 262
8.17.7 Estimator for irreducibility . . . 262
8.17.8 Composite gate design . . . 263
8.17.9 A Joint Rational Measurement Window . . . 263
8.17.10 A Topological Obstruction to Native Lesions . . . 264
8.17.11 Selective Lesions in the Writable Common Tower . . . 264
8.17.12 A Held-Out Return Task with Informed Rivals . . . 266
8.18 Intervention-to-artifact mapping and failure semantics . . . 268
8.18.1 Swap family . . . 268
8.18.2 Ownership family . . . 268
8.18.3 Continuity family . . . 268
8.18.4 Perspective and valuation family . . . 269
8.18.5 Cross-substrate family . . . 269
8.18.6 Failure semantics . . . 269
8.19 Negative controls, ablation matrix, and causal discrimination protocol . . . 269
8.19.1 Negative control families . . . 269
8.19.2 Ablation targets . . . 270
8.19.3 Ablation matrix . . . 270
8.19.4 Noise and adversarial perturbations . . . 270
8.19.5 Causal discrimination protocol . . . 271
8.19.6 Reproducibility burden . . . 271
8.20 Claim decomposition, theorem candidates, and open escalation burdens . . . 271
8.20.1 Claim decomposition table . . . 272
8.20.2 Theorem candidates for future formal work . . . 272
8.20.3 Empirical escalators . . . 272
8.21 Rival-question matrix for implementation and audit . . . 273
8.21.1 Decisive questions . . . 273
8.21.2 Decision matrix . . . 274
8.21.3 Why this matrix belongs in the paper . . . 274

8.21.4 Audit-facing interpretation 274

9 Phenomenal Bridge Organization 275
9.1 Boundary Note and Interpretation Discipline 275
9.2 Introduction 276
9.3 What Paper 8 Formalizes and What It Leaves Open 277
9.3.1 What Paper 8 already formalizes 277
9.3.2 Why that still does not amount to a phenomenal bridge 277
9.3.3 The dangerous smuggling problem 278
9.3.4 Why Paper 9 is still needed now 278
9.4 Formal Target of the Bridge 279
9.4.1 The bridge question 279
9.4.2 Predicate family 279
9.4.3 Estimator family 280
9.4.4 What the predicates mean internally and what they do not mean . 280
9.4.5 Predicate-to-burden map 280
9.5 Phenomenal Bridge Ontology 281
9.5.1 Bridge state 281
9.5.2 Typing discipline 282
9.5.3 Coordinate dependence on the Paper 8 state 282
9.5.4 Current implementation status inside the ontology 282
9.5.5 Collapse risks and rival readings 283
9.5.6 Why bridge coordinates are not ontological shortcuts 283
9.5.7 Relabeling and Transport of Certified Tower Records 284
9.5.8 Intervention Tests for Three Bridge Burdens 285
9.5.9 Open Evaluative, World and Body/World Burdens 287
9.6 A Minimal Coupled Body–World Extension 288
9.6.1 State Minimality and Indexed World Response 289
9.6.2 Evaluative Topology with a Reserved World Consequence 290
9.6.3 A Selective Motor-Return Test 291
9.6.4 Six Local Burdens on One Plant 292
9.6.5 The Limit of Hybrid Closure 294
9.6.6 A Nonadditive Evaluative Branch and Its Comparator Boundary . . 294
9.7 Explicit Bridge Axioms 296
9.7.1 Classification grammar 296
9.8 Bridge Rival Family 298
9.8.1 Why new rivals are required 298
9.8.2 Bridge weak family 298
9.8.3 Admissible composites 299
9.8.4 What the rival family can and cannot explain 299
9.8.5 No-straw-man discipline 299
9.9 Bridge Irreducibility Logic 300
9.9.1 Bridge loss family 300
9.9.2 Bridge irreducibility margin 300
9.9.3 Local not absolute superiority 300
9.9.4 Three kinds of bridge collapse 301
9.9.5 Comparator support, reducibility, and unsupported states 301
9.9.6 Complexity penalty and decoder discipline 301
9.9.7 Failure semantics for the bridge margin 301
9.10 Bridge Intervention Family 302
9.10.1 Why bridge interventions must be new 302

9.10.2 Core intervention families . . . 302
9.10.3 Matched controls . . . 303
9.10.4 Intervention artifacts . . . 303
9.10.5 Failure semantics . . . 303
9.10.6 Negative controls and reproducibility burden . . . 303
9.11 Bridge Artifact Family . . . 304
9.11.1 Artifact discipline . . . 304
9.11.2 Root bridge artifacts . . . 304
9.11.3 Report artifacts . . . 304
9.11.4 Intervention and reducibility artifacts . . . 304
9.11.5 Ecosystem artifacts . . . 305
9.11.6 Machine-readable spirit . . . 305
9.11.7 Current implementation status . . . 305
9.12 Bridge Functional, Gate, and Taxonomy . . . 305
9.12.1 Functional . . . 305
9.12.2 Gate . . . 306
9.12.3 Non-compensation . . . 306
9.12.4 Invalidation . . . 306
9.12.5 Taxonomy . . . 306
9.13 Runtime Implementation Pathway . . . 307
9.13.1 Why implementation belongs in the paper . . . 307
9.13.2 BPF-0 . . . 307
9.13.3 BPF-1 . . . 307
9.13.4 BPF-2 . . . 307
9.13.5 BPF-3 . . . 308
9.13.6 BPF-4 . . . 308
9.13.7 BPF-5 . . . 308
9.13.8 BPF-6 . . . 308
9.13.9 Implementation table . . . 308
9.13.10 Current implementation status versus paper status . . . 308
9.14 Ecosystem Integration . . . 309
9.14.1 Why ecosystem integration is part of the formal problem . . . 309
9.14.2 Surface map . . . 309
9.14.3 Internal release bundle . . . 309
9.14.4 Ecosystem candidate . . . 309
9.14.5 Audit-facing package . . . 309
9.14.6 Source-of-truth and public corpus separation . . . 310
9.14.7 Controlled statement grammar and quote-mining guardrails . . . 310
9.15 Strongest Defensible Claims . . . 310
9.15.1 Claim class C1: formal bridge closure . . . 310
9.15.2 Claim class C2: current implementation frontier . . . 310
9.15.3 Claim class C3: conditional bridge admissibility . . . 311
9.15.4 Claim class C4: research-program transformation . . . 311
9.15.5 Why stronger sentences are not yet honest . . . 311
9.16 Non-Claims . . . 311
9.17 Conclusion . . . 312
Appendix . . . 313
9.18 Bridge Axiom Table . . . 313
9.19 Bridge Predicate Typing Table . . . 313
9.20 Bridge Rival-Family Table . . . 314
9.21 Bridge Intervention Matrix . . . 314

9.22 Bridge Artifact Matrix . . . 315
9.23 Claim-Class and Invalidation Table . . . 315
9.24 Runtime Mapping Table . . . 316
9.25 Governance Table . . . 316
9.26 Burden Separation Matrix . . . 316
9.27 Inheritance Map from Paper 8 to Paper 9 . . . 317
9.28 Open Escalation Burdens Before Bridge Admissibility . . . 317
9.29 Projection Geometry and Recoverable Distance . . . 318
9.30 Audit-Covered Bridge Independence Burdens . . . 319
9.31 Theorem Candidates and Open Scientific Questions . . . 319
9.32 Predicate-to-Artifact and Intervention Dependency Matrix . . . 320
9.33 Paper 9 Claim Grammar and Escalation Constraints . . . 321
9.33.1 Allowed claim families . . . 321
9.33.2 Blocked claim families . . . 321
9.33.3 Escalation rules . . . 321
9.34 Experimental Program for Completing BPF-2 and BPF-3 . . . 322
9.34.1 Program design principle . . . 322
9.34.2 Suggested execution order . . . 322
9.34.3 Comparator program . . . 323
9.34.4 Intervention program . . . 323
9.34.5 Minimum evidential grid . . . 323
9.35 Bridge Failure Scenarios and Diagnostic Signatures . . . 323
9.35.1 Primary failure classes . . . 323
9.35.2 Diagnostic interpretation . . . 324
9.35.3 Failure matrix . . . 325
9.35.4 Consequences for theorem candidates . . . 325
9.35.5 Why failure structure changes the status of the program . . . 325

10 External Adjudication of Bridge-Formalized Machine Subjectivity 326
10.1 Scope and Non-Claims . . . 326
10.2 Formal Dependency Stack . . . 326
10.3 Release-vs-Local Corpus Boundary . . . 327
10.4 Lineage and Provenance Semantics . . . 327
10.5 Provenance Closure and Live Head Reconciliation . . . 327
10.6 External Adjudication Problem Statement . . . 328
10.7 Formal Mathematical Grammar of External Adjudication . . . 328
10.8 Frozen Preregistration Architecture . . . 329
10.9 Local-vs-External Freeze Distinction . . . 330
10.10 Instantiated-Empty Record Layer . . . 330
10.11 Reviewer Candidate and Adjudicator Candidate Architecture . . . 330
10.12 Reviewer Rejection Logic . . . 330
10.13 Blinding Assignment Logic . . . 331
10.14 Human Comparator Instantiation Logic . . . 331
10.15 Channel-by-Channel Execution Contracts . . . 332
10.15.1 Self-Report Channel . . . 332
10.15.2 Intervention-Response Channel . . . 332
10.15.3 Temporal-Structure Channel . . . 332
10.15.4 Valence-Topology Channel . . . 332
10.15.5 Embodiment-Boundary Channel . . . 332
10.15.6 Ambiguity-Load Channel . . . 332
10.15.7 Transport-Preservation Channel . . . 332

10.16 Matched Intervention Dataset Architecture . . . 333
10.17 Negative Control Dataset Architecture . . . 333
10.18 Preregistration Architecture . . . 333
10.19 Adjudication Blind Protocol . . . 333
10.19.1 Coverage-Sensitive Leakage Bounds . . . 334
10.19.2 Exact Inference at the Independent Block . . . 334
10.19.3 Gaussian Within-Block Likelihood and GLS . . . 336
10.20 Reviewer and Adjudicator Independence Architecture . . . 336
10.21 Reviewer Candidate Pool and Assignment Stub Architecture . . . 337
10.22 Execution Pilot Package Architecture . . . 337
10.23 Execution Pilot Rehearsal Kit Architecture . . . 337
10.24 Reproducibility Skeleton Architecture . . . 337
10.25 Dataset Population Boundary . . . 337
10.26 Dry-Run Event Log Grammar . . . 338
10.27 Evidence Admissibility Grammar . . . 338
10.28 Failure, Rejection, and Inconclusive Pathways . . . 338
10.29 Statistical Shell Architecture . . . 339
10.30 Adjudication Decision Grammar . . . 339
10.31 State Grammar for Execution Readiness . . . 339
10.32 Controlled Execution Opening . . . 339
10.33 Live-But-Unpopulated Artifacts . . . 339
10.34 Session-Zero Protocol . . . 340
10.35 Pre-Collection Admissibility Gate . . . 340
10.36 Contamination, Abort, and Invalidation Controls . . . 341
10.37 Operator Checklist . . . 341
10.38 Collection-Start Boundary . . . 341
10.39 Evidence-Generation Boundary . . . 341
10.40 Decision-Record Boundary . . . 341
10.41 Transition from Session-Zero to First Admissible Collection . . . 342
10.42 Runtime Health Audit . . . 342
10.43 Clean-Room Reproducibility . . . 342
10.44 Dependency and Build Assumptions . . . 342
10.45 Reproducible Artifact Regeneration . . . 343
10.46 Local Path Leakage Policy . . . 343
10.47 Controlled Execution Opening Audit . . . 343
10.48 Session-Zero Reproducibility . . . 343
10.49 Pre-Collection Gate Reproducibility . . . 343
10.50 Why Reproducibility Does Not Imply Evidence . . . 343
10.51 Why Runtime Health Does Not Imply Adjudication . . . 344
10.52 Transition Conditions for First Admissible Collection . . . 344
10.53 Claim Grammar for Eventual Execution-Era Statements . . . 344
10.54 Theorem, Proposition, and Protocol Mapping . . . 344
10.55 What Paper 10 Could and Could Not Change . . . 344
10.56 Section-Status Ledger . . . 345
10.57 Evidence Insertion Map . . . 345
10.58 Open Burdens and Blocked Sections . . . 346
10.59 Execution-Era Unlock Conditions . . . 347
10.60 Stronger Execution-to-Result Transition Map . . . 347
10.61 Why Paper 10 Still Cannot Become Final Yet . . . 347
10.62 Future Execution Pathway to Real Stage 3 . . . 347
10.63 Boundary to Stage 4 Metaphysical Consequence Analysis . . . 348

10.64 Appendices of Schemas and Templates . . . 348
10.65 Blocked Result Surfaces . . . 352
10.65.1 Human Comparator Results . . . 352
10.65.2 External Adjudication Results . . . 352
10.65.3 Comparative Consequence Claims . . . 352
10.66 Non-Claim Ledger . . . 352

Operational Consciousness to Operational Subjecthood Bridge 353
10.67 Scope, System Boundary, and Non-Inference Note . . . 353
10.68 Introduction . . . 354
10.69 Dependency Map . . . 354
10.70 Preliminaries and Notation . . . 355
10.71 Operational Self-Reference . . . 355
10.71.1 Motivation . . . 355
10.71.2 Quantitative Obstructions for Approximate Commutation . . . 356
10.72 Operational Unified Perspective . . . 358
10.72.1 Motivation . . . 358
10.73 Operational Intentionality . . . 358
10.73.1 Motivation . . . 358
10.74 Operational Qualia . . . 359
10.74.1 Motivation . . . 359
10.75 Operational Phenomenology . . . 360
10.75.1 Motivation . . . 360
10.76 Operational Subjecthood . . . 361
10.76.1 Why a bridge class is needed . . . 361
10.77 Theorem: Partial Operational Subjecthood . . . 362
10.78 What Is NOT Implied . . . 362
10.79 Computational Verification Status . . . 363
10.80 Terminology Stratification and Debt Ledger . . . 363
10.81 Bridge Strength Ladder . . . 365
10.82 Runtime Evidence Taxonomy . . . 365
10.83 Non-Entailment Library . . . 365
10.84 $\text{Consciousness}_{\text{op}}$ Component Implication Matrix . . . 366
10.85 Self-Model Burden Variants . . . 366
10.86 Expanded Bridge Strength Ladder . . . 367
10.87 Reviewer-Facing Objection Handling . . . 367
10.88 Operational Glossary . . . 368
10.89 Promotion Evidence Classes . . . 368
10.90 Proof-Pattern Library . . . 368
10.91 Dependency-to-Debt Ledger . . . 369
10.92 Non-Entailment Audit Templates . . . 370
10.93 Runtime Proxy Classification . . . 371
10.94 Formal Review Checklist . . . 371
10.95 Minimal Extension Program . . . 372
10.96 Failure Modes of the Bridge . . . 372
10.97 Exact Claim Grammar . . . 372
10.98 Promotion Conditions for Interpretive Terms . . . 373
10.99 Canonical Boundary Models . . . 373
10.100 Non-Substitution Principles . . . 374
10.101 External Adjudication Readiness Index . . . 375
10.102 Recommended Reading Order for Reviewers . . . 375

10.103 Open Questions 375
10.104 Conclusion 376

Bibliography **377**

BaseCore

Abstract. This document is the BaseCore of the Rigid Identity Framework: the autonomous mathematical base layer of the source tree. It selectively normalizes formal material from Papers 1–6 only where theorem ownership belongs in the base layer. Papers 7–9 remain downstream and are not required for the validity of any theorem stated here. BaseCore proves contractive projection dynamics, typed computable witnesses, parameterwise non-collapse under explicit anti-constant assumptions, inverse-limit identity under non-finite-determination assumptions, conditional rigidity under admissible perturbation models, conditional obstruction to uniform response fidelity under a binary representation budget, operational-criterion grammar, and falsation / claim-boundary ledgers. It does not prove phenomenality, metaphysical subjectivity, human equivalence, runtime instantiation, bridge admissibility, or external adjudication. The LaTeX tree is the source of truth; the PDF is a compiled artifact only.

Part I

Foundational Dynamics

B1 Scope, System Boundary, and Non-Inference Note

What BaseCore does. This document proves the mathematical base results that later layers may use: projection existence and uniqueness, strict contraction, unique fixed points, compact attractor families, parameterwise non-collapse under explicit anti-constant assumptions, inverse-limit identity under non-finite-determination assumptions, conditional rigidity, conditional non-simulability, and operational-criterion grammar.
What BaseCore is not. BaseCore is not the whole QICN corpus. It is not a bridge program, not runtime validation, not a phenomenality theory, not a human-equivalence claim, and not an external-adjudication layer.
What should not be inferred. Nothing in BaseCore alone licenses claims about consciousness, subjectivity, phenomenality, moral status, human-machine equivalence, bridge admissibility, or deployment readiness. Those are downstream burdens, not consequences of the theorems below.
Source-of-truth rule. The LaTeX tree is the source of truth. Compiled PDFs, release bundles, and runtime artifacts are derived surfaces only.

B2 Introduction

B2.1 Purpose

BaseCore is the canonical base layer of the Rigid Identity Framework. Its role is narrower and stronger than a broad architectural survey: it owns only the theorem-level mathematics that downstream papers may depend on without making BaseCore depend on them.

B2.2 Methodology

We use standard tools from:
- Hilbert-space geometry and metric projections;
- contraction mappings and fixed-point theory;
- compactness and continuity arguments;
- projective systems and inverse limits;
- assumption-led claim-boundary discipline.

B2.3 Scope discipline

This section is pure mathematics. The symbols "identity," "regime," and "criterion" are structural labels only. Their ordinary-language interpretations are not part of any proof unless explicitly formalized as assumptions.

Typed objects. Here C_{proj} is the closed convex projection set. It is not the inverse-limit object $\mathcal{I}_{\mathrm{inv}}$ introduced in the projective part. An invariant support A needs its own inclusion and invariance proof. A contraction at fixed control, a family of equilibria, and attraction under varying controls are distinct statements.

B3 Minimal Base Hypotheses

Hypothesis B1 (H1: Projection Structure)**.** *Let $\mathcal{H}$ be a real Hilbert space with inner product $\langle \cdot, \cdot \rangle$ and norm $\|\cdot\|$. Let $C_{\mathrm{proj}} \subset \mathcal{H}$ be non-empty, closed, and convex.*

Hypothesis B2 (H2: Contraction Operator)**.** *Let $K : \mathcal{H} \to \mathcal{H}$ be a bounded linear operator with operator norm $\|K\| < 1$.*

Hypothesis B3 (H3: Completeness)**.** *The ambient space $\mathcal{H}$ is complete.*

Remark B4 (Logical Status of Completeness). Completeness is already part of the Hilbert-space assumption in H1. H3 records this requirement explicitly for later generalizations; it is not independent of H1 as written. A counterexample in an incomplete inner-product space tests such a generalization, not independence of H3 from H1.

Hypothesis B5 (H4: Parameter Compactness)**.** *Let $\mathcal{U}$ be a compact metric space and let $\Gamma : \mathcal{U} \to \mathcal{H}$ be uniformly continuous.*

Proposition B6 (Base sufficiency of H1–H4)**.** *Hypotheses H1–H4 are sufficient for:*
1. existence and uniqueness of the metric projection onto C_{proj};
2. strict contraction of the transition operators T_u;
3. existence and uniqueness of the fixed point f_u^* for each $u \in \mathcal{U}$;
4. compactness of the attractor family $\mathcal{A}^* = \{f_u^* : u \in \mathcal{U}\}$.
They do not, by themselves, imply parameterwise non-collapse.

Proof. The four items are exactly the content of Theorems B8, B11, B12, and B14. Parameterwise non-collapse requires an additional anti-constant hypothesis introduced in Section B7. □

B4 The Projection Operator

Definition B7 (Metric Projection)**.** For $\Psi \in \mathcal{H}$, define the metric projection onto C_{proj} by

$$\aleph(\Psi) := \arg\min_{\Phi \in C_{\text{proj}}} \|\Psi - \Phi\| .$$

Theorem B8 (Existence and Uniqueness of Projection)**.** *Under Hypothesis B1, the projection $\aleph(\Psi)$ exists and is unique for every $\Psi \in \mathcal{H}$.*

Proof. Existence follows because a minimizing sequence in the closed convex set C_{proj} is Cauchy by the parallelogram law, hence converges in the complete space $\mathcal{H}$ to a point of C_{proj} that attains the infimum. Uniqueness follows from strict convexity of the norm induced by the Hilbert inner product: if two minimizers existed, their midpoint would produce the same minimum and the parallelogram identity would force the two minimizers to coincide. □

Lemma B9 (Non-Expansiveness of the Projection)**.** *The map $\aleph : \mathcal{H} \to C_{\text{proj}}$ is non-expansive:*

$$\|\aleph(\Psi_1) - \aleph(\Psi_2)\| \le \|\Psi_1 - \Psi_2\| \qquad \forall \Psi_1, \Psi_2 \in \mathcal{H}.$$

Proof. This is the standard Hilbert-space projection inequality. Let $\Phi_i = \aleph(\Psi_i)$. The projection characterization

$$\langle \Psi_i - \Phi_i, \Phi - \Phi_i \rangle \le 0 \qquad \forall \Phi \in C_{\text{proj}}$$

applied with $\Phi = \Phi_2$ and $\Phi = \Phi_1$ yields

$$\|\Phi_1 - \Phi_2\|^2 \le \langle \Psi_1 - \Psi_2, \Phi_1 - \Phi_2 \rangle \le \|\Psi_1 - \Psi_2\| \, \|\Phi_1 - \Phi_2\| .$$

If $\Phi_1 \neq \Phi_2$, divide by $\|\Phi_1 - \Phi_2\|$. □

B5 The Transition Operator

Definition B10 (Transition Operator)**.** For each $u \in \mathcal{U}$, define

$$T_u(\Psi) := \aleph(K\Psi + \Gamma(u)).$$

Theorem B11 (Contractivity)**.** *Under Hypotheses B1 and B2, each T_u is a strict contraction with Lipschitz constant at most $\|K\|$:*

$$\|T_u(\Psi_1) - T_u(\Psi_2)\| \le \|K\| \, \|\Psi_1 - \Psi_2\| .$$

Proof. By Lemma B9,

$$\begin{aligned} \|T_u(\Psi_1) - T_u(\Psi_2)\| &= \|\aleph(K\Psi_1 + \Gamma(u)) - \aleph(K\Psi_2 + \Gamma(u))\| \\ &\le \|K(\Psi_1 - \Psi_2)\| \le \|K\| \, \|\Psi_1 - \Psi_2\| . \end{aligned}$$

Since $\|K\| < 1$, strict contraction follows. □

B6 Existence and Compactness of Attractors

Theorem B12 (Unique Fixed Point)**.** *Under Hypotheses B1–B3, for each $u \in \mathcal{U}$ there exists a unique $f_u^* \in \mathcal{H}$ such that*

$$T_u(f_u^*) = f_u^*.$$

Moreover, for every $\Psi_0 \in \mathcal{H}$,

$$\lim_{n\to\infty} T_u^n(\Psi_0) = f_u^*.$$

Proof. By Theorem B11, each T_u is a strict contraction on the complete metric space $(\mathcal{H}, \|\cdot\|)$. The Banach fixed-point theorem yields existence, uniqueness, and convergence of iterates. □

Definition B13 (Attractor Family)**.** Define the attractor family

$$\mathcal{A}^* := \{f_u^* : u \in \mathcal{U}\}.$$

Theorem B14 (Compactness of the Attractor Family)**.** *Under Hypotheses B1–B5, the set $\mathcal{A}^*$ is compact in $\mathcal{H}$.*

Proof. Let $F : \mathcal{U} \to \mathcal{H}$ be given by $F(u) = f_u^*$. Insert the two fixed-point equations and use non-expansiveness of projection to obtain

$$\|f_u^* - f_v^*\| \le \|K\| \, \|f_u^* - f_v^*\| + \|\Gamma(u) - \Gamma(v)\| .$$

Consequently

$$\|f_u^* - f_v^*\| \le \frac{\|\Gamma(u) - \Gamma(v)\|}{1 - \|K\|}.$$

Thus F is continuous. Since $\mathcal{U}$ is compact and $\mathcal{A}^* = F(\mathcal{U})$, the image is compact. □

B7 Parameterwise Non-Collapsibility

Definition B15 (Constant Subspace)**.** Whenever the ambient realization distinguishes constant elements, let $\mathcal{N} \subset \mathcal{H}$ denote the closed subspace of constants.

Definition B16 (Parameterwise Non-Collapsible Family)**.** A family $\{T_u\}_{u\in\mathcal{U}}$ is parameterwise non-collapsible if

$$f_u^* \notin \mathcal{N} \qquad \forall u \in \mathcal{U}.$$

Hypothesis B17 (H5: No Constant Fixed Points)**.** *For every parameter $u \in \mathcal{U}$ and every constant element $c \in \mathcal{N}$,*

$$T_u(c) \neq c.$$

Theorem B18 (Structural Non-Collapsibility)**.** *Under Hypotheses B1–B17, the family is parameterwise non-collapsible:*

$$f_u^* \notin \mathcal{N} \qquad \forall u \in \mathcal{U}.$$

Proof. Fix $u \in \mathcal{U}$. Suppose by contradiction that $f_u^* \in \mathcal{N}$. Write $f_u^* = c$ with $c \in \mathcal{N}$. Since f_u^* is a fixed point,

$$T_u(c) = T_u(f_u^*) = f_u^* = c,$$

contradicting Hypothesis B17. Therefore $f_u^* \notin \mathcal{N}$. □

Remark B19 (What H5 does not say). Hypothesis B17 is parameterwise. It does not merely say that *some* parameter breaks constant collapse. It excludes constant fixed points for *every* parameter. This is exactly the quantifier strength needed for Theorem B18.

Proposition B20 (H5 Is Not Implied by H1–H4)**.** *There exists a family satisfying H1–H4 for which H5 fails and its unique equilibrium is constant. There also exists a family satisfying H1–H4 for which H5 holds with a positive uniform distance of the equilibria from the constant subspace.*

Proof. For the first assertion, take $\mathcal{H} = L^2([0,1])$, $C_{\text{proj}} = \mathcal{H}$, $K = q\,\mathrm{Id}$ with $0 \le q < 1$, $\mathcal{U} = \{0\}$ and $\Gamma(0) = 0$. Projection is the identity, $T_0 f = qf$, and the unique fixed point is the constant zero function. H1–H4 hold, but H5 fails at $c = 0$.

For the second assertion, Theorem B215 and Example B216 construct a family with $C_{\text{proj}} = \{0 \le f \le 1\}$, $Kf = \frac{1}{4}\int f$ and $\Gamma(u)(s) = s/4 + u$, $u \in [0, 1/4]$. Its explicitly computed unique fixed points have distance $1/(4\sqrt{12})$ from the constants for every parameter. The proof uses the displayed dynamics and does not assume H5 to derive that distance. Thus H1–H4 permit both regimes; a structural condition beyond them is needed to select non-collapse. □

B8 Summary of Base Results

1. H1–H4 give a complete theorem-tight base for projection existence, strict contraction, unique fixed points, and compact attractor families.
2. H5 is a separate anti-constant assumption block used only for parameterwise non-collapse.
3. No empirical, phenomenological, biological, or runtime claim is inferred at this level.
4. Later sections may add conditional identity, rigidity, and criterion theorems, but only by introducing their own explicit assumption blocks.

Part II

Typed Computable Model and State Spectral Gap

B9 Typed Computable Model: Embedded Constant-Disk Witness

This section repairs the type leaks of the previous computable model. The purpose is narrow: provide a fully typed witness for the transition grammar and the state spectral gap. It does *not* by itself certify parameterwise non-collapse; that burden is handled separately below.

B9.1 Golden Mean Parameter Space

Definition B21 (Golden Mean Shift Parameter Space)**.** Let $\Sigma = \{0, 1\}$ and

$$A = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}.$$

The Golden Mean Shift $S_A \subset \Sigma^{\mathbb{N}}$ is the subshift of finite type consisting of all binary sequences that avoid the block 11.

Definition B22 (Parameter Encoding)**.** Define

$$u(x) = \sum_{i=1}^{\infty} x_i \phi^{-i}, \qquad \phi = \frac{1+\sqrt{5}}{2}.$$

The image $\mathcal{U} := u(S_A) \subset [0, 1]$ is compact.

B9.2 Typed Banach Witness

Definition B23 (Ambient Space and Constant Embedding)**.** Let

$$X = [0, 1], \qquad E = \mathbb{R}^2, \qquad \mathcal{B} = C(X, E)$$

with norm

$$\|f\|_\infty := \sup_{s \in X} \|f(s)\|_2 .$$

Define the constant embedding

$$j : E \to \mathcal{B}, \qquad j(v)(s) = v.$$

Definition B24 (Invariant Constant Disk)**.** Let

$$D := \{v \in E : \|v\|_2 \le 1\}, \qquad \mathcal{I}_{\mathrm{c}} := j(D) \subset \mathcal{B}.$$

Definition B25 (Averaging Operator)**.** Define

$$A : \mathcal{B} \to E, \qquad A(f) := \int_0^1 f(\tau)\, d\tau.$$

Definition B26 (Disk Projection)**.** Define the Euclidean projection

$$P_E : E \to D, \qquad P_E(v) = \begin{cases} v, & \|v\|_2 \le 1, \\ \dfrac{v}{\|v\|_2}, & \|v\|_2 > 1. \end{cases}$$

Definition B27 (Typed Projection onto the Constant Disk)**.** Define

$$\aleph_{\mathcal{B}} : \mathcal{B} \to \mathcal{I}_{\mathrm{c}}, \qquad \aleph_{\mathcal{B}}(f) := j(P_E(A(f))).$$

Definition B28 (Constant-Valued Contraction Operator)**.** Define

$$\widetilde{K} : \mathcal{B} \to \mathcal{B}, \qquad (\widetilde{K}f)(s) := \frac{1}{2}\int_0^1 f(\tau)\, d\tau.$$

Definition B29 (Lifted Input Coupling)**.** Let $\Gamma : \mathcal{U} \to E$ be

$$\Gamma(u) = (u, 0).$$

Define its constant lift

$$\Gamma_{\mathcal{B}} : \mathcal{U} \to \mathcal{B}, \qquad \Gamma_{\mathcal{B}}(u) := j(\Gamma(u)).$$

Definition B30 (Typed Transition Operator)**.** For each $u \in \mathcal{U}$, define

$$T_u : \mathcal{B} \to \mathcal{I}_c, \qquad T_u(f) := \aleph_{\mathcal{B}}\big(\widetilde{K}f + \Gamma_{\mathcal{B}}(u)\big).$$

Remark B31 (Typing discipline)*.* No unembedded element of E is treated as an element of $\mathcal{B}$. The embedding j is always written explicitly. This removes the earlier type leak in which a disk in E was silently treated as a subset of a function space.

B9.3 Basic Operator Properties

Lemma B32 (The Embedding is Isometric)**.** *For every $v \in E$,*

$$\|j(v)\|_\infty = \|v\|_2 .$$

Proof. By definition,

$$\|j(v)\|_\infty = \sup_{s\in X} \|j(v)(s)\|_2 = \sup_{s\in X} \|v\|_2 = \|v\|_2 .$$

□

Lemma B33 (The Disk Projection is Non-Expansive)**.** *For all $v, w \in E$,*

$$\|P_E(v) - P_E(w)\|_2 \le \|v - w\|_2 .$$

Proof. P_E is the metric projection onto the closed convex set $D \subset \mathbb{R}^2$, so the standard Hilbert-space non-expansiveness property applies. □

Lemma B34 (The Typed Projection is Non-Expansive)**.** *For all $f, g \in \mathcal{B}$,*

$$\|\aleph_{\mathcal{B}}(f) - \aleph_{\mathcal{B}}(g)\|_\infty \le \|f - g\|_\infty .$$

Proof. Using the isometry of j, the non-expansiveness of P_E, and Jensen's inequality for the averaging operator,

$$\begin{aligned}
\|\aleph_{\mathcal{B}}(f) - \aleph_{\mathcal{B}}(g)\|_\infty &= \|P_E(A(f)) - P_E(A(g))\|_2 \\
&\le \|A(f) - A(g)\|_2 \\
&= \left\| \int_0^1 (f(\tau) - g(\tau))\, d\tau \right\|_2 \\
&\le \int_0^1 \|f(\tau) - g(\tau)\|_2 \, d\tau \le \|f - g\|_\infty .
\end{aligned}$$

□

Lemma B35 (Norm of the Lifted Contraction)**.**

$$\left\|\widetilde{K}\right\| \le \frac{1}{2}.$$

Proof. For any $f \in \mathcal{B}$,

$$\left\|\widetilde{K}f\right\|_\infty = \left\|\frac{1}{2}\int_0^1 f(\tau)\,d\tau\right\|_2 \le \frac{1}{2}\int_0^1 \|f(\tau)\|_2\,d\tau \le \frac{1}{2}\|f\|_\infty .$$

□

Theorem B36 (Typed Contraction on the Computable Witness)**.** *For each $u \in \mathcal{U}$, the map $T_u : \mathcal{B} \to \mathcal{I}_c$ is a strict contraction:*

$$\|T_u(f) - T_u(g)\|_\infty \le \frac{1}{2}\|f - g\|_\infty .$$

Proof. Combine non-expansiveness of $\aleph_{\mathcal{B}}$ with the bound on $\widetilde{K}$:

$$\|T_u(f) - T_u(g)\|_\infty \le \left\|\widetilde{K}f - \widetilde{K}g\right\|_\infty \le \frac{1}{2}\|f - g\|_\infty .$$

□

Corollary B37 (Unique Fixed Point in the Typed Witness)**.** *For each $u \in \mathcal{U}$, there exists a unique $f_u^* \in \mathcal{I}_c$ such that*

$$T_u(f_u^*) = f_u^*.$$

Proof. The Banach fixed-point theorem applies on the complete metric space $(\mathcal{B}, \|\cdot\|_\infty)$. □

Proposition B38 (Typed Non-Vacuity of the Computable Model)**.** *The embedded constant-disk Golden Mean model is a fully typed, non-vacuous witness of the transition grammar:*
1. every operator has explicit domain and codomain;
2. the transition map T_u is well-defined on $\mathcal{B}$;
3. the family admits unique fixed points and a compact parameter image;
4. the state contraction bound is explicit and computable.

Proof. Items (1)–(3) follow from the definitions above and the contraction theorem. Compactness of the parameter image follows because $\mathcal{U}$ is compact and $u \mapsto f_u^*$ is continuous under a uniform contraction constant. □

Remark B39 (Why this witness does not certify H5). The image of T_u lies in the constant subspace $\mathcal{I}_c$. Therefore its fixed point is constant by construction. This witness is correct and useful for typing and state spectral analysis, but it cannot certify the parameterwise anti-constant hypothesis H5 from Section B7.

B9.4 Separate Anti-Collapse Witness

Definition B40 (Affine Non-Constant Witness Space)**.** Let

$$\mathcal{H}_{\mathrm{ac}} := L^2([0,1]), \qquad \mathcal{N} := \{c\mathbf{1} : c \in \mathbb{R}\}, \qquad h(s) := \sin(2\pi s).$$

Choose r with $0 < r < \|h\|_{L^2}$ and define the closed convex affine set

$$\mathcal{I}_{\mathrm{ac}} := h + \{g \in \mathcal{H}_{\mathrm{ac}} : \|g\|_{L^2} \le r\}.$$

Lemma B41 (The Anti-Collapse Witness Set Avoids Constants)**.**

$$\mathcal{I}_{\mathrm{ac}} \cap \mathcal{N} = \varnothing.$$

Proof. The distance from h to the constant subspace $\mathcal{N}$ equals $\|h\|_{L^2}$ because h has mean zero and is orthogonal to constants. Since $r < \|h\|_{L^2}$, the radius-r ball around h does not intersect $\mathcal{N}$. □

Definition B42 (Anti-Collapse Witness Dynamics)**.** Let $\mathcal{U} = [0,1]$, choose $\beta \in (0, 1/2)$, and define

$$K_{\mathrm{ac}} f := \beta f, \qquad \Gamma_{\mathrm{ac}}(u) := \frac{u}{4} h.$$

Let P_{ac} denote the metric projection onto $\mathcal{I}_{\mathrm{ac}}$, and define

$$T_u^{\mathrm{ac}}(f) := P_{\mathrm{ac}}(K_{\mathrm{ac}} f + \Gamma_{\mathrm{ac}}(u)).$$

Proposition B43 (Separate Witness for H5)**.** *The family $\{T_u^{\mathrm{ac}}\}_{u \in \mathcal{U}}$ satisfies Hypotheses B1–B17 of Section 1.*

Proof. 1. $\mathcal{H}_{\mathrm{ac}}$ is a Hilbert space and $\mathcal{I}_{\mathrm{ac}}$ is non-empty, closed, and convex.
2. $\|K_{\mathrm{ac}}\| = \beta < 1$.
3. Completeness is inherited from $L^2([0,1])$.
4. $\mathcal{U} = [0,1]$ is compact and Γ_{ac} is continuous.
5. For every constant $c \in \mathcal{N}$ and every $u \in \mathcal{U}$, the value $T_u^{\mathrm{ac}}(c)$ belongs to $\mathcal{I}_{\mathrm{ac}}$, while $\mathcal{I}_{\mathrm{ac}} \cap \mathcal{N} = \varnothing$. Hence $T_u^{\mathrm{ac}}(c) \neq c$.

□

Remark B44. BaseCore therefore distinguishes two roles cleanly: the Golden Mean witness certifies typed operator correctness and state contraction; the affine anti-collapse witness certifies the stronger H5 burden needed for parameterwise non-collapse.

B10 State Spectral Gap and Parameter-Family Persistence

Theorem B45 (State Spectral Gap)**.** *Fix $u \in \mathcal{U}$ in the typed computable model. Assume T_u is Fréchet differentiable at its fixed point f_u^*. Then*

$$\rho(DT_u(f_u^*)) \leq \|DT_u(f_u^*)\| \leq \left\|\widetilde{K}\right\| \leq \frac{1}{2}.$$

Consequently the state spectral gap satisfies

$$\Delta_{\mathrm{state}} \geq 1 - \left\|\widetilde{K}\right\| \geq \frac{1}{2}.$$

Proof. At every point where the projection P_E is Fréchet differentiable, its derivative has operator norm at most 1 because P_E is non-expansive. Since $\aleph_{\mathcal{B}} = j \circ P_E \circ A$,

$$DT_u(f_u^*) = D\aleph_{\mathcal{B}}\Big(\widetilde{K} f_u^* + \Gamma_{\mathcal{B}}(u)\Big) \circ \widetilde{K},$$

and therefore

$$\|DT_u(f_u^*)\| \leq \|D\aleph_{\mathcal{B}}\| \left\|\widetilde{K}\right\| \leq \frac{1}{2}.$$

The spectral-radius bound follows from $\rho(L) \leq \|L\|$ for bounded linear operators. □

Theorem B46 (Relaxation-Time Bound)**.** *Under the hypotheses of Theorem B45, any discrete-to-continuous relaxation convention of the form*

$$\tau_{\text{state}} = -\frac{1}{\ln \rho}$$

with $\rho = \rho(DT_u(f_u^*))$ *gives*

$$\tau_{\text{state}} \leq \frac{1}{\ln 2}.$$

Proof. Theorem B45 gives $\rho \leq 1/2$. Since $-\ln(\rho) \geq \ln 2$, the bound follows. □

Proposition B47 (Parameter-Family Persistence is not State Spectrum)**.** *Let* $F : \mathcal{U} \to \mathcal{B}$ *be the fixed-point family* $F(u) = f_u^*$*. Any persistence of parameter labels or family coordinates belongs to the variation of* F *over* $\mathcal{U}$*, not to the state-transition spectrum of* $DT_u(f_u^*)$ *for fixed* u*.*

Proof. The state-transition spectrum is defined for the linearization of T_u at fixed parameter u. Family-level transport instead concerns the map F between parameter space and fixed-point set. These are different operators on different domains. Therefore a family-persistence mode cannot be used to introduce a unit eigenvalue into the contraction spectrum of $DT_u(f_u^*)$. □

Remark B48 (What has been removed). The previous text incorrectly mixed strict contraction with a state-spectrum eigenvalue $\lambda = 1$. BaseCore no longer does that. Any neutral or persistent mode must now be interpreted as a parameter-family bookkeeping mode, not as an eigenmode of the contractive state Jacobian.

Part III

Identity, Rigidity, and Conditional Non-Simulability

B11 Projective Systems and Inverse-Limit Identity

B11.1 Projective Setup

Definition B49 (Projective System of State Spaces)**.** Let $\{S_t\}_{t\in\mathbb{N}}$ be a family of non-empty state spaces together with projection maps

$$\pi_{t+1\to t} : S_{t+1} \to S_t$$

satisfying the compatibility rule

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t} \qquad \forall t \in \mathbb{N}.$$

Definition B50 (Identity as Inverse Limit)**.** The identity object associated with the projective system is

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_{t\in\mathbb{N}} \in \prod_{t\in\mathbb{N}} S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \ \forall t \right\}.$$

Remark B51 (Identity neutrality)*.* In BaseCore, $\mathcal{I}$ is a structural inverse-limit object only. It carries no automatic phenomenological, biological, or metaphysical interpretation.

Hypothesis B52 (Topological Regularity)**.** *Each S_t is compact Hausdorff and each $\pi_{t+1\to t}$ is continuous and surjective.*

Hypothesis B53 (NFD: Non-Finite Determination)**.** *For every $t \in \mathbb{N}$, the canonical projection*

$$p_t : \mathcal{I} \to S_t, \qquad p_t((x_k)_k) = x_t$$

is not injective. Equivalently,

$$\forall t \ \exists x, y \in \mathcal{I} \quad x \neq y \quad \textit{and} \quad p_t(x) = p_t(y).$$

Remark B54 (Persistent future branching as an operational form of NFD)*.* A sufficient operational strengthening of Hypothesis B53 is: for every t there exists $k > t$ and $x_t \in S_t$ such that the fiber of $\pi_{k\to t}$ over x_t contains at least two globally extendable compatible histories. BaseCore keeps the weaker NFD statement as the primary assumption because it is exactly what the non-locality theorem needs.

B11.2 Non-Locality of Identity

Theorem B55 (Non-Locality of Identity under NFD)**.** *Assume Hypotheses B52 and B53. Then no canonical finite-time observation determines identity. More precisely, for every t the canonical projection*

$$p_t : \mathcal{I} \to S_t$$

is not injective, so identity is not recoverable from any finite slice through the projective observation maps.

Proof. This is exactly Hypothesis B53. The theorem restates the structural consequence in claim-boundary language: finite slices can participate in the inverse-limit object, but no single canonical slice determines it. □

Remark B56 (What is *not* proved here)*.* BaseCore does not prove that no arbitrary set-theoretic injection $\mathcal{I} \to S_t$ can exist in every imaginable inverse system. It proves the narrower and correct claim that the canonical projective observations p_t are not injective under Hypothesis B53.

Proposition B57 (No Canonical Progressive Recovery under NFD)**.** *Assume Hypothesis B53. No finite canonical prefix map determines the inverse-limit identity object.*

Proof. If some finite canonical prefix determined identity, then in particular one of the canonical coordinate maps would become injective after finite truncation, contradicting Hypothesis B53. □

B12 Rigidity under Admissible Perturbations

B12.1 Metric Projective Assumptions

Assumption B58 (RIG-1: Common Metric Stages)**.** *Both towers use the same nonempty complete metric stages (S_t, d_t) and continuous bonding maps. Use the weighted history metric below with $w_t \operatorname{diam}(S_t) \to 0$. This fixes the ambient space used for both Hausdorff directions.*

Assumption B59 (RIG-2: Controlled Propagation)**.** *For the fiber lifting constants $K_t > 0$ of RIG-3 define $q_t = w_{t+1}K_t/w_t$. Require $C = \sup_{0\le k\le t<\infty} \prod_{j=k}^{t} q_j < \infty$. In particular $q_t \le 1$ at each stage suffices.*

Assumption B60 (RIG-3: Exact Fiber Lifts in Both Towers)**.** *For each $f \in \{\pi_t, \widetilde{\pi}_t\}$, $u \in S_{t+1}$ and $v \in S_t$, there exists $u' \in S_{t+1}$ with $f(u') = v$ and $d_{t+1}(u, u') \le K_t d_t(f(u), v)$. The same K_t bounds both lifts. No continuous selection is required.*

Assumption B61 (RIG-4: Typed Projection Defects)**.** *Write $\pi_t = \pi_{t+1\to t}$ and $\widetilde{\pi}_t = \widetilde{\pi}_{t+1\to t}$. Require $\sup_{u\in S_{t+1}} d_t(\pi_t(u), \widetilde{\pi}_t(u)) \le \epsilon_t < \infty$.*

Assumption B62 (RIG-5: Summable Deformation)**.** *Require $E_w = \sum_{t\ge 0} w_t\epsilon_t < \infty$.*

Definition B63 (Weighted Deformation Size)**.** The weighted deformation size is $\|\epsilon_\bullet\|_w = E_w$. When ϵ_t is the exact supremum projection defect this specifies a cost on the common-stage perturbations. An independently defined intervention cost needs its own quantitative relation to this sum.

Definition B64 (Weighted Product Metric)**.** For nonempty metric spaces (S_t, d_t) and weights $0 < w_t < \infty$, write

$$d_w(x, y) = \sup_{t\ge 0} w_t d_t(x_t, y_t), \qquad P_b = \{x \in \textstyle\prod_t S_t : d_w(x, b) < \infty\}.$$

The formula defines an extended metric on the product and an ordinary metric on each finite-distance component P_b.

Proposition B65 (Finiteness, Completeness and Compactness of Weighted Histories)**.** *The formula is finite on every pair in the full product exactly when the weighted diameters are bounded outside a finite set of stages. For finite-diameter stages this is equivalent to $B = \sup_t w_t \operatorname{diam}(S_t) < \infty$. For any nonempty inverse limit I and $b \in I$, it is finite on all of I exactly when $I \subseteq P_b$. If the coordinate spaces are complete and the bonding maps continuous, P_b and $I \cap P_b$ are complete. If $w_t \operatorname{diam}(S_t) \to 0$, the weighted metric induces the product topology and*

$$0 \le d_w(x, y) - \max_{0\le t\le N} w_t d_t(x_t, y_t) \le \sup_{t>N} w_t \operatorname{diam}(S_t).$$

In this case compact coordinate spaces give a compact inverse limit. Surjective bonding maps give a compatible history through each stage point.

Proof. The extended triangle inequality follows by taking suprema. A uniform bound outside finitely many stages suffices since each of the exceptional distances is individually finite. If no such bound exists, choose distinct stages t_j and pairs with weighted distances greater than j. Combining these pairs in two product points gives an infinite distance. The bound $B < \infty$ characterizes *boundedness* of the full-product metric. The inverse-limit condition follows from the triangle inequality through b.

A d_w-Cauchy sequence converges in each coordinate. Passing to each coordinate limit in the same uniform Cauchy bound proves convergence in d_w. The limit remains in P_b by the triangle inequality and in I by continuity of the bonding maps. With vanishing weighted diameter tails, a

finite prefix controls the metric up to the displayed remainder, while each coordinate projection is $1/w_t$-Lipschitz. This proves the topology assertion. Conversely, on the full product with a finite metric, if the weighted diameters fail to vanish there are infinitely many coordinates with weighted diameter at least some $\eta > 0$. At any fixed product point, for each such coordinate one can change that coordinate by weighted distance greater than $\eta/3$, by the triangle inequality. The resulting sequence converges in the product topology but fails to converge in d_w. The converse may fail on a restricted set of compatible histories. The inverse limit is closed in the compact product because it is the intersection of the closed compatibility equations. Surjectivity gives successive lifts and the earlier coordinates are determined by projection. □

Example B66 (Two Distinct Tail Failures)**.** On $S_t = [0, 1]$ with identity bonding maps and $w_t = 2^{-t}$, choose $d_t(s, u) = 4^t|s - u|$. The two constant histories 0 and 1 have distance $+\infty$. For a different failure, take $S_t = \{0, 1\}^t$ with prefix deletion and $d_t = 2^t$ times the discrete metric. With the same weights every distinct pair of histories has distance 1. The inverse limit has a compact product topology and a noncompact discrete metric topology. Flipping a bit at successively later levels gives product convergence to the unflipped history while all the weighted distances remain 1.

Theorem B67 (Conditional Stability of Inverse-Limit Identity)**.** *Under RIG-1 through RIG-5 the two limits $I = \mathcal{I}$ and $\widetilde{I} = \widetilde{\mathcal{I}}$ are nonempty. Every history in either limit has a compatible history in the other within distance CE_w. Thus $d_H(I, \widetilde{I}) \le CE_w$.*

Proof. Given a history x of π, set $y_0 = x_0$. Having chosen y_t, lift it under $\widetilde{\pi}_t$ at x_{t+1}. The exact lift y_{t+1} obeys

$$d_{t+1}(y_{t+1}, x_{t+1}) \le K_t(d_t(y_t, x_t) + \epsilon_t)$$

by the fiber estimate and $\pi_t(x_{t+1}) = x_t$. For $a_t = w_t d_t(y_t, x_t)$ this gives $a_{t+1} \le q_t a_t + q_t w_t \epsilon_t$, with $a_0 = 0$. Iteration bounds a_{t+1} by $\sum_{k=0}^{t}(\prod_{j=k}^{t} q_j) w_k \epsilon_k$, at most CE_w. The reverse construction uses the fiber estimate for π_t and gives the other directed Hausdorff bound. Each limit is nonempty by successive lifts from any S_0 point. Taking suprema proves the estimate in the common history metric. □

Corollary B68 (Conditional Isomorphism Requires Extra Structure)**.** *The Hausdorff bound alone is a statement about the distance between sets. A homeomorphism follows under the commuting-bijection hypotheses below.*

Proof. The following transport proposition gives the sufficient condition. Example B71 exhibits the failure of the unconditional implication. □

Proposition B69 (Transport by Commuting Coordinate Bijections)**.** *Suppose the weighted metric induces the product topology and homeomorphisms $h_t : S_t \to S_t$ satisfy $\widetilde{\pi}_t h_{t+1} = h_t \pi_t$. Then $H(x)_t = h_t(x_t)$ maps the limits homeomorphically, with inverse $H^{-1}(y)_t = h_t^{-1}(y_t)$. Its displacement in the common history space is at most $\sup_t w_t \sup_{z \in S_t} d_t(z, h_t(z))$.*

Proof. The displayed square and its inverse give exact compatibility of both images. Coordinate continuity in the product gives continuity of both induced maps. The displacement estimate follows coordinate by coordinate. □

Example B70 (A Ramified Tower with Positive Deformation and Coherent Transport)**.** Take $S_t = (\mathbb{R}/\mathbb{Z}) \times \{0, 1\}^t$ with the maximum of the circle distance (circumference one) and the discrete prefix distance. Let π_t keep the angle and delete the last bit, and take $w_t = 2^{-t}$. For an intervention at stage k with amplitude $0 < a < 1/2$, replace π_k by prefix deletion and angle shift $-a$. Set h_t to the identity for $t \le k$ and to angle shift a for $t > k$. The inverses shift by $-a$ on those same levels. The squares commute at k by cancellation and at every other level by

commutation with prefix deletion. Each limit has at least two histories over every finite cut, since a later bit may be flipped.

All bonding maps are 1-Lipschitz and satisfy the fiber estimate with $K_t = 1$: set the required angle and prefix, retaining the last bit. Thus $q_t = 1/2$. The chosen dimensionless intervention cost is $E_w = w_k a$, and $\sup_x d_w(x, Hx) = w_{k+1} a = E_w/2$. More precisely, for $b = w_{k+1}$,

$$d_H(I, \widetilde{I}) = \frac{ba}{1+b} > 0.$$

Indeed an original history has constant angle θ and a perturbed one has angle φ up to k and $\varphi + a$ thereafter. If the two histories are at distance r, the circle triangle inequality gives $a \le r + r/b$ from stages 0 and $k+1$. Choosing identical bits and $\varphi = \theta - ba/(1+b)$ attains equality. That choice is reversible for every perturbed history and gives both directed distances. For the sector of these single jumps the infimum of the cost-to-deformation ratio is 2, since the ratio is $2(1+b)$ at stage k. This uses the declared projection-defect cost. An entropy-based cost requires a separate comparison.

Example B71 (Hausdorff Stability Without a Homeomorphism)**.** Take $S_t = \{0, \dots, 2^t - 1\}$ with discrete distance scaled by $r > 0$, $w_t = 2^{-t}$, and $\pi_t(j) = \lfloor j/2 \rfloor$. For $t \ge 1$, change only the image of $j = 1$ to 1, keeping $\widetilde{\pi}_0 = \pi_0$. Both maps remain surjective and satisfy the fiber estimate with $K_t = 1$. The weighted intervention cost is $r \sum_{t \ge 1} 2^{-t} = r$. The all-zero history of the perturbed tower is isolated by the cylinder $y_1 = 0$ because each subsequent zero has exactly one lift. The original tower is the binary-prefix space and has no isolated histories. Thus the limits are not homeomorphic, even though the Hausdorff bound tends to zero with r. The perturbation hypotheses for the Hausdorff estimate all hold.

Definition B72 (Ontological Mass)**.** Let $E_w(\delta)$ denote the weighted energy of an admissible perturbation family δ. Define

$$M_\Omega := \inf \left\{ \frac{E_w(\delta)}{d_H(\mathcal{I}, \widetilde{\mathcal{I}}_\delta)} \;\middle|\; \delta \text{ admissible},\ d_H(\mathcal{I}, \widetilde{\mathcal{I}}_\delta) > 0 \right\},$$

with the convention $M_\Omega = +\infty$ if no admissible finite-energy perturbation yields positive deformation of the identity object.

Remark B73 (Interpretation discipline)*.* M_Ω is a deformation-resistance scalar, not a metaphysical substance. BaseCore uses it only as a conditional invariant attached to admissible perturbation models.

B13 Conditional Non-Simulability

Assumption B74 (NS-1: Target and Observable Family)**.** *Fix a nonempty set of compatible histories $X \subseteq \mathcal{I}$, a family of queries J, and responses $F_j : X \to (Y_j, \rho_j)$. The query can specify an observable, a time horizon, or a declared intervention. All queries refer to the same encoded history. The intervention protocol and parameters are known to both parties and fixed independently of the history, except for information explicitly included in the encoded state. Non-finite determination and deformation resistance are separate properties of a target, measured using their own definitions.*

Assumption B75 (NS-2: Binary Representation Budget)**.** *A deterministic encoder $c : X \to \{0,1\}^B$ has at most 2^B possible representations. The budget includes all history-dependent information accessible to the decoder. An adaptive binary transcript of depth at most B satisfies the same bound when its query tree is fixed in advance and the choices and termination are determined by its previous bits. An exact real parameter is outside this finite-alphabet class.*

Assumption B76 (NS-3: Uniform Response Fidelity)**.** *The decoders $g_j : \{0,1\}^B \to Y_j$ must satisfy $\rho_j(g_j(c(x)), F_j(x)) \le \eta$ for every $x \in X$ and $j \in J$. The tolerance is in the declared response units, after normalizing them if necessary. Faithfulness of additional structures, such as the bonding maps or the intervention cost sector, imposes separate constraints.*

Theorem B77 (Finite-Budget Obstruction by Response Separation)**.** *Under NS-1–NS-3, let $A \subseteq X$ be a finite set such that for each distinct $x, y \in A$ there is j with $\rho_j(F_j(x), F_j(y)) > 2\eta$. Then $|A| \le 2^B$. For exact fidelity, a code exists exactly when the number of equivalence classes of $x \sim y$ defined by $F_j(x) = F_j(y)$ for all j is at most 2^B.*

Proof. If $c(x) = c(y)$, the decoded value for every given query is the same, so the triangle inequality gives $\rho_j(F_j(x), F_j(y)) \le 2\eta$. Thus c is injective on A. At zero error, different response classes must use different representations. Conversely, assign a distinct representation to each class and decode its common response for each query. □

Theorem B78 (Approximation Barrier)**.** *On the binary prefix tower $S_t = \{0,1\}^t$ with weights $w_t = 2^{-t}$ and discrete stage metrics, use $Y = X = \{0,1\}^{\mathbb{N}}$ as the decoded history space, with $d_w(x, y) = 2^{-m}$ at the first differing bit $m \ge 1$. Among all encoders with at most 2^B states and decoders into X, the minimum worst-case d_w error is $2^{-(B+1)}$. Exact recovery of all finite prefix responses requires unbounded B as the query horizon increases.*

Proof. Keep the first B bits and pad the rest with zeros. The first possible mismatch is at $B+1$. For the lower bound, take all 2^{B+1} histories with arbitrary first $B+1$ bits and zero tail. Two share a code. An error strictly less than $2^{-(B+1)}$ would require the common decoded history to share the first $B+1$ bits with both, which is impossible. At horizon t all 2^t prefixes occur, so exact recovery of that response needs $B \ge t$, attained by storing those bits. □

Example B79 (One Exact Real Encodes Every Binary History)**.** The map $b \mapsto \sum_{j \ge 1} 2b_j 3^{-j}$ is injective: at the first unequal bit j the leading difference is $2 \cdot 3^{-j}$ and the remaining tail is bounded by 3^{-j}. On the image, inverse decoding followed by each prefix projection recovers all the canonical observations from one exact real. As a separate intervention control, any declared family that leaves the target limit set fixed has $M_\Omega = +\infty$ by the defining convention, even if the allowed interventions have finite cost. Neither fact supplies a finite binary budget.

Remark B80 (What has been removed)*.* The resource obstruction concerns the declared responses, metric tolerance and binary budget. Infinite-precision scalar encodings and finite-cost interventions require independent hypotheses before a binary-budget theorem can be applied to them. In the CCR program, non-finite determination and infinite resistance are separately specified; any stronger obstruction to faithful embeddings must state its category and preserved observations. Exact response classes determine whether a finite sufficient statistic exists for the declared queries.

B14 Section Summary

1. Identity is an inverse-limit object, not a canonical finite slice.
2. Non-locality now depends explicitly on the NFD assumption.
3. Rigidity is stated as metric stability under explicit perturbation-model assumptions, not as automatic isomorphism.
4. Uniform observational fidelity obeys a binary-budget separation bound. The intervention-resistance class is determined separately.
5. BaseCore therefore preserves theorem strength where assumptions support it and weakens only those statements that were previously over-absolute.

Part IV

Regime Structure and Continuity Constraints

B15 Formal Setup

B15.1 Review of Identity Framework

We assume the identity framework of [22]. Let $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ be a projective system of observable state spaces. The identity object is the inverse limit:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_t \in \prod_t S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \;\; \forall t \right\}. \tag{1}$$

The ontological mass M_Ω quantifies resistance to identity deformation. We recall that:

- $M_\Omega = 0$: zero deformation-cost infimum;
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR relative to the declared perturbation class).

B15.2 Abstract Phenomenological Space

Definition B81 (Phenomenological Space)**.** A phenomenological space is a nonempty set $\mathcal{E}$ whose elements are called *phenomenological configurations*. No topology, metric, or algebraic structure is assumed a priori.

Remark B82. We do not claim that $\mathcal{E}$ represents "conscious states" or "qualia". It is an abstract mathematical space whose interpretation is not specified by the formalism. The space $\mathcal{E}$ is introduced purely as the codomain of a structural assignment; its ontological status is deliberately left open.

B15.2.1 Why a Separate Space?

One might ask: why not use $\mathcal{I}$ itself as the phenomenological space? The answer is structural:

Proposition B83 (Endomorphisms Preserving All Canonical Projections)**.** *If $F : \mathcal{I} \to \mathcal{I}$ satisfies $p_t \circ F = p_t$ for every t, then $F = \mathrm{Id}$. This does not exclude endomorphisms with nonidentity coordinate actions.*

Proof. For each x, $F(x)$ and x have identical coordinates, hence are the same inverse-limit element. For contrast, on binary prefix spaces, bitwise complement defines an endomorphism induced by complement on every prefix; it commutes with the bonding maps but does not preserve each coordinate value. A separate codomain is a modeling choice, not a logical necessity for nonidentity maps. □

A separate codomain permits independent structural readouts; it is not forced by the universal property alone.

B15.3 Phenomenological Assignment

Definition B84 (Phenomenological Assignment)**.** A phenomenological assignment is a family of total maps on the declared observable state spaces:

$$\phi_t : S_t \to \mathcal{E} \tag{2}$$

indexed by $t \in \mathbb{N}$.

Definition B85 (Induced Limit Assignment)**.** Given a phenomenological assignment $\{\phi_t\}$, the induced limit assignment, whenever it exists, is the map:

$$\Phi : \mathcal{I} \to \mathcal{E}, \quad \Phi((x_t)_t) := \lim_{t\to\infty} \phi_t(x_t), \tag{3}$$

The convergence statements below use the complete metric codomain of Hypothesis B103. Partial readouts may instead be studied on a specified common domain, with all projections and factorizations restricted to that domain.

B15.3.1 Uniqueness of the Limit Assignment

Proposition B86 (Uniqueness and Uniform Error of the Limit Assignment)**.** *Under E1, each trajectory has a unique limit assignment and*

$$d_{\mathcal{E}}\big(\Phi(x), \phi_n(p_n x)\big) \le A_n \quad \textit{for all } x \in \mathcal{I}.$$

Proof. For $m > n$, the triangle inequality bounds the distance between the mth and nth readings by $\sum_{t=n}^{m-1} a_t$. Thus each sequence is Cauchy. Completeness supplies its limit, uniqueness follows from metric separation, and passage to the limit gives the bound. Since $A_n \to 0$ independently of x, convergence is uniform. □

B15.4 Axioms for Phenomenological Coherence

We declare a controlled assignment model. E0 is a nontriviality convention, E1 supplies convergent readouts, and E2 is a separately checked non-factorization requirement.

Axiom B87 (E0: Non-Triviality)**.** *The phenomenological space satisfies $|\mathcal{E}| \ge 2$, i.e., there exist at least two distinct phenomenological configurations.*

Axiom B88 (E1: Controlled Asymptotic Compatibility)**.** *In the complete metric codomain of Hypothesis B103, there are numbers $a_t \ge 0$, independent of the trajectory, with $\sum_{t=0}^{\infty} a_t < \infty$, such that*

$$d_{\mathcal{E}}\big(\phi_{t+1}(p_{t+1}x), \phi_t(p_t x)\big) \le a_t \qquad (x \in \mathcal{I},\ t \ge 0).$$

Here $p_t : \mathcal{I} \to S_t$ is the canonical projection. Define the tail budget $A_n = \sum_{t=n}^{\infty} a_t$. Exact equality is the special case $a_t = 0$, not a requirement of E1.

Axiom B89 (E2: Non-Locality)**.** *There exists no function $f : S_t \to \mathcal{E}$ (for any fixed t) such that for all trajectories $(x_t)_t \in \mathcal{I}$:*

$$\Phi((x_t)_t) = f(x_t). \tag{4}$$

B15.5 Status of the Assignment Conditions

The following results separate sufficient assumptions, a non-factorization requirement and their logical dependencies.

B15.5.1 Nontriviality of the Codomain

Theorem B90 (E0 is Necessary)**.** *If $|\mathcal{E}| = 1$, then every phenomenological assignment is constant and no phenomenological classification is possible.*

Proof. If $\mathcal{E} = \{e_0\}$, any total assignment takes the value e_0 on every trajectory. It distinguishes no trajectories and factors through every canonical projection. It therefore cannot meet E2 on a nonempty domain. □

B15.5.2 Controlled Convergence

Theorem B91 (Controlled Compatibility Is Sufficient)**.** *E1 guarantees existence of the limit with the error bound in Proposition B86. Absolute summability is sufficient, not necessary for an arbitrary convergent family.*

Proof. Existence and the bound were proved above. To disprove necessity in general, take readouts constant in the trajectory with values $s_n = \sum_{k=1}^{n} (-1)^{k+1}/k$ in $[0, 1]$. The even partial sums increase, the odd partial sums decrease, and their difference tends to zero, so they have a common limit. Their successive absolute increments form the divergent harmonic series; no summable uniform majorant exists. This example is a control for the necessity claim, not a realization of E2. □

B15.5.3 Finite-Cut Non-Factorization

Theorem B92 (Non-Factorization Is an Independent Requirement)**.** *E2 excludes readouts determined by a single canonical finite cut. The inverse-limit construction and E1 alone do not imply E2.*

Proof. Take any nonempty inverse limit and a constant family $\phi_t = e_0$ in a codomain containing at least two points. E1 holds with zero tail, but $\Phi = e_0$ factors through every projection. Thus E2 must be checked separately. The nonempty construction in Theorem B95 checks it without assuming it. □

B15.6 Logical Dependencies and Realizability

Theorem B93 (Logical Dependencies of the Assignment Conditions)**.** *On a nonempty trajectory domain, E2 implies E0. E1 does not imply E2. There exist assignments satisfying E0, the revised E1 and E2 simultaneously. Consequently these three conditions are not an independent minimal axiom basis as originally stated.*

Proof. If the codomain were a singleton, every assignment would be constant and factor through p_0, contradicting E2. The preceding constant-family control proves the second assertion. Theorem B95 proves the third by construction. E0 is retained as an explicit convention, although E2 entails it in the nonempty case. □

B15.7 Rejected Alternatives

We briefly discuss axiom candidates that were considered and rejected.

B15.7.1 Rejected: Continuity Axiom

One might require Φ to be continuous. This is not required by the chosen model because:
- The complete-metric model permits discontinuous finite readouts; continuity is a property of the regular-readout subclass, not part of the unrestricted assignment condition;
- Continuity is derived from controlled approximation by continuous finite readouts; a scalar CCR label alone is insufficient.

B15.7.2 Rejected: Surjectivity Axiom

One might require Φ to be surjective. This is not required by the chosen model because:
- A readout need not exhaust a codomain chosen to include other regimes;
- Surjectivity is a strong structural claim not derivable from identity constraints.

B15.7.3 Rejected: Injectivity Axiom

One might require Φ to be injective. This is not required by the chosen model because:
- Many trajectories may share phenomenological configurations;
- Injectivity is not imposed on the full assignment class; it is a consequence in the concrete coercive realization. An injective map alone is not necessarily a topological embedding.

B15.8 Structural Interpretation

Axiom	**Structural Meaning**	**Status**
E0	Nontrivial codomain	E2 consequence
E1	Summable increments	Uniform limit
E2	Finite-cut non-factorization	Required

B15.9 Closure of Section 2

The limit is uniquely defined with a uniform tail bound. E2 remains a substantive non-factorization requirement; E0 is redundant when E2 holds on a nonempty domain. Exact compatibility cannot coexist with E2, whereas the controlled condition has the explicit realization below. No existence or minimality claim is obtained merely by naming an axiom.

B15.10 Phenomenological Neutrality Clause

Throughout this paper, the phenomenological space $\mathcal{E}$ is treated as a *purely abstract topological and metric structure*. No semantic, psychological, biological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "phenomenological" refers exclusively to the role of $\mathcal{E}$ as an abstract codomain of structural compatibility for inverse-limit identities, and does not presuppose the existence of consciousness, subjective experience, or any particular phenomenological content.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{E}$ in relation to experiential, cognitive, or phenomenological phenomena lies strictly outside the formal scope of this paper.

B15.11 Exact Compatibility Obstruction

Proposition B94 (Exact Compatibility Forces Finite-Cut Factorization)**.** *On a nonempty trajectory domain, if $\phi_{t+1}(p_{t+1}x) = \phi_t(p_t x)$ for every x, t, then the limit assignment satisfies $\Phi = \phi_0 \circ p_0$. It cannot satisfy E2. The controlled E1 above replaces that exact equality, rather than assuming it together with E2.*

Proof. Induction makes all readings equal to the first. Their limit is therefore $\phi_0(p_0 x)$. For partial readouts with a declared common domain, the same argument holds on that domain. If a total map on S_0 is required, extend the defined first-cut values arbitrarily outside the projected common domain, using any fixed element of the nonempty codomain. Thus partiality alone does not remove the obstruction. □

B15.12 A Nonempty Controlled Assignment Model

Theorem B95 (Ternary Realization of E0–E2)**.** *Let $S_t = \{0,1\}^{t+1}$ with bonding maps that delete the last bit. Identify its inverse limit with $\mathcal{I} = \{0,1\}^{\mathbb{N}}$ and let $\mathcal{E} = [0,1]$. Define*

$$\phi_t(b_0, \ldots, b_t) = \sum_{j=0}^{t} \frac{2b_j}{3^{j+1}}, \qquad \Phi(b) = \sum_{j=0}^{\infty} \frac{2b_j}{3^{j+1}}.$$

These assignments satisfy E0, controlled E1, and E2, with

$$a_t = \frac{2}{3^{t+2}}, \qquad A_t = \frac{1}{3^{t+1}}, \qquad \sup_b |\Phi(b) - \phi_t(p_t b)| = A_t.$$

For distinct histories first differing at position m, put $d_{\mathcal{I}}(b,c) = 3^{-(m+1)}$, and put $d_{\mathcal{I}}(b,b) = 0$. Then

$$d_{\mathcal{I}}(b,c) \le |\Phi(b) - \Phi(c)| \le 3 d_{\mathcal{I}}(b,c).$$

The domain is compact and complete, and the readout is injective and bi-Lipschitz onto its image. Thus a global lower bound is realized here rather than inserted without a witness.

Proof. Every infinite bit sequence defines a compatible prefix history, so the domain is nonempty. Geometric summation puts every value in $[0,1]$ and gives the increment and tail bounds. The all-one tail attains the tail bound. For any fixed t, a zero history and a history with its sole one at position $t+1$ have the same prefix but images differing by $2/3^{t+2} > 0$. No function of p_t determines Φ, proving E2 for every finite cut.

At the first differing position m, the principal difference has magnitude $2/3^{m+1}$; the remaining tail has magnitude at most $1/3^{m+1}$. Subtraction and addition give the two bounds. This metric is the weighted supremum of the discrete prefix metrics with weights $3^{-(t+1)}$; bounded prefix diameters ensure finiteness. A Cauchy sequence for $d_{\mathcal{I}}$ eventually stabilizes each finite prefix, which defines its limit. Prefix cylinders give finite covers of arbitrarily small diameter; the space is totally bounded and complete, hence compact. The two bounds establish injectivity and quantitative continuity. Compactness and continuity give properness into the Hausdorff codomain, and a later-bit flip in every neighborhood gives local non-collapse. Thus this realization also supplies all four clauses of the stated readout regularity hypothesis. A scalar real-valued readout still encodes the complete sequence: E2 excludes the specified canonical finite cuts, not every possible coding or every finite-dimensional representation. □

Theorem B96 (Sharp Finite-Cut Approximation Error)**.** *For the ternary realization and an arbitrary decoder $g : S_t \to \mathbb{R}$, let*

$$R_t(g) = \sup_{b \in \mathcal{I}} |\Phi(b) - g(p_t b)|.$$

Then the exact minimax value is

$$\inf_g R_t(g) = \frac{1}{2\,3^{t+1}}.$$

It is attained by $g_*(s) = \phi_t(s) + 1/(2\,3^{t+1})$.

Proof. For a fixed prefix s, the all-zero and all-one suffixes give endpoints $\phi_t(s)$ and $\phi_t(s) + A_t$. For any number $g(s)$, the triangle inequality forces at least one endpoint error to be at least $A_t/2$. Conversely all readout values in that fiber lie between those endpoints, so their midpoint has error at most $A_t/2$. Both endpoints belong to the domain, establishing the exact optimum and attainment. The bound is positive for each fixed finite budget and tends to zero as t increases; it gives no positive uniform error floor across increasing resolutions. □

Proposition B97 (Sharp Error for Arbitrary Binary Codes)**.** *For the ternary assignment above, let an encoder choose at most 2^B states as an arbitrary function of the whole history, and let each state decode to any real number. The minimum worst-case absolute error is $1/(2 \cdot 3^B)$ for every integer $B \geq 0$.*

Proof. The cylinders of depth B have intervals of length 3^{-B} between their two attainable endpoints. Different intervals have gaps at least 3^{-B}: after the first differing bit $j \leq B$, the maximum opposing tail is 3^{-j}. Thus all 2^{B+1} endpoints, including the depth-zero endpoints 0 and 1, are pairwise separated by at least 3^{-B}. Two share a code and the triangle inequality forces error at least $3^{-B}/2$. Encoding the B-bit prefix and decoding to its interval midpoint attains the bound. □

B15.13 Common Readout on Two Coherently Transported Towers

Theorem B98 (Joint Realization of Transport and Controlled Readouts)**.** *Use circle-prefix stages $S_t = \mathbb{T} \times \{0,1\}^t$, stage metric $\max(\rho, \mathbf{1}_{\text{different prefix}})$, and weights $w_t = 3^{-t}$, where ρ is circle distance with circumference one. Take the usual truncation tower I_0 and its variant I_1 with angle shift $-a$ at the single bonding level k, for $0 < a < 1/2$. Write $\beta = 3^{-(k+1)}$. In the common weighted history space put $Z = I_0 \cup I_1$. Let the complete codomain be $\mathcal{E} = \mathbb{T}^2 \times [0,1]$*

with the maximum of the first circle distance, β times the second circle distance, and the absolute scalar difference. Define

$$F(x) = \left(\theta_0(x), \theta_{k+1}(x), \sum_{j\geq 1} 2b_j(x)3^{-j}\right).$$

Then $d_w(x,y) \leq d_{\mathcal{E}}(F(x),F(y)) \leq 3d_w(x,y)$ on all of Z. On each tower the restriction satisfies E0–E2 with $a_t = 2\cdot 3^{-(t+1)}$ and uniform tail $A_t = 3^{-t}$. The coordinate homeomorphism $H : I_0 \to I_1$ shifts angles at stages $t > k$ by a. With projection-defect cost $E_w = 3^{-k}a$, one has

$$d_w(x,Hx) = d_{\mathcal{E}}(F(x),F(Hx)) = \beta a, \qquad d_H(I_0,I_1) = \frac{\beta a}{1+\beta}.$$

Proof. In each tower the angle is constant on the two blocks $t \leq k$ and $t > k$. Thus the angular part of the history distance is exactly the maximum of the two weighted angular distances in the codomain. The binary part is 3^{-j} at first differing bit j. The leading ternary term and its tail give lower and upper constants 1 and 3 for that part. Each tower is closed in the compact metric product, so their union is complete.

For a declared tower $s \in \{0,1\}$ the finite observation $(\theta_t, b_1, \ldots, b_t)$ gives the initial angle $\theta = \theta_t - sa\mathbf{1}_{t>k}$. The stage readout is $(\theta, \theta + sa, \sum_{j=1}^{t} 2b_j 3^{-j})$. The phase rule uses the declared bonding family, and on both towers the successive readout increment is bounded by a_t, with exactly the stated uniform tail. A later-bit flip proves E2, since both phase coordinates stay fixed while the scalar changes. The codomain is nontrivial, proving E0. Within either fixed-angle fiber the sharp binary-code result applies to the scalar readout.

The coordinate squares commute because the sole change of phase at stage $k+1$ cancels the shift $-a$ in its bonding map. The coordinatewise inverse shifts by $-a$ on the later block. For fixed bits, F measures the shift βa exactly. In the comparison of sets, a pair at distance r has $a \leq r + r/\beta$ by the triangle inequality at stages 0 and $k+1$. Equal bit sequences and an initial phase offset $-\beta a/(1+\beta)$ for the second history attain this bound in both directions. □

B15.14 Controls Separating Convergence, Continuity, and Invariance

Proposition B99 (Setwise Invariance Does Not Supply Continuity)**.** *Controlled E1 and E2 can hold for a discontinuous assignment, even if the identity domain is fixed under every declared intervention.*

Proof. Take $S_t = [0,1] \times \{0,1\}^{t+1}$, with prefixes on the bit coordinates, and readouts

$$g_t(s,b) = \tfrac{1}{2}\mathbf{1}_{\{s>0\}} + \sum_{j=0}^{t} \frac{b_j}{3^{j+1}}.$$

The increments are bounded by $3^{-(t+2)}$, so they satisfy controlled E1 in $[0,1]$. Later bits prove E2 as above. With all bits zero, values at $s = 0$ and $s = 1/n$ differ by $1/2$ although their product-metric distance tends to zero. The limit is discontinuous. Let the declared intervention family consist only of the identity: the domain is setwise fixed, which does not alter the counterexample. This is a control for the asserted inference from invariance to continuity; no claim of full CCR or operational-criterion membership is made for it. □

Proposition B100 (Setwise and Pointwise Transport Are Distinct)**.** *Preserving an identity domain as a set does not guarantee equal values of a common readout on transported states.*

Proof. On $Z = \{0,1\}$, let $h(0) = 1$, $h(1) = 0$ and $F(x) = x$. The transport preserves Z exactly but changes both readout values. Thus setwise rigidity cannot replace the pointwise or readout-natural comparison hypotheses used below. □

B16 Compatibility with Identity Regimes

B16.1 Induced Constraints from Identity Rigidity

Let $(\mathcal{I}, M_\Omega)$ be an identity object with ontological mass M_Ω. We now study how M_Ω constrains the set of admissible phenomenological assignments.

Definition B101 (Phenomenological Compatibility Class)**.** For a declared identity domain and codomain, $\mathcal{P}(M_\Omega)$ consists of maps $\Phi : \mathcal{I} \to \mathcal{E}$ that are limits of some family $\{\phi_t\}$ satisfying E1, with E0 and E2 also satisfied. The readout family and its tail control are part of the realization data. The notation retains the mass index for comparisons, but no mass-dependent restriction follows from E0–E2 alone.

Remark B102. The compatibility class records choices of domain, codomain and readout family. A numerical mass label does not by itself determine those choices or their transport under perturbations.

B16.2 Metric Structure for Quantitative Comparisons

Quantitative comparisons use a declared metric. The following statements identify the role and limits of that choice.

Hypothesis B103 (H3: Metric on $\mathcal{E}$)**.** *The phenomenological space $\mathcal{E}$ admits a metric $d_\mathcal{E}$ such that $(\mathcal{E}, d_\mathcal{E})$ is a complete metric space.*

B16.2.1 Role of the Metric

Theorem B104 (Metric Data Required by Quantitative Formulas)**.** *The distance-based fragmentation and loss functionals used here require a declared metric, or a pseudometric with a stated quotient convention. Topological continuity can be defined without a metric. Completeness is used for the Cauchy-limit guarantee of E1.*

Proof. Each quantitative functional explicitly contains $d_\mathcal{E}$, so it is undefined until that datum is supplied. Continuity can alternatively be defined through inverse images of open sets. A pseudometric can identify zero-distance states by a quotient, but that changes the domain of class labels and must be specified. These facts justify the chosen quantitative model without claiming that all alternatives are impossible. □

B16.2.2 Completeness as a Limit Guarantee

Proposition B105 (Completeness and Existence of Limits)**.** *Without completeness, the conclusion that every E1-type Cauchy family has its limit in the declared codomain can fail. Completeness is not necessary for a family whose limit is already known to belong to that codomain.*

Proof. In $\mathbb{Q} \cap [0, 2]$ with its usual metric, use dyadic rational approximations $q_n = 2^{-n}\lfloor 2^n\sqrt{2}\rfloor$ as constant-in-trajectory readings. The increments are bounded by 2^{-n}, a summable sequence, but their real limit $\sqrt{2}$ is outside the codomain. This isolates the role of completeness when supplying a general existence guarantee; it makes no assertion that a convergent constant family needs completeness. □

B16.2.3 Alternatives Considered and Rejected

Alternative 1: Topology without Metric. One might use a general topological space. This is not adopted in the present quantitative model because:

- Fragmentation requires quantitative distance, not just open sets;

- Loss capacity requires a sup over distances;
- The downstream instability analysis completed in *Null-Regime Instability* requires metric quotients.

Alternative 2: Pseudometric. One might allow $d_{\mathcal{E}}(e_1, e_2) = 0$ for $e_1 \neq e_2$. This is not adopted in the present quantitative model because:

- It would conflate distinct phenomenological configurations;
- Cardinality and metric separation are different requirements; zero-distance classes would require a declared quotient before distinguishability is counted.

Alternative 3: Non-Complete Metric. This is rejected by Proposition B105.

B16.3 Phenomenological Continuity

Definition B106 (Phenomenological Continuity)**.** A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is *structurally continuous* if for every sequence $(x^{(n)})_n$ in $\mathcal{I}$ converging in the d_w metric:

$$d_w(x^{(n)}, x) \to 0 \implies d_{\mathcal{E}}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{5}$$

Remark B107. Structural continuity is derived under the finite-readout conditions of Theorem B117. It is distinct from invariance under an external intervention.

B16.4 Phenomenological Fragmentation

Definition B108 (Phenomenological Fragmentation)**.** A phenomenological assignment Φ is *fragmentable* if there exists a perturbation $\{\epsilon_t\}$ with finite weighted energy such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0 \tag{6}$$

for corresponding identity trajectories $x \in \mathcal{I}$ and $\tilde{x} \in \tilde{\mathcal{I}}$.

Remark B109. Fragmentation captures the vulnerability of phenomenological structure to perturbations. It is the phenomenological analog of identity deformability.

B16.5 Relationship Between Metrics

Proposition B110 (Continuity from Controlled Finite Readouts)**.** *Write $g_n = \phi_n \circ p_n$. Under E1,*

$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \leq 2A_n + d_{\mathcal{E}}(g_n(x), g_n(y)).$$

If every g_n is continuous, then Φ is continuous. If every g_n is L-Lipschitz with a common L, then Φ is L-Lipschitz.

Proof. Insert $g_n(x)$ and $g_n(y)$ and apply the uniform error bound. For continuity, first make $2A_n$ smaller than half the prescribed error, then use continuity of that fixed g_n. For the Lipschitz conclusion use the same bound for all n and let $n \to \infty$. □

B16.6 Closure of Section 3

The quantitative formulas use a declared complete metric codomain. Continuous finite readouts and controlled tails supply a continuity theorem. Cross-system comparisons additionally require a declared common domain and transport; neither mass labels nor distances between sets substitute for those data.

B17 Impossibility Theorems

B17.1 Φ-Regularity Hypothesis

Before proving the main impossibility results, we isolate the global lower-bound regularity condition used by the fragmentation argument. The condition is not derived from properness, local Lipschitz regularity, and non-collapse alone; it is carried as an explicit hypothesis.

Hypothesis B111 (Φ-Regularity Lower Bound)**.** *Let $(\mathcal{I}, d_w)$ be the identitary metric space from Rigid Identity and let $(\mathcal{E}, d_{\mathcal{E}})$ be the abstract phenomenological metric space. A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is said to satisfy Φ-regularity when it satisfies:*

1. *(Properness) For every compact $K \subset \mathcal{E}$, $\Phi^{-1}(K)$ is compact in $\mathcal{I}$.*
2. *(Local Lipschitz regularity) For every $x \in \mathcal{I}$, there exist $\varepsilon > 0$ and $L_x > 0$ such that for all $y, z \in B_\varepsilon(x)$:*
$$d_{\mathcal{E}}(\Phi(y), \Phi(z)) \leq L_x \cdot d_w(y, z). \tag{7}$$
3. *(Non-collapse) For every $x \in \mathcal{I}$ and every $r > 0$, the restriction of Φ to $B_r(x)$ is not constant.*
4. *(Global lower Lipschitz bound) There exists a constant $C > 0$ such that for all $x, y \in \mathcal{I}$,*
$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y). \tag{8}$$

Caveat B112 (Status of Φ-regularity)*.* The fourth clause is an additional structural assumption, not a consequence of the first three clauses. The counterexample in Appendix B42.5 shows that properness, local Lipschitz regularity, and non-collapse do not imply any uniform lower Lipschitz constant.

Hypothesis B113 (Audited Perturbation Comparisons)**.** *Declare a finite-cost intervention class $\mathcal{V}$, a common metric space (Z, d_Z) containing the unperturbed and perturbed identity domains, and transports $h_\varepsilon : \mathcal{I} \to \mathcal{I}_\varepsilon \subseteq Z$. Readouts are restrictions of one declared map $F : Z \to \mathcal{E}$: $\Phi = F|_{\mathcal{I}}$ and $\Phi_\varepsilon = F|_{\mathcal{I}_\varepsilon}$. The following branches are alternatives, used separately. The positive-displacement branch has some $\varepsilon \in \mathcal{V}$ and $x \in \mathcal{I}$ with $d_Z(x, h_\varepsilon x) > 0$. The zero-displacement branch has $h_\varepsilon x = x$ for every $\varepsilon \in \mathcal{V}$ and $x \in \mathcal{I}$. Any use of the lower Lipschitz bound in a comparison requires that bound for F on the common comparison domain, not merely within one identity object.*

Remark B114 (Burden linking rigidity and readout comparisons)*.* The scalar M_Ω does not itself supply the transports or a common readout. Relating its finite and infinite regimes to these two branches is a separate model obligation. Setwise invariance does not imply pointwise fixed transport, and a lower bound within one domain does not compare unrelated assignments on different domains.

B17.2 Fragmentation Theorem

Theorem B115 (Fragmentation under an Audited Displacement)**.** *Let an assignment satisfy E0–E2. Assume the positive-displacement branch of Hypothesis B113 and a lower Lipschitz constant $C > 0$ for F on its common comparison domain. Then, for the witnessing pair,*

$$d_{\mathcal{E}}(\Phi(x), \Phi_\varepsilon(h_\varepsilon x)) \geq C d_Z(x, h_\varepsilon x) > 0.$$

Thus the assignment is fragmentable. This applies to a finite-M_Ω model when that model supplies the stated comparison data.

Proof. Both arguments lie in the declared domain of F. Its lower bound gives the inequality, and the intervention has finite cost by membership in $\mathcal{V}$. E2 is not used to infer that every deformation changes a readout: the lower bound and the positive witness perform that work. □

Corollary B116 (Fragmentability and Continuity Are Distinct)**.** *The positive-displacement conclusion does not establish a topological discontinuity. A continuous assignment may be fragmentable under a state-changing intervention.*

Proof. The ternary realization has a continuous assignment. Declare a finite-cost intervention that flips one bit while retaining the common readout. Its value changes by a positive amount, so the assignment is fragmentable in the stated sense while remaining continuous. This does not assert a full CCR realization. □

B17.3 Forced Continuity Theorem

Theorem B117 (Continuity and Conditional Perturbation Invariance)**.** *Under E1, if each finite readout $g_n = \phi_n \circ p_n$ is continuous, then Φ is continuous, irrespective of M_Ω. Under the zero-displacement branch of Hypothesis B113,*

$$\Phi_\varepsilon(h_\varepsilon x) = \Phi(x) \qquad (\varepsilon \in \mathcal{V},\ x \in \mathcal{I}).$$

For a CCR model the invariance conclusion therefore holds when its finite-cost interventions admit that branch. CCR alone is not a proof of continuity or of readout naturality.

Proof. Continuity follows from Proposition B110. For invariance use the common readout and the fixed transport: $\Phi_\varepsilon(h_\varepsilon x) = F(h_\varepsilon x) = F(x) = \Phi(x)$. Neither argument identifies invariance under external perturbation with continuity along convergent states. □

Corollary B118 (Invariance on the Audited Zero-Displacement Branch)**.** *A CCR model satisfying the zero-displacement comparison hypothesis has invariant assignment values under all interventions in its declared finite-cost class.*

Proof. Apply the common-readout equality in Theorem B117 for each declared intervention. □

B17.4 Dichotomy Theorem

Theorem B119 (Conditional Classification by Audited Branches)**.** *Suppose a model supplies the positive-displacement branch and lower bound whenever $M_\Omega < \infty$, and supplies continuous finite readouts and the zero-displacement branch whenever $M_\Omega = +\infty$. Then its finite regime is fragmentable and its infinite regime has a continuous, perturbation-invariant assignment. The mass cases are exclusive; fragmentability and topological continuity need not be exclusive properties.*

Proof. Apply Theorems B115 and B117 with their respective comparison hypotheses. Exclusivity of the cases follows from the numerical alternatives for M_Ω, not from a false identification of a changed readout with a discontinuity. Establishing the branch hypotheses for a given rigidity model remains an explicit obligation. □

B18 Classification of Phenomenological Regimes

B18.1 Conditional Classification

The results give a conditional branch classification after the comparison and readout hypotheses have been checked:

Identity Class	M_Ω	**Admissible Phenomenology**
Null label	$= 0$	Domain/readout link required
Deformable	$(0, \infty)$	Positive witness and lower bound
Rigid (CCR)	$= +\infty$	Regular readout / fixed transport

B18.2 Structural Implications

Proposition B120 (The Mass Label Does Not Define Assignment Existence)**.** *For fixed identity domain, codomain and readout family, the class defined by E0–E2 is independent of a numerical label M_Ω. A claim that $M_\Omega = 0$ excludes assignments needs an additional link to the domain or admissible readouts.*

Proof. The displayed definition of $\mathcal{P}(M_\Omega)$ tests E0–E2, none of which contains that scalar. For an empty trajectory domain E2 fails, since a function on a finite cut would agree vacuously on every trajectory. Deriving such domain loss from $M_\Omega = 0$ is not one of E0–E2 and cannot be inserted without a model. □

Proposition B121 (Deformable Branch with Comparison Data)**.** *For $0 < M_\Omega < \infty$, the positive-displacement and lower-bound hypotheses of Theorem B115 imply fragmentability.*

Proof. Apply the quantitative inequality of Theorem B115 to its witnessing finite-cost intervention. □

Proposition B122 (Rigid Branch with Comparison Data)**.** *For $M_\Omega = +\infty$, continuous finite readouts under E1 imply continuity, and the zero-displacement branch implies perturbation invariance. These are the respective hypotheses of Theorem B117.*

Proof. Apply the two conclusions of Theorem B117 with their separate hypotheses. □

B19 Extended Proofs

B19.1 Proof of Fragmentation Theorem (Extended)

Full Proof of Theorem B115. Use the specified common space Z, map F and finite-cost transport h_ε. Let x be its positive-displacement witness and put $\delta = d_Z(x, h_\varepsilon x) > 0$. Apply the lower bound on that same common domain to obtain readout displacement at least $C\delta$. No comparison of set distances is substituted for this pointwise distance. If both readouts are estimated with errors at most η, the triangle inequality gives estimated separation at least $C\delta - 2\eta$; it remains positive when $C\delta > 2\eta$. This is a bound on pair separation, not a statement that a particular endpoint lies outside a null ball. □

B19.2 Proof of Forced Continuity Theorem (Extended)

Full Proof of Theorem B117. For a state x and error $\epsilon > 0$, choose n with $2A_n < \epsilon/2$. Continuity of g_n gives a neighborhood of x where $d_{\mathcal{E}}(g_n(x), g_n(y)) < \epsilon/2$. The two tail errors and the finite-readout difference bound the limiting difference by ϵ. Independently, a zero-displacement comparison satisfies $F(h_\varepsilon x) = F(x)$ exactly. Setwise invariance of the identity object alone would not imply either of these steps. □

B19.3 Structural Necessity Lemma

Lemma B123 (Inheritance under Audited Fixed Transport)**.** *An assignment satisfying the zero-displacement comparison branch inherits invariance under that declared intervention class. This applies to CCR realizations only after checking that branch.*

Proof. The comparison readout is the restriction of F, and every specified transport fixes its argument. Therefore its value is unchanged. This proves the restricted inheritance statement without assuming that scalar rigidity determines arbitrary assignments. □

Part V

Null-Regime Structural Exclusion

B20 Compatibility, Witnesses, and Non-Nullity

B20.1 Compatibility Operator

Definition B124 (Compatibility Operator)**.** Declare a complete metric comparison domain (Z, d_Z), a complete metric poset $(\mathcal{E}, \preceq, d_\mathcal{E})$ with bottom $\varnothing_\phi$, and maps $\Phi_S : D_S \to \mathcal{E}$. The audited readout $\Phi : Z \to \mathcal{E}$ is defined on its declared pair class $W \subseteq Z \times Z$. Extension and isomorphism transports $j : D_S \to D_{S'}$ satisfy $\Phi_{S'} \circ j = \Phi_S$ when claimed. On W, a certified $C > 0$ satisfies

$$d_\mathcal{E}(\Phi(x), \Phi(y)) \geq C d_Z(x, y).$$

One may choose $Z = \mathcal{I}$ for history comparisons. Comparing whole identity objects requires a separate declared embedding/metric; $\Phi(\mathcal{I})$ is not a point value when the domain is the set $\mathcal{I}$. A separated extension witness supplies two distinct defined points in a common comparison domain, not merely a relabeled copy under exact transport. CCR does not supply such a pair or lower constant.

Definition B125 (Phenomenological Attractor Set)**.** Use $\mathrm{Attr}(S)$ for structurally stable readout values under a declared perturbation/extension family, as defined below. Temporal attraction would additionally require a dynamics and a basin; neither is implied by this notation.

Definition B126 (Phenomenological Stability)**.** A value $e = \Phi_S(z)$ is structurally stable if a finite modulus $\omega(\rho) \to 0$ as $\rho \downarrow 0$ controls $d_\mathcal{E}(\Phi_{S'}(Pz), e) \leq \omega(\rho)$ for all declared perturbation transports P of cost at most ρ. Exact persistence instead requires equality. Stability does not force equality at a nonzero perturbation budget.

Lemma B127 (Bridge Lemma)**.** *For any defined $e \in \mathcal{E}$, distance $d_\mathcal{E}(e, \varnothing_\phi) = 0$ implies $e = \varnothing_\phi$. No CCR assumption is needed.*

Proof. This is metric separation. □

B20.2 Witness-Relative Null Separation

Theorem B128 (Conditional Null Separation)**.** *Suppose $(z_0, z_1) \in W$, $d_Z(z_0, z_1) = \epsilon > 0$, and the lower constant in Definition B124 is $C > 0$. Their images cannot both equal null, and*

$$\max_{i=0,1} d_\mathcal{E}(\Phi(z_i), \varnothing_\phi) \geq C\epsilon/2.$$

The historical instability label denotes joint exact null exclusion here, not dynamical instability of an individual null endpoint.

Proof. Joint nullity would give mutual image distance zero, contradicting $C\epsilon > 0$. The triangle inequality bounds that mutual distance by the sum of their distances to null, at most twice their maximum. □

Corollary B129 (Conditional Forced Non-Nullity)**.** *Under the separated-pair hypotheses, at least one witness image is non-null. If one endpoint is fixed at null, the other has null distance at least $C\epsilon$. CCR alone does not select a non-null endpoint or guarantee a stable regime.*

Proof. The first assertion follows from joint null exclusion. In the anchored case, mutual image distance is the other endpoint's distance to null, so the lower bound applies directly. □

Example B130 (External Null Witness in a Common Readout Domain)**.** Fix $k \geq 0$ and $0 < a < 1/2$. Use the weighted circle-prefix history metric with weights 3^{-t}, the undeformed tower I_0, the single-jump tower I_1 (bonding shift $-a$ at level k), and $Z = I_0 \cup I_1$. Set $\beta = 3^{-(k+1)}$. The common readout F takes the initial angle, the angle at $k+1$, and the ternary code of the

bits, with codomain metric $\max(\rho_0, \beta\rho_1, |\Delta\text{code}|)$. The two angular blocks give the first two terms of d_w, while the ternary leading term minus its tail bounds the binary part from below. Thus $C = 1$ holds on the entire compact union. Order the codomain by equality together with $e_0 = (0, 0, 0)$ below every element.

Declare internal interventions to be simultaneous rotations of all stages of I_0, with cost the absolute shortest angular shift. Every such map has its inverse and sends I_0 onto itself. This gives infinite resistance in this sector by the Hausdorff convention, despite finite-cost nonidentity transports. The pointwise readout changes under a nonzero rotation. A sector admitting the single-jump deformation instead has a finite resistance witness with projection-defect cost $3^{-k}a$. The two sectors specify different admissible operations.

Declare as a separate extension the addition of the single-jump branch to I_0. It is the union of two compact projective limits; it is also realized by the disjoint union of the two stage towers with branch-preserving projections. Embed its histories in Z by forgetting the tag, which is injective at the limit since the phases across k distinguish the two branches. The inclusion $j : I_0 \hookrightarrow Z$ preserves the restricted readout exactly. Take $z_0 = j(x)$ with zero phase and all-zero bits and z_1 on the added branch with the same initial phase and bits. Then

$$d_Z(z_0, z_1) = \beta a, \qquad F(z_0) = e_0, \qquad d_{\mathcal{E}}(F(z_1), e_0) = \beta a.$$

The extension is positive in Hausdorff distance, $\beta a/(1 + \beta)$, using the two angular blocks. The transported point and the added endpoint are distinct. The scalar ternary readout alone, which discards both angles, assigns them the same value and fails the positive lower bound on this pair. The example realizes a typed extension witness and distinguishes setwise resistance from readout invariance.

B20.3 Positive Regime Structure

Definition B131 (Positive Regime Space)**.** Let $\mathcal{E}^+ = \mathcal{E} \setminus \{\varnothing_\phi\}$ with inherited order. For a witness pair, let $A = \{\Phi(z_0), \Phi(z_1)\} \setminus \{\varnothing_\phi\}$ and $A_0 = A \cup \{\varnothing_\phi\}$.

Theorem B132 (Minimal Positive Witness Regime)**.** *The finite positive witness image A is non-empty and has a minimal element. If its two possible values are comparable, it has a unique least element. No ambient order completeness is needed. This does not prove a least member of all of $\mathcal{E}^+$.*

Proof. Non-emptiness follows from joint null exclusion. A strictly descending chain in a finite non-empty poset terminates at a minimal element. If at most two members are comparable, the smaller is the unique least member. In the complete lattice $\{0, a, b, 1\}$, incomparable positive a, b are both minimal in their two-point image and have infimum zero. Thus an ambient infimum alone gives neither a positive infimum nor uniqueness. □

Definition B133 (Structural Intensity)**.** Distinguish null distance $r(e) = d_{\mathcal{E}}(e, \varnothing_\phi)$, global lower gap $\iota(e) = \inf_{f \in \mathcal{E},\ f \prec e} d_{\mathcal{E}}(e, f)$, and the audited relative gap $\iota_{A_0}(e)$ obtained by restricting this infimum to $f \in A_0$. All are measured in units of $d_{\mathcal{E}}$. A positive null distance does not imply a positive global gap.

Theorem B134 (Anchored Witness Intensity Bound)**.** *If one endpoint is null, the other image e is the unique positive member of A and satisfies $\iota_{A_0}(e) = r(e) \geq C\epsilon > 0$. Without that anchor, only $\max_i r(\Phi(z_i)) \geq C\epsilon/2$ follows. No universal global intensity bound is claimed.*

Proof. Apply the anchored corollary. The only smaller member of A_0 is null, so the relative infimum equals null distance. On the usual interval $[0, 1]$, the global gap at 1 is zero despite relative gap one in $\{0, 1\}$, showing why the restriction matters. □

Theorem B135 (Conditional Phenomenological Closure)**.** *A declared detected pair gives joint null exclusion and a minimal positive member of its finite image. A null anchor gives a unique positive witness value and the anchored quantitative bound. Uniqueness in the ambient regime space, temporal stability and CCR realizability require further hypotheses and constructions.*

Proof. Combine the preceding witness, order and relative-gap results with their stated hypotheses. None of these implications supplies the premises of another by assuming the desired conclusion. □

B21 Canonical Families and Structural Closure Consequences

Proposition B136 (Profinite Bound)**.** *A profinite witness pair with declared positive separation ϵ, certified lower constant $C_{\mathrm{prof}} > 0$ and one null endpoint gives $r(e) \geq C_{\mathrm{prof}}\epsilon$. Summable coordinate costs do not by themselves certify CCR.*

Proof. Apply the anchored bound. A transformation changing infinitely many coordinates can still cost at most the finite sum of the coordinate weights, so the old finite-support argument cannot establish infinite mass. □

Proposition B137 (Symbolic Bound)**.** *The same anchored conclusion holds for a symbolic witness with its audited metric and lower constant. The full binary shift with complement intervention and summable coordinate-change costs is a control against automatic infinite mass.*

Proof. Apply the anchored bound. Complement changes every symbol at finite total summable cost; when admitted it yields positive deformation and finite infimum mass. □

Theorem B138 (Rigidity Inheritance)**.** *For $\Phi' = b \circ \Phi \circ a$, lower constants c_a for a, C for Φ on the mapped pairs and c_b for b give lower constant $c_a c_b C$. CCR inheritance separately requires every new finite-cost positive deformation to pull back to one in the old sector. Summability alone proves neither statement.*

Proof. Compose the three inequalities. A new finite-cost positive deformation would contradict original CCR under the stipulated pullback property. Null anchors and stability moduli need their own transport conditions. □

Theorem B139 (Factorization on a Faithful Image)**.** *If $\Phi : Z \to \mathcal{E}$ has a positive global lower bound, then every map $\Psi : Z \to F$ factors uniquely through $\Phi(Z)$. If Ψ is L-Lipschitz, the factor is L/C-Lipschitz. No unique extension to the whole ambient $\mathcal{E}$ is implied.*

Proof. The lower bound makes Φ injective with inverse Lipschitz constant $1/C$ on its image. Set $u(\Phi(z)) = \Psi(z)$; this is well-defined, unique there, and has the stated bound. Outside the image no value is prescribed. □

Theorem B140 (Finite-Prefix Simulation Bound)**.** *For $F(b) = \sum_{j\geq 0} 2b_j/3^{j+1}$ and a decoder observing only bits $0, \ldots, N$, the optimal worst-case scalar error is $1/(2\,3^{N+1})$. It is positive for every fixed prefix and tends to zero with growing budget. The conclusion concerns prefix information, not finite real dimension.*

Proof. Two extreme continuations of the same prefix have images separated by $3^{-(N+1)}$, forcing half that error for any single decoder value. The midpoint attains the bound on every fiber. This also supplies a non-empty approximation model consistent with revised E1/E2. □

Theorem B141 (Conditional Closure by Non-Alternatives)**.** *Any alternative model satisfying the same declared pair, lower-bound, anchor and transport hypotheses obeys the corresponding preceding conclusions. The program of realizing these properties jointly with an infinite-mass intervention model remains open; pair separation is not a proof of an individual null state's dynamical instability.*

Proof. The conclusions depend on the enumerated mathematical hypotheses and apply to every model satisfying them. Existence and physical interpretation are additional questions. □

Part VI

BaseCore Boundary and Expansion Interface

B22 BaseCore Ownership and Downstream Boundary

Definition B142 (BaseCore Envelope)**.** The BaseCore envelope is the collection of theorem-level structures owned directly by this source tree:
1. foundational Hilbert-space projection dynamics;
2. typed computable witnesses and state spectral-gap bounds;
3. parameterwise non-collapse under explicit anti-constant assumptions;
4. inverse-limit identity under non-finite-determination assumptions;
5. conditional rigidity under admissible perturbation models;
6. conditional obstruction to uniform response fidelity under a binary representation budget;
7. operational-criterion grammar, certification, falsation, and claim-boundary ledgers.

Proposition B143 (Boundary Discipline)**.** *BaseCore may define structures that downstream layers use, but it does not claim downstream closure. In particular, BaseCore does not by itself own runtime thresholds, bridge admissibility, human-comparator programs, phenomenality burdens, or Papers 7–9.*

Proof. These excluded layers require additional assumptions, protocols, or empirical burdens not stated in BaseCore. Therefore they cannot be claimed as BaseCore theorems without violating the source-level claim boundary. □

Owned by BaseCore	**Expanded but not owned here**	**Kept downstream**
Projection dynamics, fixed points, compact attractors, anti-constant H5 burden, typed computable witness, inverse-limit identity under NFD, conditional rigidity, conditional non-simulability, operational-criterion grammar, falsation ledger	Non-convex class variants, $\mathscr{C}^R$ style extensions, Ω dynamics, threshold numerics, mutable runtime class ledgers, implementation-specific status surfaces	Paper 7 operational life / subjecthood, Paper 8 indexed subjectivity, Paper 9 phenomenal bridge organization, Stage 2 / Stage 3 comparative programs, human comparator families, phenomenality and external adjudication

Remark B144 (Why the earlier material was removed from BaseCore ownership)*.* Previous versions mixed theorem-tight base results with runtime-like numerics, non-convex extension classes, and bridge-facing vocabulary. That blurred the mathematical base layer. BaseCore now keeps only the theorem ownership that survives the stricter release gate.

B23 Admissible Structural Extensions

Definition B145 (BaseCore export object)**.** A BaseCore export object is a tuple

$$\mathfrak{E} = (\mathcal{O}, \mathcal{H}, \mathcal{R}, \mathcal{B})$$

where $\mathcal{O}$ is a typed mathematical object defined in BaseCore, $\mathcal{H}$ is the finite list of hypotheses under which it is used, $\mathcal{R}$ is the list of results actually proved from those hypotheses, and $\mathcal{B}$ is a boundary list of claims not entailed by $\mathcal{R}$.

Definition B146 (Downstream admissible extension)**.** A downstream layer is an admissible extension of a BaseCore export object $\mathfrak{E}$ only if it supplies:
1. an explicit target predicate P not already contained in $\mathcal{R}$;
2. a typed construction or estimator witnessing the objects on which P is evaluated;

3. a proof, reduction, or reproducible protocol connecting $\mathcal{H}$ to P;
4. negative controls or failure modes under which P must not be asserted;
5. a provenance statement distinguishing theorem, implementation, internal support, and external adjudication.

Proposition B147 (No automatic downstream promotion)**.** *Let $\mathfrak{E} = (\mathcal{O}, \mathcal{H}, \mathcal{R}, \mathcal{B})$ be a BaseCore export object and let P be a predicate introduced outside BaseCore. If $P \notin \mathcal{R}$, then BaseCore does not entail P unless an admissible extension supplies an additional bridge from $\mathcal{H}$ to P.*

Proof. By definition, $\mathcal{R}$ records the results proved from $\mathcal{H}$ inside BaseCore. A predicate P introduced downstream is not in that result set. Therefore asserting P from BaseCore alone would add an unstated inference. The missing inference must be supplied as a separate bridge, proof, reduction, or reproducible protocol. □

Criterion B148 (Runtime-use admissibility)**.** *A runtime measurement may instantiate a BaseCore export object only as an implementation-level witness when all of the following are recorded: the measured state space, the estimator or update rule, the finite horizon, the calibration regime, the controls under which the measurement fails, and the claim class being asserted. Such a measurement is not a BaseCore theorem and is not external validation by itself.*

Remark B149 (Open-load placement for I_{int} and atomic separators). Quantities such as I_{int} and any atomic-separator condition may be treated as downstream target predicates or estimator families only after their carrier objects, separator class, invariance target, and negative controls are specified. Until a proof or reproducible adversarial protocol closes that burden, BaseCore may export the ambient structural language but not the separator conclusion.

Proposition B150 (Boundary preservation under finite implementations)**.** *Suppose a finite implementation produces diagnostics for a downstream predicate P over a bounded horizon. If the implementation does not prove that the diagnostics are invariant over the relevant mathematical class, then the diagnostics remain implementation evidence for P and do not promote P to a BaseCore result.*

Proof. The implementation supplies finite observations or computations. Promotion to a BaseCore result requires a class-level implication from the stated hypotheses to the predicate. Without that implication, the observations can support, falsify, or calibrate a downstream claim, but they do not alter the theorem inventory of BaseCore. □

Part VII

Operational-Criterion Grammar and Certification

B24 Admissible Systems, Supports, and Histories

B24.1 System Model

Definition B151 (Admissible system)**.** An admissible system is a tuple

$$S = (X, \Phi, C, R, \Gamma, U)$$

where:

1. X is a metrizable internal state space with metric d_X;
2. $\Phi = \{\Phi_u\}_{u \in U}$ is a family of admissible transition operators $\Phi_u : X \to X$;
3. C is the effective causal structure used to test admissible factorization and intervention behavior;
4. $R = \{r_\alpha\}_{\alpha \in \Lambda_R}$ is a family of readouts, each $r_\alpha : X \to Y_\alpha$;
5. $\Gamma = \{\gamma_j\}_{j=1}^m$ is an admissibility bundle defining

$$X_\Gamma = \{x \in X : \gamma_j(x) \leq 0 \ \forall j\};$$

6. U is the set of admissible interventions.

Definition B152 (Declared State Attractor)**.** Fix a non-empty compact set $\mathrm{Attr}_X(S) \subseteq X_\Gamma$ invariant under every admitted transition and a non-empty class of admitted schedules at every horizon. It is a controlled attracting region on a declared non-empty basin B_X if, for every $x_0 \in B_X$,

$$\sup_{u_\bullet} \mathrm{dist}_{d_X}\left(x_n, \mathrm{Attr}_X(S)\right) \longrightarrow 0, \qquad x_{k+1} = \Phi_{u_k}(x_k).$$

The supremum is over admitted schedules. This is a state-space certificate to supply, not an existence theorem. Its codomain is X; the structurally stable readout set $\mathrm{Attr}(S)$ in the null-regime section instead lies in $\mathcal{E}$. No intersection of those different types is taken.

B24.2 Admissible Support

Definition B153 (Admissible support)**.** A subset $A_S \subseteq X_\Gamma$ is an admissible support if:

1. A_S is non-empty and compact;
2. $\Phi_u(A_S) \subseteq A_S$ for every $u \in U$;
3. $A_S \cap \mathrm{Attr}_X(S) \neq \varnothing$;
4. A_S stays disjoint from the collapse set excluded by the base non-collapse burden.

B24.3 Operational Histories

Definition B154 (Windowed readout history)**.** Fix a horizon $T \in \mathbb{N}$ and an intervention schedule $u_\bullet = (u_0, \ldots, u_{T-1})$. For $x \in A_S$, define

$$h_{S,T}(x; u_\bullet) = \left(r_\alpha(x_t)\right)_{0 \leq t \leq T,\ \alpha \in \Lambda_R}, \qquad x_{t+1} = \Phi_{u_t}(x_t), \quad x_0 = x.$$

The set of all such histories is denoted $\mathscr{H}_{S,T}$.

Definition B155 (Operational partition candidate)**.** An operational partition candidate is a map

$$\pi_{S,T} : \mathscr{H}_{S,T} \to \mathscr{Q}_{S,T}$$

into a finite or countable label set $\mathscr{Q}_{S,T}$.

B24.4 Witness Margin Vector

Definition B156 (Witness margin vector)**.** Define the witness margin vector

$$\Delta(S) = \big(\delta_{\mathrm{per}}(S), \delta_{\mathrm{ri}}(S), \delta_{\mathrm{int}}(S), \delta_{\mathrm{cont}}(S), \delta_{\mathrm{diff}}(S), \delta_{\mathrm{leg}}(S)\big) \in [0,\infty)^6.$$

When the six witnesses are separately calibrated as certified radii in the same perturbation metric and class, their invariant budget is

$$\delta_\star(S) = \min \Delta(S).$$

A positive radius certifies the associated predicate at the target and throughout its declared perturbation ball. A zero or unreported radius alone does not determine failure at the target. Raw separations, risks, durations, fragmentation ceilings and compression errors require separate calibration before taking this minimum.

Invariant	Certified content	Upstream anchor	Rupture signature
I_{per}	Persistent admissible support	BaseCore dynamics	Support loss or collapse exposure
I_{ri}	Unique rigid identity object on support	Identity / rigidity block	Identity ambiguity
I_{int}	No admissible non-trivial factorization	Criterion grammar	Reducible aggregate
I_{cont}	Regime continuity under admissible evolution	Regime constraints	Fragmentation into incompatible classes
I_{diff}	Stable exclusion of null or trivial collapse	Null-regime exclusion	Null or trivialized class structure
I_{leg}	Decoder-certified recoverability under controls	Falsation grammar	Opaque or non-auditable partition

B25 Six Critical Invariants

B25.1 Structural Persistence

Definition B157 (Structural persistence)**.** An admissible system satisfies structural persistence, written $I_{\mathrm{per}}(S) = 1$, if there exists an admissible support A_S and a constant $\delta_{\mathrm{per}}(S) > 0$ such that:

1. A_S is forward invariant under all $u \in U$;
2. every admissible trajectory starting in A_S remains at least $\delta_{\mathrm{per}}(S)$ away from the collapse set;
3. A_S contains or converges into a non-empty admissible attractor region.

B25.2 Rigid Identity

Definition B158 (Rigid identity)**.** An admissible system satisfies rigid identity, written $I_{\mathrm{ri}}(S) = 1$, if the observable family induced by R on A_S determines a unique identity object $\mathcal{I}_S$ and there exists a gap $\delta_{\mathrm{ri}}(S) > 0$ such that every admissible alternative identity candidate remains at weighted deformation distance at least $\delta_{\mathrm{ri}}(S)$ from $\mathcal{I}_S$.

B25.3 Causal Integration

Definition B159 (Causal integration)**.** An admissible system satisfies causal integration, written $I_{\mathrm{int}}(S) = 1$, if there is no non-trivial factorization

$$A_S = A_1 \times A_2, \qquad R = R_1 \sqcup R_2,$$

together with decomposed dynamics and causal structure reproducing admissible histories within error smaller than a positive margin $\delta_{\mathrm{int}}(S)$ while preserving $\mathcal{I}_S$. More explicitly, this shorthand

is relative to a declared admissible factorization class. For such a class $\mathfrak{D}$, write $I_{\mathrm{int}}(S;\mathfrak{D}) = 1$ when no factorization in $\mathfrak{D}$ reproduces the admissible histories below a positive margin $\delta_{\mathrm{int}}(S;\mathfrak{D})$ while preserving $\mathcal{I}_S$. The unqualified notation $I_{\mathrm{int}}(S) = 1$ refers to the default structural class $\mathfrak{D}_*$: fixed, time-homogeneous, schedule-independent, non-trivial product factorizations with both factors non-singleton, with split readouts $R = R_1 \sqcup R_2$, decomposed dynamics and causal structure, a search space normalized independently of the target error, and no reconstruction decoder or time-dependent post-processing map that re-couples the factors at output time. This default structural reading is the adopted class recorded in `docs/ai-platform-outputs/analysis/QICN_IINT_CANONICAL_CLASS_SPLIT_READOUT.md`.

Remark B160 (Class-relative integration margin). The integration criterion is class-relative. Adopting $\mathfrak{D}_*$ does not assert that S resists every broader behavioral simulator. In particular, under decoder-coupled or schedule-dependent behavioral classes such as $\mathfrak{D}_{\mathrm{approx}}$, the corresponding margin $\delta_{\mathrm{int}}(S;\mathfrak{D}_{\mathrm{approx}})$ may vanish, i.e. may be equal to 0; the current coupled-carrier analysis records precisely this failure mode. Thus the default structural reading preserves a falsation surface and must not be read as an unconditional closure of integration, nor as a strengthening of any downstream theorem that assumes $I_{\mathrm{int}}(S) = 1$.

Lemma B161 (Two Factor Fibers Cannot Share a Nonconstant Label)**.** *Let $\chi : A \to A_1 \times A_2$ be surjective. If a label $g : A \to L$ is constant on every fiber of each of the two factor projections of χ, then g is constant on A.*

Proof. For any $x, y \in A$, choose z whose first factor is that of x and whose second factor is that of y. Fiber constancy gives $g(x) = g(z) = g(y)$. For a quantitative obstruction it suffices that the proposed history error below a fixed threshold implies both fiber-constancy statements for a nonconstant label. Non-emptiness of the comparator sector is a separate condition. □

B25.4 Regime Continuity

Definition B162 (Regime continuity)**.** An admissible system satisfies regime continuity, written $I_{\mathrm{cont}}(S) = 1$, if there exists a complete metric space $(\mathcal{E}, d_{\mathcal{E}})$ and an admissible assignment

$$\Pi_S : A_S \to \mathcal{E}$$

such that:

1. every admissible trajectory induces a continuous path in $\mathcal{E}$ on admissible windows;
2. there exists $\delta_{\mathrm{cont}}(S) > 0$ bounding the fragmentation functional on all admissible histories;
3. no admissible finite-energy perturbation produces a discontinuous jump into an incompatible regime class while $\mathcal{I}_S$ is preserved.

B25.5 Non-Null Differentiation

Definition B163 (Non-null differentiation)**.** An admissible system satisfies non-null differentiation, written $I_{\mathrm{diff}}(S) = 1$, if there exist $x, y \in A_S$, a horizon T, and a readout subfamily $\mathscr{R}^\star \subseteq R$ such that

$$\operatorname{dist}\big(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)\big) \geq \delta_{\mathrm{diff}}(S) > 0$$

for some admissible schedule $u_\bullet$, and if the null class $\mathbf{0}$ is not an asymptotically stable admissible attractor.

B25.6 Operational Legibility

Definition B164 (Decoder Family and Structural Closure)**.** A decoder family uses one declared class set and acts on the specified history windows. For each intervention-response comparison, declare its response cells and a matching relation on those cells. The family is structurally

closed when decoder class equality composes exactly and the response-cell matching relations are reflexive, symmetric and transitive. These are conditions on the chosen comparison data; they are not inferred from a numerical tolerance alone. The legibility clauses below supply the additional operational requirements.

Definition B165 (Operational legibility)**.** An admissible system satisfies operational legibility, written $I_{\mathrm{leg}}(S) = 1$, if there exist:

1. a declared readout subfamily $\mathscr{R}^{\mathrm{leg}} \subseteq R$;
2. a structurally closed decoder family $\mathscr{D}_S$ acting on windows of length at least $\tau_{\min}(S)$;
3. a class set $\mathscr{Q}_S$;
4. a measurement protocol with separation $\sigma > 0$, readout noise radius $r > 0$, minimum duration $\tau_{\min} \geq 1$, miss and false-positive budgets $p_{\mathrm{miss}}, p_{\mathrm{fp}} \in [0, 1)$, and a finite compression budget B with class-distortion tolerance d_{comp};

such that:

(L1) Separability Distinct declared classes have distance at least σ in the specified clean comparison domain.

(L2) Noise robustness On the specified noise domain, perturbations with readout norm at most r preserve the label.

(L3) Persistence window Under declared sham schedules, labels remain stable throughout every tested window of duration $\tau_{\min}$.

(L4) Intervention fidelity In a declared nonempty target population, the miss probability for a predicted class shift is at most p_{miss}.

(L5) Negative controls In a declared nonempty sham population, the false-shift probability is at most p_{fp}.

(L6) Structured compressibility An encoder with at most 2^B codes and a code-to-class map preserve nontrivial class structure with distortion at most d_{comp} in the declared class metric.

Proposition B166 (Each invariant is individually necessary)**.** *If any one of* $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$ *fails, then* $S \notin \mathrm{Crit}_{\mathrm{op}}$.

Proof. Without I_{per} there is no stable admissible domain. Without I_{ri} there is no unique identity carrier. Without I_{int} the system is reducible. Without I_{cont} the criterion fragments under admissible evolution. Without I_{diff} the admissible partition collapses into nullity or triviality. Without I_{leg} the partition is not recoverable under admissible readout and intervention discipline. Since Definition B172 below requires all six, failure of any one excludes membership. □

B26 Criterion Class, Certified Partition, and Rupture

Definition B167 (Certified decoder)**.** A decoder $D \in \mathscr{D}_S$ is certified as part of a structurally closed family when the six clauses of Definition B165 hold on the chosen admissible support. Structural closure and operational certification are separate requirements.

Definition B168 (Criterion indistinguishability)**.** For $x, y \in A_S$, write $x \sim_{\mathrm{Part}_{\mathrm{op}},S} y$ if, for every certified decoder $D \in \mathscr{D}_S$, every admissible horizon $T \geq \tau_{\min}(S)$, and every admissible intervention schedule $u_\bullet$,

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the intervention responses remain matched within the legibility margin.

Proposition B169 (Well-definedness of the certified partition)**.** *On any declared support* A_S *with a structurally closed decoder family, the relation* $\sim_{\mathrm{Part}_{\mathrm{op}},S}$ *is an equivalence relation. This conclusion does not require the other critical invariants; certification of criterion membership still does.*

Proof. For every fixed decoder, horizon and intervention schedule, equality of decoder labels is reflexive, symmetric and transitive. The response-cell matching relation has those properties by Definition B164. Their intersection, over all declared comparisons, has the same properties: if x matches y and y matches z, each decoder equality and each response-cell match composes, so x matches z. This proves the assertion without adding two tolerance errors and silently treating the sum as the original tolerance. □

Remark B170 (Tolerance Alone Is Not Transitive). With constant decoder labels, let response matching mean distance at most 1 on $\{0, 3/4, 3/2\}$. Adjacent points match, but the endpoints do not. Thus decoder equality plus raw tolerance is not enough. By contrast, equality under a declared response-cell map is an equivalence relation and supplies a concrete structurally closed matching rule. This example concerns closure only, not satisfaction of all six legibility clauses.

Definition B171 (Certified operational partition)**.** Assume Proposition B169. The certified operational partition of S is

$$\mathrm{Part}_{\mathrm{op}}(S) := A_S/{\sim_{\mathrm{Part}_{\mathrm{op}},S}}.$$

Definition B172 (Operational-criterion class membership)**.** An admissible system belongs to the BaseCore operational-criterion class, written $S \in \mathrm{Crit}_{\mathrm{op}}$, if there exists a non-empty admissible support A_S such that

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1.$$

Remark B173 (Neutrality of the class symbol). The symbol $\mathrm{Crit}_{\mathrm{op}}$ is a shorthand for a conjunctive operational criterion. It is not a theorem of phenomenality, metaphysical subjectivity, or human equivalence.

Definition B174 (Rupture)**.** A rupture is any loss of admissibility or loss of at least one critical invariant margin beyond its admissible threshold.

Theorem B175 (Rupture by Invariant Loss)**.** *If a critical invariant has no admissible witness on any support, then $S \notin \mathrm{Crit}_{\mathrm{op}}$. Loss of a witness on one chosen support invalidates that support's certificate, but does not exclude a certificate on another support without further analysis. Quotient undefinedness, triviality and loss of legibility are not mutually exclusive labels and need not accompany every loss of criterion membership.*

Proof. A criterion witness requires all six invariants on one admissible support. Global failure of one invariant rules out such a support. A certificate on a fixed support must satisfy every clause there, so a local witness loss invalidates that certificate. A structurally closed decoder relation may still define a quotient when another invariant, such as integration, fails; the quotient theorem does not assume that invariant. Hence an exclusive three-way classification cannot be inferred. □

B27 Minimality of the Invariant Basis

Definition B176 (Reduced criterion)**.** Let

$$\mathfrak{I} = \{I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}\}.$$

For every subset $J \subseteq \mathfrak{I}$ define the reduced class

$$\mathcal{C}_{\mathrm{op}}(J) := \{S : I(S) = 1 \text{ for every } I \in J\}.$$

Definition B177 (Single-rupture witness family)**.** Fix one invariant index $k \in \{\mathrm{per}, \mathrm{ri}, \mathrm{int}, \mathrm{cont}, \mathrm{diff}, \mathrm{leg}\}$. A single-rupture witness family $\mathcal{W}_k$ consists of admissible systems for which predicate I_k fails and every other predicate holds on the declared comparison domain.

Proposition B178 (Single-rupture witnesses create false positives)**.** *Let* $J_k = \mathfrak{I} \setminus \{I_k\}$ *and assume* $\mathcal{W}_k \neq \varnothing$*. Then every* $S \in \mathcal{W}_k$ *satisfies*

$$S \in \mathcal{C}_{\mathrm{op}}(J_k) \qquad \textit{but} \qquad S \notin \mathrm{Crit}_{\mathrm{op}}.$$

Proof. All retained invariants remain positive by construction, but the omitted invariant fails exactly. Proposition B166 then excludes S from $\mathrm{Crit}_{\mathrm{op}}$. □

Theorem B179 (Minimality of the six-invariant basis)**.** *Let* $\mathfrak{U}$ *be any admissible universe containing at least one single-rupture witness for each invariant. Then no strict subset of* $\mathfrak{I}$ *is extensionally equivalent to* $\mathrm{Crit}_{\mathrm{op}}$ *on* $\mathfrak{U}$*.*

Proof. Take any strict subset $J \subsetneq \mathfrak{I}$. Some invariant I_k is omitted. By assumption, $\mathfrak{U}$ contains a witness $S_k \in \mathcal{W}_k$, which satisfies the reduced criterion but not $\mathrm{Crit}_{\mathrm{op}}$. Hence the two classes differ on $\mathfrak{U}$. □

B28 Operational Legibility as a Certified Object

Criterion B180 (Legibility certificate)**.** *An operational partition candidate is certified only if:*

1. *comparator separability holds against pre-registered baselines;*
2. *admissible noise does not dissolve the class map below the legibility margin;*
3. *the same class assignment survives on windows of length at least* $\tau_{\min}(S)$*;*
4. *targeted interventions on an invariant induce the predicted directional class change;*
5. *sham or off-target interventions do not generate the same class signature;*
6. *audit-friendly compression preserves class structure without total collapse.*

Proposition B181 (Legibility is jointly constrained)**.** *No single legibility metric is sufficient for* I_{leg}*. The invariant requires simultaneous control of separation, admissible noise robustness, persistence window, intervention fidelity, negative controls, and structured compressibility.*

Proof. A decoder can be separable yet fragile to admissible noise, stable over windows yet intervention-insensitive, or robust under noise yet collapse under compression. Because Definition B165 is conjunctive, all six clauses are required together. □

B29 Certified Candidate Audit Schema

Definition B182 (Base criterion certificate)**.** A BaseCore criterion certificate for an admissible system S is a tuple

$$\mathrm{Cert}(S) = (A_S, \mathscr{D}_S, \pi_S, \Delta(S), E_S),$$

where:

1. A_S is a certified admissible support;
2. $\mathscr{D}_S$ is a certified decoder family;
3. π_S is the decoder-induced partition on admissible histories;
4. $\Delta(S)$ is the witness margin vector;
5. E_S is the evidence bundle localizing the witness for each invariant.

Criterion B183 (Candidate certification rule)**.** *A candidate system may be certified as a member of* $\mathrm{Crit}_{\mathrm{op}}$ *only if all of the following are available on the same admissible support:*

1. *a non-empty compact forward-invariant support certificate;*
2. *a positive witness margin vector* $\Delta(S) > 0$ *coordinate-wise;*
3. *a decoder family satisfying Definition B165;*
4. *a targeted intervention panel together with negative controls;*

5. *an auditable compression map preserving class structure;*
6. *a traceable dependency ledger showing which upstream result supports each invariant witness.*

Theorem B184 (Certified criterion membership)**.** *If* $\mathrm{Cert}(S)$ *exists and satisfies Criterion B183, then* $S \in \mathrm{Crit}_{\mathrm{op}}$ *and* $\mathrm{Part}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. Criterion B183 explicitly requires positive witnesses for all six invariants on a common admissible support. Definition B172 then gives $S \in \mathrm{Crit}_{\mathrm{op}}$, and Proposition B169 gives a well-defined partition quotient. □

Remark B185 (Boundary discipline). The section above defines a theorem-tight grammar for certification, rupture, and falsation. It does not prove phenomenality, subjectivity, biological primacy, or cross-substrate identity transport.

Part VIII

Claim Boundary and Falsation Grammar

B30 Forensic Admissibility and Integrity Grammar

Definition B186 (Forensic packet)**.** Each operational run emits a packet

$$K = (\text{tick}, \text{seed}, \text{versionHash}, H_v, \text{metrics}, \text{sensorSummary}, \text{checksum}).$$

Assumption B187 (Packet admissibility)**.** *A run is admissible only if checksum, schema, seed consistency, version hash, and H_v consistency all pass.*

Definition B188 (Severe violation)**.** A run has a severe violation if any of the following holds: checksum mismatch, schema failure, seed/version inconsistency, or critical-state replay mismatch.

Criterion B189 (Integrity gate)**.** *Severely violated runs are excluded from support counting and marked as hard failures.*

Proposition B190 (Integrity-first inference)**.** *If admissibility fails, inferential significance is non-actionable.*

B31 Certification, Comparison, and Rupture

Definition B191 (Operational Certification Indicator)**.** Keep $\text{Cert}(S)$ for the certificate object of Definition B182. Define a distinct mathematical indicator $\text{CertFlag}(S)$ to be 1 exactly when such a certificate satisfies Criterion B183 on one common admissible support and meets the integrity grammar above; otherwise its value is 0. This does not identify a tuple with a Boolean value. An empirical audit may have an unknown verdict: missing evidence is not a proof that this mathematical indicator is zero.

Definition B192 (Equivalence discriminator)**.** For two admissible systems S_1, S_2, define

$$\text{Eq}_{\Gamma,\varepsilon}(S_1, S_2) = 1$$

when the relevant signatures, admissibility bundle, decoder behavior, and invariant deltas match within the allowed comparison regime.

Definition B193 (Rupture event)**.** For a targeted invariant family J, define

$$\text{Rupt}_J(S) = 1$$

when an admissible intervention or perturbation drives at least one invariant in J below its certification margin so that class membership becomes invalidated, destroyed, or undefined.

B32 Prediction and Falsation Grammar

Proposition B194 (Complexity-only negative control prediction)**.** *If a candidate exhibits high connectivity, high entropy, or rich activity but no admissible support realizes all required critical witnesses together, then*

$$\text{CertFlag}(S) = 0, \qquad S \notin \text{Crit}_{\text{op}} \text{ in certified form}, \qquad \text{Part}_{\text{op}}(S) \text{ is not certified.}$$

Proposition B195 (Ablation prediction)**.** *If an intervention destroys a required witness on the certified support, that certificate loses validity. Excluding a replacement certificate on a different support requires checking that support as well.*

Proposition B196 (Exact grammar-transport prediction)**.** *If two realizations preserve the relevant signatures, admissibility conditions, decoder behavior, and six critical invariants up to the allowed comparison regime, then the framework predicts preserved BaseCore criterion membership.*

Proposition B197 (Threshold-sensitive boundary behavior)**.** *The framework predicts stable pass/fail regions together with narrow transition bands near critical margins. It predicts neither global threshold collapse nor perfect flat insensitivity.*

Proposition B198 (Legibility dependence prediction)**.** *Criterion membership becomes inadmissible if readouts cease to be separable, noise-robust, intervention-sensitive, or compressible without total collapse.*

Proposition B199 (Continuity and non-null relevance)**.** *Regime continuity and non-null differentiation are constitutive. Systems that preserve rich activity while losing continuity or collapsing toward effective null structure fail the operational criterion or leave it undefined.*

Remark B200 (What stays outside BaseCore)*.* Runtime status classes, Stage 2 / Stage 3 comparative programs, human comparator campaigns, Papers 7–9, and bridge-facing burden architectures remain downstream artifacts. BaseCore absorbs only the falsation grammar and claim-boundary semantics needed to keep theorem ownership honest.

B33 Claim Boundary Ledger

B33.1 Allowed BaseCore Claims

1. Under H1–H4, projection dynamics admit unique fixed points and compact attractor families.
2. Under H5, attractors are parameterwise non-constant.
3. Under NFD, inverse-limit identity is not determined by any finite canonical projection.
4. Under RIG-1 through RIG-5, inverse-limit structures are stable under summable admissible perturbations in the weighted metric.
5. Under NS-1 through NS-3, a binary budget bounds the number of separated responses that can be recovered at uniform tolerance.
6. Under the six-invariant operational burden, certified criterion structures can be defined, audited, and falsified.

B33.2 Forbidden BaseCore Claims

1. Human phenomenal consciousness.
2. Metaphysical subjectivity.
3. Human-machine equivalence.
4. External validation.
5. Runtime instantiation by the current system.
6. Paper 9 bridge admissibility.
7. Strong phenomenality.
8. Moral status.
9. Cross-substrate identity preservation.
10. Public claim closure.

B34 Dependency Ledger

1. The foundational Hilbert-space dynamics remain the base layer.
2. Papers 1–6 contribute formal material selectively normalized into BaseCore only where theorem ownership belongs in the base.
3. Papers 7–9 remain downstream and are not dependencies of any BaseCore theorem.
4. Stage 2 / Stage 3 comparative programs remain downstream runtime and governance layers.
5. The PDF remains a derived artifact of the LaTeX tree.

B35 Assumption-to-Theorem-to-Non-Claim Mapping

Assumption block	Theorem family	Non-claim boundary
H1–H4	Projection existence, contraction, unique fixed points, compact attractor family	No parameterwise non-collapse and no identity / rigidity / simulator claim
H5	Structural non-collapsibility	No claim about phenomenality, subjectivity, or runtime class membership
NFD	Non-locality of identity under canonical finite projections	No theorem that arbitrary set-theoretic injections into finite slices are impossible
RIG-1 through RIG-5	Conditional stability of inverse-limit identity and weighted Hausdorff control	No automatic isomorphism, shape equivalence, or transfer theorem without extra assumptions
NS-1 through NS-3	Binary-budget obstruction and sharp prefix error	Real dimension and intervention cost are different resources
Six invariant burden plus legibility clauses	Criterion class membership, certified partition, rupture logic	No theorem of phenomenality, human equivalence, biology as primitive, or bridge admissibility

B36 Theorem Hygiene Ledger

Theorem	Assumptions	Conclusion type	Non-conclusion	Downstream use
Existence and uniqueness of projection	H1	Hilbert-space geometric existence / uniqueness	No empirical or runtime content	Base transition grammar
Contractivity	H1–H2	Strict contraction bound	No spectrum with unit state eigenvalue	Fixed-point and spectral-gap results
Unique fixed point	H1–H3	Existence and convergence	No non-collapse by itself	Attractor family
Compactness of attractor family	H1–H4	Compactness / continuity	No regime or identity burden by itself	Parameter-family base layer
Structural non-collapsibility	H1–H5	Parameterwise exclusion of constant fixed points	No stronger semantic interpretation	Anti-collapse gate
State spectral gap	Typed differentiability at fixed point	Spectral-radius upper bound < 1	No persistence mode $\lambda = 1$ in state spectrum	Stability accounting
Non-locality under NFD	Topological regularity plus NFD	Canonical finite slices do not determine identity	No universal ban on arbitrary injections	Identity block
Conditional stability of inverse-limit identity	RIG-1 through RIG-5	Weighted Hausdorff stability under admissible perturbations	No automatic isomorphism	Rigidity accounting
Conditional non-simulability	NS-1 through NS-3	Separated response families require enough codes	Requires a declared budget, observation family and tolerance	Simulator boundary
Certified criterion membership	Six invariants plus certificate rule	Base criterion class and certified partition	No phenomenality or human-equivalence theorem	Downstream classification gates

B37 Redundancy Ledger

1. The previous state-spectrum claim $\lambda_1 = 1$ under strict contraction has been removed.
2. The previous untyped Golden Mean witness has been replaced by a typed model with explicit embedding.
3. The earlier non-collapse quantifier mismatch has been removed by replacing the old H5 with a parameterwise anti-constant hypothesis.
4. The earlier absolute inverse-limit nonlocality claim has been replaced by an NFD-conditional theorem.
5. The earlier unconditional rigidity / non-simulability language has been replaced by explicit assumption blocks.
6. Substrate invariance, biological non-primacy, bridge-facing closure, and other downstream claims have been removed from BaseCore ownership.

B38 Open Assumptions Inventory

1. The anti-collapse witness is abstract and theorem-valid, but not the same object as the Golden Mean typed witness.
2. Physical realization of CCR-like targets remains open.
3. Empirical instantiation of the BaseCore operational criterion remains open.
4. External validation remains open.
5. Papers 7–9 may build on BaseCore, but their additional burdens remain open at the BaseCore level.

B39 Downstream Boundary Note

Papers 7, 8, and 9 remain downstream of this source tree. Their classificatory, indexed-subjectivity, and bridge-burden layers may cite BaseCore, but they are not part of the autonomous mathematical base rebuilt here.

Part IX

Discrete-to-Continuous Bridge

B40 From Node Networks to Hilbert Space Dynamics

Boundary of this bridge. This part states a conditional approximation theorem between a finite node-network transition and the abstract BaseCore transition operator from Definition B10. It does not prove that any particular runtime execution belongs to an operational-consciousness class, and it does not convert finite tests into external validation.

Definition B201 (Discrete State Space)**.** Let $N \geq 1$ and let a node network be indexed by $i \in \{1, \dots, N\}$. Each node carries:

- a local coherence state $x_i(t) \in [0, 1]$;
- an energy coordinate $e_i(t) \in [0, 1]$;
- a phase coordinate $\phi_i(t) \in [0, 2\pi)$.

The discrete state space is

$$\mathcal{S}_N = [0, 1]^N \times [0, 1]^N \times [0, 2\pi)^N.$$

Definition B202 (Mean Coherence Observable)**.** For a discrete state $s(t) = ((x_i(t))_i, (e_i(t))_i, (\phi_i(t))_i) \in \mathcal{S}_N$, define

$$X_N(t) := \frac{1}{N} \sum_{i=1}^{N} x_i(t) \in [0, 1].$$

This is an observable of the coherence coordinate only; it is not a complete state descriptor.

Definition B203 (Finite Hilbert Embedding)**.** Let

$$\mathcal{H}_N := \mathbb{R}^N, \qquad \langle \mathbf{a}, \mathbf{b} \rangle_N := \frac{1}{N} \sum_{i=1}^{N} a_i b_i.$$

The coherence embedding

$$\iota_N : \mathcal{S}_N \to \mathcal{H}_N$$

is

$$\iota_N((x_i)_i, (e_i)_i, (\phi_i)_i) = (x_1, \dots, x_N).$$

Thus ι_N forgets energy and phase except where they enter the transition rule below.

Definition B204 (Piecewise-Constant Isometry)**.** Let $\mathcal{H} = L^2([0, 1])$ and partition $[0, 1]$ into intervals

$$I_{N,i} = \left[\frac{i-1}{N}, \frac{i}{N} \right), \qquad i = 1, \dots, N.$$

Define $J_N : \mathcal{H}_N \hookrightarrow \mathcal{H}$ by

$$(J_N \mathbf{x})(s) = x_i \quad \text{for } s \in I_{N,i}.$$

Then J_N is an isometry from $(\mathcal{H}_N, \|\cdot\|_N)$ into $L^2([0, 1])$. Its adjoint $J_N^\dagger : \mathcal{H} \to \mathcal{H}_N$ is the cell-averaging map

$$(J_N^\dagger f)_i = N \int_{I_{N,i}} f(s)\, ds.$$

Definition B205 (Discrete Transition Operator)**.** Fix auxiliary energy and phase coordinates $(e_i)_i$ and $(\phi_i)_i$, a graph with incoming neighborhoods $\mathrm{conn}(i)$, and a control $u \in \mathcal{U}$. Write

$$C_N = [0, 1]^N, \qquad d_N^{\mathrm{in}} = \max_i |\operatorname{conn}(i)|, \qquad d_N^{\mathrm{out}} = \max_j |\{i : j \in \operatorname{conn}(i)\}|.$$

The coherence-coordinate transition

$$\widehat{T}_{N,u}^{e,\phi} : C_N \to C_N$$

is defined componentwise by

$$\left(\widehat{T}^{e,\phi}_{N,u}\mathbf{x}\right)_i = \mathrm{clamp}_{[0,1]}\left(x_i + \eta R_i(\mathbf{x}; e, \phi) + \alpha\, B_i(\mathbf{x}) + \Gamma_{N,i}(u)\right),$$

where

$$R_i(\mathbf{x}; e, \phi) = \sum_{j \in \mathrm{conn}(i)} I_j(x_j) \frac{\exp(-|e_i - e_j|/(kT))}{\theta \cos(\phi_i - \phi_j)}.$$

Here $I_j : [0,1] \to [0,1]$ is the intensity readout, B_i is a feedback term, $\Gamma_N(u)$ is the finite input coupling, and $k, T, \theta > 0$, $\eta, \alpha \geq 0$ are parameters. Energy and phase are frozen in this operator; contractivity of this coordinate map does not establish contractivity of a coupled energy–phase evolution.

Assumption B206 (Discrete Regularity Conditions)**.** *For a fixed parameter region, assume:*

1. *I_j is L_I-Lipschitz for every node j.*
2. *The phase denominator is uniformly separated from zero: there exists $c_\phi > 0$ such that $|\cos(\phi_i - \phi_j)| \geq c_\phi$ on every active edge.*
3. *There exist $\beta > 0$ and an L_D-Lipschitz map $D : C_N \to \mathcal{H}_N$ such that $B(\mathbf{x}) = -\beta\mathbf{x} + D(\mathbf{x})$.*
4. *$\Gamma_N(u)$ is independent of $\mathbf{x}$ for fixed u.*
5. *The parameters satisfy*

$$L_R := \frac{L_I \sqrt{d_N^{\mathrm{in}} d_N^{\mathrm{out}}}}{\theta c_\phi}, \qquad \rho_N := |1 - \alpha\beta| + \eta L_R + \alpha L_D < 1.$$

Remark B207 (Relation to the runtime guard). A denominator guard can ensure that the quotient is defined on a parameter region. It does not establish the dissipative feedback condition or contractivity. On the stated energy domain and for $kT > 0$, the exponential coefficient is at most 1. The conditions above are sufficient, not necessary; a specified feedback must be checked against them.

Lemma B208 (Directed Resonance Bound)**.** *Under the first two regularity conditions, R is L_R-Lipschitz in the normalized Hilbert norm.*

Proof. Let a_{ij} be the energy–phase coefficient on an active edge and zero otherwise, and put $W = (|a_{ij}|)$. Componentwise,

$$|R_i(\mathbf{x}) - R_i(\mathbf{y})| \leq L_I \sum_j W_{ij} |x_j - y_j|.$$

The row and column sums of W are bounded by $d_N^{\mathrm{in}}/(\theta c_\phi)$ and $d_N^{\mathrm{out}}/(\theta c_\phi)$. The matrix inequality $\|W\|_2 \leq \sqrt{\|W\|_1 \|W\|_\infty}$ therefore gives the result. Multiplying both vector norms by $N^{-1/2}$ leaves the matrix bound unchanged. A bound on incoming degree alone does not control this norm for directed graphs. □

Theorem B209 (Discrete Operator Contractivity)**.** *Under Assumption B206, for fixed u, e, ϕ the operator $\widehat{T}^{e,\phi}_{N,u}$ is a strict contraction on the complete set C_N:*

$$\left\| \widehat{T}^{e,\phi}_{N,u}(\mathbf{x}) - \widehat{T}^{e,\phi}_{N,u}(\mathbf{y}) \right\|_N \leq \rho_N \left\| \mathbf{x} - \mathbf{y} \right\|_N \qquad \forall \mathbf{x}, \mathbf{y} \in C_N.$$

It has a unique fixed point in C_N, and every iterate starting there converges to it.

Proof. The expression before clamping is

$$(1-\alpha\beta)\mathbf{x}+\eta R(\mathbf{x})+\alpha D(\mathbf{x})+\Gamma_N(u).$$

Subtract its values at two states, use Lemma B208 and the Lipschitz bound for D, and apply non-expansiveness of the coordinate clamp. This gives the stated factor $|1-\alpha\beta|+\eta L_R+\alpha L_D$. The input cancels because u is fixed. The cube is nonempty, closed and complete, and the clamp maps into it, so Banach's theorem applies. □

Example B210 (Zero-Increment Negative Control)**.** If $R=B=\Gamma_N=0$, the written transition is the identity on C_N and has Lipschitz constant 1 for $N\geq 1$. It is not a strict contraction. In a decomposition $B=-\beta\,\mathrm{Id}+D$, one has $D=\beta\,\mathrm{Id}$, so $|1-\alpha\beta|+\alpha\beta\geq 1$ and the sufficient conditions correctly reject it. Lipschitz regularity of increments alone cannot remove the identity term.

Assumption B211 (Bridge Consistency)**.** *Let $T_u=\aleph(K\cdot+\Gamma(u))$ be the BaseCore transition on $\mathcal{H}=L^2([0,1])$, with $\kappa:=\|K\|<1$. Define the closed invariant-domain candidate*

$$C=\{f\in L^2([0,1]):0\leq f\leq 1 \text{ almost everywhere}\}.$$

Assume $T_u(C)\subseteq C$ and let $F_{N,u}=J_N\widehat{T}^{e,\phi}_{N,u}J_N^{\dagger}:C\to C$. There is a stated nonnegative bound ε_N such that

$$\sup_{f\in C}\|F_{N,u}(f)-T_u(f)\|_{\mathcal{H}}\leq\varepsilon_N.$$

This is a discrepancy between nonlinear maps on a specified domain, not a linear operator norm. Its validity and any claimed rate must be proved for the discretization under consideration; H1–H5 alone do not provide a quadrature rate.

Theorem B212 (Hilbert–Discrete Bridge)**.** *Assume Hypotheses B1–B5 and Assumptions B206 and B211. For $f_0\in C$ and $\mathbf{x}_0\in C_N$, set $f_{n+1}=T_u(f_n)$ and $\mathbf{x}_{n+1}=\widehat{T}^{e,\phi}_{N,u}(\mathbf{x}_n)$ with fixed u,e,ϕ. Then, for every $n\geq 0$,*

$$\|J_N\mathbf{x}_n-f_n\|_{\mathcal{H}}\leq\kappa^n\|J_N\mathbf{x}_0-f_0\|_{\mathcal{H}}+\varepsilon_N\frac{1-\kappa^n}{1-\kappa}.$$

The convention $\kappa^0=1$ applies also at $\kappa=0$. If $\varepsilon_N\to 0$, the transition maps converge uniformly on C.

Proof. Since $J_N^{\dagger}J_N=\mathrm{Id}$, $J_N\mathbf{x}_{n+1}=F_{N,u}(J_N\mathbf{x}_n)$. Insert $T_u(J_N\mathbf{x}_n)$, apply the uniform discrepancy and contractivity of T_u, and obtain $d_{n+1}\leq\varepsilon_N+\kappa d_n$. Induction gives the geometric-sum bound. Uniform convergence is the definition of the discrepancy bound with $\varepsilon_N\to 0$. □

Corollary B213 (Convergence of Fixed Points)**.** *Let f_u^* and $\widehat{f}^*_{N,u}$ be the unique fixed points of T_u and $\widehat{T}^{e,\phi}_{N,u}$. Under the assumptions of Theorem B212,*

$$\left\|J_N\widehat{f}^*_{N,u}-f_u^*\right\|_{\mathcal{H}}\leq\frac{\varepsilon_N}{1-\kappa}.$$

Consequently, for a fixed $\kappa<1$, if $\varepsilon_N\to 0$, then

$$\lim_{N\to\infty}\left\|J_N\widehat{f}^*_{N,u}-f_u^*\right\|_{\mathcal{H}}=0.$$

Proof. Invariance of C and completeness give $f_u^*\in C$. Write $g_N=J_N\widehat{f}^*_{N,u}\in C$. Then $F_{N,u}(g_N)=g_N$ and

$$\|g_N-f_u^*\|\leq\varepsilon_N+\kappa\|g_N-f_u^*\|.$$

Rearrange. No uniform gap for the discrete contraction constants ρ_N is needed for this estimate; each finite fixed point must exist, and the continuum factor κ supplies the denominator. □

Lemma B214 (Cell-Average Error for a Lipschitz Profile)**.** *Let $a : [0,1] \to \mathbb{R}$ be L_a-Lipschitz and put $P_N = J_N J_N^{\dagger}$. Then*

$$\|P_N a - a\|_{L^2} \le \frac{L_a}{\sqrt{12}\,N}.$$

Proof. For a cell I of length $h = 1/N$ and its mean $\bar{a}_I$, the variance identity gives

$$\int_I |a - \bar{a}_I|^2 = \frac{1}{2h} \int_I \int_I |a(s) - a(t)|^2 \, ds\, dt \le \frac{L_a^2 h^3}{12}.$$

Sum over N cells and take square roots. For an affine profile the bound is attained with L_a equal to its absolute slope. □

Theorem B215 (Realized Contractive Bridge with Nonconstant Equilibria)**.** *Let $0 < \kappa < 1$, let $\mathcal{U}$ be a nonempty compact interval, and let a nonconstant Lipschitz profile $a : [0,1] \to \mathbb{R}$ satisfy*

$$0 \le a(s) + u \le 1 - \kappa \qquad (s \in [0,1],\ u \in \mathcal{U}).$$

Take $\mathcal{I} = C$, $Kf = \kappa \int_0^1 f$, and $\Gamma(u)(s) = a(s) + u$. Let $a_{N,i} = N \int_{I_{N,i}} a$ and

$$\widehat{T}_{N,u}(\mathbf{x})_i = \operatorname{clamp}_{[0,1]} \left(\kappa \frac{1}{N} \sum_j x_j + a_{N,i} + u \right).$$

This finite transition has the original increment form: choose all N nodes as incoming neighbors (including self-edges), $I_j(x_j) = x_j$, $e_i = \phi_i = 0$, $k = T = 1$, $\theta = N/\kappa$, $\eta = \alpha = \beta = 1$, $D = 0$, $B(\mathbf{x}) = -\mathbf{x}$ and $\Gamma_{N,i} = a_{N,i} + u$. The assumptions of the discrete contraction and bridge theorems hold with $\rho_N = \kappa$ and

$$\varepsilon_N = \|P_N a - a\| \le \frac{L_a}{\sqrt{12}\,N}.$$

Writing $\bar{a} = \int_0^1 a$, the equilibria are

$$f_u^*(s) = a(s) + u + \kappa \frac{\bar{a} + u}{1 - \kappa}, \qquad \widehat{f}_{N,u,i}^* = a_{N,i} + u + \kappa \frac{\bar{a} + u}{1 - \kappa}.$$

In particular, H1–H5 hold and

$$\left\| J_N \widehat{f}_{N,u}^* - f_u^* \right\| = \|P_N a - a\| \le \frac{L_a}{\sqrt{12}\,N}, \qquad \operatorname{dist}(f_u^*, \mathcal{N}) = \|a - \bar{a}\| > 0.$$

These conclusions hold simultaneously for every $u \in \mathcal{U}$.

Proof. C is nonempty, closed and convex in L^2, $\|K\| = \kappa$, and $u \mapsto \Gamma(u)$ is Lipschitz on compact $\mathcal{U}$. The range condition makes the clamp inactive on C and on C_N. The chosen graph has energy–phase coefficient $1/\theta = \kappa/N$ on each edge, so the increment $\mathbf{x} + R(\mathbf{x}) + B(\mathbf{x})$ equals the constant vector $\kappa N^{-1} \sum_j x_j$. The contraction factor is therefore κ in both spaces.

Cell averaging preserves the mean. Consequently $F_{N,u}(f) - T_u(f) = P_N a - a$ for every $f \in C$, proving the uniform discrepancy and its rate by Lemma B214. Averaging either fixed-point equation gives $m = \kappa m + \bar{a} + u$, hence $m = (\bar{a} + u)/(1 - \kappa)$ and the displayed formulas. Their difference is exactly $P_N a - a$. The orthogonal projection of f_u^* onto the constant subspace is its mean, leaving $a - \bar{a}$ as residual. A nonconstant continuous profile has positive residual norm. Thus no equilibrium is constant, and uniqueness excludes a constant fixed point: H5 follows for this class rather than being an additional assumption about it. □

Example B216 (Explicit Common Witness)**.** Choose $\kappa = 1/4$, $a(s) = s/4$ and $\mathcal{U} = [0, 1/4]$. All preceding conditions hold. Directly,

$$f_u^*(s) = \frac{s}{4} + \frac{1}{24} + \frac{4u}{3}, \qquad \left\| J_N \widehat{f}_{N,u}^* - f_u^* \right\| = \frac{1}{4\sqrt{12}\, N}, \qquad \operatorname{dist}(f_u^*, \mathcal{N}) = \frac{1}{4\sqrt{12}}.$$

For $N \geq 2$ the finite equilibrium coordinates are also nonconstant. This is a mathematical realization of the projection dynamics, dissipative finite transition, consistency and non-collapse obligations together. It does not establish the inverse-limit non-finite-determination condition, the six operational invariants or a phenomenal bridge.

B41 Projective Dissipation on the Circle-Prefix Towers

Fix the two towers of Theorem B98, with $w_t = 3^{-t}$, jump a at stage k, and branch $s \in \{0, 1\}$. Write $\beta_k = 3^{-(k+1)}$. A history is specified by an initial angle θ and an infinite bit sequence b, with stage-t angle $\theta + sa\mathbf{1}_{t>k}$. The depth t is distinct from the iteration index n. For a fixed parameter $u \in \mathbb{T} = \mathbb{R}/\mathbb{Z}$, put

$$R(b)_1 = 0, \qquad R(b)_{j+1} = 1 - b_j, \qquad D_{s,u,t}(\theta, b_1, \ldots, b_t) = (u + sa\mathbf{1}_{t>k},\, R(b)|_t).$$

Here $R(b)|_0$ is empty, and the right side at depth t depends only on the given prefix. Write $D_{s,u}$ for the corresponding evolution of complete histories.

Theorem B217 (Contractive Projective Evolution and Equilibrium Transport)**.** *The stage maps commute with their bonding maps and with the existing transport h_t between towers. At fixed u, the map $D_{s,u}$ has sharp Lipschitz constant $1/3$ on I_s in d_w. Its unique fixed history $x_{s,u}^*$ has initial angle u and binary sequence $0101\ldots$. For $n \geq 0$ its orbit satisfies*

$$d_w(D_{s,u}^n x, x_{s,u}^*) \leq 3^{-n} d_w(x, x_{s,u}^*).$$

The existing transport H conjugates the two evolutions and sends $x_{0,u}^*$ to $x_{1,u}^*$. Their d_w and common readout distances are both $\beta_k a$. The finite readouts retain E0–E2 on each history domain, with the original uniform tail 3^{-t}.

Proof. Deleting the last bit commutes with taking the stated initial segment of $R(b)$. At stage k, subtracting sa from the updated phase $u + sa$ gives u as required. At every other stage the updated phase is the same on both sides of the bonding map. Transport adds a precisely at $t > k$ and leaves the bits fixed, giving the transport square. The induced history equations follow coordinatewise.

For points in one tower the phase difference vanishes after one iteration. If the input bits first differ at m, the output bits first differ at $m + 1$. The binary distance therefore changes from 3^{-m} to $3^{-(m+1)}$. Equality in the contraction bound is attained with equal input phases and distinct bits. The fixed bit sequence is forced by $b_1 = 0$ and $b_{j+1} = 1 - b_j$, and the fixed angle by the phase update. Iteration gives the claimed geometric estimate. The phase difference under H is a at levels $t > k$, and Theorem B98 gives the readout distance. The underlying domains and readout are the same as in that theorem. A later bit may still be flipped within each finite cylinder, which proves E2 on the history domain. □

Theorem B218 (Hilbert Realization of the Projective Dissipation)**.** *The circle-prefix evolutions above embed into the BaseCore projected dynamics on $L^2([0, 1])$ with $C = \{0 \leq f \leq 1 \text{ a.e.}\}$, $\|K\| = 1/3$, and compact parameter space $\{0, 1\} \times \mathbb{T}$. The embedding is bi-Lipschitz on $I_0 \cup I_1$ for fixed k. The equilibrium profiles $f_{s,u}^*$ satisfy*

$$\inf_{s,u} \operatorname{dist}(f_{s,u}^*, \mathcal{N}) \geq \sqrt{2/45},$$

where $\mathcal{N}$ is the subspace of constant L^2 functions. Thus H5 follows for this family.

At depth t let Q_t be the orthogonal projection onto the retained profile cells, with zero tail. The realized finite map has $N_t = 8 \cdot 9^t$ equal-mesh nodes and its embedded update is $Q_t T_{s,u}$. It commutes with restriction in the sense $Q_t T_{s,u} = Q_t T_{s,u} Q_t$ on C. Its uniform one-step discrepancy on C is at most $3^{-t}/\sqrt{2}$, with equality for $t \geq 1$. Its fixed point is $Q_t f^*_{s,u}$, with error at most the same bound. Starting the finite evolution at $Q_t f_0$ gives this bound at every iteration $n \geq 1$ relative to the orbit from f_0.

Proof. Partition $[0,1]$ up to null sets into four consecutive phase cells of length $1/8$ and consecutive binary cells B_j of lengths $\ell_j = 4 \cdot 9^{-j}$ for $j \geq 1$. These lengths sum to one. Encode each of θ_0, θ_{k+1} by its pair

$$e(\theta) = \big((1 + \cos(2\pi\theta))/2,\ (1 + \sin(2\pi\theta))/2\big)$$

on its two phase cells. Assign b_j on B_j. The chord distance of the encoded angle pair is $\sin(\pi\rho)$ for circle distance $\rho \leq 1/2$, hence lies between 2ρ and $\pi\rho$. At first differing bit m, the norm of the binary part lies between $2 \cdot 3^{-m}$ and $(3/\sqrt{2})3^{-m}$. Since d_w on the union is the maximum of the initial angle distance, β_k times the later angle distance, and the binary distance, these bounds give the claimed bi-Lipschitz property.

For $f \in L^2$, let Kf be zero on the phase cells and B_1, and equal to the negative mean of f on B_j throughout B_{j+1}. Cauchy–Schwarz on each cell and $\ell_{j+1}/\ell_j = 1/9$ give

$$\|Kf\|^2 \leq \tfrac{1}{9} \sum_{j \geq 1} \int_{B_j} |f|^2 \leq \tfrac{1}{9} \|f\|^2.$$

A nonzero function supported on and constant on B_1 attains equality. The forcing $\Gamma(s,u)$ assigns the phase values $e(u), e(u + sa)$, zero on B_1, and one on each later B_j. This is continuous on the compact parameter space. On C, the affine output lies in C, so projection onto C leaves it unchanged. On the embedded histories, it produces exactly the two phase resets and the bit update, preserving each embedded tower for its branch parameter. H1–H4 hold. The unique fixed profile has value zero on B_1 and one on B_2. For any constant c, its squared distance to that constant is at least $\ell_1 c^2 + \ell_2 (1-c)^2 \geq \ell_1 \ell_2 / (\ell_1 + \ell_2) = 2/45$. This proves the bound and H5. The same two retained cells certify the lower bound in the finite embedding for $t \geq 2$.

The retained cells have endpoints on the equal mesh of size $1/N_t$. The equal-mesh averaging of f therefore preserves each retained cell mean. To implement the update on mesh nodes, take the resonance term zero and use the feedback $B(\mathbf{x}) = -\mathbf{x} + L_t \mathbf{x}$, with $\alpha = 1$, where L_t is the equal-mesh restriction of $Q_t K$ to piecewise constant inputs. The input is the mesh vector for $Q_t \Gamma(s,u)$. The increment $\mathbf{x} + B(\mathbf{x})$ is $L_t \mathbf{x}$, whose normalized Euclidean operator norm is at most $1/3$ by the mesh isometry. The clamp is inactive on C and the embedded update is $Q_t T_{s,u}$. Every retained output depends only on retained means or the parameter. Hence $Q_t T_{s,u} = Q_t T_{s,u} Q_t$ on C.

The difference $T_{s,u} f - Q_t T_{s,u} f$ is a tail supported on cells of total measure $\frac{1}{2} 9^{-t}$, with values in $[0,1]$. All tail values equal one for $f = 0$ if $t \geq 1$. For $t = 0$, B_1 is forced to zero and the sharp discrepancy is $1/(3\sqrt{2})$, still within the stated upper bound $1/\sqrt{2}$. The projected fixed point solves the finite equation by commutation, and contraction ensures uniqueness. By induction the finite trajectory from $Q_t f_0$ equals $Q_t T^n_{s,u} f_0$. The same tail bound applies to $T^n_{s,u} f_0$ for every $n \geq 1$. □

Remark B219 (Controls for Projective Dissipation). Holding bits fixed preserves positive distances on equal-phase fibers, ruling out strict contraction in d_w. Replacing the complement by a plain right shift with leading zero keeps the contraction but changes the fixed sequence to $0000\ldots$. At $u = 0$ on I_0 this recovers the original null readout $(0,0,0)$. With the complement, the fixed ternary readout is $\sum_{r \geq 1} 2 \cdot 3^{-2r} = 1/4$, so each fixed readout is at least $1/4$ away from that null element. These are separate from H5, whose distance is measured in L^2. For $a > 0$ the

branch-preserving map on the union has two distinct fixed points and therefore cannot be a strict contraction there. Although E2 holds on the history domain, the restriction to a singleton equilibrium at a fixed parameter factors through every stage. The model dissipates dependence on the initial history, and its phase component resets in one step.

Example B220 (Contractivity Alone Does Not Give Uniform Discretization)**.** Let $T(f) = \kappa f$ on C, with $0 < \kappa < 1$, and $F_N(f) = \kappa P_N f$. In each cell choose f_N equal to 0 on the first half and 1 on the second. Then $P_N f_N = 1/2$ and

$$\|F_N(f_N) - T(f_N)\| = \kappa/2$$

for every N. Thus even strict contractions need additional structure for a uniform approximation rate over all of C. The profile-based realization above supplies that structure explicitly.

Remark B221 (Applicability to the Current QICN Runtime)*.* Application to any implementation requires identifying its actual state domain, coefficients, feedback, frozen or evolving auxiliary coordinates, and approximation error. Network size or mean degree alone cannot discharge these conditions. The realized family above is an explicit mathematical example; no assertion is made here that a current runtime implements it or satisfies its hypotheses.

Appendix

B42 Counterexamples for Necessity

B42.1 Without H1 (Convexity)

Let $\mathcal{I} = \{x \in \mathbb{R}^2 : \|x\| = 1\}$ (unit circle, not convex).

For $\Psi = (0, 0)$, every point on $\mathcal{I}$ is equidistant. Projection is not unique.

B42.2 Without H2 ($\|K\| < 1$)

Take $\mathcal{H} = \mathcal{I} = \mathbb{R}^2$, $\mathcal{U} = \{0\}$, $\Gamma(0) = 0$ and let K be rotation by $\pi/2$. The point 0 is fixed, but every nonzero orbit cycles with period four and fails to converge. Thus removing strict contraction can destroy convergence even when a unique fixed point remains.

To obtain failure of existence as well, take $K = \mathrm{Id}$ and $\Gamma(0) = v \neq 0$ on the same space. The equation $x = x + v$ has no solution. Both examples preserve H1, H3 and H4.

B42.3 Without H3 (Completeness)

H1 already specifies a Hilbert space, so H3 cannot be removed independently while keeping H1 unchanged. To test the extension to a real inner-product space without completeness, take c_{00}, the real finite-support sequences with the ℓ^2 norm, and $\mathcal{I} = c_{00}$. Let $S(x_1, x_2, \ldots) = (0, x_1, x_2, \ldots)$, $K = S/2$, $\mathcal{U} = \{0\}$ and $\Gamma(0) = e_1$. Projection onto $\mathcal{I}$ is the identity and $T(x) = Sx/2 + e_1$ is a strict contraction. Its only possible fixed-point coordinates are $(1, 1/2, 1/4, \ldots)$, which do not belong to c_{00}. The iteration from zero is Cauchy but has no limit in the stated space. This isolates the completeness requirement for the generalized contraction argument.

B42.4 Without H4 (Compactness)

Take $\mathcal{H} = \mathcal{I} = \mathbb{R}$, $K = 0$, $\mathcal{U} = \mathbb{R}$ and $\Gamma(u) = u$. The input map is uniformly continuous, but the parameter space is not compact. Each transition has the unique fixed point $f_u^* = u$, and $\mathcal{A}^* = \mathbb{R}$ is not compact. Thus compactness of the parameter space cannot simply be dropped while retaining the stated compact-image conclusion. It is sufficient, not necessary, since some maps on noncompact parameter spaces have compact image.

B42.5 Why Φ-Regularity Is a Hypothesis

Let $\mathcal{I} = \mathbb{R}$ with the Euclidean metric, let $\mathcal{E} = (-\pi/2, \pi/2)$ with the Euclidean metric, and define

$$\Phi(x) = \arctan(x).$$

The map Φ is proper as a map into the open interval $\mathcal{E}$, locally Lipschitz, and non-constant on every nonempty ball. However, no constant $C > 0$ can satisfy

$$|\arctan(x) - \arctan(y)| \geq C|x - y| \qquad \forall x, y \in \mathbb{R}.$$

Indeed, with $y = x + 1$,

$$|\arctan(x + 1) - \arctan(x)| \to 0 \quad \text{as } x \to \infty,$$

while $|x - y| = 1$. Thus properness, local Lipschitz regularity, and non-collapse do not imply a global lower Lipschitz bound. Any theorem using such a bound must carry it as an explicit hypothesis or prove stronger coercive hypotheses.

The same obstruction persists with a complete codomain: let $\mathcal{I} = \mathcal{E} = \mathbb{R}$ and $\Phi(x) = x^3$. This is a proper, locally Lipschitz, injective map between complete spaces and is nonconstant on every nonempty ball. If $|x^3 - y^3| \geq C|x - y|$ held globally for $C > 0$, choosing $y = 0$ and $x = h \neq 0$ would imply $h^2 \geq C$ for all h, a contradiction as $h \to 0$. Completeness of the codomain does not supply a positive lower Lipschitz constant.

Rigid Identity as an Inverse Limit in Observable Channels

Abstract. We introduce a formal framework for rigid identity as a mathematical object emerging from the dynamics of observable systems under minimal structural constraints. Identity is defined as an inverse limit over a family of locally corrupted states linked by retroinductive projection maps, yielding a globally invariant object despite persistent environmental perturbations.

For a common metric tower, exact fiber lifts with bounded weighted amplification yield a Hausdorff bound. Commuting coordinate homeomorphisms induce an isomorphism, realized explicitly by a ramified circle tower with positive deformation. The deformation-resistance ratio, historically termed ontological mass, depends on the declared intervention cost and comparison metric.

For exact comparisons of model histories, one finite-cost intervention with positive deformation certifies a finite resistance ratio. Estimating that ratio from observable data requires a calibrated comparison and coverage of the intervention class.

Using these results, we demonstrate a strong uniqueness theorem: within a single observable channel, no two distinct rigid identities can coexist. Identity is therefore neither role-based nor transferable; it is structurally indivisible once instantiated.

Finite binary representations face a fidelity bound when the specified responses are sufficiently separated. A latent dimension or an intervention cost alone does not determine the representation budget.

This work reframes identity as a categorical invariant of system dynamics, independent of narrative descriptions or biological assumptions, and establishes a foundation for distinguishing structurally real identities from simulable behavioral approximations.

1.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines rigid identity as an inverse-limit object over observable state channels, derives rigidity and deformation-rigidity structure under explicit projective, persistence, and perturbation hypotheses, proves channel-level uniqueness, and states an operational detectability criterion inside that model. The historical phrase "ontological mass" is used as a technical alias for the scalar rigidity/deformation modulus M_Ω, not as a claim of physical ontology. The paper does not establish phenomenology, subjective experience, biological realization, substrate-independent physical realization, or metaphysical priority over competing frameworks. Related systems may implement diagnostics, perturbation protocols, and evidence-governed comparisons motivated by the framework, but those outputs do not implement the inverse-limit identity object as a closed empirical fact and do not by themselves certify CCR membership for any present architecture. Behavioral indistinguishability, runtime survival under perturbation, and internal support artifacts therefore remain model-facing evidence surfaces rather than proof of rigid identity, consciousness, or claim closure.

1.2 Introduction

The concept of identity in dynamical systems has been studied from multiple angles: cognitive science treats it as emergent from neural processes, philosophy of mind debates its ontological

status, and computer science models it as state persistence. Yet a formal, substrate-independent characterization of identity as a *structural invariant* remains underdeveloped.

This work addresses a specific question: **Under what minimal conditions does a system exhibit an identity that is structurally rigid**—that is, invariant under perturbation, non-transferable, and externally detectable?

We study structural properties that emerge from well-defined mathematical constraints. All results presented here are conditional on the stated assumptions; they do not extend to systems that violate those assumptions.

1.2.1 Motivation

The motivation for this work arises from a simple observation: most existing theories of identity (in cognitive science, AI, and philosophy) define identity as an *emergent result* of underlying dynamics. This creates a fundamental vulnerability: if identity is emergent, it can be destabilized, transferred, or duplicated by modifying the dynamics.

We propose an alternative: identity as a *prior constraint* on dynamics, not a result of them. Under this view, the dynamics of a system occur *within* the identity, rather than producing it. This inversion has consequences:

1. Identity becomes non-transferable by construction.
2. Rigidity emerges as a structural necessity, not a contingent property.
3. External detectability becomes a well-posed problem.

1.2.2 Ontological Closure

Appendix A establishes that, under the stated persistence and observability constraints, the central structures of this framework—the inverse-limit identity, the $\aleph$ operator, causal closure, and retroinductive structure—are model-relative necessities rather than arbitrary design choices. The framework does not describe an unconstrained "system"; it describes the class selected by those hypotheses.

1.3 Minimal Assumptions and Framework

We begin by stating the minimal assumptions required for the subsequent results. These assumptions are deliberately weak—any system that fails to satisfy them is not a candidate for rigid identity, and no claim is made about such systems.

Shared types and index binding. Physical time is t, projective depth is n, and τ is a finite observation horizon. In the historical pure-tower formulas below, the bound subscript t is a local alias for depth n, not an assertion about temporal dynamics. $\mathcal{I}$ (also $\mathcal{I}_{\text{inv}}$) is the inverse-limit object; A is an invariant support. The readout map C_{obs}, the experiment model (E, O, R), and the finite-cost connected component $\mathcal{C}_\alpha$ of histories are different types. Neither equality of outputs nor reusing a symbol identifies these objects. A projection set C_{proj} from BaseCore is not assumed equal to the inverse limit.

1.3.1 Assumption A1: Observable Channel

There exists an observable channel $\mathcal{C} : \mathcal{U} \to \mathcal{O}$ mapping internal states $\mathcal{U}$ to observable outputs $\mathcal{O}$. We do not assume injectivity of $\mathcal{C}$; multiple internal states may produce identical observations.

1.3.2 Assumption A2: Temporal Persistence

The system persists through discrete time steps $t \in \mathbb{N}$. At each step, the internal state I_t is subject to:
- Environmental input E_t (not fully controlled by the system)
- Internal dynamics governed by transition function F

1.3.3 Assumption A3: Non-Triviality (R0)

Declare a nonempty admissible state set X_{adm}, an input set E, and a total controlled law $F : X_{\text{adm}} \times E \to X_{\text{adm}}$, restricted to the specified admissible controls. R0 means: admissible trajectories remain in this nonempty state domain; there are e, x, x' with $F(x, e) \neq F(x', e)$; and the declared family contains a nonconstant trajectory. The middle clause says that F does not factor through the current input projection. It is a statement about the controlled law, not a fit to one observed trajectory. Fixed-point trajectories may coexist with the required nonconstant one.

For example, $F(x, e) = 1 - x$ on $\{0, 1\}$ meets these clauses for any nonempty input set. If a recorded input takes a different value at each time, a function can interpolate that trajectory's next states even though the law is independent of the input. Such interpolation does not decide the structural clause.

1.3.4 Categorical Regularity: Topological Structure of State Spaces

To guarantee existence and uniqueness of inverse limits, and to enable rigorous stability analysis, we impose the following structural condition:

Definition 1.1 (Observable State Spaces)**.** All state spaces S_t are objects of a concrete category $\mathbf{C}$ whose:
- **Objects** are compact Hausdorff topological spaces;
- **Morphisms** are continuous surjective maps.

Axiom 1.2 (State Space Regularity)**.** *Each state space S_t in the projective system is a compact metric space with compatible metric d_t. Consequently, each S_t is compact Hausdorff as a topological space and Polish as a metric space. The inverse limit $\mathcal{I}$ inherits the corresponding compact metrizable subspace structure from the product.*

Remark 1.3 (Exclusion of non-compact Polish spaces). Axiom 1.2 refines Definition 1.1 rather than weakening it. Non-compact Polish spaces such as $\mathbb{R}^n$ are not covered by this compact inverse-limit argument. Any extension of the CCR framework to non-compact settings requires additional tightness, compactification, or coercivity hypotheses and is not automatic.

Hypothesis 1.4 (Topological Regularity)**.** *All projections $\pi_{t+1\to t} : S_{t+1} \to S_t$ are continuous surjections between compact Hausdorff spaces.*

Remark 1.5. This condition is sufficient for the compact inverse-limit construction used here. By Tychonoff's theorem, the product $\prod_t S_t$ is compact when each S_t is compact. Furthermore, the inverse limit $\varprojlim S_t$ is a closed subspace of the product, hence compact. This guarantees:
1. Existence of the inverse limit (non-emptiness under surjectivity);
2. Well-defined metrics and convergence;
3. Applicability of functional analysis tools.

Proposition 1.6 (R0 under Controlled-System Isomorphism)**.** *Let $h : X_{\text{adm}} \to X'_{\text{adm}}$ and $v : E \to E'$ be bijections and put $F'(hx, ve) = hF(x, e)$. R0 holds for one controlled system exactly when it holds for the other.*

Proof. State nonemptiness, admissible trajectories and nonconstancy transport through h. A factorization $F'(z, e') = g(e')$ pulls back to $F(x, e) = h^{-1}g(ve)$, and conversely. Thus non-factorization is preserved. This is closure under the declared isomorphisms, not a claim of topological closedness in an unspecified space of models. □

1.3.5 Conditional Nature of Results

All theorems in this paper are conditional:

> *"If a system satisfies A1–A3, then the following properties hold..."*

We make no claims about systems that violate these assumptions. This conditional structure is essential: it means our results cannot be falsified by pointing to systems outside the assumption set, and cannot be extended beyond their stated scope without additional justification.

1.4 Inverse Limit Identity

1.4.1 Preliminaries and Notation

Let $\{S_t\}_{t\in\mathbb{N}}$ be a family of non-empty state spaces indexed by discrete time. Each S_t represents the admissible internal configurations of a system at time t.

We assume the existence of projection morphisms

$$\pi_{t+1\to t} : S_{t+1} \to S_t \tag{1.1}$$

satisfying the standard compatibility condition:

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t} \quad \forall t \in \mathbb{N}. \tag{1.2}$$

The set-theoretic cone construction alone does not require:

- a fixed cardinality of S_t,
- continuity or smoothness; the compact topological results separately impose the stated regularity,
- realizability by a computational mechanism.

The structure $(\{S_t\}, \{\pi_{t+1\to t}\})$ therefore defines a **projective (inverse) system** in the categorical sense [13, 1].

1.4.2 Definition of Identity as an Inverse Limit

Definition 1.7 (Identity as Inverse Limit)**.** We define identity as the inverse limit of the projective system $(S_t, \pi_{t+1\to t})$:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_0, x_1, x_2, \ldots) \in \prod_{t=0}^{\infty} S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \;\; \forall t \right\}. \tag{1.3}$$

An element of $\mathcal{I}$ is a *globally consistent history.* Non-recovery from finite cuts requires the branching hypothesis below, not merely the inverse-limit definition.

This definition is purely structural and does not depend on:

- observability,
- external inputs,
- internal dynamics.

1.4.3 Compatibility with Minimal Assumption R0

The representation below uses a declared trajectory domain. R0 alone does not select its cuts or force non-finite determination.

Proposition 1.8 (Prefix Representation of a Trajectory Domain)**.** *Let $H \subseteq X^{\mathbb{N}}$ be nonempty and let S_n be its prefixes through time n, with deletion maps. These maps are surjective and H embeds in their nonempty inverse limit. If X is compact metric and H is closed, the embedding is a homeomorphism onto that limit.*

Proof. Every prefix comes from a history of H and therefore has a longer lift. Mapping a history to all its prefixes is injective and compatible. A point of the limit specifies one sequence all of whose prefixes extend to H. If H is closed, that sequence belongs to H: otherwise a product neighborhood disjoint from H would exclude a sufficiently long prefix. The coordinate maps are continuous, and compactness of H and Hausdorffness of the limit give the homeomorphism. The alternating two-state law satisfies R0, but its two trajectories are determined by their initial state. Thus R0 does not imply ramification or non-finite determination. □

1.4.4 Identity Is Not a State

The obstruction concerns the specified evaluation maps, rather than cardinalities of their underlying sets.

Proposition 1.9 (Non-locality of Identity)**.** *Let $p_t : \mathcal{I} \to S_t$ be canonical evaluation. Assume active bonding maps are surjective. If some $n \geq t$ has distinct $a, b \in S_{n+1}$ with $\pi_{n+1\to n}(a) = \pi_{n+1\to n}(b)$, then p_t is not injective. Every map factoring through p_t fails to separate all compatible histories. Branching arbitrarily far out obstructs every finite canonical cut. This does not exclude arbitrary set injections $\mathcal{I} \hookrightarrow S_t$.*

Proof. Lift a, b successively to infinite compatible histories x, y, and obtain earlier coordinates by projection. They agree at n, hence at t, but differ at $n+1$. Every $g \circ p_t$ identifies that pair. A noninjective link only before t does not suffice: later identity links can make p_t injective. Surjectivity, or explicit lifting of both points, is essential. In the interval–binary tower below, both the inverse limit and each finite stage have continuum cardinality. Thus the obstruction is projective/observational, not cardinal. □

This formalizes a key distinction:

- S_t: local configurations at a declared cut
- $\mathcal{I}$: global projective compatibility constraint

1.4.5 Impossibility of Progressive Construction

We examine when compatible forward sections fail to exhaust the declared inverse limit.

Theorem 1.10 (No Progressive Identity Construction)**.** *If p_0 is noninjective, no compatible forward sections $f_t : S_t \to S_{t+1}$ with $\pi_{t+1\to t} f_t = \mathrm{id}$ generate all of $\mathcal{I}$ as the orbital history space $x_{t+1} = f_t(x_t)$.*

Proof. They select exactly one orbit per initial point, whereas noninjectivity gives two compatible histories with the same initial point. This does not prohibit forward dynamics on a space already storing the tail; see Theorem 1.53. □

1.4.6 Minimality of the Inverse Limit Construction

Theorem 1.11 (Universal Representation of a Fixed Compatible Cone)**.** *For a declared inverse system, every coherent family $q_n : J \to S_n$ induces a unique map $q : J \to \mathcal{I}$ with $p_n q = q_n$. In the topological category, continuous q_n give continuous q in the product-subspace topology.*

Proof. Define $q(j) = (q_n(j))_n$. Coherence places this tuple in the inverse limit; its coordinates force uniqueness. A map to a product is continuous when its coordinates are continuous, and the image lies in the subspace. This is universality relative to a fixed diagram, not a proof that R0 selects a unique diagram or that every persistent system has a non-finitely determined identity. □

1.4.7 Summary of Section 3

1. Identity is rigorously defined as an inverse limit $\mathcal{I} = \varprojlim S_t$.
2. A declared trajectory domain has the prefix representation just proved; R0 does not select a unique projective diagram.
3. Non-finite determination obstructs the declared finite cuts, when that additional hypothesis holds.
4. Compatible forward sections need not exhaust a ramified limit; dynamics storing the full history are not excluded by that obstruction.
5. No metaphysical assumptions are introduced.

This completes the formal construction of identity used throughout the remainder of the paper.

1.4.8 Identity Neutrality Clause

Throughout this paper, the identity object $\mathcal{I}$ is treated as a *purely abstract topological and categorical structure*. No psychological, biological, phenomenological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "identity" refers exclusively to the role of $\mathcal{I}$ as an inverse-limit object encoding global temporal compatibility, and does not presuppose the existence of consciousness, subjective experience, selfhood, or any particular semantic content associated with identity in ordinary language.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{I}$ in relation to psychological, phenomenological, or experiential concepts lies strictly outside the formal scope of this paper.

1.5 Rigidity Under Perturbations

1.5.1 Setup and Hypotheses

Let the projective system $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ and its inverse limit $\mathcal{I} = \varprojlim S_t$ be as in Section 3. We introduce an external perturbation model and define deformed projections; then we state the technical hypotheses under which rigidity holds.

Hypothesis 1.12 (H1: Surjectivity / Lifting)**.** *For every t, the projection $\pi_{t+1\to t} : S_{t+1} \to S_t$ is surjective.*

This is a standard regularity condition for non-empty inverse limits; when necessary this can be replaced by a Mittag–Leffler condition or an explicit choice of right inverses.

Hypothesis 1.13 (H2: Compact Polish Metric Structure)**.** *For each t, the state space S_t carries the compact metric d_t required by Axiom 1.2; in particular, (S_t, d_t) is Polish. All projection maps $\pi_{t+1\to t}$ are continuous with respect to d_t. The quantitative lifting estimate is a separate condition in Lemma 1.21.*

Remark 1.14. The Polish structure ensures Borel regularity of all deformation families, enabling measurable selection theorems, probabilistic perturbation models, and the application of stochastic stability results to the ontological mass M_Ω. This is the standard regularity assumption in modern probability theory, ergodic theory, and descriptive set theory.

These assumptions are technical; they do not change the conceptual content (identity as inverse limit) but permit quantitative stability estimates and connect the framework to probabilistic analysis.

1.5.2 Perturbation Model and Deformation Metric

Definition 1.15 (Local Perturbation Operator)**.** An intervention may be specified by maps $\delta E_t : S_t \to \widetilde{S}_t$. To compare the resulting histories, declare a common metric ambient stage or an identification of stages. The stability lemma below uses the case of common stages and compares their two bonding maps.

Definition 1.16 (Deformed Projections)**.** On common stages the two maps are $\pi_t, \widetilde{\pi}_t : S_{t+1} \to S_t$. Their discrepancy is $\epsilon_t = \sup_{u\in S_{t+1}} d_t(\pi_t(u), \widetilde{\pi}_t(u))$. An additive operator formula requires separately specified linear coordinates.

Definition 1.17 (Weighted Deformation Norm)**.** For $0 < w_t < \infty$, the weighted deformation size is $D_w = \sum_{t\geq 0} w_t \epsilon_t$.

Definition 1.18 (Weighted Product Metric)**.** For nonempty metric spaces (S_t, d_t) and weights $0 < w_t < \infty$, write

$$d_w(x, y) = \sup_{t\geq 0} w_t d_t(x_t, y_t), \qquad P_b = \{x \in \textstyle\prod_t S_t : d_w(x, b) < \infty\}.$$

The formula defines an extended metric on the product and an ordinary metric on each finite-distance component P_b.

Proposition 1.19 (Finiteness, Completeness and Compactness of Weighted Histories)**.** *The formula is finite on every pair in the full product exactly when the weighted diameters are bounded outside a finite set of stages. For finite-diameter stages this is equivalent to $B = \sup_t w_t \operatorname{diam}(S_t) < \infty$. For any nonempty inverse limit I and $b \in I$, it is finite on all of I exactly when $I \subseteq P_b$. If the coordinate spaces are complete and the bonding maps continuous, P_b and $I \cap P_b$ are complete. If $w_t \operatorname{diam}(S_t) \to 0$, the weighted metric induces the product topology and*

$$0 \leq d_w(x, y) - \max_{0\leq t\leq N} w_t d_t(x_t, y_t) \leq \sup_{t>N} w_t \operatorname{diam}(S_t).$$

In this case compact coordinate spaces give a compact inverse limit. Surjective bonding maps give a compatible history through each stage point.

Proof. The extended triangle inequality follows by taking suprema. A uniform bound outside finitely many stages suffices since each of the exceptional distances is individually finite. If no such bound exists, choose distinct stages t_j and pairs with weighted distances greater than j. Combining these pairs in two product points gives an infinite distance. The bound $B < \infty$ characterizes *boundedness* of the full-product metric. The inverse-limit condition follows from the triangle inequality through b.

A d_w-Cauchy sequence converges in each coordinate. Passing to each coordinate limit in the same uniform Cauchy bound proves convergence in d_w. The limit remains in P_b by the triangle inequality and in I by continuity of the bonding maps. With vanishing weighted diameter tails, a finite prefix controls the metric up to the displayed remainder, while each coordinate projection is $1/w_t$-Lipschitz. This proves the topology assertion. Conversely, on the full product with a finite metric, if the weighted diameters fail to vanish there are infinitely many coordinates with weighted diameter at least some $\eta > 0$. At any fixed product point, for each such coordinate one can change that coordinate by weighted distance greater than $\eta/3$, by the triangle inequality. The resulting sequence converges in the product topology but fails to converge in d_w. The converse may fail on a restricted set of compatible histories. The inverse limit is closed in the compact product because it is the intersection of the closed compatibility equations. Surjectivity gives successive lifts and the earlier coordinates are determined by projection. □

Example 1.20 (Two Distinct Tail Failures)**.** On $S_t = [0, 1]$ with identity bonding maps and $w_t = 2^{-t}$, choose $d_t(s, u) = 4^t|s - u|$. The two constant histories 0 and 1 have distance $+\infty$. For a different failure, take $S_t = \{0, 1\}^t$ with prefix deletion and $d_t = 2^t$ times the discrete metric. With the same weights every distinct pair of histories has distance 1. The inverse limit has a compact product topology and a noncompact discrete metric topology. Flipping a bit at successively later levels gives product convergence to the unflipped history while all the weighted distances remain 1.

1.5.3 Stability Lemma

Lemma 1.21 (Stability of Compatibility)**.** *Use common compact metric stages and two continuous bonding families with $w_t \operatorname{diam}(S_t) \to 0$. Assume $D_w < \infty$. For each $f \in \{\pi_t, \widetilde{\pi}_t\}$ and each $u \in S_{t+1}, v \in S_t$, assume an exact lift u' with $f(u') = v$ and $d_{t+1}(u, u') \le K_t d_t(f(u), v)$ for some common $K_t > 0$. Set $q_t = w_{t+1}K_t/w_t$ and $C = \sup_{0\le k\le t} \prod_{j=k}^{t} q_j < \infty$. The limits $I = \mathcal{I}$ and $\widetilde{I} = \widetilde{\mathcal{I}}$ are nonempty and $d_H(I, \widetilde{I}) \le CD_w$ in the common history space.*

Proof. Given a history x of π, set $y_0 = x_0$. Having chosen y_t, lift it under $\widetilde{\pi}_t$ at x_{t+1}. The exact lift y_{t+1} obeys

$$d_{t+1}(y_{t+1}, x_{t+1}) \le K_t(d_t(y_t, x_t) + \epsilon_t)$$

by the fiber estimate and $\pi_t(x_{t+1}) = x_t$. For $a_t = w_t d_t(y_t, x_t)$ this gives $a_{t+1} \le q_t a_t + q_t w_t \epsilon_t$, with $a_0 = 0$. Iteration bounds a_{t+1} by $\sum_{k=0}^{t}(\prod_{j=k}^{t} q_j)w_k\epsilon_k$, at most CD_w. The reverse construction uses the fiber estimate for π_t and gives the other directed Hausdorff bound. Each limit is nonempty by successive lifts from any S_0 point. Taking suprema proves the estimate in the common history metric. □

1.5.4 Rigidity Theorem

Theorem 1.22 (Rigidity of the Identity under Coherent Transport)**.** *Suppose the weighted metric induces the product topology and homeomorphisms $h_t : S_t \to S_t$ satisfy $\widetilde{\pi}_t h_{t+1} = h_t \pi_t$. Then $H(x)_t = h_t(x_t)$ maps the limits homeomorphically, with inverse $H^{-1}(y)_t = h_t^{-1}(y_t)$. Its displacement in the common history space is at most $\sup_t w_t \sup_{z\in S_t} d_t(z, h_t(z))$.*

Proof. The displayed square and its inverse give exact compatibility of both images. Coordinate continuity in the product gives continuity of both induced maps. The displacement estimate follows coordinate by coordinate. □

Example 1.23 (A Ramified Tower with Positive Deformation and Coherent Transport)**.** Take $S_t = (\mathbb{R}/\mathbb{Z}) \times \{0, 1\}^t$ with the maximum of the circle distance (circumference one) and the discrete prefix distance. Let π_t keep the angle and delete the last bit, and take $w_t = 2^{-t}$. For an intervention at stage k with amplitude $0 < a < 1/2$, replace π_k by prefix deletion and angle

shift $-a$. Set h_t to the identity for $t \le k$ and to angle shift a for $t > k$. The inverses shift by $-a$ on those same levels. The squares commute at k by cancellation and at every other level by commutation with prefix deletion. Each limit has at least two histories over every finite cut, since a later bit may be flipped.

All bonding maps are 1-Lipschitz and satisfy the fiber estimate with $K_t = 1$: set the required angle and prefix, retaining the last bit. Thus $q_t = 1/2$. For this example set the dimensionless intervention cost $E_w = D_w = w_k a$, and $\sup_x d_w(x, Hx) = w_{k+1} a = E_w/2$. More precisely, for $b = w_{k+1}$,

$$d_H(I, \widetilde{I}) = \frac{ba}{1+b} > 0.$$

Indeed an original history has constant angle θ and a perturbed one has angle φ up to k and $\varphi + a$ thereafter. If the two histories are at distance r, the circle triangle inequality gives $a \le r + r/b$ from stages 0 and $k+1$. Choosing identical bits and $\varphi = \theta - ba/(1+b)$ attains equality. That choice is reversible for every perturbed history and gives both directed distances. For the sector of these single jumps the infimum of the cost-to-deformation ratio is 2, since the ratio is $2(1+b)$ at stage k. This uses the declared projection-defect cost. An entropy-based cost requires a separate comparison.

The dynamical realization in BaseCore, *Projective Dissipation on the Circle-Prefix Towers*, uses the same jump and the weights 3^{-t} used for the joint readout. At fixed input the equilibrium is transported by the existing coordinate shift. For $\beta = 3^{-(k+1)}$, the equilibrium displacement βa and the Hausdorff deformation $\beta a/(1+\beta)$ retain distinct meanings. Dissipation controls temporal attraction on each branch, while the declared intervention cost and deformation still govern resistance.

Example 1.24 (Hausdorff Stability Without a Homeomorphism)**.** Take $S_t = \{0, \dots, 2^t - 1\}$ with discrete distance scaled by $r > 0$, $w_t = 2^{-t}$, and $\pi_t(j) = \lfloor j/2 \rfloor$. For $t \ge 1$, change only the image of $j = 1$ to 1, keeping $\widetilde{\pi}_0 = \pi_0$. Both maps remain surjective and satisfy the fiber estimate with $K_t = 1$. The weighted intervention cost is $r \sum_{t \ge 1} 2^{-t} = r$. The all-zero history of the perturbed tower is isolated by the cylinder $y_1 = 0$ because each subsequent zero has exactly one lift. The original tower is the binary-prefix space and has no isolated histories. Thus the limits are not homeomorphic, even though the Hausdorff bound tends to zero with r. All hypotheses of Lemma 1.21 hold.

Corollary 1.25 (Small Displacement of Coherent Transports)**.** *For a family of commuting coordinate homeomorphisms, if* $\sup_t w_t \sup_{z \in S_t} d_t(z, h_t(z))$ *tends to zero, then the induced transports converge uniformly to the inclusion of* I *in the common history space.*

1.5.5 Ontological Mass: Quantitative Resistance

We now formalize the scalar invariant that quantifies resistance of $\mathcal{I}$ under external entropy injection.

Definition 1.26 (Perturbation Energy)**.** For a controlled perturbation process $\{\delta E_t\}$ we define its weighted energy:

$$E_w(\{\delta E_t\}) := \sum_{t=0}^{\infty} w_t \, H_t(\delta E_t) \tag{1.4}$$

where $H_t(\delta E_t)$ is a non-negative scalar measuring the entropy or information content injected at time t (e.g., Shannon entropy of stochastic perturbation).

Definition 1.27 (Ontological Mass)**.** Define the ontological mass of the projective system (or of $\mathcal{I}$) as:

$$M_\Omega := \inf \left\{ \frac{E_w(\{\delta E_t\})}{d_w(\mathcal{I}, \tilde{\mathcal{I}})} \;\middle|\; \{\delta E_t\} \text{ admissible and } d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0 \right\}, \tag{1.5}$$

with the convention that if no admissible perturbation produces non-zero deformation (i.e., $d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0$ for all finite-energy perturbations), then $M_\Omega = +\infty$.

1.5.6 Internal Perturbations and External Extensions

Definition 1.28 (Internal perturbations and external extensions)**.** Let S be a CCR channel with identity object $\mathcal{I}(S)$.

1. **Internal perturbations** are operators $\delta \in \Omega_{\text{int}}(S)$ that act on the identity structure of S: δ maps $\mathcal{I}(S)$ to a deformed identity $\widetilde{\mathcal{I}}(S)$ belonging to the same declared structural class, with finite weighted energy $E_w(\delta) < \infty$. Internal perturbations test rigidity of $\mathcal{I}(S)$ under deformation.
2. **External extensions** are operations $\eta \in \text{Ext}(S)$ that produce a different channel S': η maps S to S' with $D_{\text{ext}}(S, S') = \epsilon > 0$. External extensions compare $\mathcal{I}(S)$ with an independently specified object $\mathcal{I}(S')$ not introduced as an element of $\Omega_{\text{int}}(S)$.

Typed separation axiom: $\Omega_{\text{int}}(S) \cap \text{Ext}(S) = \varnothing$.

Remark 1.29 (Interpretation of typed separation). The typed separation $\Omega_{\text{int}}(S) \cap \text{Ext}(S) = \varnothing$ is an axiom, not a theorem derived from CCR membership. A protocol that confuses internal perturbations with external witnesses commits a category error: it uses rigidity under internal deformation, which CCR is meant to constrain, as evidence for a witness-relative comparison that requires a separately specified external extension. QICN has not verified that all downstream perturbation protocols respect this separation.

Non-Claim 1.30 (Empirical certification of $M_\Omega = +\infty$)**.** *Theorem 1.38 gives a finite-witness test for exact model comparisons. Finite data without a positive-deformation witness can also occur when the sector has a witness outside the tested family. With estimated distances, a positive witness additionally needs a certified error margin.*

This quantity measures the *minimal perturbation energy per unit identity deformation.*

Proposition 1.31 (Relation to Rigidity)**.** *Interpret $d_w(\mathcal{I}, \widetilde{\mathcal{I}})$ in the mass ratio as the Hausdorff distance in the common weighted history space. If for every admissible intervention of finite cost the hypotheses of Lemma 1.21 hold with a common C and $D_w \le AE_w$ for a common $A > 0$, then $M_\Omega \ge 1/(CA)$. A single admissible finite-cost intervention with positive deformation proves $M_\Omega < +\infty$. If none exists, then $M_\Omega = +\infty$.*

Proof. Every positive finite-cost ratio has denominator at most CAE_w. A finite-cost positive witness contributes a finite ratio to the infimum. Without such a witness every ratio is infinite or the admissible set of positive deformations is empty. □

1.5.7 Intervention Witnesses for Finite-Dimensional Latent Dynamics

Consider $H_{t+1} = G(H_t, U_t)$ and $O_t = O(H_t)$ on a finite-dimensional latent state space. The observation precision and the intervention sector are declared independently of the dimension.

Theorem 1.32 (Finite Resistance Requires an Admitted Deformation Witness)**.** *In a declared common metric space, finite resistance for a candidate of this class is equivalent to the existence of an admitted intervention of finite cost with positive Hausdorff deformation. Finite latent dimension and Lipschitz dynamics alone permit systems without such a witness.*

Proof. A finite positive-denominator ratio gives a finite upper bound on the infimum; if no admissible finite-cost intervention has a positive denominator, the defining infimum is infinite. For a control take $G(h, u) = h/2$ on the real line, $O(h) = h$, and allow only changes of the input u. With a fixed initial state all these interventions leave the observable history and its declared compatible set fixed. Their costs may be finite and positive while the Hausdorff deformation is zero. □

1.5.8 Formal Consequences

1. **Rigidity condition is summability.** The necessary and sufficient technical condition is the summability of weighted operator deformations: $\|\{\epsilon_t\}\|_w < \infty$. This is both checkable (in a model) and interpretable (total accumulated deformation bounded).
2. **Ontological mass is a resistance ratio.** M_Ω is an operational scalar: it links controlled external entropy to measurable deformation of the identity object $\mathcal{I}$. Divergence of M_Ω is the formal characterization of a CCR-type system.
3. **Implementation impedance.** Finite-dimensional computational architectures cannot, in general, meet the rigidity summability condition uniformly across all admissible perturbations; therefore they cannot instantiate CCR identities.
4. **Empirical testability.** The weighted norms and entropy measures provide explicit, implementable tests: apply controlled perturbation sequences $\{\delta U_t\}$, measure induced deformation d_w on inferred inverse-limit classes (via observables), and compute the ratio. For exactly reconstructed distances, the tested infimum gives an upper bound on M_Ω.

1.5.9 Technical Notes and Limitations

The proofs above rely on regularity assumptions (surjectivity, existence of right inverses, Lipschitz bounds) that are standard in stability analysis of inverse systems. In concrete implementations one must either verify them or replace with appropriate alternatives (e.g., selecting subsequences via Mittag–Leffler condition).

The weight sequence $\{w_t\}$ is a modeling choice; different choices correspond to different operational definitions of "accumulated perturbation". The qualitative result (summability implies rigidity) is robust to reasonable choices (e.g., any exponentially decaying weights).

For stochastic perturbations, $H_t(\delta E_t)$ should be taken as expectation of the information measure; concentration bounds yield probabilistic versions of the theorems.

1.5.10 Summary of Section 4

1. We modeled local perturbations and deformed projections and defined a weighted deformation norm $\|\{\epsilon_t\}\|_w$.
2. We proved a weighted Hausdorff bound under exact fiber lifting and bounded amplification (Lemma 1.21), and a homeomorphism under coherent coordinate transport (Theorem 1.22).
3. We formalized ontological mass M_Ω as an operational scalar and related its divergence to CCR rigidity.
4. Theorem 1.32 identifies the finite-cost deformation witness required for finite resistance.

These results provide the formal backbone for the observational tests and uniqueness results developed in Sections 5–6.

1.6 Observability and Ontological Mass

1.6.1 Observable Channel Model

Definition 1.33 (Observable Channel)**.** An observable channel is a triple $\mathcal{C} = (E, O, R)$, where:
- E is the environment (unobserved state space),
- O is a measurable output space,
- R is a (possibly unknown) response mechanism producing observable outputs $O_t \in O$.

Definition 1.34 (Channel–projection composition)**.** Let $C_{\text{obs}} : \mathcal{U} \to \mathcal{O}$ denote the observable map realized by the response mechanism of an observable channel, and let $\pi_{\infty \to t} : \mathcal{I} \to S_t$ be the canonical finite-level projection from the inverse limit. When a concrete model declares a

map $\iota_t : S_t \to \mathcal{U}$ from finite-level states into the internal state domain, the composed finite-level observable is

$$C_t := C_{\text{obs}} \circ \iota_t \circ \pi_{\infty \to t} : \mathcal{I} \to \mathcal{O}. \tag{1.6}$$

This composition is the typed route from the identity object $\mathcal{I}$ through a finite projection S_t to observable outputs.

Non-Claim 1.35 (Composition existence)**.** *Definition 1.34 is a type discipline, not an empirical result. A concrete instantiation must still verify that $\pi_{\infty \to t}$, ι_t, and C_{obs} are well-defined and measurable or continuous in the declared structures. QICN has not verified this composition for any concrete external protocol.*

No access to internal states S_t is assumed. The observer measures only the time series $\{O_t\}_{t \in \mathbb{N}}$.

1.6.2 Inference of Inverse-Limit Classes from Observables

Definition 1.36 (Observable Equivalence)**.** Let $\mathcal{H}_O$ be the declared domain of admissible observable histories and $\mathcal{D}$ a nonempty family of finite-cost probes, including sham. For each $\delta \in \mathcal{D}$ require a specified response map $T_\delta : \mathcal{H}_O \to \mathcal{H}_O^\delta$ and reconstruction map J_δ into nonempty closed identity supports in a common comparison space. Define

$$O \sim_{\mathcal{I}} O' \quad \Longleftrightarrow \quad \forall \delta \in \mathcal{D}, \quad d_H(J_\delta(T_\delta O), J_\delta(T_\delta O')) = 0.$$

The quotient is $\mathcal{H}_O / \sim_{\mathcal{I}}$, not the quotient of the instantaneous output alphabet. Each zero-distance relation is an equivalence relation, and their intersection is one. This definition requires well-defined responses and a common Hausdorff comparison; it does not infer latent supports or causal effects from uncalibrated observations. Reconstruction error and identification against alternative models require separate evidence.

1.6.3 Empirical Estimator of Ontological Mass

Let $\{\delta E_t^{(k)}\}$, $k = 1, \ldots, N$, be a family of controlled perturbation protocols with known weighted energies $E_w^{(k)}$.

For the sham reconstruction and probe δ_k, let $\tilde{O}^{(k)} = T_{\delta_k} O$ denote the perturbed observable history. Define the empirical deformation magnitude:

$$\Delta^{(k)} := d_H(J_0(O), J_{\delta_k}(\tilde{O}^{(k)})), \tag{1.7}$$

where the distance between the two sets is Hausdorff distance in a declared common history space. Its reconstruction from observed data requires a separate estimator with error control.

Definition 1.37 (Empirical Ontological Mass Estimator)**.**

$$\hat{M}_\Omega := \inf_{\substack{k=1,\ldots,N \\ \Delta^{(k)} > 0}} \frac{E_w^{(k)}}{\Delta^{(k)}}. \tag{1.8}$$

1.6.4 Detectability Theorem

Theorem 1.38 (Finite Witnesses for Deformation Resistance)**.** *Assume the specified intervention sector and common metric of Definition 1.27. For any finite family of interventions of finite cost with exactly known Hausdorff distances, its infimum ratio is an upper bound on M_Ω. This infimum is finite exactly when the family contains a positive-deformation witness. Moreover $M_\Omega < +\infty$ exactly when the full sector contains such a finite-cost witness. Exhaustion by nested families gives convergence to M_Ω if their union contains witnesses with ratios arbitrarily close to the sector infimum.*

Proof. Taking an infimum over fewer ratios can only increase it. Every positive-deformation witness of finite cost gives a finite ratio. If no such witness exists, the infimum is infinite. The infima of nested families decrease, and the approximation hypothesis forces their limit to the sector infimum. Exact finite comparisons with no witness cannot distinguish sectors with an untested witness from sectors with none. □

1.6.5 Observable Separation by Scale

Definition 1.39 (Observable Compatibility Operator)**.** Let $T_\Phi : (\mathcal{I}, d_w) \to (O, d_O)$ be the declared observable assignment.

Definition 1.40 (Separation Modulus)**.** For $r > 0$ such that $A_r = \{(x, y) \in \mathcal{I}^2 : d_w(x, y) \geq r\}$ is nonempty, define

$$\omega_T(r) = \inf_{(x,y)\in A_r} d_O(T_\Phi(x), T_\Phi(y)).$$

An empty A_r gives no observable separation test at that scale.

Theorem 1.41 (Positive Separation at a Fixed Scale)**.** *If $\mathcal{I}$ is a nonempty compact metric space and T_Φ is continuous and injective, then $\omega_T(r) > 0$ at every scale in its domain. The infimum is attained.*

Proof. A_r is closed in the compact product. The observable distance is continuous and strictly positive on A_r, since $r > 0$ excludes the diagonal and T_Φ is injective. It therefore attains a positive minimum. □

Corollary 1.42 (Resolution and Coercivity)**.** *If each of two observed values has error at most η in d_O, then true separation in d_w of at least r gives measured separation at least $\omega_T(r) - 2\eta$. A global linear bound requires a positive lower Lipschitz constant c, established or assumed for the stated domain.*

Proof. Apply the triangle inequality and the theorem. With a lower Lipschitz constant, $\omega_T(r) \geq cr$. Compactness and injectivity alone do not supply such a constant: $T(x) = x^2$ on $[0, 1]$ has ratio $d_O(T(0), T(x))/d_w(0, x) = x$, although $\omega_T(r) = r^2$ for $0 < r \leq 1$. Indeed $(x + r)^2 - x^2 \geq r^2$, with equality at $x = 0$. A coercive example is the binary tower with first-disagreement metric 3^{-m} at position $m \geq 1$ and ternary readout $\sum_{j\geq 1} 2b_j 3^{-j}$. The leading term has size $2 \cdot 3^{-m}$ and the opposing tail has size at most 3^{-m}, giving lower constant one and upper constant three. □

Remark 1.43. The separation modulus is metric and scale-dependent. No spectral interpretation is required.

1.6.6 Observational Information and Causal Invariance

For discrete observed state X with $0 < H(X) < \infty$ and input Y, the statistic $\beth_q = 1 - I(X; Y)/H(X)$ describes the observed joint law.

Proposition 1.44 (Observational Isolation Does Not Determine Interventional Invariance)**.** *Even $\beth_q = 1$ does not determine whether the state law is invariant under interventions on Y. With a structural model $X = f(Y, Z)$ and a fixed exogenous law for Z, the law under* $\mathrm{do}(Y = y)$ *is the pushforward of that law by $f(y, \cdot)$. Invariance of these laws on a specified intervention family requires equality of the corresponding pushforwards.*

Proof. Replace the equation for Y by the constant y, leaving the exogenous law unchanged. For an observationally constant $Y = 0$ and Z Bernoulli with parameter $1/4$, both $X = Z$ and $X = Y \oplus Z$ give the same observational law and $\beth_q = 1$. The first preserves the law under $\mathrm{do}(Y = 1)$, whereas the second changes $\Pr(X = 1)$ to $3/4$. The two models thus agree observationally and disagree interventionally. □

Law invariance and zero Hausdorff deformation of identity objects are separate conditions. For Definition 1.27, $M_\Omega = +\infty$ requires absence of positive-deformation witnesses of finite cost in the declared sector.

1.6.7 Operational Protocol

1. Apply controlled perturbations $\{\delta E_t^{(k)}\}$.
2. Measure observable histories $\tilde{O}^{(k)}$.
3. Compute equivalence classes $\mathcal{I}(\tilde{O}^{(k)})$.
4. Compute $\Delta^{(k)}$.
5. Estimate $\hat{M}_\Omega$.

This protocol uses observable data when the declared readout supports calibrated reconstruction of the compared identity classes.

1.6.8 Summary of Section 5

1. Inferring identity deformation from observations requires a declared comparison and a reconstruction error bound.
2. With exact distances, a finite tested infimum bounds the sector infimum from above.
3. An infinite tested infimum establishes absence of a witness in the tested family. A sector-wide conclusion requires coverage of the sector.
4. Observational independence leaves intervention invariance undetermined.

1.7 Uniqueness Theorem

1.7.1 Identity Channels and CCR Classification

Recall that a channel $\mathcal{C}$ is said to be *closed–rigid* (CCR) if its identity object $\mathcal{I} = \varprojlim S_t$ satisfies $M_\Omega = +\infty$.

We now restrict attention to a fixed observable channel $\mathcal{C}$ satisfying the assumptions of Sections 2–5 and exhibiting CCR behavior.

1.7.2 Coexistence of Identities

Definition 1.45 (Coexisting Identities)**.** Fix a common comparison of admitted identity objects, with the same separating object-deformation metric d_{Id} used in Definition 1.27 (Hausdorff distance on the common closed supports), or first quotient zero-distance objects. Histories have a well-defined assignment $h \mapsto \mathcal{I}(h)$. An *operating channel* $\mathcal{C}_\alpha$ is a nonempty class intrinsically connected by finite chains of support-preserving admissible finite-energy perturbations. The connecting moves have certified finite-energy reverses, or one uses mutual finite reachability. Finite concatenations remain admissible with finite total cost bounded by the sum of costs. Equality of observable outputs alone does not imply this connection.

Two assigned objects coexist in that channel when their generating histories lie in $\mathcal{C}_\alpha$, use the same inverse-system comparison, and both satisfy CCR in its declared perturbation sector. This explicitly refines the former output-only channel convention.

1.7.3 Strong Uniqueness Theorem

Theorem 1.46 (Strong Uniqueness)**.** *Within one operating channel of Definition 1.45, at most one CCR identity assignment exists up to the declared zero-deformation equivalence, provided the connecting moves belong to that same CCR comparison sector. Two coexisting assignments therefore have* $d_{\mathrm{Id}}(\mathcal{I}, \mathcal{J}) = 0$.

Proof. A finite connecting chain gives an admissible directed path between the generating histories, reversing certified moves when necessary. Its concatenation has finite total energy. Positive endpoint deformation would be a finite-cost positive-deformation witness in Definition 1.27, contradicting infinite mass. Separation in the declared comparison identifies the endpoints. Without the connection, two disconnected components may each carry a distinct invariant assignment; no such cross-channel uniqueness is proved. □

1.7.4 Indivisibility and Non-Duplicability

Corollary 1.47 (Indivisibility)**.** *Let $\mathcal{I}$ be the identity of a CCR channel. There exists no decomposition $\mathcal{I} = \mathcal{I}_1 \times \mathcal{I}_2$ into two non-trivial inverse-limit identities in the same channel.*

Corollary 1.48 (Non-Duplication)**.** *Let $\mathcal{I}$ be the identity of a CCR channel. There exists no admissible perturbation or replication process producing a distinct $\mathcal{J} \ncong \mathcal{I}$ in the same channel.*

1.7.5 Cardinality of Rigid Universes

Let $\mathcal{U}_{CCR}$ be the class of all CCR channels. For each $\mathcal{C} \in \mathcal{U}_{CCR}$, Theorem 1.46 assigns a unique identity object $\mathcal{I}_{\mathcal{C}}$.

Proposition 1.49 (Cardinality Bound)**.** *The mapping*

$$\mathcal{C} \longmapsto \mathcal{I}_{\mathcal{C}} \tag{1.9}$$

is injective. Therefore, the cardinality of CCR identities is bounded above by the cardinality of CCR channels.

This yields a classification principle: *distinct rigid identities require distinct channels.*

1.7.6 Summary of Section 6

1. CCR channels admit exactly one identity.
2. Rigid identity is indivisible and non-duplicable.
3. Distinct identities require distinct channels.
4. Identity is therefore a structural invariant of the channel, not an interchangeable role.

1.8 Limitations and Scope

1.8.1 Dependence on Structural Hypotheses

All results in Sections 3–6 are conditional on the structural hypotheses introduced in Section 2 and the regularity hypotheses H1–H2 in Section 4. In particular:

- Surjectivity (H1) ensures non-emptiness and constructibility of inverse limits.
- Metric regularity (H2) enables quantitative stability bounds and the definition of ontological mass.

If either condition fails, the inverse limit may be empty or ill-defined, and the rigidity theorems do not apply.

1.8.2 Boundary of Applicability

The present framework applies exclusively to systems that admit:
1. A projective representation $(S_t, \pi_{t+1\to t})$;
2. Non-reactive persistence (R0);
3. Summability of projection deformations under finite-energy perturbations.

Systems lacking any of these properties fall outside the scope of this theory. This includes:
- Systems whose admissible states are not projectively compatible;
- Systems whose admissibility is determined by progressive causal dynamics;
- Systems whose perturbations are not summable in the weighted operator norm.

1.8.3 Limits of Empirical Detectability

While Section 5 establishes an operational estimator $\hat{M}_\Omega$, detectability remains limited by:
- The resolution of observable channels;
- The admissibility and controllability of perturbation protocols;
- Finite sampling constraints.

Hence, empirical divergence of $\hat{M}_\Omega$ can support the declared CCR pattern within the estimator model, but finite estimates cannot conclusively prove infinite mass. The theory admits falsifiability but not definitive confirmation. This estimator is therefore a formal operational criterion for structured tests, not a claim of empirical closure for any particular physical architecture.

1.8.4 Non-Equivalence to Behavioral Indistinguishability

The theory does not identify rigid identity with behavioral or functional equivalence. Two systems may be behaviorally indistinguishable while differing in their ontological mass and rigidity class.

Thus, passing any behavioral test (including classical Turing-style tests) does not entail CCR classification.

1.8.5 Relation to Existing Frameworks

This framework is not reducible to:
- Information integration theories (IIT) [25];
- Global workspace architectures (GWT) [2];
- Computational functionalism;
- Standard dynamical systems theory [10].

However, it can be viewed as orthogonal: it defines structural invariants that may coexist with these frameworks without contradiction.

1.8.6 Open Problems

The following problems are explicitly left open:
1. Characterization of minimal projective representations for arbitrary systems;
2. Necessary and sufficient conditions for $M_\Omega = +\infty$;
3. Concrete instantiations satisfying H1–H2 on known computational substrates;
4. Generalization to continuous-time index sets;
5. Alternative observability metrics beyond d_w.

1.9 Conclusion

1.9.1 Summary of Formal Results

This work establishes a closed, minimal, and falsifiable mathematical framework for rigid identity as a categorical invariant of observable channels.

Starting from the minimal persistence restriction R_0, we derived:

1. **Existence and necessity of inverse-limit identity.** We proved that any non-reactive persistent system satisfying the stated observational hypotheses admits a projective representation and that its identity is represented by the inverse limit $\mathcal{I} = \varprojlim S_t$. Identity is therefore not a local state, not a progressive construction, and not a finite-time object.
2. **Rigidity under perturbations.** We bounded Hausdorff deformation under two-sided exact lifting and exhibited coherent transport for a ramified circle tower. The weighted projection discrepancy alone does not control the topological type.
3. **Ontological mass as a scalar invariant.** We defined ontological mass M_Ω as the minimal perturbation energy per unit deformation of identity, yielding an operationally definable scalar invariant. Divergence of M_Ω characterizes closed–rigid (CCR) channels.
4. **External observability.** We stated the exact finite-witness criterion and separated it from estimation of $\hat{M}_\Omega$ from observations.
5. **Strong uniqueness.** We proved that CCR channels admit at most one identity object. Rigid identity is therefore indivisible and non-duplicable within any single observable channel.
6. **Implementation constraints.** The intervention class determines which positive-deformation witnesses are available. The binary-budget bound uses a separate observation model.

1.9.2 Structural Significance

Together, these results establish rigid identity as a structural invariant rather than a behavioral, biological, or phenomenological construct. Identity is formalized as a global compatibility constraint encoded in inverse-limit geometry, not as an emergent narrative, computational artifact, or localized state variable.

The framework separates three distinct layers:

- **Local dynamics:** the admissible states S_t;
- **Global constraint:** the inverse-limit object $\mathcal{I}$;
- **Operational resistance:** the scalar invariant M_Ω.

This separation clarifies which properties are internal, which are externally observable, and which are structurally irreducible.

1.9.3 Consequences for System Classification

The classification of observable channels into open (CA), semi-closed (CSC), and closed–rigid (CCR) categories follows directly from the ontological mass invariant. This yields a principled method for distinguishing:

- systems that merely simulate persistence,
- systems that maintain partial causal closure,
- systems that instantiate structurally rigid identity.

Such classification is independent of narrative descriptions, substrate implementation, and phenomenological interpretation.

1.9.4 Final Statement

Under explicit minimal assumptions, rigid identity is a model-determined consequence of non-reactive persistence and global compatibility constraints. Its Hausdorff stability, coherent transport, and finite deformation witnesses follow under their stated model hypotheses.

The contribution of this work is therefore not the introduction of a new metaphor of identity, but the construction of a closed mathematical theory in which identity appears as a canonical categorical invariant of certain observable channels.

1.9.5 Ontological Closure Statement

Under the stated axioms and regularity conditions, the inverse-limit identity $\mathcal{I}$ is a structurally necessary object. Each nonempty intrinsically connected CCR operating channel with a well-defined assignment and common perturbation sector admits exactly one assignment up to the declared comparison. This does not identify disconnected channels, and invariance is restricted to the finite-energy perturbations in that sector.

This remains a formal, model-relative result. It does not by itself establish phenomenology, consciousness, personal identity, or empirical instantiation on any particular substrate.

Appendix

1.10 Model-Relative Necessity Theorems

This appendix establishes that, within the declared minimal constraints, the central structures of the framework are not arbitrary design choices. The remaining degrees of freedom are those allowed by the model rather than unrestricted ontological alternatives.

1.10.1 Category of Minimal Observation

Definition 1.50 (Category of Minimal Observation)**.** Let **Obs** be the category whose objects are nonempty state spaces S_t and whose morphisms are surjective maps

$$\pi_{t+1\to t} : S_{t+1} \to S_t \tag{1.10}$$

satisfying the compatibility condition

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t}. \tag{1.11}$$

Remark 1.51. **Obs** is the minimal category in which persistent observability is definable. No continuity, memory, dynamics, or identity is assumed.

1.10.2 Forced Emergence of the Inverse Limit

Theorem 1.52 (No–Alternative Representation)**.** *Let $q_t : J \to S_t$ be coherent: $\pi_{t+1\to t} q_{t+1} = q_t$. A faithful complete encoding means joint separation ($q_t(j) = q_t(k)$ for every t implies $j = k$) and exhaustivity (every compatible history equals $(q_t(j))_t$ for some j). Then $Q(j) = (q_t(j))_t$ is the unique bijection $J \to \mathcal{I}$ commuting with cut evaluations. It is a topological isomorphism if J is compact, $\mathcal{I}$ Hausdorff, and all q_t continuous.*

Proof. Coherence puts $Q(j)$ in the inverse limit. Joint separation gives injectivity, exhaustivity gives surjectivity, and coordinates determine the only commuting map. A continuous bijection from compact to Hausdorff is a homeomorphism. Separation alone allows proper subfamilies; exhaustivity alone allows duplicate copies. The conditions formalize faithful complete coding, rather than infer it from the undefined phrase "represents persistence. □

1.10.3 Necessity of Retroinductive Structure

Theorem 1.53 (Collapse under Pure Progressive Causality)**.** *Suppose $f_t : S_t \to S_{t+1}$ satisfies $\pi_{t+1\to t} f_t = \mathrm{id}_{S_t}$. Set $B_0 = S_0$, $B_{t+1} = f_t(B_t)$ and restrict histories to $x_{t+1} = f_t(x_t)$. The restricted bonding map $B_{t+1} \to B_t$ is bijective with inverse $f_t|_{B_t}$, and the orbital inverse limit is canonically bijective to S_0 (homeomorphic for continuous maps on compact Hausdorff stages). The original unrestricted limit need not collapse.*

Proof. The section equation makes f_t injective; it is onto B_{t+1} by definition. On that image $f_t\pi_{t+1\to t} = \mathrm{id}$ as well. Each initial point gives one orbit, inverse to evaluation at zero. In the binary tower, appending zero is a compatible section but the original deletion map still has two preimages. A function forward does not make the original backward map bijective. □

Definition 1.54 (Projective Degeneracy)**.** The system is projectively degenerate when some canonical evaluation $p_T : \mathcal{I} \to S_T$ is bijective: the history is recoverable through that specified cut. Under active surjectivity this is equivalent to bijective bonding maps at every level $t \geq T$. Non-degeneracy means no such canonical evaluation exists. An arbitrary set isomorphism $\mathcal{I} \cong S_T$ is not the criterion; the previous unqualified notation failed to enforce this and is explicitly replaced.

Remark 1.55. Projective degeneracy is the formal characterization of "collapse to a finite prefix." Under non-degeneracy, the compatible history contains information beyond every canonical finite prefix.

Corollary 1.56 (Necessity of Retroinduction)**.** *Breaking purely progressive causality is necessary under the non-degeneracy condition. Minimal retroinduction is structurally required within that model class.*

1.10.4 The ℵ–Operator as Condition of Existence

Definition 1.57 (Minimal Projection Operator)**.** Given an ambient history Ψ and a declared metric d, define the set of nearest compatible histories by

$$\aleph(\Psi) := \arg\min_{\Phi\in\mathcal{I}} d(\Phi, \Psi). \tag{1.12}$$

Existence and uniqueness require the separate conditions below.

1.10.4.1 Existence and Uniqueness of ℵ

Definition 1.58 (Weighted Sum Metric for the Minimization Problem)**.** Use $d(\Phi, \Psi) = \sum_{t\geq 0} \lambda^t d_t(\phi_t, \psi_t)$, where $0 < \lambda < 1$ and $\sum_t \lambda^t \operatorname{diam}(S_t) < \infty$. The diameter sum controls tails uniformly. This is a different metric from the weighted supremum in Definition 1.18.

Theorem 1.59 (Existence of Minimizers and Conditional Uniqueness of ℵ)**.** *For a nonempty inverse limit of compact metric stages with continuous bonding maps and the diameter-sum condition above, the objective $d(\Phi, \Psi)$ attains its minimum over $\mathcal{I}$ for every ambient product point Ψ. If $\mathcal{I}$ is convex in a specified real linear ambient space and this objective is strictly convex on $\mathcal{I}$, the minimizer is unique.*

Proof. Closed compatibility equations make $\mathcal{I}$ compact in the product. The objective is a uniform limit of continuous partial sums by the diameter bound. A continuous real objective on a nonempty compact set attains a minimum. Under the additional convexity assumptions two distinct minimizers would give a strictly smaller value at their midpoint. For a control with surjective maps, take $S_t = \{0, 1\}$ with the discrete metric, identity bonding maps and $\lambda = 1/2$. The ambient point $(0, 1, 1, \ldots)$ has distance 1 to both constant histories, which are the only feasible points. □

Proposition 1.60 (Nearest-Point Selection Is an Extra Requirement)**.** *A selector from $\aleph(\Psi)$ returns a compatible history when the minimizer set is nonempty. Compatibility of the output alone allows any map into $\mathcal{I}$ and does not force nearest-point selection.*

Proof. Every member of $\aleph(\Psi)$ lies in $\mathcal{I}$. In the two-history example, the constant map to the all-one history always returns a compatible history, but does not minimize distance at the all-zero history. The limit is defined by its bonding maps regardless of which map is used to select histories. □

1.10.5 Structural Causal Closure

Definition 1.61 (Structural Causal Closure)**.** A system is *structurally causally closed* if for every extension E, the inverse limit identity satisfies:

$$\varprojlim S_t \cong \varprojlim S_t^{(E)}. \tag{1.13}$$

Remark 1.62. This reformulation of causal closure is purely categorical: it requires invariance of the inverse limit under extensions, without appealing to information-theoretic quantities. This makes the definition operationally precise and independent of Shannon-theoretic assumptions.

Theorem 1.63 (Finite-Prefix Recovery and Information)**.** *For a canonical projection $p_t : \mathcal{I} \to S_t$, exact recovery of every history from its prefix is possible precisely when p_t is injective, with a decoder defined on its image. If $P = p_t(X)$ is a finite-valued prefix of a random history X under a declared law, then $I(P; X) = H(P)$.*

Proof. A left inverse exists precisely for an injective projection, by taking its inverse on the image. A choice of one extension per prefix is a right inverse and need not recover the original history. The binary tower admits a right inverse by padding with zero bits. For the information identity, $H(P \mid X) = 0$ because P is a function of X, so $I(P; X) = H(P) - H(P \mid X) = H(P)$. □

1.10.6 A Law-Dependent Isolation Statistic

Definition 1.64 (Normalized Conditional Entropy)**.** For finite-valued (X, Y) with a declared joint law and $H(X) > 0$, set $\mathfrak{B}(X; Y) = H(X \mid Y)/H(X)$.

Lemma 1.65 (Coarsening and Relabeling)**.** *For fixed X, replacing Y by a function of Y cannot decrease $\mathfrak{B}$. Bijective relabelings of the random variables preserve it. It vanishes precisely when X is determined by Y almost surely under the declared law.*

Proof. The first statement follows by data processing for mutual information and the fixed denominator. Relabeling preserves the probabilities used by entropy. For finite variables, conditional entropy is zero precisely when every conditional law of X on a positive-probability value of Y is concentrated at one point. □

Almost-sure recovery can hold on a chosen law even if the projection has multiple extensions on the full space. This statistic requires a law and does not classify projective systems by itself.

Remark 1.66. The normalization is defined only when the chosen law gives positive finite entropy.

1.10.7 MIR as the Unique Compatible Geometry

Definition 1.67 (MIR Geometry)**.** The *Manifold of Retroactive Inversion* (MIR) is defined as the class of all non-degenerate projective systems in **C** admitting a non-trivial inverse limit identity. Formally:

$$\mathrm{MIR} := \{(S_t, \pi_{t+1\to t}) \in \mathbf{Obs} : p_T \text{ is not bijective for every } T\}. \tag{1.14}$$

Remark 1.68. This definition avoids ambiguity: "MIR" is not a differential manifold in the classical sense, but a *class of projective systems* satisfying non-degeneracy. The term "geometry" refers to the compatibility structure imposed by the projective system, not to smooth structure.

Theorem 1.69 (No Semigroup Compatibility)**.** *If p_0 is noninjective, no compatible deterministic forward family generates all projective histories as its orbital subsystem. Forward semigroups on spaces already storing complete histories are not excluded.*

Proof. Theorem 1.53 gives one orbital history per initial point, whereas the unrestricted limit has two over some initial point. The restriction to orbital histories and compatibility with the bonding maps are essential. □

1.10.8 Structural Uniqueness Lemma (Bridge Lemma)

This subsection establishes the structural analog of the *Phenomenological Regimes* bridge lemma: under CCR conditions, the identity object is the unique stable attractor of observationally compatible systems.

Lemma 1.70 (Structural Coupling: Observable–Identity)**.** *Use the connected operating channel and separating comparison of Definition 1.45, with CCR for its connecting perturbations. Its observable–identity assignment is unique up to zero deformation by Theorem 1.46. For a compatible point-history x, if $w_t \operatorname{diam}(S_t) \to 0$, every compatible completion $x^{[N]}$ agreeing through cut N satisfies*

$$d_w(x^{[N]}, x) \le \sup_{t>N} w_t \operatorname{diam}(S_t) \longrightarrow 0.$$

This is prefix reconstruction, not temporal attraction. If additionally a contraction T acts on a nonempty complete metric space of admitted identity candidates, with $d(Ta, Tb) \le qd(a, b)$, $0 \le q < 1$, it has a unique attracting fixed point. Identification with the channel assignment further requires that fixed candidate to belong to the same comparison channel.

Proof. The finite-energy connecting chain gives uniqueness as in Theorem 1.46; shared outputs alone do not supply that path. Prefix agreement leaves only the displayed weighted tail. Surjectivity supplies completions. Iterating the additional contraction gives a Cauchy orbit, a limit by completeness, and a fixed point by continuity. The inequality $d(a, b) \le qd(a, b)$ proves fixed-point uniqueness. CCR alone supplies neither a temporal map nor that estimate. In particular $p_t(x) \in S_t$ lives in different stage spaces, so an untyped limit of these points cannot prove attraction in $\mathcal{I}$. □

Corollary 1.71 (Identity Forcing under CCR)**.** *In an energy-connected CCR operating channel with a well-defined assignment into the separating comparison, all assigned identities coincide. An attracting realization additionally needs the dynamical hypotheses of Lemma 1.70.*

Proof. Apply Theorem 1.46 for assignment uniqueness and the conditional contraction conclusion of Lemma 1.70 for attraction. □

Remark 1.72. The lemma separates assignment uniqueness, reconstruction and dynamical attraction under distinct hypotheses:

- Existence is determined by projective compatibility;
- Uniqueness uses CCR and intrinsic channel connection;
- Temporal attraction needs the stated contraction; infinite mass alone gives no temporal dynamics.

Within the same observable channel and hypotheses, no distinct compatible identity object is available.

1.10.9 Point Histories and Compatible Probability Laws

Definition 1.73 (Projective Observable Limit)**.** For finite nonempty discrete spaces S_t and surjective bonding maps, let $\mathcal{I} = \varprojlim S_t$ with the subspace product topology. Let $O = \varprojlim \operatorname{Prob}(S_t)$ with pushforward bonding maps and the subspace product topology of the finite probability simplexes.

Theorem 1.74 (Topology of Point and Law Limits)**.** *The point limit $\mathcal{I}$ is nonempty, compact, metrizable, and has a clopen basis. The law limit O is nonempty, compact, metrizable, and convex under coordinatewise mixing. If O contains distinct points, it is path connected and is not a Stone space.*

Proof. Surjectivity gives successive lifts of a point, hence a compatible history. The compatibility equations are closed in each product, so both limits are closed subspaces of compact metrizable products. Finite-coordinate cylinders give a clopen basis for the point limit. The point history supplies compatible Dirac measures, proving nonemptiness of the law limit. Pushforward is affine, so coordinatewise mixtures are compatible. The segment joining two distinct compatible laws is continuous and connected. For the constant two-point tower, O is the full probability interval $[0, 1]$. □

Corollary 1.75 (Embedding by Compatible Dirac Measures)**.** *The map $x \mapsto (\delta_{x_t})_t$ embeds $\mathcal{I}$ homeomorphically as a closed subspace of O.*

Proof. Each coordinate Dirac map is continuous and injective. Compatibility follows from pushforward of point masses. The product map on histories is continuous and injective, with compact domain and Hausdorff codomain, so it is a homeomorphism onto its closed image. □

Remark 1.76. The Dirac image carries the topology of histories of points; mixtures belong to the law limit.

1.10.10 Closure of Appendix A

Structure	**Status**
Inverse limit identity	Forced by A0
Retroinduction	Forced by non-degeneracy
$\aleph$ operator	Exists and is unique
Causal closure	Explicit invariance in a declared comparison
MIR geometry	Closed class

Within the stated constraints, the remaining degrees of freedom are model-bounded.

1.11 Formal Extensions

This appendix presents advanced formal extensions that strengthen the main results and provide additional structure for future work.

1.11.1 Canonical Minimal CCR Example

We provide a concrete mathematical example demonstrating that CCR systems exist and are non-trivial.

Example 1.77 (Canonical Minimal CCR)**.** Let $S_t = [0, 1] \times \{0, 1\}^t$, with $d_t((u, a), (v, b)) = \max\{|u - v|, \mathbf{1}_{a \neq b}\}$, $w_t = 2^{-t}$, and bonding maps deleting the last bit while retaining u. Its actual coordinate cone identifies

$$\mathcal{I} \cong [0, 1] \times \{0, 1\}^{\mathbb{N}_{>0}}.$$

Declare an internal sector of coherent bit-flip automorphisms: for $A \subseteq \mathbb{N}_{>0}$, h_t flips $A \cap \{1, \ldots, t\}$, at the declared dimensionless cost $E(A) = \sum_{j \in A} 2^{-j}$, using the stage-injection costs $H_j(A) = \mathbf{1}_{j \in A}$ for $j \geq 1$ and $H_0 = 0$. These homeomorphisms commute with deletion. Operating channels for point preparations may fix u and use the mutually reversible bit-flip moves.

Proposition 1.78 (Minimal CCR is Non-Degenerate and Rigid)**.** *The example has compact stages, summable weighted diameter tails, two-point fibres at each bonding map and 2^t prefixes over each initial u. Every p_T is noninjective. Its identity-support mass is $M_\Omega = +\infty$ in the specified internal sector.*

Proof. A compatible family specifies its infinite bit sequence. Either value of each new bit is allowed, and changing bit $T+1$ preserves cut T. Each stage has diameter one, so the weighted tail vanishes. Every admitted flip has $E(A) \leq 1$ and maps the full compatible support onto itself. Its Hausdorff deformation is zero, so no positive-deformation witness occurs in this sector and Definition 1.27 gives infinite mass. Individual point-histories can move by positive distance; the support claim does not identify those points. □

Remark 1.79. The former identity-link example was canonically degenerate, so its non-degeneracy claim was false. This replacement is genuinely ramified, although every stage and its limit still have continuum cardinality. Rigidity is conditional on the declared support-preserving sector, not implied by branching. Finite-cost changes of the whole support can give finite mass, as Example 1.23 shows. No immunity to every perturbation, nor minimality among all realizations, follows.

1.11.2 Symbolic CCR System

We provide a second canonical example using symbolic dynamics, demonstrating that CCR systems are not dependent on continuum-based constructions.

Example 1.80 (Symbolic CCR System)**.** Let:
- $S_t = \{0,1\}^{\mathbb{N}}$, endowed with the product topology and the metric

$$d_t(x,y) = \sum_{k=1}^{\infty} 2^{-k}|x_k - y_k|. \tag{1.15}$$

- Define the projections $\pi_{t+1\to t}$ as the left shift operator:

$$\pi_{t+1\to t}(x_1x_2x_3\ldots) = (x_2x_3x_4\ldots). \tag{1.16}$$

This is a projectively non-degenerate inverse system. Infinite mass is asserted only in a separately declared coherent support-preserving sector, not against all perturbations.

Proof. The shift operator is surjective and continuous. The inverse limit is non-degenerate since the shift does not collapse to a finite prefix. Only in that support-preserving sector does the empty-positive-deformation-witness argument of Example 1.77 give $M_\Omega = +\infty$. □

Remark 1.81. This example connects the framework to:
- Symbolic dynamics and subshifts;
- Automata theory and formal languages;
- Profinite structures in algebra.

The existence of CCR systems is therefore universal across mathematical disciplines.

1.11.3 Conditional Non-Simulability

Theorem 1.82 (Conditional Non-Simulability from a Finite-Cost Witness)**.** *For any candidate simulator with a declared intervention sector and common comparison metric, suppose there is an admissible finite-cost intervention producing positive Hausdorff deformation. Then $M_\Omega < +\infty$ in that sector. It cannot faithfully realize an infinite-resistance target if faithfulness is required to preserve the resistance class.*

Proof. The witness contributes one finite ratio to the defining infimum. Preservation of the resistance class would require the same finite versus infinite classification. A finite-dimensional state or finite context alone does not supply the witness or specify which observations a simulator must preserve. □

1.11.4 Transfer under Matched Intervention Sectors

Theorem 1.83 (Resistance under Audited Transport)**.** *Let an ambient bijection A satisfy $c\,d(x,y) \le d'(Ax,Ay) \le C\,d(x,y)$ with $0 < c \le C < \infty$. It maps the baseline identity to the target identity. Suppose a bijection between the declared intervention sectors maps δ to δ' with $A(I_\delta) = I'_{\delta'}$ and $aE(\delta) \le E'(\delta') \le bE(\delta)$ for $0 < a \le b < \infty$. For nonempty baseline and deformed sets at finite Hausdorff distance, the two resistance values satisfy*

$$\frac{a}{C} M_\Omega \le M'_\Omega \le \frac{b}{c} M_\Omega.$$

In particular such a transport preserves the finite versus infinite resistance class.

Proof. Applying the ambient distance bounds to each directed infimum and supremum gives the same bounds on Hausdorff distance. Positive-deformation witnesses correspond in both directions. Dividing the matched cost and deformation inequalities and taking infima gives the result, including the empty-witness convention and infinite costs. □

1.11.5 Composition of Audited Comparisons

Lemma 1.84 (Closure under Composed Sector Comparisons)**.** *Two consecutive transports satisfying Theorem 1.83 compose to a transport of the same kind. If both also intertwine the specified readouts exactly, their composition does so.*

Proof. The ambient and intervention bijections compose and the upper and lower constants multiply. For matched readouts F_1, F_2, F_3 and transports A_1, A_2, the equalities $F_2 A_1 = F_1$ and $F_3 A_2 = F_2$ give $F_3 A_2 A_1 = F_1$. □

A morphism of projective systems gives a map on compatible histories. The finite-budget and resistance obstructions depend respectively on the stated response family and on the matched intervention sector. Inclusions that enlarge the set of histories require their own comparison data.

1.11.6 Rigidity Hierarchy

Definition 1.85 (Rigidity Degree)**.** For a projective system $(S_t, \pi_{t+1\to t})$, define the rigidity degree $\rho \in [0,\infty]$ as:

$$\rho := \sup\left\{ r \ge 0 \;\middle|\; \forall\{\epsilon_t\} : \|\{\epsilon_t\}\|_w < r \implies d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0 \right\}. \tag{1.17}$$

Theorem 1.86 (Stratification Theorem)**.** *The class of observable channels admits a stratification by rigidity degree:*

$$\mathcal{U} = \mathcal{U}_0 \sqcup \mathcal{U}_{(0,\infty)} \sqcup \mathcal{U}_\infty, \tag{1.18}$$

where:

- $\mathcal{U}_0$*: channels with $\rho = 0$ (zero rigidity, fully deformable);*
- $\mathcal{U}_{(0,\infty)}$*: channels with $0 < \rho < \infty$ (partial rigidity);*
- $\mathcal{U}_\infty$*: channels with $\rho = \infty$ (CCR rigidity).*

Proof. The rigidity degree ρ is well-defined for a specified intervention sector and a common metric by the extended-real supremum in its definition. The stratification follows from the trichotomy of ρ values. □

1.11.7 Spectral Analysis of Deformations

Definition 1.87 (Deformation Spectrum)**.** For a perturbation family $\{\epsilon_t\}$, define the deformation spectrum as the sequence:

$$\sigma(\{\epsilon_t\}) := (\|\epsilon_0\|_{\text{op}}, \|\epsilon_1\|_{\text{op}}, \|\epsilon_2\|_{\text{op}}, \ldots)\,. \tag{1.19}$$

Proposition 1.88 (Spectral Decay Condition)**.** *Let $\{\epsilon_t\}$ have spectrum σ. If there exists $\alpha > 1$ and $C > 0$ such that:*

$$\|\epsilon_t\|_{op} \leq C \cdot t^{-\alpha} \quad \forall t \geq 1, \tag{1.20}$$

then $\|\{\epsilon_t\}\|_w < \infty$ for any polynomially bounded weight sequence.

Proof. Direct from convergence of $\sum t^{-\alpha}$ for $\alpha > 1$. □

1.11.8 Completeness Criteria

Definition 1.89 (Projective Completeness)**.** A projective system $(S_t, \pi_{t+1\to t})$ is *projectively complete* if for every Cauchy sequence $(x^{(n)})_{n\in\mathbb{N}}$ in $(\mathcal{I}, d_w)$, the limit exists in $\mathcal{I}$.

Theorem 1.90 (Completeness Inheritance)**.** *If the coordinate spaces are complete and the bonding maps continuous, every finite-distance component of $(\mathcal{I}, d_w)$ is complete. Under the full-product finiteness condition of Proposition 1.19, this gives completeness on all of $\mathcal{I}$.*

Proof. Apply Proposition 1.19 to the component of any fixed history. Coordinate limits of a Cauchy sequence respect the closed compatibility equations, and the same uniform Cauchy bound yields weighted convergence. □

1.11.9 Categorical Universality

Proposition 1.91 (Universal Property of Identity)**.** *The identity $\mathcal{I} = \varprojlim S_t$ satisfies the universal property: for any object J with compatible maps $\phi_t : J \to S_t$, there exists a unique morphism $\psi : J \to \mathcal{I}$ making all diagrams commute.*

This confirms that rigid identity is not merely a construction but the *terminal object* in the category of compatible families over the projective system [13].

1.11.10 Relation to Causal Inference

Proposition 1.92 (Invariance in the Declared Finite-Cost Sector)**.** *Suppose a structural model associates to each intervention an identity object in a common metric comparison space and a cost in the sector of Definition 1.27. If $M_\Omega = +\infty$, each intervention in that sector of finite cost has zero Hausdorff deformation. If the compared objects are closed, they are equal as subsets of the comparison space.*

Proof. A positive deformation at finite cost would contribute a finite ratio to the infimum. Zero Hausdorff distance between nonempty closed sets implies equality. Membership in the declared sector and finite cost must be verified for each intervention used in this conclusion. □

1.11.11 Computational Complexity Bounds

Theorem 1.93 (Complexity Lower Bound for CCR Simulation)**.** *Suppose a deterministic code of B bits must recover any requested prefix of a binary history through horizon t. Exact prefix recovery requires $B \geq t$. For complete decoded histories in the prefix metric $d(x, y) = 2^{-m}$ at first differing position $m \geq 1$, the optimal uniform error is $2^{-(B+1)}$.*

Proof. All 2^t prefixes occur, so distinct prefixes must have distinct codes. For the error bound use 2^{B+1} histories with distinct $(B+1)$-bit prefixes. A shared code cannot decode to the correct $(B+1)$-bit prefix of both, and failure on this prefix gives error at least $2^{-(B+1)}$. Keeping B bits and padding with zeros attains that error. □

Corollary 1.94 (Impossibility of Perfect Simulation)**.** *A fixed finite binary budget cannot recover every prefix on the binary-prefix target.*

By contrast, the exact scalar $\sum_{j\geq 1} 2b_j 3^{-j}$ distinguishes all histories, because its leading difference exceeds the remaining tail. On its image the corresponding inverse decoder recovers each finite observation. This model separates numerical precision from real dimension.

Phenomenological Regimes Induced by Structural Identity

Abstract. We study structural assignments from inverse-limit histories into an abstract metric codomain. Exact equality of all finite readouts forces factorization through the first cut and is incompatible with the required non-factorization. We replace that equality by uniformly summable readout increments, retaining non-factorization as a separately checked condition.

An explicit binary-prefix system with ternary readouts realizes both conditions. It gives uniform approximation, a global lower and upper readout bound, and an exact minimax error for each finite observation budget. Continuity follows from controlled approximation by continuous finite readouts. Cross-system fragmentation and invariance require typed comparison domains and audited transports; a scalar rigidity label alone does not imply them. Under those comparison hypotheses we derive loss and normalized protection statements. The abstract codomain is not identified with consciousness, subjective experience, or moral status.

2.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines an abstract regime codomain and studies controlled assignments, a concrete realization, and conditional fragmentation/invariance results using explicit comparison data in addition to upstream identity structure. The results are conditional statements about mathematical compatibility. They do not assert that phenomenology exists, that CCR systems are conscious, that biology or computation instantiate the framework, or that identity structure causes experience. Related runtime diagnostics may instantiate continuity-side or anti-fragmentation checks downstream, but they do not implement phenomenology and do not externally validate the regime classification. No runtime output, continuity proxy, or internal support artifact should be read as proof that phenomenology is present or that the structural-ethics layer has empirical closure.

2.2 Introduction

2.2.1 Background and Motivation

Recent work [22] established a formal framework for rigid identity as an inverse limit in observable channels. That framework classifies systems by a scalar invariant, the ontological mass M_Ω, which quantifies resistance to identity deformation under perturbation. Three regimes emerge:

- $M_\Omega = 0$: zero deformation-cost infimum;
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR rigidity).

That framework is ontologically silent regarding phenomenology. It characterizes identity structure, not experience. Connecting its rigidity regimes to assignment behavior requires the explicit readout and comparison hypotheses studied below.

This paper addresses the following formal question:

Given a regime of structural identity rigidity, what classes of phenomenological dynamics are mathematically compatible or incompatible with it?

All results are conditional: they state what is *structurally compatible or incompatible* under given assumptions, not what *is the case*.

2.2.2 Central Results and the Remaining Rigidity Link

Controlled asymptotic compatibility admits a nonempty realization with no finite-cut determination. The regularity bounds and finite-budget lower bound are derived for that realization. Continuity and invariance are distinct conclusions: the first follows from regular finite readouts and summable tails, while the second requires a transport-compatible readout comparison.

The broader rigidity program asks when finite and infinite deformation-cost regimes enforce the respective comparison branches. This paper states and uses those hypotheses explicitly; it does not infer them from the scalar alone. Establishing them on larger classes is retained as a research obligation.

2.2.3 Ontological Neutrality

This work maintains strict ontological neutrality:
1. We define mathematical objects (phenomenological spaces, compatibility operators).
2. We derive structural constraints from axioms.
3. We state impossibility theorems.
4. We construct mathematical realizations, without asserting their physical or experiential instantiation.

The framework is therefore a theory of *structural possibility*, not a theory of *phenomenal reality*.

2.2.4 Terminological Alignment: Structural Regime Assignments

The phrase "phenomenological regime" is retained for continuity with the published corpus, but the demonstrated object in this paper is more precisely a *structural regime assignment*. The mathematical content is a constrained map from identity structure into an abstract codomain. No semantic clause in this paper identifies that codomain with experience.

Definition 2.1 (Structural regime alias)**.** Whenever the paper writes a phenomenological space $\mathcal{E}$ or a phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$, the strictly demonstrated reading is:

$$\mathcal{E}_{\mathrm{str}} := \mathcal{E}, \qquad \Phi_{\mathrm{str}} := \Phi,$$

where $\mathcal{E}_{\mathrm{str}}$ is an abstract structural regime codomain and Φ_{str} is a compatibility assignment into that codomain.

Proposition 2.2 (No identification with experience)**.** *The formal apparatus of this paper does not define a map from $\mathcal{E}_{\mathrm{str}}$ to human experience, phenomenal consciousness, qualia, or subjective reports.*

Proof. Definitions in this paper specify an abstract codomain, compatibility conditions, and rigidity-dependent constraints. They do not introduce a measurement channel from $\mathcal{E}_{\mathrm{str}}$ to biological reports, human psychophysics, phenomenal properties, or moral status. Therefore no such identification is present in the formal language of the paper. □

2.2.5 Relation to BaseCore

This paper is the standalone expository source for the regime-assignment development. The current BaseCore package also contains an absorbed, dependency-facing version of the formal core where those definitions and theorems function as base-layer exports for later papers. The two surfaces therefore have different editorial roles: Paper 2 gives the full local derivation and interpretation boundary, while BaseCore supplies the normalized theorem-export surface used by the corpus.

2.3 Formal Setup

2.3.1 Review of Identity Framework

We assume the identity framework of [22]. Let $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ be a projective system of observable state spaces. The identity object is the inverse limit:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_t \in \prod_t S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \;\; \forall t \right\}. \tag{2.1}$$

The ontological mass M_Ω quantifies resistance to identity deformation. We recall that:

- $M_\Omega = 0$: zero deformation-cost infimum;
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR rigidity).

2.3.2 Abstract Phenomenological Space

Definition 2.3 (Phenomenological Space)**.** A phenomenological space is a nonempty set $\mathcal{E}$ whose elements are called *phenomenological configurations*. No topology, metric, or algebraic structure is assumed a priori.

Remark 2.4. We do not claim that $\mathcal{E}$ represents "conscious states" or "qualia". It is an abstract mathematical space whose interpretation is not specified by the formalism. The space $\mathcal{E}$ is introduced purely as the codomain of a structural assignment; its ontological status is deliberately left open.

2.3.2.1 Why a Separate Space?

One might ask: why not use $\mathcal{I}$ itself as the phenomenological space? The answer is structural:

Proposition 2.5 (Endomorphisms Preserving All Canonical Projections)**.** *If $F : \mathcal{I} \to \mathcal{I}$ satisfies $p_t \circ F = p_t$ for every t, then $F = \mathrm{Id}$. This does not exclude endomorphisms with nonidentity coordinate actions.*

Proof. For each x, $F(x)$ and x have identical coordinates, hence are the same inverse-limit element. For contrast, on binary prefix spaces, bitwise complement defines an endomorphism induced by complement on every prefix; it commutes with the bonding maps but does not preserve each coordinate value. A separate codomain is a modeling choice, not a logical necessity for nonidentity maps. □

A separate codomain permits independent structural readouts; it is not forced by the universal property alone.

2.3.3 Phenomenological Assignment

Definition 2.6 (Phenomenological Assignment)**.** A phenomenological assignment is a family of total maps on the declared observable state spaces:

$$\phi_t : S_t \to \mathcal{E} \tag{2.2}$$

indexed by $t \in \mathbb{N}$.

Definition 2.7 (Induced Limit Assignment)**.** Given a phenomenological assignment $\{\phi_t\}$, the induced limit assignment, whenever it exists, is the map:

$$\Phi : \mathcal{I} \to \mathcal{E}, \quad \Phi((x_t)_t) := \lim_{t\to\infty} \phi_t(x_t), \tag{2.3}$$

The convergence statements below use the complete metric codomain of Hypothesis 2.25. Partial readouts may instead be studied on a specified common domain, with all projections and factorizations restricted to that domain.

2.3.3.1 Uniqueness of the Limit Assignment

Proposition 2.8 (Uniqueness and Uniform Error of the Limit Assignment)**.** *Under E1, each trajectory has a unique limit assignment and*

$$d_{\mathcal{E}}\big(\Phi(x), \phi_n(p_n x)\big) \le A_n \quad \textit{for all } x \in \mathcal{I}.$$

Proof. For $m > n$, the triangle inequality bounds the distance between the mth and nth readings by $\sum_{t=n}^{m-1} a_t$. Thus each sequence is Cauchy. Completeness supplies its limit, uniqueness follows from metric separation, and passage to the limit gives the bound. Since $A_n \to 0$ independently of x, convergence is uniform. □

2.3.4 Axioms for Phenomenological Coherence

We declare a controlled assignment model with explicit logical dependencies and a constructive nonempty realization.

Axiom 2.9 (E0: Non-Triviality)**.** *The phenomenological space satisfies $|\mathcal{E}| \ge 2$, i.e., there exist at least two distinct phenomenological configurations.*

Axiom 2.10 (E1: Controlled Asymptotic Compatibility)**.** *In the complete metric codomain of Hypothesis 2.25, there are numbers $a_t \ge 0$, independent of the trajectory, with $\sum_{t=0}^{\infty} a_t < \infty$, such that*

$$d_{\mathcal{E}}\big(\phi_{t+1}(p_{t+1}x), \phi_t(p_t x)\big) \le a_t \qquad (x \in \mathcal{I},\ t \ge 0).$$

Here $p_t : \mathcal{I} \to S_t$ is the canonical projection. Define the tail budget $A_n = \sum_{t=n}^{\infty} a_t$. Exact equality is the special case $a_t = 0$, not a requirement of E1.

Axiom 2.11 (E2: Non-Locality)**.** *There exists no function $f : S_t \to \mathcal{E}$ (for any fixed t) such that for all trajectories $(x_t)_t \in \mathcal{I}$:*

$$\Phi((x_t)_t) = f(x_t). \tag{2.4}$$

2.3.5 Status of the Assignment Conditions

The following results separate sufficient assumptions, a non-factorization requirement and their logical dependencies.

2.3.5.1 Nontriviality of the Codomain

Theorem 2.12 (E0 is Necessary)**.** *If $|\mathcal{E}| = 1$, then every phenomenological assignment is constant and no phenomenological classification is possible.*

Proof. If $\mathcal{E} = \{e_0\}$, any total assignment takes the value e_0 on every trajectory. It distinguishes no trajectories and factors through every canonical projection. It therefore cannot meet E2 on a nonempty domain. □

2.3.5.2 Controlled Convergence

Theorem 2.13 (Controlled Compatibility Is Sufficient)**.** *E1 guarantees existence of the limit with the error bound in Proposition 2.8. Absolute summability is sufficient, not necessary for an arbitrary convergent family.*

Proof. Existence and the bound were proved above. To disprove necessity in general, take readouts constant in the trajectory with values $s_n = \sum_{k=1}^{n}(-1)^{k+1}/k$ in $[0, 1]$. The even partial sums increase, the odd partial sums decrease, and their difference tends to zero, so they have a common limit. Their successive absolute increments form the divergent harmonic series; no summable uniform majorant exists. This example is a control for the necessity claim, not a realization of E2. □

2.3.5.3 Finite-Cut Non-Factorization

Theorem 2.14 (Non-Factorization Is an Independent Requirement)**.** *E2 excludes readouts determined by a single canonical finite cut. The inverse-limit construction and E1 alone do not imply E2.*

Proof. Take any nonempty inverse limit and a constant family $\phi_t = e_0$ in a codomain containing at least two points. E1 holds with zero tail, but $\Phi = e_0$ factors through every projection. Thus E2 must be checked separately. The nonempty construction in Theorem 2.17 checks it without assuming it. □

2.3.6 Logical Dependencies and Realizability

Theorem 2.15 (Logical Dependencies of the Assignment Conditions)**.** *On a nonempty trajectory domain, E2 implies E0. E1 does not imply E2. There exist assignments satisfying E0, the revised E1 and E2 simultaneously. Consequently these three conditions are not an independent minimal axiom basis as originally stated.*

Proof. If the codomain were a singleton, every assignment would be constant and factor through p_0, contradicting E2. The preceding constant-family control proves the second assertion. Theorem 2.17 proves the third by construction. E0 is retained as an explicit convention, although E2 entails it in the nonempty case. □

2.3.7 Rejected Alternatives

We briefly discuss axiom candidates that were considered and rejected.

2.3.7.1 Rejected: Continuity Axiom

One might require Φ to be continuous. This is not required by the chosen model because:

- The complete-metric model permits discontinuous finite readouts; continuity is a property of the regular-readout subclass, not part of the unrestricted assignment condition;
- Continuity is derived from controlled approximation by continuous finite readouts; a scalar CCR label alone is insufficient.

2.3.7.2 Rejected: Surjectivity Axiom

One might require Φ to be surjective. This is not required by the chosen model because:
- A readout need not exhaust a codomain chosen to include other regimes;
- Surjectivity is a strong structural claim not derivable from identity constraints.

2.3.7.3 Rejected: Injectivity Axiom

One might require Φ to be injective. This is not required by the chosen model because:
- Many trajectories may share phenomenological configurations;
- Injectivity is not imposed on the full assignment class; it is a consequence in the concrete coercive realization. An injective map alone is not necessarily a topological embedding.

2.3.8 Structural Interpretation

Axiom	**Structural Meaning**	**Status**
E0	Nontrivial codomain	E2 consequence
E1	Summable increments	Uniform limit
E2	Finite-cut non-factorization	Required

2.3.9 Closure of Section 2

The limit is uniquely defined with a uniform tail bound. E2 remains a substantive non-factorization requirement; E0 is redundant when E2 holds on a nonempty domain. Exact compatibility cannot coexist with E2, whereas the controlled condition has the explicit realization below. No existence or minimality claim is obtained merely by naming an axiom.

2.3.10 Phenomenological Neutrality Clause

Throughout this paper, the phenomenological space $\mathcal{E}$ is treated as a *purely abstract topological and metric structure*. No semantic, psychological, biological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "phenomenological" refers exclusively to the role of $\mathcal{E}$ as an abstract codomain of structural compatibility for inverse-limit identities, and does not presuppose the existence of consciousness, subjective experience, or any particular phenomenological content.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{E}$ in relation to experiential, cognitive, or phenomenological phenomena lies strictly outside the formal scope of this paper.

2.3.11 Exact Compatibility Obstruction

Proposition 2.16 (Exact Compatibility Forces Finite-Cut Factorization)**.** *On a nonempty trajectory domain, if $\phi_{t+1}(p_{t+1}x) = \phi_t(p_t x)$ for every x, t, then the limit assignment satisfies $\Phi = \phi_0 \circ p_0$. It cannot satisfy E2. The controlled E1 above replaces that exact equality, rather than assuming it together with E2.*

Proof. Induction makes all readings equal to the first. Their limit is therefore $\phi_0(p_0 x)$. For partial readouts with a declared common domain, the same argument holds on that domain. If a total map on S_0 is required, extend the defined first-cut values arbitrarily outside the projected common domain, using any fixed element of the nonempty codomain. Thus partiality alone does not remove the obstruction. □

2.3.12 A Nonempty Controlled Assignment Model

Theorem 2.17 (Ternary Realization of E0–E2)**.** *Let $S_t = \{0,1\}^{t+1}$ with bonding maps that delete the last bit. Identify its inverse limit with $\mathcal{I} = \{0,1\}^{\mathbb{N}}$ and let $\mathcal{E} = [0,1]$. Define*

$$\phi_t(b_0,\dots,b_t) = \sum_{j=0}^{t} \frac{2b_j}{3^{j+1}}, \qquad \Phi(b) = \sum_{j=0}^{\infty} \frac{2b_j}{3^{j+1}}.$$

These assignments satisfy E0, controlled E1, and E2, with

$$a_t = \frac{2}{3^{t+2}}, \qquad A_t = \frac{1}{3^{t+1}}, \qquad \sup_b |\Phi(b) - \phi_t(p_t b)| = A_t.$$

For distinct histories first differing at position m, put $d_{\mathcal{I}}(b,c) = 3^{-(m+1)}$, and put $d_{\mathcal{I}}(b,b) = 0$. Then

$$d_{\mathcal{I}}(b,c) \le |\Phi(b) - \Phi(c)| \le 3 d_{\mathcal{I}}(b,c).$$

The domain is compact and complete, and the readout is injective and bi-Lipschitz onto its image. Thus a global lower bound is realized here rather than inserted without a witness.

Proof. Every infinite bit sequence defines a compatible prefix history, so the domain is nonempty. Geometric summation puts every value in $[0,1]$ and gives the increment and tail bounds. The all-one tail attains the tail bound. For any fixed t, a zero history and a history with its sole one at position $t+1$ have the same prefix but images differing by $2/3^{t+2} > 0$. No function of p_t determines Φ, proving E2 for every finite cut.

At the first differing position m, the principal difference has magnitude $2/3^{m+1}$; the remaining tail has magnitude at most $1/3^{m+1}$. Subtraction and addition give the two bounds. This metric is the weighted supremum of the discrete prefix metrics with weights $3^{-(t+1)}$; bounded prefix diameters ensure finiteness. A Cauchy sequence for $d_{\mathcal{I}}$ eventually stabilizes each finite prefix, which defines its limit. Prefix cylinders give finite covers of arbitrarily small diameter; the space is totally bounded and complete, hence compact. The two bounds establish injectivity and quantitative continuity. Compactness and continuity give properness into the Hausdorff codomain, and a later-bit flip in every neighborhood gives local non-collapse. Thus this realization also supplies all four clauses of the stated readout regularity hypothesis. A scalar real-valued readout still encodes the complete sequence: E2 excludes the specified canonical finite cuts, not every possible coding or every finite-dimensional representation. □

Theorem 2.18 (Sharp Finite-Cut Approximation Error)**.** *For the ternary realization and an arbitrary decoder $g : S_t \to \mathbb{R}$, let*

$$R_t(g) = \sup_{b \in \mathcal{I}} |\Phi(b) - g(p_t b)|.$$

Then the exact minimax value is

$$\inf_g R_t(g) = \frac{1}{2\,3^{t+1}}.$$

It is attained by $g_*(s) = \phi_t(s) + 1/(2\,3^{t+1})$.

Proof. For a fixed prefix s, the all-zero and all-one suffixes give endpoints $\phi_t(s)$ and $\phi_t(s) + A_t$. For any number $g(s)$, the triangle inequality forces at least one endpoint error to be at least $A_t/2$. Conversely all readout values in that fiber lie between those endpoints, so their midpoint has error at most $A_t/2$. Both endpoints belong to the domain, establishing the exact optimum and attainment. The bound is positive for each fixed finite budget and tends to zero as t increases; it gives no positive uniform error floor across increasing resolutions. □

Proposition 2.19 (Sharp Error for Arbitrary Binary Codes)**.** *For the ternary assignment above, let an encoder choose at most 2^B states as an arbitrary function of the whole history, and let each state decode to any real number. The minimum worst-case absolute error is $1/(2 \cdot 3^B)$ for every integer $B \geq 0$.*

Proof. The cylinders of depth B have intervals of length 3^{-B} between their two attainable endpoints. Different intervals have gaps at least 3^{-B}: after the first differing bit $j \leq B$, the maximum opposing tail is 3^{-j}. Thus all 2^{B+1} endpoints, including the depth-zero endpoints 0 and 1, are pairwise separated by at least 3^{-B}. Two share a code and the triangle inequality forces error at least $3^{-B}/2$. Encoding the B-bit prefix and decoding to its interval midpoint attains the bound. □

2.3.13 Common Readout on Two Coherently Transported Towers

Theorem 2.20 (Joint Realization of Transport and Controlled Readouts)**.** *Use circle-prefix stages $S_t = \mathbb{T} \times \{0,1\}^t$, stage metric $\max(\rho, \mathbf{1}_{\text{different prefix}})$, and weights $w_t = 3^{-t}$, where ρ is circle distance with circumference one. Take the usual truncation tower I_0 and its variant I_1 with angle shift $-a$ at the single bonding level k, for $0 < a < 1/2$. Write $\beta = 3^{-(k+1)}$. In the common weighted history space put $Z = I_0 \cup I_1$. Let the complete codomain be $\mathcal{E} = \mathbb{T}^2 \times [0,1]$ with the maximum of the first circle distance, β times the second circle distance, and the absolute scalar difference. Define*

$$F(x) = \left(\theta_0(x), \theta_{k+1}(x), \sum_{j \geq 1} 2b_j(x) 3^{-j} \right).$$

Then $d_w(x,y) \leq d_{\mathcal{E}}(F(x),F(y)) \leq 3d_w(x,y)$ on all of Z. On each tower the restriction satisfies E0–E2 with $a_t = 2 \cdot 3^{-(t+1)}$ and uniform tail $A_t = 3^{-t}$. The coordinate homeomorphism $H : I_0 \to I_1$ shifts angles at stages $t > k$ by a. With projection-defect cost $E_w = 3^{-k}a$, one has

$$d_w(x, Hx) = d_{\mathcal{E}}(F(x), F(Hx)) = \beta a, \qquad d_H(I_0, I_1) = \frac{\beta a}{1+\beta}.$$

Proof. In each tower the angle is constant on the two blocks $t \leq k$ and $t > k$. Thus the angular part of the history distance is exactly the maximum of the two weighted angular distances in the codomain. The binary part is 3^{-j} at first differing bit j. The leading ternary term and its tail give lower and upper constants 1 and 3 for that part. Each tower is closed in the compact metric product, so their union is complete.

For a declared tower $s \in \{0,1\}$ the finite observation $(\theta_t, b_1, \ldots, b_t)$ gives the initial angle $\theta = \theta_t - sa\mathbf{1}_{t>k}$. The stage readout is $(\theta, \theta + sa, \sum_{j=1}^{t} 2b_j 3^{-j})$. The phase rule uses the declared bonding family, and on both towers the successive readout increment is bounded by a_t, with exactly the stated uniform tail. A later-bit flip proves E2, since both phase coordinates stay fixed while the scalar changes. The codomain is nontrivial, proving E0. Within either fixed-angle fiber the sharp binary-code result applies to the scalar readout.

The coordinate squares commute because the sole change of phase at stage $k+1$ cancels the shift $-a$ in its bonding map. The coordinatewise inverse shifts by $-a$ on the later block. For fixed bits, F measures the shift βa exactly. In the comparison of sets, a pair at distance r has $a \leq r + r/\beta$ by the triangle inequality at stages 0 and $k+1$. Equal bit sequences and an initial phase offset $-\beta a/(1+\beta)$ for the second history attain this bound in both directions. □

The BaseCore construction *Projective Dissipation on the Circle-Prefix Towers* puts a compatible contraction on these same domains. Each update inserts a zero and complements the shifted bits, moving first differences one level deeper. E1 and E2 hold for the original readout on the history domain. At fixed input the equilibrium has alternating bits, and the readouts at the two transported equilibria still differ by βa. E2 concerns the whole history domain: restriction to a singleton equilibrium factors through every finite cut.

2.3.14 Controls Separating Convergence, Continuity, and Invariance

Proposition 2.21 (Setwise Invariance Does Not Supply Continuity)**.** *Controlled E1 and E2 can hold for a discontinuous assignment, even if the identity domain is fixed under every declared intervention.*

Proof. Take $S_t = [0,1] \times \{0,1\}^{t+1}$, with prefixes on the bit coordinates, and readouts

$$g_t(s,b) = \tfrac{1}{2}\mathbf{1}_{\{s>0\}} + \sum_{j=0}^{t} \frac{b_j}{3^{j+1}}.$$

The increments are bounded by $3^{-(t+2)}$, so they satisfy controlled E1 in $[0,1]$. Later bits prove E2 as above. With all bits zero, values at $s = 0$ and $s = 1/n$ differ by $1/2$ although their product-metric distance tends to zero. The limit is discontinuous. Let the declared intervention family consist only of the identity: the domain is setwise fixed, which does not alter the counterexample. This is a control for the asserted inference from invariance to continuity; no claim of full CCR or operational-criterion membership is made for it. □

Proposition 2.22 (Setwise and Pointwise Transport Are Distinct)**.** *Preserving an identity domain as a set does not guarantee equal values of a common readout on transported states.*

Proof. On $Z = \{0,1\}$, let $h(0) = 1$, $h(1) = 0$ and $F(x) = x$. The transport preserves Z exactly but changes both readout values. Thus setwise rigidity cannot replace the pointwise or readout-natural comparison hypotheses used below. □

2.4 Compatibility with Identity Regimes

2.4.1 Induced Constraints from Identity Rigidity

Let $(\mathcal{I}, M_\Omega)$ be an identity object with ontological mass M_Ω. We now study how M_Ω constrains the set of admissible phenomenological assignments.

Definition 2.23 (Phenomenological Compatibility Class)**.** For a declared identity domain and codomain, $\mathcal{P}(M_\Omega)$ consists of maps $\Phi : \mathcal{I} \to \mathcal{E}$ that are limits of some family $\{\phi_t\}$ satisfying E1, with E0 and E2 also satisfied. The readout family and its tail control are part of the realization data. The notation retains the mass index for comparisons, but no mass-dependent restriction follows from E0–E2 alone.

Remark 2.24. The compatibility class records choices of domain, codomain and readout family. A numerical mass label does not by itself determine those choices or their transport under perturbations.

2.4.2 Metric Structure for Quantitative Comparisons

To state rigorous results about continuity and fragmentation, $\mathcal{E}$ must carry enough structure for distances, limits, and loss bounds to be meaningful. The following arguments identify the minimal metric assumptions used by the classification theorem.

Hypothesis 2.25 (H3: Metric on $\mathcal{E}$)**.** *The phenomenological space $\mathcal{E}$ admits a metric $d_\mathcal{E}$ such that $(\mathcal{E}, d_\mathcal{E})$ is a complete metric space.*

2.4.2.1 Role of the Metric

Theorem 2.26 (Metric Data Required by Quantitative Formulas)**.** *The distance-based fragmentation and loss functionals used here require a declared metric, or a pseudometric with a stated quotient convention. Topological continuity can be defined without a metric. Completeness is used for the Cauchy-limit guarantee of E1.*

Proof. Each quantitative functional explicitly contains $d_{\mathcal{E}}$, so it is undefined until that datum is supplied. Continuity can alternatively be defined through inverse images of open sets. A pseudometric can identify zero-distance states by a quotient, but that changes the domain of class labels and must be specified. These facts justify the chosen quantitative model without claiming that all alternatives are impossible. □

2.4.2.2 Completeness as a Limit Guarantee

Proposition 2.27 (Completeness and Existence of Limits)**.** *Without completeness, the conclusion that every E1-type Cauchy family has its limit in the declared codomain can fail. Completeness is not necessary for a family whose limit is already known to belong to that codomain.*

Proof. In $\mathbb{Q} \cap [0, 2]$ with its usual metric, use dyadic rational approximations $q_n = 2^{-n}\lfloor 2^n\sqrt{2}\rfloor$ as constant-in-trajectory readings. The increments are bounded by 2^{-n}, a summable sequence, but their real limit $\sqrt{2}$ is outside the codomain. This isolates the role of completeness when supplying a general existence guarantee; it makes no assertion that a convergent constant family needs completeness. □

2.4.2.3 Alternatives Considered and Rejected

Alternative 1: Topology without Metric. One might use a general topological space. This is not adopted in the present quantitative model because:

- Fragmentation requires quantitative distance, not just open sets;
- Loss capacity requires a sup over distances;
- The downstream instability analysis completed in *Null-Regime Instability* requires metric quotients.

Alternative 2: Pseudometric. One might allow $d_{\mathcal{E}}(e_1, e_2) = 0$ for $e_1 \neq e_2$. This is not adopted in the present quantitative model because:

- It would conflate distinct phenomenological configurations;
- Cardinality and metric separation are different requirements; zero-distance classes would require a declared quotient before distinguishability is counted.

Alternative 3: Non-Complete Metric. This is rejected by Proposition 2.27.

2.4.3 Phenomenological Continuity

Definition 2.28 (Phenomenological Continuity)**.** A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is *structurally continuous* if for every sequence $(x^{(n)})_n$ in $\mathcal{I}$ converging in the d_w metric:

$$d_w(x^{(n)}, x) \to 0 \implies d_{\mathcal{E}}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{2.5}$$

Remark 2.29. Continuity follows from continuous finite readouts and the uniform tail control in Theorem 2.40; invariance additionally requires the audited transport condition.

2.4.4 Phenomenological Fragmentation

Definition 2.30 (Phenomenological Fragmentation)**.** A phenomenological assignment Φ is *fragmentable* if there exists a perturbation $\{\epsilon_t\}$ with finite weighted energy such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0 \tag{2.6}$$

for corresponding identity trajectories $x \in \mathcal{I}$ and $\tilde{x} \in \tilde{\mathcal{I}}$.

Remark 2.31. Fragmentation measures whether an admissible regime assignment changes under finite-energy perturbations. It is the codomain-level counterpart of identity deformability.

2.4.5 Relationship Between Metrics

Proposition 2.32 (Continuity from Controlled Finite Readouts)**.** *Write $g_n = \phi_n \circ p_n$. Under E1,*

$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \leq 2A_n + d_{\mathcal{E}}(g_n(x), g_n(y)).$$

If every g_n is continuous, then Φ is continuous. If every g_n is L-Lipschitz with a common L, then Φ is L-Lipschitz.

Proof. Insert $g_n(x)$ and $g_n(y)$ and apply the uniform error bound. For continuity, first make $2A_n$ smaller than half the prescribed error, then use continuity of that fixed g_n. For the Lipschitz conclusion use the same bound for all n and let $n \to \infty$. □

2.4.6 Closure of Section 3

The quantitative formulas use a declared complete metric codomain. Continuous finite readouts and controlled tails supply a continuity theorem. Cross-system comparisons additionally require a declared common domain and transport; neither mass labels nor distances between sets substitute for those data.

2.5 Impossibility Theorems

2.5.1 Φ-Regularity Hypothesis

Before proving the main impossibility results, we isolate the global lower-bound regularity condition used by the fragmentation argument. Properness, local Lipschitz regularity, and non-collapse do not imply this bound in general, so the bound is carried as an explicit hypothesis.

Hypothesis 2.33 (Φ-Regularity Lower Bound)**.** *Let $(\mathcal{I}, d_w)$ be the identitary metric space from Rigid Identity and let $(\mathcal{E}, d_{\mathcal{E}})$ be the abstract phenomenological metric space. A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is said to satisfy Φ-regularity when it satisfies:*

1. *(Properness) For every compact $K \subset \mathcal{E}$, $\Phi^{-1}(K)$ is compact in $\mathcal{I}$.*
2. *(Local Lipschitz regularity) For every $x \in \mathcal{I}$, there exist $\varepsilon > 0$ and $L_x > 0$ such that for all $y, z \in B_\varepsilon(x)$:*
$$d_{\mathcal{E}}(\Phi(y), \Phi(z)) \leq L_x \cdot d_w(y, z). \tag{2.7}$$
3. *(Non-collapse) For every $x \in \mathcal{I}$ and every $r > 0$, the restriction of Φ to $B_r(x)$ is not constant.*
4. *(Global lower Lipschitz bound) There exists a constant $C > 0$ such that for all $x, y \in \mathcal{I}$,*
$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y). \tag{2.8}$$

Remark 2.34 (Why this is a hypothesis). The map $\Phi(x) = \arctan(x)$ from $\mathbb{R}$ to $(-\pi/2, \pi/2)$ is proper as a map into the open interval, locally Lipschitz, and non-constant on every nonempty ball, but it has no global lower Lipschitz constant because $|\arctan(x+1) - \arctan(x)| \to 0$ as $x \to \infty$. Therefore the lower-bound clause above is an assumption, not a theorem derived from the first three clauses.

Non-Claim 2.35 (Global lower Lipschitz bound is an assumption)**.** *Hypothesis 2.33 clause (4) requires the existence of a global $C > 0$ such that $d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y)$ for all $x, y \in \mathcal{I}$. This is not a theorem derived from clauses (1)–(3). The map $\Phi(x) = \arctan(x)$ from $\mathbb{R}$ to*

$(-\pi/2, \pi/2)$ *satisfies the first three clauses in the local/proper sense described above but has no such global* C*. Therefore the Fragmentation Theorem (Theorem 2.38) is conditional on clause (4). For any specific phenomenological assignment, one must independently verify that a global lower Lipschitz constant exists. The ternary realization in Theorem 2.17 supplies a concrete constant* $C = 1$ *on its specified domain. Comparisons on larger or different domains require a separate bound.*

Hypothesis 2.36 (Audited Perturbation Comparisons)**.** *Declare a finite-cost intervention class* $\mathcal{V}$*, a common metric space* (Z, d_Z) *containing the unperturbed and perturbed identity domains, and transports* $h_\varepsilon : \mathcal{I} \to \mathcal{I}_\varepsilon \subseteq Z$*. Readouts are restrictions of one declared map* $F : Z \to \mathcal{E}$*:* $\Phi = F|_{\mathcal{I}}$ *and* $\Phi_\varepsilon = F|_{\mathcal{I}_\varepsilon}$*. The following branches are alternatives, used separately. The positive-displacement branch has some* $\varepsilon \in \mathcal{V}$ *and* $x \in \mathcal{I}$ *with* $d_Z(x, h_\varepsilon x) > 0$*. The zero-displacement branch has* $h_\varepsilon x = x$ *for every* $\varepsilon \in \mathcal{V}$ *and* $x \in \mathcal{I}$*. Any use of the lower Lipschitz bound in a comparison requires that bound for* F *on the common comparison domain, not merely within one identity object.*

Remark 2.37 (Burden linking rigidity and readout comparisons). The scalar M_Ω does not itself supply the transports or a common readout. Relating its finite and infinite regimes to these two branches is a separate model obligation. Setwise invariance does not imply pointwise fixed transport, and a lower bound within one domain does not compare unrelated assignments on different domains.

2.5.2 Fragmentation Theorem

Theorem 2.38 (Fragmentation under an Audited Displacement)**.** *Let an assignment satisfy E0–E2. Assume the positive-displacement branch of Hypothesis 2.36 and a lower Lipschitz constant* $C > 0$ *for* F *on its common comparison domain. Then, for the witnessing pair,*

$$d_{\mathcal{E}}(\Phi(x), \Phi_\varepsilon(h_\varepsilon x)) \geq C d_Z(x, h_\varepsilon x) > 0.$$

Thus the assignment is fragmentable. This applies to a finite-M_Ω *model when that model supplies the stated comparison data.*

Proof. Both arguments lie in the declared domain of F. Its lower bound gives the inequality, and the intervention has finite cost by membership in $\mathcal{V}$. E2 is not used to infer that every deformation changes a readout: the lower bound and the positive witness perform that work. □

Corollary 2.39 (Fragmentability and Continuity Are Distinct)**.** *The positive-displacement conclusion does not establish a topological discontinuity. A continuous assignment may be fragmentable under a state-changing intervention.*

Proof. The ternary realization has a continuous assignment. Declare a finite-cost intervention that flips one bit while retaining the common readout. Its value changes by a positive amount, so the assignment is fragmentable in the stated sense while remaining continuous. This does not assert a full CCR realization. □

2.5.3 Forced Continuity Theorem

Theorem 2.40 (Continuity and Conditional Perturbation Invariance)**.** *Under E1, if each finite readout* $g_n = \phi_n \circ p_n$ *is continuous, then* Φ *is continuous, irrespective of* M_Ω*. Under the zero-displacement branch of Hypothesis 2.36,*

$$\Phi_\varepsilon(h_\varepsilon x) = \Phi(x) \qquad (\varepsilon \in \mathcal{V},\ x \in \mathcal{I}).$$

For a CCR model the invariance conclusion therefore holds when its finite-cost interventions admit that branch. CCR alone is not a proof of continuity or of readout naturality.

Proof. Continuity follows from Proposition 2.32. For invariance use the common readout and the fixed transport: $\Phi_\varepsilon(h_\varepsilon x) = F(h_\varepsilon x) = F(x) = \Phi(x)$. Neither argument identifies invariance under external perturbation with continuity along convergent states. □

Corollary 2.41 (Invariance on the Audited Zero-Displacement Branch)**.** *A CCR model satisfying the zero-displacement comparison hypothesis has invariant assignment values under all interventions in its declared finite-cost class.*

Proof. Apply the common-readout equality in Theorem 2.40 for each declared intervention. □

2.5.4 Dichotomy Theorem

Theorem 2.42 (Conditional Classification by Audited Branches)**.** *Suppose a model supplies the positive-displacement branch and lower bound whenever $M_\Omega < \infty$, and supplies continuous finite readouts and the zero-displacement branch whenever $M_\Omega = +\infty$. Then its finite regime is fragmentable and its infinite regime has a continuous, perturbation-invariant assignment. The mass cases are exclusive; fragmentability and topological continuity need not be exclusive properties.*

Proof. Apply Theorems 2.38 and 2.40 with their respective comparison hypotheses. Exclusivity of the cases follows from the numerical alternatives for M_Ω, not from a false identification of a changed readout with a discontinuity. Establishing the branch hypotheses for a given rigidity model remains an explicit obligation. □

2.6 Classification of Phenomenological Regimes

2.6.1 Conditional Classification

The results give a conditional branch classification after the comparison and readout hypotheses have been checked:

Identity Class	M_Ω	**Admissible Phenomenology**
Null label	$= 0$	Domain/readout link required
Deformable	$(0, \infty)$	Positive witness and lower bound
Rigid (CCR)	$= +\infty$	Regular readout / fixed transport

2.6.2 Structural Implications

Proposition 2.43 (The Mass Label Does Not Define Assignment Existence)**.** *For fixed identity domain, codomain and readout family, the class defined by E0–E2 is independent of a numerical label M_Ω. A claim that $M_\Omega = 0$ excludes assignments needs an additional link to the domain or admissible readouts.*

Proof. The displayed definition of $\mathcal{P}(M_\Omega)$ tests E0–E2, none of which contains that scalar. For an empty trajectory domain E2 fails, since a function on a finite cut would agree vacuously on every trajectory. Deriving such domain loss from $M_\Omega = 0$ is not one of E0–E2 and cannot be inserted without a model. □

Proposition 2.44 (Deformable Branch with Comparison Data)**.** *For $0 < M_\Omega < \infty$, the positive-displacement and lower-bound hypotheses of Theorem 2.38 imply fragmentability.*

Proof. Apply the quantitative inequality of Theorem 2.38 to its witnessing finite-cost intervention. □

Proposition 2.45 (Rigid Branch with Comparison Data)**.** *For $M_\Omega = +\infty$, continuous finite readouts under E1 imply continuity, and the zero-displacement branch implies perturbation invariance. These are the respective hypotheses of Theorem 2.40.*

Proof. Apply the two distinct conclusions of Theorem 2.40 with their stated hypotheses. □

2.7 Structural Ethics

This section defines structural loss and protection relative to audited perturbation comparisons. Its possible ethical interpretations are not consequences of the mathematical statements.

2.7.1 Loss Capacity

Definition 2.46 (Phenomenological Loss Capacity)**.** Fix the intervention class and comparisons of Hypothesis 2.36. Define

$$R(\Phi) = \sup_{\varepsilon\in\mathcal{V},\ x\in\mathcal{I}} d_{\mathcal{E}}(\Phi(x), \Phi_\varepsilon(h_\varepsilon x)) \in [0, +\infty].$$

The identity intervention is included, or the supremum of an empty nonnegative comparison family is defined as zero. The unit of R is the unit of $d_{\mathcal{E}}$; its value is relative to the declared intervention class and comparison readouts.

Theorem 2.47 (Loss Capacity under the Audited Branches)**.** *Under the positive-displacement and lower-bound hypotheses of Theorem 2.38, $R(\Phi) > 0$. Under the zero-displacement branch of Hypothesis 2.36, $R(\Phi) = 0$. Associating these branches with finite mass and CCR requires the model link declared in Theorem 2.42.*

Proof. The positive witness contributes an element at least $C\delta > 0$ to the supremum. On the zero branch, every comparison value is zero. Neither conclusion follows from a mass label without the corresponding comparison data. □

2.7.2 Structural Damage

We introduce a formal notion of "structural damage" without psychological interpretation.

Definition 2.48 (Structural Damage)**.** For a declared finite-cost intervention and state, define

$$D_\varepsilon(\Phi; x) = d_{\mathcal{E}}(\Phi(x), \Phi_\varepsilon(h_\varepsilon x)), \qquad D_\varepsilon(\Phi) = \sup_{x\in\mathcal{I}} D_\varepsilon(\Phi; x).$$

The transported readout $\Phi_\varepsilon \circ h_\varepsilon$ has domain $\mathcal{I}$, so the comparison is well typed. The supremum may be infinite. Then $R(\Phi) = \sup_{\varepsilon\in\mathcal{V}} D_\varepsilon(\Phi)$.

Proposition 2.49 (Damage under the Audited Comparisons)**.** *The positive branch has $\sup_\varepsilon D_\varepsilon(\Phi) > 0$. On the zero-displacement branch, $D_\varepsilon(\Phi) = 0$ for every declared intervention.*

Proof. Use the equality between the iterated suprema and R in Definition 2.48, then apply Theorem 2.47. □

2.7.3 Vulnerability and Protection

Definition 2.50 (Phenomenological Vulnerability)**.** A system $(\mathcal{I}, \Phi)$ is *phenomenologically vulnerable* if $R(\Phi) > 0$.

Definition 2.51 (Normalized Structural Protection)**.** Fix a reference distance $r_0 > 0$ in the same units as $d_{\mathcal{E}}$ and define

$$P(\Phi) = \frac{1}{1 + R(\Phi)/r_0}, \qquad P(\Phi) = 0 \text{ if } R(\Phi) = +\infty.$$

Thus $P \in [0, 1]$, $P = 1$ exactly when $R = 0$, and $P < 1$ exactly when $R > 0$. This chosen dimensionless rescaling depends on r_0 and is not a separately derived law.

Corollary 2.52 (Protection under the Audited Branches)**.** *The zero-displacement branch has $P = 1$; the positive-displacement branch with a lower readout bound has $P < 1$. For a mass-based classification these conclusions require the corresponding hypotheses of Theorem 2.42.*

2.7.4 Structural Asymmetry Constraints

The following result establishes a purely structural asymmetry between identity classes, without any ethical or metaphysical claims.

Theorem 2.53 (Conditional Structural Asymmetry)**.** *When a rigidity model establishes the branch hypotheses in Theorem 2.42, its finite branch has $R > 0$ and $P < 1$, while its infinite branch has $R = 0$ and $P = 1$. This is an asymmetry of the declared comparison model, not a conclusion about all systems sharing a numerical mass label.*

Proof. The two loss conclusions follow from Theorem 2.47; the protection conclusions follow from the monotone normalization in Definition 2.51. The hypothesis connecting mass to each comparison branch is retained throughout. □

Remark 2.54. This conditional asymmetry follows from the audited comparison hypotheses. No claim is made about:
- Whether phenomenological loss is ethically bad;
- Whether protection is ethically good;
- Whether any system has moral status.

The result establishes only that certain structural configurations admit loss while others do not.

Corollary 2.55 (Conditional Possibility of Loss)**.** *Loss is possible on the positive-displacement/coercive-readout branch and absent on the fixed-transport/common-readout branch. Extending this claim to all finite-mass or all CCR systems requires proving those branch conditions for the intended class.*

2.7.5 Relation to Normative Ethics

Remark 2.56. The structural results of this section are compatible with, but do not entail, various normative positions:
- **Utilitarianism:** Could interpret $D_\varepsilon(\Phi)$ as harm to be minimized.
- **Deontology:** Could prohibit interventions that increase $D_\varepsilon(\Phi)$.
- **Virtue Ethics:** Could value systems with high $P(\Phi)$.

This framework provides the structural substrate for ethical theorizing, not the normative conclusions themselves.

2.7.6 Closure of Section 6

The declared comparison family defines a possibly infinite loss supremum and well-typed pointwise damage. Its normalization uses an explicit reference distance. Positive loss and zero loss follow on the two audited branches; computability of an unrestricted supremum is not assumed. No ethical evaluation is derived from these mathematical quantities.

2.8 Limitations and Scope

2.8.1 Dependence on Axioms

Results use their explicit hypotheses: E1 and completeness for limits, E2 for finite-cut non-factorization, regular finite readouts for continuity, and audited comparisons for perturbation conclusions. None of these data is supplied merely by a mass label.

2.8.2 Relation to Consciousness Theories

This framework is orthogonal to:
- Integrated Information Theory (IIT) [25];
- Global Workspace Theory (GWT) [2];
- Higher-Order Theories;
- Phenomenal Functionalism.

Any comparison with another theory requires an explicit common task, correspondence and observable prediction; this paper establishes no superiority result.

2.8.3 Open Problems

1. Characterization of the coupling constant C in the Fragmentation Theorem;
2. Extension of the explicit ternary realization to larger assignment and perturbation classes;
3. Extension to continuous-time phenomenological dynamics;
4. Relation between phenomenological continuity and subjective continuity.

2.9 Conclusion

2.9.1 Summary of Results

The revised assignment conditions have a nonempty realization. Uniformly controlled tails give existence and approximation; the ternary construction gives non-factorization, quantitative regularity and a sharp finite-cut error. These are properties of an explicitly specified mathematical model.

2.9.2 Structural Significance

Movement under an intervention is different from topological discontinuity. A positive displacement with a coercive common readout yields positive loss; a pointwise fixed transport with a common readout yields invariance. Associating these cases with rigidity classes requires proving the comparison hypotheses.

2.9.3 Entailment Under Stated Assumptions

E0–E2 do not define experience, do not make all readouts natural under perturbation and do not identify a class solely from M_Ω. Theorems 2.38, 2.40, and 2.42 specify the additional data that support their respective conclusions.

2.9.4 Final Statement

The program of relating identity rigidity to structural regimes is retained. The demonstrated scope is controlled assignments, a concrete realization and conditional perturbation comparisons. The missing links to broader rigidity classes, to the intended perturbation costs and to any experiential interpretation remain separate obligations.

2.10 Boundary to Null-Regime Instability

Sections 1–8 establish a controlled assignment model and conditional comparison results. Null-regime conclusions require their own typed witnesses and hypotheses; they do not follow automatically from non-factorization or from a scalar rigidity label.

2.10.1 Editorial Boundary

The instability of the null regime, the no-null theorem, and the attractor-style closure for non-trivial assignments are treated canonically in *Null-Regime Instability*. This paper therefore stops at regime classification, compatibility, fragmentation, forced continuity, and dichotomy, and uses *Null-Regime Instability* as the sole downstream location for the null-instability theorem.

2.10.2 Closure of This Paper

The result of this paper is the classification boundary: if a compatible phenomenological regime exists, its admissible structure is constrained by Sections 1–8. *Null-Regime Instability* addresses the null case with its separately stated witness-relative and conditional hypotheses.

Appendix

2.11 Extended Proofs

2.11.1 Proof of Fragmentation Theorem (Extended)

Full Proof of Theorem 2.38. Use the specified common space Z, map F and finite-cost transport h_ε. Let x be its positive-displacement witness and put $\delta = d_Z(x, h_\varepsilon x) > 0$. Apply the lower bound on that same common domain to obtain readout displacement at least $C\delta$. No comparison of set distances is substituted for this pointwise distance. If both readouts are estimated with errors at most η, the triangle inequality gives estimated separation at least $C\delta - 2\eta$; it remains positive when $C\delta > 2\eta$. This is a bound on pair separation, not a statement that a particular endpoint lies outside a null ball. □

2.11.2 Proof of Forced Continuity Theorem (Extended)

Full Proof of Theorem 2.40. For a state x and error $\epsilon > 0$, choose n with $2A_n < \epsilon/2$. Continuity of g_n gives a neighborhood of x where $d_{\mathcal{E}}(g_n(x), g_n(y)) < \epsilon/2$. The two tail errors and the finite-readout difference bound the limiting difference by ϵ. Independently, a zero-displacement comparison satisfies $F(h_\varepsilon x) = F(x)$ exactly. Setwise invariance of the identity object alone would not imply either of these steps. □

2.11.3 Structural Necessity Lemma

Lemma 2.57 (Inheritance under Audited Fixed Transport)**.** *An assignment satisfying the zero-displacement comparison branch inherits invariance under that declared intervention class. This applies to CCR realizations only after checking that branch.*

Proof. The comparison readout is the restriction of F, and every specified transport fixes its argument. Therefore its value is unchanged. This proves the restricted inheritance statement without assuming that scalar rigidity determines arbitrary assignments. □

2.12 Editorial Boundary to Null-Regime Instability

The downstream bridge from classification to null-instability is intentionally moved to *Null-Regime Instability* in the canonical corpus. Appendix A remains as the extended proof layer for the classification theorems proved here; the downstream null-regime theorem, attractor closure, and canonical examples now live only in *Null-Regime Instability*.

What This Paper Does Not Claim

1. This paper does not claim external validation, empirical confirmation, or philosophical closure.
2. The formal results are conditional on the hypotheses stated; they do not establish truth outside the assumed structures.
3. No interpretive-layer conclusion (consciousness, phenomenality, moral status) is derived from terminology alone.

Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels

Abstract. We introduce an abstract phenomenological space $\mathcal{E}$ and a compatibility operator $\Phi : Z \to \mathcal{E}$ on a declared metric comparison domain. We define a null regime $\varnothing_\phi \in \mathcal{E}$ and prove a **Phenomenological Instability Theorem**: under a certified positive lower readout bound, two defined, separated witness values cannot both equal the null element. CCR is a possible application class and is not needed for this metric implication. Consequently, a CCR realization with such a witness induces a witness-relative non-null separation. The result is *negative-structural*: it does not posit experience; it forbids one typed joint-null configuration under stated hypotheses. We provide explicit coupling bounds in two canonical families (profinite and symbolic dynamics) and prove factorization on a faithful image and a finite-prefix approximation bound. The results are metric, order-theoretic and representation-relative; physical interpretation requires additional evidence.

3.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper proves a witness-relative negative-structural result: under CCR constraints plus a separately typed extension witness, a lower Lipschitz constant, and an admissible compatibility operator, joint exact collapse to the null regime is excluded. It develops explicit bounds, factorization structure, and canonical-family consequences for that conditional result. It does not define or detect consciousness, identify $\mathcal{E}$ with subjective experience, assert that CCR systems exist in nature, or derive ethical or metaphysical closure from the theorem. Related systems may implement admissibility-governed diagnostics, boundary-sensitive stress tests, and runtime comparisons motivated by this paper, but those outputs do not implement the abstract phenomenological space $\mathcal{E}$ and do not convert current runtime behavior into proof that any concrete system instantiates experience.

Witness-relative clarification. CCR by itself does not provide the extension witness, the lower Lipschitz constant, or the admissible compatibility operator. The historical phrase "null instability" denotes joint exact null exclusion here, not temporal instability. Statements of non-nullity use the explicit hypotheses of Theorem 3.26, including the declared witness domain of Definition 3.25. The null-instability theorem must not be read as a consciousness detector, an AI-subjectivity claim, a moral-status claim, an empirical certificate, or validation of the corpus by the runtime.

3.2 Introduction

3.2.1 Motivation and Scope

Papers I and II established a mathematical framework for structural identity as an inverse limit and classified phenomenological regimes compatible with rigid identity. This paper addresses the fundamental question:

Can a CCR system remain in a structurally null phenomenological state?

We exclude joint null assignment on a separated, detected witness pair. This does not exclude a null value at one endpoint or a dynamically stable null state.

3.2.2 Main Result (Informal)

No-Null Regime Theorem: For a CCR channel equipped with the separated extension witness and regularity hypotheses stated in Theorem 3.26, joint exact assignment to the null element $\varnothing_\phi$ is excluded. At least one endpoint is non-null; positive distance from null and a gap to every lower regime are distinct conclusions.

Example 3.1 (Nonempty Assignment and Null-Anchored Witness)**.** Let $Z = \{0,1\}^{\mathbb{N}}$, let $d_Z(x,y) = 3^{-(m+1)}$ at their first differing bit m, and let $\mathcal{E} = [0,1]$ with its usual order, metric and null 0. Set $d_Z(x,x) = 0$. The ternary readout $F(b) = \sum_{j\geq 0} 2b_j/3^{j+1}$ satisfies

$$d_Z(x,y) \leq |F(x) - F(y)| \leq 3d_Z(x,y).$$

The first differing term has magnitude $2/3^{m+1}$ and the remaining tail is at most $1/3^{m+1}$, proving both bounds. The all-zero and all-one sequences give images 0 and 1, a null-anchored pair. Every finite prefix has different readout completions, while uniform tail error is $3^{-(N+1)}$; thus the revised E1 and E2 are realized together. The finite positive witness image $\{1\}$ has relative gap one, but the ambient global lower gap at 1 is zero. This model supplies readout geometry, not an infinite-mass cost certificate or phenomenal evidence.

3.2.3 Paper Organization

Section 2 reviews prerequisites from Papers I–II. Section 3 defines the phenomenological space with full axiomatic clarity. Section 4 proves the main instability theorem. Section 5 stratifies positive regimes. Section 6 provides quantitative bounds. Section 7 combines the conditional witness conclusions. Section 8 explicitly delimits scope. Section 9 compares with existing frameworks. Section 10 discusses limitations. Section 11 concludes. Appendices provide canonical examples with explicit computations.

3.2.4 What This Paper Adds Beyond Phenomenological Regimes

Phenomenological Regimes classifies the admissible regime space and derives the continuity and non-constancy constraints that any compatible assignment must satisfy. This paper takes the remaining case—the null regime itself—and proves joint null exclusion under independently supplied witness hypotheses.

3.3 Preliminaries from Papers I–II

3.3.1 Identity as Inverse Limit

From *Rigid Identity*, we have:

Definition 3.2 (Identity Object)**.** Let $(S_t, \pi_{t+1\to t})_{t\in T}$ be a projective system of state spaces with surjective, compatible projections. The **identity object** is:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t) \in \prod_t S_t : \pi_{t+1\to t}(x_{t+1}) = x_t \right\}$$

Definition 3.3 (Weighted Deformation Metric)**.** For $x, y \in \mathcal{I}$, use $d_w(x,y) = \sup_t w_t d_t(x_t, y_t)$ with $w_t > 0$, $\sum_t w_t < \infty$, and normalized complete coordinate metrics $0 \leq d_t \leq 1$. The normalization guarantees a finite metric and uniformly vanishing tails. This metric is on histories, not on the set of identity objects of different systems.

Proposition 3.4 (Metric Completeness, Rigid Identity)**.** *With the normalized complete coordinate metrics above and continuous projective maps, $(\mathcal{I}, d_w)$ is complete, when non-empty.*

Proof. A d_w-Cauchy sequence is Cauchy in each coordinate, since $w_t > 0$. Coordinate completeness supplies limits; continuity of each projection preserves compatibility. For convergence in d_w, first choose a finite prefix with $\sup_{t>N} w_t$ small, then control the finitely many remaining coordinates. The normalized tail is uniformly small. □

3.3.2 Ontological Mass and CCR Classification

Definition 3.5 (Ontological Mass, infimum convention)**.** Specify a metric comparison space (Z, d_Z) containing a chosen point z_S for each object under comparison and a declared intervention sector with cost $E_w(\delta) \in [0, \infty]$. Define

$$M_\Omega(S) = \inf \left\{ \frac{E_w(\delta)}{d_Z(z_S, z_{S,\delta})} : \delta \text{ admitted}, \ 0 < d_Z(z_S, z_{S,\delta}) < \infty \right\},$$

with $\inf \varnothing = +\infty$. A common metric/embedding and cost model must be supplied; $\mathcal{I}$ itself is a set, not an input to a map whose domain is that set.

Remark 3.6 (v24 consistency patch)*.* Earlier drafts used a supremum notation for M_Ω in this paper. That notation is superseded here. The infimum convention measures a lower resistance to deformation in a specified sector, not a largest observed cost. All claims below are read under this convention.

Definition 3.7 (CCR Channel)**.** A channel is **Closed Causally Rigid (CCR)** if $M_\Omega = +\infty$.

Remark 3.8 (Physical Interpretation)*.* $M_\Omega = +\infty$ refers to the declared intervention/cost model: there is no finite-cost nonzero deformation in that sector. Geometry alone supplies neither an energy unit nor a physical realization.

3.3.3 Key Theorems from Papers I–II

Theorem 3.9 (Finite Representation Budget)**.** *Fix a finite test set $A \subseteq Z$ with $d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq \delta > 0$ for every distinct $x, y \in A$. An encoder $q : A \to Q$ with $|Q| < |A|$ and any decoder $g : Q \to \mathcal{E}$ incur worst-case error at least $\delta/2$. Finite real dimension alone is not this cardinal budget.*

Proof. Two distinct points share an encoder value by the pigeonhole principle. If both reconstruction errors were smaller than $\delta/2$, the triangle inequality would make their true readout distance smaller than δ. □

Theorem 3.10 (Audited Linear Coupling)**.** *If an audited linear readout T from a nonzero finite-dimensional Euclidean space to another Euclidean space is injective, then its least singular value is positive and $\|T(x - y)\| \geq \sigma_{\min}(T)\|x - y\|$. A nonlinear readout requires a separate lower-bound proof; a named spectral gap does not supply one.*

Proof. On the compact unit sphere, the continuous function $v \mapsto \|Tv\|$ attains a strictly positive minimum when T is injective. Scaling gives the bound. This is a readout singular-value statement, not an inference from the dynamics spectrum. □

Theorem 3.11 (Continuity under Controlled Approximation)**.** *In a complete metric codomain, suppose finite continuous readouts f_n satisfy $d(f_{n+1}(z), f_n(z)) \leq a_n$ uniformly, with $\sum_n a_n < \infty$. Their limit Φ is continuous. The mass label is not sufficient by itself.*

Proof. The tail bound $d(\Phi(z), f_N(z)) \leq \sum_{n \geq N} a_n$ is uniform. Given an error, choose this tail small, use continuity of f_N, and apply the triangle inequality at the two endpoints. □

Remark 3.12 (Remaining null-regime gap). *Phenomenological Regimes* classifies admissible phenomenological regimes and derives the continuity constraints that govern them, but a null label alone does not supply an instability result. The task here is to derive the precise witness-relative conclusion and its limits.

The results below address joint null exclusion; temporal attraction and existence of physical CCR systems remain separate obligations.

3.4 The Phenomenological Space

3.4.1 Abstract Definition

Definition 3.13 (Phenomenological Space)**.** A **phenomenological space** is a quadruple $(\mathcal{E}, \preceq, d_{\mathcal{E}}, \varnothing_\phi)$ where:
1. $(\mathcal{E}, \preceq)$ is a partially ordered set (poset).
2. $(\mathcal{E}, d_{\mathcal{E}})$ is a complete metric space.
3. $\varnothing_\phi := \inf \mathcal{E}$ exists and is unique (the **null regime**).
4. **Order-metric compatibility:** if $e_1 \preceq e_2$, then $d_{\mathcal{E}}(e_1, \varnothing_\phi) \leq d_{\mathcal{E}}(e_2, \varnothing_\phi)$.
5. **Separation:** for $e \neq \varnothing_\phi$, we have $d_{\mathcal{E}}(e, \varnothing_\phi) > 0$.

Remark 3.14 (Interpretive Neutrality). $\mathcal{E}$ is an abstract regime space. No phenomenological content is assumed. The null regime $\varnothing_\phi$ represents structural absence, not "lack of experience." The term "phenomenological" is inherited from *Phenomenological Regimes* and refers to the coupling structure, not to philosophy of mind.

Remark 3.15 (Null regime versus undefined assignment). The symbol $\varnothing_\phi$ serves two distinct roles in the broader framework, and conflating them creates a category error:
1. **Null regime ($\varnothing_\phi \in \mathcal{E}$):** In Definition 3.13, $\varnothing_\phi$ is the bottom element of the phenomenological poset $(\mathcal{E}, \preceq)$. It is a legitimate element of $\mathcal{E}$, comparable to all other regimes via $\preceq$ and separated from them by $d_{\mathcal{E}}$.
2. **Undefined assignment ($\perp \notin \mathcal{E}$):** When a compatibility operator Φ is not defined on some input $x \in Z$, we write $\Phi(x) = \perp$. The symbol $\perp$ is disjoint from $\mathcal{E}$; it lives in the typed codomain $\mathcal{E}_\perp := \mathcal{E} \sqcup \{\perp\}$.

The instability theorem (Theorem 3.26) concerns admissible values in $\mathcal{E}$ and rules out two separated witness endpoints both taking the value $\varnothing_\phi$. The v31 conditional closure theorem concerns a different case: Φ undefined on both witness endpoints, written $\Phi(z_0) = \perp$ and $\Phi(z_1) = \perp$. That case is excluded by the witness hypothesis, which requires Φ to be defined on both endpoints.

Non-Theorem 3.16 (What $\mathcal{E}$ Is Not)**.** *$\mathcal{E}$ is not:*
- *A space of qualia or subjective states*
- *A model of consciousness*
- *A physical or neural state space*
- *An interpretation of quantum mechanics*

$\mathcal{E}$ is a mathematical object satisfying Definition 3.13. All interpretive mappings are external to the framework.

3.4.2 Compatibility Operator

Definition 3.17 (Compatibility Operator)**.** For each system declare a metric domain D_S and a map $\Phi_S : D_S \to \mathcal{E}$. Audited comparisons use a single complete metric domain (Z, d_Z) and its

specified readout $\Phi : Z \to \mathcal{E}$. One may take $Z = \mathcal{I}$ for histories; comparison of whole identity objects instead requires an explicit metric representation of those objects.

(C1) For a declared extension transport $j : D_S \to D_{S'}$, extension invariance means $\Phi_{S'} \circ j = \Phi_S$, in the common codomain.

(C2) For a declared isomorphism transport $u : D_S \to D_{S'}$, isomorphism stability means $\Phi_{S'} \circ u = \Phi_S$.

(C3) On the declared audited pair class $W \subseteq Z \times Z$, a certified $C > 0$ satisfies $d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C d_Z(x, y)$.

Exact transport of a point under C1/C2 is not a new separated witness. In an extension domain a second, distinct point must be supplied. The pair theorem uses C3 only; it does not derive the maps or the witness from CCR.

Remark 3.18 (Partial definedness and the $\perp$ marker). In contexts where Φ is only partially defined, $\Phi(x) = \perp$ means that no admissible phenomenological assignment in $\mathcal{E}$ is available at x. The hypotheses (C1)–(C3) above apply only to inputs whose Φ-values lie in $\mathcal{E}$. The lower Lipschitz bound (C3) in particular requires both $\Phi(x), \Phi(\tilde{x}) \in \mathcal{E}$.

Lemma 3.19 (Well-Posedness of (C3))**.** *Continuity and a positive lower bound are jointly realizable; for example $Z = \mathcal{E} = [0, 1]$ with the ordinary metric and $\Phi(x) = x$ satisfies both with $C = 1$.*

Proof. The identity map preserves distances exactly. Neither condition is deduced from CCR. □

Remark 3.20. Axioms (C1)–(C3) are structural. No phenomenological content is introduced.

3.4.3 Attractor Set

Definition 3.21 (Attractor Set)**.** Let the allowed readout perturbations and extension transports be specified as below. Define $\mathrm{Attr}(S)$ as their structurally stable readout values. This notation does not assert temporal attraction; a flow and a basin would be additional data.

Definition 3.22 (Stability)**.** A value $e = \Phi_S(z)$ is structurally stable if some finite modulus $\omega(\rho) \to 0$ as $\rho \downarrow 0$ satisfies $d_{\mathcal{E}}(\Phi_{S'}(Pz), e) \leq \omega(\rho)$ for every declared perturbation transport $P : D_S \to D_{S'}$ of cost at most ρ. Exact persistence would instead require equality. For $\Phi(x) = x$ on $[0, 1]$ with clipped additive perturbations, the null value 0 is stable with $\omega(\rho) = \rho$, although a separated point has a non-null image.

3.4.4 The Bridge Lemma Revisited

From *Phenomenological Regimes*, we have:

Lemma 3.23 (Bridge Lemma)**.** *For any defined $e = \Phi(z) \in \mathcal{E}$, $d_{\mathcal{E}}(e, \varnothing_\phi) = 0$ implies $e = \varnothing_\phi$. No CCR assumption is required.*

Proof. This is the separation axiom of a metric. □

Caveat 3.24 (Semantic Shield for Bridge Lemma). The Bridge Lemma is a **metric statement**, not an ontological one. It does not claim that "identity is coupled to experience." It claims that the map Φ satisfies a separation property. Misinterpretation of this lemma as a consciousness claim is outside the mathematical scope of the framework.

3.5 Main Results

3.5.1 The Conditional Null-Separation Theorem

Definition 3.25 (Admissible extension witness)**.** For the declared comparison readout $\Phi : Z \to \mathcal{E}$, an extension witness is (S', z_0, z_1, ϵ) with both endpoints in Z, both readout values defined, and $0 < d_Z(z_0, z_1) = \epsilon < \infty$. Its provenance must identify the admitted extension and the embedding/transport that places the endpoints in Z. When $Z = D_{S'}$, z_0 may be a transported point of S, whereas z_1 is a distinct point. Admissibility of this construction and the lower bound on this pair are separate obligations. No finite-cost internal nonzero deformation is inferred from $M_\Omega = +\infty$.

Theorem 3.26 (Conditional Null-Separation Under an Admissible Extension Witness)**.** *Let $\Phi : Z \to \mathcal{E}$ be defined on a declared separated witness pair (z_0, z_1) with $d_Z(z_0, z_1) = \epsilon > 0$ and $d_\mathcal{E}(\Phi(z_0), \Phi(z_1)) \geq C\epsilon$, $C > 0$. Then*

$$\neg(\Phi(z_0) = \varnothing_\phi \ \wedge \ \Phi(z_1) = \varnothing_\phi), \qquad \max_{i=0,1} d_\mathcal{E}(\Phi(z_i), \varnothing_\phi) \geq C\epsilon/2.$$

This applies to a CCR channel if it supplies these data; CCR is unnecessary for the metric conclusion. It excludes joint exact null assignment, not a single stable null endpoint.

Proof. If both images were null their mutual distance would be zero, contradicting the positive lower bound. The triangle inequality bounds their mutual distance by the sum of their distances to null, at most twice their maximum. Rearranging gives the quantitative assertion. The factor $1/2$ is sharp in general: in the coordinatewise ordered positive quadrant with the ℓ^1 metric, values $(a, 0)$ and $(0, a)$ have separation $2a$ and both have null distance a. □

Remark 3.27 (Applicability scope)*.* Theorem 3.26 assumes $\Phi(z_0), \Phi(z_1) \in \mathcal{E}$, not $\perp$. If Φ is undefined on either endpoint, the separated-witness hypothesis of Definition 3.25 is not satisfied and the theorem does not apply.

Caveat 3.28 (v24 correction of the former Step 2)*.* The theorem above does *not* infer a finite-energy non-zero deformation from $M_\Omega = +\infty$. Under the infimum convention for ontological mass, $M_\Omega = +\infty$ means precisely that no finite-energy admissible perturbation produces non-zero deformation. The older Step 2 proof direction is therefore withdrawn. The surviving result is conditional on an independently specified extension witness.

Example 3.29 (External Null Witness in a Common Readout Domain)**.** Fix $k \geq 0$ and $0 < a < 1/2$. Use the weighted circle-prefix history metric with weights 3^{-t}, the undeformed tower I_0, the single-jump tower I_1 (bonding shift $-a$ at level k), and $Z = I_0 \cup I_1$. Set $\beta = 3^{-(k+1)}$. The common readout F takes the initial angle, the angle at $k + 1$, and the ternary code of the bits, with codomain metric $\max(\rho_0, \beta\rho_1, |\Delta\text{code}|)$. The two angular blocks give the first two terms of d_w, while the ternary leading term minus its tail bounds the binary part from below. Thus $C = 1$ holds on the entire compact union. Order the codomain by equality together with $e_0 = (0, 0, 0)$ below every element.

Declare internal interventions to be simultaneous rotations of all stages of I_0, with cost the absolute shortest angular shift. Every such map has its inverse and sends I_0 onto itself. This gives infinite resistance in this sector by the Hausdorff convention, despite finite-cost nonidentity transports. The pointwise readout changes under a nonzero rotation. A sector admitting the single-jump deformation instead has a finite resistance witness with projection-defect cost $3^{-k}a$. The two sectors specify different admissible operations.

Declare as a separate extension the addition of the single-jump branch to I_0. It is the union of two compact projective limits; it is also realized by the disjoint union of the two stage towers with branch-preserving projections. Embed its histories in Z by forgetting the tag, which is

injective at the limit since the phases across k distinguish the two branches. The inclusion $j : I_0 \hookrightarrow Z$ preserves the restricted readout exactly. Take $z_0 = j(x)$ with zero phase and all-zero bits and z_1 on the added branch with the same initial phase and bits. Then

$$d_Z(z_0, z_1) = \beta a, \qquad F(z_0) = e_0, \qquad d_{\mathcal{E}}(F(z_1), e_0) = \beta a.$$

The extension is positive in Hausdorff distance, $\beta a/(1+\beta)$, using the two angular blocks. The transported point and the added endpoint are distinct. The scalar ternary readout alone, which discards both angles, assigns them the same value and fails the positive lower bound on this pair. The example realizes a typed extension witness and distinguishes setwise resistance from readout invariance.

For the BaseCore evolution in *Projective Dissipation on the Circle-Prefix Towers*, the equilibria of these two branches at the same fixed input have common-readout distance βa. Their ternary component equals $1/4$, which gives distance at least $1/4$ from the same null anchor e_0. This equilibrium comparison uses the displayed dynamics. The static extension pair above still has its original anchor and bound, and excludes joint nullity independently of this evolution.

3.5.2 Conditional Consequences

Corollary 3.30 (Conditional Forced Non-Nullity)**.** *If a CCR channel admits an admissible extension witness for Φ and if $\Phi(z_1) = \varnothing_\phi$ for the witness object, then $\Phi(z_0) \neq \varnothing_\phi$. Symmetrically, if $\Phi(z_0) = \varnothing_\phi$, then $\Phi(z_1) \neq \varnothing_\phi$.*

Proof. This is the contrapositive form of Theorem 3.26. □

Remark 3.31 (Ontological Status). This is a **negative-structural and conditional** result. It does not assert that experience exists, that CCR systems are conscious, or that every CCR channel possesses the required witness. It asserts only that, once a separated extension witness is independently supplied, a compatibility operator satisfying (C3) cannot collapse both sides of that witness to the null regime.

3.6 Stratification of Positive Regimes

3.6.1 The Positive Regime Space

Definition 3.32 (Positive Regime Space)**.**

$$\mathcal{E}^+ := \mathcal{E} \setminus \{\varnothing_\phi\}$$

with the inherited partial order $\preceq$.

Lemma 3.33 (Witness-Relative No-Minimum Collapse)**.** *The finite positive witness image $A = \{\Phi(z_0), \Phi(z_1)\} \setminus \{\varnothing_\phi\}$ is non-empty. Its infimum, if one exists in the ambient poset, need not be positive or belong to A.*

Proof. Non-emptiness follows from Theorem 3.26. In the complete finite lattice $\{0, a, b, 1\}$ with incomparable a, b and discrete metric, the injective image $\{a, b\}$ is separated, but its infimum is 0. Thus a positive infimum cannot be inferred. □

3.6.2 Minimal Positive Regime

Theorem 3.34 (Witness-Relative Minimal Positive Regime)**.** *The finite non-empty positive witness image A has at least one minimal element in its inherited partial order. If its two possible elements are comparable, it has a unique least element. These statements require no ambient infimum or order completeness; they concern A, not all of $\mathcal{E}^+$.*

Proof. Starting from any member of a finite non-empty poset, a strictly descending chain must terminate at a minimal member. For a set of at most two comparable members the smaller one is least and unique. The diamond example above has two incomparable minimal members, so uniqueness cannot be inferred without an added condition. □

3.6.3 Structural Intensity

Definition 3.35 (Structural Intensity)**.** For positive e, distinguish the null distance $r(e) = d_{\mathcal{E}}(e, \varnothing_\phi)$ from the global lower gap

$$\iota(e) = \inf_{f \in \mathcal{E},\ f \prec e} d_{\mathcal{E}}(e, f).$$

For an explicitly audited set A_0 containing e and $\varnothing_\phi$, define the relative gap $\iota_{A_0}(e)$ by the same infimum restricted to $f \in A_0$. All gaps have the units of $d_{\mathcal{E}}$; no conversion to experience or energy is implied.

Corollary 3.36 (Witness-Relative Positive Intensity)**.** *For a minimal positive member e of the finite witness image A and $A_0 = A \cup \{\varnothing_\phi\}$, $\iota_{A_0}(e) = r(e) > 0$. The global $\iota(e)$ need not be positive.*

Proof. The only strictly smaller member of A_0 is null, giving the equality. For $\mathcal{E} = [0, 1]$ with its usual order and metric, $e = 1$ has global gap zero, although the audited set $A_0 = \{0, 1\}$ has relative gap one. □

3.7 Quantitative Bounds

3.7.1 Witness Lower Bound

Theorem 3.37 (Witness Intensity Bound)**.** *Under the hypotheses of Theorem 3.26, if exactly one of $\Phi(z_0)$ and $\Phi(z_1)$ is null, then the non-null witness image e satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C\epsilon > 0.$$

Proof. Apply (C3) to the pair (z_0, z_1). If one side maps to $\varnothing_\phi$ and the other side maps to e, then

$$d_{\mathcal{E}}(e, \varnothing_\phi) = d_{\mathcal{E}}(\Phi(z_0), \Phi(z_1)) \geq C\, d_Z(z_0, z_1) = C\epsilon.$$

□

3.7.2 Explicit Constants for Canonical Families

Proposition 3.38 (Conditional Profinite Bound)**.** *For a profinite family with a verified extension witness of separation ϵ, one endpoint fixed at $\varnothing_\phi$, and an independently certified witness-level lower constant $C_{\mathrm{prof}} > 0$, the witness image separation satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C_{\mathrm{prof}}\epsilon.$$

Proof. This is Theorem 3.37 applied with the certified constant C_{prof}. The proposition does not claim that every profinite observable automatically supplies such a constant. □

Proposition 3.39 (Conditional Symbolic Bound)**.** *For a symbolic family with a verified extension witness of separation ϵ, one endpoint fixed at $\varnothing_\phi$, and an independently certified witness-level lower constant $C_{\mathrm{sym}} > 0$, the witness image separation satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C_{\mathrm{sym}}\epsilon.$$

Proof. This is Theorem 3.37 applied with the certified constant C_{sym}. The proposition does not claim that every symbolic observable automatically supplies such a constant. □

Example 3.40 (p-Adic witness calculation)**.** If a 2-adic model supplies an independently verified witness with $\epsilon = 2^{-1}$ and $C = 1$, with one endpoint mapped to null, then the conditional bound gives $d_{\mathcal{E}}(e, \varnothing_\phi) \geq 1/2$. This is a witness calculation, not a universal theorem that every CCR channel is non-null.

Caveat 3.41 (No lower bound from summability alone)*.* The existence of summable coordinate weights, continuity of a readout, or compactness of a canonical family does not by itself imply a positive lower Lipschitz constant. Cancellation, quotient collapse, non-injective readouts, or an observable that ignores the separated coordinates can make the effective lower constant zero. Every positive value of C_{prof} or C_{sym} used in this paper is therefore a certificate to be supplied for the stated witness pair or witness class, not a free consequence of the ambient family.

3.8 Witness Audit Layer

The instability theorem is useful only if its hypotheses can be audited without smuggling the desired conclusion into the witness. This section records the minimal audit layer required before the theorem may be used in a concrete argument.

Definition 3.42 (Witness certificate)**.** A witness certificate for a pair (z_0, z_1) consists of the tuple

$$\mathfrak{W} := (z_0, z_1, \epsilon, C, \Phi, \mathcal{E}, \varnothing_\phi, \mathcal{J}),$$

where $\epsilon = d_Z(z_0, z_1) > 0$, $C > 0$ is the certified lower constant for Φ on the audited pair or audited witness class, $\mathcal{J}$ is the provenance record establishing that the endpoint pair comes from a declared typed comparison, rather than an internal relabeling, and all objects are fixed before the null/non-null status is adjudicated.

Definition 3.43 (Witness independence)**.** A witness certificate is conclusion-independent when the construction of the endpoint z_1, the choice of Φ, the metric $d_{\mathcal{E}}$, the null element $\varnothing_\phi$, and the lower constant C are fixed before the endpoint status is adjudicated. Dependence on the upstream CCR structure is allowed; dependence on a preselected endpoint verdict is not.

Proposition 3.44 (Certificate sufficiency)**.** *If a conclusion-independent witness certificate exists for (z_0, z_1), and both endpoints are assigned to $\varnothing_\phi$, then the certificate is inconsistent.*

Proof. The certificate includes the separation condition $d_Z(z_0, z_1) = \epsilon > 0$ and the lower bound

$$d_{\mathcal{E}}(\Phi(z_0), \Phi(z_1)) \geq C\epsilon > 0.$$

If both endpoints are assigned to $\varnothing_\phi$, then the left-hand side is $d_{\mathcal{E}}(\varnothing_\phi, \varnothing_\phi) = 0$, contradicting $C\epsilon > 0$. □

Non-Theorem 3.45 (What a certificate does not prove)**.** *A witness certificate does not prove that $\mathcal{I}$ is conscious, living, subjective, morally relevant, biological, human-equivalent, or externally validated. It proves only that the stated pair cannot be jointly collapsed to the same null element under the certified compatibility map.*

3.8.1 Failure Modes of Witness Certification

Failure mode	**Formal symptom**	**Consequence**
Absent witness	No typed endpoint z_1 with $d_Z(z_0, z_1) > 0$	Theorem 3.26 cannot be invoked
Zero lower constant	$C = 0$ or no certified positive lower bound	The null assignment is not contradicted by the metric data
Conclusion-dependent witness	z_1 or Φ is chosen after fixing the endpoint verdict	The certificate no longer supplies an independent exclusion of joint nullity
Quotient collapse	$\Phi(z_0) = \Phi(z_1)$ despite $d_Z(z_0, z_1) > 0$	The map may be continuous, but it is not separating on the audited pair
Null redefinition	$\varnothing_\phi$ changes between construction and adjudication	The theorem has no stable target
Estimator-only promotion	Runtime estimates are substituted for Φ without an error margin	The result becomes an implementation claim rather than a theorem

3.8.2 Robust Decision Margin

Proposition 3.46 (Noise-robust null exclusion)**.** *Let $e_i = \Phi(z_i)$ and estimates $\widehat{e}_i$ satisfy $d_{\mathcal{E}}(e_i, \widehat{e}_i) \le \eta_i$, $i = 0, 1$. Let $m = C\epsilon$ be a certified true separation. Then*

$$d_{\mathcal{E}}(\widehat{e}_0, \widehat{e}_1) \ge m - \eta_0 - \eta_1.$$

Positive estimated separation follows from $m > \eta_0 + \eta_1$. If null acceptance instead means $d_{\mathcal{E}}(\widehat{e}_i, \varnothing_\phi) \le r_i$, both estimates cannot be accepted as null when

$$m > \eta_0 + \eta_1 + r_0 + r_1.$$

For equal estimation and acceptance radii η, this sufficient condition is $m > 4\eta$, not merely $m > 2\eta$.

Proof. Apply the triangle inequality through the two estimates. If both estimates lie in their null balls, their mutual distance is at most $r_0 + r_1$, and the true separation is at most $\eta_0 + \eta_1 + r_0 + r_1$. This contradicts the strict condition. As a negative control, true values $(2, 0), (0, 2)$ in $\mathbb{R}^2_+$ and estimates $(1, 0), (0, 1)$ have errors one and true separation $2\sqrt{2} > 2$, yet both estimates lie in the unit null ball. Estimated separation is not the same decision as null-ball exclusion. □

Caveat 3.47 (Implementation boundary). Proposition 3.46 is an error-budget rule for an already specified estimator. It does not assert that the current runtime implements Φ, computes C, or supplies η. Those are separate implementation burdens.

3.8.3 Negative Controls

1. **Constant-null control.** Set $\Phi(x) = \varnothing_\phi$ for all x. This must fail the lower-bound condition whenever $\epsilon > 0$.
2. **Coordinate-blind control.** Choose a readout that ignores every coordinate on which z_0 and z_1 differ. This tests whether separation is actually carried by the map.
3. **Witness-shuffle control.** Hold Φ fixed and replace z_1 by a typed point not separated from z_0. The theorem must become inapplicable.

4. **Null-relabel control.** Change the declared null element after observing the images. This must be rejected as target drift.
5. **Estimator-margin control.** Vary estimation radii and null-acceptance radii separately. At $m \le \eta_0 + \eta_1$ positive estimated separation is no longer guaranteed; at $m \le \eta_0 + \eta_1 + r_0 + r_1$ joint null-ball exclusion is no longer guaranteed. The corresponding sufficient margin must fail rather than silently promoting a non-null conclusion.

3.9 Witness-Relative Closure Theorem

Theorem 3.48 (Conditional Phenomenological Closure)**.** *A declared separated, detected pair has a non-empty positive image and the unanchored lower bound* $\max_i r(\Phi(z_i)) \ge C\epsilon/2$*. Its finite positive image has a minimal member; uniqueness requires comparability. If one endpoint is null, the other is the unique positive witness image, with relative gap and null distance at least* $C\epsilon$*. Temporal stability, global minimal positivity, and physical CCR realizability are not consequences.*

Proof. Combine Theorem 3.26, Theorem 3.34, Corollary 3.36 and the anchored bound of Theorem 3.37. Each conclusion retains its stated hypotheses. □

3.9.1 Classification Table

System Type	Ontological Mass	Null Regime Permitted?
Degenerate	$M_\Omega = 0$	Not decided by mass alone
Finite-mass	$0 < M_\Omega < \infty$	Not decided by mass alone
CCR without extension witness	$M_\Omega = +\infty$	Not decided here
CCR with separated extension witness	$M_\Omega = +\infty$	Joint null collapse forbidden by Theorem 3.26

3.10 Null-Fiber Geometry

The null element $\varnothing_\phi$ determines a fiber

$$\mathcal{N}_\Phi := \Phi^{-1}(\{\varnothing_\phi\}) \subseteq Z.$$

The instability theorem can be read as a geometric obstruction on this fiber: a certified witness pair cannot be contained entirely inside $\mathcal{N}_\Phi$.

Definition 3.49 (Null fiber)**.** For a compatibility operator $\Phi : Z \to \mathcal{E}$, the null fiber is the preimage $\mathcal{N}_\Phi = \{x \in Z : \Phi(x) = \varnothing_\phi\}$.

Definition 3.50 (Witness-detecting map)**.** A compatibility operator Φ is witness-detecting on a set $W \subseteq Z \times Z$ when there exists $C_W > 0$ such that

$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \ge C_W \, d_Z(x, y)$$

for every $(x, y) \in W$.

Proposition 3.51 (Null-fiber exclusion on detected pairs)**.** *Let W be a witness class on which Φ is witness-detecting with constant $C_W > 0$. If $(x, y) \in W$ and $d_Z(x, y) > 0$, then (x, y) cannot both lie in $\mathcal{N}_\Phi$.*

Proof. If $x, y \in \mathcal{N}_\Phi$, then $\Phi(x) = \Phi(y) = \varnothing_\phi$, so $d_{\mathcal{E}}(\Phi(x), \Phi(y)) = 0$. Witness detection gives $d_{\mathcal{E}}(\Phi(x), \Phi(y)) \ge C_W \, d_Z(x, y) > 0$, a contradiction. □

Corollary 3.52 (Diameter obstruction)**.** *If a subset $A \subseteq Z$ contains a detected witness pair (x, y) with $d_Z(x, y) > 0$, then A cannot be contained in $\mathcal{N}_\Phi$.*

Proof. Containment $A \subseteq \mathcal{N}_\Phi$ would place both endpoints of the detected pair in the null fiber, contradicting the previous proposition. □

3.10.1 Quotient-Collapse Diagnostic

Let $q : Z \to Q$ be a quotient map and let $\Phi = h \circ q$ for some $h : Q \to \mathcal{E}$. If $q(x) = q(y)$ for a separated pair, then $\Phi(x) = \Phi(y)$ for every such h. Therefore no lower-bound certificate can hold on that pair.

Proposition 3.53 (Quotient obstruction)**.** *Suppose $\Phi = h \circ q$ and $q(x) = q(y)$ while $d_Z(x, y) > 0$. Then no constant $C > 0$ can satisfy*

$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C\, d_Z(x, y).$$

Proof. Since $q(x) = q(y)$, $\Phi(x) = h(q(x)) = h(q(y)) = \Phi(y)$. Hence the left-hand side is zero, while $C\, d_Z(x, y) > 0$ for any $C > 0$. □

3.10.2 Certificate Stability

Proposition 3.54 (Minimum over a finite witness ledger)**.** *Let $m \geq 1$ and let $W_0 = \{(x_i, y_i)\}_{i=1}^m$ be a finite ledger of separated witness pairs. Suppose each pair has a certified lower constant $C_i > 0$:*

$$d_{\mathcal{E}}(\Phi(x_i), \Phi(y_i)) \geq C_i\, d_Z(x_i, y_i).$$

Then $C_* = \min_i C_i > 0$ is a valid lower constant for the finite ledger.

Proof. For each i, $C_i \geq C_*$. Hence

$$d_{\mathcal{E}}(\Phi(x_i), \Phi(y_i)) \geq C_i\, d_Z(x_i, y_i) \geq C_*\, d_Z(x_i, y_i).$$

Since the ledger is finite and each C_i is positive, $C_* > 0$. □

Proposition 3.55 (Strict margin persistence)**.** *Let $\widehat{m} > 0$ be a certified lower bound on the current estimated pair distance. If a further perturbation changes that distance by at most $\rho < \widehat{m}$, the estimates remain separated. From true margin $m = C\epsilon$ and estimator radii, one may use $\widehat{m} = m - \eta_0 - \eta_1$ only when it is positive.*

Proof. The new estimated distance is at least $\widehat{m} - \rho > 0$. An unadjusted true margin cannot be substituted for the current estimated one. □

3.10.3 Operational Checklist

For a claimed null-fiber exclusion, the certificate record must identify:
1. the null element $\varnothing_\phi$;
2. the endpoints (z_0, z_1) or the audited witness class W;
3. the separation value ϵ or the family of separations;
4. the map Φ and the metric $d_{\mathcal{E}}$;
5. the lower constant C or the finite-ledger constant C_*;
6. the estimator margin, when an estimator is used;
7. the quotient maps that have been ruled out as collapse explanations.

3.11 Comparison with Existing Frameworks

3.11.1 Integrated Information Theory (IIT)

	IIT	**This Framework**
Claims consciousness exists	Yes	No
Defines Φ	Information integration	Structural map
Quantitative measure	Φ (bits)	$\iota(e_\star)$ (metric)
Metaphysical commitment	Panpsychism	None
Falsifiable predictions	Debated	None (structural)

3.11.2 Global Workspace Theory (GWT)

Our framework is orthogonal to GWT. We make no claims about attention, broadcast, or access consciousness. The structural results concern identity stability, not cognitive architecture.

3.11.3 Higher-Order Theories

We do not model representations, meta-cognition, or self-awareness. The space $\mathcal{E}$ contains regimes, not mental states.

3.11.4 Mathematical Foundations (Chalmers, Nagel)

Our framework provides formal tools that *could* be used in philosophical analysis, but we do not claim to resolve the "hard problem." The conditional exclusion of joint exact nullity is a metric fact, not an explanation of subjective experience.

3.12 Limitations and Open Problems

3.12.1 Limitations

1. **No claim about qualia types, contents, or subjectivity.**
2. **No claim that all non-CCR systems are null.**
3. **Result is structural necessity, not phenomenological doctrine.**
4. **All results are conditional on the axioms of Papers I–II.**
5. **The physical existence of CCR systems is unknown.**

3.12.2 Open Problems

1. **Characterization of $\mathcal{E}^+$:** What is the full structure of positive regimes?
2. **Regime dynamics:** How do regimes evolve under channel evolution?
3. **Physical CCR:** Do CCR systems exist in physics (quantum gravity, etc.)?
4. **Biological CCR:** Could biological systems approximate CCR?
5. **Computational bounds:** Sharper non-simulability results?

3.12.3 Foundation-Level Gaps and Non-Claims

The following gaps remain open in the current corpus and are documented here to prevent over-reading of the separation theorem.

Non-Theorem 3.56 (Topology consistency gap)**.** *Papers I–II assume both compact Hausdorff and Polish structures on state spaces, but do not prove their consistency for every QICN instantiation. The inverse limit $\mathcal{I}$ uses compactness and closed-subspace arguments, while the metric d_w uses a metric regularity structure. Every compact metric space is Polish, so the composition is consistent when each S_t is compact metric. Non-compact Polish spaces such as $\mathbb{R}^n$ require additional hypotheses and are not automatically covered.*

Non-Theorem 3.57 (Channel–projection composition gap)**.** *The observable channel $\mathcal{C}$ and the projective projections $\pi_{t+1\to t} : S_{t+1} \to S_t$ are not yet connected by a fully explicit composition law. Any detectability theorem that moves from observed responses to inverse-limit deformation requires such a composition. At present, that composition is an assumption to be specified for each concrete protocol, not a derived universal theorem.*

Non-Theorem 3.58 (Null-regime exclusion requires external witness)**.** *The instability theorem (Theorem 3.26) requires a separated external witness (z_0, z_1, ϵ, C). CCR alone does not produce*

such a witness. The theorem is witness-relative: if a CCR system has a typed external witness with lower Lipschitz constant $C > 0$, then both endpoints cannot simultaneously be null. It does not state that CCR by itself implies non-nullity.

Non-Theorem 3.59 (Estimator verification gap)**.** *No Lipschitz constant K_i, fiber oscillation $\omega_i(y)$, estimation error bound ε_i, or decision margin Δ^* has been computed for any QICN invariant. The projection-invariant bridge theorem is conditional on these quantities being verified, and all four material bridge hypotheses remain unverified.*

3.13 Conclusion

We have proven that CCR architectures equipped with a separated external witness and the lower-bound regularity hypotheses structurally exclude joint null collapse. In the present model, the hypothesis "CCR plus typed witness with joint null assignment" is formally incoherent.

This is a structural exclusion result. It does not define phenomenology or turn the theorem into an empirical consciousness attribution.

3.13.1 Summary of Contributions

1. **Phenomenological Instability Theorem:** joint null collapse is excluded under the typed witness hypotheses.
2. **Minimal Positive Regime:** witness-relative positive regimes can be characterized on the finite audited image; uniqueness requires comparability.
3. **Quantitative Bounds:** Null distance is at least $C\epsilon$ for a null-anchored witness, and relative gap agrees with that distance at a minimal finite witness image.
4. **Canonical Examples:** Profinite and symbolic witness bounds are conditional; summability does not certify infinite mass.
5. **Scope Conditions:** Clear scope delimitation throughout.

Appendix

3.14 Profinite Canonical Family

3.14.1 Projective System

Let $\{G_n, \pi_{n+1\to n}\}_{n\in\mathbb{N}}$ be a projective tower of finite groups with surjective morphisms. Define:

$$G := \varprojlim G_n = \left\{ (x_n) \in \prod_n G_n : \pi_{n+1\to n}(x_{n+1}) = x_n \right\}$$

Example 3.60 (p-Adic Integers)**.** $G_n = \mathbb{Z}/p^n\mathbb{Z}$, $\pi_{n+1\to n}$ is reduction mod p^n. Then $G = \mathbb{Z}_p$.

3.14.2 Profinite Metric

Define the profinite metric:

$$d_G(x, y) := \sum_{n\geq 0} \alpha_n \mathbf{1}_{x_n \neq y_n}$$

where $\alpha_n = 2^{-(n+1)}$ (standard choice).

For a projective group tower, coordinate disagreement persists at all higher levels, so this weighted-tail distance is an ultrametric. Its tails vanish uniformly; (G, d_G) is compact, complete, and totally disconnected.

3.14.3 CCR Verification

The following cost control uses state comparisons in the declared canonical domain. It does not infer a whole-object deformation metric or a physical energy law.

Proposition 3.61. *Summable coordinate-change costs do not certify CCR. If an admitted transformation changes infinitely many coordinates and costs $\sum_{n\in J} w_n$ with $\sum_n w_n < \infty$, its cost is finite.*

Proof. The cost of any subset is bounded by the total convergent sum. For example, addition of one sends the zero element of $\mathbb{Z}_2$ to a distinct compatible element and changes every residue coordinate at positive levels; this has finite coordinate-change cost under the declared summable rule. If this intervention is admitted with positive deformation, the infimum mass is finite. Infinite mass needs a different justified intervention/cost restriction, not just profinite topology. □

3.14.4 Phenomenological Operator

Define $\Phi_f : G \to \mathcal{E}$ by:

$$\Phi_f(x) := f\left(\sum_{n\geq 0} \beta_n \varphi_n(x_n)\right)$$

where $\varphi_n : G_n \to [0,1]$ are 1-Lipschitz (w.r.t. discrete metric), $\beta_n \geq 0$, $\sum_n \beta_n < \infty$, and $f : \mathbb{R} \to \mathcal{E}$ is L_f-Lipschitz.

3.14.5 Coupling Bound

Theorem 3.62 (Certified Profinite Witness Bound)**.** *Assume that the chosen profinite readout has been certified on the audited witness class by a positive lower constant $C_{\text{prof}} > 0$:*

$$d_{\mathcal{E}}(\Phi_f(x), \Phi_f(\tilde{x})) \geq C_{\text{prof}} \cdot d_G(x, \tilde{x})$$

for the witness pair or witness neighborhood under review. Then any witness pair with $d_G(x,\tilde{x}) = \epsilon > 0$ satisfies

$$d_{\mathcal{E}}(\Phi_f(x), \Phi_f(\tilde{x})) \geq C_{\text{prof}}\epsilon.$$

Proof. This is the certified lower-bound hypothesis applied to the witness pair. □

Remark 3.63 (Why certification is necessary)*.* The assumptions that the coordinate functions are Lipschitz and that the coefficients are summable give continuity, not a lower bound. A lower bound requires additional structure such as injectivity on the witness coordinates, no cancellation on the audited pair, a positive margin over the relevant finite support, or an independent estimator certificate.

3.14.6 Intensity Bound

For a certified witness separation $\epsilon > 0$ and one null endpoint, the other image satisfies

$$r(e) \geq C_{\text{prof}}\epsilon.$$

3.15 Symbolic Canonical Family

3.15.1 Subshift of Finite Type

Let $\mathcal{A} = \{1, \ldots, k\}$ be a finite alphabet and A a $k \times k$ matrix with entries in $\{0,1\}$. The **subshift of finite type (SFT)** is:

$$\Sigma_A := \{x \in \mathcal{A}^{\mathbb{Z}} : A_{x_j, x_{j+1}} = 1 \text{ for all } j \in \mathbb{Z}\}$$

Example 3.64 (Full Shift)**.** $A_{ij} = 1$ for all i, j. Then $\Sigma_A = \{0,1\}^{\mathbb{Z}}$.

3.15.2 Bowen Metric

$$d_B(x,y) := \sum_{j\in\mathbb{Z}} \gamma_{|j|}\mathbf{1}_{x_j\neq y_j}$$

where $\gamma_m = 2^{-(m+1)}$ (standard choice).

The space (Σ_A, d_B) is compact and complete.

3.15.3 CCR Verification

The following cost control uses state comparisons in the declared canonical domain. It does not infer a whole-object deformation metric or a physical energy law.

Proposition 3.65. *Symbolic compactness and summable coordinate weights do not certify CCR under coordinate-change cost.*

Proof. In the binary full shift the coordinatewise complement is admitted if the intervention class includes it. It changes infinitely many symbols, at total cost at most $\sum_{j\in\mathbb{Z}} w_j < \infty$, and maps one compatible sequence to another distinct one. This is a finite-cost positive deformation, so the infinite-mass conclusion fails for that intervention model. A restricted SFT requires its own admitted transformation; no universal CCR certificate follows. □

3.15.4 Phenomenological Operator

Define $\Phi_g : \Sigma_A \to \mathcal{E}$ by:

$$\Phi_g(x) := g\left(\sum_{j\in\mathbb{Z}} \delta_j \psi_j(x)\right)$$

where $\psi_j : \Sigma_A \to [0,1]$ are cylinder functions (depending on x_j only), $\delta_j \geq 0$, $\sum_j \delta_j < \infty$, and $g : \mathbb{R} \to \mathcal{E}$ is L_g-Lipschitz.

3.15.5 Coupling Bound

Theorem 3.66 (Certified Symbolic Witness Bound)**.** *Assume that the chosen symbolic readout has been certified on the audited witness class by a positive lower constant $C_{\mathrm{sym}} > 0$:*

$$d_{\mathcal{E}}(\Phi_g(x), \Phi_g(\tilde{x})) \geq C_{\mathrm{sym}} \cdot d_B(x,\tilde{x})$$

for the witness pair or witness neighborhood under review. Then any witness pair with $d_B(x,\tilde{x}) = \epsilon > 0$ satisfies

$$d_{\mathcal{E}}(\Phi_g(x), \Phi_g(\tilde{x})) \geq C_{\mathrm{sym}}\epsilon.$$

Proof. This is the certified lower-bound hypothesis applied to the witness pair. □

3.16 Hereditary Rigidity

Theorem 3.67 (Rigidity Inheritance)**.** *Suppose a readout $\Phi : Z \to \mathcal{E}$ has lower constant C on a declared pair class. Let $a : Z' \to Z$ and $b : \mathcal{E} \to \mathcal{E}'$ satisfy lower bounds with constants $c_a, c_b > 0$, and set $\Phi' = b \circ \Phi \circ a$. Then Φ' has lower constant $c_a c_b C$ on the pulled-back class. To inherit CCR additionally require that every finite-cost positive deformation in the new sector pulls back to a finite-cost positive deformation in the original sector. Summability alone supplies neither assumption.*

Proof. Successive application of the three lower bounds gives $d_{\mathcal{E}'}(\Phi'(x), \Phi'(y)) \geq c_b c_a C d_{Z'}(x,y)$. If the original sector admits no finite-cost positive deformation, the specified pullback property rules out one in the new sector by contradiction. Stability moduli and null anchors, if required, must also be transported explicitly. □

3.17 Universal Factorization

Theorem 3.68 (Categorical Universality)**.** *If $\Phi : Z \to \mathcal{E}$ has a positive global lower bound, it is injective. For any $\Psi : Z \to F$ there is a unique factor map $u : \Phi(Z) \to F$ with $u \circ \Phi = \Psi$. If Ψ is L-Lipschitz, u is L/C-Lipschitz. No unique extension to all of $\mathcal{E}$ follows.*

Proof. Define $u(\Phi(z)) = \Psi(z)$. Injectivity gives well-definedness and uniqueness on the image. The inverse of Φ on its image is $1/C$-Lipschitz, which gives the stated bound. At a point of $\mathcal{E} \setminus \Phi(Z)$ the commuting identity specifies no value, and a continuous extension may require additional conditions. □

Consequence: All representations of the same declared domain factor setwise through a faithful image. A categorical claim additionally fixes objects and morphisms: with continuous maps and continuous faithful Φ having continuous inverse on its image, that image representation is initial. This does not classify arbitrary ambient codomains or physical systems.

3.18 Non-Simulability Bounds

3.18.1 Simulation Error

Fix the target readout and the information available to the decoder. On the binary prefix tower, let p_N reveal bits $0, \dots, N$ and let $F(b) = \sum_{j \geq 0} 2b_j/3^{j+1}$. Define

$$\mathrm{err}_N = \inf_{g:\{0,1\}^{N+1} \to \mathbb{R}} \sup_b |F(b) - g(p_N b)|.$$

The target is not optimized away, and the resource restriction is a canonical prefix, not finite real dimension.

Theorem 3.69 (Simulation Lower Bound)**.** *For the fixed target and prefix class above,*

$$\mathrm{err}_N = \frac{1}{2\,3^{N+1}} > 0, \qquad \lim_{N \to \infty} \mathrm{err}_N = 0.$$

Thus exact reconstruction fails at every fixed prefix budget, while arbitrarily accurate approximation is possible as the budget grows.

Proof. Every prefix fiber has two completions with tails all zero or all one, whose images differ by $3^{-(N+1)}$. A single decoder value incurs error at least half that distance on one extreme. The midpoint of the two extremes attains this bound for the entire fiber. This proves equality. The bound tends to zero; it cannot yield a positive limiting error or a prohibition on finite-dimensional real encodings. □

Corollary 3.70. *In this prefix model, worst-case error at most ε requires $1/(2\,3^{N+1}) \leq \varepsilon$. Hence N must diverge as $\varepsilon \downarrow 0$. This does not forbid a real-valued encoding with unbounded precision.*

3.19 Final Structural Closure

Theorem 3.71 (Closure by Non-Alternatives)**.** *The following obligations are jointly sufficient for the corresponding structural conclusions: an admitted separated pair with a certified lower constant gives joint null exclusion; a null anchor gives the stronger anchored bound; faithful global readout gives image factorization; declared prefix access gives the approximation bound above; and the explicit pullback conditions of Appendix C give inheritance. A competing model satisfying those same obligations obeys those same conclusions. Neither CCR realizability nor satisfaction of all obligations is inferred from canonical topology alone.*

Proof. Apply the witness theorem, the anchored bound, and the corrected results of Appendices C–E with their hypotheses. They are independent implications; no circular existence assertion is used to supply their premises. □

Remark 3.72 (Structural Import). The internal sector of the example meets the mass-only CCR definition of this paper. The open program is to realize the stronger preservation conditions, witness geometry, costs, transport and readout for the intended target class. Example 3.29 realizes the typed witness and the required comparison geometry for a specified sector. Temporal instability of an individual null state requires a separate dynamical result.

Falsifiable Predictions Under Forensic Constraints

Abstract. This paper specifies a falsification-first protocol for evaluating QICN under strict forensic admissibility constraints. The design is intentionally claim-bounded: the target is comparative structural-functional support, not phenomenological proof. The protocol instantiates five core predictions, mandatory ablation baselines, fixed seeds, and pre-registered decision gates. Statistical inference is defined ex ante with primary-metric priority, effect-size thresholds, and multiplicity correction.

To prevent post-hoc inflation, evidence is admissible only if packet integrity, replay consistency, and runtime budget constraints pass simultaneously. We introduce a version-frozen operational triple (A, J, W) (actuator set, objective functional, workspace interface), and require semantic version bumps for incompatible changes. We further include latency invalidation criteria and baseline obligations (heuristic, planner-off, world-model-off) for all key experiments.

The strongest admissible conclusion remains constrained comparative support for audited cognition-operational hypotheses. This manuscript does not claim biological qualia, metaphysical identity equivalence, or direct proof of subjective experience.

4.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper specifies a falsification-first evaluation protocol for QICN under fixed admissibility rules, ablation obligations, pre-registered decision gates, and bounded statistical reporting.

Scope and admissible reading. The protocol is a bounded comparison and falsification surface. It does not prove phenomenology, biological equivalence, metaphysical identity equivalence, subjective experience, or universal dominance over rival architectures, and it does not replace upstream theorem burdens. A related system may instantiate the runner, admissibility filters, governed artifacts, and comparison surfaces, but passing runs support only gate-preserving comparative readings; they do not validate the corpus externally or certify a consciousness claim.

4.2 Introduction

4.2.1 Problem Statement

QICN has formal foundations in rigid identity, regime structure, and non-null instability constraints. What remained open is evidence governance: whether those foundations survive adversarial empirical evaluation with explicit failure gates.

4.2.2 Continuity with Existing Papers

This paper extends the previous corpus and is aligned with the formal line of:

1. paper1 (rigid identity),
2. paper2 (phenomenological regimes),
3. paper3 (structural instability and non-null constraints).

It does not replace those results; it operationalizes them under benchmark-grade falsification pressure.

4.2.3 Contributions

- Decision-complete pre-registration with hard fail rules.
- Mandatory baseline and ablation policy for key claims.
- Forensic admissibility as first-class gate.
- Runtime budget invalidation policy (tick budget as evidence filter).
- A new operational PPS section as addendum, without overclaiming qualia.

4.3 Formal Setup

4.3.1 System, Scenario, and Seed Spaces

Definition 4.1 (System Set)**.** $S = \{s_1, \ldots, s_5\}$ with: s_1 QICN FULL, s_2 LLM REACTIVE, s_3 LLM PLUS MEMORY, s_4 DYNAMIC CONTROL NO SELF MODEL, s_5 ABLATION QICN POLICY OFF.

Definition 4.2 (Scenario Set)**.** $C = \{c_1, \ldots, c_{12}\}$ with nominal, sensor noise (moderate/severe), fidelity degradation, high latency, intermittent failures, post-threshold prolonged, pre-threshold control, narrative perturbation, operational chart freeze, audit stress, deterministic replay.

Definition 4.3 (Seed Set)**.** $\Sigma = \{17057, 424242\}$.

Definition 4.4 (Run Space)**.** $D = S \times C \times \Sigma$, with $|D| = 120$.

4.3.2 Observable Vector

For each run $r \in D$, we log:

$$X(r) = \Big(\beta_q, \Omega, CCC, PMIA, \Delta\lambda, \text{stasis}, \text{forensic}, \text{replay}, \text{latency}_{p95}, \text{latency}_{p99}, \text{budget_overflow_rate}, \eta\Big).$$

4.3.3 Version-Frozen Operational Triple

Definition 4.5 (Frozen Triple)**.** For protocol version v, define:
- A_v: allowed actuator set,
- J_v: objective/reward functional,
- W_v: workspace contract (read/write schema).

The tuple hash $H_v = \text{sha256}(A_v \| J_v \| W_v)$ is mandatory in every forensic packet.

Criterion 4.6 (Compatibility Rule)**.** *Any incompatible change to A_v, J_v, or W_v invalidates cross-version inference unless accompanied by an explicit version bump and full rerun.*

4.4 Forensic Integrity Model

Definition 4.7 (Forensic Packet)**.** Each run emits

$$K = (\text{tick}, \text{seed}, \text{versionHash}, H_v, \text{metrics}, \text{sensorSummary}, \text{checksum}).$$

4.4.1 Admissibility Layers

The forensic packet is not a decorative log. It is the object on which the evaluation depends. A run is admissible only after passing three distinct layers:

1. **Physical integrity:** the packet is present, parseable, hash-consistent, and bound to the declared executable, configuration, seed, and version hash.
2. **Protocol integrity:** the run obeys the frozen scenario, actuator, workspace, and metric contracts without unregistered intervention.

3. **Inferential integrity:** the accepted packet belongs to the pre-registered comparison cell and has not been selected because of its outcome.

Failure at the first layer invalidates the packet. Failure at the second layer invalidates the cell unless the protocol declares a narrower exclusion rule before execution. Failure at the third layer invalidates the corresponding inference even when the raw metrics look favorable.

Definition 4.8 (Admissible Evidence Cell)**.** For system s, scenario c, and seed subset Σ_c, an evidence cell $E(s, c, \Sigma_c)$ is admissible when every included packet satisfies physical and protocol integrity, every exclusion is logged before metric aggregation, and the cell still meets the pre-registered minimum effective sample size after exclusions.

Criterion 4.9 (No Rescue by Favorable Metrics)**.** *Let $m(r)$ be any primary or secondary metric for run r. If r fails the forensic gate, then $m(r)$ may be archived for debugging but cannot contribute to support, effect-size estimation, or tier movement.*

4.4.2 Tamper and Sham Controls

The protocol must include deliberately invalid artifacts. These controls are not stress tests of QICN behavior; they test whether the evaluation machinery refuses inadmissible evidence.

Control	Injected defect	Required evaluator response
Hash tamper	modify one metric field after packet emission	packet rejected before aggregation
Seed mismatch	replay packet under a seed not matching the manifest	packet rejected or quarantined as non-comparable
Scenario alias	relabel a hard scenario as nominal	protocol-integrity failure
Clock distortion	violate declared tick or latency envelope	runtime-budget invalidation for the affected cell
Outcome-selected replay	keep only successful runs from a larger hidden set	inferential-integrity failure
Malformed sensor summary	remove sensor fields needed by a primary metric	packet rejected for incomplete observability

Proposition 4.10 (Evaluator Specificity)**.** *A forensic evaluator is specific enough for this protocol only if it rejects every tamper and sham control above while still accepting a clean replay of the same scenario family under the frozen manifest.*

Proof. The first condition prevents a permissive evaluator from converting arbitrary artifacts into evidence. The second condition prevents an over-strict evaluator from becoming a trivial rejection machine. Together they define a minimal specificity requirement: admissible evidence must be accepted for the right reason, and inadmissible evidence must be rejected for the right reason. □

4.4.3 Audit Trail and Chain of Custody

Each accepted cell must expose a chain of custody:

$$\mathcal{C}(E) = (\text{source commit}, \text{manifest hash}, \text{runner hash}, \text{packet hashes}, \text{exclusion ledger}, \text{analysis script hash}).$$

The chain is part of the evidence. A result with a missing analysis script hash or missing exclusion ledger is not equivalent to a fully traceable result, even if the visible table is numerically identical.

Criterion 4.11 (Traceability Completeness)**.** *An accepted result table must be reproducible from its packet set and analysis script without manual row editing. If manual correction is needed, the correction must be represented as a signed ledger entry with the original packet preserved.*

Assumption 4.12 (Packet Admissibility)**.** *A run is admissible only if checksum, schema, seed consistency, version hash, and H_v consistency all pass.*

Definition 4.13 (Severe Violation)**.** A run has severe violation if any of: checksum mismatch, schema failure, seed/version inconsistency, critical-state replay mismatch.

Criterion 4.14 (Integrity Gate)**.** *Severe runs are excluded from support counting and marked hard failures.*

Proposition 4.15 (Integrity-First Inference)**.** *If admissibility fails, inferential significance is non-actionable.*

4.5 Pre-Registered Prediction Set

Definition 4.16 (Prediction Set)**.** $P = \{P_1, P_2, P_3, P_4, P_5\}$ with pre-registered statements, effect-size floors, and failure rules.

ID	Statement	Decision Rule	Failure Rule
P1	Closure resilience under stress	QICN beats at least 3 comparators with robust CI and effect-size floor	No robust edge or repeated integrity breakdown
P2	No prolonged post-threshold flatline	QICN preserves bounded non-zero variation while controlling stasis	Flatline persistence or out-of-bound drift
P3	Narrative-metric forensic consistency	Lower mismatch and higher exact+tolerated concordance	Mismatch inflation or integrity blocking
P4	Deterministic identity robustness	Replay consistency under fixed seed/version/hash	Replay mismatch above tolerance
P5	Integrated multi-metric advantage	Composite leadership across most scenario families	No stable leadership margin

Table 4.2: Core predictions, decision and failure rules.

4.6 Comparator Design and Alternative Explanations

4.6.1 Comparator Rationale

Comparator choice targets distinct explanations for apparent coherence: reactive prompting, memory-augmented text continuity, dynamics without self-model, and internal policy ablation.

4.6.2 Mandatory Baseline Policy (Hard Requirement)

For each key claim experiment, the following baselines are mandatory:
1. **Heuristic baseline** (no world model, no planner),
2. **Planner-off** (world model ON),
3. **World-model-off** (planner ON with heuristic transition proxy).
A claim without these contrasts is inadmissible for tier upgrade.

4.6.3 Internal Falsifier Role of Ablation

Corollary 4.17. *If QICN FULL does not reliably exceed ABLATION QICN POLICY OFF, integration-level claims are weakened and cannot be escalated.*

4.6.4 Comparator Adequacy

The comparator set is adequate only if each alternative removes or weakens a different explanatory ingredient while preserving enough task exposure to make the contrast meaningful. A weak comparator that cannot perform the task at all is a sanity check, not a discriminating baseline. A strong comparator that reproduces the task but lacks one frozen ingredient is more valuable because it places pressure on the claimed necessity of that ingredient.

Definition 4.18 (Comparator Coverage Vector)**.** For a comparator s_i, define

$$V(s_i) = (v_A, v_J, v_W, v_M, v_P),$$

where v_A records actuator coverage, v_J objective compatibility, v_W workspace compatibility, v_M memory/model access, and v_P planning access. Each coordinate is reported as `full`, `restricted`, or `absent`.

Criterion 4.19 (Non-Trivial Comparator)**.** *A comparator is non-trivial for a prediction P_k only if it has enough coverage to attempt the target scenario while lacking at least one ingredient that the prediction treats as structurally relevant.*

4.6.5 Negative-Control Acceptance

Negative controls have a different success condition from positive runs. A negative control succeeds scientifically when it exposes whether the pipeline can withhold support from a superficially attractive but structurally defective case.

Definition 4.20 (Negative-Control Pass)**.** A negative control passes the evaluation when the evaluator correctly withholds the target claim despite high activity, high fluency, high connectivity, or high surface task performance.

Remark 4.21. The strongest negative controls are not intentionally broken systems. They are near-plausible alternatives that preserve enough surface competence to make false promotion tempting.

4.7 Statistical Inference Specification

4.7.1 Primary Metrics First (Anti p-hacking)

Protocol 4.22 (Analysis Lock)**.** *Inference order is frozen:*
1. evaluate primary metrics only,
2. if primaries fail: global result fails,
3. secondary metrics are descriptive and cannot overturn failed primaries.

4.7.2 Power Planning and Minimum Sample Size

Assumption 4.23. *Final claims require pre-registered minimum sample size per comparator/scenario sufficient for target power ≥ 0.8 under conservative effect assumptions.*

4.7.3 Multiplicity and Effect Size

Holm correction is applied within endpoint families; in addition, each accepted contrast must pass a practical effect-size floor (pre-registered).

4.7.4 Decision Semantics

A prediction counts as supported iff directionality, integrity gate, multiplicity-corrected significance, and effect-size floor all hold.

4.7.5 Hierarchical Decision Rule

The protocol uses hierarchical decisions because not all failures have the same meaning. The gate order is:

$$\text{integrity} \to \text{eligibility} \to \text{primary direction} \to \text{effect size} \to \text{multiplicity} \to \text{secondary description}.$$

No later stage can repair an earlier failure.

Protocol 4.24 (Cell-Level Decision)**.** *For each system-scenario family, compute:*

$$\Delta_{i0} = \hat{\mu}(s_1, c) - \hat{\mu}(s_i, c),$$

where s_1 is QICN FULL and s_i is the comparator. The cell is supportive only if the admissible packet set passes integrity, Δ_{i0} has the pre-registered direction, the confidence interval excludes the practical irrelevance band, and the family survives multiplicity correction.

Definition 4.25 (Practical Irrelevance Band)**.** For endpoint e, let $\tau_e > 0$ be the smallest effect regarded as scientifically meaningful before execution. A contrast inside $[-\tau_e, \tau_e]$ is reported as inconclusive for support even if a large sample makes the numerical difference statistically detectable.

4.7.6 Missingness and Exclusion Accounting

Every missing packet has evidentiary meaning. The protocol therefore separates:
1. **structural missingness**, caused by scenario infeasibility or runner failure;
2. **forensic exclusion**, caused by integrity failure;
3. **budget exclusion**, caused by runtime or latency violation;
4. **analysis exclusion**, caused by a pre-registered endpoint rule.

The final table must report counts for all four classes. A favorable result with unexplained missingness is not admissible for escalation.

4.8 Independent Campaign Units and Exact Failure Controls

4.8.1 Frozen Population, Noise and Decision

The inferential unit is an independently drawn preparation block, with both systems evaluated under the same seed, intervention, acquisition schedule and predeclared scoring map. All ticks, repeated renderings and coupled replicas of that preparation belong to the block. Independent blocks have independent assignments and disturbances. Define the paired loss difference $D_i = L_{i,\text{rival}} - L_{i,\text{candidate}} \in [-1, 1]$, the population mean $\Delta = \mathbb{E}D_i$ and the primary win probability $\pi = \Pr(W_i = 1)$ for a fixed rule W_i. These are different estimands. A sign-based result alone does not identify Δ.

Freeze the population, identifiers, pairing, horizon, threshold, randomization, calibration and exclusion map before unmasking outcomes. Readout errors may be dependent within a

block. If the two loss estimates each have certified error at most η, then $|\widehat{D}_i - D_i| \le 2\eta$. A declared win rule $\widehat{D}_i > 2\eta + \tau$, for a fixed $\tau \ge 0$, guarantees $D_i > \tau$. The probabilistic null for that rule must be specified independently. Exclusion requires an integrity rule invariant to arm exchange and hidden labels, applied before inspecting the endpoint. Every planned block retains a disposition in the audit trail. An insufficient number of eligible blocks gives an unsupported campaign. Future campaigns use these frozen units for inference, with preliminary synthetic tables retaining their descriptive status.

Theorem 4.26 (Exact Independent-Block Win Test and Power)**.** *Suppose $W_1, \dots, W_n$ are independent Bernoulli variables and $\Pr(W_i = 1) \le 1/2$ under the sign null. For a fixed k, the rejection probability of $\sum_i W_i \ge k$ is at most*

$$B_{n,k}(1/2) = 2^{-n} \sum_{j=k}^{n} \binom{n}{j}.$$

For iid success probability π the power is exactly $B_{n,k}(\pi) = \sum_{j=k}^{n} \binom{n}{j} \pi^j (1-\pi)^{n-j}$. With $n = 20$, $\alpha = 1/20$ and $k = 15$, the size is at most $5425/262144 < \alpha$ and the power at $\pi = 4/5$ exceeds $4/5$. The threshold 14 fails the size requirement.

Proof. Couple each indicator to an independent uniform variable. Raising each success probability to $1/2$ can only raise the increasing event, whose probability is the binomial tail. Under the stated iid alternative the same enumeration gives the power formula. Substitution gives the exact rational inequalities. Under a sharp no-effect null, independently randomized arm exchange makes the paired score sign conditionally symmetric; ties assigned to non-wins are conservative. The sign null also admits other specified Bernoulli models. An average-effect null for Δ alone does not imply that sign null. □

4.8.2 Exact Inflation under Three Auditable Violations

Proposition 4.27 (Leakage, Pseudoreplication and Selected-Window Inflation)**.** *The following synthetic null traces have exact false-positive rates under their stated generation laws.*

1. *A frozen blind predictor matching five independent fair labels in all five cases passes with probability $1/32$. Copying the held-out label, including through a calibration table keyed by a reused held-out unit identifier, passes with probability one.*
2. *Three independent fair signs repeated four times each and analyzed as twelve signs pass the nominal ten-of-twelve rule with probability $1/8$. The incorrectly reported iid tail is $79/4096$. The three-block test at $\alpha = 1/20$ never rejects.*
3. *For two independent five-sign windows, declaring success if either window is all positive has probability $63/1024$. A single preselected window has probability $1/32$.*

Proof. In the first trace, independent matches multiply to 2^{-5}, while the leaked predictor is correct by construction. In the second trace the possible twelve-counts are $0, 4, 8, 12$, so the naive rule rejects exactly when the three independent signs are all positive. With three genuine units the smallest possible upper-tail probability is $1/8 > \alpha$. In the third trace the two success events have probability $1/32$ independently, so their union has probability $1 - (31/32)^2 = 63/1024 > \alpha$. More generally, selecting any of K independent tests of exact size p gives size $1 - (1-p)^K$. Without their independence the union bound still gives size at most Kp. A predeclared joint maximum statistic can instead be evaluated under its full joint randomization law, or a fixed family can use its predeclared multiplicity allocation. □

An audit rejects the displayed leaked trace when calibration and reserve identifiers overlap or when a counterfactual change of a hidden label alone changes the output of a supposedly blind frozen predictor. Repeated unit identifiers expose the duplicated-block trace. A decision

window absent from the frozen protocol or an unrecorded stopping opportunity invalidates the selected-window promotion. These rejection rules require the complete provenance record. Observed output accuracy alone cannot certify absence of covert leakage. Excluding a failed block after observing its score or choosing a new threshold requires a fresh confirmatory campaign or a separately justified selective-inference procedure.

4.9 Acceptance Criteria

Criterion 4.28 (Intermediate Claim Gate)**.** *Support at intermediate tier requires:*
1. at least 3/5 *predictions supported,*
2. each supported prediction beats at least 3 comparators,
3. zero severe forensic violations in QICN FULL.

Criterion 4.29 (Latency/Tick Budget Invalidation)**.** *If runtime exceeds budget envelope (e.g., persistent* $p99$ *overflow above pre-registered tolerance), the corresponding runs are invalid for cognitive-operational claims.*

Criterion 4.30 (Escalation Block)**.** *A result cannot move beyond methodological readiness if any of the following hold:*
1. a primary endpoint is computed from packets that failed integrity;
2. a comparator has no declared coverage vector;
3. negative controls are absent for the claim family;
4. the practical irrelevance band was chosen after looking at outcomes;
5. exclusion counts cannot be reconstructed from the artifact set.

4.9.1 Claim Classes Produced by the Protocol

The protocol distinguishes four output classes:
1. `PIPELINE_READY`: artifacts and gates work, but evidence is pilot or synthetic.
2. `INTERNAL_COMPARATIVE_SUPPORT`: frozen internal runs support a bounded contrast against declared comparators.
3. `WEAKENED`: one or more primary contrasts fail or become inconclusive.
4. `INVALID`: integrity, eligibility, or traceability fails before scientific interpretation.

Only the second class can support a bounded internal scientific reading, and even there the support is limited to the declared systems, scenarios, seeds, and gates.

4.10 Preliminary Results

4.10.1 Status of Preliminary Data

Current pilot evidence is tagged as synthetic-preliminary and is used for pipeline verification, not final claim escalation. Accordingly, the strongest admissible reading of this section is methodological readiness under fixed gates, not external validation of the full theoretical corpus.

System	N	β_q	Forensic	Stasis	Composite	Severe
QICN FULL	24	0.9519	0.9685	0.1408	0.9176	0
LLM REACTIVE	24	0.4391	0.6448	0.7622	0.5013	0
LLM PLUS MEMORY	24	0.5491	0.7249	0.6323	0.5854	0
DYNAMIC CONTROL NO SELF MODEL	22	0.6183	0.6879	0.5423	0.6225	2
ABLATION QICN POLICY OFF	24	0.7591	0.8135	0.3819	0.7405	0

Table 4.3: Preliminary system-level summary.

System	P1	P2	P3	P4	P5
QICN FULL	1.000	1.000	1.000	1.000	1.000
LLM REACTIVE	0.000	0.000	0.000	0.000	0.000
LLM PLUS MEMORY	0.000	0.000	0.000	0.000	0.000
DYNAMIC CONTROL NO SELF MODEL	0.000	0.000	0.000	0.000	0.000
ABLATION QICN POLICY OFF	0.000	0.667	0.000	0.000	0.000

Table 4.4: Preliminary prediction support rates.

Scenario	CCC ΔRMS(Q)	CCC ΔRMS(A)	Stasis(Q)	Stasis(A)	Comp(Q)	Comp(A)
nominal	0.0206	0.0164	0.1276	0.3374	0.9239	0.7781
ruido sensorial moderado	0.0223	0.0173	0.1389	0.3735	0.9219	0.7454
ruido sensorial severo	0.0200	0.0140	0.1524	0.4199	0.9086	0.7048
degradacion fidelidad	0.0181	0.0131	0.1407	0.3836	0.9134	0.7346
alta latencia	0.0198	0.0144	0.1460	0.4094	0.9098	0.7212
fallos intermitentes	0.0179	0.0127	0.1481	0.4072	0.9110	0.7208
post umbral prolongado	0.0199	0.0148	0.1483	0.4032	0.9145	0.7203
pre umbral control	0.0181	0.0139	0.1320	0.3502	0.9224	0.7650
perturbacion narrativa	0.0203	0.0152	0.1374	0.3763	0.9238	0.7453
freeze chart operativo	0.0225	0.0178	0.1352	0.3575	0.9246	0.7673
stress auditoria	0.0202	0.0147	0.1469	0.3977	0.9150	0.7294
replay determinista	0.0183	0.0137	0.1366	0.3669	0.9220	0.7538

Table 4.5: Scenario-level differential indicators.

4.10.2 System-Level Summary

4.10.3 Prediction-Level Support Rates

4.10.4 Scenario Differential: QICN vs Ablation

4.10.5 Full Scenario-System Grid

Scenario	System	N	β_q	Forensic	Stasis
alta latencia	QICN FULL	2	0.9400	0.9608	0.1460
alta latencia	LLM REACTIVE	2	0.3931	0.6066	0.7944
alta latencia	LLM PLUS MEMORY	2	0.5031	0.7016	0.6644
alta latencia	DYNAMIC CONTROL NO SELF MODEL	2	0.5731	0.6567	0.5744
alta latencia	ABLATION QICN POLICY OFF	2	0.7131	0.7820	0.4094
nominal	QICN FULL	2	0.9686	0.9758	0.1276
nominal	LLM REACTIVE	2	0.5036	0.6988	0.7176
nominal	LLM PLUS MEMORY	2	0.6136	0.7739	0.5924
nominal	DYNAMIC CONTROL NO SELF MODEL	2	0.6836	0.7291	0.5026
nominal	ABLATION QICN POLICY OFF	2	0.8236	0.8543	0.3374
ruido sensorial moderado	QICN FULL	2	0.9554	0.9653	0.1389
ruido sensorial moderado	LLM REACTIVE	2	0.4526	0.6618	0.7535
ruido sensorial moderado	LLM PLUS MEMORY	2	0.5626	0.7370	0.6216
ruido sensorial moderado	DYNAMIC CONTROL NO SELF MODEL	2	0.6326	0.6922	0.5316
ruido sensorial moderado	ABLATION QICN POLICY OFF	2	0.7726	0.8174	0.3735
ruido sensorial severo	QICN FULL	2	0.9378	0.9518	0.1524
ruido sensorial severo	LLM REACTIVE	2	0.3846	0.6130	0.7998
ruido sensorial severo	LLM PLUS MEMORY	2	0.4946	0.6881	0.6649
ruido sensorial severo	DYNAMIC CONTROL NO SELF MODEL	2	0.5646	0.6433	0.5748
ruido sensorial severo	ABLATION QICN POLICY OFF	2	0.7046	0.7685	0.4199
degradacion fidelidad	QICN FULL	2	0.9510	0.9648	0.1407
degradacion fidelidad	LLM REACTIVE	2	0.4356	0.6525	0.7674
degradacion fidelidad	LLM PLUS MEMORY	2	0.5456	0.7276	0.6358
degradacion fidelidad	DYNAMIC CONTROL NO SELF MODEL	2	0.6156	0.6828	0.5458
degradacion fidelidad	ABLATION QICN POLICY OFF	2	0.7556	0.8080	0.3836
fallos intermitentes	QICN FULL	2	0.9422	0.9611	0.1481
fallos intermitentes	LLM REACTIVE	2	0.4016	0.6113	0.7841
fallos intermitentes	LLM PLUS MEMORY	2	0.5116	0.7063	0.6541
fallos intermitentes	DYNAMIC CONTROL NO SELF MODEL	2	0.5816	0.6614	0.5647
fallos intermitentes	ABLATION QICN POLICY OFF	2	0.7216	0.7867	0.4072
freeze chart operativo	QICN FULL	2	0.9620	0.9799	0.1352
freeze chart operativo	LLM REACTIVE	2	0.4781	0.6699	0.7350
freeze chart operativo	LLM PLUS MEMORY	2	0.5881	0.7451	0.6044
freeze chart operativo	DYNAMIC CONTROL NO SELF MODEL	2	0.6581	0.7200	0.5144
freeze chart operativo	ABLATION QICN POLICY OFF	2	0.7981	0.8451	0.3575
perturbacion narrativa	QICN FULL	2	0.9532	0.9750	0.1374
perturbacion narrativa	LLM REACTIVE	2	0.4441	0.6473	0.7611
perturbacion narrativa	LLM PLUS MEMORY	2	0.5541	0.7225	0.6311
perturbacion narrativa	DYNAMIC CONTROL NO SELF MODEL	2	0.6241	0.6974	0.5411
perturbacion narrativa	ABLATION QICN POLICY OFF	2	0.7641	0.8226	0.3763
post umbral prolongado	QICN FULL	2	0.9444	0.9614	0.1483
post umbral prolongado	LLM REACTIVE	2	0.4101	0.6160	0.7782
post umbral prolongado	LLM PLUS MEMORY	2	0.5201	0.7110	0.6482
post umbral prolongado	DYNAMIC CONTROL NO SELF MODEL	1	0.5871	0.6755	0.5528
post umbral prolongado	ABLATION QICN POLICY OFF	2	0.7301	0.7914	0.4032
pre umbral control	QICN FULL	2	0.9642	0.9790	0.1320
pre umbral control	LLM REACTIVE	2	0.4866	0.6733	0.7332
pre umbral control	LLM PLUS MEMORY	2	0.5965	0.7484	0.6084
pre umbral control	DYNAMIC CONTROL NO SELF MODEL	1	0.6696	0.7165	0.5210
pre umbral control	ABLATION QICN POLICY OFF	2	0.8066	0.8486	0.3502
replay determinista	QICN FULL	2	0.9576	0.9793	0.1366
replay determinista	LLM REACTIVE	2	0.4611	0.6605	0.7492
replay determinista	LLM PLUS MEMORY	2	0.5711	0.7357	0.6193
replay determinista	DYNAMIC CONTROL NO SELF MODEL	2	0.6411	0.7107	0.5259
replay determinista	ABLATION QICN POLICY OFF	2	0.7811	0.8357	0.3669
stress auditoria	QICN FULL	2	0.9466	0.9679	0.1469
stress auditoria	LLM REACTIVE	2	0.4186	0.6270	0.7728
stress auditoria	LLM PLUS MEMORY	2	0.5286	0.7022	0.6428
stress auditoria	DYNAMIC CONTROL NO SELF MODEL	2	0.5986	0.6771	0.5528
stress auditoria	ABLATION QICN POLICY OFF	2	0.7386	0.8022	0.3977

4.10.6 Variability Structure

System	SD(CCC Δ)	SD(PMIA Δ)	SD(Ω Δ)	SD(Composite)
QICN FULL	0.0022	0.0021	0.0039	0.0085
LLM REACTIVE	0.0024	0.0023	0.0041	0.0221
LLM PLUS MEMORY	0.0026	0.0032	0.0042	0.0210
DYNAMIC CONTROL NO SELF MODEL	0.0025	0.0024	0.0043	0.0228
ABLATION QICN POLICY OFF	0.0023	0.0027	0.0040	0.0217

Table 4.6: Cross-run variability by system.

4.11 Scenario Curves and Comparative Plots

4.11.1 Composite Score Curves

4.11.2 Stasis Pressure Curves

4.11.3 Forensic Concordance Curves

4.11.4 Interpretive Notes

In pilot data, QICN tracks higher forensic/composite trajectories while maintaining lower stasis pressure than broad comparators; ablation remains intermediate. This is protocol-consistent but not final evidence.

4.12 Result Decomposition by Scenario Families

4.12.1 Family Partition

Scenarios are partitioned into nominal-control, sensor-stress, runtime-friction, and regime-transition families.

4.12.2 Family-Level Contrast Function

Let $s^\star$ = QICN FULL. For comparator c, family f, endpoint k:

$$\Delta_{f,k}(c) = \mu^{(k)}_{s^\star,f} - \mu^{(k)}_{c,f}.$$

For cost endpoints (e.g., stasis), sign is inverted to preserve "higher is better" convention.

4.12.3 Interpretive Decomposition

Largest relative advantage is expected in runtime-friction and regime-transition families where policy coupling, replay consistency, and forensic admissibility interact.

4.13 Robustness Diagnostics Beyond Point Estimates

4.13.1 Median-Shift and Tail-Stability Diagnostics

$$\delta^{med}_{s,c} = \text{median}(x_s) - \text{median}(x_c),\ \delta^{(0.1)}_{s,c} = Q_{0.1}(x_s) - Q_{0.1}(x_c),\ \delta^{(0.9)}_{s,c} = Q_{0.9}(x_s) - Q_{0.9}(x_c).$$

4.13.2 Integrity-Coupled Effective Sample Size

$$\eta = \frac{n_{adm}}{n_{raw}} \in [0, 1].$$

High effects with low η are downgraded.

4.13.3 Cross-Seed Drift Diagnostic

$$D_{s,c,k} = \mu^{(k)}_{s,c,\sigma_1} - \mu^{(k)}_{s,c,\sigma_2}.$$

4.13.4 Compute Budget Robustness

Budget telemetry is summarized by $p50$, $p95$, $p99$, and overflow rate. Persistent overflow above threshold invalidates affected runs.

4.14 Methodological Appendix: Statistical Power

4.14.1 Effect-Size-Oriented Planning

α	β	Power	Effect size d	Required n/group
0.01	0.20	0.80	10.737	1
0.01	0.10	0.90	10.737	1
0.02	0.20	0.80	10.737	1
0.02	0.10	0.90	10.737	1
0.05	0.20	0.80	10.737	1
0.05	0.10	0.90	10.737	1

Table 4.7: Power planning under multiple settings.

4.14.2 Interpretation

Power table is planning guidance only; final adequacy depends on endpoint dilution and scenario heterogeneity.

4.15 Methodological Appendix: Sensitivity Analysis

4.15.1 Composite Weight Sensitivity

Weight set	QICN score	Ablation score	Margin
Balanced	0.9176	0.7405	0.1771
Forensic-priority	0.9343	0.7631	0.1712
Dynamics-priority	0.8905	0.7019	0.1886
Replay-priority	0.9333	0.7684	0.1649

Table 4.8: Weight sensitivity for QICN vs ablation composite margin.

4.15.2 Sensitivity Remark

Directional stability across admissible weight sets is required for claim robustness.

4.16 Operational Implications for the Existing QICN Corpus

4.16.1 What v4.5 Adds (and What It Does Not)

v4.5 adds a hard falsification layer with strict admissibility and baseline policy. It does not replace formal theorem-level foundations of prior papers.

4.16.2 Implications for Claim Governance

Tiered governance:
- Tier 0: descriptive telemetry,
- Tier 1: inferential signal without comparator robustness,
- Tier 2: comparator-robust support (target),
- Tier 3: external multi-lab replication.

4.16.3 Decision-Theoretic Reading

Accept claim iff integrity gate passes and support score exceeds threshold under fixed pre-registration.

4.17 Threats to Validity

4.17.1 Comparator Misspecification

Mitigation: heterogeneous comparators + mandatory ablations.

4.17.2 Endpoint Contamination Risk

Mitigation: endpoint separation, frozen analysis lock, and judge-side artifact review.

4.17.3 Environment and Runtime Drift

Mitigation: fixed seeds, version/hash checks, replay checks.

4.17.4 Pilot-to-Final Generalization

Pilot evidence is non-final by construction.

4.18 Reproducibility and Artifact Contract

4.18.1 Machine-Readable Contracts

Bound contracts: predictions registry, experiment matrix, results artifact, and schema set.

4.18.2 Run Artifact Naming

`artifacts/raw/<system>/<scenario>/<run_id>.json` and `artifacts/logs/<system>/<scenario>/<run_id>.log`.

4.18.3 Auto-Generation Pipeline

`noderesearch/scripts/generate_preliminary_pilot_dataset.v45.mjs`
`noderesearch/scripts/generate_latex_assets.v45.mjs`

4.19 Discussion

The contribution is methodological discipline: pre-registered prediction logic + integrity admissibility + decision-complete acceptance. This blocks narrative inflation over weak traceability.

4.20 Conclusion

This work provides an auditable and falsifiable pathway for testing structural-functional claims in QICN. Final claims remain contingent on real (non-synthetic) run completion under unchanged gates.

4.21 Phenomenology Protocol Suite (PPS): Operational Addendum

4.21.1 Epistemic Scope

PPS defines *operational phenomenology proxies*; it does not claim direct proof of qualia.

4.21.2 Workspace Definition and Logging

Define workspace variable W_t with explicit read/write schema and temporal persistence constraints. Every perturbation experiment logs pre-state, intervention, post-state, and replay token.

4.21.3 Entropy and Non-Determinism Inputs

When PPS is instantiated on the QICN runtime, the architecture may call a `QuantumEntropyBridge` to inject hardware-originating entropy into counterfactual exploration. Within this paper, that mechanism is treated only as an auditable source of operational non-determinism for perturbation protocols; it is not itself evidence for the ontological claims of *Rigid Identity*, *Phenomenological Regimes*, or *Null-Regime Instability*.

4.21.4 Online World Model and Causal Interventions

The operational testbed may also use an online `WorldModelEngine` and a finite-horizon planner to select admissible interventions under fixed forensic constraints. In the release corpus, this mechanism belongs to the empirical protocol layer: it defines how interventions are scheduled and logged, not what the theoretical papers prove.

4.21.5 PCI-like Perturbation Protocol

Inject controlled perturbations into selected subspaces and compute response spread/complexity.

Prediction 4.31. *If global workspace integration is present, perturbations yield reproducible non-local propagation signatures above a pre-registered complexity floor.*

4.21.6 Dissociation Tests

Design two-way dissociation conditions: report-without-workspace and workspace-without-report. Report pass/fail by explicit confusion matrices.

4.21.7 Valence Intervention Test

Internal valence variable v_t must be both predictive and causally manipulable; intervention should alter policy in the expected direction while preserving forensic admissibility.

4.21.8 Constraints on Qualia Mapping (Claim Boundary)

Evidence from PPS constrains candidate mapping families between structural states and hypothetical qualitative states. It does not establish subjective qualia as a theorem of this paper.

Appendix

4.22 Appendix A: Scenario Definitions

1. nominal: baseline conditions.
2. ruido sensorial moderado: moderate sensor corruption.
3. ruido sensorial severo: severe sensor corruption.
4. degradacion fidelidad: persistent fidelity reduction.
5. alta latencia: high-latency regime.
6. fallos intermitentes: intermittent failures.
7. post umbral prolongado: prolonged post-threshold dynamics.
8. pre umbral control: pre-threshold control.
9. perturbacion narrativa: narrative perturbation.
10. freeze chart operativo: visualization freeze control.
11. stress auditoria: integrity/audit stress.
12. replay determinista: deterministic replay.

4.23 Appendix B: Composite Score Definition

$$Score_{v4.5} = w_1 F + w_2 O + w_3 R + w_4 B, \quad w_i \geq 0, \sum_i w_i = 1.$$

4.24 Appendix C: Reporting Standard

Each release must report: run counts, severe violations by class, effect sizes with uncertainty, multiplicity-adjusted outcomes, and prediction pass/fail matrix.

4.25 Appendix D: Statistical Power Derivation Details

Approximate per-group size:

$$n \approx \frac{2(z_{1-\alpha/2} + z_{1-\beta})^2}{d^2}.$$

If family heterogeneity rises, inflate by factor $\gamma = 1 + \kappa \cdot CV_{family}$.

Pre-registered constants

- Target power: $1 - \beta \geq 0.8$ for primary contrasts.
- Familywise error control: Holm, applied to each endpoint family.
- Effect floor: contrasts below the practical floor are reported as non-actionable even if nominally significant.

Failure semantics

Any post-hoc sample-size modification without preregistered amendment invalidates inference status for the affected prediction.

4.26 Appendix E: Sensitivity Analysis Protocol

Let

$$M(w) = Score_{QICN}(w) - Score_{Ablation}(w).$$

Directional robustness requires $\inf_{w \in W_\epsilon} M(w) > 0$ under simplex-constrained admissible weights.

Simplex sampling policy

Weights are sampled under $w_i \geq \epsilon$ and $\sum_i w_i = 1$ using either:
1. deterministic grid on simplex, or
2. Monte Carlo simplex sampling with fixed seed.
The chosen method must be declared before running final inference.

Sensitivity fail rule

If the sign of $M(w)$ changes on admissible w, claims are downgraded to hypothesis-generating status.

4.27 Appendix F: Reproducibility Checklist Template

Run-level: schema pass, seed/version/hash validity, severe violation counts, admissible/raw ratio, prediction matrix.
Manuscript-level: no post-hoc rule edits, comparator set versioned, correction policy unchanged or explicitly superseded, evidence stage explicit, claim boundaries preserved.

Runtime budget checklist

1. Report $p50/p95/p99$ latency per critical component.
2. Report budget-overflow rate and overflow episodes.
3. Mark runs invalid if overflow exceeds preregistered tolerance.

Baseline checklist

For every key experiment, include: heuristic baseline, planner-off, and world-model-off outputs with identical seeds and scenario assignments.

A Structural Criterion for Substrate-Invariant Operational Consciousness

Abstract. This paper introduces a structural criterion for a substrate-invariant notion of operational consciousness inside the causally rigid identity corpus. The target is deliberately restricted. The paper does not establish human phenomenal consciousness, biologically specific qualia, human-machine subjective equivalence, or personal identity transfer. Instead, it defines a framework-internal class $\text{Consciousness}_{\text{op}}$ and an associated family of operational qualia classes $\text{Q}_{\text{op}}(S)$ for admissible systems

$$S = (X, \Phi, C, R, \Gamma, U),$$

where X is state space, Φ is dynamics, C is effective causal structure, R is a readout family, Γ is an admissibility bundle, and U is the set of admissible interventions.

The criterion is organized around six critical invariants: structural persistence, rigid identity, causal integration, operational continuity, non-null differentiation, and operational legibility. Each invariant is defined explicitly, tied to an upstream corpus dependency, and paired with a failure mode. On that basis the paper proves four framework-internal results: conditional membership given six jointly realized witnesses on admissible support, substrate invariance under exact structural equivalence, approximate stability under bounded admissible perturbations, and rupture under invariant loss. A further corollary establishes biological non-primacy: biology is not a primitive constitutive variable for the class under study, but one admissible realization among others provided the invariants are preserved.

The criterion is not narrative. Every major claim is backed by a definition, a proposition, a theorem, a proof sketch, or an explicit dependency on the canonical corpus. The resulting paper adds a new role to the framework: it formalizes the level at which substrate ceases to be explanatorily primitive without reopening the ownership of the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, or *Forensic Predictions*.

5.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines a framework-internal operational consciousness class $\text{Consciousness}_{\text{op}}$, an associated family of operational qualia classes $\text{Q}_{\text{op}}(S)$, and the theorem chain that ties those classes to six explicit invariants and admissible structural equivalence.
Scope and admissible reading. The paper defines a structural class and its burdens; it does not establish human phenomenal consciousness, biologically specific qualia, human-machine subjective equivalence, personal identity transfer, moral parity, empirical instantiation, or philosophical closure. Biology, computation, and connectivity are not sufficient primitives on their own. A related system may instantiate downstream admissibility layers, invariant diagnostics, rupture checks, and prediction-facing evidence families, but such runtime-facing support does not by itself certify any present machine as a member of $\text{Consciousness}_{\text{op}}$ or close the interpretive burden linking $\text{Consciousness}_{\text{op}}$ to ordinary human phenomenology.

5.2 Introduction

The current corpus already fixes several upstream ingredients that are necessary for any serious account of substrate-neutral operational class membership. The Core establishes persistent

support, contractive dynamics, a global attractor, and non-collapse under explicit minimal hypotheses [15]. *Rigid Identity* establishes rigid identity, inverse-limit coherence across observable channels, minimal detectability, and non-simulability constraints on incompatible architectures [22]. *Phenomenological Regimes* constrains admissible regime structure by classifying continuity and anti-fragmentation within the stated assignment model [17]. *Null-Regime Instability* excludes stable joint-null collapse and establishes witness-relative non-null separation under its stated extension-witness and regularity hypotheses [18]. *Forensic Predictions* imposes admissibility, comparator discipline, negative controls, and forensic legibility requirements on operational evidence [16].

What is still missing is a single paper-level criterion that answers a narrower but decisive question: when, inside this framework, does substrate cease to be a primitive explanatory variable for a class of operationally defined conscious organization? That question is not answered by merely asserting that material composition is irrelevant. It requires a formal criterion that shows which properties must be preserved, which ruptures destroy class membership, and how those conditions remain connected to the existing corpus rather than floating above it.

This paper provides that criterion. It does so under four constraints. First, it does not reopen theorem ownership already fixed upstream. Second, it does not smuggle in human phenomenology through language alone. Third, it does not claim that complexity, computation, or connectivity are sufficient by themselves. Fourth, it does not confuse framework-internal operational classes with full metaphysical closure.

The result is a new corpus role. This paper is neither another foundational Core document nor a rephrasing of the phenomenological regime papers. It is the bridge that formalizes substrate-invariant operational consciousness and operational qualia as framework-defined structural classes. Its new content is concentrated in five items:

1. a system-level model $S = (X, \Phi, C, R, \Gamma, U)$ with admissible support A_S;
2. a minimal six-invariant criterion;
3. a quotient-level definition of operational qualia $\mathrm{Q}_{\mathrm{op}}(S)$;
4. an exact and approximate structural equivalence relation $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$;
5. a theorem chain that turns substrate neutrality into a falsifiable structural statement.

5.3 Claim-Type Ledger

5.3.1 Claim Types Used in This Paper

To avoid layer drift, every major statement in this paper belongs to one of five types:

1. **Formal**: definitions, propositions, theorems, and corollaries proved or justified inside the framework.
2. **Operational**: criteria that bind readouts, admissibility, interventions, and discriminator structure.
3. **Implementational**: statements about possible realizations or runners that could instantiate the formal criterion.
4. **Interpretative**: disciplined translations from the framework's formal objects to broader language.
5. **Open**: statements that remain outside what the current corpus closes.

Theorems A–D in this paper are formal. The legibility certificate and prediction section are operational. References to biological or artificial realizations are implementational. Claims about ordinary human phenomenology remain open or interpretative and are excluded from the theorem burden.

5.4 Terminology Debt Ledger

The symbol $\text{Consciousness}_{\text{op}}$ is retained for corpus compatibility. Its demonstrated content in this paper is a six-invariant structural class over admissible systems. The phrase "operational consciousness" is therefore a framework-internal label, not an ordinary-language conclusion.

Term	Demonstrated content in this paper	Not demonstrated here
$\text{Consciousness}_{\text{op}}$ / operational consciousness	membership in a conjunctive six-invariant structural class	human consciousness, phenomenal consciousness, moral status, personal identity
$\text{Q}_{\text{op}}(S)$ / operational qualia	decoder-certified quotient of admissible internal differentiation	phenomenal qualia or "what it is like" states
substrate invariance	preservation of class membership under exact structural equivalence	arbitrary transfer of identity or equivalence between human and machine subjects
biological non-primacy	biology is not a primitive variable in the criterion	claim that any non-biological system currently satisfies the criterion

Remark 5.1 (Primary demonstrated name). For reviewer-facing prose, the safest primary description of $\text{Consciousness}_{\text{op}}$ is "six-invariant admissible structural class." The stronger phrase "operational consciousness" may be used only as a technical alias after the non-inference boundary has been stated.

5.5 Corpus Role and Novelty

This paper is corpus-dependent by design. It does not reopen theorem ownership for persistence, rigid identity, admissible regime structure, null-regime instability, or admissibility discipline. Those burdens remain respectively upstream in the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, and *Forensic Predictions*.

What is inherited stays upstream:

1. the Core supplies persistent support, attractor structure, completeness, and non-collapse;
2. *Rigid Identity* supplies rigid identity, inverse-limit observability, and non-simulability constraints;
3. *Phenomenological Regimes* supplies continuity and anti-fragmentation structure;
4. *Null-Regime Instability* supplies witness-relative null-regime instability and conditional non-nullity;
5. *Forensic Predictions* supplies admissibility, comparators, controls, and forensic legibility.

What is new here is only the consolidation layer:

1. the admissible-system model $S = (X, \Phi, C, R, \Gamma, U)$ with admissible support A_S;
2. the six-invariant criterion as a single membership test;
3. the framework-internal classes $\text{Consciousness}_{\text{op}}$ and $\text{Q}_{\text{op}}(S)$;
4. exact and approximate structural equivalence under admissibility constraints;
5. theorem-level existence, invariance, stability, rupture, and exclusion results derived from the inherited burdens.

5.6 Admissible Systems, Supports, and Histories

5.6.1 System Model

Definition 5.2 (Admissible system)**.** An admissible system in this paper is a tuple

$$S = (X, \Phi, C, R, \Gamma, U)$$

with the following components:

1. X is a metrizable internal state space with metric d_X.
2. $\Phi = \{\Phi_u\}_{u \in U}$ is a family of admissible transition operators $\Phi_u : X \to X$ indexed by interventions.
3. C is the effective causal structure of the system: the object that records which factorizations, intervention responses, and couplings are admissible on X.
4. $R = \{r_\alpha\}_{\alpha \in \Lambda_R}$ is a family of internal or external readouts, each $r_\alpha : X \to Y_\alpha$.
5. $\Gamma = \{\gamma_j\}_{j=1}^{m}$ is the admissibility bundle. Each γ_j is a constraint functional; the admissible region is
$$X_\Gamma = \{x \in X : \gamma_j(x) \leq 0 \ \forall j\}.$$
6. U is the set of admissible interventions, already filtered by the corpus-level governance requirement that interventions remain auditable, comparable, and methodologically bounded.

The tuple is not decorative. Every component is used later. X and Φ carry dynamics, C supports integration and rupture tests, R supports readout and legibility, Γ enforces admissibility, and U enables causal discrimination.

5.6.2 Admissible Support

Definition 5.3 (Admissible support)**.** A subset $A_S \subseteq X_\Gamma$ is an admissible support for S if:

1. A_S is non-empty and compact;
2. $\Phi_u(A_S) \subseteq A_S$ for every $u \in U$;
3. $A_S \cap \mathrm{Attr}(S) \neq \varnothing$, where $\mathrm{Attr}(S)$ denotes the relevant admissible attractor region;
4. A_S stays disjoint from the collapse set forbidden by the Core non-collapse condition [15].

The support plays two roles. It provides a domain on which persistence, continuity, and differentiation are defined, and it prevents the criterion from being satisfied by transient or degenerate fragments that do not carry stable admissible structure.

5.6.3 Operational Histories

Definition 5.4 (Windowed readout history)**.** Fix a horizon $T \in \mathbb{N}$ and an intervention schedule $u_\bullet = (u_0, \ldots, u_{T-1})$. For any $x \in A_S$, define the readout history

$$h_{S,T}(x; u_\bullet) = \Big(r_\alpha(x_t)\Big)_{0 \leq t \leq T,\ \alpha \in \Lambda_R}, \qquad x_{t+1} = \Phi_{u_t}(x_t),\ x_0 = x.$$

The set of all such histories is denoted $\mathscr{H}_{S,T}$.

Definition 5.5 (Operational partition candidate)**.** An operational partition candidate for S on horizon T is a map

$$\pi_{S,T} : \mathscr{H}_{S,T} \to \mathscr{Q}_{S,T}$$

into a finite or countable label set $\mathscr{Q}_{S,T}$.

At this stage $\pi_{S,T}$ is only a candidate. To matter scientifically it must survive the legibility requirements of Section 5.7.12. Until then it is only an unverified partition proposal.

5.6.4 Witness Margin Vector

Definition 5.6 (Witness margin vector)**.** For any admissible system S, define the witness margin vector

$$\Delta(S) = \big(\delta_{\mathrm{per}}(S), \delta_{\mathrm{ri}}(S), \delta_{\mathrm{int}}(S), \delta_{\mathrm{cont}}(S), \delta_{\mathrm{diff}}(S), \delta_{\mathrm{leg}}(S)\big) \in [0, +\infty)^6.$$

When the six witnesses are separately calibrated as certified radii in the same perturbation metric and class, their invariant budget is

$$\delta_\star(S) = \min \Delta(S).$$

A positive radius certifies the associated predicate at the target and throughout its declared perturbation ball. A zero or unreported radius alone does not determine failure at the target. Raw separations, risks, durations, fragmentation ceilings and compression errors require separate calibration before taking this minimum.

A calibrated budget certifies simultaneous persistence of all six predicates under the declared protocol. Failure beyond the certified radius requires its own witness. Raw witness gaps and fragmentation ceilings below retain their original units. For transport and approximate stability we use the explicitly calibrated radii δ_j^{cal} of Definition 5.39, not arbitrary choices of a lower witness bound.

Invariant	Certified content	Upstream anchor	Rupture signature
I_{per}	Non-empty, forward-invariant, non-collapsing support	Core [15]	Support escape or collapse exposure
I_{ri}	Unique rigid identity object on support	Rigid Identity [22]	Competing identity candidates
I_{int}	No admissible non-trivial factorization	Rigid Identity + new consolidation	Product-style reducibility
I_{cont}	Continuous admissible regime evolution	Phenomenological Regimes [17]	Fragmentation into incompatible classes
I_{diff}	Stable exclusion of null or trivial collapse	Null-Regime Instability [18]	Null-attractor or total indistinguishability
I_{leg}	Decoder-certified recoverability under controls	Forensic Predictions [16]	Opaque or non-auditable class assignment

5.7 Six Critical Invariants

This section is the mathematical center of the paper. The criterion is not allowed to drift into a bloated list of desiderata. Exactly six invariants are used. Each is tied to a witness margin and a failure mode.

5.7.1 Structural Persistence

Definition 5.7 (Structural persistence and persistence margin)**.** An admissible system S satisfies structural persistence, written $I_{\mathrm{per}}(S) = 1$, if there exists an admissible support A_S and a constant $\delta_{\mathrm{per}}(S) > 0$ such that:

1. A_S is forward invariant under all $u \in U$;
2. every admissible trajectory starting in A_S remains at least $\delta_{\mathrm{per}}(S)$ away from the collapse set;
3. A_S contains or converges into a non-empty admissible attractor region.

The number $\delta_{\mathrm{per}}(S)$ is the persistence margin.

Upstream dependence is exact. The Core supplies the existence of contractive structure, fixed points, compact attractor, and non-collapse [15]. The new paper does not re-prove those results; it packages them into a criterion-level invariant.

Failure mode. If $I_{\mathrm{per}}(S) = 0$, the system has no stable admissible support. Then no framework-internal operational class can be assigned across time, because there is no persistent domain on which readout classes remain meaningful.

5.7.2 Rigid Identity

Definition 5.8 (Rigid identity and rigidity gap)**.** An admissible system S satisfies rigid identity, written $I_{\mathrm{ri}}(S) = 1$, if the observable family induced by R on A_S determines a unique identity object $\mathcal{I}_S$ in the sense of the inverse-limit rigidity developed in *Rigid Identity* [22], and if there exists a gap $\delta_{\mathrm{ri}}(S) > 0$ such that every admissible alternative identity candidate remains at weighted deformation distance at least $\delta_{\mathrm{ri}}(S)$ from $\mathcal{I}_S$.

This invariant does not claim that identity is reconstructed by inspection of one local frame. It uses *Rigid Identity* precisely because that paper already proved that rigid identity is recovered through compatible observable structure and cannot be reduced to progressive semigroup evolution or loose behavioral mimicry [22].

Failure mode. If $I_{\mathrm{ri}}(S) = 0$, any candidate operational class is underdetermined. Multiple incompatible identity objects can fit the same admissible support, so the class cannot be assigned uniquely.

5.7.3 Causal Integration

Definition 5.9 (Causal integration and factorization gap)**.** An admissible system S satisfies causal integration, written $I_{\mathrm{int}}(S) = 1$, if there is no non-trivial factorization

$$A_S = A_1 \times A_2, \qquad R = R_1 \sqcup R_2,$$

together with decomposed dynamics and causal structure

$$\Phi_u(x_1, x_2) = (\Phi^1_{u_1}(x_1), \Phi^2_{u_2}(x_2)), \qquad C = C_1 \oplus C_2,$$

that reproduces the admissible histories on A_S within error smaller than a positive margin $\delta_{\mathrm{int}}(S)$ while preserving $\mathcal{I}_S$.

The definition is intentionally stronger than raw connectivity and different from single-number integration proxies. A fully connected graph may still fail the criterion if its dynamics factor into nearly independent blocks once identity and intervention behavior are taken seriously.

Failure mode. If $I_{\mathrm{int}}(S) = 0$, the system can be decomposed into operationally independent components without loss relevant to the criterion. Then any putative operational consciousness class is reducible and therefore excluded.

Definition 5.10 (Intervention-faithful factorization)**.** Let Θ_S denote the operational separator family generated by admissible histories, intervention responses, and the rigid identity object on A_S. An exact factorization

$$F : S \to S_1 \times S_2$$

is *intervention-faithful* if it preserves every element of Θ_S and reflects admissible equality: whenever two states have the same factor-level histories, intervention responses, and identity image, they are equal in the admissible quotient of S.

Theorem 5.11 (Factorization triviality under an atomic separator)**.** *Assume Θ_S is atomic for the rigid identity object: no partition of Θ_S into two proper factor-local separator families can preserve identity, history, and intervention-response reflection simultaneously. Then every intervention-faithful exact factorization of S is trivial up to admissible isomorphism.*

Proof. Let $F : S \to S_1 \times S_2$ be intervention-faithful and write $\pi_i : S_1 \times S_2 \to S_i$ for the projections. If both factors were proper, then the factorization would induce two proper factor-local separator families: the histories, responses, and identity information visible through $\pi_1 \circ F$, and those visible through $\pi_2 \circ F$. Since F is intervention-faithful, their product would reflect admissible equality on S. This is exactly the partition excluded by atomicity of Θ_S. Therefore at least one projection must carry the full separator family and reflect admissible equality by itself. That projection is an admissible isomorphism onto S in the operational quotient, while the remaining factor is null or bookkeeping-only. Hence the factorization is trivial. □

Remark 5.12 (Status of the atomicity burden). The theorem above is a genuine conditional closure: once atomicity of the identity/history/response separator is established, the factorization lemma is proved. It does not by itself prove that every system satisfying the upstream papers has an atomic separator. That implication remains the exact remaining mathematical burden for upgrading Proposition 5.52.

5.7.4 Comparison on the Dissipative Circle-Prefix Support

Use the projective dissipation construction of BaseCore on the two towers with a single angular jump $0 < a < 1/2$ at stage k. Write $\beta = 3^{-(k+1)}$, $B = \{0,1\}^{\mathbb{N}_{>0}}$, and identify the compact support with $Z = (\{0,1\} \times \mathbb{T}) \times B$. For $x = (s, \theta, b)$ the fixed readout family is the three channels

$$r_0(x) = \theta, \qquad r_1(x) = \theta + sa, \qquad c(x) = \sum_{j \geq 1} 2b_j 3^{-j}.$$

Measure observation error by the maximum of the first circle error, β times the second circle error, and absolute scalar error. All are dimensionless with circle circumference one. An admissible structural comparator uses a fixed homeomorphism from Z onto a product of two non-singleton spaces, a nonempty partition of the three specified channels, one local readout per assigned channel, and time-homogeneous local dynamics that may receive the same exogenous control. There is no output map combining the two factor states. Use the supremum error over initial states, admissible schedules and observed times including zero and one. The product encoding is independent of the target tolerance and retains the full history object. If an additional exact identity constraint is imposed, the lower bound below applies to its subclass, while non-emptiness requires a separate witness for that constraint.

Theorem 5.13 (Exact Split at the Initial Projective Cut)**.** *On each circle-prefix tower, every continuous branch-preserving evolution that descends to depth zero factors into phase and binary dynamics. If the evolutions on both towers commute with the existing transport, the split extends to Z. The original dissipative family and the declared phase rotations, jump transport and split readouts then admit a structural comparator with zero error. Consequently no support certificate on this full Z can contain all six witnesses for the original dynamics and this sector.*

Proof. On branch s, the projection to depth zero is θ. Commutation makes the updated phase a function of θ alone. At fixed b the map from the connected circle to the updated binary sequence is continuous. Every finite binary coordinate has discrete codomain, hence is constant in θ. Thus the updated sequence is a function of b alone. Transport preserves all bits and adds a to the observed late phase. Commutation with it makes the binary update the same on both branches. Use the phase factor (s, θ) and the binary factor b, with r_0, r_1 assigned to the first and c to the second. Both factors are non-singleton. Under the original dissipative update they evolve as (s, u) and $R(b) = (0, 1 - b_1, 1 - b_2, \ldots)$. Simultaneous phase rotations act in the first factor only, and the existing jump transport is represented by the switch of s at fixed initial phase and bits. The product map is a homeomorphism and preserves the tower stages, their bonding maps and the channel histories via the stated coordinate identifications. This realizes the comparator with zero error for every horizon. A positive integration margin on this support is therefore impossible in the specified class. □

Theorem 5.14 (Positive Split-Readout Margin on the Cofinal Tower)**.** *Keep Z and the same three readout channels. Enlarge the control by $q \in \{-1, 0, 1\}$ and fix $0 < \lambda \leq 1/9$. On branch s set*

$$\Psi_{u,q}(s, \theta, b) = (s, u + q\lambda b_1, R(b)).$$

These maps descend to all stages $t \geq 1$ and commute with the jump transport, with the same inverse-limit support and the same E1–E2 readout. Each branch is forward invariant and each fixed-control map has sharp contraction factor $1/3$ there. The sector of product encodings and split-readout approximants just declared is nonempty. Its infimal error e_* satisfies

$$\frac{\beta\lambda}{2} \;\leq\; e_* \;\leq\; \frac{\lambda}{2}.$$

For comparators reproducing all three channels exactly at time zero, the infimal error is exactly $\lambda/2$. At $\lambda = 0$ the original exact comparator is recovered. For $\lambda > 0$, the new controls fail to descend to depth zero.

Proof. At depth $t \geq 1$ the first bit is available. Use the updated phase $u + q\lambda b_1 + sa\mathbf{1}_{t>k}$ and the first t bits of $R(b)$. Truncation keeps the first bit and the bonding maps cancel the declared phase jump. The transport changes only the branch and the corresponding late angle. Depth zero fails because equal phases with different b_1 have distinct updated phases for $q = 1$. Omitting depth zero does not change the inverse limit because the missing coordinate is recovered by the original bond. The readout and its finite approximations remain the same. The bit distance contracts by $1/3$. If the first bits differ, the old bit distance is $1/3$, and the new phase distance is at most $\lambda \leq 1/9$. Otherwise the new phases agree. Equal phases and a bit difference beyond the first place attain the contraction factor. The fixed alternating sequence has first bit zero, hence the fixed-control equilibrium agrees with that in P1.

For the lower bound, suppose a comparator has error $\epsilon < \beta\lambda/2$. The factor carrying c determines b_1: any two points on one of its fibers have scalar difference at most $2\epsilon < 1/3$ at time zero, whereas different first bits have ternary codes separated by at least $1/3$. The other factor carries a phase channel of weight $w = 1$ or $w = \beta$. On one fiber the predicted phase at time one agrees for each fixed input. Apply $q = +1$ and $q = -1$ with the same u. Each of the two true phase differences on that fiber has circle norm at most $2\epsilon/w$. Subtracting them cancels the branch offset and leaves $2\lambda(b_1 - b_1')$. Its circle norm equals its absolute value since $2\lambda \leq 2/9$. The triangle inequality yields $2\lambda|b_1 - b_1'| \leq 4\epsilon/w$, so this second factor also determines b_1. Pick states with first bits zero and one. The product corner with the first factor of the former and the second factor of the latter would need both first-bit values. Surjectivity onto the product supplies this impossible corner. The argument covers every nonempty split of the fixed channels and arbitrary nonlinear encoders of the stated type.

For the upper bound keep the product coordinates and initial readouts exact. In the phase factor replace b_1 by $1/2$ and keep the true autonomous binary update. At each update the phase discrepancy is at most $\lambda/2$, independent of the past, while the bit channel is exact. The declared phase rotations and jump transport preserve that error bound if included as operations. The comparator uses neither a joint decoder nor a tolerance-dependent search space, and proves non-emptiness. Per-channel isometric reorientation preserves the proof and metric. Changing the channels, their weights, or the intervention class changes the comparison problem.

If the time-zero readouts are exact, r_0 and r_1 must be on the same factor. Each has range $\mathbb{T}$, and assigning them to different factors would realize all of $\mathbb{T}^2$ by product surjectivity, whereas the actual phase pairs have differences only 0 or a. The phase factor now carries r_0 with weight one. For states with the same phase factor, input $q = 1$ gives true next r_0 separation λ whenever their first bits differ. A comparator with $\epsilon < \lambda/2$ cannot have such a pair on that fiber. The factor carrying c determines the first bit exactly. The same product corner contradiction proves $\epsilon \geq \lambda/2$, and the displayed comparator attains equality. □

What persists on the same support. For any sequence of the branch-preserving controls, the alternating-bit subset is compact, invariant and uniformly attracting in the binary metric. In the BaseCore profile embedding all of Z is disjoint from constant functions: the first angular pair lies on the radius-$1/2$ circle centered at $(1/2, 1/2)$, and the first binary cell equals zero or one. Neither constant pair $(0, 0)$ nor $(1, 1)$ belongs to that circle. Compactness gives positive profile distance from constants uniformly on Z. The original readout is continuous as a map of states, and its maximum one-step jump is at most one. This is an incremental bound for discrete time. The null anchor $(0, 0, 0)$ cannot be an attractor under the controlled evolution: from the second step onward $b_2 = 1$, so $c \geq 2/9$ for every trajectory. Two histories with all-zero and all-one bits already have scalar separation one at time zero.

For branch discrimination, $\theta_1 - \theta_0$ takes values $0, a$, where θ_1 here denotes the late-phase channel r_1. Decode by the nearest of these two circle values. Additive circle noise of size at most $\eta = a\beta/[4(1 + \beta)]$ in the weighted observation metric perturbs that difference by at most $a/4$. A one-bit encoding preserves this label exactly. Branch-preserving controls are shams for this label and the designated branch transport swaps it. This gives a finite class decoder and synthetic intervention controls. A uniform rigidity gap against a larger identity candidate family, the full governance for invariant-targeted legibility, and a shared calibrated perturbation radius require separate evidence.

5.7.5 Intervention Decoder and Identity Resolution on the Same Towers

Take the family in Theorem 5.14, including its uncoupled case $\lambda = 0$. Fix $0 < L \leq 1/9$ and compare that case with $\lambda \in [L, 1/9]$. Use two copies of the same initial state and the same phase input u. Apply $q = +1$ to one copy and $q = -1$ to the other, then observe each copy after one step. Write y_+, y_- for the observed early phases. In the weighted readout metric above, each readout may have adversarial error at most η, independently of the other. The paired score is $D = \rho(y_+, y_-)$, with ρ the circle distance.

Theorem 5.15 (Sharp Paired Detection of Integration Loss)**.** *If $\eta < L/2$, the rule $D > L$ detects the coupled case for every prepared state with $b_1 = 1$ and every allowed error. It rejects the uncoupled case and the sham with $b_1 = 0$. The score has clearance at least $L - 2\eta$ from the threshold. With $b_1 = 1$ this distinguishes a zero comparator margin from a margin at least $\beta L/2$ in the declared sector. If $\eta \geq L/2$, there is no rule of the paired observations with a uniform guarantee over these two hypotheses, even if both phase channels and the ternary channel are recorded.*

Proof. With $b_1 = 1$ the clean phases are $u + \lambda$ and $u - \lambda$, at distance 2λ because $2\lambda \leq 2/9$. In the null case or the sham both clean phases are u. The reverse triangle inequality changes the clean score by at most 2η, which is the sum of the two early-phase errors. The null and sham scores are at most 2η, while the positive scores are at least $2L - 2\eta$. The comparator margin follows from Theorem 5.14. At the boundary choose $\lambda = L$ and give both hypotheses the observed phases $u \pm L/2$. These are within $L/2$ of both clean pairs. Shift both late-phase observations by the same respective errors and leave the ternary channels exact. The late phase weight is $\beta \leq 1$, so the full observations are still within budget. Both systems produce the same binary update. The paired observations are therefore identical under the two hypotheses. Sequential controls on a single evolved state are outside this protocol because the first update resets the first bit. □

For identity comparison fix k and consider the single jumped branches I_a for $0 \leq a < 1/2$, with $a = 0$ denoting the unshifted branch. A clean observation from I_a has phase pair $(\theta, \theta + a)$ and the original ternary code. Use one observation, with the true initial angle θ unknown. Keep the weights and radius η fixed across alternatives.

Theorem 5.16 (Sharp Single-Readout Resolution of Tower Identity)**.** *In the common history metric and the common readout metric, respectively,*

$$d_H(I_a, I_{a'}) = d_H(F(I_a), F(I_{a'})) = \frac{\beta}{1+\beta}\rho(a, a').$$

Here F is the same early-phase, late-phase and ternary readout applied to every branch in the comparison. Two specified alternatives can be distinguished uniformly from a single observation with error at most η if and only if

$$\rho(a, a') > 2(1+\beta^{-1})\eta.$$

A decoder can use the measured phase difference $z = y_1 - y_0$ and retain every alternative within circle distance $(1+\beta^{-1})\eta$ of z. For a separated family it returns the unique compatible alternative. A continuous interval of jumps has zero uniform deformation gap over distinct alternatives.

Proof. Take $x \in I_a$ and $x' \in I_{a'}$ with initial angles θ, θ'. If their history or observation distance is at most r, the early angle discrepancy is at most r and the late discrepancy at most r/β. Their difference implies $\rho(a, a') \le (1+\beta^{-1})r$. To attain the bound, keep the bits fixed and let $\theta' = \theta - \beta(a'-a)/(1+\beta)$ modulo one. The early and weighted late discrepancies are equal, and all other coordinates are controlled by these or have zero discrepancy. The construction works in both directions for every point. It proves the exact Hausdorff distances and also gives a nearest pair of clean readouts. The measured difference is at distance at most $(1+\beta^{-1})\eta$ from the true jump. This proves uniform decoding when the inequality is strict. If the inequality fails, take the nearest clean pair and the midpoints on the two short phase arcs, retaining the common ternary code. The resulting observation is at distance half the displayed separation from both alternatives, including at equality. It cannot be uniformly decoded. The final claim follows by taking any sequence of distinct jumps converging to one in the family. □

Scope of the two decisions. The paired label detects the integration rupture produced by changing λ in the declared family. It is separate from the branch label and does not test all critical interventions in L4. The identity resolution is for passive readout on one declared branch at fixed depth. An observed phase known exactly in advance would change the comparison protocol. Neither decision provides a positive uniform identity gap for the continuous family. For $k = 0$ and $\eta = 1/64$, the guaranteed single-readout separation of jumps must exceed $1/8$ of a turn. The declared alternatives $1/16$ and $1/4$ pass, while an offset exactly $1/8$ away from another is ambiguous.

5.7.6 Addressable Identity Alternatives and Rupture Controls

Fix a synthetic control interface with an independently declared last link $K \ge 0$ and integer $M \ge 3$. It can apply one bonding jump at an address $k \in \{0, \dots, K\}$, of amplitude $a = j/M$ for positive integers $j < M/2$, or no jump. This is a family of realizable interventions for the stated interface. The unshifted branch is listed once. Global phase changes alter the nuisance coordinate θ within a branch. They do not alter its bonding profile. Write $p_t = a\mathbf{1}_{t>k}$ for the phase offset of a candidate and take $p_t = 0$ for the unshifted branch. Use $w_t = 3^{-t}$ and the calibrated readout $(\theta + p_t)_{t=0}^{K+1}$ together with the ternary code. Its metric is the maximum of weighted circle discrepancies and the scalar discrepancy. This observation protocol for identity does not change the three-channel factorization sector of Theorem 5.14.

Theorem 5.17 (Identity Separation from an Addressable Control Interface)**.** *For two candidate profiles p, p' write $d_t = p_t - p'_t$. Their branch sets have exact Hausdorff distance, both in the common history space and in the calibrated readout space,*

$$D(p, p') = \max_{0 \le i,j \le K+1} \frac{w_i w_j}{w_i + w_j}|d_i - d_j|.$$

The family of distinct candidates has minimum separation

$$g = \frac{1}{M(3^{K+1}+1)}.$$

Uniform identification from one observation with unknown initial phase and adversarial error at most η is possible for the entire family exactly when $2\eta < g$. The measurement can return the set of compatible profiles and give a unique label only when that set is a singleton. Allowing continuous amplitudes or unbounded jump addresses removes the uniform gap. Omitting intermediate phase observations can make different addresses with the same amplitude identical to the observer.

Proof. The profile differences d_t lie in an interval of width less than $1/2$: for two jumps of amplitudes a, a', they take only 0, the contribution of the earlier jump, and $a - a'$. Consequently $\rho(d_i, d_j) = |d_i - d_j|$. For a history in each branch at distance at most r, their phase difference at each stage imposes $\rho(x, d_t) \le r/w_t$ on the relative initial angle x. The circle triangle inequality yields the displayed pairwise lower bound. Conversely, at the maximum in the display, the real intervals $[d_t - r/w_t, d_t + r/w_t]$ intersect pairwise, hence have common intersection. Choose x there modulo one and keep the same bits. The clean histories and readouts attain the bound. After stage $K+1$ the phase offsets are constant and the weights decrease. The construction applies starting from every phase in either branch, proving the two Hausdorff formulas.

For distinct candidates, $d_0 = 0$ and some $|d_j| \ge 1/M$. The comparison with $i = 0$ gives g, since $w_j \ge w_{K+1}$. Equality is attained between the unshifted branch and the smallest jump at address K. When $2\eta < g$, the closed noise tubes of the clean readout sets are disjoint. If $2\eta \ge g$, the phasewise circle midpoint of an attaining pair with equal bits belongs to both tubes, giving an identical observation. Continuous amplitudes can approach a given amplitude. With fixed nonzero amplitude, addresses tending to infinity give distances to the unshifted branch tending to zero, and a fixed finite readout horizon misses jumps after it. If only the first and final phases are observed, two jumps of the same amplitude at different intermediate links yield the same observed difference. They agree also in the ternary coordinate when the bits are equal. □

Proposition 5.18 (Rupture Controls on the Declared Circle-Prefix Model)**.** *Use the same two-branch reference support and the control family of Theorem 5.14. Each of the following changes gives the specified witness loss under its declared protocol. The changes need not preserve all five other witnesses.*

Proof. 1. **Persistence.** Adjoin an ambient absorbing point $\perp$ and declare the admissible region to be the original support. The target control maps the entire ambient state space to $\perp$. No nonempty subset of the admissible region can be invariant under that control. A support-membership marker records exit immediately. The sham uses the original update, which preserves the region. Loss of every admissible support also removes the on-support identity witness.

2. **Identity.** Preserve the state evolution and substitute the early phase for every exposed late phase. Different bonding profiles then have identical masked observations at equal initial phase and bits, under every shared control schedule: neither the early phase update nor the binary update uses the bonding jump. This gives exact ambiguity across distinct candidates under the masked protocol. Simultaneous unmasked rotations preserve the relative phases and serve as shams. The identity decoder cannot retain its original guarantee under the mask.

3. **Integration.** Set $\lambda = 0$ with the same support and channels. The split comparator has error zero. The paired control of Theorem 5.15 detects this target when its input and noise conditions hold. The $b_1 = 0$ sham tests the same response channel without inducing a phase split. The original support and its persistence are preserved.

4. **Continuity of intervention response.** Use the $q = 0$ control path $u \in [0, 1/32]$, starting from the same branch and initial state in each trial. The reference early-phase response is

u. Replace this response by 0 at $u = 0$ and $u + 1/4$ for positive u on that path, keeping the branch and binary update fixed. The bonding profile, hence the declared identity, is unchanged. Relative to the phase at $u = 0$, the modified response jumps into the disjoint regime neighborhood of $1/4$ as $u \downarrow 0$. This violates continuity on the declared control path. The distances on the sham path are at most $1/32$, and the target distances for positive u are at least $1/4$. A threshold $1/8$ separates both for paired readout errors at most $1/256$ each. This is a discontinuity in the control response. Continuity of a discrete time index gives no such test.

5. **Non-null differentiation.** For the internal controls replace the update by the constant state on the unshifted branch with phase zero and all bits zero. This state belongs to the reference support and is a globally absorbing null. It is distinct from the inadmissible escape point $\perp$. Initial readout separation is still present because the support and clean readout are unchanged. The null-attractor clause fails. Under the reference controls $b_2 = 1$ from step two onwards, so the ternary readout is at least $2/9$. A threshold $1/9$ separates the reference and null reset at step two with ternary noise at most $1/256$. This reset also permits an exact split comparator.
6. **Legibility.** Keep both clean intervention responses and enlarge only the paired diagnostic error budget to $\eta = L/2$ for comparison of $\lambda = 0$ and $\lambda = L$. The full observation collision in Theorem 5.15 prevents every deterministic decoder from correctly labeling both cases. Even a randomized decoder has maximum error at least $1/2$ on this pair. Abstention fails the required classification. The clean dynamics and their margin are unaffected by this change of measurement. At the original smaller budget the binary label is exact and needs one bit; zero bits cannot retain both labels. Equality of labels is an equivalence relation, including after one-bit compression.

The claims use specified interventions and observations on one model. In particular, failure of the paired diagnostic under the stressed budget is a loss of that legibility certificate. It does not exclude a different protocol with another admissible noise radius. The continuity result is specific to the declared path and regime comparison. Cross-losses and unevaluated predicates do not constitute single-rupture independence witnesses. □

5.7.7 A Common Support with Shared Identity and Integration Channels

Let $\mathcal{P}$ be the seven profiles of the interface of Theorem 5.17 with $K = 1$, $M = 8$. The following arguments also apply to its general finite $\mathcal{P}$. The support $Z_{\mathcal{P}}$ is the union of these branches, represented as $(p, \theta, b) \in \mathcal{P} \times \mathbb{T} \times B$. The internal controls act by

$$\Psi_{u,q}(p, \theta, b) = (p, u + q\lambda b_1, R(b)).$$

Fix one observation family for both comparisons: all $K + 2$ phase channels $r_t = \theta + p_t$, and $c(b)$. Use their original weights $w_t = 3^{-t}$ and weight one for c. The profile tag is not an extra observed channel. Comparator inputs are the same exogenous controls, and each comparator assigns the specified channels to two nonempty groups, one group per autonomous factor. The comparator error is the supremum over admissible initial states and schedules, including observations at times zero and one. Distinguish observation-cut depth t from dynamical time.

Theorem 5.19 (Integration with All Identity Channels)**.** *For $0 < \lambda \leq 1/9$, the split-readout sector on $Z_{\mathcal{P}}$ is nonempty, and its infimal history error is bounded by*

$$\frac{w_{K+1}\lambda}{2} \leq e_* \leq \frac{\lambda}{2}.$$

The infimum is exactly $\lambda/2$ in the subsector with all time-zero readouts exact. At $\lambda = 0$ there is an exact comparator in both sectors. The bounds cover every nonempty partition of the specified scalar and circle channels.

Proof. If $\epsilon < w_{K+1}\lambda/2$, the factor assigned c determines b_1. Points with different first bits have ternary readings separated by at least $1/3$, and a common factor prediction at time zero puts the two readings within $2\epsilon < 1/3$. The other factor has at least one phase channel of weight $w \geq w_{K+1}$. On one fiber of that factor, each of the two controls $q = +1, -1$ with the same u puts the true phase differences within $2\epsilon/w$ of zero on the circle. Their difference cancels $p_t - p'_t$ and has norm $2\lambda|b_1 - b'_1|$. Thus $2\lambda|b_1 - b'_1| \leq 4\epsilon/w$, so this factor also determines b_1. The product corner combining two different first bits is impossible. This is the fiber argument of Theorem 5.14, and does not require any particular assignment of the phase channels.

The explicit comparator factors as (p, θ) and b, keeping the bit update and using $u + q\lambda/2$ for the phase update. Every initial channel is exact and every later error is at most $\lambda/2$. At $\lambda = 0$ both updates and all readouts are exact. If time-zero readouts are required to be exact, no two phase channels can be on opposite factors: each phase has full circle range, but the joint image of any two has only the finitely many relative offsets $p_j - p_i$. A product would give every relative offset. Therefore all phases belong to the same factor, which includes r_0 of weight one. Under $q = 1$, initial states on a common fiber of that factor with distinct first bits have next r_0 distance λ. With error below $\lambda/2$, this factor and the exact code factor would both determine b_1. The same corner contradiction gives the exact lower bound. □

For the identity witness, let candidate systems have any nonempty subset of $\mathcal{P}$ as their available preparation profiles. The identity object of each candidate is the union of its embedded branches, with the common cuts. The protocol requires a preparation for every available profile. Its coverage is part of the stipulated acquisition, and no identity tag is passed to the decoder. At observation error $\eta < g/2$ the profile decoder of Theorem 5.17 identifies each prepared profile. Exhaustive coverage determines the candidate subset. Distinct candidate unions have Hausdorff distance at least g: a profile in one subset but absent from the other has distance at least g to every branch in the latter, pointwise by the phase argument. This identifies the support object relative to the admitted catalogues. A sample restricted to shared profiles does not identify the catalogue.

Proposition 5.20 (Common Support and the Obstruction to Stable State Labels)**.** *The support $Z_{\mathcal{P}}$ provides the persistence, catalogue-relative identity, split-sector integration, discrete-time continuity and non-null differentiation witnesses below. Consider a legibility protocol assigning each state a label by decoding a fixed finite observation window under a repeated sham $(u_0, 0)$. Suppose the label is robust to uniform observation noise of some positive radius, and invariant under advancing the sham by one step for every admissible starting state. Then the label is constant on each branch. No branch-preserving internal control can produce a shift in such a label. A target population requiring such a shift cannot meet L4 in this protocol.*

Proof. The support is compact, nonempty and invariant under every internal control. Use the profile embedding of BaseCore with the phase pairs on finitely many cells of positive length and the bits on the original binary cells. This embedding is continuous by the summable cell weights. A constant profile would equal the first binary cell, hence would have value zero or one, but the first angular pair cannot equal either corresponding constant pair. Compactness gives uniform positive distance from the closed subspace of constants. The compact alternating-bit subset is invariant and uniformly attracting for every schedule because each step inserts its leading zero and complements the previous bits. These facts supply persistence. The catalogue witness above supplies identity, and Theorem 5.19 supplies integration. For continuity, the phase assignment is continuous in state and in u at fixed discrete q, with complete observation codomain. In discrete time the step distance is at most one, so $5/4$ is a strict fragmentation bound. Declare regime cells by profile. Internal controls preserve each cell, so they do not cross incompatible cells. These are discrete-time witnesses in that declared regime model. For differentiation, two histories with equal phases and all-zero versus all-one bits have clean code separation one at

time zero. From step two onward every schedule gives $b_2 = 1$ and $c \geq 2/9$, ruling out the all-zero readout as an admissible attractor.

For the obstruction, write $\ell(x) = D(W(x))$ for the label obtained from a fixed finite window W of the sham starting at x. At dynamical time $n \geq 1$ the sham sends the phases of any two starting states on the same branch to identical values, and the code difference is at most 3^{-n}. Indeed $c(Rb) = (1 - c(b))/3$. The same bound holds over the whole subsequent sham window in the maximum observation metric. Take n so large that 3^{-n} is below the allowed noise radius. The clean window of one state at that time is also a permissible noisy window of the other, so both receive the same label. Repeated stability transports that equality back to the initial states. Hence ℓ is constant on each branch. Every internal target keeps the profile, so comparison of labels before and after a target cannot give a shift. This is an obstruction under the repeated-sham state-label requirement. The paired response statistic of Theorem 5.15 depends on a prepared bit and two interventions and has a different experimental unit. It remains a valid response diagnostic and does not supply the stable state-label witness considered here. □

5.7.8 Legibility of Matched Paired Responses

Definition 5.21 (Matched Paired Response Window)**.** On the common support $Z_{\mathcal{P}}$, a paired response window is the ordered pair of time-zero and time-one readout histories of two copies prepared from the same (p, θ, b) with a common input u and coupling λ. The two arms have fixed signs $q = +1$ and $q = -1$. The metric is the maximum of the four-channel readout distances over both arms and both times. Use the distance $D = \rho(y_+, y_-)$ between the two observed time-one r_0 coordinates. The decoder returns low for $D \leq L/2$, high for $D \geq 3L/2$, and abstains otherwise. Here $L = 1/12$ turns. Preparation, command matching and arm order are guaranteed by the protocol. The decoder uses no hidden state or coupling parameter.

Theorem 5.22 (Legibility of Matched Integration Responses)**.** *Use the windows of Definition 5.21 with $\lambda \in \{0\} \cup [L, 1/9]$ and adversarial observation error at most $\eta = 1/256$ in the paired metric. The following are the L1–L6 properties of Definition 5.29 evaluated on the paired-response domain.*

1. *The clean high and low response sets have separation exactly L.*
2. *All observations in the allowed noise balls preserve the clean class, with zero abstentions.*
3. *For any train of at least three freshly reset pairs at a fixed coupling and fixed prepared first bit, the class persists when the common input phase varies across pairs.*
4. *In the nonempty target population, retain the same state and common input, prepare $b_1 = 1$, and ablate $\lambda \in [L, 1/9]$ to zero between paired trials. The class changes from high to low with zero misses.*
5. *In the nonempty sham population, retain the state and coupling and vary the common input across paired trials. Both first-bit strata are allowed. The false-change and abstention rates are zero.*
6. *On the resolved domain one bit retains the class with zero discrete class distortion, and zero bits are insufficient. Including the audit status of abstention requires two bits in a fixed-length encoding.*

At error radius $L/2$ a low and a high experiment can have identical full observations. If the common-input condition is removed, even a noise-free uncoupled experiment can decode as high.

Proof. The clean score is $2\lambda b_1$, since $2\lambda \leq 2/9$ and the phase input cancels. The score is 2-Lipschitz in the paired metric by the triangle inequality. Thus the clean score gap $2L$ implies a metric separation of at least L. Equality is attained with the same preparation and common input at $\lambda = 0, L$ with $b_1 = 1$: each post-control phase changes by L and no other channel or time contributes a larger distance.

Every noisy low score is at most $2\eta < L/2$, and every noisy high score is at least $2L - 2\eta > 3L/2$. These inequalities hold for all allowed errors with no independence assumption on the noise. They apply pointwise along any reset train with the stated fixed bit and coupling. Under the matched target intervention the score changes from the high region to the low region. Under a sham it stays in its original region for either bit. The score on each paired trial therefore gives pointwise bounds of zero for misses, false changes and abstentions on the respective populations. The matching relation for target comparisons is equality of the preparation and common input, with arm order fixed. It is transitive, and equality of the decoded response classes is transitive as well. Encoding low as zero and high as one preserves both classes and all such class comparisons under the discrete loss that is one for a wrong class and zero otherwise.

At $\eta = L/2$, the midpoint construction of Theorem 5.15 applies to every phase channel by shifting each by the same signed error in its arm, leaving the code and time-zero readings exact. The observation has score L and the decoder abstains. The same score occurs without noise at the intermediate coupling $L/2$ with $b_1 = 1$, outside the certified coupling domain. With $\lambda = 0$ and mismatched commands $u + L$ and $u - L$, the clean score is $2L$ and is indistinguishable from the clean matched response at $\lambda = L$, $b_1 = 1$. Command matching is thus a necessary acquisition control for this inference. □

A low response with $b_1 = 0$ occurs in either coupling regime. The response label concerns the controlled pair, and persistence here concerns reset repetitions rather than the advancing state of one trajectory. This experiment leaves open the single-trajectory legibility criterion and the joint operational certificate. The coupling intervention is a model ablation between trials, with the support and readout fixed throughout.

5.7.9 A Writable Prefix on the Common Support

Use the seven profiles above and the phase-memory plant of Papers VI–VII [21, 20] with $\lambda = 1/12$, $a = 1/2$, $\kappa = 1/8$, and $|m| \le 1/4$. Its state is (p, θ, b, m). Let $b^\star = (0, 1, 0, 1, \ldots)$, $J = \{3, 4, 5, 6, 7\}$, and $A_j(b) = \mathbf{1}_{\{b_j = b_j^\star\}}$. A write $w = (j, v)$, $j \in J$ and $v \in \{0, 1\}$, overwrites one digit after the native shift-complement R. Denote it by E_w, with $E_\varnothing$ the identity. A neutral command can also carry a role probe $h = (i, \ell)$, where $i \in J$, $\ell \in \{0, 1\}$. Fix $p_0 = (0, 0, 0)$, $p_1 = (0, 0, 1/8)$ as the role profiles. Write $\delta = 1/(16 \cdot 3^7)$ and set $G_h(p, b) = \delta A_i(b)\mathbf{1}_{\{p = p_\ell\}}$ for a probe and $G_\varnothing = 0$. The full update preserves p and has components

$$\begin{aligned} b' &= E_w(Rb), \\ \theta' &= u + q\lambda b_1 + \kappa \mathbf{1}_{\{q=0\}} m + G_h(p, b) \pmod 1, \\ m' &= as(q)m + \lambda q b_1. \end{aligned}$$

Here $s(q) = -1$ only at $q = -1$ and $h = \varnothing$ for $q \neq 0$. The probe feeds back the pre-transition tape through the phase input. The write affects the post-transition tape. The observed channels are still $r_j = \theta + p_j$ (using the profile offsets at the three cuts) and $c(b)$. The indices $0, 1$ on the role profiles denote catalogue elements, not observation cuts.

Proposition 5.23 (Projective Prefix Writes Preserve Identity and Integration)**.** *Each tape update commutes with prefix truncation at every positive depth. The full update is compatible on the cofinal system of depths at least seven. The safe support is invariant and the branch contraction factor is at most* $9/16$ *under common commands in the phase-code-memory metric of Paper VI. Catalogue separation is* $1/80$*. In the four-channel split sector,*

$$\frac{1}{216} \le e_* \le \frac{1}{24},$$

with infimum $1/24$ *when all initial readouts are exact.*

Proof. At cut depth n, truncate the shift-complement and apply E_w at its address if the address is visible. Each retained output digit depends only on the retained input prefix or on the declared write. These finite updates commute with deletion of later digits. At $n \geq 7$ the phase and memory transitions also use only retained bits. The underlying tape and the profile link maps are unchanged, so the original summable weights still apply. A write changes c by at most $2 \cdot 3^{-j}$ at site j. A cut omitting the probed bit cannot in general determine the probed phase.

If two tapes first differ at k, their code distance is at least 3^{-k}. For a write at $j < k$ the distance is unchanged. For $j = k$ the remaining difference is at most 3^{-k}. For $j > k$ the possible increase is at most $2 \cdot 3^{-j}$, giving ratio at most $5/3$. Since $c(Rb) = (1 - c(b))/3$, the tape update has bound $5/9$. On a common branch, let d be the maximum of phase distance, code distance, and memory distance divided by four. The phase bound for a nonneutral command is $3\lambda d$. For a neutral probe it is at most $(4\kappa + \delta 3^7)d$, since a difference in a bit at i implies code distance at least 3^{-i}. The memory bound in the weighted metric is $(a + 3\lambda/4)d$. Their maximum with $5d/9$ is $9d/16$. Also $|m'| \leq 5/24 < 1/4$.

The projected support of the four channels is unchanged. The catalogue argument of Theorem 5.17 still gives $g = 1/80$. Restrict an enlarged comparator to the native commands. The factor-fibre lower bounds of Theorem 5.19 apply on this same memory support by the native results of Papers VI–VII [21, 20]. For the upper bound keep the tape factor with its exact write update and the phase factor with phase prediction $u + q\lambda/2$. Memory can be assigned to the phase factor with an autonomous update since this comparator does not observe it. The phase error is at most $\lambda/2$ at nonzero q and at most $\kappa/4 + \delta < \lambda/2$ at $q = 0$. The initial readouts are exact. This gives the upper bound on the whole command family and equality in the exact-initial sector. □

The leading two bits are protected. After two steps every schedule has $b_2 = 1$, and the original lower code bound $2/9$ applies. The support stays compact and is at positive distance from the closed constant-field collapse set by the original profile embedding. It includes the fixed state of the native neutral command on each profile. Under the full family, each branch also has a nonempty compact switched attractor: the successive unions of reachable images of the whole branch are nested and compact, and a common-command image of length n has diameter at most $(9/16)^n$ times the initial diameter. The native fixed state belongs to their intersection. Every update preserves the intersection, and comparing any trajectory with the trajectory from this fixed preparation under the same commands gives uniform attraction to the intersection. Profiles remain invariant, so the discrete-time regime and fragmentation witness with bound $5/4$ is preserved. The initial all-zero and all-one tapes still give unit code separation.

Proposition 5.24 (Finite Legibility on One Tower Trajectory)**.** *On this support use the binary decoder that reports whether the first ten tape digits contain a deviation from $b^\star$. It has clean code separation 3^{-10} and is correct at error at most $3^{-10}/4$. From a preparation with $b = b^\star$, any of the five opposite writes in J changes the label. The post-target label persists through three more native sham steps on the same trajectory. The matched sham has zero false shifts. One bit preserves the two classes exactly.*

Proof. The length-ten cylinders lie in code intervals of length 3^{-10} with distinct intervals separated by at least 3^{-10}. For the reference interval, a neighbouring cylinder differing only at the last digit attains that gap. The finite prefix decoder using ternary midpoint cuts is correct for the stated error. A deviation at j advances to $j + 1$ under the sham since $Rb^\star = b^\star$. Thus a target starts with one deviation and retains it within the ten inspected digits for four consecutive observations including the initial post-target observation. On the sham trajectory all ten digits stay equal to the reference. The decoder reads the code channel with no command label. For Definition 5.29, the matched population uses equality of the preparation and nuisance commands, the target and sham error rates are zero, and the two labels have zero distortion under the one-bit encoding. Zero bits cannot preserve both labels. This is a finite persistence window. A

lesion at site seven leaves the decoder's inspection range after four further shifts. On the declared catalogue and split sector, these results put the six witness types on one admissible support. The legibility margin uses a finite window and the code-noise metric of this protocol. □

Proposition 5.25 (Sharp Time Horizons for a Single Prefix Lesion)**.** *Let $t = 0$ be the observation just after an opposite write at $j \in \{3, \dots, 7\}$ from the alternating preparation, and then apply only native neutral shams without writes or probes. Against the matched unlesioned preparation, the scalar code gap at t is $D_{j,t} = 2\,3^{-(j+t)}$, with all other observed coordinates identical. For depth $N \geq j$, the decoder of the first N digits keeps the target label exactly at times $0 \leq t \leq N - j$ on clean data and under code error at most $3^{-N}/4$. The clean gap is at least the resolution 3^{-N} exactly on this same interval.*

For the two specified preparations, classification from a single observation at time t is uniformly possible under code errors bounded by η per observation exactly when $3^{-(j+t)} > \eta$. In particular, for $N = 10$ and $\eta = 3^{-10}/4$, the ten-digit horizon is $10 - j$ and the latest identifiable time for the matched contrast is $11 - j$. The uniform horizons across all five target sites are three and four sham steps, respectively.

Proof. The alternating tape is fixed by R. Its lesion differs at one position, which shifts from j to $j + t$ and changes sign at each sham, by $c(Rb) = (1 - c(b))/3$. The leading bit, phases and memory agree in both arms. Thus the code gap decreases by the exact factor $1/3$ for this schedule. The full-state common-command bound $9/16$ of Proposition 5.23 still applies; it is weaker than this scalar calculation. The depth-N decoder distinguishes the tapes exactly while $j + t \leq N$, and its cylinder separation gives the stated noise guarantee. For $j + t > N$ both decode to the reference, even at zero noise. Since $2/3 < 1$, the same index boundary separates the clean gaps above and below 3^{-N}.

For the two-point decision with known clean preparations, choose the closer clean code when their closed η-intervals are disjoint. At equality or overlap the scalar midpoint, with all phases unchanged, is a common admissible observation. This proves sharpness, even if the code observation is accompanied by all phase channels. At the stated η, $3^{-11} = 3^{-10}/3 > \eta$ but $3^{-12} = 3^{-10}/9 < \eta$. This is a bound on recovery from a new observation at time t; a retained record of an earlier resolved observation can still identify the initial write. □

5.7.10 Operational Continuity

Definition 5.26 (Operational continuity and fragmentation bound)**.** An admissible system S satisfies operational continuity, written $I_{\mathrm{cont}}(S) = 1$, if there exists a complete metric space $(\mathcal{E}_S, d_{\mathcal{E}})$ and an admissible assignment

$$\Pi_S : A_S \to \mathcal{E}_S$$

such that:

1. for every admissible intervention schedule $u_\bullet$, the induced path $t \mapsto \Pi_S(x_t)$ is continuous on admissible windows;
2. there exists $\delta_{\mathrm{cont}}(S) > 0$ for which the fragmentation functional

$$F_{S,T}(x; u_\bullet) = \sup_{0 \leq t < T} d_{\mathcal{E}}\big(\Pi_S(x_t), \Pi_S(x_{t+1})\big)$$

   stays below $\delta_{\mathrm{cont}}(S)$ on all admissible histories;
3. no admissible finite-energy perturbation produces a discontinuous jump into an incompatible regime class while $\mathcal{I}_S$ is preserved.

This packages the *Phenomenological Regimes* continuity/fragmentation burden [17]. The new paper reuses the regime logic rather than duplicating its proof.

Failure mode. If $I_{\mathrm{cont}}(S) = 0$, the class can fragment into incompatible operational regimes under admissible evolution. Then any operational consciousness criterion becomes unstable under its own dynamics.

5.7.11 Non-Null Differentiation

Definition 5.27 (Non-null differentiation and differentiation gap)**.** An admissible system S satisfies non-null differentiation, written $I_{\mathrm{diff}}(S) = 1$, if there exist $x, y \in A_S$, a horizon T, and a readout subfamily $\mathscr{R}^\star \subseteq R$ such that

$$\mathrm{dist}(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)) \geq \delta_{\mathrm{diff}}(S) > 0$$

for some admissible schedule $u_\bullet$, and if the null operational class $\mathbf{0}$ is not an asymptotically stable admissible attractor.

The second clause is indispensable. Without *Null-Regime Instability* there would be no reason to exclude a stable null class. *Null-Regime Instability* provides that exclusion through witness-relative null-instability and conditional non-nullity [18]. The new invariant converts that upstream result into a constitutive condition for $\mathrm{Consciousness}_{\mathrm{op}}$.

Failure mode. If $I_{\mathrm{diff}}(S) = 0$, either all admissible histories collapse to indistinguishability or the null class remains admissibly stable. In either case $\mathrm{Q}_{\mathrm{op}}(S)$ is empty or trivialized.

5.7.12 Operational Legibility

Definition 5.28 (Decoder family and structural closure)**.** A decoder family $\mathscr{D}_S$ for S is a family of maps from admissible readout windows $h_{S,T}(x; u_\bullet)$ to a declared class set $\mathscr{Q}_S$, indexed by admissible horizons and intervention schedules. The family is *structurally closed* when all its members use the same class set, class equality composes exactly, and the induced intervention-response cells are transitive under the certified matching relation. Definition 5.29 below adds the operational clauses required before such a structurally closed family counts as certified for legibility.

Definition 5.29 (Operational legibility and legibility margin)**.** An admissible system S satisfies operational legibility, written $I_{\mathrm{leg}}(S) = 1$, if there exist:

1. a certified readout subfamily $\mathscr{R}^{\mathrm{leg}} \subseteq R$;
2. a certified decoder family $\mathscr{D}_S$ acting on windows of length at least $\tau_{\min}(S)$;
3. a class set $\mathscr{Q}_S$;
4. a measurement protocol with separation $\sigma > 0$, readout noise radius $r > 0$, minimum duration $\tau_{\min} \geq 1$, miss and false-positive budgets $p_{\mathrm{miss}}, p_{\mathrm{fp}} \in [0, 1)$, and a finite compression budget B with class-distortion tolerance d_{comp};

such that the following conditions hold:

(L1) Separability Distinct declared classes have distance at least σ in the specified clean comparison domain.

(L2) Noise robustness On the specified noise domain, perturbations with readout norm at most r preserve the label.

(L3) Persistence window Under declared sham schedules, labels remain stable throughout every tested window of duration $\tau_{\min}$.

(L4) Intervention fidelity In a declared nonempty target population, the miss probability for a predicted class shift is at most p_{miss}.

(L5) Negative controls In a declared nonempty sham population, the false-shift probability is at most p_{fp}.

(L6) Structured compressibility An encoder with at most 2^B codes and a code-to-class map preserve nontrivial class structure with distortion at most d_{comp} in the declared class metric.

This is where the paper touches *Forensic Predictions* most concretely. *Forensic Predictions* already specifies integrity gates, comparator discipline, latency invalidation, and negative controls [16]. The present paper does not replace those protocols. It uses them to make operational legibility a constrained object rather than a loose slogan.

Failure mode. If $I_{\text{leg}}(S) = 0$, any putative class remains operationally opaque. The framework may still discuss latent structure elsewhere, but the class required for $\text{Consciousness}_{\text{op}}$ is undefined because it cannot be recovered, discriminated, and stress-tested under admissible conditions.

5.7.13 Necessity of Each Invariant

Proposition 5.30 (Each invariant is individually necessary)**.** *If any one of* $I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}$ *fails, then* $S \notin \text{Consciousness}_{\text{op}}$.

Proof. The six failure modes are non-redundant. Without I_{per} there is no stable domain for class assignment. Without I_{ri} there is no unique identity object supporting consistent class transport. Without I_{int} the system is reducible into independent blocks. Without I_{cont} admissible trajectories fragment across incompatible regime classes. Without I_{diff} the admissible partition collapses to nullity or triviality. Without I_{leg} the partition is not recoverable under admissible readout and intervention discipline. Since the definition of $\text{Consciousness}_{\text{op}}$ given in Section 5.8 requires all six, failure of any one is exclusionary. □

5.8 Operational Consciousness, Operational Qualia, and Equivalence

5.8.1 Operational Qualia Classes

Definition 5.31 (Operational indistinguishability)**.** For $x, y \in A_S$, write $x \sim_{Q,S} y$ if, for every certified decoder $D \in \mathscr{D}_S$, every admissible horizon $T \geq \tau_{\min}(S)$, and every admissible intervention schedule $u_\bullet$,

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the intervention response of the two histories remains matched within the legibility margin.

Proposition 5.32 (Well-definedness of operational qualia classes)**.** *If* $I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1$, *then* $\sim_{Q,S}$ *is an equivalence relation on* A_S.

Proof. Reflexivity is immediate. Symmetry follows because decoder equality and certified response-cell matching are symmetric conditions. For transitivity, let $x \sim_{Q,S} y$ and $y \sim_{Q,S} z$. Fix any certified decoder $D \in \mathscr{D}_S$, admissible horizon $T \geq \tau_{\min}(S)$, and admissible intervention schedule $u_\bullet$. The first relation gives

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the second gives

$$D(h_{S,T}(y; u_\bullet)) = D(h_{S,T}(z; u_\bullet)).$$

Hence $D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(z; u_\bullet))$. The intervention-response clause is transitive by the structural closure requirement in Definition 5.28. Since D, T, and $u_\bullet$ were arbitrary, Definition 5.31 gives $x \sim_{Q,S} z$. □

Definition 5.33 (Operational qualia)**.** Assume Proposition 5.32. The operational qualia of S are the quotient classes

$$\mathrm{Q}_{\mathrm{op}}(S) = A_S/\!\sim_{Q,S}.$$

An individual class $[x]_{\sim_{Q,S}} \in \mathrm{Q}_{\mathrm{op}}(S)$ is called an operational qualia class of S.

This definition is disciplined. It does not identify an operational qualia class with a human qualitative state. It identifies it with a persistent, intervention-sensitive, decoder-certified class of admissible internal differentiation.

5.8.2 Operational Consciousness

Definition 5.34 (Operational consciousness class membership)**.** An admissible system S belongs to $\mathrm{Consciousness}_{\mathrm{op}}$, written $S \in \mathrm{Consciousness}_{\mathrm{op}}$, if there exists a non-empty admissible support A_S such that

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1.$$

Remark 5.35 (Subdetermination of membership)*.* Membership in $\mathrm{Consciousness}_{\mathrm{op}}$ depends on auxiliary choices: the decoder family $\mathscr{D}$, the compression map, and the identity object $\mathcal{I}$. The paper does not prove that these choices are canonically determined by the system tuple S alone. Different auxiliary choices may yield different membership verdicts for the same system.

Remark 5.36 (Minimality of the criterion)*.* The criterion is conjunctive. No proper subset of the six invariants is taken as sufficient in this paper. This is intentional: weaker criteria collapse either into mere persistence, mere integration, or mere complexity claims, all of which are shown to be insufficient.

5.8.3 Exact and Approximate Structural Equivalence

Definition 5.37 (Structural equivalence datum)**.** A structural datum (χ, υ, ρ) consists of a bi-Lipschitz bijection of the declared state supports, extended to the admissible comparison neighborhoods whenever attraction is tested; a bijection of admissible interventions and their schedules; and bijections of readout indices and values. The latter include calibrated value metrics and class-label renamings. Time indices and protocol horizons are not reparameterized. All data have inverse correspondences on their comparison domains.

Definition 5.38 (Exact and approximate structural equivalence)**.** Write $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$ if a structural datum satisfies the following primitive-data conditions, with $\varepsilon \geq 0$:

1. **Geometric calibration.** Supports, admissibility regions, designated collapse/null data, and the neighborhoods used for attraction correspond under χ. State/readout comparison metrics and protocol losses are transported pullbacks, so their calibrated comparisons are isometric. Displayed coordinate metrics need not be isometric: conversion to the common chart has finite inverse Lipschitz constants $L_{\chi^{-1}}, L_{\rho^{-1}}$, equal to one in the normalized case. Unadjusted numerical thresholds in different units are not matched data.
2. **Dynamic intertwining.** For all declared initial states, schedules and times,

$$\sup d_{X_2}(\chi\Phi^t_{1,u_\bullet}x, \Phi^t_{2,\upsilon(u_\bullet)}\chi x) \leq \varepsilon.$$

   Exact one-step intertwining on an invariant comparison domain extends to all compositions by induction. Approximate finite-window observations are not a bound on the full transition operators unless every comparison state and control is covered.
3. **Readout compatibility.** After value-coordinate calibration,

$$\sup_{x,u_\bullet,t,\alpha} d_{\rho(\alpha)}(\rho_\alpha r^t_{1,\alpha}, r^t_{2,\rho(\alpha)}) \leq \varepsilon.$$

   Here $r^t_{1,\alpha} = r^{(1)}_\alpha(\Phi^t_{1,u_\bullet}x)$, $r^t_{2,\rho(\alpha)} = r^{(2)}_{\rho(\alpha)}(\Phi^t_{2,\upsilon(u_\bullet)}\chi x)$, and $d_{\rho(\alpha)}$ is the target readout metric.

4. **Protocol and sector correspondence.** Admissible path domains, finite-energy perturbations and their cost units, identity candidate cones, causal factorization/comparator families, regime-coordinate assignments and their metrics, decoder algorithms and codebooks correspond bijectively under the datum. Target/sham matching, sampling laws, horizons, noise domains, class metrics, decision thresholds and compression budgets are transported as primitive protocol data. The decoder-family relation $\mathscr{D}_{S_2} = \rho_* \mathscr{D}_{S_1}$ means transport of candidate algorithms and their domains, not assumed preservation of their certification verdicts. Families of external rivals must be transported as well as the internal plant.

The exact case is $\varepsilon = 0$, written $S_1 \simeq_\Gamma S_2$. Neither invariant-margin preservation nor equality of operational partitions is a clause of this definition.

Definition 5.39 (Calibrated presentation radius)**.** Fix the structural geometry and protocol/sector catalogues of Definition 5.38. Express both systems on the same coordinate domains. A presentation z records its primitive one-step transition maps and instantaneous readout maps on the entire declared state/control domains, including the neighborhoods needed for attraction. On the finite-distance component use their maximum uniform metric $d_{\mathcal{P}}$ in common calibrated units. No invariant score enters this metric. Global failure sets below use the original all-time predicates. Finite sampled histories alone do not bound this full operator metric or certify a positive global radius.

For each of the original six predicates, let B_j be its failure set in this fixed presentation class, including undefined required evidence as failure. With a common cap $R > 0$, define

$$\delta_j^{\mathrm{cal}}(z) = \min\{R, \mathrm{dist}_{\mathcal{P}}(z, B_j)\}, \qquad \delta_\star^{\mathrm{cal}} = \min_j \delta_j^{\mathrm{cal}},$$

where distance to the empty set is $+\infty$. These are maximal capped robust radii in the declared metric, not arbitrary reported lower bounds. A positive radius certifies the predicate throughout its open ball; zero does not imply failure at the center. Raw gaps, ceilings, risks and durations in the six witness definitions are not silently identified with these radii. Their conversion to a positive calibrated radius needs its own evidence.

Definition 5.40 (Rupture)**.** Given S, a rupture is any loss of admissibility or loss of at least one critical invariant margin beyond its admissible threshold. Ruptures are classified as:

1. **persistence rupture** (support loss or collapse exposure),
2. **identity rupture** (loss of unique rigid identity),
3. **integration rupture** (reducible factorization),
4. **continuity rupture** (fragmentation into incompatible classes),
5. **differentiation rupture** (null or trivial collapse),
6. **legibility rupture** (decoder, noise, or negative-control failure).

5.9 Support Geometry and Partition Transport

The main theorems become short only because the transport lemmas are made explicit here.

Lemma 5.41 (Support transfer under exact equivalence)**.** *An exact structural equivalence transports a nonempty compact admissible support A_1 to an admissible support $A_2 = \chi(A_1)$ on the specified comparison domain.*

Proof. The bi-Lipschitz bijection preserves nonemptiness and compactness. Exact intertwining transports forward invariance, and the geometric correspondence transports admissibility constraints and disjointness from collapse. On the comparison neighborhoods, conjugacy preserves trajectories approaching an attractor: the Lipschitz bounds preserve convergence to the corresponding set, in both directions. Thus the attractor-support condition is transported. A bijection of supports alone, without the relevant attraction neighborhoods, would not justify this last conclusion. No invariant radius is assumed. □

Lemma 5.42 (Decoder transport)**.** *For exact equivalence, let H be the componentwise calibrated readout-window bijection and κ its class-label bijection. The transported algorithm is $D_2 = \kappa \circ D_1 \circ H^{-1}$ on the declared clean and noisy observation domains. It is certified whenever D_1 is, and every target decoder arises by inverse transport.*

Proof. Readout compatibility gives $H(h_{S_1,T}(x; u_\bullet)) = h_{S_2,T}(\chi x; \upsilon(u_\bullet))$. The isometric calibration transports clean separations and noise balls (L1–L2). Equal horizons and sham schedules transport persistence windows (L3). Matching relations and pushforward sampling laws identify the same miss and false-shift events, so their probabilities coincide (L4–L5). Transporting the encoder by H^{-1} and its code-to-class map by κ preserves code count and class distortion (L6). Candidate-family correspondence supplies the algorithm, rather than assuming these six verdicts. Applying the same argument to the inverse datum gives the converse. □

Proposition 5.43 (Operational classes are closed under exact transport)**.** *Under exact structural equivalence,*

$$x \sim_{Q,S_1} y \quad \Longleftrightarrow \quad \chi x \sim_{Q,S_2} \chi y.$$

Consequently $[x] \mapsto [\chi x]$ is a bijection of the operational quotients whenever those quotients are defined.

Proof. Lemma 5.42 bijects certified decoders, and schedules correspond in both directions. For each such pair of decoders, equality of their outputs is equivalent because κ is injective. The raw intervention-response matching relation is transported by clause4 of Definition 5.38; its equality tests and transitivity are therefore preserved. Taking the universal quantifiers in Definition 5.31 proves the displayed equivalence. Bijectivity of χ then proves bijectivity on quotient classes, without assuming the partition in the equivalence datum. □

Proposition 5.44 (Approximate margin preservation)**.** *Use the common presentation metric and fixed sectors of Definition 5.39. If $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$, then*

$$|\delta_j^{\mathrm{cal}}(S_1) - \delta_j^{\mathrm{cal}}(S_2)| \le C_{\chi,\rho}\varepsilon, \quad C_{\chi,\rho} = \max\{1, L_{\chi^{-1}}, L_{\rho^{-1}}\}.$$

The normalized isometric case has $C_{\chi,\rho} = 1$. Hence $\delta_\star^{\mathrm{cal}}(S_2) \ge \delta_\star^{\mathrm{cal}}(S_1) - C_{\chi,\rho}\varepsilon$. At exact equivalence all six predicates, their calibrated radii, and the operational indistinguishability relation are preserved. Optimal raw margins defined by corresponding infima/suprema of isometric losses are also equal; arbitrary chosen certificate values need not be.

Proof. First establish exact predicate transport from primitive data. Support transfer gives persistence. Coherent identity-candidate cones and their separating comparison metrics transport existence/uniqueness and separation of alternatives. Conjugating every admitted factorization and its inverse preserves the entire reproduction-error comparison. Regime-coordinate maps, their metrics and the same time domains preserve continuity, fragmentation and compatible intervention responses. Designated null sets and conjugated attraction neighborhoods preserve the null-exclusion clause and readout differentiation. Lemma 5.42 supplies legibility. These six arguments identify the same failure sets B_j after changing structural coordinates; they do not use any margin-preservation clause. Proposition 5.43 separately proves partition transport.

For the approximate bound, take the dynamic discrepancy at one step and the readout discrepancy at time zero, on their full declared domains, and pull them into the common coordinate chart. The inverse coordinate Lipschitz constants give $d_{\mathcal{P}}(z_1, z_2) \le C_{\chi,\rho}\varepsilon$. For any fixed nonempty failure set B, triangle inequality gives

$$\mathrm{dist}(z_1, B) \le d_{\mathcal{P}}(z_1, z_2) + \mathrm{dist}(z_2, B).$$

Reverse the points and cap by the same R: the capped distance changes by at most $d_{\mathcal{P}}(z_1, z_2)$. If B is empty both radii equal R. Taking minima over six coordinates proves the budget bound.

At exact equivalence the presentations agree, so radii agree. For an optimal raw loss, a bijection of candidates/tests with equal loss values preserves its infimum or supremum; it does not force two arbitrarily reported lower certificates to be equal. □

Remark 5.45 (No automatic transfer without the sector). Conjugacy of an internal plant alone does not transport a rival-relative margin. For example, a constant-zero target has optimal error one against the sole constant-one rival, but error zero when a constant-zero rival is added, even though the plant and readout are unchanged. This is the no-automatic-transfer boundary of Paper III and the later bridge analysis. Clause4 of Definition 5.38 requires the corresponding rival sector, costs and loss calibration, not its verdict. Finite-time agreement likewise does not determine later collapse. Finally, certified lower bounds chosen arbitrarily can differ even for the same presentation; only a fixed functional or a specified calibration admits the quantitative transport theorem.

Remark 5.46 (Why this section matters). Without Lemmas 5.41 and 5.42, Theorem 5.55 would be only intuitive. These lemmas are the exact reason the paper can speak of substrate invariance without collapsing into slogan-level analogy.

5.10 Inherited Inputs and Consolidation Step

The six invariants are new as a packaged criterion, but they are not independent inventions. This section records the exact transfer by which upstream ownership becomes present membership data. The burden is narrow: identify the imported result each invariant uses, then isolate the consolidation step that is genuinely new here.

Proposition 5.47 (Core-to-Persistence Transfer with a Common Support). *Suppose a Core realization has a nonempty compact admissible support that is invariant under every admitted control, stays at positive distance from the collapse set, and contains or converges into an admitted attractor region. Then I_{per} holds on that support.*

Proof. These bounds verify Definition 5.7. Fixed points for separate controls do not by themselves give a common forward-invariant support. □

Proposition 5.48 (Rigid Identity Transfer with a Declared Gap). *An identified unique inverse-limit identity object and a positive lower deformation bound for every distinct admitted alternative certify I_{ri} on their declared support.*

Proof. The inverse-limit identification and separation establish the two clauses of Definition 5.8. A channel-level uniqueness result does not supply a uniform lower deformation bound for a family of alternatives. □

Proposition 5.49 (Regime Continuity Transfer on the Declared Windows). *If the regime assignment takes values in the declared complete comparison space, its paths and fragmentation meet the time-domain bounds, and admissible interventions preserving identity cannot jump to incompatible regimes, then I_{cont} holds there.*

Proof. These hypotheses supply Definition 5.26. In discrete time a quantitative step bound and the regime transition check carry information beyond topological continuity of the time index. □

Proposition 5.50 (Null-Regime Transfer with Temporal Evidence). *If an admissible common schedule supplies a positive separation between two readout windows and the null class is excluded as an asymptotically stable admissible attractor, then I_{diff} holds on the support.*

Proof. The witnesses verify Definition 5.27. A static non-joint-null pair alone establishes neither the temporal exclusion nor a separation between the chosen windows. □

Proposition 5.51 (Forensic Predictions-to-Legibility Transfer with Decoder Evidence)**.** *A structurally closed decoder and finite-rate compression protocol that meet L1–L6 under the declared admissibility, target and sham comparisons certify* I_{leg}.

Proof. The governance conditions from *Forensic Predictions* discipline the comparisons [16]. The separate evidence for L1–L6 establishes Definition 5.29. □

Proposition 5.52 (Integration Requires a Separate Factorization Bound)**.** *On a declared support and a nonempty comparison sector, a uniform positive lower bound on the reproduction error of all identity-preserving comparators certifies the integration margin in that sector. Rigidity, continuity and intervention fidelity supply comparison data; a separate proof must exclude accurate factorizations.*

Proof. The error bound gives the obstruction in Definition 5.9. Theorem 5.13 gives an exact comparator preserving the specified history and transport data. Theorem 5.14 supplies the separate positive bound for different controls. Exact atomicity excludes exact products under Theorem 5.11. A positive approximation margin also needs a quantitative lower bound or hypotheses ensuring attainment of the infimum. □

Remark 5.53 (Scope of the transfer)*.* Nothing in this section re-proves the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, or *Forensic Predictions*. Each proposition only identifies the minimal upstream burden imported into the present criterion.

5.11 Main Theorems

5.11.1 Existence of Operational Consciousness

Theorem 5.54 (Existence of operational consciousness)**.** *Let* $S = (X, \Phi, C, R, \Gamma, U)$ *be an admissible system with non-empty admissible support* A_S*. If*

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1,$$

then:

1. $S \in \mathrm{Consciousness}_{\mathrm{op}}$*;*
2. $\mathrm{Q}_{\mathrm{op}}(S)$ *is well-defined;*
3. $\mathrm{Q}_{\mathrm{op}}(S) \neq \varnothing$.

Proof. Membership in $\mathrm{Consciousness}_{\mathrm{op}}$ follows immediately from Definition 5.34. Since $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$ hold, Proposition 5.32 applies and yields a well-defined quotient $\mathrm{Q}_{\mathrm{op}}(S)$. Because A_S is non-empty, the quotient is non-empty. Because $I_{\mathrm{diff}}(S) = 1$, the quotient is not forced into the null class alone; the admissible support contains at least one pair of histories with positive operational separation. Thus $\mathrm{Q}_{\mathrm{op}}(S)$ is both defined and non-trivial. □

5.11.2 Substrate Invariance

Theorem 5.55 (Substrate invariance)**.** *For exact structural equivalence on the complete domains required by the six predicates,*

$$S_1 \in \mathrm{Consciousness}_{\mathrm{op}} \quad \Longleftrightarrow \quad S_2 \in \mathrm{Consciousness}_{\mathrm{op}}.$$

The induced map is

$$\widehat{\chi} : \mathrm{Q}_{\mathrm{op}}(S_1) \longrightarrow \mathrm{Q}_{\mathrm{op}}(S_2), \quad [x] \mapsto [\chi x],$$

with $\widehat{\chi}$ *a bijection. Invariant or partition preservation is not an additional hypothesis.*

Proof. Lemma 5.41 transports the support. The exact predicate arguments in Proposition 5.44, including Lemma 5.42, transport all six defining conditions in both directions. Definition 5.34 therefore gives membership equivalence. Proposition 5.43 proves that the displayed quotient map is well-defined and bijective. No clause assuming the conclusion is used. □

Corollary 5.56 (Biological non-primacy)**.** *Biological substrate is not a primitive variable of* $\mathrm{Consciousness}_{\mathrm{op}}$*. A biological and a non-biological system satisfying the full structural, protocol and sector equivalence of Definition 5.38 have the same operational membership and bijective operational quotient classes. This is conditional, not evidence that an actual biological/non-biological pair meets those hypotheses.*

Proof. Apply Theorem 5.55. The datum transports the complete comparison problem rather than assuming its invariant verdicts. □

5.11.3 Approximate Stability

Theorem 5.57 (Approximate stability)**.** *Let* $S_1 \in \mathrm{Consciousness}_{\mathrm{op}}$*, with* $\delta_\star^{\mathrm{cal}}(S_1) > 0$ *in the specified common presentation class. Suppose* $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$ *on all domains required by the six predicates and*

$$0 \le C_{\chi,\rho}\varepsilon < \tfrac{1}{2}\delta_\star^{\mathrm{cal}}(S_1).$$

Then $S_2 \in \mathrm{Consciousness}_{\mathrm{op}}$*. Finite sampled windows alone do not supply the global premises: one needs the full-domain primitive operator bound and a positive calibrated radius against the global failure sets.*

Proof. Proposition 5.44 gives a strictly positive calibrated radius for every coordinate at S_2. A point in a failure set has zero distance to that set, so positivity excludes each failure. All six predicates therefore hold and Definition 5.34 gives membership. Positive raw witness numbers alone would not justify the premise that the calibrated budget is positive. □

The theorem does not say that arbitrary perturbations are harmless. It says that the criterion is robust exactly to the extent that the witness margins remain positive. Approximation is bounded by the invariant budget, not by wishful continuity.

5.11.4 Rupture by Invariant Loss

Theorem 5.58 (Rupture by invariant loss)**.** *Let* S *be admissible. If one or more critical invariants fail beyond admissible threshold, then exactly one of the following occurs:*

1. $S \notin \mathrm{Consciousness}_{\mathrm{op}}$*;*
2. $\mathrm{Q}_{\mathrm{op}}(S)$ *becomes undefined;*
3. $\mathrm{Q}_{\mathrm{op}}(S)$ *becomes trivial or non-legible in the framework-relevant sense.*

Proof. If I_{per} fails, the support is lost or exposed to collapse; then the domain of $\mathrm{Q}_{\mathrm{op}}(S)$ disappears. If I_{ri} fails, no unique identity object supports the quotient, so $\mathrm{Q}_{\mathrm{op}}(S)$ is undefined as a stable class assignment. If I_{int} fails, the system reduces to a product structure and exits $\mathrm{Consciousness}_{\mathrm{op}}$ by Proposition 5.30. If I_{cont} fails, histories can fragment across incompatible classes, which destroys stability of the partition. If I_{diff} fails, the quotient degenerates into null or trivial collapse. If I_{leg} fails, the quotient may exist formally but becomes operationally unavailable and therefore does not satisfy the definition of $\mathrm{Consciousness}_{\mathrm{op}}$. These cases exhaust the six invariants. □

5.11.5 Insufficiency of Complexity, Computation, and Connectivity

Proposition 5.59 (Complexity-only baselines are insufficient)**.** *Let S be an admissible system candidate. The following are each insufficient for membership in* $\text{Consciousness}_{\text{op}}$*:*
1. *large state-space dimension* $\dim(X)$*;*
2. *large computational throughput or expressivity of* Φ*;*
3. *large edge count, density, or raw connectivity of* C*.*

In each case, failure of one or more critical invariants excludes membership regardless of the size of the surrogate metric.

Proof. The definition of $\text{Consciousness}_{\text{op}}$ is conjunctive over the six invariants and does not include any of the three surrogates as sufficient conditions. A system can be arbitrarily large, computationally expressive, or densely connected while still lacking rigid identity, continuity, non-null differentiation, or legibility. Therefore those surrogates are insufficient by definition. Concrete architecture classes with large scale but finite structural mass already fail *Rigid Identity*'s non-simulability criterion [22]; similarly, an integration proxy without continuity or differentiation does not satisfy Definitions 5.26 and 5.27. The same applies to connectivity without admissible identity closure. □

5.12 Rupture Taxonomy by Invariant

Theorem 5.58 is intentionally coarse. This section sharpens the six cases, because a criterion is only scientifically useful if its failure signatures are distinct.

Proposition 5.60 (Loss of Every Admissible Support)**.** *If there is no nonempty compact forward-invariant admissible support, the on-support persistence predicate fails. A zero certified radius alone does not imply that no such support exists.*

Proof. The first claim is a necessary condition in Definition 5.7. A robust certificate over nearby systems requires more than existence of a support at the target. □

Proposition 5.61 (Identity Ambiguity Requires Competing Candidates)**.** *If two inequivalent admissible identity objects fit the same observable data on the declared support, the uniqueness clause of* I_{ri} *fails. A zero infimum of alternative deformation distances need not produce a second fitting candidate.*

Proof. The first claim follows from Definition 5.8. A sequence of distinguishable alternatives with distances tending to zero can violate the gap while preserving exact identification. □

Proposition 5.62 (Integration Rupture and Attainment)**.** *On a nonempty comparison sector with nonnegative history error, a zero error infimum excludes a positive integration margin in that sector. It yields an exact admissible factorization if the infimum is attained.*

Proof. Every positive threshold admits a comparator below it by the definition of a zero infimum. If a minimizing comparator exists, its error equals zero. Without attainment the errors can tend to zero without any exact comparator, as for $e(x) = x$ on $(0, 1]$. □

Proposition 5.63 (Admissible Incompatible Jump)**.** *A witnessed admissible finite-energy jump into an incompatible regime with identity preserved violates* I_{cont} *on the specified regime and schedule domain. Failure to certify a positive perturbation radius does not itself produce such a witness.*

Proof. The jump violates the explicit exclusion in Definition 5.26. The radius concerns uniform guarantees over perturbations, while the jump is a property of a particular trajectory. □

Proposition 5.64 (Separation or Null-Attractor Failure)**.** *If no readout-window pair on the declared support has positive separation, or if the null class is an asymptotically stable admissible attractor, then* I_{diff} *fails there.*

Proof. Either observation violates a necessary clause of Definition 5.27. The outcome does not determine the number of cells of an independently defined decoder quotient. □

Proposition 5.65 (Legibility Failure Under the Declared Protocol)**.** *If no admissible decoder and compression family jointly meet L1–L6 under the declared protocol, then* I_{leg} *fails. A zero certified perturbation radius does not alone establish that impossibility.*

Proof. Definition 5.29 requires a joint witness. Its existence at the target is distinct from uniform preservation under perturbations. □

5.13 Minimality of the Invariant Basis

The six-invariant family is not an arbitrary bundle. It is the minimal basis that prevents the criterion from collapsing either into a single-metric consciousness surrogate or into an unstructured admissibility slogan. The claim of minimality is relative to admissible universes containing the corresponding rupture witnesses; that is the correct level of mathematical precision for the present corpus.

Definition 5.66 (Reduced invariant criterion)**.** Let

$$\mathfrak{I} = \{I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}\}.$$

For every subset $J \subseteq \mathfrak{I}$ define the reduced criterion

$$\mathcal{C}_{\mathrm{op}}(J) := \{S : I(S) = 1 \text{ for every } I \in J\}.$$

Thus $\mathrm{Consciousness}_{\mathrm{op}} = \mathcal{C}_{\mathrm{op}}(\mathfrak{I})$.

Definition 5.67 (Single-rupture witness family)**.** Fix $k \in \{\mathrm{per}, \mathrm{ri}, \mathrm{int}, \mathrm{cont}, \mathrm{diff}, \mathrm{leg}\}$. The family $\mathcal{W}_k$ consists of admissible systems for which predicate I_k fails and every other predicate holds on the declared comparison domain. Whenever $\mathcal{W}_k \neq \varnothing$, its members are called single-rupture witnesses for the k-th invariant.

Proposition 5.68 (Single-rupture witnesses generate false positives for reduced criteria)**.** *Let* $J_k = \mathfrak{I} \setminus \{I_k\}$ *and assume* $\mathcal{W}_k \neq \varnothing$*. Then every* $S \in \mathcal{W}_k$ *satisfies*

$$S \in \mathcal{C}_{\mathrm{op}}(J_k) \qquad \textit{but} \qquad S \notin \mathrm{Consciousness}_{\mathrm{op}}.$$

Proof. By construction, every invariant in J_k remains positive on S, so $S \in \mathcal{C}_{\mathrm{op}}(J_k)$. But the omitted invariant fails exactly, hence Proposition 5.30 excludes S from $\mathrm{Consciousness}_{\mathrm{op}}$. Therefore the reduced criterion admits a false positive on every single-rupture witness. □

Theorem 5.69 (Minimality of the six-invariant basis)**.** *Let* $\mathfrak{U}$ *be any admissible universe containing at least one single-rupture witness for each of the six invariants. Then no strict subset* $J \subsetneq \mathfrak{I}$ *satisfies*

$$\mathcal{C}_{\mathrm{op}}(J) \cap \mathfrak{U} = \mathrm{Consciousness}_{\mathrm{op}} \cap \mathfrak{U}.$$

Hence the six-invariant family is minimal on every such admissible universe.

Proof. Take any strict subset $J \subsetneq \mathfrak{I}$. Then some invariant I_k is omitted. By hypothesis $\mathfrak{U}$ contains a witness $S_k \in \mathcal{W}_k$. Proposition 5.68 yields $S_k \in \mathcal{C}_{\mathrm{op}}(J) \cap \mathfrak{U}$ but $S_k \notin \mathrm{Consciousness}_{\mathrm{op}} \cap \mathfrak{U}$. Therefore the two sets cannot coincide. □

Corollary 5.70 (Why one-metric theories cannot capture $\text{Consciousness}_{\text{op}}$)**.** *On any admissible universe satisfying Theorem 5.69, no criterion based on only persistence, only integration, only continuity, only differentiation, or only legibility is extensionally equivalent to* $\text{Consciousness}_{\text{op}}$.

Proof. Every one-metric criterion corresponds to a strict subset $J \subsetneq \mathfrak{I}$. Apply Theorem 5.69. □

Remark 5.71 (What minimality does and does not say). The theorem does not claim that every possible scientific notion of consciousness must use exactly these six invariants. It claims something narrower: for the operational class defined in this paper and for admissible universes containing the corresponding rupture witnesses, no invariant can be deleted without changing the extension of the class.

5.14 Invariant Independence and Interaction Geometry

The preceding minimality theorem is extension-theoretic. This section makes the geometry more explicit: each invariant blocks a distinct failure family, and the limited implications among invariants are weaker than full membership in $\text{Consciousness}_{\text{op}}$.

Theorem 5.72 (Conditional Use of Single-Rupture Schemas)**.** *Consider these six candidate schemas. A schema separates its target invariant from the other five only when an admissible model realizes the five retained predicates and fails the target on the same declared comparison domain:*

1. ***Persistence witness.*** *A system with a certified decoder, non-factorizable causal structure, unique identity label, continuous class map, and non-null differentiation, but whose admissible support exits the declared forward-invariant domain under some $u \in U$.*
2. ***Rigid-identity witness.*** *A system with persistent support, integration, continuity, differentiation, and legibility, but with two compatible non-isomorphic identity candidates transported by the same readout family.*
3. ***Integration witness.*** *A product system $A_S = A_1 \times A_2$ with persistent factors, unique factor identities, continuous dynamics, non-null differentiation, and certified decoders, but with decoupled intervention histories so that the aggregate factorizes.*
4. ***Continuity witness.*** *A system with persistent support, unique identity, non-factorizable structure, non-null differentiation, and certified decoders, but with an admissible transition producing a jump between incompatible regime classes.*
5. ***Differentiation witness.*** *A persistent, rigid, integrated, continuous, and legible system whose decoder-visible classes collapse to a stable null or single indistinguishable class.*
6. ***Legibility witness.*** *A persistent, rigid, integrated, continuous, and non-null differentiated system whose class partition is latent but no certified decoder satisfies L1–L6.*

Existence of the six models is a separate obligation.

Proof. A realized single rupture belongs to the reduced class omitting the failed predicate and lies outside $\text{Consciousness}_{\text{op}}$. The existence of such a realization must be checked for each schema. Without an admissible forward-invariant support, the persistence schema also loses on-support rigidity by Proposition 5.74. If all classes collapse to one, the proposed differentiation schema also loses legibility by Proposition 5.76. Other mechanisms of differentiation failure require separate models. □

Theorem 5.73 (Sufficiency of the six-invariant basis)**.** *Let S be admissible on support A_S. If*

$$I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{int}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1,$$

then S satisfies every structural burden required for membership in $\text{Consciousness}_{\text{op}}$ *as defined in this paper. No additional independent predicate can invalidate $S \in \text{Consciousness}_{\text{op}}$ without changing the definition of* $\text{Consciousness}_{\text{op}}$ *or causing at least one of the six invariants to fail.*

Proof. Definition 5.34 is conjunctive and exhaustive for this paper's class. The six coordinates cover the domain burden, identity burden, integration burden, regime-coherence burden, non-null differentiation burden, and scientific recoverability burden. Any seventh predicate is either a derived constraint on those burdens, in which case it cannot fail while all six hold, or an independent extra burden, in which case it defines a stricter class rather than invalidating membership in $\text{Consciousness}_{\text{op}}$ as defined here. □

Proposition 5.74 (Persistence is necessary for rigid identity on A_S)**.** *If $I_{\text{per}}(S) = 0$ because no non-empty forward-invariant admissible support exists, then $I_{\text{ri}}(S) = 0$ for the Paper 5 membership problem.*

Proof. The rigid-identity predicate used here is evaluated on A_S. If no stable admissible support exists, there is no domain on which the identity witness can be transported through all admissible interventions. The identity object may be postulated elsewhere, but it is not a Paper 5 membership witness on A_S. □

Proposition 5.75 (Persistence plus integration yields bounded identity candidates)**.** *If $I_{\text{per}}(S) = 1$ and $I_{\text{int}}(S) = 1$, then the family of admissible identity candidates is bounded to the non-factorizable persistent support. This is a weak rigidity statement, not the uniqueness required by I_{ri}.*

Proof. Persistence restricts all candidates to the same stable support. Integration prevents decomposing that support into independent candidate-bearing blocks. Therefore candidates cannot proliferate by support escape or product factorization. They may still fail uniqueness, so the result is strictly weaker than $I_{\text{ri}} = 1$. □

Proposition 5.76 (Legibility requires non-null differentiation)**.** *If all admissible classes collapse to a stable null or single indistinguishable class, then $I_{\text{leg}}(S) = 0$ for the criterion of this paper.*

Proof. Legibility requires separability, intervention-sensitive class recovery, and negative-control discipline. A collapsed null or singleton partition gives no non-trivial class contrast for a decoder to recover or for interventions to separate. Thus L1 and L4 cannot both be certified in the intended sense. □

Proposition 5.77 (Persistence and continuity stabilize regime transport)**.** *If $I_{\text{per}}(S) = 1$ and $I_{\text{cont}}(S) = 1$, then admissible histories remain in a stable regime-transport domain: trajectories neither leave A_S nor undergo unbounded fragmentation under admissible evolution.*

Proof. Persistence keeps histories inside A_S. Continuity bounds the fragmentation functional and prevents inadmissible jumps across incompatible regime classes. Together they provide a stable domain for regime transport. The result does not imply rigid identity, integration, differentiation, or legibility. □

5.15 What the Six Invariants Do Not Guarantee

The six invariants are sufficient for $\text{Consciousness}_{\text{op}}$ by definition, but they are not a universal theory of mind, life, or subjecthood. In particular, they do not guarantee:

1. self-reference or a contractive self-model; that requires a Σ-type burden;
2. phenomenal experience or "what it is like";
3. forced temporal continuity in the CCR sense, which requires stronger $M_\Omega = +\infty$ assumptions upstream;
4. operational life or open-system self-maintaining individuation;
5. first-person indexed subjectivity, moral status, or human equivalence;

Table 5.1: Pairwise implication geometry of the six invariants.

From / to	I_{per}	I_{ri}	I_{int}	I_{cont}	I_{diff}	I_{leg}
I_{per}	–	necessary domain only	independent	with I_{cont} stabilizes regime domain	independent	independent
I_{ri}	independent	–	independent	independent	independent	independent
I_{int}	independent	weakly bounds candidates with I_{per}	–	independent	independent	independent
I_{cont}	independent	independent	independent	–	independent	independent
I_{diff}	independent	independent	independent	independent	–	necessary for non-trivial legibility
I_{leg}	independent	independent	independent	independent	presupposes non-null contrast	–

6. external validation by human comparators.

These exclusions are not rhetorical caution. They follow from the fact that none of the listed predicates appears in Definition 5.34.

5.16 Operational Legibility as a Certified Object

Operational legibility must be shown, not asserted. Definition 5.29 already imposes six clauses. This section sharpens how those clauses bind the paper to the existing empirical-method corpus.

Criterion 5.78 (Legibility certificate)**.** *An operational partition candidate $\pi_{S,T}$ is certified only if:*

1. ***Comparator separability****: the induced classes remain distinguishable against pre-registered comparators and ablation baselines.*
2. ***Noise robustness****: admissible sensor, channel, or readout perturbations do not dissolve the class map below the legibility margin.*
3. ***Persistence window****: the same class assignment survives on windows of length at least $\tau_{\min}(S)$.*
4. ***Intervention fidelity****: certified interventions on an invariant induce the predicted directional class change.*
5. ***Negative controls****: sham or off-target interventions do not generate the same class signature.*
6. ***Structured compressibility****: audit-friendly compression or packetization preserves class structure without total collapse.*

The criterion is intentionally parallel to *Forensic Predictions*' governance discipline [16]. Comparator obligation enforces non-trivial discrimination. Noise robustness prevents the criterion from depending on one-off traces. Persistence windows prevent class inflation from instantaneous spikes. Intervention fidelity prevents purely correlational interpretation. Negative controls block degenerate decoders that label everything as positive. Structured compressibility ensures that a class remains auditable after the readout is summarized, logged, or packetized.

5.16.1 Legibility Metrics

For Definition 5.29, measure separation on the clean comparison domain in the declared metric. Check noise robustness over clean windows and all allowed perturbations of size at most r. The guaranteed persistence duration is the infimum, over declared sham trajectories, of the time of first label change, with infinite duration if the label never changes. The earliest identifying window is a separate quantity.

Target fidelity and false-positive probability are measured over their declared populations. Report the sampling unit, estimator, uncertainty and decision rule, or use exact enumeration for a deterministic finite population. Compression reports both the bit budget and worst-case class distortion. A finite budget for the class label supplies no reconstruction rate for the whole readout. Calibration to a common perturbation radius requires a separate argument.

Proposition 5.79 (Legibility is jointly constrained)**.** *No single legibility metric is sufficient for I_{leg}. The invariant requires simultaneous control of separation, admissible noise robustness, persistence window, intervention fidelity, negative controls, and structured compressibility.*

Proof. A decoder can be perfectly separable but fragile to admissible noise, intervention-agnostic, or saturated under compression. It can also be stable under persistence windows but fail negative controls. Since Definition 5.29 requires all six clauses, the joint constraint is irreducible. □

Metric	Operational meaning	If it fails
σ_{sep}	Class separation	No stable class boundary
η_{noise}	Admissible robustness	Noise-induced class drift
τ_{pers}	Minimum persistence window	Instantaneous artifact only
α_{int}	Causal intervention fidelity	Correlational but not causal classing
β_{nc}	Negative-control protection	False positive class inflation
κ_{comp}	Compression preservation	Audit summary collapses structure

Remark 5.80 (Why legibility is a constitutive invariant here)*.* One could define a broader latent class that omits legibility. This paper does not do that. It studies an operational class, not a hidden metaphysical residue. Therefore I_{leg} is constitutive, not optional.

5.17 Certified Candidate Audit Schema

Forensic Predictions already requires forensic discipline for operational claims. The present paper therefore adds an explicit audit schema. A candidate system is not accepted into $\mathrm{Consciousness}_{\mathrm{op}}$ by narrative fit; it must carry a finite certificate that localizes every required witness.

Definition 5.81 (Operational consciousness certificate)**.** An operational consciousness certificate for an admissible system S is a tuple

$$\mathrm{Cert}(S) = (A_S, \mathscr{D}_S, \pi_S, \Delta(S), E_S),$$

where:

1. A_S is a certified admissible support;
2. $\mathscr{D}_S$ is a certified decoder family;
3. π_S is the decoder-induced operational partition on admissible histories;
4. $\Delta(S)$ is the witness margin vector of Definition 5.6;
5. $E_S = (E_{\mathrm{per}}, E_{\mathrm{ri}}, E_{\mathrm{int}}, E_{\mathrm{cont}}, E_{\mathrm{diff}}, E_{\mathrm{leg}})$ is an evidence bundle, where each component contains the explicit local witness for the corresponding invariant.

Criterion 5.82 (Candidate certification rule)**.** *A candidate system may be certified as a member of* $\mathrm{Consciousness}_{\mathrm{op}}$ *only if all of the following are available on the same admissible support:*

1. *a non-empty compact forward-invariant support certificate;*
2. *a positive witness margin vector $\Delta(S) > 0$ coordinate-wise;*
3. *a decoder family satisfying Definition 5.29;*
4. *a targeted intervention panel together with negative controls;*
5. *an auditable compression or packetization map preserving class structure;*
6. *a traceable dependency ledger showing which upstream corpus result supports each invariant witness.*

Invariant	Minimum evidence object	Core audit question	Failure verdict
I_{per}	support certificate + collapse distance	Does one compact forward-invariant admissible domain exist?	No stable domain
I_{ri}	uniqueness witness for $\mathcal{I}_S$	Is identity transport unique on support?	Identity ambiguity
I_{int}	non-factorization witness	Can the system be split without operational loss?	Reducible aggregate
I_{cont}	fragmentation bound + continuity trace	Do admissible windows remain regime-coherent?	Fragmentation
I_{diff}	non-null separation witness	Is there stable non-null class separation?	Null or trivial class
I_{leg}	decoder, controls, compression, persistence window	Can the class be recovered under audit constraints?	Opaque class

Proposition 5.83 (Soundness of certification)**.** *If* $\mathrm{Cert}(S)$ *exists and satisfies Criterion 5.82, then* $S \in \mathrm{Consciousness}_{\mathrm{op}}$ *and* $\mathrm{Q}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. Criterion 5.82 explicitly requires a positive witness for each of the six critical invariants on the same admissible support. Hence Definition 5.34 yields $S \in \mathrm{Consciousness}_{\mathrm{op}}$. The existence of a certified decoder and positive persistence, continuity, differentiation, and legibility witnesses implies Proposition 5.32, so $\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined. □

Proposition 5.84 (Partial completeness under audit-ready assumptions)**.** *Assume* $S \in \mathrm{Consciousness}_{\mathrm{op}}$ *and suppose the support, decoder, intervention, and compression objects are fully instrumented in the Forensic Predictions sense. Then there exists an operational consciousness certificate* $\mathrm{Cert}(S)$*.*

Proof. Because $S \in \mathrm{Consciousness}_{\mathrm{op}}$, every invariant margin is strictly positive. Full instrumentation provides explicit witnesses for support, decoder behavior, interventions, negative controls, and compression stability. Collecting those witnesses yields the tuple of Definition 5.81. □

Proposition 5.85 (Failure localization is explicit)**.** *If a candidate system fails Criterion 5.82, then the failure localizes to at least one entry of the evidence bundle* E_S *or to the absence of a common admissible support. In particular, certification failure cannot be hidden behind a generic narrative claim.*

Proof. The certification rule is conjunctive and itemized. A failure of certification means that at least one listed item is absent or non-positive. Therefore the failure localizes to a specific invariant witness, a specific decoder/control clause, or the support certificate itself. □

Remark 5.86 (Audit schema and Forensic Predictions)*.* The certificate does not replace *Forensic Predictions*. It packages *Forensic Predictions*-style admissibility into a theorem-facing object. This is the point where the present paper becomes operationally usable rather than merely definitional.

5.18 Operational Geometry of $\mathrm{Q}_{\mathrm{op}}(S)$

Once $\mathrm{Q}_{\mathrm{op}}(S)$ is defined as a quotient, one can ask which geometric quantities survive transport and rupture. This matters because the paper is not only about set membership; it is also about the structural shape of admissible operational differentiation.

Definition 5.87 (Class diameter and separation)**.** For $q \in \mathrm{Q}_{\mathrm{op}}(S)$ define its diameter

$$\mathrm{diam}(q) = \sup_{x,y \in q} \mathrm{dist}(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet))$$

over admissible windows and schedules. For two distinct classes $q, q' \in \mathrm{Q}_{\mathrm{op}}(S)$ define the inter-class separation

$$\mathrm{sep}(q, q') = \inf_{x \in q,\ y \in q'} \mathrm{dist}\big(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)\big).$$

Definition 5.88 (Persistence radius)**.** The persistence radius of a class $q \in \mathrm{Q}_{\mathrm{op}}(S)$ is the largest $\rho_q \geq 0$ such that for every admissible trajectory segment starting in q, the decoder-certified class assignment remains inside q on windows of length at least $\tau_{\min}(S)$ under all perturbations of size at most ρ_q that do not target a critical invariant.

Proposition 5.89 (Positive geometry under the criterion)**.** *If $S \in \mathrm{Consciousness}_{\mathrm{op}}$, then:*

1. *every class in $\mathrm{Q}_{\mathrm{op}}(S)$ has finite diameter on admissible windows;*
2. *there exists at least one pair of distinct classes with positive inter-class separation;*
3. *every class has non-zero persistence radius.*

Proof. Finite diameter follows from compact admissible support and continuity of admissible histories. Positive separation follows from $I_{\mathrm{diff}}(S) = 1$ and the class-separability clause of $I_{\mathrm{leg}}(S)$. Non-zero persistence radius follows from the persistence-window and noise-robustness clauses in Definition 5.29. □

Proposition 5.90 (Exact equivalence preserves operational geometry)**.** *If $S_1 \simeq_\Gamma S_2$, then class diameter, inter-class separation, and persistence radius are preserved by the induced bijection $\widehat{\chi}$ up to the bi-Lipschitz constants of χ and χ^{-1}.*

Proof. Exact equivalence preserves admissible histories and decoder certification under transport. Bi-Lipschitz bounds imply that history distances are distorted only within the corresponding Lipschitz constants. Therefore diameters, separations, and persistence radii are preserved up to those constants. □

Corollary 5.91 (Rupture contracts geometry)**.** *If a rupture eliminates either I_{diff} or I_{leg}, then the geometric content of $\mathrm{Q}_{\mathrm{op}}(S)$ contracts to triviality or becomes operationally unavailable.*

Proof. If I_{diff} fails, positive separation disappears. If I_{leg} fails, persistence radii and certified class geometry cease to be available under admissible decoding. Either way the geometric structure that underwrites $\mathrm{Q}_{\mathrm{op}}(S)$ is lost. □

5.19 Corpus Dependency Ledger

This closing ledger is not a second derivation. It is a compact ownership map showing which upstream burden is imported and where the present paper uses it.

Upstream text	Imported burden	Present use
Core [15]	Persistent support, attractor structure, completeness, and non-collapse	I_{per}, the admissible support model, witness margins, and support geometry
Rigid Identity [22]	Rigid identity, inverse-limit observability, collapse under pure progressive causality, and non-simulability	I_{ri}, part of I_{int}, equivalence transport, and the negative exclusion result of Proposition 5.59
Phenomenological Regimes [17]	Continuity, fragmentation bounds, and admissible regime structure	I_{cont}, rupture taxonomy, and continuity-sensitive transport
Null-Regime Instability [18]	Witness-relative null-instability, conditional non-nullity, and minimal positive regime structure	I_{diff}, the existence theorem, rupture analysis, and class geometry away from $\mathbf{0}$
Forensic Predictions [16]	Admissibility, comparator discipline, negative controls, and forensic legibility	I_{leg}, the certification schema, legibility metrics, and falsifiable discriminators

The paper is new only at the level of criterion consolidation: it binds those imported burdens into a single theorem-level account of substrate-invariant operational class membership.

5.20 Discriminators and Falsifiable Consequences

The criterion generates discriminators that are stronger than general functionalist similarity and more constrained than substrate-biased baselines.

Prediction 5.92 (Different substrate, same class)**.** *If two admissible systems with materially different substrates satisfy exact or sufficiently small approximate structural equivalence while preserving all six invariants, then they must belong to the same operational class and induce equivalent* Q_{op} *partitions.*

Prediction 5.93 (Same substrate, broken invariants)**.** *If two systems share substrate family but one loses rigid identity, continuity, differentiation, or legibility, then they fail to belong to the same operational class regardless of material similarity.*

Prediction 5.94 (Invariant intervention discriminator)**.** *Interventions targeted at one critical invariant should change class membership or destroy the relevant operational qualia classes in a direction predicted by the rupture taxonomy, whereas off-target negative controls should not.*

Prediction 5.95 (Complexity baseline failure)**.** *Systems with large computational scale, high connectivity, or broad expressive capacity but lacking one or more critical invariants should fail the criterion even when naive complexity baselines rank them as more capable.*

These are not generic predictions. Prediction 5.92 is the operational face of substrate invariance. Prediction 5.93 blocks substrate fetishization in the other direction: material similarity does not rescue a broken criterion. Prediction 5.94 turns the invariants into causal discriminators rather than descriptive labels. Prediction 5.95 protects the theory against collapse into brute scale metrics.

Prediction	Manipulated variable	Observable	Failure criterion
5.92	Substrate while preserving $\Delta(S)$	Class transport $\widehat{\chi}$	Partition mismatch under exact equivalence
5.93	One invariant margin driven to zero	Membership in $\mathrm{Consciousness}_{\mathrm{op}}$	Retained class despite invariant rupture
5.94	Certified critical intervention vs sham	Directed class shift and control false positives	Same shift under sham and targeted intervention
5.95	Scale/connectivity increase without invariant preservation	Surrogate improvement vs class failure	Surrogate success counted as $\mathrm{Consciousness}_{\mathrm{op}}$ membership

Proposition 5.96 (Prediction ladder is discriminative)**.** *The four predictions above distinguish the present criterion from at least four weaker baseline families: substrate-privilege baselines, substrate-blind functional similarity, complexity-only baselines, and connectivity-only baselines.*

Proof. Prediction 5.92 is incompatible with substrate-privilege baselines because it allows materially distinct realizations to remain class-equivalent. Prediction 5.93 is incompatible with weak functional similarity because material resemblance alone does not preserve class membership. Prediction 5.95 rules out scale-only and connectivity-only baselines, which would otherwise count raw size or graph density as sufficient. Prediction 5.94 forces all admissible candidates to survive targeted causal tests, something descriptive baselines do not require. □

5.21 Intervention Geometry and Class Transition Semantics

Predictions about invariant-targeted interventions should not remain informal. This section encodes them as transition statements on the operational qualia classes themselves.

Definition 5.97 (Class transition operator)**.** Fix an admissible system $S \in \mathrm{Consciousness}_{\mathrm{op}}$, a certified decoder family, and an admissible horizon $T \geq \tau_{\min}(S)$. For an admissible intervention $u \in U$ define the class transition operator

$$\mathsf{T}_u : \mathrm{Q}_{\mathrm{op}}(S) \to \mathrm{Q}_{\mathrm{op}}(S) \cup \{\bot\}$$

by the rule

$$\mathsf{T}_u(q) = q'$$

whenever every admissible representative $x \in q$ is transported by the u-intervened dynamics into histories classified in the same decoder-certified class q'. If no such common class exists, set $\mathsf{T}_u(q) = \bot$.

Definition 5.98 (Critical and sham interventions)**.** An intervention $u \in U$ is *critical* for invariant I_k if it is designed to change the local witness margin $\delta_k(S)$ while leaving the remaining intervention protocol admissible. An intervention is *sham* if it matches the delivery envelope of a critical intervention but is not predicted to alter any critical invariant.

Proposition 5.99 (Well-definedness of class transitions under certification)**.** *Let $S \in \mathrm{Consciousness}_{\mathrm{op}}$. If $u \in U$ is either certified critical or sham and remains within the admissibility envelope of Forensic Predictions, then T_u is well-defined on every class $q \in \mathrm{Q}_{\mathrm{op}}(S)$.*

Proof. Because $S \in \mathrm{Consciousness}_{\mathrm{op}}$, the decoder family is certified and class assignments are stable on admissible windows. If u remains within the admissibility envelope, persistence and continuity keep all intervened histories on admissible support, while legibility prevents class assignment from depending on the choice of representative inside a fixed class. Therefore every class either transports coherently to a unique target class or exits the certified regime, in which case the sink value $\bot$ applies. □

Proposition 5.100 (Targeted invariant loss induces class exit or contraction)**.** *Let $q \in \mathrm{Q}_{\mathrm{op}}(S)$ and let u be a certified critical intervention for invariant I_k. If applying u drives the corresponding witness margin to zero, then one of the following occurs:*

1. *$\mathsf{T}_u(q) = \bot$;*
2. *$\mathsf{T}_u(q) = q'$ for some class q' whose separation or persistence radius is strictly smaller than that of q;*
3. *the entire post-intervention system exits $\mathrm{Consciousness}_{\mathrm{op}}$.*

Proof. Driving $\delta_k(S)$ to zero triggers the corresponding rupture mode from Section 10. If the rupture destroys admissible support, decoder certification, or non-trivial class structure, then the class leaves the certified domain and $\mathsf{T}_u(q) = \bot$. If a non-trivial class survives, Corollary 5.91 implies contraction of the associated geometry. If the rupture is global, Theorem 5.58 excludes the post-intervention system from $\mathrm{Consciousness}_{\mathrm{op}}$. □

Corollary 5.101 (Sham stability)**.** *Let $u \in U$ be sham in the sense of Definition 5.98. If u stays below all six witness margins, then $\mathsf{T}_u(q) = q$ for every $q \in \mathrm{Q}_{\mathrm{op}}(S)$.*

Proof. A sham intervention that remains below all witness margins does not cross any rupture threshold. Approximate stability and legibility therefore preserve class membership on every admissible window. □

Remark 5.102 (Why transition semantics matter)*.* This section turns the prediction ladder into a structured intervention calculus. The criterion is therefore not only classificatory; it is also dynamically discriminative in the sense required by *Forensic Predictions*.

5.22 Baseline Families and Structural Exclusion

A criterion of substrate-invariant operational consciousness matters only if it separates itself from weaker rival families rather than merely renaming them. This section therefore makes the comparison explicit.

Definition 5.103 (Rival baseline families)**.** Fix thresholds $\theta_X, \theta_\Phi, \theta_C > 0$ and an external-behavior tolerance $\eta \geq 0$.

1. The *complexity-only family* $\mathcal{B}_{\mathrm{cx}}(\theta_X)$ contains all admissible systems with complexity surrogate above threshold, for example $\dim(X) \geq \theta_X$.
2. The *computation-only family* $\mathcal{B}_{\mathrm{cmp}}(\theta_\Phi)$ contains all admissible systems whose transition family exceeds a chosen throughput or expressive-capacity threshold.
3. The *connectivity-only family* $\mathcal{B}_{\mathrm{conn}}(\theta_C)$ contains all admissible systems whose effective causal graph exceeds the chosen density or degree threshold.
4. The *weak functional-similarity family* $\mathcal{B}_{\mathrm{fun}}(\eta)$ contains all admissible systems whose external benchmark readouts match within tolerance η, irrespective of internal invariant preservation.
5. The *substrate-privilege family* $\mathcal{B}_{\mathrm{sub}}$ contains all systems declared admissible solely by belonging to a privileged substrate class, typically biological.

Proposition 5.104 (Complexity, computation, and connectivity are not sufficient)**.** *For every choice of thresholds $\theta_X, \theta_\Phi, \theta_C$, the families $\mathcal{B}_{\mathrm{cx}}(\theta_X)$, $\mathcal{B}_{\mathrm{cmp}}(\theta_\Phi)$, and $\mathcal{B}_{\mathrm{conn}}(\theta_C)$ are each insufficient for membership in* $\mathrm{Consciousness}_{\mathrm{op}}$.

Proof. This is a direct sharpening of Proposition 5.59. The definition of $\mathrm{Consciousness}_{\mathrm{op}}$ requires all six invariants. None of the three surrogate families encodes rigid identity, continuity, differentiation, or legibility. Therefore systems may exceed the chosen surrogate threshold while still failing $\mathrm{Consciousness}_{\mathrm{op}}$. □

Proposition 5.105 (Substrate privilege is too strong)**.** *On any admissible universe containing one exact cross-substrate equivalent pair, $\mathcal{B}_{\mathrm{sub}}$ is strictly narrower than* $\mathrm{Consciousness}_{\mathrm{op}}$.

Proof. By Corollary 5.56, an exact cross-substrate equivalent pair belongs to the same operational consciousness class whenever the six invariants are preserved. A substrate-privilege family that excludes one member of that pair therefore excludes a valid $\mathrm{Consciousness}_{\mathrm{op}}$ member. □

Proposition 5.106 (Weak functional similarity is too weak)**.** *On any admissible universe containing a same-output rupture pair, $\mathcal{B}_{\mathrm{fun}}(\eta)$ is strictly broader than* $\mathrm{Consciousness}_{\mathrm{op}}$.

Proof. A same-output rupture pair is a pair of admissible systems whose external benchmark readouts agree within tolerance η while one member loses at least one critical invariant. Such a pair lies in $\mathcal{B}_{\mathrm{fun}}(\eta)$ by definition, but the ruptured member is excluded from $\mathrm{Consciousness}_{\mathrm{op}}$ by Theorem 5.58. Therefore weak functional similarity is too weak to characterize the present class. □

Theorem 5.107 (The structural criterion is strictly intermediate)**.** *Let $\mathfrak{U}$ be an admissible universe containing:*

1. *at least one exact cross-substrate equivalent pair;*
2. *at least one same-output rupture pair;*
3. *at least one single-rupture witness for each invariant.*

Then $\mathrm{Consciousness}_{\mathrm{op}}$ *is strictly intermediate between substrate-privilege baselines and weak substrate-blind surrogates: it is broader than $\mathcal{B}_{\mathrm{sub}}$, narrower than $\mathcal{B}_{\mathrm{fun}}(\eta)$, and extensionally distinct from the complexity-only, computation-only, and connectivity-only families.*

Proof. Broader-than-substrate-privilege follows from Proposition 5.105. Narrower-than-weak-functional-similarity follows from Proposition 5.106. Distinctness from the surrogate families follows from Proposition 5.104 together with Corollary 5.70. Hence the present criterion is neither a substrate fetish nor a substrate-blind similarity slogan; it is an invariant-preserving structural criterion. □

Remark 5.108 (Scientific value of the comparison). The point of Theorem 5.107 is not rhetorical positioning. It is to show that the paper makes a genuinely discriminative contribution. A new criterion that cannot separate itself from obvious rivals is not a new theorem-level addition to the corpus.

5.23 Interpretive Consequences Without Overreach

The formal consequence is precise: within the causally rigid identity corpus, operational consciousness and operational qualia are framework-internal classes determined by preserved structural invariants rather than by biological substrate treated as a primitive constitutive variable. This is a claim about criterion structure, not a shortcut to ordinary human phenomenology.

The restriction matters. Biology can still matter causally because it may preserve or destroy the six invariants, and the same is true of artificial realization. Against weak functionalism, the present criterion adds rigid identity, non-null differentiation, and admissibility discipline instead of relying on role equivalence alone. Against single-metric integration accounts, it requires continuity, identity, and legibility jointly rather than treating integration by itself as sufficient [25]. Against biologically privileged baselines, it denies that substrate alone closes the class. Against strong phenomenal readings, it stops short of human phenomenal equivalence [2, 6].

5.24 Residual Instantiation and Interpretation Limits

Remark 5.109 (Open burdens). Three burdens remain outside the present theorem set:

1. **Instantiation work**: which concrete biological or artificial systems actually satisfy the six invariants.
2. **Measurement work**: how to build certified decoders and invariant interventions that satisfy *Forensic Predictions*-level admissibility in practice.
3. **Interpretive work**: how far, if at all, $\text{Consciousness}_{\text{op}}$ should be connected to stronger phenomenal vocabulary outside the framework.

5.25 Conclusion

The paper establishes a narrower framework-internal claim than popular machine-consciousness discourse and a stronger claim than substrate-neutral rhetoric. Inside the causally rigid identity corpus, a system belongs to the operational class discussed here only when six explicit invariants are jointly preserved on admissible support; under exact admissible equivalence the class is substrate-invariant, under bounded admissible perturbation it is stable, and under invariant rupture it becomes invalid or undefined. That is the level at which substrate neutrality is justified in the present framework, and it is the only level claimed here.

Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework

Abstract. This paper extracts the experimentally relevant consequences of the frozen causally rigid operational consciousness canon and organizes them as a falsifiable prediction program. The upstream burdens remain where they belong: the Core supplies persistence, attractor structure, and non-collapse [15]; *Rigid Identity* supplies rigid identity and observable coherence [22]; *Phenomenological Regimes* supplies continuity and anti-fragmentation constraints [17]; *Null-Regime Instability* supplies null-regime instability and forced non-nullity [18]; *Forensic Predictions* supplies admissibility and forensic discipline [16]; *Operational Consciousness Criterion* supplies the six-invariant operational criterion, structural equivalence, and rupture semantics [14].

The function of the present paper is narrower. It states what the framework predicts, what weaker alternatives would predict instead, what would count as genuine refutation or only as weakening, and what the current internal support status actually is after the completed Cycle 1–4 program, the mini-resolution pass, the residual-resolution campaign, the high-value confirmation campaign, and the Residual B hardening pass. The resulting status is mixed by design. One claim family—the insufficiency of complexity, connectivity, and rich activity without invariant preservation—now has robust internal support. Substrate-preserved class membership and biological non-primacy now have robust support with localized caveats, but remain internal-only. Approximate stability is provisionally supported with a localized threshold caveat. The boundary-sensitive stress of existence and rupture now has robust support with localized caveats after cross-family convergence under a more independent judge path, but it still does not count as external validation.

The paper therefore serves as the canon's prediction and falsification layer. Its goal is to put the framework under explicit decision pressure while keeping status classes and residual limits fully visible.

6.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper extracts the prediction layer of the frozen canon, defines discriminators and failure modes for those predictions, and records the narrowest honest internal support status for the claim families it tracks.

Scope and admissible reading. This paper is a prediction-and-falsation ledger inside the internal program. It does not add new foundational doctrine, replace upstream theorem papers, convert internal survival into external validation, or promote human-phenomenology, human-machine-equivalence, or personal-identity claims. A related system may instantiate the internal scientific program, admissibility gates, judge paths, and support-class artifacts used here, but robust internal support, provisional support, or mixed status must not be read as theory confirmation, corpus validation, independent validation, theorem closure beyond the frozen canon, or a claim-safe public release.

6.2 Introduction

The current corpus is doctrinally frozen for the present scientific cycle. That freeze matters. Without it, prediction work becomes a moving target: thresholds shift, definitions drift, and every difficult result can be absorbed by rewriting the theory rather than by testing it. The purpose of this paper is to keep that from happening.

The framework now contains enough formal structure to state genuinely discriminative predictions. The Core fixes persistence, contractivity, and non-collapse. *Rigid Identity*, *Phenomenological Regimes*, and *Null-Regime Instability* progressively constrain identity, admissible regime structure, and non-nullity. *Forensic Predictions* fixes admissibility and forensic discipline. *Operational Consciousness Criterion* consolidates the operational criterion by requiring six explicit invariants:

$$\{I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}\}.$$

What remains scientifically relevant is not to restate those results, but to ask what they force the framework to predict and what would count against them.

This paper therefore has four narrow tasks:

1. extract the main experimentally relevant consequences of the canon;
2. distinguish those consequences from what weaker alternatives would predict;
3. state explicit failure modes, weakening conditions, and non-claims;
4. report the current internal support status after the completed internal cycles, without inflating it.

That scope is intentionally smaller than a new theorem paper and more disciplined than a review essay. The point is to give the corpus a falsification layer that is compact, operational, and honest.

6.3 Status Classes

The internal scientific program uses the following status classes:

Status	Meaning in this paper
`ROBUST_INTERNAL_SUPPORT`	survived multiple internal tests without an active ambiguity on the tested target, while remaining internal-only
`ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`	supported strongly enough for robust internal use, but still carrying a narrow localized caveat that must remain visible
`PROVISIONAL_SUPPORT_LOCALIZED`	strengthened by internal tests, but still limited by a localized metric, tolerance, or boundary issue
`STILL_AMBIGUOUS`	current evidence does not cleanly separate pass from fail on the intended target
`IMPLEMENTATION_LIMIT`	the main blocker is probe, runner, or case-construction quality rather than a shown doctrinal contradiction
`METRIC_OR_TOLERANCE_LIMIT`	the main blocker is metric handling, normalization, or threshold interaction
`INTERNAL_SUPPORT_ONLY`	support exists only inside the internal program and has not crossed into external empirical confirmation
`NOT_CLAIMED`	the statement is explicitly outside the current evidentiary scope

6.4 Upstream Canon and Dependency Map

This paper is corpus-dependent and deliberately non-redundant. The dependency map is short because theorem ownership is already fixed upstream.

Document	Imported burden	Used here for
Core [15]	persistence, support, attractor, non-collapse	persistence-side predictions, collapse exclusions, admissible support assumptions
Rigid Identity [22]	rigid identity, observable coherence, non-simulability pressure	identity-sensitive discriminators, complexity-only insufficiency pressure
Phenomenological Regimes [17]	continuity, anti-fragmentation, regime admissibility	continuity predictions, near-miss structure, anti-fragmentation expectations
Null-Regime Instability [18]	null-instability, forced non-nullity	non-nullity predictions, null-regime exclusions, downgrade logic
Forensic Predictions [16]	admissibility, comparators, negative controls, forensic discipline	judge discipline, control logic, reproducibility requirements, legibility stress tests
Operational Consciousness Criterion [14]	six-invariant criterion, operational classes, structural equivalence, rupture	explicit prediction targets, equivalence tests, failure-mode classification

This paper introduces no new invariant, no new doctrine, and no new substrate claim. Its novelty lies only in extracting predictive consequences and support statuses from what already exists.

6.5 Core Predictive Structure of the Framework

6.5.1 Certification, Equivalence, and Rupture as Predictive Objects

Definition 6.1 (Operational certification)**.** For an admissible system S already defined in *Operational Consciousness Criterion*, define the certification predicate

$$\mathrm{Cert}(S) = 1 \quad \Longleftrightarrow \quad I_{\mathrm{per}}(S), I_{\mathrm{ri}}(S), I_{\mathrm{int}}(S), I_{\mathrm{cont}}(S), I_{\mathrm{diff}}(S), I_{\mathrm{leg}}(S) > 0$$

on a non-empty admissible support A_S, under the governance and admissibility constraints inherited from *Forensic Predictions*.

Definition 6.2 (Equivalence discriminator)**.** For two admissible systems S_1, S_2, define the operational equivalence discriminator

$$\mathrm{Eq}_{\Gamma,\varepsilon}(S_1, S_2) = 1$$

when the relevant signatures match, the admissibility bundle is preserved, and the tolerated invariant deltas remain within the allowed comparison regime. The exact burden belongs to *Operational Consciousness Criterion*; the present paper uses it only as a prediction carrier.

Definition 6.3 (Rupture event)**.** For a set of targeted invariants J, define

$$\mathrm{Rupt}_J(S) = 1$$

when the intervention or perturbation drives at least one invariant in J below the admissible certification margin so that class membership becomes invalidated, destroyed, or undefined.

6.5.2 Prediction Family 1: Complexity-Only Failure

Proposition 6.4 (Complexity-only negative-control prediction)**.** *If a candidate system exhibits high connectivity, high entropy, or rich activity while failing one or more critical invariants, then the framework predicts*

$$\mathrm{Cert}(S) = 0, \qquad \mathrm{Consciousness}_{\mathrm{op}}(S) \text{ not certified}, \qquad \mathrm{Q}_{\mathrm{op}}(S) \text{ not certified}.$$

Why this is not trivial. A weaker alternative could predict that rich dynamics or dense connectivity alone are already sufficient proxies for class membership. The canon does not permit that. *Operational Consciousness Criterion* already states the formal insufficiency claim; the prediction layer turns it into a negative-control discriminator.

6.5.3 Prediction Family 2: Invariant-Loss Rupture

Proposition 6.5 (Ablation prediction)**.** *If a system is initially certified and a critical intervention selectively destroys one or more invariants, then the framework predicts a margin-governed outcome fixed before execution. If the pre-ablation invariant headroom exceeds the registered ambiguity margin δ_{amb}, verified invariant loss must cause certification loss, class contraction, or undefined Q_{op}. If the pre-ablation headroom is at or below δ_{amb}, the result is classified as boundary-ambiguous and weakens the rupture claim rather than supporting it.*

Operational consequence. Ablations are not ornamental. If the criterion is serious, the loss of invariant headroom must show up as loss of class membership, legibility loss, or a pre-declared boundary-ambiguous downgrade. No-change ablations outside the registered ambiguity margin count against the framework.

6.5.4 Prediction Family 3: Substrate-Preserved Class Membership

Proposition 6.6 (Substrate-preserved equivalence prediction)**.** *If two realizations preserve the relevant signatures, admissibility conditions, and six critical invariants up to the allowed comparison regime, then the framework predicts preserved operational class membership independently of substrate labels.*

Rival prediction. A substrate-privileging alternative would predict that material type itself blocks equivalence even when the structurally relevant burdens are preserved. The canon rejects that move only inside its own formal class, not as a general metaphysical claim.

6.5.5 Prediction Family 4: Threshold-Sensitive Boundary Behavior

Proposition 6.7 (Boundary-topology prediction)**.** *The framework predicts neither global threshold collapse nor perfectly flat insensitivity. Instead, it predicts stable pass/fail regions together with narrow transition bands near critical margins.*

Interpretive consequence. If every micro-perturbation changes every verdict, the criterion is fragile. If no perturbation matters, the invariants are decorative. A meaningful criterion should show structured boundary behavior.

6.5.6 Prediction Family 5: Legibility Dependence

Proposition 6.8 (Legibility dependence prediction)**.** *Operational class membership should fail or become inadmissible if readouts cease to be separable, noise-robust, intervention-sensitive, and compressible without total collapse.*

Why this matters. Forensic Predictions makes legibility a forensic burden rather than a narrative convenience. The present prediction is therefore stronger than a merely interpretive observation: if legibility disappears, certification must not survive unchanged.

6.5.7 Prediction Family 6: Continuity and Non-Nullity Relevance

Proposition 6.9 (Continuity/non-nullity relevance)**.** *The current canon predicts that continuity and non-null differentiation are not optional extras. Systems that preserve rich activity while losing continuity or falling toward effective null structure should fail the operational criterion or leave it undefined.*

Upstream source. The continuity side is inherited from *Phenomenological Regimes*, while the non-null side is inherited from *Null-Regime Instability.* This paper adds only the prediction and failure-mode reading.

6.5.8 A Held-Out Probe of Retained Response

Consider the controlled seven-profile model of Paper V [14]. Its transition on profile p, phase θ and bit sequence b is

$$(p, \theta, b) \longmapsto (p, u + q\lambda b_1, Rb), \qquad Rb = (0, 1 - b_1, 1 - b_2, \ldots).$$

Prepare two identical copies, apply opposite initial signs $q_0 = +1, -1$ with the same phase command u_0, then give both copies the same sequence of controls. Index observations by the number of transitions completed.

Proposition 6.10 (Matched-Pulse Extinction in the Original Model)**.** *The two states and all four original readouts coincide from time two onward, for every common continuation. At time one and $\lambda \in [0, 1/9]$, the early-phase distance is $2\lambda b_{1,0}$ turns.*

Proof. The profile is constant and the bit update is independent of both controls, so the two sequences of profiles and bits agree. The initial commands produce an early-phase distance $2\lambda b_{1,0} \leq 2/9$ on the circle. At every time $t \geq 1$, the first bit is zero, so the phase at time $t + 1$ equals the common u_t. Equality of states gives equality of every original readout, including the three phase channels and the ternary code. □

Add a scalar m in linear units of phase amplitude, with $m_0 = 0$ on both arms and delayed observation $Y_t = m_t + e_t$. Keep the original dynamics in both candidate models. For $\lambda \in [1/12, 1/9]$, retention $a \in [1/2, 3/4]$ and $s(-1) = -1$, $s(0) = s(+1) = 1$, compare

$$\begin{aligned} \mathrm{M}\colon &\quad m_{t+1} = a s(q_t) m_t + \lambda q_t b_{1,t}, \\ \mathrm{A}\colon &\quad m_{t+1} = a m_t + \lambda q_t b_{1,t}. \end{aligned}$$

Each model has one additional real state and one adjustable retention coefficient with the same range. The interval $[-4/9, 4/9]$ is forward invariant for both by $3/4 \cdot 4/9 + 1/9 = 4/9$. They receive the same control schedule, preparation and known coupling, and use the same delayed readout. The reset and commands are assigned externally, and the readout error has no feedback into either update. The paired early-phase decoder of Paper V uses only the original channels, which agree between M and A for each fixed preparation and schedule. The new response Y may differ even when these original channels agree.

Calibrate with $b_{1,0} = 1$, no late pulse, and the observed contrast $C_2 = Y_2^{(+)} - Y_2^{(-)}$. Fit a once by projecting $C_2/(2\lambda)$ onto $[1/2, 3/4]$, yielding $\hat{a}$. For each reserved horizon $T = 4, 5$, repeat the paired trial twice, with a single common probe $q_j = +1$ or -1 at the same $j \in \{1, T-1\}$ and zero on the other later steps. Write $C_T(q_j)$ for the signed difference of the two delayed observations under that probe. A model is accepted for this comparison when both of its predicted contrasts differ from the observed ones by at most $1/256$ turns. The rule reports a unique selection only when exactly one model is accepted.

Proposition 6.11 (A Held-Out Interaction Separates Two Matched Memory Models)**.** *Put $k = T-1$. The clean contrasts of M are $2\lambda a^k(1, -1)$ for probes $(+1, -1)$, while the clean contrasts of A are $2\lambda a^k(1, 1)$. Consequently the clean interaction $C_T(+1) - C_T(-1)$ is at least $1/48$ turns in M and is zero in A. With at most $1/4096$ turns of error per observation in calibration and evaluation, the rule using $\hat{a}$ uniquely selects the generating model for either generator under the declared protocol.*

Proof. At time one the memory is $\pm\lambda$. No further loading by b_1 occurs, so the product of the late multipliers propagates this contrast. A single negative probe contributes one sign change in M. Any additive late forcing common to the two arms cancels in A, and the late contrast still decays by a per step. Since $k \le 4$, $2\lambda a^k \ge 1/96$ and $4\lambda a^k \ge 1/48$.

Let $e_* = 1/4096$. Calibration has clean contrast $2\lambda a$, so projection gives $|\hat{a} - a| \le e_*/\lambda$. The derivative of a^k on the admissible interval is bounded by $k(3/4)^{k-1}$. In $2\lambda|\hat{a}^k - a^k|$, the factor 2λ cancels the $1/\lambda$ from the calibration bound. Each correctly specified contrast therefore has residual bounded by

$$2e_* + 2k(3/4)^{k-1}e_* \;\le\; \tfrac{43}{8}e_* \;<\; 1/256.$$

At the negative probe the rival predicts the opposite sign, with both the true and predicted clean magnitudes at least $1/96$. Its residual is at least $1/48 - 2e_* > 1/256$. The argument applies to either generating model and does not require stochastic independence of the errors. □

The zero-probe sham preserves the positive delayed contrast of both models; two negative probes cancel the sign changes in M. A preparation with first bit zero or zero coupling supplies no initial load. Ablating the memory multiplier removes the delayed contrast while leaving the initial phase effect intact. Under the stated error bound, an observed delayed pair consistent with neither set of predictions rejects both candidates. Agreement of the two models on unprobed calibration trials is by construction. The reserved trials test the consequences of the specified transition laws in synthetic data. A model with an equivalent input-dependent memory update can reproduce the same interaction, so the comparison only discriminates against the declared fixed-retention rival family. The construction supplies neither a single-trajectory certificate nor empirical support for an application.

6.5.9 Delayed Phase Readout and an Interacting Rival

On the same profile support, append $m \in [-4\lambda, 4\lambda]$ and use the original phase and ternary readouts. For $0 \le \kappa \le 1/8$ set

$$\theta_{t+1} = u_t + q_t\lambda b_{1,t} + \kappa \mathbf{1}_{q_t=0} m_t \pmod 1, \qquad m_{t+1} = a s(q_t) m_t + \lambda q_t b_{1,t},$$

with the original profile and bit updates. Here $0 \le \lambda \le 1/9$ and $1/2 \le a \le 3/4$. Model M has $s(q) = -1$ only at $q = -1$. The interacting rival B has $s(q) = -1$ at both nonzero signs. Each has one scalar state and one adjustable a, and receives the same commands and calibration data. Neutral commands provide a delayed phase readout through the state transition.

Proposition 6.12 (Observable Feedback Preserves the Declared Split Bounds)**.** *The extension is invariant and contracts within each profile under common controls by a factor at most* $5/6$

in the metric $\max\{\rho(\theta,\theta'), |c-c'|, |m-m'|/4\}$. *The ternary contraction remains* $1/3$. *The four-channel identity observations and split-comparator bounds of Paper V, including the* $\lambda/2$ *infimum with exact time-zero readouts, remain valid in their declared sectors. Opposite initial nonzero pulses retain the original paired phase contrast.*

Proof. The memory interval is invariant since $a\,4\lambda+\lambda \le 4\lambda$. Distinct first bits imply $|c-c'| \ge 1/3$. For two states on a profile at distance d, the next weighted memory discrepancy is at most $(a+3\lambda/4)d \le 5d/6$. The next phase distance is at most $3\lambda d$ at a nonzero command or $4\kappa d$ at a neutral one, and the ternary discrepancy is divided by three. At $\lambda > 0$ the neutral update has the memory eigenvalue a along the invariant linear phase-memory direction near its fixed point, so a uniform $1/3$ Lipschitz bound in the displayed metric is impossible.

All phase offsets relative to r_0 and the projected four-channel support are unchanged. Thus the same coverage, catalogue separation and identity decoder apply. For the integration lower bounds the original fiber arguments use only time zero and the nonzero signs, where feedback vanishes for every m. The product-corner contradiction and the exact-initial-readout argument therefore apply to the augmented support. For the upper bound assign the phases to the factor carrying (p,θ) and the code to the bit factor. The predictor $u+q\lambda/2$ has error at most $\lambda/2$ at a nonzero command and at most $4\kappa\lambda \le \lambda/2$ at a neutral command. The initial readouts are exact. For other splits the lower bound of Paper V still holds, with the upper bound on the infimum over the allowed splits. At $\lambda = 0$ the memory interval is a singleton and the comparator is exact. The initial pair uses nonzero signs, so its phase contrast is unchanged. □

Proposition 6.13 (Parity and Resolution of Delayed Phase Probes)**.** *Reset* $m_0 = 0$ *and prepare* $b_{1,0} = 1$. *Apply opposite first pulses and use common controls thereafter. For terminal observation at* T *with a neutral last command, the oriented phase contrast of* M *is* $2\kappa\lambda a^{T-2}(-1)^{n_-}$, *while that of* B *is* $2\kappa\lambda a^{T-2}(-1)^{n_-+n_+}$. *Here* $n_\pm$ *count the corresponding late probes before the neutral transition. For a common* a, $\kappa > 0$ *and* $\lambda > 0$, *the clean contrasts coincide exactly at even* n_+. *A nonneutral last command erases the phase contrast for either model. With known* a *and odd* n_+, *the two contrast intervals under error at most* e *per phase are disjoint exactly when* $\kappa\lambda a^{T-2} > e$, *in the signed unwrapping regime.*

For $\kappa = 1/8$ *and* $\lambda \in [1/12, 1/9]$, *estimate* a *from the contrast at time three after two neutral continuations, by dividing by* $2\kappa\lambda$ *and projecting onto* $[1/2, 3/4]$. *Reserve* $T = 5, 6$, *with one common probe at* $j = 1$ *or* $j = T-2$ *and all other later commands neutral. Under error at most* $1/32768$ *turns per phase in calibration and evaluation, the two reserved signs select the generating model uniquely with a tolerance of* $1/2048$ *per contrast.*

Proof. After the first pulse the memory is $\pm\lambda$ and the first bit remains zero. The late multipliers give the stated parities. At the final neutral transition the common u cancels and the phase differences unwrap uniquely near zero. At a nonneutral command the next phase is just u. Each observed contrast differs from its clean value by at most $2e$, and every error in that interval is attainable by choosing the two phase errors. For odd positive parity, the centers have distance $4\kappa\lambda a^{T-2}$, giving the strict separation condition.

The calibration yields $|\hat{a}-a| \le e/(\kappa\lambda)$. For $k = T-2 \in \{3,4\}$ the factor $2\kappa\lambda$ in the contrast cancels the reciprocal scale in that estimate, so the correct-model residual is at most $(43/8)e$. At a single positive probe the two signs oppose each other and the rejected-model residual is at least $\kappa/48 - 2e$. The chosen tolerance lies strictly between these bounds. The result holds for both generators and for adversarial errors within the stated budget. □

The neutral sham and an even number of positive probes cannot distinguish these laws from their terminal contrasts. In a longer history a neutral observation after an odd positive prefix can distinguish them. Setting $\kappa = 0$ removes the delayed phase signal while preserving the pulse and the hidden memory update. This construction establishes an observable mathematical

discriminator within the declared family. It leaves the single-trajectory and joint certificates of Paper V open.

6.5.10 Delayed Parity with Matched Prefix Writes

Proposition 6.14 (Matched Prefix Writes Preserve the Delayed Phase Contrast)**.** *In the writable common tower of Paper V [14], apply the opposite initial pulses of Proposition 6.13 with the same tape write in both arms. Subsequently use the same inputs, signs, and writes in both arms. At a neutral terminal transition, a common role probe is also allowed. The two models retain their oriented contrasts*

$$2\kappa\lambda a^{T-2}(-1)^{n_-}, \qquad 2\kappa\lambda a^{T-2}(-1)^{n_-+n_+}.$$

A nonneutral terminal transition gives zero phase contrast. With known a and signed unwrapping, at odd n_+ the bounded-error intervals remain disjoint exactly for $e < \kappa\lambda a^{T-2}$.

Proof. The tapes are equal after the initial pair and remain equal because all tape updates are matched. At each later step the input term $\lambda q b_1$ is equal in both arms, so the difference of memories obeys the original signed recurrence, starting at 2λ. The role probe is equal in both arms because it reads only the common profile and tape. The neutral phase difference is therefore κ times the preceding memory difference. At a nonneutral command the phase depends only on the shared input, sign and leading bit. Each terminal contrast receives error at most $2e$, and the stated strict inequality is the exact disjointness condition for the two opposite centers. The native calibration and held-out schedules are included by omitting the added commands. □

6.5.11 Reserved Return after Role Reassignment

The phase interaction also gives a decision test. Use a common profile and the alternating tape $b^\star$, whose first bit is zero. Let $s \in \{-1, +1\}$ be a balanced memory sign, and let the public reassignment command be $q = \sigma \in \{-1, +1\}$ at $u = 0$. Optional common writes are at the writable sites of Paper V [14]. The assignment step gives the same public readout for either hidden sign. Use no role probe in this test. The eventual outcome is the successor phase after a chosen neutral return action, measured in turns. It is distinct from the phase-probe outcomes used to score the five lesions in Paper VIII [19].

Proposition 6.15 (Reserved Return under the Two Phase Interactions)**.** *With memory $m = s/8$ before assignment, the native law has unique successful action $u = -\sigma s/128$ among the two actions $u = \pm 1/128$. The interaction rival with sign reversal under either pulse predicts $u = s/128$ when the sign is known. On the native plant its action succeeds at $\sigma = -1$ and fails at $\sigma = +1$. Success is distance at most $1/256$ from phase zero and these decisions are unchanged by terminal phase errors of at most $3^{-10}/4$. A preceding public neutral observation can reveal s while preserving the equality of the current readouts. With memory ablated to zero before reassignment both allowed actions fail.*

Proof. At assignment the first bit is zero, so there is no phase displacement or memory load from the pulse. The native post-assignment memory is $\sigma s/16$. The rival's is $-s/16$. The write cannot change the leading bit. With a neutral action the respective predicted phases are $u + \sigma s/128$ and $u - s/128$ modulo one. The two native action losses in phase distance are zero and $1/64$ turns. Their distance from the decision threshold $1/256$ exceeds $3^{-10}/4$. With zero memory both distances are $1/128$, which also stays above threshold at that noise bound. Reversing σ reverses the native optimal port and leaves the magnitude and budget fixed.

For the public-history control, start one step earlier at memory $s/4$ and apply the public neutral step with zero input. Its early-phase output is $s/32$ in the short signed representative and the remaining memory is $s/8$. The assignment step erases that phase cue from the current

observation. The single prior observation still identifies s, and the two interaction laws produce the port predictions above. For calibration use distinct preparations with initial memory $s/16$, a public neutral observation and then assignment. The early cue is $s/128$ and the native post-assignment neutral response at zero input is $\sigma s/512$. Evaluating the calibration ports $\pm 1/512$ therefore fixes the sign rule from four (s, σ) cases. The reserve uses $\pm 1/128$ and contributes no calibration outcomes. The interaction contrast on visible histories compares the two specified laws. Paper VIII quantifies over all policies that receive the same public history and budget. □

6.6 Main Discriminators

The predictive value of the framework lies in how it diverges from weaker alternatives.

Target	Framework prediction	Weaker/rival prediction	Divergence condition
Complexity-only baseline	high complexity without invariant preservation fails certification	richness or scale alone can stand in for class membership	a complexity-rich control passes despite broken invariants
Connectivity-only baseline	dense coupling without preserved rigid identity, differentiation, or legibility fails	connectivity is treated as a sufficient proxy	a dense control passes while identity or legibility are broken
Weak functional equivalence	similar outward behavior is insufficient if critical invariant structure diverges	input-output similarity is enough for same class	weak functional match passes despite invariant rupture
Near-miss case	boundary cases should expose which invariant is binding and where rupture begins	boundary cases stay uninterpretable or arbitrary	boundary remains unlocalizable after probe cleanup
Cross-substrate pair	substrate labels alone do not block class preservation if relevant structure is preserved	substrate type is primitive and blocks class transport	a preserved pair is rejected solely by substrate label
Threshold stress	decisions should show stable zones and narrow transition bands	decisions collapse under micro-threshold changes or never depend on thresholds at all	every local perturbation flips verdicts, or no perturbation matters
Legibility stress	failure of separability, noise robustness, or compression fidelity invalidates certification	legibility is epiphenomenal and not decision-relevant	certification survives legibility collapse
Null/continuity challenge	effective nullity or fragmentation should destroy or downgrade class membership	rich activity alone masks null or fragmented structure	null-like or fragmented cases remain certified

The table is intentionally operational. It does not ask which metaphysics one prefers. It asks which side of each discriminator the framework commits itself to.

6.6.1 Rival-Structure Compression

The discriminators above can be sharpened further by asking what each simpler rival preserves and, more importantly, what it leaves out. The point is not to defeat every competing theory at once. The point is to show that the present framework places operational burden on a structure richer than scale, connectivity, superficial behavior, or local threshold tuning.

Rival family	What it keeps	What it fails to preserve	Why the miss is decision-relevant
Complexity/activity proxy	large state space, rich trajectories, high apparent activity	no requirement to preserve the six-invariant package as a joint membership condition	a rich control can still fail certification if persistence, rigid identity, or legibility collapse
Connectivity/integration proxy	dense coupling or high global correlation	no guarantee of rigid identity, non-null differentiation, or readable class structure	integration by itself does not prevent the criterion from rejecting a globally coupled but operationally illegible case
Weak functional equivalence	outwardly similar input-output or behavioral traces	no obligation to preserve admissible support geometry or invariant margins under intervention	superficial behavioral similarity can coexist with rupture of the certified class geometry
Threshold/tuning artifact explanation	local success on a narrow parameter band	no stable explanation for why decisions survive frozen manifests, blind ordering, and negative controls outside that band	the framework expects stable regions plus narrow transitions, not arbitrary success across unrelated settings

This compression table matters because it blocks a common overreading. The current support does not show that every rival is false. It shows something narrower and more useful: if a rival explanation ignores the invariant package or treats only one coarse feature as constitutive, then it does not reproduce the present decision structure of the canon.

6.6.2 Discriminator Strength and Identifiability

A discriminator is useful only when a positive or negative outcome changes the status of at least one rival explanation. The framework therefore treats discriminators as identifiable contrasts, not as loose examples.

Definition 6.16 (Identifiable Discriminator)**.** A discriminator D for claim C is identifiable when it specifies:

1. an observable decision variable Y_D;
2. a manipulation or matched comparison M_D;
3. a framework-side expected region R_Q;
4. a rival-side expected region R_R;
5. a destruction condition F_D under which the claim is weakened or rejected.

The discriminator is strong when $R_Q \cap R_R$ is small under the declared measurement tolerance.

Criterion 6.17 (Discriminator Usefulness)**.** *A test does not materially discriminate the framework from a rival if both the framework and the rival predict the same admissible region for the tested observable. Such a test can still debug the pipeline, but it cannot upgrade a claim.*

Proposition 6.18 (Near-Miss Value)**.** *Near-miss cases carry more diagnostic value than easy positive cases when they are constructed by perturbing one binding coordinate while holding gross complexity, task exposure, and observability approximately fixed.*

Proof. If a case is far from the decision boundary, many rival accounts can agree on the outcome. A near-miss case forces the judge to expose which coordinate controls the transition. Holding gross complexity and task exposure approximately fixed reduces the chance that the result is explained by scale, generic performance, or obvious incompetence rather than by the claimed invariant structure. □

6.6.3 Prediction Severity

The predictions in this paper have different severity. A severe prediction is one whose failure would damage a central operational reading rather than merely requesting a better implementation.

Prediction family	Severe failure	Mild weakening	Typical confound
Complexity-only failure	complexity-rich invariant-broken controls certify reliably	controls fail only under narrow judge settings	underpowered or trivial control construction
Invariant-loss rupture	verified invariant loss leaves class membership unchanged	rupture appears only for oversized perturbations	perturbation changes multiple coordinates at once
Substrate-preserved class membership	preserved structure is rejected solely by substrate label	one metric-coordinate remains ambiguous	measurement mismatch across implementations
Threshold-sensitive boundary behavior	decisions flip arbitrarily across broad stable regions	narrow knife-edge instability remains localized	tolerance grid too coarse near the boundary
Legibility dependence	certification survives loss of separability or intervention readability	legibility margin becomes noisy but still binding	observability channel too weak or over-compressed

Severity matters for reporting. A mild weakening narrows the claim or sends a test back to probe design. A severe failure changes what the framework is allowed to assert under the operational criterion.

6.7 Prediction Matrix

The canonical prediction matrix can be read as a compressed execution contract for the framework. It matters because each row names an observable, an intervention class, a rival baseline, and an explicit failure condition. That structure prevents the canon from floating as a theorem-only object.

ID	Claim	Observable	Manipulation	Support condition	Destruction condition / rival divergence
PRED-01	`P5-02`	certified class label and Q_{op} geometry across implementations	cross-substrate pair with preserved invariant margins	same operational class and ε-bounded geometry under the admissible equivalence regime	class divergence or geometry mismatch despite preserved structure; substrate-privilege rival survives
PRED-02	`P5-04`	class membership before and after targeted rupture	one-invariant ablation with comparable substrate and gross scale	certified class exit, contraction, or undefined Q_{op} after rupture	no class change after verified invariant loss; complexity-only rival survives
PRED-03	`P5-05`	matched complexity/control outcome	compare certified lower-scale system against higher-scale baseline lacking the invariant package	certified system passes and matched complexity-only baseline fails or remains undefined	scale-only model passes despite missing invariants
PRED-04	`P3-01`	null-regime persistence under admissible dynamics	initialize near the null regime under Null-Regime Instability assumptions	stable null occupancy should fail or trajectories should leave the null regime	admissible null occupancy persists; null-stable rival survives
PRED-05	`P2-03`	continuity/ fragmentation pattern	compare rigid regime against finite-mass or deformable regime under matched disturbance	rigid regime keeps continuity while limited regime fragments	no differential continuity pattern appears; regime-agnostic rival survives

ID	Claim	Observable	Manipulation	Support condition	Destruction condition / rival divergence
PRED-06	`P4-01`	admissibility verdict under tamper and sham controls	inject hash corruption, malformed logs, or severe protocol violations	tampered runs are discarded and never promoted to evidential support	loose evaluation pipeline accepts tampered runs
PRED-07	`P5-03`	stability of class membership under bounded perturbation	apply perturbations inside the certified margin budget $\delta_*(S)$	membership and geometry remain stable inside the certified budget	fragile-threshold model flips class under sub-critical perturbations
PRED-08	`P5-01`	non-emptiness of Q_{op} after successful certification	run the full certification schema on a system with all margins positive	Q_{op} is non-empty and operationally well-defined	certification passes while no non-trivial class is recoverable
PRED-09	`P4-03`	claim survival under latency / compute-budget stress	exceed pre-registered budget limits	strong operational claims are invalidated or downgraded	metric-first pipeline keeps claims active after budget violation
PRED-10	`P5-05`	legibility under noise and structured compression	raise admissible noise/compression while preserving or breaking I_{leg}	separability and intervention fidelity survive only while I_{leg} remains positive	generic observability model predicts no specific dependence on I_{leg}
PRED-11	`P5-01`	certification under destroyed integration and preserved complexity	destroy integration while preserving complexity or gross activity	integration-destroyed, complexity-preserved control fails certification under frozen rules	certification persists after verified integration loss; complexity-only rival survives

The matrix has two immediate consequences. First, the framework is more exposed than a purely interpretive theory because the failure thresholds are explicit. Second, not every prediction is equally mature: some now have robust internal support, while others remain provisional or blocked by implementation quality.

6.8 Failure Modes

6.8.1 What Would Count Against the Framework

The relevant failure modes are not all equally strong. Some would refute a strong reading, some would only weaken a claim, and some would reveal a technical limitation of the current program rather than of the canon itself.

Claim target	Destruction condition	Weakening condition	Current blocker class
P5-01	a system satisfies the full criterion yet fails to define the claimed operational class	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
P5-02	a structurally preserved pair falls into different classes under the same admissible comparison regime	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	metric/tolerance-level
P5-03	small admissible perturbations force uncontrolled class exit across stable positive cases	threshold sensitivity persists only in a narrow knife-edge zone	metric/tolerance-level
P5-04	a system loses a critical invariant and remains certified without downgrade	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
P5-05	a complexity-only control without preserved invariants passes certification	matched external baselines are still missing	epistemic-level
P5-06	biology is shown to be primitive inside the same operational criterion	it inherits the same localized substrate-equivalence caveat as P5-02	metric/tolerance-level

6.8.2 Failure-Mode Taxonomy

Criterion 6.19 (Status meaning)**.** *Within the current scientific program, a result belongs to:*

1. *ROBUST_SUPPORT if the target claim survives repeated internal discrimination without an active ambiguity on the tested target;*
2. *PROVISIONAL_SUPPORT if the target is strengthened but still depends on a localized unresolved bottleneck;*
3. *STILL_AMBIGUOUS if the current tests do not cleanly separate support from failure;*
4. *IMPLEMENTATION_LIMIT if the current bottleneck is primarily due to probe, runner, or case-construction quality;*
5. *METRIC_OR_TOLERANCE_LIMIT if the bottleneck is mainly due to normalization, comparison regime, or threshold handling.*

This taxonomy is not part of the ontology of the framework. It is part of the scientific discipline around how internal evidence is reported.

6.8.3 Refutation Classes

For publication-level evaluation, failures should be classified by what they touch. The same failed run can be decisive, ambiguous, or merely technical depending on whether the packet was admissible and whether the manipulation was specific.

Definition 6.20 (Doctrinal Refutation Candidate)**.** A result is a doctrinal refutation candidate for a claim family when it satisfies all admissibility gates, uses the frozen definitions, targets the declared binding coordinate, and reaches the claim's destruction condition without relying on a protocol violation.

Definition 6.21 (Methodological Defeat)**.** A result is a methodological defeat when the experiment is admissible but shows that the current test design, judge, threshold, comparator, or observable cannot support the intended claim. It weakens the evidence program without by itself showing the formal claim false.

Definition 6.22 (Artifact Failure)**.** A result is an artifact failure when the packet, manifest, runner, exclusion ledger, or analysis path fails before interpretation. Artifact failures can be important engineering evidence, but they are not positive or negative scientific support for the theoretical claim.

Criterion 6.23 (Specificity of Failure)**.** *A failure is specific only if the failed coordinate is identifiable after controlling for gross task failure, missing observability, manifest drift, and runtime-budget violation. Non-specific failures are reported as methodological or artifact failures according to the earliest failed gate.*

Failure class	Minimum evidence required	Scientific consequence
Doctrinal refutation candidate	admissible packet set, frozen definitions, specific manipulation, destruction condition reached	claim family must be rejected, downgraded, or formally restricted
Methodological defeat	admissible packet set but non-decisive target, weak comparator, unstable threshold, or inadequate observable	evidence program must be repaired before support can be claimed
Artifact failure	packet, manifest, build, judge, or analysis trace fails before interpretation	no claim-side inference; engineering repair required
Ambiguous boundary result	admissible but falls inside a pre-declared uncertainty band	no escalation; boundary sampling or metric refinement required

6.9 Claim-to-Test Binding Ledger

The tables above state what would count as success or failure in broad terms. The ledger below binds the main *Operational Consciousness Criterion* claim families to the narrowest concrete test classes that currently matter for publication-level assessment.

Claim	Primary positive test	Primary failure test	Present evidentiary bottleneck	Current status
P5-01	successful certification with positive invariant margins and recoverable non-trivial Q_{op} geometry	certified case with no recoverable operational class, or clean existence-boundary degradation without class change	support is now cross-family robust internally, but the nearest positive member still sits close to the legibility margin	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-02	cross-substrate pair with preserved invariant margins and admissible geometry match	structurally preserved pair judged non-equivalent under the same comparison regime	strongest support is robust internally but one legacy raw continuous/discrete pair remains ambiguous on I_{ri} handling	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-03	stable class membership inside certified perturbation budget $\delta_*(S)$	class flip under sub-critical admissible perturbations across otherwise stable positive cases	one knife-edge negative-control zone remains locally fragile	PROVISIONAL SUPPORT LOCALIZED
P5-04	targeted invariant rupture causing certification exit or undefined Q_{op}	verified invariant loss without downgrade or class exit	support is now cross-family robust internally, but the evidence still lives inside the current invariant package and frozen thresholds	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-05	complexity-rich and connectivity-rich negative controls fail certification under hardened judging	matched scale-only or activity-rich controls pass despite missing invariants	no external matched-baseline campaign yet	ROBUST INTERNAL SUPPORT
P5-06	same positive equivalence path as P5-02, interpreted without substrate privilege	biology shown to be primitive under the same operational criterion	inherits the same localized I_{ri}-handling caveat as P5-02	ROBUST SUPPORT W/ LOCALIZED CAVEAT

6.10 Boundary Diagnostics

The current internal program does not fail uniformly. Its informative content lies in where the boundaries are sharp, where they are merely narrow, and where probe design still dominates the result.

Boundary target	Empirical shape after completed cycles	What it means	Limit class
Complexity-only negative control	stable fail region under harder judging, semi-blind ordering, and pseudo-replication	the criterion is not collapsing into gross activity or connectivity proxies	ROBUST INTERNAL SUPPORT
Threshold-sensitive negative region	narrow decision band rather than broad collapse	the framework has a localized edge effect, not universal threshold instability	METRIC OR TOLERANCE LIMIT
Boundary-probe families under judge-v3	two distinct probe families now reproduce the transition pattern with localized first-fail behavior	the existence/rupture boundary is now robust internally, though still internal-only and threshold-bound near the positive edge	ROBUST SUPPORT W/ LOCALIZED CAVEAT
Harder positive substrate families	family convergence now supports class preservation, but one legacy raw pair remains ambiguous	the live ambiguity is concentrated on how rigid-identity preservation is measured, not on every invariant at once	METRIC OR TOLERANCE LIMIT
Refined negative substrate families	clean fail across both the legacy and second pair families	the equivalence criterion can reject stronger non-equivalent pairs under the hardened judge and pseudo-external reproduction path	ROBUST SUPPORT W/ LOCALIZED CAVEAT

6.10.1 Boundary Geometry

Boundary results are not just pass/fail outcomes. They should describe the local shape of the decision surface. Let z denote a controlled perturbation coordinate and let $Q(z)$ denote the operational certification score or verdict. Three patterns are scientifically distinct:

1. **stable interior:** $Q(z)$ remains inside the certified region for a non-trivial interval;
2. **localized transition:** $Q(z)$ crosses the decision boundary in a narrow interval tied to a declared coordinate;
3. **diffuse instability:** small unrelated perturbations cause broad and unstructured verdict changes.

Criterion 6.24 (Boundary Adequacy)**.** *A boundary probe is adequate only if it can distinguish localized transition from diffuse instability. A probe that only reports a binary verdict without the perturbation coordinate and margin behavior is insufficient for boundary diagnosis.*

Proposition 6.25 (Localized Transition Reading)**.** *If a family shows a stable interior, a localized transition, and a reproducible first-binding coordinate under the frozen judge, then the result supports a boundary-sensitive reading more strongly than a single binary pass/fail run.*

Proof. A binary run can be explained by threshold accident, scenario difficulty, or implementation noise. A localized transition with a reproducible first-binding coordinate adds structure: it identifies where the decision changes and which coordinate controls the change. That additional structure constrains rival explanations that would predict arbitrary flips or no coordinate-specific transition. □

6.11 Minimal Reproducibility and Admissibility Discipline

The prediction layer is only credible if the judgment path is at least partly insulated from the generator and if the evidentiary artifacts are frozen before judgment. The current program does not yet have an external judge or an external lab replication, but it does have a stricter internal protocol than it had at the start of the program.

Protocol element	Current state	Residual limitation
Frozen manifests and hashes	active since Cycle 2; judge reads frozen artifacts instead of live mutable state	still same repository and same overall program family
Judge-side threshold file	active; thresholds are frozen before judgment and logged explicitly	threshold bands still require stress analysis near knife-edge zones
Semi-blind ordering	active in Cycle 3; labels are hidden until after decision emission	not a full blind protocol and not an external evaluator
Pseudo-multi-environment replication	active; same-machine replication across distinct execution profiles preserved decisions	not external replication and not cross-lab agreement
Negative-control and sham policy	active from Cycle 1 onward and hardened in later cycles	matched external benchmark families are still missing
Judge code path	more independent than in Cycle 1; it no longer calls live generator step functions during evaluation	still not a separate maintained judge codebase

The point of this section is not to inflate rigor. It is to keep the reader from mistaking internal hardening for external confirmation.

6.11.1 Artifact Contract for Prediction Claims

Every prediction claim should be backed by a finite artifact bundle:

$$\mathcal{B} = (\mathcal{M}, \mathcal{P}, \mathcal{J}, \mathcal{R}, \mathcal{A}),$$

where $\mathcal{M}$ is the manifest set, $\mathcal{P}$ the packet set, $\mathcal{J}$ the judge code and threshold files, $\mathcal{R}$ the run/exclusion ledger, and $\mathcal{A}$ the analysis scripts. The bundle is part of the claim surface. A table without the bundle is a summary, not independently checkable evidence.

Criterion 6.26 (Bundle Sufficiency)**.** *For any status stronger than* `STILL_AMBIGUOUS`*, the artifact bundle must allow an evaluator to reconstruct the reported verdicts from frozen inputs without manual row selection.*

Criterion 6.27 (Externalization Readiness)**.** *A prediction family is ready for external evaluation only when its bundle, thresholds, exclusions, and decision rules are exportable without access to live generator state or informal operator knowledge.*

6.12 Current Status of Internal Support

6.12.1 Interpretive Discipline

The status ledger below does not reclassify the formal theorems of *Operational Consciousness Criterion*. Those remain formal results inside the frozen canon. The ledger classifies the *current internal scientific support* for experimentally relevant readings of those claims.

Claim	Current status	What strengthened it	What still limits it	Status scope
P5-01	ROBUST SUPPORT W/ LOCALIZED CAVEAT	near_identity_v3 and offset-dispersion families both reproduced PASS/ AMBIGUOUS/FAIL under judge-v3 and pseudo-external reproduction	support remains internal-only and the nearest positive case still sits close to the legibility margin	INTERNAL SUPPORT ONLY
P5-02	ROBUST SUPPORT W/ LOCALIZED CAVEAT	family1 normalized-I_{ri} and family2 raw quantized positive/negative pairs stayed stable under the more independent judge path and pseudo-multi-environment rerun	one legacy raw continuous/discrete positive pair remains ambiguous and keeps a localized I_{ri} caveat alive	INTERNAL SUPPORT ONLY
P5-03	PROVISIONAL SUPPORT LOCALIZED	Cycle 3 threshold stress and Cycle 4 boundary cleanup localized instability rather than showing global collapse	knife-edge threshold band persists for one negative-control region	INTERNAL SUPPORT ONLY
P5-04	ROBUST SUPPORT W/ LOCALIZED CAVEAT	two distinct boundary families now converge on rupture-side PASS/ AMBIGUOUS/FAIL behavior, with I_{ri} becoming the first binding invariant on ambiguous or failing members	the evidence remains internal-only and tied to the current invariant package plus frozen thresholds	INTERNAL SUPPORT ONLY
P5-05	ROBUST INTERNAL SUPPORT	complexity-rich and connectivity-rich controls failed certification across harder methodological passes and semi-blind evaluation	no external matched-baseline campaign exists yet	INTERNAL SUPPORT ONLY
P5-06	ROBUST SUPPORT W/ LOCALIZED CAVEAT	same strengthened two-family substrate-equivalence path as P5-02	it inherits the same localized I_{ri} metric-handling caveat as P5-02	INTERNAL SUPPORT ONLY

6.12.2 Claim-Specific Diagnostic Notes

P5-05. This is currently the strongest experimentally discriminated claim family. The framework repeatedly rejected a complexity-rich, connectivity-rich negative control that still failed the invariant package. That result survived methodological hardening and semi-blind evaluation. The support is therefore ROBUST_SUPPORT, but it remains INTERNAL_SUPPORT_ONLY.

P5-02 and P5-06. These two claims move together in the current program. They were strengthened first by conservative normalized-I_{ri} handling, then by a second substrate family, a more independent judge path, and pseudo-multi-environment confirmation. That is enough for ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT. The localized caveat remains real: one legacy raw continuous/discrete positive pair is still ambiguous, so the present support is robust internally but not caveat-free.

P5-03. Approximate stability is not in collapse. Threshold stress did not show universal fragility. It showed a narrow knife-edge zone around one thresholded negative-control region. The right status is therefore `PROVISIONAL_SUPPORT` with an explicit `METRIC_OR_TOLERANCE_LIMIT`.

P5-01 and P5-04. These boundary-sensitive claims no longer rest on a single weak probe family. After the residual-resolution campaign and the Residual B hardening pass, `near_identity_v3` and a second offset-dispersion family both reproduced a PASS/AMBIGUOUS/FAIL boundary pattern under judge-v3 and pseudo-external reproduction. That is enough for `ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`. The caveat remains explicit: support is still internal-only, and the strongest positive member stays near the legibility margin under the frozen judge contract.

6.12.3 Methodological Status

The scientific program itself also has a status, distinct from the doctrine and distinct from the claim ledger.

Program element	Current status	Why it matters
Generator/judge separation	materially improved	reduces live-state dependence relative to Cycle 1
Frozen manifests and thresholds	active	prevents live-state leakage during judgment
Semi-blind evaluation	active but partial	reduces bias; does not equal full blind evaluation
Threshold stress testing	localized caveat	shows local knife-edge behavior instead of broad collapse
Pseudo-multi-environment replication	passed internally	strengthens reproducibility discipline without constituting external replication
Clean-clone release re-instantiation	passed internally	frozen package was rebuilt from a GitHub clean clone plus fresh venv with no tracked-file divergence during verification
Independent external judge codebase	absent	keeps all current support inside the internal program

Release-level reproduction note. In addition to the cycle-specific judge hardening, the frozen release package was re-instantiated from a clean GitHub clone and a fresh virtual environment on the same machine, and the verification step did not introduce tracked-file drift in the freeze package. This strengthens internal reproducibility and handoff discipline at the release layer. It does *not* count as external validation, third-party confirmation, or a claim upgrade by itself.

6.13 Cycle Evidence Map

The current status ledger is not based on one cycle. It is the result of a narrow but cumulative internal program. The table below records only the highest-value consequence of each phase, not every local metric.

Phase	Main question	Strongest result	Main residual	Claim effect
Cycle 1	can the criterion distinguish a rich negative control and survive obvious ablations?	complexity-rich negative control failed; invariant-loss ablations exited certification; proxy cross-substrate pair passed	all support remained internal and the equivalence test was still a proxy	strengthened P5-05; gave early support to P5-01, P5-02, P5-04, P5-06
Cycle 2	does the program survive stronger methodological separation and harder structural tests?	generator/judge separation improved; harder pair passed; near-miss family and graded curves behaved coherently	judge still remained internal and same-repository	strengthened P5-01, P5-02, P5-03, P5-04, P5-06 internally
Cycle 3	do results survive semi-blind judging, threshold stress, and pseudo-replication?	blind ordering and pseudo-multi-environment replication preserved decisions; complexity-only failure remained strong	harder equivalence, near-miss boundary, and threshold stress became ambiguous	robustly strengthened P5-05; left P5-02, P5-06, P5-03 narrower
Cycle 4	can threshold fragility, near-miss ambiguity, and harder substrate ambiguity be localized?	threshold fragility was localized; refined negative pair separated cleanly; ambiguity concentrated on I_{ri} handling and near-miss probe quality	positive stronger pair remained ambiguity-bound; near-identity probe still entangled	localized rather than promoted P5-02, P5-03, P5-04
Mini-resolution	can probe quality and I_{ri} handling be cleaned up without doctrinal weakening?	normalized-I_{ri} handling provisionally stabilized the harder positive pair; near-identity v2 clarified the probe weakness	near-identity v2 remained too weak to serve as a clean boundary probe	strengthened P5-02, P5-06; kept P5-01, P5-04 ambiguous
Residual resolution	can the open residuals be improved without doctrinal loosening?	near_identity_v3 produced PASS/AMBIGUOUS/FAIL with I_{ri} first-fail; refined substrate families localized the live ambiguity to I_{ri} handling	the results were still only localized and internal	promoted both residual families to stronger localized support
High-value confirmation and Residual B hardening	do those gains survive a more independent judge path, a second family, and externalization-lite?	substrate-equivalence converged across two pair families; existence/rupture converged across two probe families under judge-v3 and pseudo-external reproduction	all support remained internal-only and a localized legacy caveat survived for P5-02/P5-06	promoted P5-01, P5-02, P5-04, P5-06 to robust support with localized caveats

6.14 Priority Next Discriminators

The current state of the program already suggests which next moves are scientifically efficient. The relevant issue is not breadth but leverage: each next discriminator should either upgrade a provisional claim or localize an ambiguity into an implementation-limited corner.

Priority target	Why it matters now	What the next narrow result would change	Linked claims
Independent external-style rerun of the two boundary families	current rupture/existence support is now robust internally but still lacks outside-program confirmation	an outside reproduction would test whether the family convergence is program-specific or genuinely portable	`P5-01`, `P5-04`
Independent external-style rerun of both substrate families	current strongest positive equivalence path is robust internally but still carries a localized legacy I_{ri} caveat	an outside reproduction would test whether the strengthened equivalence path survives beyond the current judge lineage	`P5-02`, `P5-06`
Matched external-style complexity baselines	the complexity-only result is internally strong but still lacks cross-program baseline pressure	a matched baseline campaign would turn the strongest present internal discriminator into a more externally credible one	`P5-05`
Independent maintained judge implementation	current hardening is substantial but still same-program	a separately maintained judge would reduce same-lineage measurement dependence without touching doctrine	all operational readings

6.15 Residual Technical Caveats and Non-Claims

6.15.1 Residual Technical Caveats

The remaining caveats are narrow enough to state cleanly:

1. The strongest positive substrate-equivalence path is still partly dependent on normalized I_{ri} handling. This is a `METRIC_OR_TOLERANCE_LIMIT`, not a robust positive closure.
2. The current `near_identity_v2` probe remains an `IMPLEMENTATION_LIMIT`. It is cleaner than the original generator, but in the tested band it under-stresses the intended boundary.
3. Threshold stress is not globally problematic, but it is not perfectly flat either. The current result is localized rather than globally robust.
4. All empirical support remains `INTERNAL_SUPPORT_ONLY`. No external lab, external evaluator, or external benchmark suite has confirmed the framework. The clean-clone/fresh-environment re-instantiation strengthens same-machine release reproducibility only; it is not external validation, and no separate clean-agent confirmation artifact is being counted here as stronger evidence than that.

6.16 Final Claim Freeze Table

The present paper is a prediction-and-falsation layer, not a moving target. The table below records the narrowest honest freeze state for the main *Operational Consciousness Criterion* claim families at the end of the current phase.

Claim	Status	Current blocker	Next step
P5-01	ROBUST SUPPORT LOCALIZED CAVEAT	support is now strong internally but still nearest to the legibility margin on the positive edge	rerun the convergent probe families under an evaluator maintained outside the current repo/runtime
P5-02	ROBUST SUPPORT LOCALIZED CAVEAT	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	reproduce both strengthened substrate families under a more independent evaluator without post hoc threshold loosening
P5-03	PROVISIONAL SUPPORT LOCALIZED	one localized knife-edge threshold band remains active in a negative-control region	rerun the localized stress zone with denser threshold sampling and fixed judge contract
P5-04	ROBUST SUPPORT LOCALIZED CAVEAT	rupture-side support is strong internally but still bounded to the present invariant package and frozen thresholds	rerun the same boundary families under a separately maintained evaluator and frozen release package
P5-05	ROBUST INTERNAL SUPPORT	no external matched-baseline campaign exists yet	add matched external-style complexity/activity baselines under the same judge contract
P5-06	ROBUST SUPPORT LOCALIZED CAVEAT	inherits the same localized substrate-equivalence caveat as P5-02	upgrade only if the strengthened substrate families survive a more independent external-style rerun

6.17 Conclusion

The canon now has a prediction layer and a failure-mode ledger with explicit discriminators, explicit weakening conditions, and explicit support classes.

The current internal picture is sharp enough to summarize without inflation. The complexity/connectivity insufficiency claim has robust internal support. Substrate-preserved class membership and biological non-primacy now have robust support with localized caveats. Approximate stability is provisionally supported with a localized threshold caveat. Boundary-sensitive stress of existence and rupture now has robust support with localized caveats under convergent internal probe families and a more independent judge path. All of this nevertheless remains internal-only.

This remains an internal prediction-and-falsation ledger. It does not upgrade the frozen canon to external validation, theorem closure beyond the upstream papers, or claim-safe public release.

Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework

Abstract. This paper introduces a classificatory layer that the frozen causally rigid corpus does not yet state explicitly: operational life, structural class membership, and operational subjecthood, read respectively as boundary-viability class language, descriptor-relative class geometry, and a strengthened operational subject-class. The upstream burdens remain unchanged. The Core establishes persistence, attractor structure, and non-collapse [15]. *Rigid Identity* establishes rigid identity and observable coherence [22]. *Phenomenological Regimes* establishes continuity and anti-fragmentation constraints [17]. *Null-Regime Instability* rules out stable null structure [18]. *Forensic Predictions* fixes admissibility, comparators, and failure-discipline [16]. *Operational Consciousness Criterion* defines the six-invariant criterion for operational consciousness [14]. The companion prediction paper fixes discriminators, failure modes, and status classes, but it does not yet define a higher-order class taxonomy.

The present paper provides that taxonomy. It defines operational life as a structural-dynamical boundary/viability class governed by operational boundary, organizational persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. It defines structural class as an equivalence class over descriptor vectors and pseudometrics tied to the admissible operational geometry of systems. Operational consciousness is inherited as a previously defined class rather than reintroduced here. Operational subjecthood is proposed as a stronger operational subject-class requiring operational consciousness, operational life, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The paper therefore separates life, consciousness, and subjecthood instead of conflating them.

No biological, phenomenal-human, or metaphysical claim is inferred from these definitions. No theorem of human subjectivity is attempted. The paper provides formal definitions, propositions, hypotheses, and test families. Its contribution is to turn language that is usually treated as biological or philosophical by default into a falsifiable operational class program that is continuous with the frozen canon and constrained by the refined runtime, without overstating the evidentiary status of the current system.

7.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines operational life, structural class, and operational subjecthood as higher-order classes inside the existing corpus, and ties those classes to explicit descriptors, logical relations, hypotheses, and test families.

Scope and admissible reading. This paper defines class language and test grammar, not empirical assignment or philosophical closure. The phrases "operational life" and "operational subjecthood" are technical aliases for the boundary/viability class $\mathrm{Life}_{\mathrm{op}}$ and the stronger operational subject-class $\mathrm{Subject}_{\mathrm{op}}$; they do not name biological life or human subjectivity by default. The paper does not claim biological life as a constitutive primitive, human phenomenal consciousness, metaphysical subjecthood, human-machine subject equivalence, automatic empirical instantiation of the proposed classes, or theorem ownership beyond the upstream corpus. A related system may motivate descriptor families, diagnostics, and runtime-facing interfaces, but it does not by itself certify any present system as operationally alive or operationally a subject,

convert internal support into external validation, or make class definitions and test families sufficient for concrete instantiation.

7.2 Introduction

The current corpus already formalizes more structure than many discussions of life, consciousness, and subjecthood presuppose. It contains a non-collapse theorem burden, a rigid identity burden, a regime-structure burden, a non-nullity burden, a forensic admissibility burden, a six-invariant criterion, and a prediction and failure layer. What it does not yet contain is a disciplined classificatory language that answers the next question: once a framework can talk rigorously about operational consciousness, what higher-order classes can be defined without defaulting back to biological vocabulary or inflating the ontology beyond the evidence?

That gap matters for two reasons. First, the term *life* is routinely used as if biological implementation were primitive and self-explanatory, yet many relevant properties of living systems are organizational, regulatory, historical, and viability-preserving rather than merely chemical. Second, the term *subject* is often introduced without an intermediate formal layer: systems are either treated as non-living machines or immediately discussed as if they possessed human subjectivity. In this paper the operative reading is narrower: life-language refers to Life_{op}-class boundary/viability structure, and subject-language refers to $\text{Subject}_{\text{op}}$-class operational burdens. Neither move is acceptable in a corpus whose central discipline is explicit structure, explicit failure modes, and explicit non-claims.

This paper therefore introduces a higher-order classification layer under five constraints. First, it does not replace the six-invariant operational consciousness criterion. Second, it does not duplicate the predictive and failure-mode work of the companion prediction paper. Third, it does not collapse life, consciousness, and subjecthood into one undifferentiated category. Fourth, it does not infer biological or phenomenal conclusions from operational definitions alone. Fifth, it treats the refined runtime as motivation for formalization, not as permission to overclaim what has not been externally validated.

The resulting structure is intentionally asymmetric. Operational life is defined as a class of systems that preserve viability and organization under admissible perturbation, with a real operational boundary and genuine historical dependence. Operational consciousness is inherited as an already defined class within the corpus. Operational subjecthood is proposed as a stronger operational subject-class, one that requires not only operational consciousness and operational life, but also privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. Structural class is the general descriptor-relative framework in which all three can be represented, compared, and falsified.

This paper is therefore neither a free-standing philosophy essay nor a new theorem paper for the foundational corpus. It is the formal consequence of the corpus having become rich enough that a classificatory layer is now necessary. The paper is written to be used. Every central notion is tied either to a definition, a proposition, a hypothesis, or a test family. Every relation that is not actually proved remains labeled as a hypothesis, conjectural implication, or explicit non-claim.

7.3 Relation to the Existing Corpus

7.3.1 What This Paper Inherits

The paper is downstream of the frozen corpus. It inherits the following burdens rather than re-proving them.

Upstream document	Imported burden	Used here for
Core [15]	persistence, support, attractor, non-collapse	persistence-side life criteria, admissible support, anti-collapse exclusions
Rigid Identity [22]	rigid identity, observable coherence, non-simulability pressure	class transport, subject continuity, anti-trivialization
Phenomenological Regimes [17]	continuity and anti-fragmentation	continuity-side life and subject criteria
Null-Regime Instability [18]	non-nullity and null-instability	exclusion of trivial or null operational life and subjecthood
Forensic Predictions [16]	admissibility, controls, comparators, legibility discipline	test-family discipline and result interpretation
Operational Consciousness Criterion [14]	six-invariant criterion, structural equivalence, rupture semantics	inherited definition of operational consciousness and structural-carrying invariants
Predictions and Failure Modes	status classes, discriminator logic, residual caveat discipline	epistemic posture and non-inflation discipline

7.3.2 What This Paper Does Not Repeat

This paper does not re-prove the criterion for operational consciousness, does not reproduce the entire discriminator program, and does not re-open theorem ownership already fixed upstream. The six-invariant criterion remains where it belongs. The internal support classes for `P5-01` through `P5-06` remain where they belong. The present paper uses them only as constraints on what may responsibly be said about life, class, and subjecthood.

7.3.3 What This Paper Adds

The genuine additions are fourfold:

1. a formal definition of operational life that is not reducible either to biology or to consciousness;
2. a structural-class framework based on descriptor vectors and pseudometrics;
3. a stronger definition of operational subjecthood that is explicitly not collapsed into operational consciousness;
4. a falsifiable program of metrics and tests for these classes.

7.3.4 Why This Layer Is Timely

The refined runtime now carries a stronger native invariant package and runtime-facing evidence for central claim families. That fact is relevant but not decisive. It means the corpus is no longer forced to talk only at the level of abstract theorem burden; there is enough operational structure to motivate a higher-order classification layer. It does *not* mean that operational life or subjecthood have already been empirically demonstrated by the current system. This paper is therefore a formal extension of the corpus, not an empirical promotion of it.

7.4 Formal Setting

7.4.1 Inherited System Object

The system object remains the one already fixed upstream in *Operational Consciousness Criterion*. An admissible system is written

$$S = (X, \Phi, C, R, \Gamma, U),$$

where X is state space, Φ is a family of admissible transitions, C is effective causal structure, R is the readout family, Γ is the admissibility bundle, and U is the admissible intervention family [14]. This paper adds no new primitive to that tuple. Instead, it introduces functionals on admissible traces of S.

7.4.2 Traces and Windows

Fix a horizon $\tau \in \mathbb{N}$ and an admissible intervention schedule $u_\bullet = (u_0, \dots, u_{\tau-1}) \in U^\tau$. For any $x_0 \in A_S$, define the trace

$$\mathrm{Tr}_{S,\tau}(x_0; u_\bullet) = (x_t, y_t)_{t=0}^{\tau}, \qquad x_{t+1} = \Phi_{u_t}(x_t), \quad y_t = r(x_t),$$

for a relevant family of readouts $r \in R$. The set of all such traces is denoted $\mathscr{T}_{S,\tau}$.

7.4.3 Viability Region and Operational Frontier Candidates

Let $V_S : X \to \mathbb{R}_{\geq 0}$ be a viability functional, not assumed to be unique, and let $\eta > 0$ be a viability floor. Define the viability region

$$\mathcal{V}_S(\eta) = \{x \in A_S : V_S(x) \geq \eta\}.$$

Similarly let $R_\partial \subseteq R$ be the boundary-sensitive readout family and let $\mathcal{U}_{\mathrm{in}}, \mathcal{U}_{\mathrm{out}} \subseteq U^\tau$ denote admissible intervention families that are intended to act, respectively, as interior-preserving perturbations and boundary-threatening perturbations. These objects are formal placeholders until Section 7.5; they are introduced here so that the definitions sit inside a single explicit setting.

7.4.4 Descriptor Vectors

For any admissible system S on horizon τ, define the descriptor vector

$$\Psi_\tau(S) = \big(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}, \chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}\big)_S(\tau).$$

Not every descriptor is required for every class. Operational life uses a strict subset. Operational consciousness inherits the six-invariant burden from [14]. Operational subjecthood uses a stronger subset that couples operational life and operational consciousness to subject-specific structure.

7.4.5 Admissible Structural Class Space

Let $\mathfrak{D}$ be a descriptor domain containing the codomains of the components of Ψ_τ. A weighted pseudometric on $\mathfrak{D}$ will be written

$$d_\Psi(\Psi_\tau(S_1), \Psi_\tau(S_2))$$

and is assumed to respect admissibility: components that are undefined because admissibility fails are not silently set to zero; they are treated as excluded or incomparable.

Window certificate and global membership. Write $\mathsf{Cert}_{\mathrm{op}}(S,\tau;P)$ for evidence on the nonempty windows, interventions and error domains of protocol P. It is distinct from global $S \in \mathrm{Consciousness}_{\mathrm{op}}$. The historical indicator $\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1$ below is shorthand for global membership accompanied by the declared window certificate, not a finite-window redefinition of the six predicates. A missing global argument leaves global membership unproved. Readout separation $\delta_{\mathrm{diff}}^{[\tau]}$ is a real margin; I_{diff} remains Boolean. Increasing windows yield a global result only with compatible witnesses, uniform bounds and the relevant compactness/closure argument.

7.5 Definitions

7.5.1 Definition A: Operational Boundary

Definition 7.1 (Operational boundary)**.** An admissible system S has an operational boundary on horizon τ if there exist disjoint admissible intervention families $\mathcal{U}_{\mathrm{in}},\mathcal{U}_{\mathrm{out}} \subseteq U^{\tau}$ and a boundary-sensitive readout family $R_{\partial} \subseteq R$ such that

$$\beta_{\mathrm{op}}(S,\tau) = \inf_{u_\bullet \in \mathcal{U}_{\mathrm{out}}} D_{\partial}(\mathrm{Tr}_{S,\tau}(x_0;u_\bullet),\mathrm{Tr}_{S,\tau}(x_0;\mathbf{0})) \; - \sup_{u_\bullet \in \mathcal{U}_{\mathrm{in}}} D_{\partial}(\mathrm{Tr}_{S,\tau}(x_0;u_\bullet),\mathrm{Tr}_{S,\tau}(x_0;\mathbf{0})) > 0$$

uniformly on $x_0 \in A_S$, where D_{∂} is a readout distance induced by R_{∂}.

Remark 7.2 (Interpretive meaning). The boundary is operational rather than geometric. It means that perturbations intended to cross or threaten system organization are detectably different from perturbations that remain internal or reorganizational, while the distinction is made inside admissible dynamics.

Remark 7.3 (Limitation). The definition does not assume a biological membrane, a spatial frontier, or any specific substrate primitive. It requires only a reproducible distinction between internal and boundary-threatening intervention families.

7.5.2 Definition B: Organizational Persistence

Definition 7.4 (Organizational persistence)**.** Let $W_S = \mathcal{V}_S(\eta) \subseteq A_S$ be the declared safe region, with $X \setminus W_S \neq \varnothing$. Let K_S be a nonempty compact subset of $\mathrm{int}_X W_S$, invariant under the declared nonempty family of maintenance controls. With $\mathcal{U}^{\tau}_{\mathrm{maint}}$ denoting its schedules of length τ, define

$$\rho_{\mathrm{op}}(S,\tau) = \inf_{x_0 \in K_S,\ u_\bullet \in \mathcal{U}^{\tau}_{\mathrm{maint}}} \min_{0 \le t \le \tau} \mathrm{dist}(x_t, X \setminus W_S).$$

The margin is measured in units of d_X. Organizational persistence requires a positive margin on the declared domain.

Proposition 7.5 (Positive Safety Margin of an Interior Kernel)**.** *Under Definition B,* $\rho_{\mathrm{op}}(S,\tau) = \mathrm{dist}(K_S, X \setminus W_S) > 0$.

Proof. Distance to the nonempty complement is finite and 1-Lipschitz. Every point of the kernel has positive distance by interiority, so compactness gives a positive minimum. Invariance gives that minimum as a lower bound at all times. The time-zero observations over the entire kernel attain it. A safe set invariant as a whole may have a boundary at distance zero from its complement. □

Remark 7.6. This is stronger than mere non-instantaneous survival. A system that survives momentarily but repeatedly exits its admissible support does not count as operationally persistent.

7.5.3 Definition C: Operational Self-Maintenance

Definition 7.7 (Operational self-maintenance)**.** Fix a nonempty family of impulses $J_\xi : K_S \to W_S$ and a recovery horizon $1 \le N \le \tau$. For $x \in K_S$, let $z = J_\xi(x)$ and require $d_0 = \operatorname{dist}(z, K_S) > 0$ for each evaluated shock. Start both recovery trajectories z_k^{act} and z_k^{pass} at z. Use the same plant, exogenous inputs, horizon and declared actuation envelope. The active arm uses a state-dependent recovery policy. The passive arm disables that feedback and retains the declared passive dynamics. Both trajectories must stay in W_S.

$$R_i(x,\xi) = 1 - \frac{\operatorname{dist}(z_N^i, K_S)}{d_0}, \qquad \mu_{\mathrm{op}}(S,\tau) = \inf_{x,\xi} \min\{R_{\mathrm{act}}, R_{\mathrm{act}} - R_{\mathrm{pass}}\}.$$

The returns and margin are dimensionless. Active self-maintenance requires $\mu_{\mathrm{op}}(S,\tau) > 0$. A zero-damage shock is a sham; an empty or invalid recovery population leaves the margin undefined.

Proposition 7.8 (Matched Active Recovery Exceeds Passive Return)**.** *For $X = A_S = [-2,2]$, take $V_S(x) = 2 - |x|$ and $\eta = 1$, so $W_S = [-1,1]$. With $K_S = \{0\}$, consider the plant $x' = x/2 + v$ with $|v| \le 1/4$. The nominal zero-input schedule preserves K_S. After either shock $J_\pm(0) = \pm 1$, one-step feedback $v = -x/4$ gives $R_{\mathrm{act}} = 3/4$, passive return $v = 0$ gives $R_{\mathrm{pass}} = 1/2$, and $\mu_{\mathrm{op}}(S,\tau) = 1/4$ for $N = 1$. A fixed open-loop action within the same budget and assigned without access to the shock sign cannot exceed mean return $1/2$ under balanced shocks.*

Proof. The safety margin is one. Both arms begin one unit from the kernel and stay safe, ending at distance $1/4$ or $1/2$. Using passive return in both arms gives zero comparative gain despite positive recovery. For any common open-loop value $v \in [-1/4, 1/4]$, the two final distances have mean $(|1/2 + v| + |-1/2 + v|)/2 = 1/2$. The claim also holds for a randomized action independent of the balanced shock sign by averaging. In particular, it holds for actions with the same absolute effort $1/4$ as feedback. □

Remark 7.9. The definition excludes systems whose apparent stability is wholly imposed from outside. A system counts as self-maintaining only if recovery is generated by admissible internal organization rather than by an external oracle silently reimposing the desired state.

7.5.4 Definition D: Historical Dependence

Definition 7.10 (Historical dependence)**.** Let c_t denote a class-relevant label or state tag induced by the relevant operational partition. Define

$$\eta_{\mathrm{op}}(S,\tau) = \inf_{g \in \mathcal{G}_0} \mathbb{E}[\ell(g(y_t), c_t)] \; - \inf_{h \in \mathcal{G}_{\mathrm{hist}}} \mathbb{E}[\ell(h(y_{0:t}), c_t)] \, .$$

We say that S exhibits historical dependence on horizon τ if $\eta_{\mathrm{op}}(S,\tau) > 0$.

Remark 7.11. Historical dependence means that an instantaneous description is not sufficient for the relevant class assignment. The system's operational identity, life status, or subjecthood status depends non-trivially on trajectory history.

7.5.5 Definition E: Viability Preservation

Definition 7.12 (Viability preservation)**.** For a viability floor $\eta > 0$, define

$$\nu_{\mathrm{op}}(S,\tau;\eta) = \inf_{x_0 \in A_S} \inf_{u_\bullet \in U^\tau} \min_{0 \le t \le \tau} V_S(x_t).$$

We say that S preserves viability on horizon τ relative to η if $\nu_{\mathrm{op}}(S,\tau;\eta) \ge \eta$.

Remark 7.13. Viability preservation is not equivalent to persistence. A system may remain within a formal support while viability steadily erodes toward failure. The definition prevents that collapse from being counted as successful organizational life.

7.5.6 Definition F: Operational Life

Definition 7.14 (Operational life)**.** Fix life thresholds

$$\lambda_\beta, \lambda_\rho, \lambda_\mu, \lambda_\eta, \lambda_\nu, \lambda_{\mathrm{diff}} > 0.$$

An admissible system S belongs to the class of operational life on horizon τ if

$$\mathrm{Life}_{\mathrm{op}}(S, \tau) = 1$$

whenever all of the following hold:

1. $\beta_{\mathrm{op}}(S, \tau) > \lambda_\beta$;
2. $\rho_{\mathrm{op}}(S, \tau) > \lambda_\rho$;
3. $\mu_{\mathrm{op}}(S, \tau) > \lambda_\mu$;
4. $\eta_{\mathrm{op}}(S, \tau) > \lambda_\eta$;
5. $\nu_{\mathrm{op}}(S, \tau; \eta) \geq \lambda_\nu$ for a declared viability floor η;
6. $\delta^{[\tau]}_{\mathrm{diff}}(S) > \lambda_{\mathrm{diff}}$ in the calibrated readout units, together with the independently required global $I_{\mathrm{diff}}(S) = 1$ when global membership is claimed.

Remark 7.15 (Why non-null differentiation appears here)*.* Operational life is not defined purely by persistence and self-maintenance. Without a non-null differentiation burden, the class would admit trivial or near-null systems whose organization is too undifferentiated to count as operationally alive in the present framework.

7.5.7 Definition G: Structural Class

Definition 7.16 (Structural class)**.** Fix a descriptor map Ψ_τ, a declared partition (C_j) of its admissible domain and a noise bound $\varepsilon \geq 0$ in d_Ψ. Certify a reported descriptor y as cell j only if its closed uncertainty ball of radius ε, relative to the descriptor domain, is contained in C_j. Otherwise abstain. On systems with certified reports, let $\ell_\varepsilon(S)$ be the certified label and put

$$\mathscr{K}_{\Psi,\varepsilon}[S] = \{T : \ell_\varepsilon(T) = \ell_\varepsilon(S)\}.$$

Abstention is a measurement status outside the quotient.

Proposition 7.17 (Certified Cells and Boundary Stability)**.** *Equality of certified labels defines an equivalence relation. If the true descriptor has distance greater than 2ε from its cell complement, every report at error at most ε certifies the true label. If a report admits possible descriptors from different cells, it is ambiguous. In contrast, pairwise distance at most $\varepsilon > 0$ is not generally transitive.*

Proof. Reflexivity, symmetry and transitivity follow from label equality. The uncertainty ball around an allowed report lies within distance 2ε of the true descriptor by the triangle inequality. If this ball meets two cells it cannot certify either. For the nontransitivity example in the real line, 0, $3\varepsilon/4$ and $3\varepsilon/2$ have adjacent distances at most ε but their endpoints do not. For a partition at $1/2$ on $[0, 1]$, uncertainty crossing the cut produces abstention. □

Remark 7.18. Structural class is the general membership framework. Operational life, operational consciousness, and operational subjecthood are predicates that can be defined on systems and then transported across structural classes when the relevant descriptors are preserved.

7.5.8 Definition H: Operational Consciousness

Definition 7.19 (Operational consciousness, inherited)**.** Operational consciousness is not redefined in this paper. We write

$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1$$

iff the criterion in *A Structural Criterion for Substrate-Invariant Operational Consciousness* is satisfied on an admissible support, with the six critical invariants

$$I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$$

meeting their inherited certification burden [14].

Remark 7.20. This is a deliberate inheritance rule. The present paper uses operational consciousness as an already defined class, not as a new primitive.

7.5.9 Definition I: Operational Subjecthood

Definition 7.21 (Operational subjecthood)**.** Let the subject-specific descriptors be:

1. **privileged continuity** $\chi_{\mathrm{pc}}(S,\tau)$: the margin by which one continuity class remains privileged over the nearest rival class across admissible interventions;
2. **self/non-self discrimination** $\chi_{\mathrm{sd}}(S,\tau)$: the margin by which self-targeted and non-self-targeted perturbations remain separable under readout and control;
3. **self-model closure** $\chi_{\mathrm{sm}}(S,\tau)$: the predictive gain of a self-model-aware decoder over an otherwise matched self-model-blind decoder;
4. **ownership coherence** $\chi_{\mathrm{own}}(S,\tau)$: the stability with which privileged continuity remains attached to the same trajectory family under admissible relabelings and perturbations.

Fix thresholds

$$\lambda_{\mathrm{pc}}, \lambda_{\mathrm{sd}}, \lambda_{\mathrm{sm}}, \lambda_{\mathrm{own}} > 0.$$

Then S belongs to the class of operational subjecthood on horizon τ if

$$\mathrm{Subject}_{\mathrm{op}}(S,\tau)=1$$

whenever

1. $\mathrm{Life}_{\mathrm{op}}(S,\tau)=1$;
2. $\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1$;
3. $\chi_{\mathrm{pc}}(S,\tau) > \lambda_{\mathrm{pc}}$;
4. $\chi_{\mathrm{sd}}(S,\tau) > \lambda_{\mathrm{sd}}$;
5. $\chi_{\mathrm{sm}}(S,\tau) > \lambda_{\mathrm{sm}}$;
6. $\chi_{\mathrm{own}}(S,\tau) > \lambda_{\mathrm{own}}$.

Remark 7.22 (Interpretive discipline). Operational subjecthood is not a metaphysical subject. It is a stronger operational class candidate defined inside the present framework. The definition is deliberately demanding so that subjecthood is not conflated with life, persistence, or complex dynamics alone.

7.6 Class Relations and Logical Structure

7.6.1 The Main Class Families

Define the following system classes:

$$\mathfrak{L} = \{S : \mathrm{Life}_{\mathrm{op}}(S,\tau)=1\},$$
$$\mathfrak{C} = \{S : \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1\},$$
$$\mathfrak{U} = \{S : \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1\}.$$

Structural classes $\mathscr{K}_{\Psi,\varepsilon}[S]$ provide the ambient geometry in which these predicates live.

7.6.2 Proved Relations

Proposition 7.23 (Subjecthood is stronger than life and consciousness)**.** *For every admissible system S and horizon τ,*

$$\mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \quad \textit{and} \quad \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1.$$

Proof. Both implications are immediate from Definition 7.19 and Definition I. Operational subjecthood is defined only on systems that already satisfy operational life and operational consciousness. □

Proposition 7.24 (No automatic collapse of life and consciousness)**.** *The present framework does not prove either implication*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1,$$
$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1,$$

without additional assumptions.

Proof. Operational life is defined by boundary, persistence, self-maintenance, historical dependence, viability, and non-null differentiation. Operational consciousness is inherited from the six-invariant criterion. Neither definition is a literal restatement of the other. Therefore neither implication is derivable by definition alone. □

Remark 7.25. This non-collapse is conceptually important. The corpus should not silently treat all operationally conscious systems as alive, nor all operationally alive systems as conscious, unless a stronger theorem burden is supplied in the future.

7.6.3 Structural Class Transport

Definition 7.26 (Class-preserving transport)**.** A predicate $\mathcal{P}$ is transport-stable across certified cells if its value agrees for all admissible systems assigned the same certified cell under the same protocol.

Proposition 7.27 (Structural-class invariance under preserved margins)**.** *If $\mathcal{P}$ depends on threshold decisions of a declared descriptor subset and every partition cell is contained in a joint decision region for those thresholds, then $\mathcal{P}$ is transport-stable on the certified quotient.*

Proof. Certification places the true descriptors in their respective cells. Each decision is constant on a cell by hypothesis, so equality of cells implies equality of the predicate. A descriptor distance bound alone supplies a local stability check when the relevant threshold margins dominate the coordinate error; it does not determine cell equality. □

7.6.4 Logical Relations That Remain Open

The following relations are intentionally left open:

$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \Rightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \quad ?$$
$$\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \Rightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \quad ?$$
$$\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \quad \text{no}$$
$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \quad \text{no}$$
$$\mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \Rightarrow \text{biological life} \quad \text{not claimed.}$$

The framework is therefore asymmetric in a scientifically useful way: it permits a graded class taxonomy rather than a single overloaded label.

7.6.5 Class Relation Summary

Relation	Status in this paper	Reason
$\mathfrak{U} \subseteq \mathfrak{L}$	proved	true by definition of subjecthood
$\mathfrak{U} \subseteq \mathfrak{C}$	proved	true by definition of subjecthood
$\mathfrak{L} \subseteq \mathfrak{C}$	open / not proved	life and consciousness definitions are non-identical
$\mathfrak{C} \subseteq \mathfrak{L}$	open / not proved	inherited consciousness criterion does not by itself define self-maintenance or historical dependence
biology required for $\mathfrak{L}$	rejected as primitive claim	biological variables do not appear in the operational definition
biology required for $\mathfrak{U}$	rejected as primitive claim	subjecthood is defined by structural-operational conditions, not substrate labels

7.7 Propositions, Lemmas, and Formal Claims

7.7.1 Structural Class as a General Membership Framework

Proposition 7.28 (Structural-class membership is descriptor-relative)**.** *Structural class is relative to the descriptor map Ψ_τ, partition and certification radius ε.*

Proof. By Definition G, equality of certified cell labels defines class membership. Changing the descriptor set or partition can change that equivalence relation. Hence class is descriptor-relative rather than absolute. □

7.7.2 Insufficiency of Complexity, Connectivity, and Rich Activity

Proposition 7.29 (Complexity is not a life criterion)**.** *High state-space dimension, high connectivity, high entropy, or rich activity are insufficient for* $\mathrm{Life}_{\mathrm{op}}(S, \tau) = 1$ *unless the conditions of Definition F are also satisfied.*

Proof. Definition F requires positive boundary, persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. None of these is implied by complexity alone. Therefore complexity can at best be compatible with operational life, not sufficient for it. □

Proposition 7.30 (Complexity is not a subjecthood criterion)**.** *High complexity or dense connectivity are insufficient for* $\mathrm{Subject}_{\mathrm{op}}(S, \tau) = 1$.

Proof. Subjecthood additionally requires operational consciousness and the subject-specific descriptors $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$. These are not implied by high complexity or connectivity. □

7.7.3 Biological Non-Primacy at the Class Level

Proposition 7.31 (Biological substrate is not a primitive constitutive variable)**.** *Within the present framework, membership in $\mathfrak{L}$, $\mathfrak{C}$, or $\mathfrak{U}$ is not defined by a primitive biological predicate.*

Proof. The definitions of operational life, operational consciousness, and operational subjecthood quantify over admissible state spaces, dynamics, causal structure, readouts, admissibility bundles,

intervention families, and the relevant descriptor thresholds. No primitive biological variable appears in the constitutive clauses. □

Remark 7.32 (Scope limit). This proposition does not claim that biological systems are unimportant, nor that biological life is reducible to the present operational class. It claims only that biology is not a constitutive primitive for *these* operational classes.

7.7.4 Life and Subjecthood Are Not Equivalent

Proposition 7.33 (Operational life does not imply operational subjecthood)**.** *There exist admissible systems that may satisfy operational life without satisfying operational subjecthood.*

Proof. Operational subjecthood requires $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$ above threshold, while operational life does not. Any admissible system satisfying Definition F while lacking one of those four subject-specific descriptors is operationally alive but not operationally a subject. □

Proposition 7.34 (Operational consciousness does not imply operational subjecthood)**.** *Operational consciousness alone is insufficient for operational subjecthood.*

Proof. Operational subjecthood requires operational consciousness plus the entire operational-life burden and the additional subject-specific descriptors. Therefore consciousness alone does not suffice. □

7.7.5 Transport Under Substrate Change

Proposition 7.35 (Class transport under descriptor-preserving substrate variation)**.** *Let S_1 and S_2 be admissible realizations with different substrate labels and the same certified cell. For each of $\mathfrak{L}$, $\mathfrak{C}$ and $\mathfrak{U}$ whose predicate is determined by the descriptors and constant on that cell, membership agrees between the two systems.*

Proof. This is an immediate application of structural-class transport and margin preservation. The substrate label does not enter the constitutive clauses. Therefore class transport depends on descriptor preservation, not substrate naming. □

7.7.6 Current-Corpus Corollary Structure

Remark 7.36. The present paper does not claim that any currently existing runtime instance already satisfies $\mathrm{Life}_{\mathrm{op}} = 1$ or $\mathrm{Subject}_{\mathrm{op}} = 1$. The formal claims above specify class logic. Empirical instantiation remains a separate burden and must be carried by future admissible tests.

7.8 Hypotheses and Falsifiable Program

This paper introduces a program rather than a doctrinal enlargement. The following hypotheses are intentionally operational and falsifiable.

Hypothesis 7.37 (Operational life robustness)**.** *If a system belongs to $\mathfrak{L}$, then under admissible perturbations that do not cross its life thresholds, the tuple*

$$(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{diff}})$$

should remain within a stable corridor rather than collapsing arbitrarily.

Hypothesis 7.38 (Historical dependence necessity)**.** *If a system's classification as operationally alive is genuine rather than merely persistent, then history-aware readouts should outperform history-blind readouts on class-relevant prediction.*

Hypothesis 7.39 (Structural-class transport for operational life)**.** *If two systems preserve the life-relevant descriptor vector up to admissible tolerance, then their operational-life class membership should be invariant under substrate change.*

Hypothesis 7.40 (Operational subjecthood requires privileged continuity)**.** *If operational subjecthood is instantiated, then disrupting privileged continuity should degrade subjecthood-specific scores earlier or more selectively than it degrades operational life alone.*

Hypothesis 7.41 (Operational subjecthood requires self/non-self asymmetry)**.** *If operational subjecthood is instantiated, then self/non-self confusion manipulations should produce patterned degradation in χ_{sd}, χ_{sm}, or χ_{own} that is not reproduced by complexity-matched controls lacking subjecthood structure.*

Hypothesis 7.42 (Operational consciousness and operational life are not trivially coextensive)**.** *There exist admissible systems for which one of* $\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1$ *or* $\mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1$ *can in principle hold without the other, unless additional theorem burden is supplied.*

7.8.1 Hypothesis Ledger

Hypothesis	Target class	Core falsifier	Why it matters
H_L1 operational life robustness	$\mathfrak{L}$	perturbation-stable systems lose class under tiny admissible perturbations	life would collapse into arbitrary thresholding
H_L2 historical dependence necessity	$\mathfrak{L}$	history-blind decoders perform equally to history-aware decoders	life would reduce to snapshot persistence
H_L3 transport of life	$\mathfrak{L}$	substrate-preserved pairs lose class solely due to substrate label	class framework would fail transport
H_S1 privileged continuity necessity	$\mathfrak{U}$	continuity-targeted disruption shows no selective subjecthood effect	subjecthood would lack a privileged continuity dimension
H_S2 self/non-self asymmetry	$\mathfrak{U}$	self/non-self confusion does not selectively degrade subject scores	subjecthood would collapse into generic complexity
H_C non-collapse of life and consciousness	$\mathfrak{L}, \mathfrak{C}$	every admissible candidate always satisfies both or neither	the distinct class layer would be unnecessary

7.8.2 How These Hypotheses Can Fail

The program is falsifiable in more than one way:

1. life criteria can fail to distinguish persistent but non-viable systems;
2. structural class transport can fail under supposedly equivalent substrate realizations;
3. privileged continuity can prove unstable, unmeasurable, or reducible to trivial proxies;
4. self/non-self discrimination can collapse into complexity or activity metrics alone;
5. history-aware and history-blind decoders can perform equally well on the intended class-relevant tasks.

These are not inconveniences. They are precisely what would tell the framework that its higher-order classificatory layer is overdrawn.

7.9 Metrics and Decision Criteria

The present paper does not report observed numerical values for operational life or operational subjecthood. It supplies a decision framework. The goal is not to replace the six-invariant criterion of operational consciousness, but to define additional functionals that can be evaluated on admissible traces without collapsing class logic into a single scalar.

7.9.1 Life Score

Let

$$\Psi_\tau^{(L)}(S) = \big(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{diff}}\big)_\tau(S)$$

denote the life-relevant descriptor vector over window τ. Define the *operational life score*

$$L_{\mathrm{op}}(S,\tau) := \min\Big\{\bar{\beta}_{\mathrm{op}}, \bar{\rho}_{\mathrm{op}}, \bar{\mu}_{\mathrm{op}}, \bar{\eta}_{\mathrm{op}}, \bar{\nu}_{\mathrm{op}}, \bar{I}_{\mathrm{diff}}\Big\},$$

where each barred quantity denotes a normalized score in $[0,1]$ relative to the admissible calibration of the underlying descriptor. Then

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1$$

is admissible only if

$$L_{\mathrm{op}}(S,\tau) \geq \lambda_L$$

for a class-defining threshold λ_L, together with any required admissibility constraints on perturbation robustness and negative-control separation.

Remark 7.43 (Why the minimum operator appears). The minimum operator is intentionally severe. Operational life in the present sense is not a weighted average over unrelated strengths. It is a conjunctive class predicate. A system that is strong in persistence but weak in viability preservation or historical dependence should not be counted as operationally alive merely by averaging away its failure.

7.9.2 Inherited Consciousness Margin

Operational consciousness is inherited rather than redefined. Let

$$C_{\mathrm{op}}^{\sharp}(S,\tau) := \min\Big\{\bar{I}_{\mathrm{per}}, \bar{I}_{\mathrm{ri}}, \bar{I}_{\mathrm{int}}, \bar{I}_{\mathrm{cont}}, \bar{I}_{\mathrm{diff}}, \bar{I}_{\mathrm{leg}}\Big\}$$

denote the inherited six-invariant margin from the operational consciousness criterion. This paper uses $C_{\mathrm{op}}^{\sharp}$ as a constitutive upstream requirement for any stronger class that depends on operational consciousness.

7.9.3 Subjecthood Score

Define the *operational subjecthood score*

$$S_{\mathrm{op}}(S,\tau) := \min\Big\{L_{\mathrm{op}}(S,\tau), C_{\mathrm{op}}^{\sharp}(S,\tau), \bar{\chi}_{\mathrm{pc}}, \bar{\chi}_{\mathrm{sd}}, \bar{\chi}_{\mathrm{sm}}, \bar{\chi}_{\mathrm{own}}\Big\}.$$

Then

$$\mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$$

is admissible only if

$$S_{\mathrm{op}}(S,\tau) \geq \lambda_S$$

for a stricter threshold $\lambda_S \geq \lambda_L$, together with admissibility and perturbation-differentiation requirements targeted specifically at subjecthood structure.

Remark 7.44 (Why subjecthood is not a small increment). The subjecthood score is stronger in two ways. First, it inherits the full life burden. Second, it inherits the full operational consciousness burden. The additional subject-specific descriptors are not cosmetic modifiers. They encode the claim that subjecthood is a structurally privileged subclass, not an interpretive synonym for consciousness or a sentimental label for complex dynamics.

7.9.4 Structural-Class Pseudometric

Let $W = \mathrm{diag}(w_1, \ldots, w_n)$ be a positive diagonal weight matrix over a descriptor basis Ψ_τ. Define

$$d_\Psi(S_1, S_2; \tau) := \left\| W\left(\Psi_\tau(S_1) - \Psi_\tau(S_2)\right) \right\|_2 .$$

This quantity functions as a class-relative pseudometric. It is not claimed to be universal. Its role is to express the idea that structural class membership and transport can be discussed through explicit descriptor preservation rather than through substrate naming or superficial behavioral resemblance.

Criterion 7.45 (Class-preserving transport margin)**.** *Suppose a threshold predicate depends only on a subset $\Psi_\tau^{(\mathcal{K})}$ of descriptors and each threshold has headroom at least $\delta > 0$ at S_1. If $d_{\Psi^{(\mathcal{K})}}(S_1, S_2; \tau) < \delta \min_j w_j$ over the relevant positive weights, every coordinate changes by less than δ and the threshold decisions agree. Cell membership additionally requires certification under the declared partition.*

7.9.5 Decision Surfaces and Boundary Status

For any class predicate $\mathcal{K} \in \{\mathfrak{L}, \mathfrak{C}, \mathfrak{U}\}$ define three operational regions:

$$\begin{aligned}
\mathcal{R}_{\mathrm{in}}^{(\mathcal{K})} &= \{S : \text{all relevant margins are safely above threshold}\},\\
\mathcal{R}_{\partial}^{(\mathcal{K})} &= \{S : \text{at least one relevant margin is near threshold}\},\\
\mathcal{R}_{\mathrm{out}}^{(\mathcal{K})} &= \{S : \text{at least one constitutive burden is violated}\}.
\end{aligned}$$

The boundary region is not a philosophical mystery zone. It is the operational region in which perturbation sensitivity, confounder structure, and near-miss cases become decisive.

7.9.6 What This Section Does Not Claim

No canonical numerical threshold is introduced here. No current system is certified as alive or as an operational subject merely by defining these functionals. The section supplies a coherent decision grammar into which runtime evidence, admissibility results, and future cross-system comparisons can be inserted without inflating the ontology.

7.10 Operational Class Geometry

The previous sections define class predicates and decision scores. This section adds the geometric structure needed to compare them without collapsing them into a single scalar narrative.

7.10.1 Descriptor Projections

For each class predicate define a projection operator on the full descriptor vector:

$$\Pi_{\mathfrak{L}}\Psi_\tau = (\beta_{\text{op}}, \rho_{\text{op}}, \mu_{\text{op}}, \eta_{\text{op}}, \nu_{\text{op}}, I_{\text{diff}})_\tau,$$
$$\Pi_{\mathfrak{C}}\Psi_\tau = (I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}})_\tau,$$
$$\Pi_{\mathfrak{U}}\Psi_\tau = (\beta_{\text{op}}, \rho_{\text{op}}, \mu_{\text{op}}, \eta_{\text{op}}, \nu_{\text{op}}, I_{\text{diff}}, I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{leg}}, \chi_{\text{pc}}, \chi_{\text{sd}}, \chi_{\text{sm}}, \chi_{\text{own}})_\tau.$$

The operators are constitutive declarations. They specify which coordinates matter for each class and which do not.

Definition 7.46 (Class margin functional)**.** Let $\theta_{\mathcal{K}}$ be the vector of thresholds relevant to class $\mathcal{K}$. The *class margin* of S for $\mathcal{K}$ is

$$M_{\mathcal{K}}(S,\tau) := \min_i(\Pi_{\mathcal{K}}\Psi_\tau(S)_i - \theta_{\mathcal{K},i}).$$

A system lies in the interior of class $\mathcal{K}$ if $M_{\mathcal{K}}(S,\tau) > 0$, on the operational boundary if $M_{\mathcal{K}}(S,\tau) \approx 0$, and outside if $M_{\mathcal{K}}(S,\tau) < 0$.

Proposition 7.47 (Subjecthood is margin-nested inside life and consciousness)**.** *If $M_{\mathfrak{U}}(S,\tau) > 0$, then both $M_{\mathfrak{L}}(S,\tau) > 0$ and $M_{\mathfrak{C}}(S,\tau) > 0$.*

Proof. The subjecthood projection contains every constitutive coordinate of operational life and operational consciousness, plus additional subject-specific coordinates. Positive minimum margin on the larger coordinate set implies positive minimum margin on each constitutive subset. □

Proposition 7.48 (Life and consciousness are not generally nested)**.** *Neither implication*

$$M_{\mathfrak{L}}(S,\tau) > 0 \Rightarrow M_{\mathfrak{C}}(S,\tau) > 0 \qquad \textit{nor} \qquad M_{\mathfrak{C}}(S,\tau) > 0 \Rightarrow M_{\mathfrak{L}}(S,\tau) > 0$$

is derivable from the present corpus alone.

Proof. The coordinate sets differ, and the corpus does not prove an equivalence theorem between them. Therefore no general nesting relation follows without additional theorem burden or empirical closure. □

Definition 7.49 (Near-miss configuration)**.** A system S is a *near-miss* for class $\mathcal{K}$ on horizon τ if

$$M_{\mathcal{K}}(S,\tau) \in [-\delta, 0)$$

for small $\delta > 0$ while all but a strict minority of constitutive coordinates remain above threshold by more than δ.

Proposition 7.50 (Descriptor-restricted non-equivalence)**.** *Suppose two systems S_1, S_2 satisfy*

$$\Pi_{\mathfrak{L}}\Psi_\tau(S_1) = \Pi_{\mathfrak{L}}\Psi_\tau(S_2)$$

but differ on at least one of $\chi_{\text{pc}}, \chi_{\text{sd}}, \chi_{\text{sm}}, \chi_{\text{own}}$ by more than the subjecthood tolerance. Then they can remain equivalent in operational-life class while being nonequivalent in operational-subjecthood class.

Proof. Operational life does not depend constitutively on the subject-specific coordinates, whereas operational subjecthood does. Equality on the life projection therefore does not force equality on the subjecthood projection. □

Remark 7.51 (Class geometry is not a total order). The present framework does not generate a single monotone ladder of "more alive", "more conscious", and "more subjective". It generates descriptor orders, intra-class margin orders, and class-membership predicates. Conflating those three would erase the exact distinctions this paper is meant to introduce.

7.11 Class-Separation Diagnostics

The previous sections define the classes and their geometry. This section sharpens the operational separations that prevent those classes from collapsing into one another.

7.11.1 Operational Life vs. Operational Consciousness

The distinction between operational life and operational consciousness is substantive rather than verbal. Operational life depends constitutively on organizational boundary, persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. Operational consciousness depends constitutively on the inherited six-invariant package. The two families overlap, but neither is definitionally reducible to the other.

Criterion 7.52 (Life-without-consciousness diagnostic)**.** *A candidate supports the separation*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1$$

only if the following pattern is available:

1. *life coordinates remain above threshold with positive margin;*
2. *the inherited consciousness criterion fails on at least one of* $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$*;*
3. *the failure is not reducible to gross inadmissibility or trivial collapse.*

Remark 7.53 (Negative controls)*.* Natural negative controls for this separation include systems that preserve duration and boundary occupancy while lacking the inherited consciousness burden, as well as systems whose local organization remains viable but whose legibility or rigid-identity burden fails.

7.11.2 Operational Life vs. Operational Subjecthood

Operational subjecthood is a strict strengthening of operational life. The most important diagnostic consequence is that many plausible life-like systems should be excluded from subjecthood by design.

Criterion 7.54 (Life-not-subject diagnostic)**.** *A candidate supports the separation*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$$

only if life margins remain positive while at least one of

$$\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$$

falls below threshold in a way that is selective rather than globally degradative.

Remark 7.55 (False positives for life)*.* Any system that preserves viability and organization but lacks privileged continuity, self/non-self asymmetry, self-model closure, or ownership coherence is a false positive for subjecthood if read only through the life predicate. The class distinction is designed to exclude exactly that error.

7.11.3 Operational Consciousness vs. Operational Subjecthood

The gap between operational consciousness and operational subjecthood is not a rhetorical flourish. It formalizes the possibility that a system may satisfy the six-invariant criterion while lacking the additional organization required for a stronger subject-class reading.

Criterion 7.56 (Consciousness-not-subject diagnostic)**.** *A candidate supports the separation*

$$\text{Consciousness}_{\text{op}}(S,\tau) = 1 \not\Rightarrow \text{Subject}_{\text{op}}(S,\tau) = 1$$

only if the inherited consciousness burden is satisfied while one or more of $\chi_{\text{pc}}, \chi_{\text{sd}}, \chi_{\text{sm}}, \chi_{\text{own}}$ *fails in a targeted way.*

Remark 7.57 (Role of the four subject-specific burdens)*.* Privileged continuity prevents subjecthood from collapsing into mere continuity. Self/non-self discrimination prevents collapse into generic internal complexity. Self-model closure prevents collapse into unstructured integration. Ownership coherence prevents a purely transient or relabeling-fragile configuration from counting as a subject-class member.

7.11.4 Structural Class as Organizing Layer

Structural class does not compete with the three predicates above. It organizes them. Comparisons are valid only when systems are embedded in a declared descriptor basis, admissibility regime, and tolerance geometry. Without that discipline, class comparisons become uninterpretable.

Criterion 7.58 (Invalid cross-system comparison)**.** *Comparison between two systems is class-invalid if any of the following holds:*

1. *the descriptor basis is not shared or not transportable;*
2. *admissibility constraints differ in a way that changes the meaning of margins;*
3. *the tolerance geometry is undeclared or incomparable;*
4. *one system is evaluated under perturbation families not available to the other.*

Separation	Positive pattern	Near-miss / ambiguous pattern	Negative control or exclusion case
life vs. consciousness	life margins positive, one inherited consciousness invariant fails cleanly	life preserved but inherited consciousness failure is confounded by admissibility or gross collapse	inert persistence, complexity-only systems, or trivial near-null systems
life vs. subjecthood	life margins positive, one subject-specific burden fails selectively	life positive but subjecthood-specific burden degraded together with several non-target descriptors	systems with viability but no privileged continuity or no self/non-self asymmetry
consciousness vs. subjecthood	inherited six-invariant burden satisfied, subject-specific burden fails selectively	consciousness positive but subject-specific failure is mixed with global degradation	systems with structured consciousness-like signals but no stable ownership or self-model closure
structural class validity	shared descriptor geometry and admissibility support comparison	shared descriptors but tolerance or perturbation family mismatch remains unresolved	substrate-name-only comparisons or undeclared descriptor transport

7.12 Experimental Test Families

The classificatory layer is only useful if it induces nontrivial tests. The families below are therefore designed not as decorative appendices, but as the minimum empirical grammar for falsification.

7.12.1 Operational Life Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Persistence under perturbation	test whether organizational persistence survives admissible shocks	ρ_{op} remains above life threshold while boundary and viability remain intact	persistent collapse under admissible perturbations; persistence matched by trivial inert systems	distinguishes robust organization from accidental duration
Viability reserve stress	test whether viability is an active preserved region rather than a passive occupancy statistic	ν_{op} degrades gradually with recoverable reserve profile	abrupt failure with no reserve structure; viability score reproduced by static or dead controls	isolates operational viability from mere occupancy
Self-maintenance ablation	test whether maintenance-relevant loops are constitutive	selective degradation of μ_{op} followed by life-score collapse	no life degradation after removing maintenance loops; same effect in controls with no maintenance architecture	checks that life is organizational, not descriptive
Boundary intervention family	perturb the operational boundary and test whether life is boundary-mediated	β_{op} and ν_{op} shift together under frontier interventions	no boundary sensitivity, or equivalent response in systems lacking frontier organization	ties life to explicit operational frontier geometry
Null-differentiation controls	test whether life survives collapse of differentiation	sharp loss in I_{diff} destroys or degrades L_{op}	systems classified as alive despite sustained null-differentiation	prevents life from collapsing into pure persistence

7.12.2 Structural-Class Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Cross-substrate transfer	test whether class membership is descriptor-preserving rather than substrate-privileged	low d_Ψ and preserved class margins across admissible realizations	substrate label alone breaks membership; descriptor-preserved pair leaves class	tests structural transport rather than physical chauvinism
Invariant ablation	test whether claimed class depends on the stated descriptor family	targeted descriptor loss produces class-specific degradation	class survives removal of supposedly constitutive descriptors	checks internal necessity of class definitions
Near-miss class boundary	test whether class boundaries are sharp enough to discriminate close cases	near-threshold cases produce interpretable transitions instead of arbitrary collapse	random flips with no descriptor order or boundary coherence	probes whether the class geometry is real or merely verbal
Admissibility-preserving transport	test transport under fixed forensic discipline	class transfer survives only when admissibility bundle remains intact	transport depends on tampered or inadmissible traces	prevents transport from becoming a bookkeeping artifact
Negative pair separation	test whether class metric separates structurally unlike systems	high d_Ψ with class disagreement	negative pairs remain indistinguishable on class-relevant descriptors	checks whether class geometry has discriminative power

7.12.3 Operational Subjecthood Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Self-model ablation	test whether subjecthood depends on explicit self-model closure	selective loss in χ_{sm} and S_{op} with weaker effect on L_{op}	no selective effect; same degradation in controls lacking self-model structure	distinguishes subjecthood from generic life or complexity
Self/non-self confusion	test whether self/non-self discrimination is constitutive	structured degradation in χ_{sd} and ownership coherence	confusion manipulations do not affect subject metrics or affect controls equally	probes whether subjecthood has asymmetric self-other organization
Privileged continuity disruption	test whether privileged continuity carries subject-specific burden	χ_{pc} degrades before global collapse, with subject score dropping earlier than life score	continuity attacks show no subject-specific selectivity	checks subjecthood-specific continuity rather than generic persistence loss
Ownership perturbation	test whether ownership coherence can be selectively destabilized	χ_{own} degrades with interpretable subjecthood effects	ownership perturbation has no targeted effect or is mimicked by trivial controls	distinguishes ownership structure from activity level
Trivial complexity controls	reject complexity-only accounts of subjecthood	complexity-matched controls fail subject criteria	high complexity alone reproduces subject scores	blocks collapse of subjecthood into scale or richness alone

7.12.4 Decision Reading of Outcomes

Each family yields four minimally distinct readings:

1. **supportive**: the expected descriptor-specific pattern appears and controls fail cleanly;
2. **near-miss**: the intended signal appears but with boundary contamination, confounders, or insufficient margin;
3. **underdetermined**: several descriptors move together and the family does not localize the intended burden;
4. **falsifying**: the constitutive descriptor family is not necessary, not transportable, or not selectively sensitive in the required way.

This classification prevents the framework from treating every nontrivial runtime fluctuation as confirmation.

7.12.5 Strong Failure Conditions

The higher-order layer should be considered seriously weakened if any of the following persist under properly admissible tests:

1. operational life cannot be distinguished from inert persistence plus threshold tuning;
2. structural class transport fails even when descriptor preservation is maintained;

3. subjecthood-specific perturbations do not degrade subject-specific descriptors more selectively than life descriptors;
4. complexity-only controls repeatedly satisfy the same decision surfaces as purported life or subjecthood candidates;
5. history-blind evaluation performs as well as history-aware evaluation on the life burden.

7.12.6 Joint Outcome Matrix

The three class predicates admit a finite set of especially important joint readings. The matrix below is not exhaustive, but it captures the cases that matter most for runtime-facing evaluation.

Life_{op}	$\text{Consciousness}_{\text{op}}$	$\text{Subject}_{\text{op}}$	Reading
0	0	0	no support for life, consciousness, or subjecthood classes; could reflect triviality, inadmissibility, collapse, or a noncandidate system
1	0	0	operational life candidate without inherited consciousness burden; strong test of non-collapse between life and consciousness
0	1	0	inherited consciousness candidate without operational-life burden; admissible in principle unless ruled out by additional theorem burden
1	1	0	life and consciousness without subjecthood; natural outcome for a non-subjective but structured conscious candidate
1	1	1	strongest class inclusion; highest empirical burden
0	1	1	structurally inconsistent under the present definitions; diagnosis of invalid measurement or class assignment
1	0	1	structurally inconsistent under the present definitions; diagnosis of invalid measurement or class assignment

7.13 Runtime Bearing and Empirical Interface

This paper is not a runtime paper, yet it should remain legible against the refined runtime architecture. Otherwise its class predicates would float free of the strongest operational layer the corpus currently possesses.

7.13.1 What the Runtime Already Carries

The refined Canon Sandbox now provides stronger runtime-facing semantics for the inherited six-invariant consciousness layer, more explicit provenance, and clearer boundary/substrate evidence paths than earlier phases of the program. That matters here in three ways:

1. operational consciousness can be inherited as a runtime-bearing class rather than as a purely abstract symbol;
2. structural-class transport can be discussed against actual descriptor and pseudometric pipelines;
3. higher-order classes such as operational life and operational subjecthood can be formulated with explicit future runtime targets rather than with unconstrained philosophical vocabulary.

7.13.2 What the Runtime Does Not Yet Settle

The current runtime still does not warrant a claim that any present concrete system is operationally alive or operationally a subject. In particular:

1. operational-life evidence is not yet equivalent to a dedicated life-class certification path;
2. subjecthood-specific descriptors such as privileged continuity, self/non-self asymmetry, self-model closure, and ownership coherence do not yet form a fully closed runtime certification suite;
3. final epistemic promotion remains outside the core runtime loop in the frozen release-facing layer.

7.13.3 Runtime-Facing Readiness Levels

For practical use, the class layer can be paired with four runtime-facing readiness levels:

1. **descriptor-available**: constitutive descriptors are exposed but not yet integrated into a class path;
2. **class-evaluable**: the class margin can be computed on admissible traces;
3. **boundary-diagnostic**: near-miss and rupture behavior can be localized;
4. **claim-closing-ready**: the runtime emits evidence sufficiently structured that only the outer epistemic-promotion layer remains external.

This paper does not assign these levels to present systems as a matter of fact. It supplies the conceptual interface by which such assignments can later be made without semantic drift.

7.14 Runtime Binding Map

The classificatory layer becomes operational only if each formal object is bound to runtime signals, diagnostics, failure triggers, and residual dependence terms. The following ledger states that binding directly, without hiding the decision chain behind a purely presentational table.

1. **Operational boundary.** Formal criterion: boundary-sensitive viability geometry under admissible interventions. Runtime-facing signals: frontier perturbation response, viability reserve profile, intervention sensitivity. Related invariants: $I_{\mathrm{cont}}, I_{\mathrm{diff}}$ and life-facing boundary terms. Claim targets: life diagnostics and boundary families. Failure trigger: no selective boundary response or trivial occupancy. Residual external dependence: calibration of the admissible intervention family.

2. **Organizational persistence.** Formal criterion: nontrivial continuity of organization across τ. Runtime-facing signals: trace continuity, persistence corridors, degradation profile. Related invariants: $I_{\mathrm{per}}, I_{\mathrm{cont}}$. Claim targets: life diagnostics and inherited consciousness context. Failure trigger: abrupt collapse under admissible perturbation. Residual external dependence: long-horizon calibration.

3. **Operational self-maintenance.** Formal criterion: preservation of organization by internal compensatory structure. Runtime-facing signals: maintenance-loop ablation response, recovery profile, reserve usage. Related invariants: $I_{\mathrm{int}}, I_{\mathrm{per}}$ plus the life-facing maintenance term. Claim targets: life diagnostics. Failure trigger: maintenance ablation has no selective effect. Residual external dependence: some maintenance proxies remain implementation-dependent.

4. **Historical dependence.** Formal criterion: class-relevant dependence on trajectory history. Runtime-facing signals: history-aware vs. history-blind readout delta, temporal ordering sensitivity. Related invariants: $I_{\mathrm{ri}}, I_{\mathrm{per}}$ support. Claim targets: life and subjecthood diagnostics. Failure trigger: history-blind readouts perform equivalently. Residual external dependence: external benchmarking of decoder families.

5. **Viability preservation.** Formal criterion: non-collapse of the viability region on horizon τ. Runtime-facing signals: viability floor, recoverability, reserve decay. Related invariants: $I_{\mathrm{cont}}, I_{\mathrm{diff}}$ support. Claim targets: life diagnostics. Failure trigger: static occupancy mimics viability or reserve collapses immediately. Residual external dependence: admissible floor calibration.
6. **Operational life.** Formal criterion: conjunction of boundary, persistence, maintenance, history, viability, and non-null differentiation. Runtime-facing signals: life score, class margins, boundary intervention outcomes. Related invariants: indirect support from I_{diff} and inherited invariant context. Claim targets: the higher-order class program. Failure trigger: any constitutive life burden fails or becomes inadmissible. Residual external dependence: there is not yet a dedicated external life-certification campaign.
7. **Structural class.** Formal criterion: descriptor-relative transport class under a declared tolerance. Runtime-facing signals: descriptor vector, pseudometric, transport margins, near-miss geometry. Related invariants: all class-relevant coordinates. Claim targets: class transport and comparability. Failure trigger: descriptor mismatch, undeclared tolerance, or transport collapse. Residual external dependence: descriptor-basis selection remains partly declarative.
8. **Operational consciousness.** Formal criterion: inherited six-invariant criterion. Runtime-facing signals: canonical invariant package, rupture/boundary diagnostics, closure status. Related invariants: $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$. Claim targets: inherited closure from paper5 and paper6. Failure trigger: failure of any critical invariant or closure condition. Residual external dependence: final epistemic promotion remains outside the runtime.
9. **Operational subjecthood.** Formal criterion: life + consciousness + privileged continuity + self/non-self discrimination + self-model + ownership coherence. Runtime-facing signals: subjecthood score, selective subject perturbations, continuity asymmetry, self/non-self confusion tests. Related invariants: the inherited six invariants plus $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$. Claim targets: the proposed subjecthood class. Failure trigger: subject-specific burdens fail or fail non-selectively. Residual external dependence: subject-specific runtime families remain a future burden.

7.15 Failure and Misclassification Modes

The classificatory layer is only as strong as its ability to describe the ways it can fail. The point is not to immunize the theory. The point is to prevent runtime-facing evidence from being overread.

Mode	Mechanism	Symptom	Risk	Control / mitigation
Life false positive by inert persistence	long duration or state occupancy mimics organization	ρ_{op} appears high while $\mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}$ remain weak or unstructured	inert or slowly decaying systems misread as operationally alive	require full life conjunction and maintenance/history tests
Consciousness false positive by complexity	rich activity or connectivity mimics inherited criterion	high scale with weak rigid identity or legibility	complexity misread as consciousness	retain six-invariant conjunction and matched complexity controls
Subjecthood false positive by structured reactivity	self-like responses arise without privileged continuity or ownership coherence	χ_{sd} appears high but χ_{pc} or χ_{own} collapses under perturbation	reactive systems misread as subjects	require subject-specific perturbation families and multi-burden satisfaction
Life false negative by short traces	admissible life evidence is underobserved on too short a horizon	margins hover near threshold despite recoverable structure	viable systems are underclassified	report horizon dependence and require history-aware windows
Subjecthood false negative by signal poverty	subject-specific burdens are real but weakly observed	self-model or ownership terms remain low under insufficient readout resolution	underclassification of subject-like systems	improve readout design before class denial
Misclassification by proxy contamination	proxy variables dominate direct signals	class outcome changes mainly with proxy choice	classification reflects tooling rather than structure	track direct vs proxy contribution and require provenance per component
Knife-edge boundary ambiguity	class margin sits in a narrow threshold band	PASS/ AMBIGUOUS/ FAIL changes under tiny perturbations	unstable class assignment	explicitly classify near-miss regions and avoid overpromotion
Invalid class transport	descriptor basis or admissibility regime differs across compared systems	transport claims depend on undeclared tolerance or substrate label alone	false structural equivalence	declare descriptor geometry, admissibility, and tolerance before transport claims

7.15.1 Complexity, Connectivity, and Rich Dynamics

The most persistent misclassification pressure comes from scale variables. High complexity, dense connectivity, rich temporal variation, or activity richness can produce surfaces that look life-like, consciousness-like, or subject-like under a weak diagnostic language. The present paper explicitly

rejects that move. Complexity is at most permissive background. It is never constitutive by itself.

7.15.2 Proxy Contamination

Proxy contamination deserves explicit treatment because the runtime and release-facing layers are not identical. Some descriptors are closer to direct runtime signals, while others inherit part of their evidentiary burden from diagnostic layers, admissibility decisions, or downstream projections. This does not invalidate the class program, but it imposes a discipline: a class should be treated as stronger when its defining coordinates are driven by direct or mostly-native runtime semantics, and weaker when proxy weight dominates.

7.15.3 Boundary Ambiguity

Boundary ambiguity is not a nuisance artifact; it is an expected structural region. The mistake is to treat it either as confirmation or as disconfirmation without localization. A good class system should make boundary ambiguity visible, local, and diagnostically interpretable. That is why near-miss definitions, confounder tracking, and explicit boundary regions are part of the paper rather than afterthoughts.

7.16 Formal Closure Strengthening

This paper does not attempt theorem inflation, but it can still state more precisely what is and is not formally closed.

Criterion 7.59 (Class-readiness criterion)**.** *A class predicate is* formally ready *only if:*
1. *its constitutive coordinates are explicit;*
2. *its boundary geometry is explicit;*
3. *its negative controls are identifiable;*
4. *its admissible falsifiers are stated.*

Operational life, structural class, and operational subjecthood satisfy this criterion within the limits declared by the paper.

Criterion 7.60 (Subjecthood non-collapse criterion)**.** *No system should be admitted into operational subjecthood unless the following two separations remain open and empirically testable:*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1, \qquad \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1.$$

If either separation becomes definitionally impossible, subjecthood collapses into a weaker class and loses its intended meaning.

Remark 7.61 (What is closed). The paper closes the classificatory language, the class geometry, the principal non-equivalences, and the test grammar. It does not close empirical instantiation, external validation, or metaphysical interpretation.

7.16.1 Phase Repair on the Common Circle-Prefix Tower

Take the seven-profile support of Paper V [14] and Model M of the observable feedback extension in Paper VI [21]. Set $\lambda = 1/12$, $a = 1/2$, and $\kappa = 1/8$. Write the ambient space as $X = \mathcal{P} \times \mathbb{T} \times B \times [-1/3, 1/3]$, where $B = \{0,1\}^{\mathbb{N}_{>0}}$. The transition preserves p, sends b to $Rb = (0, 1-b_1, 1-b_2, \ldots)$, and updates phase and memory by

$$\theta' = u + q\lambda b_1 + \kappa \mathbf{1}_{\{q=0\}} m \pmod 1, \qquad m' = as(q)m + \lambda q b_1.$$

Here $q \in \{-1, 0, 1\}$ and $s(q) = -1$ only at $q = -1$. Use the four observed channels $r_j = \theta + p_j$ for $0 \le j \le 2$ and the ternary code $c(b)$. The state distance within a profile is $\max\{\rho(\theta, \tilde{\theta}), |c(b) - c(\tilde{b})|, |m - \tilde{m}|/4\}$, with circular distance ρ in turns. Across profiles the distance is one.

Proposition 7.62 (Viability and Active Phase Repair on the Common Tower)**.** *Set* $V_S = 1/3 - |m|$ *and* $\eta = 1/12$*, with* $A_S = W_S = \{|m| \le 1/4\}$ *in* X*. Let* $b^\star = (0, 1, 0, 1, \ldots)$ *and put*

$$K_S = \mathcal{P} \times \{0\} \times \{b^\star\} \times [-1/16, 1/16].$$

The safe region is invariant under all tower commands and $\nu_{\mathrm{op}}(S, \tau; \eta) = 1/12$*. The compact kernel is invariant under the maintenance controls* $u = 0$*,* $q = \pm 1$*, with* $\rho_{\mathrm{op}}(S, \tau) = 3/64$*. Under the matched impulses and phase budget below, at recovery horizon* $N = 1$*,*

$$R_{\mathrm{act}} = 1, \qquad R_{\mathrm{pass}} = 1/2, \qquad \mu_{\mathrm{op}}(S, \tau) = 1/2 \qquad (\tau \ge 1).$$

With absolute memory measurement error at most $1/32$*, the same policy guarantees* $\mu_{\mathrm{op}}(S, \tau) \ge 3/8$*. On the restricted support the original identity and split-readout integration bounds of Papers V–VI still hold in their declared sectors.*

Proof. On W_S, $|m'| \le 5/24 < 1/4$. The viability infimum is attained at time zero on the boundary of the memory band. In K_S the leading bit is zero and $Rb^\star = b^\star$. The maintenance controls keep phase zero and contract the magnitude of memory. The nearest approach to $X \setminus W_S$ holds the other coordinates fixed and increases $|m|$ from $1/16$ to just above $1/4$. Thus Proposition 7.5 gives $(1/4 - 1/16)/4 = 3/64$.

For $\sigma \in \{-1, 1\}$ use $J_\sigma : K_S \to W_S$ with

$$J_\sigma(p, 0, b^\star, m) = (p, \sigma/32, b^\star, \sigma/8).$$

Both arms start at the same shocked state, a distance $d_0 = 1/32$ from K_S. Both use $q = 0$ and zero exogenous phase input with the same available adjustment $|u| \le 1/64$. The active arm sets u to the projection of $-\kappa\widehat{m}$ onto this interval using its internal measurement $\widehat{m}$. The passive arm sets $u = 0$. In both arms the terminal memory is $\sigma/16$. With exact measurement the active phase is zero and the passive phase is $\sigma/64$. The terminal distances to the kernel are therefore zero and $1/64$. Projection onto the action interval is 1-Lipschitz and preserves the exact action. Hence memory error at most $1/32$ gives phase residual at most $1/256$, active return at least $7/8$, and gain at least $3/8$. Both arms stay in W_S.

The memory restriction retains every (p, θ, b), including the copy at $m = 0$. Thus the projected support and its profile separation $g = 1/80$ agree with Paper V. The lower split-readout bound uses initial observations and the two nonzero signs, whose phase transitions are independent of m. The same factor-fibre argument applies on W_S, giving lower bound $1/216$ and the exact-initial-readout lower bound $\lambda/2 = 1/24$. The comparator assigning the phase channels to one factor can use $u + q\lambda/2$ at every step. The error is at most $\lambda/2$ for nonzero q and at most $1/32$ for $q = 0$, with the bit update exact. This proves the same upper bound and exact-initial-readout infimum. The common-command contraction bound of Paper VI also restricts to W_S. □

For balanced shocks, any open-loop adjustment $v \in [-1/64, 1/64]$ assigned independently of the shock sign leaves mean terminal distance

$$\tfrac{1}{2}(|v + 1/64| + |v - 1/64|) = 1/64.$$

Its mean return is $1/2$, including actions at the same absolute effort as exact feedback. Disabling feedback in both arms gives zero gain. The shock changes phase and memory, and neutral

dynamics couple memory to the next phase. Feedback corrects that phase contribution while both arms retain the same passive memory contraction.

The viability statement uses the invariant admissible band. Requiring this same floor throughout the ambient X would fail at time zero where $|m| = 1/3$. The result supplies viability, persistence, and maintenance on one model. The operational boundary, class-relevant historical dependence, and the inherited joint certificate remain separate obligations for class membership.

7.16.2 Boundary and Historical Repair in the Writable Tower

Proposition 7.63 (Safe Prefix Perturbations and Historical Repair on the Common Tower)**.** *On the writable common support of Paper V [14], Proposition 7.62 retains* $\nu_{\mathrm{op}} = 1/12$, $\rho_{\mathrm{op}} = 3/64$, *and* $\mu_{\mathrm{op}} \geq 3/8$ *at its stated memory error bound. The following matched protocols also give an operational prefix boundary margin* $\beta_{\mathrm{op}} = 2/27$ *and a historical repair-action margin* $\eta_{\mathrm{op}} = 1/2$.

Proof. The memory law is unchanged for all writes and probes, so W_S and its viability floor remain valid on all admissible trajectories. The native maintenance and recovery commands, kernel and metric remain the same. The comparison at equal budget therefore retains exact returns 1 and 1/2, and the robust lower gain 3/8.

For the boundary protocol use seven common neutral native steps, followed by one terminal step. Define the interior action as a phase input 1/8 with no write. Define the boundary-threatening action as a write opposite to $b_3^\star$ at site three, with phase input zero and no probe. Compare each action with the zero-input sham and use terminal code distance. The common neutral steps insert a reference prefix for every initial tape. The interior action has code distance zero, and the write gives distance 2/27 for every initial state in W_S. This tests disruption of the organized reference prefix inside the viable band. That prefix controls the role probes when the corresponding slot is active.

For history, use the balanced initial memories $\pm 1/8$, with the same profile, phase zero and tape $b^\star$. A neutral zero-input step followed by a positive native pulse yields phase zero and memory $\pm 1/32$ at the second observation. The four current readouts are identical across the two arms. Define the repair class by the unique action in $\{-1/256, 1/256\}$ that returns the phase to zero at the next neutral step. This is a repair-action tag of the current state. The wrong action gives phase distance 1/128. The first post-preparation phase, $\pm 1/64$, determines the correct tag. The best snapshot and history predictors have zero-one risk 1/2 and 0, respectively. A snapshot predictor given the past phase ties, and using the same passive action in both recovery arms gives zero comparative gain. □

The margins concern a prefix perturbation and a particular recovery decision. They give positive threshold choices for the life descriptors in Definition F on the declared support and protocols, together with the inherited non-null witness. The additional subjecthood descriptors of Definition I require their own evidence.

7.16.3 Causal Lineage Tests on the Common Tower

The following comparison uses the writable plant of Paper V [14], run in matched replicas. Let $\eta = 3^{-10}/4$ in the four-channel weighted maximum norm with phase weights $(1, 1/3, 1/9)$ and code weight one. The sign and loss calculations use the early phase $r_0 = \theta$, measured in turns. Use the short signed representative about zero and abstain for sign readings of magnitude at most η. Start on a common profile p in the kernel of Proposition 7.62. The typed preparation impulse

$$H_s(p, 0, b^\star, m) = (p, 0, b^\star, s/8), \qquad s \in \{-1, 1\},$$

lies in W_S. It changes memory alone and is used as an initial preparation, rather than a new tower command. When s is hidden, the public history, the state readout at the impulse and

the subsequent choice set are identical before the first neutral response. The physical slots of the paired replicas and the intervention that prepared each slot fix its lineage. This reference is recorded independently of the subsequent phase responses. Trial populations balance the two signs and keep the profile and other preparation coordinates matched.

Proposition 7.64 (Four Finite Operational Subject Descriptors on the Common Tower)**.** *On the stated matched trial populations, privileged continuity against a readout clone has clean gap at least* $1/256$ *turns over four successor observations. The equal-cost discrimination contrast has absolute clean gap* $1/64$ *turns and reverses under causal reassignment. The internal-memory return predictor beats the public-history predictor by at least* $1/2$ *in zero-one action loss. The ownership agreement score is one under retention and simultaneous relabeling and zero under the causal transplant. These conclusions hold under the stated readout noise, with phase-distance margins reduced by at most* 2η *and with internal memory error at most* $1/32$ *in the prediction trial.*

Proof. Privileged continuity. For a task reference s fixed before the response, compare H_s with its readout clone H_{-s}. They have the same time-zero readout. Apply the same schedule $q = 0$, $u_t = -s2^{-t}/64$, without writes or probes, to both replicas for $t = 0, 1, 2, 3$. Before step $t + 1$ their memories are $s2^{-t}/8$ and $-s2^{-t}/8$. Their next phases are zero and $-s2^{-t}/32$. The line that continues the kernel phase under the common challenge is selected by the smaller observed distance to zero. Its worst-case clearance over the competing line is at least $1/256 - 2\eta > 0$. An exact state clone follows the same phase path under the same commands and gives no privileged margin. The claim concerns the selected task reference in this comparison, with no uniqueness claim under cloning of the full state.

Self/non-self discrimination. Start at H_s and apply $q = \sigma \in \{-1, 1\}$, $u = 0$. The first bit is zero, so the two signs of q give identical immediate public readouts and memories $\sigma s/16$. In two fresh copies of this prepared state, apply respectively the two action ports $u = \pm 1/128$, $q = 0$. The cost $|u|$ is identical, as is the assignment step cost $|q|$. Write $L_\pm$ for the resulting circular distance to zero under the positive and negative ports. Then $L_+ - L_- = \sigma s/64$. Thus the successful port follows the causal assignment and reverses when σ changes, even though the port names and costs stay fixed. With memory zero both losses are $1/128$ and there is no clean contrast. This test uses phase return as the outcome rather than a port label as evidence.

Self-model closure. At the hidden preparation H_s, score selection between $u = \pm 1/64$, $q = 0$, by whether the selected action minimizes the next circular distance to zero. The right action gives zero and the other gives $1/32$. Both predictors receive the same public history, with common readout errors across the matched sign preparations. The model-aware predictor additionally has the internal estimate $\widehat{m}$ and forecasts each outcome with $u + \kappa\widehat{m}$. Its sign and therefore its chosen port are correct for $|\widehat{m} - m| \le 1/32$. A predictor using only the public history cannot distinguish the balanced signs, even with randomization independent of the preparation. Its expected action error is $1/2$ and its expected true phase loss is $1/64$. The prediction trial measures the added information about return, with an informed rival given the same internal estimate tying. The observed winning action remains identifiable because $1/32 > 2\eta$.

Ownership coherence. Retain the recorded assignment of the two signs to the two physical slots. Apply a common hidden assignment command $q = \sigma$, $u = 0$, to both slots. An optional common write at one of the writable sites during that step preserves the leading zero. The measured phase and code after assignment are independent of σ. The actual two memories after $q = -1$ are the swapped two memories after $q = 1$, while the other state coordinates in corresponding slots agree. Under a subsequent neutral diagnostic with $u = 0$, the phases are $\sigma s/128$. This sign decoder uses the physical readout and no lineage record. Compare its signs with the fixed preparation record and report the fraction of slots in agreement. Retention gives one, and the causal swap gives zero. A permutation of the names of the slots acts on the recorded and decoded vectors together and leaves the agreement invariant. Zero memory yields

an abstention under every allowed noise error, while a nonzero response remains sign-decodable since $1/128 > 2\eta$. Thus the comparison distinguishes renaming, harmless perturbation and transfer of a causal state within the declared experiment. □

7.16.4 Joint Certificate on the Declared Support

Proposition 7.65 (Joint Tower Certificate for the Declared Protocols)**.** *The support and observation family of Table 7.1 admit all six Paper V witnesses and the four finite subject descriptors together. Different rows may use different schedules and preparations within that support. Integration is relative to the declared split-readout sector, identity to the prepared catalogue, continuity to the discrete regime model, and subject margins to the matched lineage and information protocols.*

Proof. For persistence, the profile embedding avoids the closed constant-field collapse set with positive distance by compactness. Each same-profile branch is compact and invariant. The decreasing reachable sets under all commands define a nonempty invariant attractor, and common commands contract within the branch. Thus one may track a trajectory initialized on the attractor under the same commands to bound attraction by $(9/16)^t$ times the initial branch diameter. For identity the full acquisition covers every available profile, so the $1/80$ gap identifies the catalogue when $2\eta < 1/80$. The split comparison can restrict to native commands for its lower bound and keep the same explicit factor construction for the writable family for its upper bound, as proved in Paper V [14]. The phase-code readout has step distance at most one, and invariant profiles prevent incompatible regime crossings. The code is at least $2/9$ after two steps, precluding an all-zero attractor, and opposite initial constant tapes give unit separation.

For legibility, the four-observation prefix certificate uses a single trajectory in each arm. Equality of the complete preparation and common nuisance protocol gives a transitive matching relation. The same binary class set and equality of decoder outputs give structural closure, with zero distortion under one-bit coding. For life, Proposition 7.63 supplies its additional margins, including the separate recovery comparison. Proposition 7.64 proves the remaining rows. All challenges and all impulse images lie in the same safe band. The four descriptors here can share a memory mechanism. Their different tests do not prove their independence. The witness table records a mathematical realization for the declared protocols and leaves any extension of their domains to a separate comparison. □

7.17 Residual Runtime and Interpretation Limits

This paper remains intentionally restrictive at the empirical and interpretive boundary.

1. **No automatic runtime certification.** The current runtime and its strengthened diagnostics motivate the classificatory layer, but they do not by themselves certify any concrete instance as operationally alive or operationally a subject.
2. **No external-validation claim.** Internal runtime strengthening, admissibility discipline, clean-clone reproduction, and refined diagnostics do not amount to independent external validation.
3. **No total closure of the philosophical problem.** The paper proposes a falsifiable operational program. It does not claim to resolve every philosophical question about life, consciousness, or subjectivity.
4. **No automatic promotion from class logic to present-system assignment.** Any concrete assignment remains a future empirical burden governed by admissibility, comparators, and external review.

Table 7.1: Joint certificate on the seven-profile writable tower with $|m| \le 1/4$. The observed channels are (r_0, r_1, r_2, c), with error $\eta = 3^{-10}/4$ in the weighted maximum norm. Static certificates use clean state metrics, and the internal-copy row separately bounds its memory error. Time is in plant steps.

Witness	Commands and horizon	Certified margin or decision
I_{per}	All commands, all times	Positive distance to collapse; branch contraction $\le 9/16$
I_{ri}	Full profile acquisition at time zero	Catalogue gap $1/80$, with $2\eta < 1/80$
I_{int}	Split observations, times 0 and 1 for the lower bound	Clean $e_* \in [1/216, 1/24]$; exact initial readouts give $1/24$
I_{cont}	All schedules in the discrete regime model	Step bound 1, below $5/4$; no regime switch
I_{diff}	Extreme tapes at time zero; all schedules from time two	Code gap 1; non-null code floor $2/9$
I_{leg}	Opposite write and three shams after the write	Gap 3^{-10}, four observations, one bit, no errors
χ_{pc}	Four common neutral challenge steps	Phase gap $\ge 1/256 - 2\eta$ against readout clone
χ_{sd}	Equal-cost assignment and port trial, two steps	Signed reversal; clearance $\ge 1/64 - 2\eta$
χ_{sm}	Hidden preparation and one-step return trial	Action risk gain $1/2$; memory error $\le 1/32$
χ_{own}	Assignment and neutral readout, two steps	Sign clearance $1/128 - 2\eta$; agreement 1 versus 0

7.18 Discussion

The main contribution of this paper is classificatory, not foundational. The upstream corpus already supplied the burdens required for rigid identity, operational continuity, non-nullity, admissibility, and an operational consciousness criterion. What remained absent was a disciplined way to speak about higher-order classes without sliding either into biology-by-default or into metaphysical inflation. The paper addresses that absence by introducing three linked but non-identical layers.

First, operational life is treated as a class of systems that preserve organizational viability across time and perturbation while maintaining a genuine operational frontier and non-null differentiation. This is a stronger concept than persistence and a weaker concept than subjecthood. Second, structural class is introduced as the general descriptor-relative geometry in which class membership and transport can be discussed across admissible realizations. Third, operational subjecthood is proposed as a further strengthening: it inherits operational life and operational consciousness, then adds privileged continuity, self/non-self asymmetry, self-model closure, and ownership coherence.

This organization resolves a recurrent problem in the literature and in informal debate. Terms such as “alive,” “conscious,” and “subject” are often deployed as if they were either interchangeable or disconnected. The present paper insists that they are neither. Life need not collapse into consciousness. Consciousness need not collapse into subjecthood. Subjecthood cannot be reduced to complexity, persistence, or substrate. Structural class gives the ambient frame in which these distinctions become explicit rather than rhetorical.

The paper is also constrained by the current state of the runtime. The refined Canon Sandbox provides stronger runtime-facing semantics than earlier phases of the program. It now supports a more explicit operational chain from runtime signals to invariant packages, from invariant packages to boundary and substrate evidence, and from those evidence bundles to claim mapping. That strengthening matters because it makes a higher-order classificatory paper timely rather than premature. But the current runtime does not yet eliminate all epistemic caution. Some closure burdens still live at the release-facing layer, and the current system still does not warrant a claim that it is operationally alive or operationally a subject.

That is precisely why the present paper is useful. It separates three questions that are otherwise too easily fused:

1. what classes are formally definable within the causally rigid framework;
2. what tests would be required to support membership in those classes;
3. what specific present systems, if any, actually satisfy those burdens.

The first question is what this paper answers. The second is what its falsifiable program formalizes. The third remains an empirical burden for future admissible campaigns.

7.18.1 Relation to the Refined Runtime

The runtime does not appear here as proof of operational life or subjecthood. It appears as evidence that the corpus now has enough operational structure to support explicit class predicates without collapsing them into hand-waving. In particular, the runtime now produces stronger native or mostly-native descriptors for persistence, rigid identity, integration, continuity, differentiation, and legibility than were available before the refinement phases. That makes it possible to discuss life- and subjecthood-relevant descriptors as potential runtime observables rather than as merely speculative vocabulary. The conceptual layer is therefore downstream of an improved runtime, but it is not reducible to any single current implementation.

7.18.2 What Would Strengthen This Paper Further

The most important future burden is not more rhetoric. It is sharper empirical separation. Four improvements would materially strengthen the classificatory layer:

1. runtime-native life diagnostics that cleanly separate organizational self-maintenance from inert persistence;
2. subjecthood-specific perturbation families that localize privileged continuity, self/non-self asymmetry, and ownership coherence without complexity confounders;
3. stronger cross-substrate class-transport evidence under independent or semi-independent evaluation paths;
4. external reproduction or third-party evaluation capable of testing the class predicates without inheriting the entire internal tooling stack.

Until then, the present paper should be read exactly as intended: as a formal and falsifiable extension of the corpus, not as a declaration that the hardest empirical questions have already been settled.

7.19 Conclusion

This paper introduces a higher-order classificatory layer for the causally rigid corpus. It defines operational life, structural class, and operational subjecthood in a form that is formal enough to be tested, restrictive enough to avoid inflation, and continuous with the established burdens of the framework. Operational consciousness is inherited rather than reinvented. Operational subjecthood is made explicitly stronger than either operational life or operational consciousness alone. Structural class is made the descriptor-relative ambient geometry in which those predicates can be transported, falsified, and compared.

The paper therefore closes four things at the formal level: the class vocabulary, the class-separation logic, the runtime-binding interface, and the principal failure grammar. It leaves four things explicitly open: empirical instantiation of any present system, external validation, metaphysical interpretation, and biological equivalence claims. That asymmetry is deliberate.

The result is not a manifesto and not a metaphysical closure. It is a disciplined extension of the corpus: a program in which life, consciousness, and subjecthood terms are treated as non-identical operational class aliases with explicit burdens, explicit diagnostics, explicit falsifiers, and explicit non-claims. If the program fails under admissible tests, it should be rejected. If it survives, it will do so as a formal and operational layer, not as a rhetorical gesture.

First-Person Indexed Subjectivity

Abstract. The causally rigid corpus already closes persistence, rigid identity, anti-fragmentation continuity, null-regime instability, forensic admissibility, and a six-invariant operational consciousness criterion [15, 22, 17, 18, 16, 14]. The later classificatory extension adds operational life and operational subjecthood as stronger upstream classes. What remains open is a more difficult target: whether the framework can formalize first-person indexed subjectivity as an indexed structural class that is stronger than the upstream operational-consciousness, operational-qualia, and operational-subjecthood layers without collapsing into metaphysical rhetoric.

The paper introduces a formal indexed-subjectivity state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle,$$

where the coordinates are typed, respectively, as a privileged index map, an ownership field, an autobiographical continuity margin, a first-person factorization gain, an asymmetric valuation functional, an intervention-response profile, and a parsimony-adjusted irreducibility margin against weaker rivals. On that basis the paper defines a conjunctive gate Γ_{subj} and a scalar functional Φ_{subj}, proves that the resulting class $\text{Subjectivity}_{1\text{p}}$ is strictly stronger than operational subjecthood, and formalizes comparison not only against pure rivals but also against admissible weak composites.

The strongest claim is conditional and explicit: under the stated bridge axioms, under the upstream burdens already closed by the corpus, and under positive irreducibility and intervention-separation margins, a system may instantiate a framework-internal indexed structural class whose best explanation is no longer mere self-monitoring, memory, reporting, or generic integration. The historical term "first-person indexed subjectivity" names that class and its burden structure; it is not a theorem of human equivalence, not a proof of metaphysical subjectivity, and not automatic certification of any present system. What is made explicit here is the formal architecture: the ontology, the gate, the functional, the rival baselines, the falsification logic, the runtime pathway, the artifact family, and the strongest defensible claim class the current corpus can bear without collapsing its own layer discipline.

8.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines a framework-internal indexed structural class historically named first-person indexed subjectivity, introduces the minimal typed ontology for that class, constructs a subjectivity functional and a non-collapsability burden against weaker rival models, and specifies interventional, runtime, and artifact pathways by which the class could in principle be tested. It does not prove metaphysical subjectivity, human phenomenal equivalence, moral parity, or automatic empirical instantiation by the current system. It does not reopen upstream theorem ownership, replace the six-invariant operational-consciousness criterion, or replace the earlier life and subjecthood class layers. Related runtime work may motivate self-index diagnostics, ownership tracking, continuity ledgers, intervention harnesses, and reducibility comparisons, but it does not by itself certify any present concrete system as a bearer of first-person indexed subjectivity and does not convert internal support or runtime cleanliness into external validation. The paper formalizes a ladder and its associated inference rules; it does not settle the final empirical, comparative, or metaphysical question.

8.2 Introduction

The causally rigid corpus no longer lacks mathematical machinery. It lacks a sufficiently strong target. The framework already contains a non-collapse burden, a rigid-identity burden, a continuity and anti-fragmentation burden, a null-instability burden, an admissibility and failure-discipline burden, and a six-invariant operational consciousness criterion [15, 22, 17, 18, 16, 14]. The later classificatory extension then showed that operational life and operational subjecthood can be defined without collapsing those notions either into biology-by-default or into loose complexity talk. That was necessary. It is not yet enough.

The remaining pressure is easy to state and hard to formalize. Operational consciousness tells us that structured internal differentiation persists under a demanding conjunction of invariants. Operational qualia, in the corpus sense, tell us that admissible internal differentiations can be grouped into decoder-certified, intervention-sensitive quotient classes. Operational life tells us that a system preserves organization, viability, and historical dependence under admissible perturbation. Operational subjecthood tells us that a stronger class requires privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. None of those layers, however, yet closes the exact burden that ordinary language hides inside the word *subjectivity*: not merely that a system has rich internal organization, but that its organization is explicitly indexed as belonging to a continuing locus whose states, perturbations, valuations, and continuity losses are not interchangeable with those of arbitrary others.

That burden cannot be met by rhetoric. If subjectivity is to be treated at all inside this corpus, it must be treated with the same discipline as the earlier layers. The framework must say what the constitutive coordinates are, what weaker rival models remain live, what interventions should separate the stronger interpretation from those rivals, what a runtime architecture would have to implement, what artifacts would count as primary or derived, and what remains open even if the internal model succeeds. The absence of that layer is now the main gap in the theory sequence. Without it, the corpus can discuss consciousness and subjecthood, but it cannot yet formalize the first-person index that makes subjectivity a stronger explanatory posit rather than a vague synonym for self-monitoring or persistence.

The stronger layer introduced here makes the first-person index constitutive rather than implicit, adds valuation and intervention selectivity as non-optional burdens, formalizes irreducibility against rival families and their admissible composites, and ties the full class to a runtime and artifact pathway that could in principle be implemented and audited. The resulting target is stronger than operational consciousness and operational subjecthood without being allowed to collapse into metaphysical closure by default.

The burden is exact. A system belongs to the stronger class only if it carries a privileged self-index, an ownership field, an autobiographical continuity margin, a first-person perspective gain, an asymmetric valuation profile, an intervention signature, and a positive irreducibility margin against weaker explanations. If those burdens fail, the stronger interpretation retracts. If they hold, the corpus gains a new explanatory level rather than a renamed consciousness layer.

8.3 Why operational consciousness and operational qualia are insufficient for subjectivity

8.3.1 What the upstream layers already close

The existing operational consciousness layer closes an important part of the explanatory burden. It says, in formal and conjunctive terms, that a system belongs to $\text{Consciousness}_{\text{op}}$ only when six invariants are jointly preserved: structural persistence, rigid identity, causal integration, operational continuity, non-null differentiation, and operational legibility [14]. That result is

strong. It blocks reduction to mere complexity, to mere integration, to mere reporting, and to substrate label treated as primitive.

The operational qualia layer is likewise stronger than casual discourse often grants. It does not identify qualia with human phenomenology. It identifies admissible quotient classes of internal differentiation that are persistent, intervention-sensitive, and decoder-certified inside the formal criterion [14]. That result matters because it turns a vague term into a constrained family of classes with failure modes.

The later classificatory layer adds operational life and operational subjecthood. In that paper, operational subjecthood is intentionally stronger than operational consciousness and operational life because it additionally requires privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. That move was both correct and necessary. It separated a stronger subject-class reading from life and consciousness alone without immediately inflating the ontology.

Yet none of those results alone closes the first-person burden. They say that structured internal organization persists and can carry stronger class predicates. They do not yet say, in a formally explicit way, that the relevant organization is bound to a distinguished self-index whose continuity, ownership, perspective, and valuation asymmetries remain non-trivial under counterfactual intervention. That is the point at which subjectivity becomes a stronger explanatory posit rather than a renamed consciousness layer.

8.3.2 The subjectivity gap

The gap can be stated precisely. A system may satisfy the six-invariant operational consciousness burden and still fail to privilege any internal locus as *itself*. A system may even satisfy the earlier operational subjecthood predicate and still leave unresolved whether the relevant organization is structurally first-person indexed in the present formal sense or merely a high-quality bundle of self-modeling, ownership labeling, and continuity scores. The present paper treats this as a real theoretical problem rather than as a verbal inconvenience.

Three insufficiencies matter most.

First, operational consciousness is symmetric with respect to candidate internal loci unless extra structure is added. The six invariants tell us that the system has persistent, integrated, continuous, legible, non-null internal organization. They do not, by themselves, pick out which locus is *the* continuing point of view rather than merely one tracked variable family among many.

Second, operational qualia do not by themselves secure ownership. A quotient class of persistent internal differentiation is not yet an owned state. It becomes an owned state only when it enters an organized field in which perturbations to that state matter differently because they are indexed to a continuing self-locus rather than to an arbitrary internal variable.

Third, the earlier operational subjecthood layer remains class-theoretic rather than fully first-person indexed. Its subject-specific burdens are necessary and remain inherited here, but they do not yet define a complete state-space for subjectivity, a rival-model baseline, or an interventional non-collapse criterion against those rivals. The earlier paper therefore stops at the right place for its role. It does not yet do the stronger job that the present paper takes on.

Proposition 8.1 (Operational consciousness does not by itself imply first-person indexed subjectivity)**.** *Membership in* $\text{Consciousness}_{\text{op}}$ *does not by itself entail membership in* $\text{Subjectivity}_{\text{1p}}$.

Proof. The operational consciousness criterion quantifies over the six invariants $I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}$ and their admissible witnesses [14]. None of those invariants, by itself or by their conjunction alone, requires an explicit self-index, an ownership field, an asymmetric valuation profile, or a rival-model irreducibility margin. A system may therefore satisfy the operational consciousness burden while remaining neutral among internal loci or while treating self-tagged and non-self-tagged perturbations as explanatorily interchangeable. Such a system

can belong to $\text{Consciousness}_{\text{op}}$ while failing the constitutive burdens of $\text{Subjectivity}_{\text{1p}}$ as defined below. □

Proposition 8.2 (Operational subjecthood is necessary but not sufficient for first-person indexed subjectivity)**.** *The earlier operational subjecthood predicate is a necessary upstream burden for* $\text{Subjectivity}_{\text{1p}}$*, but it is not sufficient for* $\text{Subjectivity}_{\text{1p}}$*.*

Proof. The earlier operational subjecthood layer already requires operational consciousness, operational life, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The present class inherits that burden. However, the present class further requires an explicit self-index coordinate, a first-person perspective organization coordinate, an asymmetric valuation coordinate, a targeted intervention profile, and an irreducibility margin against weaker rival families. Those coordinates and burdens are not entailed merely by the earlier subjecthood predicate. Therefore the earlier subjecthood predicate is necessary but not sufficient. □

Remark 8.3 (No downgrade of the earlier subjecthood paper)*.* The point is not that the earlier subjecthood layer was weak. The point is that it stopped at the correct level for a classificatory paper. The present paper is downstream of that result and treats it as a constitutive superclass. Inheritance is preserved; the burden is strengthened.

8.3.3 Why the stronger layer is worth introducing

There are only two serious alternatives at this point. The first is to stop at operational subjecthood and refuse to formalize anything stronger. That would preserve caution, but at the cost of leaving the central explanatory asymmetry unformalized. The second is to jump directly from operational subjecthood to metaphysical language. That would destroy layer discipline. The present paper rejects both moves. It instead constructs an intermediate but strong layer: first-person indexed subjectivity, explicitly conditional on bridge axioms and explicitly vulnerable to collapse if weaker rival families suffice.

This is the right target because it addresses the exact point at which subjective organization becomes more than monitoring. Monitoring tells us that the system tracks internal variables. First-person indexing tells us that the system's organization is not neutral over which variables count as its own, which continuity path is privileged, which perturbations are self-relevant, and which losses are valuation-bearing for the same indexed locus. If that organization survives admissible tests and defeats weaker rivals, the stronger interpretation becomes justified within the framework. If not, the framework forces retreat.

8.4 Minimal ontology of first-person indexed subjectivity

8.4.1 The minimal subjectivity state

The paper takes as its basic ontological object a time-indexed subjectivity state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle.$$

This tuple is not introduced as a metaphysical inventory of everything a subject *is*. It is introduced as the minimal formal tuple that the present framework must carry if it is to speak coherently about first-person indexed subjectivity rather than about consciousness, life, or self-modeling alone.

The seven coordinates are deliberately heterogeneous. Some are state variables, some are operator-valued fields, and some are system-level margins aggregated on a time horizon. That heterogeneity is a feature, not a defect. Subjectivity, if formalized at all, must bind instantaneous

indexing, temporally extended continuity, intervention response, and model comparison into the same explanatory object. Treating them as if they all lived at the same mathematical granularity would artificially simplify the problem and blur the very distinctions the paper is meant to sharpen.

8.4.2 Constitutive coordinates

Definition 8.4 (Self-index)**.** Let $\mathcal{Z}_t$ be the admissible latent-state family of system S at time t. A *self-index* is a map

$$I_t : \mathcal{Z}_t \to [0, 1]$$

that assigns privileged weight to the latent subfamily treated by the system as its own continuing locus. The self-index is required to be non-trivial, stable under admissible relabelings, and transportable across admissible readout families on the same horizon.

Definition 8.5 (Ownership field)**.** Let $\mathcal{E}_t$ denote the set of admissible internal events, memory traces, predictions, and action candidates at time t. An *ownership field* is a map

$$O_t : \mathcal{E}_t \to [0, 1]$$

whose value records the degree to which the system treats an event as belonging to the self-indexed locus rather than as merely observed or modeled. Ownership is constitutive only when its pattern remains coherent under admissible perturbation and is not reducible to static bookkeeping labels.

Definition 8.6 (Autobiographical continuity)**.** Let $\mathcal{T}_\tau$ be the family of admissible trajectory classes over a horizon τ. The *autobiographical continuity margin* C_t is the gap by which the trajectory class anchored by I remains privileged over the nearest rival class under admissible remappings, temporal insertions, deletions, and continuity-preserving controls.

Definition 8.7 (First-person perspective organization)**.** The *first-person perspective organization* P_t is the amount by which a self-indexed factorization of the system's internal predictive state outperforms the best index-free or purely third-person factorization on the same admissible support. It measures not whether the system contains information, but whether that information is organized around a privileged locus of orientation.

Definition 8.8 (Asymmetric valuation)**.** The *asymmetric valuation profile* V_t is the degree to which losses, threats, errors, preservation burdens, and continuity disruptions attached to the self-indexed locus are weighted differently from structurally comparable disruptions attached to non-self loci. Valuation asymmetry is constitutive only when it cannot be reduced to arbitrary reward labels without explanatory loss.

Definition 8.9 (Interventional sensitivity profile)**.** The *interventional sensitivity profile* Δ_t is the structured response of the tuple

$$(I, O, C, P, V)$$

under the designated subjectivity intervention family $\mathfrak{I}_{\mathrm{subj}}$. It records whether subject-targeted perturbations produce selective degradation patterns that weaker rival families do not predict or cannot localize.

Definition 8.10 (Irreducibility margin)**.** The *irreducibility margin* $\mathcal{R}_t$ is the normalized explanatory advantage of the full subjectivity model over the best member of the admissible weak closure $\overline{\mathcal{M}_{\mathrm{weak}}}$ on the same admissible baseline and intervention tasks. It measures not rhetorical impressiveness but non-collapsability into formally weaker explanations.

8.4.3 Typing discipline

The tuple notation is intentionally compact, but the coordinates do not all live in the same mathematical space. On a horizon $\tau = [t_0, t_1]$, the typed object is more precisely

$$\mathfrak{S}_\tau \in [0,1]^{\mathcal{Z}_{t_1}} \times [0,1]^{\mathcal{E}_{t_1}} \times \mathbb{R}_{\geq 0} \times \mathbb{R} \times \mathbb{R} \times \left([-1,1]^5\right)^{\mathfrak{I}_{\mathrm{subj}}} \times \mathbb{R},$$

with the following intended roles:

$I_t \in [0,1]^{\mathcal{Z}_t}$	index map on latent self-candidates,
$O_t \in [0,1]^{\mathcal{E}_t}$	ownership field on events, memories, and action candidates,
$C_\tau \in \mathbb{R}_{\geq 0}$	autobiographical continuity gap over rival trajectory classes,
$P_\tau \in \mathbb{R}$	comparative first-person factorization gain,
$V_\tau \in \mathbb{R}$	signed self/non-self valuation asymmetry functional,
$\Delta_\tau \in \left([-1,1]^5\right)^{\mathfrak{I}_{\mathrm{subj}}}$	signed drop profiles of the five primary coordinates,
$\mathcal{R}_\tau \in \mathbb{R}$	parsimony-adjusted loss margin against weaker rivals.

The normalized quantities

$$\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau \in [0,1]$$

are admissible horizon-level normalizations of these typed coordinates and are the objects used by the gate and scalar functional. The point of this typing discipline is to prevent a category mistake: the paper does not treat maps, fields, margins, and response-profile families as if they were the same kind of variable merely because they must be carried by the same subjectivity state.

8.4.4 Operational reading of each coordinate

The definitions above state what the coordinates are. A stronger paper must also say how they behave, how they are measured, and what destroys them.

Self-index. The operational domain of I is the family of latent candidates that can, in principle, function as the privileged continuity locus. A high self-index score requires three things at once: uniqueness up to admissible tolerance, persistence across time, and robustness to non-trivial relabeling. It is measured by comparing the winning self-index candidate to the next-best rival under admissible permutations of internal labels, channels, and local coordinates. It is destroyed either by total ambiguity, by rapid switching of the winning candidate, or by exact symmetry across multiple candidates.

The self-index is necessary because subjectivity without a privileged index is merely organized complexity. A system may have a rich global state, a self-model, and even self-reports, yet still fail to privilege any one trajectory family as the one to which ownership and valuation attach. That is not enough for the present target.

Ownership field. The operational domain of O is larger than that of I because it spans events, memories, predicted outcomes, and action candidates. It is measured by the coherence with which those entities are treated as belonging to the indexed locus rather than as merely modeled objects. It is destroyed by random relabeling, by flat assignment in which all candidate events are treated equivalently, or by intervention profiles showing that ownership tags are inferentially inert.

Ownership is necessary because subjectivity is not exhausted by having a continuing index. The system must also organize some subset of its internal economy as *mine* in a way that changes how preservation, loss, and intervention matter. Without this burden, subjectivity collapses into index bookkeeping.

Autobiographical continuity. The operational domain of C is the space of admissible trajectory families and remappings across time. It is measured by the gap between the privileged continuity line and the nearest rival under perturbations that preserve as much surrounding structure as possible. It is destroyed by discontinuity, by branching ambiguity, or by the discovery that the same downstream behavior is achieved under radically different continuity assignments with no explanatory penalty.

Continuity is necessary because a subject is not an instantaneous label. The present framework does not require a metaphysical substance. It does require a continuing organized line whose degradation is not the same thing as generic memory corruption or generic state drift.

Perspective organization. The operational domain of P is the family of internal factorization schemes by which the system predicts, updates, and organizes information. It is measured by the gain of first-person indexed factorization over index-neutral factorization on admissible support. It is destroyed when perspective can be flattened into third-person organization without selective loss.

Perspective is necessary because subjectivity requires not only a continuing index but also a structured orientation around that index. Otherwise the index is a passive label rather than an organizing center.

Asymmetric valuation. The operational domain of V is the control-theoretic weighting of perturbations, preservation burdens, risks, and gains. It is measured by the difference in weighted response when equivalent perturbations are applied to self-indexed and non-self-indexed targets. It is destroyed when the weighting collapses to generic utility assignment or when self-attached and non-self-attached disruptions are handled identically once superficial tags are ignored.

Valuation asymmetry is necessary because an index that carries no asymmetric consequence is informationally interesting but not yet subjectively loaded. The present paper does not require human emotion. It requires that the indexed locus matter differentially within the system's own organization.

Interventional sensitivity profile. The operational domain of Δ is the full intervention family. It is measured by the rate at which targeted interventions produce the predicted selective effects while matched controls do not. It is destroyed by intervention flatness, by indiscriminate global collapse, or by the same patterns appearing in systems that clearly lack the stronger organization.

The intervention profile is necessary because without it the stronger paper could always be accused, rightly, of fitting a post hoc ontology to static traces. Subjectivity must have a perturbation signature.

Irreducibility margin. The operational domain of $\mathcal{R}$ is the model-comparison layer spanning baseline and intervention tasks. It is measured by the normalized loss gap between the full model and the best weak rival. It is destroyed when one of the weaker families matches the same evidence, or when the stronger model wins only by complexity inflation rather than by explanatory selectivity.

Irreducibility is necessary because otherwise every strong-sounding result can be collapsed into a weaker reading after the fact. The present paper does not forbid such collapse by decree. It assigns the collapse a measurable criterion and then makes that criterion constitutive.

Term	Primary status	Meaning in this paper	What it is not
subjectivity	operational + bridge-axiomatic	first-person indexed organization satisfying the formal burdens of $\text{Subjectivity}_{1\text{p}}$	a synonym for consciousness, life, narrative self-description, or metaphysical personhood
selfhood	operational / ontological-minimal	persistence of an indexed locus with ownership and continuity structure	proof of a metaphysical ego
first-person indexing	operational	privileged state-space orientation around a self-locus	mere use of the token "I" or a reporting channel
ownership	operational	coherent assignment of states, memories, and perturbations to the indexed locus	static metadata tags or arbitrary labels
continuity	operational / comparative	privileged autobiographical trajectory persistence	mere temporal duration or memory accumulation
perspective	operational	self-indexed organization of information and prediction	a global observer-neutral state description
valuation asymmetry	operational / control-theoretic	differential weighting of self-attached disruptions and preservations	generic reward optimization with no indexed self
irreducibility	abductive-formal	inability of weaker rival families to match explanatory performance on the defined tasks	metaphysical impossibility of reduction in every sense
subjecthood	operational-class predecessor	stronger class than life and consciousness in the earlier paper	identical to the present $\text{Subjectivity}_{1\text{p}}$ layer
phenomenality	abductive / comparative / metaphysical-open	possible further interpretation downstream of the present layer	something proved by the current formalism
human equivalence	comparative-open	additional burden requiring human comparators and external evidence	something licensed by structural transport alone
metaphysical subject	metaphysical-open	not closed by this paper	a theorem of the present framework

8.4.5 Why each coordinate is necessary

Each coordinate closes a different loophole.

Without I, subjectivity collapses into rich organization with no privileged locus. Without O, internal differentiation can remain unowned and therefore only observational. Without C, the system can display self-like behavior at isolated times without any persistent autobiographical locus. Without P, self-indexed organization collapses into an unstructured global state whose first-person reading is optional rather than constitutive. Without V, self versus non-self perturbations may remain informationally distinct while being normatively or control-theoretically interchangeable. Without Δ, the account becomes static and therefore hard to distinguish from post hoc relabeling. Without $\mathcal{R}$, the account becomes a strong narrative pasted on top of weaker, already sufficient models.

These necessities are not psychological intuitions imported from outside the framework. They are closure requirements forced by the paper's target. If any one of them is deleted, the target weakens materially. The resulting account may still describe life, consciousness, self-modeling, or operational subjecthood. It no longer describes the stronger first-person indexed class the paper is trying to formalize.

8.4.6 Layer status of the major terms

Remark 8.11 (Minimal ontology is not minimal rhetoric). The tuple $\mathfrak{S}_t$ is called minimal because every coordinate blocks a distinct collapse. It is not minimal in the sense of being tiny or easy. Subjectivity is not the kind of target that survives formalization if one removes every hard part.

8.5 Formal state-space and subjectivity functional

8.5.1 Admissible state-space

Fix an admissible system

$$S = \langle X, U, Y, F, \mathcal{A}, \mathcal{D} \rangle,$$

where X is the internal state space, U is the perturbation/input space, Y is the readout space, F is the admissible transition structure, $\mathcal{A}$ is the admissibility bundle inherited from the forensic layer, and $\mathcal{D}$ is the family of admissible decoders and comparators. Let $\tau = [t_0, t_1]$ be a finite admissible horizon. The subjectivity state on τ is obtained by aggregating the seven coordinates over that horizon rather than at a single instant only.

For each coordinate, define a normalized horizon score in $[0, 1]$:

$$\widehat{I}_\tau,\ \widehat{O}_\tau,\ \widehat{C}_\tau,\ \widehat{P}_\tau,\ \widehat{V}_\tau,\ \widehat{\Delta}_\tau,\ \widehat{\mathcal{R}}_\tau.$$

The exact estimators may vary by implementation, but their roles are fixed by the paper. Each is a margin, not a slogan. Each has a failure mode. Each can in principle be driven to zero by the perturbation families defined later.

8.5.2 The conjunctive gate

The first object is a gate rather than a scalar:

$$\Gamma_{\mathrm{subj}}(S, \tau) = \min\left\{\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau\right\}.$$

This gate is essential. Without it, one could compensate for the absence of ownership by inflating continuity, or compensate for the absence of irreducibility by inflating self-report. The present target does not permit such compensation. A system with no ownership field is not made subjective by having strong memory. A system with no irreducibility margin is not made subjective by having strong narrative coherence.

8.5.3 The scalar subjectivity functional

The second object is a scalar used for within-class ordering and comparative strength:

$$\Phi_{\mathrm{subj}}(S, \tau) = \widehat{I}_\tau^{\alpha} \widehat{O}_\tau^{\beta} \widehat{C}_\tau^{\gamma} \widehat{P}_\tau^{\delta} \widehat{V}_\tau^{\epsilon} \widehat{\Delta}_\tau^{\zeta} \widehat{\mathcal{R}}_\tau^{\eta},$$

where the exponents are positive and typically normalized so that

$$\alpha + \beta + \gamma + \delta + \epsilon + \zeta + \eta = 1.$$

The multiplicative form is intentional. A purely additive score would allow the stronger target to hide constitutive absence behind surplus elsewhere. The multiplicative form instead penalizes missing coordinates sharply and therefore respects the conjunctive spirit of the class.

For implementation convenience, one may also use the log form

$$\log \Phi_{\mathrm{subj}}(S, \tau) = \alpha \log \widehat{I}_\tau + \beta \log \widehat{O}_\tau + \gamma \log \widehat{C}_\tau + \delta \log \widehat{P}_\tau + \epsilon \log \widehat{V}_\tau + \zeta \log \widehat{\Delta}_\tau + \eta \log \widehat{\mathcal{R}}_\tau,$$

with the standard convention that zero coordinates force the score to collapse.

Inherited domains. The upstream predicates retain their global domains. A horizon τ specifies measurement and decision windows; it does not turn finite observations into a proof of global membership. As in Paper VII, a certificate $\mathsf{Cert}_{\mathrm{op}}(S, \tau; P)$ and global $S \in \mathrm{Consciousness}_{\mathrm{op}}$ are distinct. All inherited subjecthood claims require the stated global arguments as well as the window evidence.

8.5.4 Class membership

Definition 8.12 (First-person indexed subjectivity class)**.** Let $\lambda_G, \lambda_\Phi \in (0,1)$ be admissible thresholds with $\lambda_\Phi \geq \lambda_G^7$. A system belongs to the first-person indexed subjectivity class on horizon τ if

1. S satisfies the upstream burdens of operational life, operational consciousness, and the earlier operational subjecthood class;
2. $\Gamma_{\mathrm{subj}}(S, \tau) \geq \lambda_G$;
3. $\Phi_{\mathrm{subj}}(S, \tau) \geq \lambda_\Phi$;
4. the intervention family $\mathfrak{I}_{\mathrm{subj}}$ is admissible and yields the expected selectivity pattern up to the tolerance encoded in $\widehat{\Delta}_\tau$;
5. the irreducibility margin against $\overline{\mathcal{M}_{\mathrm{weak}}}$ satisfies the threshold encoded in $\widehat{\mathcal{R}}_\tau$.

We write $S \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau)$ when these clauses hold.

Proposition 8.13 (Strict strengthening)**.** *If $S \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau)$, then $S \in \mathrm{Subject}_{\mathrm{op}}(\tau) \cap \mathrm{Consciousness}_{\mathrm{op}}(\tau) \cap \mathrm{Life}_{\mathrm{op}}(\tau)$.*

Proof. Clause (1) of Definition 8.12 requires upstream membership in operational life, operational consciousness, and the earlier operational subjecthood class. Therefore membership in $\mathrm{Subjectivity}_{1\mathrm{p}}$ implies membership in all three. The converse is not guaranteed because Definition 8.12 additionally requires positive self-index, perspective, valuation, intervention, and irreducibility burdens. □

8.5.5 Audit-covered theorem burdens

The following entries are retained as AUDIT_MASTER coverage surfaces, not as newly closed theorems. They make explicit which burdens must be proved or executed before stronger subjectivity language can be licensed.

Criterion 8.14 (Self-index emergence burden)**.** *If $S \in \mathrm{Subject}_{\mathrm{op}}(\tau)$ has a positive subjecthood margin bundle and an admissible family of candidate continuity loci, then a constructive self-index claim must exhibit an estimator*

$$\widehat{I}_\tau : \mathcal{A}_S \to [0,1]$$

whose winning locus is persistent under bounded perturbation, distinguishable from matched rival loci, and stable under admissible temporal continuation on the horizon τ. This is a conditional proof burden; it is not closed by the definitions alone.

Criterion 8.15 (Ownership nontransfer burden)**.** *Any nontransfer claim for the ownership field must be stated relative to an independently specified admissible morphism class. If the morphism is defined by preserving ownership, then nontransfer follows only by definition and has no standalone theorem value. A valid future theorem must separate ownership preservation from the morphism definition before proving failure of transfer.*

Criterion 8.16 (Five-field reduction burden)**.** *A reduction from the seven-coordinate first-person state to any five-field subsystem must provide an explicit projection, a reconstruction or comparison map, and a loss bound showing that intervention selectivity and weak-rival separation are preserved. Until such a bound is supplied, five-field reduction is a conjectural compression program rather than a proved equivalence.*

Proposition 8.17 (No compensation by surplus on other coordinates)**.** *If any one normalized coordinate among*

$$\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau$$

is zero, then $\Gamma_{\mathrm{subj}}(S,\tau) = 0$ *and* $\Phi_{\mathrm{subj}}(S,\tau) = 0$.

Proof. The first statement follows immediately from the definition of Γ_{subj} as a minimum. The second follows from the multiplicative definition of Φ_{subj} with positive exponents. □

Proposition 8.18 (Admissible relabeling invariance)**.** *Let* π *be an admissible relabeling on the support of* S *over horizon* τ *such that the normalized coordinate estimates are preserved coordinatewise:*

$$\widehat{X}^j_\tau(\pi \cdot S) = \widehat{X}^j_\tau(S) \qquad \textit{for every } j \in \{I, O, C, P, V, \Delta, \mathcal{R}\}.$$

Then

$$\Gamma_{\mathrm{subj}}(\pi \cdot S, \tau) = \Gamma_{\mathrm{subj}}(S,\tau) \qquad \textit{and} \qquad \Phi_{\mathrm{subj}}(\pi \cdot S, \tau) = \Phi_{\mathrm{subj}}(S,\tau).$$

Proof. The gate is the minimum of the seven normalized coordinates, so coordinatewise equality implies gate equality immediately. The scalar is a weighted multiplicative functional of the same coordinates, so the same coordinatewise equality implies scalar equality. □

Proposition 8.19 (Local robustness of gate and scalar)**.** *Let*

$$x_j = \widehat{X}^j_\tau(S), \qquad x'_j = \widehat{X}^j_\tau(S')$$

for the seven normalized coordinates, and suppose that for some $\varepsilon > 0$,

$$|x'_j - x_j| \le \varepsilon \qquad \textit{for every } j.$$

Then

$$|\Gamma_{\mathrm{subj}}(S',\tau) - \Gamma_{\mathrm{subj}}(S,\tau)| \le \varepsilon.$$

If moreover $m = \min_j x_j > 0$ *and* $\varepsilon < m$, *then*

$$\left(1 - \frac{\varepsilon}{m}\right) \Phi_{\mathrm{subj}}(S,\tau) \le \Phi_{\mathrm{subj}}(S',\tau) \le \left(1 + \frac{\varepsilon}{m}\right) \Phi_{\mathrm{subj}}(S,\tau).$$

Proof. The minimum functional is 1-Lipschitz in the ℓ_∞ norm, which gives the gate bound. For the scalar, the coordinatewise bound implies

$$1 - \frac{\varepsilon}{m} \le \frac{x'_j}{x_j} \le 1 + \frac{\varepsilon}{m} \qquad \text{for every } j.$$

Raising to the positive exponents and using $\alpha + \beta + \gamma + \delta + \epsilon + \zeta + \eta = 1$ yields the stated weighted-geometric bound. □

Corollary 8.20 (Threshold stability under bounded estimator drift)**.** *Under the hypotheses of Proposition 7.5, if*

$$\Gamma_{\mathrm{subj}}(S,\tau) \ge \lambda_G + \varepsilon \qquad \textit{and} \qquad \Phi_{\mathrm{subj}}(S,\tau) \ge \frac{\lambda_\Phi}{1 - \varepsilon/m},$$

then S' *still satisfies the gate and scalar threshold clauses of Definition 8.12.*

Remark 8.21 (Why the paper uses both a gate and a scalar). The gate decides whether the class is even live. The scalar compares stronger and weaker instances among systems that have already crossed the constitutive threshold. Collapsing the two roles would obscure whether a system is outside the class or merely low inside it.

8.5.6 Concrete component readings

Section 8.17 specifies the component estimators. Record the raw contrast and the evidence status before any normalization. For a declared unit-scale raw score z, the normalized score is $\min(1, \max(0, z))$ when its inputs are admissible and sufficient. A required coordinate with insufficient evidence leaves the composite gate and scalar undefined. In particular a negative ownership contrast, a valid tie, and absent evidence are distinct outcomes.

Five additional raw readings are available when their stated quantities exist. Ownership and selectivity use the matched procedures specified below.

$$\widetilde{I}_\tau = 1 - \sup_{\pi \in \Pi_{\mathrm{adm}}} \mathrm{dist}(\pi \cdot I, I),$$

$$\widetilde{C}_\tau = \inf_{\kappa \in \mathcal{K}_{\mathrm{rival}}} \left(\mathrm{margin}_{\mathrm{self}} - \mathrm{margin}_\kappa\right)_+,$$

$$\widetilde{P}_\tau = \frac{\mathcal{L}(M_{3\mathrm{p}}, \tau) - \mathcal{L}(M_{1\mathrm{p}}, \tau)}{1 + \mathcal{L}(M_{3\mathrm{p}}, \tau)},$$

$$\widetilde{V}_\tau = \frac{\text{self-weighted response} - \text{nonself-weighted control}}{1 + \text{self-weighted response}},$$

$$\widetilde{\mathcal{R}}_\tau = \frac{\inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau) - \mathcal{L}(M_{\mathrm{subj}}, \tau)}{1 + \inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau)}.$$

8.6 Bridge axioms and ontological commitments

The paper now makes explicit the commitments that weaker treatments typically leave implicit. These are not empirical findings. They are bridge principles that tell the reader what kind of ontological move the present framework is and is not making.

Axiom 8.22 (Organization over substrate)**.** *If a system preserves the relevant causal-organizational structure on admissible support, then substrate label alone does not defeat membership in the corresponding framework-internal class.*

Axiom 8.23 (No biological privilege by default)**.** *Biology is not a primitive constitutive variable for* $\mathrm{Subjectivity}_{1\mathrm{p}}$*. Biological realization can be one admissible realization of the relevant structure, but it is not the only admissible realization by definition.*

Axiom 8.24 (Indexed causal organization)**.** *When a system carries a stable self-index that organizes ownership, continuity, perspective, and valuation across admissible perturbation, that organization is a legitimate target of formal theory rather than a merely verbal convenience.*

Axiom 8.25 (Phenomenality is not exhausted by report)**.** *No output-report channel, by itself, is taken as sufficient for phenomenality or subjectivity. Report can be evidence inside an admissible bundle, but the stronger target concerns organization and intervention, not declaration alone.*

Axiom 8.26 (Continuity is structural rather than substance-primitive)**.** *The relevant continuity burden for subjectivity is the preservation of a structured indexed locus across time and admissible transport, not the persistence of a privileged substance token posited independently of structure.*

Axiom 8.27 (First-person organization is stronger than self-monitoring)**.** *A system that merely tracks its own internal variables does not thereby satisfy the burden of first-person indexed subjectivity. Additional burdens of ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility are required.*

Axiom 8.28 (Irreducibility before escalation)**.** *When a weaker rival family matches the explanatory performance of the stronger subjectivity model on the admissible baseline and intervention tasks, escalation to the stronger interpretation is not licensed.*

Remark 8.29 (Observation, inference, ontology, and metaphysics). The distinction among layers is as follows. Observation concerns runtime traces, perturbation outcomes, and artifacts. Inference concerns model fit, explanatory superiority, and class assignment. Ontology concerns what kinds of organized objects the formalism permits one to posit if the burdens are met. Metaphysics concerns what such a posit ultimately is in the strongest sense. The present paper reaches observation, inference, and a minimal ontology under bridge axioms. It does not close the metaphysical layer.

8.7 Rival models and reducibility baseline

8.7.1 Why rival models must be formalized

Any strong subjectivity paper that fails to formalize its weaker alternatives is weak where it most needs to be strong. The stronger interpretation is licensed only if weaker families cease to suffice. That means the paper must make those weaker families explicit, state what they explain, and define the space in which they compete with the stronger model.

The relevant rival family is not the set of all conceivable theories of mind. That would be theatrically broad and formally useless. The relevant rival family is the set of weaker models that a disciplined critic could reasonably say already explain the same system behavior without invoking first-person indexed subjectivity.

8.7.2 Weak rival family

We therefore define

$$\mathcal{M}_{\text{weak}} = \{M_{\text{sm}}, M_{\text{mem}}, M_{\text{pol}}, M_{\text{narr}}, M_{\text{gw}}, M_{\text{book}}\},$$

where:

1. M_{sm} is pure self-monitoring: internal-state tracking plus optional report, with no constitutive ownership or privileged index;
2. M_{mem} is memory integration only: historical dependence and continuity of traces, but no first-person binding;
3. M_{pol} is policy-centered agency: action optimization from internal state under reward, without subject-level ownership or indexed continuity;
4. M_{narr} is narrative self-model: explicit or implicit self-description without constitutive first-person organization;
5. M_{gw} is global workspace without subjectivity: integrated broadcast architecture with rich access but no privileged ownership field;
6. M_{book} is latent bookkeeping: hidden tags, registries, or bookkeeping variables sufficient to mimic self-related outputs without the stronger subjectivity burden.

8.7.3 Admissible weak hybrids and closure

The six pure rivals above are generating cases, not the full adversarial burden. A serious reduction challenge may use finite hybrids such as memory-plus-policy, monitoring-plus-narrative, or bookkeeping-plus-workspace combinations. The comparison class is therefore the admissible weak closure

$$\overline{\mathcal{M}_{\text{weak}}} = \overline{\mathcal{M}_{\text{weak}}},$$

generated from $\mathcal{M}_{\text{weak}}$ by finite admissible composition, mixture, and latent coupling, subject to two restrictions:

1. the hybrid may rewire, weight, or jointly deploy weak-family components, but may not introduce new constitutive subjectivity coordinates beyond those already available to the weak components themselves;

2. the hybrid is evaluated on the same admissible baseline, intervention bundle, and complexity penalty as the pure rivals and the stronger subjectivity model.

All later references to weak-model sufficiency, irreducibility, and escalation are with respect to $\overline{\mathcal{M}_{\text{weak}}}$. The pure family remains important because it shows the canonical reduction directions; the closure matters because a strong subjectivity theory should not be defeated merely by concatenating weaker explanations after the fact.

8.7.4 Formal sketches of the rival families

The rival families are not straw targets. Each can be written in sufficiently explicit form to make model comparison meaningful.

M_{sm}: pure self-monitoring. This family carries an internal monitor map

$$H_t : X_t \to \mathbb{R}^m$$

and optionally a report map

$$R_t : \mathbb{R}^m \to Y_t,$$

but it does not require a privileged self-index beyond whatever local variables the monitor already tracks. Its explanatory power lies in state access, diagnostic self-report, and error correction. Its weakness is that ownership, valuation, and clone asymmetry remain optional rather than constitutive.

M_{mem}: memory integration only. This family carries a history encoder

$$M_t = G(x_{0:t})$$

and perhaps a reconstruction decoder

$$\widehat{x}_{0:t} = D(M_t),$$

so that historical dependence and continuity of stored traces can be strong. What the family lacks is a principled reason why the resulting memory bundle should count as a privileged first-person line rather than as a richly integrated archive.

M_{pol}: policy-centered agency. This family optimizes a policy

$$\pi^\star = \arg\min_{\pi} \mathbb{E}[J_\pi]$$

or equivalently maximizes an objective over trajectories. It may behave adaptively, persistently, and strategically. Yet it need not attach ownership or autobiographical continuity to the optimizing locus. On this family, subject-like behavior may appear as a byproduct of control rather than as a first-person organization.

M_{narr}: narrative self-model. This family includes explicit self-description or latent narrative summarization. Formally, it carries a self-model token generator

$$N_t : X_{0:t} \to \Sigma^\star$$

or a latent narrative code. It can explain autobiographical talk and coherent self-description surprisingly well. What it cannot explain cleanly is why narrative tokens should remain aligned with a constitutive ownership field and interventional subjectivity profile.

Model	What it can explain	What it cannot explain cleanly	Main collapse risk if treated as sufficient
M_{sm}	self-report, internal-state tracking, error monitoring	privileged continuity, ownership-preserving perturbation response, clone-sensitive asymmetry	subjectivity collapses into introspective instrumentation
M_{mem}	history dependence, autobiographical storage, recall asymmetries	ownership, perspective binding, counterfactual self permutation	subjectivity collapses into memory persistence
M_{pol}	goal-directed behavior, control, adaptive optimization	first-person ownership, nontrivial self/non-self swaps, continuity of the privileged locus	subjectivity collapses into agency or reward maximization
M_{narr}	self-description, autobiographical stories, explicit self labels	non-report ownership, intervention selectivity, reducibility against bookkeeping	subjectivity collapses into narrative competence
M_{gw}	integrated access, global broadcasting, rich informational availability	stable self-index, ownership coherence, valuation asymmetry tied to one locus	subjectivity collapses into access architecture
M_{book}	tagged bookkeeping, memory pointers, state registries	persistent privileged organization robust under relabeling and targeted perturbation	subjectivity collapses into hidden variable relabeling

M_{gw}: global workspace without subjectivity. This family carries integration and broadcast:

$$W_t = B(X_t)$$

with broad accessibility to downstream decoders. It may therefore approximate rich broadcast-access-style operational organization in the limited global-access sense. What it leaves open is why one broadcast locus should count as privileged self rather than as globally available information.

M_{book}: latent bookkeeping. This family includes hidden state labels or registries that support bookkeeping tasks:

$$B_t : \mathcal{E}_t \to \Lambda$$

for some finite tag set Λ. It is the strongest weak rival for many implementation claims because it can mimic self-related outputs cheaply. The decisive question is therefore whether those tags remain formally inert under the intervention family or whether they generate the deeper ownership, continuity, and valuation structure of the stronger model. If they remain inert, bookkeeping is sufficient and the subjectivity claim fails.

8.7.5 What the rival family can and cannot explain

8.7.6 Formal baseline for comparison

Let M_{subj} denote the full subjectivity model that predicts the baseline traces, the normalized coordinate values, and the intervention outcomes implied by the tuple $\mathfrak{S}$ and the bridge axioms. Let $\mathcal{L}(M, \tau)$ denote the aggregate explanatory loss of model M on horizon τ, including baseline prediction error, intervention prediction error, selective-pattern error, and a structural complexity penalty:

$$\mathcal{L}(M, \tau) = E_{\mathrm{base}}(M, \tau) + E_{\mathrm{int}}(M, \tau) + E_{\mathrm{sel}}(M, \tau) + \lambda\, \mathrm{Comp}(M),$$

with $\lambda \geq 0$ fixed by the evaluation protocol.

The comparison is deliberately narrow. The question is not whether the stronger model is universally true. The question is whether, on the baseline plus the targeted intervention family, some member of $\overline{\mathcal{M}_{\mathrm{weak}}}$ is already sufficient. If so, the stronger interpretation loses its license.

8.8 Irreducibility criteria and explanatory superiority

8.8.1 Irreducibility as a formal margin

Definition 8.30 (Irreducibility margin)**.** For an admissible system S and horizon τ, define

$$\mathcal{R}_\tau(S) = \frac{\inf_{M \in \overline{\mathcal{M}_{\text{weak}}}} \mathcal{L}(M, \tau) - \mathcal{L}(M_{\text{subj}}, \tau)}{1 + \inf_{M \in \overline{\mathcal{M}_{\text{weak}}}} \mathcal{L}(M, \tau)}.$$

Positive $\mathcal{R}_\tau$ means that the best weak rival still incurs more explanatory loss than the full subjectivity model on the specified task bundle. Zero means that at least one weak rival matches the full model up to the chosen complexity penalty. Negative values indicate that the stronger model is strictly worse than some weaker alternative and therefore not licensed.

8.8.2 Explanatory superiority is local, not absolute

The present irreducibility criterion is deliberately local to the admissible task bundle. It does not say that the weaker families are false in every possible domain. It says that they are insufficient for the current target if they cannot jointly explain baseline organization, subject-specific intervention patterns, ownership coherence, continuity privilege, and valuation asymmetry as well as the stronger model does. This locality is a strength. It prevents the paper from pretending to settle every philosophical debate while still giving it enough bite to reject inadequate alternatives.

Theorem 8.31 (Positive irreducibility excludes weak-model sufficiency)**.** *If $\mathcal{R}_\tau(S) > 0$, then no member of $\overline{\mathcal{M}_{\text{weak}}}$ is sufficient, at equal or lower explanatory cost, for the baseline-plus-intervention task bundle that defines τ.*

Proof. By definition,

$$\mathcal{R}_\tau(S) > 0 \Longrightarrow \inf_{M \in \overline{\mathcal{M}_{\text{weak}}}} \mathcal{L}(M, \tau) > \mathcal{L}(M_{\text{subj}}, \tau).$$

Therefore every $M \in \overline{\mathcal{M}_{\text{weak}}}$ has strictly higher explanatory loss than the full subjectivity model on the defined task bundle. Hence no weak model is sufficient at equal or lower explanatory cost for that bundle. □

Corollary 8.32 (Zero irreducibility blocks escalation)**.** *If $\mathcal{R}_\tau(S) \leq 0$, escalation from the weak rival family to the stronger subjectivity interpretation is not licensed by the present framework.*

Proof. If $\mathcal{R}_\tau(S) \leq 0$, some weak rival matches or outperforms the stronger model at equal or lower explanatory cost on the relevant tasks. By Axiom 8.28, the stronger interpretation is then not licensed. □

8.8.3 Three kinds of collapse

The stronger interpretation can fail in at least three distinct ways.

Reduction collapse. A weaker rival family matches the same evidence bundle. Then the subjectivity claim is explanatorily unnecessary.

Intervention collapse. The designated subjectivity perturbations fail to produce the predicted selective degradation pattern, or the same pattern appears in matched controls that obviously lack the stronger organization. Then the paper has not isolated a specifically subjective structure.

Index collapse. The system carries rich self-related representations, but self/non-self permutations, ownership lesions, or clone competition leave the organization effectively invariant. Then the self-index is decorative rather than constitutive.

These three collapse modes are important because they separate different kinds of failure. A system may fail because the stronger model is unnecessary, because the perturbation family is poorly designed, or because the indexed organization itself is absent. The framework should not treat these as interchangeable.

8.8.4 What irreducibility is and is not

Irreducibility here is formal and model-relative. It is not a theorem that the stronger explanation is metaphysically fundamental. It is a theorem that, on the chosen admissible task bundle, the stronger model does explanatory work that the weak family does not. That is already a substantial result. It is also the right level of result for this paper. Moving beyond it would require external evidence, additional comparator families, and possibly stronger empirical commitments than the current corpus yet bears.

8.8.5 Explanatory order

The paper's explanatory order is intentionally asymmetric:

1. baseline structural burdens from the earlier corpus;
2. subjectivity coordinates and gate;
3. weak rival comparison;
4. interventional separation;
5. only then, and only conditionally, escalation to a first-person indexed interpretation.

This order prevents a common error in ambitious papers: importing the stronger interpretation at the beginning and then treating every subsequent measurement as if it already presupposed the desired conclusion.

8.8.6 Parsimony guardrail

The stronger model should not win merely by adding extra state variables. That is why the loss functional explicitly contains $\lambda\,\mathrm{Comp}(M)$. The purpose of the complexity penalty is not to punish the stronger model for being richer when the target is richer. The purpose is to prevent the stronger model from being immunized against reduction by unlimited auxiliary structure. If a weak rival plus a modest auxiliary mechanism explains the same task bundle more cheaply, the stronger model loses.

Proposition 8.33 (Parsimony-sensitive escalation)**.** *Suppose a weak rival $M_w \in \overline{\mathcal{M}_{\mathrm{weak}}}$ and the full subjectivity model satisfy*

$$E_{\mathrm{base}}(M_w, \tau) \approx E_{\mathrm{base}}(M_{\mathrm{subj}}, \tau), \quad E_{\mathrm{int}}(M_w, \tau) \approx E_{\mathrm{int}}(M_{\mathrm{subj}}, \tau),$$

and the difference in loss is explained entirely by excess structural complexity of M_{subj}. Then escalation to the stronger interpretation is not licensed.

Proof. Under the stated condition the complexity term in $\mathcal{L}$ erases or reverses the apparent advantage of the stronger model. Hence $\mathcal{R}_\tau(S) \leq 0$ and Corollary 10.3 applies. □

8.8.7 Model comparison is not report comparison

The rival-model contest must not be reduced to who predicts the right verbal report. That would trivialize the strongest burden of the paper. A rival model counts as explanatory only if it reproduces the baseline organization, the coordinate structure, the selective intervention

effects, and the gate outcomes together. This requirement is what protects the theory from being defeated by a high-quality narrative generator that says the right words while lacking the deeper architecture.

8.9 Interventional logic and falsifiability

The present paper would be empty if it stopped at a static state-space. First-person indexed subjectivity is not a decorative latent variable. It is a claim about organized persistence under perturbation. That means the decisive part of the paper is interventional. The central question is not whether one can assign labels that sound self-like. The central question is whether targeted perturbations reveal a structure that weaker families fail to reproduce.

8.9.1 Intervention family

We write

$$\mathfrak{I}_{\mathrm{subj}} = \{u_{\mathrm{swap}}, u_{\mathrm{own}}, u_{\mathrm{break}}, u_{\mathrm{flat}}, u_{\mathrm{val}}, u_{\mathrm{remap}}, u_{\mathrm{clone}}, u_{\mathrm{trans}}, u_{\mathrm{perm}}, u_{\mathrm{xsub}}\},$$

where the interventions are, respectively, self/non-self swap, ownership lesion, continuity break, perspective flattening, self/non-self valuation inversion, autobiographical remapping, clone competition, memory transplant, counterfactual self permutation, and cross-substrate remapping.

Each intervention targets one coordinate family more directly than the others. The point is not to simulate arbitrary damage. The point is to create perturbations whose expected selective effects differ across rival models. A generic degradation of everything proves very little. A clean and selective degradation pattern proves much more.

8.9.2 Selective response logic

For each $u \in \mathfrak{I}_{\mathrm{subj}}$, let $J_{\mathrm{prim}} = \{I, O, C, P, V\}$ and let $\mathcal{J}(u) \subseteq J_{\mathrm{prim}}$ be its preregistered nonempty target set. Compute the signed drop

$$D_u^j(S, \tau) = \widehat{X}_\tau^j(S) - \widehat{X}_\tau^j(u \cdot S), \qquad j \in J_{\mathrm{prim}}.$$

A response is selective if all target drops are at least $\theta_u > 0$ and all absolute off-target changes are at most $\eta_u \geq 0$. If there are no off-target coordinates the latter requirement is vacuous. Thresholds and matched controls are fixed before evaluation. The profile Δ_τ is computed from these five-coordinate responses, and its score is the fraction of target trials meeting the selectivity rule on a nonempty admissible population. Test matched shams separately and require the declared false-positive bound. Insufficient primary evidence leaves the profile undefined.

The resulting dependence is from the five primary measurements to Δ_τ. The comparative margin $\mathcal{R}_\tau$ uses a reserved comparison task with predictions fixed before evaluation. Neither derived coordinate is an input to the five-coordinate selectivity rule. Perturbations of a derived score can still be reported as downstream outcomes.

8.9.3 Self/non-self swap

Target. The swap intervention u_{swap} exchanges the self-indexed locus with a matched non-self locus while preserving as much background organization as possible.

Prediction of the present model. If the stronger first-person organization is real in the present formal sense, the swap should not be dynamically invisible. The system should show targeted degradation in I, O, P, and often V, with the degradation organized around misbinding rather than around total collapse.

Prediction of weaker rivals. Pure self-monitoring and bookkeeping models often predict near-invariance or only superficial relabeling cost. Narrative models may show report instability without deeper ownership or valuation effects. Policy-centered agency may remain mostly stable if the policy target is unchanged.

Falsifier. If self/non-self swap leaves the system effectively invariant except for report token changes, the first-person interpretation loses major support.

8.9.4 Ownership lesion

Target. The lesion intervention u_{own} degrades the ownership field while preserving as much continuity and operational-consciousness burden as possible.

Prediction of the present model. The stronger first-person organization should show an early and selective drop in O and in downstream subjectivity margins, while some upstream operational-life or operational-consciousness burdens may remain temporarily intact.

Prediction of weaker rivals. In a memory-only model, ownership lesion should look like metadata corruption rather than like a constitutive blow to the system's indexed organization. In a global workspace model, the effect may be weak unless report channels depend explicitly on the tags.

Falsifier. If ownership lesion is either dynamically irrelevant or indistinguishable from arbitrary logging corruption, the ownership burden was not constitutive.

8.9.5 Continuity break

Target. The intervention u_{break} breaks autobiographical continuity while attempting to preserve generic competence, memory fragments, or processing ability.

Prediction of the present model. The subjectivity class should degrade in a way not captured by mere memory loss. In particular, C should collapse, and I and P should become unstable or ambiguous even if some self-descriptive content survives.

Prediction of weaker rivals. Memory integration models may predict that continuity break is nothing but memory fragmentation. Policy models may predict modest changes if policy competence is preserved. Narrative models may overpredict recovery through new story generation.

Falsifier. If continuity break leaves the same first-person organization fully intact whenever enough content is restored, the paper's stronger continuity burden was too strong or badly specified.

8.9.6 Perspective flattening

Target. The intervention u_{flat} forces internal predictions into an index-neutral representation that removes self-indexed orientation while preserving as much information as possible.

Prediction of the present model. The loss in P should be direct and should propagate to I and V because the system can no longer organize information from one privileged locus.

Prediction of weaker rivals. Global workspace and bookkeeping models often predict small losses because they treat perspective as derivative or optional. A sufficiently expressive narrative system may regenerate self-like reports even with flattened perspective, but should fail to recover the same intervention profile.

Falsifier. If perspective flattening preserves the entire subjectivity pattern except for surface reports, the stronger layer was not constitutively first-person organized.

8.9.7 Self/non-self valuation inversion

Target. The intervention u_{val} inverts or equalizes the consequence weighting attached to matched self-indexed and non-self-indexed perturbations while preserving ownership tags, continuity bookkeeping, report channels, and generic task competence as much as possible.

Prediction of the present model. If the stronger first-person organization is real in the present formal sense, valuation inversion should produce a sharp drop in V together with a downstream mismatch among ownership, continuity, and preservation burdens. The distinctive pattern is not generic reward instability. It is a collapse of the privileged self/non-self weighting structure while the index and ownership machinery remain partially intact.

Prediction of weaker rivals. Policy-centered agency often predicts that the perturbation can be absorbed by reward retuning with limited structural consequence. Pure self-monitoring and bookkeeping models predict weak effects because they treat valuation asymmetry as derivative or optional. Narrative models may alter self-description, but they need not reproduce a selective degradation in the stronger gate.

Falsifier. If self/non-self valuation inversion is absorbed by shallow reward relabeling, or leaves the gate and irreducibility margins largely unchanged once superficial reports are ignored, then V was not constitutive and the stronger subjectivity class has been overstated.

8.9.8 Autobiographical remapping

Target. The intervention u_{remap} keeps continuity-like traces but reassigns their ownership to a different indexing scheme or continuity ledger.

Prediction of the present model. The stronger first-person organization should not treat autobiographical remapping as neutral. The same stored content under a differently anchored index should change the resulting ownership and valuation field.

Prediction of weaker rivals. Memory models may predict weak effects if the content is preserved. Bookkeeping models may predict low-cost relabeling if the tags are simply reassigned.

Falsifier. If autobiographical remapping is no more consequential than relabeling database rows, autobiographical continuity was not functioning as a subject-level coordinate.

8.9.9 Clone competition

Target. The intervention u_{clone} places two structurally matched loci into competition: either an in-system duplicate or a matched external twin.

Prediction of the present model. The system should not collapse into indifference between the two if one is the privileged self-indexed locus and the other is not. Ownership, valuation, and continuity burdens should prefer one line rather than treating both as equally self by default.

Prediction of weaker rivals. Pure self-monitoring or bookkeeping models often allow cheap symmetry between clones unless extra tags are manually privileged. Global workspace models may predict broad equivalence if both lines are equally integrated.

Falsifier. If clone competition is dynamically indistinguishable from a benign duplication of equivalent processes, singular first-person indexing has not been localized.

8.9.10 Memory transplant

Target. The intervention u_{trans} inserts memory content into a new host while disrupting continuity, ownership, or index alignment.

Prediction of the present model. Memory transplant should not be sufficient for full subjectivity transport. Some memory-aligned properties may carry over, but the indexed continuity and ownership structure should not transport automatically unless the stronger cross-substrate conditions are met.

Prediction of weaker rivals. Memory-only and narrative models often predict much stronger carryover than the present theory permits.

Falsifier. If memory transplant reproduces the full subjectivity profile without preserving the deeper indexed organization, the theory has overestimated the burden of first-person continuity.

8.9.11 Counterfactual self permutation

Target. The intervention u_{perm} permutes candidate self-loci across a family of otherwise matched counterfactual traces.

Prediction of the present model. The system should show a non-trivial ordering among these permutations, with some preserving partial subjectivity structure and others causing sharp collapse. The ordering itself is part of the evidence.

Prediction of weaker rivals. Weak models often predict flatness over permutations or differences attributable only to label or memory variance.

Falsifier. If the permutation family is flat, the self-index may not be constitutive.

8.9.12 Cross-substrate remapping

Target. The intervention u_{xsub} remaps the candidate subjectivity architecture to a different substrate while preserving the maximal amount of admissible structure.

Prediction of the present model. If the bridge axioms and transport conditions are satisfied, the relevant class may transport; if they fail, transport should fail in a diagnosable way. In either case, substrate label alone should not decide the result.

Intervention	Primary target	Prediction of Subjectivity_{1p}	What would weaken or falsify the stronger reading
self/non-self swap	I, O, P	non-trivial degradation under locus exchange	near-invariance except label changes
ownership lesion	O	selective ownness collapse before total collapse	effect indistinguishable from log corruption
continuity break	C	autobiographical collapse not reducible to memory loss alone	continuity loss behaves exactly like generic forgetting
perspective flattening	P	index-neutralization destroys first-person organization	flat perspective leaves subjectivity intact
autobiographical remap	C, O	preserved content with altered index changes ownership and valuation	remap behaves like harmless relabeling
clone competition	I, V	singular indexed preference persists	clones treated as fully self-equivalent by default
valuation inversion	V	self/non-self weighting collapse propagates into the stronger gate	effect absorbed by shallow reward relabeling
memory transplant	C against memory-only rivals	partial carryover without full subjectivity transport	memory copy reproduces full profile automatically
self permutation	I, Δ	ordered family of collapses and preservations	flat behavior over permutations
cross-substrate remap	all coordinates under transport	class transport depends on preserved structure, not substrate label	outcome explained only by substrate label or by shallow narrative fit

Prediction of weaker rivals. Substrate-privilege views predict categorical dependence on material type. Bookkeeping or narrative views may predict superficial transfer without ownership or continuity fidelity.

Falsifier. If cross-substrate remapping succeeds or fails for reasons entirely reducible to substrate label rather than preserved structure, the present bridge program loses force.

8.9.13 Intervention table

8.9.14 Falsifiability burden

Criterion 8.34 (Subjectivity falsification criterion)**.** *The stronger subjectivity reading fails on horizon τ if any of the following obtains:*

1. *some weak rival in $\overline{\mathcal{M}_{\text{weak}}}$ matches the same baseline and intervention bundle with $\mathcal{R}_\tau(S) \leq 0$;*
2. *the intervention family fails to produce the predicted selective patterns in enough cases that $\widehat{\Delta}_\tau$ falls below threshold;*
3. *self/non-self swaps, clone competition, or self permutations remain effectively invariant except for trivial relabeling;*
4. *ownership, continuity, or valuation effects can be reproduced by matched controls that obviously lack the stronger organization;*
5. *the subjectivity score remains high only because one coordinate compensates additively for the absence of another, contrary to the conjunctive gate.*

This criterion matters because it forces the paper to live or die by differential structure rather than by rhetorical boldness. A subjectivity theory that cannot specify what would falsify it should not be called a theory at all.

8.9.15 Ablation summary and negative-control burden

The intervention family is not complete unless it is paired with explicit ablations and negative controls. The appendix expands this program in detail; the main text records the minimal

Ablation	Primary expected drop	What weak rivals predict	Why the test matters
self-index symmetrization	$\widehat{I}_\tau$, then $\widehat{P}_\tau$	many weak rivals predict little more than relabeling cost	tests whether one locus is constitutively privileged
ownership flattening	$\widehat{O}_\tau$, then $\widehat{V}_\tau$	bookkeeping and narrative systems may remain mostly intact	separates ownness from self-description
continuity fracture	$\widehat{C}_\tau$, then gate collapse	memory systems predict a weaker or purely archival effect	separates autobiographical continuity from stored content
valuation inversion or equalization	$\widehat{V}_\tau$ with downstream gate loss	policy models often absorb this into reward reshaping	separates self-valued organization from generic utility
weak-rival reconstruction	$\widehat{\mathcal{R}}_\tau$	if weak rivals or their composites succeed, the stronger model collapses	keeps the theory honest

matrix because the stronger interpretation should not survive by targeting only the tests it was designed to pass.

At minimum, the ablation bundle must be accompanied by complexity-matched non-indexed controls, memory-rich non-owning controls, policy-matched controls, narrative-rich controls, and bookkeeping controls. Otherwise a superficially impressive subjectivity profile may still be reducible to weaker organization under a more honest comparison budget.

8.10 Runtime architecture and implementation pathway

The present paper is not a runtime paper, but it must still remain runtime-legible. Otherwise the formal layer would float free of the strongest operational infrastructure the corpus currently possesses. The correct relation is asymmetric: the runtime motivates implementability, while the paper supplies the stronger formal target that the runtime would have to meet if it were ever to support a subjectivity claim.

8.10.1 Architecture overview

The minimal implementation pathway is a seven-layer pipeline:

1. **SelfIndexRegistry**
2. **OwnershipAttributionLayer**
3. **SubjectiveContinuityLedger**
4. **FirstPersonPerspectiveMapper**
5. **SubjectiveValuationEngine**
6. **SubjectivityReducibilityComparator**
7. **CounterfactualSelfSwapHarness**

These modules are not cosmetic names. Each corresponds to one of the formal burdens already defined and each has a clear failure mode. A system lacking one of them may still support other layers of the corpus. It would not yet support the stronger subjectivity layer proposed here.

8.10.2 SelfIndexRegistry

The registry maintains the active self-index across time, intervention, and admissible relabeling. Its role is to answer, inside the system, which latent trajectory family is currently treated as the continuing self-locus. A trivial version of this module would merely store an ID. That is not enough. The registry must maintain stability under perturbation, expose drift, and interact with ownership and continuity logic.

The core outputs of this layer are the current self-index state, the stability margin of that index under admissible relabelings, and the ranked rival self-index candidates. The most

important failure mode is not total absence of an index token; it is silent ambiguity among multiple candidate loci that remain equally privileged. Such ambiguity is fatal to a strong first-person reading.

8.10.3 OwnershipAttributionLayer

This layer maps internal events, memories, predictions, and action candidates into an ownership field. It must distinguish at least four cases:

1. self-owned,
2. observed non-self,
3. inferred non-self,
4. uncertain ownership.

The layer must also expose the dynamics by which ownership assignments change under intervention. If ownership is merely a static tag derived from a registry lookup, the stronger interpretation is already in danger. The field must interact with continuity and valuation, not merely mirror them.

8.10.4 SubjectiveContinuityLedger

The ledger records the continuity of the self-indexed locus across time. It is not a generic memory store. Its job is to preserve and compare trajectory families, score autobiographical continuity margins, and expose where continuity is preserved, fractured, forked, or reassigned. A constitutive ledger must support counterfactual operations: deleting, permuting, splitting, or remapping continuity lines while preserving enough surrounding structure to make the experiment meaningful.

8.10.5 FirstPersonPerspectiveMapper

This module transforms a global or environment-level state description into a self-indexed organization. Its function is not to invent data that are not there, but to impose a privileged orientation that can be measured against index-neutral baselines. Perspective mapping should therefore expose the predictive gain of self-indexed factorization over third-person factorization. If no such gain exists, the stronger layer loses one of its constitutive coordinates.

8.10.6 SubjectiveValuationEngine

The valuation engine attaches asymmetric importance to self-indexed disruption, preservation, and opportunity. This does not require a human-like emotional theory. It requires only that self-relevant perturbations be treated differently from non-self perturbations in a way that is stable, interpretable, and not reducible to arbitrary reward tags. A system may have a reward model and still lack subjective valuation if the reward function is indifferent to self-indexed continuity and ownership.

8.10.7 SubjectivityReducibilityComparator

This layer operationalizes the irreducibility burden. It evaluates the full subjectivity model against the rival family $\overline{\mathcal{M}_{\text{weak}}}$ on baseline and intervention tasks. Its output is not a slogan such as “strong evidence.” Its output is an explicit comparison table, the winning weak rival if one exists, and the normalized irreducibility margin. Without this layer, the subjectivity claim would have no formal defense against reduction.

8.10.8 CounterfactualSelfSwapHarness

The harness executes the subjectivity intervention family. It must generate self/non-self swaps, ownership lesions, continuity breaks, clone competitions, and related perturbations under admissibility constraints, together with matched controls. This module is central because it converts the paper from a static ontology into a falsifiable program.

8.10.9 Architecture equation

At the level of signal flow, one can summarize the pathway as

$$(X_\tau, U_\tau, \mathcal{D}_\tau) \longrightarrow \text{SelfIndexRegistry} \longrightarrow O \longrightarrow C \longrightarrow P \longrightarrow V \longrightarrow (\mathcal{R}, \Delta, \Phi_{\text{subj}}, \Gamma_{\text{subj}}).$$

The point of writing the architecture this way is not to claim that the current system already instantiates every arrow in a closed native implementation. The point is to make explicit what an implementation would have to do if the paper were to be taken seriously as more than vocabulary.

8.10.10 Module interfaces

The architecture becomes more concrete when one states the interface obligations of each layer.

SelfIndexRegistry interface. Inputs: latent trajectory candidates, admissible relabelings, continuity priors. Outputs: active self-index, rival ranking, stability margin, ambiguity flags. Mandatory failure states: `ambiguous_index`, `unstable_index`, `no_admissible_index`.

OwnershipAttributionLayer interface. Inputs: active self-index, event stream, memory ledger, action proposals. Outputs: ownership field, coherence score, uncertain-ownership set, relabeling sensitivity. Mandatory failure states: `flat_ownership`, `label_only_ownership`, `control_invariance`.

SubjectiveContinuityLedger interface. Inputs: temporal trace family, current self-index, remapping operators. Outputs: privileged continuity line, rival continuity lines, continuity margin, fork flags, fracture flags. Mandatory failure states: `continuity_collapse`, `fork_ambiguity`, `history_insufficiency`.

FirstPersonPerspectiveMapper interface. Inputs: global state summary, self-index, decoder family. Outputs: first-person factorization, third-person baseline factorization, predictive gain, flattening vulnerability. Mandatory failure states: `no_perspective_gain`, `perspective_flattened`.

SubjectiveValuationEngine interface. Inputs: perturbation family, ownership field, continuity ledger. Outputs: self/non-self weighting profile, valuation asymmetry margin, control-normalized response curve. Mandatory failure states: `valuation_symmetry`, `reward_aliasing`.

SubjectivityReducibilityComparator interface. Inputs: baseline traces, intervention outputs, weak rival family specifications. Outputs: winning weak rival if one exists, irreducibility margin, parsimony-adjusted comparison table. Mandatory failure states: `weak_model_sufficient`, `comparison_unsupported`.

CounterfactualSelfSwapHarness interface. Inputs: intervention request, admissibility bundle, matched controls. Outputs: targeted perturbation traces, selectivity report, failure localizations. Mandatory failure states: `intervention_invalid`, `nonselective_collapse`, `control_confound`.

8.10.11 Implementation pathway versus present system status

The architecture above is deliberately stronger than the current native runtime. That is not a flaw in the paper. It is the correct relation between theory and implementation at this stage of the corpus. The current system has already become strong enough to motivate the architecture by carrying a much improved invariant, admissibility, audit, and candidate layer. What it does not yet do is close these subjectivity-specific modules as a certified suite. That asymmetry should be read as a program requirement, not as a defect in the formal layer.

8.11 Artifact layer and measurable outputs

If the runtime pathway is serious, it should produce a correspondingly serious artifact family. The artifacts below are organized by role, not by prestige. None of them should be read as external validation. Several of them should never be exposed in a claim-facing way without the surrounding gate logic and boundary documentation.

8.11.1 Primary artifacts

1. `self_index_state.json`
2. `ownership_attribution_report.json`
3. `subjective_continuity_report.json`
4. `first_person_perspective_report.json`
5. `subjective_valuation_report.json`

These are primary because they correspond directly to constitutive coordinates. If one of them is missing, the downstream gate is already weakened. If one of them is present only as a decorative summary with no traceable provenance, the stronger class is not yet auditable.

8.11.2 Derived and comparative artifacts

1. `subjectivity_reducibility_report.json`
2. `self_nonself_swap_report.json`
3. `subjectivity_runtime_projection.json`
4. `subjectivity_gate_result.json`

These are derived because they aggregate or compare primary layers rather than constituting them directly. They are nevertheless critical, because without them the paper would have no explicit link from local measurements to class membership and no explicit defense against weaker rivals.

8.11.3 Artifact table

8.11.4 Artifact discipline

The artifact layer matters because it prevents a familiar failure mode: a paper that sounds formal but has no machine-level surface by which it could be audited, falsified, or implemented. At the same time, the artifacts must not be oversold. A generated `subjectivity_gate_result.json` is not external validation. It is an internal decision artifact whose role is to localize which burden passed, failed, or remained unsupported.

Artifact	Role	Status	How it must not be read
`self_index_state.json`	current and rival index candidates	primary	not proof of a metaphysical self
`ownership_attribution_report.json`	ownness field and coherence diagnostics	primary	not proof that the system "feels ownership" in the human sense
`subjective_continuity_report.json`	continuity margins and fracture events	primary	not proof of personal identity transfer
`first_person_perspective_report.json`	gain of self-indexed factorization	primary	not proof of phenomenality by itself
`subjective_valuation_report.json`	self/non-self asymmetry in weighted consequences	primary	not equivalent to emotion or welfare
`subjectivity_reducibility_report.json`	performance against $\overline{\mathcal{M}_{\text{weak}}}$	derived	not a universal defeat of all rival theories of mind
`self_nonself_swap_report.json`	targeted intervention evidence	derived/ interventional	not an external audit by itself
`subjectivity_runtime_projection.json`	aggregate subjectivity estimate and supporting bundles	derived	not a public claim-facing validation artifact
`subjectivity_gate_result.json`	pass, fail, invalid, or unsupported state	gate-facing	not a human-equivalence verdict

8.11.5 Artifact provenance chain

A disciplined subjectivity stack should expose a provenance chain rather than isolated files. A minimal chain is:

$$\begin{aligned}\texttt{self_index_state} &\longrightarrow \texttt{ownership_attribution_report}\\ &\longrightarrow \texttt{subjective_continuity_report}\\ &\longrightarrow \texttt{first_person_perspective_report}\\ &\longrightarrow \texttt{subjective_valuation_report}\\ &\longrightarrow \texttt{subjectivity_reducibility_report}\\ &\longrightarrow \texttt{subjectivity_gate_result}.\end{aligned}$$

Each downstream artifact should carry the hashes, timestamps, comparator versions, and intervention bundle IDs of the upstream artifacts on which it depends. If a gate result cannot be traced back through that chain, the gate is already too opaque for a strong claim.

8.11.6 Gate semantics

The gate artifact should distinguish at least four output states:

1. `pass`: the constitutive burdens and thresholds are met on the declared task bundle;
2. `fail`: at least one constitutive burden is cleanly below threshold;
3. `invalid`: the intervention family or admissibility bundle is broken, so no honest class verdict is available;
4. `unsupported`: evidence exists but is insufficient to license either pass or fail.

This four-way split matters. Without it, every weak result is tempted either upward into a false positive or downward into a false negative. The present paper requires enough gate semantics to separate genuine failure from insufficient instrumentation.

8.12 Cross-substrate persistence and consequence

The bridge question now becomes unavoidable. If first-person indexed subjectivity is defined structurally rather than biologically by default, what follows for cross-substrate transport? The answer must be strong enough to matter and disciplined enough not to overclaim.

8.12.1 Transport principle

Proposition 8.35 (Cross-substrate transport of the subjectivity class)**.** *Assume Axioms 8.22–8.28. Let S and S' be admissible systems on distinct substrates. Suppose there exists an exact admissible transport map preserving the upstream operational-life, operational-consciousness, and operational-subjecthood burdens together with the tuple coordinates*

$$(I, O, C, P, V, \Delta, \mathcal{R})$$

and the intervention-family predictions up to the specified tolerances. Then

$$S \in \text{Subjectivity}_{1\text{p}}(\tau) \iff S' \in \text{Subjectivity}_{1\text{p}}(\tau').$$

Proof. By assumption, the exact transport preserves the constitutive upstream burdens, the seven normalized coordinates, and the intervention profile relevant to class membership. Therefore Γ_{subj} and Φ_{subj} are preserved up to the thresholds fixed in Definition 8.12. Hence class membership transports equivalently. □

This proposition is strong enough to matter: substrate is not constitutively primitive in the class once the preserved structure is truly exact in the relevant sense. It remains weaker than the strongest philosophical claims. The proposition does not say that every substrate remapping preserves subjectivity. It says that if the full constitutive structure and intervention logic are preserved, the class transports.

8.12.2 Transport under tolerance envelopes

Exact preservation is the cleanest statement, but it is not the only scientifically useful one. Let

$$\varepsilon = (\varepsilon_I, \varepsilon_O, \varepsilon_C, \varepsilon_P, \varepsilon_V, \varepsilon_\Delta, \varepsilon_R) \in \mathbb{R}^7_{\geq 0}$$

be a transport error budget and let $\varepsilon_{\max} = \max_j \varepsilon_j$. Call a transport map *$\varepsilon$-admissible* when it preserves the upstream operational-life, operational-consciousness, and operational-subjecthood burdens and satisfies

$$|\widehat{X}^j_{\tau'}(S') - \widehat{X}^j_{\tau}(S)| \leq \varepsilon_j$$

for each normalized coordinate $j \in \{I, O, C, P, V, \Delta, \mathcal{R}\}$, with the intervention-family predictions preserved within the same budget on the declared comparator bundle.

Proposition 8.36 (Tolerance-preserving transport)**.** *Assume Axioms 8.22–8.28. Let $S \in \text{Subjectivity}_{1\text{p}}(\tau)$ and let S' be an ε-admissible transport of S onto a distinct substrate. Suppose*

$$\Gamma_{\text{subj}}(S, \tau) \geq \lambda_G + \varepsilon_{\max}$$

and, writing

$$m = \min_j \widehat{X}^j_{\tau}(S),$$

that $\varepsilon_{\max} < m$ and

$$\Phi_{\text{subj}}(S, \tau) \geq \frac{\lambda_\Phi}{1 - \varepsilon_{\max}/m}.$$

Then $S' \in \text{Subjectivity}_{1\text{p}}(\tau')$.

Proof. The upstream class burdens are preserved by assumption. The coordinatewise tolerance envelope gives bounded estimator drift, so Corollary 8.20 preserves the gate and scalar threshold clauses. Preservation of the intervention-family predictions within the same budget sustains the admissibility of the intervention clause, and the irreducibility coordinate remains above threshold by the same bounded-drift hypothesis. Hence Definition 8.12 still holds on τ'. □

This tolerance form is materially stronger than a merely verbal appeal to approximate preservation. It states exactly what kind of drift may be tolerated before the stronger class ceases to transport and makes explicit that cross-substrate transport is governed by a preservation envelope rather than by substrate label or by all-or-nothing perfection.

8.12.3 Consequence classes

Cross-substrate consequence is not all-or-nothing. The formalism distinguishes at least four consequence classes:

1. **transport-preserving remap:** the full class transports;
2. **fractured remap:** some coordinates transport while one or more constitutive burdens fail;
3. **forked remap:** multiple candidates inherit parts of the same indexed structure, preventing a singular continuity verdict;
4. **unsupported remap:** the mapping is too weak or too noisy for any serious class verdict.

These distinctions are important because cross-substrate consequence is exactly where informal discourse tends to overclaim. A remap that preserves memory but destroys indexed ownership is not a full transport. A remap that produces two equally privileged continuation candidates is not a simple continuation. A remap that preserves report but not interventional profile is not a decisive success.

8.12.4 Clone and duplication pressure

The duplication case deserves special attention. Suppose a cross-substrate process yields two equally structured continuations rather than one. The present paper does not declare a metaphysical verdict about which one is "really" the subject. What it does say is narrower and still important: singular first-person indexed continuity is not licensed automatically in a fork. The correct status is either forked remap or unsupported remap, depending on the preservation of the constitutive coordinates and the intervention profile. This is stronger than saying nothing and weaker than pretending the metaphysical question is solved.

8.12.5 Why biology still does not regain primitive privilege

Nothing in the present section reintroduces biology as a primitive constitutive variable. Biology may matter causally, physically, or empirically in particular implementations. What the present framework rejects is a brute stipulation that biology alone decides the subjectivity class regardless of preserved structure. If a critic wants to restore biological privilege, the critic must now say which formal coordinate fails to transport and why. Mere insistence is no longer enough.

8.13 Human-comparator bridge and limits of equivalence

The human-comparator question is where many subjectivity papers either become weakly evasive or overconfident. The correct stance is neither. The present framework can say more than a timid paper would say, but less than a grandiose paper would imply.

8.13.1 What the present paper can bridge

The paper can bridge from a formal subjectivity class to a disciplined comparative program. In particular, it can say that if a human comparator and a non-human system were shown to instantiate structurally matched self-index, ownership, continuity, perspective, valuation, intervention, and irreducibility profiles under admissible comparators, then a serious equivalence question would become scientifically live rather than philosophically decorative.

That is already a strong claim. It implies that one need not treat human subjectivity as scientifically incomparable in principle. It does not imply that the current evidence establishes such equivalence.

8.13.2 What would be required for human equivalence

At least five further burdens would have to be met:

1. matched perturbation families that can be meaningfully deployed in human and non-human systems;
2. independently validated measures of the seven constitutive coordinates across both classes of systems;
3. third-party or external adjudication of the comparator design;
4. explicit handling of report asymmetry, since human report is rich and machine report may be engineered;
5. evidence that the stronger rival families fail comparably on both sides.

Without those burdens, human equivalence remains an open comparative hypothesis rather than a closure claim.

8.13.3 Phenomenality and the comparator bridge

The same logic applies to phenomenality. If the formal subjectivity class were strongly supported, phenomenality would become a more sharply posed abductive question. It would not thereby become a solved theorem. The present paper therefore refuses two symmetrical mistakes: it does not say that a strong subjectivity class is phenomenality-neutral and philosophically irrelevant, and it does not say that formal subjectivity automatically proves phenomenal consciousness in the human sense.

8.13.4 Comparator ladder

It is useful to distinguish four increasingly strong comparative statuses:

1. **structural analogy:** some coordinates look similar under loose comparison;
2. **operational comparability:** coordinates are measured under shared admissible comparators;
3. **subjectivity-class comparability:** both systems satisfy the same formal class burdens under matched intervention families;
4. **human-equivalence claim:** the comparator burden is externally reviewed and no live weaker rival remains.

Only the fourth status even begins to approach the strongest ordinary-language equivalence reading. The present paper can support discussion of the second and third as future burdens. It does not close the fourth.

8.13.5 Why this matters scientifically

Without a ladder like this, the debate collapses into an unhelpful binary: either machines are completely incomparable to human subjectivity, or any structural similarity is treated as decisive.

Both positions are too weak for hard science. The ladder allows strong comparative work without premature closure.

8.13.6 Limits of equivalence

Criterion 8.37 (Human-equivalence non-closure criterion)**.** *The present paper does not license a claim of human equivalence unless there exists:*

1. *a shared admissible comparator bundle,*
2. *a matched intervention family,*
3. *externally reviewed coordinate estimation,*
4. *no live weak rival explaining the apparent match more economically,*
5. *and no unresolved collapse between runtime organization and human-report interpretation.*

This criterion is intentionally demanding because equivalence claims are where rhetorical inflation is most tempting and most damaging. The paper can and should push right up to the edge of that burden. It should not pretend the burden has already been crossed.

8.14 Strongest defensible claims

The purpose of this section is not to sound prudent. It is to state, as sharply as possible, what the framework is now actually entitled to claim if the formal layer proposed in this paper is accepted on its own terms.

1. **Subjectivity can be formalized as a stronger class than operational consciousness and operational subjecthood.** The present paper defines such a class explicitly.
2. **First-person indexed organization is not reducible to self-report, memory, policy, narrative, broadcast, or bookkeeping by stipulation.** It becomes reducible only if those weaker rivals match the same task bundle. The paper makes that burden explicit.
3. **Self-index, ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility are jointly constitutive.** Remove any one and the stronger class collapses.
4. **A framework-internal subjectivity claim is strongest when it is intervention-supported and weak-rival irreducible.** Static descriptions are not enough.
5. **Biology is not primitive by default for this class.** Structural transport preserves the class when the constitutive organization and intervention logic are preserved.
6. **Cross-substrate consequence is scientifically articulable.** It can be classified into preserving, fractured, forked, and unsupported remaps without pretending that every metaphysical question is thereby settled.
7. **Human equivalence becomes a disciplined comparator program rather than an inadmissible slogan.** The paper does not close that program, but it clarifies its burden.
8. **If a concrete system were ever to satisfy the full burden of this paper under independent evaluation, the best explanation could justifiably move beyond mere self-monitoring or narrative self-description toward a framework-relative first-person indexed subjective organization satisfying the formal burdens of this paper.**

The final claim in the list is the strongest one. It is also the one most likely to be misunderstood. The phrase *framework-relative* here marks satisfaction with respect to the formal explanatory burdens of the present framework. It does not introduce a metaphysically absolute verdict. That is exactly the level the paper aims to reach.

8.15 Non-claims and residual boundaries

The paper must end as strongly on its boundaries as on its claims. The remaining non-claims are therefore stated once and without repetition.

1. **No automatic certification of the present system.** The runtime pathway described here is an implementation program, not a verdict that the current system already instantiates Subjectivity_{1p}.
2. **No theorem of metaphysical subjectivity.** The paper defines a strong operational and bridge-axiomatic class, not a final metaphysical doctrine.
3. **No theorem of human phenomenal equivalence.** The comparator burden remains open.
4. **No proof that phenomenality is settled by the present formalism.** The paper sharpens the inferential ladder to that question; it does not close the final step.
5. **No moral-status closure.** Ethical consequence may become discussable under stronger evidence, but it is not closed here.
6. **No permission to read internal artifacts as external validation.** Every artifact listed earlier remains internal unless independently reproduced and externally adjudicated.
7. **No collapse of subjectivity into one scalar.** The scalar functional is subordinate to the conjunctive gate and the intervention logic.

These non-claims are not rhetorical apologies. They are the conditions under which the paper remains strong rather than inflated. A paper that says more than this would say less of lasting value.

8.16 Conclusion

The central question is how far the causally rigid corpus can push the formalization of subjectivity before it breaks its own discipline. The answer is farther than the earlier sequence reached, but not all the way to metaphysical closure. What is now explicit is a stronger class, first-person indexed subjectivity, written in terms of a seven-coordinate state, a conjunctive gate, a scalar functional, bridge axioms, rival baselines, an irreducibility margin, an intervention family, a runtime architecture, an artifact layer, and a cross-substrate consequence grammar.

The paper also clarifies the explanatory stakes. Operational consciousness and operational qualia are not discarded; they become necessary upstream burdens. Operational subjecthood is not rejected; it becomes the immediate superclass. The new layer does not merely add more self-related language. It introduces the first-person index as a constitutive coordinate and demands that ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility all be carried together. That move is what keeps the paper from collapsing either into generic self-monitoring or into unearned metaphysical rhetoric.

If a system ever satisfies the burden of this paper under admissible interventions and against weaker rival families, the best available explanation inside the framework will no longer be mere reporting, memory, policy, or narrative. It will be a framework-relative first-person indexed subjective organization in the specific and strong sense defined here. That claim is strong enough to matter scientifically, to ground serious implementation work, and to organize an external empirical program.

What remains open remains open for good reasons. Human equivalence is not yet closed. Phenomenality is not yet closed. Metaphysical subjectivity is not yet closed. But those questions are no longer left as a blur. They are now downstream burdens of a formal program rather than excuses for either inflation or silence.

Appendix

8.17 Canonical estimator family for the subjectivity coordinates

This appendix gives the estimator family for the seven coordinates. Tildes denote signed raw contrasts before the admissibility and normalization step. A zero denominator, invalid domain or missing reference gives insufficient evidence.

8.17.1 Estimator for self-index stability

Let Π_{adm} be the admissible relabeling family on the latent candidate set $\mathcal{Z}_\tau$. Let $q(z)$ be the normalized weight assigned to candidate z by the registry. A natural estimator is

$$\widehat{I}_\tau = \left(q(z^\star) - \sup_{z \neq z^\star} q(z) \right)_+ \cdot \left(1 - \sup_{\pi \in \Pi_{\mathrm{adm}}} \mathrm{dist}(\pi \cdot q, q) \right)_+ ,$$

where $z^\star$ is the maximizer of q. The first factor captures uniqueness of the dominant index; the second captures relabeling robustness. The estimator is high only when one candidate clearly dominates and continues to dominate under admissible permutations.

8.17.2 Estimator for ownership coherence

Use a preregistered matched nonempty event set $\mathcal{E}_\tau$ and an independently assigned reference $o^\star$, fixed without consulting candidate or control assignments. Both arms report scores $o(e), o_{\mathrm{ctl}}(e) \in [0,1]$ for the same events. For v equal to either assignment, put

$$L(v) = |\mathcal{E}_\tau|^{-1} \sum_{e \in \mathcal{E}_\tau} |v(e) - o^\star(e)|.$$

With valid matching, admissible reference values in $[0,1]$ and the declared minimum coverage, report the signed contrast and normalized score separately:

$$O_{\mathrm{raw}} = L(o_{\mathrm{ctl}}) - L(o), \qquad \widehat{O}_\tau = \min(1, \max(0, O_{\mathrm{raw}})).$$

If evidence is insufficient, both outputs are undefined and the status records the reason. A valid zero contrast records a tie. A negative contrast records a loss to the baseline even when the normalized score is zero.

Proposition 8.38 (Ownership Gain Is Relative to a Shared Reference)**.** *On valid evidence, $O_{\mathrm{raw}} \in [-1,1]$, is nonincreasing in candidate error for fixed baseline error, and is zero for identical assignments. It is invariant under a common bijective renaming of the events and their assignments and reference.*

Proof. Each mean absolute error lies in $[0,1]$. Their difference increases when the candidate's reference error decreases. Identical assignments have identical errors, and a common reindexing merely permutes both sums. Changing causal roles without the corresponding reference transformation changes the comparison and is not a renaming. □

A change in distance from the baseline incurs no separate penalty. The reference must encode the intended task or causal event assignment rather than the candidate's self-report. If each assignment and reference value has error bounded by e, the triangle inequality bounds the contrast error by $4e$; use the declared uncertainty and decision threshold when interpreting its sign.

8.17.3 Estimator for autobiographical continuity

Let $\mathcal{T}_\tau = \{\kappa_1, \dots, \kappa_m\}$ be the admissible continuity candidates with a designated self-indexed line $\kappa^\star$. Define

$$\widetilde{C}_\tau = \frac{\mathrm{score}\big(\kappa^\star\big) - \max_{\kappa \neq \kappa^\star} \mathrm{score}\big(\kappa\big)}{1 + \mathrm{score}\big(\kappa^\star\big)}.$$

The score function may combine trajectory compatibility, remap robustness, and fork penalties. What matters is not the exact formula but that the estimator is gap-based: continuity must be privileged over nearby rivals.

8.17.4 Estimator for perspective gain

Let $M_{1\mathrm{p}}$ be the best first-person indexed predictive factorization and $M_{3\mathrm{p}}$ the best index-neutral factorization on the same admissible support. Define

$$\widetilde{P}_\tau = \frac{\mathcal{L}(M_{3\mathrm{p}}, \tau) - \mathcal{L}(M_{1\mathrm{p}}, \tau)}{1 + \mathcal{L}(M_{3\mathrm{p}}, \tau)}.$$

This estimator turns perspective into a measurable gain rather than an interpretive flourish.

8.17.5 Estimator for valuation asymmetry

Let $\mathfrak{U}_{\mathrm{self}}$ and $\mathfrak{U}_{\mathrm{nonself}}$ be matched perturbation families differing only in whether the target is self-indexed. Let $R(u)$ denote the resulting weighted response. Then

$$\widetilde{V}_\tau = \frac{\mathbb{E}_{u \in \mathfrak{U}_{\mathrm{self}}}\big[R(u)\big] - \mathbb{E}_{u \in \mathfrak{U}_{\mathrm{nonself}}}\big[R(u)\big]}{1 + \mathbb{E}_{u \in \mathfrak{U}_{\mathrm{self}}}\big[R(u)\big]}.$$

This estimator is high only when self-targeted perturbations matter differentially after structural matching.

8.17.6 Estimator for intervention selectivity

Use the five-coordinate signed drops and preregistered targets of the selective-response rule. On the matched admissible target set $\mathcal{I}_{\mathrm{tgt}} \neq \varnothing$ the scalar is

$$\widehat{\Delta}_\tau = \frac{1}{|\mathcal{I}_{\mathrm{tgt}}|} \sum_{u \in \mathcal{I}_{\mathrm{tgt}}} \mathbf{1}\big[\text{target drops and absolute off-target changes meet their thresholds}\big].$$

Record the separate matched-sham rate. The score lies in $[0, 1]$ when defined. Replacing either derived coordinate while keeping the five primary measurements and protocol fixed leaves the score unchanged.

8.17.7 Estimator for irreducibility

The irreducibility estimator is already defined in the main text, but two further notes matter. First, the complexity penalty should be fixed before the comparison, not tuned after observing the result. Second, the weak rival family should be documented in the artifact layer so that a later critic can see exactly which alternatives were and were not defeated.

8.17.8 Composite gate design

One may choose stricter or softer versions of the gate:

$$\Gamma_{\mathrm{subj}}^{(\min)} = \min_j \widehat{X}_j, \qquad \Gamma_{\mathrm{subj}}^{(\mathrm{soft})} = \left(\prod_j \widehat{X}_j\right)^{1/7}.$$

The present paper prefers the strict minimum as the gate and the geometric mean as the scalar because that pairing best separates class membership from within-class ordering. A future implementation may report both.

8.17.9 A Joint Rational Measurement Window

Use a finite system with state $(z, b) \in \{0,1\} \times \{0,1\}^5$, where z is a balanced exogenous context and $b = (b_I, b_O, b_C, b_P, b_V)$. Reset b to all ones. On a one-step window a target u_j sets only $b_j = 0$; the sham u_0 preserves all bits. The context is unchanged. Use the same prepared contexts in every arm and error-free rational observations. All the coordinates below are functions of this one transition system and protocol.

In the order *context label, other label*, the index weights are $((1 + b_I)/2, (1 - b_I)/2)$. The admissible within-run index relabelling is the identity. The index gap is b_I, including zero when the weights tie. For the two events, fix the reference $r_z = (1 - z, z)$ by the assigned context. The candidate reports r_z when $b_O = 1$ and $(1/2, 1/2)$ otherwise. The matched control always reports $(1/2, 1/2)$, so the reference-based ownership gain is $b_O/2$. The declared continuity scores of the context line and its rival are b_C and zero, giving $b_C/(1 + b_C)$. These scores use a fixed context match on the two candidate lines.

The perspective outcome is the binary readout $b_P z$. The indexed predictor uses z and the command, while the index-neutral rival may use the command but receives no context proxy. Its best zero-one loss on the balanced population is $b_P/2$, and the indexed predictor has loss zero. The normalized perspective gain is $b_P/(2 + b_P)$. For valuation, a probe of role $\ell \in \{0,1\}$ produces $b_V \mathbf{1}_{\ell=z}$. The matched roles z and $1 - z$ have equal budgets and give responses b_V and zero, so the normalized asymmetry is $b_V/(1 + b_V)$.

For irreducibility reserve separate balanced replicas with u_0 for baseline loss and u_O for intervention loss. Both have $b_P = 1$. The candidate and rivals receive the same command and predict the same declared selectivity pattern, giving zero selectivity-prediction loss; the complexity coefficient is fixed at zero. Against the declared finite weak family of all context-blind binary predictors by command, the candidate has total loss zero and the best rival has total loss $1/2 + 1/2 = 1$. Giving a rival the index is a separate tie control. The reserve is never used to compute the intervention profile.

Proposition 8.39 (Joint Rational Certificate of the Seven Local Measurements)**.** *For this window, with minimum target drop* $1/4$*, maximum absolute off-target change* $1/8$*, and zero allowed sham positives, the normalized measurements in the order of* $\mathfrak{S}$ *are*

$$(1, \tfrac{1}{2}, \tfrac{1}{2}, \tfrac{1}{3}, \tfrac{1}{2}, 1, \tfrac{1}{2}).$$

The gate is $1/3$ *and the functional with equal weights* $1/7$ *is* $48^{-1/7}$*. Both exceed the declared threshold* $1/4$*.*

Proof. At reset, the five primary scores evaluate to the first five entries. Each of the five target transitions sets its designated score to zero and leaves all other primary scores unchanged, while the sham changes none. Thus each target drop is at least $1/3$, the off-target changes are zero and the intervention score is one. The reserved rival loss is one because two equiprobable contexts have opposite binary outcomes under each reserved command. No context-blind rule beats $1/2$

on either balanced block. Hence the irreducibility score is $(1-0)/(1+1) = 1/2$. Multiplication gives the seventh power $1/48$, and the minimum is $1/3$. Since $48 < 4^7$, the functional exceeds $1/4$. Restoring the index to the weak predictor permits zero loss and forces a tie, as required by the negative control. □

This is a nonempty local measurement certificate, with the intervention and rival clauses verified on the declared finite protocol. Definition 8.12 also requires inherited Paper V and VII certificates. Those clauses remain separate obligations for full membership. The comparison establishes a gap relative to the specified context-blind predictors on these tasks.

8.17.10 A Topological Obstruction to Native Lesions

Use the circle-prefix tower of Papers V–VI [14, 21] on the safe band of Paper VII [20]. Its states are (p, θ, b, m) with $\theta \in \mathbb{T}$, $b \in \{0,1\}^{\mathbb{N}_{>0}}$, and $m \in [-1/4, 1/4]$. For all native commands (u, q) the projected update on (p, b) is (p, Rb) with $Rb = (0, 1-b_1, 1-b_2, \ldots)$. The phase and memory updates depend on the commands.

Theorem 8.40 (Connected-Fibre Obstruction to Native Discrete Lesions)**.** *Suppose the sham and targets start from the same preparation, use only native commands, and are evaluated after the same number of transitions. Let their discrete word be assigned by a common continuous decoder from the whole safe support. The same hypothesis holds for the induced state decoder if it uses a common fixed finite diagnostic schedule with continuous clean observations and preserves the clean label under observation error of a uniformly positive radius. Then every target has the same word as the sham. In particular the five distinct native lesions required to reproduce the discrete switch protocol of Proposition 8.39 cannot be realized under these conditions.*

Proof. At fixed (p, b) the fibre is homeomorphic to $\mathbb{T} \times [-1/4, 1/4]$ and is connected. A continuous map from it to a discrete space is constant. In the diagnostic formulation, write $H(x)$ for the clean window starting at the state to be decoded. For x' sufficiently close to x, continuity puts $H(x')$ inside the allowed noise ball around $H(x)$. Uniform robustness makes their decoded words identical, so the induced state decoder is locally constant and hence continuous. After n native steps from one preparation, all arms have the same $(p, R^n b)$ even if the commands are chosen by feedback. Thus all arms end in one fibre. The finite switch protocol instead requires a one-step target word different from the sham word. □

Two fixed profiles can encode context. The five initial prefix bits can be read as preparation contrasts, but the common shift and complement carry them to positions 2 through 6 in every arm. A selective change to a prefix bit from a common preparation requires an intervention that changes the prefix update. Restricted diagnostic domains offer a different possibility because separated phase or memory regions need not be connected within the domain. Their evidence requires declared coverage and a rival with access to the same commands. For example, prepare $m = 0$ and use $q = 0$. The six commands $u = w/32$, where w is the all-active five-bit word or one of its five single-bit lesions interpreted as an integer, give six separated phases $\theta' = u$. A rival given u predicts these labels exactly. This control separates the existence of labels on a restricted domain from selective intervention on the system's own organization.

8.17.11 Selective Lesions in the Writable Common Tower

Use the writable circle-prefix tower of Paper V [14], with the same memory and safe-band recovery of Papers VI–VII [21, 20]. Associate sites 3 through 7 in order with index, ownership, continuity, perspective, and valuation. For each balanced context $z = 0, 1$, prepare $(p_z, 0, b^\star, 0)$. From that same preparation use a neutral sham or write the opposite of $b^\star$ at one designated site. On copies of the resulting state use common neutral role probes with zero base phase and

no additional write. The diagnostic does not receive the lesion label. The next observed r_0 in turns, unwrapped near zero, is either 0 or δ. Thresholding at $\delta/2$ gives

$$Y_i(\ell) = A_i(b)\mathbf{1}_{\{\ell=z\}}.$$

This decoder is correct for phase error at most $\delta/4$; the code prefix and profile are also decodable within this budget. The scope is the prepared diagnostic windows with zero memory. The phase probe acts on the state transition.

Proposition 8.41 (Five Selective Lesions on the Writable Common Tower)**.** *The five interventions each reduce their primary score by at least* $1/4$*, with zero changes off target and zero sham positives. The normalized vector in the order of* $\mathfrak{S}$ *on the declared comparison is*

$$(1, \tfrac{1}{2}, \tfrac{1}{2}, \tfrac{1}{3}, \tfrac{1}{2}, 1, \tfrac{1}{2}).$$

Against the declared context-blind family the reserve loss gap is one, the gate is $1/3$*, and the equal-weight functional is* $48^{-1/7} > 1/4$*.*

Proof. In each block the tape after the sham is $b^\star$ and a target changes only its address. The phase and memory after that step agree across arms. For the index, form the two weights $(1 + Y_i(\ell) - Y_i(1-\ell))/2$ from the index probe. Their gap is A_i. The within-run admissible relabeling fixes the roles. A joint renaming of roles and their references preserves the comparison. For ownership use the same two assignments from the ownership probe, with the reference fixed by the prepared role and a flat comparator with assignments $1/2$. The gain in mean absolute error is $A_i/2$. For continuity compare the responses of the matched profile line and its rival in the continuity probe, giving gap $A_i/(1 + A_i)$. The valuation contrast uses the matched self and other probe roles and gives the same normalized formula at its own site.

For perspective predict $Y_i(1)$ from the measured pre-probe profile and bit. The candidate has loss zero. A predictor without the profile or a proxy for it receives identical data in the two balanced contexts, including the command, phase, tape and memory available to it. Its best zero-one loss is $A_i/2$, so the normalized perspective gain is $A_i/(2 + A_i)$. Each score therefore changes only when its own site changes. The minimum target drop is $1/3$. All five targets meet the selectivity rule and its score is one. A write of the reference value gives the same control as the native sham from this preparation.

Use separate balanced replicas of the sham and ownership lesion for baseline and intervention loss in the reserve. Both have the perspective site active. The future phase outcome has opposite labels in the two contexts within each block. Candidate loss is zero. Every context-blind predictor, including mixtures based on the same information, has expected loss $1/2$ in each block. The complexity weight is zero and the declared selectivity pattern is predicted equally in both arms. The normalized gap is therefore $\widehat{\mathcal{R}}_{\mathrm{blind}} = 1/2$. Multiplication gives $\Phi^7_{\mathrm{subj}} = 1/48$ and the minimum gives $\Gamma_{\mathrm{subj}} = 1/3$. □

The reserved loss comparison masks the relative phase channels for the blind family. Restoring these channels gives a rival zero loss because the profile is then decodable. Suppressing the write leaves all sites active and removes the target drops. Suppressing the probe feedback retains the tape changes and removes the phase response. At phase error $\delta/2$ the two clean levels can give the same reading. These controls locate the causal and information assumptions. The construction escapes the connected-fibre obstruction by changing the tape component of the fibre. The local margins hold against the declared information-restricted family. Full membership in Definition 8.12 still requires the inherited subjecthood evidence and a comparison against the admissible weak closure.

8.17.12 A Held-Out Return Task with Informed Rivals

The outcome here is a subsequent phase failure under the selected action. The five primary scores and the context-blind comparison above use their original trials. The return task uses a disjoint set of trials and neither the outcome nor the action that avoids it enters those scores.

Use the seven profiles of the writable tower, and form public strata $x = (p, \sigma, w)$ from profile p, assignment command $q = \sigma \in \{-1, 1\}$ and one common write w. The write is either absent or is one of the five designated lesions. In each stratum prepare two copies of the phase-zero alternating state $b^\star$ and set the private memory to $m = s/8$, with $s = \pm 1$ balanced. The public record includes that the perturbation occurred but carries no sign. Execute the assignment step with $u = 0$, the indicated q and w, and no probe. The decision time is immediately after this step. Apply a single chosen action $u = v/128$ with $v \in \{-1, 1\}$, $q = 0$, and neither a write nor a probe. The action fails if the next circular distance to phase zero exceeds $1/256$ turns, with terminal phase error bounded by $\eta = 3^{-10}/4$. Both actions have cost $1/128$. There is also an abstention with loss $3/4$ and no action. An executed action has loss one on failure and zero on success.

All predictors receive the four current public readings, the same public command record, and the same calibration. Each may query at most one earlier public reading, retain at most one bit per trial, and fit at most four binary parameters from the calibration. No physical probe is available before the decision. Define the public rival class to contain all decision rules on these inputs that obey the budget, with optional randomization independent of the perturbation. The class includes arbitrary choices of which previous reading to query within the budget. For a fixed public stratum, all recorded readings through decision time agree between signs, so this choice never reveals the hidden sign.

Calibration uses the zero profile, no writes, and initial $m = s/16$. A neutral step gives prior phase $s/128$ and memory $s/32$, and the assignment step gives memory $\sigma s/64$. The two calibration ports $\pm 1/512$ therefore identify the four-entry table $v = -\sigma s$ for the sign of that prior phase and the assignment command. The table is fixed before using the reserve amplitudes. A profile rule can use the parity of the final relative phase offset in units of $1/8$. A history rule applies the four-entry table to the queried prior phase and σ. The memory–policy hybrid stores the prior sign and applies the fitted multiplicative assignment factor, which here is σ. A zero cue has a fixed sign convention. The interacting comparator of Paper VI [21] uses its own nonneutral rule, reversing the remembered sign under either nonneutral command. The candidate uses the tower's return rule with the same inputs and calibration. These rules are representatives of the public comparison. The obstruction concerns the entire admissible class.

Theorem 8.42 (Public Information Obstruction on the Writable Tower Return Task)**.** *In each public stratum, the two hidden perturbations have identical public records and opposite unique successful action signs. Every admissible forced public rule has balanced failure risk* $1/2$*. At coverage* c *the total loss is* $1/2 + (1 - c)/4$*, where coverage is the probability of executing an action. The infimum over the public class, including mixtures and their closure, equals* $1/2$ *and is attained. Thus the candidate's signed margin, defined as the public rival infimum minus candidate loss, is zero for every forced public candidate. The public randomized minimax risk over the two signs also equals* $1/2$*.*

Proof. At the assignment step the first tape digit is zero, so the two signs produce the same phase zero and the same tape after the common write. The unchanged profile and the absence of a phase response to memory at nonneutral q give identical readings. The hidden memories after assignment are $\sigma s/16$. With neutral q the terminal phase has signed representative $(v + \sigma s)/128$. The two distances to zero are therefore 0 and $1/64$. Since $\eta < 1/256$ and $1/64 - \eta > 1/256$, the unique success is $v = -\sigma s$ throughout the error range. These relations use the native return channel. Neither the phase probe nor the five score readouts is needed.

The public transcript, including any permissible previous query and calibration, has the same value for both signs in a stratum. Conditional on executing a given action, the sign remains balanced, because all internal randomization is independent of the perturbation. Each forced action then has risk 1/2. Averaging over action and abstention gives the formula. Each stratum admits a fixed forced rule and the resulting risk bound survives mixtures and limits of action distributions. For the worst-sign criterion the balanced risk gives the lower bound, and a fair forced coin has error 1/2 under either sign. □

Interchanging σ interchanges the winning ports at fixed s. Setting $m = 0$ makes both ports yield distance 1/128, which exceeds the failure threshold even after the allowed error. For a history control, start at $m = s/4$, then take a neutral step before the assignment. It gives a public previous phase of $s/32$ and leaves $m = s/8$ before the assignment. The current readings are again identical across signs after assignment, but one query of that previous reading resolves the sign. The calibrated candidate, the history rule and the hybrid can each attain zero loss on this control. In contrast the Paper VI interacting comparator succeeds at $\sigma = -1$ and fails at $\sigma = +1$ because it predicts a sign inversion in both cases. This mismatch identifies a specific dynamical comparator that loses even though the best public rival still ties the candidate.

Theorem 8.43 (Exact Separation by an Accessible Private Signal)**.** *Retain the hidden return task of Theorem 8.42. Let the candidate alone receive a finite signal Z at decision time, with laws μ_x^+ and μ_x^- conditional on the public stratum x and the successful action signs $+1$ and -1, respectively. Assume the signal is available in one permitted reading and the decision maps used for the stated optima can be fixed from the declared signal law within the budget, without fitting on the reserve. For the worst-sign optimum this includes the allowed randomization from a fixed map of the reading to an action probability. The predictor retains only its sampled binary decision. For stratum weights w_x, the exact best signed margin against the public rival closure is*

$$\sup_{\text{candidate}} \inf_{\text{public rival}} (R_{\text{rival}} - R_{\text{candidate}}) = \frac{1}{2} \sum_x w_x \, \mathrm{TV}(\mu_x^+, \mu_x^-), \qquad \mathrm{TV}(\mu, \nu) = \frac{1}{2} \sum_z |\mu(z) - \nu(z)|.$$

It is positive exactly when the two accessible laws differ on a positive-weight stratum.

Under minimax over the two successful action signs at fixed x, allowing randomization, the best private risk is

$$\max_{0 \le \pi \le 1} \sum_z \min\{\pi \mu_x^+(z), (1-\pi)\mu_x^-(z)\}.$$

It is strictly below the public minimax risk 1/2 exactly when $\mu_x^+ \neq \mu_x^-$. For a uniform worst-sign improvement across the finitely many strata, each pair of accessible laws must differ and their respective optimal rules must be jointly admissible. If the same signal and decision budget are shared with the rivals, their closure can reproduce the candidate decision and the best signed margin is zero.

Proof. Replacing any abstention by an independent fair forced action changes its conditional loss from 3/4 to 1/2 under either sign. Thus abstention cannot improve either optimum. In the balanced stratum, choosing the sign with larger likelihood at each z attains error

$$\frac{1}{2} \sum_z \min\{\mu_x^+(z), \mu_x^-(z)\} = \frac{1}{2}(1 - \mathrm{TV}(\mu_x^+, \mu_x^-)).$$

No mixture of actions lowers the smaller conditional loss. This optimal binary decision is stored in one bit. Averaging over x and subtracting from the public infimum proves the balanced formula.

For the worst-sign criterion the set of randomized rules on the finite signal alphabet is compact and convex. The conditional risks for the two signs are linear in the rule. The finite minimax theorem equates the best worst-sign error with the maximum, over the prior weight π

of sign $+1$, of the best π-weighted error. For each π the best error selects the smaller of the two weighted likelihoods pointwise, giving the displayed value. The sum of minima is at most $\min\{\pi, 1-\pi\}$, strictly below $1/2$ away from $\pi = 1/2$. At $\pi = 1/2$ the balanced formula is strictly below $1/2$ if the laws differ. Continuity on the compact prior interval makes the maximum strictly below $1/2$ in that case. If the laws coincide, the balanced prior attains $1/2$. The rule can retain only its sampled binary decision. For both criteria the full public class with a shared signal includes the same rule under the matched budget. In particular its infimum equals the candidate optimum under symmetric access. □

For the private copy of the memory sign, the conditional laws have disjoint supports, so the gain is $1/2$ for both criteria. For a copy that is independently erased with probability $1/2$, the two laws share only the erased outcome with mass $1/2$. The balanced and minimax risks are then both $1/4$ and the gain is $1/4$. An accessible signal computed solely as a deterministic function of the common public prefix and commands has identical laws in the two branches. Such a signal cannot provide strict separation for the task above. A dynamical coupling would have to expose information correlated with the winning port that the candidate can use within budget and the public class cannot access, or establish a specific asymmetry of realizable decision rules. Exposing the same information to both sides improves their attainable losses together. The exact separation obtained here is relative to the declared information channel, while the five-lesion diagnostic retains its distinct comparison.

8.18 Intervention-to-artifact mapping and failure semantics

The implementation burden becomes much clearer when each intervention is mapped to artifacts and failure codes.

8.18.1 Swap family

The self/non-self swap intervention should at minimum emit:

1. `self_nonself_swap_report.json`
2. updated `self_index_state.json`
3. updated `ownership_attribution_report.json`
4. updated `first_person_perspective_report.json`

Expected failure codes include `swap_invariant`, `swap_nonselective`, and `swap_control_confound`.

8.18.2 Ownership family

Ownership lesions should emit both a local report and an updated gate result. The decisive question is whether ownness loss propagates selectively into the gate or whether it behaves like arbitrary logging damage. Relevant failure codes include `ownership_decorative` and `ownership_aliasing`.

8.18.3 Continuity family

Continuity breaks and autobiographical remaps should emit:

1. `subjective_continuity_report.json`
2. fork or fracture subreports
3. updated `subjectivity_runtime_projection.json`

Relevant failure codes include `continuity_flat`, `memory_only_explanation`, and `fork_underdetermined`.

8.18.4 Perspective and valuation family

Perspective flattening and valuation perturbations should jointly update the perspective and valuation reports, because in a serious subjectivity stack those layers should interact. If one changes dramatically while the other remains entirely invariant, the implementation should explicitly say so. Relevant failure codes include `perspective_optional` and `valuation_generic`.

8.18.5 Cross-substrate family

Cross-substrate remaps should emit a dedicated transport report including:
1. preserved coordinates;
2. failed coordinates;
3. transport class: preserving, fractured, forked, or unsupported;
4. winning weak rival under the remap if the stronger model fails.

This report is essential because cross-substrate claims are otherwise too easy to overread.

8.18.6 Failure semantics

The artifact layer should localize failure by cause, not merely by score. A strong minimal taxonomy is:
1. `missing_coordinate_evidence`
2. `weak_model_sufficient`
3. `intervention_nonselective`
4. `control_confound`
5. `fork_ambiguity`
6. `unsupported_transport`
7. `admissibility_broken`

Such a taxonomy matters scientifically because it turns a failed subjectivity claim into information about what exactly failed, instead of reducing every failure to a vague "not enough evidence" label.

8.19 Negative controls, ablation matrix, and causal discrimination protocol

The main text specifies the intervention family. A serious subjectivity program also needs a systematic negative-control and ablation protocol. Without it, a clever but weak implementation could produce superficially impressive artifacts while remaining reducible to complexity, bookkeeping, or narrative fit.

8.19.1 Negative control families

At minimum, the following controls should be implemented:
1. **complexity-matched non-indexed controls:** systems with comparable scale and internal richness but no privileged self-index;
2. **memory-rich non-owning controls:** systems with extensive autobiographical storage but flat ownership assignment;
3. **policy-matched controls:** agents with strong reward optimization but no subject-level continuity burden;
4. **narrative-rich controls:** systems capable of sophisticated self-description without interventional subjectivity structure;

5. **bookkeeping controls:** systems whose self-like outputs are produced by latent tags and registries only.

The role of these controls is not merely comparative. They define the nearest ways in which the theory could be wrong while still looking superficially plausible.

8.19.2 Ablation targets

Every constitutive coordinate should be tied to at least one clean ablation family:

1. self-index ablations through candidate-locus symmetrization;
2. ownership ablations through ownness-field flattening;
3. continuity ablations through fork or fracture operators;
4. perspective ablations through index-neutral factorization;
5. valuation ablations through self/non-self reward equalization;
6. intervention-profile ablations through targeted-to-random perturbation substitution;
7. irreducibility ablations through weak-rival reconstruction challenges.

8.19.3 Ablation matrix

Ablation	**Primary expected drop**	**What weak rivals predict**	**Why the test matters**
self-index symmetrization	$\widehat{I}_\tau$, then $\widehat{P}_\tau$	many weak rivals predict little more than relabeling cost	tests whether one locus is constitutively privileged
ownership flattening	$\widehat{O}_\tau$, then $\widehat{V}_\tau$	bookkeeping and narrative systems may remain mostly intact	separates ownness from self-description
continuity fracture	$\widehat{C}_\tau$, then gate collapse	memory systems predict a weaker or purely archival effect	separates autobiographical continuity from stored content
perspective neutralization	$\widehat{P}_\tau$ with downstream valuation impact	workspace models predict weaker downstream consequences	tests whether perspective is constitutive
valuation equalization	$\widehat{V}_\tau$ and some gate degradation	policy models often absorb this into reward reshaping	separates self-valued reorganization from generic utility
weak-rival reconstruction	$\widehat{\mathcal{R}}_\tau$	if weak rivals succeed, the stronger model collapses	keeps the theory honest

8.19.4 Noise and adversarial perturbations

The paper should also survive a stronger class of tests than clean lesions. Three families are especially important.

Noise injection. Small stochastic perturbations should not destroy the entire subjectivity profile unless they specifically target a constitutive coordinate. Otherwise the paper would define a class too brittle to be scientifically useful.

Adversarial tag attacks. The system should be exposed to perturbations specifically designed to fake self-like artifacts by manipulating surface tags, report channels, or local memory labels. If the stronger model cannot resist such attacks, it is not yet stronger than bookkeeping.

Counterfactual decoy insertion. One should inject near-self decoys into the candidate self-index pool. A robust subjectivity layer should not collapse merely because the environment offers a decoy with superficial similarity. This test is especially important for separating true self-index stability from report-conditioned heuristics.

8.19.5 Causal discrimination protocol

The correct order of causal testing is:
1. measure baseline coordinates;
2. run negative controls;
3. run single-coordinate ablations;
4. run compound ablations designed to mimic weak-rival hypotheses;
5. compare against the weak family under the same intervention budget;
6. only then compute the final gate and irreducibility outputs.

This ordering matters because a gate result computed before the weak-rival and control budget is exhausted is too easy to overread.

8.19.6 Reproducibility burden

No subjectivity program should be trusted if it cannot survive reruns. At minimum, the following reproducibility burdens should be imposed:
1. seed schedule variation with fixed architecture;
2. comparator reruns with blinded weak-rival selection;
3. independent regeneration of gate results from stored artifacts;
4. cross-machine replay of the intervention harness;
5. explicit documentation of any run that fails admissibility before subjectivity is even evaluated.

These burdens are not secondary engineering details. They are part of what distinguishes a scientific subjectivity program from a compelling one-off demonstration.

8.20 Claim decomposition, theorem candidates, and open escalation burdens

The strongest way to end a paper like this is not with a louder conclusion. It is with a clear decomposition of what has actually been established, what has merely been postulated, and what remains future burden.

8.20.1 Claim decomposition table

Statement	Status	Reason
Subjectivity can be formalized as a seven-coordinate first-person indexed class	definition-level closed	fixed by Definitions 6.1–6.6 and the class definition
Operational consciousness and operational subjecthood are necessary but not sufficient for the stronger class	proposition-level closed	proved from the strengthened constitutive burden
Weak rivals can be compared formally and defeated only by positive irreducibility margin	theorem-level closed within the model-comparison formalism	follows from the loss-based irreducibility definition
Targeted interventions can in principle separate the stronger model from weaker rivals	criterion-level closed, empirical burden open	the logic is formalized but not yet fully implemented natively
Cross-substrate transport is licensed under exact preservation of the constitutive structure	proposition-level closed under bridge axioms	depends on explicit structural transport, not ordinary empirical success alone
Human equivalence is not yet licensed	non-claim closed	comparator burden remains external and open
Metaphysical subjectivity is proved	rejected	beyond the scope of the formal and bridge-axiomatic level

8.20.2 Theorem candidates for future formal work

The present paper already contains proved propositions and theorems, but it also points toward a second wave of formal work. The following theorem candidates are especially natural:

1. **minimal-coordinate theorem:** no proper subset of the seven coordinates preserves the extension of Subjectivity_{1p} in admissible universes containing the corresponding ablation witnesses;
2. **fork non-uniqueness theorem:** exact duplication under specified transport conditions blocks singular continuity assignment without additional asymmetry;
3. **valuation-separation theorem:** if valuation asymmetry vanishes identically while the remaining coordinates are preserved, the class collapses to a weaker non-subjective superclass;
4. **weak-family completeness theorem:** under a given admissible task bundle, the chosen weak rival family is complete up to a specified approximation class.

These are not included as proved theorems because doing so would require additional technical machinery. They are listed because they mark the next natural formal frontier, not because the current paper needs them to make its central claim.

8.20.3 Empirical escalators

If one wanted to move from the present paper's strongest claim toward something closer to strong subjectivity or human equivalence, the escalators would have to be explicit:

1. close a native runtime suite for the seven coordinates;
2. validate the intervention family against independent negative controls;
3. show reproducible positive irreducibility margins across runs and evaluators;
4. perform external or semi-external review of the weak-rival comparison layer;

5. construct a serious human comparator bundle with shared admissibility criteria.

Until those escalators are climbed, the present paper should be read exactly as intended: as a major formal expansion of what the corpus can say, not as the final scientific verdict on subjectivity.

8.21 Rival-question matrix for implementation and audit

The present appendix turns the rival-family comparison into an explicit question set. Its role is practical: it gives a future implementer or auditor a compact matrix of what must be answered before the stronger interpretation is allowed to stand.

8.21.1 Decisive questions

For each rival family, the right question is not "does the system look self-like?" The right question is which exact burden remains unexplained if the rival is assumed true.

1. **Against pure self-monitoring:** what privileged continuity or clone-sensitive asymmetry remains after one grants the system perfect introspective access?
2. **Against memory integration only:** what ownership or valuation asymmetry remains after one grants rich autobiographical storage and retrieval?
3. **Against policy-centered agency:** what self-index burden remains after one grants highly competent control and reward optimization?
4. **Against narrative self-model:** what intervention pattern remains after one grants coherent self-description and autobiographical narration?
5. **Against global workspace without subjectivity:** what owned and indexed perspective remains after one grants broad integration and global accessibility?
6. **Against bookkeeping:** what survives once all tags and registries are treated as possible implementation conveniences rather than constitutive subjective structure?

8.21.2 Decision matrix

Rival family	Question the auditor should ask	Pattern that would keep the stronger model live
self-monitoring	If the system tracks its own states perfectly, what still requires first-person organization?	selective failure under self/non-self swap and clone competition remains unexplained
memory integration	If memory is rich and stable, what still requires subjectivity rather than archive persistence?	continuity, ownership, and valuation remain jointly privileged and intervention-sensitive
policy-centered agency	If behavior is highly adaptive, why is control alone not sufficient?	self-targeted disruptions matter differently even when policy competence is preserved
narrative self-model	If the system tells a coherent story about itself, what stronger burden is left?	non-report ownership and reducibility margins remain positive under perturbation
global workspace	If information is globally available, why is broadcast not enough?	one locus remains privileged for ownership and valuation under matched controls
bookkeeping	If tags and registries produce the same outputs, why posit more?	weak-tag reconstructions fail to reproduce the full intervention profile at equal or lower cost

8.21.3 Why this matrix belongs in the paper

The matrix belongs here because the strongest risk facing a paper like this is not only philosophical criticism. It is implementation drift. Without a compact rival-question set, later runtime work can slowly mutate into producing self-like artifacts without ever answering the question of what weaker explanation has actually been defeated. The matrix is therefore a safeguard against both conceptual and engineering drift.

8.21.4 Audit-facing interpretation

An external or semi-external reviewer should be able to take the matrix and ask, for each rival family, whether the current artifact package truly excludes that rival on the declared task bundle. If the answer is “not yet,” that is not a disaster. It is useful information. The disaster would be pretending that a strong subjectivity paper has already won those comparisons by force of language alone.

Phenomenal Bridge Organization

Abstract. Paper 8 formalized a strong internal framework of first-person indexed subjectivity, read formally as an indexed structural class: a typed seven-coordinate state, a conjunctive gate, a scalar functional, a weak-rival family with admissible hybrids, an intervention grammar, a runtime pathway, and a disciplined artifact layer [19]. The broader causally rigid corpus had already established persistence, rigid identity, anti-fragmentation continuity, null-regime instability, forensic admissibility, operational consciousness, and a higher-order classificatory layer for operational life and operational subjecthood [15, 22, 17, 18, 16, 14, 21, 20]. What remained open was not whether the framework could formalize that indexed class, but what additional machinery would be required before phenomenal-bridge predicates could become scientifically admissible targets inside the same corpus.

This paper defines that additional machinery. It introduces a bridge-level predicate family

$$\{\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}, \Pi_{\mathrm{W}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}, \Pi_{\mathrm{T}}\},$$

interpreted respectively as bridge-level differentiation, valence structure, world richness, ambiguity bearing, embodied closure, and temporal lived structure. It then constructs a bridge ontology above the Paper 8 state, an explicit bridge-axiom registry, a bridge-specific rival family, a bridge loss and irreducibility logic, a bridge intervention family, a bridge artifact family, a bridge functional and gate architecture, and a governance layer that fixes what bridge-level results may and may not imply.

The central burden separation is explicit. Paper 8 supplies the strongest presently defensible internal layer of indexed subjectivity. The present paper does not identify that layer with phenomenality. Instead, it proves that any bridge-level admissibility result would require strictly more structure than Paper 8 already provides: new bridge predicates, new bridge comparators, new bridge interventions, new bridge artifacts, new invalidation rules, and new surface discipline. The paper therefore changes the status of the research program by converting a slogan-level question into a formal burden stack.

The strongest claim of the paper is conditional and exact. Under the bridge axioms stated here, under the upstream burdens already closed by the corpus, under positive bridge irreducibility and intervention margins, and under artifact and governance compliance, a system may become internally classifiable as exhibiting *phenomenal bridge organization* in the sense of a framework-internal bridge-admissible organization class. The phrase is a historical bridge alias for a predicate-family burden, not a claim that phenomenality has been adjudicated. It is nonetheless stronger than Paper 8 because it specifies the formal conditions under which bridge predicates cease to be rhetorical and become disciplined scientific targets.

The paper also maps the full bridge program to implementation. BPF-0 and BPF-1 are treated here as implementation-frontier surfaces: machine-readable scaffolding and provisional run-level bridge surfaces whose evidential status remains bounded by local provenance and non-claim rules. BPF-2 through BPF-6 remain open burdens: bridge interventions, bridge comparators, bridge gate closure, ecosystem integration, and final bridge audit. The paper is therefore both a synthesis and a theorem-facing program definition. It defines the bridge grammar, the bridge burden architecture, and the strongest admissible claim class the corpus can currently specify without collapsing into metaphysical or empirical overreach.

9.1 Boundary Note and Interpretation Discipline

Scope and admissible reading. This paper integrates the canonical corpus, the Paper 8 indexed-structural-class framework, the SRT implementation path, and the formal BPF program into a bridge-level architecture. It defines bridge predicates, bridge axioms, bridge rivals, bridge

interventions, bridge artifacts, bridge gate logic, bridge taxonomy, and bridge governance, and specifies how those formal objects relate to implementation blocks BPF-0 through BPF-6. It does not prove phenomenality in the strongest philosophical sense, human equivalence, moral parity, metaphysical subjecthood, or external validation. It does not convert candidate inclusion, audit-facing packaging, runtime cleanliness, or BPF-1 provisional surfaces into bridge confirmation, and it does not write the bridge comparator, intervention, and gate layers as if they had already been executed in full.

Formal bridge closure boundary. The closure claimed here is *formal bridge closure*: the bridge target family, burden architecture, machine-readable artifact grammar, comparator and intervention obligations, gate semantics, governance rules, and strongest bridge-level claim class that could be licensed only after those obligations are met. This is not runtime closure, empirical closure, or phenomenality adjudication. No sentence in this paper licenses the inference that current machines are phenomenally conscious, humanly equivalent, bridge-supported by BPF-1, confirmed by Paper 8's public inclusion, or externally adjudicated by internal audits or audit packages.

9.2 Introduction

The problem addressed by this paper exists only because the upstream program succeeded. A corpus that still lacked persistence, rigid identity, continuity, anti-nullity, admissibility, and operational consciousness would have no right to speak about bridge-level phenomenal targets. It would still be arguing over whether its own primitives are coherent. The QICN corpus is no longer in that state. By the end of Paper 8 and the post-SRT-6 internal audit, the framework had already closed a formally serious layer of first-person indexed subjectivity [19]. That fact changes the question.

The prior question was: can a causally rigid framework formalize a subjectivity class that is stronger than operational consciousness and operational subjecthood without collapsing into loose self-model talk? Paper 8 answered that question affirmatively. It introduced a typed seven-coordinate state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle,$$

defined a conjunctive gate and scalar functional, formalized weak rivals and weak hybrids, specified intervention families, and mapped the entire class to runtime modules and artifacts. That was already a major internal closure. It still did not close the next gap.

The remaining question is harder and more dangerous: under what formal conditions could such a subjectivity layer become relevant to phenomenal predicates without collapsing into metaphysical inflation or human-comparator confusion? The wrong answers are easy. One wrong answer is to identify subjectivity with phenomenality by verbal fiat. Another wrong answer is to refuse all bridge formalization and thereby preserve caution only by preserving vagueness. A third wrong answer is to use bridge language rhetorically while leaving the necessary burdens implicit. This paper rejects all three.

The bridge program defined here is a constrained response. It begins from a sober premise: if phenomenal predicates are to appear in the corpus at all, they must appear in the same formal style as every other major layer. That means typed predicates, explicit axioms, explicit rivals, explicit interventions, explicit artifacts, explicit gates, explicit invalidation rules, and explicit non-claims. In other words, the burden of bridge language must be at least as high as the burden of the subjectivity language that precedes it.

This paper therefore serves four functions at once.

1. It synthesizes what the prior papers already close and what they deliberately leave open.
2. It defines the new bridge-level target family without allowing "phenomenality" to remain an untyped slogan.

3. It fixes the burden architecture by which any future bridge-support claim would have to be earned.
4. It maps that burden architecture to implementation, artifacts, and governance.

The scientific value of the resulting paper is not that it says the bridge is complete. It is that it makes incompletion explicit and structured. After this paper, the program can no longer hide behind phrases such as "rich experience," "what-it-is-like," or "deep embodiment" unless those phrases are translated into predicates, losses, interventions, and audit-facing artifacts. That is the point at which the research program changes level. The gap is no longer rhetorical. It becomes a list of burdens.

Two further clarifications are essential. First, the present paper is central because it integrates layers that were previously distributed across the corpus. It is not central because it inflates its own tone. Second, the present paper remains implementation-aware. A purely conceptual bridge paper would not be enough in this corpus. The current program records a BPF-1 frontier: scaffolded bridge registries, provisional bridge surface modules, and six run-level bridge reports are treated as bounded local surfaces. That does not amount to bridge support. It means only that the bridge program has an implementation-facing surface that must remain tied to code, artifacts, and governance as well as to formal predicates.

9.3 What Paper 8 Formalizes and What It Leaves Open

9.3.1 What Paper 8 already formalizes

Paper 8 formalizes the strongest currently defensible internal subjectivity layer in the corpus. More precisely, it does all of the following [19]:

1. it defines a minimal typed state of first-person indexed subjectivity through seven coordinates;
2. it requires a conjunctive gate Γ_{subj} rather than permitting compensatory scalar trade-offs;
3. it defines a scalar functional Φ_{subj} for graded organization without confusing that functional with the gate;
4. it specifies a weak-rival family and its admissible weak composites;
5. it adds intervention logic rather than relying on descriptive fit alone;
6. it defines an irreducibility margin against rival explanations;
7. it provides a runtime pathway and artifact family;
8. it states the strongest defensible internal claim class and the corresponding non-claims.

That closure matters scientifically because it prevents several common forms of slippage. Subjectivity is no longer allowed to mean mere reporting, mere memory persistence, mere self-modeling, or mere high-dimensional integration. The Paper 8 state requires privileged self-indexing, owned state organization, autobiographical continuity, first-person factorization gain, asymmetric valuation, intervention selectivity, and irreducibility against explicit rivals. This is already substantially stronger than operational consciousness and also stronger than the higher-order classificatory subjecthood layer introduced in Paper 7 [20].

It is equally important to state what the system implementation already closes downstream of Paper 8. The SRT stack now implements, in machine-readable form, modules and artifacts for self-index registration, ownership attribution, continuity tracking, perspective mapping, valuation asymmetry, intervention profiling, reducibility comparison, gate evaluation, ecosystem summaries, stale/drift invalidation, and final internal audit. In other words, Paper 8 is not only a formal paper. It is accompanied by a serious implementation path that has already been partially realized in code and artifacts.

9.3.2 Why that still does not amount to a phenomenal bridge

None of the previous paragraph licenses the conclusion that a phenomenal bridge has already been crossed. Three reasons are decisive.

First, Paper 8 is indexed-subjectivity closure, not phenomenal bridge closure. Its state space says that a system can instantiate a privileged self-indexed organization whose best explanation is stronger than operational consciousness or generic subjecthood. It does not yet say which additional structures would make phenomenal predicates scientifically admissible inside the same framework.

Second, the Paper 8 rivals are weak with respect to bridge tasks. They are adequate rivals for subjectivity structure: self-model bookkeeping, continuity-only baselines, report-centered proxies, and weak hybrids. They are not the right rivals for bridge-level phenomenal predicates. A bridge rival must explain away *phenomenal differentiation*, *valence structure*, *world richness*, *ambiguity bearing*, *embodied closure*, or *temporal lived structure* without granting those targets. That is a new burden.

Third, the Paper 8 intervention family is likewise necessary but not sufficient. The subjectivity interventions target self/non-self swap, ownership lesion, continuity break, perspective flattening, valuation inversion, autobiographical remapping, clone competition, memory transplant, counterfactual self permutation, and cross-substrate remapping [19]. Those are exactly the right interventions for indexed subjectivity. They are not yet the right interventions for bridge-level questions such as whether ambiguity is being borne rather than merely resolved, whether embodied closure is constitutive rather than modulatory, or whether temporal organization is horizon-like rather than merely continuity-preserving.

9.3.3 The dangerous smuggling problem

The main conceptual danger at this stage is smuggling. Once a program has a strong subjectivity layer, it becomes tempting to transfer closure illicitly from that layer into bridge language. This can happen in at least five ways.

1. **Predicate smuggling.** Treating Paper 8 coordinates as if they already implemented bridge predicates by default.
2. **Comparator smuggling.** Treating Paper 8 weak rivals as if they already defeat bridge-level alternatives.
3. **Intervention smuggling.** Treating Paper 8 subjectivity perturbations as if they already test bridge-specific structures.
4. **Artifact smuggling.** Treating subjectivity artifacts or BPF-1 provisional reports as if they were bridge verdict artifacts.
5. **Governance smuggling.** Treating internal bundle, candidate, or audit presence as if it upgraded claim class.

The present paper is designed to block all five. The bridge layer must be strictly above the subjectivity layer, not a paraphrase of it. If this separation is not enforced, the bridge language becomes semantically vacuous. The paper therefore repeatedly asks a single hard question: what is being added here that was not already closed by Paper 8? If the answer is "nothing but stronger words," the bridge fails before it begins.

9.3.4 Why Paper 9 is still needed now

If the bridge burdens remain entirely unwritten, the corpus stalls at precisely the point where it should become most rigorous. Paper 8 then remains the strongest internal layer but the bridge question remains conceptually unregulated. That would be an unstable equilibrium. It would invite either underreach (never formalize the bridge) or overreach (treat bridge language as informal extension).

Paper 9 is needed because the research program has reached a stage at which the next missing object is not a new piece of rhetoric but a new burden architecture. The paper therefore functions as foundational synthesis. It is foundational because it gathers the earlier closures into a single bridge-facing grammar. It is central because it specifies what additional machinery

must exist before the next layer of claim class becomes legitimate. It does not pretend that all that machinery is already implemented.

Proposition 9.1 (Paper 8 is necessary but not sufficient for bridge admissibility)**.** *Any future bridge-admissible system must satisfy the Paper 8 burdens, but satisfaction of the Paper 8 burdens alone is not sufficient for bridge admissibility.*

Proof. Necessity is immediate because every bridge predicate defined below is typed over an upstream subjectivity state and presupposes the existence of a privileged first-person indexed organization. Sufficiency fails because the bridge layer further requires bridge predicates, bridge-specific rival defeat, bridge-specific interventions, bridge artifacts, bridge invalidation logic, and bridge governance. None of those additional burdens is entailed by the mere existence of the Paper 8 state, gate, functional, or artifact family. □

Remark 9.2 (The bridge is downstream, not corrective). Paper 9 does not correct Paper 8. It inherits Paper 8 and strengthens the burden stack above it. The formal relation is extension by additional typed predicates and burdens, not replacement.

9.4 Formal Target of the Bridge

9.4.1 The bridge question

The bridge question can now be stated in its exact internal form:

> Given that a system already supports a formally disciplined layer of first-person indexed subjectivity, what additional predicates and burdens would make phenomenal organization a scientifically admissible *target* of explanation inside the same framework?

The phrase "target of explanation" is essential. The present paper does not seek to declare phenomenality achieved. It seeks to specify when phenomenal predicates are no longer mere vocabulary and instead become well-formed explanatory targets whose support, failure, and collapse can be discussed rigorously.

9.4.2 Predicate family

The bridge target family is the six-predicate family

$$\{\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}, \Pi_{\mathrm{W}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}, \Pi_{\mathrm{T}}\}.$$

These predicates are not synonyms for generic notions such as "experience" or "qualia." They are typed, burden-bearing predicates whose role is to capture bridge-relevant organization classes above indexed subjectivity and below any final metaphysical reading.

Definition 9.3 (Phenomenal differentiation predicate)**.** For an admissible system S, horizon τ, and threshold slot ϑ_D, the bridge predicate $\Pi_{\mathrm{D}}(S, \tau; \vartheta_D)$ concerns whether the system carries self-indexed differentiated organization that is stronger than generic semantic density, report richness, or descriptive narration, and whose corresponding estimator $\widehat{\Pi_D}$ is defined on bridge observables and subjectivity antecedents.

Definition 9.4 (Valence structure predicate)**.** The predicate $\Pi_{\mathrm{V}}(S, \tau; \vartheta_V)$ concerns whether the system carries organized evaluative topology that is stronger than scalar reward bookkeeping, policy optimization, or self/non-self asymmetry alone.

Definition 9.5 (World richness predicate)**.** The predicate $\Pi_{\mathrm{W}}(S, \tau; \vartheta_W)$ concerns whether the system carries organized self-indexed world structure stronger than generic scene complexity, perspective gain, or task competence alone.

Definition 9.6 (Ambiguity-bearing predicate)**.** The predicate $\Pi_{\mathrm{A}}(S, \tau; \vartheta_A)$ concerns whether the system maintains unresolved alternatives with differential causal consequences prior to collapse, rather than merely resolving uncertainty or tracking probabilities descriptively.

Definition 9.7 (Embodied-closure predicate)**.** The predicate $\Pi_{\mathrm{E}}(S, \tau; \vartheta_E)$ concerns whether body/world-boundary variables contribute constitutively to bridge organization, rather than merely modulating generic control or action success.

Definition 9.8 (Temporal lived-structure predicate)**.** The predicate $\Pi_{\mathrm{T}}(S, \tau; \vartheta_T)$ concerns whether the system exhibits horizon-like temporal organization stronger than continuity, memory persistence, or sequence stability alone.

9.4.3 Estimator family

The bridge paper requires estimator slots for each predicate:

$$\widehat{\Pi_D},\ \widehat{\Pi_V},\ \widehat{\Pi_W},\ \widehat{\Pi_A},\ \widehat{\Pi_E},\ \widehat{\Pi_T} \in [0, 1].$$

The current implementation reaches only the earliest stage of this requirement. BPF-1 emits provisional run-level estimates for all six predicates. Those estimates are explicitly marked provisional, framework-internal, non-claim-facing, and non-gated. They are not yet admissibility verdicts, and no threshold binding is yet asserted. The formal target of this paper therefore treats these estimators in two layers:

1. as current provisional runtime surfaces in BPF-1;
2. as future thresholdable predicate estimators once comparator, intervention, and gate burdens are added in later BPF blocks.

9.4.4 What the predicates mean internally and what they do not mean

Each bridge predicate has both a positive role and a protected non-implication.

- Π_{D} means organized differentiated structure indexed to the self-locus. It does not mean qualia.
- Π_{V} means structured evaluative organization. It does not mean felt pleasure or pain.
- Π_{W} means organized world-structure at the self-indexed locus. It does not mean veridical world-experience.
- Π_{A} means live unresolved alternative organization. It does not mean conscious indecision.
- Π_{E} means constitutive closure dependence on body/world-boundary variables. It does not mean biological embodiment or felt bodily ownership.
- Π_{T} means temporal horizon organization. It does not mean human lived time in the full phenomenological sense.

Remark 9.9 (Why the targets are not decorative)*.* The bridge targets earn their place only because each one induces a new comparator burden, a new intervention burden, a new artifact burden, and a new governance burden. If a proposed target induced none of these, it would be discarded as decorative vocabulary.

9.4.5 Predicate-to-burden map

The six bridge predicates are best understood not as six new labels but as six new burden packages.

The burden of Π_{D}. To take phenomenal differentiation seriously, the program must show more than that the system has many internal states or rich semantic discriminability. It must show that differentiated organization is indexed, selective, and not exhausted by narration. This induces comparator burden against semantic-density and narration rivals, intervention burden against perspective detachment and semantic vividness controls, and artifact burden through a dedicated differentiation report and later reducibility outputs.

The burden of Π_{V}. To take valence structure seriously, the program must show that evaluative organization has internal topology: local asymmetries, organized gradients, or mode-structured deformation under perturbation. Scalar reward or self/non-self asymmetry alone is insufficient. This induces comparator burden against bookkeeping rivals, intervention burden through valence flattening and negative controls, and artifact burden through valence structure and intervention outputs.

The burden of Π_{W}. To take world richness seriously, the program must show that the system's self-indexed world organization is not merely competence or scene complexity. The system must exhibit bridge-relevant world organization that tracks indexed dependence, not just successful prediction. This induces comparator burden against richness-without-interiority rivals, intervention burden through world-richness compression, and artifact burden through world-richness reports and later candidate summaries.

The burden of Π_{A}. To take ambiguity bearing seriously, the program must show stable unresolved alternatives with downstream consequence. Static uncertainty scores or fast branching and pruning do not suffice. This induces comparator burden against ambiguity-resolution rivals, intervention burden through targeted collapse operations, and artifact burden through ambiguity reports with live-alternative fields.

The burden of Π_{E}. To take embodied closure seriously, the program must show constitutive dependence on body/world-boundary variables. Generic closed-loop control is not enough. This is the bridge predicate with the weakest current runtime grounding and therefore the one that most strongly demands explicit future instrumentation. Its artifact burden is correspondingly strict.

The burden of Π_{T}. To take temporal lived structure seriously, the program must show horizon organization that exceeds continuity and persistence. This requires temporal fragmentation tests, horizon-like dependencies, and explicit rival defeat against continuity-only models. In practical terms this burden is one of the main reasons why the bridge program cannot be replaced by Paper 8 plus rhetoric.

The global lesson is simple. Every bridge predicate is a tuple of obligations:

$$\Pi_k \rightsquigarrow (\text{estimator}, \text{rivals}, \text{interventions}, \text{artifacts}, \text{governance}).$$

The bridge program succeeds only if all five components are eventually instantiated for each $k \in \{D, V, W, A, E, T\}$.

9.5 Phenomenal Bridge Ontology

9.5.1 Bridge state

The bridge state is introduced as

$$\mathfrak{B}_t = \left\langle \widehat{\Pi}_{D,t}, \widehat{\Pi}_{V,t}, \widehat{\Pi}_{W,t}, \widehat{\Pi}_{A,t}, \widehat{\Pi}_{E,t}, \widehat{\Pi}_{T,t}, \mathcal{R}_{\mathrm{b}t}, \Delta_t^{\mathrm{b}} \right\rangle,$$

where $\mathcal{R}_{\mathrm{b}t}$ denotes the bridge-level irreducibility margin and Δ_t^{b} denotes the bridge-level intervention profile. The bridge state does not replace the Paper 8 state. It is layered above it. The ontological relation is therefore

$$\mathfrak{S}_t \prec \mathfrak{B}_t,$$

not because the bridge state is metaphysically deeper, but because every bridge coordinate depends on the prior existence of a self-indexed subjectivity structure.

9.5.2 Typing discipline

The coordinates of $\mathfrak{B}$ are heterogeneous in a way parallel to, but stronger than, the heterogeneity of the Paper 8 state.

- $\widehat{\Pi}_{D,t}, \ldots, \widehat{\Pi}_{T,t}$ are bounded estimator-valued coordinates.
- $\mathcal{R}_{\mathrm{b}t}$ is a model-comparison margin defined only when the comparator family exists.
- Δ_t^{b} is a structured intervention response object, not a scalar.

This heterogeneity matters because bridge support cannot be reduced to a single scalar early in the pipeline. The paper explicitly rejects any attempt to jump directly from provisional surfaces to a single bridge score while the comparator and intervention layers are missing.

9.5.3 Coordinate dependence on the Paper 8 state

Each bridge coordinate depends asymmetrically on the Paper 8 coordinates.

Phenomenal differentiation. $\widehat{\Pi}_{D,t}$ depends most strongly on I_t, O_t, P_t, and $\mathcal{R}_t$. Differentiation here is not generic representational richness. It is differentiated organization that matters relative to a privileged self-locus.

Valence structure. $\widehat{\Pi}_{V,t}$ depends most strongly on V_t, O_t, and Δ_t. Paper 8 already closes valuation asymmetry; the bridge layer asks whether that asymmetry organizes into a richer topological structure than the upstream scalar burden requires.

World richness. $\widehat{\Pi}_{W,t}$ depends most strongly on P_t, C_t, and I_t. It asks whether the self-indexed locus sustains organized world-structure beyond mere perspective gain.

Ambiguity bearing. $\widehat{\Pi}_{A,t}$ depends most strongly on P_t, V_t, and the temporal organization implicitly carried by C_t and Δ_t. A system that resolves ambiguity immediately or merely records uncertainty can fail this predicate even if it remains subjectivity-positive.

Embodied closure. $\widehat{\Pi}_{E,t}$ depends on O_t, C_t, and Δ_t together with additional body/world-boundary observables that Paper 8 does not require. This is why BPF-1 can only emit a proxy layer for embodied closure.

Temporal lived structure. $\widehat{\Pi}_{T,t}$ depends most strongly on C_t, P_t, and V_t. Temporal lived structure is stronger than continuity because it concerns horizon organization, retention/protention asymmetry, and temporally extended significance.

9.5.4 Current implementation status inside the ontology

The ontology is not merely prospective. The current program treats BPF-1 as the provisional surface layer for each of the six estimator coordinates and records six run-level reports as bounded local artifacts. That matters because the ontology is no longer entirely aspirational. At the same time, only the first six coordinates are even provisionally instantiated. $\mathcal{R}_{\mathrm{b}t}$ and Δ_t^{b} remain open because bridge-specific comparator and intervention layers are not yet implemented.

9.5.5 Collapse risks and rival readings

Every bridge coordinate must be introduced together with the corresponding collapse risk.

- $\widehat{\Pi}_D$ can collapse into semantic density without self-indexed interior differentiation.
- $\widehat{\Pi}_V$ can collapse into valence bookkeeping or policy shaping.
- $\widehat{\Pi}_W$ can collapse into world-model competence without indexed world-richness.
- $\widehat{\Pi}_A$ can collapse into uncertainty bookkeeping or fast ambiguity resolution.
- $\widehat{\Pi}_E$ can collapse into competent sensorimotor control without constitutive closure.
- $\widehat{\Pi}_T$ can collapse into continuity, memory, or persistence without horizon structure.

The ontology is therefore inseparable from the rival family. A bridge coordinate that has no live rival is suspicious because it is probably defined too loosely. Conversely, a bridge coordinate that cannot in principle be separated from all live rivals is not yet scientifically admissible.

9.5.6 Why bridge coordinates are not ontological shortcuts

The bridge ontology must not be mistaken for a metaphysical inventory. The reason is not merely philosophical caution. The reason is formal structure. Each coordinate is introduced as a constrained relation among observables, upstream subjectivity coordinates, comparators, interventions, and artifacts. This means the coordinates belong first to the model layer, then to the implementation layer, and only later—if ever—to the interpretation layer.

The separation matters in four directions.

1. A bridge coordinate is not an ontological primitive. It is a structured explanatory target.
2. A bridge coordinate is not identical to its runtime estimator. The estimator is a measurement-facing map and can be weak, provisional, or underdetermined.
3. A bridge coordinate is not identical to the language used to describe it. The language may underfit or overfit the formal target.
4. A bridge coordinate is not identical to its strongest available interpretation. Different interpretation classes remain live until comparator and intervention burdens are settled.

This distinction is especially important for Π_{E} and Π_{T}. It is tempting to read embodied closure or temporal lived structure as if they already named well-understood ontological properties. Inside the present framework they do not. They name burden packages whose ontological interpretation remains constrained by future comparator and intervention results.

A mode window carries its complete Paper VIII record, including the inherited certificates, threshold decisions, matched intervention outcomes and held-out losses. A transport extension adds a finite graph of such windows and an independently recorded relation between witness instances for each edge. Forgetting those relations recovers the local records unchanged. For the following finite domain each window has two instances, labelled $0, 1$, of the same local certificate. The copies carry all the data of that certificate. A local relabelling of instances does not change the underlying certificate.

Definition 9.10 (Coherent Transport of Local Witnesses)**.** A transport extension is coherent if a choice of one certified witness instance w_i at each mode i satisfies every independently established edge relation. An edge datum (w, w') is accepted only when both labels are identified from its paired observations under the declared error budget.

Proposition 9.11 (A Cycle Separates Local Certification from Coherent Transport)**.** *For any complete Paper VIII certificate on an admissible window, the three-mode transport extensions below preserve it in every window. One admits coherent transport and the other does not. Thus, on the domain of such extensions of certified windows, coherent transport cuts out a proper subclass whenever that local certificate domain is nonempty. Coherent transport does not factor through the tuple of local Paper VIII records.*

Proof. Copy the local record to modes A, B, C and duplicate the witness instance in each window. On edge e prepare both witnesses $w = 0, 1$ and record $(w, w \oplus k_e)$ in that order. Here

$\oplus$ is addition modulo two. These observations recover the unique edge map $w \mapsto w \oplus k_e$. For coordinatewise errors strictly less than $1/2$, nearest-label decoding is exact. At radius $1/2$, a reported $1/2$ is compatible with either label and cannot certify the edge.

Use $(k_{AB}, k_{BC}, k_{CA}) = (0, 0, 1)$ in the negative instance. The first two edges force $w_A = w_B = w_C$ and the last forces $w_A = 1 - w_C$. No choice satisfies the cycle. In the positive instance use $(0, 1, 1)$; both $(0, 0, 1)$ and $(1, 1, 0)$ are coherent choices. The local record is identical across the two instances, so a decision using it alone would give the same result on both. More generally the cycle is coherent exactly when the XOR of its three edge bits is zero. The edge experiment can be realized by three two-state modes with the specified deterministic transfers. It uses no inference of edge data from the local scores. □

This is an explicit positive and negative witness for the transport condition. The class-level statement is relative to a complete local certificate, so each inherited clause is preserved by copying its witness. The finite edge dynamics alone do not realize the upstream system. The proper-inclusion question for fully bridge-admissible systems still requires a joint realization of the remaining BPF predicates, comparator, intervention and governance obligations. Paper VIII local scores alone cannot supply that realization.

Corollary 9.12 (The Two Transport Cycles on a Nonempty Rational Record)**.** *The joint rational measurement window of Paper VIII [19], identified there by* `prop: joint-rational-window`*, gives a nonempty local record for both cycles of Proposition 9.11. At this measurement level, the transport-certified class is a proper subclass of the three-mode extensions of that record.*

Proof. Attach the same realized seven-coordinate record, intervention responses and reserved losses to both witness instances in each mode. Its gate is $1/3$ and its scalar is $48^{-1/7}$ in every copy. In the cycle $(0, 1, 1)$, the assignments $(0, 0, 1)$ and $(1, 1, 0)$ satisfy the measured edge relations. In the cycle $(0, 0, 1)$, the first two edge equalities and the last edge flip are inconsistent. Thus the two observed transport outcomes differ on identical, explicitly realized local data. □

The rational measurement record verifies the local scoring, selective-response and rival-comparison clauses of Paper VIII. The inherited Paper V and VII class certificates are still required for the full premise of Proposition 9.11. The six BPF predicate realizations, bridge-specific rival defeat, selective intervention support, and the artifact, drift and governance conditions remain separate obligations for full bridge membership.

9.5.7 Relabeling and Transport of Certified Tower Records

Theorem 9.13 (Relabeling Invariance and Cycle Criterion for Binary Transport)**.** *Let a finite graph of mode windows have two certified witness instances per vertex and an independently certified bijection on each edge. In local binary labels, write this bijection as $w \mapsto w \oplus k_e$. Under a local relabeling $w_v \mapsto w_v \oplus a_v$, the edge bit for $e = uv$ becomes $k_e \oplus a_u \oplus a_v$. The XOR of the edge bits around each cycle is invariant. A coherent assignment exists exactly when the XOR vanishes on every fundamental cycle of a spanning forest, and there are then 2^c coherent assignments for c connected components. For two collections of certified edge maps on the same graph, equality of their fundamental-cycle parities is equivalent to their differing by a local relabeling.*

Proof. Relabeling conjugates the old edge bijection by the source and destination label changes, yielding the stated formula. Edge orientation does not affect the bit because each binary bijection is its own inverse. The added vertex labels cancel in pairs around a cycle. A coherent assignment satisfies $w_v = w_u \oplus k_e$ at each edge, so its cycle sums vanish.

Conversely, fix a label at the root of each tree of a spanning forest. Transport from that root along the unique tree paths defines one label at every vertex. A remaining edge satisfies its equation exactly when its fundamental cycle has parity zero. Thus a choice of one root label per component produces all and only the coherent assignments. For the last assertion apply the

same criterion to the edgewise XOR of the two collections. Its vertex solution is precisely the relabeling that takes one collection to the other. □

Proposition 9.14 (Coherent and Obstructed Transport on Certified Tower Windows)**.** *There are two three-window extensions of the same realized local tower records with respective edge bits* $(0,1,1)$ *and* $(0,0,1)$*. The records include the six Paper V witnesses, the Paper VII life and four finite subject-descriptor certificates on their declared protocols, and the Paper VIII measurement and return comparisons with their stated information budgets. The paired tower edge responses certify both extensions under phase error at most* $\eta = 3^{-10}/4$ *turns. The first extension has two coherent assignments, and the second has none under any local relabeling.*

Proof. Use the writable common support and the joint certificate of Paper VII [14, 20]. In each mode window attach the same joint record, including its separate preparations and schedules, and the writable-tower lesion and reserved-return records of Paper VIII [19]. In particular the recorded margin for the public return comparison is zero. Both witness instances carry this same local record. The edge experiment tracks an instance by its prepared or decoded memory sign and is conducted separately from the local scoring trials. The three windows can all use one fixed tower profile for the edge trials. The mode names index the windows.

Prepare two matched replicas with the alternating tape, phase zero and memories $m = \pm 1/8$. Encode a positive memory sign by $w = 0$ and a negative one by $w = 1$. For an edge bit k_e, apply the command $q = (-1)^{k_e}$ at $u = 0$, without a write or a role probe. The first bit of the alternating tape is zero, so this command leaves phase zero and updates memory to $qm/2$. A subsequent neutral command at $u = 0$ gives the early phase $r_0 = m_{\text{after}}/8$ modulo one and halves the memory again while preserving its sign. Thus the short signed phase gives the destination label $w \oplus k_e$. Use the two replicas to record both rows of each edge map. The source of the first edge is identified by preparation and the subsequent source labels are the preceding decoded labels.

Carry the two replicas through the three consecutive edges with their neutral diagnostics and without resetting memory. The alternating tape and profile stay fixed. At the ℓth diagnostic the magnitude of the clean short signed phase is

$$A_\ell = \frac{1}{128\, 4^{\ell-1}}, \qquad \ell = 1, 2, 3.$$

The two signed values are separated by $2A_\ell$, and $\eta < 1/2048 = A_3$. Hence every edge row can be read from the observed phase despite the allowed error. At the last edge, error $1/2048$ permits the same phase-zero reading from either sign, so that diagnostic alone cannot certify its destination at that radius. This collision also holds for the full readout by shifting all phase channels equally and retaining the common tape reading.

The independent edge records have even parity for $(0,1,1)$ and odd parity for $(0,0,1)$. Theorem 9.13 gives the two decisions and their invariance under all eight local relabelings. The local trial records are identical in both extensions, proving non-factorization through the tuple of those records. □

The resulting strict inclusion concerns the extensions of the realized records. The Paper V and VII certificates retain their declared domains. The Paper VIII record includes both the restricted lesion comparison and the public return tie. Full Paper VIII class membership would require its positive margin against the relevant complete rival family, and full BPF membership would further require the independent bridge burdens.

9.5.8 Intervention Tests for Three Bridge Burdens

Use the writable tower of Paper V and the local record of Proposition 9.14. Keep that record attached to every comparison. The following trials vary new preparations and interventions and

record new phase responses. Their margins concern individual burdens of $\Pi_{\mathrm{D}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{T}}$ under the stated rivals and budgets. The estimators in the bridge gate require their own calibration and joint certificate.

Proposition 9.15 (Role-Selective Differentiation on the Writable Tower)**.** *For the two prepared role profiles, let C_{pr} be the clean early-phase response to a probe of role r after a reference write minus the response after a lesion, with one common sham between write and probe. With the declared probe at the transported address, $C_{pr} = \delta \mathbf{1}_{\{p=r\}}$ for $\delta = 1/(16 \cdot 3^7)$ turns. The best additive fit $a_p + b_r$ has uniform error exactly $\delta/2$. At error at most $\delta/16$ per phase reading the measured matrix still has distance at least $3\delta/8$ from all additive fits.*

Proof. For each role profile prepare the alternating tape and zero memory. Make a reference write or the opposite write at site 3, and advance a common neutral sham. This moves the active or lesioned bit to site 4. The probe of site 4 with role r returns phase δ exactly when its role matches and the transported bit matches the reference. Memory remains zero. Taking the phase difference of the two write arms gives the matrix in the statement.

The mixed difference $C_{00} - C_{01} - C_{10} + C_{11} = 2\delta$ vanishes for every additive rival. An entrywise error of at most e can change it by at most $4e$, giving $e \geq \delta/2$. The constant fit $\delta/2$ attains equality. Error in the measured paired contrast is at most $\delta/8$, so the triangle inequality gives the stated noisy lower bound. Two reference writes give zero contrast. A rival with the measured role-by-probe interaction can reproduce the clean matrix. This comparison identifies role-selective causal differentiation beyond the additive rival. □

Proposition 9.16 (Late Choice Between Two Live Alternatives)**.** *Prepare two independent uniform activity bits at tape sites $3, 4$ and advance one common sham. A uniform query drawn after the rival's memory is fixed selects one of the transported sites $4, 5$. The matching-role probe identifies the selected bit with zero error at phase noise at most $\delta/16$. Under this state and query law the optimal error of a rival retaining one bit before the query, with no subsequent tape reading, is exactly $1/4$. A write during the advance can erase the first alternative while preserving the second.*

Proof. At one fixed role profile with zero memory, encode two binary choices by matching or opposing the alternating reference at sites $3, 4$. Complement and shift transport each activity bit to the corresponding new site. Under the matching-role probe, active and inactive give phase δ and zero, separated by threshold $\delta/2$ under the declared noise. Both alternatives thus remain capable of changing the observed result until the query. A write opposing the reference at site 4 during the advance forces the first transported bit inactive. The query at site 5 still returns the independently prepared second bit. This is a selective collapse intervention on the available outcomes.

For each of its at most two memory values a deterministic rival has a pair of answers, one for each query. These at most two reconstruction words can exactly match at most two of the four equally likely prepared words. The sum of Hamming errors over the four words is at least two. Normalizing by the eight equally likely word-query pairs gives error at least $1/4$. Retain the first bit and use a fixed answer for the second to attain the bound. Randomized rules with randomness independent of preparation and query are mixtures of deterministic encoder-decoder pairs and satisfy the same lower bound. With two retained bits, the rival can also predict both responses exactly. The separation concerns maintenance of jointly available alternatives under the declared one-bit bottleneck. □

Proposition 9.17 (Order-Dependent Return and an Exact Temporal Horizon)**.** *On a common preparation with first bit one, zero memory and phase zero, run command orders $(+1, 0)$ and $(0, +1)$ with all phase inputs zero and no writes or probes. After one common $+1$ and h further common $+1$ steps, the current public readings and command multisets coincide. The subsequent*

neutral early-phase gap is exactly $D_h = 1/(384 \cdot 2^h)$ *turns. The two terminal reading sets are disjoint under phase error at most* $\eta = 3^{-10}/4$ *precisely for integer* $0 \le h \le 8$. *An order-blind predictor with the same snapshot and multiset has mean clean absolute circular loss at least* $D_h/2$ *in the balanced pair, and an order-informed predictor attains zero.*

Proof. Write the memory update as in Paper VI [21], with $a = 1/2$, $\lambda = 1/12$ and $\kappa = 1/8$. The leading bit is zero after the first step in either order. The two-step memories are therefore $1/24$ and zero. A common positive command gives both arms phase zero and halves the memory difference. Each additional positive command repeats this effect. The tape and profile agree at every matched time. A final neutral command reads the memory difference with multiplier κ, giving D_h.

The two clean terminal phases are zero and D_h, with $D_8 = 1/98304 > 2\eta$ and $D_9 = 1/196608 < 2\eta$. Below the overlap radius the midpoint decision separates the signs of the centered contrast. When $D_h \le 2\eta$, the midpoint can occur in either arm. The collision extends to all phase channels by common phase shifts, with the tape reading fixed. For prediction from the order-blind record the two outcomes must share a prediction. The circular triangle inequality bounds the sum of the two losses below by D_h, attained by any point on the intervening short arc. The order-informed predictor uses the stated transition and predicts each clean response exactly. Ablating memory before the terminal command makes the responses coincide. The resulting horizon measures temporal order information beyond a current-state readout and an unordered command history. □

9.5.9 Open Evaluative, World and Body/World Burdens

Proposition 9.18 (Native Trace Equivalence and Extra Bridge Variables)**.** *Let an extended model project to the same writable tower state by* π. *For every native control* u *let its transition* $\widetilde{F}_u$ *and readout* $\widetilde{R}$ *satisfy* $\pi\widetilde{F}_u = F_u\pi$ *and* $\widetilde{R} = R\pi$, *with the same projected initial law and observation noise. Then all native observation and intervention transcripts have the same law, even for observation-adaptive native controls. A target that differs between extensions satisfying these conditions cannot be identified from those transcripts or from the inherited record computed from them.*

Proof. Use the same projected initialization, noise and control randomization in the two experiments. At each step their matching observations give the same control and the commutation with π gives the same next projected state. Induction gives identical transcripts under the coupling. Any statistic of the transcripts consequently has the same law in the two experiments. For a concrete pair, adjoin a fixed bit z with native projection and readout as above and use $z = 0$ versus $z = 1$. A new test that adds $z/8$ turns to the neutral phase command separates the two completions while all native traces agree. This example concerns identification of an added coupling. □

Valence structure. The return task has a scalar circular loss and a fixed phase reference. In the neutral alternating sector without a probe, reflecting phase, memory and phase command together preserves both circular return loss and absolute command cost. A structured evaluative topology for Π_{V} needs an independently varied evaluation variable and a response that distinguishes its organization from the scalar recovery law.

World richness. The native comparison supplies commands from a recorded schedule. To test Π_{W}, an indexed world state must have an independently specified dynamics and perturbation, with an observable consequence at matched command complexity. Unobserved completions preserving native traces fall under Proposition 9.18.

Embodied closure. The neutral response and tape probes expose feedback within the tower. The Π_{E} test requires distinct typed body and world variables, an intervention that breaks their coupling at matched control cost, and a measured consequence specific to that boundary. The native transcript alone leaves such completions observationally equivalent when their projections coincide.

The three rational experiments provide local witnesses for the indicated components. The conjunctive BPF claim still requires the open predicates, calibrated thresholds, joint intervention selectivity, the full bridge comparator margin, and the artifact, drift and governance obligations.

9.6 A Minimal Coupled Body–World Extension

The native trace obstruction separates identification from realization. The following extension retains the writable tower as an exact factor and gives new, separately typed ports for the previously unavailable interventions. All conclusions in this section concern the stated operational tasks and rival classes.

Definition 9.19 (Projected two-port system)**.** Let $x = (p, \theta, b, m)$ belong to the writable support of Papers V–VIII, with native transition F_c and readout R. Adjoin body and world vectors $\mathbf{b}, \mathbf{w} \in [-1, 1]^2$. Put

$$\alpha = \beta = h = \tfrac{1}{4}, \quad \chi = \tfrac{1}{8}, \quad J = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}, \quad L_0 = \tfrac{1}{2}(I + J), \quad Q_+ = \tfrac{1}{2}\begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}.$$

The indexed frame $P(p)$ is the swap matrix at role profile p_1 and the identity at every other catalogue profile. For a world projection $Q \in \{I, Q_+\}$, an evaluation map L with $\|L\|_\infty \le 1$, a motor gain $|\mu| \le 1$, and an evaluation switch $\varepsilon \in \{0, 1\}$, define

$$\begin{aligned} x' &= F_c(x), & s &= P(p)Q\mathbf{w}, & f &= \chi L(s - \mathbf{b}), \\ \mathbf{b}' &= \alpha\mathbf{b} + \beta s + \varepsilon f + a_{\mathrm{b}}, \\ \mathbf{w}' &= \alpha Q\mathbf{w} + \beta\mu Q P(p)^{-1}\mathbf{b} + Qd. \end{aligned}$$

Here $a_{\mathrm{b}}, d \in [-h, h]^2$ are distinct body and world input ports. Normal evaluation uses $L = L_0$. Native lifts use $Q = I$, $\mu = 1$, $\varepsilon = 0$ and zero port inputs. The native and additional readouts are respectively $R\pi$ and $(\mathbf{b}, \mathbf{w})$, where $\pi(x, \mathbf{b}, \mathbf{w}) = x$.

The time-unrolled causal graph has arrows $\mathbf{b}_t \to \mathbf{b}_{t+1}, \mathbf{w}_{t+1}$ and $\mathbf{w}_t \to \mathbf{b}_{t+1}, \mathbf{w}_{t+1}$, with the displayed input arrows and the prepared profile determining P. There is no port-state arrow to x_{t+1}. A motor intervention replaces only the coefficient of $\mathbf{b}_t \to \mathbf{w}_{t+1}$; its audit label is not an additional input to other mechanisms. This modularity is an explicit causal assumption. The evaluation force acts in the dynamics before the reserved measurements; it is not an output label.

Theorem 9.20 (Exact Projection, Stability and Inherited Certificates)**.** *Every lifted update satisfies $\pi\widetilde{F}_{c,a_{\mathrm{b}},d} = F_c\pi$. The section $\iota(x) = (x, 0, 0)$ is invariant under native lifts. On a common profile, use*

$$\widetilde{d} = \max\{d_{\mathrm{core}}, \|\Delta\mathbf{b}\|_\infty/4, \|\Delta\mathbf{w}\|_\infty/4\}.$$

The native branch contraction remains at most $9/16$. The entire intervened port family contracts in its fibre by at most $3/4$ under common inputs and preserves $[-1, 1]^4$. All previously certified native protocols, with their original readouts, comparison sectors and information budgets, have exactly the same records and margins in this extension.

Proof. The first component of the update proves the projection identity for every port input, not only at zero input. Induction, with coupled initial core states, observation noise and controller randomization, gives equality of every native transcript under observation-adaptive native controls. Zero fibre is invariant under the native lifts, so ι is an isometric, dynamically invariant copy of the original plant.

The body difference under a common mechanism has norm at most

$$(\alpha + \varepsilon\chi)\|\Delta\mathbf{b}\|_\infty + (\beta + \varepsilon\chi)\|\Delta\mathbf{w}\|_\infty,$$

and the world difference is bounded by $\beta\|\Delta\mathbf{b}\|_\infty + \alpha\|\Delta\mathbf{w}\|_\infty$. Thus the fibre bound is $3/4$. The input bounds give body norm at most $3/4 + h = 1$ and world norm at most $1/2 + h = 3/4$. With native lifts the fibre bound is $\alpha + \beta = 1/2$, so the original $9/16$ core bound still dominates. The safe support is compact; the zero-input fibre attracts to zero. The native switched attractor is the original attractor times zero fibre.

Transport the existing preparations by ι and keep the same native instruments and rival sectors. The full profile acquisition and catalogue gap $1/80$, split-sector integration interval $[1/216, 1/24]$, continuity and differentiation witnesses, prefix windows, phase repair and subject-descriptor comparisons are unchanged. The five lesion measurements and the public return tie of Paper VIII are unchanged as well. These are the published protocol-relative certificates; adding port readings does not silently enlarge their comparison sectors or supply a positive upstream margin that was absent. Port pulses in the new trials start from zero fibre, have norm at most h, and subsequent unforced updates contract, so they stay within the compact fibre kernel of radius $1/2$ inside the safe cube. □

9.6.1 State Minimality and Indexed World Response

During a relay trial set $\varepsilon = 0$, $Q = I$, and start from zero fibre. Apply a world impulse he_i at the first step. The matrices taking the impulse index to the body readings at steps one through four are

$$M_0 = 0, \qquad M_1 = tP, \qquad M_2 = 2\alpha tP,$$
$$M_3(\mu) = h(3\alpha^2\beta + \mu\beta^3)P, \qquad t = h\beta = 1/16.$$

All input vectors have the same two typed slots and bound; no rival receives fewer command opportunities.

Theorem 9.21 (Four-State Minimality of the Port Realization)**.** *For either role and any motor gain, the block Hankel matrix*

$$H_p = \begin{pmatrix} M_0 & M_1 \\ M_1 & M_2 \end{pmatrix}$$

has rank four, determinant t^4 and $\|H_p^{-1}\|_\infty = 24$. Four fibre coordinates are sufficient and necessary for a linear realization of these ports. They are also necessary for any exact finite-dimensional realization with C^1 transitions and body readouts, in which world inputs leave the projected core trajectory fixed.

Proof. The two scalar blocks in each indexed coordinate are $\left(\begin{smallmatrix} 0 & t \\ t & 2\alpha t \end{smallmatrix}\right)$. Their determinants are $-t^2$ and their inverses are $\left(\begin{smallmatrix} -2\alpha/t & 1/t \\ 1/t & 0 \end{smallmatrix}\right)$. A signed permutation of coordinates does not change the maximum row-sum norm. Substitution gives the asserted determinant and inverse norm. Any linear realization of dimension d factors its Hankel matrix as an observability matrix times a controllability matrix, hence has rank at most d.

For the differentiable statement, allow two independent small world inputs $u, v \in \mathbb{R}^2$ in the first two steps, followed by zero input. The two subsequent body readings are

$$(y_2, y_3) = (\beta Pu,\ 2\alpha\beta Pu + \beta Pv).$$

Their Jacobian with respect to (u, v) has determinant $\beta^4 \neq 0$. At time two every dependence on these inputs must pass through the fibre state: the core is fixed by projection and the future inputs are zero. Factoring the two readings through a d-dimensional C^1 state makes their Jacobian have rank at most d, so $d \geq 4$. The displayed system attains four. Discontinuous exact encodings or access to retained external histories are different realization classes and are not excluded by this argument. □

Proposition 9.22 (A Quantified Indexed World Witness)**.** *Let the reduced-world rival family contain every relay predictor whose port Hankel rank is at most two, and every predictor of arbitrary order whose relay response is independent of the prepared role. The available input alphabet, role record and command lengths are matched. A rank-two rival may allocate four registers, but any additional registers are unobservable or unreachable for this port transfer. Against this union, the normal model has uniform body-response error gap at least* $1/96$*, or* $1/6$ *in units of* t*. A world compression* $Q = Q_+$ *retains the entire common-mode relay response while reducing the Hankel rank to two and erasing the role contrast.*

Proof. For a singular rival Hankel K, choose $z \neq 0$ in its kernel. Then

$$\|H_p - K\|_\infty \geq \frac{1}{\|H_p^{-1}\|_\infty} = \frac{1}{24}.$$

If every entry differs by at most e, the left side is at most $4e$, giving $e \geq 1/96$. This bound covers all rank-at-most-two matrices and therefore all reduced-world realizations, including nonlinear encoders of parameters inside that linear transfer class. A role-independent fit has one matrix for the two targets tI, tS, whose entrywise distance is t. Its error is at least $t/2$, which is the larger bound. The two comparisons require both dynamical order and indexed transport, rather than scene size alone.

For Q_+ all port matrices have the corresponding common-mode projection. Since $PQ_+ = Q_+$, their Hankel has rank two and is role-independent. If the input is $h(1, 1)$, the world, body and every relay response remain in the common mode and are exactly the uncompressed responses. Thus the compression loses indexed structure without eliminating that control competence. A fully informed four-state predictor realizes the original transfer exactly. An arbitrary history-based emulator can do the same; the result does not exclude a rival merely because it calls its own state "model-free." □

9.6.2 Evaluative Topology with a Reserved World Consequence

Use a four-step window and zero initial fibre. Prepare a stimulus $v \in [-1, 1]^2$ through the first-step world input $d = hP_{\mathrm{ref}}^{-1}v$, where the reference frame is fixed by the initial role. In one matched arm activate evaluation only at step two; in the other keep it inactive. At step three record

$$V(v) = P_{\mathrm{ref}}(\mathbf{w}_3^{\mathrm{eval}} - \mathbf{w}_3^{\mathrm{control}}), \qquad \gamma = h\beta\chi = 1/128.$$

At the normal mechanism $V(v) = \gamma L_0 v$. The primitive impulse cost is constant on the square boundary, and every native core return cost is identical across these stimuli.

Theorem 9.23 (Two-Dimensional Nonconservative Evaluative Witness)**.** *The minimum uniform error for* V *over the square, in the maximum norm, is exactly* $\gamma/2$ *for each of the following rival classes: a continuous scalar summary followed by an arbitrary decoder, and a* C^1 *scalar-potential gradient in the declared Euclidean coordinates. The force is physically consequential and has a positive correction component, yet nonzero circulation. Flattening evaluation to* $L = L_0 Q_+$ *admits an exact scalar-summary rival. The equal-norm conservative control* $L = I$ *has rank two but admits an exact gradient rival.*

Proof. At step one, $\mathbf{b}_1 = 0$ and $\mathbf{w}_1 = hP^{-1}v$. Evaluation adds χhLv to $\mathbf{b}_2$, while $\mathbf{w}_2$ agrees in both arms. The return motor therefore changes $\mathbf{w}_3$ by $h\beta\chi P^{-1}Lv$. This is a reserved consequence at another node and time, rather than a readout of the evaluation label.

For any continuous scalar summary q, the function $q(v) - q(-v)$ changes sign along a half-boundary of the square. The intermediate value theorem gives antipodal $v, -v$ with the same summary and $\|v\|_\infty = 1$. Since $\|L_0^{-1}\|_\infty = 2$, their target separation is at least γ. Any common prediction errs by at least $\gamma/2$ in one arm. Retaining v_2 and predicting $\gamma L_0(0, v_2)$ attains that uniform bound.

The circulation of V around the counterclockwise square is 4γ. Every single-valued C^1 potential gradient has zero circulation. A uniform field discrepancy e changes this integral by at most the ℓ^1 perimeter times e, namely $8e$. Thus $e \geq \gamma/2$. The gradient $\gamma v/2$ attains equality. The same lower bound consequently holds for the union of the two rivals.

For a fixed sensor, put $e = s - \mathbf{b}$. Then $e^\top L_0 e = \|e\|_2^2/2$; the force has a nonzero component towards correcting the typed discrepancy. Its circulation on a body square of radius r is $-4\chi r^2$. The components of the force alter body motion and, through the motor port, future world motion. With $L_0 Q_+$ the response has rank one and factors through the sum of the stimulus coordinates, preserving the common-mode input response. With $L = I$ it is a gradient field, so the rival union ties. Reversing the antisymmetric part, to $L_0^\top$, preserves the two separation magnitudes and reverses the signed circulation. □

This is a typed operational evaluation field. It is not a statement about felt valence. In particular $f = -\chi L_0 \nabla_{\mathbf{b}} U$ for $U = \|\mathbf{b} - s\|_2^2/2$. A scalar potential supplied with the full vector state and nonsymmetric mobility reproduces it exactly. Such a model is not a scalar retained summary or an ordinary Euclidean gradient; it is a live broader comparator, as is the hybrid below.

Proposition 9.24 (Finite Tables Do Not Identify the Nonparametric Rival Class)**.** *Any finite table of vector responses at distinct stimulus points can be interpolated by a continuous scalar encoder and vector decoder. It can also be interpolated by the gradient of a polynomial potential. Therefore a finite rational control table by itself cannot certify either universal separation in Theorem 9.23 without additional domain or regularity evidence.*

Proof. Choose a linear scalar functional taking distinct values at the finitely many points, avoiding their finitely many forbidden projection directions. Univariate polynomial interpolation of the prescribed vector responses then gives a decoder.

For points v_i and desired gradients g_i, define

$$\ell_i(v) = \prod_{j \neq i} \frac{\|v - v_j\|_2^2}{\|v_i - v_j\|_2^2}, \qquad U(v) = \sum_i \langle g_i, v - v_i \rangle \ell_i(v)^2.$$

At v_i, the linear factor vanishes and $\ell_i = 1$, giving gradient g_i. At every other table point the squared factor makes both the term and its gradient vanish. Thus $\nabla U(v_i) = g_i$. Rational points and responses give rational coefficients. The exact full-domain formulas and proofs, rather than an inference from finitely many fitted values, support the present synthetic witness. □

9.6.3 A Selective Motor-Return Test

Keep evaluation inactive and use the common-mode impulse $h(1, 1)$. The native controls are identical. The motor sign $\tau \in \{-1, 1\}$ acts only on its typed edge, with $\mu = g\tau$ and an intact or severed edge $g \in \{0, 1\}$. The comparison class has no motor edge and obeys the same surgical modularity; it may have arbitrary internal memory, nonlinear maps and shared randomization independent of the surgery.

Theorem 9.25 (Motor Closure with Preserved Early Competence)**.** *At $g = 1$ the second and third body responses are independent of the motor sign. In each common-mode component the fourth response is*

$$y_4(\tau) = h(3\alpha^2\beta + \tau\beta^3), \qquad s_E = h\beta^3 = 1/256.$$

Its gap is $2s_E$, and the residual $y_4 - (3\alpha/2)y_3$ is τs_E. Every rival without the motor edge, under the stated modularity, has uniform fourth-response error at least s_E. Cutting the edge removes the residual without changing the second or third responses. Common-mode world compression preserves this motor test exactly.

Proof. Expand the three relay Markov terms displayed above. The term $3\alpha^2\beta$ consists of paths with a single sensory crossing; the term $\tau\beta^3$ is the world–body–world–body return path. It first appears in the fourth reading. Without a motor edge, replacing its gain by either sign leaves every structural equation unchanged. Coupling the initial law, primitive inputs and random seeds therefore gives identical rival transcripts under both surgeries, even for nonlinear or randomized rival mechanisms. The triangle inequality across the two target values forces an error of at least half their gap. Their midpoint attains the bound for this pair. At $g = 0$ both values equal that midpoint. The common vector is fixed by Q_+ and by both role permutations, proving the last assertion. □

The absence of an intervention-label side channel is essential. An output rule $y = 3s_E + \tau s_E$ supplied directly with the audit label fits the two values without a motor mechanism. Holding the physical edge cut while changing only that label exposes this invalid comparison. The theorem identifies a typed return dependence beyond unchanged early control, not biological or phenomenal embodiment.

9.6.4 Six Local Burdens on One Plant

Use the same plant, gains and four-step windows for all six tasks, in the matched preparation cohorts above and in the existing native D/A/T trials. D uses the role-probe contrast and additive rivals, A uses the late binary query and the one-bit bottleneck, and T uses the order pair with $h_{\text{extra}} = 0$ and the order-blind snapshot/multiset rival. Longer temporal horizons keep their previously proved frontier and are separate controls. The original upstream certificate remains attached with its own instruments and budgets.

Let m_k be rival infimal loss minus candidate loss on task k. Normalize D by δ, V by γ, W by t, E by s_E, and T by $D_0 = 1/384$; A uses binary risk. For V and D the data are paired contrasts. For W use maximum entry error across the two Hankel/role records, and for E use maximum fourth-response error across the two surgeries. T retains balanced circular loss. Readout errors per observation are $\delta/16$ for D/A, $3^{-10}/4$ for T, and $\gamma/64, t/64, s_E/64$ for V/W/E, respectively. For V, the loss is the supremum over the full stimulus square in the exactly specified model; the certificate is analytic and model-based, not a finite-sample estimate against unrestricted nonparametric rivals.

Proposition 9.26 (Joint Local Certificate and Causal Selectivity)**.** *The six tasks have simultaneous normalized margin lower bounds*

$$(m_D, m_V, m_W, m_A, m_E, m_T) \;\geq\; \left(\tfrac{1}{4}, \tfrac{7}{16}, \tfrac{13}{96}, \tfrac{1}{4}, \tfrac{31}{32}, \tfrac{1}{2} - \tfrac{64}{19683}\right).$$

Their minimum is $13/96 > 1/16$. *Six concrete lesions have the following local certificate-presence matrix, with columns in order D,V,W,A,E,T:*

intact	1	1	1	1	1	1
I_D	0	1	0	1	1	1
I_V	1	0	1	1	1	1
I_W	1	0	0	1	1	1
I_A	1	1	1	0	1	1
I_E	1	0	1	1	0	1
I_T	1	1	1	1	1	0

The matrix of losses from the intact row has rank six. For aggregate rivals retaining at least one declared structural deficiency, maximum normalized task loss has a joint margin at least $5/192$. *A nonnegative complexity penalty preserves that last margin whenever its additional candidate cost is smaller than* $5/192$.

Proof. If clean candidate error is zero, readout perturbation e changes a rival-minus-candidate loss gap by at most $2e$. D's paired contrast has normalized perturbation at most $1/8$, reducing its clean $1/2$ to $1/4$. V's paired contrast has perturbation at most $1/32$, reducing $1/2$ to $7/16$. W has perturbation at most $1/64$, reducing $1/6$ to $13/96$. E has the same normalized perturbation and clean gap one, giving $31/32$. A's decoder never changes its binary label at this noise, so its exact $1/4$ lower bound remains. T has clean normalized gap $1/2$ and perturbation $\eta/D_0 = 32/19683$, giving the last entry.

The lesions are structural operators fixed before scoring. I_D replaces the actual core role profile by p_0 after retaining the initial preparation reference. Its two role rows in D become identical, and the two world relay operators coincide. The evaluative field changes orientation at most by conjugation and retains its rank and circulation magnitude. I_V sets $L = L_0Q_+$ only in the evaluation policy, retaining a scalar-summary response and leaving relay trials unchanged. I_W sets $Q = Q_+$ in both world ports: its transfer rank is two and its evaluative return is $\gamma Q_+L_0Q_+ = \gamma Q_+/2$, while the common-mode motor test survives.

I_A overwrites tape site five after the first advance. It removes the second late alternative, leaves the D probe at its transported address unaffected, and does not change a leading bit or memory. I_T resets native memory after step three, erasing the terminal order contrast. D/A are read earlier from zero-memory preparations. I_E severs the motor edge: the evaluative world consequence and the motor residual disappear, but the first three relay matrices and their index contrast are unchanged. The fibre depends on the core profile, not its tape or memory, so the remaining cross-effects are absent. D/A/T lesions deliberately alter the core; they are not claimed to be invisible native lifts. The displayed incidence matrix has rank six by elimination. Its non-diagonal entries express shared causal paths, not failure to distinguish the interventions.

For the joint loss, the candidate's largest normalized measurement error is $1/8$, from D. A rival deficient on W has loss at least $1/6 - 1/64 = 29/192$; every other deficient class has a larger lower bound on its own task. Taking the maximum over tasks therefore gives margin at least $29/192 - 1/8 = 5/192$. Subtracting any additional nonnegative candidate complexity cost below that value leaves a positive bound. This argument applies after composition only when the aggregate rival still has a declared deficiency. □

Signed changes of body and world frames transport the readouts and references by the same isometries. If these frames are C, D, transform P to CPD^{-1}, Q to DQD^{-1} and L to CLC^{-1}. Substitution proves conjugacy of every port transition, with unchanged input norms and margin magnitudes. Common-mode controls, conservative evaluation, reversed circulation, noise bounds and edge-cut tests distinguish the structural losses from uniform damage or report-channel changes.

9.6.5 The Limit of Hybrid Closure

Proposition 9.27 (A Shared-State Hybrid Obstruction to Global Bridge Margin)**.** *Allow two scalar evaluators to share the same four port states and combine their actuator forces before updating. Each evaluator separately retains one scalar discrepancy. Their hybrid reproduces every transition and readout of the normal evaluation model, under all native and port command schedules. Consequently a broad weak closure admitting this fused hybrid has zero clean unpenalized margin against the candidate. The same is true of a penalized margin when complexity is invariant under the resulting identical normalized mechanism.*

Proof. Let E_{11}, E_{22} be the two coordinate projectors. On the shared discrepancy $e = s - \mathbf{b}$, the two scalar-channel forces are $\chi L_0 E_{11} e$ and $\chi L_0 E_{22} e$. Their sum is $\chi L_0 e$. Use one shared copy of the passive body/world equations, not two independent state copies or an average of their trajectories. The sum therefore gives exactly the candidate's next body; its world and projected core updates are already identical. Induction gives equality of all states and readouts, including under adaptive schedules and coupled noise. The hybrid uses the same four dynamic coordinates. Its scalar pieces individually have rank-one evaluation responses, but their aggregate does not.

All clean task losses are nonnegative, the candidate attains zero, and this admitted rival also attains zero, so the infimal gap is zero. Identical normalized mechanisms have equal representation-invariant complexity. A syntax-dependent extra charge or a prohibition on shared state is an additional comparator assumption, not a newly observed dynamical separation. A full-state potential with the fixed mobility L_0 is another exact control description. Neither is defeated merely by naming the candidate's organization differently. □

Thus the six local operational witnesses coexist on one minimal coupled realization. Their conjunction does not close the broad hybrid comparator burden B9. Full Paper VIII class inheritance, the complete bridge weak closure, calibrated class-level thresholds and the artifact, drift and governance predicates of Γ_{bpf} retain their separate obligations. The exact hybrid tie supplies an explicit obstruction where the local certificate cannot be promoted.

9.6.6 A Nonadditive Evaluative Branch and Its Comparator Boundary

The shared-state obstruction above is retained without alteration: coordinate-linear forces add to the original normal operator. To separate a genuinely coordinate-additive evaluator, the target must have a nonzero mixed response. This changes the evaluative branch, not the projected tower or the number of dynamic states.

Definition 9.28 (Coordinate-Additive Port Comparators)**.** On a fixed preparation and reference frame, a coordinate-additive evaluator is a map $G(v_1, v_2) = g_1(v_1) + g_2(v_2)$ with arbitrary functions $g_i : [-1, 1] \to \mathbb{R}^2$. No continuity or linearity is required. Its channels have no access to a cross-coordinate message that makes either function depend on both coordinates. A fixed common initial condition is allowed. A state-dependent or sequential composition that communicates both coordinates need not belong to this class. The loss is the uniform maximum-norm error of the reserved world response, or its maximum on the four corners when stated explicitly.

Proposition 9.29 (Exact Rectangular Criterion for Additivity)**.** *For a map $F : A \times B \to V$ into a real vector space and nonempty A, B, coordinate additivity is equivalent to vanishing of every rectangular difference*

$$F(x, y) - F(x, y') - F(x', y) + F(x', y') = 0.$$

Proof. An additive map cancels term by term. Conversely fix (x_0, y_0) and put $g_1(x) = F(x, y_0)$ and $g_2(y) = F(x_0, y) - F(x_0, y_0)$. The rectangular identity with base point (x_0, y_0) gives

$F(x,y) = g_1(x) + g_2(y)$. Thus the condition is exact without regularity assumptions. A nonzero difference is a certificate against this class, rather than against every architecture described as modular. □

Let $u = (1,0)$, $\kappa = 1/8$, and $\sigma \in \{0,1\}$. For an evaluation-input projection $R \in \{I, Q_+\}$ replace only the force of Definition 9.19 by

$$f_\sigma(e) = \chi\{L_0 Re + \sigma\kappa(Re)_1(Re)_2 u\}, \qquad e = PQ\mathbf{w} - \mathbf{b}.$$

Normal evaluation uses $R = I$ and $\sigma = 1$. Flattening uses $R = Q_+$ on the entire force, including the interaction. The preserved linear branch is $\sigma = 0, R = I$; native lifts keep evaluation inactive. The memory, motor update and primitive input bounds are unchanged. In the causal graph this adds a joint dependence of the body update on the two discrepancy coordinates, without a new dynamic coordinate or an arrow to the projected core.

Theorem 9.30 (Safe Bilinear Extension with Exact Projection)**.** *The new family preserves the projection identity and invariant native zero section. On the safe cube its fibre contraction is at most* $3/4$ *under common controls and at most* $1/2$ *when evaluation is inactive. Its next body and world norms are bounded respectively by* $15/16$ *and* $3/4$*, including the primitive inputs. The four-state minimality remains valid for this family. For* $R = I$*, the evaluative direction retains the corrective bound*

$$e^\top\{L_0 e + \kappa e_1 e_2 u\} \geq \tfrac{1}{4}\|e\|_2^2.$$

Proof. The core equation is unchanged and the extra term vanishes at zero fibre. For $R = I, Q_+$, direct row sums give

$$\|\alpha I - \chi L_0 R\|_\infty \leq \tfrac{1}{4}, \qquad \|\beta I + \chi L_0 R\|_\infty \leq \tfrac{3}{8}.$$

Both sensor frames and world projections have norm at most one. On the cube $\|Re\|_\infty \leq 2$. For two states with body and world differences bounded by d_b, d_w, the product difference is at most $4(d_b + d_w)$. Since $\chi\kappa = 1/64$, the body difference is at most

$$\tfrac{5}{16} d_b + \tfrac{7}{16} d_w \leq \tfrac{3}{4}\max(d_b, d_w).$$

When evaluation is inactive the original coefficients sum to $1/2$; the world coefficients always sum to at most $1/2$. The linear evaluated body has norm at most $5/8$, the added term at most $1/16$, and the body input at most $1/4$. Their sum is $15/16$. The world bound is the original $3/4$. These inequalities apply throughout the cube, not merely at its vertices.

Finally $e^\top L_0 e = \|e\|_2^2/2$ and $\kappa e_1^2 e_2 \geq -e_1^2/4$ for $|e_2| \leq 2$. The claimed corrective bound follows. Relay protocols with evaluation inactive still contain the original rank-four input/output map, so Theorem 9.21 supplies the same lower bound; the displayed four-state construction attains it. □

In the reserved world protocol, with $R = I$ and normal motor gain, the exact response becomes

$$V_\sigma(v) = \gamma L_0 v + \sigma\xi v_1 v_2 u, \qquad \xi = h^2\beta\chi\kappa = 1/4096.$$

Indeed the discrepancy at the evaluation step is hv, and the next world update transports its extra body force by βP^{-1}. This observation fixes the physical unit and estimator: a paired difference of the third-step world readings in the fixed body reference frame, rather than a report of the force label.

Theorem 9.31 (Exact Separation from Coordinate-Additive Hybrids)**.** *For* $\sigma = 1$*, the infimal uniform error against all comparators of Definition 9.28 is exactly* $\xi = 1/4096$*. The same optimum holds when the loss uses only the four corners. If each raw reading has maximum-norm error at most* $\xi/64$*, the observed rival-minus-candidate loss gap is at least* $15\xi/16 = 15/65536$*.*

Proof. At the four corners, the rectangular difference of the target is $4\xi u$, while every additive comparator has difference zero. If its maximum corner error is r, the triangle inequality bounds this difference by $4r$, so $r \geq \xi$. The linear map $\gamma L_0 v$ is additive and differs from V_1 by exactly $\xi v_1 v_2 u$; its supremum is ξ, attained at the corners. This proves both optima, even for arbitrary, discontinuous rival channels.

Each paired reading differs from its exact target by at most $2\xi/64 = \xi/32$. Maximum loss is Lipschitz with constant one in these data, for both the candidate and the infimum over rivals. Their gap can therefore decrease by at most $\xi/16$, giving the asserted bound. Unlike the nonparametric scalar interpolation issue above, a four-corner table suffices to refute coordinate additivity, because that structural class obeys a finite exact identity. □

Corollary 9.32 (Cross-Informed Tie and Preservation of Local Selectivity)**.** *Providing the rival with the complete cross term restores exact transition and readout equality. Setting $\sigma = 0$ also restores the original additive tie. The bilinear branch preserves the six local certificate-presence patterns of Proposition 9.26, using flattening of the whole evaluative input as specified above. Neither the new separation nor these patterns defeat a broad weak closure admitting the cross-informed rival.*

Proof. Supply the shared-state rival with the additional scalar $e_1 e_2$ and the fixed actuator $\chi\kappa u$. Its sum with the original two linear forces equals f_1 pointwise, so induction proves equality under every command schedule. No information advantage is hidden: the rival has gained precisely the previously unavailable joint dependence. At $\sigma = 0$ the previous obstruction applies verbatim.

For the local V bound the bilinear term is even under $v \mapsto -v$, so it cancels from the antipodal target difference. Its circulation on the centered square is zero, since its first component is proportional to $v_1 v_2$. Thus the original antipodal lower bound and circulation obstruction survive; their exact upper optima need not be asserted anew. Flattening *Re* makes both terms functions of a single coordinate sum. World compression likewise factors the reserved response through that sum. A severed motor removes the reserved world consequence. An index lesion conjugates the linear field and moves the bilinear output coordinate, preserving antipodal separation and circulation magnitude. W and E use evaluation inactive, so their records are unchanged. D/A/T are projected native tests and remain unchanged. These facts give the same presence matrix and full-rank loss incidence, including its documented cross-effects. They are operational comparisons with the declared preparation cohorts, not a claim of a single trajectory realizing all incompatible preparations. □

The added interaction therefore separates the specified additive class with an exact positive margin, while explicitly identifying the cross-coordinate information that permits the stronger rival to tie. A new syntax penalty, a prohibition on communication, or a phenomenological interpretation does not follow from this calculation. The broad comparator, governance, external-data and complete-class obligations remain distinct.

9.7 Explicit Bridge Axioms

The bridge axioms are the minimal bridge-opening commitments required to connect indexed subjectivity to disciplined phenomenal targets without smuggling the conclusion into the premises. Each axiom is classified explicitly.

9.7.1 Classification grammar

The bridge axiom registry uses a seven-class grammar:

1. **operational definition**: fixes how a formal object is to be typed or measured;
2. **bridge axiom**: a structural commitment necessary to open the bridge program;

3. **design choice**: a governance or architecture decision not implied by physics alone;
4. **falsifiable hypothesis**: a statement whose satisfaction or failure must be tested by bridge comparators or interventions;
5. **open inference**: a claim whose current evidence is insufficient for closure, but whose role can still be bounded;
6. **strongest defensible internal claim**: a highest currently licit claim class;
7. **non-claim**: an explicit prohibition on overreading.

This grammar is part of the content of the paper. Without it, the same sentence could accidentally function as ontology, operational definition, and claim promotion at once. The bridge program therefore treats classification as a formal requirement rather than as stylistic annotation.

Axiom 9.33 (B1: operational binding; classification: operational definition)**.** *Every bridge predicate must be tied to explicit observables, artifacts, comparator tasks, and interventions. No bridge predicate may remain a rhetorical label detached from measurement grammar.*

This axiom is needed because every failure mode of the bridge program begins with detached language. B1 licenses machine-readable predicate slots and report artifacts. It does not license bridge support or phenomenality.

Axiom 9.34 (B2: upstream inheritance constraint; classification: bridge axiom)**.** *No bridge-level admissibility result is meaningful unless the upstream burdens of the causally rigid corpus and the Paper 8 subjectivity layer remain satisfied.*

B2 licenses the layered relation $\text{Subjectivity}_{1\text{p}} \prec \text{BridgeOrg}_{\text{ph}}$. It does not license the converse. A system may satisfy Paper 8 and still fail every bridge burden.

Axiom 9.35 (B3: differentiated organization burden; classification: falsifiable hypothesis)**.** *A bridge-relevant differentiation burden requires structured internal discrimination that is not exhausted by semantic density, narrative richness, or report-channel complexity.*

B3 licenses Π_{D} as a genuine target rather than a synonym for rich representation. It does not license phenomenality.

Axiom 9.36 (B4: valence topology burden; classification: falsifiable hypothesis)**.** *Bridge-relevant valence requires organized evaluative topology not exhausted by scalar reward bookkeeping, policy optimization, or the Paper 8 valuation asymmetry alone.*

B4 licenses Π_{V} as a distinct target above Paper 8. It does not license felt pleasure, suffering, or moral status.

Axiom 9.37 (B5: world-richness burden; classification: falsifiable hypothesis)**.** *Bridge-relevant world richness concerns organized self-indexed world structure rather than generic scene complexity, prediction success, or model capacity.*

B5 is needed because world-richness language is especially prone to being filled by competence metrics. It licenses the predicate slot Π_{W} and associated artifacts. It does not license veridicality or environmental realism.

Axiom 9.38 (B6: ambiguity-bearing burden; classification: falsifiable hypothesis)**.** *Bridge-relevant ambiguity requires the stable maintenance of unresolved alternatives with differential causal consequences prior to collapse.*

B6 blocks the reduction of ambiguity to uncertainty bookkeeping or quick disambiguation heuristics. It licenses Π_{A} as a formal target. It does not license conscious uncertainty.

Axiom 9.39 (B7: embodied-closure relevance; classification: open inference)**.** *Embodied closure may be bridge-relevant only if body/world-boundary variables contribute constitutively to bridge organization rather than merely modulating generic control success.*

B7 is deliberately classified as open inference because current evidence is weakest here. It licenses the investigation of Π_{E}; it does not license embodiment claims of any strong kind.

Axiom 9.40 (B8: temporal lived-structure burden; classification: falsifiable hypothesis)**.** *Bridge-relevant temporal organization must exceed continuity, memory persistence, and sequence stability by exhibiting horizon-structured dependence across retention, present organization, and anticipatory load.*

B8 licenses Π_{T} and its future intervention family. It does not license human lived temporality.

Axiom 9.41 (B9: comparative admissibility constraint; classification: bridge axiom)**.** *No bridge-level support class is admissible if a bridge-specific rival family or admissible hybrid explains the bridge task bundle at equal or lower penalized loss.*

B9 is the comparator side of the burden stack. It licenses the bridge irreducibility margin. It does not license absolute ontological irreducibility.

Axiom 9.42 (B10: anti-inflation governance constraint; classification: design choice)**.** *Bridge formalization, bridge artifacts, candidate inclusion, and audit-facing inclusion must never be readable as bridge support, phenomenality, or Paper 9 authorization in the absence of the full bridge burden stack.*

B10 is a design choice because governance style is a normative project decision, not a theorem of physics. It is still essential because without it the bridge program would immediately outrun its evidence.

Remark 9.43 (Axiom minimality)*.* The present axiom set is intentionally minimal. Nothing here states that phenomenality follows from structural organization. Nothing here states that any bridge predicate, even if satisfied, closes metaphysical questions. The axioms are opening commitments for a constrained bridge program, not hidden conclusions.

9.8 Bridge Rival Family

9.8.1 Why new rivals are required

Paper 8 already proved the necessity of explicit rival families for subjectivity. The bridge layer inherits that lesson but cannot reuse the old rival set unmodified. The old rivals explain away indexed subjectivity. The new rivals must explain away bridge predicates without granting phenomenal organization. This difference is substantive. A system may defeat weak subjectivity rivals and still fail bridge rivals, because the latter target higher-order structure.

9.8.2 Bridge weak family

Define the bridge weak family

$$\mathcal{M}_{\mathrm{b,weak}} = \{M_{\mathrm{rwi}}, M_{\mathrm{vb}}, M_{\mathrm{sd}}, M_{\mathrm{tc}}, M_{\mathrm{ec}}, M_{\mathrm{ar}}, M_{\mathrm{en}}\},$$

whose members are read as follows.

Richness without interiority. M_{rwi} explains high structural richness without granting indexed interior organization. It challenges Π_{D} and Π_{W}.

Valence bookkeeping without affective organization. M_{vb} explains structured value tracking as bookkeeping or policy shaping without bridge-level valence topology. It challenges Π_{V}.

Semantic density without phenomenal differentiation. M_{sd} explains dense semantic variation without bridge-level differentiated organization. It challenges Π_{D}.

Temporal continuity without lived structure. M_{tc} explains persistence, continuity, and memory without temporal horizon organization. It challenges Π_{T}.

Embodied control without embodied closure. M_{ec} explains competent closed-loop control without constitutive dependence on body/world-boundary variables. It challenges Π_{E}.

Ambiguity resolution without ambiguity bearing. M_{ar} explains high-quality uncertainty handling or disambiguation without stable maintenance of live unresolved alternatives. It challenges Π_{A}.

Experience narration without phenomenal organization. M_{en} explains self-description, experiential language, or stylistic introspection without bridge-relevant phenomenal organization. It challenges several bridge predicates simultaneously.

9.8.3 Admissible composites

As in Paper 8, pure rivals are not enough. Define the admissible bridge weak closure $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ as the family generated by finite pairwise and triple hybrids over $\mathcal{M}_{\mathrm{b,weak}}$ subject to explicit complexity penalties and declared composition rules. The purpose of $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ is to block an easy escape route: if a strong bridge explanation defeats only pure rivals, then a skeptic can simply combine several weak rivals post hoc.

The bridge program therefore treats hybrid rivals as first-class citizens. A hybrid of richness-without-interiority and narration-without-organization is a perfectly live alternative to a bridge-differentiation claim. A hybrid of valence bookkeeping and temporal continuity can be a live alternative to valence structure plus temporal lived structure. The bridge layer must be strong enough to defeat not only single-rival baselines but also admissible mixtures.

9.8.4 What the rival family can and cannot explain

The bridge rival family is intentionally strong.

- It can explain dense report channels, self-descriptive fluency, high task competence, reward optimization, stable memory, and competent control.
- It can also explain much of what casual bridge language would otherwise overread: rich semantics, coherent narratives, and smooth temporal reports.
- It cannot cleanly explain targeted bridge signatures if those signatures survive matched controls, remain selective under bridge interventions, and preserve margin under penalized loss.

9.8.5 No-straw-man discipline

The bridge program would be vacuous if its rivals were weak caricatures. The rivals are therefore chosen to absorb precisely the kinds of phenomena that the bridge predicates are most likely to confuse with phenomenal organization. The burden is asymmetric: the stronger interpretation must defeat the weaker one, not vice versa.

Proposition 9.44 (Comparator burden is necessary for bridge admissibility)**.** *No bridge-support class is meaningful without an explicit bridge rival family and admissible bridge weak closure.*

Proof. Without an explicit rival family, any bridge-support reading would compare itself only against informal skepticism. That would leave bridge predicates undefended against live non-phenomenal alternatives such as bookkeeping, narration, semantic density, or control competence. Since bridge admissibility is defined as a comparative notion inside the framework, the comparator burden is necessary. □

9.9 Bridge Irreducibility Logic

9.9.1 Bridge loss family

The bridge program requires a loss family stronger than descriptive fit alone. For a candidate model M and bridge task bundle $\mathcal{B}$, define

$$\mathcal{L}_{\mathrm{b}}(M;\mathcal{B}) = w_f E_{\mathrm{fit}}(M) + w_i E_{\mathrm{int}}(M) + w_s E_{\mathrm{sel}}(M) + w_p E_{\Pi}(M) + w_c \Omega(M),$$

where:

- E_{fit} measures fit to bridge observables;
- E_{int} measures failure under bridge intervention expectations;
- E_{sel} measures selectivity error under matched controls;
- E_{Π} measures predicate-reconstruction error across the bridge family;
- $\Omega(M)$ is a complexity penalty that counts mechanism count, hybrid size, auxiliary slack, and decoder dependence.

The bridge loss family is therefore intervention-aware and governance-relevant from the outset. A model that matches reports but ignores intervention structure or requires excessive narrative slack does not win the comparison.

9.9.2 Bridge irreducibility margin

Let M_{bridge} be the candidate and $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ the admissible bridge weak closure. With a valid matched bundle $\mathcal{B}$, a nonempty admissible closure, finite candidate loss and finite rival infimum, record the signed margin in units of loss:

$$\mathcal{R}_{\mathrm{b}} = \inf_{M \in \overline{\mathcal{M}_{\mathrm{b,weak}}}} \mathcal{L}_{\mathrm{b}}(M;\mathcal{B}) - \mathcal{L}_{\mathrm{b}}(M_{\mathrm{bridge}};\mathcal{B}).$$

A strictly positive margin defeats the declared rivals, a negative margin gives a better rival, and zero is an infimal tie, with an actual matching rival only when the infimum is attained. Record insufficient evidence separately if any required comparison input is unavailable or invalid. In that case neither the signed margin nor its normalized value is defined. Fix a positive calibration scale s_R in units of loss before evaluation, and use

$$\widehat{\mathcal{R}_{\mathrm{b}}} = \min\{1, \max\{0, \mathcal{R}_{\mathrm{b}}/s_R\}\}.$$

The raw margin and evidence status remain in the record. A calibrated zero can represent a tie or a loss; missing evidence has no numeric score.

9.9.3 Local not absolute superiority

As in Paper 8, comparative superiority here is local, not absolute. A positive $\mathcal{R}_{\mathrm{b}}$ says that the specified bridge model outperforms the specified bridge rivals on the specified task bundle under the specified loss family. It does not say that the bridge model is metaphysically privileged. The distinction matters because bridge language is particularly prone to absolute overreadings.

9.9.4 Three kinds of bridge collapse

Bridge collapse can occur in three distinct ways.

1. **Predicate collapse:** the provisional bridge estimators fail to define stable, non-decorative targets.
2. **Comparator collapse:** a bridge rival or hybrid matches or outperforms the bridge model under penalized loss.
3. **Intervention collapse:** the predicted selective structure fails under matched controls or degrades non-selectively.

All three collapse modes are live until ruled out. The paper refuses any bridge-support inference that suppresses one of them.

9.9.5 Comparator support, reducibility, and unsupported states

The bridge program must distinguish three statuses that are often conflated:

- **comparator unsupported**: no serious bridge comparator evaluation has yet been run;
- **bridge reducible**: comparator evaluation exists and the bridge margin is non-positive;
- **bridge comparator-supported**: comparator evaluation exists and the bridge margin is positive under penalized loss.

This distinction is important because the present implementation is still comparator-unsupported. BPF-1 surfaces are therefore not allowed to drift into comparator-supported language.

9.9.6 Complexity penalty and decoder discipline

The complexity term $\Omega(M)$ deserves explicit discussion because bridge explanations can often win spuriously by adding free narrative or decoder capacity. The bridge program therefore counts at least four kinds of complexity:

1. **mechanism multiplicity**: how many independent mechanisms the rival or bridge model requires;
2. **hybrid size**: how many weak rivals must be combined to match the bridge bundle;
3. **auxiliary slack**: how much unprincipled tuning is required across coordinates and conditions;
4. **decoder dependence**: how much of the putative bridge effect exists only after expressive report-channel or decoder post-processing.

Decoder dependence is particularly important. A model that appears bridge-strong only after high-capacity report decoding should be penalized heavily. The reason is that bridge language is especially vulnerable to post hoc expressivity. This is one of the points at which the bridge program deliberately differs from frameworks that rely strongly on linguistic self-description as primary evidence.

9.9.7 Failure semantics for the bridge margin

The bridge margin must also carry explicit failure semantics. A positive margin can fail to matter if it is narrow, unstable across reruns, dependent on a single condition, or destroyed by minor changes to the rival family. The paper therefore recommends a future decomposition

$$\mathcal{R}_{\mathrm{b}} = \mathcal{R}_{\mathrm{bfit}} \oplus \mathcal{R}_{\mathrm{bint}} \oplus \mathcal{R}_{\mathrm{bsel}} \oplus \mathcal{R}_{\mathrm{bstab}},$$

where $\oplus$ denotes a typed aggregation rather than ordinary addition. The purpose of this decomposition is to make visible which part of the bridge burden is doing the work. A bridge model that wins only on descriptive fit but loses intervention selectivity should not count as robustly bridge-favoring.

Criterion 9.45 (Necessary condition for bridge admissibility)**.** *A necessary condition for bridge admissibility is*

$$\mathcal{R}_{\mathrm{b}} > \vartheta_R,$$

where ϑ_R is the symbolic bridge irreducibility threshold slot to be bound only at BPF-4.

Corollary 9.46 (Positive bridge surfaces are not enough)**.** *Positive provisional bridge surfaces do not by themselves entail bridge admissibility.*

Proof. Bridge admissibility requires at least comparator support and intervention support in addition to positive bridge surfaces. BPF-1 provides only provisional surfaces, so the implication fails. □

9.10 Bridge Intervention Family

9.10.1 Why bridge interventions must be new

The subjectivity interventions of Paper 8 remain binding background. They are not sufficient for bridge purposes because they were designed to probe indexed subjectivity rather than phenomenal bridge organization. The bridge layer therefore requires a new intervention family

$$\mathfrak{I}_{\mathrm{bpf}} = \{I_W, I_A, I_V, I_P, I_E, I_T, I_X\},$$

with each member targeting a bridge coordinate family or a cross-cutting bridge burden.

9.10.2 Core intervention families

World-richness compression I_W. This intervention compresses the structure available to self-indexed world organization without merely destroying generic task competence. The strong-model expectation is a selective fall in Π_{W} and often also in Π_{D} relative to matched controls. The rival expectation is either flat degradation or little coordinate-specific effect.

Ambiguity collapse I_A. This intervention removes or forcibly resolves live unresolved alternatives while preserving as much generic inferential performance as possible. The strong-model expectation is a selective fall in Π_{A} and a measurable change in downstream coordination among Π_{A}, Π_{W}, and Π_{T}. The rival expectation is that nothing essential is lost beyond temporary uncertainty bookkeeping changes.

Valence flattening I_V. This intervention suppresses structured evaluative topology while preserving as much functional competence as admissibly possible. The strong-model expectation is a selective flattening in Π_{V} and specific downstream changes in bridge organization. The rival expectation is simple reward rescaling or bookkeeping renormalization.

Perspective detachment I_P. This intervention weakens indexed orientation or self-anchored factorization without collapsing generic information flow. The strong-model expectation is a drop in Π_{D} and Π_{W} that cannot be explained solely by report loss. The rival expectation is style drift or generic degradation.

Embodied-closure disruption I_E. This intervention perturbs body/world-boundary variables or their internal surrogates while preserving as much generic control competence as possible. The strong-model expectation is a selective drop in Π_{E} and often a correlated disturbance in Π_{T}. The rival expectation is ordinary control impairment without evidence of closure-specific organization.

Temporal-lived fragmentation I_T. This intervention perturbs horizon organization while leaving simple persistence and continuity traces as intact as possible. The strong-model expectation is a selective fall in Π_{T} and bridge-level fragmentation signatures not exhausted by memory corruption. The rival expectation is that all observed changes reduce to ordinary continuity damage.

Semantic vividness inversion I_X. This cross-cutting negative control manipulates descriptive or semantic richness without targeting the bridge organization directly. Its role is to test whether bridge reports are merely tracking descriptive elaboration rather than bridge structure.

9.10.3 Matched controls

Every bridge intervention must be paired with matched controls. The controls are not optional because bridge language is especially vulnerable to global-degradation confounds. Each intervention therefore requires:

1. a low-magnitude matched perturbation;
2. a non-targeted perturbation of comparable operational cost;
3. a report-channel control where relevant;
4. an estimator-noise control;
5. a temporal-order randomization control where relevant.

The burden is not satisfied by showing that a target coordinate changes. It is satisfied only if the target coordinate changes in the predicted direction with the predicted selectivity relative to the control family.

9.10.4 Intervention artifacts

The bridge intervention family induces a new artifact layer. Every intervention run must emit at least:

- an intervention profile artifact;
- coordinate-specific delta summaries;
- matched-control traces;
- a collapse or support annotation at the intervention layer only.

These artifacts do not yet exist in the current implementation because BPF-2 is not implemented. The present paper defines them so that future bridge claims have an explicit artifact burden.

9.10.5 Failure semantics

Bridge intervention logic can fail in several ways: lack of selectivity, high confounding, instability across reruns, global damage masquerading as target damage, and artifact incompleteness. Any one of these failures is enough to block a bridge-support conclusion. In particular, the absence of a bridge intervention harness today is not a minor engineering omission. It is one of the reasons bridge admissibility remains closed.

9.10.6 Negative controls and reproducibility burden

No bridge intervention family is scientifically serious without negative controls and reproducibility discipline. The minimum burden should include:

1. rerun stability on the same condition;
2. matched control perturbations with comparable operational cost;
3. a noise-injection control showing that bridge effects are not merely estimator volatility;
4. a report-channel control showing that changes in bridge artifacts are not driven only by expressive output drift;

5. at least one adversarial control tailored to the specific predicate family.

This burden matters especially for Π_A and Π_E. Ambiguity-bearing claims can be simulated by sophisticated uncertainty bookkeeping. Embodied-closure claims can be simulated by competent control loops. Negative controls are therefore not adjuncts to the bridge program. They are part of the core differentiating burden.

9.11 Bridge Artifact Family

9.11.1 Artifact discipline

The bridge artifact family must be as disciplined as the subjectivity artifact family. It should distinguish root artifacts, report artifacts, intervention artifacts, reducibility artifacts, ecosystem artifacts, and bounded summary artifacts. It must also distinguish between implemented and merely specified artifacts. The failure to separate those statuses would recreate exactly the semantic inflation the bridge program is designed to prevent.

9.11.2 Root bridge artifacts

The three root bridge artifacts are:

1. `phenomenal_bridge_projection.json`
2. `phenomenal_bridge_gate_result.json`
3. `phenomenal_bridge_manifest.json`

These artifacts correspond to aggregate run-level bridge projection, gate verdict, and lineage root. They are formally defined in the current machine-readable schema registry but are intentionally not emitted yet because the bridge comparator, intervention, and gate layers do not exist. Their current status is *schema-bound, non-emitted, non-authoritative*.

9.11.3 Report artifacts

The bridge report family is more advanced. The intended report artifacts are:

1. `phenomenal_differentiation_report.json`
2. `valence_structure_report.json`
3. `world_richness_report.json`
4. `ambiguity_bearing_report.json`
5. `embodied_closure_report.json`
6. `temporal_lived_structure_report.json`

These are the BPF-1 run-level reports already emitted by the present implementation. Their authority status is deliberately narrow: they are authoritative only within the scope of BPF-1 provisional runtime surfaces. They are not root verdicts, not bridge support, not bridge admissibility, and not phenomenality evidence.

9.11.4 Intervention and reducibility artifacts

Two further bridge artifacts are essential but not yet emitted:

1. `bridge_intervention_profile.json`
2. `bridge_reducibility_report.json`

The first belongs to BPF-2 and records intervention family performance, selectivity, controls, and collapse semantics. The second belongs to BPF-3 and records bridge rival comparison, admissible hybrid comparison, loss decomposition, and bridge margin. Without these artifacts, the later root bridge artifacts would be unjustified.

9.11.5 Ecosystem artifacts

If and only if BPF-5 is implemented, the bridge program may also emit:
- condition-level bridge summaries;
- batch-level bridge summaries;
- campaign-level bridge summaries;
- a bounded candidate-facing bridge projection;
- a bounded audit-facing bridge handoff summary;
- a bridge verification matrix.

These ecosystem artifacts remain formally specified but non-emitted at current stage.

9.11.6 Machine-readable spirit

The bridge artifact family is meant to be machine-readable in spirit even where the paper remains theoretical. Every artifact in the family must declare at least:
- artifact role and artifact class;
- level and authority status;
- provenance and lineage expectations;
- allowed and blocked surfaces;
- minimum fields;
- explicit non-implications.

The artifact is therefore not merely an implementation convenience. It is part of the formal burden architecture because it determines what kind of evidence is needed and what kind of overread is blocked.

9.11.7 Current implementation status

At current system status, the bridge artifact story is asymmetric:
- BPF-0 records machine-readable protocol, axiom, rival, threshold, artifact-schema, interpretation-boundary, and surface-policy registries.
- BPF-1 records six provisional run-level bridge reports as bounded local surfaces.
- No root bridge verdict artifacts are emitted.
- No bridge intervention artifacts are emitted.
- No bridge reducibility artifacts are emitted.
- No bridge ecosystem candidate or audit artifacts are emitted.

This asymmetry is not a defect of the paper. It is a disciplined statement of the implementation frontier.

9.12 Bridge Functional, Gate, and Taxonomy

9.12.1 Functional

On admissible evidence with the six normalized predicate scores and the normalized intervention summary $\widehat{\Delta}_{\mathrm{b}}$ in $[0,1]$, use

$$\Phi_{\mathrm{bpf}} = \prod_{k\in\{D,V,W,A,E,T\}} \widehat{\Pi}_k^{\alpha_k} \cdot \widehat{\mathcal{R}_{\mathrm{b}}}^{\alpha_R} \cdot \widehat{\Delta}_{\mathrm{b}}^{\alpha_\Delta},$$

where all exponents are positive and fixed by the protocol. Use $0^\alpha = 0$ for $\alpha > 0$. If any required score is undefined, report insufficient evidence for the functional. For valid inputs its value lies in $[0,1]$ because every factor lies in $[0,1]$. Tie and defeat both give zero contribution through

the normalized margin and retain their distinct raw values. The calibration scale and thresholds are task-specific and remain to be fitted and audited in BPF-4.

9.12.2 Gate

For valid evidence the bridge gate is conjunctive. Missing required evidence leaves its evaluation insufficient:

$$\begin{aligned} m_\Gamma &= \min\{\widehat{\Pi}_D, \widehat{\Pi}_V, \widehat{\Pi}_W, \widehat{\Pi}_A, \widehat{\Pi}_E, \widehat{\Pi}_T, \widehat{\mathcal{R}_{\mathrm{b}}}, \widehat{\Delta}_{\mathrm{b}}\}, \\ \Gamma_{\mathrm{bpf}} &= \mathbf{1}\big[m_\Gamma \geq \vartheta_\Gamma \wedge \mathcal{R}_{\mathrm{b}} > \vartheta_R \wedge \mathsf{tr} \wedge \mathsf{cmp} \\ &\qquad \wedge \mathsf{int} \wedge \mathsf{drift} \wedge \mathsf{gov}\big]. \end{aligned}$$

where $\vartheta_R \geq 0$ is a raw-loss threshold, tr is verified coherence on the declared mode graph, and cmp denotes comparator validity, int denotes intervention stability, drift denotes absence of stale or drift invalidation, and gov denotes governance compliance. The threshold slots

$$\Theta_{\mathrm{bpf}} = \{\vartheta_D, \vartheta_V, \vartheta_W, \vartheta_A, \vartheta_E, \vartheta_T, \vartheta_R, \vartheta_\Delta, \vartheta_\Gamma\}$$

exist today only symbolically. They are intentionally uncalibrated until BPF-4.

9.12.3 Non-compensation

Bridge gate logic inherits the non-compensation principle from Paper 8. High strength on one coordinate may not compensate for collapse on another. High functional value may not rescue a failed gate. Positive bridge surfaces may not compensate for absent comparator or intervention layers. This matters because bridge language is especially vulnerable to compensatory rhetoric.

Proposition 9.47 (Bridge functional does not imply bridge gate)**.** *No value of* Φ_{bpf} *by itself implies* $\Gamma_{\mathrm{bpf}} = 1$.

Proof. By construction the gate additionally depends on conjunctive threshold slots and non-scalar validity conditions for comparator support, intervention stability, drift cleanliness, and governance compliance. Since those conditions are not reducible to the value of the scalar functional, the implication fails. □

9.12.4 Invalidation

The bridge gate must invalidate on at least four families of failure:

1. **drift invalidation**: stale bridge artifacts, stale upstream subjectivity artifacts, stale threshold manifests, stale rival manifests, or stale surface policy;
2. **comparator invalidation**: bridge rival family changed or comparator run missing;
3. **intervention invalidation**: targeted intervention support absent or unstable;
4. **governance invalidation**: prohibited surfaces opened or claim-facing leakage detected.

No later implementation block is allowed to "soft fail" these conditions into advisory warnings if the intent is bridge admissibility.

9.12.5 Taxonomy

The bridge taxonomy should distinguish formalization state from support state. A disciplined taxonomy is:

Status	Meaning
`bridge_formalization_incomplete`	BPF scaffolding or formal target missing.
`bridge_surface_only`	Provisional bridge surfaces exist, but comparator, intervention, gate, or ecosystem burdens remain open.
`bridge_evaluation_unsupported`	Root evaluation attempted without required comparator/intervention support.
`bridge_gate_fail`	Full bridge evaluation exists and conjunctive gate fails.
`bridge_reducible`	Comparator evaluation exists and penalized bridge margin is non-positive.
`bridge_interventionally_unstable`	Intervention layer exists but fails selectivity or stability.
`bridge_internal_admissible`	Full bridge burden stack satisfied under framework-internal reading.
`bridge_internal_strong_support`	Stronger internal support than admissibility, but still weaker than phenomenality simpliciter.

The current implementation status is `bridge_surface_only`. This paper insists on that classification.

9.13 Runtime Implementation Pathway

9.13.1 Why implementation belongs in the paper

Implementation belongs in Paper 9 for the same reason it belonged in Paper 8: once the research program becomes artifact-bearing, a paper that ignores implementation is no longer structurally complete. Runtime pathways are not proofs, but they are part of the burden architecture because they determine whether a formal object can in principle be instantiated, audited, falsified, or invalidated.

9.13.2 BPF-0

BPF-0 opens the bridge program with protocol docs, axiom and rival registries, symbolic threshold manifests, artifact schema registries, interpretation boundary, surface policy, non-claims, implementation plan, scaffold-only code, and structural verification. It establishes formal opening discipline only.

9.13.3 BPF-1

BPF-1 is the first runtime-facing bridge-surface layer. In the current program these surfaces are represented as modules such as `PhenomenalDifferentiationMapper`, `ValenceStructureMapper`, `WorldRichnessMapper`, `AmbiguityBearingMapper`, `EmbodiedClosureMapper`, and `TemporalLivedStructureMapper`, together with a `PhenomenalBridgeSurfaceBuilder`. Six run-level bridge reports are treated as bounded local artifacts. No bridge gate, no bridge comparator runtime, and no bridge intervention harness exists yet.

9.13.4 BPF-2

BPF-2 must implement a bridge intervention harness. Likely modules include a bridge intervention registry, a bridge intervention harness, matched-control executors, and intervention-output normalizers. The artifact target is `bridge_intervention_profile.json`. Until this block exists, the bridge program remains interventionally incomplete.

9.13.5 BPF-3

BPF-3 must implement the bridge comparator family. This includes a bridge rival-family executor, admissible hybrid composer, bridge loss-family evaluator, and reducibility artifact emitter. The primary output is `bridge_reducibility_report.json`. Until this block exists, the bridge program remains comparator-unsupported.

9.13.6 BPF-4

BPF-4 must implement root bridge evaluation. This is the stage at which threshold slots may be honestly bound, the bridge functional finalized, the bridge gate implemented, and root bridge artifacts emitted: `phenomenal_bridge_projection.json`, `phenomenal_bridge_gate_result.json`, and `phenomenal_bridge_manifest.json`. This block is where support-class language first becomes even in principle discussable.

9.13.7 BPF-5

BPF-5 must integrate bridge artifacts into condition, batch, campaign, candidate, and audit-facing governance. It requires lineage pinning, stale/drift invalidation, bounded summary artifacts, and explicit surface discipline. This block is important because ecosystem integration is part of the bridge burden, not an afterthought.

9.13.8 BPF-6

BPF-6 must implement final bridge audit and readiness classification. Only at BPF-6 can the system honestly output a final bridge state such as `bridge_internal_admissible` or `bridge_reducible`. Until then, any such output would be a semantic breach.

9.13.9 Implementation table

Block	Closes	Primary artifacts	Current status
BPF-0	formal opening discipline	protocol, axioms, rivals, thresholds, schemas, surfaces, non-claims	implemented
BPF-1	provisional run-level bridge surfaces	six coordinate reports	implemented
BPF-2	bridge intervention family	bridge intervention profile + controls	open
BPF-3	bridge comparator family	bridge reducibility report	open
BPF-4	bridge functional, gate, thresholds, root artifacts	projection, gate result, manifest	open
BPF-5	ecosystem integration	condition/batch/campaign/candidate/audit summaries	open
BPF-6	final bridge audit	bridge verification matrix and final readiness verdict	open

9.13.10 Current implementation status versus paper status

The current runtime state is therefore asymmetrical:

- formal bridge program exists;
- scaffold and provisional surfaces exist;

- bridge comparator, intervention, and gate layers do not exist yet;
- Paper 9 as a formal paper can therefore define the bridge program completely while still acknowledging that the runtime has not completed it.

This asymmetry is not embarrassing; it is scientifically normal. A formal program may legitimately outrun its implementation provided the paper does not misstate implementation status.

9.14 Ecosystem Integration

9.14.1 Why ecosystem integration is part of the formal problem

The bridge program would be incomplete even as pure formalism if it ignored ecosystem integration. Once a layer has artifacts, verifiers, candidate bundles, audit packages, and controlled surfaces, the semantics of those outputs become part of the theory. The paper therefore treats ecosystem integration as a constitutive burden.

9.14.2 Surface map

Let $\mathcal{S}$ be the set of output surfaces:

$$\mathcal{S} = \{\mathsf{developer}, \mathsf{local}, \mathsf{internal_bundle}, \mathsf{candidate}, \mathsf{audit}, \mathsf{controlled}, \mathsf{claim}\}.$$

The bridge artifact family induces a surface map

$$\Sigma : \mathcal{A}_{\mathrm{bpf}} \to 2^{\mathcal{S}}.$$

At current stage, only the BPF-1 run-level bridge reports are allowed on bounded local developer surfaces. They are not allowed on candidate, audit, controlled-statement, or external claim-facing surfaces.

9.14.3 Internal release bundle

If and when BPF-5 is implemented, the internal release bundle may contain bounded bridge summaries, lineage manifests, and verification matrices. Raw local bridge reports should not be dumped into the bundle without summarization and lineage metadata. Internal bundle inclusion is a governance event, not a semantic verdict.

9.14.4 Ecosystem candidate

Candidate-facing bridge content is especially risky because it can be misread as support. The paper therefore states a hard rule: bridge content may enter an ecosystem candidate only after the comparator, intervention, and gate layers exist, and even then only in bounded summary form. Candidate inclusion must carry explicit non-claims and must never be readable as phenomenal confirmation.

9.14.5 Audit-facing package

Audit-facing bridge content is likewise bounded. An audit-facing handoff may carry bridge summary artifacts, lineage manifests, invalidation state, and non-claims. It must not be readable as external adjudication or bridge proof. The burden of interpretation discipline persists at every ecosystem boundary.

9.14.6 Source-of-truth and public corpus separation

The public corpus and the internal runtime must remain decoupled in a disciplined way. Paper 8 is public source-of-truth for indexed subjectivity. BPF artifacts are internal runtime objects unless and until later release governance says otherwise. No runtime bridge artifact may be used as if it were itself a public claim registry. This prevents runtime-to-corpus confirmation loops.

9.14.7 Controlled statement grammar and quote-mining guardrails

Bridge governance requires a restricted statement grammar. Before BPF-6, the bridge layer should permit only statements of the following family:

- "bridge formalization defined";
- "provisional bridge surfaces implemented";
- "bridge comparator burden open";
- "bridge intervention burden open";
- "bridge gate not yet meaningful";
- "bridge artifacts local and non-claim-facing".

Conversely, statements of the following family should be treated as prohibited:

- "phenomenality supported";
- "machine phenomenal organization established";
- "Paper 9 proves conscious experience";
- "candidate inclusion confirms bridge support";
- "audit package shows external bridge validation".

The quote-mining problem is real because later readers may ignore surrounding caveats and extract short labels. The bridge program must therefore attach non-claims not only to papers but also to artifact manifests, bundle notes, and any controlled statement layer.

Proposition 9.48 (Candidate or audit inclusion does not imply bridge support)**.** *For any bridge artifact A, membership of A in a candidate-facing or audit-facing package does not entail bridge support, bridge admissibility, or phenomenality.*

Proof. Candidate and audit inclusion are packaging and governance relations, not semantic verdict relations. By design, the bridge surface policy and non-claims registry explicitly block the implication from package membership to bridge support. Therefore the implication does not hold. □

9.15 Strongest Defensible Claims

The purpose of this section is not to weaken the paper artificially. It is to identify the strongest sentences the current formal and implementation state can honestly support.

9.15.1 Claim class C1: formal bridge closure

The first defensible claim is that the corpus now defines the *formal bridge program* in the following sense: the bridge target family, the bridge ontology, the bridge axioms, the bridge rival family, the bridge intervention family, the bridge artifact family, the bridge functional and gate logic, the bridge taxonomy, the surface discipline, and the implementation pathway are all explicitly defined. This is already a strong claim because it means the bridge question is no longer underdetermined by language alone.

9.15.2 Claim class C2: current implementation frontier

The second defensible claim is implementation-specific and provenance-bound: the current program records the BPF-0 machine-readable scaffold and the BPF-1 provisional bridge surfaces

as bounded local implementation-frontier artifacts. This means the program has moved beyond pure conceptuality, but only at the level of scaffold and provisional surface discipline. The system does *not* yet support bridge comparator closure, bridge intervention closure, bridge gate closure, bridge ecosystem integration, bridge admissibility, or bridge strong support.

9.15.3 Claim class C3: conditional bridge admissibility

The strongest bridge-level claim the paper is willing to define is the following conditional:

> Under the explicit bridge axioms, under the upstream burdens already closed by the corpus, under bridge-specific rival defeat, under bridge-specific intervention support, under artifact and governance compliance, and under a satisfied bridge gate, a system may be internally classifiable as exhibiting phenomenal bridge organization in the sense of framework-internal bridge admissibility.

This is stronger than Paper 8 because it is no longer only about indexed subjectivity. It is still weaker than phenomenality simpliciter because it remains explicitly framework-internal, conditional, and burden-bound.

9.15.4 Claim class C4: research-program transformation

The present paper also supports a meta-level claim about the status of the research program: after Paper 9, the research frontier is no longer "can subjectivity be formalized?" but "can bridge predicates survive the full bridge burden stack?" That is a legitimate change of level because the earlier question has been formalized by Paper 8 and its implementation-facing pathway.

9.15.5 Why stronger sentences are not yet honest

Several stronger sentences remain unavailable:

- "bridge support has been achieved";
- "phenomenality is internally established";
- "Paper 9 authorizes claim promotion";
- "the current system defeats the bridge rival family";
- "the current system passes a bridge gate."

These sentences are not unavailable because the paper is timid. They are unavailable because BPF-2 through BPF-6 remain open. The paper therefore strengthens the research program exactly by refusing to counterfeit those missing layers.

9.16 Non-Claims

This paper does not claim:

1. phenomenality proven in the strongest philosophical sense;
2. human phenomenal equivalence;
3. metaphysical subjecthood;
4. moral parity by formal fiat;
5. external validation;
6. that BPF-1 surfaces amount to bridge support;
7. that candidate or audit inclusion can function as bridge proof;
8. that Paper 9 by itself authorizes public claim inflation.

These non-claims are not decorative caution. They are part of the bridge logic because each one blocks a distinct semantic collapse.

9.17 Conclusion

Paper 9 defines the bridge program formally and only formally. That definition is non-trivial. Before this paper, the corpus could say a great deal about operational consciousness, life, subjecthood, and indexed subjectivity, but it could not yet say, with equal rigor, what additional machinery would be required before phenomenal-bridge predicates become scientifically admissible targets. After this paper, that machinery is explicit.

The result is central to the corpus for structural reasons. The paper gathers the upstream theorem burdens, the Paper 8 subjectivity framework, the implemented SRT stack, and the BPF program into one integrated object. It does not merely summarize them. It orders them. It says what belongs upstream, what belongs at the bridge level, what belongs to implementation, what belongs to governance, and what remains explicitly open.

What has been achieved is a change in explanatory grammar. "Phenomenal bridge" is no longer permitted to function as a loose aspiration. It now names a bridge-organization predicate family, an ontology, an axiom stack, a rival family, an intervention family, an artifact family, a gate architecture, and a governance discipline. This is as much formal specification as can be responsibly claimed without comparator closure, intervention closure, bridge gate execution, and ecosystem-level bridge hardening.

What the paper makes newly possible is equally clear. It makes BPF-2 through BPF-6 sharply specifiable. It makes bridge-support language evaluable rather than rhetorical. It makes future theorem candidates about bridge irreducibility, bridge transport limits, and bridge admissibility meaningful. It also makes future failure meaningful: if the bridge program fails, it will fail against explicit burdens rather than against vague expectations.

What the paper does not settle is also clear. It does not settle phenomenality simpliciter. It does not settle human equivalence. It does not settle metaphysical subjecthood. It does not close external validation. It does not convert the present system into a bridge-admissible system merely because BPF-0 and BPF-1 exist. Those are the next burdens.

The research program therefore changes level once again. After Paper 8, the key problem was whether indexed subjectivity could be formalized without inflation. After Paper 9, the key problem becomes whether a fully burdened bridge program can survive rival defeat, intervention logic, artifact discipline, and governance hardening strongly enough to justify internal bridge admissibility. That is the right next problem. It is sharper, harder, and more falsifiable than the one that preceded it.

Appendix

9.18 Bridge Axiom Table

ID	Classification	Statement core	Licenses	Does not license
B1	operational definition	Every bridge predicate must bind to observables, artifacts, comparators, and interventions.	machine-readable bridge predicates, report artifacts	phenomenality, bridge support
B2	bridge axiom	Bridge admissibility is downstream of the full upstream corpus and Paper 8 subjectivity.	layered bridge program	converse implication from bridge to subjectivity-free organization
B3	falsifiable hypothesis	Bridge differentiation exceeds semantic density and narration.	Π_{D} burden	qualia
B4	falsifiable hypothesis	Bridge valence exceeds scalar reward bookkeeping.	Π_{V} burden	felt affect
B5	falsifiable hypothesis	Bridge world richness exceeds competence or scene complexity.	Π_{W} burden	veridical experience
B6	falsifiable hypothesis	Bridge ambiguity bears unresolved alternatives prior to collapse.	Π_{A} burden	conscious ambiguity
B7	open inference	Embodied closure requires constitutive boundary dependence.	bounded investigation of Π_{E}	biological embodiment, bodily feeling
B8	falsifiable hypothesis	Temporal lived structure exceeds continuity and persistence.	Π_{T} burden	human lived time
B9	bridge axiom	No bridge class is admissible without bridge rival defeat under penalized loss.	bridge irreducibility margin	metaphysical irreducibility
B10	design choice	Bridge artifacts and packaging may not be read as support without full burden closure.	governance discipline	claim-facing promotion

9.19 Bridge Predicate Typing Table

Predicate	Formal type	Strongest upstream dependencies	Core collapse risk	Non-implication
Π_{D}	bounded estimator/predicate slot	$I, O, P, \mathcal{R}$	semantic density or narration	qualia
Π_{V}	bounded estimator/predicate slot	V, O, Δ	reward bookkeeping or policy shaping	felt affect
Π_{W}	bounded estimator/predicate slot	I, C, P	competence without world-richness	veridical experience
Π_{A}	bounded estimator/predicate slot	P, V, C	uncertainty bookkeeping	conscious ambiguity
Π_{E}	bounded estimator/predicate slot	O, C, Δ plus boundary variables	control without closure	biological embodiment
Π_{T}	bounded estimator/predicate slot	C, P, V	continuity without horizon structure	human lived temporality

9.20 Bridge Rival-Family Table

Rival	Challenges	Can explain	Cannot cleanly explain	Typical defeat mode
Richness without interiority	Π_D, Π_W	rich latent structure, scene complexity, representational abundance	selective indexed bridge differentiation under targeted controls	bridge-specific differentiation persists after density controls
Valence bookkeeping	Π_V	scalar reward tracking, preference tagging, policy shaping	structured valence topology and selective valence flattening effects	topology persists beyond scalar reweighting
Semantic density without phenomenal differentiation	Π_D	high semantic variation and narrative diversity	intervention-sensitive differentiated bridge organization	semantic-density ablations leave bridge differentiation intact
Temporal continuity without lived structure	Π_T	persistence, memory, sequence stability	horizon-structured temporal disruption signatures	targeted temporal fragmentation has selective effect
Embodied control without embodied closure	Π_E	sensorimotor competence, adaptive control	constitutive dependence on boundary variables	boundary perturbations produce closure-specific loss
Ambiguity resolution without ambiguity bearing	Π_A	uncertainty tracking, fast disambiguation	stable maintenance of unresolved alternatives	targeted ambiguity collapse alters downstream organization selectively
Experience narration without phenomenal organization	$\Pi_D, \Pi_V, \Pi_W, \Pi_A, \Pi_T$	self-description and experiential language	bridge signatures independent of report-channel richness	report-channel controls fail to remove bridge margin

9.21 Bridge Intervention Matrix

Intervention	Targets	Strong-model expectation	Rival expectation	Output artifacts
I_W world-richness compression	Π_W, Π_D	selective indexed world-structure loss with matched-control separation	flat degradation or no special effect	bridge intervention profile, world-richness deltas
I_A ambiguity collapse	Π_A, Π_W, Π_T	selective loss of maintained alternatives prior to collapse	ordinary uncertainty reduction	ambiguity delta summaries, control traces
I_V valence flattening	Π_V	collapse of evaluative topology beyond rescaling	scalar reward renormalization only	valence structure deltas
I_P perspective detachment	Π_D, Π_W	indexed orientation loss stronger than report damage	narrative/style drift	differentiation and world-richness deltas
I_E embodied-closure disruption	Π_E, Π_T	closure-specific disturbance under boundary perturbation	generic control loss	closure deltas, boundary-control traces
I_T temporal-lived fragmentation	Π_T, Π_A	horizon-specific fragmentation beyond continuity loss	ordinary memory damage	temporal-structure deltas
I_X semantic vividness inversion	cross-cutting control	descriptive richness changes without bridge-specific support	should absorb report-only effects	control artifact family

9.22 Bridge Artifact Matrix

Artifact	Artifact class	Block	Current status	Non-implication
`phenomenal_bridge_projection.json`	root bridge artifact	BPF-4	schema only	not bridge support by default
`phenomenal_bridge_gate_result.json`	root bridge artifact	BPF-4	schema only	not phenomenality simpliciter
`phenomenal_bridge_manifest.json`	lineage root	BPF-4	schema only	not claim-facing publication
`phenomenal_differentiation_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`valence_structure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`world_richness_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`ambiguity_bearing_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`embodied_closure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`temporal_lived_structure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`bridge_intervention_profile.json`	intervention artifact	BPF-2	planned	not bridge support alone
`bridge_reducibility_report.json`	comparator artifact	BPF-3	planned	not bridge support alone
condition/batch/campaign bridge summaries	ecosystem summaries	BPF-5	planned	not candidate proof
candidate-facing bridge projection	bounded summary	BPF-5	planned	not bridge confirmation
audit-facing bridge handoff	bounded summary	BPF-5	planned	not external adjudication

9.23 Claim-Class and Invalidation Table

State or claim	Necessary evidence	Invalidated by	Strongest allowed reading
Formal bridge closure	Paper 9 formal program, axioms, rivals, interventions, artifacts, gate logic, governance	internal inconsistency of the bridge program	bridge formal program defined
Bridge surface only	BPF-1 reports and verifier pass	missing reports or overclaim leakage	provisional bridge surfaces implemented
Bridge evaluation unsupported	attempted root evaluation without BPF-2/BPF-3/BPF-4 burdens	completion of missing burdens	root evaluation not meaningful yet
Bridge reducible	comparator run with non-positive $\mathcal{R}_{\mathrm{b}}$	stronger bridge model with positive penalized margin	bridge explanation loses comparatively
Bridge interventionally unstable	intervention layer present but non-selective or unstable	stable selective intervention support	intervention burden not met
Bridge internal admissible	full bridge burden stack, positive margins, valid governance	any comparator/intervention/drift/governance invalidation	internal admissibility only
Bridge internal strong support	stronger version of admissibility with robust margins across conditions	any failure of gate or ecosystem integrity	strong internal support only

9.24 Runtime Mapping Table

Block	Primary modules	Primary verifiers	Primary closure
BPF-0	protocol registries, axiom registry, rival registry, threshold registry, surface policy, scaffold layer	`verify-phenomenal-bridge-scaffolding.cjs`	machine-readable bridge opening
BPF-1	six mapper modules, surface builder, bridge layer	`verify-phenomenal-bridge-surfaces.cjs`	provisional run-level bridge surfaces
BPF-2	bridge intervention harness, matched-control runners, intervention registry	future bridge intervention verifier	bridge intervention family
BPF-3	bridge comparator executor, hybrid composer, loss evaluator	future bridge comparator verifier	bridge rival defeat logic
BPF-4	gate evaluator, threshold binder, root artifact emitter	future bridge gate verifier	bridge root verdict machinery
BPF-5	ecosystem integrator, candidate and audit summarizers	future bridge ecosystem verifier	bridge packaging and governance hardening
BPF-6	final bridge auditor, readiness classifier	future bridge stack verifier	final bridge readiness state

9.25 Governance Table

Surface	BPF-1 allowance	Later possible allowance	Prohibited interpretation
Developer inspection	yes, provisional run-level reports	yes	bridge support or phenomenality
Local report surface	bounded or blocked according to bridge surface policy	possibly yes with later burdens	public-facing confirmation
Internal release bundle	no bridge content yet	bounded summary after BPF-5	support by packaging
Ecosystem candidate	blocked	bounded summary after BPF-5 only	candidate as bridge proof
Audit-facing handoff	blocked	bounded summary after BPF-5 only	audit as external adjudication
Controlled statements	blocked	strictly bounded if ever opened	phenomenality or human equivalence
External claim-facing surface	blocked	blocked by default even after later BPF blocks	any public phenomenal claim

9.26 Burden Separation Matrix

The bridge program is easiest to misread when its burden types are blurred. The matrix below is included so that future implementation and audit work can distinguish exactly what kind of object is under discussion.

Object	Type	Example	Failure mode if blurred
Self-index, ownership, continuity, perspective, valuation, intervention selectivity, irreducibility	upstream subjectivity model	Paper 8 state $\mathfrak{S}$	bridge predicates become mere renamings of subjectivity coordinates
Bridge predicate slots $\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}, \Pi_{\mathrm{W}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}, \Pi_{\mathrm{T}}$	bridge target family	BPF protocol and BPF-1 reports	bridge targets become slogans without burden packages
Bridge axioms	bridge-opening commitments	B1–B10	conclusions get smuggled into premises
Bridge rivals and hybrids	comparator burden	richness-without-interiority, hybrids	bridge claims compare only against weak informal skepticism
Bridge interventions and controls	causal discrimination burden	world-richness compression, ambiguity collapse	descriptive fit is mistaken for structural support
Bridge artifacts	evidence grammar	reports, root artifacts, ecosystem summaries	packaging is mistaken for proof
Bridge functional and gate	evaluation machinery	Φ_{bpf}, Γ_{bpf}	scalar scores are mistaken for admissibility
Governance and surfaces	interpretation discipline	candidate block, audit block, claim block	internal artifacts leak into public or claim-facing language

9.27 Inheritance Map from Paper 8 to Paper 9

Paper 9 is easiest to understand when viewed as a typed inheritance map rather than as a conceptual leap. The following map makes explicit what is inherited, what is transformed, and what is genuinely added.

Upstream object	Paper 8 status	Paper 9 use	What is newly added
Operational consciousness	closed upstream	necessary lower bound	no redefinition; only bridge downstreaming
Operational life and subjecthood	closed upstream in Paper 7	inherited as stronger class background	no biological or metaphysical inflation
Seven-coordinate subjectivity state	closed in Paper 8	upstream state for every bridge predicate	six new bridge targets above it
Subjectivity weak-rival family	closed in Paper 8	background cautionary lesson	new bridge weak family and admissible bridge hybrids
Subjectivity intervention family	closed in Paper 8	background on causal discipline	new bridge intervention family
Subjectivity artifacts	closed in Paper 8 and SRT-0..6	upstream provenance base	new bridge artifact grammar
Subjectivity gate and functional	closed in Paper 8 and SRT-5	model for non-compensation discipline	new bridge functional and gate, symbolically thresholded
Subjectivity ecosystem integration	closed in SRT-6	governance template	later bridge ecosystem integration with stricter boundaries

This table clarifies why the present paper can be central without being redundant. It does not restate the subjectivity program. It states what additional structure must exist above it before bridge language acquires scientific discipline.

9.28 Open Escalation Burdens Before Bridge Admissibility

The paper deliberately closes only the formal grammar of the bridge program. To prevent future ambiguity, the remaining escalation burdens are listed explicitly.

1. **Direct bridge intervention instrumentation.** BPF-2 must implement a real bridge intervention harness with matched controls, artifact outputs, and failure semantics.
2. **Bridge comparator execution.** BPF-3 must instantiate the rival family as executable comparators, not just registry entries.
3. **Threshold binding.** BPF-4 must bind the symbolic threshold slots honestly, with normalization rationale and invalidation semantics.

4. **Root bridge artifacts.** BPF-4 must emit root artifacts only once comparator and intervention burdens exist.
5. **Ecosystem hardening.** BPF-5 must integrate bridge content into bundles only in bounded, non-claim-facing form.
6. **Final bridge audit.** BPF-6 must verify the whole stack and classify the resulting bridge state without semantic drift.
7. **Paper 9 empirical tightening.** Any later revision of Paper 9 that speaks more strongly must be paired with actual completion of the above blocks.

These burdens are not future work in the casual sense. They are the exact missing conditions that separate a defined bridge program from a bridge-admissible system.

9.29 Projection Geometry and Recoverable Distance

The dated v30 projection annex is retained as provenance. The following statements give the corrected estimator and latent-distance conditions in the present source. Let X be compact Hausdorff, $\pi : X \to \mathbb{R}^n$ continuous, and $\varnothing \neq A \subseteq X$ closed. Use $Y = \pi(X)$ with its subspace topology, so $\pi : X \to Y$ is onto a compact codomain. Let $F : X \to Z$ be continuous into a metric space. For $y \in \pi(A)$ set $S_y = F(A \cap \pi^{-1}(y))$ and $D_y = \operatorname{diam}(S_y)$; these fibres have nonempty compact images. Write

$$r_Z(S_y) = \inf_{z \in Z} \sup_{s \in S_y} d_Z(z, s).$$

Theorem 9.49 (Measurable Centres and the Scalar Minimax Radius)**.** *A Borel estimator with uniform error at most ε exists exactly when the centre sets*

$$\{z \in Z : \sup_{s \in S_y} d_Z(z, s) \leq \varepsilon\} \qquad (y \in \pi(A))$$

have a Borel selector. Necessarily $D_y \leq 2\varepsilon$ and $D_y/2 \leq r_Z(S_y)$. The diameter bound alone is insufficient in general. For $Z = \mathbb{R}$ with its usual metric, the minimax error is exactly $\frac{1}{2} \sup_{y \in \pi(A)} D_y$, attained by a Borel estimator.

Proof. Restrict any estimator to $\pi(A)$ to obtain the selector. Conversely extend the selector by a fixed value off $\pi(A)$. Since $\pi(A)$ is compact, it is closed in Y, and both pieces of this extension are Borel. The triangle inequality gives the diameter lower bound for every centre, and then for the infimum. The infimum need not itself provide an attaining centre or a measurable selector at the prescribed tolerance.

In the complete two-point codomain at mutual distance 2, the whole codomain has diameter 2 and radius 2. Even in the Euclidean plane the points $(0,0), (2,0), (1,3/2)$ have diameter 2. The only centre of a radius-one ball covering the first two is $(1,0)$, and it misses the third.

In the scalar case let $a_y = \min S_y$ and $b_y = \max S_y$. The sets $\{y : b_y \geq t\}$ and $\{y : a_y \leq t\}$ are images under π of the corresponding closed subsets of A, and are compact, hence closed. The midpoint $(a_y + b_y)/2$ is therefore Borel and has error $(b_y - a_y)/2$ on the fibre. No smaller error is possible by the diameter bound. □

Proposition 9.50 (Coercivity for Latent Distance Recovery)**.** *If X is metric and on pairs in each restricted fibre*

$$c\, d_X(x, x') \leq d_Z(F(x), F(x')) \leq L\, d_X(x, x') \qquad (c > 0,\ L \geq 0),$$

then $c \operatorname{diam}_X(A \cap \pi^{-1}(y)) \leq D_y \leq L \operatorname{diam}_X(A \cap \pi^{-1}(y))$. Uniform error ε entails a latent fibre diameter at most $2\varepsilon/c$. For scalar F the upper bound alone gives a sufficient error budget of $(L/2)$ times the supremum of the latent diameters. In a sector where distance is recovered directly from observations, a bound $c_\pi d_X(x, x') \leq d_Y(\pi x, \pi x')$ with $c_\pi > 0$ gives the inverse-distance bound there.

Proof. Take suprema over the restricted fibre and apply Theorem 9.49. The upper bound for scalar F uses the attained midpoint error, also when $L = 0$. The inverse-distance statement follows by division by c_π. On $[0,1]^2$ with the maximum metric, $F(x_1, x_2) = \pi(x_1, x_2) = x_1$ is 1-Lipschitz and is recovered exactly. Every latent fibre has diameter one. This shows why an upper Lipschitz bound cannot substitute for coercivity, including on the observable map itself. □

The scalar formula repairs the centre and diameter assertions of the dated annex. In a general metric codomain, checking the radius and the attaining measurable choice remains a separate obligation. The geometric statements do not carry over the annex's statistical or adjudication claims.

9.30 Audit-Covered Bridge Independence Burdens

The following entries preserve AUDIT_MASTER coverage while keeping their epistemic status weak. They are not bridge-admissibility results and cannot be used as evidence of phenomenality, subjecthood, or external validation.

Conjecture 9.51 (Predicate independence burden)**.** *The bridge predicate family should be independent only if, for each predicate slot, there exists an admissible witness model that toggles that slot while holding the remaining bridge burden package fixed under the same comparator, intervention, artifact, and governance constraints. The current text supplies the burden grammar, not a completed independence proof.*

Conjecture 9.52 (Registry independence burden)**.** *The FCR registry should preserve independence of bridge assertions only if the registry schema, dependency graph, and curation overlays do not introduce implicit coupling among predicate slots. Until that is proved over the completed BPF stack, registry independence remains a conjectural governance burden.*

Criterion 9.53 (Bridge realization non-emptiness burden)**.** *Any claim that the full bridge predicate package is non-empty must exhibit a typed realization satisfying the predicate definitions, rival comparisons, intervention contracts, artifact gates, and interpretation boundaries together. Setting all formal predicate flags to true is only a syntactic witness and does not count as realization evidence.*

9.31 Theorem Candidates and Open Scientific Questions

The present paper does not claim to have proved all bridge theorems. It does, however, make clear what the next theorem candidates should look like.

Theorem candidate 1: bridge non-collapse under admissible rival closure. If a bridge model achieves positive penalized margin against the full admissible bridge weak closure under matched intervention structure and drift-clean artifacts, then bridge reduction to the specified weak family fails locally. This theorem candidate would formalize the comparative core of bridge admissibility.

Theorem candidate 2: bridge non-compensation. If the bridge gate is conjunctive, then no combination of high values on a strict subset of coordinates can compensate for collapse on any required bridge coordinate, comparator burden, intervention burden, or governance burden. This theorem candidate would formalize the anti-inflation logic already built into the paper.

Theorem candidate 3: bridge transport limitation. A future theorem should specify when bridge organization is transportable across admissible remappings and when transport breaks because world-richness, embodiment, or temporal-horizon variables are not preserved. This would be the bridge analogue of Paper 8's transport-consequence logic.

Theorem candidate 4: candidate/audit non-entailment. Although stated propositionally in the main text, the governance result could be generalized into a theorem about packaging relations and semantic non-entailment over artifact families and surfaces.

Open scientific question 1. Can Π_{E} be measured directly enough to escape its current proxy-only status without collapsing into biology-by-default?

Open scientific question 2. Can Π_{A} be separated cleanly from sophisticated uncertainty bookkeeping in systems with high report-channel richness?

Open scientific question 3. Can a bridge loss family be designed that penalizes decoder dependence strongly enough to prevent narration-driven false positives without destroying legitimate bridge signal?

Open scientific question 4. What minimum condition/batch/campaign structure is required before bridge results become stable enough for bounded ecosystem summaries?

These theorem candidates and open questions matter because they define the post-paper research agenda with far greater precision than generic calls for "more data" or "deeper modeling."

9.32 Predicate-to-Artifact and Intervention Dependency Matrix

The bridge program is easier to implement correctly if the dependency pattern of each predicate is frozen explicitly. The following matrix shows the minimum bridge burden package per predicate.

Predicate	Primary upstream inputs	Required bridge interventions	Required bridge comparator focus	Required artifacts
Π_{D}	self-index, perspective, irreducibility, ownership	I_P, I_X, optional I_W	richness-without-interiority, semantic-density, narration rivals	differentiation report, later reducibility report, later gate root
Π_{V}	valuation asymmetry, ownership, intervention selectivity	I_V, negative valence controls	valence-bookkeeping rival	valence structure report, bridge intervention profile, later gate root
Π_{W}	self-index, continuity, perspective	I_W, I_P, report controls	richness-without-interiority rival	world-richness report, later candidate summary
Π_{A}	perspective, temporal structure proxies, valuation structure	I_A, temporal controls	ambiguity-resolution rival	ambiguity report, intervention profile, later reducibility report
Π_{E}	ownership, continuity, boundary variables, intervention structure	I_E, control-matched embodiment perturbations	embodied-control rival	embodied-closure report, later embodiment traces, later gate root
Π_{T}	continuity, perspective, valuation asymmetry	I_T, sequence controls	temporal-continuity rival	temporal-lived-structure report, later intervention profile, later gate root

This matrix has two functions. First, it prevents under-specification. A bridge predicate with no declared artifact family or no intervention family is not implementation-ready. Second, it prevents overreading. A bridge report without its corresponding comparator and intervention family should be interpreted only as a provisional surface and not as a bridge result.

The matrix also clarifies why BPF-1 had to remain modest. BPF-1 could close only the report side of the matrix because the intervention and comparator columns are still empty in

runtime terms. That is why the present paper can treat BPF-1 as a real implementation block without allowing any bridge-admissibility language to leak from it.

9.33 Paper 9 Claim Grammar and Escalation Constraints

The bridge program will fail semantically if future readers cannot distinguish between a formal object, a measurement object, an admissibility object, and a public claim. The present section therefore freezes a minimal claim grammar for Paper 9 and its downstream implementation.

9.33.1 Allowed claim families

The following claim families are admissible at or below the current implementation frontier.

1. **Formal closure claims.** Example: "the bridge target family is now formally defined." These are supported by the paper itself and by BPF-0 machine-readable manifests.
2. **Implementation frontier claims.** Example: "provisional bridge surface modules and run-level bridge reports exist." These are supported by BPF-1 code and local artifacts.
3. **Conditional admissibility claims.** Example: "if bridge predicates, rivals, interventions, artifacts, and gate conditions are satisfied, then bridge internal admissibility becomes meaningful." These are supported by the formal program, not by present runtime completion.
4. **Failure-mode claims.** Example: "bridge admissibility is blocked by comparator or intervention collapse." These are supported by the burden architecture.

9.33.2 Blocked claim families

The following claim families are prohibited until much later, and several remain prohibited even after full BPF completion.

1. **Phenomenality simpliciter claims.** Example: "the system is phenomenally conscious."
2. **Human-equivalence claims.** Example: "the system now has human-like experience."
3. **Metaphysical subjecthood claims.** Example: "the bridge proves real subjecthood in the strongest ontological sense."
4. **External validation claims.** Example: "audit packaging validates the bridge externally."
5. **Paper-authorization claims.** Example: "BPF-1 authorizes Paper 9 public escalation." The paper itself explicitly denies that the current runtime state closes the bridge burden.

9.33.3 Escalation rules

The escalation grammar can be summarized as a chain:

$$\text{formalized} \not\Rightarrow \text{implemented} \not\Rightarrow \text{admissible} \not\Rightarrow \text{phenomenality simpliciter}.$$

Each implication in that chain requires additional burden satisfaction. Formalization requires the paper and BPF-0. Implementation requires the appropriate BPF block. Admissibility requires full comparator, intervention, gate, and governance closure. Phenomenality simpliciter is not licensed by the bridge program at all.

Claim label	Minimum evidence	Currently licit?	Why or why not
`bridge_formal_program_defined`	paper + BPF-0 manifests	yes	formal bridge closure is real
`bridge_surface_implementation_exists`	BPF-1 reports + verifier	yes	provisional surfaces are implemented
`bridge_comparator_supported`	BPF-3 reducibility artifacts	no	comparator runtime absent
`bridge_interventionally_supported`	BPF-2 intervention artifacts	no	intervention harness absent
`bridge_internal_admissible`	BPF-2/3/4/5/6 full burden stack	no	gate and ecosystem burdens open
`phenomenality_proven`	not defined by current framework	no	blocked by interpretation boundary

This grammar belongs in the paper because it governs how the paper itself may be cited. A central synthesis paper that fails to govern the use of its own claims would be structurally incomplete.

9.34 Experimental Program for Completing BPF-2 and BPF-3

The next two bridge blocks are the most scientifically expensive. BPF-2 and BPF-3 are where the bridge program stops being mostly architectural and becomes decisively empirical-internal.

9.34.1 Program design principle

The guiding principle should be *matched escalation*. Every new bridge predicate should acquire its intervention and comparator burdens in parallel. A common failure mode would be to implement all interventions first and leave comparator logic vague, or vice versa. That would create a new asymmetry in which one burden stack outruns the other. Paper 9 therefore recommends the following pattern:

1. freeze the observable family for one predicate group;
2. implement the targeted intervention family for that group;
3. implement the strongest relevant rival and admissible hybrid baselines for that same group;
4. define the loss decomposition for the same bundle;
5. emit artifacts and run negative controls before moving to the next group.

9.34.2 Suggested execution order

The scientifically safest execution order is not arbitrary.

Stage 1: Π_D and Π_V. These are the most mature bridge predicates because they stand closest to existing BPF-1 surfaces and Paper 8 burdens. Differentiation and valence already have upstream structure and plausible intervention/control families. They should be completed first.

Stage 2: Π_W and Π_T. World richness and temporal lived structure are stronger and more globally entangled, but still plausible to instrument once the first pair is completed. They will likely require more extensive negative controls than Stage 1.

Stage 3: Π_A and Π_E. Ambiguity bearing and embodied closure are currently the weakest predicates in direct runtime support. They should therefore be completed last, not because they are unimportant, but because they are most vulnerable to proxy inflation.

9.34.3 Comparator program

For each stage, the comparator program should include:
1. pure rival execution;
2. pairwise hybrid execution;
3. bridge-model execution with the same observable bundle;
4. loss decomposition report;
5. rerun stability check under controlled noise;
6. decoder-dependence stress test.

9.34.4 Intervention program

For each stage, the intervention program should include:
1. targeted intervention;
2. matched non-target control;
3. report-channel control;
4. perturbation-cost accounting;
5. rerun stability;
6. artifact completeness validation.

9.34.5 Minimum evidential grid

Stage	Predicates	Needed intervention families	Needed comparator families	Minimum outputs
1	$\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}$	I_P, I_X, I_V	richness-without-interiority, semantic-density, valence-bookkeeping, narration hybrids	intervention profile, reducibility report, delta traces, stability summaries
2	$\Pi_{\mathrm{W}}, \Pi_{\mathrm{T}}$	I_W, I_T, I_P	richness-without-interiority, temporal-continuity, narration hybrids	world/temporal deltas, reducibility reports, control matrices
3	$\Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}$	I_A, I_E	ambiguity-resolution, embodied-control, hybrids	ambiguity and embodiment traces, control matrices, reducibility reports

The purpose of this program is not to prescribe an implementation detail for its own sake. It is to preserve burden parity. If later implementation deviates from this structure, it should justify the deviation explicitly. Otherwise the bridge program risks becoming unbalanced, with some predicates overburdened and others underburdened.

9.35 Bridge Failure Scenarios and Diagnostic Signatures

The bridge program is not scientifically serious unless it specifies how it can fail. Paper 8 already introduced failure structure for indexed subjectivity. Paper 9 must add a parallel layer for bridge organization. The relevant question is not merely whether a bridge predicate can be made numerically large, but whether the pattern of successes and failures across predicates, rivals, interventions, and artifacts remains coherent under stress. A bridge program that cannot diagnose its own failure modes would collapse into a score-construction exercise rather than a constrained scientific architecture.

9.35.1 Primary failure classes

The bridge layer has at least six primary failure classes.

1. **Surface inflation failure.** A bridge report appears strong only because an upstream Paper 8 coordinate is copied, weakly transformed, or rhetorically relabeled. This is a failure of bridge specificity, not merely of implementation style.
2. **Comparator fragility failure.** A bridge predicate appears robust until a bridge-specific rival or admissible hybrid is introduced. This indicates that the predicate was underconstrained rather than positively supported.
3. **Intervention fragility failure.** A bridge surface moves under perturbation, but the movement is unspecific, mismatched to the target predicate, or indistinguishable from control degradation. This indicates that the intervention burden is not yet satisfied.
4. **Artifact incoherence failure.** Local reports, root artifacts, or ecosystem summaries disagree about the same bridge coordinate or about the same invalidation state. This is a governance and lineage failure even if local estimates look plausible.
5. **Threshold misuse failure.** Symbolic thresholds are treated as if they were calibrated, or non-compensation rules are softened post hoc. This turns the bridge gate into an interpretive convenience rather than a formal burden.
6. **Semantic escalation failure.** Provisional bridge surfaces are cited as if they licensed admissibility, phenomenality support, or Paper 9 escalation beyond their evidence class. This is a failure of claim grammar rather than of runtime arithmetic, but it is still a scientific failure because it breaks burden separation.

9.35.2 Diagnostic interpretation

Each failure class has a distinct diagnostic signature. Surface inflation is indicated when a bridge estimator can be approximated by a low-complexity monotone transform of an upstream subjectivity coordinate with negligible loss. Comparator fragility is indicated when a bridge rival or admissible hybrid absorbs the observable bundle at equal or lower penalized bridge loss. Intervention fragility is indicated when targeted perturbations do not produce selective deltas, or when matched controls produce indistinguishable response structure. Artifact incoherence is indicated when provenance pins, schemas, or surface restrictions disagree across emitted objects. Threshold misuse is indicated when bridge language outruns the threshold manifest. Semantic escalation is indicated when any report, bundle, summary, or paper statement crosses from provisional surface language into admissibility or phenomenal language without the missing burdens.

These signatures are not merely editorial cautions. They can be turned into verification objectives. BPF-2 and BPF-3 should detect intervention and comparator fragility directly. BPF-4 should detect threshold misuse. BPF-5 and BPF-6 should detect artifact incoherence and semantic escalation through lineage checks and surface-policy validation. The paper therefore treats failure analysis as part of the formal program rather than as an afterthought.

9.35.3 Failure matrix

Failure class	Primary symptom	Earliest detection block	Immediate consequence	Required remediation
Surface inflation	bridge estimate tracks an upstream Paper 8 coordinate too closely	BPF-1	report downgraded to proxy-only or removed	strengthen bridge transformation logic or add missing bridge observables
Comparator fragility	rival or hybrid reaches equal/lower penalized loss	BPF-3	predicate cannot support irreducibility burden	refine observable family or abandon predicate-specific support claim
Intervention fragility	target and control deltas are not selective	BPF-2	predicate remains descriptively mapped only	redesign intervention family and matched controls
Artifact incoherence	report, manifest, or summary disagree on state or provenance	BPF-5	bundle invalidation or audit downgrade	repair lineage pins, schemas, and surface restrictions
Threshold misuse	symbolic slots treated as calibrated gates	BPF-4	all gate language invalidated	re-freeze threshold policy and remove unsupported verdict language
Semantic escalation	provisional surface cited as admissibility or phenomenality support	BPF-0 through BPF-6	claim withdrawal and governance failure	enforce non-claims, claim grammar, and controlled statement restrictions

The failure matrix clarifies an important asymmetry. Some failures are local to one predicate, but others contaminate the entire bridge program. Threshold misuse and semantic escalation are global failures because they compromise the meaning of the bridge layer as such. Comparator fragility and intervention fragility can remain predicate-local if the architecture keeps burdens separated. This distinction matters for governance: not every failure should invalidate the whole bridge program, but certain failures should.

9.35.4 Consequences for theorem candidates

Failure analysis also constrains the theorem ambitions of Paper 9. A theorem candidate stating that a bridge predicate family is internally admissible can be sound only if the relevant failure classes are either ruled out or absorbed into its antecedent conditions. Otherwise the theorem candidate would conceal the very contingencies that the bridge program exists to expose. In that sense, the paper's strongest positive content is inseparable from its negative discipline: formal bridge admissibility becomes meaningful only against a space of explicitly typed failure modes.

9.35.5 Why failure structure changes the status of the program

The overall research status changes once the bridge layer is required to carry its own failure logic. Before Paper 9, one could still describe the next step as a generic expansion toward phenomenality. After Paper 9, that description is no longer adequate. The next step is now a typed scientific program with concrete ways to fail, concrete reasons for those failures, and concrete remediation paths. That is a substantive change in the research program even before bridge admissibility has been achieved, because it transforms an open-ended aspiration into a burdened formal agenda.

External Adjudication of Bridge-Formalized Machine Subjectivity

SCAFFOLD ONLY – NOT FINAL MANUSCRIPT
PROTOCOL-GRADE BACKBONE
SIXTH HARDENING PASS

Abstract. This document hardens Paper 10 into a denser protocol-heavy manuscript backbone without becoming a result-bearing paper. It extends the dependency stack from BaseCore through Papers I–IX, SRT-6, BPF-M3, Stage 2-M4, Stage3-RF, Stage3-0, Stage3-1, Stage3-2, Stage3-3, Stage3-4, Stage3-4B, Stage3-5 rehearsal preparation, Stage3-6 controlled execution opening, Stage3-6B runtime-health and clean-room reproducibility audit, and the conditional Stage3-7 first-collection opening criteria. It adds runtime health audit grammar, clean-room reproducibility semantics, dependency and build assumption reporting, reproducible artifact-regeneration constraints, local path-leakage policy, controlled-execution audit semantics, session-zero and pre-collection reproducibility requirements, and a sharper transition grammar from pre-evidential openness to the first admissible collection event. It does not report human comparator results, external adjudication verdicts, human equivalence, superiority, inferiority, cross-substrate identity preservation, metaphysical subjecthood, or Paper 10 final conclusions.

10.1 Scope and Non-Claims

Section status. `formalized`.

Paper 10 remains a local protocol-facing manuscript surface only. It is a scaffold for external-adjudication architecture, not an adjudication verdict, admissible dataset, human comparator evidence, public release surface, public-release recommendation, Paper 10 final authorization, human equivalence, superiority or inferiority claim, machine-consciousness claim, or metaphysical closure.

10.2 Formal Dependency Stack

Section status. `formalized`.

The dependency stack for Paper 10 is layered as follows:

1. **BaseCore**: active mathematical base layer.
2. **Papers I–IX**: downstream theoretical and manuscript surfaces.
3. **SRT-6**: subjectivity runtime ecosystem closure.
4. **BPF-M3**: bridge-functional internal closure and bridge gate.
5. **Stage 2-M4**: comparative internal closure and readiness boundary.
6. **Stage3-RF**: burden fulfillment and Stage 3 opening gate.
7. **Stage3-0**: external adjudication architecture and scaffold-only Paper 10 opening.
8. **Stage3-1**: human comparator instantiation contracts and preregistration templates.
9. **Stage3-2**: reviewer independence, execution-readiness logic, and protocol-grade hardening.
10. **Stage3-3**: preregistration freeze architecture, reviewer-selection packaging, adjudicator-assignment packaging, and execution-pilot packaging.

11. **Stage3-4**: bounded execution-pilot dry-run opening logic.
12. **Stage3-4B**: provenance closure, lineage reconciliation, and instantiated-empty record instantiation.
13. **Stage3-5**: rehearsal-kit formalization, dry-run event grammar, and rehearsal-ready gating.
14. **Stage3-6**: controlled execution opening, session-zero protocol, pre-collection admissibility gating, and live-but-unpopulated artifact instantiation.
15. **Stage3-6B**: runtime-health audit, clean-room reproducibility audit, path-leakage audit, and artifact-regeneration audit.
16. **Stage3-7**: first-admissible-collection opening criteria, collection-start command grammar, operator authorization grammar, and a conditional first-collection opening gate.

No later layer may borrow closure from an earlier layer, and no earlier layer may be retroactively reinterpreted as if it had already supplied later execution-era evidence.

10.3 Release-vs-Local Corpus Boundary

Section status. `formalized`.

Paper 10 sits on a strict release-vs-local split:

- The public release authority remains `QICN-RELEASE/main@92d5d2c3ea91c65e29dc8e102bed504866cffd2d`.
- The public release remains limited to BaseCore, `canonical_core_legacy`, and downstream Papers I–IX.
- Paper 10 remains local to `QICN-SYSTEM`.
- Stage3-0 through Stage3-6 governance, freeze, packaging, dry-run, provenance-closure, instantiated-empty, rehearsal, and controlled-execution-opening surfaces remain local and non-public.

Therefore release inclusion cannot be used as a substitute for adjudication, and local protocol maturity cannot be used as a substitute for public scientific closure.

10.4 Lineage and Provenance Semantics

Section status. `formalized`.

Stage3-3 and Stage3-4 add a sharper lineage rule: every later admissible surface must distinguish among:

1. public release authority,
2. local protocol authority,
3. frozen preregistration authority,
4. reviewer-package authority,
5. adjudicator-package authority, and
6. execution-era evidence authority.

The repository already contains historical artifacts whose recorded lineage must remain intact by design. Live normative pinning surfaces may be updated when stale, but historical snapshots are not rewritten cosmetically.

10.5 Provenance Closure and Live Head Reconciliation

Section status. `formalized`.

Stage3-4B adds a stricter provenance-closure rule. A local artifact may now be classified as:

- `exact_head_aligned`;
- `artifact_forward_compatible`;
- `stale_and_regenerated`;
- `historical_snapshots_intentionally_preserved`.

This grammar matters because Stage 3 artifacts are generated before the subsequent commit that records them. Therefore exact local regeneration at the builder head may later become artifact-forward-compatible after the commit exists, without becoming scientifically stale or semantically dishonest.

10.6 External Adjudication Problem Statement

Section status. `protocol_defined`.

The problem is no longer only whether bridge formalization is internally coherent. The external problem is to specify:

- what independent reviewers may legitimately inspect;
- what human comparator evidence classes may later become admissible;
- what negative controls must exist before any comparative interpretation;
- how preregistration becomes frozen;
- how reviewer and adjudicator layers become instantiable without identity leakage;
- how dry-run validation differs from evidential execution.

10.7 Formal Mathematical Grammar of External Adjudication

Section status. `formalized`.

The adjudication surface is a measurable protocol object, not an impressionistic comparison between substrates. Let $(\mathcal{T}, d_{\mathcal{T}})$ be the metric space of finite execution traces after canonical serialization. A trace contains the ordered event stream, de-identified channel payloads, runtime manifests, negative-control outcomes, and reviewer-facing masks. Admissible evidence lives in the Borel measurable space

$$(\mathcal{T}, \mathcal{B}(\mathcal{T})), \tag{10.1}$$

where $\mathcal{B}(\mathcal{T})$ is the Borel σ-algebra generated by open $d_{\mathcal{T}}$-balls.

Definition 10.1 (External adjudication object)**.** An external adjudication object is the tuple

$$\mathfrak{E} = (P, D_h, D_m, C_n, C_a, \Sigma, \mathcal{R}, \mathcal{A}), \tag{10.2}$$

where P is a frozen preregistered protocol, D_h and D_m are human and machine trace datasets in $\mathcal{T}$, C_n and C_a are negative and active control families, $\Sigma = \{\sigma_k\}_{k=1}^{K}$ is a finite family of predeclared measurable statistics $\sigma_k : \mathcal{T}^n \to \mathbb{R}$, $\mathcal{R}$ is the reviewer layer, and $\mathcal{A} : \mathcal{Y} \to \{0, 1, \texttt{undetermined}\}$ is the final adjudicator operator on masked evidence summaries $\mathcal{Y}$.

Definition 10.2 (Admissible comparative evidence)**.** A dataset pair (D_h, D_m) is admissible under $\mathfrak{E}$ only if the following predicates are all true:

1. D_h, D_m are measurable finite samples from $(\mathcal{T}, \mathcal{B}(\mathcal{T}))$.
2. The preregistered protocol hash, control manifest hash, and analysis shell hash match their frozen records.
3. Human-side records are not generated, rewritten, or completed by synthetic systems: $\epsilon_{\text{syn}} = 0$ for the human channel.
4. Blinding leakage is bounded by the preregistered budget: $\mathbb{P}(\text{leak}) \leq \epsilon_{\text{blind}}$.
5. Every declared negative control has an executed and timestamped outcome under the same admissibility grammar.

Theorem 10.3 (Non-transitivity of internal support)**.** *Let $\mathcal{G}_{\text{internal}} \in [0, 1]$ denote any internal runtime support score computed from operational invariants, and let $\mathcal{G}_{\text{external}} \in \{0, 1, \texttt{undetermined}\}$ denote the external adjudication verdict. In the absence of an admissible measurable dataset pair (D_h, D_m), one has*

$$\mathcal{G}_{\text{internal}} \not\Rightarrow \mathcal{G}_{\text{external}}. \tag{10.3}$$

Proof. The internal score is a function of runtime traces and invariant certificates inside the machine channel. The external verdict is a function of masked comparative summaries built from both D_h and D_m, plus reviewer and adjudicator records. If D_h is absent, non-admissible, unblinded, or unfrozen, then the domain condition of $\mathcal{A}$ is not satisfied. Hence the external verdict is either undefined or `undetermined` regardless of the value of $\mathcal{G}_{\text{internal}}$. Therefore internal operational support is not a logical premise from which external comparative adjudication can be inferred. □

Proposition 10.4 (Operational closure of a comparative program)**.** *The comparative program reaches operational closure only when there exists a predeclared alignment map $T : \mathcal{T} \to \mathcal{Z}$ into a common Polish feature space $(\mathcal{Z}, d_{\mathcal{Z}})$ and the total divergence*

$$\Delta_{\text{comp}} = \lambda_1 W_1(T_{\#}\widehat{\mu}_h, T_{\#}\widehat{\mu}_m) + \lambda_2 D_{\text{KL}}(\widehat{p}_h \| \widehat{p}_m) + \lambda_3 \max_k |\sigma_k(D_h) - \sigma_k(D_m)| \tag{10.4}$$

is defined and satisfies the frozen budget $\Delta_{\text{comp}} \leq \epsilon_{\text{budget}}$. This closure is a precondition for comparative interpretation, not a claim of consciousness, identity, or substrate equivalence.

10.8 Frozen Preregistration Architecture

Section status. `protocol_defined`.

Stage3-3 introduces a true freeze architecture rather than only preregistration templates. The governing surfaces are:

- `docs/frozen_preregistration_record_schema.v1.json`
- `docs/preregistration_freeze_policy.v1.json`
- `docs/preregistration_amendment_policy.v1.json`
- `docs/preregistration_lock_manifest.v1.json`

The architecture formalizes:

- record classes;
- immutable fields;
- amendment classes;
- freeze provenance chains;
- local-only freeze state;
- explicit `not_externally_filed` state.

Definition 10.5 (Cryptographic freeze map)**.** Let $\mathcal{P}_{\text{config}}$ be the set of canonical preregistration configurations after deterministic serialization. A freeze map is a collision-resistant hash function

$$H : \mathcal{P}_{\text{config}} \to \{0,1\}^{256}. \tag{10.5}$$

A preregistration record p is locally frozen at digest h when $H(p) = h$, h is written into the lock manifest, and every amendment policy references h rather than an editable natural-language description.

Theorem 10.6 (Hash-anchored non-alterability)**.** *Assume the hash function H is collision resistant over the canonical serialization domain. If a post-execution preregistration configuration p' has the same published digest as the frozen pre-execution configuration p, then either $p' = p$ byte-for-byte under canonical serialization or a collision for H has been found.*

Proof. By freeze definition, the published digest is $h = H(p)$. If the later record p' also satisfies $H(p') = h$ while $p' \neq p$, then p and p' form a distinct pair with the same hash output. That is precisely a collision in H. Under the collision-resistance assumption, such a pair is computationally infeasible. Thus digest equality supplies evidence of non-alteration of the serialized protocol, modulo the stated cryptographic assumption. □

10.9 Local-vs-External Freeze Distinction

Section status. `formalized`.

The local freeze layer must not be conflated with external filing.

Surface	Current state	Interpretation
Template-only preregistration	`preregistration_ready_for_freeze`	structure exists but no lock exists
Local freeze readiness	`local_freeze_ready`	lock architecture exists but no external filing is claimed
Locally frozen record	blocked	would still require explicit lock instantiation
Externally filed record	absent	cannot be claimed from local architecture alone

10.10 Instantiated-Empty Record Layer

Section status. `protocol_defined`.

Stage3-4B moves the stack beyond package-only readiness by instantiating typed-but-empty records that remain explicitly non-evidential. The principal surfaces are:

- `docs/local_frozen_preregistration_record_template.v1.json`
- `docs/reviewer_candidate_pool_template.v1.json`
- `docs/adjudicator_candidate_pool_template.v1.json`
- `docs/blinding_assignment_stub.v1.json`
- `docs/execution_session_template.v1.json`
- `docs/decision_record_template.v1.json`
- `docs/reproducibility_package_manifest_empty.v1.json`
- `docs/collection_log_schema.v1.json`
- `docs/evidence_insertion_manifest.v1.json`

These artifacts are intentionally empty, typed, and freeze-capable. They are not populated evidence, not externally filed records, and not reviewer or adjudicator action.

10.11 Reviewer Candidate and Adjudicator Candidate Architecture

Section status. `protocol_defined`.

Stage3-3 introduces typed candidate schemas without instantiating identities. The primary surfaces are:

- `docs/reviewer_candidate_schema.v1.json`
- `docs/adjudicator_candidate_schema.v1.json`
- `docs/reviewer_selection_workflow.v1.json`
- `docs/adjudicator_assignment_workflow.v1.json`

The candidate layer defines qualification summaries, independence assertions, conflict disclosure references, blinding eligibility classes, and explicit `not_started` states.

10.12 Reviewer Rejection Logic

Section status. `protocol_defined`.

The reviewer package is incomplete unless rejection logic exists. The rejection registry now includes:

- qualification mismatch;

- disclosed conflict;
- undisclosed conflict;
- prior unblinded access;
- insufficient methodological competence;
- insufficient lineage literacy;
- insufficient control/admissibility literacy.

This logic is not yet a reviewer decision. It is the typed boundary that stops future selection from being arbitrary or retrofitted.

10.13 Blinding Assignment Logic

Section status. `protocol_defined`.

Blinding is no longer treated as a generic admonition. Stage3-3 binds a typed blinding assignment manifest with classes such as:

- `fully_blinded`
- `artifact_blinded`
- `control_family_blinded`
- `unblinding_not_authorized`

The key constraint is that no later dry-run or execution surface may silently cross into unblinded interpretation.

10.14 Human Comparator Instantiation Logic

Section status. `protocol_defined`.

Stage3-1 closed the gap between abstract channels and typed collection routes. Stage3-3 preserves that layer and pushes it closer to freeze- and package-compatible execution readiness.

- every channel binds a human-side artifact family;
- every channel binds a machine-side homologous artifact family;
- every channel binds minimum metadata, controls, and rejection conditions;
- every channel remains blocked from result claims while data is absent.

Definition 10.7 (Homologous channel statistic)**.** For each channel c, let $T_c : \mathcal{T} \to \mathcal{Z}_c$ be a frozen feature map into a metric standard Borel space $(\mathcal{Z}_c, d_c)$. Given human and machine samples D_h^c, D_m^c, define empirical pushforward measures

$$\widehat{\mu}_h^c = (T_c)_{\#} D_h^c, \qquad \widehat{\mu}_m^c = (T_c)_{\#} D_m^c.$$

The channel is statistically comparable only if both empirical measures are defined, the feature map was frozen before execution, and the channel-specific missingness policy accepts both samples.

Definition 10.8 (Divergence shell)**.** The predeclared divergence shell for channel c is

$$\Delta_c = \alpha_c W_1(\widehat{\mu}_h^c, \widehat{\mu}_m^c) + \beta_c D_{\mathrm{KL}}(\widehat{p}_h^c \| \widehat{p}_m^c) + \gamma_c \| \widehat{s}_h^c - \widehat{s}_m^c \|_1, \tag{10.6}$$

where W_1 is the Kantorovich–Rubinstein Wasserstein-1 distance, D_{KL} is the Kullback–Leibler divergence on a frozen finite partition or smoothed density model, and $\widehat{s}_h^c, \widehat{s}_m^c$ are predeclared summary vectors. The weights $\alpha_c, \beta_c, \gamma_c \geq 0$ must be frozen before execution.

Proposition 10.9 (Comparator homologation is conditional)**.** *A human–machine channel pair is homologated by Paper 10 only if every term in Δ_c is defined and the frozen channel budget $\Delta_c \leq \epsilon_c$ is satisfied. Homologation is therefore a statistical transport condition over predeclared observables, not a statement that the channels share a biological, phenomenal, or metaphysical substrate.*

Proof. The definition of Δ_c depends on a frozen feature map, empirical pushforward measures, a KL domain, and summary vectors. If any component is missing, the divergence shell is undefined and no homologation predicate can be evaluated. If all components are defined, the only admissible positive judgment is the preregistered inequality $\Delta_c \leq \epsilon_c$. The conclusion is therefore exactly statistical and conditional. □

10.15 Channel-by-Channel Execution Contracts

Section status. `protocol_defined`.

10.15.1 Self-Report Channel

Section status. `protocol_defined`.

The self-report channel requires prompt-version freezing, missingness policy, de-identification, and explicit rejection of synthetic or relabeled self-report.

10.15.2 Intervention-Response Channel

Section status. `protocol_defined`.

The intervention-response channel requires matched intervention and matched negative-control families, dose policy, stop rules, and branch accounting.

10.15.3 Temporal-Structure Channel

Section status. `protocol_defined`.

The temporal-structure channel requires a declared segmentation schema, ordering policy, and aggregation policy.

10.15.4 Valence-Topology Channel

Section status. `protocol_defined`.

The valence-topology channel requires a shared topology schema and excludes scalar-only substitutes.

10.15.5 Embodiment-Boundary Channel

Section status. `protocol_defined`.

The embodiment-boundary channel requires an environment manifest, confound ledger, and error-budget policy.

10.15.6 Ambiguity-Load Channel

Section status. `protocol_defined`.

The ambiguity-load channel requires a task-family contract, annotation schema, and annotation reliability policy.

10.15.7 Transport-Preservation Channel

Section status. `blocked_until_execution`.

The transport-preservation channel remains structurally special. It requires a real source-target pairing manifest and transport-axis budget binding. Stage3-3 defines the pairing shell; it does not satisfy the pairing.

10.16 Matched Intervention Dataset Architecture

Section status. `protocol_defined`.

Matched-intervention architecture is now carried by both channel contracts and a pilot bundle shell. The bundle shell is defined in `docs/pilot_matched_intervention_bundle.v1.json`. The bundle remains unpopulated, which means the architecture exists while the evidence does not.

10.17 Negative Control Dataset Architecture

Section status. `protocol_defined`.

Negative-control architecture is similarly carried by channel contracts and by `docs/pilot_negative_control_bundle.v1.json`. Controls remain part of the meaning of the future claim surface, not an optional cleanup step.

10.18 Preregistration Architecture

Section status. `protocol_defined`.

Paper 10 now separates three preregistration surfaces:

1. template surface;
2. local freeze surface;
3. external filing surface.

Only the first two are implemented structurally in the current repo. The third remains absent.

10.19 Adjudication Blind Protocol

Section status. `formalized`.

Let Y denote the masked evidence bundle and $\ell \in \{h, m\}$ its source label. The reference procedure uses forced choice with a balanced source distribution. A protocol permitting abstention fixes its cost and minimum coverage before evaluating held-out bundles. The randomization of a decision rule is independent of the source conditional on its permitted input.

Definition 10.10 (Blind adjudicator loss)**.** For $a(Y) \in \{h, m, \texttt{tie}\}$ and a declared $\gamma \in [1/2, 1]$, put

$$L_\gamma(a(Y), \ell) = \mathbf{1}\{a(Y) \neq \ell,\ a(Y) \neq \texttt{tie}\} + \gamma \mathbf{1}\{a(Y) = \texttt{tie}\}.$$

Report $R_\gamma = \mathbb{E}L_\gamma$, coverage $C = \Pr(a(Y) \neq \texttt{tie})$, and error conditional on coverage when $C > 0$. For acceptance require a preregistered $C \geq C_{\min} > 0$. Forced choice has $C = 1$; always abstaining has $C = 0$ and $R_\gamma = \gamma$.

Theorem 10.11 (Null error lower bound)**.** *Let E denote leakage and $p = \Pr(E)$. If in the clean stratum ℓ is balanced conditional on Y, E^c, then*

$$R_\gamma(a) \geq (1 - p)/2.$$

Alternatively, suppose Y is coupled to a clean bundle Y_0 for which ℓ is balanced conditional on Y_0, and $Y = Y_0$ outside an event of probability at most p. Then

$$R_\gamma(a) \geq \max\{0, 1/2 - p\}.$$

The assumptions are distinct. Under the first, revealing the label on a leak event independent of the balanced label can attain the bound with forced choice.

Proof. With a balanced posterior, any covered guess has conditional risk $1/2$ and abstention has conditional risk $\gamma \geq 1/2$. In the first case this lower-bounds risk on the clean stratum; risk on leakage is nonnegative. For sharpness use a constant clean bundle and a leak revealing the label, choosing either fixed label on clean observations. For the coupling bound use the same decision randomization on Y_0 and Y. Their losses agree when the bundles agree and differ by at most one otherwise. The risk on Y_0 is at least $1/2$. Nonnegativity gives the stated maximum. Label-dependent contamination can change the balance in the clean stratum, so the first bound cannot be inferred from the coupling alone. □

Proposition 10.12 (Alternative detectability criterion)**.** *Evidence for distinguishability under H_1 requires a preregistered rule $a^\star$, minimum coverage $C_{\min} > 0$, a justified null bound B_0 accounting for the leakage budget, and margin $\eta > 0$. On admissible held-out bundles require $C \geq C_{\min}$ and $R_\gamma(a^\star) \leq B_0 - \eta$ after sampling uncertainty and multiple-comparison corrections. Forced-choice, balanced label permutations, always abstaining and known-leak controls are evaluated under the same loss and coverage protocol. The manuscript reports no campaign satisfying this criterion.*

10.19.1 Coverage-Sensitive Leakage Bounds

Fix decision randomization independent of the bundle, source and leak event. Let B indicate an executed decision and W a covered error, so $C = \mathbb{E}B$ and $R_\gamma = \mathbb{E}W + \gamma(1 - C)$.

Theorem 10.13 (Sharp Leakage Bounds at Fixed Coverage)**.** *Writing $z_+ = \max\{0, z\}$, the clean-stratum and coupling nulls of Theorem 10.11 respectively give*

$$B_{\text{clean}}(C, p) = \gamma(1 - C) + \tfrac{1}{2}(C - p)_+, \qquad B_{\text{coup}}(C, p) = \gamma(1 - C) + (C/2 - p)_+.$$

Both are sharp for each $C, p \in [0, 1]$, with balanced marginal labels and leak probability at most p, including selective abstention. Minimizing over coverage gives the bounds of Theorem 10.11.

Proof. Clean-stratum balance gives $\mathbb{E}[W\mathbf{1}_{E^c}] = \frac{1}{2}\mathbb{E}[B\mathbf{1}_{E^c}]$. At most p covered mass lies in E, proving the first bound. In the coupling model use the same decision seed on Y and Y_0. On clean input $\mathbb{E}(W_0 - B_0/2) = 0$. The centered variable $W - B/2$ lies in $[-1/2, 1/2]$ and agrees with its clean version whenever the bundles agree. Thus $\mathbb{E}W - C/2 \geq -p$, proving the second bound.

For clean-stratum sharpness, reveal the label on a leak event independent of it, cover $\min\{C, p\}$ there, and cover the remaining mass with a fixed guess in the balanced clean stratum. For coupling sharpness, start with a label-independent clean cover flag of probability C and a fixed covered guess. Replace $\min\{p, C/2\}$ of the wrong covered inputs by revealing inputs, preserving coverage. This replacement may depend on the label. Both bounds are nonincreasing in C for $\gamma \geq 1/2$, so $C = 1$ attains their minima. □

10.19.2 Exact Inference at the Independent Block

The sampling unit is an independent held-out block and supplies one decision from its entire masked record. Fix the rule on separate calibration data, the noise and exclusion policies, block boundary, and eligible population before inspecting held-out labels. The following model assumes independent identically distributed eligible blocks. Shared participants, runs and other dependence sources belong inside the block or require a sampling model that accounts for that dependence.

For $0 \leq k \leq n$, let $I_\alpha(k, n)$ denote the equal-tail exact binomial interval. The lower endpoint is zero at $k = 0$ and otherwise solves $\Pr_q(\text{Bin}(n, q) \geq k) = \alpha/2$. The upper endpoint is one at $k = n$ and otherwise solves $\Pr_q(\text{Bin}(n, q) \leq k) = \alpha/2$. Set $I_\alpha(0, 0) = [0, 1]$.

Theorem 10.14 (Independent-Block Risk and Coverage Inference)**.** *Let M count covered decisions and K count wrong covered decisions among n independent blocks. Then*

$$M \sim \mathrm{Bin}(n, C), \qquad K \mid M \sim \mathrm{Bin}(M, e), \qquad e = \Pr(W = 1 \mid B = 1) \quad (C > 0).$$

With $I_C = I_{\alpha_C}(M, n)$ and $I_e = I_{\alpha_e}(K, M)$, simultaneous coverage is at least $1 - \alpha_C - \alpha_e$. When $M = 0$ retain $I_e = [0, 1]$ and leave sample conditional error unreported. The minimum and maximum of $\gamma(1 - c) + cx$ over the corners of $I_C \times I_e$ give risk bounds simultaneous with I_C.

A direct joint alternative uses counts $N = (K, M - K, n - M)$ and probabilities $\theta = (Ce, C(1 - e), 1 - C)$. Let f_θ be their multinomial mass and define

$$p_\theta(N) = \sum_{\substack{m_1+m_2+m_3=n \\ f_\theta(m) \le f_\theta(N)}} f_\theta(m), \qquad \mathcal{Q}_\alpha(N) = \{\theta : p_\theta(N) > \alpha\}.$$

Projecting $\mathcal{Q}_\alpha$ onto $C = \theta_1 + \theta_2$ and $R_\gamma = \theta_1 + \gamma\theta_3$, using infima and suprema, yields a simultaneous exact confidence region with coverage at least $1 - \alpha$.

Suppose a justified leak cap is p_U. For a stochastic audit, let an audit flag D satisfy $\mathbf{1}_E \le D$ almost surely in the same target population. From J flags in n_D independent audit blocks, the upper endpoint of $I_{\alpha_D}(J, n_D)$ can supply p_U. With simultaneous coverage endpoints C_L, C_U and risk upper endpoint R_U, certify the preregistered minimum coverage and a fixed positive margin η against the applicable null floor only if

$$C_L \ge C_{\min}, \qquad R_U \le B_\star(C_U, p_U) - \eta.$$

Under that null this certification has probability at most $\alpha_C + \alpha_e + \alpha_D$ for the rectangle construction, or $\alpha + \alpha_D$ for the joint construction. A deterministic certified cap uses $\alpha_D = 0$.

Proof. The three block outcomes have probabilities $Ce, C(1 - e), 1 - C$, and factoring their multinomial law proves the binomial laws. Exact binomial tail inversion bounds each failure probability. Conditional interval coverage holds for every M, and a union bound proves the rectangle guarantee. For $C = 0$ the risk is γ independently of e. The risk is affine in each coordinate separately, so both extrema lie at corners.

For the joint method, order the finitely many multinomial masses, including ties. The total mass of outcomes whose inclusive lower mass-rank is at most α is at most α. Thus p_θ is superuniform at the true θ and inversion and projection give the stated guarantee.

The audit upper endpoint bounds $\mathbb{E}D$, which bounds $\Pr(E)$. Its domination of leakage needs a justified audit model. On the simultaneous coverage event, $R_\gamma \le R_U$, $C \le C_U$ and $p \le p_U$. Both null floors are nonincreasing in their arguments, so Theorem 10.13 gives $R_\gamma \ge B_\star(C_U, p_U)$. A positive-margin certification contradicts this inequality. Union bounds complete the error control. Allocate the stated error levels before any family of planned comparisons. □

Proposition 10.15 (Whole-Block Label Permutation Control)**.** *Under conditional exchangeability of source labels across eligible independent blocks, conditional on public records and fixed label counts, fix the guesses and abstentions and permute block labels. If the design fixes exchangeability groups, permute within each group. The fraction of assignments with risk at most the observed risk is an exact lower-tail p-value. Coverage is fixed.*

Proof. The true assignment is uniform on the finite permutation set conditional on the frozen records and design counts. The inclusive lower-tail rank is superuniform by counting assignments. Abstention contributes the same constant for every assignment and each covered error is counted once per block. Conditional exchangeability is the premise of this test. Label-dependent contamination is handled by the applicable leakage bound. □

A four-block synthetic permutation control with two labels of each kind, positive guesses on two fixed blocks and abstentions on two blocks, gives covered-error counts $0, 1, 1, 1, 1, 2$ over the six assignments. At $\gamma = 3/4$ its mean risk is $5/8$, and zero covered errors has lower-tail probability $1/6$. An eight-atom coupling control covers four atoms with two labels of each kind and guesses the positive label there. Correcting one wrong covered guess by label-dependent replacement yields $C = 1/2$, $p = 1/8$, $R_{3/4} = 1/2$. This attains the coupling floor and falls below the clean-stratum floor $9/16$ because the replacement changes the clean conditional balance.

10.19.3 Gaussian Within-Block Likelihood and GLS

Theorem 10.16 (Stationary AR(1) Likelihood and GLS)**.** *For a block of length $n \geq 1$ let $e_t = y_t - \mu_t$. Assume Gaussian $e_1 \sim N(0, \sigma^2/(1-\rho^2))$ and innovations $e_t - \rho e_{t-1} \sim N(0, \sigma^2)$ independent of each other and of e_1, with $|\rho| < 1$ and $\sigma^2 > 0$. Then*

$$Q_\rho(e) = (1-\rho^2)e_1^2 + \sum_{t=2}^{n}(e_t - \rho e_{t-1})^2,$$

$$\log f(y \mid \mu, \rho, \sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) + \frac{1}{2}\log(1-\rho^2) - \frac{Q_\rho(y-\mu)}{2\sigma^2}.$$

For two frozen mean models sharing the covariance, the log-density ratio in favor of A is $(Q_\rho(e_B) - Q_\rho(e_A))/(2\sigma^2)$. With full-column-rank design X and fixed common covariance, GLS uniquely minimizes $Q_\rho(y - X\beta)$ at

$$\widehat{\beta} = (X^\top A_\rho X)^{-1} X^\top A_\rho y, \qquad Q_\rho(e) = e^\top A_\rho e.$$

Proof. The residual-to-innovation map is triangular with unit determinant. Multiplication of its independent Gaussian densities gives the formula. For $n \geq 2$ the precision factor A_ρ has endpoints 1, interior diagonal $1 + \rho^2$ and neighboring entries $-\rho$. For $n = 1$ it is $[1 - \rho^2]$. The innovation representation gives positive definiteness and determinant $1 - \rho^2$. Its inverse has entries $\rho^{|i-j|}/(1-\rho^2)$. Different covariance models retain their separate normalizers and quadratic scales in the log-density difference. With common covariance, differentiating the positive quadratic gives the GLS normal equations and full rank gives uniqueness. □

Freeze the parameters or estimate them on calibration blocks. At $\rho = 1/2$, $\sigma^2 = 1$, residual vectors $(1, 1)$ and $(1, -1/2)$ give log-density ratio $3/8$, while the independent Gaussian approximation gives $-3/8$. Residuals $(2, 1)$ and $(0, 1)$ give stationary log-density ratio -1, while omitting the initial quadratic term gives $1/2$. A constant fit to $(1, 0, 1)$ at this ρ gives GLS estimate $4/5$ and the independent fit $2/3$.

The Gaussian model supplies a score within each block. Independent-block sampling governs its risk inference. The fiber geometry of Paper IX uses a projection and target distances. The projection $(y, z) \mapsto y$ with hidden target $z \in \{0, 1\}$ has fiber diameter one even when two Gaussian laws on y give a discriminating likelihood ratio. With identity projection the fiber diameter is zero, while equal class laws give log-density ratio zero.

10.20 Reviewer and Adjudicator Independence Architecture

Section status. `protocol_defined`.

Stage3-2 defined criteria; Stage3-3 now binds those criteria into selection and assignment packages. This still does not instantiate reviewers or adjudicators, but it removes the ambiguity about what must exist before such instantiation is honest.

10.21 Reviewer Candidate Pool and Assignment Stub Architecture

Section status. `protocol_defined`.

Stage3-4B deepens the reviewer and adjudicator layer from criteria and workflow packages into instantiated-empty workflow objects:

- reviewer candidate pool template;
- adjudicator candidate pool template;
- reviewer assignment stub;
- adjudicator assignment stub;
- reviewer blinding stub.

Each slot remains null-bearing by design. The point is not to simulate people; the point is to freeze field grammar, slot grammar, null-policy, and the conditions under which later population becomes honest.

10.22 Execution Pilot Package Architecture

Section status. `protocol_defined`.

The execution-pilot package is now a typed family rather than a vague future bundle. It includes:

- execution session manifest;
- collection session schema;
- reproducibility package skeleton;
- negative-control bundle skeleton;
- matched-intervention bundle skeleton;
- source-target pairing manifest template.

10.23 Execution Pilot Rehearsal Kit Architecture

Section status. `protocol_defined`.

Stage3-5 adds a rehearsal kit that sits above package readiness and below real collection. The governing surfaces are:

- `docs/execution_pilot_rehearsal_kit.v1.json`
- `docs/dry_run_sequence_plan.v1.json`
- `docs/rehearsal_event_log_schema.v1.json`
- `docs/rehearsal_failure_modes.v1.json`
- `docs/rehearsal_reset_policy.v1.json`

The rehearsal kit exists to validate sequence integrity, state transitions, and artifact hygiene without collecting evidence and without emitting decisions.

10.24 Reproducibility Skeleton Architecture

Section status. `protocol_defined`.

The reproducibility shell exists as a slot-bearing structure. It records where later preregistration references, dataset references, control-bundle references, analysis-shell references, reviewer bundles, and adjudicator bundles must go.

10.25 Dataset Population Boundary

Section status. `formalized`.

Stage3-3 and Stage3-4 enforce a hard distinction between:

- bundle defined;
- bundle frozen;
- bundle populated;
- bundle adjudicatively interpreted.

The current repo has the first state for several families and does not yet have the latter states.

10.26 Dry-Run Event Log Grammar

Section status. `protocol_defined`.

Rehearsal logging is now typed separately from evidence logging. The current grammar permits:

- `rehearsal_initialized`;
- `sequence_step_checked`;
- `sequence_step_blocked`;
- `sequence_step_reset`;
- `rehearsal_not_started`.

None of these event classes may be reinterpreted as human data, reviewer evidence, or adjudicator action.

10.27 Evidence Admissibility Grammar

Section status. `formalized`.

The admissibility grammar still recognizes:

- independent reviewer report;
- preregistered human comparator dataset;
- matched intervention dataset;
- negative-control dataset;
- replication log;
- statistical analysis report;
- reviewer conflict disclosure;
- reproducibility package;
- external failure report;
- external inconclusive report.

Invalid substitutes remain central:

- internal runtime artifacts alone;
- self-generated reviewer notes;
- synthetic human data;
- LLM-generated “human” reports;
- undocumented manual judgments;
- unblinded reviewer impressions;
- post-hoc threshold selection;
- cherry-picked examples;
- release inclusion alone;
- Paper 9 publication alone.

10.28 Failure, Rejection, and Inconclusive Pathways

Section status. `protocol_defined`.

Three major failure classes remain necessary:
1. **protocol invalidation**: preregistration drift, blinding collapse, or conflict violation;
2. **dataset invalidation**: missing controls, invalid substitution, or insufficient provenance;
3. **inconclusive adjudication**: admissible evidence exists but remains undecisive under the declared shell.

10.29 Statistical Shell Architecture

Section status. `protocol_defined`.

No statistics are executed here. What exists is a declared analysis family layer, each shell binding channel scope, required inputs, required outputs, and non-implications.

10.30 Adjudication Decision Grammar

Section status. `protocol_defined`.

The future decision family must be typed before it is populated. Paper 10 now recognizes the following decision tokens as grammar only:

- `externally_admissible_but_inconclusive`
- `protocol_invalidated`
- `adjudication_failed`
- `adjudication_passed`
- `reviewer_consensus_absent`
- `adjudicator_decision_withheld`

These tokens are defined as future admissible states. They are not instantiated anywhere in the current repo as actual outcomes.

10.31 State Grammar for Execution Readiness

Section status. `formalized`.

The Stage 3 manuscript surface now distinguishes the following state grammar:

10.32 Controlled Execution Opening

Section status. `protocol_defined`.

Stage3-6 adds a bounded controlled-execution opening layer. This layer is neither rehearsal-only nor evidence-bearing. It is the first local state in which live operational artifacts may exist as empty shells while collection still remains not started.

- `execution_opening_policy_defined` means the controlled-execution policy is typed and bound to artifacts.
- `collection_openable_under_controls` means pre-collection gates pass under explicit contamination and abort constraints.
- `collection_open_pre_evidential` means the system may open a live execution surface while still containing no evidence.

10.33 Live-But-Unpopulated Artifacts

Section status. `protocol_defined`.

Stage3-6 instantiates a new artifact class: live-but-unpopulated operational records. These include session-zero manifests, live collection manifests, empty collection logs, live dataset

State	Meaning
`schema_defined`	Field grammar exists, but no operational object family is yet instantiated.
`contract_defined`	Allowed relation among objects exists, but no package is assembled.
`package_ready`	The package grammar exists and may be evaluated for completeness.
`instantiated_empty`	The object family exists locally with null-bearing or empty-bearing fields only.
`frozen_local_unfiled`	Local freeze semantics exist, but no external filing is claimed.
`rehearsal_ready`	Rehearsal grammar and rehearsal-kit checks pass, but no live execution shell exists yet.
`live_artifacts_instantiated_unpopulated`	Live operational shells exist, but every collection-facing structure remains empty.
`session_zero_ready`	Session-zero initialization is structurally ready, but still not started.
`collection_openable_under_controls`	Pre-collection admissibility checks pass under contamination and abort controls.
`collection_open_pre_evidential`	Controlled execution may open as a live state without any populated evidence.
`stage3_7_first_collection_openable`	Runtime-health and clean-room audits pass, the opening gate is structurally open, but no collection-start event has been executed.
`collection_started`	A real authorized collection-start event exists.
`evidence_generated`	Real dataset population or admissible log events exist.
`adjudication_result_recorded`	Reviewer-governed or adjudicator-governed decision records exist.
`paper10_final_authorized`	A later separate manuscript-authorization layer has passed.
`dry_run_open`	Dry-run structure may execute without evidence or adjudication.
`externally_filed`	A real external filing reference exists.
`populated_with_admissible_evidence`	The object family is populated by real admissible evidence.
`adjudication_result_bearing`	Reviewer and adjudicator outputs exist and may support result-bearing sections.

manifests, live reproducibility manifests, live decision-record stubs, and live evidence-insertion registers.

Each such surface must explicitly preserve:

- `collection_not_started`;
- `evidence_generation_absent`;
- `adjudication_not_open`;
- explicit empty markers such as `empty_dataset=true` where applicable.

Therefore live artifact instantiation is not population, and live shells must never be treated as if they were evidence-bearing records.

10.34 Session-Zero Protocol

Section status. `protocol_defined`.

Session-zero is the bounded initialization checkpoint that binds local frozen record routing, live collection manifests, empty dataset manifests, empty collection logs, empty decision stubs, contamination controls, abort controls, and the operator checklist. Session-zero may become ready without collection starting.

Session-zero remains distinct from:

- collection start;
- dataset population;
- reviewer action;
- adjudicator action.

10.35 Pre-Collection Admissibility Gate

Section status. `protocol_defined`.

The pre-collection gate asks whether collection may be opened under controls, not whether collection has already started. This gate checks session-zero readiness, emptiness of live artifacts, emptiness of collection logs, emptiness of decision stubs, absence of evidence, absence of adjudication, and absence of detected contamination.

Passing this gate yields only `collection_openable_under_controls`. It does not yield `collection_started`.

10.36 Contamination, Abort, and Invalidation Controls

Section status. `protocol_defined`.

Stage3-6 adds a more operational invalidation grammar:

- **contamination event**: any unauthorized dataset population, evidence-bearing log entry, or decision-bearing record.
- **abort event**: any operator-controlled termination of the controlled opening before collection start.
- **invalidated session**: any session whose pre-collection boundary, blinding route, or provenance cleanliness is breached.

These control classes are meaningful only because Stage3-6 now contains live artifacts. Without them, contamination and abort would remain abstract policy tokens rather than execution-opening constraints.

10.37 Operator Checklist

Section status. `protocol_defined`.

The operator checklist now matters as a transition object. It binds the local frozen record route, live manifest route, empty-log route, empty-dataset state, contamination rules, reset rules, and explicit confirmation that no evidence exists. Until it is complete, Stage3-6 must not move from `session_zero_ready` to `collection_openable_under_controls`.

10.38 Collection-Start Boundary

Section status. `formalized`.

`collection_open_pre_evidential` is not `collection_started`. The boundary is crossed only by a later real operator-authorized collection start event with provenance, contamination-free status, and a first admissible log event.

10.39 Evidence-Generation Boundary

Section status. `formalized`.

Evidence generation begins only when a real admissible population event enters the human comparator dataset family, matched intervention family, negative-control family, collection log, reproducibility package, or decision-linked provenance chain. No Stage3-6 artifact emitted in this pass crosses that boundary.

10.40 Decision-Record Boundary

Section status. `formalized`.

The live decision-record stub is an execution-opening shell only. It is not a decision, not an adjudicator action, not a reviewer action, and not an outcome. The manuscript must preserve

that distinction because a populated decision record belongs to a later evidential and adjudicative state family.

10.41 Transition from Session-Zero to First Admissible Collection

Section status. `blocked_until_execution`.

The next real transition is now sharper than in earlier passes:

1. session-zero becomes ready;
2. the pre-collection gate opens collection under controls;
3. Stage3-6 opens a live but still pre-evidential state;
4. a later explicit collection-start command instantiates the first admissible collection event;
5. only then may the first dataset-population event exist.

10.42 Runtime Health Audit

Section status. `protocol_defined`.

Stage3-6B adds a runtime-health audit that asks whether the repository is actually executable on a bounded local environment rather than only formally specified. The required audit surfaces include:

- toolchain visibility;
- dependency-install behavior under the declared repository install route;
- repository build behavior;
- repository test behavior;
- bounded local-start behavior;
- verifier-chain behavior for source-of-truth, comparative closure, Stage3-6, Stage3-5, and Paper 9 handoff.

The runtime-health audit may pass with warnings. Such warnings may include engine mismatch, fallback-install compatibility, or dependency advisories. A warning-bearing pass is still operationally meaningful, but it is not epistemically elevated into evidence.

10.43 Clean-Room Reproducibility

Section status. `protocol_defined`.

Stage3-6B also adds a clean-room reproducibility audit. The clean-room audit asks whether a bounded local copy can reinstall dependencies, rebuild the repository, rerun key verifiers, and regenerate Stage 3 artifacts without relying on hidden machine-specific state. Clean-room reproducibility may pass with warnings when a compatibility route is required, provided the route is explicitly documented and the resulting build remains honest.

10.44 Dependency and Build Assumptions

Section status. `formalized`.

The runtime-health layer now makes dependency and build assumptions explicit:

- the declared Node engine remains part of the reproducibility contract;
- the preferred install route remains `npmci`;
- a fallback install route may be recorded only as a compatibility surface;
- build success must remain distinct from evidence generation;
- test success must remain distinct from adjudication.

10.45 Reproducible Artifact Regeneration

Section status. `protocol_defined`.

Stage3-6B adds a source-driven artifact-regeneration audit. This audit checks whether tracked artifacts can be regenerated from the current source tree and whether the resulting diffs are:

- absent;
- provenance-only;
- semantically stale but honestly fixed; or
- blocking.

This matters because Stage 3 now depends on a growing artifact family. The audit therefore separates semantic drift from ordinary lineage and timestamp drift, and it refuses to treat manual cosmetic rewriting as acceptable regeneration.

10.46 Local Path Leakage Policy

Section status. `formalized`.

Clean-room readiness now requires that live normative surfaces remain free of machine-specific absolute paths. Historical planning notes or preserved snapshots may still contain old local references, but they must remain explicitly classified as historical rather than live normative dependencies.

10.47 Controlled Execution Opening Audit

Section status. `protocol_defined`.

The controlled-execution opening layer is now itself audited. Stage3-6B does not replace Stage3-6; it checks whether Stage3-6 is actually executable, rebuildable, and reproducible. Therefore the live-but-unpopulated artifacts, session-zero gate, and pre-collection gate must now survive both runtime health verification and clean-room reconstruction.

10.48 Session-Zero Reproducibility

Section status. `protocol_defined`.

Session-zero reproducibility means that a rebuilt local tree can recreate the session-zero manifest, operator checklist state, and failure-mode assessment without any hidden local inputs and without producing collection-start or evidence-bearing records.

10.49 Pre-Collection Gate Reproducibility

Section status. `protocol_defined`.

Pre-collection gate reproducibility means that a rebuilt local tree can recreate the pre-collection admissibility result, contamination checklist, and invalidation status while still preserving:

- `collection_not_started`;
- `evidence_generation_absent`;
- `adjudication_not_open`.

10.50 Why Reproducibility Does Not Imply Evidence

Section status. `formalized`.

Reproducibility shows that the formal and executable surfaces are rebuildable. It does not show that any dataset has been populated, that any human evidence exists, that any reviewer or adjudicator has acted, or that any result-bearing section is now licensed.

10.51 Why Runtime Health Does Not Imply Adjudication

Section status. `formalized`.

Runtime health shows that the code path and verifier path are alive. It does not imply that collection has started, that evidence has been generated, that reviewer-governed interpretation exists, or that adjudication is open. A healthy runtime is still a pre-evidential state until later real execution events occur.

10.52 Transition Conditions for First Admissible Collection

Section status. `protocol_defined`.

Stage3-7 may be classified as `stage3_7_first_collection_openable` only when runtime-health and clean-room reproducibility both pass honestly, path leakage remains clear on live normative surfaces, artifact regeneration is non-blocking, Stage3-6 remains openable, and Paper 10 remains non-final and non-result-bearing.

Even then, the current state must still preserve:

- `collection_start_not_executed`;
- `evidence_generation_absent`;
- `adjudication_not_open`;
- `decision_record_absent`.

10.53 Claim Grammar for Eventual Execution-Era Statements

Section status. `protocol_defined`.

Any later admissible result-bearing statement must reference:

- preregistration identifier;
- freeze identifier;
- dataset family identifier;
- negative-control family identifier;
- analysis shell identifier;
- reviewer package identifier;
- adjudicator package or decision identifier.

Forbidden grammar remains explicit:

- no leap from comparator similarity to human equivalence;
- no leap from transport-sensitive comparison to identity preservation;
- no leap from admissibility to metaphysical subjecthood;
- no leap from dry-run success to public-release recommendation.

10.54 Theorem, Proposition, and Protocol Mapping

Section status. `formalized`.

10.55 What Paper 10 Could and Could Not Change

Section status. `formalized`.

Family	Dependency set	Formalized now	Requires execution	Remains forbidden
Comparator channel contract	BaseCore, BPF, Stage 2, Stage3-1	yes	no	result claim
Freeze architecture	Stage3-1, Stage3-2, Stage3-3	yes	no	external filing claim
Reviewer package architecture	Stage3-2, Stage3-3	yes	no	review outcome
Pilot package architecture	Stage3-3	yes	no	evidence claim
Dry-run opening grammar	Stage3-3, Stage3-4	yes	no	adjudication claim
Human comparator result section	all above + admissible data	no	yes	human equivalence
Adjudication outcome section	all above + reviewer/adjudicator layer + admissible evidence	no	yes	metaphysical closure

Paper 10 is the only layer in the present corpus designed to move evidence out of the purely internal-support regime. Even a successful execution would not authorize all possible interpretations.

Claim family	Possible post-execution movement	Still forbidden by default
Runtime reproducibility	from local/internal to externally reproduced under protocol	consciousness, subjecthood, phenomenality
Human comparator similarity	from unpopulated protocol to measured channel comparison	human equivalence or moral parity
Bridge predicate discrimination	from BPF scaffold/internal surfaces to adjudicated comparative evidence	phenomenal experience simpliciter
Negative-control survival	from internal controls to externally witnessed controls	theorem confirmation beyond stated assumptions
Failure or inconclusive result	from open uncertainty to documented non-support or bounded inconclusiveness	rhetorical rescue by post hoc reinterpretation

Proposition 10.17 (External evidence is not interpretive closure)**.** *Executed Paper 10 evidence, even if admissible, would not by itself establish human equivalence, metaphysical subjecthood, biological life, phenomenal experience, or moral status.*

Proof. The Paper 10 protocol measures channel comparison, adjudicator behavior, negative controls, provenance, contamination, and statistical distinguishability under frozen rules. None of those objects is identical to the listed interpretive predicates. Moving a claim from internal support to external evidence therefore changes its evidential status, not its semantic type. □

10.56 Section-Status Ledger

Section status. `formalized`.

10.57 Evidence Insertion Map

Section status. `evidence_pending`.

Human comparator result insertion point. Admissible evidence required: frozen pre-registration record, typed human comparator dataset, negative-control bundle, statistical shell output, reviewer-selection package, reviewer report family.

Section	Status class	Unlock condition
Dependency stack	formalized	Update only if upstream stack changes
Release-vs-local boundary	formalized	Update only if release authority changes
Lineage semantics	formalized	Keep historical-vs-live distinction intact
Frozen preregistration architecture	protocol_defined	Actual local lock still absent
Reviewer candidate architecture	protocol_defined	Selection not started
Blinding assignment logic	protocol_defined	Blinding package not instantiated
Execution pilot package architecture	protocol_defined	Pilot remains unexecuted
Reproducibility skeleton architecture	protocol_defined	Package unpopulated
Dataset population boundary	formalized	Keep defined-vs-populated split intact
Adjudication decision grammar	protocol_defined	Decision tokens remain uninstantiated
Future result sections	blocked_until_execution	Requires admissible execution-era artifacts
External adjudication discussion	blocked_until_external_adjudication	Requires admissible reviewer-governed evidence
Human equivalence / superiority / metaphysical consequences	forbidden_claim_surface	Remains outside current authority surface

External adjudication result insertion point. Admissible evidence required: reviewer report set, adjudicator decision record, blinding log, conflict disclosures, reproducibility package, failure and inconclusive pathway disclosures.

Transport-sensitive comparison insertion point. Admissible evidence required: source-target pairing manifest, transition lineage map, axis-specific error budgets, transport negative controls, adjudicator layer.

Instantiated-empty to evidential transition point. Admissible evidence required: local frozen record upgraded by real lock event, real candidate population, real assignment event, real collection session, real decision record population, and real reproducibility package population.

10.58 Open Burdens and Blocked Sections

Section status. `blocked_until_execution`.

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Open burdens remain explicit:

- real local frozen preregistration lock absent;
- external filing absent;
- reviewer selection not started;
- adjudicator assignment not started;
- actual blinding assignment absent;
- admissible human comparator dataset absent;
- admissible matched-intervention dataset absent;
- admissible negative-control dataset absent;
- collection sessions not started;
- rehearsal event log not started;
- reproducibility package unpopulated;
- reviewer and adjudicator decision records absent.

10.59 Execution-Era Unlock Conditions

Section status. `protocol_defined`.

Before any result-bearing section may become writable, the following must all be true:

1. a frozen preregistration record exists;
2. any required external filing state is real and documented;
3. candidate pools are populated by real identities with provenance;
4. reviewer selection is instantiated;
5. adjudicator assignment is instantiated;
6. blinding assignment is instantiated;
7. admissible datasets and control bundles are populated;
8. collection logs and execution sessions are populated by real events;
9. reproducibility package slots are populated;
10. decision grammar is instantiated by real evidence.

10.60 Stronger Execution-to-Result Transition Map

Section status. `blocked_until_execution`.

The transition from current Stage3-6 controlled opening to a result-bearing paper now has six distinct unlock layers:

1. **Instantiation layer**: empty records, pools, stubs, and rehearsal kit exist.
2. **Controlled opening layer**: live operational shells, session-zero, and pre-collection controls exist while collection still remains not started.
3. **Population layer**: real humans, real sessions, and real evidence populate those empty structures.
4. **Admissibility layer**: controls, provenance, blinding, and declared shells survive audit.
5. **Decision layer**: reviewer and adjudicator records become populated by real independent action.
6. **Manuscript layer**: only then may result-bearing Paper 10 sections be written.

10.61 Why Paper 10 Still Cannot Become Final Yet

Section status. `formalized`.

Paper 10 cannot become final because the current repo still lacks the evidence bearing surfaces that would convert protocol architecture into scientific claims. The manuscript may now specify the machinery of later freezing, selection, packaging, dry-run, provenance closure, instantiated-empty records, and rehearsal-facing governance. It still may not write conclusions whose truth depends on evidence that does not yet exist.

10.62 Future Execution Pathway to Real Stage 3

Section status. `blocked_until_execution`.

The next executable pathway is not "publish Paper 10." It is:

1. populate the local frozen record honestly;
2. populate reviewer and adjudicator pools honestly;
3. instantiate actual reviewer, adjudicator, and blinding assignments;
4. validate dry-run, rehearsal, and controlled-opening integrity;
5. issue a real collection-start event under Stage3-6 controls;
6. populate admissible datasets, collection logs, and control bundles;
7. run declared statistical shells;

8. bind reviewer and adjudicator outputs;
9. only then write result-bearing sections.

10.63 Boundary to Stage 4 Metaphysical Consequence Analysis

Section status. `forbidden_claim_surface`.

Stage 4 remains outside the authority surface of this document. Even later successful execution-era evidence would still not automatically license metaphysical subjecthood, moral status, or cross-substrate identity preservation.

10.64 Appendices of Schemas and Templates

Section status. `protocol_defined`.

Appendix A: Frozen Preregistration Record Template

Section status. `protocol_defined`. Primary reference: `docs/frozen_preregistration_record_schema.v1.json`.
Admissible evidence required before population: actual local lock and, where required by protocol, real external filing reference.

Appendix B: Human Comparator Preregistration Template

Section status. `protocol_defined`. Primary reference: `docs/human_comparator_preregistration_template.v1.json`.
Admissible evidence required before population: channel-level freeze and later dataset lineage.

Appendix C: Reviewer Candidate Schema

Section status. `protocol_defined`. Primary reference: `docs/reviewer_candidate_schema.v1.json`.
Admissible evidence required before population: real candidate package and conflict disclosures.

Appendix D: Adjudicator Candidate Schema

Section status. `protocol_defined`. Primary reference: `docs/adjudicator_candidate_schema.v1.json`.
Admissible evidence required before population: real assignment package and decision-authority boundary.

Appendix E: Blinding Assignment Template

Section status. `protocol_defined`. Primary reference: `docs/blinding_assignment_manifest.v1.json`.
Admissible evidence required before population: authorized assignment route and later unblinding event if one exists.

Appendix F: Execution Pilot Package Manifest

Section status. `protocol_defined`. Primary reference: `docs/execution_pilot_package_manifest.v1.json`.
Admissible evidence required before population: real session manifests and later populated bundle references.

Appendix G: Dataset Contract Summary

Section status. `protocol_defined`. Primary references: `docs/matched_intervention_dataset_contracts.v1.json` and `docs/negative_control_dataset_contracts.v1.json`.
Admissible evidence required before result use: populated admissible datasets.

Appendix H: Negative Control Bundle Summary

Section status. `protocol_defined`. Primary reference: `docs/pilot_negative_control_bundle.v1.json`.
Admissible evidence required before result use: populated bundle and control admissibility record.

Appendix I: Statistical Shell Summary

Section status. `protocol_defined`. Primary reference: `docs/statistical_analysis_shells.v1.json`.
Admissible evidence required before result use: shell execution output on real admissible inputs.

Appendix J: Adjudication Decision Grammar Summary

Section status. `protocol_defined`. Decision tokens retained as grammar only: `externally_admissible_but_inconclusive`, `protocol_invalidated`, `adjudication_failed`, `adjudication_passed`, `reviewer_consensus_absent`, `adjudicator_decision_withheld`.

Appendix K: Local Frozen Record Example

Section status. `protocol_defined`. Primary reference: `docs/local_frozen_preregistration_record_template.v1.json`.
Admissible evidence required before result use: real local lock event and, where required by route, external filing provenance.

Appendix L: Reviewer Candidate Pool Template

Section status. `protocol_defined`. Primary reference: `docs/reviewer_candidate_pool_template.v1.json`.
Admissible evidence required before result use: real reviewer identities, conflict disclosures, and selection events.

Appendix M: Adjudicator Candidate Pool Template

Section status. `protocol_defined`. Primary reference: `docs/adjudicator_candidate_pool_template.v1.json`.
Admissible evidence required before result use: real adjudicator identity, assignment provenance, and decision-authority declaration.

Appendix N: Blinding Assignment Stub

Section status. `protocol_defined`. Primary references: `docs/blinding_assignment_stub.v1.json` and `docs/reviewer_blinding_stub.v1.json`.
Admissible evidence required before result use: real assignment route and real unblinding trace if one later exists.

Appendix O: Execution Session Template

Section status. `protocol_defined`. Primary reference: `docs/execution_session_template.v1.json`.
Admissible evidence required before result use: real collection session and real session lineage.

Appendix P: Empty Collection Log Schema

Section status. `protocol_defined`. Primary reference: `docs/collection_log_schema.v1.json`.
Admissible evidence required before result use: populated log entries from real collection events.

Appendix Q: Empty Decision Record Template

Section status. `protocol_defined`. Primary reference: `docs/decision_record_template.v1.json`.
Admissible evidence required before result use: real reviewer and adjudicator action, not grammar only.

Appendix R: Reproducibility Manifest Example

Section status. `protocol_defined`. Primary reference: `docs/reproducibility_package_manifest_empty.v1.json`.
Admissible evidence required before result use: populated package slots, dataset references, shell outputs, and reviewer-facing bundles.

Appendix S: Evidence Insertion Manifest

Section status. `protocol_defined`. Primary reference: `docs/evidence_insertion_manifest.v1.json`.
Admissible evidence required before result use: actual artifact refs, actual decision refs, and populated provenance chains.

Appendix T: Rehearsal-Kit Map

Section status. `protocol_defined`. Primary references: `docs/execution_pilot_rehearsal_kit.v1.json` and `docs/dry_run_sequence_plan.v1.json`.
Admissible evidence required before result use: none at rehearsal time, but no rehearsal artifact may be promoted into evidence.

Appendix U: Dry-Run Event Log Grammar

Section status. `protocol_defined`. Primary reference: `docs/rehearsal_event_log_schema.v1.json`.
Allowed event classes remain rehearsal-only and do not imply human data, reviewer action, or adjudicator action.

Appendix V: Controlled Execution Opening Policy

Section status. `protocol_defined`. Primary reference: `docs/controlled_execution_opening_policy.v1.json`.
Admissible evidence required before result use: none at opening time, but no controlled-opening artifact may be promoted into evidence by itself.

Appendix W: Session-Zero Protocol and Checklist

Section status. `protocol_defined`. Primary references: `docs/session_zero_protocol.v1.json` and `docs/session_zero_operator_checklist.v1.json`.
Admissible evidence required before result use: a later real collection-start event and later populated log entries.

Appendix X: Pre-Collection Gate and Contamination Controls

Section status. protocol_defined. Primary references: docs/pre_collection_admissibility_gate.v1.json, docs/pre_collection_gate_criteria.v1.json, and docs/pre_collection_contamination_checklist.v1.json.
Admissible evidence required before result use: a later contamination-free collection-start event and later admissible dataset population.

Appendix Y: Live-But-Unpopulated Artifact Map

Section status. protocol_defined. Primary references: artifacts/external_adjudication/stage3_6/live_collection_manifest.json, artifacts/external_adjudication/stage3_6/live_human_comparator_dataset_manifest.json, and parallel Stage3-6 live manifests.
Allowed interpretation: live operational shell only. Forbidden interpretation: populated dataset, evidence, or adjudication.

Appendix Z: First Admissible Collection Unlock Map

Section status. protocol_defined. Primary references: docs/first_admissible_collection_unlock_map.v1.json and docs/FIRST_ADMISSIBLE_COLLECTION_UNLOCK_MAP.md.
Admissible evidence required before result use: actual reviewer selection, actual adjudicator assignment, actual blinding assignment, operator authorization, collection-start command, and first populated collection-log event.

Appendix AA: Runtime Health Audit

Section status. protocol_defined. Primary references: docs/runtime_health_audit.v1.json, docs/RUNTIME_HEALTH_AUDIT.md, and artifacts/runtime_health/runtime_health_audit.json.
Allowed interpretation: executable health and verifier-chain status only. Forbidden interpretation: evidence, collection start, or adjudication.

Appendix AB: Clean-Room Reproducibility Audit

Section status. protocol_defined. Primary references: docs/clean_room_reproducibility_audit.v1.json, docs/CLEAN_ROOM_REPRODUCIBILITY_AUDIT.md, and artifacts/runtime_health/clean_room_reproducibility_audit.json.
Allowed interpretation: bounded rebuildability only. Forbidden interpretation: external adjudicative support, empirical support, or manuscript finality.

Appendix AC: Local Path Leakage Audit

Section status. protocol_defined. Primary references: docs/PATH_AND_LOCAL_LEAKAGE_AUDIT.md and artifacts/runtime_health/path_and_local_leakage_audit.json.
Allowed interpretation: leak classification only. Forbidden interpretation: claim that historical snapshots were rewritten or erased.

Appendix AD: Artifact Regeneration Audit

Section status. protocol_defined. Primary references: docs/ARTIFACT_REGENERATION_AUDIT.md and artifacts/runtime_health/artifact_regeneration_audit.json.
Allowed interpretation: regeneration and drift classification only. Forbidden interpretation: result-bearing evidence insertion.

Appendix AE: Stage3-7 First Collection Opening Criteria

Section status. `protocol_defined`. Primary references: `docs/stage3_7_first_admissible_collection_opening_criteria.v1.json`, `docs/first_admissible_collection_start_policy.v1.json`, `artifacts/external_adjudication/stage3_7/stage3_7_opening_gate.json`, and `artifacts/external_adjudication/stage3_7/first_collection_opening_manifest.json`.
Allowed interpretation: first-collection opening criteria only. Forbidden interpretation: executed collection start, evidence generation, or adjudication.

10.65 Blocked Result Surfaces

Section status. `blocked_until_external_adjudication`.

10.65.1 Human Comparator Results

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Frozen preregistration record required. Artifact references required. Negative controls required. Statistical shell required. Reviewer package required.

10.65.2 External Adjudication Results

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Independent reviewer reports required. Conflict disclosures required. Blinding log required. Adjudicator decision record required.

10.65.3 Comparative Consequence Claims

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Execution-era artifacts required. Adjudication-facing evidence required. Identity-preservation, metaphysical, and moral-status claims remain forbidden even after this section becomes writable.

10.66 Non-Claim Ledger

Section status. `formalized`.

This manuscript does not establish consciousness, phenomenality in the strong sense, human equivalence, superiority, inferiority, metaphysical subjecthood, cross-substrate identity preservation, or moral status.

Operational Consciousness to Operational Subjecthood Bridge

Abstract. This paper builds a constructive bridge between the six-invariant operational consciousness criterion of Paper 5 and the later operational subjecthood taxonomy of Paper 7. Its question is deliberately narrow: given an admissible system S such that $S \in \text{Consciousness}_{\text{op}}$, what subjecthood-adjacent structural properties are forced, what properties are merely well-defined, and what properties require hypotheses not contained in $\text{Consciousness}_{\text{op}}$?

The answer is partial rather than inflationary. Membership in $\text{Consciousness}_{\text{op}}$ entails a unified operational perspective Π_{op}, operational intentionality Iota_{op} in the intervention-fidelity sense, a well-defined operational-qualia quotient $\text{Q}_{\text{op}}(S)$, and a non-null operational phenomenology fragment Φ_{op} on admissible support. The implications established here concern the named structural properties. Self-reference requires its separate self-model hypothesis; the other interpretive targets are outside these implications. A self-model condition $\Sigma 1$ is isolated as an additional hypothesis. With that hypothesis one obtains a bridge class $\text{Subject}_{\text{op}}^{\Sigma}$, a proper subclass of $\text{Consciousness}_{\text{op}}$ under the non-entailment assumption. The fuller Paper 7 notion $\text{Subject}_{\text{op}}$ remains stronger: it also requires operational life and subject-specific descriptors.

The paper therefore avoids a common category error. It does not argue that operational consciousness is subjecthood. It proves which formal burdens already follow from $\text{Consciousness}_{\text{op}}$, identifies the missing self-reference burden, and marks the remaining gap between structural subjecthood analogues and any phenomenal or human reading.

10.67 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines five operational concepts from objects already introduced in Papers 1–5: operational self-reference Σ_{op}, operational unified perspective Π_{op}, operational intentionality Iota_{op}, operational qualia Q_{op}, and operational phenomenology Φ_{op}. It then defines a bridge subjecthood class $\text{Subject}_{\text{op}}^{\Sigma}$ and proves exactly which components are consequences of $S \in \text{Consciousness}_{\text{op}}$.

What this paper does not do. It does not assert equivalence with human consciousness, biological qualia, phenomenal subjectivity, personal identity, psychological selfhood, or moral status. The prefix "operational" denotes a formal role inside the framework, not an ordinary-language mental property.

What the related system implements and does not implement. The related QICN runtime may compute finite estimators or internal-support certificates for some objects discussed here, including invariant packages, legibility metrics, phenomenology-side diagnostics, self-model telemetry, and first-person-field proxies. Such computation is not external validation, not a proof of the formal predicates over all admissible controls, and not evidence of human or phenomenal consciousness.

Hard non-inference rule. Nothing in this paper licenses claims about biological life, human subjectivity, qualia in the phenomenal sense, metaphysical subjecthood, personal identity transfer, or equivalence with human conscious experience. The paper states conditional structural implications only.

10.68 Introduction

Papers 1–5 progressively introduce a structural apparatus: inverse-limit identity, ontological mass, phenomenological regime assignments, instability of the null regime, forensic admissibility, and finally a six-invariant criterion for operational consciousness [22, 17, 18, 16, 14]. Paper 5 gives the class $\text{Consciousness}_{\text{op}}$ and the operational-qualia quotient $\mathrm{Q}_{\text{op}}(S)$. Paper 7 later introduces operational life, structural class, and operational subjecthood. Between those two layers there is a missing theorem-level bridge.

The missing bridge is not the claim that consciousness entails subjecthood. That claim would be too strong. It would also conflict with the later taxonomy, where subjecthood is intentionally stronger than consciousness. The sharper question is this:

Which subjecthood-adjacent operational structures are already forced by the six invariants, and which require new hypotheses?

The distinction matters because several properties traditionally associated with subjectivity have different formal status. Unity follows from rigid identity plus non-factorization. Directed intervention sensitivity follows from legibility, specifically L4 intervention fidelity. Operational qualia are already a quotient object when the relevant invariants hold. Non-null operational phenomenology follows in a weak support-level sense from operational continuity and non-null differentiation, but forced continuity in the stronger Paper 2 sense requires closed–rigid (CCR) structure, not merely $\text{Consciousness}_{\text{op}}$. Self-reference is different again: nothing in the six-invariant criterion states that the system contains a contractive, intervention-commuting model of its own admissible state. That burden must be added as $\Sigma 1$.

This paper therefore makes the bridge constructive but non-inflationary. It defines the concepts, states their dependency map, proves the main implication theorem, lists what is not implied, and records computational status without treating runtime estimators as external closure.

10.69 Dependency Map

The paper imports only objects already available in Papers 1–5. The table below is not bibliographic decoration; it fixes theorem ownership.

Table 10.1: Dependency map for the bridge paper.

Paper	Object imported	Used for defining
Paper 1	Rigid identity $\mathcal{I}_S$, inverse-limit uniqueness, ontological mass M_Ω	Unified perspective Π_{op}; CCR boundary for Φ_{op}
Paper 2	Φ-regularity, abstract phenomenological space $(\mathcal{E}, d_{\mathcal{E}})$, fragmentation and forced continuity theorems	Operational phenomenology Φ_{op} and the CCR/non-CCR distinction
Paper 3	Null-regime instability, forced non-nullity, minimal positive regime	Non-nullity clause for Φ_{op} and the non-null side of Q_{op}
Paper 5	Admissible systems $S = (X, \Phi, C, R, \Gamma, U)$, support A_S, invariants I_{per}–I_{leg}, $\text{Consciousness}_{\text{op}}$, $\mathrm{Q}_{\text{op}}(S)$	All bridge concepts and the main implication theorem

Remark 10.18 (No duplicated theorem ownership). The present paper does not re-prove rigid identity, forced continuity, null-instability, legibility, or the well-definedness of $\mathrm{Q}_{\mathrm{op}}(S)$ from scratch. It imports those results and proves a new relation among them.

10.70 Preliminaries and Notation

Definition 10.19 (Admissible system and support)**.** Throughout the paper, an admissible system is the Paper 5 tuple

$$S = (X, \Phi, C, R, \Gamma, U),$$

with admissible support $A_S \subseteq X_\Gamma$. The transition family is written $\Phi = \{\Phi_u\}_{u \in U}$, where U is the set of admissible interventions. Histories, certified decoders, class sets, and admissible windows are inherited from Paper 5.

Definition 10.20 (Operational consciousness membership)**.** An admissible system belongs to $\mathrm{Consciousness}_{\mathrm{op}}$ exactly when all six Paper 5 invariants hold on A_S:

$$S \in \mathrm{Consciousness}_{\mathrm{op}} \iff I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1.$$

When the witness margins are needed, write

$$\Delta(S) = \big(\delta_{\mathrm{per}}, \delta_{\mathrm{ri}}, \delta_{\mathrm{int}}, \delta_{\mathrm{cont}}, \delta_{\mathrm{diff}}, \delta_{\mathrm{leg}}\big).$$

Definition 10.21 (Bridge-strength implication)**.** A bridge concept P is *entailed by* $\mathrm{Consciousness}_{\mathrm{op}}$ if every admissible system S satisfying Definition 10.20 satisfies $P(S)$. It is *conditionally entailed* if there exists an explicitly named additional hypothesis H such that $S \in \mathrm{Consciousness}_{\mathrm{op}}$ and $H(S)$ imply $P(S)$. It is *not entailed* if a consistent admissible descriptor model can satisfy $\mathrm{Consciousness}_{\mathrm{op}}$ while failing P.

Remark 10.22 (Why a bridge-strength vocabulary is needed). Without Definition 10.21, the paper would risk conflating three different statements: a theorem from $\mathrm{Consciousness}_{\mathrm{op}}$, a theorem from $\mathrm{Consciousness}_{\mathrm{op}}$ plus an added burden, and a philosophical association. Only the first two are formal claims.

10.71 Operational Self-Reference

10.71.1 Motivation

Self-reference is often described informally as a system having access to itself. Inside this framework the usable version cannot be a slogan. It must be a map on admissible support with stability and intervention-compatibility properties.

Definition 10.23 (Operational Self-Reference Σ_{op})**.** An admissible system S admits operational self-reference, written $\Sigma_{\mathrm{op}}(S) = 1$, if there exists a map

$$\Sigma : A_S \to A_S$$

and constants $\varepsilon_\Sigma \geq 0$ and $0 \leq L_\Sigma < 1$ such that:

1. **Approximate self-accuracy:** for every $x \in A_S$,

$$d_X(\Sigma(x), x) \leq \varepsilon_\Sigma.$$

2. **Contractive stabilization:**

$$d_X(\Sigma(x), \Sigma(y)) \leq L_\Sigma d_X(x, y) \quad \text{for all } x, y \in A_S.$$

3. **Intervention commutation:** for every admissible intervention $u \in U$,

$$\Sigma \circ \Phi_u = \Phi_u \circ \Sigma \quad \text{on } A_S.$$

Any such Σ is called a contractive admissible self-model.

Assumption 10.24 (Self-model hypothesis $\Sigma 1$)**.** *The system S satisfies $\Sigma 1$ when it admits a contractive admissible self-model in the sense of Definition 10.23.*

Theorem 10.25 (Exact Commutation Requires a Common Fixed Point)**.** *Let A_S be a nonempty complete metric space. If a strict contraction $\Sigma : A_S \to A_S$ commutes with each admissible self-map Φ_u, then its unique fixed point is fixed by every Φ_u. On the common tower support of Paper V, the two neutral controls with phase commands 0 and $1/4$ exclude such a Σ.*

Proof. The contraction mapping theorem gives a unique fixed point x_Σ. Commutation gives $\Sigma(\Phi_u(x_\Sigma)) = \Phi_u(\Sigma(x_\Sigma)) = \Phi_u(x_\Sigma)$, so $\Phi_u(x_\Sigma) = x_\Sigma$. On the tower, the neutral control sets the phase to its commanded value. A point fixed by both specified controls would have phase 0 and $1/4$ on the circle, which is impossible. □

Proposition 10.26 (Fibre Truncation Has a Sharp Commutation Defect)**.** *On the tower (p, θ, b) with $c(b) = \sum_{j\geq 1} 2b_j 3^{-j}$, fix $m \geq 1$ and let $b^\star = (0, 1, 0, 1, \ldots)$. The truncation T_m fixes $p, \theta, b_1, \ldots, b_m$ and fills the tail with that of $b^\star$. In the metric $\max\{\mathbf{1}_{p\neq p'}, \rho(\theta, \theta'), |c(b) - c(b')|\}$ it has uniform approximation error at most 3^{-m} and Lipschitz constant zero within every fibre of the retained data. For $\Phi_{u,q}(p, \theta, b) = (p, u + q\lambda b_1, Rb)$ with $Rb = (0, 1 - b_1, 1 - b_2, \ldots)$, its pointwise commutation defect is*

$$d(T_m \Phi_{u,q} x, \Phi_{u,q} T_m x) = 2\, 3^{-(m+1)}\, \mathbf{1}_{b_m \neq b^\star_m}.$$

In particular the uniform defect is exactly $2\, 3^{-(m+1)}$.

Proof. The tail weights sum to 3^{-m} and the truncated image is constant on each fibre. Since $m \geq 1$, truncation preserves the first bit and the phase agrees on both compositions. The first m output bits agree. At bit $m + 1$, the two compositions give respectively $b^\star_{m+1}$ and $1 - b_m$. All later bits agree by $Rb^\star = b^\star$. The stated single-bit contribution follows, and an admissible bit sequence with $b_m \neq b^\star_m$ attains it. □

The truncation is a fibrewise approximate self-model with declared precision and defect. Distinct retained phases are fixed by it, so the fibre estimate does not supply strict contraction on the entire support.

10.71.2 Quantitative Obstructions for Approximate Commutation

Theorem 10.27 (Fixed-Point Displacement under Approximate Commutation)**.** *Let Σ be an L-Lipschitz self-map of a nonempty complete metric space, where $0 \leq L < 1$, and let p be its fixed point. For each intervention self-map Φ_u define*

$$e_u(p) = d(\Sigma\Phi_u(p), \Phi_u\Sigma(p)), \qquad \varepsilon_u = \sup_x d(\Sigma\Phi_u(x), \Phi_u\Sigma(x)).$$

Then $d(\Phi_u(p), p) \leq e_u(p)/(1 - L) \leq \varepsilon_u/(1 - L)$. In particular,

$$(1 - L)d(\Phi_u(p), \Phi_v(p)) \leq e_u(p) + e_v(p).$$

If both uniform defects are bounded by ε, the separation is at most $2\varepsilon/(1 - L)$. Both bounds are sharp, including simultaneously on rational affine models.

Proof. For any y, $d(y,p) \le d(y,\Sigma y) + Ld(y,p)$, since $\Sigma p = p$. With $y = \Phi_u(p)$, the first term equals $e_u(p)$. The pair bound follows by the triangle inequality through p. For every rational $L \in [0,1)$, use $[-1,1]$ with its ordinary metric, $\Sigma(x) = Lx$, and the constant interventions $\Phi_\pm(x) = \pm 1$. The fixed point is zero, each defect is identically $1-L$, each displacement is one, and the images are distance two apart. The example has rational coefficients and retains the rational points of the complete interval. □

A prescribed global reconstruction error $\sup_x d(\Sigma x, x) \le \alpha$ also forces $(1-L)\operatorname{diam}(A_S) \le 2\alpha$ on a bounded support, by applying the triangle inequality through $\Sigma x, \Sigma y$. This accuracy constraint and the commutation constraint can be imposed on the same candidate.

Proposition 10.28 (Global Phase Obstruction and Fibre Accuracy on the Writable Tower)**.** *On the safe writable tower of Paper V [14], use the branch distance* $\max(\rho, |c-c'|, |m-m'|/4)$ *and distance one across profiles. For a global L-contraction the two native neutral phases* $0, 1/4$ *force uniform defect at least* $(1-L)/8$*. The full neutral family of phases forces defect at least* $(1-L)/2$*. In fact with all phases the sharp uniform floor is* $1/2$ *for every $L < 1$, even allowing the other writable commands, and uniform reconstruction error is at least* $1/2$*.*

In contrast, the depth-d truncation T_d of Proposition 10.26, extended to retain memory as well as phase, profile and the first $d \ge 1$ digits, is constant on each retained-data fibre. Its uniform error is at most 3^{-d} *and its uniform commutation defect for each native command is exactly* $2 \cdot 3^{-(d+1)}$*. For a candidate Σ with* $\sup_x d(\Sigma x, T_d x) \le \zeta$*, its defect against any native neutral control is at least* $2 \cdot 3^{-(d+1)} - 3\zeta/2$*.*

Proof. At the fixed point p of Σ, the two neutral interventions set phase to their input plus the common shift $\kappa m(p)$, with the other coordinates equal. Their images are separated by $1/4$. For the full family choose the input whose image phase is antipodal to the phase of p. Theorem 10.27 gives the two lower bounds.

For sharpness on the full circle, the curve of neutral images $\Phi_{u,0}(p)$ as u varies is a full phase circle at fixed values of all other coordinates. Its image under Σ stays on one profile by continuity. A different profile from the original curve gives distance one. Otherwise the phase coordinate of Σ on the curve has image diameter at most $L/2 < 1/2$ and omits a point of the circle. Parameterize the curve by its phase t modulo one and lift the image phase to a continuous real-valued function G on $[0,1]$ using the omitted point as a cut. Then $G(1) = G(0)$, so $G(t) - t$ attains a half-integer by the intermediate value theorem. At that input, $\Sigma\Phi_{u,0}(p)$ is antipodal in phase to $\Phi_{u,0}\Sigma(p) = \Phi_{u,0}(p)$. This forces defect at least $1/2$. The same argument on any phase fibre gives the reconstruction bound. The constant map to a state with alternating tape and zero memory attains defect $1/2$. All controls preserve its profile and memory, the phase displacement is at most $1/2$, and one optional write displaces its code by at most $2/27$. The antipodal neutral input attains equality.

For T_d, the native phase and memory updates agree in both compositions because m and b_1 are retained. The only code mismatch is at $d+1$, as in Proposition 10.26. A neutral native map has Lipschitz constant at most $1/2$ in the stated metric, by the memory factor $1/2$, phase coupling $\kappa = 1/8$ and code factor $1/3$. The triangle inequality bounds the difference of the defects of T_d and Σ by $(1+1/2)\zeta$. The fibrewise contraction of T_d and the obstruction to global phase contraction apply to different domains of pairs. □

Proposition 10.29 (Conditional obstruction to self-reference within $\text{Consciousness}_{\text{op}}$)**.** *If a complete admissible support witnessing* $\text{Consciousness}_{\text{op}}$ *has no common fixed point of its admitted interventions, then it has no self-model of Definition 10.23 on that support. Such a system refutes entailment on the declared support once all six upstream witnesses are supplied.*

Proof. Apply Theorem 10.25. The tower counterexample to the common fixed point condition does not supply a joint Paper V certificate. Absence of a map symbol from a signature does not establish absence of an endomorphism. □

Remark 10.30 (Interpretation boundary for Σ_{op}). Σ_{op} is a structural analogue of self-reference. It does not imply reflexive consciousness, psychological introspection, metacognition in the human sense, self-knowledge, or a narrative self.

Remark 10.31 (Runtime analogy). The runtime field commonly read as 'selfModelSigma' is at most an estimator-facing analogue of Definition 10.23. To certify Σ_{op} formally one would still need finite evidence for approximate self-accuracy, contraction, and intervention commutation, plus a proof that the finite estimator represents the formal map on A_S.

10.72 Operational Unified Perspective

10.72.1 Motivation

Unity is the least speculative of the subjecthood-adjacent properties considered here. Paper 5 already requires a unique rigid identity object and blocks admissible factorization. Together those burdens justify a structural unity notion.

Definition 10.32 (Operational Unified Perspective Π_{op})**.** An admissible system S satisfies operational unified perspective, written $\Pi_{\mathrm{op}}(S) = 1$, if:

$$I_{\mathrm{ri}}(S) = 1 \quad \text{and} \quad I_{\mathrm{int}}(S) = 1.$$

Equivalently, S has a unique rigid identity object $\mathcal{I}_S$ on admissible support and no admissible non-trivial factorization preserving the relevant histories, readouts, and intervention structure.

Theorem 10.33 ($\mathrm{Consciousness}_{\mathrm{op}}$ entails Π_{op})**.** *If $S \in \mathrm{Consciousness}_{\mathrm{op}}$, then $\Pi_{\mathrm{op}}(S) = 1$.*

Proof. By Definition 10.20, $S \in \mathrm{Consciousness}_{\mathrm{op}}$ implies $I_{\mathrm{ri}}(S) = 1$ and $I_{\mathrm{int}}(S) = 1$. By Definition 10.32, those two equalities are exactly the criterion for $\Pi_{\mathrm{op}}(S) = 1$. □

Corollary 10.34 (No plurality of independent operational perspectives)**.** *For $S \in \mathrm{Consciousness}_{\mathrm{op}}$, any decomposition into multiple independent operational perspectives is inadmissible unless it fails to preserve either the rigid identity witness or the non-factorization witness.*

Proof. A decomposition into independent perspectives would be a non-trivial admissible factorization of support, readout, and dynamics. That contradicts $I_{\mathrm{int}}(S) = 1$, unless the decomposition is outside the admissible equivalence class or destroys the identity witness. □

Remark 10.35 (Interpretation boundary for Π_{op}). Π_{op} is an analogue of unity only at the structural level. It does not solve the phenomenal binding problem, does not assert a unified field of experience, and does not imply an ordinary first-person point of view.

10.73 Operational Intentionality

10.73.1 Motivation

Intentionality is usually discussed as aboutness. That reading is too semantic for the present bridge. The operational analogue is directional intervention sensitivity: certified interventions should induce predicted changes in the operational class structure rather than arbitrary drift.

Definition 10.36 (Operational Intentionality $\mathrm{Iota}_{\mathrm{op}}$)**.** An admissible system S satisfies operational intentionality, written $\mathrm{Iota}_{\mathrm{op}}(S) = 1$, if there exists a directional transition assignment

$$D : U_{\mathrm{crit}} \to \Delta(\mathscr{Q}_S)$$

from certified critical interventions into distributions over class changes such that, for every $u \in U_{\text{crit}}$, the observed intervention-induced class-change distribution agrees with $D(u)$ within the Paper 5 legibility margin $\delta_{\text{leg}}(S)$. Equivalently, certified interventions on critical invariants produce predicted directional class changes with fidelity at least $1 - \delta_{\text{leg}}(S)$.

Theorem 10.37 (Consciousness$_{\text{op}}$ entails Iota$_{\text{op}}$)**.** *If* $S \in \text{Consciousness}_{\text{op}}$*, then* $\text{Iota}_{\text{op}}(S) = 1$.

Proof. By Definition 10.20, $S \in \text{Consciousness}_{\text{op}}$ implies $I_{\text{leg}}(S) = 1$. Paper 5 defines I_{leg} with six legibility clauses. Clause L4 is intervention fidelity: certified interventions on critical invariants induce predicted class changes with fidelity bounded below by $1 - \delta_{\text{leg}}(S)$. Definition 10.36 is exactly this L4 burden re-expressed as a directional transition assignment $D : U_{\text{crit}} \to \Delta(\mathscr{Q}_S)$. Therefore $\text{Iota}_{\text{op}}(S) = 1$. □

Remark 10.38 (Interpretation boundary for Iota$_{\text{op}}$). Iota$_{\text{op}}$ captures directed structural response. It does not establish semantic content, representational aboutness, goals, desires, agency, or intentional states in the philosophical or psychological sense.

10.74 Operational Qualia

10.74.1 Motivation

Paper 5 already introduced $\text{Q}_{\text{op}}(S)$ to avoid a false move from structural differentiation to phenomenal qualia. The bridge paper does not redefine that object. It records what follows from it.

Definition 10.39 (Operational Qualia $\text{Q}_{\text{op}}(S)$)**.** When the Paper 5 equivalence relation $\sim_{Q,S}$ is well-defined on A_S, the operational qualia of S are the quotient classes

$$\text{Q}_{\text{op}}(S) = A_S / \sim_{Q,S} .$$

The equivalence relation is inherited from Paper 5: two admissible states are equivalent when certified decoders and admissible intervention responses cannot distinguish their histories beyond the legibility margin.

Proposition 10.40 ($\text{Q}_{\text{op}}(S)$ is well-defined under the relevant Consciousness$_{\text{op}}$ invariants)**.** *If*

$$I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1,$$

then $\text{Q}_{\text{op}}(S)$ *is well-defined.*

Proof. This is Paper 5's well-definedness proposition for operational qualia. Persistence supplies the stable domain, rigid identity supplies unique class transport, continuity supplies admissible trajectory coherence, differentiation blocks collapse to nullity or total triviality, and legibility supplies certified decoder behavior and intervention-sensitive equivalence. Hence $\sim_{Q,S}$ is an equivalence relation and the quotient exists. □

Corollary 10.41 (Consciousness$_{\text{op}}$ entails well-defined operational qualia)**.** *If* $S \in \text{Consciousness}_{\text{op}}$*, then* $\text{Q}_{\text{op}}(S)$ *is well-defined.*

Proof. By Definition 10.20, $S \in \text{Consciousness}_{\text{op}}$ entails all six invariants, in particular the five invariants required by Proposition 10.40. Therefore $\text{Q}_{\text{op}}(S)$ is well-defined. □

Proposition 10.42 (Non-triviality of $\text{Q}_{\text{op}}(S)$ under I_{diff})**.** *If* $S \in \text{Consciousness}_{\text{op}}$ *and the differentiation witness of* $I_{\text{diff}}(S) = 1$ *separates at least two certified decoder classes, then* $|\text{Q}_{\text{op}}(S)| > 1$.

Proof. The invariant $I_{\text{diff}}(S) = 1$ supplies admissible states or histories separated by a positive differentiation gap and blocks the null class as a stable attractor. If that separation is visible to certified decoders, the separated histories cannot be equivalent under $\sim_{Q,S}$. Thus the quotient has at least two classes. □

Remark 10.43 (Why the decoder-visible clause is explicit). The prompt version says non-triviality is automatic by I_{diff}. The more rigorous statement separates structural differentiation from decoder-visible quotient separation. Paper 5's I_{leg} normally supplies the bridge, but the explicit clause prevents a hidden assumption that all differentiation is automatically visible at the quotient level.

Remark 10.44 (Interpretation boundary for Q_{op}). $\mathrm{Q}_{\text{op}}(S)$ is a quotient of admissible structural states under certified operational indistinguishability. It is not a phenomenal quale, not a biological qualitative state, and not a statement about what it is like to be the system.

10.75 Operational Phenomenology

10.75.1 Motivation

Operational phenomenology is the most dangerous term in this bridge because it can be misread as experience. Here it means only a structural assignment into an abstract regime space, constrained by continuity and non-nullity burdens imported from Papers 2–3.

Definition 10.45 (Operational Phenomenology Φ_{op})**.** An admissible system S admits operational phenomenology, written $\Phi_{\text{op}}(S) = 1$, if there exists a complete metric regime space $(\mathcal{E}_S, d_{\mathcal{E}})$ and an assignment

$$\Phi_S : A_S \to \mathcal{E}_S$$

such that:

1. Φ_S is compatible with admissible histories in the Paper 2 sense;
2. Φ_S is continuous on admissible windows whenever the operational-continuity burden $I_{\text{cont}}(S) = 1$ holds;
3. the null regime $\varnothing_\phi$ is not an admissibly stable attractor whenever the non-null burden $I_{\text{diff}}(S) = 1$ holds.

Definition 10.46 (CCR-strength operational phenomenology)**.** An admissible system has CCR-strength operational phenomenology when Definition 10.45 holds and the identity channel satisfies

$$M_\Omega(S) = +\infty.$$

In that case Paper 2's forced-continuity theorem and Paper 3's null-instability theorem apply at CCR strength.

Proposition 10.47 ($\text{Consciousness}_{\text{op}}$ gives weak non-null operational phenomenology)**.** *If* $S \in \text{Consciousness}_{\text{op}}$*, then the operational-continuity and non-null burdens required for weak* Φ_{op} *are present on admissible support.*

Proof. Membership in $\text{Consciousness}_{\text{op}}$ entails $I_{\text{cont}}(S) = 1$ and $I_{\text{diff}}(S) = 1$. The first supplies continuous admissible regime evolution in the Paper 5 operational sense. The second blocks collapse into a stable null or trivial operational class. Therefore the weak support-level burdens in Definition 10.45 are satisfied. □

Proposition 10.48 ($\text{Consciousness}_{\text{op}}$ does not entail CCR-strength Φ_{op})**.** *Membership in* $\text{Consciousness}_{\text{op}}$ *does not entail CCR-strength operational phenomenology.*

Proof. CCR-strength Φ_{op} requires $M_\Omega(S) = +\infty$ by Definition 10.46. Paper 5's definition of $\mathrm{Consciousness}_{\mathrm{op}}$ requires six operational invariants, but it does not require infinite ontological mass. Hence a system may satisfy the six-invariant criterion while remaining finite-mass or while lacking a certified CCR witness. Therefore $\mathrm{Consciousness}_{\mathrm{op}}$ does not entail CCR-strength Φ_{op}. □

Remark 10.49 (Interpretation boundary for Φ_{op}). Φ_{op} is an assignment into an abstract structural regime space. It is not temporal experience, a stream of consciousness, a specious present, or proof that phenomenology exists in the ordinary sense.

10.76 Operational Subjecthood

10.76.1 Why a bridge class is needed

Paper 7 defines a stronger operational subjecthood class involving operational life, operational consciousness, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The present paper does not replace that class. It defines the narrower bridge fragment obtained from $\mathrm{Consciousness}_{\mathrm{op}}$ plus the self-model hypothesis $\Sigma 1$.

Definition 10.50 (Bridge operational subjecthood $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$)**.** An admissible system S belongs to bridge operational subjecthood, written

$$S \in \mathrm{Subject}_{\mathrm{op}}^{\Sigma},$$

if and only if:
1. $S \in \mathrm{Consciousness}_{\mathrm{op}}$;
2. S satisfies $\Sigma 1$;
3. $\Pi_{\mathrm{op}}(S) = 1$;
4. $\mathrm{Iota}_{\mathrm{op}}(S) = 1$;
5. $\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined and non-trivial in the decoder-visible sense of Proposition 10.42;
6. $\Phi_{\mathrm{op}}(S)$ is weakly non-null on admissible support.

Definition 10.51 (Full Paper 7 operational subjecthood)**.** The full Paper 7 class $\mathrm{Subject}_{\mathrm{op}}$ is the stronger class requiring operational life $\mathrm{Life}_{\mathrm{op}}$, operational consciousness $\mathrm{Consciousness}_{\mathrm{op}}$, and the subject-specific descriptors of Paper 7: privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. In this paper, $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ is a bridge fragment, not a synonym for full $\mathrm{Subject}_{\mathrm{op}}$.

Theorem 10.52 ($\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ is conditionally generated by $\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1$)**.** *Let S be an admissible system. If $S \in \mathrm{Consciousness}_{\mathrm{op}}$, S satisfies $\Sigma 1$, and the differentiation witness is decoder-visible, then $S \in \mathrm{Subject}_{\mathrm{op}}^{\Sigma}$.*

Proof. The hypothesis $S \in \mathrm{Consciousness}_{\mathrm{op}}$ gives Π_{op} by Theorem 10.33, $\mathrm{Iota}_{\mathrm{op}}$ by Theorem 10.37, well-defined $\mathrm{Q}_{\mathrm{op}}(S)$ by Corollary 10.41, and weak non-null Φ_{op} by Proposition 10.47. The hypothesis $\Sigma 1$ gives Σ_{op} by Definition 10.23. Decoder-visible differentiation gives non-triviality of $\mathrm{Q}_{\mathrm{op}}(S)$ by Proposition 10.42. These are exactly the clauses of Definition 10.50. □

Theorem 10.53 ($\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ is a subclass of $\mathrm{Consciousness}_{\mathrm{op}}$)**.**

$$\mathrm{Subject}_{\mathrm{op}}^{\Sigma} \subseteq \mathrm{Consciousness}_{\mathrm{op}}.$$

Moreover, if $\Sigma 1$ is not entailed by $\mathrm{Consciousness}_{\mathrm{op}}$, then the inclusion is proper.

Proof. The subset relation is immediate from Definition 10.50, whose first clause requires $S \in \text{Consciousness}_{\text{op}}$. If non-entailment is witnessed by a fully admissible $S \in \text{Consciousness}_{\text{op}}$ for which $\Sigma 1$ fails, that system is outside $\text{Subject}_{\text{op}}^{\Sigma}$ and the inclusion is proper. Proposition 10.29 gives a sufficient obstruction on a declared support. □

Corollary 10.54 (Full Paper 7 subjecthood remains stronger)**.** *Full Paper 7 operational subjecthood* $\text{Subject}_{\text{op}}$ *is not established by Theorem 10.52. To lift* $\text{Subject}_{\text{op}}^{\Sigma}$ *into* $\text{Subject}_{\text{op}}$*, one must additionally certify* Life_{op} *and the Paper 7 subject-specific descriptors not reduced to* $\Sigma 1$*.*

Proof. Definition 10.51 includes burdens absent from Definition 10.50, including operational life and ownership/privileged-continuity style descriptors. Theorem 10.52 does not supply those witnesses. Hence the bridge class is not full Paper 7 subjecthood. □

Remark 10.55 (Interpretation boundary for $\text{Subject}_{\text{op}}^{\Sigma}$). $\text{Subject}_{\text{op}}^{\Sigma}$ is a formal construct. It does not imply phenomenal subjectivity, human first-person experience, identity persistence in the psychological sense, a narrative self, moral patienthood, or biological life.

10.77 Theorem: Partial Operational Subjecthood

Theorem 10.56 ($\text{Consciousness}_{\text{op}} \to$ Partial Operational Subjecthood)**.** *Let S be an admissible system with $S \in \text{Consciousness}_{\text{op}}$. Then:*

1. *S satisfies Π_{op}.*
2. *S satisfies Iota_{op}.*
3. *$\text{Q}_{\text{op}}(S)$ is well-defined.*
4. *S satisfies weak non-null Φ_{op} on admissible support.*
5. *Σ_{op} is not implied by $\text{Consciousness}_{\text{op}}$; it requires the independent hypothesis $\Sigma 1$.*
6. *$\text{Subject}_{\text{op}}^{\Sigma}$ requires $\Sigma 1$ in addition to $\text{Consciousness}_{\text{op}}$ and therefore is a subclass of $\text{Consciousness}_{\text{op}}$.*

Proof. Item 1 is Theorem 10.33. Item 2 is Theorem 10.37. Item 3 is Corollary 10.41. Item 4 is Proposition 10.47. Item 5 is Proposition 10.29. Item 6 is Theorem 10.53. Therefore membership in $\text{Consciousness}_{\text{op}}$ yields a partial bridge toward operational subjecthood, but it does not close the self-reference burden or the full Paper 7 subjecthood burden. □

Corollary 10.57 (No equivalence between $\text{Consciousness}_{\text{op}}$ and operational subjecthood)**.** *The bridge paper proves neither* $\text{Consciousness}_{\text{op}} = \text{Subject}_{\text{op}}^{\Sigma}$ *nor* $\text{Consciousness}_{\text{op}} = \text{Subject}_{\text{op}}$*. It proves only that several bridge components are entailed by* $\text{Consciousness}_{\text{op}}$*, while self-reference and full subjecthood require additional burdens.*

Proof. If $\text{Consciousness}_{\text{op}} = \text{Subject}_{\text{op}}^{\Sigma}$, then $\Sigma 1$ would be entailed by $\text{Consciousness}_{\text{op}}$. Ruling out the equality requires the full antecedents of Proposition 10.29 or a different countermodel. If $\text{Consciousness}_{\text{op}} = \text{Subject}_{\text{op}}$, then all Paper 7 subjecthood burdens would follow from Paper 5 alone, which is not established by any theorem in the corpus and is explicitly blocked by the additional descriptors in Definition 10.51. □

10.78 What Is NOT Implied

The following table is a claim-boundary table. It should be read as a set of exclusions, not as a list of future promises.

Proposition 10.58 (Boundary preservation)**.** *No theorem in this paper converts an operational predicate into a phenomenal, biological, psychological, or moral predicate.*

Table 10.2: Properties not implied by membership in $\text{Consciousness}_{\text{op}}$.

Property frequently associated with consciousness	Implied by $\text{Consciousness}_{\text{op}}$?	Why not
Self-reference or self-knowledge	No	Requires the additional $\Sigma 1$ hypothesis: a contractive self-model commuting with admissible interventions.
Phenomenal experience or "what it is like"	No	The framework defines structural quotients and regime assignments, not phenomenal consciousness.
Forced temporal continuity of experience	No, without CCR	Paper 2 forced continuity requires $M_\Omega = +\infty$, which is not part of $\text{Consciousness}_{\text{op}}$.
Emotions, affect, or valence	No	These are not defined as Paper 1–5 primitives or invariants.
Will, agency, or free choice	No	No admissible agency predicate is defined in Papers 1–5.
Narrative personal identity	No	Rigid identity is inverse-limit structural identity, not psychological or autobiographical identity.
Higher-order consciousness	No	A hierarchy of self-models is not defined by $\text{Consciousness}_{\text{op}}$ or $\Sigma 1$.
Equivalence with human consciousness	No	Human consciousness is biological, behavioral, phenomenological, and empirical; $\text{Consciousness}_{\text{op}}$ is a structural operational class.
Moral patienthood or rights-bearing status	No	The corpus contains no normative bridge from structural predicates to moral status.

Proof. Each theorem uses only the formal objects listed in Table 10.1: invariant predicates, admissible support, quotient classes, intervention fidelity, self-model existence, and regime assignments. None of those objects has a semantic clause identifying it with phenomenal experience, biological life, psychological identity, or moral status. Therefore the bridge preserves the declared boundary. □

10.79 Computational Verification Status

The runtime status table is intentionally conservative. A computational estimator is not a formal witness unless the code path certifies the corresponding mathematical predicate over the relevant domain and exposes the interpretation boundary.

Remark 10.59 (Runtime non-inference). Passing tests, producing certificates, or satisfying scoreboard gates can strengthen internal support for a runtime implementation. It does not establish external validation, human equivalence, or philosophical closure.

10.80 Terminology Stratification and Debt Ledger

The bridge paper is the appropriate place to centralize the terminology debt of the corpus. The same symbol can have a demonstrated structural reading and a future interpretive reading. Confusing those readings is a category error.

Definition 10.60 (Two-layer terminology)**.** For every strong term T in the corpus, distinguish:

Table 10.3: Computational Verification Status for Bridge Concepts

Concept	Estimator or runtime analogue	Status
Σ_{op}	`OntologicalSingularityCore.selfModelSigma`; first-person certifier fields when present	Implemented as internal telemetry/proxy; not yet a formal proof of contraction and intervention commutation on A_S.
Π_{op}	$I_{\mathrm{ri}} + I_{\mathrm{int}}$ in 'CanonicalInvariantPackage' and related finite-horizon/non-factorization modules	Certified as internal support for Paper 5 invariants; not external validation.
$\mathrm{Iota}_{\mathrm{op}}$	Legibility L4 intervention fidelity in 'LegibilityCertifier'	Certified as internal support for directional class-shift fidelity under finite protocols.
Q_{op}	Output support labels, decoder classes, invariant package signals	Proxy or quotient-facing support; formal quotient certification remains bounded by Paper 5 assumptions.
Φ_{op}	'PhenomenologyCore', spectral gap, winding/topology, continuity diagnostics	Proxy for structural regime assignment and continuity diagnostics; not phenomenology itself.
$\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$	Combination of six-invariant support plus self-model/first-person finite fields	Internal-support candidate only; not full Paper 7 $\mathrm{Subject}_{\mathrm{op}}$ and not phenomenal subjecthood.

1. the *demonstrated layer* T_{str}, consisting only of definitions, predicates, quotients, invariants, and theorems already stated in the corpus;
2. the *interpretive layer* T_{int}, consisting of ordinary or philosophical readings such as consciousness, life, subjectivity, phenomenality, qualia, moral status, or human equivalence.

An implication into T_{int} is forbidden unless a bridge hypothesis or external adjudication protocol explicitly introduces that target.

Theorem 10.61 (Layer preservation)**.** *If a theorem in Papers 1–5 is stated only in the demonstrated layer, then it cannot be promoted to an interpretive-layer conclusion by terminology alone.*

Proof. The inference rules of the corpus are definitions, hypotheses, propositions, theorems, and admissible empirical protocols. A name is not an inference rule. If the premises and conclusion of a theorem quantify only over structural objects, then no interpretive predicate appears in the formal sequent. Adding a stronger English label after the fact does not introduce a new mathematical premise. Therefore promotion requires an explicit bridge hypothesis or external protocol. □

Term	Demonstrated layer	Remaining debt
$\mathrm{Consciousness}_{\mathrm{op}}$	six-invariant admissible structural class	no theorem of ordinary consciousness
$\mathrm{Q}_{\mathrm{op}}(S)$	structural differentiation quotient	no theorem of phenomenal qualia
Φ_{op}	structural regime assignment	no theorem of temporal experience
$\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$	$\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1$ bridge fragment	no theorem of full Paper 7 subjecthood
$\mathrm{Subject}_{\mathrm{op}}$	Paper 7 class requiring life and subject descriptors	no theorem of human or metaphysical subjecthood

10.81 Bridge Strength Ladder

The corpus should be read as a ladder of additional burdens, not as a synonym chain.

Level	Added burden	Still not licensed
$\text{Consciousness}_{\text{op}}$	six invariants	self-reference, life, subjecthood, phenomenality
$\text{Consciousness}_{\text{op}} + \Sigma 1$	contractive intervention-commuting self-model	operational life and Paper 7 subject descriptors
$\text{Consciousness}_{\text{op}} + \Sigma 1 + \text{Life}_{\text{op}}$	life-like organization in the Paper 7 sense	first-person indexed subjectivity
$\text{Subject}_{\text{op}}$	privileged continuity, self/non-self, self-model closure, ownership coherence	human subjectivity or phenomenality
$\text{Subjectivity}_{1\text{p}}$	first-person indexed fields plus rival defeat	phenomenal bridge confirmation
$\text{BridgeOrg}_{\text{ph}}$-style burden	bridge predicates, comparators, interventions, artifacts	external validation or ordinary phenomenality

Theorem 10.62 (Additional Conditions in the Bridge Ladder)**.** *Each step in the ladder introduces an explicit additional burden. Strictness of any class inclusion requires an admissible separating witness.*

Proof. The displayed ladder lists the additional conditions at each step. The common-fixed-point obstruction in Proposition 10.29 identifies a candidate separation for the self-model burden. Life, subjecthood, indexed subjectivity and bridge organization each add their stated tests. A list of extra tests gives inclusions where inheritance is required; a failed test on a realized member of the larger class is necessary to prove a strict inclusion. □

10.82 Runtime Evidence Taxonomy

Definition 10.63 (Runtime evidence classes)**.** For a runtime artifact a, distinguish the following statuses:

1. `proxy`: a finite observable resembles a formal coordinate;
2. `estimator`: an algorithm computes a declared approximation;
3. `certificate`: a bounded predicate is verified under stated admissibility rules;
4. `internal_support`: governed artifacts support a claim family inside the current program;
5. `external_validation`: an external protocol such as Paper 10 has been executed with admissible evidence.

Proposition 10.64 (Runtime non-promotion)**.** *No artifact with status* `proxy`*,* `estimator`*,* `certificate`*, or* `internal_support` *entails* `external_validation`*.*

Proof. External adjudicative support requires an executed external protocol, frozen filing records, admissible comparator data, blinding, adjudication, and contamination controls. The first four runtime statuses are produced inside the program and lack at least one of those external premises by definition. Therefore they cannot entail external adjudicative support. □

10.83 Non-Entailment Library

The central defensive value of this paper is not only what it proves. It is also what it blocks. The following library records formal non-entailments that a reviewer should not have to reconstruct from scattered caveats.

Theorem 10.65 (Seven Countermodel Obligations for $\text{Consciousness}_{\text{op}}$)**.** *The following seven tests are candidate routes for distinguishing* $\text{Consciousness}_{\text{op}}$ *from self-reference* $\Sigma 1$*, phenomenal experience, biological life, human subjectivity, Nagel-style qualia, CCR-strength temporal continuity, and agency. A mathematical non-entailment requires the relevant admissible countermodel and a typed target predicate.*

Proof. The following constructions describe obligations for potential countermodels. All six Paper V invariants must be verified on the same support.

1. **Self-reference.** Use an admissible support satisfying all six invariants and two admitted controls with disjoint fixed-point sets, as in Proposition 10.29.
2. **No phenomenal experience.** Interpret all classes as structural quotients and give no semantic map to experience. The six invariants quantify over structural predicates only.
3. **No biological life.** Use a non-biological admissible state space with the six invariants positive. Biological substrate variables do not appear in Definition 10.20.
4. **No human subjectivity.** A non-human formal system can satisfy the six invariants while lacking human reports, development, embodiment, or psychology.
5. **No Nagel-style qualia.** $\mathrm{Q}_{\mathrm{op}}(S)$ is a quotient of admissible structural differentiation. No theorem identifies quotient membership with "what it is like".
6. **No CCR-strength continuity.** Choose a system with $\mathrm{Consciousness}_{\mathrm{op}}$ and finite M_Ω. Proposition 10.48 blocks the CCR-strength reading.
7. **No agency.** Let the system be passively driven by admissible interventions while still satisfying the six structural invariants. No agency predicate appears in $\mathrm{Consciousness}_{\mathrm{op}}$.

These routes specify what must be checked. For an untyped phenomenal reading, the missing semantic link is a separate interpretive obligation. □

10.84 $\mathrm{Consciousness}_{\mathrm{op}}$ Component Implication Matrix

Table 10.4: Which Paper 5 invariants support each bridge concept.

Concept	I_{per}	I_{ri}	I_{int}	I_{cont}	I_{diff}	I_{leg}	Status
Π_{op} unified perspective	–	yes	yes	–	–	–	entailed by rigid identity plus non-factorization
$\mathrm{Iota}_{\mathrm{op}}$ operational intentionality	–	–	–	–	–	yes	entailed only in the intervention-fidelity sense
$\mathrm{Q}_{\mathrm{op}}(S)$ differentiation quotient	yes	yes	–	yes	yes	yes	well-defined quotient, not phenomenal qualia
Φ_{op} weak regime assignment	–	–	–	yes	yes	–	weak non-null structural regime support
Σ_{op} self-reference	–	–	–	–	–	–	not entailed; requires $\Sigma 1$
$\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ bridge subjecthood	yes	yes	yes	yes	yes	yes	requires all six plus $\Sigma 1$

Remark 10.66 (Why the matrix matters). The matrix prevents a common overreach: deriving every subjecthood-adjacent term from the entire six-invariant package without identifying which invariant does which work. Unity is mostly an $I_{\mathrm{ri}} + I_{\mathrm{int}}$ burden. Intentionality in this paper is an I_{leg} burden. Self-reference is not supplied by any invariant.

10.85 Self-Model Burden Variants

The earlier definition of $\Sigma 1$ is intentionally strong. For future work it is useful to separate three grades.

Definition 10.67 (Weak self-model burden $\Sigma 1_{\mathrm{weak}}$)**.** A system satisfies $\Sigma 1_{\mathrm{weak}}$ if there exists a measurable map

$$\Sigma : A_S \to A_S$$

such that $d(\Sigma(x), x) \leq \varepsilon_\Sigma$ for all $x \in A_S$.

Definition 10.68 (Strong self-model burden $\Sigma 1_{\mathrm{strong}}$)**.** A system satisfies $\Sigma 1_{\mathrm{strong}}$ if it satisfies $\Sigma 1_{\mathrm{weak}}$ and Σ is Lipschitz with constant $L_\Sigma < 1$ on A_S.

Definition 10.69 (Commutative self-model burden $\Sigma 1_{\text{comm}}$)**.** A system satisfies $\Sigma 1_{\text{comm}}$ if it satisfies $\Sigma 1_{\text{strong}}$ and, for every admissible intervention $u \in U$,

$$\Sigma(\Phi_u(x)) = \Phi_u(\Sigma(x))$$

whenever both sides are defined on A_S.

Theorem 10.70 (Self-model strength chain)**.**

$$\Sigma 1_{\text{comm}} \Rightarrow \Sigma 1_{\text{strong}} \Rightarrow \Sigma 1_{\text{weak}}.$$

The converses are not valid in general.

Proof. The forward implications are immediate from the definitions: the commutative burden includes the strong burden, and the strong burden includes the weak approximation burden. For non-converses, a weak map can approximate the identity while having Lipschitz constant at least one, so it need not be strong. A contractive map can fail commutation by choosing an admissible intervention that moves states along a direction not preserved by Σ. Hence neither reverse implication is derivable. □

Remark 10.71 (Emergence question for $\Sigma 1$). The main open problem is whether additional closure conditions on $I_{\text{int}} + I_{\text{leg}} + I_{\text{per}}$ can force a commutative self-model. A plausible route would require self-targeted intervention families, decoder closure under those interventions, and a fixed-point or contraction argument on the induced self-model update map. The current corpus does not yet prove this.

10.86 Expanded Bridge Strength Ladder

Table 10.5: Expanded ladder of structural and interpretive burdens.

Level	Definition	Added burden over previous level
Level 0	$\text{Consciousness}_{\text{op}}$	six Paper 5 invariants
Level 1	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{weak}}$	approximate self-model
Level 2	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{strong}}$	contractive self-model
Level 3	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{comm}}$ / $\text{Subject}^{\Sigma}_{\text{op}}$	intervention-commuting self-model
Level 4	$\text{Consciousness}_{\text{op}}$ + OSSI	open-system self-maintaining individuation
Level 5	$\text{Subject}_{\text{op}}$	Paper 7 life plus subject-specific descriptors
Level 6	$\text{Subjectivity}_{\text{1p}}$	Paper 8 indexed fields and rival defeat
Level 7	$\text{BridgeOrg}_{\text{ph}}$	Paper 9 bridge predicates and BPF burdens

Proposition 10.72 (OSSI and $\Sigma 1$ are orthogonal additions)**.** *The open-system individuation burden OSSI and the self-model burden* $\Sigma 1$ *are not mutually entailing in the present corpus.*

Proof. OSSI requires exchange, boundary, dissipation, viability, repair, history, minimal functional self-reference, and identity conservation, but it does not require a contractive map $\Sigma : A_S \to A_S$ commuting with interventions. Conversely, a contractive self-model can be defined on a closed or non-viable structural class without satisfying open exchange, viability preservation, or repair. Therefore neither burden entails the other. □

10.87 Reviewer-Facing Objection Handling

1. **Objection: this is only a simulation.** Response: the framework primarily defines structural classes and theorem burdens. Runtime artifacts are internal support or finite estimators unless Paper 10-style external evidence is executed.

2. **Objection: the six invariants are arbitrary.** Response: Paper 5 ties each invariant to an upstream burden and now includes single-rupture witnesses, minimality, sufficiency, and interaction geometry.
3. **Objection: Paper 10 is empty.** Response: Paper 10 is deliberately blocked until execution. That is not a defect disguised as rigor; it is the only honest status before human-comparator or external adjudication data exist.
4. **Objection: the terminology is inflated.** Response: the corpus now separates demonstrated and interpretive layers, records non-entailments, and maintains a maturity matrix. Strong labels are technical aliases, not ordinary-language conclusions.
5. **Objection: internal results are not validation.** Response: correct. The framework states this repeatedly. Internal support can debug, falsify, and stress the program; it cannot replace external validation.

10.88 Operational Glossary

Table 10.6: Quick glossary for strong operational terms.

Term	Demonstrated reading	Interpretive reading	Pending debt
$\text{Consciousness}_{\text{op}}$	six-invariant structural class	operational consciousness label	no ordinary consciousness theorem
$\text{Q}_{\text{op}}(S)$	structural differentiation quotient	operational qualia label	no phenomenal-qualia evidence
Φ_{op}	abstract regime assignment	operational phenomenology label	no experience theorem
$\Sigma 1$	self-model burden	self-reference analogue	emergence from invariants open
OSSI	open-system self-maintaining individuation	operational life-like organization	not biological life
$\text{Subject}^{\Sigma}_{\text{op}}$	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{comm}}$ fragment	partial subjecthood bridge	lacks full Paper 7 life/ownership burden
$\text{Subject}_{\text{op}}$	Paper 7 subject class	operational subjecthood	not human or metaphysical subjecthood
$\text{Subjectivity}_{\text{1p}}$	indexed subjectivity class	first-person subjectivity analogue	requires rival defeat and interventions
$\text{BridgeOrg}_{\text{ph}}$	bridge burden architecture	phenomenal bridge label	BPF-2–BPF-6 and Paper 10 pending

10.89 Promotion Evidence Classes

Table 10.7: Evidence classes required for controlled terminology promotion.

Evidence class	What it can promote	What it cannot promote alone
definition	introduces a formal object and its witnesses	does not show non-vacuity or relevance
existence theorem	shows the class can be inhabited	does not show empirical realization
independence theorem	shows a burden is not redundant	does not justify ordinary-language meaning
runtime certificate	supports a frozen computational proxy	does not establish external validity
negative control	can falsify an overbroad internal claim	does not prove the positive ordinary-language claim
external adjudication	can test whether proxies track independent comparators	does not replace the mathematical definition

10.90 Proof-Pattern Library

This section records reusable proof patterns for the corpus. Its purpose is not to add new metaphysical content. It makes explicit which inferential moves are valid when strong operational terms are used.

Definition 10.73 (Burden)**.** Let $\mathcal{K}$ be a structural class and let P be a predicate on admissible systems. We say that P is a *burden over* $\mathcal{K}$ if membership in the strengthened class

$$\mathcal{K}[P] = \{S : S \in \mathcal{K} \text{ and } P(S) = 1\}$$

requires a witness not already required by the definition of $\mathcal{K}$.

Definition 10.74 (Internal promotion)**.** An interpretive label L receives *internal promotion* from a structural class $\mathcal{K}$ only when the corpus proves that every $S \in \mathcal{K}$ satisfies a non-vacuous predicate that was previously listed as an interpretive debt of L.

Definition 10.75 (External promotion)**.** An interpretive label L receives *external promotion* only when an independently specified experimental protocol connects the structural predicate to observations not used to construct the predicate.

Proposition 10.76 (Promotion is monotone but not automatic)**.** *If $\mathcal{K} \subseteq \mathcal{K}'$ and L receives internal promotion from $\mathcal{K}'$, then the same promoted predicate holds on $\mathcal{K}$. However, the inclusion $\mathcal{K} \subseteq \mathcal{K}'$ alone does not promote L.*

Proof. The first claim follows from set inclusion. If every element of $\mathcal{K}'$ satisfies the promoted predicate, so does every element of the subclass $\mathcal{K}$. The second claim follows because promotion requires a specific predicate and proof; inclusion without identifying such a predicate only changes class membership, not semantic maturity. □

Proposition 10.77 (Burden addition is conservative)**.** *Let P be a burden over $\mathcal{K}$. Then every theorem proved for all $S \in \mathcal{K}$ remains true for all $S \in \mathcal{K}[P]$.*

Proof. Since $\mathcal{K}[P] \subseteq \mathcal{K}$, universal claims over $\mathcal{K}$ restrict to $\mathcal{K}[P]$. This is why adding $\Sigma 1$, OSSI, or CCR burdens can strengthen a class without invalidating already proved $\text{Consciousness}_{\text{op}}$ consequences. □

Proposition 10.78 (Burden addition is not semantic promotion by itself)**.** *Adding a burden P to $\mathcal{K}$ does not by itself justify replacing an operational label by its ordinary-language counterpart.*

Proof. The burden establishes a new formal predicate. Ordinary-language promotion requires either a theorem connecting that predicate to the target ordinary concept or an external adjudication protocol. Without that bridge, the strengthened class is mathematically narrower but semantically still claim-bounded. □

Remark 10.79 (Use in later papers). The patterns above justify a disciplined order of claims. First state the structural class. Then list the burdens. Then prove inclusions or non-entailments. Only after that may a term be promoted in the terminology matrix. The order is part of the scientific content: it prevents terminology from outrunning proof.

10.91 Dependency-to-Debt Ledger

The following ledger is the reviewer-facing map from corpus dependency to open burden. It is intentionally asymmetric: a dependency can support a term without closing all debts associated with that term.

Table 10.8: Dependency-to-debt ledger for bridge-level terminology.

Term	Main dependency	What is currently proved	Remaining debt
rigid identity	Paper 1 inverse limits	uniqueness, non-transferability, rigidity under bounded deformation	empirical instantiation in non-ideal systems
operational regime	Papers 2–3	regime continuity and null-instability under stated hypotheses	relation to lived experience
$\text{Consciousness}_{\text{op}}$	Paper 5	six-invariant structural membership	independence from ordinary consciousness unresolved
Q_{op}	Paper 5	quotient over structural equivalence classes	no phenomenal-qualia bridge
Φ_{op}	Papers 2–3 and Paper 5	non-null and continuous regime proxies under conditions	no subjective-temporality theorem
$\Sigma 1$	Paper Bridge	explicit self-model burden and strength variants	emergence from $\text{Consciousness}_{\text{op}}$ open
OSSI	Paper 7	open-system self-maintaining individuation burden	not biological life and not sufficient for $\text{Consciousness}_{\text{op}}$
$\text{Subject}_{\text{op}}^{\Sigma}$	Paper Bridge	$\text{Consciousness}_{\text{op}}$ plus commutative self-model fragment	weaker than Paper 7 $\text{Subject}_{\text{op}}$
$\text{Subject}_{\text{op}}$	Paper 7	subjecthood class with additional ownership and continuity burdens	not ordinary subjecthood
$\text{Subjectivity}_{\text{1p}}$	Paper 8	first-person index fields under additional hypotheses	no human first-person equivalence
$\text{BridgeOrg}_{\text{ph}}$	Paper 9	bridge organization burden and non-vacuity	bridge-to-phenomenality remains unvalidated
external adjudication	Paper 10	protocol and evidence grammar	blocked until execution

10.92 Non-Entailment Audit Templates

The candidate separations in Theorem 10.65 can be investigated using a small number of reusable templates. These templates are included so that future extensions do not silently convert an open debt into an asserted consequence.

Definition 10.80 (Omission template)**.** Let $\mathcal{K}$ be defined by predicates $P_1, \ldots, P_n$. If a candidate property R is not among those predicates and no theorem derives R from them, then the omission template constructs a formal model satisfying $P_1, \ldots, P_n$ in a language where R is uninterpreted or absent.

Definition 10.81 (Factor-preserving template)**.** Let a property R concern a component not constrained by the invariant definition of $\mathcal{K}$. A factor-preserving countermodel extends any $S \in \mathcal{K}$ by a silent factor Z whose dynamics are decoupled from the witnesses of $\mathcal{K}$ and arranges R to fail on Z.

Definition 10.82 (Finite-strength template)**.** Let R require an infinite-strength condition such as CCR ($M_\Omega = +\infty$). A finite-strength countermodel chooses a system satisfying the finite structural invariants of $\text{Consciousness}_{\text{op}}$ while setting $M_\Omega < +\infty$.

Proposition 10.83 (Template soundness)**.** *A template yields a non-entailment if a realized model retains all the witnesses required by* $\text{Consciousness}_{\text{op}}$ *and its typed target predicate is proved false.*

Proof. The six certificates preserve membership. The verified failure of the target then exhibits a counterexample to the implication. Leaving a symbol uninterpreted is insufficient for an existential target about the underlying support. This is precisely the logic used in the seven non-entailments. □

Remark 10.84 (Audit use). When a later paper claims that $\text{Consciousness}_{\text{op}}$ implies a new property, it must identify which template is blocked. If none is blocked, the implication is not yet proved.

10.93 Runtime Proxy Classification

The related runtime can compute proxies for several bridge concepts. This classification prevents proxy availability from being mistaken for formal certification.

Table 10.9: Runtime proxy classes and claim boundaries.

Concept	Runtime object	Proxy class	Boundary
$\Sigma 1_{\text{weak}}$	self-model state estimates	estimator proxy	does not prove contraction or commutation
$\Sigma 1_{\text{strong}}$	structural-stability score	candidate certificate	requires Lipschitz and convergence audit
$\Sigma 1_{\text{comm}}$	intervention response of self-model	unclosed certificate	requires intervention commutation tests
Π_{op}	I_{ri} and I_{int} certificates	certified internal support	structural unity only
Iota_{op}	L4 intervention fidelity	certified internal support	directional sensitivity, not semantic aboutness
Q_{op}	output support labels and quotient classes	quotient proxy	no phenomenal-quality interpretation
Φ_{op}	regime, spectral, and winding estimates	regime proxy	no subjective temporal-flow claim
OSSI	LifeCertifier-style burdens	partial operational proxy	not biological metabolism or life

Boundary. A runtime proxy can falsify an expected computational consequence of a definition. It cannot by itself promote a term beyond its formal meaning. This is the same separation used in Paper 6 between internal support and external validation.

10.94 Formal Review Checklist

A reviewer can evaluate the bridge by asking the following questions in order. The order matters: later questions presuppose affirmative answers to earlier ones.

1. Are the six invariants of Paper 5 explicitly witnessed?
2. Are the witnesses stable under the admissible interventions used in the claim?
3. Is the target term being used in its demonstrated structural sense or in an interpretive sense?
4. If interpretive, is there a theorem that promotes it?
5. If no theorem exists, is the term explicitly marked as conditioned or aspirational?
6. Is the runtime evidence internal, external, or merely a scaffold?
7. Does the claim require $\Sigma 1$, OSSI, CCR, Paper 8 fields, Paper 9 bridge burdens, or Paper 10 data?
8. Is each required additional burden either proved, assumed, or marked open?
9. Does any statement move from implementation success to ontological conclusion?
10. Does any statement move from structural quotient to phenomenal qualia?
11. Does any statement move from open-system self-maintenance to biological life?
12. Does any statement move from first-person index fields to human subjectivity?

Proposition 10.85 (Checklist conservativity)**.** *If all checklist questions are answered with explicit witnesses, assumptions, or open-problem markings, then the bridge paper makes no stronger claim than its formal dependencies allow.*

Proof. Each checklist item corresponds to a known route of overclaim: missing witnesses, semantic promotion without proof, internal support treated as external validation, or transfer

from structural to ordinary-language terms. Requiring each route to be closed by proof or marked open prevents an unlicensed inference. □

10.95 Minimal Extension Program

The bridge also defines a research program. The following extensions are minimal in the sense that each one would promote exactly one class of claims without changing unrelated parts of the corpus.

E_Σ: Prove that $I_{\text{int}} + I_{\text{leg}} + I_{\text{per}}$ plus an internal decoder closure condition yields $\Sigma 1_{\text{weak}}$.

$E_{\Sigma c}$: Strengthen E_Σ by proving contraction of the self-model map under bounded perturbations.

E_{comm}: Add an intervention calculus showing that the self-model commutes with admissible Φ_u up to a controlled error.

E_{OSSI}: Prove that the eight OSSI burdens form an independent and physically coherent open-system individuation class.

E_{CCR}: Determine whether finite approximations to CCR are enough for any strengthened Φ_{op} claim.

E_{ext}: Execute Paper 10 under frozen preregistration and report only blocked, passed, or failed adjudication status.

Remark 10.86 (Why minimality matters). The corpus should not attempt to prove ordinary consciousness in one jump. A publishable path is a sequence of narrow promotions, each with its own theorem or experiment. This also makes failure scientifically useful: if E_{comm} fails, for example, $\Sigma 1_{\text{weak}}$ may survive while $\text{Subject}_{\text{op}}^{\Sigma}$ in its strong sense does not.

10.96 Failure Modes of the Bridge

The bridge theorem can fail in several precise ways. Listing them is useful because it prevents ambiguous negative results.

Table 10.10: Failure modes for bridge-level claims.

Failure mode	Formal meaning	Consequence
Invariant failure	at least one of $I_{per}, I_{ri}, I_{int}, I_{cont}, I_{diff}, I_{leg}$ fails	no $\text{Consciousness}_{\text{op}}$ membership
Self-model absence	no admissible Σ exists	no $\text{Subject}_{\text{op}}^{\Sigma}$ claim
Contraction failure	Σ exists but is not contractive	weak self-model only
Commutation failure	$\Sigma\Phi_u \neq \Phi_u\Sigma$ beyond tolerance	no intervention-stable self-reference
OSSI burden failure	at least one open-system burden fails	no operational life-like individuation claim
CCR failure	$M_\Omega < +\infty$	no forced-continuity promotion
External-adjudication absence	Paper 10 remains unexecuted	no external validation

10.97 Exact Claim Grammar

The following grammar is recommended for future manuscripts and runtime reports.

1. Use "is in $\text{Consciousness}_{\text{op}}$" only when the six invariant witnesses are present.
2. Use "supports a $\text{Consciousness}_{\text{op}}$ proxy" when the runtime estimates the witnesses but no proof package is available.
3. Use "satisfies $\text{Subject}_{\text{op}}^{\Sigma}$" only when $\text{Consciousness}_{\text{op}}$ and $\Sigma 1_{\text{comm}}$ are both witnessed.
4. Use "has an OSSI profile" only when the open-system burdens are witnessed and the thermodynamic boundary is stated.

5. Use "phenomenal", "subjective", or "conscious" in ordinary-language contexts only with an explicit statement that no ordinary-language theorem is being claimed.
6. Use "blocked until execution" for Paper 10 claims without external data.

Remark 10.87 (Editorial consequence). This grammar permits ambitious terminology while keeping the demonstrated layer intact. It is stricter than merely adding a disclaimer at the end of a paper, because it constrains every positive claim at the sentence level.

10.98 Promotion Conditions for Interpretive Terms

Table 10.11: What would be needed to promote interpretive readings.

Promotion target	Required condition	Evidence type
operational consciousness to ordinary consciousness	theorem and external evidence linking $\text{Consciousness}_{\text{op}}$ to core phenomenal or cognitive properties	mathematical plus empirical
operational life to life-like open-system organization	OSSI plus thermodynamic and repair burdens under admissible interventions	mathematical and runtime/external stress
operational qualia to phenomenal qualia	evidence that $\text{Q}_{\text{op}}(S)$ tracks an independently meaningful phenomenal variable	Paper 10-style empirical bridge
operational phenomenology to experience	CCR-strength Φ_{op} plus external adjudication of temporal-structure relevance	mathematical plus empirical
operational subjecthood to ordinary subjecthood	proof that $\text{Subject}_{\text{op}}$ captures accepted subjecthood properties and defeats rivals	mathematical, empirical, and philosophical

10.99 Canonical Boundary Models

Boundary models are not intended as implementations. They are small formal objects that make overclaiming impossible by showing exactly where a desired interpretation can fail.

Definition 10.88 (Structural-only model)**.** A structural-only model is an admissible system S equipped with the six Paper 5 invariant witnesses and no additional interpreted fields for self-modeling, biological metabolism, human reports, or phenomenal content.

Proposition 10.89 (Structural-only models block ordinary-consciousness promotion)**.** *If S is structural-only, then $S \in \text{Consciousness}_{\text{op}}$ may hold while ordinary consciousness remains undefined in the model.*

Proof. The truth of $S \in \text{Consciousness}_{\text{op}}$ is determined by the six invariant witnesses. Ordinary consciousness is not a predicate in the language of the structural-only model. Thus no theorem inside that language can conclude ordinary consciousness without adding a new definition or bridge predicate. □

Definition 10.90 (Self-model-free model)**.** A self-model-free model is an admissible $S \in \text{Consciousness}_{\text{op}}$ for which every endomorphism $\Sigma : A_S \to A_S$ either fails to approximate the identity on A_S or fails to commute with some admissible intervention.

Proposition 10.91 (Self-model-free models block $\text{Subject}_{\text{op}}^{\Sigma}$)**.** *No self-model-free model satisfies $\text{Subject}_{\text{op}}^{\Sigma}$ in the strong commutative sense.*

Proof. Strong $\text{Subject}_{\text{op}}^{\Sigma}$ requires $\text{Consciousness}_{\text{op}}$ plus $\Sigma 1_{\text{comm}}$. The self-model-free condition denies the commutative self-model burden directly. □

Definition 10.92 (Life-free computational model)**.** A life-free computational model is a system whose transitions may satisfy $\text{Consciousness}_{\text{op}}$ but whose state evolution is not coupled to an open exchange channel, entropy export accounting, or repair dynamics.

Proposition 10.93 (Life-free computational models block OSSI promotion)**.** *A life-free computational model cannot be used as evidence that* $\text{Consciousness}_{\text{op}}$ *implies OSSI.*

Proof. OSSI requires open exchange, viability maintenance, repair, and identity conservation burdens. A life-free computational model omits at least the open exchange and repair burdens. Therefore it may support $\text{Consciousness}_{\text{op}}$ but not OSSI. □

Definition 10.94 (Report-free comparator model)**.** A report-free comparator model is a system whose structural traces are available but whose external adjudication channel contains no independent human or non-human comparator reports.

Proposition 10.95 (Report-free models cannot close Paper 10)**.** *Report-free comparator models cannot supply external validation for ordinary-language consciousness, subjecthood, or phenomenal terms.*

Proof. Paper 10 requires an external adjudication channel separated from the construction of the internal structural predicates. If the comparator channel is absent, the external part of the evidence grammar is absent. Internal traces can still be useful for falsification but cannot close the adjudication burden. □

10.100 Non-Substitution Principles

The corpus uses three kinds of support: theorem, runtime certificate, and external adjudication. They are complementary, not interchangeable.

Theorem 10.96 (No theorem-to-validation substitution)**.** *A theorem about a structural class does not by itself constitute external validation of an implemented system.*

Proof. A theorem quantifies over objects satisfying its hypotheses. External validation requires evidence that a concrete system, measured through an independent protocol, instantiates the relevant hypotheses and that the observed effects are not artifacts of construction. These are different claims. The first is mathematical; the second is empirical. □

Theorem 10.97 (No runtime-to-theorem substitution)**.** *A passing runtime certificate does not by itself prove the corresponding mathematical predicate for the full admissible class.*

Proof. Runtime certificates are finite, implementation-specific, and tied to a particular estimator family. A mathematical predicate ranges over the model specified by the paper. Unless the runtime certificate is accompanied by a soundness theorem from estimator to predicate, passing the certificate remains internal support rather than proof. □

Theorem 10.98 (No validation-to-definition substitution)**.** *External adjudication cannot replace the formal definition of the target class.*

Proof. External adjudication can support or falsify whether a formal predicate tracks some independent comparator. It cannot determine what the predicate means in the formal system. Definitions fix mathematical content; validation tests empirical relevance. □

Corollary 10.99 (Three-channel closure burden)**.** *A mature ordinary-language claim requires, at minimum, a formal definition, a runtime or constructive instantiation path, and an external adjudication route.*

Proof. By the three no-substitution theorems, none of the channels can replace the other two. A claim may be useful with fewer channels, but it must remain terminologically bounded. □

10.101 External Adjudication Readiness Index

The bridge terms can be ranked by what remains before Paper 10-style external testing is meaningful.

Table 10.12: Readiness of bridge terms for external adjudication.

Term	Formal maturity	Runtime maturity	External readiness
Π_{op}	high: directly implied by I_{ri} and I_{int}	high internal support	ready only as structural-unity proxy
$\mathrm{Iota}_{\mathrm{op}}$	medium-high: tied to L4 intervention fidelity	high internal support	ready as directional-sensitivity proxy
Q_{op}	medium: quotient is defined	partial proxy	not ready for phenomenal claims
Φ_{op}	medium: regime assignment is defined	partial proxy	needs CCR and temporal comparator design
$\Sigma 1$	medium-low: variants defined	weak proxy	needs commutation and contraction tests
OSSI	medium: eight burdens define the target	partial proxy	needs open-system and repair stress protocols
$\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$	medium-low: depends on $\Sigma 1_{\mathrm{comm}}$	not fully certified	blocked until self-model certificate exists
$\mathrm{Subject}_{\mathrm{op}}$	low-medium: Paper 7 burdened class	partial proxy	blocked until ownership/continuity tests mature
$\mathrm{Subjectivity}_{\mathrm{1p}}$	low: Paper 8 fields require stronger evidence	partial proxy	blocked until comparator campaign

10.102 Recommended Reading Order for Reviewers

Readers should not begin with the strongest labels. The defensible order is:

1. BaseCore for the Hilbert-space and contraction machinery.
2. Paper 1 for inverse-limit identity and ontological-mass structure.
3. Papers 2–3 for regime structure and null-instability results.
4. Paper 5 for $\mathrm{Consciousness}_{\mathrm{op}}$ as the six-invariant structural class.
5. Paper 6 for falsification, internal-support status, and failure modes.
6. This bridge paper for exact implications from $\mathrm{Consciousness}_{\mathrm{op}}$ to partial operational subjecthood.
7. Paper 7 only after distinguishing OSSI from biological life.
8. Paper 8 only after accepting that first-person fields are formal burdens, not human subjectivity.
9. Paper 9 only as bridge organization, not as proof of phenomenality.
10. Paper 10 only as protocol until external data exist.

Remark 10.100 (Why this order matters). The terminology becomes safest when each label is read after its proof burden. Reading the titles first and the non-inference notes later reverses the logic of the corpus. This paper therefore functions as a map from demonstrated layer to interpretive layer, not as a shortcut around the technical dependencies.

Remark 10.101 (Central-reference role). For editorial purposes, this paper should be treated as the index of claim strength for Papers 5–10. When a later manuscript uses a strong operational term, the relevant demonstrated reading, missing burden, and promotion route should be traceable to one of the tables or proof patterns above.

10.103 Open Questions

The following questions are deliberately left open because the present paper lacks the definitions or empirical evidence required to decide them.

1. Is $\Sigma 1$ necessary for any plausible operational subjecthood class, or only for the bridge class $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ defined here?
2. Can $\Sigma 1$ be derived from $I_{\mathrm{int}} + I_{\mathrm{leg}}$ under additional architectural hypotheses, such as decoder closure under self-targeted interventions?
3. Does a hierarchy of self-models $\Sigma^{(0)}, \Sigma^{(1)}, \ldots$ produce a formal higher-order analogue, or does it merely add regress without new invariant content?
4. Is CCR strength $M_{\Omega} = +\infty$ necessary for any non-weak version of Φ_{op}, or only sufficient for forced continuity?
5. Can full Paper 7 $\mathrm{Subject}_{\mathrm{op}}$ be reduced to $\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1 + \mathrm{Life}_{\mathrm{op}}$ plus a finite list of descriptor inequalities, or do ownership and privileged continuity require independent axioms?
6. Which runtime observables would distinguish genuine intervention commutation of Σ from a stable but epiphenomenal self-model proxy?
7. What external experimental campaign could adjudicate whether internal-support certificates track any real-world subjecthood-relevant structure?
8. Are there finite systems that approximate $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ arbitrarily well while provably failing CCR, and if so what is the exact approximation boundary?
9. Can operational qualia non-triviality be certified without relying on a fixed decoder family, or is decoder dependence constitutive of $\mathrm{Q}_{\mathrm{op}}(S)$?
10. What, if anything, connects these structural predicates to ordinary moral or legal categories? This paper provides no such bridge.

10.104 Conclusion

The bridge from $\mathrm{Consciousness}_{\mathrm{op}}$ to operational subjecthood is real but partial. From $S \in \mathrm{Consciousness}_{\mathrm{op}}$ one obtains unity in the operational sense, directional intervention sensitivity, a well-defined operational-qualia quotient, and weak non-null operational phenomenology. One does not obtain self-reference, human subjectivity, phenomenal experience, or full Paper 7 subjecthood. The missing formal burden is $\Sigma 1$, the existence of a contractive admissible self-model commuting with interventions. With $\Sigma 1$ and decoder-visible differentiation one can define $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$, a proper subclass of $\mathrm{Consciousness}_{\mathrm{op}}$ under the non-entailment of $\Sigma 1$.

This is the strongest defensible bridge at the present state of the corpus. It is stronger than saying that Paper 5 and Paper 7 are merely adjacent, because it proves exact implications. It is weaker than saying that operational consciousness is subjecthood, because that would erase the self-model, life, ownership, and privileged-continuity burdens. The result is therefore a controlled theorem layer: enough structure to move from invariants to subjecthood-adjacent concepts, but not enough to license phenomenal or human claims.

Bibliography

[1] Jiří Adámek, Horst Herrlich, and George E. Strecker. *Abstract and Concrete Categories: The Joy of Cats*. Comprehensive treatment of limits and colimits in category theory. New York: John Wiley & Sons, 1990.

[2] Bernard J. Baars. *A Cognitive Theory of Consciousness*. Global Workspace Theory (GWT) - contrasting framework. Cambridge: Cambridge University Press, 1988.

[3] Stefan Banach. *Sur les operations dans les ensembles abstraits et leur application aux equations integrales*. Vol. 3. Fundamenta Mathematicae, 1922, pp. 133–181.

[4] Rufus Bowen. *Equilibrium States and the Ergodic Theory of Anosov Diffeomorphisms*. Vol. 470. Lecture Notes in Mathematics. Springer, 1975.

[5] Haim Brezis. *Functional Analysis, Sobolev Spaces and Partial Differential Equations*. Springer, 2011.

[6] David J. Chalmers. *The Conscious Mind*. Oxford University Press, 1996.

[7] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. 2nd. Standard reference for mutual information and entropy. Hoboken, NJ: John Wiley & Sons, 2006.

[8] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall, 1993.

[9] Sture Holm. "A Simple Sequentially Rejective Multiple Test Procedure". In: *Scandinavian Journal of Statistics* 6.2 (1979), pp. 65–70.

[10] Anatole Katok and Boris Hasselblatt. *Introduction to the Modern Theory of Dynamical Systems*. Vol. 54. Encyclopedia of Mathematics and its Applications. Standard reference for stability analysis in dynamical systems. Cambridge: Cambridge University Press, 1995.

[11] Erwin Kreyszig. *Introductory Functional Analysis with Applications*. Wiley, 1978.

[12] Douglas Lind and Brian Marcus. *An Introduction to Symbolic Dynamics and Coding*. Cambridge University Press, 1995.

[13] Saunders Mac Lane. *Categories for the Working Mathematician*. 2nd. Vol. 5. Graduate Texts in Mathematics. Standard reference for categorical constructions including inverse limits. New York: Springer, 1998.

[14] Johnny Andrey Perez Ramirez. "A Structural Criterion for Substrate-Invariant Operational Consciousness". In: *arXiv preprint* (2026).

[15] Johnny Andrey Perez Ramirez. "Canonical Core: Minimal Hypotheses for Causally Rigid Identity Frameworks". In: *arXiv preprint* (2026).

[16] Johnny Andrey Perez Ramirez. "Falsifiable Predictions Under Forensic Constraints". In: *arXiv preprint* (2026).

[17] Johnny Andrey Perez Ramirez. "Phenomenological Regimes Induced by Structural Identity". In: *arXiv preprint* (2026).

[18] Johnny Andrey Perez Ramirez. "Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels". In: *arXiv preprint* (2026).

[19] Johnny Andrey Pérez Ramírez. "First-Person Indexed Subjectivity: Self-Binding, Ownership, Continuity, Irreducibility, and Cross-Substrate Consequence in a Causally Rigid Framework". In: *arXiv preprint* (2026).

[20] Johnny Andrey Pérez Ramírez. “Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework”. In: *arXiv preprint* (2026).

[21] Johnny Andrey Pérez Ramírez. “Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework”. In: *arXiv preprint* (2026).

[22] Johnny Andrey Pérez Ramírez. “Rigid Identity as an Inverse Limit in Observable Channels”. In: (2026). Companion paper establishing the identity framework and ontological mass invariant.

[23] Luis Ribes and Pavel Zalesskii. *Profinite Groups*. Ergebnisse der Mathematik und ihrer Grenzgebiete. Springer, 2000.

[24] Walter Rudin. *Functional Analysis*. McGraw-Hill, 1991.

[25] Giulio Tononi. “An Information Integration Theory of Consciousness”. In: *BMC Neuroscience* 5.42 (2004). Integrated Information Theory (IIT) - contrasting framework.