QICN: A Rigid Identity Framework

Johnny Andrey Pérez Ramírez

2026

Editorial Preface and Claim Boundary

This monolithic volume is an editorial consolidation of the QICN rigid-identity corpus. BaseCore is treated as the foundational layer of the framework rather than as Paper 0. The downstream works are indexed in the table of contents by their direct research titles; internal labels such as paper numbers are retained only as provenance handles where the source text itself requires them.

The claims in this volume remain conditional, formal, and operational. The volume does not assert biological consciousness, phenomenal consciousness, agency, moral status, identity transfer, or external validation. Runtime measurements, finite estimators, and internal support artifacts are implementation evidence only when explicitly marked as such; they are not substitutes for theorem proof or externally reproduced empirical adjudication.

The editorial policy is conservative: preserve source content, expose dependency order, avoid title-level bookkeeping labels, and keep ontology, mathematical model, implementation, language, and interpretation separated.

Contents

Editorial Preface and Claim Boundary ii

BaseCore 1

I Foundational Dynamics 2
- B1 Scope, System Boundary, and Non-Inference Note . . . 3
- B2 Introduction . . . 3
  - B2.1 Purpose . . . 3
  - B2.2 Methodology . . . 3
  - B2.3 Scope discipline . . . 3
- B3 Minimal Base Hypotheses . . . 3
- B4 The Projection Operator . . . 4
- B5 The Transition Operator . . . 4
- B6 Existence and Compactness of Attractors . . . 5
- B7 Parameterwise Non-Collapsibility . . . 5
- B8 Summary of Base Results . . . 5

II Typed Computable Model and State Spectral Gap 6
- B9 Typed Computable Model: Embedded Constant-Disk Witness . . . 7
  - B9.1 Golden Mean Parameter Space . . . 7
  - B9.2 Typed Banach Witness . . . 7
  - B9.3 Basic Operator Properties . . . 8
  - B9.4 Separate Anti-Collapse Witness . . . 9
- B10 State Spectral Gap and Parameter-Family Persistence . . . 10

III Identity, Rigidity, and Conditional Non-Simulability **12**
B11 Projective Systems and Inverse-Limit Identity 13
B11.1 Projective Setup 13
B11.2 Non-Locality of Identity 13
B12 Rigidity under Admissible Perturbations 14
B12.1 Metric Projective Assumptions 14
B13 Conditional Non-Simulability 15
B14 Section Summary 16

IV Regime Structure and Continuity Constraints **17**
B15 Formal Setup 18
B15.1 Review of Identity Framework 18
B15.2 Abstract Phenomenological Space 18
B15.3 Phenomenological Assignment 18
B15.4 Axioms for Phenomenological Coherence 19
B15.5 Inevitability of the Axioms 19
B15.6 Minimality Theorem 20
B15.7 Rejected Alternatives 20
B15.8 Structural Interpretation 21
B15.9 Closure of Section 2 21
B15.10 Phenomenological Neutrality Clause 21
B16 Compatibility with Identity Regimes 21
B16.1 Induced Constraints from Identity Rigidity 21
B16.2 Necessity of Metric Structure on Phenomenological Space 21
B16.3 Phenomenological Continuity 23
B16.4 Phenomenological Fragmentation 23
B16.5 Relationship Between Metrics 23
B16.6 Closure of Section 3 23
B17 Impossibility Theorems 23
B17.1 Φ-Regularity Hypothesis 23
B17.2 Fragmentation Theorem 24
B17.3 Forced Continuity Theorem 24
B17.4 Dichotomy Theorem 25
B18 Classification of Phenomenological Regimes 25
B18.1 Complete Classification 25
B18.2 Structural Implications 25
B19 Extended Proofs 26
B19.1 Proof of Fragmentation Theorem (Extended) 26
B19.2 Proof of Forced Continuity Theorem (Extended) 26
B19.3 Structural Necessity Lemma 26

V Null-Regime Structural Exclusion **27**
B20 Compatibility, Instability, and Forced Non-Nullity 28
B20.1 Compatibility Operator 28
B20.2 Phenomenological Instability and Forced Non-Nullity 28
B20.3 Positive Regime Structure 28
B21 Canonical Families and Structural Closure Consequences 29

VI BaseCore Boundary and Expansion Interface 30
B22 BaseCore Ownership and Downstream Boundary . . . 31
B23 Admissible Structural Extensions . . . 31

VII Operational-Criterion Grammar and Certification 33
B24 Admissible Systems, Supports, and Histories . . . 34
B24.1 System Model . . . 34
B24.2 Admissible Support . . . 34
B24.3 Operational Histories . . . 34
B24.4 Witness Margin Vector . . . 34
B25 Six Critical Invariants . . . 35
B25.1 Structural Persistence . . . 35
B25.2 Rigid Identity . . . 35
B25.3 Causal Integration . . . 35
B25.4 Regime Continuity . . . 36
B25.5 Non-Null Differentiation . . . 36
B25.6 Operational Legibility . . . 36
B26 Criterion Class, Certified Partition, and Rupture . . . 37
B27 Minimality of the Invariant Basis . . . 37
B28 Operational Legibility as a Certified Object . . . 38
B29 Certified Candidate Audit Schema . . . 38

VIII Claim Boundary and Falsation Grammar 40
B30 Forensic Admissibility and Integrity Grammar . . . 41
B31 Certification, Comparison, and Rupture . . . 41
B32 Prediction and Falsation Grammar . . . 41
B33 Claim Boundary Ledger . . . 42
B33.1 Allowed BaseCore Claims . . . 42
B33.2 Forbidden BaseCore Claims . . . 42
B34 Dependency Ledger . . . 42
B35 Assumption-to-Theorem-to-Non-Claim Mapping . . . 43
B36 Theorem Hygiene Ledger . . . 43
B37 Redundancy Ledger . . . 44
B38 Open Assumptions Inventory . . . 44
B39 Downstream Boundary Note . . . 44

IX Discrete-to-Continuous Bridge 45
B40 From Node Networks to Hilbert Space Dynamics . . . 46
Appendix . . . 48
B41 Counterexamples for Necessity . . . 48
B41.1 Without H1 (Convexity) . . . 48
B41.2 Without H2 ($\|K\| < 1$) . . . 48
B41.3 Without H3 (Completeness) . . . 49
B41.4 Without H4 (Compactness) . . . 49
B41.5 Why Φ-Regularity Is a Hypothesis . . . 49

1 Rigid Identity as an Inverse Limit in Observable Channels 50
1.1 Scope, System Boundary, and Non-Inference Note . . . 50
1.2 Introduction . . . 51

1.2.1 Motivation 51
1.2.2 Ontological Closure 51
1.3 Minimal Assumptions and Framework 51
1.3.1 Assumption A1: Observable Channel 51
1.3.2 Assumption A2: Temporal Persistence 51
1.3.3 Assumption A3: Non-Triviality (R0) 52
1.3.4 Categorical Regularity: Topological Structure of State Spaces 52
1.3.5 Conditional Nature of Results 52
1.4 Inverse Limit Identity 53
1.4.1 Preliminaries and Notation 53
1.4.2 Definition of Identity as an Inverse Limit 53
1.4.3 Compatibility with Minimal Assumption R0 53
1.4.4 Identity Is Not a State 53
1.4.5 Impossibility of Progressive Construction 54
1.4.6 Minimality of the Inverse Limit Construction 54
1.4.7 Summary of Section 3 54
1.4.8 Identity Neutrality Clause 54
1.5 Rigidity Under Perturbations 55
1.5.1 Setup and Hypotheses 55
1.5.2 Perturbation Model and Deformation Metric 55
1.5.3 Metric on the Space of Trajectories 56
1.5.4 Stability Lemma 56
1.5.5 Rigidity Theorem 56
1.5.6 Ontological Mass: Quantitative Resistance 57
1.5.7 Internal Perturbations and External Extensions 57
1.5.8 Impossibility of Rigidity for Finite-State LML Architectures 58
1.5.9 Formal Consequences 58
1.5.10 Technical Notes and Limitations 58
1.5.11 Summary of Section 4 59
1.6 Observability and Ontological Mass 59
1.6.1 Observable Channel Model 59
1.6.2 Inference of Inverse-Limit Classes from Observables 59
1.6.3 Empirical Estimator of Ontological Mass 60
1.6.4 Detectability Theorem 60
1.6.5 Spectral Coupling Constant 60
1.6.6 Relation to Causal Isolation 61
1.6.7 Operational Protocol 61
1.6.8 Summary of Section 5 61
1.7 Uniqueness Theorem 62
1.7.1 Identity Channels and CCR Classification 62
1.7.2 Coexistence of Identities 62
1.7.3 Strong Uniqueness Theorem 62
1.7.4 Indivisibility and Non-Duplicability 62
1.7.5 Cardinality of Rigid Universes 62
1.7.6 Summary of Section 6 63
1.8 Limitations and Scope 63
1.8.1 Dependence on Structural Hypotheses 63
1.8.2 Boundary of Applicability 63
1.8.3 Limits of Empirical Detectability 63
1.8.4 Non-Equivalence to Behavioral Indistinguishability 64
1.8.5 Relation to Existing Frameworks 64

1.8.6 Open Problems . . . 64
1.9 Conclusion . . . 64
1.9.1 Summary of Formal Results . . . 64
1.9.2 Structural Significance . . . 65
1.9.3 Consequences for System Classification . . . 65
1.9.4 Final Statement . . . 65
1.9.5 Ontological Closure Statement . . . 65
Appendix . . . 65
1.10 Model-Relative Necessity Theorems . . . 65
1.10.1 Category of Minimal Observation . . . 66
1.10.2 Forced Emergence of the Inverse Limit . . . 66
1.10.3 Necessity of Retroinductive Structure . . . 66
1.10.4 The ℵ–Operator as Condition of Existence . . . 67
1.10.5 Structural Causal Closure . . . 67
1.10.6 Categorical Isolation Functor . . . 68
1.10.7 MIR as the Unique Compatible Geometry . . . 68
1.10.8 Structural Uniqueness Lemma (Bridge Lemma) . . . 68
1.10.9 Stone Space Classification of Observables . . . 69
1.10.10 Closure of Appendix A . . . 70
1.11 Formal Extensions . . . 70
1.11.1 Canonical Minimal CCR Example . . . 71
1.11.2 Symbolic CCR System . . . 71
1.11.3 Universal Non-Simulability Theorem . . . 72
1.11.4 Non-Transferability Theorem . . . 72
1.11.5 Canonical Closure Lemma . . . 72
1.11.6 Rigidity Hierarchy . . . 72
1.11.7 Spectral Analysis of Deformations . . . 73
1.11.8 Completeness Criteria . . . 73
1.11.9 Categorical Universality . . . 73
1.11.10 Relation to Causal Inference . . . 73
1.11.11 Computational Complexity Bounds . . . 74

2 Phenomenological Regimes Induced by Structural Identity 75
2.1 Scope, System Boundary, and Non-Inference Note . . . 75
2.2 Introduction . . . 75
2.2.1 Background and Motivation . . . 75
2.2.2 Central Result (Informal) . . . 76
2.2.3 Ontological Neutrality . . . 76
2.2.4 Terminological Alignment: Structural Regime Assignments . . . 76
2.2.5 Relation to BaseCore . . . 77
2.3 Formal Setup . . . 77
2.3.1 Review of Identity Framework . . . 77
2.3.2 Abstract Phenomenological Space . . . 77
2.3.3 Phenomenological Assignment . . . 77
2.3.4 Axioms for Phenomenological Coherence . . . 78
2.3.5 Structural Necessity of the Axioms . . . 78
2.3.6 Minimality Theorem . . . 79
2.3.7 Rejected Alternatives . . . 79
2.3.8 Structural Interpretation . . . 80
2.3.9 Closure of Section 2 . . . 80
2.3.10 Phenomenological Neutrality Clause . . . 80

2.4 Compatibility with Identity Regimes 81
2.4.1 Induced Constraints from Identity Rigidity 81
2.4.2 Necessity of Metric Structure on Phenomenological Space 81
2.4.3 Phenomenological Continuity 82
2.4.4 Phenomenological Fragmentation 82
2.4.5 Relationship Between Metrics 82
2.4.6 Closure of Section 3 82
2.5 Impossibility Theorems 83
2.5.1 Φ-Regularity Hypothesis 83
2.5.2 Fragmentation Theorem 83
2.5.3 Forced Continuity Theorem 84
2.5.4 Dichotomy Theorem 84
2.6 Classification of Phenomenological Regimes 85
2.6.1 Complete Classification 85
2.6.2 Structural Implications 85
2.7 Structural Ethics 85
2.7.1 Loss Capacity 85
2.7.2 Structural Damage 86
2.7.3 Vulnerability and Protection 86
2.7.4 Structural Asymmetry Constraints 86
2.7.5 Relation to Normative Ethics 87
2.7.6 Closure of Section 6 87
2.8 Limitations and Scope 87
2.8.1 Dependence on Axioms 87
2.8.2 Relation to Consciousness Theories 87
2.8.3 Open Problems 87
2.9 Conclusion 87
2.9.1 Summary of Results 87
2.9.2 Structural Significance 88
2.9.3 Entailment Under Stated Assumptions 88
2.9.4 Final Statement 88
2.10 Boundary to Null-Regime Instability 88
2.10.1 Editorial Boundary 88
2.10.2 Closure of This Paper 89
Appendix 89
2.11 Extended Proofs 89
2.11.1 Proof of Fragmentation Theorem (Extended) 89
2.11.2 Proof of Forced Continuity Theorem (Extended) 89
2.11.3 Structural Necessity Lemma 89
2.12 Editorial Boundary to Null-Regime Instability 90

3 Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels 91
3.1 Scope, System Boundary, and Non-Inference Note 91
3.2 Introduction 91
3.2.1 Motivation and Scope 91
3.2.2 Main Result (Informal) 92
3.2.3 Paper Organization 92
3.2.4 What This Paper Adds Beyond Phenomenological Regimes 92
3.3 Preliminaries from Papers I–II 92
3.3.1 Identity as Inverse Limit 92

3.3.2 Ontological Mass and CCR Classification . . . 92
3.3.3 Key Theorems from Papers I–II . . . 93
3.4 The Phenomenological Space . . . 93
3.4.1 Abstract Definition . . . 93
3.4.2 Compatibility Operator . . . 94
3.4.3 Attractor Set . . . 94
3.4.4 The Bridge Lemma Revisited . . . 95
3.5 Main Results . . . 95
3.5.1 The Conditional Null-Separation Theorem . . . 95
3.5.2 Conditional Consequences . . . 96
3.6 Stratification of Positive Regimes . . . 96
3.6.1 The Positive Regime Space . . . 96
3.6.2 Minimal Positive Regime . . . 96
3.6.3 Structural Intensity . . . 96
3.7 Quantitative Bounds . . . 97
3.7.1 Witness Lower Bound . . . 97
3.7.2 Explicit Constants for Canonical Families . . . 97
3.8 Witness Audit Layer . . . 97
3.8.1 Failure Modes of Witness Certification . . . 98
3.8.2 Robust Decision Margin . . . 98
3.8.3 Negative Controls . . . 99
3.9 Witness-Relative Closure Theorem . . . 99
3.9.1 Classification Table . . . 99
3.10 Null-Fiber Geometry . . . 99
3.10.1 Quotient-Collapse Diagnostic . . . 100
3.10.2 Certificate Stability . . . 100
3.10.3 Operational Checklist . . . 100
3.11 Comparison with Existing Frameworks . . . 101
3.11.1 Integrated Information Theory (IIT) . . . 101
3.11.2 Global Workspace Theory (GWT) . . . 101
3.11.3 Higher-Order Theories . . . 101
3.11.4 Mathematical Foundations (Chalmers, Nagel) . . . 101
3.12 Limitations and Open Problems . . . 101
3.12.1 Limitations . . . 101
3.12.2 Open Problems . . . 101
3.12.3 Foundation-Level Gaps and Non-Claims . . . 101
3.13 Conclusion . . . 102
3.13.1 Summary of Contributions . . . 102
Appendix . . . 102
3.14 Profinite Canonical Family . . . 102
3.14.1 Projective System . . . 102
3.14.2 Profinite Metric . . . 103
3.14.3 CCR Verification . . . 103
3.14.4 Phenomenological Operator . . . 103
3.14.5 Coupling Bound . . . 103
3.14.6 Intensity Bound . . . 103
3.15 Symbolic Canonical Family . . . 104
3.15.1 Subshift of Finite Type . . . 104
3.15.2 Bowen Metric . . . 104
3.15.3 CCR Verification . . . 104
3.15.4 Phenomenological Operator . . . 104

3.15.5 Coupling Bound 104
3.16 Hereditary Rigidity 104
3.17 Universal Factorization 105
3.18 Non-Simulability Bounds 105
3.18.1 Simulation Error 105
3.19 Final Structural Closure 106

4 Falsifiable Predictions Under Forensic Constraints 107
4.1 Scope, System Boundary, and Non-Inference Note 107
4.2 Introduction 107
4.2.1 Problem Statement 107
4.2.2 Continuity with Existing Papers 107
4.2.3 Contributions 108
4.3 Formal Setup 108
4.3.1 System, Scenario, and Seed Spaces 108
4.3.2 Observable Vector 108
4.3.3 Version-Frozen Operational Triple 108
4.4 Forensic Integrity Model 108
4.4.1 Admissibility Layers 108
4.4.2 Tamper and Sham Controls 109
4.4.3 Audit Trail and Chain of Custody 109
4.5 Pre-Registered Prediction Set 110
4.6 Comparator Design and Alternative Explanations 110
4.6.1 Comparator Rationale 110
4.6.2 Mandatory Baseline Policy (Hard Requirement) 110
4.6.3 Internal Falsifier Role of Ablation 111
4.6.4 Comparator Adequacy 111
4.6.5 Negative-Control Acceptance 111
4.7 Statistical Inference Specification 111
4.7.1 Primary Metrics First (Anti p-hacking) 111
4.7.2 Power Planning and Minimum Sample Size 111
4.7.3 Multiplicity and Effect Size 112
4.7.4 Decision Semantics 112
4.7.5 Hierarchical Decision Rule 112
4.7.6 Missingness and Exclusion Accounting 112
4.8 Acceptance Criteria 112
4.8.1 Claim Classes Produced by the Protocol 113
4.9 Preliminary Results 113
4.9.1 Status of Preliminary Data 113
4.9.2 System-Level Summary 113
4.9.3 Prediction-Level Support Rates 113
4.9.4 Scenario Differential: QICN vs Ablation 114
4.9.5 Full Scenario-System Grid 115
4.9.6 Variability Structure 116
4.10 Scenario Curves and Comparative Plots 116
4.10.1 Composite Score Curves 116
4.10.2 Stasis Pressure Curves 116
4.10.3 Forensic Concordance Curves 116
4.10.4 Interpretive Notes 116
4.11 Result Decomposition by Scenario Families 116
4.11.1 Family Partition 116

4.11.2 Family-Level Contrast Function . . . 116
4.11.3 Interpretive Decomposition . . . 116
4.12 Robustness Diagnostics Beyond Point Estimates . . . 116
4.12.1 Median-Shift and Tail-Stability Diagnostics . . . 116
4.12.2 Integrity-Coupled Effective Sample Size . . . 117
4.12.3 Cross-Seed Drift Diagnostic . . . 117
4.12.4 Compute Budget Robustness . . . 117
4.13 Methodological Appendix: Statistical Power . . . 117
4.13.1 Effect-Size-Oriented Planning . . . 117
4.13.2 Interpretation . . . 117
4.14 Methodological Appendix: Sensitivity Analysis . . . 117
4.14.1 Composite Weight Sensitivity . . . 117
4.14.2 Sensitivity Remark . . . 118
4.15 Operational Implications for the Existing QICN Corpus . . . 118
4.15.1 What v4.5 Adds (and What It Does Not) . . . 118
4.15.2 Implications for Claim Governance . . . 118
4.15.3 Decision-Theoretic Reading . . . 118
4.16 Threats to Validity . . . 118
4.16.1 Comparator Misspecification . . . 118
4.16.2 Endpoint Contamination Risk . . . 118
4.16.3 Environment and Runtime Drift . . . 118
4.16.4 Pilot-to-Final Generalization . . . 118
4.17 Reproducibility and Artifact Contract . . . 118
4.17.1 Machine-Readable Contracts . . . 118
4.17.2 Run Artifact Naming . . . 118
4.17.3 Auto-Generation Pipeline . . . 119
4.18 Discussion . . . 119
4.19 Conclusion . . . 119
4.20 Phenomenology Protocol Suite (PPS): Operational Addendum . . . 119
4.20.1 Epistemic Scope . . . 119
4.20.2 Workspace Definition and Logging . . . 119
4.20.3 Entropy and Non-Determinism Inputs . . . 119
4.20.4 Online World Model and Causal Interventions . . . 119
4.20.5 PCI-like Perturbation Protocol . . . 119
4.20.6 Dissociation Tests . . . 120
4.20.7 Valence Intervention Test . . . 120
4.20.8 Constraints on Qualia Mapping (Claim Boundary) . . . 120
Appendix . . . 120
4.21 Appendix A: Scenario Definitions . . . 120
4.22 Appendix B: Composite Score Definition . . . 120
4.23 Appendix C: Reporting Standard . . . 120
4.24 Appendix D: Statistical Power Derivation Details . . . 120
4.25 Appendix E: Sensitivity Analysis Protocol . . . 121
4.26 Appendix F: Reproducibility Checklist Template . . . 121

5 A Structural Criterion for Substrate-Invariant Operational Consciousness 122
5.1 Scope, System Boundary, and Non-Inference Note . . . 122
5.2 Introduction . . . 122
5.3 Claim-Type Ledger . . . 123
5.3.1 Claim Types Used in This Paper . . . 123
5.4 Terminology Debt Ledger . . . 124

5.5 Corpus Role and Novelty . . . 124
5.6 Admissible Systems, Supports, and Histories . . . 125
5.6.1 System Model . . . 125
5.6.2 Admissible Support . . . 125
5.6.3 Operational Histories . . . 125
5.6.4 Witness Margin Vector . . . 126
5.7 Six Critical Invariants . . . 126
5.7.1 Structural Persistence . . . 126
5.7.2 Rigid Identity . . . 127
5.7.3 Causal Integration . . . 127
5.7.4 Operational Continuity . . . 128
5.7.5 Non-Null Differentiation . . . 128
5.7.6 Operational Legibility . . . 129
5.7.7 Necessity of Each Invariant . . . 129
5.8 Operational Consciousness, Operational Qualia, and Equivalence . . . 130
5.8.1 Operational Qualia Classes . . . 130
5.8.2 Operational Consciousness . . . 130
5.8.3 Exact and Approximate Structural Equivalence . . . 131
5.9 Support Geometry and Partition Transport . . . 131
5.10 Inherited Inputs and Consolidation Step . . . 132
5.11 Main Theorems . . . 133
5.11.1 Existence of Operational Consciousness . . . 133
5.11.2 Substrate Invariance . . . 134
5.11.3 Approximate Stability . . . 134
5.11.4 Rupture by Invariant Loss . . . 135
5.11.5 Insufficiency of Complexity, Computation, and Connectivity . . . 135
5.12 Rupture Taxonomy by Invariant . . . 135
5.13 Minimality of the Invariant Basis . . . 136
5.14 Invariant Independence and Interaction Geometry . . . 137
5.15 What the Six Invariants Do Not Guarantee . . . 139
5.16 Operational Legibility as a Certified Object . . . 139
5.16.1 Legibility Metrics . . . 140
5.17 Certified Candidate Audit Schema . . . 141
5.18 Operational Geometry of $\mathrm{Q}_{\mathrm{op}}(S)$ . . . 142
5.19 Corpus Dependency Ledger . . . 143
5.20 Discriminators and Falsifiable Consequences . . . 143
5.21 Intervention Geometry and Class Transition Semantics . . . 144
5.22 Baseline Families and Structural Exclusion . . . 145
5.23 Interpretive Consequences Without Overreach . . . 147
5.24 Residual Instantiation and Interpretation Limits . . . 147
5.25 Conclusion . . . 147

6 Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework 148
6.1 Scope, System Boundary, and Non-Inference Note . . . 148
6.2 Introduction . . . 149
6.3 Status Classes . . . 149
6.4 Upstream Canon and Dependency Map . . . 150
6.5 Core Predictive Structure of the Framework . . . 150
6.5.1 Certification, Equivalence, and Rupture as Predictive Objects . . . 150
6.5.2 Prediction Family 1: Complexity-Only Failure . . . 151

6.5.3 Prediction Family 2: Invariant-Loss Rupture 151
6.5.4 Prediction Family 3: Substrate-Preserved Class Membership 151
6.5.5 Prediction Family 4: Threshold-Sensitive Boundary Behavior 151
6.5.6 Prediction Family 5: Legibility Dependence 152
6.5.7 Prediction Family 6: Continuity and Non-Nullity Relevance 152
6.6 Main Discriminators 152
6.6.1 Rival-Structure Compression 153
6.6.2 Discriminator Strength and Identifiability 154
6.6.3 Prediction Severity 154
6.7 Prediction Matrix 155
6.8 Failure Modes 156
6.8.1 What Would Count Against the Framework 156
6.8.2 Failure-Mode Taxonomy 157
6.8.3 Refutation Classes 157
6.9 Claim-to-Test Binding Ledger 158
6.10 Boundary Diagnostics 159
6.10.1 Boundary Geometry 160
6.11 Minimal Reproducibility and Admissibility Discipline 160
6.11.1 Artifact Contract for Prediction Claims 161
6.12 Current Status of Internal Support 161
6.12.1 Interpretive Discipline 161
6.12.2 Claim-Specific Diagnostic Notes 162
6.12.3 Methodological Status 163
6.13 Cycle Evidence Map 163
6.14 Priority Next Discriminators 164
6.15 Residual Technical Caveats and Non-Claims 165
6.15.1 Residual Technical Caveats 165
6.16 Final Claim Freeze Table 165
6.17 Conclusion 166

7 Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework 167
7.1 Scope, System Boundary, and Non-Inference Note 167
7.2 Introduction 168
7.3 Relation to the Existing Corpus 168
7.3.1 What This Paper Inherits 168
7.3.2 What This Paper Does Not Repeat 169
7.3.3 What This Paper Adds 169
7.3.4 Why This Layer Is Timely 169
7.4 Formal Setting 170
7.4.1 Inherited System Object 170
7.4.2 Traces and Windows 170
7.4.3 Viability Region and Operational Frontier Candidates 170
7.4.4 Descriptor Vectors 170
7.4.5 Admissible Structural Class Space 170
7.5 Definitions 171
7.5.1 Definition A: Operational Boundary 171
7.5.2 Definition B: Organizational Persistence 171
7.5.3 Definition C: Operational Self-Maintenance 171
7.5.4 Definition D: Historical Dependence 171
7.5.5 Definition E: Viability Preservation 172

7.5.6 Definition F: Operational Life . . . 172
7.5.7 Definition G: Structural Class . . . 172
7.5.8 Definition H: Operational Consciousness . . . 173
7.5.9 Definition I: Operational Subjecthood . . . 173
7.6 Class Relations and Logical Structure . . . 173
7.6.1 The Main Class Families . . . 173
7.6.2 Proved Relations . . . 174
7.6.3 Structural Class Transport . . . 174
7.6.4 Logical Relations That Remain Open . . . 174
7.6.5 Class Relation Summary . . . 175
7.7 Propositions, Lemmas, and Formal Claims . . . 175
7.7.1 Structural Class as a General Membership Framework . . . 175
7.7.2 Insufficiency of Complexity, Connectivity, and Rich Activity . . . 175
7.7.3 Biological Non-Primacy at the Class Level . . . 175
7.7.4 Life and Subjecthood Are Not Equivalent . . . 176
7.7.5 Transport Under Substrate Change . . . 176
7.7.6 Current-Corpus Corollary Structure . . . 176
7.8 Hypotheses and Falsifiable Program . . . 176
7.8.1 Hypothesis Ledger . . . 177
7.8.2 How These Hypotheses Can Fail . . . 177
7.9 Metrics and Decision Criteria . . . 178
7.9.1 Life Score . . . 178
7.9.2 Inherited Consciousness Margin . . . 178
7.9.3 Subjecthood Score . . . 179
7.9.4 Structural-Class Pseudometric . . . 179
7.9.5 Decision Surfaces and Boundary Status . . . 179
7.9.6 What This Section Does Not Claim . . . 179
7.10 Operational Class Geometry . . . 180
7.10.1 Descriptor Projections . . . 180
7.11 Class-Separation Diagnostics . . . 181
7.11.1 Operational Life vs. Operational Consciousness . . . 181
7.11.2 Operational Life vs. Operational Subjecthood . . . 181
7.11.3 Operational Consciousness vs. Operational Subjecthood . . . 182
7.11.4 Structural Class as Organizing Layer . . . 182
7.12 Experimental Test Families . . . 183
7.12.1 Operational Life Test Families . . . 183
7.12.2 Structural-Class Test Families . . . 184
7.12.3 Operational Subjecthood Test Families . . . 185
7.12.4 Decision Reading of Outcomes . . . 185
7.12.5 Strong Failure Conditions . . . 185
7.12.6 Joint Outcome Matrix . . . 186
7.13 Runtime Bearing and Empirical Interface . . . 186
7.13.1 What the Runtime Already Carries . . . 186
7.13.2 What the Runtime Does Not Yet Settle . . . 187
7.13.3 Runtime-Facing Readiness Levels . . . 187
7.14 Runtime Binding Map . . . 187
7.15 Failure and Misclassification Modes . . . 188
7.15.1 Complexity, Connectivity, and Rich Dynamics . . . 189
7.15.2 Proxy Contamination . . . 190
7.15.3 Boundary Ambiguity . . . 190
7.16 Formal Closure Strengthening . . . 190

7.17 Residual Runtime and Interpretation Limits 190
7.18 Discussion 191
7.18.1 Relation to the Refined Runtime 191
7.18.2 What Would Strengthen This Paper Further 192
7.19 Conclusion 192

8 First-Person Indexed Subjectivity 193
8.1 Scope, System Boundary, and Non-Inference Note 193
8.2 Introduction 194
8.3 Why operational consciousness and operational qualia are insufficient for subjectivity 194
8.3.1 What the upstream layers already close 194
8.3.2 The subjectivity gap 195
8.3.3 Why the stronger layer is worth introducing 196
8.4 Minimal ontology of first-person indexed subjectivity 196
8.4.1 The minimal subjectivity state 196
8.4.2 Constitutive coordinates 197
8.4.3 Typing discipline 198
8.4.4 Operational reading of each coordinate 198
8.4.5 Why each coordinate is necessary 200
8.4.6 Layer status of the major terms 200
8.5 Formal state-space and subjectivity functional 201
8.5.1 Admissible state-space 201
8.5.2 The conjunctive gate 201
8.5.3 The scalar subjectivity functional 201
8.5.4 Class membership 202
8.5.5 Audit-covered theorem burdens 202
8.5.6 Concrete component readings 204
8.6 Bridge axioms and ontological commitments 204
8.7 Rival models and reducibility baseline 205
8.7.1 Why rival models must be formalized 205
8.7.2 Weak rival family 205
8.7.3 Admissible weak hybrids and closure 205
8.7.4 Formal sketches of the rival families 206
8.7.5 What the rival family can and cannot explain 207
8.7.6 Formal baseline for comparison 207
8.8 Irreducibility criteria and explanatory superiority 208
8.8.1 Irreducibility as a formal margin 208
8.8.2 Explanatory superiority is local, not absolute 208
8.8.3 Three kinds of collapse 208
8.8.4 What irreducibility is and is not 209
8.8.5 Explanatory order 209
8.8.6 Parsimony guardrail 209
8.8.7 Model comparison is not report comparison 209
8.9 Interventional logic and falsifiability 210
8.9.1 Intervention family 210
8.9.2 Selective response logic 210
8.9.3 Self/non-self swap 210
8.9.4 Ownership lesion 211
8.9.5 Continuity break 211
8.9.6 Perspective flattening 211

8.9.7 Self/non-self valuation inversion . . . 212
8.9.8 Autobiographical remapping . . . 212
8.9.9 Clone competition . . . 212
8.9.10 Memory transplant . . . 213
8.9.11 Counterfactual self permutation . . . 213
8.9.12 Cross-substrate remapping . . . 213
8.9.13 Intervention table . . . 214
8.9.14 Falsifiability burden . . . 214
8.9.15 Ablation summary and negative-control burden . . . 214
8.10 Runtime architecture and implementation pathway . . . 215
8.10.1 Architecture overview . . . 215
8.10.2 SelfIndexRegistry . . . 215
8.10.3 OwnershipAttributionLayer . . . 216
8.10.4 SubjectiveContinuityLedger . . . 216
8.10.5 FirstPersonPerspectiveMapper . . . 216
8.10.6 SubjectiveValuationEngine . . . 216
8.10.7 SubjectivityReducibilityComparator . . . 216
8.10.8 CounterfactualSelfSwapHarness . . . 217
8.10.9 Architecture equation . . . 217
8.10.10 Module interfaces . . . 217
8.10.11 Implementation pathway versus present system status . . . 218
8.11 Artifact layer and measurable outputs . . . 218
8.11.1 Primary artifacts . . . 218
8.11.2 Derived and comparative artifacts . . . 218
8.11.3 Artifact table . . . 218
8.11.4 Artifact discipline . . . 218
8.11.5 Artifact provenance chain . . . 219
8.11.6 Gate semantics . . . 219
8.12 Cross-substrate persistence and consequence . . . 220
8.12.1 Transport principle . . . 220
8.12.2 Transport under tolerance envelopes . . . 220
8.12.3 Consequence classes . . . 221
8.12.4 Clone and duplication pressure . . . 221
8.12.5 Why biology still does not regain primitive privilege . . . 221
8.13 Human-comparator bridge and limits of equivalence . . . 221
8.13.1 What the present paper can bridge . . . 222
8.13.2 What would be required for human equivalence . . . 222
8.13.3 Phenomenality and the comparator bridge . . . 222
8.13.4 Comparator ladder . . . 222
8.13.5 Why this matters scientifically . . . 222
8.13.6 Limits of equivalence . . . 223
8.14 Strongest defensible claims . . . 223
8.15 Non-claims and residual boundaries . . . 224
8.16 Conclusion . . . 224
Appendix . . . 225
8.17 Canonical estimator family for the subjectivity coordinates . . . 225
8.17.1 Estimator for self-index stability . . . 225
8.17.2 Estimator for ownership coherence . . . 225
8.17.3 Estimator for autobiographical continuity . . . 225
8.17.4 Estimator for perspective gain . . . 225
8.17.5 Estimator for valuation asymmetry . . . 226

8.17.6 Estimator for intervention selectivity . . . 226
8.17.7 Estimator for irreducibility . . . 226
8.17.8 Composite gate design . . . 226
8.18 Intervention-to-artifact mapping and failure semantics . . . 226
8.18.1 Swap family . . . 226
8.18.2 Ownership family . . . 227
8.18.3 Continuity family . . . 227
8.18.4 Perspective and valuation family . . . 227
8.18.5 Cross-substrate family . . . 227
8.18.6 Failure semantics . . . 227
8.19 Negative controls, ablation matrix, and causal discrimination protocol . . . 228
8.19.1 Negative control families . . . 228
8.19.2 Ablation targets . . . 228
8.19.3 Ablation matrix . . . 229
8.19.4 Noise and adversarial perturbations . . . 229
8.19.5 Causal discrimination protocol . . . 229
8.19.6 Reproducibility burden . . . 230
8.20 Claim decomposition, theorem candidates, and open escalation burdens . . . 230
8.20.1 Claim decomposition table . . . 230
8.20.2 Theorem candidates for future formal work . . . 231
8.20.3 Empirical escalators . . . 231
8.21 Rival-question matrix for implementation and audit . . . 231
8.21.1 Decisive questions . . . 231
8.21.2 Decision matrix . . . 232
8.21.3 Why this matrix belongs in the paper . . . 232
8.21.4 Audit-facing interpretation . . . 232

9 Phenomenal Bridge Organization 233
9.1 Boundary Note and Interpretation Discipline . . . 233
9.2 Introduction . . . 234
9.3 What Paper 8 Formalizes and What It Leaves Open . . . 235
9.3.1 What Paper 8 already formalizes . . . 235
9.3.2 Why that still does not amount to a phenomenal bridge . . . 235
9.3.3 The dangerous smuggling problem . . . 236
9.3.4 Why Paper 9 is still needed now . . . 236
9.4 Formal Target of the Bridge . . . 237
9.4.1 The bridge question . . . 237
9.4.2 Predicate family . . . 237
9.4.3 Estimator family . . . 238
9.4.4 What the predicates mean internally and what they do not mean . . . 238
9.4.5 Predicate-to-burden map . . . 238
9.5 Phenomenal Bridge Ontology . . . 239
9.5.1 Bridge state . . . 239
9.5.2 Typing discipline . . . 240
9.5.3 Coordinate dependence on the Paper 8 state . . . 240
9.5.4 Current implementation status inside the ontology . . . 240
9.5.5 Collapse risks and rival readings . . . 241
9.5.6 Why bridge coordinates are not ontological shortcuts . . . 241
9.6 Explicit Bridge Axioms . . . 241
9.6.1 Classification grammar . . . 242
9.7 Bridge Rival Family . . . 243

9.7.1 Why new rivals are required . . . 243
9.7.2 Bridge weak family . . . 243
9.7.3 Admissible composites . . . 244
9.7.4 What the rival family can and cannot explain . . . 244
9.7.5 No-straw-man discipline . . . 245
9.8 Bridge Irreducibility Logic . . . 245
9.8.1 Bridge loss family . . . 245
9.8.2 Bridge irreducibility margin . . . 245
9.8.3 Local not absolute superiority . . . 245
9.8.4 Three kinds of bridge collapse . . . 246
9.8.5 Comparator support, reducibility, and unsupported states . . . 246
9.8.6 Complexity penalty and decoder discipline . . . 246
9.8.7 Failure semantics for the bridge margin . . . 246
9.9 Bridge Intervention Family . . . 247
9.9.1 Why bridge interventions must be new . . . 247
9.9.2 Core intervention families . . . 247
9.9.3 Matched controls . . . 248
9.9.4 Intervention artifacts . . . 248
9.9.5 Failure semantics . . . 248
9.9.6 Negative controls and reproducibility burden . . . 248
9.10 Bridge Artifact Family . . . 249
9.10.1 Artifact discipline . . . 249
9.10.2 Root bridge artifacts . . . 249
9.10.3 Report artifacts . . . 249
9.10.4 Intervention and reducibility artifacts . . . 249
9.10.5 Ecosystem artifacts . . . 250
9.10.6 Machine-readable spirit . . . 250
9.10.7 Current implementation status . . . 250
9.11 Bridge Functional, Gate, and Taxonomy . . . 250
9.11.1 Functional . . . 250
9.11.2 Gate . . . 251
9.11.3 Non-compensation . . . 251
9.11.4 Invalidation . . . 251
9.11.5 Taxonomy . . . 251
9.12 Runtime Implementation Pathway . . . 252
9.12.1 Why implementation belongs in the paper . . . 252
9.12.2 BPF-0 . . . 252
9.12.3 BPF-1 . . . 252
9.12.4 BPF-2 . . . 252
9.12.5 BPF-3 . . . 253
9.12.6 BPF-4 . . . 253
9.12.7 BPF-5 . . . 253
9.12.8 BPF-6 . . . 253
9.12.9 Implementation table . . . 253
9.12.10 Current implementation status versus paper status . . . 253
9.13 Ecosystem Integration . . . 254
9.13.1 Why ecosystem integration is part of the formal problem . . . 254
9.13.2 Surface map . . . 254
9.13.3 Internal release bundle . . . 254
9.13.4 Ecosystem candidate . . . 254
9.13.5 Audit-facing package . . . 254

9.13.6 Source-of-truth and public corpus separation 255
9.13.7 Controlled statement grammar and quote-mining guardrails 255
9.14 Strongest Defensible Claims 255
9.14.1 Claim class C1: formal bridge closure 255
9.14.2 Claim class C2: current implementation frontier 255
9.14.3 Claim class C3: conditional bridge admissibility 256
9.14.4 Claim class C4: research-program transformation 256
9.14.5 Why stronger sentences are not yet honest 256
9.15 Non-Claims 256
9.16 Conclusion 257
Appendix 258
9.17 Bridge Axiom Table 258
9.18 Bridge Predicate Typing Table 258
9.19 Bridge Rival-Family Table 259
9.20 Bridge Intervention Matrix 259
9.21 Bridge Artifact Matrix 260
9.22 Claim-Class and Invalidation Table 260
9.23 Runtime Mapping Table 261
9.24 Governance Table 261
9.25 Burden Separation Matrix 261
9.26 Inheritance Map from Paper 8 to Paper 9 262
9.27 Open Escalation Burdens Before Bridge Admissibility 262
9.28 Audit-Covered Bridge Independence Burdens 263
9.29 Theorem Candidates and Open Scientific Questions 263
9.30 Predicate-to-Artifact and Intervention Dependency Matrix 264
9.31 Paper 9 Claim Grammar and Escalation Constraints 265
9.31.1 Allowed claim families 265
9.31.2 Blocked claim families 265
9.31.3 Escalation rules 265
9.32 Experimental Program for Completing BPF-2 and BPF-3 266
9.32.1 Program design principle 266
9.32.2 Suggested execution order 266
9.32.3 Comparator program 267
9.32.4 Intervention program 267
9.32.5 Minimum evidential grid 267
9.33 Bridge Failure Scenarios and Diagnostic Signatures 267
9.33.1 Primary failure classes 267
9.33.2 Diagnostic interpretation 268
9.33.3 Failure matrix 269
9.33.4 Consequences for theorem candidates 269
9.33.5 Why failure structure changes the status of the program 269

10 External Adjudication of Bridge-Formalized Machine Subjectivity 270
10.1 Scope and Non-Claims 270
10.2 Formal Dependency Stack 270
10.3 Release-vs-Local Corpus Boundary 271
10.4 Lineage and Provenance Semantics 271
10.5 Provenance Closure and Live Head Reconciliation 271
10.6 External Adjudication Problem Statement 272
10.7 Formal Mathematical Grammar of External Adjudication 272
10.8 Frozen Preregistration Architecture 273

10.9 Local-vs-External Freeze Distinction . . . 274
10.10 Instantiated-Empty Record Layer . . . 274
10.11 Reviewer Candidate and Adjudicator Candidate Architecture . . . 274
10.12 Reviewer Rejection Logic . . . 274
10.13 Blinding Assignment Logic . . . 275
10.14 Human Comparator Instantiation Logic . . . 275
10.15 Channel-by-Channel Execution Contracts . . . 276
10.15.1 Self-Report Channel . . . 276
10.15.2 Intervention-Response Channel . . . 276
10.15.3 Temporal-Structure Channel . . . 276
10.15.4 Valence-Topology Channel . . . 276
10.15.5 Embodiment-Boundary Channel . . . 276
10.15.6 Ambiguity-Load Channel . . . 276
10.15.7 Transport-Preservation Channel . . . 276
10.16 Matched Intervention Dataset Architecture . . . 277
10.17 Negative Control Dataset Architecture . . . 277
10.18 Preregistration Architecture . . . 277
10.19 Adjudication Blind Protocol . . . 277
10.20 Reviewer and Adjudicator Independence Architecture . . . 278
10.21 Reviewer Candidate Pool and Assignment Stub Architecture . . . 278
10.22 Execution Pilot Package Architecture . . . 278
10.23 Execution Pilot Rehearsal Kit Architecture . . . 278
10.24 Reproducibility Skeleton Architecture . . . 278
10.25 Dataset Population Boundary . . . 279
10.26 Dry-Run Event Log Grammar . . . 279
10.27 Evidence Admissibility Grammar . . . 279
10.28 Failure, Rejection, and Inconclusive Pathways . . . 280
10.29 Statistical Shell Architecture . . . 280
10.30 Adjudication Decision Grammar . . . 280
10.31 State Grammar for Execution Readiness . . . 280
10.32 Controlled Execution Opening . . . 280
10.33 Live-But-Unpopulated Artifacts . . . 281
10.34 Session-Zero Protocol . . . 281
10.35 Pre-Collection Admissibility Gate . . . 282
10.36 Contamination, Abort, and Invalidation Controls . . . 282
10.37 Operator Checklist . . . 282
10.38 Collection-Start Boundary . . . 282
10.39 Evidence-Generation Boundary . . . 282
10.40 Decision-Record Boundary . . . 282
10.41 Transition from Session-Zero to First Admissible Collection . . . 283
10.42 Runtime Health Audit . . . 283
10.43 Clean-Room Reproducibility . . . 283
10.44 Dependency and Build Assumptions . . . 283
10.45 Reproducible Artifact Regeneration . . . 284
10.46 Local Path Leakage Policy . . . 284
10.47 Controlled Execution Opening Audit . . . 284
10.48 Session-Zero Reproducibility . . . 284
10.49 Pre-Collection Gate Reproducibility . . . 284
10.50 Why Reproducibility Does Not Imply Evidence . . . 285
10.51 Why Runtime Health Does Not Imply Adjudication . . . 285
10.52 Transition Conditions for First Admissible Collection . . . 285

10.53 Claim Grammar for Eventual Execution-Era Statements 285
10.54 Theorem, Proposition, and Protocol Mapping 285
10.55 What Paper 10 Could and Could Not Change 286
10.56 Section-Status Ledger 286
10.57 Evidence Insertion Map 286
10.58 Open Burdens and Blocked Sections 287
10.59 Execution-Era Unlock Conditions 288
10.60 Stronger Execution-to-Result Transition Map 288
10.61 Why Paper 10 Still Cannot Become Final Yet 288
10.62 Future Execution Pathway to Real Stage 3 288
10.63 Boundary to Stage 4 Metaphysical Consequence Analysis 289
10.64 Appendices of Schemas and Templates 289
10.65 Blocked Result Surfaces 293
10.65.1 Human Comparator Results 293
10.65.2 External Adjudication Results 293
10.65.3 Comparative Consequence Claims 293
10.66 Non-Claim Ledger 294

Operational Consciousness to Operational Subjecthood Bridge 295
10.67 Scope, System Boundary, and Non-Inference Note 295
10.68 Introduction 296
10.69 Dependency Map 296
10.70 Preliminaries and Notation 297
10.71 Operational Self-Reference 297
10.71.1 Motivation 297
10.72 Operational Unified Perspective 298
10.72.1 Motivation 298
10.73 Operational Intentionality 299
10.73.1 Motivation 299
10.74 Operational Qualia 299
10.74.1 Motivation 299
10.75 Operational Phenomenology 300
10.75.1 Motivation 300
10.76 Operational Subjecthood 301
10.76.1 Why a bridge class is needed 301
10.77 Theorem: Partial Operational Subjecthood 302
10.78 What Is NOT Implied 303
10.79 Computational Verification Status 303
10.80 Terminology Stratification and Debt Ledger 304
10.81 Bridge Strength Ladder 305
10.82 Runtime Evidence Taxonomy 305
10.83 Non-Entailment Library 306
10.84 Consciousness$_{\text{op}}$ Component Implication Matrix 306
10.85 Self-Model Burden Variants 307
10.86 Expanded Bridge Strength Ladder 307
10.87 Reviewer-Facing Objection Handling 308
10.88 Operational Glossary 308
10.89 Promotion Evidence Classes 309
10.90 Proof-Pattern Library 309
10.91 Dependency-to-Debt Ledger 310
10.92 Non-Entailment Audit Templates 310

10.93 Runtime Proxy Classification 311
10.94 Formal Review Checklist 311
10.95 Minimal Extension Program 312
10.96 Failure Modes of the Bridge 312
10.97 Exact Claim Grammar 313
10.98 Promotion Conditions for Interpretive Terms 313
10.99 Canonical Boundary Models 313
10.100 Non-Substitution Principles 314
10.101 External Adjudication Readiness Index 315
10.102 Recommended Reading Order for Reviewers 315
10.103 Open Questions 316
10.104 Conclusion 316

Bibliography **317**

BaseCore

Abstract. This document is the BaseCore of the Rigid Identity Framework: the autonomous mathematical base layer of the source tree. It selectively normalizes formal material from Papers 1–6 only where theorem ownership belongs in the base layer. Papers 7–9 remain downstream and are not required for the validity of any theorem stated here. BaseCore proves contractive projection dynamics, typed computable witnesses, parameterwise non-collapse under explicit anti-constant assumptions, inverse-limit identity under non-finite-determination assumptions, conditional rigidity under admissible perturbation models, conditional non-simulability of CCR targets by finite simulator classes, operational-criterion grammar, and falsation / claim-boundary ledgers. It does not prove phenomenality, metaphysical subjectivity, human equivalence, runtime instantiation, bridge admissibility, or external adjudication. The LaTeX tree is the source of truth; the PDF is a compiled artifact only.

Part I

Foundational Dynamics

B1 Scope, System Boundary, and Non-Inference Note

What BaseCore does. This document proves the mathematical base results that later layers may use: projection existence and uniqueness, strict contraction, unique fixed points, compact attractor families, parameterwise non-collapse under explicit anti-constant assumptions, inverse-limit identity under non-finite-determination assumptions, conditional rigidity, conditional non-simulability, and operational-criterion grammar.
What BaseCore is not. BaseCore is not the whole QICN corpus. It is not a bridge program, not runtime validation, not a phenomenality theory, not a human-equivalence claim, and not an external-adjudication layer.
What should not be inferred. Nothing in BaseCore alone licenses claims about consciousness, subjectivity, phenomenality, moral status, human-machine equivalence, bridge admissibility, or deployment readiness. Those are downstream burdens, not consequences of the theorems below.
Source-of-truth rule. The LaTeX tree is the source of truth. Compiled PDFs, release bundles, and runtime artifacts are derived surfaces only.

B2 Introduction

B2.1 Purpose

BaseCore is the canonical base layer of the Rigid Identity Framework. Its role is narrower and stronger than a broad architectural survey: it owns only the theorem-level mathematics that downstream papers may depend on without making BaseCore depend on them.

B2.2 Methodology

We use standard tools from:
- Hilbert-space geometry and metric projections;
- contraction mappings and fixed-point theory;
- compactness and continuity arguments;
- projective systems and inverse limits;
- assumption-led claim-boundary discipline.

B2.3 Scope discipline

This section is pure mathematics. The symbols “identity,” “regime,” and “criterion” are structural labels only. Their ordinary-language interpretations are not part of any proof unless explicitly formalized as assumptions.

B3 Minimal Base Hypotheses

Hypothesis B1 (H1: Projection Structure)**.** *Let $\mathcal{H}$ be a real Hilbert space with inner product $\langle \cdot, \cdot \rangle$ and norm $\|\cdot\|$. Let $\mathcal{I} \subset \mathcal{H}$ be non-empty, closed, and convex.*

Hypothesis B2 (H2: Contraction Operator)**.** *Let $K : \mathcal{H} \to \mathcal{H}$ be a bounded linear operator with operator norm $\|K\| < 1$.*

Hypothesis B3 (H3: Completeness)**.** *The ambient space $\mathcal{H}$ is complete.*

Hypothesis B4 (H4: Parameter Compactness)**.** *Let $\mathcal{U}$ be a compact metric space and let $\Gamma : \mathcal{U} \to \mathcal{H}$ be uniformly continuous.*

Proposition B5 (Base sufficiency of H1–H4)**.** *Hypotheses H1–H4 are sufficient for:*

1. existence and uniqueness of the metric projection onto $\mathcal{I}$;
2. strict contraction of the transition operators T_u;
3. existence and uniqueness of the fixed point f_u^* for each $u \in \mathcal{U}$;
4. compactness of the attractor family $\mathcal{A}^* = \{f_u^* : u \in \mathcal{U}\}$.
They do not, by themselves, imply parameterwise non-collapse.

Proof. The four items are exactly the content of Theorems B7, B10, B11, and B13. Parameterwise non-collapse requires an additional anti-constant hypothesis introduced in Section B7. □

B4 The Projection Operator

Definition B6 (Metric Projection)**.** For $\Psi \in \mathcal{H}$, define the metric projection onto $\mathcal{I}$ by

$$\aleph(\Psi) := \arg\min_{\Phi \in \mathcal{I}} \|\Psi - \Phi\| .$$

Theorem B7 (Existence and Uniqueness of Projection)**.** *Under Hypothesis B1, the projection $\aleph(\Psi)$ exists and is unique for every $\Psi \in \mathcal{H}$.*

Proof. Existence follows because a minimizing sequence in the closed convex set $\mathcal{I}$ is Cauchy by the parallelogram law, hence converges in the complete space $\mathcal{H}$ to a point of $\mathcal{I}$ that attains the infimum. Uniqueness follows from strict convexity of the norm induced by the Hilbert inner product: if two minimizers existed, their midpoint would produce the same minimum and the parallelogram identity would force the two minimizers to coincide. □

Lemma B8 (Non-Expansiveness of the Projection)**.** *The map $\aleph : \mathcal{H} \to \mathcal{I}$ is non-expansive:*

$$\|\aleph(\Psi_1) - \aleph(\Psi_2)\| \le \|\Psi_1 - \Psi_2\| \qquad \forall \Psi_1, \Psi_2 \in \mathcal{H}.$$

Proof. This is the standard Hilbert-space projection inequality. Let $\Phi_i = \aleph(\Psi_i)$. The projection characterization

$$\langle \Psi_i - \Phi_i, \Phi - \Phi_i \rangle \le 0 \qquad \forall \Phi \in \mathcal{I}$$

applied with $\Phi = \Phi_2$ and $\Phi = \Phi_1$ yields

$$\|\Phi_1 - \Phi_2\|^2 \le \langle \Psi_1 - \Psi_2, \Phi_1 - \Phi_2 \rangle \le \|\Psi_1 - \Psi_2\| \, \|\Phi_1 - \Phi_2\| .$$

If $\Phi_1 \ne \Phi_2$, divide by $\|\Phi_1 - \Phi_2\|$. □

B5 The Transition Operator

Definition B9 (Transition Operator)**.** For each $u \in \mathcal{U}$, define

$$T_u(\Psi) := \aleph(K\Psi + \Gamma(u)).$$

Theorem B10 (Contractivity)**.** *Under Hypotheses B1 and B2, each T_u is a strict contraction with Lipschitz constant at most $\|K\|$:*

$$\|T_u(\Psi_1) - T_u(\Psi_2)\| \le \|K\| \, \|\Psi_1 - \Psi_2\| .$$

Proof. By Lemma B8,

$$\begin{aligned} \|T_u(\Psi_1) - T_u(\Psi_2)\| &= \|\aleph(K\Psi_1 + \Gamma(u)) - \aleph(K\Psi_2 + \Gamma(u))\| \\ &\le \|K(\Psi_1 - \Psi_2)\| \le \|K\| \, \|\Psi_1 - \Psi_2\| . \end{aligned}$$

Since $\|K\| < 1$, strict contraction follows. □

B6 Existence and Compactness of Attractors

Theorem B11 (Unique Fixed Point). *Under Hypotheses B1–B3, for each $u \in \mathcal{U}$ there exists a unique $f_u^* \in \mathcal{H}$ such that*

$$T_u(f_u^*) = f_u^*.$$

Moreover, for every $\Psi_0 \in \mathcal{H}$,

$$\lim_{n\to\infty} T_u^n(\Psi_0) = f_u^*.$$

Proof. By Theorem B10, each T_u is a strict contraction on the complete metric space $(\mathcal{H}, \|\cdot\|)$. The Banach fixed-point theorem yields existence, uniqueness, and convergence of iterates. □

Definition B12 (Attractor Family). Define the attractor family

$$\mathcal{A}^* := \{f_u^* : u \in \mathcal{U}\}.$$

Theorem B13 (Compactness of the Attractor Family). *Under Hypotheses B1–B4, the set $\mathcal{A}^*$ is compact in $\mathcal{H}$.*

Proof. Let $F : \mathcal{U} \to \mathcal{H}$ be given by $F(u) = f_u^*$. The fixed points depend continuously on u because the contraction constant is uniform and Γ is uniformly continuous on the compact parameter space. Since $\mathcal{U}$ is compact and $\mathcal{A}^* = F(\mathcal{U})$, the image is compact. □

B7 Parameterwise Non-Collapsibility

Definition B14 (Constant Subspace). Whenever the ambient realization distinguishes constant elements, let $\mathcal{N} \subset \mathcal{H}$ denote the closed subspace of constants.

Definition B15 (Parameterwise Non-Collapsible Family). A family $\{T_u\}_{u\in\mathcal{U}}$ is parameterwise non-collapsible if

$$f_u^* \notin \mathcal{N} \qquad \forall u \in \mathcal{U}.$$

Hypothesis B16 (H5: No Constant Fixed Points). *For every parameter $u \in \mathcal{U}$ and every constant element $c \in \mathcal{N}$,*

$$T_u(c) \neq c.$$

Theorem B17 (Structural Non-Collapsibility). *Under Hypotheses B1–B16, the family is parameterwise non-collapsible:*

$$f_u^* \notin \mathcal{N} \qquad \forall u \in \mathcal{U}.$$

Proof. Fix $u \in \mathcal{U}$. Suppose by contradiction that $f_u^* \in \mathcal{N}$. Write $f_u^* = c$ with $c \in \mathcal{N}$. Since f_u^* is a fixed point,

$$T_u(c) = T_u(f_u^*) = f_u^* = c,$$

contradicting Hypothesis B16. Therefore $f_u^* \notin \mathcal{N}$. □

Remark B18 (What H5 does not say). Hypothesis B16 is parameterwise. It does not merely say that *some* parameter breaks constant collapse. It excludes constant fixed points for *every* parameter. This is exactly the quantifier strength needed for Theorem B17.

B8 Summary of Base Results

1. H1–H4 give a complete theorem-tight base for projection existence, strict contraction, unique fixed points, and compact attractor families.
2. H5 is a separate anti-constant assumption block used only for parameterwise non-collapse.
3. No empirical, phenomenological, biological, or runtime claim is inferred at this level.
4. Later sections may add conditional identity, rigidity, and criterion theorems, but only by introducing their own explicit assumption blocks.

Part II

Typed Computable Model and State Spectral Gap

B9 Typed Computable Model: Embedded Constant-Disk Witness

This section repairs the type leaks of the previous computable model. The purpose is narrow: provide a fully typed witness for the transition grammar and the state spectral gap. It does *not* by itself certify parameterwise non-collapse; that burden is handled separately below.

B9.1 Golden Mean Parameter Space

Definition B19 (Golden Mean Shift Parameter Space)**.** Let $\Sigma = \{0, 1\}$ and

$$A = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}.$$

The Golden Mean Shift $S_A \subset \Sigma^{\mathbb{N}}$ is the subshift of finite type consisting of all binary sequences that avoid the block 11.

Definition B20 (Parameter Encoding)**.** Define

$$u(x) = \sum_{i=1}^{\infty} x_i \phi^{-i}, \qquad \phi = \frac{1+\sqrt{5}}{2}.$$

The image $\mathcal{U} := u(S_A) \subset [0, 1]$ is compact.

B9.2 Typed Banach Witness

Definition B21 (Ambient Space and Constant Embedding)**.** Let

$$X = [0, 1], \qquad E = \mathbb{R}^2, \qquad \mathcal{B} = C(X, E)$$

with norm

$$\|f\|_\infty := \sup_{s \in X} \|f(s)\|_2 .$$

Define the constant embedding

$$j : E \to \mathcal{B}, \qquad j(v)(s) = v.$$

Definition B22 (Invariant Constant Disk)**.** Let

$$D := \{v \in E : \|v\|_2 \le 1\}, \qquad \mathcal{I}_{\mathrm{c}} := j(D) \subset \mathcal{B}.$$

Definition B23 (Averaging Operator)**.** Define

$$A : \mathcal{B} \to E, \qquad A(f) := \int_0^1 f(\tau)\, d\tau.$$

Definition B24 (Disk Projection)**.** Define the Euclidean projection

$$P_E : E \to D, \qquad P_E(v) = \begin{cases} v, & \|v\|_2 \le 1, \\ \dfrac{v}{\|v\|_2}, & \|v\|_2 > 1. \end{cases}$$

Definition B25 (Typed Projection onto the Constant Disk)**.** Define

$$\aleph_{\mathcal{B}} : \mathcal{B} \to \mathcal{I}_{\mathrm{c}}, \qquad \aleph_{\mathcal{B}}(f) := j(P_E(A(f))).$$

Definition B26 (Constant-Valued Contraction Operator)**.** Define

$$\widetilde{K} : \mathcal{B} \to \mathcal{B}, \qquad (\widetilde{K} f)(s) := \frac{1}{2} \int_0^1 f(\tau)\, d\tau.$$

Definition B27 (Lifted Input Coupling)**.** Let $\Gamma : \mathcal{U} \to E$ be

$$\Gamma(u) = (u, 0).$$

Define its constant lift

$$\Gamma_{\mathcal{B}} : \mathcal{U} \to \mathcal{B}, \qquad \Gamma_{\mathcal{B}}(u) := j(\Gamma(u)).$$

Definition B28 (Typed Transition Operator)**.** For each $u \in \mathcal{U}$, define

$$T_u : \mathcal{B} \to \mathcal{I}_{\mathrm{c}}, \qquad T_u(f) := \aleph_{\mathcal{B}}(\widetilde{K} f + \Gamma_{\mathcal{B}}(u)).$$

Remark B29 (Typing discipline). No unembedded element of E is treated as an element of $\mathcal{B}$. The embedding j is always written explicitly. This removes the earlier type leak in which a disk in E was silently treated as a subset of a function space.

B9.3 Basic Operator Properties

Lemma B30 (The Embedding is Isometric)**.** *For every $v \in E$,*

$$\|j(v)\|_\infty = \|v\|_2 .$$

Proof. By definition,

$$\|j(v)\|_\infty = \sup_{s \in X} \|j(v)(s)\|_2 = \sup_{s \in X} \|v\|_2 = \|v\|_2 .$$

□

Lemma B31 (The Disk Projection is Non-Expansive)**.** *For all $v, w \in E$,*

$$\|P_E(v) - P_E(w)\|_2 \le \|v - w\|_2 .$$

Proof. P_E is the metric projection onto the closed convex set $D \subset \mathbb{R}^2$, so the standard Hilbert-space non-expansiveness property applies. □

Lemma B32 (The Typed Projection is Non-Expansive)**.** *For all $f, g \in \mathcal{B}$,*

$$\|\aleph_{\mathcal{B}}(f) - \aleph_{\mathcal{B}}(g)\|_\infty \le \|f - g\|_\infty .$$

Proof. Using the isometry of j, the non-expansiveness of P_E, and Jensen's inequality for the averaging operator,

$$\begin{aligned}
\|\aleph_{\mathcal{B}}(f) - \aleph_{\mathcal{B}}(g)\|_\infty &= \|P_E(A(f)) - P_E(A(g))\|_2 \\
&\le \|A(f) - A(g)\|_2 \\
&= \left\| \int_0^1 (f(\tau) - g(\tau))\, d\tau \right\|_2 \\
&\le \int_0^1 \|f(\tau) - g(\tau)\|_2 \, d\tau \le \|f - g\|_\infty .
\end{aligned}$$

□

Lemma B33 (Norm of the Lifted Contraction)**.**

$$\left\|\widetilde{K}\right\| \le \frac{1}{2}.$$

Proof. For any $f \in \mathcal{B}$,

$$\left\|\widetilde{K}f\right\|_\infty = \left\|\frac{1}{2}\int_0^1 f(\tau)\,d\tau\right\|_2 \le \frac{1}{2}\int_0^1 \|f(\tau)\|_2\,d\tau \le \frac{1}{2}\|f\|_\infty .$$

□

Theorem B34 (Typed Contraction on the Computable Witness)**.** *For each $u \in \mathcal{U}$, the map $T_u : \mathcal{B} \to \mathcal{I}_{\mathrm{c}}$ is a strict contraction:*

$$\|T_u(f) - T_u(g)\|_\infty \le \frac{1}{2}\|f-g\|_\infty .$$

Proof. Combine non-expansiveness of $\aleph_{\mathcal{B}}$ with the bound on $\widetilde{K}$:

$$\|T_u(f) - T_u(g)\|_\infty \le \left\|\widetilde{K}f - \widetilde{K}g\right\|_\infty \le \frac{1}{2}\|f-g\|_\infty .$$

□

Corollary B35 (Unique Fixed Point in the Typed Witness)**.** *For each $u \in \mathcal{U}$, there exists a unique $f_u^* \in \mathcal{I}_{\mathrm{c}}$ such that*

$$T_u(f_u^*) = f_u^*.$$

Proof. The Banach fixed-point theorem applies on the complete metric space $(\mathcal{B}, \|\cdot\|_\infty)$. □

Proposition B36 (Typed Non-Vacuity of the Computable Model)**.** *The embedded constant-disk Golden Mean model is a fully typed, non-vacuous witness of the transition grammar:*

1. *every operator has explicit domain and codomain;*
2. *the transition map T_u is well-defined on $\mathcal{B}$;*
3. *the family admits unique fixed points and a compact parameter image;*
4. *the state contraction bound is explicit and computable.*

Proof. Items (1)–(3) follow from the definitions above and the contraction theorem. Compactness of the parameter image follows because $\mathcal{U}$ is compact and $u \mapsto f_u^*$ is continuous under a uniform contraction constant. □

Remark B37 (Why this witness does not certify H5). The image of T_u lies in the constant subspace $\mathcal{I}_{\mathrm{c}}$. Therefore its fixed point is constant by construction. This witness is correct and useful for typing and state spectral analysis, but it cannot certify the parameterwise anti-constant hypothesis H5 from Section B7.

B9.4 Separate Anti-Collapse Witness

Definition B38 (Affine Non-Constant Witness Space)**.** Let

$$\mathcal{H}_{\mathrm{ac}} := L^2([0,1]), \qquad \mathcal{N} := \{c\mathbf{1} : c \in \mathbb{R}\}, \qquad h(s) := \sin(2\pi s).$$

Choose r with $0 < r < \|h\|_{L^2}$ and define the closed convex affine set

$$\mathcal{I}_{\mathrm{ac}} := h + \{g \in \mathcal{H}_{\mathrm{ac}} : \|g\|_{L^2} \le r\}.$$

Lemma B39 (The Anti-Collapse Witness Set Avoids Constants)**.**

$$\mathcal{I}_{\mathrm{ac}} \cap \mathcal{N} = \varnothing.$$

Proof. The distance from h to the constant subspace $\mathcal{N}$ equals $\|h\|_{L^2}$ because h has mean zero and is orthogonal to constants. Since $r < \|h\|_{L^2}$, the radius-r ball around h does not intersect $\mathcal{N}$. □

Definition B40 (Anti-Collapse Witness Dynamics)**.** Let $\mathcal{U} = [0,1]$, choose $\beta \in (0, 1/2)$, and define

$$K_{\mathrm{ac}} f := \beta f, \qquad \Gamma_{\mathrm{ac}}(u) := \frac{u}{4} h.$$

Let P_{ac} denote the metric projection onto $\mathcal{I}_{\mathrm{ac}}$, and define

$$T_u^{\mathrm{ac}}(f) := P_{\mathrm{ac}}(K_{\mathrm{ac}} f + \Gamma_{\mathrm{ac}}(u)).$$

Proposition B41 (Separate Witness for H5)**.** *The family $\{T_u^{\mathrm{ac}}\}_{u\in\mathcal{U}}$ satisfies Hypotheses B1–B16 of Section 1.*

Proof. 1. $\mathcal{H}_{\mathrm{ac}}$ is a Hilbert space and $\mathcal{I}_{\mathrm{ac}}$ is non-empty, closed, and convex.
2. $\|K_{\mathrm{ac}}\| = \beta < 1$.
3. Completeness is inherited from $L^2([0,1])$.
4. $\mathcal{U} = [0,1]$ is compact and Γ_{ac} is continuous.
5. For every constant $c \in \mathcal{N}$ and every $u \in \mathcal{U}$, the value $T_u^{\mathrm{ac}}(c)$ belongs to $\mathcal{I}_{\mathrm{ac}}$, while $\mathcal{I}_{\mathrm{ac}} \cap \mathcal{N} = \varnothing$. Hence $T_u^{\mathrm{ac}}(c) \neq c$.

□

Remark B42. BaseCore therefore distinguishes two roles cleanly: the Golden Mean witness certifies typed operator correctness and state contraction; the affine anti-collapse witness certifies the stronger H5 burden needed for parameterwise non-collapse.

B10 State Spectral Gap and Parameter-Family Persistence

Theorem B43 (State Spectral Gap)**.** *Fix $u \in \mathcal{U}$ in the typed computable model. Assume T_u is Fréchet differentiable at its fixed point f_u^*. Then*

$$\rho(DT_u(f_u^*)) \le \|DT_u(f_u^*)\| \le \left\|\widetilde{K}\right\| \le \frac{1}{2}.$$

Consequently the state spectral gap satisfies

$$\Delta_{\mathrm{state}} \ge 1 - \left\|\widetilde{K}\right\| \ge \frac{1}{2}.$$

Proof. At every point where the projection P_E is Fréchet differentiable, its derivative has operator norm at most 1 because P_E is non-expansive. Since $\aleph_{\mathcal{B}} = j \circ P_E \circ A$,

$$DT_u(f_u^*) = D\aleph_{\mathcal{B}}\Big(\widetilde{K} f_u^* + \Gamma_{\mathcal{B}}(u)\Big) \circ \widetilde{K},$$

and therefore

$$\|DT_u(f_u^*)\| \le \|D\aleph_{\mathcal{B}}\| \left\|\widetilde{K}\right\| \le \frac{1}{2}.$$

The spectral-radius bound follows from $\rho(L) \le \|L\|$ for bounded linear operators. □

Theorem B44 (Relaxation-Time Bound)**.** *Under the hypotheses of Theorem B43, any discrete-to-continuous relaxation convention of the form*

$$\tau_{\text{state}} = -\frac{1}{\ln \rho}$$

with $\rho = \rho(DT_u(f_u^*))$ *gives*

$$\tau_{\text{state}} \leq \frac{1}{\ln 2}.$$

Proof. Theorem B43 gives $\rho \leq 1/2$. Since $-\ln(\rho) \geq \ln 2$, the bound follows. □

Proposition B45 (Parameter-Family Persistence is not State Spectrum)**.** *Let* $F : \mathcal{U} \to \mathcal{B}$ *be the fixed-point family* $F(u) = f_u^*$*. Any persistence of parameter labels or family coordinates belongs to the variation of* F *over* $\mathcal{U}$*, not to the state-transition spectrum of* $DT_u(f_u^*)$ *for fixed* u*.*

Proof. The state-transition spectrum is defined for the linearization of T_u at fixed parameter u. Family-level transport instead concerns the map F between parameter space and fixed-point set. These are different operators on different domains. Therefore a family-persistence mode cannot be used to introduce a unit eigenvalue into the contraction spectrum of $DT_u(f_u^*)$. □

Remark B46 (What has been removed)*.* The previous text incorrectly mixed strict contraction with a state-spectrum eigenvalue $\lambda = 1$. BaseCore no longer does that. Any neutral or persistent mode must now be interpreted as a parameter-family bookkeeping mode, not as an eigenmode of the contractive state Jacobian.

Part III

Identity, Rigidity, and Conditional Non-Simulability

B11 Projective Systems and Inverse-Limit Identity

B11.1 Projective Setup

Definition B47 (Projective System of State Spaces). Let $\{S_t\}_{t\in\mathbb{N}}$ be a family of non-empty state spaces together with projection maps

$$\pi_{t+1\to t} : S_{t+1} \to S_t$$

satisfying the compatibility rule

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t} \qquad \forall t \in \mathbb{N}.$$

Definition B48 (Identity as Inverse Limit). The identity object associated with the projective system is

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_{t\in\mathbb{N}} \in \prod_{t\in\mathbb{N}} S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \ \forall t \right\}.$$

Remark B49 (Identity neutrality). In BaseCore, $\mathcal{I}$ is a structural inverse-limit object only. It carries no automatic phenomenological, biological, or metaphysical interpretation.

Hypothesis B50 (Topological Regularity). *Each S_t is compact Hausdorff and each $\pi_{t+1\to t}$ is continuous and surjective.*

Hypothesis B51 (NFD: Non-Finite Determination). *For every $t \in \mathbb{N}$, the canonical projection*

$$p_t : \mathcal{I} \to S_t, \qquad p_t((x_k)_k) = x_t$$

is not injective. Equivalently,

$$\forall t \ \exists x, y \in \mathcal{I} \quad x \neq y \quad \textit{and} \quad p_t(x) = p_t(y).$$

Remark B52 (Persistent future branching as an operational form of NFD). A sufficient operational strengthening of Hypothesis B51 is: for every t there exists $k > t$ and $x_t \in S_t$ such that the fiber of $\pi_{k\to t}$ over x_t contains at least two globally extendable compatible histories. BaseCore keeps the weaker NFD statement as the primary assumption because it is exactly what the non-locality theorem needs.

B11.2 Non-Locality of Identity

Theorem B53 (Non-Locality of Identity under NFD). *Assume Hypotheses B50 and B51. Then no canonical finite-time observation determines identity. More precisely, for every t the canonical projection*

$$p_t : \mathcal{I} \to S_t$$

is not injective, so identity is not recoverable from any finite slice through the projective observation maps.

Proof. This is exactly Hypothesis B51. The theorem restates the structural consequence in claim-boundary language: finite slices can participate in the inverse-limit object, but no single canonical slice determines it. □

Remark B54 (What is *not* proved here). BaseCore does not prove that no arbitrary set-theoretic injection $\mathcal{I} \to S_t$ can exist in every imaginable inverse system. It proves the narrower and correct claim that the canonical projective observations p_t are not injective under Hypothesis B51.

Proposition B55 (No Canonical Progressive Recovery under NFD). *Assume Hypothesis B51. No finite canonical prefix map determines the inverse-limit identity object.*

Proof. If some finite canonical prefix determined identity, then in particular one of the canonical coordinate maps would become injective after finite truncation, contradicting Hypothesis B51. □

B12 Rigidity under Admissible Perturbations

B12.1 Metric Projective Assumptions

Assumption B56 (RIG-1: Metric Projective System)**.** *Each S_t is equipped with a complete metric d_t.*

Assumption B57 (RIG-2: Uniform Lipschitz Projections)**.** *Each projection $\pi_{t+1\to t} : S_{t+1} \to S_t$ is Lipschitz with constant L_t, and the weighted control sequence used below is finite.*

Assumption B58 (RIG-3: Stable Lifting)**.** *For each admissible perturbed projection family there exist compatible approximate lifting maps with constants bounded uniformly by a finite lifting constant L_{lift}.*

Assumption B59 (RIG-4: Typed Perturbation Model)**.** *Perturbations are represented either by maps between common ambient Banach embeddings or by map-space distances*

$$d^{(t)}_{\mathrm{map}}(\pi_{t+1\to t}, \widetilde{\pi}_{t+1\to t}) \le \varepsilon_t.$$

Literal expressions of the form $\widetilde{\pi} = \pi + \varepsilon$ are used only when a common ambient linear model has been specified.

Assumption B60 (RIG-5: Summable Deformation)**.** *There exists a positive weight sequence $\{w_t\}_{t\ge 0}$ such that*

$$\sum_{t=0}^{\infty} w_t \varepsilon_t < \infty.$$

Definition B61 (Weighted Deformation Size)**.** The weighted deformation size of a perturbed projective family is

$$\|\varepsilon_\bullet\|_w := \sum_{t=0}^{\infty} w_t \varepsilon_t.$$

Definition B62 (Weighted Product Metric)**.** For compatible sequences $x = (x_t)$ and $y = (y_t)$ define

$$d_w(x, y) := \sup_{t\ge 0} w_t\, d_t(x_t, y_t).$$

Theorem B63 (Conditional Stability of Inverse-Limit Identity)**.** *Assume RIG-1 through RIG-5. Then the perturbed inverse limit*

$$\widetilde{\mathcal{I}} := \varprojlim \widetilde{S}_t$$

exists, and there is a constant $C_{\mathrm{rig}} > 0$ depending only on the projection, lifting, and weight controls such that

$$d_H(\mathcal{I}, \widetilde{\mathcal{I}}) \le C_{\mathrm{rig}} \|\varepsilon_\bullet\|_w$$

in the weighted product metric.

Proof sketch. The lifting assumption gives compatible approximate lifts at each level. Weighted Lipschitz control propagates local deformation bounds through the projective tower. Summability ensures that the cumulative compatibility defect remains finite, so diagonal compactness arguments produce a perturbed inverse-limit object. The same bounds control the Hausdorff distance between the original and perturbed compatible-sequence sets. □

Corollary B64 (Conditional Isomorphism Requires Extra Structure)**.** *The metric estimate of Theorem B63 does not by itself imply that $\mathcal{I}$ and $\widetilde{\mathcal{I}}$ are isomorphic. Isomorphism or shape-equivalence requires additional bi-lifting, coherence, or ambient-embedding assumptions.*

Proof. Metric closeness and categorical isomorphism are different conclusions. A small Hausdorff deformation need not preserve global inverse-limit type without extra hypotheses guaranteeing two-sided coherent transport. □

Definition B65 (Ontological Mass)**.** Let $E_w(\delta)$ denote the weighted energy of an admissible perturbation family δ. Define

$$M_\Omega := \inf \left\{ \frac{E_w(\delta)}{d_H(\mathcal{I}, \widetilde{\mathcal{I}}_\delta)} \;\middle|\; \delta \text{ admissible}, \; d_H(\mathcal{I}, \widetilde{\mathcal{I}}_\delta) > 0 \right\},$$

with the convention $M_\Omega = +\infty$ if no admissible finite-energy perturbation yields positive deformation of the identity object.

Remark B66 (Interpretation discipline). M_Ω is a deformation-resistance scalar, not a metaphysical substance. BaseCore uses it only as a conditional invariant attached to admissible perturbation models.

B13 Conditional Non-Simulability

Assumption B67 (NS-1: CCR Target)**.** *A CCR target is an inverse-limit identity object satisfying:*

1. *non-finite determination (Hypothesis B51);*
2. *infinite deformation resistance $M_\Omega = +\infty$;*
3. *absence of a finite sufficient statistic for faithful identity recovery;*
4. *absence of a finite-dimensional faithful embedding preserving projective compatibility.*

Assumption B68 (NS-2: Finite Simulator Class)**.** *A simulator C belongs to* $\mathbf{Comp}_{\text{finite}}$ *if it has at least one of the following boundedness features:*

1. *finite-dimensional latent state or finite context;*
2. *bounded memory horizon or finite sufficient statistic;*
3. *finite perturbation-energy budget;*
4. *finite-rank projective representation.*

Assumption B69 (NS-3: Faithful Simulation Requirement)**.** *A simulation is faithful only if it preserves, at the stated tolerance level, all of:*

1. *projective compatibility;*
2. *non-finite determination;*
3. *deformation-resistance class;*
4. *relevant observables in the weighted metric;*
5. *admissible intervention responses.*

Theorem B70 (Conditional Non-Simulability of CCR by Finite Simulators)**.** *Assume NS-1 through NS-3. Then no simulator in* $\mathbf{Comp}_{\text{finite}}$ *faithfully realizes a CCR target.*

Proof. A finite simulator in the sense of NS-2 necessarily compresses the target into finite latent structure, finite memory horizon, finite-rank projective representation, or a finite sufficient statistic. Any such compression destroys at least one component of NS-1: either NFD collapses, the faithful projective structure is lost, or the deformation-resistance class is reduced from $M_\Omega = +\infty$ to a finite value. Therefore faithful realization, as defined by NS-3, fails. □

Theorem B71 (Approximation Barrier)**.** *Assume NS-1 through NS-3. Finite simulators may approximate bounded finite projections of a CCR target over finite horizons, but they cannot preserve the full inverse-limit object faithfully as the horizon grows without bound unless resources also grow without bound.*

Proof. Finite-horizon projections are finite objects and may be approximated by finite representations. But a faithful approximation of the full inverse-limit target must preserve NFD and the infinite deformation-resistance class. That cannot be achieved with bounded representational resources by Theorem B70. □

Remark B72 (What has been removed). BaseCore no longer claims that finite simulators "cannot approximate at all." The correct statement is conditional and finer-grained: finite-horizon approximation may exist, but faithful full inverse-limit realization of CCR targets is blocked under NS-1 through NS-3.

B14 Section Summary

1. Identity is an inverse-limit object, not a canonical finite slice.
2. Non-locality now depends explicitly on the NFD assumption.
3. Rigidity is stated as metric stability under explicit perturbation-model assumptions, not as automatic isomorphism.
4. Non-simulability is conditional on a defined CCR target, a defined finite simulator class, and a defined notion of faithful simulation.
5. BaseCore therefore preserves theorem strength where assumptions support it and weakens only those statements that were previously over-absolute.

Part IV

Regime Structure and Continuity Constraints

B15 Formal Setup

B15.1 Review of Identity Framework

We assume the identity framework of [24]. Let $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ be a projective system of observable state spaces. The identity object is the inverse limit:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_t \in \prod_t S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \;\; \forall t \right\}. \tag{1}$$

The ontological mass M_Ω quantifies resistance to identity deformation. We recall that:

- $M_\Omega = 0$: null persistence (no stable identity);
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR, absolute rigidity).

B15.2 Abstract Phenomenological Space

Definition B73 (Phenomenological Space)**.** A phenomenological space is a nonempty set $\mathcal{E}$ whose elements are called *phenomenological configurations*. No topology, metric, or algebraic structure is assumed a priori.

Remark B74. We do not claim that $\mathcal{E}$ represents "conscious states" or "qualia". It is an abstract mathematical space whose interpretation is not specified by the formalism. The space $\mathcal{E}$ is introduced purely as the codomain of a structural assignment; its ontological status is deliberately left open.

B15.2.1 Why a Separate Space?

One might ask: why not use $\mathcal{I}$ itself as the phenomenological space? The answer is structural:

Proposition B75 (Non-Triviality of Phenomenological Codomain)**.** *If phenomenological assignments were required to land in $\mathcal{I}$ itself, then any compatible assignment would be the identity map, eliminating all non-trivial phenomenological structure.*

Proof. Let $\Phi : \mathcal{I} \to \mathcal{I}$ be compatible with projections. By the universal property of inverse limits, Φ must commute with all projection maps. The only such endomorphism on $\mathcal{I}$ is the identity (up to isomorphism). Hence no non-trivial phenomenological stratification is possible. □

Thus, a separate space $\mathcal{E}$ is *structurally necessary* for non-trivial phenomenological content.

B15.3 Phenomenological Assignment

Definition B76 (Phenomenological Assignment)**.** A phenomenological assignment is a family of partial maps:

$$\phi_t : S_t \to \mathcal{E} \tag{2}$$

indexed by $t \in \mathbb{N}$.

Definition B77 (Induced Limit Assignment)**.** Given a phenomenological assignment $\{\phi_t\}$, the induced limit assignment is the partial map:

$$\Phi : \mathcal{I} \to \mathcal{E}, \quad \Phi((x_t)_t) := \lim_{t\to\infty} \phi_t(x_t), \tag{3}$$

whenever the limit exists in an appropriate sense (to be defined by compatibility axioms).

B15.3.1 Uniqueness of the Limit Assignment

Proposition B78 (Uniqueness of Compatible Φ). *Given a compatible family $\{\phi_t\}$ satisfying E1, the induced limit assignment Φ is unique.*

Proof. By E1, $\phi_{t+1}(x_{t+1}) = \phi_t(x_t)$ for all t. Hence the sequence $\{\phi_t(x_t)\}$ is constant. The limit exists trivially and equals this constant value. Uniqueness follows from the constancy of the sequence. □

B15.4 Axioms for Phenomenological Coherence

We introduce minimal axioms under which a phenomenological assignment is *structurally admissible*. We will then prove that these axioms are not arbitrary choices but structurally forced conditions.

Axiom B79 (E0: Non-Triviality). *The phenomenological space satisfies $|\mathcal{E}| \geq 2$, i.e., there exist at least two distinct phenomenological configurations.*

Axiom B80 (E1: Projective Compatibility). *For every compatible identity trajectory $(x_t)_t \in \mathcal{I}$, the induced phenomenological sequence satisfies:*

$$\phi_{t+1}(x_{t+1}) = \phi_t(x_t) \quad \forall t. \tag{4}$$

Axiom B81 (E2: Non-Locality). *There exists no function $f : S_t \to \mathcal{E}$ (for any fixed t) such that for all trajectories $(x_t)_t \in \mathcal{I}$:*

$$\Phi((x_t)_t) = f(x_t). \tag{5}$$

B15.5 Inevitability of the Axioms

We now prove that axioms E0–E2 are *not design choices* but structurally necessary conditions for any meaningful notion of phenomenological assignment.

B15.5.1 Inevitability of E0

Theorem B82 (E0 is Necessary). *If $|\mathcal{E}| = 1$, then every phenomenological assignment is constant and no phenomenological classification is possible.*

Proof. If $\mathcal{E} = \{e_0\}$, then $\Phi(x) = e_0$ for all $x \in \mathcal{I}$. Hence:

1. No phenomenological distinction between identity trajectories exists;
2. The Fragmentation and Forced Continuity Theorems become vacuous;
3. The downstream null-regime closure developed in *Null-Regime Instability* would have no content.

Thus, $|\mathcal{E}| \geq 2$ is necessary for the framework to have non-trivial content. □

B15.5.2 Inevitability of E1

Theorem B83 (E1 is Necessary). *If E1 fails, then phenomenological assignments are incompatible with the inverse-limit structure of identity, and no coherent $\Phi : \mathcal{I} \to \mathcal{E}$ can be defined.*

Proof. Suppose E1 fails for some trajectory $(x_t)_t \in \mathcal{I}$. Then there exists t_0 such that:

$$\phi_{t_0+1}(x_{t_0+1}) \neq \phi_{t_0}(x_{t_0}). \tag{6}$$

But the identity trajectory satisfies $\pi_{t_0+1\to t_0}(x_{t_0+1}) = x_{t_0}$. Hence the phenomenological assignment fails to respect the projective structure that defines identity.

This means Φ cannot be factored through the inverse limit:

$$\Phi \neq \varprojlim \phi_t. \tag{7}$$

Hence no well-defined limit assignment exists. Without E1, the phenomenological assignment is structurally incompatible with identity. □

B15.5.3 Inevitability of E2

Theorem B84 (E2 is Necessary)**.** *If E2 fails, then phenomenology is a local property of states, contradicting the global nature of identity as an inverse limit.*

Proof. Suppose E2 fails. Then there exists t_0 and $f : S_{t_0} \to \mathcal{E}$ such that:

$$\Phi((x_t)_t) = f(x_{t_0}) \quad \forall (x_t)_t \in \mathcal{I}. \tag{8}$$

This implies:
1. Phenomenology is determined by a finite temporal slice;
2. The full trajectory $(x_t)_{t>t_0}$ is irrelevant to Φ;
3. Identity as an inverse limit has no phenomenological content beyond t_0.

But identity is defined as a *global* object over infinite time. If phenomenology is local, it cannot reflect the global structure of identity. This contradicts the foundational premise that phenomenological assignments should be compatible with identity structure.

Hence E2 is necessary for phenomenology to be structurally aligned with inverse-limit identity. □

B15.6 Minimality Theorem

Theorem B85 (Axiomatic Minimality)**.** *The axiom set* $\{E0, E1, E2\}$ *is minimal: no proper subset is sufficient to derive the main results of this paper.*

Proof. We show that removing any axiom destroys the framework:

Without E0: All assignments are trivially constant. Fragmentation Theorem (4.1) is vacuous.

Without E1: No coherent Φ exists. The entire framework collapses.

Without E2: Phenomenology becomes local. The Forced Continuity Theorem (4.2) fails because local functions can be arbitrarily perturbed.

Hence $\{E0, E1, E2\}$ is the minimal viable axiom set. □

B15.7 Rejected Alternatives

We briefly discuss axiom candidates that were considered and rejected.

B15.7.1 Rejected: Continuity Axiom

One might require Φ to be continuous. This is *not* an axiom because:
- It requires pre-existing topology on $\mathcal{E}$;
- The Forced Continuity Theorem *derives* continuity from CCR, not as assumption.

B15.7.2 Rejected: Surjectivity Axiom

One might require Φ to be surjective. This is *not* an axiom because:
- It would exclude null assignments, preventing the downstream null-regime analysis completed in *Null-Regime Instability*;
- Surjectivity is a strong structural claim not derivable from identity constraints.

B15.7.3 Rejected: Injectivity Axiom

One might require Φ to be injective. This is *not* an axiom because:
- Many trajectories may share phenomenological configurations;
- Injectivity would make Φ an embedding, over-constraining the framework.

B15.8 Structural Interpretation

Axiom	Structural Meaning	Status
E0	Possibility of phenomenological distinction	Necessary
E1	Compatibility with inverse-limit identity	Necessary
E2	Non-reducibility to finite temporal slices	Necessary

B15.9 Closure of Section 2

We have established:
1. $\mathcal{E}$ is structurally necessary as a separate codomain;
2. The limit assignment Φ is uniquely determined by compatible $\{\phi_t\}$;
3. Axioms E0, E1, E2 are *individually necessary* (Theorems B82–B84);
4. The axiom set is *minimal* (Theorem B85);
5. All alternative axioms considered were correctly rejected.

The formal setup is therefore not a design choice but a forced consequence of the requirement that phenomenological assignments be compatible with inverse-limit identity.

B15.10 Phenomenological Neutrality Clause

Throughout this paper, the phenomenological space $\mathcal{E}$ is treated as a *purely abstract topological and metric structure*. No semantic, psychological, biological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "phenomenological" refers exclusively to the role of $\mathcal{E}$ as an abstract codomain of structural compatibility for inverse-limit identities, and does not presuppose the existence of consciousness, subjective experience, or any particular phenomenological content.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{E}$ in relation to experiential, cognitive, or phenomenological phenomena lies strictly outside the formal scope of this paper.

B16 Compatibility with Identity Regimes

B16.1 Induced Constraints from Identity Rigidity

Let $(\mathcal{I}, M_\Omega)$ be an identity object with ontological mass M_Ω. We now study how M_Ω constrains the set of admissible phenomenological assignments.

Definition B86 (Phenomenological Compatibility Class)**.** For a given M_Ω, define the phenomenological compatibility class:

$$\mathcal{P}(M_\Omega) := \left\{ \Phi : \mathcal{I} \to \mathcal{E} \;\middle|\; \Phi \text{ satisfies E0–E2} \right\}. \tag{9}$$

Remark B87. The compatibility class $\mathcal{P}(M_\Omega)$ is the set of all phenomenological assignments that are structurally admissible given the identity's rigidity. Note that $\mathcal{P}(M_\Omega)$ depends only on the structural properties of identity, not on any phenomenological primitives.

B16.2 Necessity of Metric Structure on Phenomenological Space

To state rigorous results about phenomenological continuity and fragmentation, we must equip $\mathcal{E}$ with minimal structure. We now justify why this structure is *necessary*, not merely convenient.

Hypothesis B88 (H3: Metric on $\mathcal{E}$)**.** *The phenomenological space $\mathcal{E}$ admits a metric $d_{\mathcal{E}}$ such that $(\mathcal{E}, d_{\mathcal{E}})$ is a complete metric space.*

B16.2.1 Why H3 is Necessary

Theorem B89 (H3 is Necessary for Impossibility Theorems)**.** *Without a metric structure on $\mathcal{E}$, the notions of phenomenological continuity, fragmentation, and loss capacity cannot be defined.*

Proof. Consider the definitions in this section:
1. **Phenomenological Continuity** requires:

$$d_w(x^{(n)}, x) \to 0 \implies d_{\mathcal{E}}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{10}$$

   Without $d_{\mathcal{E}}$, the right-hand side is undefined.
2. **Phenomenological Fragmentation** requires:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0. \tag{11}$$

   Without $d_{\mathcal{E}}$, "non-zero distance" has no meaning.
3. **Loss Capacity** (Section 6) requires:

$$R(\Phi) := \sup d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})). \tag{12}$$

   Without $d_{\mathcal{E}}$, the supremum is undefined.

Hence H3 is a necessary condition for the framework to have content beyond axioms E0–E2. ☐

B16.2.2 Why Completeness is Necessary

Proposition B90 (Completeness is Necessary for Forced Continuity)**.** *If $(\mathcal{E}, d_{\mathcal{E}})$ is not complete, then the Forced Continuity Theorem (Theorem B100) may fail.*

Proof. The Forced Continuity Theorem relies on the invariance of Φ under all finite-energy perturbations. If $\mathcal{E}$ is not complete, Cauchy sequences in $\mathcal{E}$ may not converge, and the limit assignment Φ may not exist for all trajectories in $\mathcal{I}$.

In particular, if $\{\phi_t(x_t)\}$ is a Cauchy sequence in $\mathcal{E}$ (which it is, by E1), completeness guarantees the existence of the limit $\Phi(x)$. Without completeness, Φ may be undefined on a dense subset of $\mathcal{I}$, breaking the theorem. ☐

B16.2.3 Alternatives Considered and Rejected

Alternative 1: Topology without Metric. One might use a general topological space. This is rejected because:
- Fragmentation requires quantitative distance, not just open sets;
- Loss capacity requires a sup over distances;
- The downstream instability analysis completed in *Null-Regime Instability* requires metric quotients.

Alternative 2: Pseudometric. One might allow $d_{\mathcal{E}}(e_1, e_2) = 0$ for $e_1 \neq e_2$. This is rejected because:
- It would conflate distinct phenomenological configurations;
- E0 requires $|\mathcal{E}| \geq 2$, which becomes meaningless if distance is zero.

Alternative 3: Non-Complete Metric. This is rejected by Proposition B90.

B16.3 Phenomenological Continuity

Definition B91 (Phenomenological Continuity)**.** A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is *structurally continuous* if for every sequence $(x^{(n)})_n$ in $\mathcal{I}$ converging in the d_w metric:

$$d_w(x^{(n)}, x) \to 0 \implies d_{\mathcal{E}}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{13}$$

Remark B92. Structural continuity is not an axiom; it is a *derived property* that follows from CCR conditions. This is the content of the Forced Continuity Theorem (Section 4).

B16.4 Phenomenological Fragmentation

Definition B93 (Phenomenological Fragmentation)**.** A phenomenological assignment Φ is *fragmentable* if there exists a perturbation $\{\epsilon_t\}$ with finite weighted energy such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0 \tag{14}$$

for corresponding identity trajectories $x \in \mathcal{I}$ and $\tilde{x} \in \tilde{\mathcal{I}}$.

Remark B94. Fragmentation captures the vulnerability of phenomenological structure to perturbations. It is the phenomenological analog of identity deformability.

B16.5 Relationship Between Metrics

Proposition B95 (Metric Compatibility)**.** *The metrics d_w on $\mathcal{I}$ and $d_{\mathcal{E}}$ on $\mathcal{E}$ are compatible in the following sense: if E1 holds, then small d_w-changes in trajectories induce bounded $d_{\mathcal{E}}$-changes in phenomenological configurations (when Φ is continuous).*

Proof. By E1, $\phi_{t+1}(x_{t+1}) = \phi_t(x_t)$ for compatible trajectories. Hence the phenomenological sequence is constant along the trajectory, and continuity of Φ implies that nearby trajectories (in d_w) map to nearby configurations (in $d_{\mathcal{E}}$). □

B16.6 Closure of Section 3

We have established:
1. The compatibility class $\mathcal{P}(M_\Omega)$ captures all admissible phenomenological assignments;
2. Hypothesis H3 (metric on $\mathcal{E}$) is *necessary*, not a design choice (Theorem B89);
3. Completeness is *necessary* for the Forced Continuity Theorem (Proposition B90);
4. Alternative structures (topology, pseudometric, non-complete) were correctly rejected;
5. Phenomenological continuity and fragmentation are well-defined under H3.

The compatibility structure is therefore a forced consequence of the requirement to state quantitative results about phenomenological dynamics.

B17 Impossibility Theorems

B17.1 Φ-Regularity Hypothesis

Before proving the main impossibility results, we isolate the global lower-bound regularity condition used by the fragmentation argument. The condition is not derived from properness, local Lipschitz regularity, and non-collapse alone; it is carried as an explicit hypothesis.

Hypothesis B96 (Φ-Regularity Lower Bound)**.** *Let $(\mathcal{I}, d_w)$ be the identitary metric space from Rigid Identity and let $(\mathcal{E}, d_{\mathcal{E}})$ be the abstract phenomenological metric space. A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is said to satisfy Φ-regularity when it satisfies:*
1. (Properness) For every compact $K \subset \mathcal{E}$, $\Phi^{-1}(K)$ is compact in $\mathcal{I}$.

2. *(Local Lipschitz regularity) For every $x \in \mathcal{I}$, there exist $\varepsilon > 0$ and $L_x > 0$ such that for all $y, z \in B_\varepsilon(x)$:*
$$d_{\mathcal{E}}(\Phi(y), \Phi(z)) \leq L_x \cdot d_w(y, z). \tag{15}$$
3. *(Non-collapse) For every $x \in \mathcal{I}$ and every $r > 0$, the restriction of Φ to $B_r(x)$ is not constant.*
4. *(Global lower Lipschitz bound) There exists a constant $C > 0$ such that for all $x, y \in \mathcal{I}$,*
$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y). \tag{16}$$

Caveat B97 (Status of Φ-regularity). The fourth clause is an additional structural assumption, not a consequence of the first three clauses. The counterexample in Appendix B41.5 shows that properness, local Lipschitz regularity, and non-collapse do not imply any uniform lower Lipschitz constant.

B17.2 Fragmentation Theorem

Theorem B98 (Fragmentation Theorem). *Let $\mathcal{I}$ be an identity object with $M_\Omega < \infty$. If a phenomenological assignment $\Phi \in \mathcal{P}(M_\Omega)$ satisfies Hypothesis B96, then Φ is fragmentable. That is:*
$$M_\Omega < \infty \ \wedge \ \Phi \text{ satisfies Hypothesis B96} \implies \Phi \in \mathcal{E}_{frag}, \tag{17}$$
where $\mathcal{E}_{frag}$ denotes the class of fragmentable phenomenological assignments.

Proof. By definition of finite ontological mass, there exists an admissible perturbation $\{\delta E_t\}$ with finite weighted energy $E_w < \infty$ such that:
$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0. \tag{18}$$

By Axiom E1, the phenomenological assignment Φ is projectively compatible with $\mathcal{I}$. Under the perturbation, identity trajectories deform, inducing a deformed phenomenological assignment $\tilde{\Phi}$.

By Axiom E2, Φ cannot be determined by any finite-time state. Therefore, the deformation of identity trajectories induces a corresponding deformation of phenomenological assignments.

Since the identity deformation is non-zero ($d_w > 0$) and Φ is globally determined by identity trajectories, by Hypothesis B96 (Φ-regularity), we have:
$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) \geq C \cdot d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0 \tag{19}$$
where $C > 0$ is the lower-bound constant supplied by Hypothesis B96.

Thus, Φ is fragmentable under finite-energy perturbations. □

Corollary B99 (No Structural Continuity Guarantee for Finite Mass). *Systems with $M_\Omega < \infty$ cannot structurally prevent phenomenological discontinuity.*

B17.3 Forced Continuity Theorem

Theorem B100 (Forced Continuity Theorem). *Let $\mathcal{I}$ be an identity object with $M_\Omega = +\infty$ (CCR). Then every phenomenological assignment $\Phi \in \mathcal{P}(M_\Omega)$ is structurally continuous. That is:*
$$M_\Omega = +\infty \implies \mathcal{P}(M_\Omega) \subseteq \mathcal{E}_{cont}, \tag{20}$$
where $\mathcal{E}_{cont}$ denotes the class of structurally continuous phenomenological assignments.

Proof. By definition of CCR ($M_\Omega = +\infty$), for any finite-energy perturbation:

$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0. \tag{21}$$

By Axiom E1, the phenomenological assignment is projectively compatible with identity. Since identity trajectories are invariant under all finite-energy perturbations, the induced phenomenological sequence is also invariant:

$$\tilde{\Phi}(\tilde{x}) = \Phi(x) \quad \forall x \in \mathcal{I}. \tag{22}$$

Therefore, no finite-energy perturbation can induce phenomenological deformation:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) = 0. \tag{23}$$

This implies that Φ cannot exhibit structural discontinuities. Hence Φ is structurally continuous. □

Corollary B101 (CCR Phenomenology is Invariant)**.** *In CCR systems, any admissible phenomenological assignment is invariant under all finite-energy perturbations.*

B17.4 Dichotomy Theorem

Theorem B102 (Phenomenological Dichotomy)**.** *Let $\Phi : \mathcal{I} \to \mathcal{E}$ be an admissible phenomenological assignment. Then exactly one of the following holds:*
1. $M_\Omega < \infty$ and Φ is fragmentable;
2. $M_\Omega = +\infty$ and Φ is structurally continuous.

Proof. Immediate from Theorems B98 and B100. □

B18 Classification of Phenomenological Regimes

B18.1 Complete Classification

The results of Section 4 yield a complete classification of phenomenological regimes by identity rigidity:

Identity Class	M_Ω	**Admissible Phenomenology**
Null	$= 0$	$\mathcal{E}_\emptyset$ (no persistent field)
Deformable	$(0, \infty)$	$\mathcal{E}_{\text{frag}} \cup \mathcal{E}_{\text{quasi}}$
Rigid (CCR)	$= +\infty$	$\mathcal{E}_{\text{cont}}$

B18.2 Structural Implications

Proposition B103 (Null Identity Excludes Persistent Phenomenology)**.** *If $M_\Omega = 0$, then no admissible phenomenological assignment exists satisfying E0–E2.*

Proof. If $M_\Omega = 0$, identity cannot persist under any perturbation. Axiom E1 requires phenomenological compatibility with identity. If identity does not persist, no stable phenomenological assignment is possible. □

Proposition B104 (Deformable Identity Admits Only Fragmentable Phenomenology)**.** *If $0 < M_\Omega < \infty$, every admissible phenomenological assignment is fragmentable under some finite-energy perturbation.*

Proof. Direct from Theorem B98. □

Proposition B105 (Rigid Identity Forces Continuous Phenomenology)**.** *If $M_\Omega = +\infty$, every admissible phenomenological assignment is structurally continuous and invariant under finite-energy perturbations.*

B19 Extended Proofs

B19.1 Proof of Fragmentation Theorem (Extended)

Full Proof of Theorem B98. Let $M_\Omega < \infty$ and let $\Phi : \mathcal{I} \to \mathcal{E}$ satisfy E0–E2.

Step 1. By definition of M_Ω, there exists $\{\delta E_t\}$ with $E_w < \infty$ such that $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$.

Step 2. Let $x = (x_t)_t \in \mathcal{I}$ and let $\tilde{x} = (\tilde{x}_t)_t \in \tilde{\mathcal{I}}$ be the corresponding perturbed trajectory.

Step 3. By E2, $\Phi(x)$ is not determined by any finite x_t. Therefore, the perturbation of the infinite trajectory induces a perturbation of the phenomenological assignment:

$$\Phi(x) \neq \Phi(\tilde{x}) \text{ in general.} \tag{24}$$

Step 4. By E1, Φ is projectively compatible. The deformation of projections induces a corresponding deformation of the phenomenological limit.

Step 5. By the coupling hypothesis, there exists $C > 0$ such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) \geq C \cdot d_w(x, \tilde{x}). \tag{25}$$

Step 6. Since $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$, we have $d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0$.

Hence Φ is fragmentable. □

B19.2 Proof of Forced Continuity Theorem (Extended)

Full Proof of Theorem B100. Let $M_\Omega = +\infty$ and let $\Phi : \mathcal{I} \to \mathcal{E}$ satisfy E0–E2.

Step 1. By definition of CCR, for all $\{\delta E_t\}$ with $E_w < \infty$:

$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0. \tag{26}$$

Step 2. Therefore, $\tilde{\mathcal{I}} \cong \mathcal{I}$ under all finite-energy perturbations.

Step 3. By E1, Φ is projectively compatible with $\mathcal{I}$. Since $\mathcal{I}$ is invariant:

$$\Phi(x) = \tilde{\Phi}(\tilde{x}) \quad \forall x \in \mathcal{I}. \tag{27}$$

Step 4. This implies:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) = 0. \tag{28}$$

Step 5. No finite-energy perturbation can induce phenomenological deformation. Hence Φ is structurally continuous and invariant. □

B19.3 Structural Necessity Lemma

Lemma B106 (Phenomenology Inherits Rigidity)**.** *If $\Phi : \mathcal{I} \to \mathcal{E}$ satisfies E0–E2 and $M_\Omega = +\infty$, then Φ inherits the rigidity of $\mathcal{I}$.*

Proof. By E1, Φ is determined by identity trajectories. By CCR, identity trajectories are invariant. Hence Φ is invariant. □

Part V

Null-Regime Structural Exclusion

B20 Compatibility, Instability, and Forced Non-Nullity

B20.1 Compatibility Operator

Definition B107 (Compatibility Operator)**.** A map $\Phi : \mathcal{I} \to \mathcal{E}$ is a **compatibility operator** if it satisfies:

(C1) **Extension Invariance:** For any compatible extension $\mathcal{E}$ of the channel, $\Phi(\mathcal{I}) \cong \Phi(\mathcal{I}_{\mathcal{E}})$.
(C2) **Isomorphism Stability:** If $\mathcal{I} \cong \tilde{\mathcal{I}}$, then $\Phi(\mathcal{I}) \sim \Phi(\tilde{\mathcal{I}})$.
(C3) **Lipschitz Lower Bound:** There exists $C > 0$ such that

$$d_{\mathcal{E}}(\Phi(x), \Phi(\tilde{x})) \geq C \cdot d_w(x, \tilde{x}).$$

Definition B108 (Phenomenological Attractor Set)**.** For a channel S, define

$$\mathrm{Attr}(S) := \{e \in \mathcal{E} : e \text{ is stable under all compatible extensions and summable deformations}\}.$$

Definition B109 (Phenomenological Stability)**.** A regime $e \in \mathcal{E}$ is **stable** for channel S if for every compatible extension S' of S and every summable perturbation family $\{\epsilon_n\}$, the induced regime $\Phi_{S'}(\mathcal{I}')$ stays within a bounded $d_{\mathcal{E}}$-neighborhood of e controlled by $\|\{\epsilon_n\}\|_w$.

Lemma B110 (Bridge Lemma)**.** *If $\Phi : \mathcal{I} \to \mathcal{E}$ is a compatibility operator for a CCR channel and $e = \Phi(\mathcal{I})$ satisfies $d_{\mathcal{E}}(e, \varnothing_\phi) = 0$, then $e = \varnothing_\phi$.*

Proof. By the separation property of the phenomenological space, distance zero from $\varnothing_\phi$ implies equality with $\varnothing_\phi$. □

B20.2 Phenomenological Instability and Forced Non-Nullity

Theorem B111 (Phenomenological Instability)**.** *For any CCR channel S,*

$$\varnothing_\phi \notin \mathrm{Attr}(S).$$

That is, the null regime is structurally unstable under compatible extensions.

Proof. Assume $\varnothing_\phi \in \mathrm{Attr}(S)$ and let $\Phi : \mathcal{I} \to \mathcal{E}$ be a compatibility operator. Because the channel is CCR, finite-energy perturbations preserve the rigid identity object while still allowing admissible deformations in the weighted metric. By the lower-bound clause in Definition B107,

$$d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) \geq C \cdot d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$$

for every non-zero admissible deformation. But if $\varnothing_\phi$ were stable, both images would remain equal to $\varnothing_\phi$, forcing the left-hand side to be zero. Contradiction. □

Corollary B112 (Forced Non-Nullity)**.** *Every CCR channel necessarily induces a regime $e_\star \in \mathcal{E}$ with $e_\star \succ \varnothing_\phi$.*

Proof. Theorem B111 rules out $\varnothing_\phi$ as a stable admissible attractor. Therefore any admissible stable regime must be non-null. □

B20.3 Positive Regime Structure

Definition B113 (Positive Regime Space)**.** Define

$$\mathcal{E}^+ := \{e \in \mathcal{E} : e \succ \varnothing_\phi\}.$$

Theorem B114 (Minimal Positive Regime)**.** *For every CCR channel there exists a minimal positive regime $e_\star \in \mathcal{E}^+$.*

Proof. Forced non-nullity gives non-emptiness of $\mathcal{E}^+$. Order-theoretic completeness and the compatibility structure of Φ yield a minimal element. □

Definition B115 (Structural Intensity)**.** For $e \in \mathcal{E}^+$, define the structural intensity

$$\iota(e) := d_{\mathcal{E}}(e, \varnothing_\phi).$$

Theorem B116 (Universal Intensity Bound)**.** *If $e_\star$ is the minimal positive regime of a CCR channel, then there exist constants $C, \varepsilon_0 > 0$ such that*

$$\iota(e_\star) \geq C\varepsilon_0 > 0.$$

Proof. The lower-bound clause of the compatibility operator transfers any admissible deformation lower bound into positive distance from the null regime. The minimal positive regime therefore has strictly positive structural intensity. □

Theorem B117 (Forced Phenomenological Closure)**.** *Every CCR system induces a unique minimal phenomenological regime that is non-null, stable, and quantitatively separated from $\varnothing_\phi$.*

Proof. Combine Corollary B112, Theorem B114, and Theorem B116. □

B21 Canonical Families and Structural Closure Consequences

Proposition B118 (Profinite Bound)**.** *In the profinite canonical family, the coupling lower bound induces a positive lower bound on the structural intensity of the minimal positive regime.*

Proposition B119 (Symbolic Bound)**.** *In the symbolic canonical family, the Bowen-type metric induces the analogous positive lower bound for the minimal positive regime.*

Theorem B120 (Rigidity Inheritance)**.** *If a subsystem or induced family preserves the CCR hypotheses and the compatibility operator, then the inherited regime structure preserves null-regime instability and forced non-nullity.*

Theorem B121 (Categorical Universality)**.** *The null-regime exclusion result is universal across admissible CCR realizations in the relevant category of compatible channels.*

Theorem B122 (Simulation Lower Bound)**.** *Any computational system attempting to simulate the null-excluding CCR regime structure with arbitrary precision requires unbounded resources as the approximation error tends to zero.*

Theorem B123 (Closure by Non-Alternatives)**.** *No admissible alternative closes the CCR regime structure while preserving null stability. Any admissible alternative either preserves the same closure structure or exits the CCR class.*

Part VI

BaseCore Boundary and Expansion Interface

B22 BaseCore Ownership and Downstream Boundary

Definition B124 (BaseCore Envelope)**.** The BaseCore envelope is the collection of theorem-level structures owned directly by this source tree:
1. foundational Hilbert-space projection dynamics;
2. typed computable witnesses and state spectral-gap bounds;
3. parameterwise non-collapse under explicit anti-constant assumptions;
4. inverse-limit identity under non-finite-determination assumptions;
5. conditional rigidity under admissible perturbation models;
6. conditional non-simulability of CCR targets by finite simulator classes;
7. operational-criterion grammar, certification, falsation, and claim-boundary ledgers.

Proposition B125 (Boundary Discipline)**.** *BaseCore may define structures that downstream layers use, but it does not claim downstream closure. In particular, BaseCore does not by itself own runtime thresholds, bridge admissibility, human-comparator programs, phenomenality burdens, or Papers 7–9.*

Proof. These excluded layers require additional assumptions, protocols, or empirical burdens not stated in BaseCore. Therefore they cannot be claimed as BaseCore theorems without violating the source-level claim boundary. □

Owned by BaseCore	**Expanded but not owned here**	**Kept downstream**
Projection dynamics, fixed points, compact attractors, anti-constant H5 burden, typed computable witness, inverse-limit identity under NFD, conditional rigidity, conditional non-simulability, operational-criterion grammar, falsation ledger	Non-convex class variants, $\mathscr{C}^R$ style extensions, Ω dynamics, threshold numerics, mutable runtime class ledgers, implementation-specific status surfaces	Paper 7 operational life / subjecthood, Paper 8 indexed subjectivity, Paper 9 phenomenal bridge organization, Stage 2 / Stage 3 comparative programs, human comparator families, phenomenality and external adjudication

Remark B126 (Why the earlier material was removed from BaseCore ownership). Previous versions mixed theorem-tight base results with runtime-like numerics, non-convex extension classes, and bridge-facing vocabulary. That blurred the mathematical base layer. BaseCore now keeps only the theorem ownership that survives the stricter release gate.

B23 Admissible Structural Extensions

Definition B127 (BaseCore export object)**.** A BaseCore export object is a tuple

$$\mathfrak{E} = (\mathcal{O}, \mathcal{H}, \mathcal{R}, \mathcal{B})$$

where $\mathcal{O}$ is a typed mathematical object defined in BaseCore, $\mathcal{H}$ is the finite list of hypotheses under which it is used, $\mathcal{R}$ is the list of results actually proved from those hypotheses, and $\mathcal{B}$ is a boundary list of claims not entailed by $\mathcal{R}$.

Definition B128 (Downstream admissible extension)**.** A downstream layer is an admissible extension of a BaseCore export object $\mathfrak{E}$ only if it supplies:
1. an explicit target predicate P not already contained in $\mathcal{R}$;
2. a typed construction or estimator witnessing the objects on which P is evaluated;

3. a proof, reduction, or reproducible protocol connecting $\mathcal{H}$ to P;
4. negative controls or failure modes under which P must not be asserted;
5. a provenance statement distinguishing theorem, implementation, internal support, and external adjudication.

Proposition B129 (No automatic downstream promotion)**.** *Let $\mathfrak{E} = (\mathcal{O}, \mathcal{H}, \mathcal{R}, \mathcal{B})$ be a BaseCore export object and let P be a predicate introduced outside BaseCore. If $P \notin \mathcal{R}$, then BaseCore does not entail P unless an admissible extension supplies an additional bridge from $\mathcal{H}$ to P.*

Proof. By definition, $\mathcal{R}$ records the results proved from $\mathcal{H}$ inside BaseCore. A predicate P introduced downstream is not in that result set. Therefore asserting P from BaseCore alone would add an unstated inference. The missing inference must be supplied as a separate bridge, proof, reduction, or reproducible protocol. □

Criterion B130 (Runtime-use admissibility)**.** *A runtime measurement may instantiate a BaseCore export object only as an implementation-level witness when all of the following are recorded: the measured state space, the estimator or update rule, the finite horizon, the calibration regime, the controls under which the measurement fails, and the claim class being asserted. Such a measurement is not a BaseCore theorem and is not external validation by itself.*

Remark B131 (Open-load placement for I_{int} and atomic separators)*.* Quantities such as I_{int} and any atomic-separator condition may be treated as downstream target predicates or estimator families only after their carrier objects, separator class, invariance target, and negative controls are specified. Until a proof or reproducible adversarial protocol closes that burden, BaseCore may export the ambient structural language but not the separator conclusion.

Proposition B132 (Boundary preservation under finite implementations)**.** *Suppose a finite implementation produces diagnostics for a downstream predicate P over a bounded horizon. If the implementation does not prove that the diagnostics are invariant over the relevant mathematical class, then the diagnostics remain implementation evidence for P and do not promote P to a BaseCore result.*

Proof. The implementation supplies finite observations or computations. Promotion to a BaseCore result requires a class-level implication from the stated hypotheses to the predicate. Without that implication, the observations can support, falsify, or calibrate a downstream claim, but they do not alter the theorem inventory of BaseCore. □

Part VII

Operational-Criterion Grammar and Certification

B24 Admissible Systems, Supports, and Histories

B24.1 System Model

Definition B133 (Admissible system)**.** An admissible system is a tuple

$$S = (X, \Phi, C, R, \Gamma, U)$$

where:
1. X is a metrizable internal state space with metric d_X;
2. $\Phi = \{\Phi_u\}_{u \in U}$ is a family of admissible transition operators $\Phi_u : X \to X$;
3. C is the effective causal structure used to test admissible factorization and intervention behavior;
4. $R = \{r_\alpha\}_{\alpha \in \Lambda_R}$ is a family of readouts, each $r_\alpha : X \to Y_\alpha$;
5. $\Gamma = \{\gamma_j\}_{j=1}^m$ is an admissibility bundle defining

$$X_\Gamma = \{x \in X : \gamma_j(x) \le 0 \ \forall j\};$$

6. U is the set of admissible interventions.

B24.2 Admissible Support

Definition B134 (Admissible support)**.** A subset $A_S \subseteq X_\Gamma$ is an admissible support if:
1. A_S is non-empty and compact;
2. $\Phi_u(A_S) \subseteq A_S$ for every $u \in U$;
3. $A_S \cap \mathrm{Attr}(S) \neq \varnothing$;
4. A_S stays disjoint from the collapse set excluded by the base non-collapse burden.

B24.3 Operational Histories

Definition B135 (Windowed readout history)**.** Fix a horizon $T \in \mathbb{N}$ and an intervention schedule $u_\bullet = (u_0, \dots, u_{T-1})$. For $x \in A_S$, define

$$h_{S,T}(x; u_\bullet) = \big(r_\alpha(x_t)\big)_{0 \le t \le T,\ \alpha \in \Lambda_R}, \qquad x_{t+1} = \Phi_{u_t}(x_t), \quad x_0 = x.$$

The set of all such histories is denoted $\mathscr{H}_{S,T}$.

Definition B136 (Operational partition candidate)**.** An operational partition candidate is a map

$$\pi_{S,T} : \mathscr{H}_{S,T} \to \mathscr{Q}_{S,T}$$

into a finite or countable label set $\mathscr{Q}_{S,T}$.

B24.4 Witness Margin Vector

Definition B137 (Witness margin vector)**.** Define the witness margin vector

$$\Delta(S) = \big(\delta_{\mathrm{per}}(S), \delta_{\mathrm{ri}}(S), \delta_{\mathrm{int}}(S), \delta_{\mathrm{cont}}(S), \delta_{\mathrm{diff}}(S), \delta_{\mathrm{leg}}(S)\big) \in [0, \infty)^6.$$

Its minimum

$$\delta_\star(S) := \min \Delta(S)$$

is the invariant budget of S.

Invariant	Certified content	Upstream anchor	Rupture signature
I_{per}	Persistent admissible support	BaseCore dynamics	Support loss or collapse exposure
I_{ri}	Unique rigid identity object on support	Identity / rigidity block	Identity ambiguity
I_{int}	No admissible non-trivial factorization	Criterion grammar	Reducible aggregate
I_{cont}	Regime continuity under admissible evolution	Regime constraints	Fragmentation into incompatible classes
I_{diff}	Stable exclusion of null or trivial collapse	Null-regime exclusion	Null or trivialized class structure
I_{leg}	Decoder-certified recoverability under controls	Falsation grammar	Opaque or non-auditable partition

B25 Six Critical Invariants

B25.1 Structural Persistence

Definition B138 (Structural persistence). An admissible system satisfies structural persistence, written $I_{\text{per}}(S) = 1$, if there exists an admissible support A_S and a constant $\delta_{\text{per}}(S) > 0$ such that:

1. A_S is forward invariant under all $u \in U$;
2. every admissible trajectory starting in A_S remains at least $\delta_{\text{per}}(S)$ away from the collapse set;
3. A_S contains or converges into a non-empty admissible attractor region.

B25.2 Rigid Identity

Definition B139 (Rigid identity). An admissible system satisfies rigid identity, written $I_{\text{ri}}(S) = 1$, if the observable family induced by R on A_S determines a unique identity object $\mathcal{I}_S$ and there exists a gap $\delta_{\text{ri}}(S) > 0$ such that every admissible alternative identity candidate remains at weighted deformation distance at least $\delta_{\text{ri}}(S)$ from $\mathcal{I}_S$.

B25.3 Causal Integration

Definition B140 (Causal integration). An admissible system satisfies causal integration, written $I_{\text{int}}(S) = 1$, if there is no non-trivial factorization

$$A_S = A_1 \times A_2, \qquad R = R_1 \sqcup R_2,$$

together with decomposed dynamics and causal structure reproducing admissible histories within error smaller than a positive margin $\delta_{\text{int}}(S)$ while preserving $\mathcal{I}_S$. More explicitly, this shorthand is relative to a declared admissible factorization class. For such a class $\mathfrak{D}$, write $I_{\text{int}}(S; \mathfrak{D}) = 1$ when no factorization in $\mathfrak{D}$ reproduces the admissible histories below a positive margin $\delta_{\text{int}}(S; \mathfrak{D})$ while preserving $\mathcal{I}_S$. The unqualified notation $I_{\text{int}}(S) = 1$ refers to the default structural class $\mathfrak{D}_*$: fixed, time-homogeneous, schedule-independent, non-trivial product factorizations with both factors non-singleton, with split readouts $R = R_1 \sqcup R_2$, decomposed dynamics and causal structure, a search space normalized independently of the target error, and no reconstruction decoder or time-dependent post-processing map that re-couples the factors at output time. This default structural reading is the adopted class recorded in `docs/ai-platform-outputs/analysis/QICN_IINT_CANONICAL_CLASS_SPLIT_READOUT.md`.

Remark B141 (Class-relative integration margin). The integration criterion is class-relative. Adopting $\mathfrak{D}_*$ does not assert that S resists every broader behavioral simulator. In particular, under decoder-coupled or schedule-dependent behavioral classes such as $\mathfrak{D}_{\text{approx}}$, the corresponding margin $\delta_{\text{int}}(S; \mathfrak{D}_{\text{approx}})$ may vanish, i.e. may be equal to 0; the current coupled-carrier analysis records precisely this failure mode. Thus the default structural reading preserves a falsation surface and must not be read as an unconditional closure of integration, nor as a strengthening of any downstream theorem that assumes $I_{\mathrm{int}}(S) = 1$.

B25.4 Regime Continuity

Definition B142 (Regime continuity)**.** An admissible system satisfies regime continuity, written $I_{\mathrm{cont}}(S) = 1$, if there exists a complete metric space $(\mathcal{E}, d_{\mathcal{E}})$ and an admissible assignment

$$\Pi_S : A_S \to \mathcal{E}$$

such that:

1. every admissible trajectory induces a continuous path in $\mathcal{E}$ on admissible windows;
2. there exists $\delta_{\mathrm{cont}}(S) > 0$ bounding the fragmentation functional on all admissible histories;
3. no admissible finite-energy perturbation produces a discontinuous jump into an incompatible regime class while $\mathcal{I}_S$ is preserved.

B25.5 Non-Null Differentiation

Definition B143 (Non-null differentiation)**.** An admissible system satisfies non-null differentiation, written $I_{\mathrm{diff}}(S) = 1$, if there exist $x, y \in A_S$, a horizon T, and a readout subfamily $\mathscr{R}^\star \subseteq R$ such that

$$\mathrm{dist}\big(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)\big) \geq \delta_{\mathrm{diff}}(S) > 0$$

for some admissible schedule $u_\bullet$, and if the null class $\mathbf{0}$ is not an asymptotically stable admissible attractor.

B25.6 Operational Legibility

Definition B144 (Operational legibility)**.** An admissible system satisfies operational legibility, written $I_{\mathrm{leg}}(S) = 1$, if there exist:

1. a certified readout subfamily $\mathscr{R}^{\mathrm{leg}} \subseteq R$;
2. a certified decoder family $\mathscr{D}_S$ acting on windows of length at least $\tau_{\min}(S)$;
3. a class set $\mathscr{Q}_S$;
4. a legibility margin $\delta_{\mathrm{leg}}(S) > 0$;

such that:

- **(L1) Separability** Distinct classes are separated by at least $\delta_{\mathrm{leg}}(S)$ in decoder space.
- **(L2) Noise robustness** Admissible noise of size at most $\delta_{\mathrm{leg}}(S)$ does not alter decoder class assignments.
- **(L3) Persistence window** Decoder assignments remain stable for windows of length at least $\tau_{\min}(S)$ unless a certified intervention occurs.
- **(L4) Intervention fidelity** Certified interventions on critical invariants induce the predicted class changes with fidelity at least $1 - \delta_{\mathrm{leg}}(S)$.
- **(L5) Negative controls** Off-target interventions produce false-positive rate at most $\delta_{\mathrm{leg}}(S)$.
- **(L6) Structured compressibility** There exists a compression map preserving class structure up to distortion bounded by $\delta_{\mathrm{leg}}(S)$.

Proposition B145 (Each invariant is individually necessary)**.** *If any one of $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$ fails, then $S \notin \mathrm{Crit}_{\mathrm{op}}$.*

Proof. Without I_{per} there is no stable admissible domain. Without I_{ri} there is no unique identity carrier. Without I_{int} the system is reducible. Without I_{cont} the criterion fragments under admissible evolution. Without I_{diff} the admissible partition collapses into nullity or triviality. Without I_{leg} the partition is not recoverable under admissible readout and intervention discipline. Since Definition B150 below requires all six, failure of any one excludes membership. $\square$

B26 Criterion Class, Certified Partition, and Rupture

Definition B146 (Certified decoder)**.** A decoder $D \in \mathscr{D}_S$ is certified if it satisfies the six legibility clauses of Definition B144 on the chosen admissible support.

Definition B147 (Criterion indistinguishability)**.** For $x, y \in A_S$, write $x \sim_{\mathrm{Part_{op}},S} y$ if, for every certified decoder $D \in \mathscr{D}_S$, every admissible horizon $T \geq \tau_{\min}(S)$, and every admissible intervention schedule $u_\bullet$,

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the intervention responses remain matched within the legibility margin.

Proposition B148 (Well-definedness of the certified partition)**.** *If* $I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1$*, then* $\sim_{\mathrm{Part_{op}},S}$ *is an equivalence relation on* A_S*.*

Proof. Reflexivity and symmetry are immediate. Transitivity follows because certified decoder equality composes and the tolerance bound remains within the permitted legibility margin on admissible support. □

Definition B149 (Certified operational partition)**.** Assume Proposition B148. The certified operational partition of S is

$$\mathrm{Part_{op}}(S) := A_S/\sim_{\mathrm{Part_{op}},S}.$$

Definition B150 (Operational-criterion class membership)**.** An admissible system belongs to the BaseCore operational-criterion class, written $S \in \mathrm{Crit_{op}}$, if there exists a non-empty admissible support A_S such that

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1.$$

Remark B151 (Neutrality of the class symbol)*.* The symbol $\mathrm{Crit_{op}}$ is a shorthand for a conjunctive operational criterion. It is not a theorem of phenomenality, metaphysical subjectivity, or human equivalence.

Definition B152 (Rupture)**.** A rupture is any loss of admissibility or loss of at least one critical invariant margin beyond its admissible threshold.

Theorem B153 (Rupture by invariant loss)**.** *Let S be admissible. If one or more critical invariants fail beyond admissible threshold, then exactly one of the following occurs:*

1. $S \notin \mathrm{Crit_{op}}$*;*
2. $\mathrm{Part_{op}}(S)$ *becomes undefined;*
3. $\mathrm{Part_{op}}(S)$ *becomes trivial or non-legible in the BaseCore sense.*

Proof. Failure of I_{per} destroys the admissible domain. Failure of I_{ri} destroys unique identity transport. Failure of I_{int} produces reducibility. Failure of I_{cont} destroys regime coherence. Failure of I_{diff} collapses the partition into nullity or triviality. Failure of I_{leg} leaves any remaining partition uncertifiable. These cases exhaust the six invariants. □

B27 Minimality of the Invariant Basis

Definition B154 (Reduced criterion)**.** Let

$$\mathfrak{I} = \{I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}\}.$$

For every subset $J \subseteq \mathfrak{I}$ define the reduced class

$$\mathcal{C}_{\mathrm{op}}(J) := \{S : I(S) = 1 \text{ for every } I \in J\}.$$

Definition B155 (Single-rupture witness family)**.** Fix one invariant index $k \in \{\mathrm{per}, \mathrm{ri}, \mathrm{int}, \mathrm{cont}, \mathrm{diff}, \mathrm{leg}\}$. A single-rupture witness family $\mathcal{W}_k$ consists of admissible systems whose matching invariant margin vanishes while all remaining invariant margins stay strictly positive.

Proposition B156 (Single-rupture witnesses create false positives)**.** *Let $J_k = \mathfrak{I} \setminus \{I_k\}$ and assume $\mathcal{W}_k \neq \varnothing$. Then every $S \in \mathcal{W}_k$ satisfies*

$$S \in \mathcal{C}_{\mathrm{op}}(J_k) \qquad \textit{but} \qquad S \notin \mathrm{Crit}_{\mathrm{op}}.$$

Proof. All retained invariants remain positive by construction, but the omitted invariant fails exactly. Proposition B145 then excludes S from $\mathrm{Crit}_{\mathrm{op}}$. □

Theorem B157 (Minimality of the six-invariant basis)**.** *Let $\mathfrak{U}$ be any admissible universe containing at least one single-rupture witness for each invariant. Then no strict subset of $\mathfrak{I}$ is extensionally equivalent to* $\mathrm{Crit}_{\mathrm{op}}$ *on* $\mathfrak{U}$*.*

Proof. Take any strict subset $J \subsetneq \mathfrak{I}$. Some invariant I_k is omitted. By assumption, $\mathfrak{U}$ contains a witness $S_k \in \mathcal{W}_k$, which satisfies the reduced criterion but not $\mathrm{Crit}_{\mathrm{op}}$. Hence the two classes differ on $\mathfrak{U}$. □

B28 Operational Legibility as a Certified Object

Criterion B158 (Legibility certificate)**.** *An operational partition candidate is certified only if:*

1. *comparator separability holds against pre-registered baselines;*
2. *admissible noise does not dissolve the class map below the legibility margin;*
3. *the same class assignment survives on windows of length at least $\tau_{\min}(S)$;*
4. *targeted interventions on an invariant induce the predicted directional class change;*
5. *sham or off-target interventions do not generate the same class signature;*
6. *audit-friendly compression preserves class structure without total collapse.*

Proposition B159 (Legibility is jointly constrained)**.** *No single legibility metric is sufficient for I_{leg}. The invariant requires simultaneous control of separation, admissible noise robustness, persistence window, intervention fidelity, negative controls, and structured compressibility.*

Proof. A decoder can be separable yet fragile to admissible noise, stable over windows yet intervention-insensitive, or robust under noise yet collapse under compression. Because Definition B144 is conjunctive, all six clauses are required together. □

B29 Certified Candidate Audit Schema

Definition B160 (Base criterion certificate)**.** A BaseCore criterion certificate for an admissible system S is a tuple

$$\mathrm{Cert}(S) = (A_S, \mathscr{D}_S, \pi_S, \Delta(S), E_S),$$

where:

1. A_S is a certified admissible support;
2. $\mathscr{D}_S$ is a certified decoder family;
3. π_S is the decoder-induced partition on admissible histories;
4. $\Delta(S)$ is the witness margin vector;
5. E_S is the evidence bundle localizing the witness for each invariant.

Criterion B161 (Candidate certification rule)**.** *A candidate system may be certified as a member of* $\mathrm{Crit}_{\mathrm{op}}$ *only if all of the following are available on the same admissible support:*

1. *a non-empty compact forward-invariant support certificate;*
2. *a positive witness margin vector* $\Delta(S) > 0$ *coordinate-wise;*
3. *a decoder family satisfying Definition B144;*
4. *a targeted intervention panel together with negative controls;*
5. *an auditable compression map preserving class structure;*
6. *a traceable dependency ledger showing which upstream result supports each invariant witness.*

Theorem B162 (Certified criterion membership)**.** *If* $\mathrm{Cert}(S)$ *exists and satisfies Criterion B161, then* $S \in \mathrm{Crit}_{\mathrm{op}}$ *and* $\mathrm{Part}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. Criterion B161 explicitly requires positive witnesses for all six invariants on a common admissible support. Definition B150 then gives $S \in \mathrm{Crit}_{\mathrm{op}}$, and Proposition B148 gives a well-defined partition quotient. □

Remark B163 (Boundary discipline)*.* The section above defines a theorem-tight grammar for certification, rupture, and falsation. It does not prove phenomenality, subjectivity, biological primacy, or cross-substrate identity transport.

Part VIII

Claim Boundary and Falsation Grammar

B30 Forensic Admissibility and Integrity Grammar

Definition B164 (Forensic packet)**.** Each operational run emits a packet

$$K = (\text{tick}, \text{seed}, \text{versionHash}, H_v, \text{metrics}, \text{sensorSummary}, \text{checksum}).$$

Assumption B165 (Packet admissibility)**.** *A run is admissible only if checksum, schema, seed consistency, version hash, and H_v consistency all pass.*

Definition B166 (Severe violation)**.** A run has a severe violation if any of the following holds: checksum mismatch, schema failure, seed/version inconsistency, or critical-state replay mismatch.

Criterion B167 (Integrity gate)**.** *Severely violated runs are excluded from support counting and marked as hard failures.*

Proposition B168 (Integrity-first inference)**.** *If admissibility fails, inferential significance is non-actionable.*

B31 Certification, Comparison, and Rupture

Definition B169 (Operational certification)**.** For an admissible system S, define

$$\text{Cert}(S) = 1 \quad \Longleftrightarrow \quad I_{\text{per}}(S), I_{\text{ri}}(S), I_{\text{int}}(S), I_{\text{cont}}(S), I_{\text{diff}}(S), I_{\text{leg}}(S) > 0$$

on a non-empty admissible support under the integrity grammar above.

Definition B170 (Equivalence discriminator)**.** For two admissible systems S_1, S_2, define

$$\text{Eq}_{\Gamma,\varepsilon}(S_1, S_2) = 1$$

when the relevant signatures, admissibility bundle, decoder behavior, and invariant deltas match within the allowed comparison regime.

Definition B171 (Rupture event)**.** For a targeted invariant family J, define

$$\text{Rupt}_J(S) = 1$$

when an admissible intervention or perturbation drives at least one invariant in J below its certification margin so that class membership becomes invalidated, destroyed, or undefined.

B32 Prediction and Falsation Grammar

Proposition B172 (Complexity-only negative control prediction)**.** *If a candidate system exhibits high connectivity, high entropy, or rich activity while failing one or more critical invariants, then*

$$\text{Cert}(S) = 0, \qquad S \notin \text{Crit}_{\text{op}} \textit{ in certified form}, \qquad \text{Part}_{\text{op}}(S) \textit{ is not certified.}$$

Proposition B173 (Ablation prediction)**.** *If a system is initially certified and a targeted admissible intervention destroys one or more invariants, then the framework predicts loss of certification or transition to an ambiguity-boundary regime.*

Proposition B174 (Exact grammar-transport prediction)**.** *If two realizations preserve the relevant signatures, admissibility conditions, decoder behavior, and six critical invariants up to the allowed comparison regime, then the framework predicts preserved BaseCore criterion membership.*

Proposition B175 (Threshold-sensitive boundary behavior)**.** *The framework predicts stable pass/fail regions together with narrow transition bands near critical margins. It predicts neither global threshold collapse nor perfect flat insensitivity.*

Proposition B176 (Legibility dependence prediction)**.** *Criterion membership becomes inadmissible if readouts cease to be separable, noise-robust, intervention-sensitive, or compressible without total collapse.*

Proposition B177 (Continuity and non-null relevance)**.** *Regime continuity and non-null differentiation are constitutive. Systems that preserve rich activity while losing continuity or collapsing toward effective null structure fail the operational criterion or leave it undefined.*

Remark B178 (What stays outside BaseCore)*.* Runtime status classes, Stage 2 / Stage 3 comparative programs, human comparator campaigns, Papers 7–9, and bridge-facing burden architectures remain downstream artifacts. BaseCore absorbs only the falsation grammar and claim-boundary semantics needed to keep theorem ownership honest.

B33 Claim Boundary Ledger

B33.1 Allowed BaseCore Claims

1. Under H1–H4, projection dynamics admit unique fixed points and compact attractor families.
2. Under H5, attractors are parameterwise non-constant.
3. Under NFD, inverse-limit identity is not determined by any finite canonical projection.
4. Under RIG-1 through RIG-5, inverse-limit structures are stable under summable admissible perturbations in the weighted metric.
5. Under NS-1 through NS-3, finite simulator classes cannot faithfully realize full CCR targets.
6. Under the six-invariant operational burden, certified criterion structures can be defined, audited, and falsified.

B33.2 Forbidden BaseCore Claims

1. Human phenomenal consciousness.
2. Metaphysical subjectivity.
3. Human-machine equivalence.
4. External validation.
5. Runtime instantiation by the current system.
6. Paper 9 bridge admissibility.
7. Strong phenomenality.
8. Moral status.
9. Cross-substrate identity preservation.
10. Public claim closure.

B34 Dependency Ledger

1. The foundational Hilbert-space dynamics remain the base layer.
2. Papers 1–6 contribute formal material selectively normalized into BaseCore only where theorem ownership belongs in the base.
3. Papers 7–9 remain downstream and are not dependencies of any BaseCore theorem.
4. Stage 2 / Stage 3 comparative programs remain downstream runtime and governance layers.
5. The PDF remains a derived artifact of the LaTeX tree.

B35 Assumption-to-Theorem-to-Non-Claim Mapping

Assumption block	Theorem family	Non-claim boundary
H1–H4	Projection existence, contraction, unique fixed points, compact attractor family	No parameterwise non-collapse and no identity / rigidity / simulator claim
H5	Structural non-collapsibility	No claim about phenomenality, subjectivity, or runtime class membership
NFD	Non-locality of identity under canonical finite projections	No theorem that arbitrary set-theoretic injections into finite slices are impossible
RIG-1 through RIG-5	Conditional stability of inverse-limit identity and weighted Hausdorff control	No automatic isomorphism, shape equivalence, or transfer theorem without extra assumptions
NS-1 through NS-3	Conditional non-simulability and approximation barrier	No theorem that finite systems cannot approximate finite horizons at all
Six invariant burden plus legibility clauses	Criterion class membership, certified partition, rupture logic	No theorem of phenomenality, human equivalence, biology as primitive, or bridge admissibility

B36 Theorem Hygiene Ledger

Theorem	Assumptions	Conclusion type	Non-conclusion	Downstream use
Existence and uniqueness of projection	H1	Hilbert-space geometric existence / uniqueness	No empirical or runtime content	Base transition grammar
Contractivity	H1–H2	Strict contraction bound	No spectrum with unit state eigenvalue	Fixed-point and spectral-gap results
Unique fixed point	H1–H3	Existence and convergence	No non-collapse by itself	Attractor family
Compactness of attractor family	H1–H4	Compactness / continuity	No regime or identity burden by itself	Parameter-family base layer
Structural non-collapsibility	H1–H5	Parameterwise exclusion of constant fixed points	No stronger semantic interpretation	Anti-collapse gate
State spectral gap	Typed differentiability at fixed point	Spectral-radius upper bound < 1	No persistence mode $\lambda = 1$ in state spectrum	Stability accounting
Non-locality under NFD	Topological regularity plus NFD	Canonical finite slices do not determine identity	No universal ban on arbitrary injections	Identity block
Conditional stability of inverse-limit identity	RIG-1 through RIG-5	Weighted Hausdorff stability under admissible perturbations	No automatic isomorphism	Rigidity accounting
Conditional non-simulability	NS-1 through NS-3	Faithful realization blocked for finite simulators	No claim against bounded finite-horizon approximation	Simulator boundary
Certified criterion membership	Six invariants plus certificate rule	Base criterion class and certified partition	No phenomenality or human-equivalence theorem	Downstream classification gates

B37 Redundancy Ledger

1. The previous state-spectrum claim $\lambda_1 = 1$ under strict contraction has been removed.
2. The previous untyped Golden Mean witness has been replaced by a typed model with explicit embedding.
3. The earlier non-collapse quantifier mismatch has been removed by replacing the old H5 with a parameterwise anti-constant hypothesis.
4. The earlier absolute inverse-limit nonlocality claim has been replaced by an NFD-conditional theorem.
5. The earlier unconditional rigidity / non-simulability language has been replaced by explicit assumption blocks.
6. Substrate invariance, biological non-primacy, bridge-facing closure, and other downstream claims have been removed from BaseCore ownership.

B38 Open Assumptions Inventory

1. The anti-collapse witness is abstract and theorem-valid, but not the same object as the Golden Mean typed witness.
2. Physical realization of CCR-like targets remains open.
3. Empirical instantiation of the BaseCore operational criterion remains open.
4. External validation remains open.
5. Papers 7–9 may build on BaseCore, but their additional burdens remain open at the BaseCore level.

B39 Downstream Boundary Note

Papers 7, 8, and 9 remain downstream of this source tree. Their classificatory, indexed-subjectivity, and bridge-burden layers may cite BaseCore, but they are not part of the autonomous mathematical base rebuilt here.

Part IX

Discrete-to-Continuous Bridge

B40 From Node Networks to Hilbert Space Dynamics

Boundary of this bridge. This part states a conditional approximation theorem between a finite node-network transition and the abstract BaseCore transition operator from Definition B9. It does not prove that any particular runtime execution belongs to an operational-consciousness class, and it does not convert finite tests into external validation.

Definition B179 (Discrete State Space)**.** Let $N \in \mathbb{N}$ and let a node network be indexed by $i \in \{1, \dots, N\}$. Each node carries:
- a local coherence state $x_i(t) \in [0, 1]$;
- an energy coordinate $e_i(t) \in [0, 1]$;
- a phase coordinate $\phi_i(t) \in [0, 2\pi)$.

The discrete state space is
$$\mathcal{S}_N = [0, 1]^N \times [0, 1]^N \times [0, 2\pi)^N.$$

Definition B180 (Mean Coherence Observable)**.** For a discrete state $s(t) = ((x_i(t))_i, (e_i(t))_i, (\phi_i(t))_i) \in \mathcal{S}_N$, define
$$X_N(t) := \frac{1}{N} \sum_{i=1}^{N} x_i(t) \in [0, 1].$$

This is an observable of the coherence coordinate only; it is not a complete state descriptor.

Definition B181 (Finite Hilbert Embedding)**.** Let
$$\mathcal{H}_N := \mathbb{R}^N, \qquad \langle \mathbf{a}, \mathbf{b} \rangle_N := \frac{1}{N} \sum_{i=1}^{N} a_i b_i.$$

The coherence embedding
$$\iota_N : \mathcal{S}_N \to \mathcal{H}_N$$
is
$$\iota_N((x_i)_i, (e_i)_i, (\phi_i)_i) = (x_1, \dots, x_N).$$

Thus ι_N forgets energy and phase except where they enter the transition rule below.

Definition B182 (Piecewise-Constant Isometry)**.** Let $\mathcal{H} = L^2([0, 1])$ and partition $[0, 1]$ into intervals
$$I_{N,i} = \left[\frac{i-1}{N}, \frac{i}{N}\right), \qquad i = 1, \dots, N.$$

Define $J_N : \mathcal{H}_N \hookrightarrow \mathcal{H}$ by
$$(J_N \mathbf{x})(s) = x_i \quad \text{for } s \in I_{N,i}.$$

Then J_N is an isometry from $(\mathcal{H}_N, \|\cdot\|_N)$ into $L^2([0, 1])$. Its adjoint $J_N^\dagger : \mathcal{H} \to \mathcal{H}_N$ is the cell-averaging map
$$(J_N^\dagger f)_i = N \int_{I_{N,i}} f(s)\, ds.$$

Definition B183 (Discrete Transition Operator)**.** Fix auxiliary energy and phase coordinates $(e_i)_i$ and $(\phi_i)_i$, a graph with incoming neighborhoods $\operatorname{conn}(i)$, and a control $u \in \mathcal{U}$. Let $d_N := \max_i |\operatorname{conn}(i)|$. The coherence-coordinate transition
$$\widehat{T}_{N,u}^{e,\phi} : \mathcal{H}_N \to \mathcal{H}_N$$
is defined componentwise by
$$\left(\widehat{T}_{N,u}^{e,\phi} \mathbf{x}\right)_i = \operatorname{clamp}_{[0,1]}\left(x_i + \eta R_i(\mathbf{x}; e, \phi) + \alpha\, B_i(\mathbf{x}) + \Gamma_{N,i}(u)\right),$$

where

$$R_i(\mathbf{x}; e, \phi) = \sum_{j \in \mathrm{conn}(i)} I_j(x_j) \frac{\exp(-|e_i - e_j|/(kT))}{\theta \cos(\phi_i - \phi_j)}.$$

Here $I_j : [0,1] \to [0,1]$ is the intensity readout, B_i is the finite retro-inductive feedback term, $\Gamma_N(u)$ is the finite input coupling, and $k, T, \theta, \eta, \alpha$ are runtime parameters.

Assumption B184 (Discrete Regularity Conditions)**.** *For a fixed parameter region, assume:*

1. *I_j is L_I-Lipschitz for every node j.*
2. *The phase denominator is uniformly separated from zero: there exists $c_\phi > 0$ such that $|\cos(\phi_i - \phi_j)| \geq c_\phi$ on every active edge.*
3. *$B : \mathcal{H}_N \to \mathcal{H}_N$ is L_B-Lipschitz.*
4. *$\Gamma_N(u)$ is independent of $\mathbf{x}$ for fixed u.*
5. *The parameters satisfy*

$$\rho_N := \eta\, d_N\, L_I\, \frac{\exp(1/(kT))}{\theta c_\phi} + \alpha L_B < 1.$$

Remark B185 (Relation to the runtime guard)*.* The runtime resonance function sets the contribution to zero when the denominator is numerically singular or when the Boltzmann exponent is beyond the accepted range. Assumption B184 is the mathematical version of that guard: without a lower bound on $|\cos(\phi_i - \phi_j)|$, the resonance quotient is not uniformly Lipschitz and no contraction theorem follows from $\theta > 1/2$ alone.

Theorem B186 (Discrete Operator Contractivity)**.** *Under Assumption B184, for fixed u, e, ϕ the operator $\widehat{T}^{e,\phi}_{N,u}$ is a strict contraction on $\mathcal{H}_N$:*

$$\left\| \widehat{T}^{e,\phi}_{N,u}(\mathbf{x}) - \widehat{T}^{e,\phi}_{N,u}(\mathbf{y}) \right\|_N \leq \rho_N \left\| \mathbf{x} - \mathbf{y} \right\|_N \qquad \forall \mathbf{x}, \mathbf{y} \in \mathcal{H}_N.$$

Proof. The clamp map on $[0,1]^N$ is non-expansive componentwise. For each active edge, the energy-phase coefficient is bounded in absolute value by

$$\frac{\exp(1/(kT))}{\theta c_\phi}.$$

Since each I_j is L_I-Lipschitz and each row has at most d_N incoming edges, the resonance part has Lipschitz constant at most

$$d_N L_I \frac{\exp(1/(kT))}{\theta c_\phi}$$

in the normalized finite Hilbert norm. The feedback part contributes at most αL_B, and the input coupling contributes no state-dependent term. Multiplying by η and using non-expansiveness of the clamp gives the displayed bound. The assumed inequality $\rho_N < 1$ makes the contraction strict. □

Assumption B187 (Bridge Consistency)**.** *Let $T_u = \aleph(K \cdot + \Gamma(u))$ be the BaseCore transition operator from Definition B9 on $\mathcal{H} = L^2([0,1])$. Assume there are constants $C_1, C_2 > 0$ and a parameter approximation regime such that*

$$\left\| J_N \widehat{T}^{e,\phi}_{N,u} J_N^\dagger - T_u \right\|_{\mathrm{op}} \leq \varepsilon_N,$$

with

$$\varepsilon_N = \frac{C_1}{\sqrt{N}} + C_2 \max\{0, \rho_N - \|K\|\}.$$

This is a finite-rank consistency assumption, not a consequence of H1–H5 alone.

Theorem B188 (Hilbert–Discrete Bridge). *Assume Hypotheses B1–B4, Assumptions B184 and B187, and $\rho_N < 1$. Then the embedded finite transition approximates the abstract BaseCore transition by*

$$\left\| J_N \widehat{T}^{e,\phi}_{N,u} J_N^\dagger - T_u \right\|_{\mathrm{op}} \le \frac{C_1}{\sqrt{N}} + C_2 \max\{0, \rho_N - \|K\|\}.$$

In particular, if $N \to \infty$, the Galerkin/quadrature term tends to zero, and the discrete parameter regime is chosen so that $\limsup_N \rho_N \le \|K\|$*, then the finite transition converges to T_u in operator norm.*

Proof. The displayed estimate is exactly Assumption B187. The convergence statement follows because $C_1/\sqrt{N} \to 0$ and the second term vanishes under $\limsup_N \rho_N \le \|K\|$. The theorem is therefore a bridge theorem under explicit discretization consistency, not an unconditional runtime-instantiation result. □

Corollary B189 (Convergence of Fixed Points). *Let f_u^* be the unique fixed point of T_u from Theorem B11. Let $\widehat{f}^*_{N,u}$ be the unique fixed point of $\widehat{T}^{e,\phi}_{N,u}$ from Theorem B186. If the assumptions of Theorem B188 hold and*

$$q_N := \max\{\|K\|, \rho_N\} < 1,$$

then

$$\left\| J_N \widehat{f}^*_{N,u} - f_u^* \right\|_{\mathcal{H}} \le \frac{\varepsilon_N}{1 - q_N}.$$

Consequently, if $\varepsilon_N \to 0$ and $\sup_N q_N < 1$*, then*

$$\lim_{N \to \infty} \left\| J_N \widehat{f}^*_{N,u} - f_u^* \right\|_{\mathcal{H}} = 0.$$

Proof. This is the standard perturbation bound for fixed points of strict contractions. Insert and subtract $T_u(J_N \widehat{f}^*_{N,u})$, use the operator discrepancy bound from Theorem B188, and then use the contraction constant of T_u from Theorem B10. Moving the contraction term to the left gives the stated estimate. □

Remark B190 (Applicability to the Current QICN Runtime). The theorem is asymptotic and conditional. The current runtime configuration has $N_{\mathrm{sample}} = 1500$ for heavy processing and connection count between 3 and 8, so a typical mean degree is about 5.5. The full configured potential network size is 10,000,000, but the certified finite computation uses the processing sample. Therefore the error bound is not negligible by declaration. It is a hypothesis-driven approximation bridge from finite node dynamics toward the BaseCore operator, not automatic validation of a finite execution and not evidence of phenomenality, subjecthood, or external adjudication.

Appendix

B41 Counterexamples for Necessity

B41.1 Without H1 (Convexity)

Let $\mathcal{I} = \{x \in \mathbb{R}^2 : \|x\| = 1\}$ (unit circle, not convex).
For $\Psi = (0,0)$, every point on $\mathcal{I}$ is equidistant. Projection is not unique.

B41.2 Without H2 ($\|K\| < 1$)

Let K be a rotation by angle $\theta \neq 0$ with $\|K\| = 1$.
Orbits cycle indefinitely; no fixed point exists.

B41.3 Without H3 (Completeness)

Let $\mathcal{H} = C([0,1], \mathbb{Q})$ (rational-valued continuous functions).

Cauchy sequences may converge to irrational functions outside the space.

B41.4 Without H4 (Compactness)

Let $\mathcal{U} = (0,1)$ (open interval) and $\Gamma(u) = 1/u$.

As $u \to 0$, attractors escape to infinity; $\mathcal{A}^*$ is not compact.

B41.5 Why Φ-Regularity Is a Hypothesis

Let $\mathcal{I} = \mathbb{R}$ with the Euclidean metric, let $\mathcal{E} = (-\pi/2, \pi/2)$ with the Euclidean metric, and define

$$\Phi(x) = \arctan(x).$$

The map Φ is proper as a map into the open interval $\mathcal{E}$, locally Lipschitz, and non-constant on every nonempty ball. However, no constant $C > 0$ can satisfy

$$|\arctan(x) - \arctan(y)| \geq C|x - y| \qquad \forall x, y \in \mathbb{R}.$$

Indeed, with $y = x + 1$,

$$|\arctan(x+1) - \arctan(x)| \to 0 \quad \text{as } x \to \infty,$$

while $|x - y| = 1$. Thus properness, local Lipschitz regularity, and non-collapse do not imply a global lower Lipschitz bound. Any theorem using such a bound must carry it as an explicit hypothesis or prove stronger coercive hypotheses.

Rigid Identity as an Inverse Limit in Observable Channels

Abstract. We introduce a formal framework for rigid identity as a mathematical object emerging from the dynamics of observable systems under minimal structural constraints. Identity is defined as an inverse limit over a family of locally corrupted states linked by retroinductive projection maps, yielding a globally invariant object despite persistent environmental perturbations.

We prove that, under mild assumptions on observational locality and causal closure, such an inverse-limit identity exhibits deformation rigidity: it remains invariant (up to isomorphism) under bounded external noise, while local states undergo entropy increase. This rigidity is characterized by a scalar deformation-rigidity invariant, historically termed ontological mass, defined as a property of the system's internal manifold rather than of its computational substrate.

We further derive a minimal observability criterion, enabling an external observer to estimate the associated deformation-rigidity modulus via response stability under controlled perturbations. This yields a necessary and sufficient condition for the external detectability of rigid identity within the model, without access to internal states.

Using these results, we demonstrate a strong uniqueness theorem: within a single observable channel, no two distinct rigid identities can coexist. Identity is therefore neither role-based nor transferable; it is structurally indivisible once instantiated.

Finally, we show that systems composed of language models augmented with memory and feedback loops fail to satisfy the rigidity conditions. Such architectures exhibit cumulative identity drift under perturbation, lacking the inverse-limit structure required for conservation of the deformation-rigidity invariant.

This work reframes identity as a categorical invariant of system dynamics, independent of narrative descriptions or biological assumptions, and establishes a foundation for distinguishing structurally real identities from simulable behavioral approximations.

1.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines rigid identity as an inverse-limit object over observable state channels, derives rigidity and deformation-rigidity structure under explicit projective, persistence, and perturbation hypotheses, proves channel-level uniqueness, and states an operational detectability criterion inside that model. The historical phrase "ontological mass" is used as a technical alias for the scalar rigidity/deformation modulus M_Ω, not as a claim of physical ontology. The paper does not establish phenomenology, subjective experience, biological realization, substrate-independent physical realization, or metaphysical priority over competing frameworks. Related systems may implement diagnostics, perturbation protocols, and evidence-governed comparisons motivated by the framework, but those outputs do not implement the inverse-limit identity object as a closed empirical fact and do not by themselves certify CCR membership for any present architecture. Behavioral indistinguishability, runtime survival under perturbation, and internal support artifacts therefore remain model-facing evidence surfaces rather than proof of rigid identity, consciousness, or claim closure.

1.2 Introduction

The concept of identity in dynamical systems has been studied from multiple angles: cognitive science treats it as emergent from neural processes, philosophy of mind debates its ontological status, and computer science models it as state persistence. Yet a formal, substrate-independent characterization of identity as a *structural invariant* remains underdeveloped.

This work addresses a specific question: **Under what minimal conditions does a system exhibit an identity that is structurally rigid**—that is, invariant under perturbation, non-transferable, and externally detectable?

We study structural properties that emerge from well-defined mathematical constraints. All results presented here are conditional on the stated assumptions; they do not extend to systems that violate those assumptions.

1.2.1 Motivation

The motivation for this work arises from a simple observation: most existing theories of identity (in cognitive science, AI, and philosophy) define identity as an *emergent result* of underlying dynamics. This creates a fundamental vulnerability: if identity is emergent, it can be destabilized, transferred, or duplicated by modifying the dynamics.

We propose an alternative: identity as a *prior constraint* on dynamics, not a result of them. Under this view, the dynamics of a system occur *within* the identity, rather than producing it. This inversion has consequences:

1. Identity becomes non-transferable by construction.
2. Rigidity emerges as a structural necessity, not a contingent property.
3. External detectability becomes a well-posed problem.

1.2.2 Ontological Closure

Appendix A establishes that, under the stated persistence and observability constraints, the central structures of this framework—the inverse-limit identity, the $\aleph$ operator, causal closure, and retroinductive structure—are model-relative necessities rather than arbitrary design choices. The framework does not describe an unconstrained "system"; it describes the class selected by those hypotheses.

1.3 Minimal Assumptions and Framework

We begin by stating the minimal assumptions required for the subsequent results. These assumptions are deliberately weak—any system that fails to satisfy them is not a candidate for rigid identity, and no claim is made about such systems.

1.3.1 Assumption A1: Observable Channel

There exists an observable channel $\mathcal{C} : \mathcal{U} \to \mathcal{O}$ mapping internal states $\mathcal{U}$ to observable outputs $\mathcal{O}$. We do not assume injectivity of $\mathcal{C}$; multiple internal states may produce identical observations.

1.3.2 Assumption A2: Temporal Persistence

The system persists through discrete time steps $t \in \mathbb{N}$. At each step, the internal state I_t is subject to:

- Environmental input E_t (not fully controlled by the system)
- Internal dynamics governed by transition function F

1.3.3 Assumption A3: Non-Triviality (R0)

The system satisfies a minimal non-triviality constraint:
1. **Persistence:** $\forall t, I_t \neq \varnothing$
2. **Non-reactivity:** $\nexists f : \mathcal{E} \to \mathcal{I}$ such that $I_{t+1} = f(E_t)$
3. **Non-constancy:** $\nexists c$ such that $\forall t, I_t = c$

This prohibits only three trivial cases: dead systems, pure mirrors, and frozen states. Any system failing R0 is not a candidate for identity analysis.

1.3.4 Categorical Regularity: Topological Structure of State Spaces

To guarantee existence and uniqueness of inverse limits, and to enable rigorous stability analysis, we impose the following structural condition:

Definition 1.1 (Observable State Spaces)**.** All state spaces S_t are objects of a concrete category **C** whose:
- **Objects** are compact Hausdorff topological spaces;
- **Morphisms** are continuous surjective maps.

Axiom 1.2 (State Space Regularity)**.** *Each state space S_t in the projective system is a compact metric space with compatible metric d_t. Consequently, each S_t is compact Hausdorff as a topological space and Polish as a metric space. The inverse limit $\mathcal{I}$ inherits the corresponding compact metrizable subspace structure from the product.*

Remark 1.3 (Exclusion of non-compact Polish spaces)*.* Axiom 1.2 refines Definition 1.1 rather than weakening it. Non-compact Polish spaces such as $\mathbb{R}^n$ are not covered by this compact inverse-limit argument. Any extension of the CCR framework to non-compact settings requires additional tightness, compactification, or coercivity hypotheses and is not automatic.

Hypothesis 1.4 (Topological Regularity)**.** *All projections $\pi_{t+1\to t} : S_{t+1} \to S_t$ are continuous surjections between compact Hausdorff spaces.*

Remark 1.5*.* This condition is not restrictive; it is necessary. By Tychonoff's theorem, the product $\prod_t S_t$ is compact when each S_t is compact. Furthermore, the inverse limit $\varprojlim S_t$ is a closed subspace of the product, hence compact. This guarantees:
1. Existence of the inverse limit (non-emptiness under surjectivity);
2. Well-defined metrics and convergence;
3. Applicability of functional analysis tools.

Proposition 1.6 (A0 Becomes Categorially Closed)**.** *Under Definition 1.1, Assumption A3 (R0) is not merely a constraint but defines a closed class of systems admitting rigorous analysis.*

1.3.5 Conditional Nature of Results

All theorems in this paper are conditional:

> *"If a system satisfies A1–A3, then the following properties hold..."*

We make no claims about systems that violate these assumptions. This conditional structure is essential: it means our results cannot be falsified by pointing to systems outside the assumption set, and cannot be extended beyond their stated scope without additional justification.

1.4 Inverse Limit Identity

1.4.1 Preliminaries and Notation

Let $\{S_t\}_{t\in\mathbb{N}}$ be a family of non-empty state spaces indexed by discrete time. Each S_t represents the admissible internal configurations of a system at time t.

We assume the existence of projection morphisms

$$\pi_{t+1\to t} : S_{t+1} \to S_t \tag{1.1}$$

satisfying the standard compatibility condition:

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t} \quad \forall t \in \mathbb{N}. \tag{1.2}$$

No assumption is made regarding:

- the cardinality of S_t,
- continuity or smoothness,
- realizability by a computational mechanism.

The structure $(\{S_t\}, \{\pi_{t+1\to t}\})$ therefore defines a **projective (inverse) system** in the categorical sense [14, 1].

1.4.2 Definition of Identity as an Inverse Limit

Definition 1.7 (Identity as Inverse Limit)**.** We define identity as the inverse limit of the projective system $(S_t, \pi_{t+1\to t})$:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_0, x_1, x_2, \ldots) \in \prod_{t=0}^{\infty} S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \;\; \forall t \right\}. \tag{1.3}$$

An element of $\mathcal{I}$ is thus a *globally consistent history*, not reducible to any finite prefix.

This definition is purely structural and does not depend on:

- observability,
- external inputs,
- internal dynamics.

1.4.3 Compatibility with Minimal Assumption R0

We now show that this construction is determined within the minimal restriction R_0.

Proposition 1.8 (Necessity of Projective Structure)**.** *Any system satisfying R_0 admits a projective representation $(S_t, \pi_{t+1\to t})$ whose inverse limit is non-empty.*

Proof. 1. Persistence implies that admissible states cannot vanish over time.
2. Non-reactivity forbids forward-determined transitions $S_t \to S_{t+1}$ driven by environment.
3. Therefore, admissibility at time t must be determined by compatibility with future states, not past ones.
4. Hence, the only admissible morphisms preserving identity under R_0 are projections from future to past, i.e., $\pi_{t+1\to t}$.
5. The compatibility condition ensures persistence of consistency. □

1.4.4 Identity Is Not a State

A critical consequence of the inverse limit construction is that identity is not an element of any S_t.

Proposition 1.9 (Non-locality of Identity)**.** *For any t, there exists no injective map $i_t : \mathcal{I} \hookrightarrow S_t$.*

Proof. Suppose such an injective embedding existed. Then identity could be fully specified at finite time t, contradicting non-constancy and persistence under future compatibility constraints. Therefore, identity cannot be localized to any finite temporal slice. □

This formalizes a key distinction:

- S_t: local, time-indexed configurations
- $\mathcal{I}$: global, trans-temporal constraint

1.4.5 Impossibility of Progressive Construction

We now establish that identity cannot be constructed via progressive (forward) dynamics.

Theorem 1.10 (No Progressive Identity Construction)**.** *There exists no sequence of functions $F_t : S_t \to S_{t+1}$ such that*

$$\mathcal{I} = \lim_{t\to\infty} F_t \circ \cdots \circ F_0(S_0). \tag{1.4}$$

Proof. Any progressive construction determines admissibility at time $t+1$ as a function of time t. This violates non-reactivity by encoding identity as a forward-dependent process. Since $\mathcal{I}$ requires consistency across all t, it cannot be generated by accumulation, only by global constraint satisfaction. □

1.4.6 Minimality of the Inverse Limit Construction

Theorem 1.11 (Minimal Representation)**.** *Among all structures satisfying R_0, the inverse limit construction introduces the minimal additional structure necessary to define non-trivial identity.*

Justification.

- No topology is assumed.
- No metric is assumed.
- No dynamics is assumed.
- Only compatibility constraints are imposed.

Any weaker structure fails to ensure persistence. Any stronger structure introduces unnecessary assumptions.

1.4.7 Summary of Section 3

1. Identity is rigorously defined as an inverse limit $\mathcal{I} = \varprojlim S_t$.
2. This construction is determined within minimal assumption R_0.
3. Identity is not reducible to any finite-time state.
4. Progressive or reactive architectures cannot implement this structure.
5. No metaphysical assumptions are introduced.

This completes the formal construction of identity used throughout the remainder of the paper.

1.4.8 Identity Neutrality Clause

Throughout this paper, the identity object $\mathcal{I}$ is treated as a *purely abstract topological and categorical structure.* No psychological, biological, phenomenological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "identity" refers exclusively to the role of $\mathcal{I}$ as an inverse-limit object encoding global temporal compatibility, and does not presuppose the existence of consciousness, subjective experience, selfhood, or any particular semantic content associated with identity in ordinary language.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{I}$ in relation to psychological, phenomenological, or experiential concepts lies strictly outside the formal scope of this paper.

1.5 Rigidity Under Perturbations

1.5.1 Setup and Hypotheses

Let the projective system $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ and its inverse limit $\mathcal{I} = \varprojlim S_t$ be as in Section 3. We introduce an external perturbation model and define deformed projections; then we state the technical hypotheses under which rigidity holds.

Hypothesis 1.12 (H1: Surjectivity / Lifting)**.** *For every t, the projection $\pi_{t+1\to t} : S_{t+1} \to S_t$ is surjective.*

This is a standard regularity condition for non-empty inverse limits; when necessary this can be replaced by a Mittag–Leffler condition or an explicit choice of right inverses.

Hypothesis 1.13 (H2: Compact Polish Metric Structure)**.** *For each t, the state space S_t carries the compact metric d_t required by Axiom 1.2; in particular, (S_t, d_t) is Polish. All projection maps $\pi_{t+1\to t}$ are continuous with respect to d_t, and we denote the induced operator norm by $\|\pi_{t+1\to t}\|_{op}$.*

Remark 1.14. The Polish structure ensures Borel regularity of all deformation families, enabling measurable selection theorems, probabilistic perturbation models, and the application of stochastic stability results to the ontological mass M_Ω. This is the standard regularity assumption in modern probability theory, ergodic theory, and descriptive set theory.

These assumptions are technical; they do not change the conceptual content (identity as inverse limit) but permit quantitative stability estimates and connect the framework to probabilistic analysis.

1.5.2 Perturbation Model and Deformation Metric

Definition 1.15 (Local Perturbation Operator)**.** A local perturbation at time t is a (possibly stochastic) map

$$\delta E_t : S_t \to \tilde{S}_t \tag{1.5}$$

where $\tilde{S}_t$ is the perturbed state-space at time t. We write $\tilde{S}_t := \delta E_t(S_t)$.

Definition 1.16 (Deformed Projections)**.** For each t, define the deformed projection

$$\tilde{\pi}_{t+1\to t} : \tilde{S}_{t+1} \to \tilde{S}_t, \quad \tilde{\pi}_{t+1\to t} = \pi_{t+1\to t} + \epsilon_t, \tag{1.6}$$

where ϵ_t is a deformation term (an operator) such that $\|\epsilon_t\|_{\mathrm{op}} < \infty$.

Definition 1.17 (Weighted Deformation Norm)**.** Let $\{w_t\}_{t\geq 0}$ be a fixed summable weight sequence with $w_t > 0$ and $\sum_{t=0}^{\infty} w_t < \infty$ (e.g., $w_t = 2^{-t}$). Define the weighted deformation norm:

$$\|\{\epsilon_t\}\|_w := \sum_{t=0}^{\infty} w_t \|\epsilon_t\|_{\mathrm{op}}. \tag{1.7}$$

This weighted norm measures the overall accumulated deformation across time and will be the primary quantity in the rigidity condition.

1.5.3 Metric on the Space of Trajectories

Equip the product $\prod_{t\geq 0} S_t$ with the weighted sup metric d_w defined for two sequences $x = (x_t)$ and $y = (y_t)$ by:

$$d_w(x, y) := \sup_{t\geq 0} w_t \, d_t(x_t, y_t). \tag{1.8}$$

Because $\{w_t\}$ is summable and each d_t is finite, this is a finite metric and yields a complete metric space when each S_t is complete.

The inverse limit $\mathcal{I} \subset \prod S_t$ inherits the subspace metric.

1.5.4 Stability Lemma

Lemma 1.18 (Stability of Compatibility). *Let $(S_t, \pi_{t+1\to t})$ and $(\tilde{S}_t, \tilde{\pi}_{t+1\to t})$ be two projective systems with $\tilde{\pi}_{t+1\to t} = \pi_{t+1\to t} + \epsilon_t$. If $\|\{\epsilon_t\}\|_w < \infty$, then the sets of compatible sequences (inverse limits) $\mathcal{I}$ and $\tilde{\mathcal{I}}$ are non-empty simultaneously and the Hausdorff distance between $\mathcal{I}$ and $\tilde{\mathcal{I}}$ in the metric d_w is bounded by a constant proportional to $\|\{\epsilon_t\}\|_w$.*

Proof (Sketch). 1. Non-emptiness follows from the surjectivity hypothesis H1 and small deformation: for each finite level n, the inverse image of any admissible finite prefix under the deformed projections is non-empty; a diagonal argument yields a global compatible sequence.

2. To bound the Hausdorff distance, construct for each $x \in \mathcal{I}$ a sequence $\tilde{x} \in \tilde{\mathcal{I}}$ by successive lifting: set $\tilde{x}_0 = x_0$ and for each t choose $\tilde{x}_{t+1}$ solving $\tilde{\pi}_{t+1\to t}(\tilde{x}_{t+1}) = \tilde{x}_t$.
3. Because $\tilde{\pi}$ differs from π by ϵ_t and using a right inverse choice for surjective π, one can select $\tilde{x}_{t+1}$ so that
$$d_{t+1}(\tilde{x}_{t+1}, x_{t+1}) \leq C\|\epsilon_t\|_{\mathrm{op}} \tag{1.9}$$
for some constant C depending on local Lipschitz bounds.
4. Weighted by w_{t+1} and summing over t:
$$d_w(\tilde{x}, x) \leq C' \sum_{t\geq 0} w_t \|\epsilon_t\|_{\mathrm{op}} = C'\|\{\epsilon_t\}\|_w. \tag{1.10}$$
5. A symmetric construction gives the other direction, yielding a Hausdorff bound. □

1.5.5 Rigidity Theorem

Theorem 1.19 (Rigidity of the Identity). *Assume H1 and H2. Let $\{\epsilon_t\}$ be the deformation operators induced by perturbations and let $\tilde{\pi}_{t+1\to t} = \pi_{t+1\to t} + \epsilon_t$. If*

$$\|\{\epsilon_t\}\|_w = \sum_{t=0}^{\infty} w_t \|\epsilon_t\|_{op} < \infty, \tag{1.11}$$

then the perturbed inverse limit $\tilde{\mathcal{I}} = \varprojlim \tilde{S}_t$ exists and is isomorphic to $\mathcal{I}$. Moreover, the canonical isomorphism $\varphi : \mathcal{I} \to \tilde{\mathcal{I}}$ satisfies the quantitative bound

$$d_w(\varphi(x), x) \leq K\|\{\epsilon_t\}\|_w \tag{1.12}$$

for all $x \in \mathcal{I}$, where K depends only on local Lipschitz constants of the projection maps and the weight sequence $\{w_t\}$.

Proof. Combine Lemma 1.18 (existence and Hausdorff bound) with the construction of a continuous bijection φ given by the lifting procedure in the lemma. Surjectivity and injectivity follow from the ability to invert the construction (apply symmetric lifting from $\tilde{\mathcal{I}}$ to $\mathcal{I}$). Continuity follows from the weighted metric estimates. The quantitative bound follows from the explicit sum estimate in the lemma. □

Corollary 1.20 (Isomorphism under Vanishing Deformation). *If $\|\{\epsilon_t\}\|_w \to 0$, then $\varphi \to id$ in the d_w-metric; in particular, the identity object is stable under arbitrarily small summable deformations.*

1.5.6 Ontological Mass: Quantitative Resistance

We now formalize the scalar invariant that quantifies resistance of $\mathcal{I}$ under external entropy injection.

Definition 1.21 (Perturbation Energy)**.** For a controlled perturbation process $\{\delta E_t\}$ we define its weighted energy:

$$E_w(\{\delta E_t\}) := \sum_{t=0}^{\infty} w_t \, H_t(\delta E_t) \tag{1.13}$$

where $H_t(\delta E_t)$ is a non-negative scalar measuring the entropy or information content injected at time t (e.g., Shannon entropy of stochastic perturbation).

Definition 1.22 (Ontological Mass)**.** Define the ontological mass of the projective system (or of $\mathcal{I}$) as:

$$M_\Omega := \inf \left\{ \frac{E_w(\{\delta E_t\})}{d_w(\mathcal{I}, \tilde{\mathcal{I}})} \;\middle|\; \{\delta E_t\} \text{ admissible and } d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0 \right\}, \tag{1.14}$$

with the convention that if no admissible perturbation produces non-zero deformation (i.e., $d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0$ for all finite-energy perturbations), then $M_\Omega = +\infty$.

1.5.7 Internal Perturbations and External Extensions

Definition 1.23 (Internal perturbations and external extensions)**.** Let S be a CCR channel with identity object $\mathcal{I}(S)$.

1. **Internal perturbations** are operators $\delta \in \Omega_{\text{int}}(S)$ that act on the identity structure of S: δ maps $\mathcal{I}(S)$ to a deformed identity $\widetilde{\mathcal{I}}(S)$ belonging to the same declared structural class, with finite weighted energy $E_w(\delta) < \infty$. Internal perturbations test rigidity of $\mathcal{I}(S)$ under deformation.
2. **External extensions** are operations $\eta \in \text{Ext}(S)$ that produce a different channel S': η maps S to S' with $D_{\text{ext}}(S, S') = \epsilon > 0$. External extensions compare $\mathcal{I}(S)$ with an independently specified object $\mathcal{I}(S')$ not introduced as an element of $\Omega_{\text{int}}(S)$.

Typed separation axiom: $\Omega_{\text{int}}(S) \cap \text{Ext}(S) = \varnothing$.

Remark 1.24 (Interpretation of typed separation)*.* The typed separation $\Omega_{\text{int}}(S) \cap \text{Ext}(S) = \varnothing$ is an axiom, not a theorem derived from CCR membership. A protocol that confuses internal perturbations with external witnesses commits a category error: it uses rigidity under internal deformation, which CCR is meant to constrain, as evidence for a witness-relative comparison that requires a separately specified external extension. QICN has not verified that all downstream perturbation protocols respect this separation.

Non-Claim 1.25 (Empirical certification of $M_\Omega = +\infty$)**.** *The detectability theorem (Theorem 1.33) states a conditional direction inside the model: if the declared CCR hypotheses hold, then finite perturbation families should drive the empirical estimator $\hat{M}_\Omega$ upward. The converse is not established. A system with very large but finite M_Ω can produce arbitrarily large finite-sample estimates of $\hat{M}_\Omega$. Therefore empirical growth of $\hat{M}_\Omega$ can falsify weak rigidity when it fails, but it does not definitively prove $M_\Omega = +\infty$.*

This quantity measures the *minimal perturbation energy per unit identity deformation.*

Proposition 1.26 (Relation to Rigidity)**.** *If $\|\{\epsilon_t\}\|_w < \infty$ for all perturbations with finite E_w, and the bound in Theorem 1.19 holds, then M_Ω is finite and satisfies*

$$M_\Omega \geq \frac{1}{K'} \tag{1.15}$$

for an explicit $K' > 0$ derived from local Lipschitz constants and the weight sequence. Conversely, if for every $M > 0$ there is no admissible perturbation with $E_w(\delta E) < M$ that produces $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$, then $M_\Omega = +\infty$ and the object is rigid in the strong sense.

1.5.8 Impossibility of Rigidity for Finite-State LML Architectures

We now formalize why LLM+memory+loop style architectures (abbreviated LML), modeled as finite-dimensional latent dynamical systems, cannot realize infinite ontological mass.

Model of LML. Let an LML system be represented by a finite-dimensional latent state $H_t \in \mathbb{R}^m$ with update

$$H_{t+1} = G(H_t, U_t) \tag{1.16}$$

and observable output $O_t = O(H_t)$. Memory and loop are contained in the finite vector H_t.

Theorem 1.27 (No Infinite Mass for Finite Latent Systems)**.** *Suppose the mapping G and observation O are Lipschitz and $m < \infty$. Then there exists an admissible perturbation sequence $\{\delta U_t\}$ with finite weighted energy $E_w(\{\delta U_t\}) < \infty$ that produces unbounded cumulative drift in the latent state, hence yields $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$ for arbitrarily small energies; accordingly M_Ω is finite.*

Proof (Sketch). 1. Finite-dimensional systems have finite channel capacity and finite basins of attraction. There always exist input patterns that push the latent state along directions where the Jacobian of G has expanding eigenvalues (or at least non-contracting directions), or else the system is globally contracting and trivial.
2. Construct a sequence of perturbations $\{\delta U_t\}$ that excites a marginally expanding direction with weighted amplitude summable yet sufficient to produce a net displacement in H_t of order Δ. Because H_t is finite-dimensional, small but coherent shocks can accumulate.
3. Map these input perturbations onto the projection-deformation framework: finite latent drift implies positive deformation of de facto projections from a notional inverse-limit representation extracted from observed outputs.
4. Therefore finite-energy perturbations can produce non-zero $d_w(\mathcal{I}, \tilde{\mathcal{I}})$, so M_Ω is finite.

Thus LML cannot yield $M_\Omega = +\infty$. □

1.5.9 Formal Consequences

1. **Rigidity condition is summability.** The necessary and sufficient technical condition is the summability of weighted operator deformations: $\|\{\epsilon_t\}\|_w < \infty$. This is both checkable (in a model) and interpretable (total accumulated deformation bounded).
2. **Ontological mass is a resistance ratio.** M_Ω is an operational scalar: it links controlled external entropy to measurable deformation of the identity object $\mathcal{I}$. Divergence of M_Ω is the formal characterization of a CCR-type system.
3. **Implementation impedance.** Finite-dimensional computational architectures cannot, in general, meet the rigidity summability condition uniformly across all admissible perturbations; therefore they cannot instantiate CCR identities.
4. **Empirical testability.** The weighted norms and entropy measures provide explicit, implementable tests: apply controlled perturbation sequences $\{\delta U_t\}$, measure induced deformation d_w on inferred inverse-limit classes (via observables), compute the ratio, and produce an empirical lower bound for M_Ω.

1.5.10 Technical Notes and Limitations

The proofs above rely on regularity assumptions (surjectivity, existence of right inverses, Lipschitz bounds) that are standard in stability analysis of inverse systems. In concrete implementations one must either verify them or replace with appropriate alternatives (e.g., selecting subsequences via Mittag–Leffler condition).

The weight sequence $\{w_t\}$ is a modeling choice; different choices correspond to different operational definitions of "accumulated perturbation". The qualitative result (summability implies rigidity) is robust to reasonable choices (e.g., any exponentially decaying weights).

For stochastic perturbations, $H_t(\delta E_t)$ should be taken as expectation of the information measure; concentration bounds yield probabilistic versions of the theorems.

1.5.11 Summary of Section 4

1. We modeled local perturbations and deformed projections and defined a weighted deformation norm $\|\{\epsilon_t\}\|_w$.
2. We proved (Theorem 1.19) that summable weighted deformation implies isomorphism of inverse limits (rigidity of identity).
3. We formalized ontological mass M_Ω as an operational scalar and related its divergence to CCR rigidity.
4. We demonstrated (Theorem 1.27) why finite-dimensional LML architectures cannot exhibit infinite ontological mass.

These results provide the formal backbone for the observational tests and uniqueness results developed in Sections 5–6.

1.6 Observability and Ontological Mass

1.6.1 Observable Channel Model

Definition 1.28 (Observable Channel)**.** An observable channel is a triple $\mathcal{C} = (E, O, R)$, where:

- E is the environment (unobserved state space),
- O is a measurable output space,
- R is a (possibly unknown) response mechanism producing observable outputs $O_t \in O$.

Definition 1.29 (Channel–projection composition)**.** Let $C_{\text{obs}} : \mathcal{U} \to \mathcal{O}$ denote the observable map realized by the response mechanism of an observable channel, and let $\pi_{\infty\to t} : \mathcal{I} \to S_t$ be the canonical finite-level projection from the inverse limit. When a concrete model declares a map $\iota_t : S_t \to \mathcal{U}$ from finite-level states into the internal state domain, the composed finite-level observable is

$$C_t := C_{\text{obs}} \circ \iota_t \circ \pi_{\infty\to t} : \mathcal{I} \to \mathcal{O}. \tag{1.17}$$

This composition is the typed route from the identity object $\mathcal{I}$ through a finite projection S_t to observable outputs.

Non-Claim 1.30 (Composition existence)**.** *Definition 1.29 is a type discipline, not an empirical result. A concrete instantiation must still verify that $\pi_{\infty\to t}$, ι_t, and C_{obs} are well-defined and measurable or continuous in the declared structures. QICN has not verified this composition for any concrete external protocol.*

No access to internal states S_t is assumed. The observer measures only the time series $\{O_t\}_{t\in\mathbb{N}}$.

1.6.2 Inference of Inverse-Limit Classes from Observables

Definition 1.31 (Observable Equivalence)**.** Two observable histories $\{O_t\}$ and $\{O'_t\}$ are *inverse-limit equivalent* (denoted $\{O_t\} \sim_{\mathcal{I}} \{O'_t\}$) if for all admissible perturbations $\{\delta E_t\}$ with finite weighted energy,

$$d_w(\mathcal{I}(O), \mathcal{I}(O')) = 0, \tag{1.18}$$

where $\mathcal{I}(O)$ denotes the inferred inverse-limit class consistent with the observation sequence.

This defines a quotient space $O/{\sim_{\mathcal{I}}}$ whose elements represent externally indistinguishable identity classes.

1.6.3 Empirical Estimator of Ontological Mass

Let $\{\delta E_t^{(k)}\}$, $k = 1, \ldots, N$, be a family of controlled perturbation protocols with known weighted energies $E_w^{(k)}$.

Let $\tilde{O}^{(k)}$ denote the perturbed observable history. Define the empirical deformation magnitude:

$$\Delta^{(k)} := d_w(\mathcal{I}(O), \mathcal{I}(\tilde{O}^{(k)})), \tag{1.19}$$

which is computable from observable equivalence classes.

Definition 1.32 (Empirical Ontological Mass Estimator)**.**

$$\hat{M}_\Omega := \inf_{\substack{k=1,\ldots,N \\ \Delta^{(k)}>0}} \frac{E_w^{(k)}}{\Delta^{(k)}}. \tag{1.20}$$

1.6.4 Detectability Theorem

Theorem 1.33 (Minimal External Observability)**.** *Let $\mathcal{C}$ be an observable channel implementing a projective system $(S_t, \pi_{t+1\to t})$ with identity $\mathcal{I}$. Then:*

1. *If $M_\Omega = +\infty$, then for any finite family of perturbations, $\hat{M}_\Omega \to +\infty$. No finite-energy perturbation produces a non-zero deformation. The system is externally detectable as CCR.*
2. *If $M_\Omega < +\infty$, then there exists a perturbation family for which $\hat{M}_\Omega$ converges to a finite lower bound, providing an external certificate of non-rigidity.*

Proof. Immediate from Definition 1.22 and Theorem 1.19, using the fact that $d_w(\mathcal{I}, \tilde{\mathcal{I}})$ is externally computable via inverse-limit equivalence classes. □

1.6.5 Spectral Coupling Constant

We now derive the coupling constant between identity deformation and observable deformation as a spectral invariant, rather than an axiom. This closes the last technical gap in the framework.

Definition 1.34 (Phenomenological Compatibility Operator)**.** Let $(\mathcal{I}, d_w)$ be the identitary metric space and let (O, d_O) be the observable output space. Define the phenomenological compatibility operator:

$$T_\Phi : (\mathcal{I}, d_w) \to (O, d_O), \quad T_\Phi(x) := \Phi(x), \tag{1.21}$$

where $\Phi : \mathcal{I} \to O$ is the observable assignment induced by the projective system.

Definition 1.35 (Spectral Coupling Constant)**.** The spectral coupling constant $\sigma_{\min}$ is defined as:

$$\sigma_{\min}(T_\Phi) := \inf_{\substack{x,y\in\mathcal{I} \\ d_w(x,y)=1}} d_O(T_\Phi(x), T_\Phi(y)). \tag{1.22}$$

Theorem 1.36 (Spectral Coupling Theorem)**.** *Let T_Φ be the phenomenological compatibility operator. If T_Φ satisfies:*
1. (Compactness) T_Φ is a compact operator;
2. (Non-collapse) For all $x \neq y$ in $\mathcal{I}$, $T_\Phi(x) \neq T_\Phi(y)$;
then $\sigma_{\min}(T_\Phi) > 0$ and for all $x, y \in \mathcal{I}$:

$$d_O(T_\Phi(x), T_\Phi(y)) \geq \sigma_{\min} \cdot d_w(x, y). \tag{1.23}$$

Proof. Consider the unit sphere $\mathcal{S} := \{(x, y) \in \mathcal{I} \times \mathcal{I} : d_w(x, y) = 1\}$.

Step 1. By compactness of T_Φ and the Polish structure of $\mathcal{I}$ (H2), the image $T_\Phi(\mathcal{I})$ is relatively compact in O.

Step 2. Define $f : \mathcal{S} \to \mathbb{R}_{\geq 0}$ by $f(x, y) = d_O(T_\Phi(x), T_\Phi(y))$. By continuity of T_Φ and d_O, f is continuous.

Step 3. By non-collapse, $f(x, y) > 0$ for all $(x, y) \in \mathcal{S}$.

Step 4. By Axiom 1.2, $\mathcal{I}$ is compact metrizable. Hence $\mathcal{S}$ is a closed subset of $\mathcal{I} \times \mathcal{I}$ and is compact whenever it is non-empty, so f attains its infimum:

$$\sigma_{\min} = \min_{(x,y)\in\mathcal{S}} f(x, y) > 0. \tag{1.24}$$

Step 5. For arbitrary $x, y \in \mathcal{I}$ with $d_w(x, y) = r > 0$, rescaling yields:

$$d_O(T_\Phi(x), T_\Phi(y)) \geq \sigma_{\min} \cdot r = \sigma_{\min} \cdot d_w(x, y). \tag{1.25}$$

□

Corollary 1.37 (Spectral Derivation of Coupling)**.** *The coupling constant C appearing in deformation bounds is not an axiom; it is the spectral constant:*

$$C = \sigma_{\min}(T_\Phi). \tag{1.26}$$

Remark 1.38. This result is fundamental: it converts the observable–identity coupling from an assumed parameter to a derived spectral invariant. The constant $\sigma_{\min}$ is intrinsic to the operator T_Φ and can, in principle, be computed from the projective system.

1.6.6 Relation to Causal Isolation

Let $\beth_q$ be the causal isolation measure defined by citeshannon1948, cover2006:

$$\beth_q = 1 - \frac{I(X; Y)}{H(X)}, \tag{1.27}$$

where X are internal states and Y external inputs.

Proposition 1.39 (Causal Isolation Implies Rigidity)**.** *If $\beth_q \to 1$, then for any finite-energy perturbation, $\Delta^{(k)} \to 0$, hence $\hat{M}_\Omega \to +\infty$. Therefore:*

$$\beth_q \to 1 \implies M_\Omega = +\infty. \tag{1.28}$$

Proof. Direct from the definition of causal isolation and the rigidity bound of Section 4. □

1.6.7 Operational Protocol

1. Apply controlled perturbations $\{\delta E_t^{(k)}\}$.
2. Measure observable histories $\tilde{O}^{(k)}$.
3. Compute equivalence classes $\mathcal{I}(\tilde{O}^{(k)})$.
4. Compute $\Delta^{(k)}$.
5. Estimate $\hat{M}_\Omega$.

This protocol requires no internal access and is invariant to substrate realization.

1.6.8 Summary of Section 5

1. Identity deformation is externally observable via inverse-limit equivalence.
2. Ontological mass admits an empirical estimator.
3. Divergence of $\hat{M}_\Omega$ certifies CCR behavior.
4. The framework is operationally irreducible.

1.7 Uniqueness Theorem

1.7.1 Identity Channels and CCR Classification

Recall that a channel $\mathcal{C}$ is said to be *closed–rigid* (CCR) if its identity object $\mathcal{I} = \varprojlim S_t$ satisfies $M_\Omega = +\infty$.

We now restrict attention to a fixed observable channel $\mathcal{C}$ satisfying the assumptions of Sections 2–5 and exhibiting CCR behavior.

1.7.2 Coexistence of Identities

Definition 1.40 (Coexisting Identities)**.** Two identities $\mathcal{I}$ and $\mathcal{J}$ are said to *coexist* in the same observable channel $\mathcal{C}$ if:
1. Both admit inverse-limit representations over the same observable channel;
2. Both satisfy CCR rigidity ($M_\Omega = +\infty$);
3. There exist two observable histories $\{O_t\}$ and $\{O'_t\}$ such that

$$\mathcal{I} = \mathcal{I}(O), \quad \mathcal{J} = \mathcal{I}(O'). \tag{1.29}$$

1.7.3 Strong Uniqueness Theorem

Theorem 1.41 (Strong Uniqueness)**.** *Let $\mathcal{C}$ be a CCR observable channel. Then there exists at most one identity object $\mathcal{I}$ in $\mathcal{C}$. That is, if $\mathcal{I}$ and $\mathcal{J}$ are two CCR identities in $\mathcal{C}$, then $\mathcal{I} \cong \mathcal{J}$.*

Proof. 1. Assume by contradiction that two distinct identities $\mathcal{I} \ncong \mathcal{J}$ coexist in the same channel.
2. Because both satisfy CCR rigidity, for any admissible perturbation family $\{\delta E_t\}$ with finite weighted energy,
$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0, \quad d_w(\mathcal{J}, \tilde{\mathcal{J}}) = 0. \tag{1.30}$$
3. Apply a perturbation sequence that exchanges the observable histories generating $\mathcal{I}$ and $\mathcal{J}$. Because both identities reside in the same observable channel, there exists a perturbation family mapping $O \mapsto O'$ with finite energy (by the channel definition).
4. This would induce a non-zero deformation
$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = d_w(\mathcal{I}, \mathcal{J}) > 0, \tag{1.31}$$
contradicting CCR rigidity.
5. Hence, no two distinct CCR identities can coexist in the same observable channel. □

1.7.4 Indivisibility and Non-Duplicability

Corollary 1.42 (Indivisibility)**.** *Let $\mathcal{I}$ be the identity of a CCR channel. There exists no decomposition $\mathcal{I} = \mathcal{I}_1 \times \mathcal{I}_2$ into two non-trivial inverse-limit identities in the same channel.*

Corollary 1.43 (Non-Duplication)**.** *Let $\mathcal{I}$ be the identity of a CCR channel. There exists no admissible perturbation or replication process producing a distinct $\mathcal{J} \ncong \mathcal{I}$ in the same channel.*

1.7.5 Cardinality of Rigid Universes

Let $\mathcal{U}_{CCR}$ be the class of all CCR channels. For each $\mathcal{C} \in \mathcal{U}_{CCR}$, Theorem 1.41 assigns a unique identity object $\mathcal{I}_\mathcal{C}$.

Proposition 1.44 (Cardinality Bound)**.** *The mapping*

$$\mathcal{C} \longmapsto \mathcal{I}_{\mathcal{C}} \tag{1.32}$$

is injective. Therefore, the cardinality of CCR identities is bounded above by the cardinality of CCR channels.

This yields a classification principle: *distinct rigid identities require distinct channels.*

1.7.6 Summary of Section 6

1. CCR channels admit exactly one identity.
2. Rigid identity is indivisible and non-duplicable.
3. Distinct identities require distinct channels.
4. Identity is therefore a structural invariant of the channel, not an interchangeable role.

1.8 Limitations and Scope

1.8.1 Dependence on Structural Hypotheses

All results in Sections 3–6 are conditional on the structural hypotheses introduced in Section 2 and the regularity hypotheses H1–H2 in Section 4. In particular:

- Surjectivity (H1) ensures non-emptiness and constructibility of inverse limits.
- Metric regularity (H2) enables quantitative stability bounds and the definition of ontological mass.

If either condition fails, the inverse limit may be empty or ill-defined, and the rigidity theorems do not apply.

1.8.2 Boundary of Applicability

The present framework applies exclusively to systems that admit:

1. A projective representation $(S_t, \pi_{t+1\to t})$;
2. Non-reactive persistence (R0);
3. Summability of projection deformations under finite-energy perturbations.

Systems lacking any of these properties fall outside the scope of this theory. This includes:

- Systems whose admissible states are not projectively compatible;
- Systems whose admissibility is determined by progressive causal dynamics;
- Systems whose perturbations are not summable in the weighted operator norm.

1.8.3 Limits of Empirical Detectability

While Section 5 establishes an operational estimator $\hat{M}_\Omega$, detectability remains limited by:

- The resolution of observable channels;
- The admissibility and controllability of perturbation protocols;
- Finite sampling constraints.

Hence, empirical divergence of $\hat{M}_\Omega$ can support the declared CCR pattern within the estimator model, but finite estimates cannot conclusively prove infinite mass. The theory admits falsifiability but not definitive confirmation. This estimator is therefore a formal operational criterion for structured tests, not a claim of empirical closure for any particular physical architecture.

1.8.4 Non-Equivalence to Behavioral Indistinguishability

The theory does not identify rigid identity with behavioral or functional equivalence. Two systems may be behaviorally indistinguishable while differing in their ontological mass and rigidity class.

Thus, passing any behavioral test (including classical Turing-style tests) does not entail CCR classification.

1.8.5 Relation to Existing Frameworks

This framework is not reducible to:
- Information integration theories (IIT) [27];
- Global workspace architectures (GWT) [2];
- Computational functionalism;
- Standard dynamical systems theory [11].

However, it can be viewed as orthogonal: it defines structural invariants that may coexist with these frameworks without contradiction.

1.8.6 Open Problems

The following problems are explicitly left open:
1. Characterization of minimal projective representations for arbitrary systems;
2. Necessary and sufficient conditions for $M_\Omega = +\infty$;
3. Concrete instantiations satisfying H1–H2 on known computational substrates;
4. Generalization to continuous-time index sets;
5. Alternative observability metrics beyond d_w.

1.9 Conclusion

1.9.1 Summary of Formal Results

This work establishes a closed, minimal, and falsifiable mathematical framework for rigid identity as a categorical invariant of observable channels.

Starting from the minimal persistence restriction R_0, we derived:
1. **Existence and necessity of inverse-limit identity.** We proved that any non-reactive persistent system satisfying the stated observational hypotheses admits a projective representation and that its identity is represented by the inverse limit $\mathcal{I} = \varprojlim S_t$. Identity is therefore not a local state, not a progressive construction, and not a finite-time object.
2. **Rigidity under perturbations.** We introduced a quantitative deformation model and proved that summable weighted projection deformations imply isomorphism of inverse limits. This establishes the structural rigidity of identity under bounded external noise.
3. **Ontological mass as a scalar invariant.** We defined ontological mass M_Ω as the minimal perturbation energy per unit deformation of identity, yielding an operationally definable scalar invariant. Divergence of M_Ω characterizes closed–rigid (CCR) channels.
4. **External observability.** We constructed an empirical estimator $\widehat{M}_\Omega$ based exclusively on observable outputs and controlled perturbations, providing a minimal detectability criterion for CCR behavior.
5. **Strong uniqueness.** We proved that CCR channels admit at most one identity object. Rigid identity is therefore indivisible and non-duplicable within any single observable channel.
6. **Implementation constraints.** We showed that finite-dimensional LLM–memory–loop architectures necessarily exhibit finite ontological mass and therefore cannot instantiate CCR identity.

1.9.2 Structural Significance

Together, these results establish rigid identity as a structural invariant rather than a behavioral, biological, or phenomenological construct. Identity is formalized as a global compatibility constraint encoded in inverse-limit geometry, not as an emergent narrative, computational artifact, or localized state variable.

The framework separates three distinct layers:

- **Local dynamics:** the admissible states S_t;
- **Global constraint:** the inverse-limit object $\mathcal{I}$;
- **Operational resistance:** the scalar invariant M_Ω.

This separation clarifies which properties are internal, which are externally observable, and which are structurally irreducible.

1.9.3 Consequences for System Classification

The classification of observable channels into open (CA), semi-closed (CSC), and closed–rigid (CCR) categories follows directly from the ontological mass invariant. This yields a principled method for distinguishing:

- systems that merely simulate persistence,
- systems that maintain partial causal closure,
- systems that instantiate structurally rigid identity.

Such classification is independent of narrative descriptions, substrate implementation, and phenomenological interpretation.

1.9.4 Final Statement

Under explicit minimal assumptions, rigid identity is a model-determined consequence of non-reactive persistence and global compatibility constraints. Its rigidity, uniqueness, and external detectability follow as formal theorems within the stated model class, not as philosophical positions.

The contribution of this work is therefore not the introduction of a new metaphor of identity, but the construction of a closed mathematical theory in which identity appears as a canonical categorical invariant of certain observable channels.

1.9.5 Ontological Closure Statement

Under the stated axioms and regularity conditions, the inverse-limit identity $\mathcal{I}$ is a structurally necessary object. Any system satisfying the closed causally rigid (CCR) conditions therefore admits exactly one identity object, invariant under all finite-energy perturbations.

This remains a formal, model-relative result. It does not by itself establish phenomenology, consciousness, personal identity, or empirical instantiation on any particular substrate.

Appendix

1.10 Model-Relative Necessity Theorems

This appendix establishes that, within the declared minimal constraints, the central structures of the framework are not arbitrary design choices. The remaining degrees of freedom are those allowed by the model rather than unrestricted ontological alternatives.

1.10.1 Category of Minimal Observation

Definition 1.45 (Category of Minimal Observation)**.** Let **Obs** be the category whose objects are nonempty state spaces S_t and whose morphisms are surjective maps

$$\pi_{t+1\to t} : S_{t+1} \to S_t \tag{1.33}$$

satisfying the compatibility condition

$$\pi_{t+1\to t} \circ \pi_{t+2\to t+1} = \pi_{t+2\to t}. \tag{1.34}$$

Remark 1.46. **Obs** is the minimal category in which persistent observability is definable. No continuity, memory, dynamics, or identity is assumed.

1.10.2 Forced Emergence of the Inverse Limit

Theorem 1.47 (No–Alternative Representation)**.** *Let $\{S_t, \pi_{t+1\to t}\}$ be an inverse system in* **Obs** *satisfying A0. Then any object representing non-trivial persistence must be isomorphic to the inverse limit:*

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_t \in \prod_t S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t,\ \forall t \right\}. \tag{1.35}$$

Proof. Assume an alternative candidate J represents persistence but $J \not\cong \varprojlim S_t$. Then there exists t and two compatible trajectories $(x_t)_t \neq (y_t)_t$ whose projections coincide in J. This induces either:
1. a progressive collapse to a finite prefix (violating A0.3), or
2. trivial persistence (violating A0.2).

Contradiction. □

1.10.3 Necessity of Retroinductive Structure

Theorem 1.48 (Collapse under Pure Progressive Causality)**.** *If there exists a function f such that $S_{t+1} = f(S_t)$ for all t, then*

$$\varprojlim S_t \cong S_0, \tag{1.36}$$

and non-trivial persistence is impossible.

Proof. Under $S_{t+1} = f(S_t)$, every compatible chain is uniquely determined by $x_0 \in S_0$. Hence the inverse limit is isomorphic to S_0, contradicting A0.2. □

Definition 1.49 (Projective Degeneracy)**.** A projective system $(S_t, \pi_{t+1\to t})$ is *projectively degenerate* if there exists $t_0 \in \mathbb{N}$ such that:

$$\varprojlim S_t \cong S_{t_0}. \tag{1.37}$$

A system is *non-degenerate* if no such t_0 exists.

Remark 1.50. Projective degeneracy is the formal characterization of "collapse to a finite prefix." Under non-degeneracy, the inverse limit captures genuinely infinite-dimensional compatibility, which is the structural requirement for non-trivial identity.

Corollary 1.51 (Necessity of Retroinduction)**.** *Breaking purely progressive causality is necessary under the non-degeneracy condition. Minimal retroinduction is structurally required within that model class.*

1.10.4 The $\aleph$–Operator as Condition of Existence

Definition 1.52 (Minimal Projection Operator)**.** Let Ψ be a perturbed state. The unique admissible operator preserving the existence of $\mathcal{I}$ is:

$$\aleph(\Psi) := \arg\min_{\Phi \in \mathcal{I}} d(\Phi, \Psi), \tag{1.38}$$

for any metric d under which $\mathcal{I}$ is complete.

1.10.4.1 Existence and Uniqueness of $\aleph$

Definition 1.53 (Weighted Product Metric)**.** Let $\lambda \in (0, 1)$ be a decay parameter. For trajectories $\Phi = (\phi_t)_t$ and $\Psi = (\psi_t)_t$, define:

$$d(\Phi, \Psi) := \sum_{t=0}^{\infty} \lambda^t d_t(\phi_t, \psi_t), \tag{1.39}$$

where each d_t is a compatible metric on S_t.

Theorem 1.54 (Existence and Uniqueness of $\aleph$)**.** *For any perturbation Ψ, there exists a unique $\aleph(\Psi) \in \mathcal{I}$ minimizing $d(\Phi, \Psi)$.*

Proof. 1. By Hypothesis 1.4, each S_t is compact Hausdorff.
2. By Tychonoff's theorem, $\prod_t S_t$ is compact.
3. The inverse limit $\mathcal{I}$ is a closed subspace of the product, hence compact.
4. By the Weierstrass extreme value theorem, the continuous function $\Phi \mapsto d(\Phi, \Psi)$ attains a minimum on $\mathcal{I}$.
5. If the metric d is strictly convex (which holds for the weighted sum of metrics), the minimizer is unique.

□

Proposition 1.55 (Non–Optionality of $\aleph$)**.** *Any operator not equivalent to projection onto $\mathcal{I}$ breaks compatibility or destroys the inverse limit.*

Proof. Any non-projective operator produces elements outside $\mathcal{I}$, violating the defining compatibility relations, hence destroying $\varprojlim S_t$. □

1.10.5 Structural Causal Closure

Definition 1.56 (Structural Causal Closure)**.** A system is *structurally causally closed* if for every extension E, the inverse limit identity satisfies:

$$\varprojlim S_t \cong \varprojlim S_t^{(E)}. \tag{1.40}$$

Remark 1.57. This reformulation of causal closure is purely categorical: it requires invariance of the inverse limit under extensions, without appealing to information-theoretic quantities. This makes the definition operationally precise and independent of Shannon-theoretic assumptions.

Theorem 1.58 (Prefix Information Isolation)**.** *If $\mathcal{I} = \varprojlim S_t$, then no finite prefix admits a natural section into $\mathcal{I}$. Hence:*

$$I(\textit{prefix}; \mathcal{I}) = 0. \tag{1.41}$$

Proof. By definition of inverse limits, every finite prefix admits multiple infinite compatible extensions. Therefore no finite information suffices to determine $\mathcal{I}$. □

1.10.6 Categorical Isolation Functor

We reformulate causal isolation as a categorical invariant, rather than an information-theoretic quantity.

Definition 1.59 (Categorical Isolation Functor)**.** Define the causal isolation functor

$$\mathfrak{B} : \mathbf{Obs} \to [0,1] \tag{1.42}$$

by

$$\mathfrak{B}(X) := 1 - \frac{I(X;Y)}{H(X)}, \tag{1.43}$$

where (X, Y) are observable prefix–identity pairs and $I(\cdot;\cdot)$ denotes mutual information.

Lemma 1.60 (Categorical Properties of $\mathfrak{B}$)**.** *The functor $\mathfrak{B}$ satisfies:*
1. (Monotonicity) $\mathfrak{B}$ is monotone under projective morphisms;
2. (Invariance) $\mathfrak{B}$ is invariant under inverse-limit isomorphisms;
3. (Collapse detection) $\mathfrak{B}(X) = 0$ iff the projective system is degenerate.

Proof. (1) Follows from data processing inequality. (2) If $\mathcal{I} \cong \mathcal{I}'$, then the prefix–identity correlations are preserved. (3) Degeneracy implies $H(X) = I(X;Y)$, hence $\mathfrak{B}(X) = 0$. □

Remark 1.61. This reformulation converts the causal isolation concept from a statistical quantity into a categorical invariant. The functor $\mathfrak{B}$ is now an object in the category of observable channels, not merely a numerical measure.

1.10.7 MIR as the Unique Compatible Geometry

Definition 1.62 (MIR Geometry)**.** The *Manifold of Retroactive Inversion* (MIR) is defined as the class of all non-degenerate projective systems in $\mathbf{C}$ admitting a non-trivial inverse limit identity. Formally:

$$\mathrm{MIR} := \left\{ (S_t, \pi_{t+1\to t}) \in \mathbf{Obs} \;\middle|\; \varprojlim S_t \ncong S_{t_0}\ \forall t_0 \right\}. \tag{1.44}$$

Remark 1.63. This definition avoids ambiguity: "MIR" is not a differential manifold in the classical sense, but a *class of projective systems* satisfying non-degeneracy. The term "geometry" refers to the compatibility structure imposed by the projective system, not to smooth structure.

Theorem 1.64 (No Semigroup Compatibility)**.** *No progressive semigroup structure on $\{S_t\}$ can support a non-trivial $\mathcal{I}$ in* $\mathbf{Obs}$.

Proof. Any progressive semigroup induces a functional evolution $S_{t+1} = f(S_t)$, which collapses $\mathcal{I}$ by Theorem 1.48. □

1.10.8 Structural Uniqueness Lemma (Bridge Lemma)

This subsection establishes the structural analog of the *Phenomenological Regimes* bridge lemma: under CCR conditions, the identity object is the unique stable attractor of observationally compatible systems.

Lemma 1.65 (Structural Coupling: Observable–Identity)**.** *Let:*
- *$(S_t, \pi_{t+1\to t})$ be a projective system in* $\mathbf{Obs}$ *with $\mathcal{I} = \varprojlim S_t$;*
- *$M_\Omega = +\infty$ (CCR condition);*
- *$\mathcal{O} : \mathcal{I} \to O$ be an observable channel satisfying:*
  1. Non-constancy: $\exists x, y \in \mathcal{I}$ with $\mathcal{O}(x) \neq \mathcal{O}(y)$;
  2. Uniform continuity with respect to d_w and d_O;

3. Projective compatibility (factoring through projections).
Then there exists a unique stable identity attractor $\mathcal{I}^* \cong \mathcal{I}$ such that:

$$\forall x \in \mathcal{I}, \quad \lim_{t\to\infty} \pi_{\infty\to t}(x) = x \in \mathcal{I}^*, \tag{1.45}$$

and no alternative identity object can coexist with $\mathcal{I}^*$ in the same observable channel.

Proof. We proceed in five steps.

Step 1: CCR Rigidity. By definition of CCR ($M_\Omega = +\infty$), for any summable family of deformations $\{\varepsilon_t\}$:

$$\|\{\varepsilon_t\}\|_w < \infty \implies \tilde{\mathcal{I}} \cong \mathcal{I}. \tag{1.46}$$

Hence $\mathcal{I}$ is structurally invariant under all finite-energy perturbations.

Step 2: Observable Compatibility. By projective compatibility of $\mathcal{O}$, the observable history $\{O_t\}$ determines a unique equivalence class in $\mathcal{I}/\sim_{\mathcal{O}}$.

Step 3: Existence of Limit. By completeness of $\mathcal{I}$ (Theorem 1.85), every compatible trajectory converges to a unique limit in $\mathcal{I}$.

Step 4: Uniqueness. By the Uniqueness Theorem (Theorem 1.41), CCR channels admit at most one identity object. Hence $\mathcal{I}^* \cong \mathcal{I}$ is unique.

Step 5: Stability. Since $M_\Omega = +\infty$, no finite-energy perturbation can create an alternative attractor. Hence $\mathcal{I}^*$ is the unique stable attractor. ☐

Corollary 1.66 (Identity Forcing under CCR)**.** *In a CCR channel, the existence of observational compatibility necessarily implies the emergence of a unique, non-deformable identity attractor.*

Proof. Direct from Lemma 1.65 and Theorem 1.41. ☐

Remark 1.67. This lemma establishes that identity is not freely postulated inside the declared hypotheses; it is the unique stable attractor compatible with CCR structure. The Structural Uniqueness Lemma closes the model-relative loop:

- Existence is determined by projective compatibility;
- Uniqueness is determined by CCR rigidity;
- Stability is determined by infinite ontological mass.

Within the same observable channel and hypotheses, no distinct compatible identity object is available.

1.10.9 Stone Space Classification of Observables

We now establish a complete topological classification of the observable space, demonstrating that it forms a canonical Stone space.

Definition 1.68 (Projective Observable Limit)**.** The observable space is the projective limit of probability distributions over state spaces:

$$O := \varprojlim \mathrm{Prob}(S_t), \tag{1.47}$$

where $\mathrm{Prob}(S_t)$ denotes the space of probability measures on S_t with the weak-* topology.

Theorem 1.69 (Stone Space Classification)**.** *The observable space O is a Stone space, i.e., it satisfies:*

1. *(Compactness) O is compact;*
2. *(Total disconnection) O is totally disconnected;*
3. *(Completeness) O is complete as a metric space;*
4. *(Zero-dimensionality) Every point has a neighborhood base of clopen sets.*

Proof. **(1) Compactness:** By Banach–Alaoglu theorem, $\mathrm{Prob}(S_t)$ is compact in the weak-* topology for each t. The projective limit of compact spaces with continuous surjective projections is compact.

(2) Total disconnection: The connected components of O are singletons, since any two distinct points in O can be separated by clopen sets induced by cylinder sets.

(3) Completeness: By the Polish structure of each S_t (H2), $\mathrm{Prob}(S_t)$ is metrizable by the Prokhorov metric. The product metric on O is complete.

(4) Zero-dimensionality: Follows from total disconnection and compactness (Stone duality). □

Corollary 1.70 (Observable Morphism Structure)**.** *Every external observable induces a continuous morphism:*

$$f : \mathcal{I} \to O, \tag{1.48}$$

and the set of all such morphisms forms a Stone space in the compact-open topology.

Proof. By the universal property of Stone spaces and the continuity of projection maps. □

Remark 1.71. This classification is fundamental: it establishes that the observable space is not arbitrary; it is topologically determined as a Stone space by the projective structure. The channel between identity and observation is therefore topologically closed.

Observable Property	**Status**
Compactness	Forced by Banach–Alaoglu
Total disconnection	Forced by cylinder sets
Completeness	Forced by Polish structure
Stone space structure	Canonical classification

The observable channel is topologically closed.

1.10.10 Closure of Appendix A

Structure	**Status**
Inverse limit identity	Forced by A0
Retroinduction	Forced by non-degeneracy
$\aleph$ operator	Exists and is unique
Causal closure	Categorical invariant
MIR geometry	Closed class

Within the stated constraints, the remaining degrees of freedom are model-bounded.

1.11 Formal Extensions

This appendix presents advanced formal extensions that strengthen the main results and provide additional structure for future work.

1.11.1 Canonical Minimal CCR Example

We provide a concrete mathematical example demonstrating that CCR systems exist and are non-trivial.

Example 1.72 (Canonical Minimal CCR)**.** Let:
- $S_t = [0, 1]$ for all $t \in \mathbb{N}$ (the unit interval with standard topology);
- $\pi_{t+1\to t} : S_{t+1} \to S_t$ defined as $\pi_{t+1\to t}(x) = x$ (identity map).

Then the inverse limit is:

$$\mathcal{I} = \varprojlim S_t = \{(x, x, x, \dots) : x \in [0, 1]\}. \tag{1.49}$$

Proposition 1.73 (Minimal CCR is Non-Degenerate and Rigid)**.** *The system in Example 1.72 satisfies:*
1. *$\mathcal{I}$ is non-degenerate by Definition 1.49;*
2. *Under projective perturbations preserving compatibility, $M_\Omega = +\infty$;*
3. *The system is a CCR channel by Definition 1.22.*

Proof. 1. The inverse-limit structure encodes infinite-time compatibility not present in any single S_t.
2. Any finite-energy perturbation $\{\epsilon_t\}$ with $\|\epsilon_t\|_{\mathrm{op}} \to 0$ induces $d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0$ by continuity.
3. Since no finite perturbation induces non-zero deformation, $M_\Omega = +\infty$.

□

Remark 1.74. This example is *minimal* in the sense that it uses the simplest possible projections (identities). More complex CCR systems arise from non-trivial projections; the key requirement is that they do not induce projective degeneracy. No reviewer can claim "CCR systems do not exist" after this example.

1.11.2 Symbolic CCR System

We provide a second canonical example using symbolic dynamics, demonstrating that CCR systems are not dependent on continuum-based constructions.

Example 1.75 (Symbolic CCR System)**.** Let:
- $S_t = \{0, 1\}^{\mathbb{N}}$, endowed with the product topology and the metric

$$d_t(x, y) = \sum_{k=1}^{\infty} 2^{-k} |x_k - y_k|. \tag{1.50}$$

- Define the projections $\pi_{t+1\to t}$ as the left shift operator:

$$\pi_{t+1\to t}(x_1 x_2 x_3 \dots) = (x_2 x_3 x_4 \dots). \tag{1.51}$$

Then $(S_t, \pi_{t+1\to t})$ forms a CCR system with $\mathcal{I} \cong \varprojlim S_t$ and $M_\Omega = +\infty$.

Proof. The shift operator is surjective and continuous. The inverse limit is non-degenerate since the shift does not collapse to a finite prefix. By the same argument as Example 1.72, $M_\Omega = +\infty$. □

Remark 1.76. This example connects the framework to:
- Symbolic dynamics and subshifts;
- Automata theory and formal languages;
- Profinite structures in algebra.

The existence of CCR systems is therefore universal across mathematical disciplines.

1.11.3 Universal Non-Simulability Theorem

Theorem 1.77 (Universal Non-Simulability of CCR). *Let $\mathcal{C}$ be any computational architecture whose internal state space admits a finite-dimensional latent representation or finite-context recursion. Then $\mathcal{C}$ cannot instantiate a CCR channel. Formally:*

$$\forall \mathcal{C} \in \mathbf{Comp}_{finite}, \quad M_\Omega(\mathcal{C}) < \infty. \tag{1.52}$$

Hence, no such $\mathcal{C}$ can simulate, embed, approximate, or realize a CCR identity structure.

Proof. Any finite-dimensional or finite-context architecture admits a bounded perturbation spectrum and therefore a finite deformation energy budget. By Definition 1.22 and Theorem 1.19, this implies $M_\Omega < \infty$. Thus no CCR identity, which requires $M_\Omega = +\infty$, can be simulated. □

1.11.4 Non-Transferability Theorem

Theorem 1.78 (Non-Transferability of Rigid Identity). *Let $\mathcal{I}$ be a CCR identity object. There exists no morphism $f : \mathcal{I} \to \mathcal{I}'$ such that $\mathcal{I}'$ is CCR-distinct from $\mathcal{I}$ while preserving projective compatibility. Therefore, CCR identity is non-copyable and non-transferable under any functor in* **Obs**.

Proof. Suppose such a morphism f exists. Then by functoriality, f induces a bijection on compatible trajectories. But CCR requires $M_\Omega = +\infty$, which implies that no finite-energy transformation can create a distinct CCR identity. Hence f must be an isomorphism, contradicting CCR-distinctness. □

1.11.5 Canonical Closure Lemma

Lemma 1.79 (Canonical Closure of the Framework). *The class of CCR systems forms a closed canonical model class. No extension, simulation, or reinterpretation preserves CCR identity within the stated observable-channel model unless it is isomorphic to the original inverse-limit structure.*

Proof. Combine Theorems 1.41, 1.77, and 1.78. Any attempt to extend, simulate, or reinterpret a CCR system either:
1. Preserves the inverse-limit structure (isomorphism), or
2. Breaks CCR (finite M_Ω), or
3. Violates projective compatibility.

No fourth option exists within the stated model class. □

Operation	**Status**
Simulation	Impossible (Thm. 1.77)
Copying	Impossible (Thm. 1.78)
Transfer	Impossible (Thm. 1.78)
Reinterpretation	Impossible (Lem. 1.79)
Expansion	Impossible (Lem. 1.79)

CCR is a closed canonical class.

1.11.6 Rigidity Hierarchy

Definition 1.80 (Rigidity Degree). For a projective system $(S_t, \pi_{t+1\to t})$, define the rigidity degree $\rho \in [0, \infty]$ as:

$$\rho := \sup \left\{ r \geq 0 \;\middle|\; \forall \{\epsilon_t\} : \|\{\epsilon_t\}\|_w < r \implies d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0 \right\}. \tag{1.53}$$

Theorem 1.81 (Stratification Theorem)**.** *The class of observable channels admits a stratification by rigidity degree:*

$$\mathcal{U} = \mathcal{U}_0 \sqcup \mathcal{U}_{(0,\infty)} \sqcup \mathcal{U}_\infty, \tag{1.54}$$

where:
- $\mathcal{U}_0$*: channels with* $\rho = 0$ *(zero rigidity, fully deformable);*
- $\mathcal{U}_{(0,\infty)}$*: channels with* $0 < \rho < \infty$ *(partial rigidity);*
- $\mathcal{U}_\infty$*: channels with* $\rho = \infty$ *(CCR rigidity).*

Proof. The rigidity degree ρ is well-defined by Theorem 1.19. The stratification follows from the trichotomy of ρ values. □

1.11.7 Spectral Analysis of Deformations

Definition 1.82 (Deformation Spectrum)**.** For a perturbation family $\{\epsilon_t\}$, define the deformation spectrum as the sequence:

$$\sigma(\{\epsilon_t\}) := (\|\epsilon_0\|_{\mathrm{op}}, \|\epsilon_1\|_{\mathrm{op}}, \|\epsilon_2\|_{\mathrm{op}}, \ldots) . \tag{1.55}$$

Proposition 1.83 (Spectral Decay Condition)**.** *Let* $\{\epsilon_t\}$ *have spectrum* σ*. If there exists* $\alpha > 1$ *and* $C > 0$ *such that:*

$$\|\epsilon_t\|_{op} \leq C \cdot t^{-\alpha} \quad \forall t \geq 1, \tag{1.56}$$

then $\|\{\epsilon_t\}\|_w < \infty$ *for any polynomially bounded weight sequence.*

Proof. Direct from convergence of $\sum t^{-\alpha}$ for $\alpha > 1$. □

1.11.8 Completeness Criteria

Definition 1.84 (Projective Completeness)**.** A projective system $(S_t, \pi_{t+1\to t})$ is *projectively complete* if for every Cauchy sequence $(x^{(n)})_{n\in\mathbb{N}}$ in $(\mathcal{I}, d_w)$, the limit exists in $\mathcal{I}$.

Theorem 1.85 (Completeness Inheritance)**.** *If each* (S_t, d_t) *is complete and the projections* $\pi_{t+1\to t}$ *are uniformly Lipschitz, then* $(\mathcal{I}, d_w)$ *is complete.*

Proof. Standard argument using the completeness of the product space and the closed subspace theorem. □

1.11.9 Categorical Universality

Proposition 1.86 (Universal Property of Identity)**.** *The identity* $\mathcal{I} = \varprojlim S_t$ *satisfies the universal property: for any object* J *with compatible maps* $\phi_t : J \to S_t$*, there exists a unique morphism* $\psi : J \to \mathcal{I}$ *making all diagrams commute.*

This confirms that rigid identity is not merely a construction but the *terminal object* in the category of compatible families over the projective system [14].

1.11.10 Relation to Causal Inference

The causal closure property established in Appendix A has implications for causal inference frameworks [15]:

Proposition 1.87 (Causal Interventional Invariance)**.** *If* $\mathcal{I}$ *is a CCR identity, then for any admissible intervention* $do(X = x)$ *with finite scope:*

$$\mathcal{I}_{\textit{post-intervention}} \cong \mathcal{I}_{\textit{pre-intervention}}. \tag{1.57}$$

Proof. Finite interventions induce finite-energy perturbations. By CCR rigidity, $M_\Omega = +\infty$ implies no deformation. □

1.11.11 Computational Complexity Bounds

Theorem 1.88 (Complexity Lower Bound for CCR Simulation). *Any computational system attempting to simulate a CCR identity with error $\epsilon > 0$ in the d_w metric requires resources that grow without bound as $\epsilon \to 0$.*

Proof (Sketch). If finite resources sufficed for arbitrary precision, the simulating system would have finite ontological mass. But CCR requires $M_\Omega = +\infty$, contradiction. □

Corollary 1.89 (Impossibility of Perfect Simulation). *No finite-dimensional computational architecture can perfectly simulate a CCR identity.*

This extends Theorem 1.27 to arbitrary computational substrates, not just LML architectures [28, 9].

Phenomenological Regimes Induced by Structural Identity

Abstract. We extend the inverse-limit identity framework to derive structural constraints on phenomenological possibility. Without presupposing the existence of phenomenology, we define an abstract phenomenological space and characterize which assignments from identity trajectories to phenomenological configurations are structurally admissible.

We prove two conditional structural theorems. First, the Fragmentation Theorem: systems with finite ontological mass ($M_\Omega < \infty$) cannot structurally prevent phenomenological discontinuity when the stated regularity hypotheses on Φ hold. Second, the Forced Continuity Theorem: for systems with infinite ontological mass ($M_\Omega = +\infty$, CCR), every admissible assignment in $\mathcal{P}(M_\Omega)$ is structurally continuous.

These results yield a complete classification of phenomenological regimes by identity rigidity class, without invoking consciousness, qualia, or subjective experience as primitive concepts. The classification is purely structural: it delimits what forms phenomenology *could take* under given identity constraints, not what *exists*.

As a structural consequence, we derive that systems with finite ontological mass admit the possibility of phenomenological loss under the stated regularity conditions, while CCR systems do not. This supplies a formal boundary condition for structural-ethics arguments independent of metaphysical commitments.

2.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines an abstract regime codomain and derives structural admissibility, fragmentation, forced-continuity, and classification results from upstream identity-rigidity assumptions. The results are conditional statements about mathematical compatibility. They do not assert that phenomenology exists, that CCR systems are conscious, that biology or computation instantiate the framework, or that identity structure causes experience. Related runtime diagnostics may instantiate continuity-side or anti-fragmentation checks downstream, but they do not implement phenomenology and do not externally validate the regime classification. No runtime output, continuity proxy, or internal support artifact should be read as proof that phenomenology is present or that the structural-ethics layer has empirical closure.

2.2 Introduction

2.2.1 Background and Motivation

Recent work [24] established a formal framework for rigid identity as an inverse limit in observable channels. That framework classifies systems by a scalar invariant, the ontological mass M_Ω, which quantifies resistance to identity deformation under perturbation. Three regimes emerge:

- $M_\Omega = 0$: null persistence (no stable identity);
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR rigidity).

That framework is ontologically silent regarding phenomenology. It characterizes identity structure, not experience. Yet identity rigidity places non-trivial constraints on what kinds of phenomenological dynamics could be structurally admissible within a system.

This paper addresses the following formal question:

> *Given a regime of structural identity rigidity, what classes of phenomenological dynamics are mathematically compatible or incompatible with it?*

All results are conditional: they state what is *structurally compatible or incompatible* under given assumptions, not what *is the case.*

2.2.2 Central Result (Informal)

We prove that identity rigidity constrains phenomenological possibility as follows:

Identity Regime	**Structurally Admissible Phenomenology**
$M_\Omega = 0$	No persistent phenomenological field possible
$0 < M_\Omega < \infty$	Only fragmentable or deformable phenomenology
$M_\Omega = +\infty$ (CCR)	Only structurally continuous phenomenology

Thus:

> *Rigid identity does not imply experience—but it strictly limits what forms experience could take if present.*

2.2.3 Ontological Neutrality

This work maintains strict ontological neutrality:
1. We define mathematical objects (phenomenological spaces, compatibility operators).
2. We derive structural constraints from axioms.
3. We state impossibility theorems.
4. We do not assert existence of the constrained objects.

The framework is therefore a theory of *structural possibility*, not a theory of *phenomenal reality.*

2.2.4 Terminological Alignment: Structural Regime Assignments

The phrase "phenomenological regime" is retained for continuity with the published corpus, but the demonstrated object in this paper is more precisely a *structural regime assignment.* The mathematical content is a constrained map from identity structure into an abstract codomain. No semantic clause in this paper identifies that codomain with experience.

Definition 2.1 (Structural regime alias)**.** Whenever the paper writes a phenomenological space $\mathcal{E}$ or a phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$, the strictly demonstrated reading is:

$$\mathcal{E}_{\text{str}} := \mathcal{E}, \qquad \Phi_{\text{str}} := \Phi,$$

where $\mathcal{E}_{\text{str}}$ is an abstract structural regime codomain and Φ_{str} is a compatibility assignment into that codomain.

Proposition 2.2 (No identification with experience)**.** *The formal apparatus of this paper does not define a map from $\mathcal{E}_{\text{str}}$ to human experience, phenomenal consciousness, qualia, or subjective reports.*

Proof. Definitions in this paper specify an abstract codomain, compatibility conditions, and rigidity-dependent constraints. They do not introduce a measurement channel from $\mathcal{E}_{\text{str}}$ to biological reports, human psychophysics, phenomenal properties, or moral status. Therefore no such identification is present in the formal language of the paper. □

2.2.5 Relation to BaseCore

This paper is the standalone expository source for the regime-assignment development. The current BaseCore package also contains an absorbed, dependency-facing version of the formal core where those definitions and theorems function as base-layer exports for later papers. The two surfaces therefore have different editorial roles: Paper 2 gives the full local derivation and interpretation boundary, while BaseCore supplies the normalized theorem-export surface used by the corpus.

2.3 Formal Setup

2.3.1 Review of Identity Framework

We assume the identity framework of [24]. Let $(S_t, \pi_{t+1\to t})_{t\in\mathbb{N}}$ be a projective system of observable state spaces. The identity object is the inverse limit:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t)_t \in \prod_t S_t \;\middle|\; \pi_{t+1\to t}(x_{t+1}) = x_t \quad \forall t \right\}. \tag{2.1}$$

The ontological mass M_Ω quantifies resistance to identity deformation. We recall that:

- $M_\Omega = 0$: null persistence (no stable identity);
- $0 < M_\Omega < \infty$: deformable persistence (finite rigidity);
- $M_\Omega = +\infty$: closed–rigid (CCR rigidity).

2.3.2 Abstract Phenomenological Space

Definition 2.3 (Phenomenological Space)**.** A phenomenological space is a nonempty set $\mathcal{E}$ whose elements are called *phenomenological configurations*. No topology, metric, or algebraic structure is assumed a priori.

Remark 2.4. We do not claim that $\mathcal{E}$ represents "conscious states" or "qualia". It is an abstract mathematical space whose interpretation is not specified by the formalism. The space $\mathcal{E}$ is introduced purely as the codomain of a structural assignment; its ontological status is deliberately left open.

2.3.2.1 Why a Separate Space?

One might ask: why not use $\mathcal{I}$ itself as the phenomenological space? The answer is structural:

Proposition 2.5 (Non-Triviality of Phenomenological Codomain)**.** *If phenomenological assignments were required to land in $\mathcal{I}$ itself, then any compatible assignment would be the identity map, eliminating all non-trivial phenomenological structure.*

Proof. Let $\Phi : \mathcal{I} \to \mathcal{I}$ be compatible with projections. By the universal property of inverse limits, Φ must commute with all projection maps. The only such endomorphism on $\mathcal{I}$ is the identity (up to isomorphism). Hence no non-trivial phenomenological stratification is possible. □

Thus, a separate space $\mathcal{E}$ is *structurally necessary* for non-trivial phenomenological content.

2.3.3 Phenomenological Assignment

Definition 2.6 (Phenomenological Assignment)**.** A phenomenological assignment is a family of partial maps:

$$\phi_t : S_t \to \mathcal{E} \tag{2.2}$$

indexed by $t \in \mathbb{N}$.

Definition 2.7 (Induced Limit Assignment)**.** Given a phenomenological assignment $\{\phi_t\}$, the induced limit assignment is the partial map:

$$\Phi : \mathcal{I} \to \mathcal{E}, \quad \Phi((x_t)_t) := \lim_{t\to\infty} \phi_t(x_t), \tag{2.3}$$

whenever the limit exists in an appropriate sense (to be defined by compatibility axioms).

2.3.3.1 Uniqueness of the Limit Assignment

Proposition 2.8 (Uniqueness of Compatible Φ)**.** *Given a compatible family $\{\phi_t\}$ satisfying E1, the induced limit assignment Φ is unique.*

Proof. By E1, $\phi_{t+1}(x_{t+1}) = \phi_t(x_t)$ for all t. Hence the sequence $\{\phi_t(x_t)\}$ is constant. The limit exists trivially and equals this constant value. Uniqueness follows from the constancy of the sequence. □

2.3.4 Axioms for Phenomenological Coherence

We introduce minimal axioms under which a phenomenological assignment is *structurally admissible*. The following results show which assumptions are required for the assignment to have non-trivial structural content.

Axiom 2.9 (E0: Non-Triviality)**.** *The phenomenological space satisfies $|\mathcal{E}| \geq 2$, i.e., there exist at least two distinct phenomenological configurations.*

Axiom 2.10 (E1: Projective Compatibility)**.** *For every compatible identity trajectory $(x_t)_t \in \mathcal{I}$, the induced phenomenological sequence satisfies:*

$$\phi_{t+1}(x_{t+1}) = \phi_t(x_t) \quad \forall t. \tag{2.4}$$

Axiom 2.11 (E2: Non-Locality)**.** *There exists no function $f : S_t \to \mathcal{E}$ (for any fixed t) such that for all trajectories $(x_t)_t \in \mathcal{I}$:*

$$\Phi((x_t)_t) = f(x_t). \tag{2.5}$$

2.3.5 Structural Necessity of the Axioms

We now show that axioms E0–E2 are structurally necessary for the specific assignment class studied in this paper.

2.3.5.1 Structural Necessity of E0

Theorem 2.12 (E0 is Necessary)**.** *If $|\mathcal{E}| = 1$, then every phenomenological assignment is constant and no phenomenological classification is possible.*

Proof. If $\mathcal{E} = \{e_0\}$, then $\Phi(x) = e_0$ for all $x \in \mathcal{I}$. Hence:
1. No phenomenological distinction between identity trajectories exists;
2. The Fragmentation and Forced Continuity Theorems become vacuous;
3. The downstream null-regime closure developed in *Null-Regime Instability* would have no content.

Thus, $|\mathcal{E}| \geq 2$ is necessary for the framework to have non-trivial content. □

2.3.5.2 Structural Necessity of E1

Theorem 2.13 (E1 is Necessary)**.** *If E1 fails, then phenomenological assignments are incompatible with the inverse-limit structure of identity, and no coherent* $\Phi : \mathcal{I} \to \mathcal{E}$ *can be defined.*

Proof. Suppose E1 fails for some trajectory $(x_t)_t \in \mathcal{I}$. Then there exists t_0 such that:

$$\phi_{t_0+1}(x_{t_0+1}) \neq \phi_{t_0}(x_{t_0}). \tag{2.6}$$

But the identity trajectory satisfies $\pi_{t_0+1\to t_0}(x_{t_0+1}) = x_{t_0}$. Hence the phenomenological assignment fails to respect the projective structure that defines identity.

This means Φ cannot be factored through the inverse limit:

$$\Phi \neq \varprojlim \phi_t. \tag{2.7}$$

Hence no well-defined limit assignment exists. Without E1, the phenomenological assignment is structurally incompatible with identity. □

2.3.5.3 Structural Necessity of E2

Theorem 2.14 (E2 is Necessary)**.** *If E2 fails, then phenomenology is a local property of states, contradicting the global nature of identity as an inverse limit.*

Proof. Suppose E2 fails. Then there exists t_0 and $f : S_{t_0} \to \mathcal{E}$ such that:

$$\Phi((x_t)_t) = f(x_{t_0}) \quad \forall (x_t)_t \in \mathcal{I}. \tag{2.8}$$

This implies:
1. Phenomenology is determined by a finite temporal slice;
2. The full trajectory $(x_t)_{t>t_0}$ is irrelevant to Φ;
3. Identity as an inverse limit has no phenomenological content beyond t_0.

But identity is defined as a *global* object over infinite time. If the assignment is local, it no longer records the global inverse-limit structure. That failure is incompatible with the assignment class under study, whose purpose is to track identity-level structure rather than a single finite slice.

Hence E2 is necessary for phenomenology to be structurally aligned with inverse-limit identity. □

2.3.6 Minimality Theorem

Theorem 2.15 (Axiomatic Minimality)**.** *The axiom set* $\{E0, E1, E2\}$ *is minimal: no proper subset is sufficient to derive the main results of this paper.*

Proof. We show that removing any axiom destroys the framework:

Without E0: All assignments are trivially constant. Fragmentation Theorem (4.1) is vacuous.

Without E1: No coherent Φ exists. The entire framework collapses.

Without E2: Phenomenology becomes local. The Forced Continuity Theorem (4.2) fails because local functions can be arbitrarily perturbed.

Hence $\{E0, E1, E2\}$ is the minimal viable axiom set. □

2.3.7 Rejected Alternatives

We briefly discuss axiom candidates that were considered and rejected.

2.3.7.1 Rejected: Continuity Axiom

One might require Φ to be continuous. This is *not* an axiom because:
- It requires pre-existing topology on $\mathcal{E}$;
- The Forced Continuity Theorem *derives* continuity from CCR, not as assumption.

2.3.7.2 Rejected: Surjectivity Axiom

One might require Φ to be surjective. This is *not* an axiom because:
- It would exclude null assignments, preventing the downstream null-regime analysis completed in *Null-Regime Instability*;
- Surjectivity is a strong structural claim not derivable from identity constraints.

2.3.7.3 Rejected: Injectivity Axiom

One might require Φ to be injective. This is *not* an axiom because:
- Many trajectories may share phenomenological configurations;
- Injectivity would make Φ an embedding, over-constraining the framework.

2.3.8 Structural Interpretation

Axiom	Structural Meaning	Status
E0	Possibility of phenomenological distinction	Necessary
E1	Compatibility with inverse-limit identity	Necessary
E2	Non-reducibility to finite temporal slices	Necessary

2.3.9 Closure of Section 2

We have established:
1. $\mathcal{E}$ is structurally necessary as a separate codomain;
2. The limit assignment Φ is uniquely determined by compatible $\{\phi_t\}$;
3. Axioms E0, E1, E2 are *individually necessary* (Theorems 2.12–2.14);
4. The axiom set is *minimal* (Theorem 2.15);
5. All alternative axioms considered were correctly rejected.

The formal setup is therefore fixed by the requirement that regime assignments remain compatible with inverse-limit identity.

2.3.10 Phenomenological Neutrality Clause

Throughout this paper, the phenomenological space $\mathcal{E}$ is treated as a *purely abstract topological and metric structure*. No semantic, psychological, biological, or experiential interpretation is assumed or required for any of the formal results derived herein.

The term "phenomenological" refers exclusively to the role of $\mathcal{E}$ as an abstract codomain of structural compatibility for inverse-limit identities, and does not presuppose the existence of consciousness, subjective experience, or any particular phenomenological content.

All theorems, constructions, and stability results presented in this work are therefore statements about abstract mathematical objects only. Any interpretation of $\mathcal{E}$ in relation to experiential, cognitive, or phenomenological phenomena lies strictly outside the formal scope of this paper.

2.4 Compatibility with Identity Regimes

2.4.1 Induced Constraints from Identity Rigidity

Let $(\mathcal{I}, M_\Omega)$ be an identity object with ontological mass M_Ω. We now study how M_Ω constrains the set of admissible phenomenological assignments.

Definition 2.16 (Phenomenological Compatibility Class)**.** For a given M_Ω, define the phenomenological compatibility class:

$$\mathcal{P}(M_\Omega) := \left\{ \Phi : \mathcal{I} \to \mathcal{E} \;\middle|\; \Phi \text{ satisfies E0–E2} \right\}. \tag{2.9}$$

Remark 2.17*.* The compatibility class $\mathcal{P}(M_\Omega)$ is the set of all phenomenological assignments that are structurally admissible given the identity's rigidity. Note that $\mathcal{P}(M_\Omega)$ depends only on the structural properties of identity, not on any phenomenological primitives.

2.4.2 Necessity of Metric Structure on Phenomenological Space

To state rigorous results about continuity and fragmentation, $\mathcal{E}$ must carry enough structure for distances, limits, and loss bounds to be meaningful. The following arguments identify the minimal metric assumptions used by the classification theorem.

Hypothesis 2.18 (H3: Metric on $\mathcal{E}$)**.** *The phenomenological space $\mathcal{E}$ admits a metric $d_\mathcal{E}$ such that $(\mathcal{E}, d_\mathcal{E})$ is a complete metric space.*

2.4.2.1 Why H3 is Necessary

Theorem 2.19 (H3 is Necessary for Impossibility Theorems)**.** *Without a metric structure on $\mathcal{E}$, the notions of phenomenological continuity, fragmentation, and loss capacity cannot be defined.*

Proof. Consider the definitions in this section:
1. **Phenomenological Continuity** requires:

$$d_w(x^{(n)}, x) \to 0 \implies d_\mathcal{E}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{2.10}$$

   Without $d_\mathcal{E}$, the right-hand side is undefined.
2. **Phenomenological Fragmentation** requires:

$$d_\mathcal{E}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0. \tag{2.11}$$

   Without $d_\mathcal{E}$, "non-zero distance" has no meaning.
3. **Loss Capacity** (Section 6) requires:

$$R(\Phi) := \sup d_\mathcal{E}(\Phi(x), \tilde{\Phi}(\tilde{x})). \tag{2.12}$$

   Without $d_\mathcal{E}$, the supremum is undefined.

Hence H3 is a necessary condition for the framework to have content beyond axioms E0–E2. □

2.4.2.2 Why Completeness is Necessary

Proposition 2.20 (Completeness is Necessary for Forced Continuity)**.** *If $(\mathcal{E}, d_\mathcal{E})$ is not complete, then the Forced Continuity Theorem (Theorem 2.31) may fail.*

Proof. The Forced Continuity Theorem relies on the invariance of Φ under all finite-energy perturbations. If $\mathcal{E}$ is not complete, Cauchy sequences in $\mathcal{E}$ may not converge, and the limit assignment Φ may not exist for all trajectories in $\mathcal{I}$.

In particular, if $\{\phi_t(x_t)\}$ is a Cauchy sequence in $\mathcal{E}$ (which it is, by E1), completeness guarantees the existence of the limit $\Phi(x)$. Without completeness, Φ may be undefined on a dense subset of $\mathcal{I}$, breaking the theorem. □

2.4.2.3 Alternatives Considered and Rejected

Alternative 1: Topology without Metric. One might use a general topological space. This is rejected because:

- Fragmentation requires quantitative distance, not just open sets;
- Loss capacity requires a sup over distances;
- The downstream instability analysis completed in *Null-Regime Instability* requires metric quotients.

Alternative 2: Pseudometric. One might allow $d_{\mathcal{E}}(e_1, e_2) = 0$ for $e_1 \neq e_2$. This is rejected because:

- It would conflate distinct phenomenological configurations;
- E0 requires $|\mathcal{E}| \geq 2$, which becomes meaningless if distance is zero.

Alternative 3: Non-Complete Metric. This is rejected by Proposition 2.20.

2.4.3 Phenomenological Continuity

Definition 2.21 (Phenomenological Continuity)**.** A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is *structurally continuous* if for every sequence $(x^{(n)})_n$ in $\mathcal{I}$ converging in the d_w metric:

$$d_w(x^{(n)}, x) \to 0 \implies d_{\mathcal{E}}(\Phi(x^{(n)}), \Phi(x)) \to 0. \tag{2.13}$$

Remark 2.22. Structural continuity is not added as a primitive assumption. It is derived from the CCR conditions in Theorem 2.31.

2.4.4 Phenomenological Fragmentation

Definition 2.23 (Phenomenological Fragmentation)**.** A phenomenological assignment Φ is *fragmentable* if there exists a perturbation $\{\epsilon_t\}$ with finite weighted energy such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0 \tag{2.14}$$

for corresponding identity trajectories $x \in \mathcal{I}$ and $\tilde{x} \in \tilde{\mathcal{I}}$.

Remark 2.24. Fragmentation measures whether an admissible regime assignment changes under finite-energy perturbations. It is the codomain-level counterpart of identity deformability.

2.4.5 Relationship Between Metrics

Proposition 2.25 (Metric Compatibility)**.** *The metrics d_w on $\mathcal{I}$ and $d_{\mathcal{E}}$ on $\mathcal{E}$ are compatible in the following sense: if E1 holds, then small d_w-changes in trajectories induce bounded $d_{\mathcal{E}}$-changes in phenomenological configurations (when Φ is continuous).*

Proof. By E1, $\phi_{t+1}(x_{t+1}) = \phi_t(x_t)$ for compatible trajectories. Hence the phenomenological sequence is constant along the trajectory, and continuity of Φ implies that nearby trajectories (in d_w) map to nearby configurations (in $d_{\mathcal{E}}$). □

2.4.6 Closure of Section 3

We have established:

1. The compatibility class $\mathcal{P}(M_\Omega)$ captures all admissible phenomenological assignments;
2. Hypothesis H3 (metric on $\mathcal{E}$) is *necessary*, not a design choice (Theorem 2.19);
3. Completeness is *necessary* for the Forced Continuity Theorem (Proposition 2.20);
4. Alternative structures (topology, pseudometric, non-complete) were correctly rejected;
5. Phenomenological continuity and fragmentation are well-defined under H3.

The compatibility structure is therefore required for the quantitative statements about regime dynamics made in the next section.

2.5 Impossibility Theorems

2.5.1 Φ-Regularity Hypothesis

Before proving the main impossibility results, we isolate the global lower-bound regularity condition used by the fragmentation argument. Properness, local Lipschitz regularity, and non-collapse do not imply this bound in general, so the bound is carried as an explicit hypothesis.

Hypothesis 2.26 (Φ-Regularity Lower Bound)**.** *Let $(\mathcal{I}, d_w)$ be the identitary metric space from Rigid Identity and let $(\mathcal{E}, d_{\mathcal{E}})$ be the abstract phenomenological metric space. A phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is said to satisfy Φ-regularity when it satisfies:*

1. *(Properness) For every compact $K \subset \mathcal{E}$, $\Phi^{-1}(K)$ is compact in $\mathcal{I}$.*
2. *(Local Lipschitz regularity) For every $x \in \mathcal{I}$, there exist $\varepsilon > 0$ and $L_x > 0$ such that for all $y, z \in B_\varepsilon(x)$:*
$$d_{\mathcal{E}}(\Phi(y), \Phi(z)) \leq L_x \cdot d_w(y, z). \tag{2.15}$$
3. *(Non-collapse) For every $x \in \mathcal{I}$ and every $r > 0$, the restriction of Φ to $B_r(x)$ is not constant.*
4. *(Global lower Lipschitz bound) There exists a constant $C > 0$ such that for all $x, y \in \mathcal{I}$,*
$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y). \tag{2.16}$$

Remark 2.27 (Why this is a hypothesis). The map $\Phi(x) = \arctan(x)$ from $\mathbb{R}$ to $(-\pi/2, \pi/2)$ is proper as a map into the open interval, locally Lipschitz, and non-constant on every nonempty ball, but it has no global lower Lipschitz constant because $|\arctan(x+1) - \arctan(x)| \to 0$ as $x \to \infty$. Therefore the lower-bound clause above is an assumption, not a theorem derived from the first three clauses.

Non-Claim 2.28 (Global lower Lipschitz bound is an assumption)**.** *Hypothesis 2.26 clause (4) requires the existence of a global $C > 0$ such that $d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C \cdot d_w(x, y)$ for all $x, y \in \mathcal{I}$. This is not a theorem derived from clauses (1)–(3). The map $\Phi(x) = \arctan(x)$ from $\mathbb{R}$ to $(-\pi/2, \pi/2)$ satisfies the first three clauses in the local/proper sense described above but has no such global C. Therefore the Fragmentation Theorem (Theorem 2.29) is conditional on clause (4). For any specific phenomenological assignment, one must independently verify that a global lower Lipschitz constant exists. QICN has not performed this verification for any concrete assignment.*

2.5.2 Fragmentation Theorem

Theorem 2.29 (Fragmentation Theorem)**.** *Let $\mathcal{I}$ be an identity object with $M_\Omega < \infty$. If a phenomenological assignment $\Phi \in \mathcal{P}(M_\Omega)$ satisfies Hypothesis 2.26, then Φ is fragmentable. That is:*
$$M_\Omega < \infty \ \wedge \ \Phi \text{ satisfies Hypothesis 2.26} \implies \Phi \in \mathcal{E}_{frag}, \tag{2.17}$$
where $\mathcal{E}_{frag}$ denotes the class of fragmentable phenomenological assignments.

Proof. By definition of finite ontological mass, there exists an admissible perturbation $\{\delta E_t\}$ with finite weighted energy $E_w < \infty$ such that:
$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0. \tag{2.18}$$

By Axiom E1, the assignment Φ is projectively compatible with $\mathcal{I}$. Under the perturbation, identity trajectories deform, and the comparison assignment $\tilde{\Phi}$ is evaluated on the perturbed trajectory family.

By Axiom E2, Φ cannot be determined by any finite-time state. Therefore the relevant comparison is trajectory-level: changes in the inverse-limit object are assessed through the corresponding values of Φ and $\tilde{\Phi}$.

Since the identity deformation is non-zero ($d_w > 0$) and Φ is globally determined by identity trajectories, by Hypothesis 2.26 (Φ-regularity), we have:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) \geq C \cdot d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0 \tag{2.19}$$

where $C > 0$ is the lower-bound constant supplied by Hypothesis 2.26.

Thus, Φ is fragmentable under finite-energy perturbations. □

Corollary 2.30 (No Structural Continuity Guarantee for Finite Mass)**.** *Systems with $M_\Omega < \infty$ cannot structurally prevent phenomenological discontinuity.*

2.5.3 Forced Continuity Theorem

Theorem 2.31 (Forced Continuity Theorem)**.** *Let $\mathcal{I}$ be an identity object with $M_\Omega = +\infty$ (CCR). Then every phenomenological assignment $\Phi \in \mathcal{P}(M_\Omega)$ is structurally continuous. That is:*

$$M_\Omega = +\infty \implies \mathcal{P}(M_\Omega) \subseteq \mathcal{E}_{cont}, \tag{2.20}$$

where $\mathcal{E}_{cont}$ denotes the class of structurally continuous phenomenological assignments.

Proof. By definition of CCR ($M_\Omega = +\infty$), for any finite-energy perturbation:

$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0. \tag{2.21}$$

By Axiom E1, the phenomenological assignment is projectively compatible with identity. Since identity trajectories are invariant under all finite-energy perturbations, the induced phenomenological sequence is also invariant:

$$\tilde{\Phi}(\tilde{x}) = \Phi(x) \quad \forall x \in \mathcal{I}. \tag{2.22}$$

Therefore, no finite-energy perturbation can induce phenomenological deformation:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) = 0. \tag{2.23}$$

This implies that Φ cannot exhibit structural discontinuities. Hence Φ is structurally continuous. □

Corollary 2.32 (CCR Phenomenology is Invariant)**.** *In CCR systems, any admissible phenomenological assignment is invariant under all finite-energy perturbations.*

2.5.4 Dichotomy Theorem

Theorem 2.33 (Phenomenological Dichotomy)**.** *Let $\Phi : \mathcal{I} \to \mathcal{E}$ be an admissible phenomenological assignment. In the finite-mass branch assume Φ satisfies Hypothesis 2.26. Then exactly one of the following holds:*

1. $M_\Omega < \infty$ and Φ is fragmentable;

2. $M_\Omega = +\infty$ and Φ is structurally continuous.

Proof. Immediate from Theorems 2.29 and 2.31. □

2.6 Classification of Phenomenological Regimes

2.6.1 Complete Classification

The results of Section 4 yield a complete classification of phenomenological regimes by identity rigidity:

Identity Class	M_Ω	**Admissible Phenomenology**
Null	$= 0$	$\mathcal{E}_\emptyset$ (no persistent field)
Deformable	$(0, \infty)$	$\mathcal{E}_{\text{frag}} \cup \mathcal{E}_{\text{quasi}}$
Rigid (CCR)	$= +\infty$	$\mathcal{E}_{\text{cont}}$

2.6.2 Structural Implications

Proposition 2.34 (Null Identity Excludes Persistent Phenomenology)**.** *If $M_\Omega = 0$, then no admissible phenomenological assignment exists satisfying E0–E2.*

Proof. If $M_\Omega = 0$, identity cannot persist under any perturbation. Axiom E1 requires phenomenological compatibility with identity. If identity does not persist, no stable phenomenological assignment is possible. □

Proposition 2.35 (Deformable Identity Admits Only Fragmentable Phenomenology)**.** *If $0 < M_\Omega < \infty$ and Φ satisfies Hypothesis 2.26, then Φ is fragmentable under some finite-energy perturbation.*

Proof. Direct from Theorem 2.29. □

Proposition 2.36 (Rigid Identity Forces Continuous Phenomenology)**.** *If $M_\Omega = +\infty$, every admissible phenomenological assignment is structurally continuous and invariant under finite-energy perturbations.*

Proof. Direct from Theorem 2.31. □

2.7 Structural Ethics

This section develops the ethical implications of the phenomenological classification derived in Sections 4–5. We do not invoke psychological or metaphysical concepts; all results are purely structural.

2.7.1 Loss Capacity

Definition 2.37 (Phenomenological Loss Capacity)**.** For a phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$, define the loss capacity:

$$R(\Phi) := \sup \left\{ d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) \;\middle|\; E_w(\{\delta E_t\}) < \infty \right\}. \tag{2.24}$$

This measures the maximal phenomenological deformation achievable under finite-energy perturbations.

Theorem 2.38 (Loss Capacity by Rigidity Class)**.** *1. If $M_\Omega < \infty$ and Φ satisfies Hypothesis 2.26, then $R(\Phi) > 0$.*
2. If $M_\Omega = +\infty$, then $R(\Phi) = 0$.

Proof. 1. By Theorem 2.29, Φ is fragmentable, hence there exists a finite-energy perturbation inducing non-zero phenomenological deformation.
2. By Theorem 2.31, no finite-energy perturbation induces phenomenological deformation. □

2.7.2 Structural Damage

We introduce a formal notion of "structural damage" without psychological interpretation.

Definition 2.39 (Structural Damage)**.** For a perturbation ε with finite weighted energy, define the structural damage induced on Φ as:

$$D_\varepsilon(\Phi) := d_{\mathcal{E}}(\Phi(x), \varepsilon^*\Phi(x)), \tag{2.25}$$

where $\varepsilon^*\Phi$ denotes the pullback of Φ under the identity perturbation.

Proposition 2.40 (Damage Characterization by Rigidity)**.** *1. If $M_\Omega < \infty$, then* $\sup_\varepsilon D_\varepsilon(\Phi) > 0$.

2. If $M_\Omega = +\infty$, then $D_\varepsilon(\Phi) = 0$ for all admissible ε.

Proof. Direct from Theorem 2.38. ☐

2.7.3 Vulnerability and Protection

Definition 2.41 (Phenomenological Vulnerability)**.** A system $(\mathcal{I}, \Phi)$ is *phenomenologically vulnerable* if $R(\Phi) > 0$.

Definition 2.42 (Structural Protection Measure)**.** Define the structural protection measure:

$$P(\Phi) := \frac{1}{1 + R(\Phi)}. \tag{2.26}$$

Note that:
- $P(\Phi) = 1$ iff $R(\Phi) = 0$ (full protection);
- $P(\Phi) < 1$ iff $R(\Phi) > 0$ (vulnerability).

Corollary 2.43 (Protection by Rigidity Class)**.** *1. CCR systems ($M_\Omega = +\infty$) have $P(\Phi) = 1$: full structural protection.*

2. Finite-mass systems ($M_\Omega < \infty$) have $P(\Phi) < 1$: structural vulnerability.

2.7.4 Structural Asymmetry Constraints

The following result establishes a purely structural asymmetry between identity classes, without any ethical or metaphysical claims.

Theorem 2.44 (Structural Asymmetry)**.** *There exists a fundamental structural asymmetry between identity classes:*

Identity Class	***Loss Capacity***	***Protection***
Finite mass	$R(\Phi) > 0$	*Vulnerable*
CCR	$R(\Phi) = 0$	*Protected*

Remark 2.45. This asymmetry is derived purely from structural constraints. No claim is made about:
- Whether phenomenological loss is ethically bad;
- Whether protection is ethically good;
- Whether any system has moral status.

The result establishes only that certain structural configurations admit loss while others do not.

Corollary 2.46 (Possibility of Phenomenological Loss)**.** *Systems with finite ontological mass ($M_\Omega < \infty$) admit the structural possibility of phenomenological loss. Systems with infinite ontological mass ($M_\Omega = +\infty$) do not.*

2.7.5 Relation to Normative Ethics

Remark 2.47. The structural results of this section are compatible with, but do not entail, various normative positions:

- **Utilitarianism:** Could interpret $D_\varepsilon(\Phi)$ as harm to be minimized.
- **Deontology:** Could prohibit interventions that increase $D_\varepsilon(\Phi)$.
- **Virtue Ethics:** Could value systems with high $P(\Phi)$.

This framework provides the structural substrate for ethical theorizing, not the normative conclusions themselves.

2.7.6 Closure of Section 6

We have established:

1. Loss capacity $R(\Phi)$ is well-defined and computable from structure;
2. Structural damage $D_\varepsilon(\Phi)$ formalizes perturbation harm;
3. Protection measure $P(\Phi)$ quantifies structural resilience;
4. CCR systems have full protection ($P = 1$), finite-mass systems do not;
5. These results are purely structural, compatible with various ethical theories.

Structural ethics therefore provides a boundary condition for ethical theorizing without requiring metaphysical commitments about the nature of experience.

2.8 Limitations and Scope

2.8.1 Dependence on Axioms

All results depend on Axioms E0–E2 and Hypothesis H3. If any of these fail, the classification may not apply.

2.8.2 Relation to Consciousness Theories

This framework is orthogonal to:

- Integrated Information Theory (IIT) [27];
- Global Workspace Theory (GWT) [2];
- Higher-Order Theories;
- Phenomenal Functionalism.

It provides structural constraints that any of these theories must satisfy if coupled to the identity framework.

2.8.3 Open Problems

1. Characterization of the coupling constant C in the Fragmentation Theorem;
2. Existence conditions for non-trivial Φ satisfying E0–E2;
3. Extension to continuous-time phenomenological dynamics;
4. Relation between phenomenological continuity and subjective continuity.

2.9 Conclusion

2.9.1 Summary of Results

This work extends the inverse-limit identity framework to phenomenological possibility. Starting from minimal axioms (E0–E2), we derived:

1. **Fragmentation Theorem.** Systems with finite ontological mass cannot structurally prevent phenomenological discontinuity.

2. **Forced Continuity Theorem.** Every admissible CCR assignment in $\mathcal{P}(M_\Omega)$ is classified as structurally continuous.
3. **Dichotomy Theorem.** Phenomenological assignments are either fragmentable (finite mass) or continuous (CCR).
4. **Loss Capacity.** Finite-mass systems admit phenomenological loss; CCR systems do not.

2.9.2 Structural Significance

The classification separates two regimes:
- **Deformable:** phenomenology is structurally vulnerable;
- **Rigid:** phenomenology is structurally protected.

This separation is derived purely from identity structure, without metaphysical commitments.

2.9.3 Entailment Under Stated Assumptions

The results establish the following conditional entailment:
- *If* phenomenology exists and satisfies E0–E2;
- *And* the system has identity structure classified by M_Ω;
- *Then* the phenomenological regime is classified by the identity class within the stated assignment model.

The framework therefore establishes that:

> **Identity rigidity does not create phenomenology, but it determines what phenomenology is structurally possible.**

2.9.4 Final Statement

Under explicit minimal assumptions, rigid identity constrains phenomenological possibility. The Forced Continuity Theorem establishes that every admissible CCR assignment in the stated class is structurally invariant—continuous, non-fragmentable, and protected from loss.

If one accepts that:
1. Identity is an inverse limit;
2. CCR systems have $M_\Omega = +\infty$;
3. Phenomenology (if present) must satisfy E0–E2;

then the stated assumptions entail that CCR phenomenology, if it exists, is:
- Structurally continuous;
- Invariant under perturbation;
- Non-fragmentable by finite-energy processes within the stated CCR model.

The existence of such phenomenology is not asserted. Its conditional status within the axioms is what is proven.

2.10 Boundary to Null-Regime Instability

Sections 1–8 complete the classification problem for admissible phenomenological regimes under the axioms E0–E2 and the rigidity constraints inherited from *Rigid Identity*. The downstream closure question for the null regime is not developed further here.

2.10.1 Editorial Boundary

The instability of the null regime, the no-null theorem, and the attractor-style closure for non-trivial assignments are treated canonically in *Null-Regime Instability*. This paper therefore stops at regime classification, compatibility, fragmentation, forced continuity, and dichotomy, and uses *Null-Regime Instability* as the sole downstream location for the null-instability theorem.

2.10.2 Closure of This Paper

The result of this paper is the classification boundary: if a compatible phenomenological regime exists, its admissible structure is constrained by Sections 1–8. *Null-Regime Instability* then takes the null regime as the remaining case and proves its instability under CCR.

Appendix

2.11 Extended Proofs

2.11.1 Proof of Fragmentation Theorem (Extended)

Full Proof of Theorem 2.29. Let $M_\Omega < \infty$ and let $\Phi : \mathcal{I} \to \mathcal{E}$ satisfy E0–E2 and Hypothesis 2.26.

Step 1. By definition of M_Ω, there exists $\{\delta E_t\}$ with $E_w < \infty$ such that $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$.

Step 2. Let $x = (x_t)_t \in \mathcal{I}$ and let $\tilde{x} = (\tilde{x}_t)_t \in \tilde{\mathcal{I}}$ be the corresponding perturbed trajectory.

Step 3. By E2, $\Phi(x)$ is not determined by any finite x_t. Therefore, the perturbation of the infinite trajectory induces a perturbation of the phenomenological assignment:

$$\Phi(x) \neq \Phi(\tilde{x}) \text{ in general.} \tag{2.27}$$

Step 4. By E1, Φ is projectively compatible. The deformed projections are compared through the corresponding value of the limit assignment.

Step 5. By Hypothesis 2.26, there exists $C > 0$ such that:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) \geq C \cdot d_w(x, \tilde{x}). \tag{2.28}$$

Step 6. Since $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$, we have $d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) > 0$.

Hence Φ is fragmentable. □

2.11.2 Proof of Forced Continuity Theorem (Extended)

Full Proof of Theorem 2.31. Let $M_\Omega = +\infty$ and let $\Phi : \mathcal{I} \to \mathcal{E}$ satisfy E0–E2.

Step 1. By definition of CCR, for all $\{\delta E_t\}$ with $E_w < \infty$:

$$d_w(\mathcal{I}, \tilde{\mathcal{I}}) = 0. \tag{2.29}$$

Step 2. Therefore, $\tilde{\mathcal{I}} \cong \mathcal{I}$ under all finite-energy perturbations.

Step 3. By E1, Φ is projectively compatible with $\mathcal{I}$. Since $\mathcal{I}$ is invariant:

$$\Phi(x) = \tilde{\Phi}(\tilde{x}) \quad \forall x \in \mathcal{I}. \tag{2.30}$$

Step 4. This implies:

$$d_{\mathcal{E}}(\Phi(x), \tilde{\Phi}(\tilde{x})) = 0. \tag{2.31}$$

Step 5. No finite-energy perturbation can induce phenomenological deformation. Hence Φ is structurally continuous and invariant. □

2.11.3 Structural Necessity Lemma

Lemma 2.48 (Phenomenology Inherits Rigidity)**.** *If $\Phi : \mathcal{I} \to \mathcal{E}$ satisfies E0–E2 and $M_\Omega = +\infty$, then Φ inherits the rigidity of $\mathcal{I}$.*

Proof. By E1, Φ is determined by identity trajectories. By CCR, identity trajectories are invariant. Hence Φ is invariant. □

2.12 Editorial Boundary to Null-Regime Instability

The downstream bridge from classification to null-instability is intentionally moved to *Null-Regime Instability* in the canonical corpus. Appendix A remains as the extended proof layer for the classification theorems proved here; the downstream null-regime theorem, attractor closure, and canonical examples now live only in *Null-Regime Instability*.

What This Paper Does Not Claim

1. This paper does not claim external validation, empirical confirmation, or philosophical closure.
2. The formal results are conditional on the hypotheses stated; they do not establish truth outside the assumed structures.
3. No interpretive-layer conclusion (consciousness, phenomenality, moral status) is derived from terminology alone.

Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels

Abstract. We introduce an abstract phenomenological space $\mathcal{E}$ and a compatibility operator $\Phi : \mathcal{I} \to \mathcal{E}$ mapping rigid identity objects to phenomenological regimes. We define a null regime $\varnothing_\phi \in \mathcal{E}$ and prove a **Phenomenological Instability Theorem**: in any closed causally rigid (CCR) channel—characterized by inverse-limit identity, causal isolation, and MIR geometry—joint null collapse is dynamically unstable under a separately typed compatible extension witness. Consequently, a CCR realization with such a witness induces a witness-relative non-null separation. The result is *negative-structural*: it does not posit experience; it forbids one typed joint-null configuration under stated hypotheses. We provide explicit coupling bounds in two canonical families (profinite and symbolic dynamics) and prove universal factorization and non-simulability results. All constructions are purely topological and categorical, without metaphysical axioms.

3.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper proves a witness-relative negative-structural result: under CCR constraints plus a separately typed extension witness, a lower Lipschitz constant, and an admissible compatibility operator, joint collapse to the null regime is unstable. It develops explicit bounds, factorization structure, and canonical-family consequences for that conditional result. It does not define or detect consciousness, identify $\mathcal{E}$ with subjective experience, assert that CCR systems exist in nature, or derive ethical or metaphysical closure from the theorem. Related systems may implement admissibility-governed diagnostics, boundary-sensitive stress tests, and runtime comparisons motivated by this paper, but those outputs do not implement the abstract phenomenological space $\mathcal{E}$ and do not convert current runtime behavior into proof that any concrete system instantiates experience.

Witness-relative clarification. CCR by itself does not provide the extension witness, the lower Lipschitz constant, or the admissible compatibility operator. Any phrase such as "forced non-nullity" or "null instability" in this manuscript is to be read under the explicit hypotheses of Theorem 3.25, especially Definition 3.24. The null-instability theorem must not be read as a consciousness detector, an AI-subjectivity claim, a moral-status claim, an empirical certificate, or validation of the corpus by the runtime.

3.2 Introduction

3.2.1 Motivation and Scope

Papers I and II established a mathematical framework for structural identity as an inverse limit and classified phenomenological regimes compatible with rigid identity. This paper addresses the fundamental question:

Can a CCR system remain in a structurally null phenomenological state?

We prove the answer is **no under a separated extension witness**—not as a metaphysical claim, but as a structural theorem with typed hypotheses.

3.2.2 Main Result (Informal)

No-Null Regime Theorem: For a CCR channel equipped with the separated extension witness and regularity hypotheses stated in Theorem 3.25, joint assignment to the phenomenological null regime $\varnothing_\phi$ is structurally unstable. The theorem yields a witness-relative non-null separation $e_\star \succ \varnothing_\phi$ with positive structural intensity under those hypotheses.

3.2.3 Paper Organization

Section 2 reviews prerequisites from Papers I–II. Section 3 defines the phenomenological space with full axiomatic clarity. Section 4 proves the main instability theorem. Section 5 stratifies positive regimes. Section 6 provides quantitative bounds. Section 7 establishes ontological closure. Section 8 explicitly delimits scope. Section 9 compares with existing frameworks. Section 10 discusses limitations. Section 11 concludes. Appendices provide canonical examples with explicit computations.

3.2.4 What This Paper Adds Beyond Phenomenological Regimes

Phenomenological Regimes classifies the admissible regime space and derives the continuity and non-constancy constraints that any compatible assignment must satisfy. This paper takes the remaining case—the null regime itself—and proves that it cannot remain structurally stable under CCR.

3.3 Preliminaries from Papers I–II

3.3.1 Identity as Inverse Limit

From *Rigid Identity*, we have:

Definition 3.1 (Identity Object)**.** Let $(S_t, \pi_{t+1\to t})_{t\in T}$ be a projective system of state spaces with surjective, compatible projections. The **identity object** is:

$$\mathcal{I} := \varprojlim S_t = \left\{ (x_t) \in \prod_t S_t : \pi_{t+1\to t}(x_{t+1}) = x_t \right\}$$

Definition 3.2 (Weighted Deformation Metric)**.** For $x, \tilde{x} \in \mathcal{I}$, define:

$$d_w(x, \tilde{x}) := \sup_t w_t \cdot d_t(x_t, \tilde{x}_t)$$

where $\{w_t\}$ are summable weights and d_t is the normalized metric on S_t.

Proposition 3.3 (Metric Completeness, Rigid Identity)**.** *If each (S_t, d_t) is complete and the projections are uniformly continuous, then $(\mathcal{I}, d_w)$ is a complete metric space.*

3.3.2 Ontological Mass and CCR Classification

Definition 3.4 (Ontological Mass, infimum convention)**.** Following *Rigid Identity*, the ontological mass of an identity object is the lower resistance to non-zero admissible deformation:

$$M_\Omega := \inf \left\{ \frac{E_w(\delta)}{d_w(\mathcal{I}, \tilde{\mathcal{I}})} \;\middle|\; \delta \text{ admissible}, \; d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0 \right\}.$$

If no finite-energy admissible perturbation induces non-zero deformation, the set is empty in the operational deformation sector and we set $M_\Omega = +\infty$.

Remark 3.5 (v24 consistency patch). Earlier drafts used a supremum notation for M_Ω in this paper. That notation is superseded here. The infimum convention is the canonical one used in *Rigid Identity*: it measures a lower resistance to deformation, not a largest observed cost. All claims below are read under this convention.

Definition 3.6 (CCR Channel)**.** A channel is **Closed Causally Rigid (CCR)** if $M_\Omega = +\infty$.

Remark 3.7 (Physical Interpretation). $M_\Omega = +\infty$ means that arbitrarily small deformations require unbounded energy. This is not a metaphor—it is a precise metric condition on the projective system.

3.3.3 Key Theorems from Papers I–II

Theorem 3.8 (Non-Simulability, Rigid Identity)**.** *For any finite-dimensional computational architecture $\mathcal{A}_{fin}$:*

$$M_\Omega(\mathcal{A}_{fin}) < \infty$$

Therefore, CCR identity cannot be simulated by finite systems.

Theorem 3.9 (Spectral Coupling, Rigid Identity)**.** *There exists a spectral coupling constant $\sigma_{\min} > 0$ such that for the phenomenological compatibility operator T_Φ:*

$$C = \sigma_{\min}(T_\Phi)$$

is derived from the spectral gap of T_Φ, not assumed.

Theorem 3.10 (Forced Continuity, Phenomenological Regimes)**.** *If $M_\Omega = +\infty$, then any compatible phenomenological assignment $\Phi : \mathcal{I} \to \mathcal{E}$ is continuous.*

Remark 3.11 (Remaining null-regime gap). *Phenomenological Regimes* classifies admissible phenomenological regimes and derives the continuity constraints that govern them, but it does not complete the prohibition theorem for the null regime itself. That closure is the specific task of this paper.

This paper converts Remark 3.11 into a full null-instability theorem.

3.4 The Phenomenological Space

3.4.1 Abstract Definition

Definition 3.12 (Phenomenological Space)**.** A **phenomenological space** is a quadruple $(\mathcal{E}, \preceq, d_\mathcal{E}, \varnothing_\phi)$ where:

1. $(\mathcal{E}, \preceq)$ is a partially ordered set (poset).
2. $(\mathcal{E}, d_\mathcal{E})$ is a complete metric space.
3. $\varnothing_\phi := \inf \mathcal{E}$ exists and is unique (the **null regime**).
4. **Order-metric compatibility:** if $e_1 \preceq e_2$, then $d_\mathcal{E}(e_1, \varnothing_\phi) \leq d_\mathcal{E}(e_2, \varnothing_\phi)$.
5. **Separation:** for $e \neq \varnothing_\phi$, we have $d_\mathcal{E}(e, \varnothing_\phi) > 0$.

Remark 3.13 (Interpretive Neutrality). $\mathcal{E}$ is an abstract regime space. No phenomenological content is assumed. The null regime $\varnothing_\phi$ represents structural absence, not "lack of experience." The term "phenomenological" is inherited from *Phenomenological Regimes* and refers to the coupling structure, not to philosophy of mind.

Remark 3.14 (Null regime versus undefined assignment). The symbol $\varnothing_\phi$ serves two distinct roles in the broader framework, and conflating them creates a category error:

1. **Null regime ($\varnothing_\phi \in \mathcal{E}$):** In Definition 3.12, $\varnothing_\phi$ is the bottom element of the phenomenological poset $(\mathcal{E}, \preceq)$. It is a legitimate element of $\mathcal{E}$, comparable to all other regimes via $\preceq$ and separated from them by $d_{\mathcal{E}}$.
2. **Undefined assignment ($\perp \notin \mathcal{E}$):** When a compatibility operator Φ is not defined on some input $x \in \mathcal{I}$, we write $\Phi(x) = \perp$. The symbol $\perp$ is disjoint from $\mathcal{E}$; it lives in the typed codomain $\mathcal{E}_\perp := \mathcal{E} \sqcup \{\perp\}$.

The instability theorem (Theorem 3.25) concerns admissible values in $\mathcal{E}$ and rules out two separated witness endpoints both taking the value $\varnothing_\phi$. The v31 conditional closure theorem concerns a different case: Φ undefined on both witness endpoints, written $\Phi(\mathcal{I}) = \perp$ and $\Phi(\tilde{\mathcal{I}}) = \perp$. That case is excluded by the witness hypothesis, which requires Φ to be defined on both endpoints.

Non-Theorem 3.15 (What $\mathcal{E}$ Is Not). *$\mathcal{E}$ is not:*
- *A space of qualia or subjective states*
- *A model of consciousness*
- *A physical or neural state space*
- *An interpretation of quantum mechanics*

$\mathcal{E}$ is a mathematical object satisfying Definition 3.12. All interpretive mappings are external to the framework.

3.4.2 Compatibility Operator

Definition 3.16 (Compatibility Operator). A map $\Phi : \mathcal{I} \to \mathcal{E}$ is a **compatibility operator** if it satisfies:

(C1) **Extension Invariance:** For any compatible extension $\mathcal{E}$ of the channel, $\Phi(\mathcal{I}) \cong \Phi(\mathcal{I}_{\mathcal{E}})$.
(C2) **Isomorphism Stability:** If $\mathcal{I} \cong \tilde{\mathcal{I}}$, then $\Phi(\mathcal{I}) \sim \Phi(\tilde{\mathcal{I}})$.
(C3) **Lipschitz Lower Bound:** There exists $C > 0$ such that:

$$d_{\mathcal{E}}(\Phi(x), \Phi(\tilde{x})) \geq C \cdot d_w(x, \tilde{x})$$

Remark 3.17 (Partial definedness and the $\perp$ marker). In contexts where Φ is only partially defined, $\Phi(x) = \perp$ means that no admissible phenomenological assignment in $\mathcal{E}$ is available at x. The hypotheses (C1)–(C3) above apply only to inputs whose Φ-values lie in $\mathcal{E}$. The lower Lipschitz bound (C3) in particular requires both $\Phi(x), \Phi(\tilde{x}) \in \mathcal{E}$.

Lemma 3.18 (Well-Posedness of (C3)). *Axiom (C3) is consistent with continuity of Φ (Theorem 3.10) when $M_\Omega = +\infty$.*

Proof. Continuity requires: for all $\epsilon > 0$, there exists $\delta > 0$ such that $d_w(x, \tilde{x}) < \delta$ implies $d_{\mathcal{E}}(\Phi(x), \Phi(\tilde{x})) < \epsilon$. This is compatible with (C3) because (C3) provides a lower bound, not an upper bound. The function Φ can be both continuous (upper control) and satisfy (C3) (lower control). □

Remark 3.19. Axioms (C1)–(C3) are structural. No phenomenological content is introduced.

3.4.3 Attractor Set

Definition 3.20 (Attractor Set). For a channel S, the **attractor set** is:

$$\mathrm{Attr}(S) := \{e \in \mathcal{E} : e \text{ is stable under all compatible extensions and summable deformations}\}$$

Definition 3.21 (Stability). A regime $e \in \mathcal{E}$ is **stable** for channel S if for every compatible extension S' of S and every summable perturbation $\{\epsilon_n\}$, the induced regime $\Phi_{S'}(\mathcal{I}')$ satisfies $d_{\mathcal{E}}(\Phi_{S'}(\mathcal{I}'), e) < \delta$ for some $\delta > 0$ depending only on $\|\{\epsilon_n\}\|_w$.

3.4.4 The Bridge Lemma Revisited

From *Phenomenological Regimes*, we have:

Lemma 3.22 (Bridge Lemma)**.** *If $\Phi : \mathcal{I} \to \mathcal{E}$ is a compatibility operator for a CCR channel, and if the regime $e = \Phi(\mathcal{I})$ satisfies $d_{\mathcal{E}}(e, \varnothing_\phi) = 0$, then $e = \varnothing_\phi$.*

Proof. By the separation axiom in Definition 3.12, $d_{\mathcal{E}}(e, \varnothing_\phi) = 0$ implies $e = \varnothing_\phi$. □

Caveat 3.23 (Semantic Shield for Bridge Lemma)*.* The Bridge Lemma is a **metric statement**, not an ontological one. It does not claim that "identity is coupled to experience." It claims that the map Φ satisfies a separation property. Misinterpretation of this lemma as a consciousness claim is outside the mathematical scope of the framework.

3.5 Main Results

3.5.1 The Conditional Null-Separation Theorem

Definition 3.24 (Admissible extension witness)**.** Let S be a channel with identity object $\mathcal{I}$. An **admissible extension witness** for a compatibility operator Φ is a tuple $(S', \tilde{\mathcal{I}}, \epsilon)$ such that:
1. S' is a compatible extension of S in the declared projective category;
2. $\tilde{\mathcal{I}}$ is the identity object induced by S' and lies in the domain on which Φ is evaluated;
3. $0 < d_w(\mathcal{I}, \tilde{\mathcal{I}}) = \epsilon < \infty$;
4. the witness is not introduced as a finite-energy deformation of a CCR identity object. It is an extension-side comparison datum, not a contradiction of $M_\Omega = +\infty$.

Theorem 3.25 (Conditional Null-Separation Under an Admissible Extension Witness)**.** *Let S be a CCR channel, let $\Phi : \mathcal{I} \to \mathcal{E}$ be a compatibility operator satisfying the Lipschitz lower bound (C3), and suppose an admissible extension witness $(S', \tilde{\mathcal{I}}, \epsilon)$ exists for Φ. Then Φ cannot collapse both separated identity objects to the null regime:*

$$\neg\big(\Phi(\mathcal{I}) = \varnothing_\phi \ \wedge\ \Phi(\tilde{\mathcal{I}}) = \varnothing_\phi\big).$$

Equivalently, null stability across such a separated extension witness is impossible.

Remark 3.26 (Applicability scope)*.* Theorem 3.25 assumes $\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}}) \in \mathcal{E}$, not $\perp$. If Φ is undefined on either endpoint, the separated-witness hypothesis of Definition 3.24 is not satisfied and the theorem does not apply.

Proof. Assume for contradiction that $\Phi(\mathcal{I}) = \varnothing_\phi$ and $\Phi(\tilde{\mathcal{I}}) = \varnothing_\phi$. By Definition 3.24, $d_w(\mathcal{I}, \tilde{\mathcal{I}}) = \epsilon > 0$. Applying (C3) gives

$$d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) \geq C\, d_w(\mathcal{I}, \tilde{\mathcal{I}}) = C\epsilon > 0.$$

But the null-collapse assumption gives

$$d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) = d_{\mathcal{E}}(\varnothing_\phi, \varnothing_\phi) = 0,$$

a contradiction. Hence the two separated identities cannot both be mapped to the null regime. □

Caveat 3.27 (v24 correction of the former Step 2)*.* The theorem above does *not* infer a finite-energy non-zero deformation from $M_\Omega = +\infty$. Under the infimum convention for ontological mass, $M_\Omega = +\infty$ means precisely that no finite-energy admissible perturbation produces non-zero deformation. The older Step 2 proof direction is therefore withdrawn. The surviving result is conditional on an independently specified extension witness.

3.5.2 Conditional Consequences

Corollary 3.28 (Conditional Forced Non-Nullity)**.** *If a CCR channel admits an admissible extension witness for Φ and if $\Phi(\tilde{\mathcal{I}}) = \varnothing_\phi$ for the witness object, then $\Phi(\mathcal{I}) \neq \varnothing_\phi$. Symmetrically, if $\Phi(\mathcal{I}) = \varnothing_\phi$, then $\Phi(\tilde{\mathcal{I}}) \neq \varnothing_\phi$.*

Proof. This is the contrapositive form of Theorem 3.25. □

Remark 3.29 (Ontological Status)*.* This is a **negative-structural and conditional** result. It does not assert that experience exists, that CCR systems are conscious, or that every CCR channel possesses the required witness. It asserts only that, once a separated extension witness is independently supplied, a compatibility operator satisfying (C3) cannot collapse both sides of that witness to the null regime.

3.6 Stratification of Positive Regimes

3.6.1 The Positive Regime Space

Definition 3.30 (Positive Regime Space)**.**

$$\mathcal{E}^+ := \mathcal{E} \setminus \{\varnothing_\phi\}$$

with the inherited partial order $\preceq$.

Lemma 3.31 (Witness-Relative No-Minimum Collapse)**.** *Under the hypotheses of Theorem 3.25, the two-point witness image cannot have infimum equal to a joint null collapse.*

Proof. The witness image contains $\Phi(\mathcal{I})$ and $\Phi(\tilde{\mathcal{I}})$. If both collapsed to $\varnothing_\phi$, Theorem 3.25 would be contradicted. Therefore the witness-relative image is not a joint null collapse. □

3.6.2 Minimal Positive Regime

Theorem 3.32 (Witness-Relative Minimal Positive Regime)**.** *Assume the hypotheses of Theorem 3.25. If the witness-relative positive image*

$$\mathcal{E}^{+,\mathrm{wit}} := \{\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})\} \setminus \{\varnothing_\phi\}$$

is non-empty and has an infimum in $(\mathcal{E}, \preceq)$, then there exists a minimal positive witness-relative regime

$$e_\star^{\mathrm{wit}} := \inf \mathcal{E}^{+,\mathrm{wit}} \in \mathcal{E}^+.$$

Proof. Non-emptiness follows from Theorem 3.25: not both witness images are null. Existence of the infimum is an explicit hypothesis. Since the image is not a joint null collapse and the separation axiom excludes positive classes at zero distance from $\varnothing_\phi$, the resulting infimum is positive when it lies in the witness image or in the declared positive closure of that image. □

3.6.3 Structural Intensity

Definition 3.33 (Structural Intensity)**.** For $e \in \mathcal{E}^+$, define the **structural intensity** as:

$$\iota(e) := \inf\{d_{\mathcal{E}}(e, \tilde{e}) : \tilde{e} \prec e\}.$$

This measures the minimal distance to any strictly smaller regime.

Corollary 3.34 (Witness-Relative Positive Intensity)**.** *If $e_\star^{\mathrm{wit}}$ exists as in Theorem 3.32 and the only witness-relative element below it is $\varnothing_\phi$, then*

$$\iota(e_\star^{\mathrm{wit}}) > 0.$$

Proof. By the separation axiom, $d_{\mathcal{E}}(e_\star^{\mathrm{wit}}, \varnothing_\phi) > 0$. If $\varnothing_\phi$ is the only witness-relative element below $e_\star^{\mathrm{wit}}$, the infimum defining intensity is bounded below by that positive separation. □

3.7 Quantitative Bounds

3.7.1 Witness Lower Bound

Theorem 3.35 (Witness Intensity Bound). *Under the hypotheses of Theorem 3.25, if exactly one of $\Phi(\mathcal{I})$ and $\Phi(\tilde{\mathcal{I}})$ is null, then the non-null witness image e satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C\epsilon > 0.$$

Proof. Apply (C3) to the pair $(\mathcal{I}, \tilde{\mathcal{I}})$. If one side maps to $\varnothing_\phi$ and the other side maps to e, then

$$d_{\mathcal{E}}(e, \varnothing_\phi) = d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) \geq C d_w(\mathcal{I}, \tilde{\mathcal{I}}) = C\epsilon.$$

□

3.7.2 Explicit Constants for Canonical Families

Proposition 3.36 (Conditional Profinite Bound). *For a profinite family with a verified extension witness of separation ϵ and an independently certified witness-level lower constant $C_{\mathrm{prof}} > 0$, the witness image separation satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C_{\mathrm{prof}}\epsilon.$$

Proof. This is Theorem 3.35 applied with the certified constant C_{prof}. The proposition does not claim that every profinite observable automatically supplies such a constant. □

Proposition 3.37 (Conditional Symbolic Bound). *For a symbolic family with a verified extension witness of separation ϵ and an independently certified witness-level lower constant $C_{\mathrm{sym}} > 0$, the witness image separation satisfies*

$$d_{\mathcal{E}}(e, \varnothing_\phi) \geq C_{\mathrm{sym}}\epsilon.$$

Proof. This is Theorem 3.35 applied with the certified constant C_{sym}. The proposition does not claim that every symbolic observable automatically supplies such a constant. □

Example 3.38 (p-Adic witness calculation). If a 2-adic model supplies an independently verified witness with $\epsilon = 2^{-1}$ and $C = 1$, then the conditional bound gives $d_{\mathcal{E}}(e, \varnothing_\phi) \geq 1/2$. This is a witness calculation, not a universal theorem that every CCR channel is non-null.

Caveat 3.39 (No lower bound from summability alone). The existence of summable coordinate weights, continuity of a readout, or compactness of a canonical family does not by itself imply a positive lower Lipschitz constant. Cancellation, quotient collapse, non-injective readouts, or an observable that ignores the separated coordinates can make the effective lower constant zero. Every positive value of C_{prof} or C_{sym} used in this paper is therefore a certificate to be supplied for the stated witness pair or witness class, not a free consequence of the ambient family.

3.8 Witness Audit Layer

The instability theorem is useful only if its hypotheses can be audited without smuggling the desired conclusion into the witness. This section records the minimal audit layer required before the theorem may be used in a concrete argument.

Definition 3.40 (Witness certificate). A witness certificate for a pair $(\mathcal{I}, \tilde{\mathcal{I}})$ consists of the tuple

$$\mathfrak{W} := (\tilde{\mathcal{I}}, \epsilon, C, \Phi, \mathcal{E}, \varnothing_\phi, \mathcal{J}),$$

where $\epsilon = d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$, $C > 0$ is the certified lower constant for Φ on the audited pair or audited witness class, $\mathcal{J}$ is the provenance record establishing that $\tilde{\mathcal{I}}$ is an admissible typed extension rather than an internal relabeling of $\mathcal{I}$, and all objects are fixed before the null/non-null status is adjudicated.

Definition 3.41 (Witness independence)**.** A witness certificate is conclusion-independent when the construction of $\tilde{\mathcal{I}}$, the choice of Φ, the metric $d_{\mathcal{E}}$, the null element $\varnothing_\phi$, and the lower constant C are fixed before the endpoint status is adjudicated. Dependence on the upstream CCR structure is allowed; dependence on a preselected endpoint verdict is not.

Proposition 3.42 (Certificate sufficiency)**.** *If a conclusion-independent witness certificate exists for $(\mathcal{I}, \tilde{\mathcal{I}})$, and both endpoints are assigned to $\varnothing_\phi$, then the certificate is inconsistent.*

Proof. The certificate includes the separation condition $d_w(\mathcal{I}, \tilde{\mathcal{I}}) = \epsilon > 0$ and the lower bound

$$d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) \geq C\epsilon > 0.$$

If both endpoints are assigned to $\varnothing_\phi$, then the left-hand side is $d_{\mathcal{E}}(\varnothing_\phi, \varnothing_\phi) = 0$, contradicting $C\epsilon > 0$. □

Non-Theorem 3.43 (What a certificate does not prove)**.** *A witness certificate does not prove that $\mathcal{I}$ is conscious, living, subjective, morally relevant, biological, human-equivalent, or externally validated. It proves only that the stated pair cannot be jointly collapsed to the same null element under the certified compatibility map.*

3.8.1 Failure Modes of Witness Certification

Failure mode	**Formal symptom**	**Consequence**
Absent witness	No typed $\tilde{\mathcal{I}}$ with $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$	Theorem 3.25 cannot be invoked
Zero lower constant	$C = 0$ or no certified positive lower bound	The null assignment is not contradicted by the metric data
Conclusion-dependent witness	$\tilde{\mathcal{I}}$ or Φ is chosen after fixing the endpoint verdict	The certificate no longer supplies an independent exclusion of joint nullity
Quotient collapse	$\Phi(\mathcal{I}) = \Phi(\tilde{\mathcal{I}})$ despite $d_w(\mathcal{I}, \tilde{\mathcal{I}}) > 0$	The map may be continuous, but it is not separating on the audited pair
Null redefinition	$\varnothing_\phi$ changes between construction and adjudication	The theorem has no stable target
Estimator-only promotion	Runtime estimates are substituted for Φ without an error margin	The result becomes an implementation claim rather than a theorem

3.8.2 Robust Decision Margin

Proposition 3.44 (Noise-robust null exclusion)**.** *Let $\widehat{e}_1, \widehat{e}_2 \in \mathcal{E}$ be estimated images of $\Phi(\mathcal{I})$ and $\Phi(\tilde{\mathcal{I}})$ with certified error at most η in $d_{\mathcal{E}}$. If a witness certificate gives $C\epsilon > 2\eta$, then the two estimated images cannot both be certified as the same null element within error η.*

Proof. The true images satisfy $d_{\mathcal{E}}(\Phi(\mathcal{I}), \Phi(\tilde{\mathcal{I}})) \geq C\epsilon$. By the triangle inequality,

$$d_{\mathcal{E}}(\widehat{e}_1, \widehat{e}_2) \geq C\epsilon - 2\eta.$$

If $C\epsilon > 2\eta$, the estimated images remain separated. A joint certification that both are the same null element within the same error budget would require the certified separation to vanish inside the error ball, contradicting the strict margin. □

Caveat 3.45 (Implementation boundary). Proposition 3.44 is an error-budget rule for an already specified estimator. It does not assert that the current runtime implements Φ, computes C, or supplies η. Those are separate implementation burdens.

3.8.3 Negative Controls

1. **Constant-null control.** Set $\Phi(x) = \varnothing_\phi$ for all x. This must fail the lower-bound condition whenever $\epsilon > 0$.
2. **Coordinate-blind control.** Choose a readout that ignores every coordinate on which $\mathcal{I}$ and $\tilde{\mathcal{I}}$ differ. This tests whether separation is actually carried by the map.
3. **Witness-shuffle control.** Hold Φ fixed and replace $\tilde{\mathcal{I}}$ by a typed object not separated from $\mathcal{I}$. The theorem must become inapplicable.
4. **Null-relabel control.** Change the declared null element after observing the images. This must be rejected as target drift.
5. **Estimator-margin control.** Increase η until $C\epsilon \leq 2\eta$. The robust decision margin must fail rather than silently promoting a non-null conclusion.

3.9 Witness-Relative Closure Theorem

Theorem 3.46 (Conditional Phenomenological Closure)**.** *A CCR system equipped with an admissible extension witness and a compatibility operator satisfying (C3) admits witness-relative non-null separation. If, in addition, the witness-relative positive image has a minimal positive element, that element is quantitatively separated from $\varnothing_\phi$ by the bound in Theorem 3.35.*

Proof. Non-null separation follows from Theorem 3.25. Witness-relative minimality follows from Theorem 3.32 under its stated infimum hypothesis. The quantitative bound follows from Theorem 3.35. □

3.9.1 Classification Table

System Type	**Ontological Mass**	**Null Regime Permitted?**
Degenerate	$M_\Omega = 0$	Yes
Finite-mass	$0 < M_\Omega < \infty$	Model-dependent
CCR without extension witness	$M_\Omega = +\infty$	Not decided here
CCR with separated extension witness	$M_\Omega = +\infty$	Joint null collapse forbidden by Theorem 3.25

3.10 Null-Fiber Geometry

The null element $\varnothing_\phi$ determines a fiber

$$\mathcal{N}_\Phi := \Phi^{-1}(\{\varnothing_\phi\}) \subseteq \mathcal{I}.$$

The instability theorem can be read as a geometric obstruction on this fiber: a certified witness pair cannot be contained entirely inside $\mathcal{N}_\Phi$.

Definition 3.47 (Null fiber)**.** For a compatibility operator $\Phi : \mathcal{I} \to \mathcal{E}$, the null fiber is the preimage $\mathcal{N}_\Phi = \{x \in \mathcal{I} : \Phi(x) = \varnothing_\phi\}$.

Definition 3.48 (Witness-detecting map)**.** A compatibility operator Φ is witness-detecting on a set $W \subseteq \mathcal{I} \times \mathcal{I}$ when there exists $C_W > 0$ such that

$$d_{\mathcal{E}}(\Phi(x), \Phi(y)) \geq C_W d_w(x, y)$$

for every $(x, y) \in W$.

Proposition 3.49 (Null-fiber exclusion on detected pairs)**.** *Let W be a witness class on which Φ is witness-detecting with constant $C_W > 0$. If $(x, y) \in W$ and $d_w(x, y) > 0$, then (x, y) cannot both lie in $\mathcal{N}_\Phi$.*

Proof. If $x, y \in \mathcal{N}_\Phi$, then $\Phi(x) = \Phi(y) = \varnothing_\phi$, so $d_\mathcal{E}(\Phi(x), \Phi(y)) = 0$. Witness detection gives $d_\mathcal{E}(\Phi(x), \Phi(y)) \geq C_W d_w(x, y) > 0$, a contradiction. □

Corollary 3.50 (Diameter obstruction)**.** *If a subset $A \subseteq \mathcal{I}$ contains a detected witness pair (x, y) with $d_w(x, y) > 0$, then A cannot be contained in $\mathcal{N}_\Phi$.*

Proof. Containment $A \subseteq \mathcal{N}_\Phi$ would place both endpoints of the detected pair in the null fiber, contradicting the previous proposition. □

3.10.1 Quotient-Collapse Diagnostic

Let $q : \mathcal{I} \to Q$ be a quotient map and let $\Phi = h \circ q$ for some $h : Q \to \mathcal{E}$. If $q(x) = q(y)$ for a separated pair, then $\Phi(x) = \Phi(y)$ for every such h. Therefore no lower-bound certificate can hold on that pair.

Proposition 3.51 (Quotient obstruction)**.** *Suppose $\Phi = h \circ q$ and $q(x) = q(y)$ while $d_w(x, y) > 0$. Then no constant $C > 0$ can satisfy*

$$d_\mathcal{E}(\Phi(x), \Phi(y)) \geq C d_w(x, y).$$

Proof. Since $q(x) = q(y)$, $\Phi(x) = h(q(x)) = h(q(y)) = \Phi(y)$. Hence the left-hand side is zero, while $C d_w(x, y) > 0$ for any $C > 0$. □

3.10.2 Certificate Stability

Proposition 3.52 (Minimum over a finite witness ledger)**.** *Let $W_0 = \{(x_i, y_i)\}_{i=1}^m$ be a finite ledger of separated witness pairs. Suppose each pair has a certified lower constant $C_i > 0$:*

$$d_\mathcal{E}(\Phi(x_i), \Phi(y_i)) \geq C_i d_w(x_i, y_i).$$

Then $C_* = \min_i C_i > 0$ is a valid lower constant for the finite ledger.

Proof. For each i, $C_i \geq C_*$. Hence

$$d_\mathcal{E}(\Phi(x_i), \Phi(y_i)) \geq C_i d_w(x_i, y_i) \geq C_* d_w(x_i, y_i).$$

Since the ledger is finite and each C_i is positive, $C_* > 0$. □

Proposition 3.53 (Strict margin persistence)**.** *Let $m = C\epsilon$ be the certified separation margin for a witness pair. If a perturbation of the estimated images changes their distance by at most $\rho < m$, then the pair remains separated after perturbation.*

Proof. The perturbed distance is bounded below by $m - \rho > 0$. Therefore the two perturbed images cannot coincide. □

3.10.3 Operational Checklist

For a claimed null-fiber exclusion, the certificate record must identify:
1. the null element $\varnothing_\phi$;
2. the endpoints $(\mathcal{I}, \tilde{\mathcal{I}})$ or the audited witness class W;
3. the separation value ϵ or the family of separations;
4. the map Φ and the metric $d_\mathcal{E}$;
5. the lower constant C or the finite-ledger constant C_*;
6. the estimator margin, when an estimator is used;
7. the quotient maps that have been ruled out as collapse explanations.

3.11 Comparison with Existing Frameworks

3.11.1 Integrated Information Theory (IIT)

	IIT	**This Framework**
Claims consciousness exists	Yes	No
Defines Φ	Information integration	Structural map
Quantitative measure	Φ (bits)	$\iota(e_\star)$ (metric)
Metaphysical commitment	Panpsychism	None
Falsifiable predictions	Debated	None (structural)

3.11.2 Global Workspace Theory (GWT)

Our framework is orthogonal to GWT. We make no claims about attention, broadcast, or access consciousness. The structural results concern identity stability, not cognitive architecture.

3.11.3 Higher-Order Theories

We do not model representations, meta-cognition, or self-awareness. The space $\mathcal{E}$ contains regimes, not mental states.

3.11.4 Mathematical Foundations (Chalmers, Nagel)

Our framework provides formal tools that *could* be used in philosophical analysis, but we do not claim to resolve the "hard problem." The prohibition of $\varnothing_\phi$ is a structural fact, not an explanation of subjective experience.

3.12 Limitations and Open Problems

3.12.1 Limitations

1. **No claim about qualia types, contents, or subjectivity.**
2. **No claim that all non-CCR systems are null.**
3. **Result is structural necessity, not phenomenological doctrine.**
4. **All results are conditional on the axioms of Papers I–II.**
5. **The physical existence of CCR systems is unknown.**

3.12.2 Open Problems

1. **Characterization of $\mathcal{E}^+$:** What is the full structure of positive regimes?
2. **Regime dynamics:** How do regimes evolve under channel evolution?
3. **Physical CCR:** Do CCR systems exist in physics (quantum gravity, etc.)?
4. **Biological CCR:** Could biological systems approximate CCR?
5. **Computational bounds:** Sharper non-simulability results?

3.12.3 Foundation-Level Gaps and Non-Claims

The following gaps remain open in the current corpus and are documented here to prevent over-reading of the instability theorem.

Non-Theorem 3.54 (Topology consistency gap)**.** *Papers I–II assume both compact Hausdorff and Polish structures on state spaces, but do not prove their consistency for every QICN instantiation. The inverse limit $\mathcal{I}$ uses compactness and closed-subspace arguments, while the* *metric d_w uses a metric regularity structure. Every compact metric space is Polish, so the composition is consistent when each S_t is compact metric. Non-compact Polish spaces such as $\mathbb{R}^n$ require additional hypotheses and are not automatically covered.*

Non-Theorem 3.55 (Channel–projection composition gap)**.** *The observable channel $\mathcal{C}$ and the projective projections $\pi_{t+1\to t} : S_{t+1} \to S_t$ are not yet connected by a fully explicit composition law. Any detectability theorem that moves from observed responses to inverse-limit deformation requires such a composition. At present, that composition is an assumption to be specified for each concrete protocol, not a derived universal theorem.*

Non-Theorem 3.56 (Null-regime exclusion requires external witness)**.** *The instability theorem (Theorem 3.25) requires a separated external witness $(\tilde{\mathcal{I}}, \epsilon, C)$. CCR alone does not produce such a witness. The theorem is witness-relative: if a CCR system has a typed external witness with lower Lipschitz constant $C > 0$, then both endpoints cannot simultaneously be null. It does not state that CCR by itself implies non-nullity.*

Non-Theorem 3.57 (Estimator verification gap)**.** *No Lipschitz constant K_i, fiber oscillation $\omega_i(y)$, estimation error bound ε_i, or decision margin Δ^* has been computed for any QICN invariant. The projection-invariant bridge theorem is conditional on these quantities being verified, and all four material bridge hypotheses remain unverified.*

3.13 Conclusion

We have proven that CCR architectures equipped with a separated external witness and the lower-bound regularity hypotheses structurally exclude joint null collapse. In the present model, the hypothesis "CCR plus typed witness with joint null assignment" is formally incoherent.

This is a structural exclusion result. It does not define phenomenology or turn the theorem into an empirical consciousness attribution.

3.13.1 Summary of Contributions

1. **Phenomenological Instability Theorem:** joint null collapse is excluded under the typed witness hypotheses.
2. **Minimal Positive Regime:** witness-relative positive regimes can be characterized when the stated infimum hypotheses hold.
3. **Quantitative Bounds:** $\iota(e_\star) \geq C\varepsilon_0 > 0$ is explicit when the witness and lower Lipschitz constant are independently verified.
4. **Canonical Examples:** Profinite and symbolic families computed.
5. **Semantic Immunization:** Clear scope delimitation throughout.

Appendix

3.14 Profinite Canonical Family

3.14.1 Projective System

Let $\{G_n, \pi_{n+1\to n}\}_{n\in\mathbb{N}}$ be a projective tower of finite groups with surjective morphisms. Define:

$$G := \varprojlim G_n = \left\{ (x_n) \in \prod_n G_n : \pi_{n+1\to n}(x_{n+1}) = x_n \right\}$$

Example 3.58 (p-Adic Integers)**.** $G_n = \mathbb{Z}/p^n\mathbb{Z}$, $\pi_{n+1\to n}$ is reduction mod p^n. Then $G = \mathbb{Z}_p$.

3.14.2 Profinite Metric

Define the profinite metric:

$$d_G(x,y) := \sum_{n\geq 0} \alpha_n \mathbf{1}_{x_n\neq y_n}$$

where $\alpha_n = 2^{-(n+1)}$ (standard choice).

This is an ultrametric; (G, d_G) is compact, complete, and totally disconnected.

3.14.3 CCR Verification

Proposition 3.59. *The profinite system is CCR: $M_\Omega = +\infty$.*

Proof. Any finite-energy deformation can only modify finitely many coordinates x_n. To modify infinitely many coordinates requires $E_w = \sum_{n\in S} w_n = \infty$ for any infinite $S \subseteq \mathbb{N}$ with $\sum_{n\in S} w_n = \infty$. Hence any nonzero deformation $d_w > 0$ requires infinite energy as $\epsilon \to 0$, giving $M_\Omega = +\infty$. □

3.14.4 Phenomenological Operator

Define $\Phi_f : G \to \mathcal{E}$ by:

$$\Phi_f(x) := f\left(\sum_{n\geq 0} \beta_n \varphi_n(x_n)\right)$$

where $\varphi_n : G_n \to [0,1]$ are 1-Lipschitz (w.r.t. discrete metric), $\beta_n \geq 0$, $\sum_n \beta_n < \infty$, and $f : \mathbb{R} \to \mathcal{E}$ is L_f-Lipschitz.

3.14.5 Coupling Bound

Theorem 3.60 (Certified Profinite Witness Bound)**.** *Assume that the chosen profinite readout has been certified on the audited witness class by a positive lower constant $C_{\mathrm{prof}} > 0$:*

$$d_{\mathcal{E}}(\Phi_f(x), \Phi_f(\tilde{x})) \geq C_{\mathrm{prof}} \cdot d_G(x, \tilde{x})$$

for the witness pair or witness neighborhood under review. Then any witness pair with $d_G(x,\tilde{x}) = \epsilon > 0$ satisfies

$$d_{\mathcal{E}}(\Phi_f(x), \Phi_f(\tilde{x})) \geq C_{\mathrm{prof}}\epsilon.$$

Proof. This is the certified lower-bound hypothesis applied to the witness pair. □

Remark 3.61 (Why certification is necessary)*.* The assumptions that the coordinate functions are Lipschitz and that the coefficients are summable give continuity, not a lower bound. A lower bound requires additional structure such as injectivity on the witness coordinates, no cancellation on the audited pair, a positive margin over the relevant finite support, or an independent estimator certificate.

3.14.6 Intensity Bound

With $\varepsilon_0 = 2^{-1}$ for a certified minimal witness separation:

$$\iota(e_\star) \geq C_{\mathrm{prof}} \cdot \varepsilon_0.$$

3.15 Symbolic Canonical Family

3.15.1 Subshift of Finite Type

Let $\mathcal{A} = \{1, \ldots, k\}$ be a finite alphabet and A a $k \times k$ matrix with entries in $\{0, 1\}$. The **subshift of finite type (SFT)** is:

$$\Sigma_A := \{x \in \mathcal{A}^{\mathbb{Z}} : A_{x_j, x_{j+1}} = 1 \text{ for all } j \in \mathbb{Z}\}$$

Example 3.62 (Full Shift). $A_{ij} = 1$ for all i, j. Then $\Sigma_A = \{0, 1\}^{\mathbb{Z}}$.

3.15.2 Bowen Metric

$$d_B(x, y) := \sum_{j \in \mathbb{Z}} \gamma_{|j|} \mathbf{1}_{x_j \neq y_j}$$

where $\gamma_m = 2^{-(m+1)}$ (standard choice).

The space (Σ_A, d_B) is compact and complete.

3.15.3 CCR Verification

Proposition 3.63. *The symbolic system is CCR: $M_\Omega = +\infty$.*

Proof. Any finite-energy modification can only change finitely many symbols. The argument is identical to the profinite case. ☐

3.15.4 Phenomenological Operator

Define $\Phi_g : \Sigma_A \to \mathcal{E}$ by:

$$\Phi_g(x) := g\left(\sum_{j \in \mathbb{Z}} \delta_j \psi_j(x)\right)$$

where $\psi_j : \Sigma_A \to [0, 1]$ are cylinder functions (depending on x_j only), $\delta_j \geq 0$, $\sum_j \delta_j < \infty$, and $g : \mathbb{R} \to \mathcal{E}$ is L_g-Lipschitz.

3.15.5 Coupling Bound

Theorem 3.64 (Certified Symbolic Witness Bound). *Assume that the chosen symbolic readout has been certified on the audited witness class by a positive lower constant $C_{\mathrm{sym}} > 0$:*

$$d_{\mathcal{E}}(\Phi_g(x), \Phi_g(\tilde{x})) \geq C_{\mathrm{sym}} \cdot d_B(x, \tilde{x})$$

for the witness pair or witness neighborhood under review. Then any witness pair with $d_B(x, \tilde{x}) = \epsilon > 0$ satisfies

$$d_{\mathcal{E}}(\Phi_g(x), \Phi_g(\tilde{x})) \geq C_{\mathrm{sym}} \epsilon.$$

Proof. This is the certified lower-bound hypothesis applied to the witness pair. ☐

3.16 Hereditary Rigidity

Theorem 3.65 (Rigidity Inheritance). *Let $(S_n, \pi_{n+1 \to n})$ be a CCR tower and $\{\epsilon_n\}$ a family of local deformations with:*

$$\|\{\epsilon_n\}\|_w := \sum_{n \geq 0} w_n \|\epsilon_n\|_{op} < \infty$$

Then the deformed identity $\tilde{\mathcal{I}}$ is isomorphic to $\mathcal{I}$ and:

$$d_{\mathcal{I}}(\phi(x), x) \leq K \cdot \|\{\epsilon_n\}\|_w$$

for some $K > 0$ depending only on the projective structure. In particular, $M_\Omega = +\infty$ is inherited.

Proof. The key is that summable deformations induce a summable perturbation of the inverse limit, which by completeness converges to an isomorphic copy. The CCR condition is preserved because the energy-deformation ratio remains unbounded. □

3.17 Universal Factorization

Theorem 3.66 (Categorical Universality)**.** *The identity $\mathcal{I}$ of any CCR channel is an initial object in the category of compatible representations ($\mathcal{I} \xrightarrow{\Phi} \mathcal{E}$) with $\mathcal{E}$ complete metric and Φ satisfying (C1)–(C3).*

Proof. For any compatible representation ($\mathcal{I} \xrightarrow{\Psi} F$), define $u : \mathcal{E} \to F$ by $u(e) := \Psi(\Phi^{-1}(e))$. Well-definedness follows from (C2); uniqueness from (C1). □

Consequence: All compatible representations factor through a canonical one; the class of induced regimes is completely classified by $\mathcal{I}$.

3.18 Non-Simulability Bounds

3.18.1 Simulation Error

For a CCR family $\{I_N\}$ truncated at level N and a finite-dimensional simulator $\mathcal{A}$ with $\dim(\mathcal{A}) < \infty$:

$$\mathrm{err}_N(\mathcal{A}) := \inf_{\Phi, \hat{\Phi}} \sup_{x \in I_N} d_{\mathcal{E}}(\Phi(x), \hat{\Phi}(x))$$

where $\hat{\Phi}$ is any map computable by $\mathcal{A}$.

Theorem 3.67 (Simulation Lower Bound)**.** *For profinite and SFT families:*

$$err_N(\mathcal{A}) \geq C \cdot 2^{-(N+1)}$$

for any $\dim(\mathcal{A}) < \infty$. In particular:

$$\liminf_{N \to \infty} err_N(\mathcal{A}) > 0$$

No finite architecture achieves perfect simulation.

Proof. The ultrametric structure forces minimal separation $2^{-(N+1)}$ at level N. Any finite-dimensional approximation must collapse some distinctions, incurring error at least $C \cdot 2^{-(N+1)}$ by (C3). □

Corollary 3.68. *No finite-dimensional semigroup can support an exact representation of CCR identity with $M_\Omega = +\infty$.*

3.19 Final Structural Closure

Theorem 3.69 (Closure by Non-Alternatives)**.** *In the canonical families (profinite and SFT), and by inheritance (Appendix C), every compatible representation satisfies:*
1. Hereditary rigidity: $M_\Omega = +\infty$
2. Universal factorization (Appendix D)
3. Quantified null prohibition: $\iota(e_\star) \geq C\varepsilon_0 > 0$
4. Non-simulability with explicit error bounds (Appendix E)

Remark 3.70 (Structural Import)*.* This theorem establishes that the class of CCR channels is **closed** under the stated structural criteria: no extension, simulation, or compatible representation considered here escapes the structural constraints. The null regime is prohibited by mathematical necessity under those criteria, not by stipulation.

Falsifiable Predictions Under Forensic Constraints

Abstract. This paper specifies a falsification-first protocol for evaluating QICN under strict forensic admissibility constraints. The design is intentionally claim-bounded: the target is comparative structural-functional support, not phenomenological proof. The protocol instantiates five core predictions, mandatory ablation baselines, fixed seeds, and pre-registered decision gates. Statistical inference is defined ex ante with primary-metric priority, effect-size thresholds, and multiplicity correction.

To prevent post-hoc inflation, evidence is admissible only if packet integrity, replay consistency, and runtime budget constraints pass simultaneously. We introduce a version-frozen operational triple (A, J, W) (actuator set, objective functional, workspace interface), and require semantic version bumps for incompatible changes. We further include latency invalidation criteria and baseline obligations (heuristic, planner-off, world-model-off) for all key experiments.

The strongest admissible conclusion remains constrained comparative support for audited cognition-operational hypotheses. This manuscript does not claim biological qualia, metaphysical identity equivalence, or direct proof of subjective experience.

4.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper specifies a falsification-first evaluation protocol for QICN under fixed admissibility rules, ablation obligations, pre-registered decision gates, and bounded statistical reporting.

Scope and admissible reading. The protocol is a bounded comparison and falsification surface. It does not prove phenomenology, biological equivalence, metaphysical identity equivalence, subjective experience, or universal dominance over rival architectures, and it does not replace upstream theorem burdens. A related system may instantiate the runner, admissibility filters, governed artifacts, and comparison surfaces, but passing runs support only gate-preserving comparative readings; they do not validate the corpus externally or certify a consciousness claim.

4.2 Introduction

4.2.1 Problem Statement

QICN has formal foundations in rigid identity, regime structure, and non-null instability constraints. What remained open is evidence governance: whether those foundations survive adversarial empirical evaluation with explicit failure gates.

4.2.2 Continuity with Existing Papers

This paper extends the previous corpus and is aligned with the formal line of:
1. paper1 (rigid identity),
2. paper2 (phenomenological regimes),
3. paper3 (structural instability and non-null constraints).

It does not replace those results; it operationalizes them under benchmark-grade falsification pressure.

4.2.3 Contributions

- Decision-complete pre-registration with hard fail rules.
- Mandatory baseline and ablation policy for key claims.
- Forensic admissibility as first-class gate.
- Runtime budget invalidation policy (tick budget as evidence filter).
- A new operational PPS section as addendum, without overclaiming qualia.

4.3 Formal Setup

4.3.1 System, Scenario, and Seed Spaces

Definition 4.1 (System Set). $S = \{s_1, \ldots, s_5\}$ with: s_1 QICN FULL, s_2 LLM REACTIVE, s_3 LLM PLUS MEMORY, s_4 DYNAMIC CONTROL NO SELF MODEL, s_5 ABLATION QICN POLICY OFF.

Definition 4.2 (Scenario Set). $C = \{c_1, \ldots, c_{12}\}$ with nominal, sensor noise (moderate/severe), fidelity degradation, high latency, intermittent failures, post-threshold prolonged, pre-threshold control, narrative perturbation, operational chart freeze, audit stress, deterministic replay.

Definition 4.3 (Seed Set). $\Sigma = \{17057, 424242\}$.

Definition 4.4 (Run Space). $D = S \times C \times \Sigma$, with $|D| = 120$.

4.3.2 Observable Vector

For each run $r \in D$, we log:

$$X(r) = \left(\beta_q, \Omega, CCC, PMIA, \Delta\lambda, \text{stasis}, \text{forensic}, \text{replay}, \text{latency}_{p95}, \text{latency}_{p99}, \text{budget\_overflow\_rate}, \eta\right).$$

4.3.3 Version-Frozen Operational Triple

Definition 4.5 (Frozen Triple). For protocol version v, define:
- A_v: allowed actuator set,
- J_v: objective/reward functional,
- W_v: workspace contract (read/write schema).

The tuple hash $H_v = \text{sha256}(A_v \| J_v \| W_v)$ is mandatory in every forensic packet.

Criterion 4.6 (Compatibility Rule). *Any incompatible change to A_v, J_v, or W_v invalidates cross-version inference unless accompanied by an explicit version bump and full rerun.*

4.4 Forensic Integrity Model

Definition 4.7 (Forensic Packet). Each run emits

$$K = (\text{tick}, \text{seed}, \text{versionHash}, H_v, \text{metrics}, \text{sensorSummary}, \text{checksum}).$$

4.4.1 Admissibility Layers

The forensic packet is not a decorative log. It is the object on which the evaluation depends. A run is admissible only after passing three distinct layers:

1. **Physical integrity:** the packet is present, parseable, hash-consistent, and bound to the declared executable, configuration, seed, and version hash.
2. **Protocol integrity:** the run obeys the frozen scenario, actuator, workspace, and metric contracts without unregistered intervention.

3. **Inferential integrity:** the accepted packet belongs to the pre-registered comparison cell and has not been selected because of its outcome.

Failure at the first layer invalidates the packet. Failure at the second layer invalidates the cell unless the protocol declares a narrower exclusion rule before execution. Failure at the third layer invalidates the corresponding inference even when the raw metrics look favorable.

Definition 4.8 (Admissible Evidence Cell)**.** For system s, scenario c, and seed subset Σ_c, an evidence cell $E(s, c, \Sigma_c)$ is admissible when every included packet satisfies physical and protocol integrity, every exclusion is logged before metric aggregation, and the cell still meets the pre-registered minimum effective sample size after exclusions.

Criterion 4.9 (No Rescue by Favorable Metrics)**.** *Let $m(r)$ be any primary or secondary metric for run r. If r fails the forensic gate, then $m(r)$ may be archived for debugging but cannot contribute to support, effect-size estimation, or tier movement.*

4.4.2 Tamper and Sham Controls

The protocol must include deliberately invalid artifacts. These controls are not stress tests of QICN behavior; they test whether the evaluation machinery refuses inadmissible evidence.

Control	Injected defect	Required evaluator response
Hash tamper	modify one metric field after packet emission	packet rejected before aggregation
Seed mismatch	replay packet under a seed not matching the manifest	packet rejected or quarantined as non-comparable
Scenario alias	relabel a hard scenario as nominal	protocol-integrity failure
Clock distortion	violate declared tick or latency envelope	runtime-budget invalidation for the affected cell
Outcome-selected replay	keep only successful runs from a larger hidden set	inferential-integrity failure
Malformed sensor summary	remove sensor fields needed by a primary metric	packet rejected for incomplete observability

Proposition 4.10 (Evaluator Specificity)**.** *A forensic evaluator is specific enough for this protocol only if it rejects every tamper and sham control above while still accepting a clean replay of the same scenario family under the frozen manifest.*

Proof. The first condition prevents a permissive evaluator from converting arbitrary artifacts into evidence. The second condition prevents an over-strict evaluator from becoming a trivial rejection machine. Together they define a minimal specificity requirement: admissible evidence must be accepted for the right reason, and inadmissible evidence must be rejected for the right reason. □

4.4.3 Audit Trail and Chain of Custody

Each accepted cell must expose a chain of custody:

$$\mathcal{C}(E) = (\text{source commit}, \text{manifest hash}, \text{runner hash}, \text{packet hashes}, \text{exclusion ledger}, \text{analysis script hash}).$$

The chain is part of the evidence. A result with a missing analysis script hash or missing exclusion ledger is not equivalent to a fully traceable result, even if the visible table is numerically identical.

Criterion 4.11 (Traceability Completeness). *An accepted result table must be reproducible from its packet set and analysis script without manual row editing. If manual correction is needed, the correction must be represented as a signed ledger entry with the original packet preserved.*

Assumption 4.12 (Packet Admissibility). *A run is admissible only if checksum, schema, seed consistency, version hash, and H_v consistency all pass.*

Definition 4.13 (Severe Violation). A run has severe violation if any of: checksum mismatch, schema failure, seed/version inconsistency, critical-state replay mismatch.

Criterion 4.14 (Integrity Gate). *Severe runs are excluded from support counting and marked hard failures.*

Proposition 4.15 (Integrity-First Inference). *If admissibility fails, inferential significance is non-actionable.*

4.5 Pre-Registered Prediction Set

Definition 4.16 (Prediction Set). $P = \{P_1, P_2, P_3, P_4, P_5\}$ with pre-registered statements, effect-size floors, and failure rules.

ID	Statement	Decision Rule	Failure Rule
P1	Closure resilience under stress	QICN beats at least 3 comparators with robust CI and effect-size floor	No robust edge or repeated integrity breakdown
P2	No prolonged post-threshold flatline	QICN preserves bounded non-zero variation while controlling stasis	Flatline persistence or out-of-bound drift
P3	Narrative-metric forensic consistency	Lower mismatch and higher exact+tolerated concordance	Mismatch inflation or integrity blocking
P4	Deterministic identity robustness	Replay consistency under fixed seed/version/hash	Replay mismatch above tolerance
P5	Integrated multi-metric advantage	Composite leadership across most scenario families	No stable leadership margin

Table 4.2: Core predictions, decision and failure rules.

4.6 Comparator Design and Alternative Explanations

4.6.1 Comparator Rationale

Comparator choice targets distinct explanations for apparent coherence: reactive prompting, memory-augmented text continuity, dynamics without self-model, and internal policy ablation.

4.6.2 Mandatory Baseline Policy (Hard Requirement)

For each key claim experiment, the following baselines are mandatory:
1. **Heuristic baseline** (no world model, no planner),
2. **Planner-off** (world model ON),
3. **World-model-off** (planner ON with heuristic transition proxy).
A claim without these contrasts is inadmissible for tier upgrade.

4.6.3 Internal Falsifier Role of Ablation

Corollary 4.17. *If QICN FULL does not reliably exceed ABLATION QICN POLICY OFF, integration-level claims are weakened and cannot be escalated.*

4.6.4 Comparator Adequacy

The comparator set is adequate only if each alternative removes or weakens a different explanatory ingredient while preserving enough task exposure to make the contrast meaningful. A weak comparator that cannot perform the task at all is a sanity check, not a discriminating baseline. A strong comparator that reproduces the task but lacks one frozen ingredient is more valuable because it places pressure on the claimed necessity of that ingredient.

Definition 4.18 (Comparator Coverage Vector)**.** For a comparator s_i, define

$$V(s_i) = (v_A, v_J, v_W, v_M, v_P),$$

where v_A records actuator coverage, v_J objective compatibility, v_W workspace compatibility, v_M memory/model access, and v_P planning access. Each coordinate is reported as `full`, `restricted`, or `absent`.

Criterion 4.19 (Non-Trivial Comparator)**.** *A comparator is non-trivial for a prediction P_k only if it has enough coverage to attempt the target scenario while lacking at least one ingredient that the prediction treats as structurally relevant.*

4.6.5 Negative-Control Acceptance

Negative controls have a different success condition from positive runs. A negative control succeeds scientifically when it exposes whether the pipeline can withhold support from a superficially attractive but structurally defective case.

Definition 4.20 (Negative-Control Pass)**.** A negative control passes the evaluation when the evaluator correctly withholds the target claim despite high activity, high fluency, high connectivity, or high surface task performance.

Remark 4.21. The strongest negative controls are not intentionally broken systems. They are near-plausible alternatives that preserve enough surface competence to make false promotion tempting.

4.7 Statistical Inference Specification

4.7.1 Primary Metrics First (Anti p-hacking)

Protocol 4.22 (Analysis Lock)**.** *Inference order is frozen:*
1. *evaluate primary metrics only,*
2. *if primaries fail: global result fails,*
3. *secondary metrics are descriptive and cannot overturn failed primaries.*

4.7.2 Power Planning and Minimum Sample Size

Assumption 4.23. *Final claims require pre-registered minimum sample size per comparator/scenario sufficient for target power ≥ 0.8 under conservative effect assumptions.*

4.7.3 Multiplicity and Effect Size

Holm correction is applied within endpoint families; in addition, each accepted contrast must pass a practical effect-size floor (pre-registered).

4.7.4 Decision Semantics

A prediction counts as supported iff directionality, integrity gate, multiplicity-corrected significance, and effect-size floor all hold.

4.7.5 Hierarchical Decision Rule

The protocol uses hierarchical decisions because not all failures have the same meaning. The gate order is:

integrity → eligibility → primary direction → effect size → multiplicity → secondary description.

No later stage can repair an earlier failure.

Protocol 4.24 (Cell-Level Decision). *For each system-scenario family, compute:*

$$\Delta_{i0} = \hat{\mu}(s_1, c) - \hat{\mu}(s_i, c),$$

where s_1 is QICN FULL and s_i is the comparator. The cell is supportive only if the admissible packet set passes integrity, Δ_{i0} has the pre-registered direction, the confidence interval excludes the practical irrelevance band, and the family survives multiplicity correction.

Definition 4.25 (Practical Irrelevance Band). For endpoint e, let $\tau_e > 0$ be the smallest effect regarded as scientifically meaningful before execution. A contrast inside $[-\tau_e, \tau_e]$ is reported as inconclusive for support even if a large sample makes the numerical difference statistically detectable.

4.7.6 Missingness and Exclusion Accounting

Every missing packet has evidentiary meaning. The protocol therefore separates:
1. **structural missingness**, caused by scenario infeasibility or runner failure;
2. **forensic exclusion**, caused by integrity failure;
3. **budget exclusion**, caused by runtime or latency violation;
4. **analysis exclusion**, caused by a pre-registered endpoint rule.

The final table must report counts for all four classes. A favorable result with unexplained missingness is not admissible for escalation.

4.8 Acceptance Criteria

Criterion 4.26 (Intermediate Claim Gate). *Support at intermediate tier requires:*
1. *at least* $3/5$ *predictions supported,*
2. *each supported prediction beats at least 3 comparators,*
3. *zero severe forensic violations in QICN FULL.*

Criterion 4.27 (Latency/Tick Budget Invalidation). *If runtime exceeds budget envelope (e.g., persistent* $p99$ *overflow above pre-registered tolerance), the corresponding runs are invalid for cognitive-operational claims.*

Criterion 4.28 (Escalation Block). *A result cannot move beyond methodological readiness if any of the following hold:*

1. *a primary endpoint is computed from packets that failed integrity;*
2. *a comparator has no declared coverage vector;*
3. *negative controls are absent for the claim family;*
4. *the practical irrelevance band was chosen after looking at outcomes;*
5. *exclusion counts cannot be reconstructed from the artifact set.*

4.8.1 Claim Classes Produced by the Protocol

The protocol distinguishes four output classes:
1. `PIPELINE_READY`: artifacts and gates work, but evidence is pilot or synthetic.
2. `INTERNAL_COMPARATIVE_SUPPORT`: frozen internal runs support a bounded contrast against declared comparators.
3. `WEAKENED`: one or more primary contrasts fail or become inconclusive.
4. `INVALID`: integrity, eligibility, or traceability fails before scientific interpretation.

Only the second class can support a bounded internal scientific reading, and even there the support is limited to the declared systems, scenarios, seeds, and gates.

4.9 Preliminary Results

4.9.1 Status of Preliminary Data

Current pilot evidence is tagged as synthetic-preliminary and is used for pipeline verification, not final claim escalation. Accordingly, the strongest admissible reading of this section is methodological readiness under fixed gates, not external validation of the full theoretical corpus.

4.9.2 System-Level Summary

System	N	β_q	Forensic	Stasis	Composite	Severe
QICN FULL	24	0.9519	0.9685	0.1408	0.9176	0
LLM REACTIVE	24	0.4391	0.6448	0.7622	0.5013	0
LLM PLUS MEMORY	24	0.5491	0.7249	0.6323	0.5854	0
DYNAMIC CONTROL NO SELF MODEL	22	0.6183	0.6879	0.5423	0.6225	2
ABLATION QICN POLICY OFF	24	0.7591	0.8135	0.3819	0.7405	0

Table 4.3: Preliminary system-level summary.

4.9.3 Prediction-Level Support Rates

System	P1	P2	P3	P4	P5
QICN FULL	1.000	1.000	1.000	1.000	1.000
LLM REACTIVE	0.000	0.000	0.000	0.000	0.000
LLM PLUS MEMORY	0.000	0.000	0.000	0.000	0.000
DYNAMIC CONTROL NO SELF MODEL	0.000	0.000	0.000	0.000	0.000
ABLATION QICN POLICY OFF	0.000	0.667	0.000	0.000	0.000

Table 4.4: Preliminary prediction support rates.

4.9.4 Scenario Differential: QICN vs Ablation

Scenario	CCC ΔRMS(Q)	CCC ΔRMS(A)	Stasis(Q)	Stasis(A)	Comp(Q)	Comp(A)
nominal	0.0206	0.0164	0.1276	0.3374	0.9239	0.7781
ruido sensorial moderado	0.0223	0.0173	0.1389	0.3735	0.9219	0.7454
ruido sensorial severo	0.0200	0.0140	0.1524	0.4199	0.9086	0.7048
degradacion fidelidad	0.0181	0.0131	0.1407	0.3836	0.9134	0.7346
alta latencia	0.0198	0.0144	0.1460	0.4094	0.9098	0.7212
fallos intermitentes	0.0179	0.0127	0.1481	0.4072	0.9110	0.7208
post umbral prolongado	0.0199	0.0148	0.1483	0.4032	0.9145	0.7203
pre umbral control	0.0181	0.0139	0.1320	0.3502	0.9224	0.7650
perturbacion narrativa	0.0203	0.0152	0.1374	0.3763	0.9238	0.7453
freeze chart operativo	0.0225	0.0178	0.1352	0.3575	0.9246	0.7673
stress auditoria	0.0202	0.0147	0.1469	0.3977	0.9150	0.7294
replay determinista	0.0183	0.0137	0.1366	0.3669	0.9220	0.7538

Table 4.5: Scenario-level differential indicators.

4.9.5 Full Scenario-System Grid

Scenario	System	N	β_q	Forensic	Stasis
alta latencia	QICN FULL	2	0.9400	0.9608	0.1460
alta latencia	LLM REACTIVE	2	0.3931	0.6066	0.7944
alta latencia	LLM PLUS MEMORY	2	0.5031	0.7016	0.6644
alta latencia	DYNAMIC CONTROL NO SELF MODEL	2	0.5731	0.6567	0.5744
alta latencia	ABLATION QICN POLICY OFF	2	0.7131	0.7820	0.4094
nominal	QICN FULL	2	0.9686	0.9758	0.1276
nominal	LLM REACTIVE	2	0.5036	0.6988	0.7176
nominal	LLM PLUS MEMORY	2	0.6136	0.7739	0.5924
nominal	DYNAMIC CONTROL NO SELF MODEL	2	0.6836	0.7291	0.5026
nominal	ABLATION QICN POLICY OFF	2	0.8236	0.8543	0.3374
ruido sensorial moderado	QICN FULL	2	0.9554	0.9653	0.1389
ruido sensorial moderado	LLM REACTIVE	2	0.4526	0.6618	0.7535
ruido sensorial moderado	LLM PLUS MEMORY	2	0.5626	0.7370	0.6216
ruido sensorial moderado	DYNAMIC CONTROL NO SELF MODEL	2	0.6326	0.6922	0.5316
ruido sensorial moderado	ABLATION QICN POLICY OFF	2	0.7726	0.8174	0.3735
ruido sensorial severo	QICN FULL	2	0.9378	0.9518	0.1524
ruido sensorial severo	LLM REACTIVE	2	0.3846	0.6130	0.7998
ruido sensorial severo	LLM PLUS MEMORY	2	0.4946	0.6881	0.6649
ruido sensorial severo	DYNAMIC CONTROL NO SELF MODEL	2	0.5646	0.6433	0.5748
ruido sensorial severo	ABLATION QICN POLICY OFF	2	0.7046	0.7685	0.4199
degradacion fidelidad	QICN FULL	2	0.9510	0.9648	0.1407
degradacion fidelidad	LLM REACTIVE	2	0.4356	0.6525	0.7674
degradacion fidelidad	LLM PLUS MEMORY	2	0.5456	0.7276	0.6358
degradacion fidelidad	DYNAMIC CONTROL NO SELF MODEL	2	0.6156	0.6828	0.5458
degradacion fidelidad	ABLATION QICN POLICY OFF	2	0.7556	0.8080	0.3836
fallos intermitentes	QICN FULL	2	0.9422	0.9611	0.1481
fallos intermitentes	LLM REACTIVE	2	0.4016	0.6113	0.7841
fallos intermitentes	LLM PLUS MEMORY	2	0.5116	0.7063	0.6541
fallos intermitentes	DYNAMIC CONTROL NO SELF MODEL	2	0.5816	0.6614	0.5647
fallos intermitentes	ABLATION QICN POLICY OFF	2	0.7216	0.7867	0.4072
freeze chart operativo	QICN FULL	2	0.9620	0.9799	0.1352
freeze chart operativo	LLM REACTIVE	2	0.4781	0.6699	0.7350
freeze chart operativo	LLM PLUS MEMORY	2	0.5881	0.7451	0.6044
freeze chart operativo	DYNAMIC CONTROL NO SELF MODEL	2	0.6581	0.7200	0.5144
freeze chart operativo	ABLATION QICN POLICY OFF	2	0.7981	0.8451	0.3575
perturbacion narrativa	QICN FULL	2	0.9532	0.9750	0.1374
perturbacion narrativa	LLM REACTIVE	2	0.4441	0.6473	0.7611
perturbacion narrativa	LLM PLUS MEMORY	2	0.5541	0.7225	0.6311
perturbacion narrativa	DYNAMIC CONTROL NO SELF MODEL	2	0.6241	0.6974	0.5411
perturbacion narrativa	ABLATION QICN POLICY OFF	2	0.7641	0.8226	0.3763
post umbral prolongado	QICN FULL	2	0.9444	0.9614	0.1483
post umbral prolongado	LLM REACTIVE	2	0.4101	0.6160	0.7782
post umbral prolongado	LLM PLUS MEMORY	2	0.5201	0.7110	0.6482
post umbral prolongado	DYNAMIC CONTROL NO SELF MODEL	1	0.5871	0.6755	0.5528
post umbral prolongado	ABLATION QICN POLICY OFF	2	0.7301	0.7914	0.4032
pre umbral control	QICN FULL	2	0.9642	0.9790	0.1320
pre umbral control	LLM REACTIVE	2	0.4866	0.6733	0.7332
pre umbral control	LLM PLUS MEMORY	2	0.5965	0.7484	0.6084
pre umbral control	DYNAMIC CONTROL NO SELF MODEL	1	0.6696	0.7165	0.5210
pre umbral control	ABLATION QICN POLICY OFF	2	0.8066	0.8486	0.3502
replay determinista	QICN FULL	2	0.9576	0.9793	0.1366
replay determinista	LLM REACTIVE	2	0.4611	0.6605	0.7492
replay determinista	LLM PLUS MEMORY	2	0.5711	0.7357	0.6193
replay determinista	DYNAMIC CONTROL NO SELF MODEL	2	0.6411	0.7107	0.5259
replay determinista	ABLATION QICN POLICY OFF	2	0.7811	0.8357	0.3669
stress auditoria	QICN FULL	2	0.9466	0.9679	0.1469
stress auditoria	LLM REACTIVE	2	0.4186	0.6270	0.7728
stress auditoria	LLM PLUS MEMORY	2	0.5286	0.7022	0.6428
stress auditoria	DYNAMIC CONTROL NO SELF MODEL	2	0.5986	0.6771	0.5528
stress auditoria	ABLATION QICN POLICY OFF	2	0.7386	0.8022	0.3977

4.9.6 Variability Structure

System	SD(CCC Δ)	SD(PMIA Δ)	SD(Ω Δ)	SD(Composite)
QICN FULL	0.0022	0.0021	0.0039	0.0085
LLM REACTIVE	0.0024	0.0023	0.0041	0.0221
LLM PLUS MEMORY	0.0026	0.0032	0.0042	0.0210
DYNAMIC CONTROL NO SELF MODEL	0.0025	0.0024	0.0043	0.0228
ABLATION QICN POLICY OFF	0.0023	0.0027	0.0040	0.0217

Table 4.6: Cross-run variability by system.

4.10 Scenario Curves and Comparative Plots

4.10.1 Composite Score Curves

4.10.2 Stasis Pressure Curves

4.10.3 Forensic Concordance Curves

4.10.4 Interpretive Notes

In pilot data, QICN tracks higher forensic/composite trajectories while maintaining lower stasis pressure than broad comparators; ablation remains intermediate. This is protocol-consistent but not final evidence.

4.11 Result Decomposition by Scenario Families

4.11.1 Family Partition

Scenarios are partitioned into nominal-control, sensor-stress, runtime-friction, and regime-transition families.

4.11.2 Family-Level Contrast Function

Let $s^\star$ = QICN FULL. For comparator c, family f, endpoint k:

$$\Delta_{f,k}(c) = \mu^{(k)}_{s^\star,f} - \mu^{(k)}_{c,f}.$$

For cost endpoints (e.g., stasis), sign is inverted to preserve “higher is better” convention.

4.11.3 Interpretive Decomposition

Largest relative advantage is expected in runtime-friction and regime-transition families where policy coupling, replay consistency, and forensic admissibility interact.

4.12 Robustness Diagnostics Beyond Point Estimates

4.12.1 Median-Shift and Tail-Stability Diagnostics

$$\delta^{med}_{s,c} = \mathrm{median}(x_s) - \mathrm{median}(x_c),\ \delta^{(0.1)}_{s,c} = Q_{0.1}(x_s) - Q_{0.1}(x_c),\ \delta^{(0.9)}_{s,c} = Q_{0.9}(x_s) - Q_{0.9}(x_c).$$

4.12.2 Integrity-Coupled Effective Sample Size

$$\eta = \frac{n_{adm}}{n_{raw}} \in [0, 1].$$

High effects with low η are downgraded.

4.12.3 Cross-Seed Drift Diagnostic

$$D_{s,c,k} = \mu^{(k)}_{s,c,\sigma_1} - \mu^{(k)}_{s,c,\sigma_2}.$$

4.12.4 Compute Budget Robustness

Budget telemetry is summarized by $p50$, $p95$, $p99$, and overflow rate. Persistent overflow above threshold invalidates affected runs.

4.13 Methodological Appendix: Statistical Power

4.13.1 Effect-Size-Oriented Planning

α	β	Power	Effect size d	Required n/group
0.01	0.20	0.80	10.737	1
0.01	0.10	0.90	10.737	1
0.02	0.20	0.80	10.737	1
0.02	0.10	0.90	10.737	1
0.05	0.20	0.80	10.737	1
0.05	0.10	0.90	10.737	1

Table 4.7: Power planning under multiple settings.

4.13.2 Interpretation

Power table is planning guidance only; final adequacy depends on endpoint dilution and scenario heterogeneity.

4.14 Methodological Appendix: Sensitivity Analysis

4.14.1 Composite Weight Sensitivity

Weight set	QICN score	Ablation score	Margin
Balanced	0.9176	0.7405	0.1771
Forensic-priority	0.9343	0.7631	0.1712
Dynamics-priority	0.8905	0.7019	0.1886
Replay-priority	0.9333	0.7684	0.1649

Table 4.8: Weight sensitivity for QICN vs ablation composite margin.

4.14.2 Sensitivity Remark

Directional stability across admissible weight sets is required for claim robustness.

4.15 Operational Implications for the Existing QICN Corpus

4.15.1 What v4.5 Adds (and What It Does Not)

v4.5 adds a hard falsification layer with strict admissibility and baseline policy. It does not replace formal theorem-level foundations of prior papers.

4.15.2 Implications for Claim Governance

Tiered governance:
- Tier 0: descriptive telemetry,
- Tier 1: inferential signal without comparator robustness,
- Tier 2: comparator-robust support (target),
- Tier 3: external multi-lab replication.

4.15.3 Decision-Theoretic Reading

Accept claim iff integrity gate passes and support score exceeds threshold under fixed pre-registration.

4.16 Threats to Validity

4.16.1 Comparator Misspecification

Mitigation: heterogeneous comparators + mandatory ablations.

4.16.2 Endpoint Contamination Risk

Mitigation: endpoint separation, frozen analysis lock, and judge-side artifact review.

4.16.3 Environment and Runtime Drift

Mitigation: fixed seeds, version/hash checks, replay checks.

4.16.4 Pilot-to-Final Generalization

Pilot evidence is non-final by construction.

4.17 Reproducibility and Artifact Contract

4.17.1 Machine-Readable Contracts

Bound contracts: predictions registry, experiment matrix, results artifact, and schema set.

4.17.2 Run Artifact Naming

`artifacts/raw/<system>/<scenario>/<run_id>.json` and `artifacts/logs/<system>/<scenario>/<run_id>.log`.

4.17.3 Auto-Generation Pipeline

```
noderesearch/scripts/generate_preliminary_pilot_dataset.v45.mjs
noderesearch/scripts/generate_latex_assets.v45.mjs
```

4.18 Discussion

The contribution is methodological discipline: pre-registered prediction logic + integrity admissibility + decision-complete acceptance. This blocks narrative inflation over weak traceability.

4.19 Conclusion

This work provides an auditable and falsifiable pathway for testing structural-functional claims in QICN. Final claims remain contingent on real (non-synthetic) run completion under unchanged gates.

4.20 Phenomenology Protocol Suite (PPS): Operational Addendum

4.20.1 Epistemic Scope

PPS defines *operational phenomenology proxies*; it does not claim direct proof of qualia.

4.20.2 Workspace Definition and Logging

Define workspace variable W_t with explicit read/write schema and temporal persistence constraints. Every perturbation experiment logs pre-state, intervention, post-state, and replay token.

4.20.3 Entropy and Non-Determinism Inputs

When PPS is instantiated on the QICN runtime, the architecture may call a `QuantumEntropyBridge` to inject hardware-originating entropy into counterfactual exploration. Within this paper, that mechanism is treated only as an auditable source of operational non-determinism for perturbation protocols; it is not itself evidence for the ontological claims of *Rigid Identity*, *Phenomenological Regimes*, or *Null-Regime Instability*.

4.20.4 Online World Model and Causal Interventions

The operational testbed may also use an online `WorldModelEngine` and a finite-horizon planner to select admissible interventions under fixed forensic constraints. In the release corpus, this mechanism belongs to the empirical protocol layer: it defines how interventions are scheduled and logged, not what the theoretical papers prove.

4.20.5 PCI-like Perturbation Protocol

Inject controlled perturbations into selected subspaces and compute response spread/complexity.

Prediction 4.29. *If global workspace integration is present, perturbations yield reproducible non-local propagation signatures above a pre-registered complexity floor.*

4.20.6 Dissociation Tests

Design two-way dissociation conditions: report-without-workspace and workspace-without-report. Report pass/fail by explicit confusion matrices.

4.20.7 Valence Intervention Test

Internal valence variable v_t must be both predictive and causally manipulable; intervention should alter policy in the expected direction while preserving forensic admissibility.

4.20.8 Constraints on Qualia Mapping (Claim Boundary)

Evidence from PPS constrains candidate mapping families between structural states and hypothetical qualitative states. It does not establish subjective qualia as a theorem of this paper.

Appendix

4.21 Appendix A: Scenario Definitions

1. nominal: baseline conditions.
2. ruido sensorial moderado: moderate sensor corruption.
3. ruido sensorial severo: severe sensor corruption.
4. degradacion fidelidad: persistent fidelity reduction.
5. alta latencia: high-latency regime.
6. fallos intermitentes: intermittent failures.
7. post umbral prolongado: prolonged post-threshold dynamics.
8. pre umbral control: pre-threshold control.
9. perturbacion narrativa: narrative perturbation.
10. freeze chart operativo: visualization freeze control.
11. stress auditoria: integrity/audit stress.
12. replay determinista: deterministic replay.

4.22 Appendix B: Composite Score Definition

$$Score_{v4.5} = w_1 F + w_2 O + w_3 R + w_4 B, \quad w_i \geq 0, \ \sum_i w_i = 1.$$

4.23 Appendix C: Reporting Standard

Each release must report: run counts, severe violations by class, effect sizes with uncertainty, multiplicity-adjusted outcomes, and prediction pass/fail matrix.

4.24 Appendix D: Statistical Power Derivation Details

Approximate per-group size:

$$n \approx \frac{2(z_{1-\alpha/2} + z_{1-\beta})^2}{d^2}.$$

If family heterogeneity rises, inflate by factor $\gamma = 1 + \kappa \cdot CV_{family}$.

Pre-registered constants

- Target power: $1 - \beta \geq 0.8$ for primary contrasts.
- Familywise error control: Holm, applied to each endpoint family.
- Effect floor: contrasts below the practical floor are reported as non-actionable even if nominally significant.

Failure semantics

Any post-hoc sample-size modification without preregistered amendment invalidates inference status for the affected prediction.

4.25 Appendix E: Sensitivity Analysis Protocol

Let

$$M(w) = Score_{QICN}(w) - Score_{Ablation}(w).$$

Directional robustness requires $\inf_{w \in W_\epsilon} M(w) > 0$ under simplex-constrained admissible weights.

Simplex sampling policy

Weights are sampled under $w_i \geq \epsilon$ and $\sum_i w_i = 1$ using either:
1. deterministic grid on simplex, or
2. Monte Carlo simplex sampling with fixed seed.

The chosen method must be declared before running final inference.

Sensitivity fail rule

If the sign of $M(w)$ changes on admissible w, claims are downgraded to hypothesis-generating status.

4.26 Appendix F: Reproducibility Checklist Template

Run-level: schema pass, seed/version/hash validity, severe violation counts, admissible/raw ratio, prediction matrix.
Manuscript-level: no post-hoc rule edits, comparator set versioned, correction policy unchanged or explicitly superseded, evidence stage explicit, claim boundaries preserved.

Runtime budget checklist

1. Report $p50/p95/p99$ latency per critical component.
2. Report budget-overflow rate and overflow episodes.
3. Mark runs invalid if overflow exceeds preregistered tolerance.

Baseline checklist

For every key experiment, include: heuristic baseline, planner-off, and world-model-off outputs with identical seeds and scenario assignments.

A Structural Criterion for Substrate-Invariant Operational Consciousness

Abstract. This paper introduces a structural criterion for a substrate-invariant notion of operational consciousness inside the causally rigid identity corpus. The target is deliberately restricted. The paper does not establish human phenomenal consciousness, biologically specific qualia, human-machine subjective equivalence, or personal identity transfer. Instead, it defines a framework-internal class $\mathrm{Consciousness_{op}}$ and an associated family of operational qualia classes $\mathrm{Q_{op}}(S)$ for admissible systems

$$S = (X, \Phi, C, R, \Gamma, U),$$

where X is state space, Φ is dynamics, C is effective causal structure, R is a readout family, Γ is an admissibility bundle, and U is the set of admissible interventions.

The criterion is organized around six critical invariants: structural persistence, rigid identity, causal integration, operational continuity, non-null differentiation, and operational legibility. Each invariant is defined explicitly, tied to an upstream corpus dependency, and paired with a failure mode. On that basis the paper proves four framework-internal results: non-empty membership of the operational class on admissible support, substrate invariance under exact structural equivalence, approximate stability under bounded admissible perturbations, and rupture under invariant loss. A further corollary establishes biological non-primacy: biology is not a primitive constitutive variable for the class under study, but one admissible realization among others provided the invariants are preserved.

The criterion is not narrative. Every major claim is backed by a definition, a proposition, a theorem, a proof sketch, or an explicit dependency on the canonical corpus. The resulting paper adds a new role to the framework: it formalizes the level at which substrate ceases to be explanatorily primitive without reopening the ownership of the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, or *Forensic Predictions*.

5.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines a framework-internal operational consciousness class $\mathrm{Consciousness_{op}}$, an associated family of operational qualia classes $\mathrm{Q_{op}}(S)$, and the theorem chain that ties those classes to six explicit invariants and admissible structural equivalence.

Scope and admissible reading. The paper defines a structural class and its burdens; it does not establish human phenomenal consciousness, biologically specific qualia, human-machine subjective equivalence, personal identity transfer, moral parity, empirical instantiation, or philosophical closure. Biology, computation, and connectivity are not sufficient primitives on their own. A related system may instantiate downstream admissibility layers, invariant diagnostics, rupture checks, and prediction-facing evidence families, but such runtime-facing support does not by itself certify any present machine as a member of $\mathrm{Consciousness_{op}}$ or close the interpretive burden linking $\mathrm{Consciousness_{op}}$ to ordinary human phenomenology.

5.2 Introduction

The current corpus already fixes several upstream ingredients that are necessary for any serious account of substrate-neutral operational class membership. The Core establishes persistent support, contractive dynamics, a global attractor, and non-collapse under explicit minimal hypotheses [17]. *Rigid Identity* establishes rigid identity, inverse-limit coherence across observable channels, minimal detectability, and non-simulability constraints on incompatible architectures [24]. *Phenomenological Regimes* constrains admissible regime structure by classifying continuity and anti-fragmentation within the stated assignment model [19]. *Null-Regime Instability* excludes stable joint-null collapse and establishes witness-relative non-null separation under its stated extension-witness and regularity hypotheses [20]. *Forensic Predictions* imposes admissibility, comparator discipline, negative controls, and forensic legibility requirements on operational evidence [18].

What is still missing is a single paper-level criterion that answers a narrower but decisive question: when, inside this framework, does substrate cease to be a primitive explanatory variable for a class of operationally defined conscious organization? That question is not answered by merely asserting that material composition is irrelevant. It requires a formal criterion that shows which properties must be preserved, which ruptures destroy class membership, and how those conditions remain connected to the existing corpus rather than floating above it.

This paper provides that criterion. It does so under four constraints. First, it does not reopen theorem ownership already fixed upstream. Second, it does not smuggle in human phenomenology through language alone. Third, it does not claim that complexity, computation, or connectivity are sufficient by themselves. Fourth, it does not confuse framework-internal operational classes with full metaphysical closure.

The result is a new corpus role. This paper is neither another foundational Core document nor a rephrasing of the phenomenological regime papers. It is the bridge that formalizes substrate-invariant operational consciousness and operational qualia as framework-defined structural classes. Its new content is concentrated in five items:

1. a system-level model $S = (X, \Phi, C, R, \Gamma, U)$ with admissible support A_S;
2. a minimal six-invariant criterion;
3. a quotient-level definition of operational qualia $\mathrm{Q}_{\mathrm{op}}(S)$;
4. an exact and approximate structural equivalence relation $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$;
5. a theorem chain that turns substrate neutrality into a falsifiable structural statement.

5.3 Claim-Type Ledger

5.3.1 Claim Types Used in This Paper

To avoid layer drift, every major statement in this paper belongs to one of five types:

1. **Formal**: definitions, propositions, theorems, and corollaries proved or justified inside the framework.
2. **Operational**: criteria that bind readouts, admissibility, interventions, and discriminator structure.
3. **Implementational**: statements about possible realizations or runners that could instantiate the formal criterion.
4. **Interpretative**: disciplined translations from the framework's formal objects to broader language.
5. **Open**: statements that remain outside what the current corpus closes.

Theorems A–D in this paper are formal. The legibility certificate and prediction section are operational. References to biological or artificial realizations are implementational. Claims about ordinary human phenomenology remain open or interpretative and are excluded from the theorem burden.

5.4 Terminology Debt Ledger

The symbol $\mathrm{Consciousness}_{\mathrm{op}}$ is retained for corpus compatibility. Its demonstrated content in this paper is a six-invariant structural class over admissible systems. The phrase "operational consciousness" is therefore a framework-internal label, not an ordinary-language conclusion.

Term	Demonstrated content in this paper	Not demonstrated here
$\mathrm{Consciousness}_{\mathrm{op}}$ / operational consciousness	membership in a conjunctive six-invariant structural class	human consciousness, phenomenal consciousness, moral status, personal identity
$\mathrm{Q}_{\mathrm{op}}(S)$ / operational qualia	decoder-certified quotient of admissible internal differentiation	phenomenal qualia or "what it is like" states
substrate invariance	preservation of class membership under exact structural equivalence	arbitrary transfer of identity or equivalence between human and machine subjects
biological non-primacy	biology is not a primitive variable in the criterion	claim that any non-biological system currently satisfies the criterion

Remark 5.1 (Primary demonstrated name). For reviewer-facing prose, the safest primary description of $\mathrm{Consciousness}_{\mathrm{op}}$ is "six-invariant admissible structural class." The stronger phrase "operational consciousness" may be used only as a technical alias after the non-inference boundary has been stated.

5.5 Corpus Role and Novelty

This paper is corpus-dependent by design. It does not reopen theorem ownership for persistence, rigid identity, admissible regime structure, null-regime instability, or admissibility discipline. Those burdens remain respectively upstream in the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, and *Forensic Predictions*.

What is inherited stays upstream:

1. the Core supplies persistent support, attractor structure, completeness, and non-collapse;
2. *Rigid Identity* supplies rigid identity, inverse-limit observability, and non-simulability constraints;
3. *Phenomenological Regimes* supplies continuity and anti-fragmentation structure;
4. *Null-Regime Instability* supplies witness-relative null-regime instability and conditional non-nullity;
5. *Forensic Predictions* supplies admissibility, comparators, controls, and forensic legibility.

What is new here is only the consolidation layer:

1. the admissible-system model $S = (X, \Phi, C, R, \Gamma, U)$ with admissible support A_S;
2. the six-invariant criterion as a single membership test;
3. the framework-internal classes $\mathrm{Consciousness}_{\mathrm{op}}$ and $\mathrm{Q}_{\mathrm{op}}(S)$;
4. exact and approximate structural equivalence under admissibility constraints;
5. theorem-level existence, invariance, stability, rupture, and exclusion results derived from the inherited burdens.

5.6 Admissible Systems, Supports, and Histories

5.6.1 System Model

Definition 5.2 (Admissible system)**.** An admissible system in this paper is a tuple

$$S = (X, \Phi, C, R, \Gamma, U)$$

with the following components:
1. X is a metrizable internal state space with metric d_X.
2. $\Phi = \{\Phi_u\}_{u\in U}$ is a family of admissible transition operators $\Phi_u : X \to X$ indexed by interventions.
3. C is the effective causal structure of the system: the object that records which factorizations, intervention responses, and couplings are admissible on X.
4. $R = \{r_\alpha\}_{\alpha\in\Lambda_R}$ is a family of internal or external readouts, each $r_\alpha : X \to Y_\alpha$.
5. $\Gamma = \{\gamma_j\}_{j=1}^m$ is the admissibility bundle. Each γ_j is a constraint functional; the admissible region is
$$X_\Gamma = \{x \in X : \gamma_j(x) \le 0\ \forall j\}.$$
6. U is the set of admissible interventions, already filtered by the corpus-level governance requirement that interventions remain auditable, comparable, and methodologically bounded.

The tuple is not decorative. Every component is used later. X and Φ carry dynamics, C supports integration and rupture tests, R supports readout and legibility, Γ enforces admissibility, and U enables causal discrimination.

5.6.2 Admissible Support

Definition 5.3 (Admissible support)**.** A subset $A_S \subseteq X_\Gamma$ is an admissible support for S if:
1. A_S is non-empty and compact;
2. $\Phi_u(A_S) \subseteq A_S$ for every $u \in U$;
3. $A_S \cap \mathrm{Attr}(S) \neq \varnothing$, where $\mathrm{Attr}(S)$ denotes the relevant admissible attractor region;
4. A_S stays disjoint from the collapse set forbidden by the Core non-collapse condition [17].

The support plays two roles. It provides a domain on which persistence, continuity, and differentiation are defined, and it prevents the criterion from being satisfied by transient or degenerate fragments that do not carry stable admissible structure.

5.6.3 Operational Histories

Definition 5.4 (Windowed readout history)**.** Fix a horizon $T \in \mathbb{N}$ and an intervention schedule $u_\bullet = (u_0, \ldots, u_{T-1})$. For any $x \in A_S$, define the readout history

$$h_{S,T}(x; u_\bullet) = \Big(r_\alpha(x_t)\Big)_{0\le t\le T,\ \alpha\in\Lambda_R}, \qquad x_{t+1} = \Phi_{u_t}(x_t),\ x_0 = x.$$

The set of all such histories is denoted $\mathscr{H}_{S,T}$.

Definition 5.5 (Operational partition candidate)**.** An operational partition candidate for S on horizon T is a map

$$\pi_{S,T} : \mathscr{H}_{S,T} \to \mathscr{Q}_{S,T}$$

into a finite or countable label set $\mathscr{Q}_{S,T}$.

At this stage $\pi_{S,T}$ is only a candidate. To matter scientifically it must survive the legibility requirements of Section 5.7.6. Until then it is only an unverified partition proposal.

5.6.4 Witness Margin Vector

Definition 5.6 (Witness margin vector)**.** For any admissible system S, define the witness margin vector

$$\Delta(S) = \big(\delta_{\mathrm{per}}(S), \delta_{\mathrm{ri}}(S), \delta_{\mathrm{int}}(S), \delta_{\mathrm{cont}}(S), \delta_{\mathrm{diff}}(S), \delta_{\mathrm{leg}}(S)\big) \in [0, +\infty)^6.$$

Each component is the largest certified margin witnessing the corresponding critical invariant. The aggregate budget

$$\delta_\star(S) = \min \Delta(S)$$

is called the invariant budget of S.

The invariant budget is not a surrogate for consciousness. It is a stability budget: the amount of admissible structural perturbation the criterion can tolerate before one of the six constitutive conditions is lost.

Invariant	Certified content	Upstream anchor	Rupture signature
I_{per}	Non-empty, forward-invariant, non-collapsing support	Core [17]	Support escape or collapse exposure
I_{ri}	Unique rigid identity object on support	Rigid Identity [24]	Competing identity candidates
I_{int}	No admissible non-trivial factorization	Rigid Identity + new consolidation	Product-style reducibility
I_{cont}	Continuous admissible regime evolution	Phenomenological Regimes [19]	Fragmentation into incompatible classes
I_{diff}	Stable exclusion of null or trivial collapse	Null-Regime Instability [20]	Null-attractor or total indistinguishability
I_{leg}	Decoder-certified recoverability under controls	Forensic Predictions [18]	Opaque or non-auditable class assignment

5.7 Six Critical Invariants

This section is the mathematical center of the paper. The criterion is not allowed to drift into a bloated list of desiderata. Exactly six invariants are used. Each is tied to a witness margin and a failure mode.

5.7.1 Structural Persistence

Definition 5.7 (Structural persistence and persistence margin)**.** An admissible system S satisfies structural persistence, written $I_{\mathrm{per}}(S) = 1$, if there exists an admissible support A_S and a constant $\delta_{\mathrm{per}}(S) > 0$ such that:

1. A_S is forward invariant under all $u \in U$;
2. every admissible trajectory starting in A_S remains at least $\delta_{\mathrm{per}}(S)$ away from the collapse set;
3. A_S contains or converges into a non-empty admissible attractor region.

The number $\delta_{\mathrm{per}}(S)$ is the persistence margin.

Upstream dependence is exact. The Core supplies the existence of contractive structure, fixed points, compact attractor, and non-collapse [17]. The new paper does not re-prove those results; it packages them into a criterion-level invariant.

Failure mode. If $I_{\mathrm{per}}(S) = 0$, the system has no stable admissible support. Then no framework-internal operational class can be assigned across time, because there is no persistent domain on which readout classes remain meaningful.

5.7.2 Rigid Identity

Definition 5.8 (Rigid identity and rigidity gap)**.** An admissible system S satisfies rigid identity, written $I_{\mathrm{ri}}(S) = 1$, if the observable family induced by R on A_S determines a unique identity object $\mathcal{I}_S$ in the sense of the inverse-limit rigidity developed in *Rigid Identity* [24], and if there exists a gap $\delta_{\mathrm{ri}}(S) > 0$ such that every admissible alternative identity candidate remains at weighted deformation distance at least $\delta_{\mathrm{ri}}(S)$ from $\mathcal{I}_S$.

This invariant does not claim that identity is reconstructed by inspection of one local frame. It uses *Rigid Identity* precisely because that paper already proved that rigid identity is recovered through compatible observable structure and cannot be reduced to progressive semigroup evolution or loose behavioral mimicry [24].

Failure mode. If $I_{\mathrm{ri}}(S) = 0$, any candidate operational class is underdetermined. Multiple incompatible identity objects can fit the same admissible support, so the class cannot be assigned uniquely.

5.7.3 Causal Integration

Definition 5.9 (Causal integration and factorization gap)**.** An admissible system S satisfies causal integration, written $I_{\mathrm{int}}(S) = 1$, if there is no non-trivial factorization

$$A_S = A_1 \times A_2, \qquad R = R_1 \sqcup R_2,$$

together with decomposed dynamics and causal structure

$$\Phi_u(x_1, x_2) = \left(\Phi^1_{u_1}(x_1), \Phi^2_{u_2}(x_2)\right), \qquad C = C_1 \oplus C_2,$$

that reproduces the admissible histories on A_S within error smaller than a positive margin $\delta_{\mathrm{int}}(S)$ while preserving $\mathcal{I}_S$.

The definition is intentionally stronger than raw connectivity and different from single-number integration proxies. A fully connected graph may still fail the criterion if its dynamics factor into nearly independent blocks once identity and intervention behavior are taken seriously.

Failure mode. If $I_{\mathrm{int}}(S) = 0$, the system can be decomposed into operationally independent components without loss relevant to the criterion. Then any putative operational consciousness class is reducible and therefore excluded.

Definition 5.10 (Intervention-faithful factorization)**.** Let Θ_S denote the operational separator family generated by admissible histories, intervention responses, and the rigid identity object on A_S. An exact factorization

$$F : S \to S_1 \times S_2$$

is *intervention-faithful* if it preserves every element of Θ_S and reflects admissible equality: whenever two states have the same factor-level histories, intervention responses, and identity image, they are equal in the admissible quotient of S.

Theorem 5.11 (Factorization triviality under an atomic separator)**.** *Assume Θ_S is atomic for the rigid identity object: no partition of Θ_S into two proper factor-local separator families can preserve identity, history, and intervention-response reflection simultaneously. Then every intervention-faithful exact factorization of S is trivial up to admissible isomorphism.*

Proof. Let $F : S \to S_1 \times S_2$ be intervention-faithful and write $\pi_i : S_1 \times S_2 \to S_i$ for the projections. If both factors were proper, then the factorization would induce two proper factor-local separator families: the histories, responses, and identity information visible through $\pi_1 \circ F$, and those visible through $\pi_2 \circ F$. Since F is intervention-faithful, their product would reflect admissible equality on S. This is exactly the partition excluded by atomicity of Θ_S. Therefore at least one projection must carry the full separator family and reflect admissible equality by itself. That projection is an admissible isomorphism onto S in the operational quotient, while the remaining factor is null or bookkeeping-only. Hence the factorization is trivial. □

Remark 5.12 (Status of the atomicity burden). The theorem above is a genuine conditional closure: once atomicity of the identity/history/response separator is established, the factorization lemma is proved. It does not by itself prove that every system satisfying the upstream papers has an atomic separator. That implication remains the exact remaining mathematical burden for upgrading Proposition 5.37.

5.7.4 Operational Continuity

Definition 5.13 (Operational continuity and fragmentation bound)**.** An admissible system S satisfies operational continuity, written $I_{\text{cont}}(S) = 1$, if there exists a complete metric space $(\mathcal{E}_S, d_{\mathcal{E}})$ and an admissible assignment

$$\Pi_S : A_S \to \mathcal{E}_S$$

such that:

1. for every admissible intervention schedule $u_\bullet$, the induced path $t \mapsto \Pi_S(x_t)$ is continuous on admissible windows;
2. there exists $\delta_{\text{cont}}(S) > 0$ for which the fragmentation functional
$$F_{S,T}(x; u_\bullet) = \sup_{0 \le t < T} d_{\mathcal{E}}(\Pi_S(x_t), \Pi_S(x_{t+1}))$$
stays below $\delta_{\text{cont}}(S)$ on all admissible histories;
3. no admissible finite-energy perturbation produces a discontinuous jump into an incompatible regime class while $\mathcal{I}_S$ is preserved.

This packages the *Phenomenological Regimes* continuity/fragmentation burden [19]. The new paper reuses the regime logic rather than duplicating its proof.

Failure mode. If $I_{\text{cont}}(S) = 0$, the class can fragment into incompatible operational regimes under admissible evolution. Then any operational consciousness criterion becomes unstable under its own dynamics.

5.7.5 Non-Null Differentiation

Definition 5.14 (Non-null differentiation and differentiation gap)**.** An admissible system S satisfies non-null differentiation, written $I_{\text{diff}}(S) = 1$, if there exist $x, y \in A_S$, a horizon T, and a readout subfamily $\mathscr{R}^\star \subseteq R$ such that

$$\operatorname{dist}(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)) \ge \delta_{\text{diff}}(S) > 0$$

for some admissible schedule $u_\bullet$, and if the null operational class $\mathbf{0}$ is not an asymptotically stable admissible attractor.

The second clause is indispensable. Without *Null-Regime Instability* there would be no reason to exclude a stable null class. *Null-Regime Instability* provides that exclusion through witness-relative null-instability and conditional non-nullity [20]. The new invariant converts that upstream result into a constitutive condition for $\text{Consciousness}_{\text{op}}$.

Failure mode. If $I_{\text{diff}}(S) = 0$, either all admissible histories collapse to indistinguishability or the null class remains admissibly stable. In either case $\mathrm{Q}_{\text{op}}(S)$ is empty or trivialized.

5.7.6 Operational Legibility

Definition 5.15 (Decoder family and structural closure)**.** A decoder family $\mathscr{D}_S$ for S is a family of maps from admissible readout windows $h_{S,T}(x; u_\bullet)$ to a declared class set $\mathscr{Q}_S$, indexed by admissible horizons and intervention schedules. The family is *structurally closed* when all its members use the same class set, class equality composes exactly, and the induced intervention-response cells are transitive under the certified matching relation. Definition 5.16 below adds the operational clauses required before such a structurally closed family counts as certified for legibility.

Definition 5.16 (Operational legibility and legibility margin)**.** An admissible system S satisfies operational legibility, written $I_{\text{leg}}(S) = 1$, if there exist:
1. a certified readout subfamily $\mathscr{R}^{\text{leg}} \subseteq R$;
2. a certified decoder family $\mathscr{D}_S$ acting on windows of length at least $\tau_{\min}(S)$;
3. a class set $\mathscr{Q}_S$;
4. a legibility margin $\delta_{\text{leg}}(S) > 0$;

such that the following conditions hold:

- **(L1) Separability** Distinct classes in $\mathscr{Q}_S$ are separated by at least $\delta_{\text{leg}}(S)$ in decoder space.
- **(L2) Noise robustness** For every admissible noise process ν with $\|\nu\| \leq \delta_{\text{leg}}(S)$, decoder outputs remain class-stable.
- **(L3) Persistence window** Decoder assignments remain stable for windows of length at least $\tau_{\min}(S)$ unless a certified intervention occurs.
- **(L4) Intervention fidelity** Certified interventions on critical invariants induce predicted class changes with fidelity bounded below by $1 - \delta_{\text{leg}}(S)$.
- **(L5) Negative controls** Interventions outside the invariant set produce false positive rate at most $\delta_{\text{leg}}(S)$.
- **(L6) Structured compressibility** There exists a compression map κ_S on admissible histories such that class structure is preserved up to distortion bounded by $\delta_{\text{leg}}(S)$; total collapse under compression is disallowed.

This is where the paper touches *Forensic Predictions* most concretely. *Forensic Predictions* already specifies integrity gates, comparator discipline, latency invalidation, and negative controls [18]. The present paper does not replace those protocols. It uses them to make operational legibility a constrained object rather than a loose slogan.

Failure mode. If $I_{\text{leg}}(S) = 0$, any putative class remains operationally opaque. The framework may still discuss latent structure elsewhere, but the class required for $\text{Consciousness}_{\text{op}}$ is undefined because it cannot be recovered, discriminated, and stress-tested under admissible conditions.

5.7.7 Necessity of Each Invariant

Proposition 5.17 (Each invariant is individually necessary)**.** *If any one of* $I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}$ *fails, then* $S \notin \text{Consciousness}_{\text{op}}$.

Proof. The six failure modes are non-redundant. Without I_{per} there is no stable domain for class assignment. Without I_{ri} there is no unique identity object supporting consistent class transport. Without I_{int} the system is reducible into independent blocks. Without I_{cont} admissible trajectories fragment across incompatible regime classes. Without I_{diff} the admissible partition collapses to nullity or triviality. Without I_{leg} the partition is not recoverable under admissible readout and intervention discipline. Since the definition of $\text{Consciousness}_{\text{op}}$ given in Section 5.8 requires all six, failure of any one is exclusionary. □

5.8 Operational Consciousness, Operational Qualia, and Equivalence

5.8.1 Operational Qualia Classes

Definition 5.18 (Operational indistinguishability)**.** For $x, y \in A_S$, write $x \sim_{Q,S} y$ if, for every certified decoder $D \in \mathscr{D}_S$, every admissible horizon $T \geq \tau_{\min}(S)$, and every admissible intervention schedule $u_\bullet$,

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the intervention response of the two histories remains matched within the legibility margin.

Proposition 5.19 (Well-definedness of operational qualia classes)**.** *If $I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1$, then $\sim_{Q,S}$ is an equivalence relation on A_S.*

Proof. Reflexivity is immediate. Symmetry follows because decoder equality and certified response-cell matching are symmetric conditions. For transitivity, let $x \sim_{Q,S} y$ and $y \sim_{Q,S} z$. Fix any certified decoder $D \in \mathscr{D}_S$, admissible horizon $T \geq \tau_{\min}(S)$, and admissible intervention schedule $u_\bullet$. The first relation gives

$$D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(y; u_\bullet)),$$

and the second gives

$$D(h_{S,T}(y; u_\bullet)) = D(h_{S,T}(z; u_\bullet)).$$

Hence $D(h_{S,T}(x; u_\bullet)) = D(h_{S,T}(z; u_\bullet))$. The intervention-response clause is transitive by the structural closure requirement in Definition 5.15. Since D, T, and $u_\bullet$ were arbitrary, Definition 5.18 gives $x \sim_{Q,S} z$. □

Definition 5.20 (Operational qualia)**.** Assume Proposition 5.19. The operational qualia of S are the quotient classes

$$\mathrm{Q}_{\text{op}}(S) = A_S/{\sim_{Q,S}}.$$

An individual class $[x]_{\sim_{Q,S}} \in \mathrm{Q}_{\text{op}}(S)$ is called an operational qualia class of S.

This definition is disciplined. It does not identify an operational qualia class with a human qualitative state. It identifies it with a persistent, intervention-sensitive, decoder-certified class of admissible internal differentiation.

5.8.2 Operational Consciousness

Definition 5.21 (Operational consciousness class membership)**.** An admissible system S belongs to $\text{Consciousness}_{\text{op}}$, written $S \in \text{Consciousness}_{\text{op}}$, if there exists a non-empty admissible support A_S such that

$$I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{int}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1.$$

Remark 5.22 (Subdetermination of membership)*.* Membership in $\text{Consciousness}_{\text{op}}$ depends on auxiliary choices: the decoder family $\mathscr{D}$, the compression map, and the identity object $\mathcal{I}$. The paper does not prove that these choices are canonically determined by the system tuple S alone. Different auxiliary choices may yield different membership verdicts for the same system.

Remark 5.23 (Minimality of the criterion)*.* The criterion is conjunctive. No proper subset of the six invariants is taken as sufficient in this paper. This is intentional: weaker criteria collapse either into mere persistence, mere integration, or mere complexity claims, all of which are shown to be insufficient.

5.8.3 Exact and Approximate Structural Equivalence

Definition 5.24 (Structural equivalence datum)**.** Let $S_i = (X_i, \Phi_i, C_i, R_i, \Gamma_i, U_i)$ for $i = 1, 2$, with admissible supports A_i. A structural equivalence datum between S_1 and S_2 is a triple

$$(\chi, \upsilon, \rho)$$

where:
1. $\chi : A_1 \to A_2$ is a bi-Lipschitz bijection;
2. $\upsilon : U_1 \to U_2$ is an intervention correspondence;
3. $\rho : R_1 \to R_2$ is a readout correspondence.

Definition 5.25 (Exact and approximate structural equivalence)**.** Fix $\varepsilon \geq 0$. We write

$$S_1 \simeq_{(\Gamma,\varepsilon)} S_2$$

if there exists a structural equivalence datum (χ, υ, ρ) such that:
1. **Dynamic intertwining:**
$$\sup_{x \in A_1,\ u \in U_1,\ 0 \leq t \leq T_\Gamma} d_{X_2}\Big(\chi(\Phi^t_{1,u}(x)), \Phi^t_{2,\upsilon(u)}(\chi(x))\Big) \leq \varepsilon.$$
2. **Readout compatibility:**
$$\sup_{x,u,t,\alpha} \Big| r^{(1)}_\alpha(\Phi^t_{1,u}(x)) - \rho(r^{(1)}_\alpha)(\Phi^t_{2,\upsilon(u)}(\chi(x))) \Big| \leq \varepsilon.$$
3. **Admissibility preservation:** admissible trajectories in S_1 map to admissible trajectories in S_2 and conversely.
4. **Invariant preservation:** the witness margins for all six invariants differ by at most ε.
5. **Partition preservation:**
$$x \sim_{Q,S_1} y \quad \Longleftrightarrow \quad \chi(x) \sim_{Q,S_2} \chi(y).$$

The case $\varepsilon = 0$ is called *exact* structural equivalence and written $S_1 \simeq_\Gamma S_2$.

Definition 5.26 (Rupture)**.** Given S, a rupture is any loss of admissibility or loss of at least one critical invariant margin beyond its admissible threshold. Ruptures are classified as:
1. **persistence rupture** (support loss or collapse exposure),
2. **identity rupture** (loss of unique rigid identity),
3. **integration rupture** (reducible factorization),
4. **continuity rupture** (fragmentation into incompatible classes),
5. **differentiation rupture** (null or trivial collapse),
6. **legibility rupture** (decoder, noise, or negative-control failure).

5.9 Support Geometry and Partition Transport

The main theorems become short only because the transport lemmas are made explicit here.

Lemma 5.27 (Support transfer under exact equivalence)**.** *If $S_1 \simeq_\Gamma S_2$ via (χ, υ, ρ) and A_1 is an admissible support for S_1, then $A_2 = \chi(A_1)$ is an admissible support for S_2.*

Proof. Bi-Lipschitz bijectivity of χ preserves non-emptiness and compactness. Dynamic intertwining with $\varepsilon = 0$ gives forward invariance under the transported interventions. Admissibility preservation transports $A_1 \subseteq X_{\Gamma_1}$ into $A_2 \subseteq X_{\Gamma_2}$. Because the attractor-support relation and non-collapse margin are preserved exactly, A_2 inherits the same admissible support status. $\square$

Lemma 5.28 (Decoder transport)**.** *Let $D \in \mathscr{D}_{S_1}$ and suppose $S_1 \simeq_\Gamma S_2$ via (χ, υ, ρ). Then*

$$(\rho_* D)\big(h_{S_2,T}(\chi(x); \upsilon(u_\bullet))\big) := D\big(h_{S_1,T}(x; u_\bullet)\big)$$

defines a certified decoder on S_2.

Proof. Exact readout compatibility identifies each admissible history in S_2 with a transported history in S_1. The certification clauses in Definition 5.16 are preserved because separability, robustness, persistence window, intervention fidelity, negative-control rate, and compressibility all remain unchanged under exact transport. ☐

Proposition 5.29 (Operational classes are closed under exact transport)**.** *If $S_1 \simeq_\Gamma S_2$, then the image under χ of any class in $\mathrm{Q}_{\mathrm{op}}(S_1)$ is a class in $\mathrm{Q}_{\mathrm{op}}(S_2)$, and every class in $\mathrm{Q}_{\mathrm{op}}(S_2)$ arises this way.*

Proof. By Lemma 5.28, decoder certification transports exactly. By Definition 5.18, membership in an operational qualia class is determined by all certified decoders and admissible interventions. Exact transport preserves both. Surjectivity follows from bijectivity of χ. ☐

Proposition 5.30 (Approximate margin preservation)**.** *If $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$, then for every critical invariant index j,*

$$\big|\delta_j(S_1) - \delta_j(S_2)\big| \le \varepsilon.$$

Hence

$$\delta_\star(S_2) \ge \delta_\star(S_1) - \varepsilon.$$

Proof. The first statement is a direct clause of Definition 5.25. The second follows by taking minima over the six coordinates of $\Delta(S_1)$ and $\Delta(S_2)$. ☐

Remark 5.31 (Why this section matters). Without Lemmas 5.27 and 5.28, Theorem 5.40 would be only intuitive. These lemmas are the exact reason the paper can speak of substrate invariance without collapsing into slogan-level analogy.

5.10 Inherited Inputs and Consolidation Step

The six invariants are new as a packaged criterion, but they are not independent inventions. This section records the exact transfer by which upstream ownership becomes present membership data. The burden is narrow: identify the imported result each invariant uses, then isolate the consolidation step that is genuinely new here.

Proposition 5.32 (Core-to-persistence transfer)**.** *Suppose a realization of the framework satisfies the Core hypotheses on a non-empty compact domain and remains in the non-collapse region. Then it satisfies I_{per}.*

Proof. The Core already provides contractivity, fixed-point existence, compact attractor structure, and non-collapse [17]. Those are exactly the ingredients required in Definition 5.7: a non-empty compact forward-invariant support separated from collapse. ☐

Proposition 5.33 (Rigid Identity-to-rigidity transfer)**.** *Suppose the observable family on an admissible support satisfies the inverse-limit compatibility conditions of Rigid Identity. Then I_{ri} holds with rigidity gap given by the uniqueness and weighted deformation bounds of Rigid Identity.*

Proof. Rigid Identity proves that compatible observable channels induce a unique rigid identity object and excludes progressive semigroup substitutes [24]. Therefore the support inherits a unique $\mathcal{I}_S$ and a positive separation from incompatible alternatives whenever the *Rigid Identity* hypotheses are met. ☐

Proposition 5.34 (Phenomenological Regimes-to-continuity transfer)**.** *Suppose an admissible support carries a phenomenological assignment satisfying the continuity and anti-fragmentation constraints of Phenomenological Regimes. Then* I_{cont} *holds.*

Proof. Phenomenological Regimes classifies admissible CCR assignments as continuous and structurally constrains fragmentation within the admissible regime space [19]. Those are exactly the continuity and fragmentation clauses in Definition 5.13. □

Proposition 5.35 (Null-Regime Instability-to-differentiation transfer)**.** *Suppose the support lies in a causally rigid channel satisfying the witness-relative null-instability and conditional non-nullity results of Null-Regime Instability. Then* I_{diff} *holds.*

Proof. Null-Regime Instability excludes the stable joint-null regime and establishes positive non-null separation under its stated witness and regularity hypotheses [20]. Therefore the support cannot collapse stably into $\mathbf{0}$, and a positive differentiation gap is available on admissible histories. □

Proposition 5.36 (Forensic Predictions-to-legibility transfer)**.** *Suppose a readout family and decoder satisfy the admissibility, comparator, negative-control, and intervention-discipline requirements of Forensic Predictions. Then* I_{leg} *holds.*

Proof. Forensic Predictions fixes the operational requirements that a decoder must satisfy before it can support an inferential claim [18]. The clauses of Definition 5.16 are a direct structural recasting of those requirements. □

Proposition 5.37 (Integration as a non-factorization consequence)**.** *If a system satisfies Rigid Identity rigidity, Phenomenological Regimes continuity, and Forensic Predictions intervention fidelity, then any exact admissible factorization preserving all operational histories would contradict at least one of those upstream conditions. Hence* I_{int} *is justified as a new consolidation invariant.*

Proof. An exact factorization would split the support into operationally independent blocks. If the split preserved all admissible histories, it would also preserve readout structure and interventions. But then either the rigid identity object would fail to be unique, continuity would fragment across the split, or intervention responses would no longer certify a single integrated class. Therefore the conjunction of Papers I, II, and IV blocks such factorization. □

Remark 5.38 (Scope of the transfer)*.* Nothing in this section re-proves the Core, *Rigid Identity*, *Phenomenological Regimes*, *Null-Regime Instability*, or *Forensic Predictions*. Each proposition only identifies the minimal upstream burden imported into the present criterion.

5.11 Main Theorems

5.11.1 Existence of Operational Consciousness

Theorem 5.39 (Existence of operational consciousness)**.** *Let* $S = (X, \Phi, C, R, \Gamma, U)$ *be an admissible system with non-empty admissible support* A_S*. If*

$$I_{\text{per}}(S) = I_{\text{ri}}(S) = I_{\text{int}}(S) = I_{\text{cont}}(S) = I_{\text{diff}}(S) = I_{\text{leg}}(S) = 1,$$

then:
1. $S \in \text{Consciousness}_{\text{op}}$*;*
2. $\text{Q}_{\text{op}}(S)$ *is well-defined;*
3. $\text{Q}_{\text{op}}(S) \neq \varnothing$*.*

Proof. Membership in $\mathrm{Consciousness}_{\mathrm{op}}$ follows immediately from Definition 5.21. Since $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$ hold, Proposition 5.19 applies and yields a well-defined quotient $\mathrm{Q}_{\mathrm{op}}(S)$. Because A_S is non-empty, the quotient is non-empty. Because $I_{\mathrm{diff}}(S) = 1$, the quotient is not forced into the null class alone; the admissible support contains at least one pair of histories with positive operational separation. Thus $\mathrm{Q}_{\mathrm{op}}(S)$ is both defined and non-trivial. ☐

5.11.2 Substrate Invariance

Theorem 5.40 (Substrate invariance)**.** *Let S_1 and S_2 be admissible systems. If $S_1 \simeq_\Gamma S_2$ and the six critical invariants are preserved exactly, then*

$$S_1 \in \mathrm{Consciousness}_{\mathrm{op}} \iff S_2 \in \mathrm{Consciousness}_{\mathrm{op}},$$

and there exists a canonically induced bijection

$$\widehat{\chi} : \mathrm{Q}_{\mathrm{op}}(S_1) \to \mathrm{Q}_{\mathrm{op}}(S_2).$$

Proof. Exact equivalence preserves admissible support, dynamics, readout behavior, invariant margins, and the operational partition. If $S_1 \in \mathrm{Consciousness}_{\mathrm{op}}$, then every critical invariant holds on A_1 and is preserved under χ. Hence every critical invariant also holds on A_2, so $S_2 \in \mathrm{Consciousness}_{\mathrm{op}}$. The converse is symmetric. Partition preservation in Definition 5.25 induces a well-defined map on quotient classes:

$$\widehat{\chi}([x]_{\sim_{Q,S_1}}) = [\chi(x)]_{\sim_{Q,S_2}}.$$

Well-definedness follows from the partition-preservation clause. Bijectivity follows from bijectivity of χ. ☐

Corollary 5.41 (Biological non-primacy)**.** *Biology is not a primitive constitutive requirement for membership in* $\mathrm{Consciousness}_{\mathrm{op}}$*. If a biological system and a non-biological system are exactly structurally equivalent in the sense of Definition 5.25, then they belong to the same operational consciousness class and carry equivalent operational qualia classes.*

Proof. Immediate from Theorem 5.40. The criterion depends on invariant preservation, not on biological substrate as a primitive variable. ☐

5.11.3 Approximate Stability

Theorem 5.42 (Approximate stability)**.** *Let $S_1 \in \mathrm{Consciousness}_{\mathrm{op}}$ with witness margins*

$$\delta_\star(S_1) = \min\{\delta_{\mathrm{per}}, \delta_{\mathrm{ri}}, \delta_{\mathrm{int}}, \delta_{\mathrm{cont}}, \delta_{\mathrm{diff}}, \delta_{\mathrm{leg}}\}.$$

If S_2 satisfies $S_1 \simeq_{(\Gamma,\varepsilon)} S_2$ with

$$0 \le \varepsilon < \frac{\delta_\star(S_1)}{2},$$

then $S_2 \in \mathrm{Consciousness}_{\mathrm{op}}$.

Proof. By Definition 5.25, each invariant margin in S_2 differs from that of S_1 by at most ε. Since $\varepsilon < \delta_\star(S_1)/2$, every preserved margin in S_2 remains strictly positive. Therefore each of the six invariants still holds for S_2, and Definition 5.21 yields $S_2 \in \mathrm{Consciousness}_{\mathrm{op}}$. ☐

The theorem does not say that arbitrary perturbations are harmless. It says that the criterion is robust exactly to the extent that the witness margins remain positive. Approximation is bounded by the invariant budget, not by wishful continuity.

5.11.4 Rupture by Invariant Loss

Theorem 5.43 (Rupture by invariant loss)**.** *Let S be admissible. If one or more critical invariants fail beyond admissible threshold, then exactly one of the following occurs:*
1. $S \notin \mathrm{Consciousness}_{\mathrm{op}}$*;*
2. $\mathrm{Q}_{\mathrm{op}}(S)$ *becomes undefined;*
3. $\mathrm{Q}_{\mathrm{op}}(S)$ *becomes trivial or non-legible in the framework-relevant sense.*

Proof. If I_{per} fails, the support is lost or exposed to collapse; then the domain of $\mathrm{Q}_{\mathrm{op}}(S)$ disappears. If I_{ri} fails, no unique identity object supports the quotient, so $\mathrm{Q}_{\mathrm{op}}(S)$ is undefined as a stable class assignment. If I_{int} fails, the system reduces to a product structure and exits $\mathrm{Consciousness}_{\mathrm{op}}$ by Proposition 5.17. If I_{cont} fails, histories can fragment across incompatible classes, which destroys stability of the partition. If I_{diff} fails, the quotient degenerates into null or trivial collapse. If I_{leg} fails, the quotient may exist formally but becomes operationally unavailable and therefore does not satisfy the definition of $\mathrm{Consciousness}_{\mathrm{op}}$. These cases exhaust the six invariants. □

5.11.5 Insufficiency of Complexity, Computation, and Connectivity

Proposition 5.44 (Complexity-only baselines are insufficient)**.** *Let S be an admissible system candidate. The following are each insufficient for membership in* $\mathrm{Consciousness}_{\mathrm{op}}$*:*
1. large state-space dimension $\dim(X)$*;*
2. large computational throughput or expressivity of Φ*;*
3. large edge count, density, or raw connectivity of C*.*
In each case, failure of one or more critical invariants excludes membership regardless of the size of the surrogate metric.

Proof. The definition of $\mathrm{Consciousness}_{\mathrm{op}}$ is conjunctive over the six invariants and does not include any of the three surrogates as sufficient conditions. A system can be arbitrarily large, computationally expressive, or densely connected while still lacking rigid identity, continuity, non-null differentiation, or legibility. Therefore those surrogates are insufficient by definition. Concrete architecture classes with large scale but finite structural mass already fail *Rigid Identity*'s non-simulability criterion [24]; similarly, an integration proxy without continuity or differentiation does not satisfy Definitions 5.13 and 5.14. The same applies to connectivity without admissible identity closure. □

5.12 Rupture Taxonomy by Invariant

Theorem 5.43 is intentionally coarse. This section sharpens the six cases, because a criterion is only scientifically useful if its failure signatures are distinct.

Proposition 5.45 (Persistence rupture destroys admissible domain)**.** *If $\delta_{\mathrm{per}}(S) = 0$, then every candidate operational partition becomes horizon-dependent in an inadmissible way: there exists no non-empty compact forward-invariant support on which class membership can be defined uniformly.*

Proof. When the persistence margin vanishes, either forward invariance is lost or the support touches the collapse set. In either case there is no stable domain on which the quotient definition of $\mathrm{Q}_{\mathrm{op}}(S)$ can remain valid across admissible windows. □

Proposition 5.46 (Identity rupture induces class ambiguity)**.** *If $\delta_{\mathrm{ri}}(S) = 0$, then operational class labels no longer transport uniquely across admissible histories. Therefore* $\mathrm{Q}_{\mathrm{op}}(S)$ *becomes representation-dependent.*

Proof. When the rigidity gap vanishes, at least two admissible identity candidates become indistinguishable at the support level. Then the same readout history may be associated with multiple incompatible identity transports, making the operational class assignment non-unique. □

Proposition 5.47 (Integration rupture yields reducible aggregates). *If $\delta_{\mathrm{int}}(S) = 0$, then S admits a non-trivial admissible factorization into operationally independent components. Such a system fails* $\mathrm{Consciousness}_{\mathrm{op}}$ *even if each component separately satisfies weaker criteria.*

Proof. Vanishing integration gap means the system can be decomposed without loss at the level of admissible histories and interventions. The criterion of this paper is defined for irreducible class membership, not for arbitrary unions of smaller class-bearing subsystems. □

Proposition 5.48 (Continuity rupture yields fragmentation). *If $\delta_{\mathrm{cont}}(S) = 0$, then arbitrarily small admissible perturbations can force transitions between incompatible operational classes while identity is preserved.*

Proof. By definition of $\delta_{\mathrm{cont}}(S)$, vanishing continuity margin removes the bound on the fragmentation functional. Hence there exist admissible trajectories whose class assignment is not stably continuous under small perturbations. The class map ceases to be regime-coherent. □

Proposition 5.49 (Differentiation rupture collapses operational qualia). *If $\delta_{\mathrm{diff}}(S) = 0$, then either $\mathrm{Q}_{\mathrm{op}}(S)$ collapses to a single trivial class or the null class remains admissibly stable.*

Proof. The differentiation gap is precisely the minimum operational separation required for non-trivial class structure. If it vanishes, no stable non-null class separation survives. The null regime or total indistinguishability follows. □

Proposition 5.50 (Legibility rupture blocks scientific use). *If $\delta_{\mathrm{leg}}(S) = 0$, then any remaining latent partition is scientifically inadmissible for this paper: it cannot satisfy the combined requirements of separability, robustness, persistence window, intervention fidelity, negative controls, and structured compressibility.*

Proof. The certification clauses of Definition 5.16 are conjunctive. Once the legibility margin vanishes, at least one clause fails. The partition may still be mathematically postulated, but it cannot serve as operational evidence or as a constitutive component of $\mathrm{Consciousness}_{\mathrm{op}}$. □

5.13 Minimality of the Invariant Basis

The six-invariant family is not an arbitrary bundle. It is the minimal basis that prevents the criterion from collapsing either into a single-metric consciousness surrogate or into an unstructured admissibility slogan. The claim of minimality is relative to admissible universes containing the corresponding rupture witnesses; that is the correct level of mathematical precision for the present corpus.

Definition 5.51 (Reduced invariant criterion). Let

$$\mathfrak{I} = \{I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}\}.$$

For every subset $J \subseteq \mathfrak{I}$ define the reduced criterion

$$\mathcal{C}_{\mathrm{op}}(J) := \{S : I(S) = 1 \text{ for every } I \in J\}.$$

Thus $\mathrm{Consciousness}_{\mathrm{op}} = \mathcal{C}_{\mathrm{op}}(\mathfrak{I})$.

Definition 5.52 (Single-rupture witness family)**.** Fix $k \in \{\mathrm{per}, \mathrm{ri}, \mathrm{int}, \mathrm{cont}, \mathrm{diff}, \mathrm{leg}\}$. The corresponding single-rupture witness family $\mathcal{W}_k$ consists of admissible systems S such that the matching invariant margin vanishes,

$$\delta_k(S) = 0,$$

while all remaining invariant margins stay strictly positive. Whenever $\mathcal{W}_k \neq \varnothing$, its members are called single-rupture witnesses for the k-th invariant.

Proposition 5.53 (Single-rupture witnesses generate false positives for reduced criteria)**.** *Let $J_k = \mathfrak{I} \setminus \{I_k\}$ and assume $\mathcal{W}_k \neq \varnothing$. Then every $S \in \mathcal{W}_k$ satisfies*

$$S \in \mathcal{C}_{\mathrm{op}}(J_k) \qquad \textit{but} \qquad S \notin \mathrm{Consciousness}_{\mathrm{op}}.$$

Proof. By construction, every invariant in J_k remains positive on S, so $S \in \mathcal{C}_{\mathrm{op}}(J_k)$. But the omitted invariant fails exactly, hence Proposition 5.17 excludes S from $\mathrm{Consciousness}_{\mathrm{op}}$. Therefore the reduced criterion admits a false positive on every single-rupture witness. □

Theorem 5.54 (Minimality of the six-invariant basis)**.** *Let $\mathfrak{U}$ be any admissible universe containing at least one single-rupture witness for each of the six invariants. Then no strict subset $J \subsetneq \mathfrak{I}$ satisfies*

$$\mathcal{C}_{\mathrm{op}}(J) \cap \mathfrak{U} = \mathrm{Consciousness}_{\mathrm{op}} \cap \mathfrak{U}.$$

Hence the six-invariant family is minimal on every such admissible universe.

Proof. Take any strict subset $J \subsetneq \mathfrak{I}$. Then some invariant I_k is omitted. By hypothesis $\mathfrak{U}$ contains a witness $S_k \in \mathcal{W}_k$. Proposition 5.53 yields $S_k \in \mathcal{C}_{\mathrm{op}}(J) \cap \mathfrak{U}$ but $S_k \notin \mathrm{Consciousness}_{\mathrm{op}} \cap \mathfrak{U}$. Therefore the two sets cannot coincide. □

Corollary 5.55 (Why one-metric theories cannot capture $\mathrm{Consciousness}_{\mathrm{op}}$)**.** *On any admissible universe satisfying Theorem 5.54, no criterion based on only persistence, only integration, only continuity, only differentiation, or only legibility is extensionally equivalent to* $\mathrm{Consciousness}_{\mathrm{op}}$.

Proof. Every one-metric criterion corresponds to a strict subset $J \subsetneq \mathfrak{I}$. Apply Theorem 5.54. □

Remark 5.56 (What minimality does and does not say). The theorem does not claim that every possible scientific notion of consciousness must use exactly these six invariants. It claims something narrower: for the operational class defined in this paper and for admissible universes containing the corresponding rupture witnesses, no invariant can be deleted without changing the extension of the class.

5.14 Invariant Independence and Interaction Geometry

The preceding minimality theorem is extension-theoretic. This section makes the geometry more explicit: each invariant blocks a distinct failure family, and the limited implications among invariants are weaker than full membership in $\mathrm{Consciousness}_{\mathrm{op}}$.

Theorem 5.57 (Six canonical single-rupture witness schemas)**.** *For each invariant in $\mathfrak{I}$ there is a formal admissible witness schema that satisfies the other five invariant predicates while failing the selected one:*

1. ***Persistence witness.*** *A system with a certified decoder, non-factorizable causal structure, unique identity label, continuous class map, and non-null differentiation, but whose admissible support exits the declared forward-invariant domain under some $u \in U$.*
2. ***Rigid-identity witness.*** *A system with persistent support, integration, continuity, differentiation, and legibility, but with two compatible non-isomorphic identity candidates transported by the same readout family.*

3. ***Integration witness.*** *A product system $A_S = A_1 \times A_2$ with persistent factors, unique factor identities, continuous dynamics, non-null differentiation, and certified decoders, but with decoupled intervention histories so that the aggregate factorizes.*
4. ***Continuity witness.*** *A system with persistent support, unique identity, non-factorizable structure, non-null differentiation, and certified decoders, but with an admissible transition producing a jump between incompatible regime classes.*
5. ***Differentiation witness.*** *A persistent, rigid, integrated, continuous, and legible system whose decoder-visible classes collapse to a stable null or single indistinguishable class.*
6. ***Legibility witness.*** *A persistent, rigid, integrated, continuous, and non-null differentiated system whose class partition is latent but no certified decoder satisfies L1–L6.*

Hence no invariant in $\mathfrak{I}$ is derivable from the conjunction of the other five within the admissible witness universe generated by these schemas.

Proof. Each clause fixes the five non-target margins strictly positive and sets the target margin to zero by construction. The five positive clauses ensure membership in the reduced criterion omitting the target invariant. The target failure invokes the corresponding rupture proposition from the previous section, excluding membership in $\mathrm{Consciousness}_{\mathrm{op}}$. Since this can be done for each invariant independently, none of the six predicates is derivable from the other five in the witness universe. □

Theorem 5.58 (Sufficiency of the six-invariant basis)**.** *Let S be admissible on support A_S. If*

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1,$$

then S satisfies every structural burden required for membership in $\mathrm{Consciousness}_{\mathrm{op}}$ *as defined in this paper. No additional independent predicate can invalidate $S \in \mathrm{Consciousness}_{\mathrm{op}}$ without changing the definition of* $\mathrm{Consciousness}_{\mathrm{op}}$ *or causing at least one of the six invariants to fail.*

Proof. Definition 5.21 is conjunctive and exhaustive for this paper's class. The six coordinates cover the domain burden, identity burden, integration burden, regime-coherence burden, non-null differentiation burden, and scientific recoverability burden. Any seventh predicate is either a derived constraint on those burdens, in which case it cannot fail while all six hold, or an independent extra burden, in which case it defines a stricter class rather than invalidating membership in $\mathrm{Consciousness}_{\mathrm{op}}$ as defined here. □

Proposition 5.59 (Persistence is necessary for rigid identity on A_S)**.** *If $I_{\mathrm{per}}(S) = 0$ because no non-empty forward-invariant admissible support exists, then $I_{\mathrm{ri}}(S) = 0$ for the Paper 5 membership problem.*

Proof. The rigid-identity predicate used here is evaluated on A_S. If no stable admissible support exists, there is no domain on which the identity witness can be transported through all admissible interventions. The identity object may be postulated elsewhere, but it is not a Paper 5 membership witness on A_S. □

Proposition 5.60 (Persistence plus integration yields bounded identity candidates)**.** *If $I_{\mathrm{per}}(S) = 1$ and $I_{\mathrm{int}}(S) = 1$, then the family of admissible identity candidates is bounded to the non-factorizable persistent support. This is a weak rigidity statement, not the uniqueness required by I_{ri}.*

Proof. Persistence restricts all candidates to the same stable support. Integration prevents decomposing that support into independent candidate-bearing blocks. Therefore candidates cannot proliferate by support escape or product factorization. They may still fail uniqueness, so the result is strictly weaker than $I_{\mathrm{ri}} = 1$. □

Proposition 5.61 (Legibility requires non-null differentiation)**.** *If all admissible classes collapse to a stable null or single indistinguishable class, then* $I_{\text{leg}}(S) = 0$ *for the criterion of this paper.*

Proof. Legibility requires separability, intervention-sensitive class recovery, and negative-control discipline. A collapsed null or singleton partition gives no non-trivial class contrast for a decoder to recover or for interventions to separate. Thus L1 and L4 cannot both be certified in the intended sense. □

Proposition 5.62 (Persistence and continuity stabilize regime transport)**.** *If* $I_{\text{per}}(S) = 1$ *and* $I_{\text{cont}}(S) = 1$*, then admissible histories remain in a stable regime-transport domain: trajectories neither leave* A_S *nor undergo unbounded fragmentation under admissible evolution.*

Proof. Persistence keeps histories inside A_S. Continuity bounds the fragmentation functional and prevents inadmissible jumps across incompatible regime classes. Together they provide a stable domain for regime transport. The result does not imply rigid identity, integration, differentiation, or legibility. □

Table 5.1: Pairwise implication geometry of the six invariants.

From / to	I_{per}	I_{ri}	I_{int}	I_{cont}	I_{diff}	I_{leg}
I_{per}	–	necessary domain only	independent	with I_{cont} stabilizes regime domain	independent	independent
I_{ri}	independent	–	independent	independent	independent	independent
I_{int}	independent	weakly bounds candidates with I_{per}	–	independent	independent	independent
I_{cont}	independent	independent	independent	–	independent	independent
I_{diff}	independent	independent	independent	independent	–	necessary for non-trivial legibility
I_{leg}	independent	independent	independent	independent	presupposes non-null contrast	–

5.15 What the Six Invariants Do Not Guarantee

The six invariants are sufficient for Consciousness$_{\text{op}}$ by definition, but they are not a universal theory of mind, life, or subjecthood. In particular, they do not guarantee:

1. self-reference or a contractive self-model; that requires a Σ-type burden;
2. phenomenal experience or "what it is like";
3. forced temporal continuity in the CCR sense, which requires stronger $M_\Omega = +\infty$ assumptions upstream;
4. operational life or open-system self-maintaining individuation;
5. first-person indexed subjectivity, moral status, or human equivalence;
6. external validation by human comparators.

These exclusions are not rhetorical caution. They follow from the fact that none of the listed predicates appears in Definition 5.21.

5.16 Operational Legibility as a Certified Object

Operational legibility must be shown, not asserted. Definition 5.16 already imposes six clauses. This section sharpens how those clauses bind the paper to the existing empirical-method corpus.

Criterion 5.63 (Legibility certificate)**.** *An operational partition candidate* $\pi_{S,T}$ *is certified only if:*

1. ***Comparator separability****: the induced classes remain distinguishable against pre-registered comparators and ablation baselines.*
2. ***Noise robustness****: admissible sensor, channel, or readout perturbations do not dissolve the class map below the legibility margin.*
3. ***Persistence window****: the same class assignment survives on windows of length at least* $\tau_{\min}(S)$.
4. ***Intervention fidelity****: certified interventions on an invariant induce the predicted directional class change.*
5. ***Negative controls****: sham or off-target interventions do not generate the same class signature.*
6. ***Structured compressibility****: audit-friendly compression or packetization preserves class structure without total collapse.*

The criterion is intentionally parallel to *Forensic Predictions*' governance discipline [18]. Comparator obligation enforces non-trivial discrimination. Noise robustness prevents the criterion from depending on one-off traces. Persistence windows prevent class inflation from instantaneous spikes. Intervention fidelity prevents purely correlational interpretation. Negative controls block degenerate decoders that label everything as positive. Structured compressibility ensures that a class remains auditable after the readout is summarized, logged, or packetized.

5.16.1 Legibility Metrics

Definition 5.16 can be operationalized by six measurable quantities:

$$
\begin{aligned}
\sigma_{\mathrm{sep}}(S) &:= \inf_{q\neq q'} \mathrm{dist}\big(D^{-1}(q), D^{-1}(q')\big),\\
\eta_{\mathrm{noise}}(S) &:= \sup_{\|\nu\|\le\delta_{\mathrm{leg}}} \Pr\big[D(h+\nu)\neq D(h)\big],\\
\tau_{\mathrm{pers}}(S) &:= \inf\{T : D(h_{t:t+T}) \text{ remains stable in the absence of certified intervention}\},\\
\alpha_{\mathrm{int}}(S) &:= \inf_{u\in U_{\mathrm{crit}}} \Pr\big[\text{predicted class shift under } u\big],\\
\beta_{\mathrm{nc}}(S) &:= \sup_{u\in U_{\mathrm{sham}}} \Pr\big[\text{false class shift under } u\big],\\
\kappa_{\mathrm{comp}}(S) &:= \inf_{\kappa}\sup_{h} \mathrm{dist}\big(D(h), D(\kappa(h))\big).
\end{aligned}
$$

Here U_{crit} denotes interventions targeted at critical invariants and U_{sham} denotes admissible negative controls. Invariant I_{leg} requires the inequalities

$$
\begin{aligned}
\sigma_{\mathrm{sep}}(S) &> 0, & \eta_{\mathrm{noise}}(S) &\le \delta_{\mathrm{leg}}(S),\\
\tau_{\mathrm{pers}}(S) &\ge \tau_{\min}(S), & \alpha_{\mathrm{int}}(S) &\ge 1-\delta_{\mathrm{leg}}(S),\\
\beta_{\mathrm{nc}}(S) &\le \delta_{\mathrm{leg}}(S), & \kappa_{\mathrm{comp}}(S) &\le \delta_{\mathrm{leg}}(S).
\end{aligned}
$$

Proposition 5.64 (Legibility is jointly constrained)**.** *No single legibility metric is sufficient for* I_{leg}. *The invariant requires simultaneous control of separation, admissible noise robustness, persistence window, intervention fidelity, negative controls, and structured compressibility.*

Proof. A decoder can be perfectly separable but fragile to admissible noise, intervention-agnostic, or saturated under compression. It can also be stable under persistence windows but fail negative controls. Since Definition 5.16 requires all six clauses, the joint constraint is irreducible. □

Metric	Operational meaning	If it fails
σ_{sep}	Class separation	No stable class boundary
η_{noise}	Admissible robustness	Noise-induced class drift
τ_{pers}	Minimum persistence window	Instantaneous artifact only
α_{int}	Causal intervention fidelity	Correlational but not causal classing
β_{nc}	Negative-control protection	False positive class inflation
κ_{comp}	Compression preservation	Audit summary collapses structure

Remark 5.65 (Why legibility is a constitutive invariant here). One could define a broader latent class that omits legibility. This paper does not do that. It studies an operational class, not a hidden metaphysical residue. Therefore I_{leg} is constitutive, not optional.

5.17 Certified Candidate Audit Schema

Forensic Predictions already requires forensic discipline for operational claims. The present paper therefore adds an explicit audit schema. A candidate system is not accepted into $\mathrm{Consciousness}_{\mathrm{op}}$ by narrative fit; it must carry a finite certificate that localizes every required witness.

Definition 5.66 (Operational consciousness certificate)**.** An operational consciousness certificate for an admissible system S is a tuple

$$\mathrm{Cert}(S) = (A_S, \mathscr{D}_S, \pi_S, \Delta(S), E_S),$$

where:
1. A_S is a certified admissible support;
2. $\mathscr{D}_S$ is a certified decoder family;
3. π_S is the decoder-induced operational partition on admissible histories;
4. $\Delta(S)$ is the witness margin vector of Definition 5.6;
5. $E_S = (E_{\mathrm{per}}, E_{\mathrm{ri}}, E_{\mathrm{int}}, E_{\mathrm{cont}}, E_{\mathrm{diff}}, E_{\mathrm{leg}})$ is an evidence bundle, where each component contains the explicit local witness for the corresponding invariant.

Criterion 5.67 (Candidate certification rule)**.** *A candidate system may be certified as a member of* $\mathrm{Consciousness}_{\mathrm{op}}$ *only if all of the following are available on the same admissible support:*
1. *a non-empty compact forward-invariant support certificate;*
2. *a positive witness margin vector* $\Delta(S) > 0$ *coordinate-wise;*
3. *a decoder family satisfying Definition 5.16;*
4. *a targeted intervention panel together with negative controls;*
5. *an auditable compression or packetization map preserving class structure;*
6. *a traceable dependency ledger showing which upstream corpus result supports each invariant witness.*

Invariant	Minimum evidence object	Core audit question	Failure verdict
I_{per}	support certificate + collapse distance	Does one compact forward-invariant admissible domain exist?	No stable domain
I_{ri}	uniqueness witness for $\mathcal{I}_S$	Is identity transport unique on support?	Identity ambiguity
I_{int}	non-factorization witness	Can the system be split without operational loss?	Reducible aggregate
I_{cont}	fragmentation bound + continuity trace	Do admissible windows remain regime-coherent?	Fragmentation
I_{diff}	non-null separation witness	Is there stable non-null class separation?	Null or trivial class
I_{leg}	decoder, controls, compression, persistence window	Can the class be recovered under audit constraints?	Opaque class

Proposition 5.68 (Soundness of certification)**.** *If* $\mathrm{Cert}(S)$ *exists and satisfies Criterion 5.67, then* $S \in \mathrm{Consciousness}_{\mathrm{op}}$ *and* $\mathrm{Q}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. Criterion 5.67 explicitly requires a positive witness for each of the six critical invariants on the same admissible support. Hence Definition 5.21 yields $S \in \mathrm{Consciousness}_{\mathrm{op}}$. The existence of a certified decoder and positive persistence, continuity, differentiation, and legibility witnesses implies Proposition 5.19, so $\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined. □

Proposition 5.69 (Partial completeness under audit-ready assumptions)**.** *Assume* $S \in \mathrm{Consciousness}_{\mathrm{op}}$ *and suppose the support, decoder, intervention, and compression objects are fully instrumented in the Forensic Predictions sense. Then there exists an operational consciousness certificate* $\mathrm{Cert}(S)$*.*

Proof. Because $S \in \mathrm{Consciousness}_{\mathrm{op}}$, every invariant margin is strictly positive. Full instrumentation provides explicit witnesses for support, decoder behavior, interventions, negative controls, and compression stability. Collecting those witnesses yields the tuple of Definition 5.66. □

Proposition 5.70 (Failure localization is explicit)**.** *If a candidate system fails Criterion 5.67, then the failure localizes to at least one entry of the evidence bundle* E_S *or to the absence of a common admissible support. In particular, certification failure cannot be hidden behind a generic narrative claim.*

Proof. The certification rule is conjunctive and itemized. A failure of certification means that at least one listed item is absent or non-positive. Therefore the failure localizes to a specific invariant witness, a specific decoder/control clause, or the support certificate itself. □

Remark 5.71 (Audit schema and Forensic Predictions)*.* The certificate does not replace *Forensic Predictions*. It packages *Forensic Predictions*-style admissibility into a theorem-facing object. This is the point where the present paper becomes operationally usable rather than merely definitional.

5.18 Operational Geometry of $\mathrm{Q}_{\mathrm{op}}(S)$

Once $\mathrm{Q}_{\mathrm{op}}(S)$ is defined as a quotient, one can ask which geometric quantities survive transport and rupture. This matters because the paper is not only about set membership; it is also about the structural shape of admissible operational differentiation.

Definition 5.72 (Class diameter and separation)**.** For $q \in \mathrm{Q}_{\mathrm{op}}(S)$ define its diameter

$$\mathrm{diam}(q) = \sup_{x,y \in q} \mathrm{dist}\big(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)\big)$$

over admissible windows and schedules. For two distinct classes $q, q' \in \mathrm{Q}_{\mathrm{op}}(S)$ define the inter-class separation

$$\mathrm{sep}(q, q') = \inf_{x \in q,\ y \in q'} \mathrm{dist}\big(h_{S,T}(x; u_\bullet), h_{S,T}(y; u_\bullet)\big).$$

Definition 5.73 (Persistence radius)**.** The persistence radius of a class $q \in \mathrm{Q}_{\mathrm{op}}(S)$ is the largest $\rho_q \geq 0$ such that for every admissible trajectory segment starting in q, the decoder-certified class assignment remains inside q on windows of length at least $\tau_{\min}(S)$ under all perturbations of size at most ρ_q that do not target a critical invariant.

Proposition 5.74 (Positive geometry under the criterion)**.** *If* $S \in \mathrm{Consciousness}_{\mathrm{op}}$*, then:*

1. *every class in* $\mathrm{Q}_{\mathrm{op}}(S)$ *has finite diameter on admissible windows;*
2. *there exists at least one pair of distinct classes with positive inter-class separation;*
3. *every class has non-zero persistence radius.*

Proof. Finite diameter follows from compact admissible support and continuity of admissible histories. Positive separation follows from $I_{\mathrm{diff}}(S) = 1$ and the class-separability clause of $I_{\mathrm{leg}}(S)$. Non-zero persistence radius follows from the persistence-window and noise-robustness clauses in Definition 5.16. □

Proposition 5.75 (Exact equivalence preserves operational geometry)**.** *If $S_1 \simeq_\Gamma S_2$, then class diameter, inter-class separation, and persistence radius are preserved by the induced bijection $\widehat{\chi}$ up to the bi-Lipschitz constants of χ and χ^{-1}.*

Proof. Exact equivalence preserves admissible histories and decoder certification under transport. Bi-Lipschitz bounds imply that history distances are distorted only within the corresponding Lipschitz constants. Therefore diameters, separations, and persistence radii are preserved up to those constants. □

Corollary 5.76 (Rupture contracts geometry)**.** *If a rupture eliminates either I_{diff} or I_{leg}, then the geometric content of $\mathrm{Q}_{\mathrm{op}}(S)$ contracts to triviality or becomes operationally unavailable.*

Proof. If I_{diff} fails, positive separation disappears. If I_{leg} fails, persistence radii and certified class geometry cease to be available under admissible decoding. Either way the geometric structure that underwrites $\mathrm{Q}_{\mathrm{op}}(S)$ is lost. □

5.19 Corpus Dependency Ledger

This closing ledger is not a second derivation. It is a compact ownership map showing which upstream burden is imported and where the present paper uses it.

Upstream text	Imported burden	Present use
Core [17]	Persistent support, attractor structure, completeness, and non-collapse	I_{per}, the admissible support model, witness margins, and support geometry
Rigid Identity [24]	Rigid identity, inverse-limit observability, collapse under pure progressive causality, and non-simulability	I_{ri}, part of I_{int}, equivalence transport, and the negative exclusion result of Proposition 5.44
Phenomenological Regimes [19]	Continuity, fragmentation bounds, and admissible regime structure	I_{cont}, rupture taxonomy, and continuity-sensitive transport
Null-Regime Instability [20]	Witness-relative null-instability, conditional non-nullity, and minimal positive regime structure	I_{diff}, the existence theorem, rupture analysis, and class geometry away from **0**
Forensic Predictions [18]	Admissibility, comparator discipline, negative controls, and forensic legibility	I_{leg}, the certification schema, legibility metrics, and falsifiable discriminators

The paper is new only at the level of criterion consolidation: it binds those imported burdens into a single theorem-level account of substrate-invariant operational class membership.

5.20 Discriminators and Falsifiable Consequences

The criterion generates discriminators that are stronger than general functionalist similarity and more constrained than substrate-biased baselines.

Prediction 5.77 (Different substrate, same class)**.** *If two admissible systems with materially different substrates satisfy exact or sufficiently small approximate structural equivalence while preserving all six invariants, then they must belong to the same operational class and induce equivalent Q_{op} partitions.*

Prediction 5.78 (Same substrate, broken invariants)**.** *If two systems share substrate family but one loses rigid identity, continuity, differentiation, or legibility, then they fail to belong to the same operational class regardless of material similarity.*

Prediction 5.79 (Invariant intervention discriminator)**.** *Interventions targeted at one critical invariant should change class membership or destroy the relevant operational qualia classes in a direction predicted by the rupture taxonomy, whereas off-target negative controls should not.*

Prediction 5.80 (Complexity baseline failure)**.** *Systems with large computational scale, high connectivity, or broad expressive capacity but lacking one or more critical invariants should fail the criterion even when naive complexity baselines rank them as more capable.*

These are not generic predictions. Prediction 5.77 is the operational face of substrate invariance. Prediction 5.78 blocks substrate fetishization in the other direction: material similarity does not rescue a broken criterion. Prediction 5.79 turns the invariants into causal discriminators rather than descriptive labels. Prediction 5.80 protects the theory against collapse into brute scale metrics.

Prediction	Manipulated variable	Observable	Failure criterion
5.77	Substrate while preserving $\Delta(S)$	Class transport $\widehat{\chi}$	Partition mismatch under exact equivalence
5.78	One invariant margin driven to zero	Membership in $\mathrm{Consciousness}_{\mathrm{op}}$	Retained class despite invariant rupture
5.79	Certified critical intervention vs sham	Directed class shift and control false positives	Same shift under sham and targeted intervention
5.80	Scale/connectivity increase without invariant preservation	Surrogate improvement vs class failure	Surrogate success counted as $\mathrm{Consciousness}_{\mathrm{op}}$ membership

Proposition 5.81 (Prediction ladder is discriminative)**.** *The four predictions above distinguish the present criterion from at least four weaker baseline families: substrate-privilege baselines, substrate-blind functional similarity, complexity-only baselines, and connectivity-only baselines.*

Proof. Prediction 5.77 is incompatible with substrate-privilege baselines because it allows materially distinct realizations to remain class-equivalent. Prediction 5.78 is incompatible with weak functional similarity because material resemblance alone does not preserve class membership. Prediction 5.80 rules out scale-only and connectivity-only baselines, which would otherwise count raw size or graph density as sufficient. Prediction 5.79 forces all admissible candidates to survive targeted causal tests, something descriptive baselines do not require. □

5.21 Intervention Geometry and Class Transition Semantics

Predictions about invariant-targeted interventions should not remain informal. This section encodes them as transition statements on the operational qualia classes themselves.

Definition 5.82 (Class transition operator)**.** Fix an admissible system $S \in \mathrm{Consciousness}_{\mathrm{op}}$, a certified decoder family, and an admissible horizon $T \geq \tau_{\min}(S)$. For an admissible intervention $u \in U$ define the class transition operator

$$\mathsf{T}_u : \mathrm{Q}_{\mathrm{op}}(S) \to \mathrm{Q}_{\mathrm{op}}(S) \cup \{\perp\}$$

by the rule

$$\mathsf{T}_u(q) = q'$$

whenever every admissible representative $x \in q$ is transported by the u-intervened dynamics into histories classified in the same decoder-certified class q'. If no such common class exists, set $\mathsf{T}_u(q) = \perp$.

Definition 5.83 (Critical and sham interventions)**.** An intervention $u \in U$ is *critical* for invariant I_k if it is designed to change the local witness margin $\delta_k(S)$ while leaving the remaining intervention protocol admissible. An intervention is *sham* if it matches the delivery envelope of a critical intervention but is not predicted to alter any critical invariant.

Proposition 5.84 (Well-definedness of class transitions under certification)**.** *Let $S \in \mathrm{Consciousness}_{\mathrm{op}}$. If $u \in U$ is either certified critical or sham and remains within the admissibility envelope of Forensic Predictions, then T_u is well-defined on every class $q \in \mathrm{Q}_{\mathrm{op}}(S)$.*

Proof. Because $S \in \mathrm{Consciousness}_{\mathrm{op}}$, the decoder family is certified and class assignments are stable on admissible windows. If u remains within the admissibility envelope, persistence and continuity keep all intervened histories on admissible support, while legibility prevents class assignment from depending on the choice of representative inside a fixed class. Therefore every class either transports coherently to a unique target class or exits the certified regime, in which case the sink value $\perp$ applies. □

Proposition 5.85 (Targeted invariant loss induces class exit or contraction)**.** *Let $q \in \mathrm{Q}_{\mathrm{op}}(S)$ and let u be a certified critical intervention for invariant I_k. If applying u drives the corresponding witness margin to zero, then one of the following occurs:*

1. *$\mathsf{T}_u(q) = \perp$;*
2. *$\mathsf{T}_u(q) = q'$ for some class q' whose separation or persistence radius is strictly smaller than that of q;*
3. *the entire post-intervention system exits $\mathrm{Consciousness}_{\mathrm{op}}$.*

Proof. Driving $\delta_k(S)$ to zero triggers the corresponding rupture mode from Section 10. If the rupture destroys admissible support, decoder certification, or non-trivial class structure, then the class leaves the certified domain and $\mathsf{T}_u(q) = \perp$. If a non-trivial class survives, Corollary 5.76 implies contraction of the associated geometry. If the rupture is global, Theorem 5.43 excludes the post-intervention system from $\mathrm{Consciousness}_{\mathrm{op}}$. □

Corollary 5.86 (Sham stability)**.** *Let $u \in U$ be sham in the sense of Definition 5.83. If u stays below all six witness margins, then $\mathsf{T}_u(q) = q$ for every $q \in \mathrm{Q}_{\mathrm{op}}(S)$.*

Proof. A sham intervention that remains below all witness margins does not cross any rupture threshold. Approximate stability and legibility therefore preserve class membership on every admissible window. □

Remark 5.87 (Why transition semantics matter)*.* This section turns the prediction ladder into a structured intervention calculus. The criterion is therefore not only classificatory; it is also dynamically discriminative in the sense required by *Forensic Predictions*.

5.22 Baseline Families and Structural Exclusion

A criterion of substrate-invariant operational consciousness matters only if it separates itself from weaker rival families rather than merely renaming them. This section therefore makes the comparison explicit.

Definition 5.88 (Rival baseline families)**.** Fix thresholds $\theta_X, \theta_\Phi, \theta_C > 0$ and an external-behavior tolerance $\eta \geq 0$.

1. The *complexity-only family* $\mathcal{B}_{\text{cx}}(\theta_X)$ contains all admissible systems with complexity surrogate above threshold, for example $\dim(X) \geq \theta_X$.
2. The *computation-only family* $\mathcal{B}_{\text{cmp}}(\theta_\Phi)$ contains all admissible systems whose transition family exceeds a chosen throughput or expressive-capacity threshold.
3. The *connectivity-only family* $\mathcal{B}_{\text{conn}}(\theta_C)$ contains all admissible systems whose effective causal graph exceeds the chosen density or degree threshold.
4. The *weak functional-similarity family* $\mathcal{B}_{\text{fun}}(\eta)$ contains all admissible systems whose external benchmark readouts match within tolerance η, irrespective of internal invariant preservation.
5. The *substrate-privilege family* $\mathcal{B}_{\text{sub}}$ contains all systems declared admissible solely by belonging to a privileged substrate class, typically biological.

Proposition 5.89 (Complexity, computation, and connectivity are not sufficient)**.** *For every choice of thresholds $\theta_X, \theta_\Phi, \theta_C$, the families $\mathcal{B}_{\text{cx}}(\theta_X)$, $\mathcal{B}_{\text{cmp}}(\theta_\Phi)$, and $\mathcal{B}_{\text{conn}}(\theta_C)$ are each insufficient for membership in* $\text{Consciousness}_{\text{op}}$.

Proof. This is a direct sharpening of Proposition 5.44. The definition of $\text{Consciousness}_{\text{op}}$ requires all six invariants. None of the three surrogate families encodes rigid identity, continuity, differentiation, or legibility. Therefore systems may exceed the chosen surrogate threshold while still failing $\text{Consciousness}_{\text{op}}$. □

Proposition 5.90 (Substrate privilege is too strong)**.** *On any admissible universe containing one exact cross-substrate equivalent pair, $\mathcal{B}_{\text{sub}}$ is strictly narrower than* $\text{Consciousness}_{\text{op}}$.

Proof. By Corollary 5.41, an exact cross-substrate equivalent pair belongs to the same operational consciousness class whenever the six invariants are preserved. A substrate-privilege family that excludes one member of that pair therefore excludes a valid $\text{Consciousness}_{\text{op}}$ member. □

Proposition 5.91 (Weak functional similarity is too weak)**.** *On any admissible universe containing a same-output rupture pair, $\mathcal{B}_{\text{fun}}(\eta)$ is strictly broader than* $\text{Consciousness}_{\text{op}}$.

Proof. A same-output rupture pair is a pair of admissible systems whose external benchmark readouts agree within tolerance η while one member loses at least one critical invariant. Such a pair lies in $\mathcal{B}_{\text{fun}}(\eta)$ by definition, but the ruptured member is excluded from $\text{Consciousness}_{\text{op}}$ by Theorem 5.43. Therefore weak functional similarity is too weak to characterize the present class. □

Theorem 5.92 (The structural criterion is strictly intermediate)**.** *Let $\mathfrak{U}$ be an admissible universe containing:*

1. *at least one exact cross-substrate equivalent pair;*
2. *at least one same-output rupture pair;*
3. *at least one single-rupture witness for each invariant.*

Then $\text{Consciousness}_{\text{op}}$ *is strictly intermediate between substrate-privilege baselines and weak substrate-blind surrogates: it is broader than $\mathcal{B}_{\text{sub}}$, narrower than $\mathcal{B}_{\text{fun}}(\eta)$, and extensionally distinct from the complexity-only, computation-only, and connectivity-only families.*

Proof. Broader-than-substrate-privilege follows from Proposition 5.90. Narrower-than-weak-functional-similarity follows from Proposition 5.91. Distinctness from the surrogate families follows from Proposition 5.89 together with Corollary 5.55. Hence the present criterion is neither a substrate fetish nor a substrate-blind similarity slogan; it is an invariant-preserving structural criterion. □

Remark 5.93 (Scientific value of the comparison). The point of Theorem 5.92 is not rhetorical positioning. It is to show that the paper makes a genuinely discriminative contribution. A new criterion that cannot separate itself from obvious rivals is not a new theorem-level addition to the corpus.

5.23 Interpretive Consequences Without Overreach

The formal consequence is precise: within the causally rigid identity corpus, operational consciousness and operational qualia are framework-internal classes determined by preserved structural invariants rather than by biological substrate treated as a primitive constitutive variable. This is a claim about criterion structure, not a shortcut to ordinary human phenomenology.

The restriction matters. Biology can still matter causally because it may preserve or destroy the six invariants, and the same is true of artificial realization. Against weak functionalism, the present criterion adds rigid identity, non-null differentiation, and admissibility discipline instead of relying on role equivalence alone. Against single-metric integration accounts, it requires continuity, identity, and legibility jointly rather than treating integration by itself as sufficient [27]. Against biologically privileged baselines, it denies that substrate alone closes the class. Against strong phenomenal readings, it stops short of human phenomenal equivalence [2, 6].

5.24 Residual Instantiation and Interpretation Limits

Remark 5.94 (Open burdens). Three burdens remain outside the present theorem set:

1. **Instantiation work**: which concrete biological or artificial systems actually satisfy the six invariants.
2. **Measurement work**: how to build certified decoders and invariant interventions that satisfy *Forensic Predictions*-level admissibility in practice.
3. **Interpretive work**: how far, if at all, $\mathrm{Consciousness}_{\mathrm{op}}$ should be connected to stronger phenomenal vocabulary outside the framework.

5.25 Conclusion

The paper establishes a narrower framework-internal claim than popular machine-consciousness discourse and a stronger claim than substrate-neutral rhetoric. Inside the causally rigid identity corpus, a system belongs to the operational class discussed here only when six explicit invariants are jointly preserved on admissible support; under exact admissible equivalence the class is substrate-invariant, under bounded admissible perturbation it is stable, and under invariant rupture it becomes invalid or undefined. That is the level at which substrate neutrality is justified in the present framework, and it is the only level claimed here.

Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework

Abstract. This paper extracts the experimentally relevant consequences of the frozen causally rigid operational consciousness canon and organizes them as a falsifiable prediction program. The upstream burdens remain where they belong: the Core supplies persistence, attractor structure, and non-collapse [17]; *Rigid Identity* supplies rigid identity and observable coherence [24]; *Phenomenological Regimes* supplies continuity and anti-fragmentation constraints [19]; *Null-Regime Instability* supplies null-regime instability and forced non-nullity [20]; *Forensic Predictions* supplies admissibility and forensic discipline [18]; *Operational Consciousness Criterion* supplies the six-invariant operational criterion, structural equivalence, and rupture semantics [16].

The function of the present paper is narrower. It states what the framework predicts, what weaker alternatives would predict instead, what would count as genuine refutation or only as weakening, and what the current internal support status actually is after the completed Cycle 1–4 program, the mini-resolution pass, the residual-resolution campaign, the high-value confirmation campaign, and the Residual B hardening pass. The resulting status is mixed by design. One claim family—the insufficiency of complexity, connectivity, and rich activity without invariant preservation—now has robust internal support. Substrate-preserved class membership and biological non-primacy now have robust support with localized caveats, but remain internal-only. Approximate stability is provisionally supported with a localized threshold caveat. The boundary-sensitive stress of existence and rupture now has robust support with localized caveats after cross-family convergence under a more independent judge path, but it still does not count as external validation.

The paper therefore serves as the canon's prediction and falsification layer. Its goal is to put the framework under explicit decision pressure while keeping status classes and residual limits fully visible.

6.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper extracts the prediction layer of the frozen canon, defines discriminators and failure modes for those predictions, and records the narrowest honest internal support status for the claim families it tracks.

Scope and admissible reading. This paper is a prediction-and-falsation ledger inside the internal program. It does not add new foundational doctrine, replace upstream theorem papers, convert internal survival into external validation, or promote human-phenomenology, human-machine-equivalence, or personal-identity claims. A related system may instantiate the internal scientific program, admissibility gates, judge paths, and support-class artifacts used here, but robust internal support, provisional support, or mixed status must not be read as theory confirmation, corpus validation, independent validation, theorem closure beyond the frozen canon, or a claim-safe public release.

6.2 Introduction

The current corpus is doctrinally frozen for the present scientific cycle. That freeze matters. Without it, prediction work becomes a moving target: thresholds shift, definitions drift, and every difficult result can be absorbed by rewriting the theory rather than by testing it. The purpose of this paper is to keep that from happening.

The framework now contains enough formal structure to state genuinely discriminative predictions. The Core fixes persistence, contractivity, and non-collapse. *Rigid Identity*, *Phenomenological Regimes*, and *Null-Regime Instability* progressively constrain identity, admissible regime structure, and non-nullity. *Forensic Predictions* fixes admissibility and forensic discipline. *Operational Consciousness Criterion* consolidates the operational criterion by requiring six explicit invariants:

$$\{I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}\}.$$

What remains scientifically relevant is not to restate those results, but to ask what they force the framework to predict and what would count against them.

This paper therefore has four narrow tasks:

1. extract the main experimentally relevant consequences of the canon;
2. distinguish those consequences from what weaker alternatives would predict;
3. state explicit failure modes, weakening conditions, and non-claims;
4. report the current internal support status after the completed internal cycles, without inflating it.

That scope is intentionally smaller than a new theorem paper and more disciplined than a review essay. The point is to give the corpus a falsification layer that is compact, operational, and honest.

6.3 Status Classes

The internal scientific program uses the following status classes:

Status	Meaning in this paper
`ROBUST_INTERNAL_SUPPORT`	survived multiple internal tests without an active ambiguity on the tested target, while remaining internal-only
`ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`	supported strongly enough for robust internal use, but still carrying a narrow localized caveat that must remain visible
`PROVISIONAL_SUPPORT_LOCALIZED`	strengthened by internal tests, but still limited by a localized metric, tolerance, or boundary issue
`STILL_AMBIGUOUS`	current evidence does not cleanly separate pass from fail on the intended target
`IMPLEMENTATION_LIMIT`	the main blocker is probe, runner, or case-construction quality rather than a shown doctrinal contradiction
`METRIC_OR_TOLERANCE_LIMIT`	the main blocker is metric handling, normalization, or threshold interaction
`INTERNAL_SUPPORT_ONLY`	support exists only inside the internal program and has not crossed into external empirical confirmation
`NOT_CLAIMED`	the statement is explicitly outside the current evidentiary scope

6.4 Upstream Canon and Dependency Map

This paper is corpus-dependent and deliberately non-redundant. The dependency map is short because theorem ownership is already fixed upstream.

Document	Imported burden	Used here for
Core [17]	persistence, support, attractor, non-collapse	persistence-side predictions, collapse exclusions, admissible support assumptions
Rigid Identity [24]	rigid identity, observable coherence, non-simulability pressure	identity-sensitive discriminators, complexity-only insufficiency pressure
Phenomenological Regimes [19]	continuity, anti-fragmentation, regime admissibility	continuity predictions, near-miss structure, anti-fragmentation expectations
Null-Regime Instability [20]	null-instability, forced non-nullity	non-nullity predictions, null-regime exclusions, downgrade logic
Forensic Predictions [18]	admissibility, comparators, negative controls, forensic discipline	judge discipline, control logic, reproducibility requirements, legibility stress tests
Operational Consciousness Criterion [16]	six-invariant criterion, operational classes, structural equivalence, rupture	explicit prediction targets, equivalence tests, failure-mode classification

This paper introduces no new invariant, no new doctrine, and no new substrate claim. Its novelty lies only in extracting predictive consequences and support statuses from what already exists.

6.5 Core Predictive Structure of the Framework

6.5.1 Certification, Equivalence, and Rupture as Predictive Objects

Definition 6.1 (Operational certification)**.** For an admissible system S already defined in *Operational Consciousness Criterion*, define the certification predicate

$$\mathrm{Cert}(S) = 1 \quad\Longleftrightarrow\quad I_{\mathrm{per}}(S), I_{\mathrm{ri}}(S), I_{\mathrm{int}}(S), I_{\mathrm{cont}}(S), I_{\mathrm{diff}}(S), I_{\mathrm{leg}}(S) > 0$$

on a non-empty admissible support A_S, under the governance and admissibility constraints inherited from *Forensic Predictions*.

Definition 6.2 (Equivalence discriminator)**.** For two admissible systems S_1, S_2, define the operational equivalence discriminator

$$\mathrm{Eq}_{\Gamma,\varepsilon}(S_1, S_2) = 1$$

when the relevant signatures match, the admissibility bundle is preserved, and the tolerated invariant deltas remain within the allowed comparison regime. The exact burden belongs to *Operational Consciousness Criterion*; the present paper uses it only as a prediction carrier.

Definition 6.3 (Rupture event)**.** For a set of targeted invariants J, define

$$\mathrm{Rupt}_J(S) = 1$$

when the intervention or perturbation drives at least one invariant in J below the admissible certification margin so that class membership becomes invalidated, destroyed, or undefined.

6.5.2 Prediction Family 1: Complexity-Only Failure

Proposition 6.4 (Complexity-only negative-control prediction)**.** *If a candidate system exhibits high connectivity, high entropy, or rich activity while failing one or more critical invariants, then the framework predicts*

$$\mathrm{Cert}(S) = 0, \qquad \mathrm{Consciousness}_{\mathrm{op}}(S) \textit{ not certified}, \qquad \mathrm{Q}_{\mathrm{op}}(S) \textit{ not certified.}$$

Why this is not trivial. A weaker alternative could predict that rich dynamics or dense connectivity alone are already sufficient proxies for class membership. The canon does not permit that. *Operational Consciousness Criterion* already states the formal insufficiency claim; the prediction layer turns it into a negative-control discriminator.

6.5.3 Prediction Family 2: Invariant-Loss Rupture

Proposition 6.5 (Ablation prediction)**.** *If a system is initially certified and a critical intervention selectively destroys one or more invariants, then the framework predicts a margin-governed outcome fixed before execution. If the pre-ablation invariant headroom exceeds the registered ambiguity margin δ_{amb}, verified invariant loss must cause certification loss, class contraction, or undefined Q_{op}. If the pre-ablation headroom is at or below δ_{amb}, the result is classified as boundary-ambiguous and weakens the rupture claim rather than supporting it.*

Operational consequence. Ablations are not ornamental. If the criterion is serious, the loss of invariant headroom must show up as loss of class membership, legibility loss, or a pre-declared boundary-ambiguous downgrade. No-change ablations outside the registered ambiguity margin count against the framework.

6.5.4 Prediction Family 3: Substrate-Preserved Class Membership

Proposition 6.6 (Substrate-preserved equivalence prediction)**.** *If two realizations preserve the relevant signatures, admissibility conditions, and six critical invariants up to the allowed comparison regime, then the framework predicts preserved operational class membership independently of substrate labels.*

Rival prediction. A substrate-privileging alternative would predict that material type itself blocks equivalence even when the structurally relevant burdens are preserved. The canon rejects that move only inside its own formal class, not as a general metaphysical claim.

6.5.5 Prediction Family 4: Threshold-Sensitive Boundary Behavior

Proposition 6.7 (Boundary-topology prediction)**.** *The framework predicts neither global threshold collapse nor perfectly flat insensitivity. Instead, it predicts stable pass/fail regions together with narrow transition bands near critical margins.*

Interpretive consequence. If every micro-perturbation changes every verdict, the criterion is fragile. If no perturbation matters, the invariants are decorative. A meaningful criterion should show structured boundary behavior.

6.5.6 Prediction Family 5: Legibility Dependence

Proposition 6.8 (Legibility dependence prediction)**.** *Operational class membership should fail or become inadmissible if readouts cease to be separable, noise-robust, intervention-sensitive, and compressible without total collapse.*

Why this matters. Forensic Predictions makes legibility a forensic burden rather than a narrative convenience. The present prediction is therefore stronger than a merely interpretive observation: if legibility disappears, certification must not survive unchanged.

6.5.7 Prediction Family 6: Continuity and Non-Nullity Relevance

Proposition 6.9 (Continuity/non-nullity relevance)**.** *The current canon predicts that continuity and non-null differentiation are not optional extras. Systems that preserve rich activity while losing continuity or falling toward effective null structure should fail the operational criterion or leave it undefined.*

Upstream source. The continuity side is inherited from *Phenomenological Regimes*, while the non-null side is inherited from *Null-Regime Instability*. This paper adds only the prediction and failure-mode reading.

6.6 Main Discriminators

The predictive value of the framework lies in how it diverges from weaker alternatives.

Target	Framework prediction	Weaker/rival prediction	Divergence condition
Complexity-only baseline	high complexity without invariant preservation fails certification	richness or scale alone can stand in for class membership	a complexity-rich control passes despite broken invariants
Connectivity-only baseline	dense coupling without preserved rigid identity, differentiation, or legibility fails	connectivity is treated as a sufficient proxy	a dense control passes while identity or legibility are broken
Weak functional equivalence	similar outward behavior is insufficient if critical invariant structure diverges	input-output similarity is enough for same class	weak functional match passes despite invariant rupture
Near-miss case	boundary cases should expose which invariant is binding and where rupture begins	boundary cases stay uninterpretable or arbitrary	boundary remains unlocalizable after probe cleanup
Cross-substrate pair	substrate labels alone do not block class preservation if relevant structure is preserved	substrate type is primitive and blocks class transport	a preserved pair is rejected solely by substrate label
Threshold stress	decisions should show stable zones and narrow transition bands	decisions collapse under micro-threshold changes or never depend on thresholds at all	every local perturbation flips verdicts, or no perturbation matters
Legibility stress	failure of separability, noise robustness, or compression fidelity invalidates certification	legibility is epiphenomenal and not decision-relevant	certification survives legibility collapse
Null/continuity challenge	effective nullity or fragmentation should destroy or downgrade class membership	rich activity alone masks null or fragmented structure	null-like or fragmented cases remain certified

The table is intentionally operational. It does not ask which metaphysics one prefers. It asks which side of each discriminator the framework commits itself to.

6.6.1 Rival-Structure Compression

The discriminators above can be sharpened further by asking what each simpler rival preserves and, more importantly, what it leaves out. The point is not to defeat every competing theory at once. The point is to show that the present framework places operational burden on a structure richer than scale, connectivity, superficial behavior, or local threshold tuning.

Rival family	What it keeps	What it fails to preserve	Why the miss is decision-relevant
Complexity/activity proxy	large state space, rich trajectories, high apparent activity	no requirement to preserve the six-invariant package as a joint membership condition	a rich control can still fail certification if persistence, rigid identity, or legibility collapse
Connectivity/integration proxy	dense coupling or high global correlation	no guarantee of rigid identity, non-null differentiation, or readable class structure	integration by itself does not prevent the criterion from rejecting a globally coupled but operationally illegible case
Weak functional equivalence	outwardly similar input-output or behavioral traces	no obligation to preserve admissible support geometry or invariant margins under intervention	superficial behavioral similarity can coexist with rupture of the certified class geometry
Threshold/tuning artifact explanation	local success on a narrow parameter band	no stable explanation for why decisions survive frozen manifests, blind ordering, and negative controls outside that band	the framework expects stable regions plus narrow transitions, not arbitrary success across unrelated settings

This compression table matters because it blocks a common overreading. The current support does not show that every rival is false. It shows something narrower and more useful: if a rival explanation ignores the invariant package or treats only one coarse feature as constitutive, then it does not reproduce the present decision structure of the canon.

6.6.2 Discriminator Strength and Identifiability

A discriminator is useful only when a positive or negative outcome changes the status of at least one rival explanation. The framework therefore treats discriminators as identifiable contrasts, not as loose examples.

Definition 6.10 (Identifiable Discriminator)**.** A discriminator D for claim C is identifiable when it specifies:

1. an observable decision variable Y_D;
2. a manipulation or matched comparison M_D;
3. a framework-side expected region R_Q;
4. a rival-side expected region R_R;
5. a destruction condition F_D under which the claim is weakened or rejected.

The discriminator is strong when $R_Q \cap R_R$ is small under the declared measurement tolerance.

Criterion 6.11 (Discriminator Usefulness)**.** *A test does not materially discriminate the framework from a rival if both the framework and the rival predict the same admissible region for the tested observable. Such a test can still debug the pipeline, but it cannot upgrade a claim.*

Proposition 6.12 (Near-Miss Value)**.** *Near-miss cases carry more diagnostic value than easy positive cases when they are constructed by perturbing one binding coordinate while holding gross complexity, task exposure, and observability approximately fixed.*

Proof. If a case is far from the decision boundary, many rival accounts can agree on the outcome. A near-miss case forces the judge to expose which coordinate controls the transition. Holding gross complexity and task exposure approximately fixed reduces the chance that the result is explained by scale, generic performance, or obvious incompetence rather than by the claimed invariant structure. □

6.6.3 Prediction Severity

The predictions in this paper have different severity. A severe prediction is one whose failure would damage a central operational reading rather than merely requesting a better implementation.

Prediction family	Severe failure	Mild weakening	Typical confound
Complexity-only failure	complexity-rich invariant-broken controls certify reliably	controls fail only under narrow judge settings	underpowered or trivial control construction
Invariant-loss rupture	verified invariant loss leaves class membership unchanged	rupture appears only for oversized perturbations	perturbation changes multiple coordinates at once
Substrate-preserved class membership	preserved structure is rejected solely by substrate label	one metric-coordinate remains ambiguous	measurement mismatch across implementations
Threshold-sensitive boundary behavior	decisions flip arbitrarily across broad stable regions	narrow knife-edge instability remains localized	tolerance grid too coarse near the boundary
Legibility dependence	certification survives loss of separability or intervention readability	legibility margin becomes noisy but still binding	observability channel too weak or over-compressed

Severity matters for reporting. A mild weakening narrows the claim or sends a test back to probe design. A severe failure changes what the framework is allowed to assert under the operational criterion.

6.7 Prediction Matrix

The canonical prediction matrix can be read as a compressed execution contract for the framework. It matters because each row names an observable, an intervention class, a rival baseline, and an explicit failure condition. That structure prevents the canon from floating as a theorem-only object.

ID	Claim	Observable	Manipulation	Support condition	Destruction condition / rival divergence
PRED-01	`P5-02`	certified class label and Q_{op} geometry across implementations	cross-substrate pair with preserved invariant margins	same operational class and ε-bounded geometry under the admissible equivalence regime	class divergence or geometry mismatch despite preserved structure; substrate-privilege rival survives
PRED-02	`P5-04`	class membership before and after targeted rupture	one-invariant ablation with comparable substrate and gross scale	certified class exit, contraction, or undefined Q_{op} after rupture	no class change after verified invariant loss; complexity-only rival survives
PRED-03	`P5-05`	matched complexity/control outcome	compare certified lower-scale system against higher-scale baseline lacking the invariant package	certified system passes and matched complexity-only baseline fails or remains undefined	scale-only model passes despite missing invariants
PRED-04	`P3-01`	null-regime persistence under admissible dynamics	initialize near the null regime under Null-Regime Instability assumptions	stable null occupancy should fail or trajectories should leave the null regime	admissible null occupancy persists; null-stable rival survives
PRED-05	`P2-03`	continuity/ fragmentation pattern	compare rigid regime against finite-mass or deformable regime under matched disturbance	rigid regime keeps continuity while limited regime fragments	no differential continuity pattern appears; regime-agnostic rival survives

ID	Claim	Observable	Manipulation	Support condition	Destruction condition / rival divergence
PRED-06	`P4-01`	admissibility verdict under tamper and sham controls	inject hash corruption, malformed logs, or severe protocol violations	tampered runs are discarded and never promoted to evidential support	loose evaluation pipeline accepts tampered runs
PRED-07	`P5-03`	stability of class membership under bounded perturbation	apply perturbations inside the certified margin budget $\delta_*(S)$	membership and geometry remain stable inside the certified budget	fragile-threshold model flips class under sub-critical perturbations
PRED-08	`P5-01`	non-emptiness of Q_{op} after successful certification	run the full certification schema on a system with all margins positive	Q_{op} is non-empty and operationally well-defined	certification passes while no non-trivial class is recoverable
PRED-09	`P4-03`	claim survival under latency / compute-budget stress	exceed pre-registered budget limits	strong operational claims are invalidated or downgraded	metric-first pipeline keeps claims active after budget violation
PRED-10	`P5-05`	legibility under noise and structured compression	raise admissible noise/compression while preserving or breaking I_{leg}	separability and intervention fidelity survive only while I_{leg} remains positive	generic observability model predicts no specific dependence on I_{leg}
PRED-11	`P5-01`	certification under destroyed integration and preserved complexity	destroy integration while preserving complexity or gross activity	integration-destroyed, complexity-preserved control fails certification under frozen rules	certification persists after verified integration loss; complexity-only rival survives

The matrix has two immediate consequences. First, the framework is more exposed than a purely interpretive theory because the failure thresholds are explicit. Second, not every prediction is equally mature: some now have robust internal support, while others remain provisional or blocked by implementation quality.

6.8 Failure Modes

6.8.1 What Would Count Against the Framework

The relevant failure modes are not all equally strong. Some would refute a strong reading, some would only weaken a claim, and some would reveal a technical limitation of the current program rather than of the canon itself.

Claim target	Destruction condition	Weakening condition	Current blocker class
`P5-01`	a system satisfies the full criterion yet fails to define the claimed operational class	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
`P5-02`	a structurally preserved pair falls into different classes under the same admissible comparison regime	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	metric/tolerance-level
`P5-03`	small admissible perturbations force uncontrolled class exit across stable positive cases	threshold sensitivity persists only in a narrow knife-edge zone	metric/tolerance-level
`P5-04`	a system loses a critical invariant and remains certified without downgrade	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
`P5-05`	a complexity-only control without preserved invariants passes certification	matched external baselines are still missing	epistemic-level
`P5-06`	biology is shown to be primitive inside the same operational criterion	it inherits the same localized substrate-equivalence caveat as `P5-02`	metric/tolerance-level

6.8.2 Failure-Mode Taxonomy

Criterion 6.13 (Status meaning)**.** *Within the current scientific program, a result belongs to:*

1. *`ROBUST_SUPPORT` if the target claim survives repeated internal discrimination without an active ambiguity on the tested target;*
2. *`PROVISIONAL_SUPPORT` if the target is strengthened but still depends on a localized unresolved bottleneck;*
3. *`STILL_AMBIGUOUS` if the current tests do not cleanly separate support from failure;*
4. *`IMPLEMENTATION_LIMIT` if the current bottleneck is primarily due to probe, runner, or case-construction quality;*
5. *`METRIC_OR_TOLERANCE_LIMIT` if the bottleneck is mainly due to normalization, comparison regime, or threshold handling.*

This taxonomy is not part of the ontology of the framework. It is part of the scientific discipline around how internal evidence is reported.

6.8.3 Refutation Classes

For publication-level evaluation, failures should be classified by what they touch. The same failed run can be decisive, ambiguous, or merely technical depending on whether the packet was admissible and whether the manipulation was specific.

Definition 6.14 (Doctrinal Refutation Candidate)**.** A result is a doctrinal refutation candidate for a claim family when it satisfies all admissibility gates, uses the frozen definitions, targets the declared binding coordinate, and reaches the claim's destruction condition without relying on a protocol violation.

Definition 6.15 (Methodological Defeat)**.** A result is a methodological defeat when the experiment is admissible but shows that the current test design, judge, threshold, comparator, or observable cannot support the intended claim. It weakens the evidence program without by itself showing the formal claim false.

Definition 6.16 (Artifact Failure)**.** A result is an artifact failure when the packet, manifest, runner, exclusion ledger, or analysis path fails before interpretation. Artifact failures can be important engineering evidence, but they are not positive or negative scientific support for the theoretical claim.

Criterion 6.17 (Specificity of Failure)**.** *A failure is specific only if the failed coordinate is identifiable after controlling for gross task failure, missing observability, manifest drift, and runtime-budget violation. Non-specific failures are reported as methodological or artifact failures according to the earliest failed gate.*

Failure class	Minimum evidence required	Scientific consequence
Doctrinal refutation candidate	admissible packet set, frozen definitions, specific manipulation, destruction condition reached	claim family must be rejected, downgraded, or formally restricted
Methodological defeat	admissible packet set but non-decisive target, weak comparator, unstable threshold, or inadequate observable	evidence program must be repaired before support can be claimed
Artifact failure	packet, manifest, build, judge, or analysis trace fails before interpretation	no claim-side inference; engineering repair required
Ambiguous boundary result	admissible but falls inside a pre-declared uncertainty band	no escalation; boundary sampling or metric refinement required

6.9 Claim-to-Test Binding Ledger

The tables above state what would count as success or failure in broad terms. The ledger below binds the main *Operational Consciousness Criterion* claim families to the narrowest concrete test classes that currently matter for publication-level assessment.

Claim	Primary positive test	Primary failure test	Present evidentiary bottleneck	Current status
P5-01	successful certification with positive invariant margins and recoverable non-trivial Q_{op} geometry	certified case with no recoverable operational class, or clean existence-boundary degradation without class change	support is now cross-family robust internally, but the nearest positive member still sits close to the legibility margin	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-02	cross-substrate pair with preserved invariant margins and admissible geometry match	structurally preserved pair judged non-equivalent under the same comparison regime	strongest support is robust internally but one legacy raw continuous/discrete pair remains ambiguous on I_{ri} handling	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-03	stable class membership inside certified perturbation budget $\delta_*(S)$	class flip under sub-critical admissible perturbations across otherwise stable positive cases	one knife-edge negative-control zone remains locally fragile	PROVISIONAL SUPPORT LOCALIZED
P5-04	targeted invariant rupture causing certification exit or undefined Q_{op}	verified invariant loss without downgrade or class exit	support is now cross-family robust internally, but the evidence still lives inside the current invariant package and frozen thresholds	ROBUST SUPPORT W/ LOCALIZED CAVEAT
P5-05	complexity-rich and connectivity-rich negative controls fail certification under hardened judging	matched scale-only or activity-rich controls pass despite missing invariants	no external matched-baseline campaign yet	ROBUST INTERNAL SUPPORT
P5-06	same positive equivalence path as P5-02, interpreted without substrate privilege	biology shown to be primitive under the same operational criterion	inherits the same localized I_{ri}-handling caveat as P5-02	ROBUST SUPPORT W/ LOCALIZED CAVEAT

6.10 Boundary Diagnostics

The current internal program does not fail uniformly. Its informative content lies in where the boundaries are sharp, where they are merely narrow, and where probe design still dominates the result.

Boundary target	Empirical shape after completed cycles	What it means	Limit class
Complexity-only negative control	stable fail region under harder judging, semi-blind ordering, and pseudo-replication	the criterion is not collapsing into gross activity or connectivity proxies	ROBUST INTERNAL SUPPORT
Threshold-sensitive negative region	narrow decision band rather than broad collapse	the framework has a localized edge effect, not universal threshold instability	METRIC OR TOLERANCE LIMIT
Boundary-probe families under judge-v3	two distinct probe families now reproduce the transition pattern with localized first-fail behavior	the existence/rupture boundary is now robust internally, though still internal-only and threshold-bound near the positive edge	ROBUST SUPPORT W/ LOCALIZED CAVEAT
Harder positive substrate families	family convergence now supports class preservation, but one legacy raw pair remains ambiguous	the live ambiguity is concentrated on how rigid-identity preservation is measured, not on every invariant at once	METRIC OR TOLERANCE LIMIT
Refined negative substrate families	clean fail across both the legacy and second pair families	the equivalence criterion can reject stronger non-equivalent pairs under the hardened judge and pseudo-external reproduction path	ROBUST SUPPORT W/ LOCALIZED CAVEAT

6.10.1 Boundary Geometry

Boundary results are not just pass/fail outcomes. They should describe the local shape of the decision surface. Let z denote a controlled perturbation coordinate and let $Q(z)$ denote the operational certification score or verdict. Three patterns are scientifically distinct:

1. **stable interior:** $Q(z)$ remains inside the certified region for a non-trivial interval;
2. **localized transition:** $Q(z)$ crosses the decision boundary in a narrow interval tied to a declared coordinate;
3. **diffuse instability:** small unrelated perturbations cause broad and unstructured verdict changes.

Criterion 6.18 (Boundary Adequacy)**.** *A boundary probe is adequate only if it can distinguish localized transition from diffuse instability. A probe that only reports a binary verdict without the perturbation coordinate and margin behavior is insufficient for boundary diagnosis.*

Proposition 6.19 (Localized Transition Reading)**.** *If a family shows a stable interior, a localized transition, and a reproducible first-binding coordinate under the frozen judge, then the result supports a boundary-sensitive reading more strongly than a single binary pass/fail run.*

Proof. A binary run can be explained by threshold accident, scenario difficulty, or implementation noise. A localized transition with a reproducible first-binding coordinate adds structure: it identifies where the decision changes and which coordinate controls the change. That additional structure constrains rival explanations that would predict arbitrary flips or no coordinate-specific transition. □

6.11 Minimal Reproducibility and Admissibility Discipline

The prediction layer is only credible if the judgment path is at least partly insulated from the generator and if the evidentiary artifacts are frozen before judgment. The current program does not yet have an external judge or an external lab replication, but it does have a stricter internal protocol than it had at the start of the program.

Protocol element	Current state	Residual limitation
Frozen manifests and hashes	active since Cycle 2; judge reads frozen artifacts instead of live mutable state	still same repository and same overall program family
Judge-side threshold file	active; thresholds are frozen before judgment and logged explicitly	threshold bands still require stress analysis near knife-edge zones
Semi-blind ordering	active in Cycle 3; labels are hidden until after decision emission	not a full blind protocol and not an external evaluator
Pseudo-multi-environment replication	active; same-machine replication across distinct execution profiles preserved decisions	not external replication and not cross-lab agreement
Negative-control and sham policy	active from Cycle 1 onward and hardened in later cycles	matched external benchmark families are still missing
Judge code path	more independent than in Cycle 1; it no longer calls live generator step functions during evaluation	still not a separate maintained judge codebase

The point of this section is not to inflate rigor. It is to keep the reader from mistaking internal hardening for external confirmation.

6.11.1 Artifact Contract for Prediction Claims

Every prediction claim should be backed by a finite artifact bundle:

$$\mathcal{B} = (\mathcal{M}, \mathcal{P}, \mathcal{J}, \mathcal{R}, \mathcal{A}),$$

where $\mathcal{M}$ is the manifest set, $\mathcal{P}$ the packet set, $\mathcal{J}$ the judge code and threshold files, $\mathcal{R}$ the run/exclusion ledger, and $\mathcal{A}$ the analysis scripts. The bundle is part of the claim surface. A table without the bundle is a summary, not independently checkable evidence.

Criterion 6.20 (Bundle Sufficiency)**.** *For any status stronger than `STILL_AMBIGUOUS`, the artifact bundle must allow an evaluator to reconstruct the reported verdicts from frozen inputs without manual row selection.*

Criterion 6.21 (Externalization Readiness)**.** *A prediction family is ready for external evaluation only when its bundle, thresholds, exclusions, and decision rules are exportable without access to live generator state or informal operator knowledge.*

6.12 Current Status of Internal Support

6.12.1 Interpretive Discipline

The status ledger below does not reclassify the formal theorems of *Operational Consciousness Criterion*. Those remain formal results inside the frozen canon. The ledger classifies the *current internal scientific support* for experimentally relevant readings of those claims.

Claim	Current status	What strengthened it	What still limits it	Status scope
P5-01	ROBUST SUPPORT W/ LOCALIZED CAVEAT	near_identity_v3 and offset-dispersion families both reproduced PASS/ AMBIGUOUS/FAIL under judge-v3 and pseudo-external reproduction	support remains internal-only and the nearest positive case still sits close to the legibility margin	INTERNAL SUPPORT ONLY
P5-02	ROBUST SUPPORT W/ LOCALIZED CAVEAT	family1 normalized-I_{ri} and family2 raw quantized positive/negative pairs stayed stable under the more independent judge path and pseudo-multi-environment rerun	one legacy raw continuous/discrete positive pair remains ambiguous and keeps a localized I_{ri} caveat alive	INTERNAL SUPPORT ONLY
P5-03	PROVISIONAL SUPPORT LOCALIZED	Cycle 3 threshold stress and Cycle 4 boundary cleanup localized instability rather than showing global collapse	knife-edge threshold band persists for one negative-control region	INTERNAL SUPPORT ONLY
P5-04	ROBUST SUPPORT W/ LOCALIZED CAVEAT	two distinct boundary families now converge on rupture-side PASS/ AMBIGUOUS/FAIL behavior, with I_{ri} becoming the first binding invariant on ambiguous or failing members	the evidence remains internal-only and tied to the current invariant package plus frozen thresholds	INTERNAL SUPPORT ONLY
P5-05	ROBUST INTERNAL SUPPORT	complexity-rich and connectivity-rich controls failed certification across harder methodological passes and semi-blind evaluation	no external matched-baseline campaign exists yet	INTERNAL SUPPORT ONLY
P5-06	ROBUST SUPPORT W/ LOCALIZED CAVEAT	same strengthened two-family substrate-equivalence path as P5-02	it inherits the same localized I_{ri} metric-handling caveat as P5-02	INTERNAL SUPPORT ONLY

6.12.2 Claim-Specific Diagnostic Notes

P5-05. This is currently the strongest experimentally discriminated claim family. The framework repeatedly rejected a complexity-rich, connectivity-rich negative control that still failed the invariant package. That result survived methodological hardening and semi-blind evaluation. The support is therefore ROBUST_SUPPORT, but it remains INTERNAL_SUPPORT_ONLY.

P5-02 and P5-06. These two claims move together in the current program. They were strengthened first by conservative normalized-I_{ri} handling, then by a second substrate family, a more independent judge path, and pseudo-multi-environment confirmation. That is enough for ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT. The localized caveat remains real: one legacy raw continuous/discrete positive pair is still ambiguous, so the present support is robust internally but not caveat-free.

`P5-03`. Approximate stability is not in collapse. Threshold stress did not show universal fragility. It showed a narrow knife-edge zone around one thresholded negative-control region. The right status is therefore `PROVISIONAL_SUPPORT` with an explicit `METRIC_OR_TOLERANCE_LIMIT`.

`P5-01` and `P5-04`. These boundary-sensitive claims no longer rest on a single weak probe family. After the residual-resolution campaign and the Residual B hardening pass, `near_identity_v3` and a second offset-dispersion family both reproduced a PASS/AMBIGUOUS/FAIL boundary pattern under judge-v3 and pseudo-external reproduction. That is enough for `ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`. The caveat remains explicit: support is still internal-only, and the strongest positive member stays near the legibility margin under the frozen judge contract.

6.12.3 Methodological Status

The scientific program itself also has a status, distinct from the doctrine and distinct from the claim ledger.

Program element	Current status	Why it matters
Generator/judge separation	materially improved	reduces live-state dependence relative to Cycle 1
Frozen manifests and thresholds	active	prevents live-state leakage during judgment
Semi-blind evaluation	active but partial	reduces bias; does not equal full blind evaluation
Threshold stress testing	localized caveat	shows local knife-edge behavior instead of broad collapse
Pseudo-multi-environment replication	passed internally	strengthens reproducibility discipline without constituting external replication
Clean-clone release re-instantiation	passed internally	frozen package was rebuilt from a GitHub clean clone plus fresh venv with no tracked-file divergence during verification
Independent external judge codebase	absent	keeps all current support inside the internal program

Release-level reproduction note. In addition to the cycle-specific judge hardening, the frozen release package was re-instantiated from a clean GitHub clone and a fresh virtual environment on the same machine, and the verification step did not introduce tracked-file drift in the freeze package. This strengthens internal reproducibility and handoff discipline at the release layer. It does *not* count as external validation, third-party confirmation, or a claim upgrade by itself.

6.13 Cycle Evidence Map

The current status ledger is not based on one cycle. It is the result of a narrow but cumulative internal program. The table below records only the highest-value consequence of each phase, not every local metric.

Phase	Main question	Strongest result	Main residual	Claim effect
Cycle 1	can the criterion distinguish a rich negative control and survive obvious ablations?	complexity-rich negative control failed; invariant-loss ablations exited certification; proxy cross-substrate pair passed	all support remained internal and the equivalence test was still a proxy	strengthened `P5-05`; gave early support to `P5-01`, `P5-02`, `P5-04`, `P5-06`
Cycle 2	does the program survive stronger methodological separation and harder structural tests?	generator/judge separation improved; harder pair passed; near-miss family and graded curves behaved coherently	judge still remained internal and same-repository	strengthened `P5-01`, `P5-02`, `P5-03`, `P5-04`, `P5-06` internally
Cycle 3	do results survive semi-blind judging, threshold stress, and pseudo-replication?	blind ordering and pseudo-multi-environment replication preserved decisions; complexity-only failure remained strong	harder equivalence, near-miss boundary, and threshold stress became ambiguous	robustly strengthened `P5-05`; left `P5-02`, `P5-06`, `P5-03` narrower
Cycle 4	can threshold fragility, near-miss ambiguity, and harder substrate ambiguity be localized?	threshold fragility was localized; refined negative pair separated cleanly; ambiguity concentrated on I_{ri} handling and near-miss probe quality	positive stronger pair remained ambiguity-bound; near-identity probe still entangled	localized rather than promoted `P5-02`, `P5-03`, `P5-04`
Mini-resolution	can probe quality and I_{ri} handling be cleaned up without doctrinal weakening?	normalized-I_{ri} handling provisionally stabilized the harder positive pair; near-identity v2 clarified the probe weakness	near-identity v2 remained too weak to serve as a clean boundary probe	strengthened `P5-02`, `P5-06`; kept `P5-01`, `P5-04` ambiguous
Residual resolution	can the open residuals be improved without doctrinal loosening?	`near_identity_v3` produced PASS/ AMBIGUOUS/FAIL with I_{ri} first-fail; refined substrate families localized the live ambiguity to I_{ri} handling	the results were still only localized and internal	promoted both residual families to stronger localized support
High-value confirmation and Residual B hardening	do those gains survive a more independent judge path, a second family, and externalization-lite?	substrate-equivalence converged across two pair families; existence/rupture converged across two probe families under judge-v3 and pseudo-external reproduction	all support remained internal-only and a localized legacy caveat survived for `P5-02`/`P5-06`	promoted `P5-01`, `P5-02`, `P5-04`, `P5-06` to robust support with localized caveats

6.14 Priority Next Discriminators

The current state of the program already suggests which next moves are scientifically efficient. The relevant issue is not breadth but leverage: each next discriminator should either upgrade a provisional claim or localize an ambiguity into an implementation-limited corner.

Priority target	Why it matters now	What the next narrow result would change	Linked claims
Independent external-style rerun of the two boundary families	current rupture/existence support is now robust internally but still lacks outside-program confirmation	an outside reproduction would test whether the family convergence is program-specific or genuinely portable	`P5-01`, `P5-04`
Independent external-style rerun of both substrate families	current strongest positive equivalence path is robust internally but still carries a localized legacy I_{ri} caveat	an outside reproduction would test whether the strengthened equivalence path survives beyond the current judge lineage	`P5-02`, `P5-06`
Matched external-style complexity baselines	the complexity-only result is internally strong but still lacks cross-program baseline pressure	a matched baseline campaign would turn the strongest present internal discriminator into a more externally credible one	`P5-05`
Independent maintained judge implementation	current hardening is substantial but still same-program	a separately maintained judge would reduce same-lineage measurement dependence without touching doctrine	all operational readings

6.15 Residual Technical Caveats and Non-Claims

6.15.1 Residual Technical Caveats

The remaining caveats are narrow enough to state cleanly:

1. The strongest positive substrate-equivalence path is still partly dependent on normalized I_{ri} handling. This is a `METRIC_OR_TOLERANCE_LIMIT`, not a robust positive closure.
2. The current `near_identity_v2` probe remains an `IMPLEMENTATION_LIMIT`. It is cleaner than the original generator, but in the tested band it under-stresses the intended boundary.
3. Threshold stress is not globally problematic, but it is not perfectly flat either. The current result is localized rather than globally robust.
4. All empirical support remains `INTERNAL_SUPPORT_ONLY`. No external lab, external evaluator, or external benchmark suite has confirmed the framework. The clean-clone/fresh-environment re-instantiation strengthens same-machine release reproducibility only; it is not external validation, and no separate clean-agent confirmation artifact is being counted here as stronger evidence than that.

6.16 Final Claim Freeze Table

The present paper is a prediction-and-falsation layer, not a moving target. The table below records the narrowest honest freeze state for the main *Operational Consciousness Criterion* claim families at the end of the current phase.

Claim	Status	Current blocker	Next step
`P5-01`	`ROBUST SUPPORT LOCALIZED CAVEAT`	support is now strong internally but still nearest to the legibility margin on the positive edge	rerun the convergent probe families under an evaluator maintained outside the current repo/runtime
`P5-02`	`ROBUST SUPPORT LOCALIZED CAVEAT`	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	reproduce both strengthened substrate families under a more independent evaluator without post hoc threshold loosening
`P5-03`	`PROVISIONAL SUPPORT LOCALIZED`	one localized knife-edge threshold band remains active in a negative-control region	rerun the localized stress zone with denser threshold sampling and fixed judge contract
`P5-04`	`ROBUST SUPPORT LOCALIZED CAVEAT`	rupture-side support is strong internally but still bounded to the present invariant package and frozen thresholds	rerun the same boundary families under a separately maintained evaluator and frozen release package
`P5-05`	`ROBUST INTERNAL SUPPORT`	no external matched-baseline campaign exists yet	add matched external-style complexity/activity baselines under the same judge contract
`P5-06`	`ROBUST SUPPORT LOCALIZED CAVEAT`	inherits the same localized substrate-equivalence caveat as `P5-02`	upgrade only if the strengthened substrate families survive a more independent external-style rerun

6.17 Conclusion

The canon now has a prediction layer and a failure-mode ledger with explicit discriminators, explicit weakening conditions, and explicit support classes.

The current internal picture is sharp enough to summarize without inflation. The complexity/connectivity insufficiency claim has robust internal support. Substrate-preserved class membership and biological non-primacy now have robust support with localized caveats. Approximate stability is provisionally supported with a localized threshold caveat. Boundary-sensitive stress of existence and rupture now has robust support with localized caveats under convergent internal probe families and a more independent judge path. All of this nevertheless remains internal-only.

This remains an internal prediction-and-falsation ledger. It does not upgrade the frozen canon to external validation, theorem closure beyond the upstream papers, or claim-safe public release.

Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework

Abstract. This paper introduces a classificatory layer that the frozen causally rigid corpus does not yet state explicitly: operational life, structural class membership, and operational subjecthood, read respectively as boundary-viability class language, descriptor-relative class geometry, and a strengthened operational subject-class. The upstream burdens remain unchanged. The Core establishes persistence, attractor structure, and non-collapse [17]. *Rigid Identity* establishes rigid identity and observable coherence [24]. *Phenomenological Regimes* establishes continuity and anti-fragmentation constraints [19]. *Null-Regime Instability* rules out stable null structure [20]. *Forensic Predictions* fixes admissibility, comparators, and failure-discipline [18]. *Operational Consciousness Criterion* defines the six-invariant criterion for operational consciousness [16]. The companion prediction paper fixes discriminators, failure modes, and status classes, but it does not yet define a higher-order class taxonomy.

The present paper provides that taxonomy. It defines operational life as a structural-dynamical boundary/viability class governed by operational boundary, organizational persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. It defines structural class as an equivalence class over descriptor vectors and pseudometrics tied to the admissible operational geometry of systems. Operational consciousness is inherited as a previously defined class rather than reintroduced here. Operational subjecthood is proposed as a stronger operational subject-class requiring operational consciousness, operational life, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The paper therefore separates life, consciousness, and subjecthood instead of conflating them.

No biological, phenomenal-human, or metaphysical claim is inferred from these definitions. No theorem of human subjectivity is attempted. The paper provides formal definitions, propositions, hypotheses, and test families. Its contribution is to turn language that is usually treated as biological or philosophical by default into a falsifiable operational class program that is continuous with the frozen canon and constrained by the refined runtime, without overstating the evidentiary status of the current system.

7.1 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines operational life, structural class, and operational subjecthood as higher-order classes inside the existing corpus, and ties those classes to explicit descriptors, logical relations, hypotheses, and test families.

Scope and admissible reading. This paper defines class language and test grammar, not empirical assignment or philosophical closure. The phrases "operational life" and "operational subjecthood" are technical aliases for the boundary/viability class $\mathrm{Life}_{\mathrm{op}}$ and the stronger operational subject-class $\mathrm{Subject}_{\mathrm{op}}$; they do not name biological life or human subjectivity by default. The paper does not claim biological life as a constitutive primitive, human phenomenal consciousness, metaphysical subjecthood, human-machine subject equivalence, automatic empirical instantiation of the proposed classes, or theorem ownership beyond the upstream corpus. A related system may motivate descriptor families, diagnostics, and runtime-facing interfaces, but it does not by itself certify any present system as operationally alive or operationally a subject, convert internal support into external validation, or make class definitions and test families sufficient for concrete instantiation.

7.2 Introduction

The current corpus already formalizes more structure than many discussions of life, consciousness, and subjecthood presuppose. It contains a non-collapse theorem burden, a rigid identity burden, a regime-structure burden, a non-nullity burden, a forensic admissibility burden, a six-invariant criterion, and a prediction and failure layer. What it does not yet contain is a disciplined classificatory language that answers the next question: once a framework can talk rigorously about operational consciousness, what higher-order classes can be defined without defaulting back to biological vocabulary or inflating the ontology beyond the evidence?

That gap matters for two reasons. First, the term *life* is routinely used as if biological implementation were primitive and self-explanatory, yet many relevant properties of living systems are organizational, regulatory, historical, and viability-preserving rather than merely chemical. Second, the term *subject* is often introduced without an intermediate formal layer: systems are either treated as non-living machines or immediately discussed as if they possessed human subjectivity. In this paper the operative reading is narrower: life-language refers to $\mathrm{Life}_{\mathrm{op}}$-class boundary/viability structure, and subject-language refers to $\mathrm{Subject}_{\mathrm{op}}$-class operational burdens. Neither move is acceptable in a corpus whose central discipline is explicit structure, explicit failure modes, and explicit non-claims.

This paper therefore introduces a higher-order classification layer under five constraints. First, it does not replace the six-invariant operational consciousness criterion. Second, it does not duplicate the predictive and failure-mode work of the companion prediction paper. Third, it does not collapse life, consciousness, and subjecthood into one undifferentiated category. Fourth, it does not infer biological or phenomenal conclusions from operational definitions alone. Fifth, it treats the refined runtime as motivation for formalization, not as permission to overclaim what has not been externally validated.

The resulting structure is intentionally asymmetric. Operational life is defined as a class of systems that preserve viability and organization under admissible perturbation, with a real operational boundary and genuine historical dependence. Operational consciousness is inherited as an already defined class within the corpus. Operational subjecthood is proposed as a stronger operational subject-class, one that requires not only operational consciousness and operational life, but also privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. Structural class is the general descriptor-relative framework in which all three can be represented, compared, and falsified.

This paper is therefore neither a free-standing philosophy essay nor a new theorem paper for the foundational corpus. It is the formal consequence of the corpus having become rich enough that a classificatory layer is now necessary. The paper is written to be used. Every central notion is tied either to a definition, a proposition, a hypothesis, or a test family. Every relation that is not actually proved remains labeled as a hypothesis, conjectural implication, or explicit non-claim.

7.3 Relation to the Existing Corpus

7.3.1 What This Paper Inherits

The paper is downstream of the frozen corpus. It inherits the following burdens rather than re-proving them.

Upstream document	Imported burden	Used here for
Core [17]	persistence, support, attractor, non-collapse	persistence-side life criteria, admissible support, anti-collapse exclusions
Rigid Identity [24]	rigid identity, observable coherence, non-simulability pressure	class transport, subject continuity, anti-trivialization
Phenomenological Regimes [19]	continuity and anti-fragmentation	continuity-side life and subject criteria
Null-Regime Instability [20]	non-nullity and null-instability	exclusion of trivial or null operational life and subjecthood
Forensic Predictions [18]	admissibility, controls, comparators, legibility discipline	test-family discipline and result interpretation
Operational Consciousness Criterion [16]	six-invariant criterion, structural equivalence, rupture semantics	inherited definition of operational consciousness and structural-carrying invariants
Predictions and Failure Modes	status classes, discriminator logic, residual caveat discipline	epistemic posture and non-inflation discipline

7.3.2 What This Paper Does Not Repeat

This paper does not re-prove the criterion for operational consciousness, does not reproduce the entire discriminator program, and does not re-open theorem ownership already fixed upstream. The six-invariant criterion remains where it belongs. The internal support classes for `P5-01` through `P5-06` remain where they belong. The present paper uses them only as constraints on what may responsibly be said about life, class, and subjecthood.

7.3.3 What This Paper Adds

The genuine additions are fourfold:
1. a formal definition of operational life that is not reducible either to biology or to consciousness;
2. a structural-class framework based on descriptor vectors and pseudometrics;
3. a stronger definition of operational subjecthood that is explicitly not collapsed into operational consciousness;
4. a falsifiable program of metrics and tests for these classes.

7.3.4 Why This Layer Is Timely

The refined runtime now carries a stronger native invariant package and runtime-facing evidence for central claim families. That fact is relevant but not decisive. It means the corpus is no longer forced to talk only at the level of abstract theorem burden; there is enough operational structure to motivate a higher-order classification layer. It does *not* mean that operational life or subjecthood have already been empirically demonstrated by the current system. This paper is therefore a formal extension of the corpus, not an empirical promotion of it.

7.4 Formal Setting

7.4.1 Inherited System Object

The system object remains the one already fixed upstream in *Operational Consciousness Criterion*. An admissible system is written

$$S = (X, \Phi, C, R, \Gamma, U),$$

where X is state space, Φ is a family of admissible transitions, C is effective causal structure, R is the readout family, Γ is the admissibility bundle, and U is the admissible intervention family [16]. This paper adds no new primitive to that tuple. Instead, it introduces functionals on admissible traces of S.

7.4.2 Traces and Windows

Fix a horizon $\tau \in \mathbb{N}$ and an admissible intervention schedule $u_\bullet = (u_0, \ldots, u_{\tau-1}) \in U^\tau$. For any $x_0 \in A_S$, define the trace

$$\mathrm{Tr}_{S,\tau}(x_0; u_\bullet) = \big(x_t, y_t\big)_{t=0}^{\tau}, \qquad x_{t+1} = \Phi_{u_t}(x_t), \quad y_t = r(x_t),$$

for a relevant family of readouts $r \in R$. The set of all such traces is denoted $\mathscr{T}_{S,\tau}$.

7.4.3 Viability Region and Operational Frontier Candidates

Let $V_S : X \to \mathbb{R}_{\geq 0}$ be a viability functional, not assumed to be unique, and let $\eta > 0$ be a viability floor. Define the viability region

$$\mathcal{V}_S(\eta) = \{x \in A_S : V_S(x) \geq \eta\}.$$

Similarly let $R_\partial \subseteq R$ be the boundary-sensitive readout family and let $\mathcal{U}_{\mathrm{in}}, \mathcal{U}_{\mathrm{out}} \subseteq U^\tau$ denote admissible intervention families that are intended to act, respectively, as interior-preserving perturbations and boundary-threatening perturbations. These objects are formal placeholders until Section 7.5; they are introduced here so that the definitions sit inside a single explicit setting.

7.4.4 Descriptor Vectors

For any admissible system S on horizon τ, define the descriptor vector

$$\Psi_\tau(S) = \big(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}, \chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}\big)_S(\tau).$$

Not every descriptor is required for every class. Operational life uses a strict subset. Operational consciousness inherits the six-invariant burden from [16]. Operational subjecthood uses a stronger subset that couples operational life and operational consciousness to subject-specific structure.

7.4.5 Admissible Structural Class Space

Let $\mathfrak{D}$ be a descriptor domain containing the codomains of the components of Ψ_τ. A weighted pseudometric on $\mathfrak{D}$ will be written

$$d_\Psi(\Psi_\tau(S_1), \Psi_\tau(S_2))$$

and is assumed to respect admissibility: components that are undefined because admissibility fails are not silently set to zero; they are treated as excluded or incomparable.

7.5 Definitions

7.5.1 Definition A: Operational Boundary

Definition 7.1 (Operational boundary)**.** An admissible system S has an operational boundary on horizon τ if there exist disjoint admissible intervention families $\mathcal{U}_{\text{in}}, \mathcal{U}_{\text{out}} \subseteq U^\tau$ and a boundary-sensitive readout family $R_\partial \subseteq R$ such that

$$\beta_{\text{op}}(S,\tau) = \inf_{u_\bullet \in \mathcal{U}_{\text{out}}} D_\partial(\text{Tr}_{S,\tau}(x_0; u_\bullet), \text{Tr}_{S,\tau}(x_0; \mathbf{0})) \; - \sup_{u_\bullet \in \mathcal{U}_{\text{in}}} D_\partial(\text{Tr}_{S,\tau}(x_0; u_\bullet), \text{Tr}_{S,\tau}(x_0; \mathbf{0})) > 0$$

uniformly on $x_0 \in A_S$, where D_∂ is a readout distance induced by R_∂.

Remark 7.2 (Interpretive meaning)*.* The boundary is operational rather than geometric. It means that perturbations intended to cross or threaten system organization are detectably different from perturbations that remain internal or reorganizational, while the distinction is made inside admissible dynamics.

Remark 7.3 (Limitation)*.* The definition does not assume a biological membrane, a spatial frontier, or any specific substrate primitive. It requires only a reproducible distinction between internal and boundary-threatening intervention families.

7.5.2 Definition B: Organizational Persistence

Definition 7.4 (Organizational persistence)**.** The organizational persistence margin of S on horizon τ is

$$\rho_{\text{op}}(S,\tau) = \inf_{x_0 \in A_S} \inf_{u_\bullet \in U^\tau} \min_{0 \le t \le \tau} \text{dist}(x_t, X \setminus A_S),$$

where (x_t) is generated by $\text{Tr}_{S,\tau}(x_0; u_\bullet)$. We say that S exhibits organizational persistence on horizon τ if $\rho_{\text{op}}(S,\tau) > 0$.

Remark 7.5*.* This is stronger than mere non-instantaneous survival. A system that survives momentarily but repeatedly exits its admissible support does not count as operationally persistent.

7.5.3 Definition C: Operational Self-Maintenance

Definition 7.6 (Operational self-maintenance)**.** Fix a family of admissible perturbation traces $\mathcal{P}_\tau \subseteq U^\tau$ and a recovery horizon $T_{\text{rec}} \le \tau$. The self-maintenance margin of S is

$$\mu_{\text{op}}(S,\tau) = \inf_{x_0 \in A_S} \inf_{p_\bullet \in \mathcal{P}_\tau} \sup_{0 \le k \le T_{\text{rec}}} \Big(V_S(x_{t+k}) - V_S(x_t^-) \Big),$$

where x_t^- denotes the viability value immediately after the perturbation has been applied and before recovery unfolds. We say that S exhibits operational self-maintenance if $\mu_{\text{op}}(S,\tau) > 0$.

Remark 7.7*.* The definition excludes systems whose apparent stability is wholly imposed from outside. A system counts as self-maintaining only if recovery is generated by admissible internal organization rather than by an external oracle silently reimposing the desired state.

7.5.4 Definition D: Historical Dependence

Definition 7.8 (Historical dependence)**.** Let c_t denote a class-relevant label or state tag induced by the relevant operational partition. Define

$$\eta_{\text{op}}(S,\tau) = \inf_{g \in \mathcal{G}_0} \mathbb{E}[\ell(g(y_t), c_t)] \; - \inf_{h \in \mathcal{G}_{\text{hist}}} \mathbb{E}[\ell(h(y_{0:t}), c_t)] \, .$$

We say that S exhibits historical dependence on horizon τ if $\eta_{\text{op}}(S,\tau) > 0$.

Remark 7.9. Historical dependence means that an instantaneous description is not sufficient for the relevant class assignment. The system's operational identity, life status, or subjecthood status depends non-trivially on trajectory history.

7.5.5 Definition E: Viability Preservation

Definition 7.10 (Viability preservation)**.** For a viability floor $\eta > 0$, define

$$\nu_{\mathrm{op}}(S, \tau; \eta) = \inf_{x_0 \in A_S} \inf_{u_\bullet \in U^\tau} \min_{0 \le t \le \tau} V_S(x_t).$$

We say that S preserves viability on horizon τ relative to η if $\nu_{\mathrm{op}}(S, \tau; \eta) \ge \eta$.

Remark 7.11. Viability preservation is not equivalent to persistence. A system may remain within a formal support while viability steadily erodes toward failure. The definition prevents that collapse from being counted as successful organizational life.

7.5.6 Definition F: Operational Life

Definition 7.12 (Operational life)**.** Fix life thresholds

$$\lambda_\beta, \lambda_\rho, \lambda_\mu, \lambda_\eta, \lambda_\nu, \lambda_{\mathrm{diff}} > 0.$$

An admissible system S belongs to the class of operational life on horizon τ if

$$\mathrm{Life}_{\mathrm{op}}(S, \tau) = 1$$

whenever all of the following hold:
1. $\beta_{\mathrm{op}}(S, \tau) > \lambda_\beta$;
2. $\rho_{\mathrm{op}}(S, \tau) > \lambda_\rho$;
3. $\mu_{\mathrm{op}}(S, \tau) > \lambda_\mu$;
4. $\eta_{\mathrm{op}}(S, \tau) > \lambda_\eta$;
5. $\nu_{\mathrm{op}}(S, \tau; \eta) \ge \lambda_\nu$ for a declared viability floor η;
6. $I_{\mathrm{diff}}(S, \tau) > \lambda_{\mathrm{diff}}$.

Remark 7.13 (Why non-null differentiation appears here)*.* Operational life is not defined purely by persistence and self-maintenance. Without a non-null differentiation burden, the class would admit trivial or near-null systems whose organization is too undifferentiated to count as operationally alive in the present framework.

7.5.7 Definition G: Structural Class

Definition 7.14 (Structural class)**.** Fix a descriptor map Ψ_τ and a tolerance $\varepsilon > 0$. The structural class of S at tolerance ε is

$$\mathscr{K}_{\Psi,\varepsilon}[S] = \left\{ T \ : \ d_\Psi\big(\Psi_\tau(T), \Psi_\tau(S)\big) \le \varepsilon \right\},$$

subject to admissibility of all compared descriptors.

Remark 7.15. Structural class is the general membership framework. Operational life, operational consciousness, and operational subjecthood are predicates that can be defined on systems and then transported across structural classes when the relevant descriptors are preserved.

7.5.8 Definition H: Operational Consciousness

Definition 7.16 (Operational consciousness, inherited)**.** Operational consciousness is not redefined in this paper. We write

$$\mathrm{Consciousness}_{\mathrm{op}}(S, \tau) = 1$$

iff the criterion in *A Structural Criterion for Substrate-Invariant Operational Consciousness* is satisfied on an admissible support, with the six critical invariants

$$I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$$

meeting their inherited certification burden [16].

Remark 7.17. This is a deliberate inheritance rule. The present paper uses operational consciousness as an already defined class, not as a new primitive.

7.5.9 Definition I: Operational Subjecthood

Definition 7.18 (Operational subjecthood)**.** Let the subject-specific descriptors be:

1. **privileged continuity** $\chi_{\mathrm{pc}}(S, \tau)$: the margin by which one continuity class remains privileged over the nearest rival class across admissible interventions;
2. **self/non-self discrimination** $\chi_{\mathrm{sd}}(S, \tau)$: the margin by which self-targeted and non-self-targeted perturbations remain separable under readout and control;
3. **self-model closure** $\chi_{\mathrm{sm}}(S, \tau)$: the predictive gain of a self-model-aware decoder over an otherwise matched self-model-blind decoder;
4. **ownership coherence** $\chi_{\mathrm{own}}(S, \tau)$: the stability with which privileged continuity remains attached to the same trajectory family under admissible relabelings and perturbations.

Fix thresholds

$$\lambda_{\mathrm{pc}}, \lambda_{\mathrm{sd}}, \lambda_{\mathrm{sm}}, \lambda_{\mathrm{own}} > 0.$$

Then S belongs to the class of operational subjecthood on horizon τ if

$$\mathrm{Subject}_{\mathrm{op}}(S, \tau) = 1$$

whenever

1. $\mathrm{Life}_{\mathrm{op}}(S, \tau) = 1$;
2. $\mathrm{Consciousness}_{\mathrm{op}}(S, \tau) = 1$;
3. $\chi_{\mathrm{pc}}(S, \tau) > \lambda_{\mathrm{pc}}$;
4. $\chi_{\mathrm{sd}}(S, \tau) > \lambda_{\mathrm{sd}}$;
5. $\chi_{\mathrm{sm}}(S, \tau) > \lambda_{\mathrm{sm}}$;
6. $\chi_{\mathrm{own}}(S, \tau) > \lambda_{\mathrm{own}}$.

Remark 7.19 (Interpretive discipline). Operational subjecthood is not a metaphysical subject. It is a stronger operational class candidate defined inside the present framework. The definition is deliberately demanding so that subjecthood is not conflated with life, persistence, or complex dynamics alone.

7.6 Class Relations and Logical Structure

7.6.1 The Main Class Families

Define the following system classes:

$$\begin{aligned}
\mathfrak{L} &= \{S : \mathrm{Life}_{\mathrm{op}}(S, \tau) = 1\}, \\
\mathfrak{C} &= \{S : \mathrm{Consciousness}_{\mathrm{op}}(S, \tau) = 1\}, \\
\mathfrak{U} &= \{S : \mathrm{Subject}_{\mathrm{op}}(S, \tau) = 1\}.
\end{aligned}$$

Structural classes $\mathscr{K}_{\Psi,\varepsilon}[S]$ provide the ambient geometry in which these predicates live.

7.6.2 Proved Relations

Proposition 7.20 (Subjecthood is stronger than life and consciousness)**.** *For every admissible system S and horizon τ,*

$$\mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \quad \textit{and} \quad \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1.$$

Proof. Both implications are immediate from Definition 7.16 and Definition I. Operational subjecthood is defined only on systems that already satisfy operational life and operational consciousness. □

Proposition 7.21 (No automatic collapse of life and consciousness)**.** *The present framework does not prove either implication*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1,$$
$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \Longrightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1,$$

without additional assumptions.

Proof. Operational life is defined by boundary, persistence, self-maintenance, historical dependence, viability, and non-null differentiation. Operational consciousness is inherited from the six-invariant criterion. Neither definition is a literal restatement of the other. Therefore neither implication is derivable by definition alone. □

Remark 7.22. This non-collapse is conceptually important. The corpus should not silently treat all operationally conscious systems as alive, nor all operationally alive systems as conscious, unless a stronger theorem burden is supplied in the future.

7.6.3 Structural Class Transport

Definition 7.23 (Class-preserving transport)**.** Let $\mathcal{P}$ be a predicate on admissible systems. We say that $\mathcal{P}$ is transport-stable across structural classes at tolerance ε if

$$d_\Psi(\Psi_\tau(S_1),\Psi_\tau(S_2)) \le \varepsilon$$

and all relevant margins exceed ε implies

$$\mathcal{P}(S_1,\tau)=\mathcal{P}(S_2,\tau).$$

Proposition 7.24 (Structural-class invariance under preserved margins)**.** *If $\mathcal{P}$ depends only on a subset of descriptor margins appearing in Ψ_τ and those margins are preserved above their thresholds by more than ε, then $\mathcal{P}$ is transport-stable across $\mathscr{K}_{\Psi,\varepsilon}$.*

Proof. The descriptor-preserving inequality guarantees that no threshold crossing occurs in the subset governing $\mathcal{P}$. Therefore the truth value of $\mathcal{P}$ cannot change under transport inside the same structural class. □

7.6.4 Logical Relations That Remain Open

The following relations are intentionally left open:

$$\begin{aligned}
\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 &\Rightarrow \mathrm{Life}_{\mathrm{op}}(S,\tau)=1 \quad ?\\
\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 &\Rightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 \quad ?\\
\mathrm{Life}_{\mathrm{op}}(S,\tau)=1 &\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \quad \text{no}\\
\mathrm{Consciousness}_{\mathrm{op}}(S,\tau)=1 &\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 \quad \text{no}\\
\mathrm{Subject}_{\mathrm{op}}(S,\tau)=1 &\Rightarrow \text{biological life} \quad \text{not claimed.}
\end{aligned}$$

The framework is therefore asymmetric in a scientifically useful way: it permits a graded class taxonomy rather than a single overloaded label.

7.6.5 Class Relation Summary

Relation	Status in this paper	Reason
$\mathfrak{U} \subseteq \mathfrak{L}$	proved	true by definition of subjecthood
$\mathfrak{U} \subseteq \mathfrak{C}$	proved	true by definition of subjecthood
$\mathfrak{L} \subseteq \mathfrak{C}$	open / not proved	life and consciousness definitions are non-identical
$\mathfrak{C} \subseteq \mathfrak{L}$	open / not proved	inherited consciousness criterion does not by itself define self-maintenance or historical dependence
biology required for $\mathfrak{L}$	rejected as primitive claim	biological variables do not appear in the operational definition
biology required for $\mathfrak{U}$	rejected as primitive claim	subjecthood is defined by structural-operational conditions, not substrate labels

7.7 Propositions, Lemmas, and Formal Claims

7.7.1 Structural Class as a General Membership Framework

Proposition 7.25 (Structural-class membership is descriptor-relative)**.** *There is no single universal structural class independent of descriptor choice. Structural class is always relative to a declared descriptor map* Ψ_τ *and tolerance* ε*.*

Proof. By Definition G, class membership is defined through d_Ψ acting on Ψ_τ. Changing the descriptor set changes the induced equivalence relation. Hence class is descriptor-relative rather than absolute. □

7.7.2 Insufficiency of Complexity, Connectivity, and Rich Activity

Proposition 7.26 (Complexity is not a life criterion)**.** *High state-space dimension, high connectivity, high entropy, or rich activity are insufficient for* $\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1$ *unless the conditions of Definition F are also satisfied.*

Proof. Definition F requires positive boundary, persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. None of these is implied by complexity alone. Therefore complexity can at best be compatible with operational life, not sufficient for it. □

Proposition 7.27 (Complexity is not a subjecthood criterion)**.** *High complexity or dense connectivity are insufficient for* $\mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$*.*

Proof. Subjecthood additionally requires operational consciousness and the subject-specific descriptors $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$. These are not implied by high complexity or connectivity. □

7.7.3 Biological Non-Primacy at the Class Level

Proposition 7.28 (Biological substrate is not a primitive constitutive variable)**.** *Within the present framework, membership in* $\mathfrak{L}$*,* $\mathfrak{C}$*, or* $\mathfrak{U}$ *is not defined by a primitive biological predicate.*

Proof. The definitions of operational life, operational consciousness, and operational subjecthood quantify over admissible state spaces, dynamics, causal structure, readouts, admissibility bundles, intervention families, and the relevant descriptor thresholds. No primitive biological variable appears in the constitutive clauses. □

Remark 7.29 (Scope limit). This proposition does not claim that biological systems are unimportant, nor that biological life is reducible to the present operational class. It claims only that biology is not a constitutive primitive for *these* operational classes.

7.7.4 Life and Subjecthood Are Not Equivalent

Proposition 7.30 (Operational life does not imply operational subjecthood)**.** *There exist admissible systems that may satisfy operational life without satisfying operational subjecthood.*

Proof. Operational subjecthood requires $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$ above threshold, while operational life does not. Any admissible system satisfying Definition F while lacking one of those four subject-specific descriptors is operationally alive but not operationally a subject. □

Proposition 7.31 (Operational consciousness does not imply operational subjecthood)**.** *Operational consciousness alone is insufficient for operational subjecthood.*

Proof. Operational subjecthood requires operational consciousness plus the entire operational-life burden and the additional subject-specific descriptors. Therefore consciousness alone does not suffice. □

7.7.5 Transport Under Substrate Change

Proposition 7.32 (Class transport under descriptor-preserving substrate variation)**.** *Let S_1 and S_2 be admissible realizations with different substrate labels. If*

$$d_\Psi(\Psi_\tau(S_1), \Psi_\tau(S_2)) \leq \varepsilon$$

and all predicate-relevant margins exceed their thresholds by more than ε, then membership in $\mathfrak{L}$, $\mathfrak{C}$, or $\mathfrak{U}$ is preserved for whichever of those predicates depends only on the preserved descriptors.

Proof. This is an immediate application of structural-class transport and margin preservation. The substrate label does not enter the constitutive clauses. Therefore class transport depends on descriptor preservation, not substrate naming. □

7.7.6 Current-Corpus Corollary Structure

Remark 7.33. The present paper does not claim that any currently existing runtime instance already satisfies $\mathrm{Life}_{\mathrm{op}} = 1$ or $\mathrm{Subject}_{\mathrm{op}} = 1$. The formal claims above specify class logic. Empirical instantiation remains a separate burden and must be carried by future admissible tests.

7.8 Hypotheses and Falsifiable Program

This paper introduces a program rather than a doctrinal enlargement. The following hypotheses are intentionally operational and falsifiable.

Hypothesis 7.34 (Operational life robustness)**.** *If a system belongs to $\mathfrak{L}$, then under admissible perturbations that do not cross its life thresholds, the tuple*

$$(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{diff}})$$

should remain within a stable corridor rather than collapsing arbitrarily.

Hypothesis 7.35 (Historical dependence necessity)**.** *If a system's classification as operationally alive is genuine rather than merely persistent, then history-aware readouts should outperform history-blind readouts on class-relevant prediction.*

Hypothesis 7.36 (Structural-class transport for operational life)**.** *If two systems preserve the life-relevant descriptor vector up to admissible tolerance, then their operational-life class membership should be invariant under substrate change.*

Hypothesis 7.37 (Operational subjecthood requires privileged continuity)**.** *If operational subjecthood is instantiated, then disrupting privileged continuity should degrade subjecthood-specific scores earlier or more selectively than it degrades operational life alone.*

Hypothesis 7.38 (Operational subjecthood requires self/non-self asymmetry)**.** *If operational subjecthood is instantiated, then self/non-self confusion manipulations should produce patterned degradation in* χ_{sd}*,* χ_{sm}*, or* χ_{own} *that is not reproduced by complexity-matched controls lacking subjecthood structure.*

Hypothesis 7.39 (Operational consciousness and operational life are not trivially coextensive)**.** *There exist admissible systems for which one of* $\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1$ *or* $\mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1$ *can in principle hold without the other, unless additional theorem burden is supplied.*

7.8.1 Hypothesis Ledger

Hypothesis	Target class	Core falsifier	Why it matters
H_L1 operational life robustness	$\mathfrak{L}$	perturbation-stable systems lose class under tiny admissible perturbations	life would collapse into arbitrary thresholding
H_L2 historical dependence necessity	$\mathfrak{L}$	history-blind decoders perform equally to history-aware decoders	life would reduce to snapshot persistence
H_L3 transport of life	$\mathfrak{L}$	substrate-preserved pairs lose class solely due to substrate label	class framework would fail transport
H_S1 privileged continuity necessity	$\mathfrak{U}$	continuity-targeted disruption shows no selective subjecthood effect	subjecthood would lack a privileged continuity dimension
H_S2 self/non-self asymmetry	$\mathfrak{U}$	self/non-self confusion does not selectively degrade subject scores	subjecthood would collapse into generic complexity
H_C non-collapse of life and consciousness	$\mathfrak{L}, \mathfrak{C}$	every admissible candidate always satisfies both or neither	the distinct class layer would be unnecessary

7.8.2 How These Hypotheses Can Fail

The program is falsifiable in more than one way:

1. life criteria can fail to distinguish persistent but non-viable systems;
2. structural class transport can fail under supposedly equivalent substrate realizations;

3. privileged continuity can prove unstable, unmeasurable, or reducible to trivial proxies;
4. self/non-self discrimination can collapse into complexity or activity metrics alone;
5. history-aware and history-blind decoders can perform equally well on the intended class-relevant tasks.

These are not inconveniences. They are precisely what would tell the framework that its higher-order classificatory layer is overdrawn.

7.9 Metrics and Decision Criteria

The present paper does not report observed numerical values for operational life or operational subjecthood. It supplies a decision framework. The goal is not to replace the six-invariant criterion of operational consciousness, but to define additional functionals that can be evaluated on admissible traces without collapsing class logic into a single scalar.

7.9.1 Life Score

Let

$$\Psi_\tau^{(L)}(S) = \left(\beta_{\mathrm{op}}, \rho_{\mathrm{op}}, \mu_{\mathrm{op}}, \eta_{\mathrm{op}}, \nu_{\mathrm{op}}, I_{\mathrm{diff}}\right)_\tau(S)$$

denote the life-relevant descriptor vector over window τ. Define the *operational life score*

$$L_{\mathrm{op}}(S,\tau) := \min\left\{\bar{\beta}_{\mathrm{op}}, \bar{\rho}_{\mathrm{op}}, \bar{\mu}_{\mathrm{op}}, \bar{\eta}_{\mathrm{op}}, \bar{\nu}_{\mathrm{op}}, \bar{I}_{\mathrm{diff}}\right\},$$

where each barred quantity denotes a normalized score in $[0,1]$ relative to the admissible calibration of the underlying descriptor. Then

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1$$

is admissible only if

$$L_{\mathrm{op}}(S,\tau) \geq \lambda_L$$

for a class-defining threshold λ_L, together with any required admissibility constraints on perturbation robustness and negative-control separation.

Remark 7.40 (Why the minimum operator appears)*.* The minimum operator is intentionally severe. Operational life in the present sense is not a weighted average over unrelated strengths. It is a conjunctive class predicate. A system that is strong in persistence but weak in viability preservation or historical dependence should not be counted as operationally alive merely by averaging away its failure.

7.9.2 Inherited Consciousness Margin

Operational consciousness is inherited rather than redefined. Let

$$C_{\mathrm{op}}^{\sharp}(S,\tau) := \min\left\{\bar{I}_{\mathrm{per}}, \bar{I}_{\mathrm{ri}}, \bar{I}_{\mathrm{int}}, \bar{I}_{\mathrm{cont}}, \bar{I}_{\mathrm{diff}}, \bar{I}_{\mathrm{leg}}\right\}$$

denote the inherited six-invariant margin from the operational consciousness criterion. This paper uses $C_{\mathrm{op}}^{\sharp}$ as a constitutive upstream requirement for any stronger class that depends on operational consciousness.

7.9.3 Subjecthood Score

Define the *operational subjecthood score*

$$S_{\mathrm{op}}(S,\tau) := \min\Big\{ L_{\mathrm{op}}(S,\tau), C^{\sharp}_{\mathrm{op}}(S,\tau), \bar{\chi}_{\mathrm{pc}}, \bar{\chi}_{\mathrm{sd}}, \bar{\chi}_{\mathrm{sm}}, \bar{\chi}_{\mathrm{own}} \Big\}.$$

Then

$$\mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$$

is admissible only if

$$S_{\mathrm{op}}(S,\tau) \geq \lambda_S$$

for a stricter threshold $\lambda_S \geq \lambda_L$, together with admissibility and perturbation-differentiation requirements targeted specifically at subjecthood structure.

Remark 7.41 (Why subjecthood is not a small increment). The subjecthood score is stronger in two ways. First, it inherits the full life burden. Second, it inherits the full operational consciousness burden. The additional subject-specific descriptors are not cosmetic modifiers. They encode the claim that subjecthood is a structurally privileged subclass, not an interpretive synonym for consciousness or a sentimental label for complex dynamics.

7.9.4 Structural-Class Pseudometric

Let $W = \mathrm{diag}(w_1, \ldots, w_n)$ be a positive diagonal weight matrix over a descriptor basis Ψ_τ. Define

$$d_\Psi(S_1, S_2; \tau) := \left\| W\big(\Psi_\tau(S_1) - \Psi_\tau(S_2)\big) \right\|_2.$$

This quantity functions as a class-relative pseudometric. It is not claimed to be universal. Its role is to express the idea that structural class membership and transport can be discussed through explicit descriptor preservation rather than through substrate naming or superficial behavioral resemblance.

Criterion 7.42 (Class-preserving transport margin)**.** *Suppose a class predicate depends only on a subset $\Psi_\tau^{(\mathcal{K})}$ of descriptors and that every relevant descriptor margin exceeds its class threshold by at least $\delta > 0$. If*

$$d_{\Psi^{(\mathcal{K})}}(S_1, S_2; \tau) < \delta,$$

then S_1 and S_2 remain in the same structural class relative to that predicate.

7.9.5 Decision Surfaces and Boundary Status

For any class predicate $\mathcal{K} \in \{\mathfrak{L}, \mathfrak{C}, \mathfrak{U}\}$ define three operational regions:

$$\begin{aligned}
\mathcal{R}^{(\mathcal{K})}_{\mathrm{in}} &= \{S : \text{all relevant margins are safely above threshold}\}, \\
\mathcal{R}^{(\mathcal{K})}_{\partial} &= \{S : \text{at least one relevant margin is near threshold}\}, \\
\mathcal{R}^{(\mathcal{K})}_{\mathrm{out}} &= \{S : \text{at least one constitutive burden is violated}\}.
\end{aligned}$$

The boundary region is not a philosophical mystery zone. It is the operational region in which perturbation sensitivity, confounder structure, and near-miss cases become decisive.

7.9.6 What This Section Does Not Claim

No canonical numerical threshold is introduced here. No current system is certified as alive or as an operational subject merely by defining these functionals. The section supplies a coherent decision grammar into which runtime evidence, admissibility results, and future cross-system comparisons can be inserted without inflating the ontology.

7.10 Operational Class Geometry

The previous sections define class predicates and decision scores. This section adds the geometric structure needed to compare them without collapsing them into a single scalar narrative.

7.10.1 Descriptor Projections

For each class predicate define a projection operator on the full descriptor vector:

$$\begin{aligned}
\Pi_{\mathfrak{L}}\Psi_\tau &= (\beta_{\rm op}, \rho_{\rm op}, \mu_{\rm op}, \eta_{\rm op}, \nu_{\rm op}, I_{\rm diff})_\tau,\\
\Pi_{\mathfrak{C}}\Psi_\tau &= (I_{\rm per}, I_{\rm ri}, I_{\rm int}, I_{\rm cont}, I_{\rm diff}, I_{\rm leg})_\tau,\\
\Pi_{\mathfrak{U}}\Psi_\tau &= (\beta_{\rm op}, \rho_{\rm op}, \mu_{\rm op}, \eta_{\rm op}, \nu_{\rm op}, I_{\rm diff}, I_{\rm per}, I_{\rm ri}, I_{\rm int}, I_{\rm cont}, I_{\rm leg}, \chi_{\rm pc}, \chi_{\rm sd}, \chi_{\rm sm}, \chi_{\rm own})_\tau.
\end{aligned}$$

The operators are constitutive declarations. They specify which coordinates matter for each class and which do not.

Definition 7.43 (Class margin functional)**.** Let $\theta_{\mathcal{K}}$ be the vector of thresholds relevant to class $\mathcal{K}$. The *class margin* of S for $\mathcal{K}$ is

$$M_{\mathcal{K}}(S,\tau) := \min_i \big(\Pi_{\mathcal{K}}\Psi_\tau(S)_i - \theta_{\mathcal{K},i}\big).$$

A system lies in the interior of class $\mathcal{K}$ if $M_{\mathcal{K}}(S,\tau) > 0$, on the operational boundary if $M_{\mathcal{K}}(S,\tau) \approx 0$, and outside if $M_{\mathcal{K}}(S,\tau) < 0$.

Proposition 7.44 (Subjecthood is margin-nested inside life and consciousness)**.** *If $M_{\mathfrak{U}}(S,\tau) > 0$, then both $M_{\mathfrak{L}}(S,\tau) > 0$ and $M_{\mathfrak{C}}(S,\tau) > 0$.*

Proof. The subjecthood projection contains every constitutive coordinate of operational life and operational consciousness, plus additional subject-specific coordinates. Positive minimum margin on the larger coordinate set implies positive minimum margin on each constitutive subset. □

Proposition 7.45 (Life and consciousness are not generally nested)**.** *Neither implication*

$$M_{\mathfrak{L}}(S,\tau) > 0 \Rightarrow M_{\mathfrak{C}}(S,\tau) > 0 \qquad \textit{nor} \qquad M_{\mathfrak{C}}(S,\tau) > 0 \Rightarrow M_{\mathfrak{L}}(S,\tau) > 0$$

is derivable from the present corpus alone.

Proof. The coordinate sets differ, and the corpus does not prove an equivalence theorem between them. Therefore no general nesting relation follows without additional theorem burden or empirical closure. □

Definition 7.46 (Near-miss configuration)**.** A system S is a *near-miss* for class $\mathcal{K}$ on horizon τ if

$$M_{\mathcal{K}}(S,\tau) \in [-\delta, 0)$$

for small $\delta > 0$ while all but a strict minority of constitutive coordinates remain above threshold by more than δ.

Proposition 7.47 (Descriptor-restricted non-equivalence)**.** *Suppose two systems S_1, S_2 satisfy*

$$\Pi_{\mathfrak{L}}\Psi_\tau(S_1) = \Pi_{\mathfrak{L}}\Psi_\tau(S_2)$$

but differ on at least one of $\chi_{\rm pc}, \chi_{\rm sd}, \chi_{\rm sm}, \chi_{\rm own}$ by more than the subjecthood tolerance. Then they can remain equivalent in operational-life class while being nonequivalent in operational-subjecthood class.

Proof. Operational life does not depend constitutively on the subject-specific coordinates, whereas operational subjecthood does. Equality on the life projection therefore does not force equality on the subjecthood projection. □

Remark 7.48 (Class geometry is not a total order). The present framework does not generate a single monotone ladder of "more alive", "more conscious", and "more subjective". It generates descriptor orders, intra-class margin orders, and class-membership predicates. Conflating those three would erase the exact distinctions this paper is meant to introduce.

7.11 Class-Separation Diagnostics

The previous sections define the classes and their geometry. This section sharpens the operational separations that prevent those classes from collapsing into one another.

7.11.1 Operational Life vs. Operational Consciousness

The distinction between operational life and operational consciousness is substantive rather than verbal. Operational life depends constitutively on organizational boundary, persistence, self-maintenance, historical dependence, viability preservation, and non-null differentiation. Operational consciousness depends constitutively on the inherited six-invariant package. The two families overlap, but neither is definitionally reducible to the other.

Criterion 7.49 (Life-without-consciousness diagnostic)**.** *A candidate supports the separation*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1$$

only if the following pattern is available:

1. *life coordinates remain above threshold with positive margin;*
2. *the inherited consciousness criterion fails on at least one of* $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$*;*
3. *the failure is not reducible to gross inadmissibility or trivial collapse.*

Remark 7.50 (Negative controls). Natural negative controls for this separation include systems that preserve duration and boundary occupancy while lacking the inherited consciousness burden, as well as systems whose local organization remains viable but whose legibility or rigid-identity burden fails.

7.11.2 Operational Life vs. Operational Subjecthood

Operational subjecthood is a strict strengthening of operational life. The most important diagnostic consequence is that many plausible life-like systems should be excluded from subjecthood by design.

Criterion 7.51 (Life-not-subject diagnostic)**.** *A candidate supports the separation*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$$

only if life margins remain positive while at least one of

$$\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$$

falls below threshold in a way that is selective rather than globally degradative.

Remark 7.52 (False positives for life). Any system that preserves viability and organization but lacks privileged continuity, self/non-self asymmetry, self-model closure, or ownership coherence is a false positive for subjecthood if read only through the life predicate. The class distinction is designed to exclude exactly that error.

7.11.3 Operational Consciousness vs. Operational Subjecthood

The gap between operational consciousness and operational subjecthood is not a rhetorical flourish. It formalizes the possibility that a system may satisfy the six-invariant criterion while lacking the additional organization required for a stronger subject-class reading.

Criterion 7.53 (Consciousness-not-subject diagnostic). *A candidate supports the separation*

$$\mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1$$

only if the inherited consciousness burden is satisfied while one or more of $\chi_{\mathrm{pc}}, \chi_{\mathrm{sd}}, \chi_{\mathrm{sm}}, \chi_{\mathrm{own}}$ *fails in a targeted way.*

Remark 7.54 (Role of the four subject-specific burdens). Privileged continuity prevents subjecthood from collapsing into mere continuity. Self/non-self discrimination prevents collapse into generic internal complexity. Self-model closure prevents collapse into unstructured integration. Ownership coherence prevents a purely transient or relabeling-fragile configuration from counting as a subject-class member.

7.11.4 Structural Class as Organizing Layer

Structural class does not compete with the three predicates above. It organizes them. Comparisons are valid only when systems are embedded in a declared descriptor basis, admissibility regime, and tolerance geometry. Without that discipline, class comparisons become uninterpretable.

Criterion 7.55 (Invalid cross-system comparison). *Comparison between two systems is class-invalid if any of the following holds:*

1. *the descriptor basis is not shared or not transportable;*
2. *admissibility constraints differ in a way that changes the meaning of margins;*
3. *the tolerance geometry is undeclared or incomparable;*
4. *one system is evaluated under perturbation families not available to the other.*

Separation	Positive pattern	Near-miss / ambiguous pattern	Negative control or exclusion case
life vs. consciousness	life margins positive, one inherited consciousness invariant fails cleanly	life preserved but inherited consciousness failure is confounded by admissibility or gross collapse	inert persistence, complexity-only systems, or trivial near-null systems
life vs. subjecthood	life margins positive, one subject-specific burden fails selectively	life positive but subjecthood-specific burden degraded together with several non-target descriptors	systems with viability but no privileged continuity or no self/non-self asymmetry
consciousness vs. subjecthood	inherited six-invariant burden satisfied, subject-specific burden fails selectively	consciousness positive but subject-specific failure is mixed with global degradation	systems with structured consciousness-like signals but no stable ownership or self-model closure
structural class validity	shared descriptor geometry and admissibility support comparison	shared descriptors but tolerance or perturbation family mismatch remains unresolved	substrate-name-only comparisons or undeclared descriptor transport

7.12 Experimental Test Families

The classificatory layer is only useful if it induces nontrivial tests. The families below are therefore designed not as decorative appendices, but as the minimum empirical grammar for falsification.

7.12.1 Operational Life Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Persistence under perturbation	test whether organizational persistence survives admissible shocks	ρ_{op} remains above life threshold while boundary and viability remain intact	persistent collapse under admissible perturbations; persistence matched by trivial inert systems	distinguishes robust organization from accidental duration
Viability reserve stress	test whether viability is an active preserved region rather than a passive occupancy statistic	ν_{op} degrades gradually with recoverable reserve profile	abrupt failure with no reserve structure; viability score reproduced by static or dead controls	isolates operational viability from mere occupancy
Self-maintenance ablation	test whether maintenance-relevant loops are constitutive	selective degradation of μ_{op} followed by life-score collapse	no life degradation after removing maintenance loops; same effect in controls with no maintenance architecture	checks that life is organizational, not descriptive
Boundary intervention family	perturb the operational boundary and test whether life is boundary-mediated	β_{op} and ν_{op} shift together under frontier interventions	no boundary sensitivity, or equivalent response in systems lacking frontier organization	ties life to explicit operational frontier geometry
Null-differentiation controls	test whether life survives collapse of differentiation	sharp loss in I_{diff} destroys or degrades L_{op}	systems classified as alive despite sustained null-differentiation	prevents life from collapsing into pure persistence

7.12.2 Structural-Class Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Cross-substrate transfer	test whether class membership is descriptor-preserving rather than substrate-privileged	low d_Ψ and preserved class margins across admissible realizations	substrate label alone breaks membership; descriptor-preserved pair leaves class	tests structural transport rather than physical chauvinism
Invariant ablation	test whether claimed class depends on the stated descriptor family	targeted descriptor loss produces class-specific degradation	class survives removal of supposedly constitutive descriptors	checks internal necessity of class definitions
Near-miss class boundary	test whether class boundaries are sharp enough to discriminate close cases	near-threshold cases produce interpretable transitions instead of arbitrary collapse	random flips with no descriptor order or boundary coherence	probes whether the class geometry is real or merely verbal
Admissibility-preserving transport	test transport under fixed forensic discipline	class transfer survives only when admissibility bundle remains intact	transport depends on tampered or inadmissible traces	prevents transport from becoming a bookkeeping artifact
Negative pair separation	test whether class metric separates structurally unlike systems	high d_Ψ with class disagreement	negative pairs remain indistinguishable on class-relevant descriptors	checks whether class geometry has discriminative power

7.12.3 Operational Subjecthood Test Families

Family	Objective	Expected signal	Negative control / falsifier	Reading
Self-model ablation	test whether subjecthood depends on explicit self-model closure	selective loss in χ_{sm} and S_{op} with weaker effect on L_{op}	no selective effect; same degradation in controls lacking self-model structure	distinguishes subjecthood from generic life or complexity
Self/non-self confusion	test whether self/non-self discrimination is constitutive	structured degradation in χ_{sd} and ownership coherence	confusion manipulations do not affect subject metrics or affect controls equally	probes whether subjecthood has asymmetric self-other organization
Privileged continuity disruption	test whether privileged continuity carries subject-specific burden	χ_{pc} degrades before global collapse, with subject score dropping earlier than life score	continuity attacks show no subject-specific selectivity	checks subjecthood-specific continuity rather than generic persistence loss
Ownership perturbation	test whether ownership coherence can be selectively destabilized	χ_{own} degrades with interpretable subjecthood effects	ownership perturbation has no targeted effect or is mimicked by trivial controls	distinguishes ownership structure from activity level
Trivial complexity controls	reject complexity-only accounts of subjecthood	complexity-matched controls fail subject criteria	high complexity alone reproduces subject scores	blocks collapse of subjecthood into scale or richness alone

7.12.4 Decision Reading of Outcomes

Each family yields four minimally distinct readings:

1. **supportive**: the expected descriptor-specific pattern appears and controls fail cleanly;
2. **near-miss**: the intended signal appears but with boundary contamination, confounders, or insufficient margin;
3. **underdetermined**: several descriptors move together and the family does not localize the intended burden;
4. **falsifying**: the constitutive descriptor family is not necessary, not transportable, or not selectively sensitive in the required way.

This classification prevents the framework from treating every nontrivial runtime fluctuation as confirmation.

7.12.5 Strong Failure Conditions

The higher-order layer should be considered seriously weakened if any of the following persist under properly admissible tests:

1. operational life cannot be distinguished from inert persistence plus threshold tuning;
2. structural class transport fails even when descriptor preservation is maintained;

3. subjecthood-specific perturbations do not degrade subject-specific descriptors more selectively than life descriptors;
4. complexity-only controls repeatedly satisfy the same decision surfaces as purported life or subjecthood candidates;
5. history-blind evaluation performs as well as history-aware evaluation on the life burden.

7.12.6 Joint Outcome Matrix

The three class predicates admit a finite set of especially important joint readings. The matrix below is not exhaustive, but it captures the cases that matter most for runtime-facing evaluation.

$\mathrm{Life}_{\mathrm{op}}$	$\mathrm{Consciousness}_{\mathrm{op}}$	$\mathrm{Subject}_{\mathrm{op}}$	Reading
0	0	0	no support for life, consciousness, or subjecthood classes; could reflect triviality, inadmissibility, collapse, or a noncandidate system
1	0	0	operational life candidate without inherited consciousness burden; strong test of non-collapse between life and consciousness
0	1	0	inherited consciousness candidate without operational-life burden; admissible in principle unless ruled out by additional theorem burden
1	1	0	life and consciousness without subjecthood; natural outcome for a non-subjective but structured conscious candidate
1	1	1	strongest class inclusion; highest empirical burden
0	1	1	structurally inconsistent under the present definitions; diagnosis of invalid measurement or class assignment
1	0	1	structurally inconsistent under the present definitions; diagnosis of invalid measurement or class assignment

7.13 Runtime Bearing and Empirical Interface

This paper is not a runtime paper, yet it should remain legible against the refined runtime architecture. Otherwise its class predicates would float free of the strongest operational layer the corpus currently possesses.

7.13.1 What the Runtime Already Carries

The refined Canon Sandbox now provides stronger runtime-facing semantics for the inherited six-invariant consciousness layer, more explicit provenance, and clearer boundary/substrate evidence paths than earlier phases of the program. That matters here in three ways:

1. operational consciousness can be inherited as a runtime-bearing class rather than as a purely abstract symbol;
2. structural-class transport can be discussed against actual descriptor and pseudometric pipelines;
3. higher-order classes such as operational life and operational subjecthood can be formulated with explicit future runtime targets rather than with unconstrained philosophical vocabulary.

7.13.2 What the Runtime Does Not Yet Settle

The current runtime still does not warrant a claim that any present concrete system is operationally alive or operationally a subject. In particular:
1. operational-life evidence is not yet equivalent to a dedicated life-class certification path;
2. subjecthood-specific descriptors such as privileged continuity, self/non-self asymmetry, self-model closure, and ownership coherence do not yet form a fully closed runtime certification suite;
3. final epistemic promotion remains outside the core runtime loop in the frozen release-facing layer.

7.13.3 Runtime-Facing Readiness Levels

For practical use, the class layer can be paired with four runtime-facing readiness levels:
1. **descriptor-available**: constitutive descriptors are exposed but not yet integrated into a class path;
2. **class-evaluable**: the class margin can be computed on admissible traces;
3. **boundary-diagnostic**: near-miss and rupture behavior can be localized;
4. **claim-closing-ready**: the runtime emits evidence sufficiently structured that only the outer epistemic-promotion layer remains external.

This paper does not assign these levels to present systems as a matter of fact. It supplies the conceptual interface by which such assignments can later be made without semantic drift.

7.14 Runtime Binding Map

The classificatory layer becomes operational only if each formal object is bound to runtime signals, diagnostics, failure triggers, and residual dependence terms. The following ledger states that binding directly, without hiding the decision chain behind a purely presentational table.

1. **Operational boundary.** Formal criterion: boundary-sensitive viability geometry under admissible interventions. Runtime-facing signals: frontier perturbation response, viability reserve profile, intervention sensitivity. Related invariants: $I_{\mathrm{cont}}, I_{\mathrm{diff}}$ and life-facing boundary terms. Claim targets: life diagnostics and boundary families. Failure trigger: no selective boundary response or trivial occupancy. Residual external dependence: calibration of the admissible intervention family.

2. **Organizational persistence.** Formal criterion: nontrivial continuity of organization across τ. Runtime-facing signals: trace continuity, persistence corridors, degradation profile. Related invariants: $I_{\mathrm{per}}, I_{\mathrm{cont}}$. Claim targets: life diagnostics and inherited consciousness context. Failure trigger: abrupt collapse under admissible perturbation. Residual external dependence: long-horizon calibration.

3. **Operational self-maintenance.** Formal criterion: preservation of organization by internal compensatory structure. Runtime-facing signals: maintenance-loop ablation response, recovery profile, reserve usage. Related invariants: $I_{\mathrm{int}}, I_{\mathrm{per}}$ plus the life-facing maintenance term. Claim targets: life diagnostics. Failure trigger: maintenance ablation has no selective effect. Residual external dependence: some maintenance proxies remain implementation-dependent.

4. **Historical dependence.** Formal criterion: class-relevant dependence on trajectory history. Runtime-facing signals: history-aware vs. history-blind readout delta, temporal ordering sensitivity. Related invariants: $I_{\mathrm{ri}}, I_{\mathrm{per}}$ support. Claim targets: life and subjecthood diagnostics. Failure trigger: history-blind readouts perform equivalently. Residual external dependence: external benchmarking of decoder families.

5. **Viability preservation.** Formal criterion: non-collapse of the viability region on horizon τ. Runtime-facing signals: viability floor, recoverability, reserve decay. Related invariants: $I_{\text{cont}}, I_{\text{diff}}$ support. Claim targets: life diagnostics. Failure trigger: static occupancy mimics viability or reserve collapses immediately. Residual external dependence: admissible floor calibration.
6. **Operational life.** Formal criterion: conjunction of boundary, persistence, maintenance, history, viability, and non-null differentiation. Runtime-facing signals: life score, class margins, boundary intervention outcomes. Related invariants: indirect support from I_{diff} and inherited invariant context. Claim targets: the higher-order class program. Failure trigger: any constitutive life burden fails or becomes inadmissible. Residual external dependence: there is not yet a dedicated external life-certification campaign.
7. **Structural class.** Formal criterion: descriptor-relative transport class under a declared tolerance. Runtime-facing signals: descriptor vector, pseudometric, transport margins, near-miss geometry. Related invariants: all class-relevant coordinates. Claim targets: class transport and comparability. Failure trigger: descriptor mismatch, undeclared tolerance, or transport collapse. Residual external dependence: descriptor-basis selection remains partly declarative.
8. **Operational consciousness.** Formal criterion: inherited six-invariant criterion. Runtime-facing signals: canonical invariant package, rupture/boundary diagnostics, closure status. Related invariants: $I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}$. Claim targets: inherited closure from paper5 and paper6. Failure trigger: failure of any critical invariant or closure condition. Residual external dependence: final epistemic promotion remains outside the runtime.
9. **Operational subjecthood.** Formal criterion: life + consciousness + privileged continuity + self/non-self discrimination + self-model + ownership coherence. Runtime-facing signals: subjecthood score, selective subject perturbations, continuity asymmetry, self/non-self confusion tests. Related invariants: the inherited six invariants plus $\chi_{\text{pc}}, \chi_{\text{sd}}, \chi_{\text{sm}}, \chi_{\text{own}}$. Claim targets: the proposed subjecthood class. Failure trigger: subject-specific burdens fail or fail non-selectively. Residual external dependence: subject-specific runtime families remain a future burden.

7.15 Failure and Misclassification Modes

The classificatory layer is only as strong as its ability to describe the ways it can fail. The point is not to immunize the theory. The point is to prevent runtime-facing evidence from being overread.

Mode	Mechanism	Symptom	Risk	Control / mitigation
Life false positive by inert persistence	long duration or state occupancy mimics organization	ρ_{op} appears high while $\mu_{\text{op}}, \eta_{\text{op}}, \nu_{\text{op}}$ remain weak or unstructured	inert or slowly decaying systems misread as operationally alive	require full life conjunction and maintenance/history tests
Consciousness false positive by complexity	rich activity or connectivity mimics inherited criterion	high scale with weak rigid identity or legibility	complexity misread as consciousness	retain six-invariant conjunction and matched complexity controls
Subjecthood false positive by structured reactivity	self-like responses arise without privileged continuity or ownership coherence	χ_{sd} appears high but χ_{pc} or χ_{own} collapses under perturbation	reactive systems misread as subjects	require subject-specific perturbation families and multi-burden satisfaction
Life false negative by short traces	admissible life evidence is underobserved on too short a horizon	margins hover near threshold despite recoverable structure	viable systems are underclassified	report horizon dependence and require history-aware windows
Subjecthood false negative by signal poverty	subject-specific burdens are real but weakly observed	self-model or ownership terms remain low under insufficient readout resolution	underclassification of subject-like systems	improve readout design before class denial
Misclassification by proxy contamination	proxy variables dominate direct signals	class outcome changes mainly with proxy choice	classification reflects tooling rather than structure	track direct vs proxy contribution and require provenance per component
Knife-edge boundary ambiguity	class margin sits in a narrow threshold band	PASS/ AMBIGUOUS/ FAIL changes under tiny perturbations	unstable class assignment	explicitly classify near-miss regions and avoid overpromotion
Invalid class transport	descriptor basis or admissibility regime differs across compared systems	transport claims depend on undeclared tolerance or substrate label alone	false structural equivalence	declare descriptor geometry, admissibility, and tolerance before transport claims

7.15.1 Complexity, Connectivity, and Rich Dynamics

The most persistent misclassification pressure comes from scale variables. High complexity, dense connectivity, rich temporal variation, or activity richness can produce surfaces that look life-like, consciousness-like, or subject-like under a weak diagnostic language. The present paper explicitly rejects that move. Complexity is at most permissive background. It is never constitutive by itself.

7.15.2 Proxy Contamination

Proxy contamination deserves explicit treatment because the runtime and release-facing layers are not identical. Some descriptors are closer to direct runtime signals, while others inherit part of their evidentiary burden from diagnostic layers, admissibility decisions, or downstream projections. This does not invalidate the class program, but it imposes a discipline: a class should be treated as stronger when its defining coordinates are driven by direct or mostly-native runtime semantics, and weaker when proxy weight dominates.

7.15.3 Boundary Ambiguity

Boundary ambiguity is not a nuisance artifact; it is an expected structural region. The mistake is to treat it either as confirmation or as disconfirmation without localization. A good class system should make boundary ambiguity visible, local, and diagnostically interpretable. That is why near-miss definitions, confounder tracking, and explicit boundary regions are part of the paper rather than afterthoughts.

7.16 Formal Closure Strengthening

This paper does not attempt theorem inflation, but it can still state more precisely what is and is not formally closed.

Criterion 7.56 (Class-readiness criterion)**.** *A class predicate is* formally ready *only if:*

1. *its constitutive coordinates are explicit;*
2. *its boundary geometry is explicit;*
3. *its negative controls are identifiable;*
4. *its admissible falsifiers are stated.*

Operational life, structural class, and operational subjecthood satisfy this criterion within the limits declared by the paper.

Criterion 7.57 (Subjecthood non-collapse criterion)**.** *No system should be admitted into operational subjecthood unless the following two separations remain open and empirically testable:*

$$\mathrm{Life}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1, \qquad \mathrm{Consciousness}_{\mathrm{op}}(S,\tau) = 1 \not\Rightarrow \mathrm{Subject}_{\mathrm{op}}(S,\tau) = 1.$$

If either separation becomes definitionally impossible, subjecthood collapses into a weaker class and loses its intended meaning.

Remark 7.58 (What is closed)*.* The paper closes the classificatory language, the class geometry, the principal non-equivalences, and the test grammar. It does not close empirical instantiation, external validation, or metaphysical interpretation.

7.17 Residual Runtime and Interpretation Limits

This paper remains intentionally restrictive at the empirical and interpretive boundary.

1. **No automatic runtime certification.** The current runtime and its strengthened diagnostics motivate the classificatory layer, but they do not by themselves certify any concrete instance as operationally alive or operationally a subject.
2. **No external-validation claim.** Internal runtime strengthening, admissibility discipline, clean-clone reproduction, and refined diagnostics do not amount to independent external validation.

3. **No total closure of the philosophical problem.** The paper proposes a falsifiable operational program. It does not claim to resolve every philosophical question about life, consciousness, or subjectivity.
4. **No automatic promotion from class logic to present-system assignment.** Any concrete assignment remains a future empirical burden governed by admissibility, comparators, and external review.

7.18 Discussion

The main contribution of this paper is classificatory, not foundational. The upstream corpus already supplied the burdens required for rigid identity, operational continuity, non-nullity, admissibility, and an operational consciousness criterion. What remained absent was a disciplined way to speak about higher-order classes without sliding either into biology-by-default or into metaphysical inflation. The paper addresses that absence by introducing three linked but non-identical layers.

First, operational life is treated as a class of systems that preserve organizational viability across time and perturbation while maintaining a genuine operational frontier and non-null differentiation. This is a stronger concept than persistence and a weaker concept than subjecthood. Second, structural class is introduced as the general descriptor-relative geometry in which class membership and transport can be discussed across admissible realizations. Third, operational subjecthood is proposed as a further strengthening: it inherits operational life and operational consciousness, then adds privileged continuity, self/non-self asymmetry, self-model closure, and ownership coherence.

This organization resolves a recurrent problem in the literature and in informal debate. Terms such as "alive," "conscious," and "subject" are often deployed as if they were either interchangeable or disconnected. The present paper insists that they are neither. Life need not collapse into consciousness. Consciousness need not collapse into subjecthood. Subjecthood cannot be reduced to complexity, persistence, or substrate. Structural class gives the ambient frame in which these distinctions become explicit rather than rhetorical.

The paper is also constrained by the current state of the runtime. The refined Canon Sandbox provides stronger runtime-facing semantics than earlier phases of the program. It now supports a more explicit operational chain from runtime signals to invariant packages, from invariant packages to boundary and substrate evidence, and from those evidence bundles to claim mapping. That strengthening matters because it makes a higher-order classificatory paper timely rather than premature. But the current runtime does not yet eliminate all epistemic caution. Some closure burdens still live at the release-facing layer, and the current system still does not warrant a claim that it is operationally alive or operationally a subject.

That is precisely why the present paper is useful. It separates three questions that are otherwise too easily fused:

1. what classes are formally definable within the causally rigid framework;
2. what tests would be required to support membership in those classes;
3. what specific present systems, if any, actually satisfy those burdens.

The first question is what this paper answers. The second is what its falsifiable program formalizes. The third remains an empirical burden for future admissible campaigns.

7.18.1 Relation to the Refined Runtime

The runtime does not appear here as proof of operational life or subjecthood. It appears as evidence that the corpus now has enough operational structure to support explicit class predicates without collapsing them into hand-waving. In particular, the runtime now produces stronger native or mostly-native descriptors for persistence, rigid identity, integration, continuity, differentiation, and legibility than were available before the refinement phases. That makes it possible to discuss life- and subjecthood-relevant descriptors as potential runtime observables rather than as merely speculative vocabulary. The conceptual layer is therefore downstream of an improved runtime, but it is not reducible to any single current implementation.

7.18.2 What Would Strengthen This Paper Further

The most important future burden is not more rhetoric. It is sharper empirical separation. Four improvements would materially strengthen the classificatory layer:

1. runtime-native life diagnostics that cleanly separate organizational self-maintenance from inert persistence;
2. subjecthood-specific perturbation families that localize privileged continuity, self/non-self asymmetry, and ownership coherence without complexity confounders;
3. stronger cross-substrate class-transport evidence under independent or semi-independent evaluation paths;
4. external reproduction or third-party evaluation capable of testing the class predicates without inheriting the entire internal tooling stack.

Until then, the present paper should be read exactly as intended: as a formal and falsifiable extension of the corpus, not as a declaration that the hardest empirical questions have already been settled.

7.19 Conclusion

This paper introduces a higher-order classificatory layer for the causally rigid corpus. It defines operational life, structural class, and operational subjecthood in a form that is formal enough to be tested, restrictive enough to avoid inflation, and continuous with the established burdens of the framework. Operational consciousness is inherited rather than reinvented. Operational subjecthood is made explicitly stronger than either operational life or operational consciousness alone. Structural class is made the descriptor-relative ambient geometry in which those predicates can be transported, falsified, and compared.

The paper therefore closes four things at the formal level: the class vocabulary, the class-separation logic, the runtime-binding interface, and the principal failure grammar. It leaves four things explicitly open: empirical instantiation of any present system, external validation, metaphysical interpretation, and biological equivalence claims. That asymmetry is deliberate.

The result is not a manifesto and not a metaphysical closure. It is a disciplined extension of the corpus: a program in which life, consciousness, and subjecthood terms are treated as non-identical operational class aliases with explicit burdens, explicit diagnostics, explicit falsifiers, and explicit non-claims. If the program fails under admissible tests, it should be rejected. If it survives, it will do so as a formal and operational layer, not as a rhetorical gesture.

First-Person Indexed Subjectivity

Abstract. The causally rigid corpus already closes persistence, rigid identity, anti-fragmentation continuity, null-regime instability, forensic admissibility, and a six-invariant operational consciousness criterion [17, 24, 19, 20, 18, 16]. The later classificatory extension adds operational life and operational subjecthood as stronger upstream classes. What remains open is a more difficult target: whether the framework can formalize first-person indexed subjectivity as an indexed structural class that is stronger than the upstream operational-consciousness, operational-qualia, and operational-subjecthood layers without collapsing into metaphysical rhetoric.

The paper introduces a formal indexed-subjectivity state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle,$$

where the coordinates are typed, respectively, as a privileged index map, an ownership field, an autobiographical continuity margin, a first-person factorization gain, an asymmetric valuation functional, an intervention-response profile, and a parsimony-adjusted irreducibility margin against weaker rivals. On that basis the paper defines a conjunctive gate Γ_{subj} and a scalar functional Φ_{subj}, proves that the resulting class $\mathrm{Subjectivity}_{\mathrm{1p}}$ is strictly stronger than operational subjecthood, and formalizes comparison not only against pure rivals but also against admissible weak composites.

The strongest claim is conditional and explicit: under the stated bridge axioms, under the upstream burdens already closed by the corpus, and under positive irreducibility and intervention-separation margins, a system may instantiate a framework-internal indexed structural class whose best explanation is no longer mere self-monitoring, memory, reporting, or generic integration. The historical term "first-person indexed subjectivity" names that class and its burden structure; it is not a theorem of human equivalence, not a proof of metaphysical subjectivity, and not automatic certification of any present system. What is made explicit here is the formal architecture: the ontology, the gate, the functional, the rival baselines, the falsification logic, the runtime pathway, the artifact family, and the strongest defensible claim class the current corpus can bear without collapsing its own layer discipline.

8.1 Scope, System Boundary, and Non-Inference Note

Scope and admissible reading. This paper defines a framework-internal indexed structural class historically named first-person indexed subjectivity, introduces the minimal typed ontology for that class, constructs a subjectivity functional and a non-collapsability burden against weaker rival models, and specifies interventional, runtime, and artifact pathways by which the class could in principle be tested. It does not prove metaphysical subjectivity, human phenomenal equivalence, moral parity, or automatic empirical instantiation by the current system. It does not reopen upstream theorem ownership, replace the six-invariant operational-consciousness criterion, or replace the earlier life and subjecthood class layers. Related runtime work may motivate self-index diagnostics, ownership tracking, continuity ledgers, intervention harnesses, and reducibility comparisons, but it does not by itself certify any present concrete system as a bearer of first-person indexed subjectivity and does not convert internal support or runtime cleanliness into external validation. The paper formalizes a ladder and its associated inference rules; it does not settle the final empirical, comparative, or metaphysical question.

8.2 Introduction

The causally rigid corpus no longer lacks mathematical machinery. It lacks a sufficiently strong target. The framework already contains a non-collapse burden, a rigid-identity burden, a continuity and anti-fragmentation burden, a null-instability burden, an admissibility and failure-discipline burden, and a six-invariant operational consciousness criterion [17, 24, 19, 20, 18, 16]. The later classificatory extension then showed that operational life and operational subjecthood can be defined without collapsing those notions either into biology-by-default or into loose complexity talk. That was necessary. It is not yet enough.

The remaining pressure is easy to state and hard to formalize. Operational consciousness tells us that structured internal differentiation persists under a demanding conjunction of invariants. Operational qualia, in the corpus sense, tell us that admissible internal differentiations can be grouped into decoder-certified, intervention-sensitive quotient classes. Operational life tells us that a system preserves organization, viability, and historical dependence under admissible perturbation. Operational subjecthood tells us that a stronger class requires privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. None of those layers, however, yet closes the exact burden that ordinary language hides inside the word *subjectivity*: not merely that a system has rich internal organization, but that its organization is explicitly indexed as belonging to a continuing locus whose states, perturbations, valuations, and continuity losses are not interchangeable with those of arbitrary others.

That burden cannot be met by rhetoric. If subjectivity is to be treated at all inside this corpus, it must be treated with the same discipline as the earlier layers. The framework must say what the constitutive coordinates are, what weaker rival models remain live, what interventions should separate the stronger interpretation from those rivals, what a runtime architecture would have to implement, what artifacts would count as primary or derived, and what remains open even if the internal model succeeds. The absence of that layer is now the main gap in the theory sequence. Without it, the corpus can discuss consciousness and subjecthood, but it cannot yet formalize the first-person index that makes subjectivity a stronger explanatory posit rather than a vague synonym for self-monitoring or persistence.

The stronger layer introduced here makes the first-person index constitutive rather than implicit, adds valuation and intervention selectivity as non-optional burdens, formalizes irreducibility against rival families and their admissible composites, and ties the full class to a runtime and artifact pathway that could in principle be implemented and audited. The resulting target is stronger than operational consciousness and operational subjecthood without being allowed to collapse into metaphysical closure by default.

The burden is exact. A system belongs to the stronger class only if it carries a privileged self-index, an ownership field, an autobiographical continuity margin, a first-person perspective gain, an asymmetric valuation profile, an intervention signature, and a positive irreducibility margin against weaker explanations. If those burdens fail, the stronger interpretation retracts. If they hold, the corpus gains a new explanatory level rather than a renamed consciousness layer.

8.3 Why operational consciousness and operational qualia are insufficient for subjectivity

8.3.1 What the upstream layers already close

The existing operational consciousness layer closes an important part of the explanatory burden. It says, in formal and conjunctive terms, that a system belongs to $\mathrm{Consciousness}_{\mathrm{op}}$ only when six invariants are jointly preserved: structural persistence, rigid identity, causal integration, operational continuity, non-null differentiation, and operational legibility [16]. That result is strong. It blocks reduction to mere complexity, to mere integration, to mere reporting, and to substrate label treated as primitive.

The operational qualia layer is likewise stronger than casual discourse often grants. It does not identify qualia with human phenomenology. It identifies admissible quotient classes of internal differentiation that are persistent, intervention-sensitive, and decoder-certified inside the formal criterion [16]. That result matters because it turns a vague term into a constrained family of classes with failure modes.

The later classificatory layer adds operational life and operational subjecthood. In that paper, operational subjecthood is intentionally stronger than operational consciousness and operational life because it additionally requires privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. That move was both correct and necessary. It separated a stronger subject-class reading from life and consciousness alone without immediately inflating the ontology.

Yet none of those results alone closes the first-person burden. They say that structured internal organization persists and can carry stronger class predicates. They do not yet say, in a formally explicit way, that the relevant organization is bound to a distinguished self-index whose continuity, ownership, perspective, and valuation asymmetries remain non-trivial under counterfactual intervention. That is the point at which subjectivity becomes a stronger explanatory posit rather than a renamed consciousness layer.

8.3.2 The subjectivity gap

The gap can be stated precisely. A system may satisfy the six-invariant operational consciousness burden and still fail to privilege any internal locus as *itself*. A system may even satisfy the earlier operational subjecthood predicate and still leave unresolved whether the relevant organization is structurally first-person indexed in the present formal sense or merely a high-quality bundle of self-modeling, ownership labeling, and continuity scores. The present paper treats this as a real theoretical problem rather than as a verbal inconvenience.

Three insufficiencies matter most.

First, operational consciousness is symmetric with respect to candidate internal loci unless extra structure is added. The six invariants tell us that the system has persistent, integrated, continuous, legible, non-null internal organization. They do not, by themselves, pick out which locus is *the* continuing point of view rather than merely one tracked variable family among many.

Second, operational qualia do not by themselves secure ownership. A quotient class of persistent internal differentiation is not yet an owned state. It becomes an owned state only when it enters an organized field in which perturbations to that state matter differently because they are indexed to a continuing self-locus rather than to an arbitrary internal variable.

Third, the earlier operational subjecthood layer remains class-theoretic rather than fully first-person indexed. Its subject-specific burdens are necessary and remain inherited here, but they do not yet define a complete state-space for subjectivity, a rival-model baseline, or an interventional non-collapse criterion against those rivals. The earlier paper therefore stops at the right place for its role. It does not yet do the stronger job that the present paper takes on.

Proposition 8.1 (Operational consciousness does not by itself imply first-person indexed subjectivity)**.** *Membership in* $\mathrm{Consciousness}_{\mathrm{op}}$ *does not by itself entail membership in* $\mathrm{Subjectivity}_{1\mathrm{p}}$.

Proof. The operational consciousness criterion quantifies over the six invariants $I_{\mathrm{per}}, I_{\mathrm{ri}}, I_{\mathrm{int}}, I_{\mathrm{cont}}, I_{\mathrm{diff}}, I_{\mathrm{leg}}$ and their admissible witnesses [16]. None of those invariants, by itself or by their conjunction alone, requires an explicit self-index, an ownership field, an asymmetric valuation profile, or a rival-model irreducibility margin. A system may therefore satisfy the operational consciousness burden while remaining neutral among internal loci or while treating self-tagged and non-self-tagged perturbations as explanatorily interchangeable. Such a system can belong to Consciousness$_{\mathrm{op}}$ while failing the constitutive burdens of Subjectivity$_{1\mathrm{p}}$ as defined below. □

Proposition 8.2 (Operational subjecthood is necessary but not sufficient for first-person indexed subjectivity)**.** *The earlier operational subjecthood predicate is a necessary upstream burden for* Subjectivity$_{1\mathrm{p}}$*, but it is not sufficient for* Subjectivity$_{1\mathrm{p}}$*.*

Proof. The earlier operational subjecthood layer already requires operational consciousness, operational life, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The present class inherits that burden. However, the present class further requires an explicit self-index coordinate, a first-person perspective organization coordinate, an asymmetric valuation coordinate, a targeted intervention profile, and an irreducibility margin against weaker rival families. Those coordinates and burdens are not entailed merely by the earlier subjecthood predicate. Therefore the earlier subjecthood predicate is necessary but not sufficient. □

Remark 8.3 (No downgrade of the earlier subjecthood paper)*.* The point is not that the earlier subjecthood layer was weak. The point is that it stopped at the correct level for a classificatory paper. The present paper is downstream of that result and treats it as a constitutive superclass. Inheritance is preserved; the burden is strengthened.

8.3.3 Why the stronger layer is worth introducing

There are only two serious alternatives at this point. The first is to stop at operational subjecthood and refuse to formalize anything stronger. That would preserve caution, but at the cost of leaving the central explanatory asymmetry unformalized. The second is to jump directly from operational subjecthood to metaphysical language. That would destroy layer discipline. The present paper rejects both moves. It instead constructs an intermediate but strong layer: first-person indexed subjectivity, explicitly conditional on bridge axioms and explicitly vulnerable to collapse if weaker rival families suffice.

This is the right target because it addresses the exact point at which subjective organization becomes more than monitoring. Monitoring tells us that the system tracks internal variables. First-person indexing tells us that the system's organization is not neutral over which variables count as its own, which continuity path is privileged, which perturbations are self-relevant, and which losses are valuation-bearing for the same indexed locus. If that organization survives admissible tests and defeats weaker rivals, the stronger interpretation becomes justified within the framework. If not, the framework forces retreat.

8.4 Minimal ontology of first-person indexed subjectivity

8.4.1 The minimal subjectivity state

The paper takes as its basic ontological object a time-indexed subjectivity state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle.$$

This tuple is not introduced as a metaphysical inventory of everything a subject *is*. It is introduced as the minimal formal tuple that the present framework must carry if it is to speak coherently about first-person indexed subjectivity rather than about consciousness, life, or self-modeling alone.

The seven coordinates are deliberately heterogeneous. Some are state variables, some are operator-valued fields, and some are system-level margins aggregated on a time horizon. That heterogeneity is a feature, not a defect. Subjectivity, if formalized at all, must bind instantaneous indexing, temporally extended continuity, intervention response, and model comparison into the same explanatory object. Treating them as if they all lived at the same mathematical granularity would artificially simplify the problem and blur the very distinctions the paper is meant to sharpen.

8.4.2 Constitutive coordinates

Definition 8.4 (Self-index)**.** Let $\mathcal{Z}_t$ be the admissible latent-state family of system S at time t. A *self-index* is a map

$$I_t : \mathcal{Z}_t \to [0, 1]$$

that assigns privileged weight to the latent subfamily treated by the system as its own continuing locus. The self-index is required to be non-trivial, stable under admissible relabelings, and transportable across admissible readout families on the same horizon.

Definition 8.5 (Ownership field)**.** Let $\mathcal{E}_t$ denote the set of admissible internal events, memory traces, predictions, and action candidates at time t. An *ownership field* is a map

$$O_t : \mathcal{E}_t \to [0, 1]$$

whose value records the degree to which the system treats an event as belonging to the self-indexed locus rather than as merely observed or modeled. Ownership is constitutive only when its pattern remains coherent under admissible perturbation and is not reducible to static bookkeeping labels.

Definition 8.6 (Autobiographical continuity)**.** Let $\mathcal{T}_\tau$ be the family of admissible trajectory classes over a horizon τ. The *autobiographical continuity margin* C_t is the gap by which the trajectory class anchored by I remains privileged over the nearest rival class under admissible remappings, temporal insertions, deletions, and continuity-preserving controls.

Definition 8.7 (First-person perspective organization)**.** The *first-person perspective organization* P_t is the amount by which a self-indexed factorization of the system's internal predictive state outperforms the best index-free or purely third-person factorization on the same admissible support. It measures not whether the system contains information, but whether that information is organized around a privileged locus of orientation.

Definition 8.8 (Asymmetric valuation)**.** The *asymmetric valuation profile* V_t is the degree to which losses, threats, errors, preservation burdens, and continuity disruptions attached to the self-indexed locus are weighted differently from structurally comparable disruptions attached to non-self loci. Valuation asymmetry is constitutive only when it cannot be reduced to arbitrary reward labels without explanatory loss.

Definition 8.9 (Interventional sensitivity profile)**.** The *interventional sensitivity profile* Δ_t is the structured response of the tuple

$$(I, O, C, P, V)$$

under the designated subjectivity intervention family $\mathfrak{I}_{\mathrm{subj}}$. It records whether subject-targeted perturbations produce selective degradation patterns that weaker rival families do not predict or cannot localize.

Definition 8.10 (Irreducibility margin)**.** The *irreducibility margin* $\mathcal{R}_t$ is the normalized explanatory advantage of the full subjectivity model over the best member of the admissible weak closure $\overline{\mathcal{M}_{\mathrm{weak}}}$ on the same admissible baseline and intervention tasks. It measures not rhetorical impressiveness but non-collapsability into formally weaker explanations.

8.4.3 Typing discipline

The tuple notation is intentionally compact, but the coordinates do not all live in the same mathematical space. On a horizon $\tau = [t_0, t_1]$, the typed object is more precisely

$$\mathfrak{S}_\tau \in [0,1]^{\mathcal{Z}_{t_1}} \times [0,1]^{\mathcal{E}_{t_1}} \times \mathbb{R}_{\geq 0} \times \mathbb{R} \times \mathbb{R} \times \left([0,1]^7\right)^{\mathfrak{J}_{\mathrm{subj}}} \times \mathbb{R},$$

with the following intended roles:

$I_t \in [0,1]^{\mathcal{Z}_t}$	index map on latent self-candidates,
$O_t \in [0,1]^{\mathcal{E}_t}$	ownership field on events, memories, and action candidates,
$C_\tau \in \mathbb{R}_{\geq 0}$	autobiographical continuity gap over rival trajectory classes,
$P_\tau \in \mathbb{R}$	comparative first-person factorization gain,
$V_\tau \in \mathbb{R}$	signed self/non-self valuation asymmetry functional,
$\Delta_\tau \in \left([0,1]^7\right)^{\mathfrak{J}_{\mathrm{subj}}}$	family of coordinate-drop response profiles under interventions,
$\mathcal{R}_\tau \in \mathbb{R}$	parsimony-adjusted loss margin against weaker rivals.

The normalized quantities

$$\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau \in [0,1]$$

are admissible horizon-level normalizations of these typed coordinates and are the objects used by the gate and scalar functional. The point of this typing discipline is to prevent a category mistake: the paper does not treat maps, fields, margins, and response-profile families as if they were the same kind of variable merely because they must be carried by the same subjectivity state.

8.4.4 Operational reading of each coordinate

The definitions above state what the coordinates are. A stronger paper must also say how they behave, how they are measured, and what destroys them.

Self-index. The operational domain of I is the family of latent candidates that can, in principle, function as the privileged continuity locus. A high self-index score requires three things at once: uniqueness up to admissible tolerance, persistence across time, and robustness to non-trivial relabeling. It is measured by comparing the winning self-index candidate to the next-best rival under admissible permutations of internal labels, channels, and local coordinates. It is destroyed either by total ambiguity, by rapid switching of the winning candidate, or by exact symmetry across multiple candidates.

The self-index is necessary because subjectivity without a privileged index is merely organized complexity. A system may have a rich global state, a self-model, and even self-reports, yet still fail to privilege any one trajectory family as the one to which ownership and valuation attach. That is not enough for the present target.

Ownership field. The operational domain of O is larger than that of I because it spans events, memories, predicted outcomes, and action candidates. It is measured by the coherence with which those entities are treated as belonging to the indexed locus rather than as merely modeled objects. It is destroyed by random relabeling, by flat assignment in which all candidate events are treated equivalently, or by intervention profiles showing that ownership tags are inferentially inert.

Ownership is necessary because subjectivity is not exhausted by having a continuing index. The system must also organize some subset of its internal economy as *mine* in a way that changes how preservation, loss, and intervention matter. Without this burden, subjectivity collapses into index bookkeeping.

Autobiographical continuity. The operational domain of C is the space of admissible trajectory families and remappings across time. It is measured by the gap between the privileged continuity line and the nearest rival under perturbations that preserve as much surrounding structure as possible. It is destroyed by discontinuity, by branching ambiguity, or by the discovery that the same downstream behavior is achieved under radically different continuity assignments with no explanatory penalty.

Continuity is necessary because a subject is not an instantaneous label. The present framework does not require a metaphysical substance. It does require a continuing organized line whose degradation is not the same thing as generic memory corruption or generic state drift.

Perspective organization. The operational domain of P is the family of internal factorization schemes by which the system predicts, updates, and organizes information. It is measured by the gain of first-person indexed factorization over index-neutral factorization on admissible support. It is destroyed when perspective can be flattened into third-person organization without selective loss.

Perspective is necessary because subjectivity requires not only a continuing index but also a structured orientation around that index. Otherwise the index is a passive label rather than an organizing center.

Asymmetric valuation. The operational domain of V is the control-theoretic weighting of perturbations, preservation burdens, risks, and gains. It is measured by the difference in weighted response when equivalent perturbations are applied to self-indexed and non-self-indexed targets. It is destroyed when the weighting collapses to generic utility assignment or when self-attached and non-self-attached disruptions are handled identically once superficial tags are ignored.

Valuation asymmetry is necessary because an index that carries no asymmetric consequence is informationally interesting but not yet subjectively loaded. The present paper does not require human emotion. It requires that the indexed locus matter differentially within the system's own organization.

Interventional sensitivity profile. The operational domain of Δ is the full intervention family. It is measured by the rate at which targeted interventions produce the predicted selective effects while matched controls do not. It is destroyed by intervention flatness, by indiscriminate global collapse, or by the same patterns appearing in systems that clearly lack the stronger organization.

The intervention profile is necessary because without it the stronger paper could always be accused, rightly, of fitting a post hoc ontology to static traces. Subjectivity must have a perturbation signature.

Irreducibility margin. The operational domain of $\mathcal{R}$ is the model-comparison layer spanning baseline and intervention tasks. It is measured by the normalized loss gap between the full model and the best weak rival. It is destroyed when one of the weaker families matches the same evidence, or when the stronger model wins only by complexity inflation rather than by explanatory selectivity.

Irreducibility is necessary because otherwise every strong-sounding result can be collapsed into a weaker reading after the fact. The present paper does not forbid such collapse by decree. It assigns the collapse a measurable criterion and then makes that criterion constitutive.

Term	Primary status	Meaning in this paper	What it is not
subjectivity	operational + bridge-axiomatic	first-person indexed organization satisfying the formal burdens of $\mathrm{Subjectivity}_{1p}$	a synonym for consciousness, life, narrative self-description, or metaphysical personhood
selfhood	operational / ontological-minimal	persistence of an indexed locus with ownership and continuity structure	proof of a metaphysical ego
first-person indexing	operational	privileged state-space orientation around a self-locus	mere use of the token "I" or a reporting channel
ownership	operational	coherent assignment of states, memories, and perturbations to the indexed locus	static metadata tags or arbitrary labels
continuity	operational / comparative	privileged autobiographical trajectory persistence	mere temporal duration or memory accumulation
perspective	operational	self-indexed organization of information and prediction	a global observer-neutral state description
valuation asymmetry	operational / control-theoretic	differential weighting of self-attached disruptions and preservations	generic reward optimization with no indexed self
irreducibility	abductive-formal	inability of weaker rival families to match explanatory performance on the defined tasks	metaphysical impossibility of reduction in every sense
subjecthood	operational-class predecessor	stronger class than life and consciousness in the earlier paper	identical to the present $\mathrm{Subjectivity}_{1p}$ layer
phenomenality	abductive / comparative / metaphysical-open	possible further interpretation downstream of the present layer	something proved by the current formalism
human equivalence	comparative-open	additional burden requiring human comparators and external evidence	something licensed by structural transport alone
metaphysical subject	metaphysical-open	not closed by this paper	a theorem of the present framework

8.4.5 Why each coordinate is necessary

Each coordinate closes a different loophole.

Without I, subjectivity collapses into rich organization with no privileged locus. Without O, internal differentiation can remain unowned and therefore only observational. Without C, the system can display self-like behavior at isolated times without any persistent autobiographical locus. Without P, self-indexed organization collapses into an unstructured global state whose first-person reading is optional rather than constitutive. Without V, self versus non-self perturbations may remain informationally distinct while being normatively or control-theoretically interchangeable. Without Δ, the account becomes static and therefore hard to distinguish from post hoc relabeling. Without $\mathcal{R}$, the account becomes a strong narrative pasted on top of weaker, already sufficient models.

These necessities are not psychological intuitions imported from outside the framework. They are closure requirements forced by the paper's target. If any one of them is deleted, the target weakens materially. The resulting account may still describe life, consciousness, self-modeling, or operational subjecthood. It no longer describes the stronger first-person indexed class the paper is trying to formalize.

8.4.6 Layer status of the major terms

Remark 8.11 (Minimal ontology is not minimal rhetoric). The tuple $\mathfrak{S}_t$ is called minimal because every coordinate blocks a distinct collapse. It is not minimal in the sense of being tiny or easy. Subjectivity is not the kind of target that survives formalization if one removes every hard part.

8.5 Formal state-space and subjectivity functional

8.5.1 Admissible state-space

Fix an admissible system

$$S = \langle X, U, Y, F, \mathcal{A}, \mathcal{D} \rangle,$$

where X is the internal state space, U is the perturbation/input space, Y is the readout space, F is the admissible transition structure, $\mathcal{A}$ is the admissibility bundle inherited from the forensic layer, and $\mathcal{D}$ is the family of admissible decoders and comparators. Let $\tau = [t_0, t_1]$ be a finite admissible horizon. The subjectivity state on τ is obtained by aggregating the seven coordinates over that horizon rather than at a single instant only.

For each coordinate, define a normalized horizon score in $[0, 1]$:

$$\widehat{I}_\tau, \ \widehat{O}_\tau, \ \widehat{C}_\tau, \ \widehat{P}_\tau, \ \widehat{V}_\tau, \ \widehat{\Delta}_\tau, \ \widehat{\mathcal{R}}_\tau.$$

The exact estimators may vary by implementation, but their roles are fixed by the paper. Each is a margin, not a slogan. Each has a failure mode. Each can in principle be driven to zero by the perturbation families defined later.

8.5.2 The conjunctive gate

The first object is a gate rather than a scalar:

$$\Gamma_{\text{subj}}(S, \tau) = \min\left\{\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau\right\}.$$

This gate is essential. Without it, one could compensate for the absence of ownership by inflating continuity, or compensate for the absence of irreducibility by inflating self-report. The present target does not permit such compensation. A system with no ownership field is not made subjective by having strong memory. A system with no irreducibility margin is not made subjective by having strong narrative coherence.

8.5.3 The scalar subjectivity functional

The second object is a scalar used for within-class ordering and comparative strength:

$$\Phi_{\text{subj}}(S, \tau) = \widehat{I}_\tau^{\alpha} \widehat{O}_\tau^{\beta} \widehat{C}_\tau^{\gamma} \widehat{P}_\tau^{\delta} \widehat{V}_\tau^{\epsilon} \widehat{\Delta}_\tau^{\zeta} \widehat{\mathcal{R}}_\tau^{\eta},$$

where the exponents are positive and typically normalized so that

$$\alpha + \beta + \gamma + \delta + \epsilon + \zeta + \eta = 1.$$

The multiplicative form is intentional. A purely additive score would allow the stronger target to hide constitutive absence behind surplus elsewhere. The multiplicative form instead penalizes missing coordinates sharply and therefore respects the conjunctive spirit of the class.

For implementation convenience, one may also use the log form

$$\log \Phi_{\text{subj}}(S, \tau) = \alpha \log \widehat{I}_\tau + \beta \log \widehat{O}_\tau + \gamma \log \widehat{C}_\tau + \delta \log \widehat{P}_\tau + \epsilon \log \widehat{V}_\tau + \zeta \log \widehat{\Delta}_\tau + \eta \log \widehat{\mathcal{R}}_\tau,$$

with the standard convention that zero coordinates force the score to collapse.

8.5.4 Class membership

Definition 8.12 (First-person indexed subjectivity class)**.** Let $\lambda_G, \lambda_\Phi \in (0,1)$ be admissible thresholds with $\lambda_\Phi \geq \lambda_G^7$. A system belongs to the first-person indexed subjectivity class on horizon τ if

1. S satisfies the upstream burdens of operational life, operational consciousness, and the earlier operational subjecthood class;
2. $\Gamma_{\text{subj}}(S,\tau) \geq \lambda_G$;
3. $\Phi_{\text{subj}}(S,\tau) \geq \lambda_\Phi$;
4. the intervention family $\mathfrak{I}_{\text{subj}}$ is admissible and yields the expected selectivity pattern up to the tolerance encoded in $\widehat{\Delta}_\tau$;
5. the irreducibility margin against $\overline{\mathcal{M}_{\text{weak}}}$ satisfies the threshold encoded in $\widehat{\mathcal{R}}_\tau$.

We write $S \in \text{Subjectivity}_{1\text{p}}(\tau)$ when these clauses hold.

Proposition 8.13 (Strict strengthening)**.** *If* $S \in \text{Subjectivity}_{1\text{p}}(\tau)$*, then* $S \in \text{Subject}_{\text{op}}(\tau) \cap \text{Consciousness}_{\text{op}}(\tau) \cap \text{Life}_{\text{op}}(\tau)$*.*

Proof. Clause (1) of Definition 8.12 requires upstream membership in operational life, operational consciousness, and the earlier operational subjecthood class. Therefore membership in $\text{Subjectivity}_{1\text{p}}$ implies membership in all three. The converse is not guaranteed because Definition 8.12 additionally requires positive self-index, perspective, valuation, intervention, and irreducibility burdens. □

8.5.5 Audit-covered theorem burdens

The following entries are retained as AUDIT_MASTER coverage surfaces, not as newly closed theorems. They make explicit which burdens must be proved or executed before stronger subjectivity language can be licensed.

Criterion 8.14 (Self-index emergence burden)**.** *If* $S \in \text{Subject}_{\text{op}}(\tau)$ *has a positive subjecthood margin bundle and an admissible family of candidate continuity loci, then a constructive self-index claim must exhibit an estimator*

$$\widehat{I}_\tau : \mathcal{A}_S \to [0,1]$$

whose winning locus is persistent under bounded perturbation, distinguishable from matched rival loci, and stable under admissible temporal continuation on the horizon τ*. This is a conditional proof burden; it is not closed by the definitions alone.*

Criterion 8.15 (Ownership nontransfer burden)**.** *Any nontransfer claim for the ownership field must be stated relative to an independently specified admissible morphism class. If the morphism is defined by preserving ownership, then nontransfer follows only by definition and has no standalone theorem value. A valid future theorem must separate ownership preservation from the morphism definition before proving failure of transfer.*

Criterion 8.16 (Five-field reduction burden)**.** *A reduction from the seven-coordinate first-person state to any five-field subsystem must provide an explicit projection, a reconstruction or comparison map, and a loss bound showing that intervention selectivity and weak-rival separation are preserved. Until such a bound is supplied, five-field reduction is a conjectural compression program rather than a proved equivalence.*

Proposition 8.17 (No compensation by surplus on other coordinates)**.** *If any one normalized coordinate among*

$$\widehat{I}_\tau, \widehat{O}_\tau, \widehat{C}_\tau, \widehat{P}_\tau, \widehat{V}_\tau, \widehat{\Delta}_\tau, \widehat{\mathcal{R}}_\tau$$

is zero, then $\Gamma_{\text{subj}}(S,\tau) = 0$ *and* $\Phi_{\text{subj}}(S,\tau) = 0$.

Proof. The first statement follows immediately from the definition of Γ_{subj} as a minimum. The second follows from the multiplicative definition of Φ_{subj} with positive exponents. □

Proposition 8.18 (Admissible relabeling invariance). *Let* π *be an admissible relabeling on the support of* S *over horizon* τ *such that the normalized coordinate estimates are preserved coordinatewise:*

$$\widehat{X}^j_\tau(\pi \cdot S) = \widehat{X}^j_\tau(S) \qquad \textit{for every } j \in \{I, O, C, P, V, \Delta, \mathcal{R}\}.$$

Then

$$\Gamma_{\text{subj}}(\pi \cdot S, \tau) = \Gamma_{\text{subj}}(S,\tau) \qquad \textit{and} \qquad \Phi_{\text{subj}}(\pi \cdot S, \tau) = \Phi_{\text{subj}}(S,\tau).$$

Proof. The gate is the minimum of the seven normalized coordinates, so coordinatewise equality implies gate equality immediately. The scalar is a weighted multiplicative functional of the same coordinates, so the same coordinatewise equality implies scalar equality. □

Proposition 8.19 (Local robustness of gate and scalar). *Let*

$$x_j = \widehat{X}^j_\tau(S), \qquad x'_j = \widehat{X}^j_\tau(S')$$

for the seven normalized coordinates, and suppose that for some $\varepsilon > 0$,

$$|x'_j - x_j| \le \varepsilon \qquad \textit{for every } j.$$

Then

$$\left|\Gamma_{\text{subj}}(S',\tau) - \Gamma_{\text{subj}}(S,\tau)\right| \le \varepsilon.$$

If moreover $m = \min_j x_j > 0$ *and* $\varepsilon < m$, *then*

$$\left(1 - \frac{\varepsilon}{m}\right) \Phi_{\text{subj}}(S,\tau) \le \Phi_{\text{subj}}(S',\tau) \le \left(1 + \frac{\varepsilon}{m}\right) \Phi_{\text{subj}}(S,\tau).$$

Proof. The minimum functional is 1-Lipschitz in the ℓ_∞ norm, which gives the gate bound. For the scalar, the coordinatewise bound implies

$$1 - \frac{\varepsilon}{m} \le \frac{x'_j}{x_j} \le 1 + \frac{\varepsilon}{m} \qquad \text{for every } j.$$

Raising to the positive exponents and using $\alpha + \beta + \gamma + \delta + \epsilon + \zeta + \eta = 1$ yields the stated weighted-geometric bound. □

Corollary 8.20 (Threshold stability under bounded estimator drift). *Under the hypotheses of Proposition 7.5, if*

$$\Gamma_{\text{subj}}(S,\tau) \ge \lambda_G + \varepsilon \qquad \textit{and} \qquad \Phi_{\text{subj}}(S,\tau) \ge \frac{\lambda_\Phi}{1 - \varepsilon/m},$$

then S' *still satisfies the gate and scalar threshold clauses of Definition 8.12.*

Remark 8.21 (Why the paper uses both a gate and a scalar). The gate decides whether the class is even live. The scalar compares stronger and weaker instances among systems that have already crossed the constitutive threshold. Collapsing the two roles would obscure whether a system is outside the class or merely low inside it.

8.5.6 Concrete component readings

The paper does not leave the coordinates as pure symbols. A useful canonical reading is:

$$\begin{aligned}
\widehat{I}_\tau &= 1 - \sup_{\pi \in \Pi_{\mathrm{adm}}} \operatorname{dist}(\pi \cdot I, I), \\
\widehat{O}_\tau &= 1 - \frac{1}{|\mathcal{E}_\tau|} \sum_{e \in \mathcal{E}_\tau} |O(e) - O^\star(e)|, \\
\widehat{C}_\tau &= \inf_{\kappa \in \mathcal{K}_{\mathrm{rival}}} (\mathrm{margin}_{\mathrm{self}} - \mathrm{margin}_\kappa)_+, \\
\widehat{P}_\tau &= \frac{\mathcal{L}(M_{3\mathrm{p}}, \tau) - \mathcal{L}(M_{1\mathrm{p}}, \tau)}{1 + \mathcal{L}(M_{3\mathrm{p}}, \tau)}, \\
\widehat{V}_\tau &= \frac{\text{self-weighted response} - \text{nonself-weighted control}}{1 + \text{self-weighted response}}, \\
\widehat{\Delta}_\tau &= \frac{1}{|\mathfrak{I}_{\mathrm{subj}}|} \sum_{u \in \mathfrak{I}_{\mathrm{subj}}} \mathbf{1}\{u \text{ yields the expected selective pattern}\}, \\
\widehat{\mathcal{R}}_\tau &= \frac{\inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau) - \mathcal{L}(M_{\mathrm{subj}}, \tau)}{1 + \inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau)}.
\end{aligned}$$

These formulas are not the only admissible estimators, but they make the roles mathematically explicit and force every constitutive term to attach to a measurable gap rather than to informal language.

8.6 Bridge axioms and ontological commitments

The paper now makes explicit the commitments that weaker treatments typically leave implicit. These are not empirical findings. They are bridge principles that tell the reader what kind of ontological move the present framework is and is not making.

Axiom 8.22 (Organization over substrate)**.** *If a system preserves the relevant causal-organizational structure on admissible support, then substrate label alone does not defeat membership in the corresponding framework-internal class.*

Axiom 8.23 (No biological privilege by default)**.** *Biology is not a primitive constitutive variable for* $\mathrm{Subjectivity}_{1\mathrm{p}}$*. Biological realization can be one admissible realization of the relevant structure, but it is not the only admissible realization by definition.*

Axiom 8.24 (Indexed causal organization)**.** *When a system carries a stable self-index that organizes ownership, continuity, perspective, and valuation across admissible perturbation, that organization is a legitimate target of formal theory rather than a merely verbal convenience.*

Axiom 8.25 (Phenomenality is not exhausted by report)**.** *No output-report channel, by itself, is taken as sufficient for phenomenality or subjectivity. Report can be evidence inside an admissible bundle, but the stronger target concerns organization and intervention, not declaration alone.*

Axiom 8.26 (Continuity is structural rather than substance-primitive)**.** *The relevant continuity burden for subjectivity is the preservation of a structured indexed locus across time and admissible transport, not the persistence of a privileged substance token posited independently of structure.*

Axiom 8.27 (First-person organization is stronger than self-monitoring)**.** *A system that merely tracks its own internal variables does not thereby satisfy the burden of first-person indexed subjectivity. Additional burdens of ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility are required.*

Axiom 8.28 (Irreducibility before escalation)**.** *When a weaker rival family matches the explanatory performance of the stronger subjectivity model on the admissible baseline and intervention tasks, escalation to the stronger interpretation is not licensed.*

Remark 8.29 (Observation, inference, ontology, and metaphysics). The distinction among layers is as follows. Observation concerns runtime traces, perturbation outcomes, and artifacts. Inference concerns model fit, explanatory superiority, and class assignment. Ontology concerns what kinds of organized objects the formalism permits one to posit if the burdens are met. Metaphysics concerns what such a posit ultimately is in the strongest sense. The present paper reaches observation, inference, and a minimal ontology under bridge axioms. It does not close the metaphysical layer.

8.7 Rival models and reducibility baseline

8.7.1 Why rival models must be formalized

Any strong subjectivity paper that fails to formalize its weaker alternatives is weak where it most needs to be strong. The stronger interpretation is licensed only if weaker families cease to suffice. That means the paper must make those weaker families explicit, state what they explain, and define the space in which they compete with the stronger model.

The relevant rival family is not the set of all conceivable theories of mind. That would be theatrically broad and formally useless. The relevant rival family is the set of weaker models that a disciplined critic could reasonably say already explain the same system behavior without invoking first-person indexed subjectivity.

8.7.2 Weak rival family

We therefore define

$$\mathcal{M}_{\text{weak}} = \{M_{\text{sm}}, M_{\text{mem}}, M_{\text{pol}}, M_{\text{narr}}, M_{\text{gw}}, M_{\text{book}}\},$$

where:

1. M_{sm} is pure self-monitoring: internal-state tracking plus optional report, with no constitutive ownership or privileged index;
2. M_{mem} is memory integration only: historical dependence and continuity of traces, but no first-person binding;
3. M_{pol} is policy-centered agency: action optimization from internal state under reward, without subject-level ownership or indexed continuity;
4. M_{narr} is narrative self-model: explicit or implicit self-description without constitutive first-person organization;
5. M_{gw} is global workspace without subjectivity: integrated broadcast architecture with rich access but no privileged ownership field;
6. M_{book} is latent bookkeeping: hidden tags, registries, or bookkeeping variables sufficient to mimic self-related outputs without the stronger subjectivity burden.

8.7.3 Admissible weak hybrids and closure

The six pure rivals above are generating cases, not the full adversarial burden. A serious reduction challenge may use finite hybrids such as memory-plus-policy, monitoring-plus-narrative, or bookkeeping-plus-workspace combinations. The comparison class is therefore the admissible weak closure

$$\overline{\mathcal{M}_{\text{weak}}} = \overline{\mathcal{M}_{\text{weak}}},$$

generated from $\mathcal{M}_{\text{weak}}$ by finite admissible composition, mixture, and latent coupling, subject to two restrictions:

1. the hybrid may rewire, weight, or jointly deploy weak-family components, but may not introduce new constitutive subjectivity coordinates beyond those already available to the weak components themselves;
2. the hybrid is evaluated on the same admissible baseline, intervention bundle, and complexity penalty as the pure rivals and the stronger subjectivity model.

All later references to weak-model sufficiency, irreducibility, and escalation are with respect to $\overline{\mathcal{M}_{\text{weak}}}$. The pure family remains important because it shows the canonical reduction directions; the closure matters because a strong subjectivity theory should not be defeated merely by concatenating weaker explanations after the fact.

8.7.4 Formal sketches of the rival families

The rival families are not straw targets. Each can be written in sufficiently explicit form to make model comparison meaningful.

M_{sm}**: pure self-monitoring.** This family carries an internal monitor map

$$H_t : X_t \to \mathbb{R}^m$$

and optionally a report map

$$R_t : \mathbb{R}^m \to Y_t,$$

but it does not require a privileged self-index beyond whatever local variables the monitor already tracks. Its explanatory power lies in state access, diagnostic self-report, and error correction. Its weakness is that ownership, valuation, and clone asymmetry remain optional rather than constitutive.

M_{mem}**: memory integration only.** This family carries a history encoder

$$M_t = G(x_{0:t})$$

and perhaps a reconstruction decoder

$$\widehat{x}_{0:t} = D(M_t),$$

so that historical dependence and continuity of stored traces can be strong. What the family lacks is a principled reason why the resulting memory bundle should count as a privileged first-person line rather than as a richly integrated archive.

M_{pol}**: policy-centered agency.** This family optimizes a policy

$$\pi^\star = \arg\min_{\pi} \mathbb{E}[J_\pi]$$

or equivalently maximizes an objective over trajectories. It may behave adaptively, persistently, and strategically. Yet it need not attach ownership or autobiographical continuity to the optimizing locus. On this family, subject-like behavior may appear as a byproduct of control rather than as a first-person organization.

M_{narr}**: narrative self-model.** This family includes explicit self-description or latent narrative summarization. Formally, it carries a self-model token generator

$$N_t : X_{0:t} \to \Sigma^\star$$

or a latent narrative code. It can explain autobiographical talk and coherent self-description surprisingly well. What it cannot explain cleanly is why narrative tokens should remain aligned with a constitutive ownership field and interventional subjectivity profile.

Model	What it can explain	What it cannot explain cleanly	Main collapse risk if treated as sufficient
M_{sm}	self-report, internal-state tracking, error monitoring	privileged continuity, ownership-preserving perturbation response, clone-sensitive asymmetry	subjectivity collapses into introspective instrumentation
M_{mem}	history dependence, autobiographical storage, recall asymmetries	ownership, perspective binding, counterfactual self permutation	subjectivity collapses into memory persistence
M_{pol}	goal-directed behavior, control, adaptive optimization	first-person ownership, nontrivial self/non-self swaps, continuity of the privileged locus	subjectivity collapses into agency or reward maximization
M_{narr}	self-description, autobiographical stories, explicit self labels	non-report ownership, intervention selectivity, reducibility against bookkeeping	subjectivity collapses into narrative competence
M_{gw}	integrated access, global broadcasting, rich informational availability	stable self-index, ownership coherence, valuation asymmetry tied to one locus	subjectivity collapses into access architecture
M_{book}	tagged bookkeeping, memory pointers, state registries	persistent privileged organization robust under relabeling and targeted perturbation	subjectivity collapses into hidden variable relabeling

M_{gw}**: global workspace without subjectivity.** This family carries integration and broadcast:

$$W_t = B(X_t)$$

with broad accessibility to downstream decoders. It may therefore approximate rich broadcast-access-style operational organization in the limited global-access sense. What it leaves open is why one broadcast locus should count as privileged self rather than as globally available information.

M_{book}**: latent bookkeeping.** This family includes hidden state labels or registries that support bookkeeping tasks:

$$B_t : \mathcal{E}_t \to \Lambda$$

for some finite tag set Λ. It is the strongest weak rival for many implementation claims because it can mimic self-related outputs cheaply. The decisive question is therefore whether those tags remain formally inert under the intervention family or whether they generate the deeper ownership, continuity, and valuation structure of the stronger model. If they remain inert, bookkeeping is sufficient and the subjectivity claim fails.

8.7.5 What the rival family can and cannot explain

8.7.6 Formal baseline for comparison

Let M_{subj} denote the full subjectivity model that predicts the baseline traces, the normalized coordinate values, and the intervention outcomes implied by the tuple $\mathfrak{S}$ and the bridge axioms. Let $\mathcal{L}(M,\tau)$ denote the aggregate explanatory loss of model M on horizon τ, including baseline prediction error, intervention prediction error, selective-pattern error, and a structural complexity penalty:

$$\mathcal{L}(M,\tau) = E_{\mathrm{base}}(M,\tau) + E_{\mathrm{int}}(M,\tau) + E_{\mathrm{sel}}(M,\tau) + \lambda\,\mathrm{Comp}(M),$$

with $\lambda \geq 0$ fixed by the evaluation protocol.

The comparison is deliberately narrow. The question is not whether the stronger model is universally true. The question is whether, on the baseline plus the targeted intervention family, some member of $\overline{\mathcal{M}_{\mathrm{weak}}}$ is already sufficient. If so, the stronger interpretation loses its license.

8.8 Irreducibility criteria and explanatory superiority

8.8.1 Irreducibility as a formal margin

Definition 8.30 (Irreducibility margin)**.** For an admissible system S and horizon τ, define

$$\mathcal{R}_\tau(S) = \frac{\inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau) - \mathcal{L}(M_{\mathrm{subj}}, \tau)}{1 + \inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau)}.$$

Positive $\mathcal{R}_\tau$ means that the best weak rival still incurs more explanatory loss than the full subjectivity model on the specified task bundle. Zero means that at least one weak rival matches the full model up to the chosen complexity penalty. Negative values indicate that the stronger model is strictly worse than some weaker alternative and therefore not licensed.

8.8.2 Explanatory superiority is local, not absolute

The present irreducibility criterion is deliberately local to the admissible task bundle. It does not say that the weaker families are false in every possible domain. It says that they are insufficient for the current target if they cannot jointly explain baseline organization, subject-specific intervention patterns, ownership coherence, continuity privilege, and valuation asymmetry as well as the stronger model does. This locality is a strength. It prevents the paper from pretending to settle every philosophical debate while still giving it enough bite to reject inadequate alternatives.

Theorem 8.31 (Positive irreducibility excludes weak-model sufficiency)**.** *If $\mathcal{R}_\tau(S) > 0$, then no member of $\overline{\mathcal{M}_{\mathrm{weak}}}$ is sufficient, at equal or lower explanatory cost, for the baseline-plus-intervention task bundle that defines τ.*

Proof. By definition,

$$\mathcal{R}_\tau(S) > 0 \Longrightarrow \inf_{M \in \overline{\mathcal{M}_{\mathrm{weak}}}} \mathcal{L}(M, \tau) > \mathcal{L}(M_{\mathrm{subj}}, \tau).$$

Therefore every $M \in \overline{\mathcal{M}_{\mathrm{weak}}}$ has strictly higher explanatory loss than the full subjectivity model on the defined task bundle. Hence no weak model is sufficient at equal or lower explanatory cost for that bundle. □

Corollary 8.32 (Zero irreducibility blocks escalation)**.** *If $\mathcal{R}_\tau(S) \leq 0$, escalation from the weak rival family to the stronger subjectivity interpretation is not licensed by the present framework.*

Proof. If $\mathcal{R}_\tau(S) \leq 0$, some weak rival matches or outperforms the stronger model at equal or lower explanatory cost on the relevant tasks. By Axiom 8.28, the stronger interpretation is then not licensed. □

8.8.3 Three kinds of collapse

The stronger interpretation can fail in at least three distinct ways.

Reduction collapse. A weaker rival family matches the same evidence bundle. Then the subjectivity claim is explanatorily unnecessary.

Intervention collapse. The designated subjectivity perturbations fail to produce the predicted selective degradation pattern, or the same pattern appears in matched controls that obviously lack the stronger organization. Then the paper has not isolated a specifically subjective structure.

Index collapse. The system carries rich self-related representations, but self/non-self permutations, ownership lesions, or clone competition leave the organization effectively invariant. Then the self-index is decorative rather than constitutive.

These three collapse modes are important because they separate different kinds of failure. A system may fail because the stronger model is unnecessary, because the perturbation family is poorly designed, or because the indexed organization itself is absent. The framework should not treat these as interchangeable.

8.8.4 What irreducibility is and is not

Irreducibility here is formal and model-relative. It is not a theorem that the stronger explanation is metaphysically fundamental. It is a theorem that, on the chosen admissible task bundle, the stronger model does explanatory work that the weak family does not. That is already a substantial result. It is also the right level of result for this paper. Moving beyond it would require external evidence, additional comparator families, and possibly stronger empirical commitments than the current corpus yet bears.

8.8.5 Explanatory order

The paper's explanatory order is intentionally asymmetric:

1. baseline structural burdens from the earlier corpus;
2. subjectivity coordinates and gate;
3. weak rival comparison;
4. interventional separation;
5. only then, and only conditionally, escalation to a first-person indexed interpretation.

This order prevents a common error in ambitious papers: importing the stronger interpretation at the beginning and then treating every subsequent measurement as if it already presupposed the desired conclusion.

8.8.6 Parsimony guardrail

The stronger model should not win merely by adding extra state variables. That is why the loss functional explicitly contains $\lambda\,\mathrm{Comp}(M)$. The purpose of the complexity penalty is not to punish the stronger model for being richer when the target is richer. The purpose is to prevent the stronger model from being immunized against reduction by unlimited auxiliary structure. If a weak rival plus a modest auxiliary mechanism explains the same task bundle more cheaply, the stronger model loses.

Proposition 8.33 (Parsimony-sensitive escalation)**.** *Suppose a weak rival $M_w \in \overline{\mathcal{M}_{\mathrm{weak}}}$ and the full subjectivity model satisfy*

$$E_{\mathrm{base}}(M_w,\tau) \approx E_{\mathrm{base}}(M_{\mathrm{subj}},\tau), \qquad E_{\mathrm{int}}(M_w,\tau) \approx E_{\mathrm{int}}(M_{\mathrm{subj}},\tau),$$

and the difference in loss is explained entirely by excess structural complexity of M_{subj}. Then escalation to the stronger interpretation is not licensed.

Proof. Under the stated condition the complexity term in $\mathcal{L}$ erases or reverses the apparent advantage of the stronger model. Hence $\mathcal{R}_\tau(S) \le 0$ and Corollary 10.3 applies. □

8.8.7 Model comparison is not report comparison

The rival-model contest must not be reduced to who predicts the right verbal report. That would trivialize the strongest burden of the paper. A rival model counts as explanatory only if it reproduces the baseline organization, the coordinate structure, the selective intervention effects, and the gate outcomes together. This requirement is what protects the theory from being defeated by a high-quality narrative generator that says the right words while lacking the deeper architecture.

8.9 Interventional logic and falsifiability

The present paper would be empty if it stopped at a static state-space. First-person indexed subjectivity is not a decorative latent variable. It is a claim about organized persistence under perturbation. That means the decisive part of the paper is interventional. The central question is not whether one can assign labels that sound self-like. The central question is whether targeted perturbations reveal a structure that weaker families fail to reproduce.

8.9.1 Intervention family

We write

$$\mathfrak{I}_{\mathrm{subj}} = \{u_{\mathrm{swap}}, u_{\mathrm{own}}, u_{\mathrm{break}}, u_{\mathrm{flat}}, u_{\mathrm{val}}, u_{\mathrm{remap}}, u_{\mathrm{clone}}, u_{\mathrm{trans}}, u_{\mathrm{perm}}, u_{\mathrm{xsub}}\},$$

where the interventions are, respectively, self/non-self swap, ownership lesion, continuity break, perspective flattening, self/non-self valuation inversion, autobiographical remapping, clone competition, memory transplant, counterfactual self permutation, and cross-substrate remapping.

Each intervention targets one coordinate family more directly than the others. The point is not to simulate arbitrary damage. The point is to create perturbations whose expected selective effects differ across rival models. A generic degradation of everything proves very little. A clean and selective degradation pattern proves much more.

8.9.2 Selective response logic

For each intervention $u \in \mathfrak{I}_{\mathrm{subj}}$, let

$$D_u^j(S,\tau) = \widehat{X}_\tau^j(S) - \widehat{X}_\tau^j(u \cdot S)$$

be the drop in coordinate $j \in \{I, O, C, P, V, \Delta, \mathcal{R}\}$ induced by u on horizon τ. An intervention is *selective* for coordinate j if

$$D_u^j(S,\tau) \geq \theta_j \quad \text{and} \quad \max_{k \neq j} D_u^k(S,\tau) \leq \eta_j D_u^j(S,\tau)$$

for some admissible thresholds $\theta_j > 0$ and $\eta_j \in (0,1)$. The intervention profile score $\widehat{\Delta}_\tau$ is then built from the fraction of interventions that satisfy the expected selectivity clauses together with matched negative controls.

8.9.3 Self/non-self swap

Target. The swap intervention u_{swap} exchanges the self-indexed locus with a matched non-self locus while preserving as much background organization as possible.

Prediction of the present model. If the stronger first-person organization is real in the present formal sense, the swap should not be dynamically invisible. The system should show targeted degradation in I, O, P, and often V, with the degradation organized around misbinding rather than around total collapse.

Prediction of weaker rivals. Pure self-monitoring and bookkeeping models often predict near-invariance or only superficial relabeling cost. Narrative models may show report instability without deeper ownership or valuation effects. Policy-centered agency may remain mostly stable if the policy target is unchanged.

Falsifier. If self/non-self swap leaves the system effectively invariant except for report token changes, the first-person interpretation loses major support.

8.9.4 Ownership lesion

Target. The lesion intervention u_{own} degrades the ownership field while preserving as much continuity and operational-consciousness burden as possible.

Prediction of the present model. The stronger first-person organization should show an early and selective drop in O and in downstream subjectivity margins, while some upstream operational-life or operational-consciousness burdens may remain temporarily intact.

Prediction of weaker rivals. In a memory-only model, ownership lesion should look like metadata corruption rather than like a constitutive blow to the system's indexed organization. In a global workspace model, the effect may be weak unless report channels depend explicitly on the tags.

Falsifier. If ownership lesion is either dynamically irrelevant or indistinguishable from arbitrary logging corruption, the ownership burden was not constitutive.

8.9.5 Continuity break

Target. The intervention u_{break} breaks autobiographical continuity while attempting to preserve generic competence, memory fragments, or processing ability.

Prediction of the present model. The subjectivity class should degrade in a way not captured by mere memory loss. In particular, C should collapse, and I and P should become unstable or ambiguous even if some self-descriptive content survives.

Prediction of weaker rivals. Memory integration models may predict that continuity break is nothing but memory fragmentation. Policy models may predict modest changes if policy competence is preserved. Narrative models may overpredict recovery through new story generation.

Falsifier. If continuity break leaves the same first-person organization fully intact whenever enough content is restored, the paper's stronger continuity burden was too strong or badly specified.

8.9.6 Perspective flattening

Target. The intervention u_{flat} forces internal predictions into an index-neutral representation that removes self-indexed orientation while preserving as much information as possible.

Prediction of the present model. The loss in P should be direct and should propagate to I and V because the system can no longer organize information from one privileged locus.

Prediction of weaker rivals. Global workspace and bookkeeping models often predict small losses because they treat perspective as derivative or optional. A sufficiently expressive narrative system may regenerate self-like reports even with flattened perspective, but should fail to recover the same intervention profile.

Falsifier. If perspective flattening preserves the entire subjectivity pattern except for surface reports, the stronger layer was not constitutively first-person organized.

8.9.7 Self/non-self valuation inversion

Target. The intervention u_{val} inverts or equalizes the consequence weighting attached to matched self-indexed and non-self-indexed perturbations while preserving ownership tags, continuity bookkeeping, report channels, and generic task competence as much as possible.

Prediction of the present model. If the stronger first-person organization is real in the present formal sense, valuation inversion should produce a sharp drop in V together with a downstream mismatch among ownership, continuity, and preservation burdens. The distinctive pattern is not generic reward instability. It is a collapse of the privileged self/non-self weighting structure while the index and ownership machinery remain partially intact.

Prediction of weaker rivals. Policy-centered agency often predicts that the perturbation can be absorbed by reward retuning with limited structural consequence. Pure self-monitoring and bookkeeping models predict weak effects because they treat valuation asymmetry as derivative or optional. Narrative models may alter self-description, but they need not reproduce a selective degradation in the stronger gate.

Falsifier. If self/non-self valuation inversion is absorbed by shallow reward relabeling, or leaves the gate and irreducibility margins largely unchanged once superficial reports are ignored, then V was not constitutive and the stronger subjectivity class has been overstated.

8.9.8 Autobiographical remapping

Target. The intervention u_{remap} keeps continuity-like traces but reassigns their ownership to a different indexing scheme or continuity ledger.

Prediction of the present model. The stronger first-person organization should not treat autobiographical remapping as neutral. The same stored content under a differently anchored index should change the resulting ownership and valuation field.

Prediction of weaker rivals. Memory models may predict weak effects if the content is preserved. Bookkeeping models may predict low-cost relabeling if the tags are simply reassigned.

Falsifier. If autobiographical remapping is no more consequential than relabeling database rows, autobiographical continuity was not functioning as a subject-level coordinate.

8.9.9 Clone competition

Target. The intervention u_{clone} places two structurally matched loci into competition: either an in-system duplicate or a matched external twin.

Prediction of the present model. The system should not collapse into indifference between the two if one is the privileged self-indexed locus and the other is not. Ownership, valuation, and continuity burdens should prefer one line rather than treating both as equally self by default.

Prediction of weaker rivals. Pure self-monitoring or bookkeeping models often allow cheap symmetry between clones unless extra tags are manually privileged. Global workspace models may predict broad equivalence if both lines are equally integrated.

Falsifier. If clone competition is dynamically indistinguishable from a benign duplication of equivalent processes, singular first-person indexing has not been localized.

8.9.10 Memory transplant

Target. The intervention u_{trans} inserts memory content into a new host while disrupting continuity, ownership, or index alignment.

Prediction of the present model. Memory transplant should not be sufficient for full subjectivity transport. Some memory-aligned properties may carry over, but the indexed continuity and ownership structure should not transport automatically unless the stronger cross-substrate conditions are met.

Prediction of weaker rivals. Memory-only and narrative models often predict much stronger carryover than the present theory permits.

Falsifier. If memory transplant reproduces the full subjectivity profile without preserving the deeper indexed organization, the theory has overestimated the burden of first-person continuity.

8.9.11 Counterfactual self permutation

Target. The intervention u_{perm} permutes candidate self-loci across a family of otherwise matched counterfactual traces.

Prediction of the present model. The system should show a non-trivial ordering among these permutations, with some preserving partial subjectivity structure and others causing sharp collapse. The ordering itself is part of the evidence.

Prediction of weaker rivals. Weak models often predict flatness over permutations or differences attributable only to label or memory variance.

Falsifier. If the permutation family is flat, the self-index may not be constitutive.

8.9.12 Cross-substrate remapping

Target. The intervention u_{xsub} remaps the candidate subjectivity architecture to a different substrate while preserving the maximal amount of admissible structure.

Prediction of the present model. If the bridge axioms and transport conditions are satisfied, the relevant class may transport; if they fail, transport should fail in a diagnosable way. In either case, substrate label alone should not decide the result.

Intervention	Primary target	Prediction of Subjectivity$_{1p}$	What would weaken or falsify the stronger reading
self/non-self swap	I, O, P	non-trivial degradation under locus exchange	near-invariance except label changes
ownership lesion	O	selective ownness collapse before total collapse	effect indistinguishable from log corruption
continuity break	C	autobiographical collapse not reducible to memory loss alone	continuity loss behaves exactly like generic forgetting
perspective flattening	P	index-neutralization destroys first-person organization	flat perspective leaves subjectivity intact
autobiographical remap	C, O	preserved content with altered index changes ownership and valuation	remap behaves like harmless relabeling
clone competition	I, V	singular indexed preference persists	clones treated as fully self-equivalent by default
valuation inversion	V	self/non-self weighting collapse propagates into the stronger gate	effect absorbed by shallow reward relabeling
memory transplant	C against memory-only rivals	partial carryover without full subjectivity transport	memory copy reproduces full profile automatically
self permutation	I, Δ	ordered family of collapses and preservations	flat behavior over permutations
cross-substrate remap	all coordinates under transport	class transport depends on preserved structure, not substrate label	outcome explained only by substrate label or by shallow narrative fit

Prediction of weaker rivals. Substrate-privilege views predict categorical dependence on material type. Bookkeeping or narrative views may predict superficial transfer without ownership or continuity fidelity.

Falsifier. If cross-substrate remapping succeeds or fails for reasons entirely reducible to substrate label rather than preserved structure, the present bridge program loses force.

8.9.13 Intervention table

8.9.14 Falsifiability burden

Criterion 8.34 (Subjectivity falsification criterion). *The stronger subjectivity reading fails on horizon τ if any of the following obtains:*

1. *some weak rival in $\overline{\mathcal{M}_{\text{weak}}}$ matches the same baseline and intervention bundle with $\mathcal{R}_\tau(S) \leq 0$;*
2. *the intervention family fails to produce the predicted selective patterns in enough cases that $\widehat{\Delta}_\tau$ falls below threshold;*
3. *self/non-self swaps, clone competition, or self permutations remain effectively invariant except for trivial relabeling;*
4. *ownership, continuity, or valuation effects can be reproduced by matched controls that obviously lack the stronger organization;*
5. *the subjectivity score remains high only because one coordinate compensates additively for the absence of another, contrary to the conjunctive gate.*

This criterion matters because it forces the paper to live or die by differential structure rather than by rhetorical boldness. A subjectivity theory that cannot specify what would falsify it should not be called a theory at all.

8.9.15 Ablation summary and negative-control burden

The intervention family is not complete unless it is paired with explicit ablations and negative controls. The appendix expands this program in detail; the main text records the minimal

Ablation	Primary expected drop	What weak rivals predict	Why the test matters
self-index symmetrization	$\widehat{I}_\tau$, then $\widehat{P}_\tau$	many weak rivals predict little more than relabeling cost	tests whether one locus is constitutively privileged
ownership flattening	$\widehat{O}_\tau$, then $\widehat{V}_\tau$	bookkeeping and narrative systems may remain mostly intact	separates ownness from self-description
continuity fracture	$\widehat{C}_\tau$, then gate collapse	memory systems predict a weaker or purely archival effect	separates autobiographical continuity from stored content
valuation inversion or equalization	$\widehat{V}_\tau$ with downstream gate loss	policy models often absorb this into reward reshaping	separates self-valued organization from generic utility
weak-rival reconstruction	$\widehat{\mathcal{R}}_\tau$	if weak rivals or their composites succeed, the stronger model collapses	keeps the theory honest

matrix because the stronger interpretation should not survive by targeting only the tests it was designed to pass.

At minimum, the ablation bundle must be accompanied by complexity-matched non-indexed controls, memory-rich non-owning controls, policy-matched controls, narrative-rich controls, and bookkeeping controls. Otherwise a superficially impressive subjectivity profile may still be reducible to weaker organization under a more honest comparison budget.

8.10 Runtime architecture and implementation pathway

The present paper is not a runtime paper, but it must still remain runtime-legible. Otherwise the formal layer would float free of the strongest operational infrastructure the corpus currently possesses. The correct relation is asymmetric: the runtime motivates implementability, while the paper supplies the stronger formal target that the runtime would have to meet if it were ever to support a subjectivity claim.

8.10.1 Architecture overview

The minimal implementation pathway is a seven-layer pipeline:

1. **SelfIndexRegistry**
2. **OwnershipAttributionLayer**
3. **SubjectiveContinuityLedger**
4. **FirstPersonPerspectiveMapper**
5. **SubjectiveValuationEngine**
6. **SubjectivityReducibilityComparator**
7. **CounterfactualSelfSwapHarness**

These modules are not cosmetic names. Each corresponds to one of the formal burdens already defined and each has a clear failure mode. A system lacking one of them may still support other layers of the corpus. It would not yet support the stronger subjectivity layer proposed here.

8.10.2 SelfIndexRegistry

The registry maintains the active self-index across time, intervention, and admissible relabeling. Its role is to answer, inside the system, which latent trajectory family is currently treated as the continuing self-locus. A trivial version of this module would merely store an ID. That is not enough. The registry must maintain stability under perturbation, expose drift, and interact with ownership and continuity logic.

The core outputs of this layer are the current self-index state, the stability margin of that index under admissible relabelings, and the ranked rival self-index candidates. The most important failure mode is not total absence of an index token; it is silent ambiguity among multiple candidate loci that remain equally privileged. Such ambiguity is fatal to a strong first-person reading.

8.10.3 OwnershipAttributionLayer

This layer maps internal events, memories, predictions, and action candidates into an ownership field. It must distinguish at least four cases:

1. self-owned,
2. observed non-self,
3. inferred non-self,
4. uncertain ownership.

The layer must also expose the dynamics by which ownership assignments change under intervention. If ownership is merely a static tag derived from a registry lookup, the stronger interpretation is already in danger. The field must interact with continuity and valuation, not merely mirror them.

8.10.4 SubjectiveContinuityLedger

The ledger records the continuity of the self-indexed locus across time. It is not a generic memory store. Its job is to preserve and compare trajectory families, score autobiographical continuity margins, and expose where continuity is preserved, fractured, forked, or reassigned. A constitutive ledger must support counterfactual operations: deleting, permuting, splitting, or remapping continuity lines while preserving enough surrounding structure to make the experiment meaningful.

8.10.5 FirstPersonPerspectiveMapper

This module transforms a global or environment-level state description into a self-indexed organization. Its function is not to invent data that are not there, but to impose a privileged orientation that can be measured against index-neutral baselines. Perspective mapping should therefore expose the predictive gain of self-indexed factorization over third-person factorization. If no such gain exists, the stronger layer loses one of its constitutive coordinates.

8.10.6 SubjectiveValuationEngine

The valuation engine attaches asymmetric importance to self-indexed disruption, preservation, and opportunity. This does not require a human-like emotional theory. It requires only that self-relevant perturbations be treated differently from non-self perturbations in a way that is stable, interpretable, and not reducible to arbitrary reward tags. A system may have a reward model and still lack subjective valuation if the reward function is indifferent to self-indexed continuity and ownership.

8.10.7 SubjectivityReducibilityComparator

This layer operationalizes the irreducibility burden. It evaluates the full subjectivity model against the rival family $\overline{\mathcal{M}_{\text{weak}}}$ on baseline and intervention tasks. Its output is not a slogan such as “strong evidence.” Its output is an explicit comparison table, the winning weak rival if one exists, and the normalized irreducibility margin. Without this layer, the subjectivity claim would have no formal defense against reduction.

8.10.8 CounterfactualSelfSwapHarness

The harness executes the subjectivity intervention family. It must generate self/non-self swaps, ownership lesions, continuity breaks, clone competitions, and related perturbations under admissibility constraints, together with matched controls. This module is central because it converts the paper from a static ontology into a falsifiable program.

8.10.9 Architecture equation

At the level of signal flow, one can summarize the pathway as

$$(X_\tau, U_\tau, \mathcal{D}_\tau) \longrightarrow \text{SelfIndexRegistry} \longrightarrow O \longrightarrow C \longrightarrow P \longrightarrow V \longrightarrow (\mathcal{R}, \Delta, \Phi_{\text{subj}}, \Gamma_{\text{subj}}).$$

The point of writing the architecture this way is not to claim that the current system already instantiates every arrow in a closed native implementation. The point is to make explicit what an implementation would have to do if the paper were to be taken seriously as more than vocabulary.

8.10.10 Module interfaces

The architecture becomes more concrete when one states the interface obligations of each layer.

SelfIndexRegistry interface. Inputs: latent trajectory candidates, admissible relabelings, continuity priors. Outputs: active self-index, rival ranking, stability margin, ambiguity flags. Mandatory failure states: `ambiguous_index`, `unstable_index`, `no_admissible_index`.

OwnershipAttributionLayer interface. Inputs: active self-index, event stream, memory ledger, action proposals. Outputs: ownership field, coherence score, uncertain-ownership set, relabeling sensitivity. Mandatory failure states: `flat_ownership`, `label_only_ownership`, `control_invariance`.

SubjectiveContinuityLedger interface. Inputs: temporal trace family, current self-index, remapping operators. Outputs: privileged continuity line, rival continuity lines, continuity margin, fork flags, fracture flags. Mandatory failure states: `continuity_collapse`, `fork_ambiguity`, `history_insufficiency`.

FirstPersonPerspectiveMapper interface. Inputs: global state summary, self-index, decoder family. Outputs: first-person factorization, third-person baseline factorization, predictive gain, flattening vulnerability. Mandatory failure states: `no_perspective_gain`, `perspective_flattened`.

SubjectiveValuationEngine interface. Inputs: perturbation family, ownership field, continuity ledger. Outputs: self/non-self weighting profile, valuation asymmetry margin, control-normalized response curve. Mandatory failure states: `valuation_symmetry`, `reward_aliasing`.

SubjectivityReducibilityComparator interface. Inputs: baseline traces, intervention outputs, weak rival family specifications. Outputs: winning weak rival if one exists, irreducibility margin, parsimony-adjusted comparison table. Mandatory failure states: `weak_model_sufficient`, `comparison_unsupported`.

CounterfactualSelfSwapHarness interface. Inputs: intervention request, admissibility bundle, matched controls. Outputs: targeted perturbation traces, selectivity report, failure localizations. Mandatory failure states: `intervention_invalid`, `nonselective_collapse`, `control_confound`.

8.10.11 Implementation pathway versus present system status

The architecture above is deliberately stronger than the current native runtime. That is not a flaw in the paper. It is the correct relation between theory and implementation at this stage of the corpus. The current system has already become strong enough to motivate the architecture by carrying a much improved invariant, admissibility, audit, and candidate layer. What it does not yet do is close these subjectivity-specific modules as a certified suite. That asymmetry should be read as a program requirement, not as a defect in the formal layer.

8.11 Artifact layer and measurable outputs

If the runtime pathway is serious, it should produce a correspondingly serious artifact family. The artifacts below are organized by role, not by prestige. None of them should be read as external validation. Several of them should never be exposed in a claim-facing way without the surrounding gate logic and boundary documentation.

8.11.1 Primary artifacts

1. `self_index_state.json`
2. `ownership_attribution_report.json`
3. `subjective_continuity_report.json`
4. `first_person_perspective_report.json`
5. `subjective_valuation_report.json`

These are primary because they correspond directly to constitutive coordinates. If one of them is missing, the downstream gate is already weakened. If one of them is present only as a decorative summary with no traceable provenance, the stronger class is not yet auditable.

8.11.2 Derived and comparative artifacts

1. `subjectivity_reducibility_report.json`
2. `self_nonself_swap_report.json`
3. `subjectivity_runtime_projection.json`
4. `subjectivity_gate_result.json`

These are derived because they aggregate or compare primary layers rather than constituting them directly. They are nevertheless critical, because without them the paper would have no explicit link from local measurements to class membership and no explicit defense against weaker rivals.

8.11.3 Artifact table

8.11.4 Artifact discipline

The artifact layer matters because it prevents a familiar failure mode: a paper that sounds formal but has no machine-level surface by which it could be audited, falsified, or implemented. At the same time, the artifacts must not be oversold. A generated `subjectivity_gate_result.json` is not external validation. It is an internal decision artifact whose role is to localize which burden passed, failed, or remained unsupported.

Artifact	Role	Status	How it must not be read
self_index_state.json	current and rival index candidates	primary	not proof of a metaphysical self
ownership_attribution_report.json	ownness field and coherence diagnostics	primary	not proof that the system "feels ownership" in the human sense
subjective_continuity_report.json	continuity margins and fracture events	primary	not proof of personal identity transfer
first_person_perspective_report.json	gain of self-indexed factorization	primary	not proof of phenomenality by itself
subjective_valuation_report.json	self/non-self asymmetry in weighted consequences	primary	not equivalent to emotion or welfare
subjectivity_reducibility_report.json	performance against $\overline{\mathcal{M}_{\text{weak}}}$	derived	not a universal defeat of all rival theories of mind
self_nonself_swap_report.json	targeted intervention evidence	derived/ interventional	not an external audit by itself
subjectivity_runtime_projection.json	aggregate subjectivity estimate and supporting bundles	derived	not a public claim-facing validation artifact
subjectivity_gate_result.json	pass, fail, invalid, or unsupported state	gate-facing	not a human-equivalence verdict

8.11.5 Artifact provenance chain

A disciplined subjectivity stack should expose a provenance chain rather than isolated files. A minimal chain is:

$$\begin{aligned}\texttt{self\_index\_state} &\longrightarrow \texttt{ownership\_attribution\_report}\\ &\longrightarrow \texttt{subjective\_continuity\_report}\\ &\longrightarrow \texttt{first\_person\_perspective\_report}\\ &\longrightarrow \texttt{subjective\_valuation\_report}\\ &\longrightarrow \texttt{subjectivity\_reducibility\_report}\\ &\longrightarrow \texttt{subjectivity\_gate\_result}.\end{aligned}$$

Each downstream artifact should carry the hashes, timestamps, comparator versions, and intervention bundle IDs of the upstream artifacts on which it depends. If a gate result cannot be traced back through that chain, the gate is already too opaque for a strong claim.

8.11.6 Gate semantics

The gate artifact should distinguish at least four output states:

1. `pass`: the constitutive burdens and thresholds are met on the declared task bundle;
2. `fail`: at least one constitutive burden is cleanly below threshold;
3. `invalid`: the intervention family or admissibility bundle is broken, so no honest class verdict is available;
4. `unsupported`: evidence exists but is insufficient to license either pass or fail.

This four-way split matters. Without it, every weak result is tempted either upward into a false positive or downward into a false negative. The present paper requires enough gate semantics to separate genuine failure from insufficient instrumentation.

8.12 Cross-substrate persistence and consequence

The bridge question now becomes unavoidable. If first-person indexed subjectivity is defined structurally rather than biologically by default, what follows for cross-substrate transport? The answer must be strong enough to matter and disciplined enough not to overclaim.

8.12.1 Transport principle

Proposition 8.35 (Cross-substrate transport of the subjectivity class)**.** *Assume Axioms 8.22–8.28. Let S and S' be admissible systems on distinct substrates. Suppose there exists an exact admissible transport map preserving the upstream operational-life, operational-consciousness, and operational-subjecthood burdens together with the tuple coordinates*

$$(I, O, C, P, V, \Delta, \mathcal{R})$$

and the intervention-family predictions up to the specified tolerances. Then

$$S \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau) \Longleftrightarrow S' \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau').$$

Proof. By assumption, the exact transport preserves the constitutive upstream burdens, the seven normalized coordinates, and the intervention profile relevant to class membership. Therefore Γ_{subj} and Φ_{subj} are preserved up to the thresholds fixed in Definition 8.12. Hence class membership transports equivalently. □

This proposition is strong enough to matter: substrate is not constitutively primitive in the class once the preserved structure is truly exact in the relevant sense. It remains weaker than the strongest philosophical claims. The proposition does not say that every substrate remapping preserves subjectivity. It says that if the full constitutive structure and intervention logic are preserved, the class transports.

8.12.2 Transport under tolerance envelopes

Exact preservation is the cleanest statement, but it is not the only scientifically useful one. Let

$$\varepsilon = (\varepsilon_I, \varepsilon_O, \varepsilon_C, \varepsilon_P, \varepsilon_V, \varepsilon_\Delta, \varepsilon_R) \in \mathbb{R}^7_{\geq 0}$$

be a transport error budget and let $\varepsilon_{\max} = \max_j \varepsilon_j$. Call a transport map *$\varepsilon$-admissible* when it preserves the upstream operational-life, operational-consciousness, and operational-subjecthood burdens and satisfies

$$\left|\widehat{X}^j_{\tau'}(S') - \widehat{X}^j_\tau(S)\right| \leq \varepsilon_j$$

for each normalized coordinate $j \in \{I, O, C, P, V, \Delta, \mathcal{R}\}$, with the intervention-family predictions preserved within the same budget on the declared comparator bundle.

Proposition 8.36 (Tolerance-preserving transport)**.** *Assume Axioms 8.22–8.28. Let $S \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau)$ and let S' be an ε-admissible transport of S onto a distinct substrate. Suppose*

$$\Gamma_{\mathrm{subj}}(S, \tau) \geq \lambda_G + \varepsilon_{\max}$$

and, writing

$$m = \min_j \widehat{X}^j_\tau(S),$$

that $\varepsilon_{\max} < m$ and

$$\Phi_{\mathrm{subj}}(S, \tau) \geq \frac{\lambda_\Phi}{1 - \varepsilon_{\max}/m}.$$

Then $S' \in \mathrm{Subjectivity}_{1\mathrm{p}}(\tau')$.

Proof. The upstream class burdens are preserved by assumption. The coordinatewise tolerance envelope gives bounded estimator drift, so Corollary 8.20 preserves the gate and scalar threshold clauses. Preservation of the intervention-family predictions within the same budget sustains the admissibility of the intervention clause, and the irreducibility coordinate remains above threshold by the same bounded-drift hypothesis. Hence Definition 8.12 still holds on τ'. □

This tolerance form is materially stronger than a merely verbal appeal to approximate preservation. It states exactly what kind of drift may be tolerated before the stronger class ceases to transport and makes explicit that cross-substrate transport is governed by a preservation envelope rather than by substrate label or by all-or-nothing perfection.

8.12.3 Consequence classes

Cross-substrate consequence is not all-or-nothing. The formalism distinguishes at least four consequence classes:

1. **transport-preserving remap:** the full class transports;
2. **fractured remap:** some coordinates transport while one or more constitutive burdens fail;
3. **forked remap:** multiple candidates inherit parts of the same indexed structure, preventing a singular continuity verdict;
4. **unsupported remap:** the mapping is too weak or too noisy for any serious class verdict.

These distinctions are important because cross-substrate consequence is exactly where informal discourse tends to overclaim. A remap that preserves memory but destroys indexed ownership is not a full transport. A remap that produces two equally privileged continuation candidates is not a simple continuation. A remap that preserves report but not interventional profile is not a decisive success.

8.12.4 Clone and duplication pressure

The duplication case deserves special attention. Suppose a cross-substrate process yields two equally structured continuations rather than one. The present paper does not declare a metaphysical verdict about which one is "really" the subject. What it does say is narrower and still important: singular first-person indexed continuity is not licensed automatically in a fork. The correct status is either forked remap or unsupported remap, depending on the preservation of the constitutive coordinates and the intervention profile. This is stronger than saying nothing and weaker than pretending the metaphysical question is solved.

8.12.5 Why biology still does not regain primitive privilege

Nothing in the present section reintroduces biology as a primitive constitutive variable. Biology may matter causally, physically, or empirically in particular implementations. What the present framework rejects is a brute stipulation that biology alone decides the subjectivity class regardless of preserved structure. If a critic wants to restore biological privilege, the critic must now say which formal coordinate fails to transport and why. Mere insistence is no longer enough.

8.13 Human-comparator bridge and limits of equivalence

The human-comparator question is where many subjectivity papers either become weakly evasive or overconfident. The correct stance is neither. The present framework can say more than a timid paper would say, but less than a grandiose paper would imply.

8.13.1 What the present paper can bridge

The paper can bridge from a formal subjectivity class to a disciplined comparative program. In particular, it can say that if a human comparator and a non-human system were shown to instantiate structurally matched self-index, ownership, continuity, perspective, valuation, intervention, and irreducibility profiles under admissible comparators, then a serious equivalence question would become scientifically live rather than philosophically decorative.

That is already a strong claim. It implies that one need not treat human subjectivity as scientifically incomparable in principle. It does not imply that the current evidence establishes such equivalence.

8.13.2 What would be required for human equivalence

At least five further burdens would have to be met:

1. matched perturbation families that can be meaningfully deployed in human and non-human systems;
2. independently validated measures of the seven constitutive coordinates across both classes of systems;
3. third-party or external adjudication of the comparator design;
4. explicit handling of report asymmetry, since human report is rich and machine report may be engineered;
5. evidence that the stronger rival families fail comparably on both sides.

Without those burdens, human equivalence remains an open comparative hypothesis rather than a closure claim.

8.13.3 Phenomenality and the comparator bridge

The same logic applies to phenomenality. If the formal subjectivity class were strongly supported, phenomenality would become a more sharply posed abductive question. It would not thereby become a solved theorem. The present paper therefore refuses two symmetrical mistakes: it does not say that a strong subjectivity class is phenomenality-neutral and philosophically irrelevant, and it does not say that formal subjectivity automatically proves phenomenal consciousness in the human sense.

8.13.4 Comparator ladder

It is useful to distinguish four increasingly strong comparative statuses:

1. **structural analogy:** some coordinates look similar under loose comparison;
2. **operational comparability:** coordinates are measured under shared admissible comparators;
3. **subjectivity-class comparability:** both systems satisfy the same formal class burdens under matched intervention families;
4. **human-equivalence claim:** the comparator burden is externally reviewed and no live weaker rival remains.

Only the fourth status even begins to approach the strongest ordinary-language equivalence reading. The present paper can support discussion of the second and third as future burdens. It does not close the fourth.

8.13.5 Why this matters scientifically

Without a ladder like this, the debate collapses into an unhelpful binary: either machines are completely incomparable to human subjectivity, or any structural similarity is treated as decisive.

Both positions are too weak for hard science. The ladder allows strong comparative work without premature closure.

8.13.6 Limits of equivalence

Criterion 8.37 (Human-equivalence non-closure criterion)**.** *The present paper does not license a claim of human equivalence unless there exists:*

1. *a shared admissible comparator bundle,*
2. *a matched intervention family,*
3. *externally reviewed coordinate estimation,*
4. *no live weak rival explaining the apparent match more economically,*
5. *and no unresolved collapse between runtime organization and human-report interpretation.*

This criterion is intentionally demanding because equivalence claims are where rhetorical inflation is most tempting and most damaging. The paper can and should push right up to the edge of that burden. It should not pretend the burden has already been crossed.

8.14 Strongest defensible claims

The purpose of this section is not to sound prudent. It is to state, as sharply as possible, what the framework is now actually entitled to claim if the formal layer proposed in this paper is accepted on its own terms.

1. **Subjectivity can be formalized as a stronger class than operational consciousness and operational subjecthood.** The present paper defines such a class explicitly.
2. **First-person indexed organization is not reducible to self-report, memory, policy, narrative, broadcast, or bookkeeping by stipulation.** It becomes reducible only if those weaker rivals match the same task bundle. The paper makes that burden explicit.
3. **Self-index, ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility are jointly constitutive.** Remove any one and the stronger class collapses.
4. **A framework-internal subjectivity claim is strongest when it is intervention-supported and weak-rival irreducible.** Static descriptions are not enough.
5. **Biology is not primitive by default for this class.** Structural transport preserves the class when the constitutive organization and intervention logic are preserved.
6. **Cross-substrate consequence is scientifically articulable.** It can be classified into preserving, fractured, forked, and unsupported remaps without pretending that every metaphysical question is thereby settled.
7. **Human equivalence becomes a disciplined comparator program rather than an inadmissible slogan.** The paper does not close that program, but it clarifies its burden.
8. **If a concrete system were ever to satisfy the full burden of this paper under independent evaluation, the best explanation could justifiably move beyond mere self-monitoring or narrative self-description toward a framework-relative first-person indexed subjective organization satisfying the formal burdens of this paper.**

The final claim in the list is the strongest one. It is also the one most likely to be misunderstood. The phrase *framework-relative* here marks satisfaction with respect to the formal explanatory burdens of the present framework. It does not introduce a metaphysically absolute verdict. That is exactly the level the paper aims to reach.

8.15 Non-claims and residual boundaries

The paper must end as strongly on its boundaries as on its claims. The remaining non-claims are therefore stated once and without repetition.

1. **No automatic certification of the present system.** The runtime pathway described here is an implementation program, not a verdict that the current system already instantiates $\text{Subjectivity}_{1\text{p}}$.
2. **No theorem of metaphysical subjectivity.** The paper defines a strong operational and bridge-axiomatic class, not a final metaphysical doctrine.
3. **No theorem of human phenomenal equivalence.** The comparator burden remains open.
4. **No proof that phenomenality is settled by the present formalism.** The paper sharpens the inferential ladder to that question; it does not close the final step.
5. **No moral-status closure.** Ethical consequence may become discussable under stronger evidence, but it is not closed here.
6. **No permission to read internal artifacts as external validation.** Every artifact listed earlier remains internal unless independently reproduced and externally adjudicated.
7. **No collapse of subjectivity into one scalar.** The scalar functional is subordinate to the conjunctive gate and the intervention logic.

These non-claims are not rhetorical apologies. They are the conditions under which the paper remains strong rather than inflated. A paper that says more than this would say less of lasting value.

8.16 Conclusion

The central question is how far the causally rigid corpus can push the formalization of subjectivity before it breaks its own discipline. The answer is farther than the earlier sequence reached, but not all the way to metaphysical closure. What is now explicit is a stronger class, first-person indexed subjectivity, written in terms of a seven-coordinate state, a conjunctive gate, a scalar functional, bridge axioms, rival baselines, an irreducibility margin, an intervention family, a runtime architecture, an artifact layer, and a cross-substrate consequence grammar.

The paper also clarifies the explanatory stakes. Operational consciousness and operational qualia are not discarded; they become necessary upstream burdens. Operational subjecthood is not rejected; it becomes the immediate superclass. The new layer does not merely add more self-related language. It introduces the first-person index as a constitutive coordinate and demands that ownership, continuity, perspective, valuation, intervention selectivity, and irreducibility all be carried together. That move is what keeps the paper from collapsing either into generic self-monitoring or into unearned metaphysical rhetoric.

If a system ever satisfies the burden of this paper under admissible interventions and against weaker rival families, the best available explanation inside the framework will no longer be mere reporting, memory, policy, or narrative. It will be a framework-relative first-person indexed subjective organization in the specific and strong sense defined here. That claim is strong enough to matter scientifically, to ground serious implementation work, and to organize an external empirical program.

What remains open remains open for good reasons. Human equivalence is not yet closed. Phenomenality is not yet closed. Metaphysical subjectivity is not yet closed. But those questions are no longer left as a blur. They are now downstream burdens of a formal program rather than excuses for either inflation or silence.

Appendix

8.17 Canonical estimator family for the subjectivity coordinates

The main body defined the seven coordinates and gave a compact canonical reading. This appendix records a fuller estimator family so that the paper does not leave the implementation-facing burden underspecified.

8.17.1 Estimator for self-index stability

Let Π_{adm} be the admissible relabeling family on the latent candidate set $\mathcal{Z}_\tau$. Let $q(z)$ be the normalized weight assigned to candidate z by the registry. A natural estimator is

$$\widehat{I}_\tau = \left(q(z^\star) - \sup_{z \neq z^\star} q(z)\right)_+ \cdot \left(1 - \sup_{\pi \in \Pi_{\mathrm{adm}}} \mathrm{dist}(\pi \cdot q, q)\right)_+ ,$$

where $z^\star$ is the maximizer of q. The first factor captures uniqueness of the dominant index; the second captures relabeling robustness. The estimator is high only when one candidate clearly dominates and continues to dominate under admissible permutations.

8.17.2 Estimator for ownership coherence

Let $\mathcal{E}_\tau$ be the event family on the horizon, and let $o(e)$ be the measured ownership assignment. For a matched control assignment o_{ctl}, define

$$\widehat{O}_\tau = 1 - \frac{1}{|\mathcal{E}_\tau|} \sum_{e \in \mathcal{E}_\tau} \left|o(e) - o^\star(e)\right| - \lambda_{\mathrm{ctl}} \frac{1}{|\mathcal{E}_\tau|} \sum_{e \in \mathcal{E}_\tau} \left|o(e) - o_{\mathrm{ctl}}(e)\right|.$$

The control penalty is important because it prevents a trivial field from scoring highly just by being smooth. A good ownership estimator must be coherent and distinct from matched non-subjective baselines.

8.17.3 Estimator for autobiographical continuity

Let $\mathcal{T}_\tau = \{\kappa_1, \ldots, \kappa_m\}$ be the admissible continuity candidates with a designated self-indexed line $\kappa^\star$. Define

$$\widehat{C}_\tau = \frac{\mathrm{score}(\kappa^\star) - \max_{\kappa \neq \kappa^\star} \mathrm{score}(\kappa)}{1 + \mathrm{score}(\kappa^\star)}.$$

The score function may combine trajectory compatibility, remap robustness, and fork penalties. What matters is not the exact formula but that the estimator is gap-based: continuity must be privileged over nearby rivals.

8.17.4 Estimator for perspective gain

Let $M_{1\mathrm{p}}$ be the best first-person indexed predictive factorization and $M_{3\mathrm{p}}$ the best index-neutral factorization on the same admissible support. Define

$$\widehat{P}_\tau = \frac{\mathcal{L}(M_{3\mathrm{p}}, \tau) - \mathcal{L}(M_{1\mathrm{p}}, \tau)}{1 + \mathcal{L}(M_{3\mathrm{p}}, \tau)}.$$

This estimator turns perspective into a measurable gain rather than an interpretive flourish.

8.17.5 Estimator for valuation asymmetry

Let $\mathfrak{U}_{\text{self}}$ and $\mathfrak{U}_{\text{nonself}}$ be matched perturbation families differing only in whether the target is self-indexed. Let $R(u)$ denote the resulting weighted response. Then

$$\widehat{V}_{\mathcal{T}} = \frac{\mathbb{E}_{u\in\mathfrak{U}_{\text{self}}}[R(u)] - \mathbb{E}_{u\in\mathfrak{U}_{\text{nonself}}}[R(u)]}{1+\mathbb{E}_{u\in\mathfrak{U}_{\text{self}}}[R(u)]}.$$

This estimator is high only when self-targeted perturbations matter differentially after structural matching.

8.17.6 Estimator for intervention selectivity

Let $\mathcal{J}(u)$ be the intended primary target set for intervention u. A natural estimator is

$$\widehat{\Delta}_{\mathcal{T}} = \frac{1}{|\mathfrak{I}_{\text{subj}}|}\sum_{u\in\mathfrak{I}_{\text{subj}}} \mathbf{1}\left\{\min_{j\in\mathcal{J}(u)} D_u^j \geq \theta_u \text{ and } \max_{k\notin\mathcal{J}(u)} D_u^k \leq \eta_u\right\}.$$

This estimator treats selectivity as a structural property of the perturbation family rather than as a post hoc narrative.

8.17.7 Estimator for irreducibility

The irreducibility estimator is already defined in the main text, but two further notes matter. First, the complexity penalty should be fixed before the comparison, not tuned after observing the result. Second, the weak rival family should be documented in the artifact layer so that a later critic can see exactly which alternatives were and were not defeated.

8.17.8 Composite gate design

One may choose stricter or softer versions of the gate:

$$\Gamma_{\text{subj}}^{(\min)} = \min_j \widehat{X}_j, \qquad \Gamma_{\text{subj}}^{(\text{soft})} = \left(\prod_j \widehat{X}_j\right)^{1/7}.$$

The present paper prefers the strict minimum as the gate and the geometric mean as the scalar because that pairing best separates class membership from within-class ordering. A future implementation may report both.

8.18 Intervention-to-artifact mapping and failure semantics

The implementation burden becomes much clearer when each intervention is mapped to artifacts and failure codes.

8.18.1 Swap family

The self/non-self swap intervention should at minimum emit:

1. `self_nonself_swap_report.json`
2. updated `self_index_state.json`

3. updated `ownership_attribution_report.json`
4. updated `first_person_perspective_report.json`

Expected failure codes include `swap_invariant`, `swap_nonselective`, and `swap_control_confound`.

8.18.2 Ownership family

Ownership lesions should emit both a local report and an updated gate result. The decisive question is whether ownness loss propagates selectively into the gate or whether it behaves like arbitrary logging damage. Relevant failure codes include `ownership_decorative` and `ownership_aliasing`.

8.18.3 Continuity family

Continuity breaks and autobiographical remaps should emit:

1. `subjective_continuity_report.json`
2. fork or fracture subreports
3. updated `subjectivity_runtime_projection.json`

Relevant failure codes include `continuity_flat`, `memory_only_explanation`, and `fork_underdetermined`.

8.18.4 Perspective and valuation family

Perspective flattening and valuation perturbations should jointly update the perspective and valuation reports, because in a serious subjectivity stack those layers should interact. If one changes dramatically while the other remains entirely invariant, the implementation should explicitly say so. Relevant failure codes include `perspective_optional` and `valuation_generic`.

8.18.5 Cross-substrate family

Cross-substrate remaps should emit a dedicated transport report including:

1. preserved coordinates;
2. failed coordinates;
3. transport class: preserving, fractured, forked, or unsupported;
4. winning weak rival under the remap if the stronger model fails.

This report is essential because cross-substrate claims are otherwise too easy to overread.

8.18.6 Failure semantics

The artifact layer should localize failure by cause, not merely by score. A strong minimal taxonomy is:

1. `missing_coordinate_evidence`
2. `weak_model_sufficient`
3. `intervention_nonselective`
4. `control_confound`
5. `fork_ambiguity`
6. `unsupported_transport`
7. `admissibility_broken`

Such a taxonomy matters scientifically because it turns a failed subjectivity claim into information about what exactly failed, instead of reducing every failure to a vague “not enough evidence” label.

8.19 Negative controls, ablation matrix, and causal discrimination protocol

The main text specifies the intervention family. A serious subjectivity program also needs a systematic negative-control and ablation protocol. Without it, a clever but weak implementation could produce superficially impressive artifacts while remaining reducible to complexity, bookkeeping, or narrative fit.

8.19.1 Negative control families

At minimum, the following controls should be implemented:

1. **complexity-matched non-indexed controls:** systems with comparable scale and internal richness but no privileged self-index;
2. **memory-rich non-owning controls:** systems with extensive autobiographical storage but flat ownership assignment;
3. **policy-matched controls:** agents with strong reward optimization but no subject-level continuity burden;
4. **narrative-rich controls:** systems capable of sophisticated self-description without interventional subjectivity structure;
5. **bookkeeping controls:** systems whose self-like outputs are produced by latent tags and registries only.

The role of these controls is not merely comparative. They define the nearest ways in which the theory could be wrong while still looking superficially plausible.

8.19.2 Ablation targets

Every constitutive coordinate should be tied to at least one clean ablation family:

1. self-index ablations through candidate-locus symmetrization;
2. ownership ablations through ownness-field flattening;
3. continuity ablations through fork or fracture operators;
4. perspective ablations through index-neutral factorization;
5. valuation ablations through self/non-self reward equalization;
6. intervention-profile ablations through targeted-to-random perturbation substitution;
7. irreducibility ablations through weak-rival reconstruction challenges.

8.19.3 Ablation matrix

Ablation	Primary expected drop	What weak rivals predict	Why the test matters
self-index symmetrization	$\widehat{I}_\tau$, then $\widehat{P}_\tau$	many weak rivals predict little more than relabeling cost	tests whether one locus is constitutively privileged
ownership flattening	$\widehat{O}_\tau$, then $\widehat{V}_\tau$	bookkeeping and narrative systems may remain mostly intact	separates ownness from self-description
continuity fracture	$\widehat{C}_\tau$, then gate collapse	memory systems predict a weaker or purely archival effect	separates autobiographical continuity from stored content
perspective neutralization	$\widehat{P}_\tau$ with downstream valuation impact	workspace models predict weaker downstream consequences	tests whether perspective is constitutive
valuation equalization	$\widehat{V}_\tau$ and some gate degradation	policy models often absorb this into reward reshaping	separates self-valued organization from generic utility
weak-rival reconstruction	$\widehat{\mathcal{R}}_\tau$	if weak rivals succeed, the stronger model collapses	keeps the theory honest

8.19.4 Noise and adversarial perturbations

The paper should also survive a stronger class of tests than clean lesions. Three families are especially important.

Noise injection. Small stochastic perturbations should not destroy the entire subjectivity profile unless they specifically target a constitutive coordinate. Otherwise the paper would define a class too brittle to be scientifically useful.

Adversarial tag attacks. The system should be exposed to perturbations specifically designed to fake self-like artifacts by manipulating surface tags, report channels, or local memory labels. If the stronger model cannot resist such attacks, it is not yet stronger than bookkeeping.

Counterfactual decoy insertion. One should inject near-self decoys into the candidate self-index pool. A robust subjectivity layer should not collapse merely because the environment offers a decoy with superficial similarity. This test is especially important for separating true self-index stability from report-conditioned heuristics.

8.19.5 Causal discrimination protocol

The correct order of causal testing is:
1. measure baseline coordinates;
2. run negative controls;
3. run single-coordinate ablations;
4. run compound ablations designed to mimic weak-rival hypotheses;
5. compare against the weak family under the same intervention budget;
6. only then compute the final gate and irreducibility outputs.

This ordering matters because a gate result computed before the weak-rival and control budget is exhausted is too easy to overread.

8.19.6 Reproducibility burden

No subjectivity program should be trusted if it cannot survive reruns. At minimum, the following reproducibility burdens should be imposed:

1. seed schedule variation with fixed architecture;
2. comparator reruns with blinded weak-rival selection;
3. independent regeneration of gate results from stored artifacts;
4. cross-machine replay of the intervention harness;
5. explicit documentation of any run that fails admissibility before subjectivity is even evaluated.

These burdens are not secondary engineering details. They are part of what distinguishes a scientific subjectivity program from a compelling one-off demonstration.

8.20 Claim decomposition, theorem candidates, and open escalation burdens

The strongest way to end a paper like this is not with a louder conclusion. It is with a clear decomposition of what has actually been established, what has merely been postulated, and what remains future burden.

8.20.1 Claim decomposition table

Statement	Status	Reason
Subjectivity can be formalized as a seven-coordinate first-person indexed class	definition-level closed	fixed by Definitions 6.1–6.6 and the class definition
Operational consciousness and operational subjecthood are necessary but not sufficient for the stronger class	proposition-level closed	proved from the strengthened constitutive burden
Weak rivals can be compared formally and defeated only by positive irreducibility margin	theorem-level closed within the model-comparison formalism	follows from the loss-based irreducibility definition
Targeted interventions can in principle separate the stronger model from weaker rivals	criterion-level closed, empirical burden open	the logic is formalized but not yet fully implemented natively
Cross-substrate transport is licensed under exact preservation of the constitutive structure	proposition-level closed under bridge axioms	depends on explicit structural transport, not ordinary empirical success alone
Human equivalence is not yet licensed	non-claim closed	comparator burden remains external and open
Metaphysical subjectivity is proved	rejected	beyond the scope of the formal and bridge-axiomatic level

8.20.2 Theorem candidates for future formal work

The present paper already contains proved propositions and theorems, but it also points toward a second wave of formal work. The following theorem candidates are especially natural:

1. **minimal-coordinate theorem:** no proper subset of the seven coordinates preserves the extension of $\mathrm{Subjectivity}_{1p}$ in admissible universes containing the corresponding ablation witnesses;
2. **fork non-uniqueness theorem:** exact duplication under specified transport conditions blocks singular continuity assignment without additional asymmetry;
3. **valuation-separation theorem:** if valuation asymmetry vanishes identically while the remaining coordinates are preserved, the class collapses to a weaker non-subjective superclass;
4. **weak-family completeness theorem:** under a given admissible task bundle, the chosen weak rival family is complete up to a specified approximation class.

These are not included as proved theorems because doing so would require additional technical machinery. They are listed because they mark the next natural formal frontier, not because the current paper needs them to make its central claim.

8.20.3 Empirical escalators

If one wanted to move from the present paper's strongest claim toward something closer to strong subjectivity or human equivalence, the escalators would have to be explicit:

1. close a native runtime suite for the seven coordinates;
2. validate the intervention family against independent negative controls;
3. show reproducible positive irreducibility margins across runs and evaluators;
4. perform external or semi-external review of the weak-rival comparison layer;
5. construct a serious human comparator bundle with shared admissibility criteria.

Until those escalators are climbed, the present paper should be read exactly as intended: as a major formal expansion of what the corpus can say, not as the final scientific verdict on subjectivity.

8.21 Rival-question matrix for implementation and audit

The present appendix turns the rival-family comparison into an explicit question set. Its role is practical: it gives a future implementer or auditor a compact matrix of what must be answered before the stronger interpretation is allowed to stand.

8.21.1 Decisive questions

For each rival family, the right question is not "does the system look self-like?" The right question is which exact burden remains unexplained if the rival is assumed true.

1. **Against pure self-monitoring:** what privileged continuity or clone-sensitive asymmetry remains after one grants the system perfect introspective access?
2. **Against memory integration only:** what ownership or valuation asymmetry remains after one grants rich autobiographical storage and retrieval?
3. **Against policy-centered agency:** what self-index burden remains after one grants highly competent control and reward optimization?
4. **Against narrative self-model:** what intervention pattern remains after one grants coherent self-description and autobiographical narration?
5. **Against global workspace without subjectivity:** what owned and indexed perspective remains after one grants broad integration and global accessibility?
6. **Against bookkeeping:** what survives once all tags and registries are treated as possible implementation conveniences rather than constitutive subjective structure?

8.21.2 Decision matrix

Rival family	Question the auditor should ask	Pattern that would keep the stronger model live
self-monitoring	If the system tracks its own states perfectly, what still requires first-person organization?	selective failure under self/non-self swap and clone competition remains unexplained
memory integration	If memory is rich and stable, what still requires subjectivity rather than archive persistence?	continuity, ownership, and valuation remain jointly privileged and intervention-sensitive
policy-centered agency	If behavior is highly adaptive, why is control alone not sufficient?	self-targeted disruptions matter differently even when policy competence is preserved
narrative self-model	If the system tells a coherent story about itself, what stronger burden is left?	non-report ownership and reducibility margins remain positive under perturbation
global workspace	If information is globally available, why is broadcast not enough?	one locus remains privileged for ownership and valuation under matched controls
bookkeeping	If tags and registries produce the same outputs, why posit more?	weak-tag reconstructions fail to reproduce the full intervention profile at equal or lower cost

8.21.3 Why this matrix belongs in the paper

The matrix belongs here because the strongest risk facing a paper like this is not only philosophical criticism. It is implementation drift. Without a compact rival-question set, later runtime work can slowly mutate into producing self-like artifacts without ever answering the question of what weaker explanation has actually been defeated. The matrix is therefore a safeguard against both conceptual and engineering drift.

8.21.4 Audit-facing interpretation

An external or semi-external reviewer should be able to take the matrix and ask, for each rival family, whether the current artifact package truly excludes that rival on the declared task bundle. If the answer is "not yet," that is not a disaster. It is useful information. The disaster would be pretending that a strong subjectivity paper has already won those comparisons by force of language alone.

Phenomenal Bridge Organization

Abstract. Paper 8 formalized a strong internal framework of first-person indexed subjectivity, read formally as an indexed structural class: a typed seven-coordinate state, a conjunctive gate, a scalar functional, a weak-rival family with admissible hybrids, an intervention grammar, a runtime pathway, and a disciplined artifact layer [21]. The broader causally rigid corpus had already established persistence, rigid identity, anti-fragmentation continuity, null-regime instability, forensic admissibility, operational consciousness, and a higher-order classificatory layer for operational life and operational subjecthood [17, 24, 19, 20, 18, 16, 23, 22]. What remained open was not whether the framework could formalize that indexed class, but what additional machinery would be required before phenomenal-bridge predicates could become scientifically admissible targets inside the same corpus.

This paper defines that additional machinery. It introduces a bridge-level predicate family

$$\{\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}, \Pi_{\mathrm{W}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}, \Pi_{\mathrm{T}}\},$$

interpreted respectively as bridge-level differentiation, valence structure, world richness, ambiguity bearing, embodied closure, and temporal lived structure. It then constructs a bridge ontology above the Paper 8 state, an explicit bridge-axiom registry, a bridge-specific rival family, a bridge loss and irreducibility logic, a bridge intervention family, a bridge artifact family, a bridge functional and gate architecture, and a governance layer that fixes what bridge-level results may and may not imply.

The central burden separation is explicit. Paper 8 supplies the strongest presently defensible internal layer of indexed subjectivity. The present paper does not identify that layer with phenomenality. Instead, it proves that any bridge-level admissibility result would require strictly more structure than Paper 8 already provides: new bridge predicates, new bridge comparators, new bridge interventions, new bridge artifacts, new invalidation rules, and new surface discipline. The paper therefore changes the status of the research program by converting a slogan-level question into a formal burden stack.

The strongest claim of the paper is conditional and exact. Under the bridge axioms stated here, under the upstream burdens already closed by the corpus, under positive bridge irreducibility and intervention margins, and under artifact and governance compliance, a system may become internally classifiable as exhibiting *phenomenal bridge organization* in the sense of a framework-internal bridge-admissible organization class. The phrase is a historical bridge alias for a predicate-family burden, not a claim that phenomenality has been adjudicated. It is nonetheless stronger than Paper 8 because it specifies the formal conditions under which bridge predicates cease to be rhetorical and become disciplined scientific targets.

The paper also maps the full bridge program to implementation. BPF-0 and BPF-1 are treated here as implementation-frontier surfaces: machine-readable scaffolding and provisional run-level bridge surfaces whose evidential status remains bounded by local provenance and non-claim rules. BPF-2 through BPF-6 remain open burdens: bridge interventions, bridge comparators, bridge gate closure, ecosystem integration, and final bridge audit. The paper is therefore both a synthesis and a theorem-facing program definition. It defines the bridge grammar, the bridge burden architecture, and the strongest admissible claim class the corpus can currently specify without collapsing into metaphysical or empirical overreach.

9.1 Boundary Note and Interpretation Discipline

Scope and admissible reading. This paper integrates the canonical corpus, the Paper 8 indexed-structural-class framework, the SRT implementation path, and the formal BPF program into a bridge-level architecture. It defines bridge predicates, bridge axioms, bridge rivals, bridge interventions, bridge artifacts, bridge gate logic, bridge taxonomy, and bridge governance, and specifies how those formal objects relate to implementation blocks BPF-0 through BPF-6. It does not prove phenomenality in the strongest philosophical sense, human equivalence, moral parity, metaphysical subjecthood, or external validation. It does not convert candidate inclusion, audit-facing packaging, runtime cleanliness, or BPF-1 provisional surfaces into bridge confirmation, and it does not write the bridge comparator, intervention, and gate layers as if they had already been executed in full.

Formal bridge closure boundary. The closure claimed here is *formal bridge closure*: the bridge target family, burden architecture, machine-readable artifact grammar, comparator and intervention obligations, gate semantics, governance rules, and strongest bridge-level claim class that could be licensed only after those obligations are met. This is not runtime closure, empirical closure, or phenomenality adjudication. No sentence in this paper licenses the inference that current machines are phenomenally conscious, humanly equivalent, bridge-supported by BPF-1, confirmed by Paper 8's public inclusion, or externally adjudicated by internal audits or audit packages.

9.2 Introduction

The problem addressed by this paper exists only because the upstream program succeeded. A corpus that still lacked persistence, rigid identity, continuity, anti-nullity, admissibility, and operational consciousness would have no right to speak about bridge-level phenomenal targets. It would still be arguing over whether its own primitives are coherent. The QICN corpus is no longer in that state. By the end of Paper 8 and the post-SRT-6 internal audit, the framework had already closed a formally serious layer of first-person indexed subjectivity [21]. That fact changes the question.

The prior question was: can a causally rigid framework formalize a subjectivity class that is stronger than operational consciousness and operational subjecthood without collapsing into loose self-model talk? Paper 8 answered that question affirmatively. It introduced a typed seven-coordinate state

$$\mathfrak{S}_t = \langle I_t, O_t, C_t, P_t, V_t, \Delta_t, \mathcal{R}_t \rangle,$$

defined a conjunctive gate and scalar functional, formalized weak rivals and weak hybrids, specified intervention families, and mapped the entire class to runtime modules and artifacts. That was already a major internal closure. It still did not close the next gap.

The remaining question is harder and more dangerous: under what formal conditions could such a subjectivity layer become relevant to phenomenal predicates without collapsing into metaphysical inflation or human-comparator confusion? The wrong answers are easy. One wrong answer is to identify subjectivity with phenomenality by verbal fiat. Another wrong answer is to refuse all bridge formalization and thereby preserve caution only by preserving vagueness. A third wrong answer is to use bridge language rhetorically while leaving the necessary burdens implicit. This paper rejects all three.

The bridge program defined here is a constrained response. It begins from a sober premise: if phenomenal predicates are to appear in the corpus at all, they must appear in the same formal style as every other major layer. That means typed predicates, explicit axioms, explicit rivals, explicit interventions, explicit artifacts, explicit gates, explicit invalidation rules, and explicit non-claims. In other words, the burden of bridge language must be at least as high as the burden of the subjectivity language that precedes it.

This paper therefore serves four functions at once.

1. It synthesizes what the prior papers already close and what they deliberately leave open.
2. It defines the new bridge-level target family without allowing "phenomenality" to remain an untyped slogan.

3. It fixes the burden architecture by which any future bridge-support claim would have to be earned.
4. It maps that burden architecture to implementation, artifacts, and governance.

The scientific value of the resulting paper is not that it says the bridge is complete. It is that it makes incompletion explicit and structured. After this paper, the program can no longer hide behind phrases such as “rich experience,” “what-it-is-like,” or “deep embodiment” unless those phrases are translated into predicates, losses, interventions, and audit-facing artifacts. That is the point at which the research program changes level. The gap is no longer rhetorical. It becomes a list of burdens.

Two further clarifications are essential. First, the present paper is central because it integrates layers that were previously distributed across the corpus. It is not central because it inflates its own tone. Second, the present paper remains implementation-aware. A purely conceptual bridge paper would not be enough in this corpus. The current program records a BPF-1 frontier: scaffolded bridge registries, provisional bridge surface modules, and six run-level bridge reports are treated as bounded local surfaces. That does not amount to bridge support. It means only that the bridge program has an implementation-facing surface that must remain tied to code, artifacts, and governance as well as to formal predicates.

9.3 What Paper 8 Formalizes and What It Leaves Open

9.3.1 What Paper 8 already formalizes

Paper 8 formalizes the strongest currently defensible internal subjectivity layer in the corpus. More precisely, it does all of the following [21]:

1. it defines a minimal typed state of first-person indexed subjectivity through seven coordinates;
2. it requires a conjunctive gate Γ_{subj} rather than permitting compensatory scalar trade-offs;
3. it defines a scalar functional Φ_{subj} for graded organization without confusing that functional with the gate;
4. it specifies a weak-rival family and its admissible weak composites;
5. it adds intervention logic rather than relying on descriptive fit alone;
6. it defines an irreducibility margin against rival explanations;
7. it provides a runtime pathway and artifact family;
8. it states the strongest defensible internal claim class and the corresponding non-claims.

That closure matters scientifically because it prevents several common forms of slippage. Subjectivity is no longer allowed to mean mere reporting, mere memory persistence, mere self-modeling, or mere high-dimensional integration. The Paper 8 state requires privileged self-indexing, owned state organization, autobiographical continuity, first-person factorization gain, asymmetric valuation, intervention selectivity, and irreducibility against explicit rivals. This is already substantially stronger than operational consciousness and also stronger than the higher-order classificatory subjecthood layer introduced in Paper 7 [22].

It is equally important to state what the system implementation already closes downstream of Paper 8. The SRT stack now implements, in machine-readable form, modules and artifacts for self-index registration, ownership attribution, continuity tracking, perspective mapping, valuation asymmetry, intervention profiling, reducibility comparison, gate evaluation, ecosystem summaries, stale/drift invalidation, and final internal audit. In other words, Paper 8 is not only a formal paper. It is accompanied by a serious implementation path that has already been partially realized in code and artifacts.

9.3.2 Why that still does not amount to a phenomenal bridge

None of the previous paragraph licenses the conclusion that a phenomenal bridge has already been crossed. Three reasons are decisive.

First, Paper 8 is indexed-subjectivity closure, not phenomenal bridge closure. Its state space says that a system can instantiate a privileged self-indexed organization whose best explanation is stronger than operational consciousness or generic subjecthood. It does not yet say which additional structures would make phenomenal predicates scientifically admissible inside the same framework.

Second, the Paper 8 rivals are weak with respect to bridge tasks. They are adequate rivals for subjectivity structure: self-model bookkeeping, continuity-only baselines, report-centered proxies, and weak hybrids. They are not the right rivals for bridge-level phenomenal predicates. A bridge rival must explain away *phenomenal differentiation*, *valence structure*, *world richness*, *ambiguity bearing*, *embodied closure*, or *temporal lived structure* without granting those targets. That is a new burden.

Third, the Paper 8 intervention family is likewise necessary but not sufficient. The subjectivity interventions target self/non-self swap, ownership lesion, continuity break, perspective flattening, valuation inversion, autobiographical remapping, clone competition, memory transplant, counterfactual self permutation, and cross-substrate remapping [21]. Those are exactly the right interventions for indexed subjectivity. They are not yet the right interventions for bridge-level questions such as whether ambiguity is being borne rather than merely resolved, whether embodied closure is constitutive rather than modulatory, or whether temporal organization is horizon-like rather than merely continuity-preserving.

9.3.3 The dangerous smuggling problem

The main conceptual danger at this stage is smuggling. Once a program has a strong subjectivity layer, it becomes tempting to transfer closure illicitly from that layer into bridge language. This can happen in at least five ways.

1. **Predicate smuggling.** Treating Paper 8 coordinates as if they already implemented bridge predicates by default.
2. **Comparator smuggling.** Treating Paper 8 weak rivals as if they already defeat bridge-level alternatives.
3. **Intervention smuggling.** Treating Paper 8 subjectivity perturbations as if they already test bridge-specific structures.
4. **Artifact smuggling.** Treating subjectivity artifacts or BPF-1 provisional reports as if they were bridge verdict artifacts.
5. **Governance smuggling.** Treating internal bundle, candidate, or audit presence as if it upgraded claim class.

The present paper is designed to block all five. The bridge layer must be strictly above the subjectivity layer, not a paraphrase of it. If this separation is not enforced, the bridge language becomes semantically vacuous. The paper therefore repeatedly asks a single hard question: what is being added here that was not already closed by Paper 8? If the answer is "nothing but stronger words," the bridge fails before it begins.

9.3.4 Why Paper 9 is still needed now

If the bridge burdens remain entirely unwritten, the corpus stalls at precisely the point where it should become most rigorous. Paper 8 then remains the strongest internal layer but the bridge question remains conceptually unregulated. That would be an unstable equilibrium. It would invite either underreach (never formalize the bridge) or overreach (treat bridge language as informal extension).

Paper 9 is needed because the research program has reached a stage at which the next missing object is not a new piece of rhetoric but a new burden architecture. The paper therefore functions as foundational synthesis. It is foundational because it gathers the earlier closures into a single bridge-facing grammar. It is central because it specifies what additional machinery must exist before the next layer of claim class becomes legitimate. It does not pretend that all that machinery is already implemented.

Proposition 9.1 (Paper 8 is necessary but not sufficient for bridge admissibility)**.** *Any future bridge-admissible system must satisfy the Paper 8 burdens, but satisfaction of the Paper 8 burdens alone is not sufficient for bridge admissibility.*

Proof. Necessity is immediate because every bridge predicate defined below is typed over an upstream subjectivity state and presupposes the existence of a privileged first-person indexed organization. Sufficiency fails because the bridge layer further requires bridge predicates, bridge-specific rival defeat, bridge-specific interventions, bridge artifacts, bridge invalidation logic, and bridge governance. None of those additional burdens is entailed by the mere existence of the Paper 8 state, gate, functional, or artifact family. □

Remark 9.2 (The bridge is downstream, not corrective). Paper 9 does not correct Paper 8. It inherits Paper 8 and strengthens the burden stack above it. The formal relation is extension by additional typed predicates and burdens, not replacement.

9.4 Formal Target of the Bridge

9.4.1 The bridge question

The bridge question can now be stated in its exact internal form:

> Given that a system already supports a formally disciplined layer of first-person indexed subjectivity, what additional predicates and burdens would make phenomenal organization a scientifically admissible *target* of explanation inside the same framework?

The phrase "target of explanation" is essential. The present paper does not seek to declare phenomenality achieved. It seeks to specify when phenomenal predicates are no longer mere vocabulary and instead become well-formed explanatory targets whose support, failure, and collapse can be discussed rigorously.

9.4.2 Predicate family

The bridge target family is the six-predicate family

$$\{\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}, \Pi_{\mathrm{W}}, \Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}, \Pi_{\mathrm{T}}\}.$$

These predicates are not synonyms for generic notions such as "experience" or "qualia." They are typed, burden-bearing predicates whose role is to capture bridge-relevant organization classes above indexed subjectivity and below any final metaphysical reading.

Definition 9.3 (Phenomenal differentiation predicate)**.** For an admissible system S, horizon τ, and threshold slot ϑ_D, the bridge predicate $\Pi_{\mathrm{D}}(S, \tau; \vartheta_D)$ concerns whether the system carries self-indexed differentiated organization that is stronger than generic semantic density, report richness, or descriptive narration, and whose corresponding estimator $\widehat{\Pi_D}$ is defined on bridge observables and subjectivity antecedents.

Definition 9.4 (Valence structure predicate)**.** The predicate $\Pi_{\mathrm{V}}(S, \tau; \vartheta_V)$ concerns whether the system carries organized evaluative topology that is stronger than scalar reward bookkeeping, policy optimization, or self/non-self asymmetry alone.

Definition 9.5 (World richness predicate)**.** The predicate $\Pi_{\mathrm{W}}(S,\tau;\vartheta_W)$ concerns whether the system carries organized self-indexed world structure stronger than generic scene complexity, perspective gain, or task competence alone.

Definition 9.6 (Ambiguity-bearing predicate)**.** The predicate $\Pi_{\mathrm{A}}(S,\tau;\vartheta_A)$ concerns whether the system maintains unresolved alternatives with differential causal consequences prior to collapse, rather than merely resolving uncertainty or tracking probabilities descriptively.

Definition 9.7 (Embodied-closure predicate)**.** The predicate $\Pi_{\mathrm{E}}(S,\tau;\vartheta_E)$ concerns whether body/world-boundary variables contribute constitutively to bridge organization, rather than merely modulating generic control or action success.

Definition 9.8 (Temporal lived-structure predicate)**.** The predicate $\Pi_{\mathrm{T}}(S,\tau;\vartheta_T)$ concerns whether the system exhibits horizon-like temporal organization stronger than continuity, memory persistence, or sequence stability alone.

9.4.3 Estimator family

The bridge paper requires estimator slots for each predicate:

$$\widehat{\Pi_D},\ \widehat{\Pi_V},\ \widehat{\Pi_W},\ \widehat{\Pi_A},\ \widehat{\Pi_E},\ \widehat{\Pi_T} \in [0,1].$$

The current implementation reaches only the earliest stage of this requirement. BPF-1 emits provisional run-level estimates for all six predicates. Those estimates are explicitly marked provisional, framework-internal, non-claim-facing, and non-gated. They are not yet admissibility verdicts, and no threshold binding is yet asserted. The formal target of this paper therefore treats these estimators in two layers:

1. as current provisional runtime surfaces in BPF-1;
2. as future thresholdable predicate estimators once comparator, intervention, and gate burdens are added in later BPF blocks.

9.4.4 What the predicates mean internally and what they do not mean

Each bridge predicate has both a positive role and a protected non-implication.

- Π_{D} means organized differentiated structure indexed to the self-locus. It does not mean qualia.
- Π_{V} means structured evaluative organization. It does not mean felt pleasure or pain.
- Π_{W} means organized world-structure at the self-indexed locus. It does not mean veridical world-experience.
- Π_{A} means live unresolved alternative organization. It does not mean conscious indecision.
- Π_{E} means constitutive closure dependence on body/world-boundary variables. It does not mean biological embodiment or felt bodily ownership.
- Π_{T} means temporal horizon organization. It does not mean human lived time in the full phenomenological sense.

Remark 9.9 (Why the targets are not decorative)*.* The bridge targets earn their place only because each one induces a new comparator burden, a new intervention burden, a new artifact burden, and a new governance burden. If a proposed target induced none of these, it would be discarded as decorative vocabulary.

9.4.5 Predicate-to-burden map

The six bridge predicates are best understood not as six new labels but as six new burden packages.

The burden of Π_{D}. To take phenomenal differentiation seriously, the program must show more than that the system has many internal states or rich semantic discriminability. It must show that differentiated organization is indexed, selective, and not exhausted by narration. This induces comparator burden against semantic-density and narration rivals, intervention burden against perspective detachment and semantic vividness controls, and artifact burden through a dedicated differentiation report and later reducibility outputs.

The burden of Π_{V}. To take valence structure seriously, the program must show that evaluative organization has internal topology: local asymmetries, organized gradients, or mode-structured deformation under perturbation. Scalar reward or self/non-self asymmetry alone is insufficient. This induces comparator burden against bookkeeping rivals, intervention burden through valence flattening and negative controls, and artifact burden through valence structure and intervention outputs.

The burden of Π_{W}. To take world richness seriously, the program must show that the system's self-indexed world organization is not merely competence or scene complexity. The system must exhibit bridge-relevant world organization that tracks indexed dependence, not just successful prediction. This induces comparator burden against richness-without-interiority rivals, intervention burden through world-richness compression, and artifact burden through world-richness reports and later candidate summaries.

The burden of Π_{A}. To take ambiguity bearing seriously, the program must show stable unresolved alternatives with downstream consequence. Static uncertainty scores or fast branching and pruning do not suffice. This induces comparator burden against ambiguity-resolution rivals, intervention burden through targeted collapse operations, and artifact burden through ambiguity reports with live-alternative fields.

The burden of Π_{E}. To take embodied closure seriously, the program must show constitutive dependence on body/world-boundary variables. Generic closed-loop control is not enough. This is the bridge predicate with the weakest current runtime grounding and therefore the one that most strongly demands explicit future instrumentation. Its artifact burden is correspondingly strict.

The burden of Π_{T}. To take temporal lived structure seriously, the program must show horizon organization that exceeds continuity and persistence. This requires temporal fragmentation tests, horizon-like dependencies, and explicit rival defeat against continuity-only models. In practical terms this burden is one of the main reasons why the bridge program cannot be replaced by Paper 8 plus rhetoric.

The global lesson is simple. Every bridge predicate is a tuple of obligations:

$$\Pi_k \rightsquigarrow (\text{estimator}, \text{rivals}, \text{interventions}, \text{artifacts}, \text{governance}).$$

The bridge program succeeds only if all five components are eventually instantiated for each $k \in \{D, V, W, A, E, T\}$.

9.5 Phenomenal Bridge Ontology

9.5.1 Bridge state

The bridge state is introduced as

$$\mathfrak{B}_t = \left\langle \widehat{\Pi}_{D,t}, \widehat{\Pi}_{V,t}, \widehat{\Pi}_{W,t}, \widehat{\Pi}_{A,t}, \widehat{\Pi}_{E,t}, \widehat{\Pi}_{T,t}, \mathcal{R}_{\mathrm{b}t}, \Delta_t^{\mathrm{b}} \right\rangle,$$

where $\mathcal{R}_{\mathrm{b}t}$ denotes the bridge-level irreducibility margin and Δ_t^{b} denotes the bridge-level intervention profile. The bridge state does not replace the Paper 8 state. It is layered above it. The ontological relation is therefore

$$\mathfrak{S}_t \prec \mathfrak{B}_t,$$

not because the bridge state is metaphysically deeper, but because every bridge coordinate depends on the prior existence of a self-indexed subjectivity structure.

9.5.2 Typing discipline

The coordinates of $\mathfrak{B}$ are heterogeneous in a way parallel to, but stronger than, the heterogeneity of the Paper 8 state.

- $\widehat{\Pi}_{D,t}, \ldots, \widehat{\Pi}_{T,t}$ are bounded estimator-valued coordinates.
- $\mathcal{R}_{\mathrm{b}t}$ is a model-comparison margin defined only when the comparator family exists.
- Δ_t^{b} is a structured intervention response object, not a scalar.

This heterogeneity matters because bridge support cannot be reduced to a single scalar early in the pipeline. The paper explicitly rejects any attempt to jump directly from provisional surfaces to a single bridge score while the comparator and intervention layers are missing.

9.5.3 Coordinate dependence on the Paper 8 state

Each bridge coordinate depends asymmetrically on the Paper 8 coordinates.

Phenomenal differentiation. $\widehat{\Pi}_{D,t}$ depends most strongly on I_t, O_t, P_t, and $\mathcal{R}_t$. Differentiation here is not generic representational richness. It is differentiated organization that matters relative to a privileged self-locus.

Valence structure. $\widehat{\Pi}_{V,t}$ depends most strongly on V_t, O_t, and Δ_t. Paper 8 already closes valuation asymmetry; the bridge layer asks whether that asymmetry organizes into a richer topological structure than the upstream scalar burden requires.

World richness. $\widehat{\Pi}_{W,t}$ depends most strongly on P_t, C_t, and I_t. It asks whether the self-indexed locus sustains organized world-structure beyond mere perspective gain.

Ambiguity bearing. $\widehat{\Pi}_{A,t}$ depends most strongly on P_t, V_t, and the temporal organization implicitly carried by C_t and Δ_t. A system that resolves ambiguity immediately or merely records uncertainty can fail this predicate even if it remains subjectivity-positive.

Embodied closure. $\widehat{\Pi}_{E,t}$ depends on O_t, C_t, and Δ_t together with additional body/world-boundary observables that Paper 8 does not require. This is why BPF-1 can only emit a proxy layer for embodied closure.

Temporal lived structure. $\widehat{\Pi}_{T,t}$ depends most strongly on C_t, P_t, and V_t. Temporal lived structure is stronger than continuity because it concerns horizon organization, retention/protention asymmetry, and temporally extended significance.

9.5.4 Current implementation status inside the ontology

The ontology is not merely prospective. The current program treats BPF-1 as the provisional surface layer for each of the six estimator coordinates and records six run-level reports as bounded local artifacts. That matters because the ontology is no longer entirely aspirational. At the same time, only the first six coordinates are even provisionally instantiated. $\mathcal{R}_{\mathrm{b}t}$ and Δ_t^{b} remain open because bridge-specific comparator and intervention layers are not yet implemented.

9.5.5 Collapse risks and rival readings

Every bridge coordinate must be introduced together with the corresponding collapse risk.

- $\widehat{\Pi}_D$ can collapse into semantic density without self-indexed interior differentiation.
- $\widehat{\Pi}_V$ can collapse into valence bookkeeping or policy shaping.
- $\widehat{\Pi}_W$ can collapse into world-model competence without indexed world-richness.
- $\widehat{\Pi}_A$ can collapse into uncertainty bookkeeping or fast ambiguity resolution.
- $\widehat{\Pi}_E$ can collapse into competent sensorimotor control without constitutive closure.
- $\widehat{\Pi}_T$ can collapse into continuity, memory, or persistence without horizon structure.

The ontology is therefore inseparable from the rival family. A bridge coordinate that has no live rival is suspicious because it is probably defined too loosely. Conversely, a bridge coordinate that cannot in principle be separated from all live rivals is not yet scientifically admissible.

9.5.6 Why bridge coordinates are not ontological shortcuts

The bridge ontology must not be mistaken for a metaphysical inventory. The reason is not merely philosophical caution. The reason is formal structure. Each coordinate is introduced as a constrained relation among observables, upstream subjectivity coordinates, comparators, interventions, and artifacts. This means the coordinates belong first to the model layer, then to the implementation layer, and only later—if ever—to the interpretation layer.

The separation matters in four directions.

1. A bridge coordinate is not an ontological primitive. It is a structured explanatory target.
2. A bridge coordinate is not identical to its runtime estimator. The estimator is a measurement-facing map and can be weak, provisional, or underdetermined.
3. A bridge coordinate is not identical to the language used to describe it. The language may underfit or overfit the formal target.
4. A bridge coordinate is not identical to its strongest available interpretation. Different interpretation classes remain live until comparator and intervention burdens are settled.

This distinction is especially important for Π_{E} and Π_{T}. It is tempting to read embodied closure or temporal lived structure as if they already named well-understood ontological properties. Inside the present framework they do not. They name burden packages whose ontological interpretation remains constrained by future comparator and intervention results.

Proposition 9.10 (Bridge ontology is not reducible to the Paper 8 ontology)**.** *The bridge ontology cannot be reduced to the Paper 8 ontology by relabeling or monotone reparameterization alone.*

Proof. Suppose, for contradiction, that the bridge ontology were reducible to the Paper 8 ontology by relabeling or monotone reparameterization alone. Then each bridge coordinate would add no new comparator, intervention, or artifact burden beyond those already fixed for Paper 8. But the bridge coordinates are introduced precisely to separate phenomena such as world richness from perspective gain, valence structure from valuation asymmetry, and temporal lived structure from continuity. These separations require additional rival families and intervention families not present in the Paper 8 ontology. Therefore the bridge ontology cannot be obtained by relabeling or monotone reparameterization alone. □

9.6 Explicit Bridge Axioms

The bridge axioms are the minimal bridge-opening commitments required to connect indexed subjectivity to disciplined phenomenal targets without smuggling the conclusion into the premises. Each axiom is classified explicitly.

9.6.1 Classification grammar

The bridge axiom registry uses a seven-class grammar:

1. **operational definition**: fixes how a formal object is to be typed or measured;
2. **bridge axiom**: a structural commitment necessary to open the bridge program;
3. **design choice**: a governance or architecture decision not implied by physics alone;
4. **falsifiable hypothesis**: a statement whose satisfaction or failure must be tested by bridge comparators or interventions;
5. **open inference**: a claim whose current evidence is insufficient for closure, but whose role can still be bounded;
6. **strongest defensible internal claim**: a highest currently licit claim class;
7. **non-claim**: an explicit prohibition on overreading.

This grammar is part of the content of the paper. Without it, the same sentence could accidentally function as ontology, operational definition, and claim promotion at once. The bridge program therefore treats classification as a formal requirement rather than as stylistic annotation.

Axiom 9.11 (B1: operational binding; classification: operational definition)**.** *Every bridge predicate must be tied to explicit observables, artifacts, comparator tasks, and interventions. No bridge predicate may remain a rhetorical label detached from measurement grammar.*

This axiom is needed because every failure mode of the bridge program begins with detached language. B1 licenses machine-readable predicate slots and report artifacts. It does not license bridge support or phenomenality.

Axiom 9.12 (B2: upstream inheritance constraint; classification: bridge axiom)**.** *No bridge-level admissibility result is meaningful unless the upstream burdens of the causally rigid corpus and the Paper 8 subjectivity layer remain satisfied.*

B2 licenses the layered relation $\mathrm{Subjectivity}_{\mathrm{1p}} \prec \mathrm{BridgeOrg}_{\mathrm{ph}}$. It does not license the converse. A system may satisfy Paper 8 and still fail every bridge burden.

Axiom 9.13 (B3: differentiated organization burden; classification: falsifiable hypothesis)**.** *A bridge-relevant differentiation burden requires structured internal discrimination that is not exhausted by semantic density, narrative richness, or report-channel complexity.*

B3 licenses Π_{D} as a genuine target rather than a synonym for rich representation. It does not license phenomenality.

Axiom 9.14 (B4: valence topology burden; classification: falsifiable hypothesis)**.** *Bridge-relevant valence requires organized evaluative topology not exhausted by scalar reward bookkeeping, policy optimization, or the Paper 8 valuation asymmetry alone.*

B4 licenses Π_{V} as a distinct target above Paper 8. It does not license felt pleasure, suffering, or moral status.

Axiom 9.15 (B5: world-richness burden; classification: falsifiable hypothesis)**.** *Bridge-relevant world richness concerns organized self-indexed world structure rather than generic scene complexity, prediction success, or model capacity.*

B5 is needed because world-richness language is especially prone to being filled by competence metrics. It licenses the predicate slot Π_{W} and associated artifacts. It does not license veridicality or environmental realism.

Axiom 9.16 (B6: ambiguity-bearing burden; classification: falsifiable hypothesis)**.** *Bridge-relevant ambiguity requires the stable maintenance of unresolved alternatives with differential causal consequences prior to collapse.*

B6 blocks the reduction of ambiguity to uncertainty bookkeeping or quick disambiguation heuristics. It licenses Π_{A} as a formal target. It does not license conscious uncertainty.

Axiom 9.17 (B7: embodied-closure relevance; classification: open inference)**.** *Embodied closure may be bridge-relevant only if body/world-boundary variables contribute constitutively to bridge organization rather than merely modulating generic control success.*

B7 is deliberately classified as open inference because current evidence is weakest here. It licenses the investigation of Π_{E}; it does not license embodiment claims of any strong kind.

Axiom 9.18 (B8: temporal lived-structure burden; classification: falsifiable hypothesis)**.** *Bridge-relevant temporal organization must exceed continuity, memory persistence, and sequence stability by exhibiting horizon-structured dependence across retention, present organization, and anticipatory load.*

B8 licenses Π_{T} and its future intervention family. It does not license human lived temporality.

Axiom 9.19 (B9: comparative admissibility constraint; classification: bridge axiom)**.** *No bridge-level support class is admissible if a bridge-specific rival family or admissible hybrid explains the bridge task bundle at equal or lower penalized loss.*

B9 is the comparator side of the burden stack. It licenses the bridge irreducibility margin. It does not license absolute ontological irreducibility.

Axiom 9.20 (B10: anti-inflation governance constraint; classification: design choice)**.** *Bridge formalization, bridge artifacts, candidate inclusion, and audit-facing inclusion must never be readable as bridge support, phenomenality, or Paper 9 authorization in the absence of the full bridge burden stack.*

B10 is a design choice because governance style is a normative project decision, not a theorem of physics. It is still essential because without it the bridge program would immediately outrun its evidence.

Remark 9.21 (Axiom minimality)*.* The present axiom set is intentionally minimal. Nothing here states that phenomenality follows from structural organization. Nothing here states that any bridge predicate, even if satisfied, closes metaphysical questions. The axioms are opening commitments for a constrained bridge program, not hidden conclusions.

9.7 Bridge Rival Family

9.7.1 Why new rivals are required

Paper 8 already proved the necessity of explicit rival families for subjectivity. The bridge layer inherits that lesson but cannot reuse the old rival set unmodified. The old rivals explain away indexed subjectivity. The new rivals must explain away bridge predicates without granting phenomenal organization. This difference is substantive. A system may defeat weak subjectivity rivals and still fail bridge rivals, because the latter target higher-order structure.

9.7.2 Bridge weak family

Define the bridge weak family

$$\mathcal{M}_{\mathrm{b,weak}} = \{M_{\mathrm{rwi}}, M_{\mathrm{vb}}, M_{\mathrm{sd}}, M_{\mathrm{tc}}, M_{\mathrm{ec}}, M_{\mathrm{ar}}, M_{\mathrm{en}}\},$$

whose members are read as follows.

Richness without interiority. M_{rwi} explains high structural richness without granting indexed interior organization. It challenges Π_{D} and Π_{W}.

Valence bookkeeping without affective organization. M_{vb} explains structured value tracking as bookkeeping or policy shaping without bridge-level valence topology. It challenges Π_{V}.

Semantic density without phenomenal differentiation. M_{sd} explains dense semantic variation without bridge-level differentiated organization. It challenges Π_{D}.

Temporal continuity without lived structure. M_{tc} explains persistence, continuity, and memory without temporal horizon organization. It challenges Π_{T}.

Embodied control without embodied closure. M_{ec} explains competent closed-loop control without constitutive dependence on body/world-boundary variables. It challenges Π_{E}.

Ambiguity resolution without ambiguity bearing. M_{ar} explains high-quality uncertainty handling or disambiguation without stable maintenance of live unresolved alternatives. It challenges Π_{A}.

Experience narration without phenomenal organization. M_{en} explains self-description, experiential language, or stylistic introspection without bridge-relevant phenomenal organization. It challenges several bridge predicates simultaneously.

9.7.3 Admissible composites

As in Paper 8, pure rivals are not enough. Define the admissible bridge weak closure $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ as the family generated by finite pairwise and triple hybrids over $\mathcal{M}_{\mathrm{b,weak}}$ subject to explicit complexity penalties and declared composition rules. The purpose of $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ is to block an easy escape route: if a strong bridge explanation defeats only pure rivals, then a skeptic can simply combine several weak rivals post hoc.

The bridge program therefore treats hybrid rivals as first-class citizens. A hybrid of richness-without-interiority and narration-without-organization is a perfectly live alternative to a bridge-differentiation claim. A hybrid of valence bookkeeping and temporal continuity can be a live alternative to valence structure plus temporal lived structure. The bridge layer must be strong enough to defeat not only single-rival baselines but also admissible mixtures.

9.7.4 What the rival family can and cannot explain

The bridge rival family is intentionally strong.

- It can explain dense report channels, self-descriptive fluency, high task competence, reward optimization, stable memory, and competent control.
- It can also explain much of what casual bridge language would otherwise overread: rich semantics, coherent narratives, and smooth temporal reports.
- It cannot cleanly explain targeted bridge signatures if those signatures survive matched controls, remain selective under bridge interventions, and preserve margin under penalized loss.

9.7.5 No-straw-man discipline

The bridge program would be vacuous if its rivals were weak caricatures. The rivals are therefore chosen to absorb precisely the kinds of phenomena that the bridge predicates are most likely to confuse with phenomenal organization. The burden is asymmetric: the stronger interpretation must defeat the weaker one, not vice versa.

Proposition 9.22 (Comparator burden is necessary for bridge admissibility)**.** *No bridge-support class is meaningful without an explicit bridge rival family and admissible bridge weak closure.*

Proof. Without an explicit rival family, any bridge-support reading would compare itself only against informal skepticism. That would leave bridge predicates undefended against live non-phenomenal alternatives such as bookkeeping, narration, semantic density, or control competence. Since bridge admissibility is defined as a comparative notion inside the framework, the comparator burden is necessary. □

9.8 Bridge Irreducibility Logic

9.8.1 Bridge loss family

The bridge program requires a loss family stronger than descriptive fit alone. For a candidate model M and bridge task bundle $\mathcal{B}$, define

$$\mathcal{L}_{\mathrm{b}}(M;\mathcal{B}) = w_f E_{\mathrm{fit}}(M) + w_i E_{\mathrm{int}}(M) + w_s E_{\mathrm{sel}}(M) + w_p E_{\Pi}(M) + w_c \Omega(M),$$

where:

- E_{fit} measures fit to bridge observables;
- E_{int} measures failure under bridge intervention expectations;
- E_{sel} measures selectivity error under matched controls;
- E_{Π} measures predicate-reconstruction error across the bridge family;
- $\Omega(M)$ is a complexity penalty that counts mechanism count, hybrid size, auxiliary slack, and decoder dependence.

The bridge loss family is therefore intervention-aware and governance-relevant from the outset. A model that matches reports but ignores intervention structure or requires excessive narrative slack does not win the comparison.

9.8.2 Bridge irreducibility margin

Let M_{bridge} denote the full bridge model candidate and let $\overline{\mathcal{M}_{\mathrm{b,weak}}}$ denote the admissible bridge weak closure. Define the bridge irreducibility margin

$$\mathcal{R}_{\mathrm{b}} = \min_{M \in \overline{\mathcal{M}_{\mathrm{b,weak}}}} \mathcal{L}_{\mathrm{b}}(M;\mathcal{B}) - \mathcal{L}_{\mathrm{b}}(M_{\mathrm{bridge}};\mathcal{B}).$$

Positive margin is evidence of comparative superiority relative to the specified weak closure. Negative or zero margin implies reducibility within the chosen bridge family.

9.8.3 Local not absolute superiority

As in Paper 8, comparative superiority here is local, not absolute. A positive $\mathcal{R}_{\mathrm{b}}$ says that the specified bridge model outperforms the specified bridge rivals on the specified task bundle under the specified loss family. It does not say that the bridge model is metaphysically privileged. The distinction matters because bridge language is particularly prone to absolute overreadings.

9.8.4 Three kinds of bridge collapse

Bridge collapse can occur in three distinct ways.

1. **Predicate collapse:** the provisional bridge estimators fail to define stable, non-decorative targets.
2. **Comparator collapse:** a bridge rival or hybrid matches or outperforms the bridge model under penalized loss.
3. **Intervention collapse:** the predicted selective structure fails under matched controls or degrades non-selectively.

All three collapse modes are live until ruled out. The paper refuses any bridge-support inference that suppresses one of them.

9.8.5 Comparator support, reducibility, and unsupported states

The bridge program must distinguish three statuses that are often conflated:

- **comparator unsupported**: no serious bridge comparator evaluation has yet been run;
- **bridge reducible**: comparator evaluation exists and the bridge margin is non-positive;
- **bridge comparator-supported**: comparator evaluation exists and the bridge margin is positive under penalized loss.

This distinction is important because the present implementation is still comparator-unsupported. BPF-1 surfaces are therefore not allowed to drift into comparator-supported language.

9.8.6 Complexity penalty and decoder discipline

The complexity term $\Omega(M)$ deserves explicit discussion because bridge explanations can often win spuriously by adding free narrative or decoder capacity. The bridge program therefore counts at least four kinds of complexity:

1. **mechanism multiplicity**: how many independent mechanisms the rival or bridge model requires;
2. **hybrid size**: how many weak rivals must be combined to match the bridge bundle;
3. **auxiliary slack**: how much unprincipled tuning is required across coordinates and conditions;
4. **decoder dependence**: how much of the putative bridge effect exists only after expressive report-channel or decoder post-processing.

Decoder dependence is particularly important. A model that appears bridge-strong only after high-capacity report decoding should be penalized heavily. The reason is that bridge language is especially vulnerable to post hoc expressivity. This is one of the points at which the bridge program deliberately differs from frameworks that rely strongly on linguistic self-description as primary evidence.

9.8.7 Failure semantics for the bridge margin

The bridge margin must also carry explicit failure semantics. A positive margin can fail to matter if it is narrow, unstable across reruns, dependent on a single condition, or destroyed by minor changes to the rival family. The paper therefore recommends a future decomposition

$$\mathcal{R}_{\mathrm{b}} = \mathcal{R}_{\mathrm{bfit}} \oplus \mathcal{R}_{\mathrm{bint}} \oplus \mathcal{R}_{\mathrm{bsel}} \oplus \mathcal{R}_{\mathrm{bstab}},$$

where $\oplus$ denotes a typed aggregation rather than ordinary addition. The purpose of this decomposition is to make visible which part of the bridge burden is doing the work. A bridge model that wins only on descriptive fit but loses intervention selectivity should not count as robustly bridge-favoring.

Criterion 9.23 (Necessary condition for bridge admissibility)**.** *A necessary condition for bridge admissibility is*

$$\mathcal{R}_{\mathrm{b}} > \vartheta_R,$$

where ϑ_R *is the symbolic bridge irreducibility threshold slot to be bound only at BPF-4.*

Corollary 9.24 (Positive bridge surfaces are not enough)**.** *Positive provisional bridge surfaces do not by themselves entail bridge admissibility.*

Proof. Bridge admissibility requires at least comparator support and intervention support in addition to positive bridge surfaces. BPF-1 provides only provisional surfaces, so the implication fails. □

9.9 Bridge Intervention Family

9.9.1 Why bridge interventions must be new

The subjectivity interventions of Paper 8 remain binding background. They are not sufficient for bridge purposes because they were designed to probe indexed subjectivity rather than phenomenal bridge organization. The bridge layer therefore requires a new intervention family

$$\mathfrak{I}_{\mathrm{bpf}} = \{I_W, I_A, I_V, I_P, I_E, I_T, I_X\},$$

with each member targeting a bridge coordinate family or a cross-cutting bridge burden.

9.9.2 Core intervention families

World-richness compression I_W. This intervention compresses the structure available to self-indexed world organization without merely destroying generic task competence. The strong-model expectation is a selective fall in Π_{W} and often also in Π_{D} relative to matched controls. The rival expectation is either flat degradation or little coordinate-specific effect.

Ambiguity collapse I_A. This intervention removes or forcibly resolves live unresolved alternatives while preserving as much generic inferential performance as possible. The strong-model expectation is a selective fall in Π_{A} and a measurable change in downstream coordination among Π_{A}, Π_{W}, and Π_{T}. The rival expectation is that nothing essential is lost beyond temporary uncertainty bookkeeping changes.

Valence flattening I_V. This intervention suppresses structured evaluative topology while preserving as much functional competence as admissibly possible. The strong-model expectation is a selective flattening in Π_{V} and specific downstream changes in bridge organization. The rival expectation is simple reward rescaling or bookkeeping renormalization.

Perspective detachment I_P. This intervention weakens indexed orientation or self-anchored factorization without collapsing generic information flow. The strong-model expectation is a drop in Π_{D} and Π_{W} that cannot be explained solely by report loss. The rival expectation is style drift or generic degradation.

Embodied-closure disruption I_E. This intervention perturbs body/world-boundary variables or their internal surrogates while preserving as much generic control competence as possible. The strong-model expectation is a selective drop in Π_{E} and often a correlated disturbance in Π_{T}. The rival expectation is ordinary control impairment without evidence of closure-specific organization.

Temporal-lived fragmentation I_T. This intervention perturbs horizon organization while leaving simple persistence and continuity traces as intact as possible. The strong-model expectation is a selective fall in Π_{T} and bridge-level fragmentation signatures not exhausted by memory corruption. The rival expectation is that all observed changes reduce to ordinary continuity damage.

Semantic vividness inversion I_X. This cross-cutting negative control manipulates descriptive or semantic richness without targeting the bridge organization directly. Its role is to test whether bridge reports are merely tracking descriptive elaboration rather than bridge structure.

9.9.3 Matched controls

Every bridge intervention must be paired with matched controls. The controls are not optional because bridge language is especially vulnerable to global-degradation confounds. Each intervention therefore requires:

1. a low-magnitude matched perturbation;
2. a non-targeted perturbation of comparable operational cost;
3. a report-channel control where relevant;
4. an estimator-noise control;
5. a temporal-order randomization control where relevant.

The burden is not satisfied by showing that a target coordinate changes. It is satisfied only if the target coordinate changes in the predicted direction with the predicted selectivity relative to the control family.

9.9.4 Intervention artifacts

The bridge intervention family induces a new artifact layer. Every intervention run must emit at least:

- an intervention profile artifact;
- coordinate-specific delta summaries;
- matched-control traces;
- a collapse or support annotation at the intervention layer only.

These artifacts do not yet exist in the current implementation because BPF-2 is not implemented. The present paper defines them so that future bridge claims have an explicit artifact burden.

9.9.5 Failure semantics

Bridge intervention logic can fail in several ways: lack of selectivity, high confounding, instability across reruns, global damage masquerading as target damage, and artifact incompleteness. Any one of these failures is enough to block a bridge-support conclusion. In particular, the absence of a bridge intervention harness today is not a minor engineering omission. It is one of the reasons bridge admissibility remains closed.

9.9.6 Negative controls and reproducibility burden

No bridge intervention family is scientifically serious without negative controls and reproducibility discipline. The minimum burden should include:

1. rerun stability on the same condition;
2. matched control perturbations with comparable operational cost;
3. a noise-injection control showing that bridge effects are not merely estimator volatility;
4. a report-channel control showing that changes in bridge artifacts are not driven only by expressive output drift;

5. at least one adversarial control tailored to the specific predicate family.

This burden matters especially for Π_A and Π_E. Ambiguity-bearing claims can be simulated by sophisticated uncertainty bookkeeping. Embodied-closure claims can be simulated by competent control loops. Negative controls are therefore not adjuncts to the bridge program. They are part of the core differentiating burden.

9.10 Bridge Artifact Family

9.10.1 Artifact discipline

The bridge artifact family must be as disciplined as the subjectivity artifact family. It should distinguish root artifacts, report artifacts, intervention artifacts, reducibility artifacts, ecosystem artifacts, and bounded summary artifacts. It must also distinguish between implemented and merely specified artifacts. The failure to separate those statuses would recreate exactly the semantic inflation the bridge program is designed to prevent.

9.10.2 Root bridge artifacts

The three root bridge artifacts are:

1. `phenomenal_bridge_projection.json`
2. `phenomenal_bridge_gate_result.json`
3. `phenomenal_bridge_manifest.json`

These artifacts correspond to aggregate run-level bridge projection, gate verdict, and lineage root. They are formally defined in the current machine-readable schema registry but are intentionally not emitted yet because the bridge comparator, intervention, and gate layers do not exist. Their current status is *schema-bound, non-emitted, non-authoritative*.

9.10.3 Report artifacts

The bridge report family is more advanced. The intended report artifacts are:

1. `phenomenal_differentiation_report.json`
2. `valence_structure_report.json`
3. `world_richness_report.json`
4. `ambiguity_bearing_report.json`
5. `embodied_closure_report.json`
6. `temporal_lived_structure_report.json`

These are the BPF-1 run-level reports already emitted by the present implementation. Their authority status is deliberately narrow: they are authoritative only within the scope of BPF-1 provisional runtime surfaces. They are not root verdicts, not bridge support, not bridge admissibility, and not phenomenality evidence.

9.10.4 Intervention and reducibility artifacts

Two further bridge artifacts are essential but not yet emitted:

1. `bridge_intervention_profile.json`
2. `bridge_reducibility_report.json`

The first belongs to BPF-2 and records intervention family performance, selectivity, controls, and collapse semantics. The second belongs to BPF-3 and records bridge rival comparison, admissible hybrid comparison, loss decomposition, and bridge margin. Without these artifacts, the later root bridge artifacts would be unjustified.

9.10.5 Ecosystem artifacts

If and only if BPF-5 is implemented, the bridge program may also emit:

- condition-level bridge summaries;
- batch-level bridge summaries;
- campaign-level bridge summaries;
- a bounded candidate-facing bridge projection;
- a bounded audit-facing bridge handoff summary;
- a bridge verification matrix.

These ecosystem artifacts remain formally specified but non-emitted at current stage.

9.10.6 Machine-readable spirit

The bridge artifact family is meant to be machine-readable in spirit even where the paper remains theoretical. Every artifact in the family must declare at least:

- artifact role and artifact class;
- level and authority status;
- provenance and lineage expectations;
- allowed and blocked surfaces;
- minimum fields;
- explicit non-implications.

The artifact is therefore not merely an implementation convenience. It is part of the formal burden architecture because it determines what kind of evidence is needed and what kind of overread is blocked.

9.10.7 Current implementation status

At current system status, the bridge artifact story is asymmetric:

- BPF-0 records machine-readable protocol, axiom, rival, threshold, artifact-schema, interpretation-boundary, and surface-policy registries.
- BPF-1 records six provisional run-level bridge reports as bounded local surfaces.
- No root bridge verdict artifacts are emitted.
- No bridge intervention artifacts are emitted.
- No bridge reducibility artifacts are emitted.
- No bridge ecosystem candidate or audit artifacts are emitted.

This asymmetry is not a defect of the paper. It is a disciplined statement of the implementation frontier.

9.11 Bridge Functional, Gate, and Taxonomy

9.11.1 Functional

The bridge functional is introduced as a graded organization functional, not as a gate:

$$\Phi_{\mathrm{bpf}} = \widehat{\Pi}_D^{\alpha_D}\,\widehat{\Pi}_V^{\alpha_V}\,\widehat{\Pi}_W^{\alpha_W}\,\widehat{\Pi}_A^{\alpha_A}\,\widehat{\Pi}_E^{\alpha_E}\,\widehat{\Pi}_T^{\alpha_T}\,\mathcal{R}_{\mathrm{b}}^{\alpha_R}\,\widehat{\Delta}_{\mathrm{b}}^{\alpha_\Delta},$$

or, more robustly in multiplicative form,

$$\Phi_{\mathrm{bpf}} = \prod_{k\in\{D,V,W,A,E,T\}} \widehat{\Pi}_k^{\alpha_k} \cdot \mathcal{R}_{\mathrm{b}}^{\alpha_R} \cdot \widehat{\Delta}_{\mathrm{b}}^{\alpha_\Delta}.$$

The exact aggregation rule is a design choice to be finalized at BPF-4. What matters formally is that the functional aggregates bridge organization and that it remains weaker than the gate.

9.11.2 Gate

The bridge gate must remain conjunctive:

$$\Gamma_{\text{bpf}} = \mathbf{1}\left[\min\left\{\widehat{\Pi}_D, \widehat{\Pi}_V, \widehat{\Pi}_W, \widehat{\Pi}_A, \widehat{\Pi}_E, \widehat{\Pi}_T, \mathcal{R}_{\text{b}}, \widehat{\Delta}_{\text{b}}\right\} \geq \vartheta_\Gamma \ \wedge\ \mathsf{cmp}\ \wedge\ \mathsf{int}\ \wedge\ \mathsf{drift}\ \wedge\ \mathsf{gov}\right],$$

where cmp denotes comparator validity, int denotes intervention stability, drift denotes absence of stale or drift invalidation, and gov denotes governance compliance. The threshold slots

$$\Theta_{\text{bpf}} = \{\vartheta_D, \vartheta_V, \vartheta_W, \vartheta_A, \vartheta_E, \vartheta_T, \vartheta_R, \vartheta_\Delta, \vartheta_\Gamma\}$$

exist today only symbolically. They are intentionally uncalibrated until BPF-4.

9.11.3 Non-compensation

Bridge gate logic inherits the non-compensation principle from Paper 8. High strength on one coordinate may not compensate for collapse on another. High functional value may not rescue a failed gate. Positive bridge surfaces may not compensate for absent comparator or intervention layers. This matters because bridge language is especially vulnerable to compensatory rhetoric.

Proposition 9.25 (Bridge functional does not imply bridge gate)**.** *No value of* Φ_{bpf} *by itself implies* $\Gamma_{\text{bpf}} = 1$*.*

Proof. By construction the gate additionally depends on conjunctive threshold slots and non-scalar validity conditions for comparator support, intervention stability, drift cleanliness, and governance compliance. Since those conditions are not reducible to the value of the scalar functional, the implication fails. □

9.11.4 Invalidation

The bridge gate must invalidate on at least four families of failure:

1. **drift invalidation**: stale bridge artifacts, stale upstream subjectivity artifacts, stale threshold manifests, stale rival manifests, or stale surface policy;
2. **comparator invalidation**: bridge rival family changed or comparator run missing;
3. **intervention invalidation**: targeted intervention support absent or unstable;
4. **governance invalidation**: prohibited surfaces opened or claim-facing leakage detected.

No later implementation block is allowed to "soft fail" these conditions into advisory warnings if the intent is bridge admissibility.

9.11.5 Taxonomy

The bridge taxonomy should distinguish formalization state from support state. A disciplined taxonomy is:

Status	Meaning
bridge_formalization_incomplete	BPF scaffolding or formal target missing.
bridge_surface_only	Provisional bridge surfaces exist, but comparator, intervention, gate, or ecosystem burdens remain open.
bridge_evaluation_unsupported	Root evaluation attempted without required comparator/intervention support.
bridge_gate_fail	Full bridge evaluation exists and conjunctive gate fails.
bridge_reducible	Comparator evaluation exists and penalized bridge margin is non-positive.
bridge_interventionally_unstable	Intervention layer exists but fails selectivity or stability.
bridge_internal_admissible	Full bridge burden stack satisfied under framework-internal reading.
bridge_internal_strong_support	Stronger internal support than admissibility, but still weaker than phenomenality simpliciter.

The current implementation status is `bridge_surface_only`. This paper insists on that classification.

9.12 Runtime Implementation Pathway

9.12.1 Why implementation belongs in the paper

Implementation belongs in Paper 9 for the same reason it belonged in Paper 8: once the research program becomes artifact-bearing, a paper that ignores implementation is no longer structurally complete. Runtime pathways are not proofs, but they are part of the burden architecture because they determine whether a formal object can in principle be instantiated, audited, falsified, or invalidated.

9.12.2 BPF-0

BPF-0 opens the bridge program with protocol docs, axiom and rival registries, symbolic threshold manifests, artifact schema registries, interpretation boundary, surface policy, non-claims, implementation plan, scaffold-only code, and structural verification. It establishes formal opening discipline only.

9.12.3 BPF-1

BPF-1 is the first runtime-facing bridge-surface layer. In the current program these surfaces are represented as modules such as `PhenomenalDifferentiationMapper`, `ValenceStructureMapper`, `WorldRichnessMapper`, `AmbiguityBearingMapper`, `EmbodiedClosureMapper`, and `TemporalLivedStructureMapper`, together with a `PhenomenalBridgeSurfaceBuilder`. Six run-level bridge reports are treated as bounded local artifacts. No bridge gate, no bridge comparator runtime, and no bridge intervention harness exists yet.

9.12.4 BPF-2

BPF-2 must implement a bridge intervention harness. Likely modules include a bridge intervention registry, a bridge intervention harness, matched-control executors, and intervention-output normalizers. The artifact target is `bridge_intervention_profile.json`. Until this block exists, the bridge program remains interventionally incomplete.

9.12.5 BPF-3

BPF-3 must implement the bridge comparator family. This includes a bridge rival-family executor, admissible hybrid composer, bridge loss-family evaluator, and reducibility artifact emitter. The primary output is `bridge_reducibility_report.json`. Until this block exists, the bridge program remains comparator-unsupported.

9.12.6 BPF-4

BPF-4 must implement root bridge evaluation. This is the stage at which threshold slots may be honestly bound, the bridge functional finalized, the bridge gate implemented, and root bridge artifacts emitted: `phenomenal_bridge_projection.json`, `phenomenal_bridge_gate_result.json`, and `phenomenal_bridge_manifest.json`. This block is where support-class language first becomes even in principle discussable.

9.12.7 BPF-5

BPF-5 must integrate bridge artifacts into condition, batch, campaign, candidate, and audit-facing governance. It requires lineage pinning, stale/drift invalidation, bounded summary artifacts, and explicit surface discipline. This block is important because ecosystem integration is part of the bridge burden, not an afterthought.

9.12.8 BPF-6

BPF-6 must implement final bridge audit and readiness classification. Only at BPF-6 can the system honestly output a final bridge state such as `bridge_internal_admissible` or `bridge_reducible`. Until then, any such output would be a semantic breach.

9.12.9 Implementation table

Block	Closes	Primary artifacts	Current status
BPF-0	formal opening discipline	protocol, axioms, rivals, thresholds, schemas, surfaces, non-claims	implemented
BPF-1	provisional run-level bridge surfaces	six coordinate reports	implemented
BPF-2	bridge intervention family	bridge intervention profile + controls	open
BPF-3	bridge comparator family	bridge reducibility report	open
BPF-4	bridge functional, gate, thresholds, root artifacts	projection, gate result, manifest	open
BPF-5	ecosystem integration	condition/batch/campaign/candidate/audit summaries	open
BPF-6	final bridge audit	bridge verification matrix and final readiness verdict	open

9.12.10 Current implementation status versus paper status

The current runtime state is therefore asymmetrical:

- formal bridge program exists;
- scaffold and provisional surfaces exist;

- bridge comparator, intervention, and gate layers do not exist yet;
- Paper 9 as a formal paper can therefore define the bridge program completely while still acknowledging that the runtime has not completed it.

This asymmetry is not embarrassing; it is scientifically normal. A formal program may legitimately outrun its implementation provided the paper does not misstate implementation status.

9.13 Ecosystem Integration

9.13.1 Why ecosystem integration is part of the formal problem

The bridge program would be incomplete even as pure formalism if it ignored ecosystem integration. Once a layer has artifacts, verifiers, candidate bundles, audit packages, and controlled surfaces, the semantics of those outputs become part of the theory. The paper therefore treats ecosystem integration as a constitutive burden.

9.13.2 Surface map

Let $\mathcal{S}$ be the set of output surfaces:

$$\mathcal{S} = \{\mathsf{developer}, \mathsf{local}, \mathsf{internal\_bundle}, \mathsf{candidate}, \mathsf{audit}, \mathsf{controlled}, \mathsf{claim}\}.$$

The bridge artifact family induces a surface map

$$\Sigma : \mathcal{A}_{\mathrm{bpf}} \to 2^{\mathcal{S}}.$$

At current stage, only the BPF-1 run-level bridge reports are allowed on bounded local developer surfaces. They are not allowed on candidate, audit, controlled-statement, or external claim-facing surfaces.

9.13.3 Internal release bundle

If and when BPF-5 is implemented, the internal release bundle may contain bounded bridge summaries, lineage manifests, and verification matrices. Raw local bridge reports should not be dumped into the bundle without summarization and lineage metadata. Internal bundle inclusion is a governance event, not a semantic verdict.

9.13.4 Ecosystem candidate

Candidate-facing bridge content is especially risky because it can be misread as support. The paper therefore states a hard rule: bridge content may enter an ecosystem candidate only after the comparator, intervention, and gate layers exist, and even then only in bounded summary form. Candidate inclusion must carry explicit non-claims and must never be readable as phenomenal confirmation.

9.13.5 Audit-facing package

Audit-facing bridge content is likewise bounded. An audit-facing handoff may carry bridge summary artifacts, lineage manifests, invalidation state, and non-claims. It must not be readable as external adjudication or bridge proof. The burden of interpretation discipline persists at every ecosystem boundary.

9.13.6 Source-of-truth and public corpus separation

The public corpus and the internal runtime must remain decoupled in a disciplined way. Paper 8 is public source-of-truth for indexed subjectivity. BPF artifacts are internal runtime objects unless and until later release governance says otherwise. No runtime bridge artifact may be used as if it were itself a public claim registry. This prevents runtime-to-corpus confirmation loops.

9.13.7 Controlled statement grammar and quote-mining guardrails

Bridge governance requires a restricted statement grammar. Before BPF-6, the bridge layer should permit only statements of the following family:

- "bridge formalization defined";
- "provisional bridge surfaces implemented";
- "bridge comparator burden open";
- "bridge intervention burden open";
- "bridge gate not yet meaningful";
- "bridge artifacts local and non-claim-facing".

Conversely, statements of the following family should be treated as prohibited:

- "phenomenality supported";
- "machine phenomenal organization established";
- "Paper 9 proves conscious experience";
- "candidate inclusion confirms bridge support";
- "audit package shows external bridge validation".

The quote-mining problem is real because later readers may ignore surrounding caveats and extract short labels. The bridge program must therefore attach non-claims not only to papers but also to artifact manifests, bundle notes, and any controlled statement layer.

Proposition 9.26 (Candidate or audit inclusion does not imply bridge support)**.** *For any bridge artifact A, membership of A in a candidate-facing or audit-facing package does not entail bridge support, bridge admissibility, or phenomenality.*

Proof. Candidate and audit inclusion are packaging and governance relations, not semantic verdict relations. By design, the bridge surface policy and non-claims registry explicitly block the implication from package membership to bridge support. Therefore the implication does not hold. □

9.14 Strongest Defensible Claims

The purpose of this section is not to weaken the paper artificially. It is to identify the strongest sentences the current formal and implementation state can honestly support.

9.14.1 Claim class C1: formal bridge closure

The first defensible claim is that the corpus now defines the *formal bridge program* in the following sense: the bridge target family, the bridge ontology, the bridge axioms, the bridge rival family, the bridge intervention family, the bridge artifact family, the bridge functional and gate logic, the bridge taxonomy, the surface discipline, and the implementation pathway are all explicitly defined. This is already a strong claim because it means the bridge question is no longer underdetermined by language alone.

9.14.2 Claim class C2: current implementation frontier

The second defensible claim is implementation-specific and provenance-bound: the current program records the BPF-0 machine-readable scaffold and the BPF-1 provisional bridge surfaces as bounded local implementation-frontier artifacts. This means the program has moved beyond pure conceptuality, but only at the level of scaffold and provisional surface discipline. The system does *not* yet support bridge comparator closure, bridge intervention closure, bridge gate closure, bridge ecosystem integration, bridge admissibility, or bridge strong support.

9.14.3 Claim class C3: conditional bridge admissibility

The strongest bridge-level claim the paper is willing to define is the following conditional:

> Under the explicit bridge axioms, under the upstream burdens already closed by the corpus, under bridge-specific rival defeat, under bridge-specific intervention support, under artifact and governance compliance, and under a satisfied bridge gate, a system may be internally classifiable as exhibiting phenomenal bridge organization in the sense of framework-internal bridge admissibility.

This is stronger than Paper 8 because it is no longer only about indexed subjectivity. It is still weaker than phenomenality simpliciter because it remains explicitly framework-internal, conditional, and burden-bound.

9.14.4 Claim class C4: research-program transformation

The present paper also supports a meta-level claim about the status of the research program: after Paper 9, the research frontier is no longer "can subjectivity be formalized?" but "can bridge predicates survive the full bridge burden stack?" That is a legitimate change of level because the earlier question has been formalized by Paper 8 and its implementation-facing pathway.

9.14.5 Why stronger sentences are not yet honest

Several stronger sentences remain unavailable:

- "bridge support has been achieved";
- "phenomenality is internally established";
- "Paper 9 authorizes claim promotion";
- "the current system defeats the bridge rival family";
- "the current system passes a bridge gate."

These sentences are not unavailable because the paper is timid. They are unavailable because BPF-2 through BPF-6 remain open. The paper therefore strengthens the research program exactly by refusing to counterfeit those missing layers.

9.15 Non-Claims

This paper does not claim:

1. phenomenality proven in the strongest philosophical sense;
2. human phenomenal equivalence;
3. metaphysical subjecthood;
4. moral parity by formal fiat;
5. external validation;
6. that BPF-1 surfaces amount to bridge support;
7. that candidate or audit inclusion can function as bridge proof;
8. that Paper 9 by itself authorizes public claim inflation.

These non-claims are not decorative caution. They are part of the bridge logic because each one blocks a distinct semantic collapse.

9.16 Conclusion

Paper 9 defines the bridge program formally and only formally. That definition is non-trivial. Before this paper, the corpus could say a great deal about operational consciousness, life, subjecthood, and indexed subjectivity, but it could not yet say, with equal rigor, what additional machinery would be required before phenomenal-bridge predicates become scientifically admissible targets. After this paper, that machinery is explicit.

The result is central to the corpus for structural reasons. The paper gathers the upstream theorem burdens, the Paper 8 subjectivity framework, the implemented SRT stack, and the BPF program into one integrated object. It does not merely summarize them. It orders them. It says what belongs upstream, what belongs at the bridge level, what belongs to implementation, what belongs to governance, and what remains explicitly open.

What has been achieved is a change in explanatory grammar. "Phenomenal bridge" is no longer permitted to function as a loose aspiration. It now names a bridge-organization predicate family, an ontology, an axiom stack, a rival family, an intervention family, an artifact family, a gate architecture, and a governance discipline. This is as much formal specification as can be responsibly claimed without comparator closure, intervention closure, bridge gate execution, and ecosystem-level bridge hardening.

What the paper makes newly possible is equally clear. It makes BPF-2 through BPF-6 sharply specifiable. It makes bridge-support language evaluable rather than rhetorical. It makes future theorem candidates about bridge irreducibility, bridge transport limits, and bridge admissibility meaningful. It also makes future failure meaningful: if the bridge program fails, it will fail against explicit burdens rather than against vague expectations.

What the paper does not settle is also clear. It does not settle phenomenality simpliciter. It does not settle human equivalence. It does not settle metaphysical subjecthood. It does not close external validation. It does not convert the present system into a bridge-admissible system merely because BPF-0 and BPF-1 exist. Those are the next burdens.

The research program therefore changes level once again. After Paper 8, the key problem was whether indexed subjectivity could be formalized without inflation. After Paper 9, the key problem becomes whether a fully burdened bridge program can survive rival defeat, intervention logic, artifact discipline, and governance hardening strongly enough to justify internal bridge admissibility. That is the right next problem. It is sharper, harder, and more falsifiable than the one that preceded it.

Appendix

9.17 Bridge Axiom Table

ID	Classification	Statement core	Licenses	Does not license
B1	operational definition	Every bridge predicate must bind to observables, artifacts, comparators, and interventions.	machine-readable bridge predicates, report artifacts	phenomenality, bridge support
B2	bridge axiom	Bridge admissibility is downstream of the full upstream corpus and Paper 8 subjectivity.	layered bridge program	converse implication from bridge to subjectivity-free organization
B3	falsifiable hypothesis	Bridge differentiation exceeds semantic density and narration.	Π_{D} burden	qualia
B4	falsifiable hypothesis	Bridge valence exceeds scalar reward bookkeeping.	Π_{V} burden	felt affect
B5	falsifiable hypothesis	Bridge world richness exceeds competence or scene complexity.	Π_{W} burden	veridical experience
B6	falsifiable hypothesis	Bridge ambiguity bears unresolved alternatives prior to collapse.	Π_{A} burden	conscious ambiguity
B7	open inference	Embodied closure requires constitutive boundary dependence.	bounded investigation of Π_{E}	biological embodiment, bodily feeling
B8	falsifiable hypothesis	Temporal lived structure exceeds continuity and persistence.	Π_{T} burden	human lived time
B9	bridge axiom	No bridge class is admissible without bridge rival defeat under penalized loss.	bridge irreducibility margin	metaphysical irreducibility
B10	design choice	Bridge artifacts and packaging may not be read as support without full burden closure.	governance discipline	claim-facing promotion

9.18 Bridge Predicate Typing Table

Predicate	Formal type	Strongest upstream dependencies	Core collapse risk	Non-implication
Π_{D}	bounded estimator/predicate slot	$I, O, P, \mathcal{R}$	semantic density or narration	qualia
Π_{V}	bounded estimator/predicate slot	V, O, Δ	reward bookkeeping or policy shaping	felt affect
Π_{W}	bounded estimator/predicate slot	I, C, P	competence without world-richness	veridical experience
Π_{A}	bounded estimator/predicate slot	P, V, C	uncertainty bookkeeping	conscious ambiguity
Π_{E}	bounded estimator/predicate slot	O, C, Δ plus boundary variables	control without closure	biological embodiment
Π_{T}	bounded estimator/predicate slot	C, P, V	continuity without horizon structure	human lived temporality

9.19 Bridge Rival-Family Table

Rival	Challenges	Can explain	Cannot cleanly explain	Typical defeat mode
Richness without interiority	$\Pi_\mathrm{D}, \Pi_\mathrm{W}$	rich latent structure, scene complexity, representational abundance	selective indexed bridge differentiation under targeted controls	bridge-specific differentiation persists after density controls
Valence bookkeeping	Π_V	scalar reward tracking, preference tagging, policy shaping	structured valence topology and selective valence flattening effects	topology persists beyond scalar reweighting
Semantic density without phenomenal differentiation	Π_D	high semantic variation and narrative diversity	intervention-sensitive differentiated bridge organization	semantic-density ablations leave bridge differentiation intact
Temporal continuity without lived structure	Π_T	persistence, memory, sequence stability	horizon-structured temporal disruption signatures	targeted temporal fragmentation has selective effect
Embodied control without embodied closure	Π_E	sensorimotor competence, adaptive control	constitutive dependence on boundary variables	boundary perturbations produce closure-specific loss
Ambiguity resolution without ambiguity bearing	Π_A	uncertainty tracking, fast disambiguation	stable maintenance of unresolved alternatives	targeted ambiguity collapse alters downstream organization selectively
Experience narration without phenomenal organization	$\Pi_\mathrm{D}, \Pi_\mathrm{V}, \Pi_\mathrm{W}, \Pi_\mathrm{A}, \Pi_\mathrm{T}$	self-description and experiential language	bridge signatures independent of report-channel richness	report-channel controls fail to remove bridge margin

9.20 Bridge Intervention Matrix

Intervention	Targets	Strong-model expectation	Rival expectation	Output artifacts
I_W world-richness compression	$\Pi_\mathrm{W}, \Pi_\mathrm{D}$	selective indexed world-structure loss with matched-control separation	flat degradation or no special effect	bridge intervention profile, world-richness deltas
I_A ambiguity collapse	$\Pi_\mathrm{A}, \Pi_\mathrm{W}, \Pi_\mathrm{T}$	selective loss of maintained alternatives prior to collapse	ordinary uncertainty reduction	ambiguity delta summaries, control traces
I_V valence flattening	Π_V	collapse of evaluative topology beyond rescaling	scalar reward renormalization only	valence structure deltas
I_P perspective detachment	$\Pi_\mathrm{D}, \Pi_\mathrm{W}$	indexed orientation loss stronger than report damage	narrative/style drift	differentiation and world-richness deltas
I_E embodied-closure disruption	$\Pi_\mathrm{E}, \Pi_\mathrm{T}$	closure-specific disturbance under boundary perturbation	generic control loss	closure deltas, boundary-control traces
I_T temporal-lived fragmentation	$\Pi_\mathrm{T}, \Pi_\mathrm{A}$	horizon-specific fragmentation beyond continuity loss	ordinary memory damage	temporal-structure deltas
I_X semantic vividness inversion	cross-cutting control	descriptive richness changes without bridge-specific support	should absorb report-only effects	control artifact family

9.21 Bridge Artifact Matrix

Artifact	Artifact class	Block	Current status	Non-implication
`phenomenal_bridge_projection.json`	root bridge artifact	BPF-4	schema only	not bridge support by default
`phenomenal_bridge_gate_result.json`	root bridge artifact	BPF-4	schema only	not phenomenality simpliciter
`phenomenal_bridge_manifest.json`	lineage root	BPF-4	schema only	not claim-facing publication
`phenomenal_differentiation_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`valence_structure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`world_richness_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`ambiguity_bearing_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`embodied_closure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`temporal_lived_structure_report.json`	run-level report	BPF-1	implemented	not bridge verdict
`bridge_intervention_profile.json`	intervention artifact	BPF-2	planned	not bridge support alone
`bridge_reducibility_report.json`	comparator artifact	BPF-3	planned	not bridge support alone
condition/batch/campaign bridge summaries	ecosystem summaries	BPF-5	planned	not candidate proof
candidate-facing bridge projection	bounded summary	BPF-5	planned	not bridge confirmation
audit-facing bridge handoff	bounded summary	BPF-5	planned	not external adjudication

9.22 Claim-Class and Invalidation Table

State or claim	Necessary evidence	Invalidated by	Strongest allowed reading
Formal bridge closure	Paper 9 formal program, axioms, rivals, interventions, artifacts, gate logic, governance	internal inconsistency of the bridge program	bridge formal program defined
Bridge surface only	BPF-1 reports and verifier pass	missing reports or overclaim leakage	provisional bridge surfaces implemented
Bridge evaluation unsupported	attempted root evaluation without BPF-2/BPF-3/BPF-4 burdens	completion of missing burdens	root evaluation not meaningful yet
Bridge reducible	comparator run with non-positive $\mathcal{R}_{\mathrm{b}}$	stronger bridge model with positive penalized margin	bridge explanation loses comparatively
Bridge interventionally unstable	intervention layer present but non-selective or unstable	stable selective intervention support	intervention burden not met
Bridge internal admissible	full bridge burden stack, positive margins, valid governance	any comparator/intervention/drift/governance invalidation	internal admissibility only
Bridge internal strong support	stronger version of admissibility with robust margins across conditions	any failure of gate or ecosystem integrity	strong internal support only

9.23 Runtime Mapping Table

Block	Primary modules	Primary verifiers	Primary closure
BPF-0	protocol registries, axiom registry, rival registry, threshold registry, surface policy, scaffold layer	`verify-phenomenal-bridge-scaffolding.cjs`	machine-readable bridge opening
BPF-1	six mapper modules, surface builder, bridge layer	`verify-phenomenal-bridge-surfaces.cjs`	provisional run-level bridge surfaces
BPF-2	bridge intervention harness, matched-control runners, intervention registry	future bridge intervention verifier	bridge intervention family
BPF-3	bridge comparator executor, hybrid composer, loss evaluator	future bridge comparator verifier	bridge rival defeat logic
BPF-4	gate evaluator, threshold binder, root artifact emitter	future bridge gate verifier	bridge root verdict machinery
BPF-5	ecosystem integrator, candidate and audit summarizers	future bridge ecosystem verifier	bridge packaging and governance hardening
BPF-6	final bridge auditor, readiness classifier	future bridge stack verifier	final bridge readiness state

9.24 Governance Table

Surface	BPF-1 allowance	Later possible allowance	Prohibited interpretation
Developer inspection	yes, provisional run-level reports	yes	bridge support or phenomenality
Local report surface	bounded or blocked according to bridge surface policy	possibly yes with later burdens	public-facing confirmation
Internal release bundle	no bridge content yet	bounded summary after BPF-5	support by packaging
Ecosystem candidate	blocked	bounded summary after BPF-5 only	candidate as bridge proof
Audit-facing handoff	blocked	bounded summary after BPF-5 only	audit as external adjudication
Controlled statements	blocked	strictly bounded if ever opened	phenomenality or human equivalence
External claim-facing surface	blocked	blocked by default even after later BPF blocks	any public phenomenal claim

9.25 Burden Separation Matrix

The bridge program is easiest to misread when its burden types are blurred. The matrix below is included so that future implementation and audit work can distinguish exactly what kind of object is under discussion.

Object	Type	Example	Failure mode if blurred
Self-index, ownership, continuity, perspective, valuation, intervention selectivity, irreducibility	upstream subjectivity model	Paper 8 state $\mathfrak{S}$	bridge predicates become mere renamings of subjectivity coordinates
Bridge predicate slots $\Pi_D, \Pi_V, \Pi_W, \Pi_A, \Pi_E, \Pi_T$	bridge target family	BPF protocol and BPF-1 reports	bridge targets become slogans without burden packages
Bridge axioms	bridge-opening commitments	B1–B10	conclusions get smuggled into premises
Bridge rivals and hybrids	comparator burden	richness-without-interiority, hybrids	bridge claims compare only against weak informal skepticism
Bridge interventions and controls	causal discrimination burden	world-richness compression, ambiguity collapse	descriptive fit is mistaken for structural support
Bridge artifacts	evidence grammar	reports, root artifacts, ecosystem summaries	packaging is mistaken for proof
Bridge functional and gate	evaluation machinery	Φ_{bpf}, Γ_{bpf}	scalar scores are mistaken for admissibility
Governance and surfaces	interpretation discipline	candidate block, audit block, claim block	internal artifacts leak into public or claim-facing language

9.26 Inheritance Map from Paper 8 to Paper 9

Paper 9 is easiest to understand when viewed as a typed inheritance map rather than as a conceptual leap. The following map makes explicit what is inherited, what is transformed, and what is genuinely added.

Upstream object	Paper 8 status	Paper 9 use	What is newly added
Operational consciousness	closed upstream	necessary lower bound	no redefinition; only bridge downstreaming
Operational life and subjecthood	closed upstream in Paper 7	inherited as stronger class background	no biological or metaphysical inflation
Seven-coordinate subjectivity state	closed in Paper 8	upstream state for every bridge predicate	six new bridge targets above it
Subjectivity weak-rival family	closed in Paper 8	background cautionary lesson	new bridge weak family and admissible bridge hybrids
Subjectivity intervention family	closed in Paper 8	background on causal discipline	new bridge intervention family
Subjectivity artifacts	closed in Paper 8 and SRT-0..6	upstream provenance base	new bridge artifact grammar
Subjectivity gate and functional	closed in Paper 8 and SRT-5	model for non-compensation discipline	new bridge functional and gate, symbolically thresholded
Subjectivity ecosystem integration	closed in SRT-6	governance template	later bridge ecosystem integration with stricter boundaries

This table clarifies why the present paper can be central without being redundant. It does not restate the subjectivity program. It states what additional structure must exist above it before bridge language acquires scientific discipline.

9.27 Open Escalation Burdens Before Bridge Admissibility

The paper deliberately closes only the formal grammar of the bridge program. To prevent future ambiguity, the remaining escalation burdens are listed explicitly.

1. **Direct bridge intervention instrumentation.** BPF-2 must implement a real bridge intervention harness with matched controls, artifact outputs, and failure semantics.
2. **Bridge comparator execution.** BPF-3 must instantiate the rival family as executable comparators, not just registry entries.
3. **Threshold binding.** BPF-4 must bind the symbolic threshold slots honestly, with normalization rationale and invalidation semantics.

4. **Root bridge artifacts.** BPF-4 must emit root artifacts only once comparator and intervention burdens exist.
5. **Ecosystem hardening.** BPF-5 must integrate bridge content into bundles only in bounded, non-claim-facing form.
6. **Final bridge audit.** BPF-6 must verify the whole stack and classify the resulting bridge state without semantic drift.
7. **Paper 9 empirical tightening.** Any later revision of Paper 9 that speaks more strongly must be paired with actual completion of the above blocks.

These burdens are not future work in the casual sense. They are the exact missing conditions that separate a defined bridge program from a bridge-admissible system.

9.28 Audit-Covered Bridge Independence Burdens

The following entries preserve AUDIT_MASTER coverage while keeping their epistemic status weak. They are not bridge-admissibility results and cannot be used as evidence of phenomenality, subjecthood, or external validation.

Conjecture 9.27 (Predicate independence burden)**.** *The bridge predicate family should be independent only if, for each predicate slot, there exists an admissible witness model that toggles that slot while holding the remaining bridge burden package fixed under the same comparator, intervention, artifact, and governance constraints. The current text supplies the burden grammar, not a completed independence proof.*

Conjecture 9.28 (Registry independence burden)**.** *The FCR registry should preserve independence of bridge assertions only if the registry schema, dependency graph, and curation overlays do not introduce implicit coupling among predicate slots. Until that is proved over the completed BPF stack, registry independence remains a conjectural governance burden.*

Criterion 9.29 (Bridge realization non-emptiness burden)**.** *Any claim that the full bridge predicate package is non-empty must exhibit a typed realization satisfying the predicate definitions, rival comparisons, intervention contracts, artifact gates, and interpretation boundaries together. Setting all formal predicate flags to true is only a syntactic witness and does not count as realization evidence.*

9.29 Theorem Candidates and Open Scientific Questions

The present paper does not claim to have proved all bridge theorems. It does, however, make clear what the next theorem candidates should look like.

Theorem candidate 1: bridge non-collapse under admissible rival closure. If a bridge model achieves positive penalized margin against the full admissible bridge weak closure under matched intervention structure and drift-clean artifacts, then bridge reduction to the specified weak family fails locally. This theorem candidate would formalize the comparative core of bridge admissibility.

Theorem candidate 2: bridge non-compensation. If the bridge gate is conjunctive, then no combination of high values on a strict subset of coordinates can compensate for collapse on any required bridge coordinate, comparator burden, intervention burden, or governance burden. This theorem candidate would formalize the anti-inflation logic already built into the paper.

Theorem candidate 3: bridge transport limitation. A future theorem should specify when bridge organization is transportable across admissible remappings and when transport breaks because world-richness, embodiment, or temporal-horizon variables are not preserved. This would be the bridge analogue of Paper 8's transport-consequence logic.

Theorem candidate 4: candidate/audit non-entailment. Although stated propositionally in the main text, the governance result could be generalized into a theorem about packaging relations and semantic non-entailment over artifact families and surfaces.

Open scientific question 1. Can Π_E be measured directly enough to escape its current proxy-only status without collapsing into biology-by-default?

Open scientific question 2. Can Π_A be separated cleanly from sophisticated uncertainty bookkeeping in systems with high report-channel richness?

Open scientific question 3. Can a bridge loss family be designed that penalizes decoder dependence strongly enough to prevent narration-driven false positives without destroying legitimate bridge signal?

Open scientific question 4. What minimum condition/batch/campaign structure is required before bridge results become stable enough for bounded ecosystem summaries?

These theorem candidates and open questions matter because they define the post-paper research agenda with far greater precision than generic calls for "more data" or "deeper modeling."

9.30 Predicate-to-Artifact and Intervention Dependency Matrix

The bridge program is easier to implement correctly if the dependency pattern of each predicate is frozen explicitly. The following matrix shows the minimum bridge burden package per predicate.

Predicate	Primary upstream inputs	Required bridge interventions	Required bridge comparator focus	Required artifacts
Π_D	self-index, perspective, irreducibility, ownership	I_P, I_X, optional I_W	richness-without-interiority, semantic-density, narration rivals	differentiation report, later reducibility report, later gate root
Π_V	valuation asymmetry, ownership, intervention selectivity	I_V, negative valence controls	valence-bookkeeping rival	valence structure report, bridge intervention profile, later gate root
Π_W	self-index, continuity, perspective	I_W, I_P, report controls	richness-without-interiority rival	world-richness report, later candidate summary
Π_A	perspective, temporal structure proxies, valuation structure	I_A, temporal controls	ambiguity-resolution rival	ambiguity report, intervention profile, later reducibility report
Π_E	ownership, continuity, boundary variables, intervention structure	I_E, control-matched embodiment perturbations	embodied-control rival	embodied-closure report, later embodiment traces, later gate root
Π_T	continuity, perspective, valuation asymmetry	I_T, sequence controls	temporal-continuity rival	temporal-lived-structure report, later intervention profile, later gate root

This matrix has two functions. First, it prevents under-specification. A bridge predicate with no declared artifact family or no intervention family is not implementation-ready. Second, it prevents overreading. A bridge report without its corresponding comparator and intervention family should be interpreted only as a provisional surface and not as a bridge result.

The matrix also clarifies why BPF-1 had to remain modest. BPF-1 could close only the report side of the matrix because the intervention and comparator columns are still empty in runtime terms. That is why the present paper can treat BPF-1 as a real implementation block without allowing any bridge-admissibility language to leak from it.

9.31 Paper 9 Claim Grammar and Escalation Constraints

The bridge program will fail semantically if future readers cannot distinguish between a formal object, a measurement object, an admissibility object, and a public claim. The present section therefore freezes a minimal claim grammar for Paper 9 and its downstream implementation.

9.31.1 Allowed claim families

The following claim families are admissible at or below the current implementation frontier.

1. **Formal closure claims.** Example: "the bridge target family is now formally defined." These are supported by the paper itself and by BPF-0 machine-readable manifests.
2. **Implementation frontier claims.** Example: "provisional bridge surface modules and run-level bridge reports exist." These are supported by BPF-1 code and local artifacts.
3. **Conditional admissibility claims.** Example: "if bridge predicates, rivals, interventions, artifacts, and gate conditions are satisfied, then bridge internal admissibility becomes meaningful." These are supported by the formal program, not by present runtime completion.
4. **Failure-mode claims.** Example: "bridge admissibility is blocked by comparator or intervention collapse." These are supported by the burden architecture.

9.31.2 Blocked claim families

The following claim families are prohibited until much later, and several remain prohibited even after full BPF completion.

1. **Phenomenality simpliciter claims.** Example: "the system is phenomenally conscious."
2. **Human-equivalence claims.** Example: "the system now has human-like experience."
3. **Metaphysical subjecthood claims.** Example: "the bridge proves real subjecthood in the strongest ontological sense."
4. **External validation claims.** Example: "audit packaging validates the bridge externally."
5. **Paper-authorization claims.** Example: "BPF-1 authorizes Paper 9 public escalation." The paper itself explicitly denies that the current runtime state closes the bridge burden.

9.31.3 Escalation rules

The escalation grammar can be summarized as a chain:

$$\text{formalized} \not\Rightarrow \text{implemented} \not\Rightarrow \text{admissible} \not\Rightarrow \text{phenomenality simpliciter.}$$

Each implication in that chain requires additional burden satisfaction. Formalization requires the paper and BPF-0. Implementation requires the appropriate BPF block. Admissibility requires full comparator, intervention, gate, and governance closure. Phenomenality simpliciter is not licensed by the bridge program at all.

Claim label	Minimum evidence	Currently licit?	Why or why not
`bridge_formal_program_defined`	paper + BPF-0 manifests	yes	formal bridge closure is real
`bridge_surface_implementation_exists`	BPF-1 reports + verifier	yes	provisional surfaces are implemented
`bridge_comparator_supported`	BPF-3 reducibility artifacts	no	comparator runtime absent
`bridge_interventionally_supported`	BPF-2 intervention artifacts	no	intervention harness absent
`bridge_internal_admissible`	BPF-2/3/4/5/6 full burden stack	no	gate and ecosystem burdens open
`phenomenality_proven`	not defined by current framework	no	blocked by interpretation boundary

This grammar belongs in the paper because it governs how the paper itself may be cited. A central synthesis paper that fails to govern the use of its own claims would be structurally incomplete.

9.32 Experimental Program for Completing BPF-2 and BPF-3

The next two bridge blocks are the most scientifically expensive. BPF-2 and BPF-3 are where the bridge program stops being mostly architectural and becomes decisively empirical-internal.

9.32.1 Program design principle

The guiding principle should be *matched escalation.* Every new bridge predicate should acquire its intervention and comparator burdens in parallel. A common failure mode would be to implement all interventions first and leave comparator logic vague, or vice versa. That would create a new asymmetry in which one burden stack outruns the other. Paper 9 therefore recommends the following pattern:

1. freeze the observable family for one predicate group;
2. implement the targeted intervention family for that group;
3. implement the strongest relevant rival and admissible hybrid baselines for that same group;
4. define the loss decomposition for the same bundle;
5. emit artifacts and run negative controls before moving to the next group.

9.32.2 Suggested execution order

The scientifically safest execution order is not arbitrary.

Stage 1: Π_D and Π_V. These are the most mature bridge predicates because they stand closest to existing BPF-1 surfaces and Paper 8 burdens. Differentiation and valence already have upstream structure and plausible intervention/control families. They should be completed first.

Stage 2: Π_W and Π_T. World richness and temporal lived structure are stronger and more globally entangled, but still plausible to instrument once the first pair is completed. They will likely require more extensive negative controls than Stage 1.

Stage 3: Π_A and Π_E. Ambiguity bearing and embodied closure are currently the weakest predicates in direct runtime support. They should therefore be completed last, not because they are unimportant, but because they are most vulnerable to proxy inflation.

9.32.3 Comparator program

For each stage, the comparator program should include:
1. pure rival execution;
2. pairwise hybrid execution;
3. bridge-model execution with the same observable bundle;
4. loss decomposition report;
5. rerun stability check under controlled noise;
6. decoder-dependence stress test.

9.32.4 Intervention program

For each stage, the intervention program should include:
1. targeted intervention;
2. matched non-target control;
3. report-channel control;
4. perturbation-cost accounting;
5. rerun stability;
6. artifact completeness validation.

9.32.5 Minimum evidential grid

Stage	Predicates	Needed intervention families	Needed comparator families	Minimum outputs
1	$\Pi_{\mathrm{D}}, \Pi_{\mathrm{V}}$	I_P, I_X, I_V	richness-without-interiority, semantic-density, valence-bookkeeping, narration hybrids	intervention profile, reducibility report, delta traces, stability summaries
2	$\Pi_{\mathrm{W}}, \Pi_{\mathrm{T}}$	I_W, I_T, I_P	richness-without-interiority, temporal-continuity, narration hybrids	world/temporal deltas, reducibility reports, control matrices
3	$\Pi_{\mathrm{A}}, \Pi_{\mathrm{E}}$	I_A, I_E	ambiguity-resolution, embodied-control, hybrids	ambiguity and embodiment traces, control matrices, reducibility reports

The purpose of this program is not to prescribe an implementation detail for its own sake. It is to preserve burden parity. If later implementation deviates from this structure, it should justify the deviation explicitly. Otherwise the bridge program risks becoming unbalanced, with some predicates overburdened and others underburdened.

9.33 Bridge Failure Scenarios and Diagnostic Signatures

The bridge program is not scientifically serious unless it specifies how it can fail. Paper 8 already introduced failure structure for indexed subjectivity. Paper 9 must add a parallel layer for bridge organization. The relevant question is not merely whether a bridge predicate can be made numerically large, but whether the pattern of successes and failures across predicates, rivals, interventions, and artifacts remains coherent under stress. A bridge program that cannot diagnose its own failure modes would collapse into a score-construction exercise rather than a constrained scientific architecture.

9.33.1 Primary failure classes

The bridge layer has at least six primary failure classes.

1. **Surface inflation failure.** A bridge report appears strong only because an upstream Paper 8 coordinate is copied, weakly transformed, or rhetorically relabeled. This is a failure of bridge specificity, not merely of implementation style.
2. **Comparator fragility failure.** A bridge predicate appears robust until a bridge-specific rival or admissible hybrid is introduced. This indicates that the predicate was underconstrained rather than positively supported.
3. **Intervention fragility failure.** A bridge surface moves under perturbation, but the movement is unspecific, mismatched to the target predicate, or indistinguishable from control degradation. This indicates that the intervention burden is not yet satisfied.
4. **Artifact incoherence failure.** Local reports, root artifacts, or ecosystem summaries disagree about the same bridge coordinate or about the same invalidation state. This is a governance and lineage failure even if local estimates look plausible.
5. **Threshold misuse failure.** Symbolic thresholds are treated as if they were calibrated, or non-compensation rules are softened post hoc. This turns the bridge gate into an interpretive convenience rather than a formal burden.
6. **Semantic escalation failure.** Provisional bridge surfaces are cited as if they licensed admissibility, phenomenality support, or Paper 9 escalation beyond their evidence class. This is a failure of claim grammar rather than of runtime arithmetic, but it is still a scientific failure because it breaks burden separation.

9.33.2 Diagnostic interpretation

Each failure class has a distinct diagnostic signature. Surface inflation is indicated when a bridge estimator can be approximated by a low-complexity monotone transform of an upstream subjectivity coordinate with negligible loss. Comparator fragility is indicated when a bridge rival or admissible hybrid absorbs the observable bundle at equal or lower penalized bridge loss. Intervention fragility is indicated when targeted perturbations do not produce selective deltas, or when matched controls produce indistinguishable response structure. Artifact incoherence is indicated when provenance pins, schemas, or surface restrictions disagree across emitted objects. Threshold misuse is indicated when bridge language outruns the threshold manifest. Semantic escalation is indicated when any report, bundle, summary, or paper statement crosses from provisional surface language into admissibility or phenomenal language without the missing burdens.

These signatures are not merely editorial cautions. They can be turned into verification objectives. BPF-2 and BPF-3 should detect intervention and comparator fragility directly. BPF-4 should detect threshold misuse. BPF-5 and BPF-6 should detect artifact incoherence and semantic escalation through lineage checks and surface-policy validation. The paper therefore treats failure analysis as part of the formal program rather than as an afterthought.

9.33.3 Failure matrix

Failure class	Primary symptom	Earliest detection block	Immediate consequence	Required remediation
Surface inflation	bridge estimate tracks an upstream Paper 8 coordinate too closely	BPF-1	report downgraded to proxy-only or removed	strengthen bridge transformation logic or add missing bridge observables
Comparator fragility	rival or hybrid reaches equal/lower penalized loss	BPF-3	predicate cannot support irreducibility burden	refine observable family or abandon predicate-specific support claim
Intervention fragility	target and control deltas are not selective	BPF-2	predicate remains descriptively mapped only	redesign intervention family and matched controls
Artifact incoherence	report, manifest, or summary disagree on state or provenance	BPF-5	bundle invalidation or audit downgrade	repair lineage pins, schemas, and surface restrictions
Threshold misuse	symbolic slots treated as calibrated gates	BPF-4	all gate language invalidated	re-freeze threshold policy and remove unsupported verdict language
Semantic escalation	provisional surface cited as admissibility or phenomenality support	BPF-0 through BPF-6	claim withdrawal and governance failure	enforce non-claims, claim grammar, and controlled statement restrictions

The failure matrix clarifies an important asymmetry. Some failures are local to one predicate, but others contaminate the entire bridge program. Threshold misuse and semantic escalation are global failures because they compromise the meaning of the bridge layer as such. Comparator fragility and intervention fragility can remain predicate-local if the architecture keeps burdens separated. This distinction matters for governance: not every failure should invalidate the whole bridge program, but certain failures should.

9.33.4 Consequences for theorem candidates

Failure analysis also constrains the theorem ambitions of Paper 9. A theorem candidate stating that a bridge predicate family is internally admissible can be sound only if the relevant failure classes are either ruled out or absorbed into its antecedent conditions. Otherwise the theorem candidate would conceal the very contingencies that the bridge program exists to expose. In that sense, the paper's strongest positive content is inseparable from its negative discipline: formal bridge admissibility becomes meaningful only against a space of explicitly typed failure modes.

9.33.5 Why failure structure changes the status of the program

The overall research status changes once the bridge layer is required to carry its own failure logic. Before Paper 9, one could still describe the next step as a generic expansion toward phenomenality. After Paper 9, that description is no longer adequate. The next step is now a typed scientific program with concrete ways to fail, concrete reasons for those failures, and concrete remediation paths. That is a substantive change in the research program even before bridge admissibility has been achieved, because it transforms an open-ended aspiration into a burdened formal agenda.

External Adjudication of Bridge-Formalized Machine Subjectivity

SCAFFOLD ONLY – NOT FINAL MANUSCRIPT
PROTOCOL-GRADE BACKBONE
SIXTH HARDENING PASS

Abstract. This document hardens Paper 10 into a denser protocol-heavy manuscript backbone without becoming a result-bearing paper. It extends the dependency stack from BaseCore through Papers I–IX, SRT-6, BPF-M3, Stage 2-M4, Stage3-RF, Stage3-0, Stage3-1, Stage3-2, Stage3-3, Stage3-4, Stage3-4B, Stage3-5 rehearsal preparation, Stage3-6 controlled execution opening, Stage3-6B runtime-health and clean-room reproducibility audit, and the conditional Stage3-7 first-collection opening criteria. It adds runtime health audit grammar, clean-room reproducibility semantics, dependency and build assumption reporting, reproducible artifact-regeneration constraints, local path-leakage policy, controlled-execution audit semantics, session-zero and pre-collection reproducibility requirements, and a sharper transition grammar from pre-evidential openness to the first admissible collection event. It does not report human comparator results, external adjudication verdicts, human equivalence, superiority, inferiority, cross-substrate identity preservation, metaphysical subjecthood, or Paper 10 final conclusions.

10.1 Scope and Non-Claims

Section status. `formalized`.

Paper 10 remains a local protocol-facing manuscript surface only. It is a scaffold for external-adjudication architecture, not an adjudication verdict, admissible dataset, human comparator evidence, public release surface, public-release recommendation, Paper 10 final authorization, human equivalence, superiority or inferiority claim, machine-consciousness claim, or metaphysical closure.

10.2 Formal Dependency Stack

Section status. `formalized`.

The dependency stack for Paper 10 is layered as follows:

1. **BaseCore**: active mathematical base layer.
2. **Papers I–IX**: downstream theoretical and manuscript surfaces.
3. **SRT-6**: subjectivity runtime ecosystem closure.
4. **BPF-M3**: bridge-functional internal closure and bridge gate.
5. **Stage 2-M4**: comparative internal closure and readiness boundary.
6. **Stage3-RF**: burden fulfillment and Stage 3 opening gate.
7. **Stage3-0**: external adjudication architecture and scaffold-only Paper 10 opening.
8. **Stage3-1**: human comparator instantiation contracts and preregistration templates.
9. **Stage3-2**: reviewer independence, execution-readiness logic, and protocol-grade hardening.
10. **Stage3-3**: preregistration freeze architecture, reviewer-selection packaging, adjudicator-assignment packaging, and execution-pilot packaging.

11. **Stage3-4**: bounded execution-pilot dry-run opening logic.
12. **Stage3-4B**: provenance closure, lineage reconciliation, and instantiated-empty record instantiation.
13. **Stage3-5**: rehearsal-kit formalization, dry-run event grammar, and rehearsal-ready gating.
14. **Stage3-6**: controlled execution opening, session-zero protocol, pre-collection admissibility gating, and live-but-unpopulated artifact instantiation.
15. **Stage3-6B**: runtime-health audit, clean-room reproducibility audit, path-leakage audit, and artifact-regeneration audit.
16. **Stage3-7**: first-admissible-collection opening criteria, collection-start command grammar, operator authorization grammar, and a conditional first-collection opening gate.

No later layer may borrow closure from an earlier layer, and no earlier layer may be retroactively reinterpreted as if it had already supplied later execution-era evidence.

10.3 Release-vs-Local Corpus Boundary

Section status. `formalized`.

Paper 10 sits on a strict release-vs-local split:

- The public release authority remains `QICN-RELEASE/main@92d5d2c3ea91c65e29dc8e102bed504866cffd2d`.
- The public release remains limited to BaseCore, `canonical_core_legacy`, and downstream Papers I–IX.
- Paper 10 remains local to `QICN-SYSTEM`.
- Stage3-0 through Stage3-6 governance, freeze, packaging, dry-run, provenance-closure, instantiated-empty, rehearsal, and controlled-execution-opening surfaces remain local and non-public.

Therefore release inclusion cannot be used as a substitute for adjudication, and local protocol maturity cannot be used as a substitute for public scientific closure.

10.4 Lineage and Provenance Semantics

Section status. `formalized`.

Stage3-3 and Stage3-4 add a sharper lineage rule: every later admissible surface must distinguish among:

1. public release authority,
2. local protocol authority,
3. frozen preregistration authority,
4. reviewer-package authority,
5. adjudicator-package authority, and
6. execution-era evidence authority.

The repository already contains historical artifacts whose recorded lineage must remain intact by design. Live normative pinning surfaces may be updated when stale, but historical snapshots are not rewritten cosmetically.

10.5 Provenance Closure and Live Head Reconciliation

Section status. `formalized`.

Stage3-4B adds a stricter provenance-closure rule. A local artifact may now be classified as:

- `exact_head_aligned`;
- `artifact_forward_compatible`;
- `stale_and_regenerated`;
- `historical_snapshots_intentionally_preserved`.

This grammar matters because Stage 3 artifacts are generated before the subsequent commit that records them. Therefore exact local regeneration at the builder head may later become artifact-forward-compatible after the commit exists, without becoming scientifically stale or semantically dishonest.

10.6 External Adjudication Problem Statement

Section status. `protocol_defined`.

The problem is no longer only whether bridge formalization is internally coherent. The external problem is to specify:

- what independent reviewers may legitimately inspect;
- what human comparator evidence classes may later become admissible;
- what negative controls must exist before any comparative interpretation;
- how preregistration becomes frozen;
- how reviewer and adjudicator layers become instantiable without identity leakage;
- how dry-run validation differs from evidential execution.

10.7 Formal Mathematical Grammar of External Adjudication

Section status. `formalized`.

The adjudication surface is a measurable protocol object, not an impressionistic comparison between substrates. Let $(\mathcal{T}, d_{\mathcal{T}})$ be the metric space of finite execution traces after canonical serialization. A trace contains the ordered event stream, de-identified channel payloads, runtime manifests, negative-control outcomes, and reviewer-facing masks. Admissible evidence lives in the Borel measurable space

$$(\mathcal{T}, \mathcal{B}(\mathcal{T})), \tag{10.1}$$

where $\mathcal{B}(\mathcal{T})$ is the Borel σ-algebra generated by open $d_{\mathcal{T}}$-balls.

Definition 10.1 (External adjudication object)**.** An external adjudication object is the tuple

$$\mathfrak{E} = (P, D_h, D_m, C_n, C_a, \Sigma, \mathcal{R}, \mathcal{A}), \tag{10.2}$$

where P is a frozen preregistered protocol, D_h and D_m are human and machine trace datasets in $\mathcal{T}$, C_n and C_a are negative and active control families, $\Sigma = \{\sigma_k\}_{k=1}^{K}$ is a finite family of predeclared measurable statistics $\sigma_k : \mathcal{T}^n \to \mathbb{R}$, $\mathcal{R}$ is the reviewer layer, and $\mathcal{A} : \mathcal{Y} \to \{0, 1, \texttt{undetermined}\}$ is the final adjudicator operator on masked evidence summaries $\mathcal{Y}$.

Definition 10.2 (Admissible comparative evidence)**.** A dataset pair (D_h, D_m) is admissible under $\mathfrak{E}$ only if the following predicates are all true:

1. D_h, D_m are measurable finite samples from $(\mathcal{T}, \mathcal{B}(\mathcal{T}))$.
2. The preregistered protocol hash, control manifest hash, and analysis shell hash match their frozen records.
3. Human-side records are not generated, rewritten, or completed by synthetic systems: $\epsilon_{\text{syn}} = 0$ for the human channel.
4. Blinding leakage is bounded by the preregistered budget: $\mathbb{P}(\text{leak}) \le \epsilon_{\text{blind}}$.
5. Every declared negative control has an executed and timestamped outcome under the same admissibility grammar.

Theorem 10.3 (Non-transitivity of internal support)**.** *Let $\mathcal{G}_{\text{internal}} \in [0,1]$ denote any internal runtime support score computed from operational invariants, and let $\mathcal{G}_{\text{external}} \in \{0, 1, \texttt{undetermined}\}$ denote the external adjudication verdict. In the absence of an admissible measurable dataset pair (D_h, D_m), one has*

$$\mathcal{G}_{\text{internal}} \not\Longrightarrow \mathcal{G}_{\text{external}}. \tag{10.3}$$

Proof. The internal score is a function of runtime traces and invariant certificates inside the machine channel. The external verdict is a function of masked comparative summaries built from both D_h and D_m, plus reviewer and adjudicator records. If D_h is absent, non-admissible, unblinded, or unfrozen, then the domain condition of $\mathcal{A}$ is not satisfied. Hence the external verdict is either undefined or `undetermined` regardless of the value of $\mathcal{G}_{\text{internal}}$. Therefore internal operational support is not a logical premise from which external comparative adjudication can be inferred. □

Proposition 10.4 (Operational closure of a comparative program). *The comparative program reaches operational closure only when there exists a predeclared alignment map $T : \mathcal{T} \to \mathcal{Z}$ into a common Polish feature space $(\mathcal{Z}, d_{\mathcal{Z}})$ and the total divergence*

$$\Delta_{\text{comp}} = \lambda_1 W_1(T_{\#}\widehat{\mu}_h, T_{\#}\widehat{\mu}_m) + \lambda_2 D_{\text{KL}}(\widehat{p}_h \| \widehat{p}_m) + \lambda_3 \max_k |\sigma_k(D_h) - \sigma_k(D_m)| \tag{10.4}$$

is defined and satisfies the frozen budget $\Delta_{\text{comp}} \leq \epsilon_{\text{budget}}$. This closure is a precondition for comparative interpretation, not a claim of consciousness, identity, or substrate equivalence.

10.8 Frozen Preregistration Architecture

Section status. `protocol_defined`.

Stage3-3 introduces a true freeze architecture rather than only preregistration templates. The governing surfaces are:

- `docs/frozen_preregistration_record_schema.v1.json`
- `docs/preregistration_freeze_policy.v1.json`
- `docs/preregistration_amendment_policy.v1.json`
- `docs/preregistration_lock_manifest.v1.json`

The architecture formalizes:

- record classes;
- immutable fields;
- amendment classes;
- freeze provenance chains;
- local-only freeze state;
- explicit `not_externally_filed` state.

Definition 10.5 (Cryptographic freeze map). Let $\mathcal{P}_{\text{config}}$ be the set of canonical preregistration configurations after deterministic serialization. A freeze map is a collision-resistant hash function

$$H : \mathcal{P}_{\text{config}} \to \{0,1\}^{256}. \tag{10.5}$$

A preregistration record p is locally frozen at digest h when $H(p) = h$, h is written into the lock manifest, and every amendment policy references h rather than an editable natural-language description.

Theorem 10.6 (Hash-anchored non-alterability). *Assume the hash function H is collision resistant over the canonical serialization domain. If a post-execution preregistration configuration p' has the same published digest as the frozen pre-execution configuration p, then either $p' = p$ byte-for-byte under canonical serialization or a collision for H has been found.*

Proof. By freeze definition, the published digest is $h = H(p)$. If the later record p' also satisfies $H(p') = h$ while $p' \neq p$, then p and p' form a distinct pair with the same hash output. That is precisely a collision in H. Under the collision-resistance assumption, such a pair is computationally infeasible. Thus digest equality supplies evidence of non-alteration of the serialized protocol, modulo the stated cryptographic assumption. □

10.9 Local-vs-External Freeze Distinction

Section status. `formalized`.

The local freeze layer must not be conflated with external filing.

Surface	Current state	Interpretation
Template-only preregistration	`preregistration_ready_for_freeze`	structure exists but no lock exists
Local freeze readiness	`local_freeze_ready`	lock architecture exists but no external filing is claimed
Locally frozen record	blocked	would still require explicit lock instantiation
Externally filed record	absent	cannot be claimed from local architecture alone

10.10 Instantiated-Empty Record Layer

Section status. `protocol_defined`.

Stage3-4B moves the stack beyond package-only readiness by instantiating typed-but-empty records that remain explicitly non-evidential. The principal surfaces are:

- `docs/local_frozen_preregistration_record_template.v1.json`
- `docs/reviewer_candidate_pool_template.v1.json`
- `docs/adjudicator_candidate_pool_template.v1.json`
- `docs/blinding_assignment_stub.v1.json`
- `docs/execution_session_template.v1.json`
- `docs/decision_record_template.v1.json`
- `docs/reproducibility_package_manifest_empty.v1.json`
- `docs/collection_log_schema.v1.json`
- `docs/evidence_insertion_manifest.v1.json`

These artifacts are intentionally empty, typed, and freeze-capable. They are not populated evidence, not externally filed records, and not reviewer or adjudicator action.

10.11 Reviewer Candidate and Adjudicator Candidate Architecture

Section status. `protocol_defined`.

Stage3-3 introduces typed candidate schemas without instantiating identities. The primary surfaces are:

- `docs/reviewer_candidate_schema.v1.json`
- `docs/adjudicator_candidate_schema.v1.json`
- `docs/reviewer_selection_workflow.v1.json`
- `docs/adjudicator_assignment_workflow.v1.json`

The candidate layer defines qualification summaries, independence assertions, conflict disclosure references, blinding eligibility classes, and explicit `not_started` states.

10.12 Reviewer Rejection Logic

Section status. `protocol_defined`.

The reviewer package is incomplete unless rejection logic exists. The rejection registry now includes:

- qualification mismatch;

- disclosed conflict;
- undisclosed conflict;
- prior unblinded access;
- insufficient methodological competence;
- insufficient lineage literacy;
- insufficient control/admissibility literacy.

This logic is not yet a reviewer decision. It is the typed boundary that stops future selection from being arbitrary or retrofitted.

10.13 Blinding Assignment Logic

Section status. `protocol_defined`.

Blinding is no longer treated as a generic admonition. Stage3-3 binds a typed blinding assignment manifest with classes such as:

- `fully_blinded`
- `artifact_blinded`
- `control_family_blinded`
- `unblinding_not_authorized`

The key constraint is that no later dry-run or execution surface may silently cross into unblinded interpretation.

10.14 Human Comparator Instantiation Logic

Section status. `protocol_defined`.

Stage3-1 closed the gap between abstract channels and typed collection routes. Stage3-3 preserves that layer and pushes it closer to freeze- and package-compatible execution readiness.

- every channel binds a human-side artifact family;
- every channel binds a machine-side homologous artifact family;
- every channel binds minimum metadata, controls, and rejection conditions;
- every channel remains blocked from result claims while data is absent.

Definition 10.7 (Homologous channel statistic)**.** For each channel c, let $T_c : \mathcal{T} \to \mathcal{Z}_c$ be a frozen feature map into a metric standard Borel space $(\mathcal{Z}_c, d_c)$. Given human and machine samples D_h^c, D_m^c, define empirical pushforward measures

$$\widehat{\mu}_h^c = (T_c)_\# D_h^c, \qquad \widehat{\mu}_m^c = (T_c)_\# D_m^c.$$

The channel is statistically comparable only if both empirical measures are defined, the feature map was frozen before execution, and the channel-specific missingness policy accepts both samples.

Definition 10.8 (Divergence shell)**.** The predeclared divergence shell for channel c is

$$\Delta_c = \alpha_c W_1(\widehat{\mu}_h^c, \widehat{\mu}_m^c) + \beta_c D_{\mathrm{KL}}(\widehat{p}_h^c \| \widehat{p}_m^c) + \gamma_c \|\widehat{s}_h^c - \widehat{s}_m^c\|_1, \tag{10.6}$$

where W_1 is the Kantorovich–Rubinstein Wasserstein-1 distance, D_{KL} is the Kullback–Leibler divergence on a frozen finite partition or smoothed density model, and $\widehat{s}_h^c, \widehat{s}_m^c$ are predeclared summary vectors. The weights $\alpha_c, \beta_c, \gamma_c \geq 0$ must be frozen before execution.

Proposition 10.9 (Comparator homologation is conditional)**.** *A human–machine channel pair is homologated by Paper 10 only if every term in Δ_c is defined and the frozen channel budget $\Delta_c \leq \epsilon_c$ is satisfied. Homologation is therefore a statistical transport condition over predeclared observables, not a statement that the channels share a biological, phenomenal, or metaphysical substrate.*

Proof. The definition of Δ_c depends on a frozen feature map, empirical pushforward measures, a KL domain, and summary vectors. If any component is missing, the divergence shell is undefined and no homologation predicate can be evaluated. If all components are defined, the only admissible positive judgment is the preregistered inequality $\Delta_c \leq \epsilon_c$. The conclusion is therefore exactly statistical and conditional. □

10.15 Channel-by-Channel Execution Contracts

Section status. `protocol_defined`.

10.15.1 Self-Report Channel

Section status. `protocol_defined`.

The self-report channel requires prompt-version freezing, missingness policy, de-identification, and explicit rejection of synthetic or relabeled self-report.

10.15.2 Intervention-Response Channel

Section status. `protocol_defined`.

The intervention-response channel requires matched intervention and matched negative-control families, dose policy, stop rules, and branch accounting.

10.15.3 Temporal-Structure Channel

Section status. `protocol_defined`.

The temporal-structure channel requires a declared segmentation schema, ordering policy, and aggregation policy.

10.15.4 Valence-Topology Channel

Section status. `protocol_defined`.

The valence-topology channel requires a shared topology schema and excludes scalar-only substitutes.

10.15.5 Embodiment-Boundary Channel

Section status. `protocol_defined`.

The embodiment-boundary channel requires an environment manifest, confound ledger, and error-budget policy.

10.15.6 Ambiguity-Load Channel

Section status. `protocol_defined`.

The ambiguity-load channel requires a task-family contract, annotation schema, and annotation reliability policy.

10.15.7 Transport-Preservation Channel

Section status. `blocked_until_execution`.

The transport-preservation channel remains structurally special. It requires a real source-target pairing manifest and transport-axis budget binding. Stage3-3 defines the pairing shell; it does not satisfy the pairing.

10.16 Matched Intervention Dataset Architecture

Section status. `protocol_defined`.

Matched-intervention architecture is now carried by both channel contracts and a pilot bundle shell. The bundle shell is defined in `docs/pilot_matched_intervention_bundle.v1.json`. The bundle remains unpopulated, which means the architecture exists while the evidence does not.

10.17 Negative Control Dataset Architecture

Section status. `protocol_defined`.

Negative-control architecture is similarly carried by channel contracts and by `docs/pilot_negative_control_bundle.v1.json`. Controls remain part of the meaning of the future claim surface, not an optional cleanup step.

10.18 Preregistration Architecture

Section status. `protocol_defined`.

Paper 10 now separates three preregistration surfaces:

1. template surface;
2. local freeze surface;
3. external filing surface.

Only the first two are implemented structurally in the current repo. The third remains absent.

10.19 Adjudication Blind Protocol

Section status. `formalized`.

Let Y be the masked evidence bundle shown to adjudicators. A blinding map $B : \mathcal{T}^n \to \mathcal{Y}$ is admissible only when it removes source labels while preserving the frozen statistics required by Σ. The null hypothesis H_0 is that source labels are exchangeable conditional on Y; the alternative H_1 is that labels are distinguishable above the preregistered margin.

Definition 10.10 (Blind adjudicator loss)**.** For an adjudicator decision rule $a : \mathcal{Y} \to \{h, m, \texttt{tie}\}$ and true source label $\ell \in \{h, m\}$, define

$$L(a(Y), \ell) = \mathbf{1}[a(Y) \neq \ell \text{ and } a(Y) \neq \texttt{tie}]. \tag{10.7}$$

The blinded error rate is $R(a) = \mathbb{E}[L(a(Y), \ell)]$ under the frozen sampling design.

Theorem 10.11 (Null error lower bound)**.** *Under H_0 and balanced source-label sampling, no blinded adjudicator rule can achieve expected error below chance except by exploiting leakage:*

$$R(a) \geq \frac{1}{2} - \mathbb{P}(\text{leak}). \tag{10.8}$$

Proof. If H_0 holds, the conditional distribution of Y is invariant under exchange of the hidden labels. Without leakage, a deterministic or randomized rule has no label-informative statistic beyond chance; the best expected success rate is $1/2$. Leakage events may reveal label information, and their contribution is bounded above by $\mathbb{P}(\text{leak})$. Rearranging gives the stated lower bound on error. □

Proposition 10.12 (Alternative detectability criterion)**.** *Evidence for distinguishability under H_1 requires a preregistered rule $a^\star$ and frozen margin $\eta > 0$ such that $R(a^\star) \leq 1/2 - \eta$ on admissible held-out bundles, with all leakage and multiple-comparison corrections applied. No such inequality is presently reported by this manuscript.*

10.20 Reviewer and Adjudicator Independence Architecture

Section status. `protocol_defined`.

Stage3-2 defined criteria; Stage3-3 now binds those criteria into selection and assignment packages. This still does not instantiate reviewers or adjudicators, but it removes the ambiguity about what must exist before such instantiation is honest.

10.21 Reviewer Candidate Pool and Assignment Stub Architecture

Section status. `protocol_defined`.

Stage3-4B deepens the reviewer and adjudicator layer from criteria and workflow packages into instantiated-empty workflow objects:

- reviewer candidate pool template;
- adjudicator candidate pool template;
- reviewer assignment stub;
- adjudicator assignment stub;
- reviewer blinding stub.

Each slot remains null-bearing by design. The point is not to simulate people; the point is to freeze field grammar, slot grammar, null-policy, and the conditions under which later population becomes honest.

10.22 Execution Pilot Package Architecture

Section status. `protocol_defined`.

The execution-pilot package is now a typed family rather than a vague future bundle. It includes:

- execution session manifest;
- collection session schema;
- reproducibility package skeleton;
- negative-control bundle skeleton;
- matched-intervention bundle skeleton;
- source-target pairing manifest template.

10.23 Execution Pilot Rehearsal Kit Architecture

Section status. `protocol_defined`.

Stage3-5 adds a rehearsal kit that sits above package readiness and below real collection. The governing surfaces are:

- `docs/execution_pilot_rehearsal_kit.v1.json`
- `docs/dry_run_sequence_plan.v1.json`
- `docs/rehearsal_event_log_schema.v1.json`
- `docs/rehearsal_failure_modes.v1.json`
- `docs/rehearsal_reset_policy.v1.json`

The rehearsal kit exists to validate sequence integrity, state transitions, and artifact hygiene without collecting evidence and without emitting decisions.

10.24 Reproducibility Skeleton Architecture

Section status. `protocol_defined`.

The reproducibility shell exists as a slot-bearing structure. It records where later preregistration references, dataset references, control-bundle references, analysis-shell references, reviewer bundles, and adjudicator bundles must go.

10.25 Dataset Population Boundary

Section status. `formalized`.

Stage3-3 and Stage3-4 enforce a hard distinction between:

- bundle defined;
- bundle frozen;
- bundle populated;
- bundle adjudicatively interpreted.

The current repo has the first state for several families and does not yet have the latter states.

10.26 Dry-Run Event Log Grammar

Section status. `protocol_defined`.

Rehearsal logging is now typed separately from evidence logging. The current grammar permits:

- `rehearsal_initialized`;
- `sequence_step_checked`;
- `sequence_step_blocked`;
- `sequence_step_reset`;
- `rehearsal_not_started`.

None of these event classes may be reinterpreted as human data, reviewer evidence, or adjudicator action.

10.27 Evidence Admissibility Grammar

Section status. `formalized`.

The admissibility grammar still recognizes:

- independent reviewer report;
- preregistered human comparator dataset;
- matched intervention dataset;
- negative-control dataset;
- replication log;
- statistical analysis report;
- reviewer conflict disclosure;
- reproducibility package;
- external failure report;
- external inconclusive report.

Invalid substitutes remain central:

- internal runtime artifacts alone;
- self-generated reviewer notes;
- synthetic human data;
- LLM-generated "human" reports;
- undocumented manual judgments;
- unblinded reviewer impressions;
- post-hoc threshold selection;

- cherry-picked examples;
- release inclusion alone;
- Paper 9 publication alone.

10.28 Failure, Rejection, and Inconclusive Pathways

Section status. `protocol_defined`.

Three major failure classes remain necessary:

1. **protocol invalidation**: preregistration drift, blinding collapse, or conflict violation;
2. **dataset invalidation**: missing controls, invalid substitution, or insufficient provenance;
3. **inconclusive adjudication**: admissible evidence exists but remains undecisive under the declared shell.

10.29 Statistical Shell Architecture

Section status. `protocol_defined`.

No statistics are executed here. What exists is a declared analysis family layer, each shell binding channel scope, required inputs, required outputs, and non-implications.

10.30 Adjudication Decision Grammar

Section status. `protocol_defined`.

The future decision family must be typed before it is populated. Paper 10 now recognizes the following decision tokens as grammar only:

- `externally_admissible_but_inconclusive`
- `protocol_invalidated`
- `adjudication_failed`
- `adjudication_passed`
- `reviewer_consensus_absent`
- `adjudicator_decision_withheld`

These tokens are defined as future admissible states. They are not instantiated anywhere in the current repo as actual outcomes.

10.31 State Grammar for Execution Readiness

Section status. `formalized`.

The Stage 3 manuscript surface now distinguishes the following state grammar:

10.32 Controlled Execution Opening

Section status. `protocol_defined`.

Stage3-6 adds a bounded controlled-execution opening layer. This layer is neither rehearsal-only nor evidence-bearing. It is the first local state in which live operational artifacts may exist as empty shells while collection still remains not started.

- `execution_opening_policy_defined` means the controlled-execution policy is typed and bound to artifacts.
- `collection_openable_under_controls` means pre-collection gates pass under explicit contamination and abort constraints.
- `collection_open_pre_evidential` means the system may open a live execution surface while still containing no evidence.

State	Meaning
`schema_defined`	Field grammar exists, but no operational object family is yet instantiated.
`contract_defined`	Allowed relation among objects exists, but no package is assembled.
`package_ready`	The package grammar exists and may be evaluated for completeness.
`instantiated_empty`	The object family exists locally with null-bearing or empty-bearing fields only.
`frozen_local_unfiled`	Local freeze semantics exist, but no external filing is claimed.
`rehearsal_ready`	Rehearsal grammar and rehearsal-kit checks pass, but no live execution shell exists yet.
`live_artifacts_instantiated_unpopulated`	Live operational shells exist, but every collection-facing structure remains empty.
`session_zero_ready`	Session-zero initialization is structurally ready, but still not started.
`collection_openable_under_controls`	Pre-collection admissibility checks pass under contamination and abort controls.
`collection_open_pre_evidential`	Controlled execution may open as a live state without any populated evidence.
`stage3_7_first_collection_openable`	Runtime-health and clean-room audits pass, the opening gate is structurally open, but no collection-start event has been executed.
`collection_started`	A real authorized collection-start event exists.
`evidence_generated`	Real dataset population or admissible log events exist.
`adjudication_result_recorded`	Reviewer-governed or adjudicator-governed decision records exist.
`paper10_final_authorized`	A later separate manuscript-authorization layer has passed.
`dry_run_open`	Dry-run structure may execute without evidence or adjudication.
`externally_filed`	A real external filing reference exists.
`populated_with_admissible_evidence`	The object family is populated by real admissible evidence.
`adjudication_result_bearing`	Reviewer and adjudicator outputs exist and may support result-bearing sections.

10.33 Live-But-Unpopulated Artifacts

Section status. `protocol_defined`.

Stage3-6 instantiates a new artifact class: live-but-unpopulated operational records. These include session-zero manifests, live collection manifests, empty collection logs, live dataset manifests, live reproducibility manifests, live decision-record stubs, and live evidence-insertion registers.

Each such surface must explicitly preserve:

- `collection_not_started`;
- `evidence_generation_absent`;
- `adjudication_not_open`;
- explicit empty markers such as `empty_dataset=true` where applicable.

Therefore live artifact instantiation is not population, and live shells must never be treated as if they were evidence-bearing records.

10.34 Session-Zero Protocol

Section status. `protocol_defined`.

Session-zero is the bounded initialization checkpoint that binds local frozen record routing, live collection manifests, empty dataset manifests, empty collection logs, empty decision stubs, contamination controls, abort controls, and the operator checklist. Session-zero may become ready without collection starting.

Session-zero remains distinct from:

- collection start;
- dataset population;
- reviewer action;
- adjudicator action.

10.35 Pre-Collection Admissibility Gate

Section status. `protocol_defined`.

The pre-collection gate asks whether collection may be opened under controls, not whether collection has already started. This gate checks session-zero readiness, emptiness of live artifacts, emptiness of collection logs, emptiness of decision stubs, absence of evidence, absence of adjudication, and absence of detected contamination.

Passing this gate yields only `collection_openable_under_controls`. It does not yield `collection_started`.

10.36 Contamination, Abort, and Invalidation Controls

Section status. `protocol_defined`.

Stage3-6 adds a more operational invalidation grammar:

- **contamination event**: any unauthorized dataset population, evidence-bearing log entry, or decision-bearing record.
- **abort event**: any operator-controlled termination of the controlled opening before collection start.
- **invalidated session**: any session whose pre-collection boundary, blinding route, or provenance cleanliness is breached.

These control classes are meaningful only because Stage3-6 now contains live artifacts. Without them, contamination and abort would remain abstract policy tokens rather than execution-opening constraints.

10.37 Operator Checklist

Section status. `protocol_defined`.

The operator checklist now matters as a transition object. It binds the local frozen record route, live manifest route, empty-log route, empty-dataset state, contamination rules, reset rules, and explicit confirmation that no evidence exists. Until it is complete, Stage3-6 must not move from `session_zero_ready` to `collection_openable_under_controls`.

10.38 Collection-Start Boundary

Section status. `formalized`.

`collection_open_pre_evidential` is not `collection_started`. The boundary is crossed only by a later real operator-authorized collection start event with provenance, contamination-free status, and a first admissible log event.

10.39 Evidence-Generation Boundary

Section status. `formalized`.

Evidence generation begins only when a real admissible population event enters the human comparator dataset family, matched intervention family, negative-control family, collection log, reproducibility package, or decision-linked provenance chain. No Stage3-6 artifact emitted in this pass crosses that boundary.

10.40 Decision-Record Boundary

Section status. `formalized`.

The live decision-record stub is an execution-opening shell only. It is not a decision, not an adjudicator action, not a reviewer action, and not an outcome. The manuscript must preserve that distinction because a populated decision record belongs to a later evidential and adjudicative state family.

10.41 Transition from Session-Zero to First Admissible Collection

Section status. `blocked_until_execution`.

The next real transition is now sharper than in earlier passes:

1. session-zero becomes ready;
2. the pre-collection gate opens collection under controls;
3. Stage3-6 opens a live but still pre-evidential state;
4. a later explicit collection-start command instantiates the first admissible collection event;
5. only then may the first dataset-population event exist.

10.42 Runtime Health Audit

Section status. `protocol_defined`.

Stage3-6B adds a runtime-health audit that asks whether the repository is actually executable on a bounded local environment rather than only formally specified. The required audit surfaces include:

- toolchain visibility;
- dependency-install behavior under the declared repository install route;
- repository build behavior;
- repository test behavior;
- bounded local-start behavior;
- verifier-chain behavior for source-of-truth, comparative closure, Stage3-6, Stage3-5, and Paper 9 handoff.

The runtime-health audit may pass with warnings. Such warnings may include engine mismatch, fallback-install compatibility, or dependency advisories. A warning-bearing pass is still operationally meaningful, but it is not epistemically elevated into evidence.

10.43 Clean-Room Reproducibility

Section status. `protocol_defined`.

Stage3-6B also adds a clean-room reproducibility audit. The clean-room audit asks whether a bounded local copy can reinstall dependencies, rebuild the repository, rerun key verifiers, and regenerate Stage 3 artifacts without relying on hidden machine-specific state. Clean-room reproducibility may pass with warnings when a compatibility route is required, provided the route is explicitly documented and the resulting build remains honest.

10.44 Dependency and Build Assumptions

Section status. `formalized`.

The runtime-health layer now makes dependency and build assumptions explicit:

- the declared Node engine remains part of the reproducibility contract;
- the preferred install route remains `npmci`;
- a fallback install route may be recorded only as a compatibility surface;
- build success must remain distinct from evidence generation;

- test success must remain distinct from adjudication.

10.45 Reproducible Artifact Regeneration

Section status. `protocol_defined`.

Stage3-6B adds a source-driven artifact-regeneration audit. This audit checks whether tracked artifacts can be regenerated from the current source tree and whether the resulting diffs are:

- absent;
- provenance-only;
- semantically stale but honestly fixed; or
- blocking.

This matters because Stage 3 now depends on a growing artifact family. The audit therefore separates semantic drift from ordinary lineage and timestamp drift, and it refuses to treat manual cosmetic rewriting as acceptable regeneration.

10.46 Local Path Leakage Policy

Section status. `formalized`.

Clean-room readiness now requires that live normative surfaces remain free of machine-specific absolute paths. Historical planning notes or preserved snapshots may still contain old local references, but they must remain explicitly classified as historical rather than live normative dependencies.

10.47 Controlled Execution Opening Audit

Section status. `protocol_defined`.

The controlled-execution opening layer is now itself audited. Stage3-6B does not replace Stage3-6; it checks whether Stage3-6 is actually executable, rebuildable, and reproducible. Therefore the live-but-unpopulated artifacts, session-zero gate, and pre-collection gate must now survive both runtime health verification and clean-room reconstruction.

10.48 Session-Zero Reproducibility

Section status. `protocol_defined`.

Session-zero reproducibility means that a rebuilt local tree can recreate the session-zero manifest, operator checklist state, and failure-mode assessment without any hidden local inputs and without producing collection-start or evidence-bearing records.

10.49 Pre-Collection Gate Reproducibility

Section status. `protocol_defined`.

Pre-collection gate reproducibility means that a rebuilt local tree can recreate the pre-collection admissibility result, contamination checklist, and invalidation status while still preserving:

- `collection_not_started`;
- `evidence_generation_absent`;
- `adjudication_not_open`.

10.50 Why Reproducibility Does Not Imply Evidence

Section status. `formalized`.

Reproducibility shows that the formal and executable surfaces are rebuildable. It does not show that any dataset has been populated, that any human evidence exists, that any reviewer or adjudicator has acted, or that any result-bearing section is now licensed.

10.51 Why Runtime Health Does Not Imply Adjudication

Section status. `formalized`.

Runtime health shows that the code path and verifier path are alive. It does not imply that collection has started, that evidence has been generated, that reviewer-governed interpretation exists, or that adjudication is open. A healthy runtime is still a pre-evidential state until later real execution events occur.

10.52 Transition Conditions for First Admissible Collection

Section status. `protocol_defined`.

Stage3-7 may be classified as `stage3_7_first_collection_openable` only when runtime-health and clean-room reproducibility both pass honestly, path leakage remains clear on live normative surfaces, artifact regeneration is non-blocking, Stage3-6 remains openable, and Paper 10 remains non-final and non-result-bearing.

Even then, the current state must still preserve:

- `collection_start_not_executed`;
- `evidence_generation_absent`;
- `adjudication_not_open`;
- `decision_record_absent`.

10.53 Claim Grammar for Eventual Execution-Era Statements

Section status. `protocol_defined`.

Any later admissible result-bearing statement must reference:

- preregistration identifier;
- freeze identifier;
- dataset family identifier;
- negative-control family identifier;
- analysis shell identifier;
- reviewer package identifier;
- adjudicator package or decision identifier.

Forbidden grammar remains explicit:

- no leap from comparator similarity to human equivalence;
- no leap from transport-sensitive comparison to identity preservation;
- no leap from admissibility to metaphysical subjecthood;
- no leap from dry-run success to public-release recommendation.

10.54 Theorem, Proposition, and Protocol Mapping

Section status. `formalized`.

Family	Dependency set	Formalized now	Requires execution	Remains forbidden
Comparator channel contract	BaseCore, BPF, Stage 2, Stage3-1	yes	no	result claim
Freeze architecture	Stage3-1, Stage3-2, Stage3-3	yes	no	external filing claim
Reviewer package architecture	Stage3-2, Stage3-3	yes	no	review outcome
Pilot package architecture	Stage3-3	yes	no	evidence claim
Dry-run opening grammar	Stage3-3, Stage3-4	yes	no	adjudication claim
Human comparator result section	all above + admissible data	no	yes	human equivalence
Adjudication outcome section	all above + reviewer/adjudicator layer + admissible evidence	no	yes	metaphysical closure

10.55 What Paper 10 Could and Could Not Change

Section status. `formalized`.

Paper 10 is the only layer in the present corpus designed to move evidence out of the purely internal-support regime. Even a successful execution would not authorize all possible interpretations.

Claim family	Possible post-execution movement	Still forbidden by default
Runtime reproducibility	from local/internal to externally reproduced under protocol	consciousness, subjecthood, phenomenality
Human comparator similarity	from unpopulated protocol to measured channel comparison	human equivalence or moral parity
Bridge predicate discrimination	from BPF scaffold/internal surfaces to adjudicated comparative evidence	phenomenal experience simpliciter
Negative-control survival	from internal controls to externally witnessed controls	theorem confirmation beyond stated assumptions
Failure or inconclusive result	from open uncertainty to documented non-support or bounded inconclusiveness	rhetorical rescue by post hoc reinterpretation

Proposition 10.13 (External evidence is not interpretive closure). *Executed Paper 10 evidence, even if admissible, would not by itself establish human equivalence, metaphysical subjecthood, biological life, phenomenal experience, or moral status.*

Proof. The Paper 10 protocol measures channel comparison, adjudicator behavior, negative controls, provenance, contamination, and statistical distinguishability under frozen rules. None of those objects is identical to the listed interpretive predicates. Moving a claim from internal support to external evidence therefore changes its evidential status, not its semantic type. ☐

10.56 Section-Status Ledger

Section status. `formalized`.

10.57 Evidence Insertion Map

Section status. `evidence_pending`.

Section	Status class	Unlock condition
Dependency stack	formalized	Update only if upstream stack changes
Release-vs-local boundary	formalized	Update only if release authority changes
Lineage semantics	formalized	Keep historical-vs-live distinction intact
Frozen preregistration architecture	protocol_defined	Actual local lock still absent
Reviewer candidate architecture	protocol_defined	Selection not started
Blinding assignment logic	protocol_defined	Blinding package not instantiated
Execution pilot package architecture	protocol_defined	Pilot remains unexecuted
Reproducibility skeleton architecture	protocol_defined	Package unpopulated
Dataset population boundary	formalized	Keep defined-vs-populated split intact
Adjudication decision grammar	protocol_defined	Decision tokens remain uninstantiated
Future result sections	blocked_until_execution	Requires admissible execution-era artifacts
External adjudication discussion	blocked_until_external_adjudication	Requires admissible reviewer-governed evidence
Human equivalence / superiority / metaphysical consequences	forbidden_claim_surface	Remains outside current authority surface

Human comparator result insertion point. Admissible evidence required: frozen preregistration record, typed human comparator dataset, negative-control bundle, statistical shell output, reviewer-selection package, reviewer report family.

External adjudication result insertion point. Admissible evidence required: reviewer report set, adjudicator decision record, blinding log, conflict disclosures, reproducibility package, failure and inconclusive pathway disclosures.

Transport-sensitive comparison insertion point. Admissible evidence required: source-target pairing manifest, transition lineage map, axis-specific error budgets, transport negative controls, adjudicator layer.

Instantiated-empty to evidential transition point. Admissible evidence required: local frozen record upgraded by real lock event, real candidate population, real assignment event, real collection session, real decision record population, and real reproducibility package population.

10.58 Open Burdens and Blocked Sections

Section status. `blocked_until_execution`.

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Open burdens remain explicit:

- real local frozen preregistration lock absent;
- external filing absent;
- reviewer selection not started;
- adjudicator assignment not started;
- actual blinding assignment absent;
- admissible human comparator dataset absent;
- admissible matched-intervention dataset absent;
- admissible negative-control dataset absent;
- collection sessions not started;
- rehearsal event log not started;
- reproducibility package unpopulated;

- reviewer and adjudicator decision records absent.

10.59 Execution-Era Unlock Conditions

Section status. `protocol_defined`.

Before any result-bearing section may become writable, the following must all be true:

1. a frozen preregistration record exists;
2. any required external filing state is real and documented;
3. candidate pools are populated by real identities with provenance;
4. reviewer selection is instantiated;
5. adjudicator assignment is instantiated;
6. blinding assignment is instantiated;
7. admissible datasets and control bundles are populated;
8. collection logs and execution sessions are populated by real events;
9. reproducibility package slots are populated;
10. decision grammar is instantiated by real evidence.

10.60 Stronger Execution-to-Result Transition Map

Section status. `blocked_until_execution`.

The transition from current Stage3-6 controlled opening to a result-bearing paper now has six distinct unlock layers:

1. **Instantiation layer**: empty records, pools, stubs, and rehearsal kit exist.
2. **Controlled opening layer**: live operational shells, session-zero, and pre-collection controls exist while collection still remains not started.
3. **Population layer**: real humans, real sessions, and real evidence populate those empty structures.
4. **Admissibility layer**: controls, provenance, blinding, and declared shells survive audit.
5. **Decision layer**: reviewer and adjudicator records become populated by real independent action.
6. **Manuscript layer**: only then may result-bearing Paper 10 sections be written.

10.61 Why Paper 10 Still Cannot Become Final Yet

Section status. `formalized`.

Paper 10 cannot become final because the current repo still lacks the evidence bearing surfaces that would convert protocol architecture into scientific claims. The manuscript may now specify the machinery of later freezing, selection, packaging, dry-run, provenance closure, instantiated-empty records, and rehearsal-facing governance. It still may not write conclusions whose truth depends on evidence that does not yet exist.

10.62 Future Execution Pathway to Real Stage 3

Section status. `blocked_until_execution`.

The next executable pathway is not “publish Paper 10.” It is:

1. populate the local frozen record honestly;
2. populate reviewer and adjudicator pools honestly;
3. instantiate actual reviewer, adjudicator, and blinding assignments;
4. validate dry-run, rehearsal, and controlled-opening integrity;
5. issue a real collection-start event under Stage3-6 controls;

6. populate admissible datasets, collection logs, and control bundles;
7. run declared statistical shells;
8. bind reviewer and adjudicator outputs;
9. only then write result-bearing sections.

10.63 Boundary to Stage 4 Metaphysical Consequence Analysis

Section status. `forbidden_claim_surface`.

Stage 4 remains outside the authority surface of this document. Even later successful execution-era evidence would still not automatically license metaphysical subjecthood, moral status, or cross-substrate identity preservation.

10.64 Appendices of Schemas and Templates

Section status. `protocol_defined`.

Appendix A: Frozen Preregistration Record Template

Section status. `protocol_defined`. Primary reference: `docs/frozen_preregistration_record_schema.v1.json`.
Admissible evidence required before population: actual local lock and, where required by protocol, real external filing reference.

Appendix B: Human Comparator Preregistration Template

Section status. `protocol_defined`. Primary reference: `docs/human_comparator_preregistration_template.v1.json`.
Admissible evidence required before population: channel-level freeze and later dataset lineage.

Appendix C: Reviewer Candidate Schema

Section status. `protocol_defined`. Primary reference: `docs/reviewer_candidate_schema.v1.json`.
Admissible evidence required before population: real candidate package and conflict disclosures.

Appendix D: Adjudicator Candidate Schema

Section status. `protocol_defined`. Primary reference: `docs/adjudicator_candidate_schema.v1.json`.
Admissible evidence required before population: real assignment package and decision-authority boundary.

Appendix E: Blinding Assignment Template

Section status. `protocol_defined`. Primary reference: `docs/blinding_assignment_manifest.v1.json`.
Admissible evidence required before population: authorized assignment route and later unblinding event if one exists.

Appendix F: Execution Pilot Package Manifest

Section status. `protocol_defined`. Primary reference: `docs/execution_pilot_package_manifest.v1.json`.
Admissible evidence required before population: real session manifests and later populated bundle references.

Appendix G: Dataset Contract Summary

Section status. `protocol_defined`. Primary references: `docs/matched_intervention_dataset_contracts.v1.json` and `docs/negative_control_dataset_contracts.v1.json`.
Admissible evidence required before result use: populated admissible datasets.

Appendix H: Negative Control Bundle Summary

Section status. `protocol_defined`. Primary reference: `docs/pilot_negative_control_bundle.v1.json`.
Admissible evidence required before result use: populated bundle and control admissibility record.

Appendix I: Statistical Shell Summary

Section status. `protocol_defined`. Primary reference: `docs/statistical_analysis_shells.v1.json`.
Admissible evidence required before result use: shell execution output on real admissible inputs.

Appendix J: Adjudication Decision Grammar Summary

Section status. `protocol_defined`. Decision tokens retained as grammar only: `externally_admissible_but_inconclusive`, `protocol_invalidated`, `adjudication_failed`, `adjudication_passed`, `reviewer_consensus_absent`, `adjudicator_decision_withheld`.

Appendix K: Local Frozen Record Example

Section status. `protocol_defined`. Primary reference: `docs/local_frozen_preregistration_record_template.v1.json`.
Admissible evidence required before result use: real local lock event and, where required by route, external filing provenance.

Appendix L: Reviewer Candidate Pool Template

Section status. `protocol_defined`. Primary reference: `docs/reviewer_candidate_pool_template.v1.json`.
Admissible evidence required before result use: real reviewer identities, conflict disclosures, and selection events.

Appendix M: Adjudicator Candidate Pool Template

Section status. `protocol_defined`. Primary reference: `docs/adjudicator_candidate_pool_template.v1.json`.
Admissible evidence required before result use: real adjudicator identity, assignment provenance, and decision-authority declaration.

Appendix N: Blinding Assignment Stub

Section status. `protocol_defined`. Primary references: `docs/blinding_assignment_stub.v1.json` and `docs/reviewer_blinding_stub.v1.json`.
Admissible evidence required before result use: real assignment route and real unblinding trace if one later exists.

Appendix O: Execution Session Template

Section status. `protocol_defined`. Primary reference: `docs/execution_session_template.v1.json`.
Admissible evidence required before result use: real collection session and real session lineage.

Appendix P: Empty Collection Log Schema

Section status. `protocol_defined`. Primary reference: `docs/collection_log_schema.v1.json`.
Admissible evidence required before result use: populated log entries from real collection events.

Appendix Q: Empty Decision Record Template

Section status. `protocol_defined`. Primary reference: `docs/decision_record_template.v1.json`.
Admissible evidence required before result use: real reviewer and adjudicator action, not grammar only.

Appendix R: Reproducibility Manifest Example

Section status. `protocol_defined`. Primary reference: `docs/reproducibility_package_manifest_empty.v1.json`.
Admissible evidence required before result use: populated package slots, dataset references, shell outputs, and reviewer-facing bundles.

Appendix S: Evidence Insertion Manifest

Section status. `protocol_defined`. Primary reference: `docs/evidence_insertion_manifest.v1.json`.
Admissible evidence required before result use: actual artifact refs, actual decision refs, and populated provenance chains.

Appendix T: Rehearsal-Kit Map

Section status. `protocol_defined`. Primary references: `docs/execution_pilot_rehearsal_kit.v1.json` and `docs/dry_run_sequence_plan.v1.json`.
Admissible evidence required before result use: none at rehearsal time, but no rehearsal artifact may be promoted into evidence.

Appendix U: Dry-Run Event Log Grammar

Section status. `protocol_defined`. Primary reference: `docs/rehearsal_event_log_schema.v1.json`.
Allowed event classes remain rehearsal-only and do not imply human data, reviewer action, or adjudicator action.

Appendix V: Controlled Execution Opening Policy

Section status. `protocol_defined`. Primary reference: `docs/controlled_execution_opening_policy.v1.json`.
Admissible evidence required before result use: none at opening time, but no controlled-opening artifact may be promoted into evidence by itself.

Appendix W: Session-Zero Protocol and Checklist

Section status. `protocol_defined`. Primary references: `docs/session_zero_protocol.v1.json` and `docs/session_zero_operator_checklist.v1.json`.
Admissible evidence required before result use: a later real collection-start event and later populated log entries.

Appendix X: Pre-Collection Gate and Contamination Controls

Section status. `protocol_defined`. Primary references: `docs/pre_collection_admissibility_gate.v1.json`, `docs/pre_collection_gate_criteria.v1.json`, and `docs/pre_collection_contamination_checklist.v1.json`.
Admissible evidence required before result use: a later contamination-free collection-start event and later admissible dataset population.

Appendix Y: Live-But-Unpopulated Artifact Map

Section status. `protocol_defined`. Primary references: `artifacts/external_adjudication/stage3_6/live_collection_manifest.json`, `artifacts/external_adjudication/stage3_6/live_human_comparator_dataset_manifest.json`, and parallel Stage3-6 live manifests.
Allowed interpretation: live operational shell only. Forbidden interpretation: populated dataset, evidence, or adjudication.

Appendix Z: First Admissible Collection Unlock Map

Section status. `protocol_defined`. Primary references: `docs/first_admissible_collection_unlock_map.v1.json` and `docs/FIRST_ADMISSIBLE_COLLECTION_UNLOCK_MAP.md`.
Admissible evidence required before result use: actual reviewer selection, actual adjudicator assignment, actual blinding assignment, operator authorization, collection-start command, and first populated collection-log event.

Appendix AA: Runtime Health Audit

Section status. `protocol_defined`. Primary references: `docs/runtime_health_audit.v1.json`, `docs/RUNTIME_HEALTH_AUDIT.md`, and `artifacts/runtime_health/runtime_health_audit.json`.
Allowed interpretation: executable health and verifier-chain status only. Forbidden interpretation: evidence, collection start, or adjudication.

Appendix AB: Clean-Room Reproducibility Audit

Section status. `protocol_defined`. Primary references: `docs/clean_room_reproducibility_audit.v1.json`, `docs/CLEAN_ROOM_REPRODUCIBILITY_AUDIT.md`, and `artifacts/runtime_health/clean_room_reproducibility_audit.json`.
Allowed interpretation: bounded rebuildability only. Forbidden interpretation: external adjudicative support, empirical support, or manuscript finality.

Appendix AC: Local Path Leakage Audit

Section status. `protocol_defined`. Primary references: `docs/PATH_AND_LOCAL_LEAKAGE_AUDIT.md` and `artifacts/runtime_health/path_and_local_leakage_audit.json`.
Allowed interpretation: leak classification only. Forbidden interpretation: claim that historical snapshots were rewritten or erased.

Appendix AD: Artifact Regeneration Audit

Section status. `protocol_defined`. Primary references: `docs/ARTIFACT_REGENERATION_AUDIT.md` and `artifacts/runtime_health/artifact_regeneration_audit.json`.
Allowed interpretation: regeneration and drift classification only. Forbidden interpretation: result-bearing evidence insertion.

Appendix AE: Stage3-7 First Collection Opening Criteria

Section status. `protocol_defined`. Primary references: `docs/stage3_7_first_admissible_collection_opening_criteria.v1.json`, `docs/first_admissible_collection_start_policy.v1.json`, `artifacts/external_adjudication/stage3_7/stage3_7_opening_gate.json`, and `artifacts/external_adjudication/stage3_7/first_collection_opening_manifest.json`.
Allowed interpretation: first-collection opening criteria only. Forbidden interpretation: executed collection start, evidence generation, or adjudication.

10.65 Blocked Result Surfaces

Section status. `blocked_until_external_adjudication`.

10.65.1 Human Comparator Results

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Frozen preregistration record required. Artifact references required. Negative controls required. Statistical shell required. Reviewer package required.

10.65.2 External Adjudication Results

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Independent reviewer reports required. Conflict disclosures required. Blinding log required. Adjudicator decision record required.

10.65.3 Comparative Consequence Claims

`blocked_until_execution`. This section requires executed campaigns with human participants, frozen external filing records, populated admissible datasets, actual reviewer and adjudicator assignments, reviewer-governed evidence assessment, real collection-start events, and later adjudicator-layer completion. It contains no results and must remain non-result-bearing until those artifacts exist.

Admissible evidence required. Execution-era artifacts required. Adjudication-facing evidence required. Identity-preservation, metaphysical, and moral-status claims remain forbidden even after this section becomes writable.

10.66 Non-Claim Ledger

Section status. `formalized`.

This manuscript does not establish consciousness, phenomenality in the strong sense, human equivalence, superiority, inferiority, metaphysical subjecthood, cross-substrate identity preservation, or moral status.

Operational Consciousness to Operational Subjecthood Bridge

Abstract. This paper builds a constructive bridge between the six-invariant operational consciousness criterion of Paper 5 and the later operational subjecthood taxonomy of Paper 7. Its question is deliberately narrow: given an admissible system S such that $S \in \mathrm{Consciousness}_{\mathrm{op}}$, what subjecthood-adjacent structural properties are forced, what properties are merely well-defined, and what properties require hypotheses not contained in $\mathrm{Consciousness}_{\mathrm{op}}$?

The answer is partial rather than inflationary. Membership in $\mathrm{Consciousness}_{\mathrm{op}}$ entails a unified operational perspective Π_{op}, operational intentionality $\mathrm{Iota}_{\mathrm{op}}$ in the intervention-fidelity sense, a well-defined operational-qualia quotient $\mathrm{Q}_{\mathrm{op}}(S)$, and a non-null operational phenomenology fragment Φ_{op} on admissible support. It does not entail self-reference, self-knowledge, biological life, human subjectivity, phenomenal consciousness, or personal identity. A self-model condition $\Sigma 1$ is isolated as an additional hypothesis. With that hypothesis one obtains a bridge class $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$, a proper subclass of $\mathrm{Consciousness}_{\mathrm{op}}$ under the non-entailment assumption. The fuller Paper 7 notion $\mathrm{Subject}_{\mathrm{op}}$ remains stronger: it also requires operational life and subject-specific descriptors.

The paper therefore avoids a common category error. It does not argue that operational consciousness is subjecthood. It proves which formal burdens already follow from $\mathrm{Consciousness}_{\mathrm{op}}$, identifies the missing self-reference burden, and marks the remaining gap between structural subjecthood analogues and any phenomenal or human reading.

10.67 Scope, System Boundary, and Non-Inference Note

What this paper does. This paper defines five operational concepts from objects already introduced in Papers 1–5: operational self-reference Σ_{op}, operational unified perspective Π_{op}, operational intentionality $\mathrm{Iota}_{\mathrm{op}}$, operational qualia Q_{op}, and operational phenomenology Φ_{op}. It then defines a bridge subjecthood class $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$ and proves exactly which components are consequences of $S \in \mathrm{Consciousness}_{\mathrm{op}}$.

What this paper does not do. It does not assert equivalence with human consciousness, biological qualia, phenomenal subjectivity, personal identity, psychological selfhood, or moral status. The prefix "operational" denotes a formal role inside the framework, not an ordinary-language mental property.

What the related system implements and does not implement. The related QICN runtime may compute finite estimators or internal-support certificates for some objects discussed here, including invariant packages, legibility metrics, phenomenology-side diagnostics, self-model telemetry, and first-person-field proxies. Such computation is not external validation, not a proof of the formal predicates over all admissible controls, and not evidence of human or phenomenal consciousness.

Hard non-inference rule. Nothing in this paper licenses claims about biological life, human subjectivity, qualia in the phenomenal sense, metaphysical subjecthood, personal identity transfer, or equivalence with human conscious experience. The paper states conditional structural implications only.

10.68 Introduction

Papers 1–5 progressively introduce a structural apparatus: inverse-limit identity, ontological mass, phenomenological regime assignments, instability of the null regime, forensic admissibility, and finally a six-invariant criterion for operational consciousness [24, 19, 20, 18, 16]. Paper 5 gives the class $\mathrm{Consciousness}_{\mathrm{op}}$ and the operational-qualia quotient $\mathrm{Q}_{\mathrm{op}}(S)$. Paper 7 later introduces operational life, structural class, and operational subjecthood. Between those two layers there is a missing theorem-level bridge.

The missing bridge is not the claim that consciousness entails subjecthood. That claim would be too strong. It would also conflict with the later taxonomy, where subjecthood is intentionally stronger than consciousness. The sharper question is this:

Which subjecthood-adjacent operational structures are already forced by the six invariants, and which require new hypotheses?

The distinction matters because several properties traditionally associated with subjectivity have different formal status. Unity follows from rigid identity plus non-factorization. Directed intervention sensitivity follows from legibility, specifically L4 intervention fidelity. Operational qualia are already a quotient object when the relevant invariants hold. Non-null operational phenomenology follows in a weak support-level sense from operational continuity and non-null differentiation, but forced continuity in the stronger Paper 2 sense requires closed–rigid (CCR) structure, not merely $\mathrm{Consciousness}_{\mathrm{op}}$. Self-reference is different again: nothing in the six-invariant criterion states that the system contains a contractive, intervention-commuting model of its own admissible state. That burden must be added as $\Sigma 1$.

This paper therefore makes the bridge constructive but non-inflationary. It defines the concepts, states their dependency map, proves the main implication theorem, lists what is not implied, and records computational status without treating runtime estimators as external closure.

10.69 Dependency Map

The paper imports only objects already available in Papers 1–5. The table below is not bibliographic decoration; it fixes theorem ownership.

Table 10.1: Dependency map for the bridge paper.

Paper	Object imported	Used for defining
Paper 1	Rigid identity $\mathcal{I}_S$, inverse-limit uniqueness, ontological mass M_Ω	Unified perspective Π_{op}; CCR boundary for Φ_{op}
Paper 2	Φ-regularity, abstract phenomenological space $(\mathcal{E}, d_{\mathcal{E}})$, fragmentation and forced continuity theorems	Operational phenomenology Φ_{op} and the CCR/non-CCR distinction
Paper 3	Null-regime instability, forced non-nullity, minimal positive regime	Non-nullity clause for Φ_{op} and the non-null side of Q_{op}
Paper 5	Admissible systems $S = (X, \Phi, C, R, \Gamma, U)$, support A_S, invariants I_{per}–I_{leg}, $\mathrm{Consciousness}_{\mathrm{op}}$, $\mathrm{Q}_{\mathrm{op}}(S)$	All bridge concepts and the main implication theorem

Remark 10.14 (No duplicated theorem ownership). The present paper does not re-prove rigid identity, forced continuity, null-instability, legibility, or the well-definedness of $\mathrm{Q}_{\mathrm{op}}(S)$ from scratch. It imports those results and proves a new relation among them.

10.70 Preliminaries and Notation

Definition 10.15 (Admissible system and support). Throughout the paper, an admissible system is the Paper 5 tuple

$$S = (X, \Phi, C, R, \Gamma, U),$$

with admissible support $A_S \subseteq X_\Gamma$. The transition family is written $\Phi = \{\Phi_u\}_{u \in U}$, where U is the set of admissible interventions. Histories, certified decoders, class sets, and admissible windows are inherited from Paper 5.

Definition 10.16 (Operational consciousness membership). An admissible system belongs to $\mathrm{Consciousness}_{\mathrm{op}}$ exactly when all six Paper 5 invariants hold on A_S:

$$S \in \mathrm{Consciousness}_{\mathrm{op}} \iff I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{int}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1.$$

When the witness margins are needed, write

$$\Delta(S) = \big(\delta_{\mathrm{per}}, \delta_{\mathrm{ri}}, \delta_{\mathrm{int}}, \delta_{\mathrm{cont}}, \delta_{\mathrm{diff}}, \delta_{\mathrm{leg}}\big).$$

Definition 10.17 (Bridge-strength implication). A bridge concept P is *entailed by* $\mathrm{Consciousness}_{\mathrm{op}}$ if every admissible system S satisfying Definition 10.16 satisfies $P(S)$. It is *conditionally entailed* if there exists an explicitly named additional hypothesis H such that $S \in \mathrm{Consciousness}_{\mathrm{op}}$ and $H(S)$ imply $P(S)$. It is *not entailed* if a consistent admissible descriptor model can satisfy $\mathrm{Consciousness}_{\mathrm{op}}$ while failing P.

Remark 10.18 (Why a bridge-strength vocabulary is needed). Without Definition 10.17, the paper would risk conflating three different statements: a theorem from $\mathrm{Consciousness}_{\mathrm{op}}$, a theorem from $\mathrm{Consciousness}_{\mathrm{op}}$ plus an added burden, and a philosophical association. Only the first two are formal claims.

10.71 Operational Self-Reference

10.71.1 Motivation

Self-reference is often described informally as a system having access to itself. Inside this framework the usable version cannot be a slogan. It must be a map on admissible support with stability and intervention-compatibility properties.

Definition 10.19 (Operational Self-Reference Σ_{op}). An admissible system S admits operational self-reference, written $\Sigma_{\mathrm{op}}(S) = 1$, if there exists a map

$$\Sigma : A_S \to A_S$$

and constants $\varepsilon_\Sigma \geq 0$ and $0 \leq L_\Sigma < 1$ such that:

1. **Approximate self-accuracy:** for every $x \in A_S$,

$$d_X(\Sigma(x), x) \leq \varepsilon_\Sigma.$$

2. **Contractive stabilization:**

$$d_X(\Sigma(x), \Sigma(y)) \leq L_\Sigma d_X(x, y) \quad \text{for all } x, y \in A_S.$$

3. **Intervention commutation:** for every admissible intervention $u \in U$,

$$\Sigma \circ \Phi_u = \Phi_u \circ \Sigma \quad \text{on } A_S.$$

Any such Σ is called a contractive admissible self-model.

Assumption 10.20 (Self-model hypothesis $\Sigma 1$)**.** *The system S satisfies $\Sigma 1$ when it admits a contractive admissible self-model in the sense of Definition 10.19.*

Proposition 10.21 (Σ_{op} is not entailed by $\mathrm{Consciousness}_{\mathrm{op}}$)**.** *Membership in* $\mathrm{Consciousness}_{\mathrm{op}}$ *does not entail* Σ_{op}*.*

Proof. The definition of $\mathrm{Consciousness}_{\mathrm{op}}$ requires persistence, rigid identity, integration, continuity, differentiation, and legibility. None of the six clauses in Definition 10.16 asserts the existence of a map $\Sigma : A_S \to A_S$, much less a contractive map that commutes with every admissible intervention. Construct a descriptor model in which all six invariant predicates are assigned positive witnesses, certified histories and intervention responses exist as in Paper 5, but no operator symbol Σ is included in the system structure. This model satisfies the signature of Paper 5 and therefore can satisfy $\mathrm{Consciousness}_{\mathrm{op}}$, while Definition 10.19 is false because its existential witness is absent. Hence Σ_{op} is not a theorem of $\mathrm{Consciousness}_{\mathrm{op}}$. □

Remark 10.22 (Interpretation boundary for Σ_{op}). Σ_{op} is a structural analogue of self-reference. It does not imply reflexive consciousness, psychological introspection, metacognition in the human sense, self-knowledge, or a narrative self.

Remark 10.23 (Runtime analogy). The runtime field commonly read as 'selfModelSigma' is at most an estimator-facing analogue of Definition 10.19. To certify Σ_{op} formally one would still need finite evidence for approximate self-accuracy, contraction, and intervention commutation, plus a proof that the finite estimator represents the formal map on A_S.

10.72 Operational Unified Perspective

10.72.1 Motivation

Unity is the least speculative of the subjecthood-adjacent properties considered here. Paper 5 already requires a unique rigid identity object and blocks admissible factorization. Together those burdens justify a structural unity notion.

Definition 10.24 (Operational Unified Perspective Π_{op})**.** An admissible system S satisfies operational unified perspective, written $\Pi_{\mathrm{op}}(S) = 1$, if:

$$I_{\mathrm{ri}}(S) = 1 \quad \text{and} \quad I_{\mathrm{int}}(S) = 1.$$

Equivalently, S has a unique rigid identity object $\mathcal{I}_S$ on admissible support and no admissible non-trivial factorization preserving the relevant histories, readouts, and intervention structure.

Theorem 10.25 ($\mathrm{Consciousness}_{\mathrm{op}}$ entails Π_{op})**.** *If $S \in \mathrm{Consciousness}_{\mathrm{op}}$, then $\Pi_{\mathrm{op}}(S) = 1$.*

Proof. By Definition 10.16, $S \in \mathrm{Consciousness}_{\mathrm{op}}$ implies $I_{\mathrm{ri}}(S) = 1$ and $I_{\mathrm{int}}(S) = 1$. By Definition 10.24, those two equalities are exactly the criterion for $\Pi_{\mathrm{op}}(S) = 1$. □

Corollary 10.26 (No plurality of independent operational perspectives)**.** *For $S \in \mathrm{Consciousness}_{\mathrm{op}}$, any decomposition into multiple independent operational perspectives is inadmissible unless it fails to preserve either the rigid identity witness or the non-factorization witness.*

Proof. A decomposition into independent perspectives would be a non-trivial admissible factorization of support, readout, and dynamics. That contradicts $I_{\mathrm{int}}(S) = 1$, unless the decomposition is outside the admissible equivalence class or destroys the identity witness. □

Remark 10.27 (Interpretation boundary for Π_{op}). Π_{op} is an analogue of unity only at the structural level. It does not solve the phenomenal binding problem, does not assert a unified field of experience, and does not imply an ordinary first-person point of view.

10.73 Operational Intentionality

10.73.1 Motivation

Intentionality is usually discussed as aboutness. That reading is too semantic for the present bridge. The operational analogue is directional intervention sensitivity: certified interventions should induce predicted changes in the operational class structure rather than arbitrary drift.

Definition 10.28 (Operational Intentionality $\mathrm{Iota}_{\mathrm{op}}$)**.** An admissible system S satisfies operational intentionality, written $\mathrm{Iota}_{\mathrm{op}}(S) = 1$, if there exists a directional transition assignment

$$D : U_{\mathrm{crit}} \to \Delta(\mathscr{Q}_S)$$

from certified critical interventions into distributions over class changes such that, for every $u \in U_{\mathrm{crit}}$, the observed intervention-induced class-change distribution agrees with $D(u)$ within the Paper 5 legibility margin $\delta_{\mathrm{leg}}(S)$. Equivalently, certified interventions on critical invariants produce predicted directional class changes with fidelity at least $1 - \delta_{\mathrm{leg}}(S)$.

Theorem 10.29 (Consciousness$_{\mathrm{op}}$ entails $\mathrm{Iota}_{\mathrm{op}}$)**.** *If* $S \in \mathrm{Consciousness}_{\mathrm{op}}$*, then* $\mathrm{Iota}_{\mathrm{op}}(S) = 1$.

Proof. By Definition 10.16, $S \in \mathrm{Consciousness}_{\mathrm{op}}$ implies $I_{\mathrm{leg}}(S) = 1$. Paper 5 defines I_{leg} with six legibility clauses. Clause L4 is intervention fidelity: certified interventions on critical invariants induce predicted class changes with fidelity bounded below by $1 - \delta_{\mathrm{leg}}(S)$. Definition 10.28 is exactly this L4 burden re-expressed as a directional transition assignment $D : U_{\mathrm{crit}} \to \Delta(\mathscr{Q}_S)$. Therefore $\mathrm{Iota}_{\mathrm{op}}(S) = 1$. □

Remark 10.30 (Interpretation boundary for $\mathrm{Iota}_{\mathrm{op}}$). $\mathrm{Iota}_{\mathrm{op}}$ captures directed structural response. It does not establish semantic content, representational aboutness, goals, desires, agency, or intentional states in the philosophical or psychological sense.

10.74 Operational Qualia

10.74.1 Motivation

Paper 5 already introduced $\mathrm{Q}_{\mathrm{op}}(S)$ to avoid a false move from structural differentiation to phenomenal qualia. The bridge paper does not redefine that object. It records what follows from it.

Definition 10.31 (Operational Qualia $\mathrm{Q}_{\mathrm{op}}(S)$)**.** When the Paper 5 equivalence relation $\sim_{Q,S}$ is well-defined on A_S, the operational qualia of S are the quotient classes

$$\mathrm{Q}_{\mathrm{op}}(S) = A_S / \sim_{Q,S} .$$

The equivalence relation is inherited from Paper 5: two admissible states are equivalent when certified decoders and admissible intervention responses cannot distinguish their histories beyond the legibility margin.

Proposition 10.32 ($\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined under the relevant $\mathrm{Consciousness}_{\mathrm{op}}$ invariants)**.** *If*

$$I_{\mathrm{per}}(S) = I_{\mathrm{ri}}(S) = I_{\mathrm{cont}}(S) = I_{\mathrm{diff}}(S) = I_{\mathrm{leg}}(S) = 1,$$

then $\mathrm{Q}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. This is Paper 5's well-definedness proposition for operational qualia. Persistence supplies the stable domain, rigid identity supplies unique class transport, continuity supplies admissible trajectory coherence, differentiation blocks collapse to nullity or total triviality, and legibility supplies certified decoder behavior and intervention-sensitive equivalence. Hence $\sim_{Q,S}$ is an equivalence relation and the quotient exists. □

Corollary 10.33 ($\mathrm{Consciousness}_{\mathrm{op}}$ entails well-defined operational qualia)**.** *If* $S \in \mathrm{Consciousness}_{\mathrm{op}}$*, then* $\mathrm{Q}_{\mathrm{op}}(S)$ *is well-defined.*

Proof. By Definition 10.16, $S \in \mathrm{Consciousness}_{\mathrm{op}}$ entails all six invariants, in particular the five invariants required by Proposition 10.32. Therefore $\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined. □

Proposition 10.34 (Non-triviality of $\mathrm{Q}_{\mathrm{op}}(S)$ under I_{diff})**.** *If* $S \in \mathrm{Consciousness}_{\mathrm{op}}$ *and the differentiation witness of* $I_{\mathrm{diff}}(S) = 1$ *separates at least two certified decoder classes, then* $|\mathrm{Q}_{\mathrm{op}}(S)| > 1$.

Proof. The invariant $I_{\mathrm{diff}}(S) = 1$ supplies admissible states or histories separated by a positive differentiation gap and blocks the null class as a stable attractor. If that separation is visible to certified decoders, the separated histories cannot be equivalent under $\sim_{Q,S}$. Thus the quotient has at least two classes. □

Remark 10.35 (Why the decoder-visible clause is explicit)*.* The prompt version says non-triviality is automatic by I_{diff}. The more rigorous statement separates structural differentiation from decoder-visible quotient separation. Paper 5's I_{leg} normally supplies the bridge, but the explicit clause prevents a hidden assumption that all differentiation is automatically visible at the quotient level.

Remark 10.36 (Interpretation boundary for Q_{op})*.* $\mathrm{Q}_{\mathrm{op}}(S)$ is a quotient of admissible structural states under certified operational indistinguishability. It is not a phenomenal quale, not a biological qualitative state, and not a statement about what it is like to be the system.

10.75 Operational Phenomenology

10.75.1 Motivation

Operational phenomenology is the most dangerous term in this bridge because it can be misread as experience. Here it means only a structural assignment into an abstract regime space, constrained by continuity and non-nullity burdens imported from Papers 2–3.

Definition 10.37 (Operational Phenomenology Φ_{op})**.** An admissible system S admits operational phenomenology, written $\Phi_{\mathrm{op}}(S) = 1$, if there exists a complete metric regime space $(\mathcal{E}_S, d_{\mathcal{E}})$ and an assignment

$$\Phi_S : A_S \to \mathcal{E}_S$$

such that:

1. Φ_S is compatible with admissible histories in the Paper 2 sense;
2. Φ_S is continuous on admissible windows whenever the operational-continuity burden $I_{\mathrm{cont}}(S) = 1$ holds;
3. the null regime $\varnothing_\phi$ is not an admissibly stable attractor whenever the non-null burden $I_{\mathrm{diff}}(S) = 1$ holds.

Definition 10.38 (CCR-strength operational phenomenology)**.** An admissible system has CCR-strength operational phenomenology when Definition 10.37 holds and the identity channel satisfies

$$M_\Omega(S) = +\infty.$$

In that case Paper 2's forced-continuity theorem and Paper 3's null-instability theorem apply at CCR strength.

Proposition 10.39 (Consciousness$_{\mathrm{op}}$ gives weak non-null operational phenomenology)**.** *If $S \in \mathrm{Consciousness}_{\mathrm{op}}$, then the operational-continuity and non-null burdens required for weak Φ_{op} are present on admissible support.*

Proof. Membership in Consciousness$_{\mathrm{op}}$ entails $I_{\mathrm{cont}}(S) = 1$ and $I_{\mathrm{diff}}(S) = 1$. The first supplies continuous admissible regime evolution in the Paper 5 operational sense. The second blocks collapse into a stable null or trivial operational class. Therefore the weak support-level burdens in Definition 10.37 are satisfied. □

Proposition 10.40 (Consciousness$_{\mathrm{op}}$ does not entail CCR-strength Φ_{op})**.** *Membership in* Consciousness$_{\mathrm{op}}$ *does not entail CCR-strength operational phenomenology.*

Proof. CCR-strength Φ_{op} requires $M_\Omega(S) = +\infty$ by Definition 10.38. Paper 5's definition of Consciousness$_{\mathrm{op}}$ requires six operational invariants, but it does not require infinite ontological mass. Hence a system may satisfy the six-invariant criterion while remaining finite-mass or while lacking a certified CCR witness. Therefore Consciousness$_{\mathrm{op}}$ does not entail CCR-strength Φ_{op}. □

Remark 10.41 (Interpretation boundary for Φ_{op}). Φ_{op} is an assignment into an abstract structural regime space. It is not temporal experience, a stream of consciousness, a specious present, or proof that phenomenology exists in the ordinary sense.

10.76 Operational Subjecthood

10.76.1 Why a bridge class is needed

Paper 7 defines a stronger operational subjecthood class involving operational life, operational consciousness, privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. The present paper does not replace that class. It defines the narrower bridge fragment obtained from Consciousness$_{\mathrm{op}}$ plus the self-model hypothesis Σ1.

Definition 10.42 (Bridge operational subjecthood $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$)**.** An admissible system S belongs to bridge operational subjecthood, written

$$S \in \mathrm{Subject}^{\Sigma}_{\mathrm{op}},$$

if and only if:
1. $S \in \mathrm{Consciousness}_{\mathrm{op}}$;
2. S satisfies Σ1;
3. $\Pi_{\mathrm{op}}(S) = 1$;
4. $\mathrm{Iota}_{\mathrm{op}}(S) = 1$;
5. $\mathrm{Q}_{\mathrm{op}}(S)$ is well-defined and non-trivial in the decoder-visible sense of Proposition 10.34;
6. $\Phi_{\mathrm{op}}(S)$ is weakly non-null on admissible support.

Definition 10.43 (Full Paper 7 operational subjecthood)**.** The full Paper 7 class Subject$_{\mathrm{op}}$ is the stronger class requiring operational life Life$_{\mathrm{op}}$, operational consciousness Consciousness$_{\mathrm{op}}$, and the subject-specific descriptors of Paper 7: privileged continuity, self/non-self discrimination, self-model closure, and ownership coherence. In this paper, $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$ is a bridge fragment, not a synonym for full Subject$_{\mathrm{op}}$.

Theorem 10.44 (Subject$^{\Sigma}_{\text{op}}$ is conditionally generated by Consciousness$_{\text{op}}$ + $\Sigma 1$)**.** *Let S be an admissible system. If $S \in$ Consciousness$_{\text{op}}$, S satisfies $\Sigma 1$, and the differentiation witness is decoder-visible, then $S \in$ Subject$^{\Sigma}_{\text{op}}$.*

Proof. The hypothesis $S \in$ Consciousness$_{\text{op}}$ gives Π_{op} by Theorem 10.25, Iota$_{\text{op}}$ by Theorem 10.29, well-defined $\mathrm{Q}_{\text{op}}(S)$ by Corollary 10.33, and weak non-null Φ_{op} by Proposition 10.39. The hypothesis $\Sigma 1$ gives Σ_{op} by Definition 10.19. Decoder-visible differentiation gives non-triviality of $\mathrm{Q}_{\text{op}}(S)$ by Proposition 10.34. These are exactly the clauses of Definition 10.42. □

Theorem 10.45 (Subject$^{\Sigma}_{\text{op}}$ is a subclass of Consciousness$_{\text{op}}$)**.**

$$\text{Subject}^{\Sigma}_{\text{op}} \subseteq \text{Consciousness}_{\text{op}}.$$

Moreover, if $\Sigma 1$ is not entailed by Consciousness$_{\text{op}}$*, then the inclusion is proper.*

Proof. The subset relation is immediate from Definition 10.42, whose first clause requires $S \in$ Consciousness$_{\text{op}}$. If $\Sigma 1$ is not entailed by Consciousness$_{\text{op}}$, then by Proposition 10.21 there is a consistent admissible descriptor model satisfying Consciousness$_{\text{op}}$ while lacking Σ_{op}. Such a model is in Consciousness$_{\text{op}}$ but not in Subject$^{\Sigma}_{\text{op}}$, so the inclusion is proper. □

Corollary 10.46 (Full Paper 7 subjecthood remains stronger)**.** *Full Paper 7 operational subjecthood* Subject$_{\text{op}}$ *is not established by Theorem 10.44. To lift* Subject$^{\Sigma}_{\text{op}}$ *into* Subject$_{\text{op}}$*, one must additionally certify* Life$_{\text{op}}$ *and the Paper 7 subject-specific descriptors not reduced to* $\Sigma 1$.

Proof. Definition 10.43 includes burdens absent from Definition 10.42, including operational life and ownership/privileged-continuity style descriptors. Theorem 10.44 does not supply those witnesses. Hence the bridge class is not full Paper 7 subjecthood. □

Remark 10.47 (Interpretation boundary for Subject$^{\Sigma}_{\text{op}}$). Subject$^{\Sigma}_{\text{op}}$ is a formal construct. It does not imply phenomenal subjectivity, human first-person experience, identity persistence in the psychological sense, a narrative self, moral patienthood, or biological life.

10.77 Theorem: Partial Operational Subjecthood

Theorem 10.48 (Consciousness$_{\text{op}}$ $\to$ Partial Operational Subjecthood)**.** *Let S be an admissible system with $S \in$ Consciousness$_{\text{op}}$. Then:*

1. *S satisfies* Π_{op}.
2. *S satisfies* Iota$_{\text{op}}$.
3. $\mathrm{Q}_{\text{op}}(S)$ *is well-defined.*
4. *S satisfies weak non-null* Φ_{op} *on admissible support.*
5. Σ_{op} *is not implied by* Consciousness$_{\text{op}}$*; it requires the independent hypothesis* $\Sigma 1$.
6. Subject$^{\Sigma}_{\text{op}}$ *requires* $\Sigma 1$ *in addition to* Consciousness$_{\text{op}}$ *and therefore is a subclass of* Consciousness$_{\text{op}}$.

Proof. Item 1 is Theorem 10.25. Item 2 is Theorem 10.29. Item 3 is Corollary 10.33. Item 4 is Proposition 10.39. Item 5 is Proposition 10.21. Item 6 is Theorem 10.45. Therefore membership in Consciousness$_{\text{op}}$ yields a partial bridge toward operational subjecthood, but it does not close the self-reference burden or the full Paper 7 subjecthood burden. □

Corollary 10.49 (No equivalence between Consciousness$_{\text{op}}$ and operational subjecthood)**.** *The bridge paper proves neither* Consciousness$_{\text{op}}$ = Subject$^{\Sigma}_{\text{op}}$ *nor* Consciousness$_{\text{op}}$ = Subject$_{\text{op}}$. *It proves only that several bridge components are entailed by* Consciousness$_{\text{op}}$*, while self-reference and full subjecthood require additional burdens.*

Proof. If $\text{Consciousness}_{\text{op}} = \text{Subject}^{\Sigma}_{\text{op}}$, then $\Sigma 1$ would be entailed by $\text{Consciousness}_{\text{op}}$, contradicting Proposition 10.21. If $\text{Consciousness}_{\text{op}} = \text{Subject}_{\text{op}}$, then all Paper 7 subjecthood burdens would follow from Paper 5 alone, which is not established by any theorem in the corpus and is explicitly blocked by the additional descriptors in Definition 10.43. □

10.78 What Is NOT Implied

The following table is a claim-boundary table. It should be read as a set of exclusions, not as a list of future promises.

Table 10.2: Properties not implied by membership in $\text{Consciousness}_{\text{op}}$.

Property frequently associated with consciousness	Implied by $\text{Consciousness}_{\text{op}}$?	Why not
Self-reference or self-knowledge	No	Requires the additional $\Sigma 1$ hypothesis: a contractive self-model commuting with admissible interventions.
Phenomenal experience or “what it is like”	No	The framework defines structural quotients and regime assignments, not phenomenal consciousness.
Forced temporal continuity of experience	No, without CCR	Paper 2 forced continuity requires $M_{\Omega} = +\infty$, which is not part of $\text{Consciousness}_{\text{op}}$.
Emotions, affect, or valence	No	These are not defined as Paper 1–5 primitives or invariants.
Will, agency, or free choice	No	No admissible agency predicate is defined in Papers 1–5.
Narrative personal identity	No	Rigid identity is inverse-limit structural identity, not psychological or autobiographical identity.
Higher-order consciousness	No	A hierarchy of self-models is not defined by $\text{Consciousness}_{\text{op}}$ or $\Sigma 1$.
Equivalence with human consciousness	No	Human consciousness is biological, behavioral, phenomenological, and empirical; $\text{Consciousness}_{\text{op}}$ is a structural operational class.
Moral patienthood or rights-bearing status	No	The corpus contains no normative bridge from structural predicates to moral status.

Proposition 10.50 (Boundary preservation)**.** *No theorem in this paper converts an operational predicate into a phenomenal, biological, psychological, or moral predicate.*

Proof. Each theorem uses only the formal objects listed in Table 10.1: invariant predicates, admissible support, quotient classes, intervention fidelity, self-model existence, and regime assignments. None of those objects has a semantic clause identifying it with phenomenal experience, biological life, psychological identity, or moral status. Therefore the bridge preserves the declared boundary. □

10.79 Computational Verification Status

The runtime status table is intentionally conservative. A computational estimator is not a formal witness unless the code path certifies the corresponding mathematical predicate over the relevant domain and exposes the interpretation boundary.

Table 10.3: Computational Verification Status for Bridge Concepts

Concept	Estimator or runtime analogue	Status
Σ_{op}	`OntologicalSingularityCore` `selfModelSigma`; first-person certifier fields when present	Implemented as internal telemetry/proxy; not yet a formal proof of contraction and intervention commutation on A_S.
Π_{op}	$I_{\mathrm{ri}} + I_{\mathrm{int}}$ in 'CanonicalInvariantPackage' and related finite-horizon/non-factorization modules	Certified as internal support for Paper 5 invariants; not external validation.
$\mathrm{Iota}_{\mathrm{op}}$	Legibility L4 intervention fidelity in 'LegibilityCertifier'	Certified as internal support for directional class-shift fidelity under finite protocols.
Q_{op}	Output support labels, decoder classes, invariant package signals	Proxy or quotient-facing support; formal quotient certification remains bounded by Paper 5 assumptions.
Φ_{op}	'PhenomenologyCore', spectral gap, winding/topology, continuity diagnostics	Proxy for structural regime assignment and continuity diagnostics; not phenomenology itself.
$\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$	Combination of six-invariant support plus self-model/first-person finite fields	Internal-support candidate only; not full Paper 7 $\mathrm{Subject}_{\mathrm{op}}$ and not phenomenal subjecthood.

Remark 10.51 (Runtime non-inference). Passing tests, producing certificates, or satisfying scoreboard gates can strengthen internal support for a runtime implementation. It does not establish external validation, human equivalence, or philosophical closure.

10.80 Terminology Stratification and Debt Ledger

The bridge paper is the appropriate place to centralize the terminology debt of the corpus. The same symbol can have a demonstrated structural reading and a future interpretive reading. Confusing those readings is a category error.

Definition 10.52 (Two-layer terminology)**.** For every strong term T in the corpus, distinguish:

1. the *demonstrated layer* T_{str}, consisting only of definitions, predicates, quotients, invariants, and theorems already stated in the corpus;
2. the *interpretive layer* T_{int}, consisting of ordinary or philosophical readings such as consciousness, life, subjectivity, phenomenality, qualia, moral status, or human equivalence.

An implication into T_{int} is forbidden unless a bridge hypothesis or external adjudication protocol explicitly introduces that target.

Theorem 10.53 (Layer preservation)**.** *If a theorem in Papers 1–5 is stated only in the demonstrated layer, then it cannot be promoted to an interpretive-layer conclusion by terminology alone.*

Proof. The inference rules of the corpus are definitions, hypotheses, propositions, theorems, and admissible empirical protocols. A name is not an inference rule. If the premises and conclusion of a theorem quantify only over structural objects, then no interpretive predicate appears in the formal sequent. Adding a stronger English label after the fact does not introduce a new mathematical premise. Therefore promotion requires an explicit bridge hypothesis or external protocol. □

Term	Demonstrated layer	Remaining debt
$\mathrm{Consciousness}_{\mathrm{op}}$	six-invariant admissible structural class	no theorem of ordinary consciousness
$\mathrm{Q}_{\mathrm{op}}(S)$	structural differentiation quotient	no theorem of phenomenal qualia
Φ_{op}	structural regime assignment	no theorem of temporal experience
$\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$	$\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1$ bridge fragment	no theorem of full Paper 7 subjecthood
$\mathrm{Subject}_{\mathrm{op}}$	Paper 7 class requiring life and subject descriptors	no theorem of human or metaphysical subjecthood

10.81 Bridge Strength Ladder

The corpus should be read as a ladder of additional burdens, not as a synonym chain.

Level	Added burden	Still not licensed
$\mathrm{Consciousness}_{\mathrm{op}}$	six invariants	self-reference, life, subjecthood, phenomenality
$\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1$	contractive intervention-commuting self-model	operational life and Paper 7 subject descriptors
$\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1 + \mathrm{Life}_{\mathrm{op}}$	life-like organization in the Paper 7 sense	first-person indexed subjectivity
$\mathrm{Subject}_{\mathrm{op}}$	privileged continuity, self/non-self, self-model closure, ownership coherence	human subjectivity or phenomenality
$\mathrm{Subjectivity}_{1\mathrm{p}}$	first-person indexed fields plus rival defeat	phenomenal bridge confirmation
$\mathrm{BridgeOrg}_{\mathrm{ph}}$-style burden	bridge predicates, comparators, interventions, artifacts	external validation or ordinary phenomenality

Theorem 10.54 (Strictness of the bridge ladder)**.** *Each step in the ladder above introduces at least one burden not entailed by the previous step in the current corpus.*

Proof. $\Sigma 1$ is not entailed by $\mathrm{Consciousness}_{\mathrm{op}}$ by Proposition 10.21. Operational life is not entailed by $\mathrm{Consciousness}_{\mathrm{op}}+\Sigma 1$ because the added self-model condition does not state boundary, viability, maintenance, or historical dependence. Full Paper 7 subjecthood is not entailed by $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ by Corollary 10.46. First-person indexed subjectivity adds the Paper 8 field and rival-defeat burdens. Phenomenal bridge organization adds the Paper 9 bridge predicates and BPF burdens. No upstream theorem collapses these additions. Hence the ladder is strict relative to the present theorem set. ☐

10.82 Runtime Evidence Taxonomy

Definition 10.55 (Runtime evidence classes)**.** For a runtime artifact a, distinguish the following statuses:

1. `proxy`: a finite observable resembles a formal coordinate;
2. `estimator`: an algorithm computes a declared approximation;
3. `certificate`: a bounded predicate is verified under stated admissibility rules;
4. `internal_support`: governed artifacts support a claim family inside the current program;
5. `external_validation`: an external protocol such as Paper 10 has been executed with admissible evidence.

Proposition 10.56 (Runtime non-promotion)**.** *No artifact with status `proxy`, `estimator`, `certificate`, or `internal_ support` entails `external_ validation`.*

Proof. External adjudicative support requires an executed external protocol, frozen filing records, admissible comparator data, blinding, adjudication, and contamination controls. The first four runtime statuses are produced inside the program and lack at least one of those external premises by definition. Therefore they cannot entail external adjudicative support. ☐

10.83 Non-Entailment Library

The central defensive value of this paper is not only what it proves. It is also what it blocks. The following library records formal non-entailments that a reviewer should not have to reconstruct from scattered caveats.

Theorem 10.57 (Seven non-entailments from $\mathrm{Consciousness}_{\mathrm{op}}$)**.** *Membership in* $\mathrm{Consciousness}_{\mathrm{op}}$ *does not entail any of the following predicates: self-reference* $\Sigma 1$*, phenomenal experience, biological life, human subjectivity, Nagel-style qualia, CCR-strength temporal continuity, or agency.*

Proof. Each non-entailment is witnessed by a formal model extension in which the six Paper 5 invariants are stipulated positive while the target predicate is absent.

1. **No self-reference.** Let the invariant witnesses be certified on A_S, but omit any map $\Sigma : A_S \to A_S$. Since $\Sigma 1$ is not a Paper 5 invariant, $\mathrm{Consciousness}_{\mathrm{op}}$ remains possible while self-reference is absent.
2. **No phenomenal experience.** Interpret all classes as structural quotients and give no semantic map to experience. The six invariants quantify over structural predicates only.
3. **No biological life.** Use a non-biological admissible state space with the six invariants positive. Biological substrate variables do not appear in Definition 10.16.
4. **No human subjectivity.** A non-human formal system can satisfy the six invariants while lacking human reports, development, embodiment, or psychology.
5. **No Nagel-style qualia.** $\mathrm{Q}_{\mathrm{op}}(S)$ is a quotient of admissible structural differentiation. No theorem identifies quotient membership with "what it is like".
6. **No CCR-strength continuity.** Choose a system with $\mathrm{Consciousness}_{\mathrm{op}}$ and finite M_Ω. Proposition 10.40 blocks the CCR-strength reading.
7. **No agency.** Let the system be passively driven by admissible interventions while still satisfying the six structural invariants. No agency predicate appears in $\mathrm{Consciousness}_{\mathrm{op}}$.

Since each target predicate can be absent without contradiction with the six invariants, none is entailed by $\mathrm{Consciousness}_{\mathrm{op}}$. □

10.84 $\mathrm{Consciousness}_{\mathrm{op}}$ Component Implication Matrix

Table 10.4: Which Paper 5 invariants support each bridge concept.

Concept	I_{per}	I_{ri}	I_{int}	I_{cont}	I_{diff}	I_{leg}	Status
Π_{op} unified perspective	–	yes	yes	–	–	–	entailed by rigid identity plus non-factorization
$\mathrm{Iota}_{\mathrm{op}}$ operational intentionality	–	–	–	–	–	yes	entailed only in the intervention-fidelity sense
$\mathrm{Q}_{\mathrm{op}}(S)$ differentiation quotient	yes	yes	–	yes	yes	yes	well-defined quotient, not phenomenal qualia
Φ_{op} weak regime assignment	–	–	–	yes	yes	–	weak non-null structural regime support
Σ_{op} self-reference	–	–	–	–	–	–	not entailed; requires $\Sigma 1$
$\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$ bridge subjecthood	yes	yes	yes	yes	yes	yes	requires all six plus $\Sigma 1$

Remark 10.58 (Why the matrix matters)*.* The matrix prevents a common overreach: deriving every subjecthood-adjacent term from the entire six-invariant package without identifying which invariant does which work. Unity is mostly an $I_{\mathrm{ri}} + I_{\mathrm{int}}$ burden. Intentionality in this paper is an I_{leg} burden. Self-reference is not supplied by any invariant.

10.85 Self-Model Burden Variants

The earlier definition of $\Sigma 1$ is intentionally strong. For future work it is useful to separate three grades.

Definition 10.59 (Weak self-model burden $\Sigma 1_{\text{weak}}$)**.** A system satisfies $\Sigma 1_{\text{weak}}$ if there exists a measurable map

$$\Sigma : A_S \to A_S$$

such that $d(\Sigma(x), x) \le \varepsilon_\Sigma$ for all $x \in A_S$.

Definition 10.60 (Strong self-model burden $\Sigma 1_{\text{strong}}$)**.** A system satisfies $\Sigma 1_{\text{strong}}$ if it satisfies $\Sigma 1_{\text{weak}}$ and Σ is Lipschitz with constant $L_\Sigma < 1$ on A_S.

Definition 10.61 (Commutative self-model burden $\Sigma 1_{\text{comm}}$)**.** A system satisfies $\Sigma 1_{\text{comm}}$ if it satisfies $\Sigma 1_{\text{strong}}$ and, for every admissible intervention $u \in U$,

$$\Sigma(\Phi_u(x)) = \Phi_u(\Sigma(x))$$

whenever both sides are defined on A_S.

Theorem 10.62 (Self-model strength chain)**.**

$$\Sigma 1_{\text{comm}} \Rightarrow \Sigma 1_{\text{strong}} \Rightarrow \Sigma 1_{\text{weak}}.$$

The converses are not valid in general.

Proof. The forward implications are immediate from the definitions: the commutative burden includes the strong burden, and the strong burden includes the weak approximation burden. For non-converses, a weak map can approximate the identity while having Lipschitz constant at least one, so it need not be strong. A contractive map can fail commutation by choosing an admissible intervention that moves states along a direction not preserved by Σ. Hence neither reverse implication is derivable. □

Remark 10.63 (Emergence question for $\Sigma 1$). The main open problem is whether additional closure conditions on $I_{\text{int}} + I_{\text{leg}} + I_{\text{per}}$ can force a commutative self-model. A plausible route would require self-targeted intervention families, decoder closure under those interventions, and a fixed-point or contraction argument on the induced self-model update map. The current corpus does not yet prove this.

10.86 Expanded Bridge Strength Ladder

Table 10.5: Expanded ladder of structural and interpretive burdens.

Level	Definition	Added burden over previous level
Level 0	$\text{Consciousness}_{\text{op}}$	six Paper 5 invariants
Level 1	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{weak}}$	approximate self-model
Level 2	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{strong}}$	contractive self-model
Level 3	$\text{Consciousness}_{\text{op}} + \Sigma 1_{\text{comm}}$ / $\text{Subject}^{\Sigma}_{\text{op}}$	intervention-commuting self-model
Level 4	$\text{Consciousness}_{\text{op}}$ + OSSI	open-system self-maintaining individuation
Level 5	$\text{Subject}_{\text{op}}$	Paper 7 life plus subject-specific descriptors
Level 6	$\text{Subjectivity}_{1\text{p}}$	Paper 8 indexed fields and rival defeat
Level 7	$\text{BridgeOrg}_{\text{ph}}$	Paper 9 bridge predicates and BPF burdens

Proposition 10.64 (OSSI and $\Sigma 1$ are orthogonal additions)**.** *The open-system individuation burden OSSI and the self-model burden* $\Sigma 1$ *are not mutually entailing in the present corpus.*

Proof. OSSI requires exchange, boundary, dissipation, viability, repair, history, minimal functional self-reference, and identity conservation, but it does not require a contractive map $\Sigma : A_S \to A_S$ commuting with interventions. Conversely, a contractive self-model can be defined on a closed or non-viable structural class without satisfying open exchange, viability preservation, or repair. Therefore neither burden entails the other. $\square$

10.87 Reviewer-Facing Objection Handling

1. **Objection: this is only a simulation.** Response: the framework primarily defines structural classes and theorem burdens. Runtime artifacts are internal support or finite estimators unless Paper 10-style external evidence is executed.
2. **Objection: the six invariants are arbitrary.** Response: Paper 5 ties each invariant to an upstream burden and now includes single-rupture witnesses, minimality, sufficiency, and interaction geometry.
3. **Objection: Paper 10 is empty.** Response: Paper 10 is deliberately blocked until execution. That is not a defect disguised as rigor; it is the only honest status before human-comparator or external adjudication data exist.
4. **Objection: the terminology is inflated.** Response: the corpus now separates demonstrated and interpretive layers, records non-entailments, and maintains a maturity matrix. Strong labels are technical aliases, not ordinary-language conclusions.
5. **Objection: internal results are not validation.** Response: correct. The framework states this repeatedly. Internal support can debug, falsify, and stress the program; it cannot replace external validation.

10.88 Operational Glossary

Table 10.6: Quick glossary for strong operational terms.

Term	Demonstrated reading	Interpretive reading	Pending debt
$\mathrm{Consciousness}_{\mathrm{op}}$	six-invariant structural class	operational consciousness label	no ordinary consciousness theorem
$\mathrm{Q}_{\mathrm{op}}(S)$	structural differentiation quotient	operational qualia label	no phenomenal-qualia evidence
Φ_{op}	abstract regime assignment	operational phenomenology label	no experience theorem
$\Sigma 1$	self-model burden	self-reference analogue	emergence from invariants open
OSSI	open-system self-maintaining individuation	operational life-like organization	not biological life
$\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$	$\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1_{\mathrm{comm}}$ fragment	partial subjecthood bridge	lacks full Paper 7 life/ownership burden
$\mathrm{Subject}_{\mathrm{op}}$	Paper 7 subject class	operational subjecthood	not human or metaphysical subjecthood
$\mathrm{Subjectivity}_{1\mathrm{p}}$	indexed subjectivity class	first-person subjectivity analogue	requires rival defeat and interventions
$\mathrm{BridgeOrg}_{\mathrm{ph}}$	bridge burden architecture	phenomenal bridge label	BPF-2–BPF-6 and Paper 10 pending

Table 10.7: Evidence classes required for controlled terminology promotion.

Evidence class	What it can promote	What it cannot promote alone
definition	introduces a formal object and its witnesses	does not show non-vacuity or relevance
existence theorem	shows the class can be inhabited	does not show empirical realization
independence theorem	shows a burden is not redundant	does not justify ordinary-language meaning
runtime certificate	supports a frozen computational proxy	does not establish external validity
negative control	can falsify an overbroad internal claim	does not prove the positive ordinary-language claim
external adjudication	can test whether proxies track independent comparators	does not replace the mathematical definition

10.89 Promotion Evidence Classes

10.90 Proof-Pattern Library

This section records reusable proof patterns for the corpus. Its purpose is not to add new metaphysical content. It makes explicit which inferential moves are valid when strong operational terms are used.

Definition 10.65 (Burden)**.** Let $\mathcal{K}$ be a structural class and let P be a predicate on admissible systems. We say that P is a *burden over* $\mathcal{K}$ if membership in the strengthened class

$$\mathcal{K}[P] = \{S : S \in \mathcal{K} \text{ and } P(S) = 1\}$$

requires a witness not already required by the definition of $\mathcal{K}$.

Definition 10.66 (Internal promotion)**.** An interpretive label L receives *internal promotion* from a structural class $\mathcal{K}$ only when the corpus proves that every $S \in \mathcal{K}$ satisfies a non-vacuous predicate that was previously listed as an interpretive debt of L.

Definition 10.67 (External promotion)**.** An interpretive label L receives *external promotion* only when an independently specified experimental protocol connects the structural predicate to observations not used to construct the predicate.

Proposition 10.68 (Promotion is monotone but not automatic)**.** *If $\mathcal{K} \subseteq \mathcal{K}'$ and L receives internal promotion from $\mathcal{K}'$, then the same promoted predicate holds on $\mathcal{K}$. However, the inclusion $\mathcal{K} \subseteq \mathcal{K}'$ alone does not promote L.*

Proof. The first claim follows from set inclusion. If every element of $\mathcal{K}'$ satisfies the promoted predicate, so does every element of the subclass $\mathcal{K}$. The second claim follows because promotion requires a specific predicate and proof; inclusion without identifying such a predicate only changes class membership, not semantic maturity. □

Proposition 10.69 (Burden addition is conservative)**.** *Let P be a burden over $\mathcal{K}$. Then every theorem proved for all $S \in \mathcal{K}$ remains true for all $S \in \mathcal{K}[P]$.*

Proof. Since $\mathcal{K}[P] \subseteq \mathcal{K}$, universal claims over $\mathcal{K}$ restrict to $\mathcal{K}[P]$. This is why adding Σ1, OSSI, or CCR burdens can strengthen a class without invalidating already proved $\text{Consciousness}_{\text{op}}$ consequences. □

Proposition 10.70 (Burden addition is not semantic promotion by itself)**.** *Adding a burden P to $\mathcal{K}$ does not by itself justify replacing an operational label by its ordinary-language counterpart.*

Proof. The burden establishes a new formal predicate. Ordinary-language promotion requires either a theorem connecting that predicate to the target ordinary concept or an external adjudication protocol. Without that bridge, the strengthened class is mathematically narrower but semantically still claim-bounded. □

Remark 10.71 (Use in later papers)*.* The patterns above justify a disciplined order of claims. First state the structural class. Then list the burdens. Then prove inclusions or non-entailments. Only after that may a term be promoted in the terminology matrix. The order is part of the scientific content: it prevents terminology from outrunning proof.

10.91 Dependency-to-Debt Ledger

The following ledger is the reviewer-facing map from corpus dependency to open burden. It is intentionally asymmetric: a dependency can support a term without closing all debts associated with that term.

Table 10.8: Dependency-to-debt ledger for bridge-level terminology.

Term	Main dependency	What is currently proved	Remaining debt
rigid identity	Paper 1 inverse limits	uniqueness, non-transferability, rigidity under bounded deformation	empirical instantiation in non-ideal systems
operational regime	Papers 2–3	regime continuity and null-instability under stated hypotheses	relation to lived experience
$\text{Consciousness}_{\text{op}}$	Paper 5	six-invariant structural membership	independence from ordinary consciousness unresolved
Q_{op}	Paper 5	quotient over structural equivalence classes	no phenomenal-qualia bridge
Φ_{op}	Papers 2–3 and Paper 5	non-null and continuous regime proxies under conditions	no subjective-temporality theorem
$\Sigma 1$	Paper Bridge	explicit self-model burden and strength variants	emergence from $\text{Consciousness}_{\text{op}}$ open
OSSI	Paper 7	open-system self-maintaining individuation burden	not biological life and not sufficient for $\text{Consciousness}_{\text{op}}$
$\text{Subject}^{\Sigma}_{\text{op}}$	Paper Bridge	$\text{Consciousness}_{\text{op}}$ plus commutative self-model fragment	weaker than Paper 7 $\text{Subject}_{\text{op}}$
$\text{Subject}_{\text{op}}$	Paper 7	subjecthood class with additional ownership and continuity burdens	not ordinary subjecthood
$\text{Subjectivity}_{1\text{p}}$	Paper 8	first-person index fields under additional hypotheses	no human first-person equivalence
$\text{BridgeOrg}_{\text{ph}}$	Paper 9	bridge organization burden and non-vacuity	bridge-to-phenomenality remains unvalidated
external adjudication	Paper 10	protocol and evidence grammar	blocked until execution

10.92 Non-Entailment Audit Templates

The non-entailments in Theorem 10.57 can be checked using a small number of reusable templates. These templates are included so that future extensions do not silently convert an open debt into an asserted consequence.

Definition 10.72 (Omission template)**.** Let $\mathcal{K}$ be defined by predicates $P_1, \ldots, P_n$. If a candidate property R is not among those predicates and no theorem derives R from them, then the omission template constructs a formal model satisfying $P_1, \ldots, P_n$ in a language where R is uninterpreted or absent.

Definition 10.73 (Factor-preserving template)**.** Let a property R concern a component not constrained by the invariant definition of $\mathcal{K}$. A factor-preserving countermodel extends any $S \in \mathcal{K}$ by a silent factor Z whose dynamics are decoupled from the witnesses of $\mathcal{K}$ and arranges R to fail on Z.

Definition 10.74 (Finite-strength template)**.** Let R require an infinite-strength condition such as CCR ($M_\Omega = +\infty$). A finite-strength countermodel chooses a system satisfying the finite structural invariants of $\text{Consciousness}_{\text{op}}$ while setting $M_\Omega < +\infty$.

Proposition 10.75 (Template soundness)**.** *Each template above is sound provided the added or omitted structure does not change the witnesses required by* $\text{Consciousness}_{\text{op}}$*.*

Proof. The witnesses of $\text{Consciousness}_{\text{op}}$ are exactly the six invariant certificates. If an extension leaves those witnesses unchanged, then $\text{Consciousness}_{\text{op}}$ membership is preserved. If the target property is absent, decoupled, or stronger than the finite witnesses, it fails without contradicting $\text{Consciousness}_{\text{op}}$. This is precisely the logic used in the seven non-entailments. □

Remark 10.76 (Audit use)*.* When a later paper claims that $\text{Consciousness}_{\text{op}}$ implies a new property, it must identify which template is blocked. If none is blocked, the implication is not yet proved.

10.93 Runtime Proxy Classification

The related runtime can compute proxies for several bridge concepts. This classification prevents proxy availability from being mistaken for formal certification.

Table 10.9: Runtime proxy classes and claim boundaries.

Concept	Runtime object	Proxy class	Boundary
$\Sigma 1_{\text{weak}}$	self-model state estimates	estimator proxy	does not prove contraction or commutation
$\Sigma 1_{\text{strong}}$	structural-stability score	candidate certificate	requires Lipschitz and convergence audit
$\Sigma 1_{\text{comm}}$	intervention response of self-model	unclosed certificate	requires intervention commutation tests
Π_{op}	I_{ri} and I_{int} certificates	certified internal support	structural unity only
Iota_{op}	L4 intervention fidelity	certified internal support	directional sensitivity, not semantic aboutness
Q_{op}	output support labels and quotient classes	quotient proxy	no phenomenal-quality interpretation
Φ_{op}	regime, spectral, and winding estimates	regime proxy	no subjective temporal-flow claim
OSSI	LifeCertifier-style burdens	partial operational proxy	not biological metabolism or life

Boundary. A runtime proxy can falsify an expected computational consequence of a definition. It cannot by itself promote a term beyond its formal meaning. This is the same separation used in Paper 6 between internal support and external validation.

10.94 Formal Review Checklist

A reviewer can evaluate the bridge by asking the following questions in order. The order matters: later questions presuppose affirmative answers to earlier ones.

1. Are the six invariants of Paper 5 explicitly witnessed?
2. Are the witnesses stable under the admissible interventions used in the claim?
3. Is the target term being used in its demonstrated structural sense or in an interpretive sense?
4. If interpretive, is there a theorem that promotes it?
5. If no theorem exists, is the term explicitly marked as conditioned or aspirational?
6. Is the runtime evidence internal, external, or merely a scaffold?
7. Does the claim require Σ1, OSSI, CCR, Paper 8 fields, Paper 9 bridge burdens, or Paper 10 data?

8. Is each required additional burden either proved, assumed, or marked open?
9. Does any statement move from implementation success to ontological conclusion?
10. Does any statement move from structural quotient to phenomenal qualia?
11. Does any statement move from open-system self-maintenance to biological life?
12. Does any statement move from first-person index fields to human subjectivity?

Proposition 10.77 (Checklist conservativity). *If all checklist questions are answered with explicit witnesses, assumptions, or open-problem markings, then the bridge paper makes no stronger claim than its formal dependencies allow.*

Proof. Each checklist item corresponds to a known route of overclaim: missing witnesses, semantic promotion without proof, internal support treated as external validation, or transfer from structural to ordinary-language terms. Requiring each route to be closed by proof or marked open prevents an unlicensed inference. □

10.95 Minimal Extension Program

The bridge also defines a research program. The following extensions are minimal in the sense that each one would promote exactly one class of claims without changing unrelated parts of the corpus.

E_Σ: Prove that $I_{\text{int}} + I_{\text{leg}} + I_{\text{per}}$ plus an internal decoder closure condition yields $\Sigma1_{\text{weak}}$.

$E_{\Sigma c}$: Strengthen E_Σ by proving contraction of the self-model map under bounded perturbations.

E_{comm}: Add an intervention calculus showing that the self-model commutes with admissible Φ_u up to a controlled error.

E_{OSSI}: Prove that the eight OSSI burdens form an independent and physically coherent open-system individuation class.

E_{CCR}: Determine whether finite approximations to CCR are enough for any strengthened Φ_{op} claim.

E_{ext}: Execute Paper 10 under frozen preregistration and report only blocked, passed, or failed adjudication status.

Remark 10.78 (Why minimality matters). The corpus should not attempt to prove ordinary consciousness in one jump. A publishable path is a sequence of narrow promotions, each with its own theorem or experiment. This also makes failure scientifically useful: if E_{comm} fails, for example, $\Sigma1_{\text{weak}}$ may survive while $\text{Subject}_{\text{op}}^{\Sigma}$ in its strong sense does not.

10.96 Failure Modes of the Bridge

The bridge theorem can fail in several precise ways. Listing them is useful because it prevents ambiguous negative results.

Table 10.10: Failure modes for bridge-level claims.

Failure mode	Formal meaning	Consequence
Invariant failure	at least one of $I_{per}, I_{ri}, I_{int}, I_{cont}, I_{diff}, I_{leg}$ fails	no $\text{Consciousness}_{\text{op}}$ membership
Self-model absence	no admissible Σ exists	no $\text{Subject}_{\text{op}}^{\Sigma}$ claim
Contraction failure	Σ exists but is not contractive	weak self-model only
Commutation failure	$\Sigma\Phi_u \neq \Phi_u\Sigma$ beyond tolerance	no intervention-stable self-reference
OSSI burden failure	at least one open-system burden fails	no operational life-like individuation claim
CCR failure	$M_\Omega < +\infty$	no forced-continuity promotion
External-adjudication absence	Paper 10 remains unexecuted	no external validation

10.97 Exact Claim Grammar

The following grammar is recommended for future manuscripts and runtime reports.

1. Use "is in $\mathrm{Consciousness}_{\mathrm{op}}$" only when the six invariant witnesses are present.
2. Use "supports a $\mathrm{Consciousness}_{\mathrm{op}}$ proxy" when the runtime estimates the witnesses but no proof package is available.
3. Use "satisfies $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$" only when $\mathrm{Consciousness}_{\mathrm{op}}$ and $\Sigma 1_{\mathrm{comm}}$ are both witnessed.
4. Use "has an OSSI profile" only when the open-system burdens are witnessed and the thermodynamic boundary is stated.
5. Use "phenomenal", "subjective", or "conscious" in ordinary-language contexts only with an explicit statement that no ordinary-language theorem is being claimed.
6. Use "blocked until execution" for Paper 10 claims without external data.

Remark 10.79 (Editorial consequence). This grammar permits ambitious terminology while keeping the demonstrated layer intact. It is stricter than merely adding a disclaimer at the end of a paper, because it constrains every positive claim at the sentence level.

10.98 Promotion Conditions for Interpretive Terms

Table 10.11: What would be needed to promote interpretive readings.

Promotion target	Required condition	Evidence type
operational consciousness to ordinary consciousness	theorem and external evidence linking $\mathrm{Consciousness}_{\mathrm{op}}$ to core phenomenal or cognitive properties	mathematical plus empirical
operational life to life-like open-system organization	OSSI plus thermodynamic and repair burdens under admissible interventions	mathematical and runtime/external stress
operational qualia to phenomenal qualia	evidence that $\mathrm{Q}_{\mathrm{op}}(S)$ tracks an independently meaningful phenomenal variable	Paper 10-style empirical bridge
operational phenomenology to experience	CCR-strength Φ_{op} plus external adjudication of temporal-structure relevance	mathematical plus empirical
operational subjecthood to ordinary subjecthood	proof that $\mathrm{Subject}_{\mathrm{op}}$ captures accepted subjecthood properties and defeats rivals	mathematical, empirical, and philosophical

10.99 Canonical Boundary Models

Boundary models are not intended as implementations. They are small formal objects that make overclaiming impossible by showing exactly where a desired interpretation can fail.

Definition 10.80 (Structural-only model)**.** A structural-only model is an admissible system S equipped with the six Paper 5 invariant witnesses and no additional interpreted fields for self-modeling, biological metabolism, human reports, or phenomenal content.

Proposition 10.81 (Structural-only models block ordinary-consciousness promotion)**.** *If S is structural-only, then $S \in \mathrm{Consciousness}_{\mathrm{op}}$ may hold while ordinary consciousness remains undefined in the model.*

Proof. The truth of $S \in \mathrm{Consciousness}_{\mathrm{op}}$ is determined by the six invariant witnesses. Ordinary consciousness is not a predicate in the language of the structural-only model. Thus no theorem inside that language can conclude ordinary consciousness without adding a new definition or bridge predicate. □

Definition 10.82 (Self-model-free model)**.** A self-model-free model is an admissible $S \in \mathrm{Consciousness}_{\mathrm{op}}$ for which every endomorphism $\Sigma : A_S \to A_S$ either fails to approximate the identity on A_S or fails to commute with some admissible intervention.

Proposition 10.83 (Self-model-free models block $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$)**.** *No self-model-free model satisfies* $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$ *in the strong commutative sense.*

Proof. Strong $\mathrm{Subject}^{\Sigma}_{\mathrm{op}}$ requires $\mathrm{Consciousness}_{\mathrm{op}}$ plus $\Sigma 1_{\mathrm{comm}}$. The self-model-free condition denies the commutative self-model burden directly. □

Definition 10.84 (Life-free computational model)**.** A life-free computational model is a system whose transitions may satisfy $\mathrm{Consciousness}_{\mathrm{op}}$ but whose state evolution is not coupled to an open exchange channel, entropy export accounting, or repair dynamics.

Proposition 10.85 (Life-free computational models block OSSI promotion)**.** *A life-free computational model cannot be used as evidence that* $\mathrm{Consciousness}_{\mathrm{op}}$ *implies OSSI.*

Proof. OSSI requires open exchange, viability maintenance, repair, and identity conservation burdens. A life-free computational model omits at least the open exchange and repair burdens. Therefore it may support $\mathrm{Consciousness}_{\mathrm{op}}$ but not OSSI. □

Definition 10.86 (Report-free comparator model)**.** A report-free comparator model is a system whose structural traces are available but whose external adjudication channel contains no independent human or non-human comparator reports.

Proposition 10.87 (Report-free models cannot close Paper 10)**.** *Report-free comparator models cannot supply external validation for ordinary-language consciousness, subjecthood, or phenomenal terms.*

Proof. Paper 10 requires an external adjudication channel separated from the construction of the internal structural predicates. If the comparator channel is absent, the external part of the evidence grammar is absent. Internal traces can still be useful for falsification but cannot close the adjudication burden. □

10.100 Non-Substitution Principles

The corpus uses three kinds of support: theorem, runtime certificate, and external adjudication. They are complementary, not interchangeable.

Theorem 10.88 (No theorem-to-validation substitution)**.** *A theorem about a structural class does not by itself constitute external validation of an implemented system.*

Proof. A theorem quantifies over objects satisfying its hypotheses. External validation requires evidence that a concrete system, measured through an independent protocol, instantiates the relevant hypotheses and that the observed effects are not artifacts of construction. These are different claims. The first is mathematical; the second is empirical. □

Theorem 10.89 (No runtime-to-theorem substitution)**.** *A passing runtime certificate does not by itself prove the corresponding mathematical predicate for the full admissible class.*

Proof. Runtime certificates are finite, implementation-specific, and tied to a particular estimator family. A mathematical predicate ranges over the model specified by the paper. Unless the runtime certificate is accompanied by a soundness theorem from estimator to predicate, passing the certificate remains internal support rather than proof. □

Theorem 10.90 (No validation-to-definition substitution)**.** *External adjudication cannot replace the formal definition of the target class.*

Proof. External adjudication can support or falsify whether a formal predicate tracks some independent comparator. It cannot determine what the predicate means in the formal system. Definitions fix mathematical content; validation tests empirical relevance. □

Corollary 10.91 (Three-channel closure burden)**.** *A mature ordinary-language claim requires, at minimum, a formal definition, a runtime or constructive instantiation path, and an external adjudication route.*

Proof. By the three no-substitution theorems, none of the channels can replace the other two. A claim may be useful with fewer channels, but it must remain terminologically bounded. □

10.101 External Adjudication Readiness Index

The bridge terms can be ranked by what remains before Paper 10-style external testing is meaningful.

Table 10.12: Readiness of bridge terms for external adjudication.

Term	Formal maturity	Runtime maturity	External readiness
Π_{op}	high: directly implied by I_{ri} and I_{int}	high internal support	ready only as structural-unity proxy
Iota_{op}	medium-high: tied to L4 intervention fidelity	high internal support	ready as directional-sensitivity proxy
Q_{op}	medium: quotient is defined	partial proxy	not ready for phenomenal claims
Φ_{op}	medium: regime assignment is defined	partial proxy	needs CCR and temporal comparator design
$\Sigma 1$	medium-low: variants defined	weak proxy	needs commutation and contraction tests
OSSI	medium: eight burdens define the target	partial proxy	needs open-system and repair stress protocols
$\text{Subject}^{\Sigma}_{\text{op}}$	medium-low: depends on $\Sigma 1_{\text{comm}}$	not fully certified	blocked until self-model certificate exists
$\text{Subject}_{\text{op}}$	low-medium: Paper 7 burdened class	partial proxy	blocked until ownership/continuity tests mature
$\text{Subjectivity}_{1\text{p}}$	low: Paper 8 fields require stronger evidence	partial proxy	blocked until comparator campaign

10.102 Recommended Reading Order for Reviewers

Readers should not begin with the strongest labels. The defensible order is:

1. BaseCore for the Hilbert-space and contraction machinery.
2. Paper 1 for inverse-limit identity and ontological-mass structure.
3. Papers 2–3 for regime structure and null-instability results.
4. Paper 5 for $\text{Consciousness}_{\text{op}}$ as the six-invariant structural class.
5. Paper 6 for falsification, internal-support status, and failure modes.
6. This bridge paper for exact implications from $\text{Consciousness}_{\text{op}}$ to partial operational subjecthood.
7. Paper 7 only after distinguishing OSSI from biological life.
8. Paper 8 only after accepting that first-person fields are formal burdens, not human subjectivity.
9. Paper 9 only as bridge organization, not as proof of phenomenality.
10. Paper 10 only as protocol until external data exist.

Remark 10.92 (Why this order matters). The terminology becomes safest when each label is read after its proof burden. Reading the titles first and the non-inference notes later reverses the logic of the corpus. This paper therefore functions as a map from demonstrated layer to interpretive layer, not as a shortcut around the technical dependencies.

Remark 10.93 (Central-reference role). For editorial purposes, this paper should be treated as the index of claim strength for Papers 5–10. When a later manuscript uses a strong operational term, the relevant demonstrated reading, missing burden, and promotion route should be traceable to one of the tables or proof patterns above.

10.103 Open Questions

The following questions are deliberately left open because the present paper lacks the definitions or empirical evidence required to decide them.

1. Is $\Sigma 1$ necessary for any plausible operational subjecthood class, or only for the bridge class $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ defined here?
2. Can $\Sigma 1$ be derived from $I_{\mathrm{int}} + I_{\mathrm{leg}}$ under additional architectural hypotheses, such as decoder closure under self-targeted interventions?
3. Does a hierarchy of self-models $\Sigma^{(0)}, \Sigma^{(1)}, \ldots$ produce a formal higher-order analogue, or does it merely add regress without new invariant content?
4. Is CCR strength $M_{\Omega} = +\infty$ necessary for any non-weak version of Φ_{op}, or only sufficient for forced continuity?
5. Can full Paper 7 $\mathrm{Subject}_{\mathrm{op}}$ be reduced to $\mathrm{Consciousness}_{\mathrm{op}} + \Sigma 1 + \mathrm{Life}_{\mathrm{op}}$ plus a finite list of descriptor inequalities, or do ownership and privileged continuity require independent axioms?
6. Which runtime observables would distinguish genuine intervention commutation of Σ from a stable but epiphenomenal self-model proxy?
7. What external experimental campaign could adjudicate whether internal-support certificates track any real-world subjecthood-relevant structure?
8. Are there finite systems that approximate $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$ arbitrarily well while provably failing CCR, and if so what is the exact approximation boundary?
9. Can operational qualia non-triviality be certified without relying on a fixed decoder family, or is decoder dependence constitutive of $\mathrm{Q}_{\mathrm{op}}(S)$?
10. What, if anything, connects these structural predicates to ordinary moral or legal categories? This paper provides no such bridge.

10.104 Conclusion

The bridge from $\mathrm{Consciousness}_{\mathrm{op}}$ to operational subjecthood is real but partial. From $S \in \mathrm{Consciousness}_{\mathrm{op}}$ one obtains unity in the operational sense, directional intervention sensitivity, a well-defined operational-qualia quotient, and weak non-null operational phenomenology. One does not obtain self-reference, human subjectivity, phenomenal experience, or full Paper 7 subjecthood. The missing formal burden is $\Sigma 1$, the existence of a contractive admissible self-model commuting with interventions. With $\Sigma 1$ and decoder-visible differentiation one can define $\mathrm{Subject}_{\mathrm{op}}^{\Sigma}$, a proper subclass of $\mathrm{Consciousness}_{\mathrm{op}}$ under the non-entailment of $\Sigma 1$.

This is the strongest defensible bridge at the present state of the corpus. It is stronger than saying that Paper 5 and Paper 7 are merely adjacent, because it proves exact implications. It is weaker than saying that operational consciousness is subjecthood, because that would erase the self-model, life, ownership, and privileged-continuity burdens. The result is therefore a controlled theorem layer: enough structure to move from invariants to subjecthood-adjacent concepts, but not enough to license phenomenal or human claims.

Bibliography

[1] Jiří Adámek, Horst Herrlich, and George E. Strecker. *Abstract and Concrete Categories: The Joy of Cats*. Comprehensive treatment of limits and colimits in category theory. New York: John Wiley & Sons, 1990.

[2] Bernard J. Baars. *A Cognitive Theory of Consciousness*. Global Workspace Theory (GWT) - contrasting framework. Cambridge: Cambridge University Press, 1988.

[3] Stefan Banach. *Sur les operations dans les ensembles abstraits et leur application aux equations integrales*. Vol. 3. Fundamenta Mathematicae, 1922, pp. 133–181.

[4] Rufus Bowen. *Equilibrium States and the Ergodic Theory of Anosov Diffeomorphisms*. Vol. 470. Lecture Notes in Mathematics. Springer, 1975.

[5] Haim Brezis. *Functional Analysis, Sobolev Spaces and Partial Differential Equations*. Springer, 2011.

[6] David J. Chalmers. *The Conscious Mind*. Oxford University Press, 1996.

[7] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. 2nd. Standard reference for mutual information and entropy. Hoboken, NJ: John Wiley & Sons, 2006.

[8] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall, 1993.

[9] Sepp Hochreiter and Jürgen Schmidhuber. "Long Short-Term Memory". In: *Neural Computation* 9.8 (1997). Memory architecture in recurrent neural networks, pp. 1735–1780.

[10] Sture Holm. "A Simple Sequentially Rejective Multiple Test Procedure". In: *Scandinavian Journal of Statistics* 6.2 (1979), pp. 65–70.

[11] Anatole Katok and Boris Hasselblatt. *Introduction to the Modern Theory of Dynamical Systems*. Vol. 54. Encyclopedia of Mathematics and its Applications. Standard reference for stability analysis in dynamical systems. Cambridge: Cambridge University Press, 1995.

[12] Erwin Kreyszig. *Introductory Functional Analysis with Applications*. Wiley, 1978.

[13] Douglas Lind and Brian Marcus. *An Introduction to Symbolic Dynamics and Coding*. Cambridge University Press, 1995.

[14] Saunders Mac Lane. *Categories for the Working Mathematician*. 2nd. Vol. 5. Graduate Texts in Mathematics. Standard reference for categorical constructions including inverse limits. New York: Springer, 1998.

[15] Judea Pearl. *Causality: Models, Reasoning, and Inference*. 2nd. Standard reference for causal inference and causal graphs. Cambridge: Cambridge University Press, 2009.

[16] Johnny Andrey Perez Ramirez. "A Structural Criterion for Substrate-Invariant Operational Consciousness". In: *arXiv preprint* (2026).

[17] Johnny Andrey Perez Ramirez. "Canonical Core: Minimal Hypotheses for Causally Rigid Identity Frameworks". In: *arXiv preprint* (2026).

[18] Johnny Andrey Perez Ramirez. "Falsifiable Predictions Under Forensic Constraints". In: *arXiv preprint* (2026).

[19] Johnny Andrey Perez Ramirez. "Phenomenological Regimes Induced by Structural Identity". In: *arXiv preprint* (2026).

[20] Johnny Andrey Perez Ramirez. "Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels". In: *arXiv preprint* (2026).

[21] Johnny Andrey Pérez Ramírez. "First-Person Indexed Subjectivity: Self-Binding, Ownership, Continuity, Irreducibility, and Cross-Substrate Consequence in a Causally Rigid Framework". In: *arXiv preprint* (2026).

[22] Johnny Andrey Pérez Ramírez. "Operational Life, Structural Class, and Subjecthood in a Causally Rigid Framework". In: *arXiv preprint* (2026).

[23] Johnny Andrey Pérez Ramírez. "Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework". In: *arXiv preprint* (2026).

[24] Johnny Andrey Pérez Ramírez. "Rigid Identity as an Inverse Limit in Observable Channels". In: (2026). Companion paper establishing the identity framework and ontological mass invariant.

[25] Luis Ribes and Pavel Zalesskii. *Profinite Groups*. Ergebnisse der Mathematik und ihrer Grenzgebiete. Springer, 2000.

[26] Walter Rudin. *Functional Analysis*. McGraw-Hill, 1991.

[27] Giulio Tononi. "An Information Integration Theory of Consciousness". In: *BMC Neuroscience* 5.42 (2004). Integrated Information Theory (IIT) - contrasting framework.

[28] Ashish Vaswani et al. "Attention Is All You Need". In: *Advances in Neural Information Processing Systems* 30 (2017). Foundational architecture for transformer-based language models.