

Predictions, Discriminators, and Failure Modes for the Causally Rigid Operational Consciousness Framework

Johnny Andrey Pérez Ramírez

Abstract

This paper does not expand the doctrine of the causally rigid operational consciousness corpus. It extracts the experimentally relevant consequences of the frozen canon and organizes them as a falsifiable prediction program. The upstream burdens remain where they belong: the Core supplies persistence, attractor structure, and non-collapse [2]; *Rigid Identity* supplies rigid identity and observable coherence [6]; *Phenomenological Regimes* supplies continuity and anti-fragmentation constraints [4]; *Null-Regime Instability* supplies null-regime instability and forced non-nullity [5]; *Forensic Predictions* supplies admissibility and forensic discipline [3]; *Operational Consciousness Criterion* supplies the six-invariant operational criterion, structural equivalence, and rupture semantics [1].

The function of the present paper is narrower. It states what the framework predicts, what weaker alternatives would predict instead, what would count as genuine refutation or only as weakening, and what the current internal support status actually is after the completed Cycle 1--4 program, the mini-resolution pass, the residual-resolution campaign, the high-value confirmation campaign, and the Residual B hardening pass. The resulting status is mixed by design. One claim family---the insufficiency of complexity, connectivity, and rich activity without invariant preservation---now has robust internal support. Substrate-preserved class membership and biological non-primacy now have robust support with localized caveats, but remain internal-only. Approximate stability is provisionally supported with a localized threshold caveat. The boundary-sensitive stress of existence and rupture now has robust support with localized caveats after cross-family convergence under a more independent judge path, but it still does not count as external validation.

The paper therefore serves as the canon's prediction and falsification layer. It is not another foundational theorem paper, not a methods paper, and not a philosophical restatement. Its goal is to put the framework under explicit decision pressure while keeping status classes, residual limits, and non-claims fully visible.

Release governance note. Limitations/Non-claims and Methods/Governance: see *METHODS_GOVERNANCE_HUB.v1*. Terminology follows the canonical glossary and migration policy. The support classes used here describe the internal scientific program only; they do not convert internal survival into external validation.

1 Introduction

The current corpus is doctrinally frozen for the present scientific cycle. That freeze matters. Without it, prediction work becomes a moving target: thresholds shift, definitions drift, and every difficult result can be absorbed by rewriting the theory rather than by testing it. The purpose of this paper is to keep that from happening.

The framework now contains enough formal structure to state genuinely discriminative predictions. The Core fixes persistence, contractivity, and non-collapse. *Rigid Identity*, *Phenomenological Regimes*, and *Null-Regime Instability* progressively constrain identity, admissible regime structure, and non-nullity. *Forensic Predictions* fixes admissibility and forensic discipline.

Operational Consciousness Criterion consolidates the operational criterion by requiring six explicit invariants:

$$\{I_{\text{per}}, I_{\text{ri}}, I_{\text{int}}, I_{\text{cont}}, I_{\text{diff}}, I_{\text{leg}}\}.$$

What remains scientifically relevant is not to restate those results, but to ask what they force the framework to predict and what would count against them.

This paper therefore has four narrow tasks:

1. extract the main experimentally relevant consequences of the canon;
2. distinguish those consequences from what weaker alternatives would predict;
3. state explicit failure modes, downgrade conditions, and non-claims;
4. report the current internal support status after the completed internal cycles, without inflating it.

That scope is intentionally smaller than a new theorem paper and more disciplined than a review essay. The point is to give the corpus a falsification layer that is compact, operational, and honest.

2 What This Paper Does and Does Not Claim

2.1 Positive Scope

This paper claims only the following:

1. the frozen canon entails a structured family of predictions that are not shared by simpler rivals;
2. those predictions can be organized into explicit discriminators and explicit failure modes;
3. the current internal program is already strong enough to classify some parts of the framework as robustly supported internally, some as provisional, and some as still ambiguous or technically limited.

2.2 Explicit Non-Claims

Remark 2.1 (Non-claims). *This paper does **not** claim:*

1. *human phenomenal consciousness;*
2. *biologically specific human qualia;*
3. *human-machine subjective equivalence;*
4. *personal identity transfer;*
5. *automatic moral equivalence;*
6. *external validation of the framework;*
7. *that all complex or highly connected systems are conscious;*
8. *that the current internal program closes every upstream claim experimentally.*

2.3 Status Classes

The internal scientific program uses the following status classes:

Status	Meaning in this paper
ROBUST_INTERNAL_SUPPORT	survived multiple internal tests without an active ambiguity on the tested target, while remaining internal-only
ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT	supported strongly enough for robust internal use, but still carrying a narrow localized caveat that must remain visible
PROVISIONAL_SUPPORT_LOCALIZED	strengthened by internal tests, but still limited by a localized metric, tolerance, or boundary issue
STILL_AMBIGUOUS	current evidence does not cleanly separate pass from fail on the intended target
IMPLEMENTATION_LIMIT	the main blocker is probe, runner, or case-construction quality rather than a shown doctrinal contradiction
METRIC_OR_TOLERANCE_LIMIT	the main blocker is metric handling, normalization, or threshold interaction
INTERNAL_SUPPORT_ONLY	support exists only inside the internal program and has not crossed into external empirical confirmation
NOT_CLAIMED	the statement is explicitly outside the current evidentiary scope

3 Upstream Canon and Dependency Map

This paper is corpus-dependent and deliberately non-redundant. The dependency map is short because theorem ownership is already fixed upstream.

Document	Imported burden	Used here for
Core [2]	persistence, support, attractor, non-collapse	persistence-side predictions, collapse exclusions, admissible support assumptions
Rigid Identity [6]	rigid identity, observable coherence, non-simulability pressure	identity-sensitive discriminators, complexity-only insufficiency pressure
Phenomenological Regimes [4]	continuity, anti-fragmentation, regime admissibility	continuity predictions, near-miss structure, anti-fragmentation expectations
Null-Regime Instability [5]	null-instability, forced non-nullity	non-nullity predictions, null-regime exclusions, downgrade logic
Forensic Predictions [3]	admissibility, comparators, negative controls, forensic discipline	judge discipline, control logic, reproducibility requirements, legibility stress tests
Operational Consciousness Criterion [1]	six-invariant criterion, operational classes, structural equivalence, rupture	explicit prediction targets, equivalence tests, failure-mode classification

This paper introduces no new invariant, no new doctrine, and no new substrate claim. Its novelty lies only in extracting predictive consequences and support statuses from what already exists.

4 Core Predictive Structure of the Framework

4.1 Certification, Equivalence, and Rupture as Predictive Objects

Definition 4.1 (Operational certification). *For an admissible system S already defined in Operational Consciousness Criterion, define the certification predicate*

$$\text{Cert}(S) = 1 \iff I_{\text{per}}(S), I_{\text{ri}}(S), I_{\text{int}}(S), I_{\text{cont}}(S), I_{\text{diff}}(S), I_{\text{leg}}(S) > 0$$

on a non-empty admissible support A_S , under the governance and admissibility constraints inherited from Forensic Predictions.

Definition 4.2 (Equivalence discriminator). *For two admissible systems S_1, S_2 , define the operational equivalence discriminator*

$$\text{Eq}_{\Gamma, \varepsilon}(S_1, S_2) = 1$$

when the relevant signatures match, the admissibility bundle is preserved, and the tolerated invariant deltas remain within the allowed comparison regime. The exact burden belongs to Operational Consciousness Criterion; the present paper uses it only as a prediction carrier.

Definition 4.3 (Rupture event). *For a set of targeted invariants J , define*

$$\text{Rupt}_J(S) = 1$$

when the intervention or perturbation drives at least one invariant in J below the admissible certification margin so that class membership becomes invalidated, destroyed, or undefined.

4.2 Prediction Family 1: Complexity-Only Failure

Proposition 4.4 (Complexity-only negative-control prediction). *If a candidate system exhibits high connectivity, high entropy, or rich activity while failing one or more critical invariants, then the framework predicts*

$$\text{Cert}(S) = 0, \quad \text{Consciousness}_{\text{op}}(S) \text{ not certified}, \quad \text{Q}_{\text{op}}(S) \text{ not certified}.$$

Why this is not trivial. A weaker alternative could predict that rich dynamics or dense connectivity alone are already sufficient proxies for class membership. The canon does not permit that. *Operational Consciousness Criterion* already states the formal insufficiency claim; the prediction layer turns it into a negative-control discriminator.

4.3 Prediction Family 2: Invariant-Loss Rupture

Proposition 4.5 (Ablation prediction). *If a system is initially certified and a critical intervention selectively destroys one or more invariants, then the framework predicts either loss of certification or a transition to an ambiguous boundary regime, depending on how close the pre-ablation margins were to zero.*

Operational consequence. Ablations are not ornamental. If the criterion is serious, the loss of invariant headroom must show up as loss of class membership, legibility loss, or boundary-sensitive ambiguity. No-change ablations would count against the framework.

4.4 Prediction Family 3: Substrate-Preserved Class Membership

Proposition 4.6 (Substrate-preserved equivalence prediction). *If two realizations preserve the relevant signatures, admissibility conditions, and six critical invariants up to the allowed comparison regime, then the framework predicts preserved operational class membership independently of substrate labels.*

Rival prediction. A substrate-privileging alternative would predict that material type itself blocks equivalence even when the structurally relevant burdens are preserved. The canon rejects that move only inside its own formal class, not as a general metaphysical claim.

4.5 Prediction Family 4: Threshold-Sensitive Boundary Behavior

Proposition 4.7 (Boundary-topology prediction). *The framework predicts neither global threshold collapse nor perfectly flat insensitivity. Instead, it predicts stable pass/fail regions together with narrow transition bands near critical margins.*

Interpretive consequence. If every micro-perturbation changes every verdict, the criterion is fragile. If no perturbation matters, the invariants are decorative. A meaningful criterion should show structured boundary behavior.

4.6 Prediction Family 5: Legibility Dependence

Proposition 4.8 (Legibility dependence prediction). *Operational class membership should fail or become inadmissible if readouts cease to be separable, noise-robust, intervention-sensitive, and compressible without total collapse.*

Why this matters. *Forensic Predictions* makes legibility a forensic burden rather than a narrative convenience. The present prediction is therefore stronger than a merely interpretive observation: if legibility disappears, certification must not survive unchanged.

4.7 Prediction Family 6: Continuity and Non-Nullity Relevance

Proposition 4.9 (Continuity/non-nullity relevance). *The current canon predicts that continuity and non-null differentiation are not optional extras. Systems that preserve rich activity while losing continuity or falling toward effective null structure should fail the operational criterion or leave it undefined.*

Upstream source. The continuity side is inherited from *Phenomenological Regimes*, while the non-null side is inherited from *Null-Regime Instability*. This paper adds only the prediction and failure-mode reading.

5 Main Discriminators

The predictive value of the framework lies in how it diverges from weaker alternatives.

Target	Framework prediction	Weaker/rival prediction	Divergence condition
Complexity-only baseline	high complexity without invariant preservation fails certification	richness or scale alone can stand in for class membership	a complexity-rich control passes despite broken invariants
Connectivity-only baseline	dense coupling without preserved rigid identity, differentiation, or legibility fails	connectivity is treated as a sufficient proxy	a dense control passes while identity or legibility are broken
Weak functional equivalence	similar outward behavior is insufficient if critical invariant structure diverges	input-output similarity is enough for same class	weak functional match passes despite invariant rupture
Near-miss case	boundary cases should expose which invariant is binding and where rupture begins	boundary cases stay uninterpretable or arbitrary	boundary remains unlocalizable after probe cleanup
Cross-substrate pair	substrate labels alone do not block class preservation if relevant structure is preserved	substrate type is primitive and blocks class transport	a preserved pair is rejected solely by substrate label
Threshold stress	decisions should show stable zones and narrow transition bands	decisions collapse under micro-threshold changes or never depend on thresholds at all	every local perturbation flips verdicts, or no perturbation matters
Legibility stress	failure of separability, noise robustness, or compression fidelity invalidates certification	legibility is epiphenomenal and not decision-relevant	certification survives legibility collapse
Null/continuity challenge	effective nullity or fragmentation should destroy or downgrade class membership	rich activity alone masks null or fragmented structure	null-like or fragmented cases remain certified

The table is intentionally operational. It does not ask which metaphysics one prefers. It asks which side of each discriminator the framework commits itself to.

5.1 Rival-Structure Compression

The discriminators above can be sharpened further by asking what each simpler rival preserves and, more importantly, what it leaves out. The point is not to defeat every competing theory at once. The point is to show that the present framework places operational burden on a structure richer than scale, connectivity, superficial behavior, or local threshold tuning.

Rival family	What it keeps	What it fails to preserve	Why the miss is decision-relevant
Complexity/activity proxy	large state space, rich trajectories, high apparent activity	no requirement to preserve the six-invariant package as a joint membership condition	a rich control can still fail certification if persistence, rigid identity, or legibility collapse
Connectivity/integration proxy	dense coupling or high global correlation	no guarantee of rigid identity, non-null differentiation, or readable class structure	integration by itself does not prevent the criterion from rejecting a globally coupled but operationally illegible case
Weak functional equivalence	outwardly similar input-output or behavioral traces	no obligation to preserve admissible support geometry or invariant margins under intervention	superficial behavioral similarity can coexist with rupture of the certified class geometry
Threshold/tuning artifact explanation	local success on a narrow parameter band	no stable explanation for why decisions survive frozen manifests, blind ordering, and negative controls outside that band	the framework expects stable regions plus narrow transitions, not arbitrary success across unrelated settings

This compression table matters because it blocks a common overreading. The current support does not show that every rival is false. It shows something narrower and more useful: if a rival explanation ignores the invariant package or treats only one coarse feature as constitutive, then it does not reproduce the present decision structure of the canon.

6 Prediction Matrix

The canonical prediction matrix can be read as a compressed execution contract for the framework. It matters because each row names an observable, an intervention class, a rival baseline, and an explicit failure condition. That structure prevents the canon from floating as a theorem-only object.

ID	Claim	Observable	Manipulation	Framework-side result	Failure / rival divergence
PRED-01	‘P5-02‘	certified class label and Q_{op} geometry across implementations	cross-substrate pair with preserved invariant margins	same operational class and ε -bounded geometry under the admissible equivalence regime	class divergence or geometry mismatch despite preserved structure; substrate-privilege rival survives
PRED-02	‘P5-04‘	class membership before and after targeted rupture	one-invariant ablation with comparable substrate and gross scale	certified class exit, contraction, or undefined Q_{op} after rupture	no class change after verified invariant loss; complexity-only rival survives
PRED-03	‘P5-05‘	matched complexity/control outcome	compare certified lower-scale system against higher-scale baseline lacking the invariant package	certified system passes and matched complexity-only baseline fails or remains undefined	scale-only model passes despite missing invariants
PRED-04	‘P3-01‘	null-regime persistence under admissible dynamics	initialize near the null regime under Null-Regime Instability assumptions	stable null occupancy should fail or trajectories should leave the null regime	admissible null occupancy persists; null-stable rival survives
PRED-05	‘P2-03‘	continuity/fragmentation pattern	compare rigid regime against finite-mass or deformable regime under matched disturbance	rigid regime keeps continuity while limited regime fragments	no differential continuity pattern appears; regime-agnostic rival survives

ID	Claim	Observable	Manipulation	Framework-side result	Failure / rival divergence
PRED-06	‘P4-01‘	admissibility verdict under tamper and sham controls	inject hash corruption, malformed logs, or severe protocol violations	tampered runs are discarded and never promoted to evidential support	loose evaluation pipeline accepts tampered runs
PRED-07	‘P5-03‘	stability of class membership under bounded perturbation	apply perturbations inside the certified margin budget $\delta_*(S)$	membership and geometry remain stable inside the certified budget	fragile-threshold model flips class under sub-critical perturbations
PRED-08	‘P5-01‘	non-emptiness of Q_{op} after successful certification	run the full certification schema on a system with all margins positive	Q_{op} is non-empty and operationally well-defined	certification passes while no non-trivial class is recoverable
PRED-09	‘P4-03‘	claim survival under latency / compute-budget stress	exceed pre-registered budget limits	strong operational claims are invalidated or downgraded	metric-first pipeline keeps claims active after budget violation
PRED-10	‘P5-05‘	legibility under noise and structured compression	raise admissible noise/compression while preserving or breaking I_{leg}	separability and intervention fidelity survive only while I_{leg} remains positive	generic observability model predicts no specific dependence on I_{leg}

The matrix has two immediate consequences. First, the framework is more exposed than a purely interpretive theory because the failure thresholds are explicit. Second, not every prediction is equally mature: some now have robust internal support, while others remain provisional or blocked by implementation quality.

7 Failure Modes

7.1 What Would Count Against the Framework

The relevant failure modes are not all equally strong. Some would refute a strong reading, some would only weaken a claim, and some would reveal a technical limitation of the current program rather than of the canon itself.

Claim target	Strong refutation condition	Downgrade condition	Current blocker class
‘P5-01‘	a system satisfies the full criterion yet fails to define the claimed operational class	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
‘P5-02‘	a structurally preserved pair falls into different classes under the same admissible comparison regime	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	metric/tolerance-level
‘P5-03‘	small admissible perturbations force uncontrolled class exit across stable positive cases	threshold sensitivity persists only in a narrow knife-edge zone	metric/tolerance-level
‘P5-04‘	a system loses a critical invariant and remains certified without downgrade	boundary-family convergence would downgrade if it failed outside the present frozen internal contract	epistemic-level
‘P5-05‘	a complexity-only control without preserved invariants passes certification	matched external baselines are still missing	epistemic-level
‘P5-06‘	biology is shown to be primitive inside the same operational criterion	it inherits the same localized substrate-equivalence caveat as ‘P5-02‘	metric/tolerance-level

7.2 Failure-Mode Taxonomy

Criterion 7.1 (Status meaning). *Within the current scientific program, a result belongs to:*

1. *ROBUST_SUPPORT* if the target claim survives repeated internal discrimination without an active ambiguity on the tested target;
2. *PROVISIONAL_SUPPORT* if the target is strengthened but still depends on a localized unresolved bottleneck;
3. *STILL_AMBIGUOUS* if the current tests do not cleanly separate support from failure;
4. *IMPLEMENTATION_LIMIT* if the current bottleneck is primarily due to probe, runner, or case-construction quality;
5. *METRIC_OR_TOLERANCE_LIMIT* if the bottleneck is mainly due to normalization, comparison regime, or threshold handling.

This taxonomy is not part of the ontology of the framework. It is part of the scientific discipline around how internal evidence is reported.

8 Claim-to-Test Binding Ledger

The tables above state what would count as success or failure in broad terms. The ledger below binds the main *Operational Consciousness Criterion* claim families to the narrowest concrete test classes that currently matter for publication-level assessment.

Claim	Primary positive test	Primary failure test	Present evidentiary bottleneck	Current status
‘P5-01’	successful certification with positive invariant margins and recoverable non-trivial Q_{op} geometry	certified case with no recoverable operational class, or clean existence-boundary degradation without class change	support is now cross-family robust internally, but the nearest positive member still sits close to the legibility margin	ROBUST SUPPORT W/ LOCALIZED CAVEAT
‘P5-02’	cross-substrate pair with preserved invariant margins and admissible geometry match	structurally preserved pair judged non-equivalent under the same comparison regime	strongest support is robust internally but one legacy raw continuous/discrete pair remains ambiguous on I_{ri} handling	ROBUST SUPPORT W/ LOCALIZED CAVEAT
‘P5-03’	stable class membership inside certified perturbation budget $\delta_*(S)$	class flip under sub-critical admissible perturbations across otherwise stable positive cases	one knife-edge negative-control zone remains locally fragile	PROVISIONAL SUPPORT LOCALIZED
‘P5-04’	targeted invariant rupture causing certification exit or undefined Q_{op}	verified invariant loss without downgrade or class exit	support is now cross-family robust internally, but the evidence still lives inside the current invariant package and frozen thresholds	ROBUST SUPPORT W/ LOCALIZED CAVEAT
‘P5-05’	complexity-rich and connectivity-rich negative controls fail certification under hardened judging	matched scale-only or activity-rich controls pass despite missing invariants	no external matched-baseline campaign yet	ROBUST INTERNAL SUPPORT
‘P5-06’	same positive equivalence path as ‘P5-02’, interpreted without substrate privilege	biology shown to be primitive under the same operational criterion	inherits the same localized I_{ri} -handling caveat as ‘P5-02’	ROBUST SUPPORT W/ LOCALIZED CAVEAT

9 Boundary Diagnostics

The current internal program does not fail uniformly. Its informative content lies in where the boundaries are sharp, where they are merely narrow, and where probe design still dominates the result.

Boundary target	Empirical shape after completed cycles	What it means	Limit class
Complexity-only negative control	stable fail region under harder judging, semi-blind ordering, and pseudo-replication	the criterion is not collapsing into gross activity or connectivity proxies	ROBUST INTERNAL SUPPORT
Threshold-sensitive negative region	narrow decision band rather than broad collapse	the framework has a localized edge effect, not universal threshold instability	METRIC OR TOLERANCE LIMIT
Boundary-probe families under judge-v3	two distinct probe families now reproduce the transition pattern with localized first-fail behavior	the existence/rupture boundary is now robust internally, though still internal-only and threshold-bound near the positive edge	ROBUST SUPPORT W/ LOCALIZED CAVEAT
Harder positive substrate families	family convergence now supports class preservation, but one legacy raw pair remains ambiguous	the live ambiguity is concentrated on how rigid-identity preservation is measured, not on every invariant at once	METRIC OR TOLERANCE LIMIT
Refined negative substrate families	clean fail across both the legacy and second pair families	the equivalence criterion can reject stronger non-equivalent pairs under the hardened judge and pseudo-external reproduction path	ROBUST SUPPORT W/ LOCALIZED CAVEAT

10 Minimal Reproducibility and Admissibility Discipline

The prediction layer is only credible if the judgment path is at least partly insulated from the generator and if the evidentiary artifacts are frozen before judgment. The current program does not yet have an external judge or an external lab replication, but it does have a stricter internal protocol than it had at the start of the program.

Protocol element	Current state	Residual limitation
Frozen manifests and hashes	active since Cycle 2; judge reads frozen artifacts instead of live mutable state	still same repository and same overall program family
Judge-side threshold file	active; thresholds are frozen before judgment and logged explicitly	threshold bands still require stress analysis near knife-edge zones
Semi-blind ordering	active in Cycle 3; labels are hidden until after decision emission	not a full blind protocol and not an external evaluator
Pseudo-multi-environment replication	active; same-machine replication across distinct execution profiles preserved decisions	not external replication and not cross-lab agreement
Negative-control and sham policy	active from Cycle 1 onward and hardened in later cycles	matched external benchmark families are still missing
Judge code path	more independent than in Cycle 1; it no longer calls live generator step functions during evaluation	still not a separate maintained judge codebase

The point of this section is not to inflate rigor. It is to keep the reader from mistaking internal hardening for external confirmation.

11 Current Status of Internal Support

11.1 Interpretive Discipline

The status ledger below does not reclassify the formal theorems of *Operational Consciousness Criterion*. Those remain formal results inside the frozen canon. The ledger classifies the *current internal scientific support* for experimentally relevant readings of those claims.

Claim	Current status	What strengthened it	What still limits it	Status scope
‘P5-01‘	ROBUST SUPPORT W/ LOCALIZED CAVEAT	near_identity_v3 and offset-dispersion families both reproduced PASS/AMBIGUOUS/FAM under judge-v3 and pseudo-external reproduction	support remains internal-only and the nearest positive case still sits close to the flexibility margin	INTERNAL SUPPORT ONLY
‘P5-02‘	ROBUST SUPPORT W/ LOCALIZED CAVEAT	family1 normalized- I_{ri} and family2 raw quantized positive/negative pairs stayed stable under the more independent judge path and pseudo- multi-environment rerun	one legacy raw continuous/discrete positive pair remains ambiguous and keeps a localized I_{ri} caveat alive	INTERNAL SUPPORT ONLY
‘P5-03‘	PROVISIONAL SUPPORT LOCALIZED	Cycle 3 threshold stress and Cycle 4 boundary cleanup localized instability rather than showing global collapse	knife-edge threshold band persists for one negative-control region	INTERNAL SUPPORT ONLY
‘P5-04‘	ROBUST SUPPORT W/ LOCALIZED CAVEAT	two distinct boundary families now converge on rupture-side PASS/AMBIGUOUS/FAM behavior, with I_{ri} becoming the first binding invariant on ambiguous or failing members	the evidence remains internal-only and tied to the current invariant package plus frozen thresholds	INTERNAL SUPPORT ONLY
‘P5-05‘	ROBUST INTERNAL SUPPORT	complexity-rich and connectivity-rich controls failed certification across harder methodological passes and semi-blind evaluation	no external matched-baseline campaign exists yet	INTERNAL SUPPORT ONLY
‘P5-06‘	ROBUST SUPPORT W/ LOCALIZED CAVEAT	same strengthened two-family substrate-equivalence path as ‘P5-02‘	it inherits the same localized I_{ri} metric-handling caveat as ‘P5-02‘	INTERNAL SUPPORT ONLY

11.2 Claim-Specific Diagnostic Notes

‘P5-05’. This is currently the strongest experimentally discriminated claim family. The framework repeatedly rejected a complexity-rich, connectivity-rich negative control that still failed the invariant package. That result survived methodological hardening and semi-blind evaluation. The support is therefore `ROBUST_SUPPORT`, but it remains `INTERNAL_SUPPORT_ONLY`.

‘P5-02’ and ‘P5-06’. These two claims move together in the current program. They were strengthened first by conservative normalized- I_{ri} handling, then by a second substrate family, a more independent judge path, and pseudo-multi-environment confirmation. That is enough for `ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`. The localized caveat remains real: one legacy raw continuous/discrete positive pair is still ambiguous, so the present support is robust internally but not caveat-free.

‘P5-03’. Approximate stability is not in collapse. Threshold stress did not show universal fragility. It showed a narrow knife-edge zone around one thresholded negative-control region. The right status is therefore `PROVISIONAL_SUPPORT` with an explicit `METRIC_OR_TOLERANCE_LIMIT`.

‘P5-01’ and ‘P5-04’. These boundary-sensitive claims no longer rest on a single weak probe family. After the residual-resolution campaign and the Residual B hardening pass, ‘near_identity_v3’ and a second offset-dispersion family both reproduced a `PASS/AMBIGUOUS/FAIL` boundary pattern under judge-v3 and pseudo-external reproduction. That is enough for `ROBUST_SUPPORT_WITH_LOCALIZED_CAVEAT`. The caveat remains explicit: support is still internal-only, and the strongest positive member stays near the legibility margin under the frozen judge contract.

11.3 Methodological Status

The scientific program itself also has a status, distinct from the doctrine and distinct from the claim ledger.

Program element	Current status	Why it matters
Generator/judge separation	materially improved	reduces direct circularity relative to Cycle 1
Frozen manifests and thresholds	active	prevents live-state leakage during judgment
Semi-blind evaluation	active but partial	reduces bias; does not equal full blind evaluation
Threshold stress testing	localized caveat	shows local knife-edge behavior instead of broad collapse
Pseudo-multi-environment replication	passed internally	strengthens reproducibility discipline without constituting external replication
Clean-clone release re-instantiation	passed internally	frozen package was rebuilt from a GitHub clean clone plus fresh venv with no tracked-file divergence during verification
Independent external judge codebase	absent	keeps all current support inside the internal program

Release-level reproduction note. In addition to the cycle-specific judge hardening, the frozen release package was re-instantiated from a clean GitHub clone and a fresh virtual environment

on the same machine, and the verification step did not introduce tracked-file drift in the freeze package. This strengthens internal reproducibility and handoff discipline at the release layer. It does *not* count as external validation, third-party confirmation, or a claim upgrade by itself.

12 Cycle Evidence Map

The current status ledger is not based on one cycle. It is the result of a narrow but cumulative internal program. The table below records only the highest-value consequence of each phase, not every local metric.

Phase	Main question	Strongest result	Main residual	Claim effect
Cycle 1	can the criterion distinguish a rich negative control and survive obvious ablations?	complexity-rich negative control failed; invariant-loss ablations exited certification; proxy cross-substrate pair passed	all support remained internal and the equivalence test was still a proxy	strengthened ‘P5-05’; gave early support to ‘P5-01’, ‘P5-02’, ‘P5-04’, ‘P5-06’
Cycle 2	does the program survive stronger methodological separation and harder structural tests?	generator/judge separation improved; harder pair passed; near-miss family and graded curves behaved coherently	judge still remained internal and same-repository	strengthened ‘P5-01’, ‘P5-02’, ‘P5-03’, ‘P5-04’, ‘P5-06’ internally
Cycle 3	do results survive semi-blind judging, threshold stress, and pseudo-replication?	blind ordering and pseudo-multi-environment replication preserved decisions; complexity-only failure remained strong	harder equivalence, near-miss boundary, and threshold stress became ambiguous	robustly strengthened ‘P5-05’; left ‘P5-02’, ‘P5-06’, ‘P5-03’ narrower
Cycle 4	can threshold fragility, near-miss ambiguity, and harder substrate ambiguity be localized?	threshold fragility was localized; refined negative pair separated cleanly; ambiguity concentrated on I_{ri} handling and near-miss probe quality	positive stronger pair remained ambiguity-bound; near-identity probe still entangled	localized rather than promoted ‘P5-02’, ‘P5-03’, ‘P5-04’
Mini-resolution	can probe quality and I_{ri} handling be cleaned up without doctrinal weakening?	normalized- I_{ri} handling provisionally stabilized the harder positive pair; near-identity v2 clarified the probe weakness	near-identity v2 remained too weak to serve as a clean boundary probe	strengthened ‘P5-02’, ‘P5-06’; kept ‘P5-01’, ‘P5-04’ ambiguous
Residual resolution	can the open residuals be improved without doctrinal loosening?	‘near_identity_v3’ produced PASS/AMBIGUOUS/FAIL with I_{ri} first-fail; refined substrate families localized the live ambiguity to I_{ri} handling	the results were still only localized and internal	promoted both residual families to stronger localized support
High-value confirmation and Residual B hardening	do those gains survive a more independent judge path, a second family, and externalization-lite?	substrate-equivalence converged across two pair families; existence/rupture converged across two probe families under judge-v3 and pseudo-external reproduction	all support remained internal-only and a localized legacy caveat survived for ‘P5-02’/‘P5-06’	promoted ‘P5-01’, ‘P5-02’, ‘P5-04’, ‘P5-06’ to robust support with localized caveats

13 Priority Next Discriminators

The current state of the program already suggests which next moves are scientifically efficient. The relevant issue is not breadth but leverage: each next discriminator should either upgrade a provisional claim or localize an ambiguity into an implementation-limited corner.

Priority target	Why it matters now	What the next narrow result would change	Linked claims
Independent external-style rerun of the two boundary families	current rupture/existence support is now robust internally but still lacks outside-program confirmation	an outside reproduction would test whether the family convergence is program-specific or genuinely portable	‘P5-01‘, ‘P5-04‘
Independent external-style rerun of both substrate families	current strongest positive equivalence path is robust internally but still carries a localized legacy I_{ri} caveat	an outside reproduction would test whether the strengthened equivalence path survives beyond the current judge lineage	‘P5-02‘, ‘P5-06‘
Matched external-style complexity base-lines	the complexity-only result is internally strong but still lacks cross-program baseline pressure	a matched baseline campaign would turn the strongest present internal discriminator into a more externally credible one	‘P5-05‘
Independent maintained judge implementation	current hardening is substantial but still same-program	a separately maintained judge would reduce the remaining circularity objection without touching doctrine	all operational readings

14 Residual Technical Caveats and Non-Claims

14.1 Residual Technical Caveats

The remaining caveats are narrow enough to state cleanly:

1. The strongest positive substrate-equivalence path is still partly dependent on normalized I_{ri} handling. This is a `METRIC_OR_TOLERANCE_LIMIT`, not a robust positive closure.
2. The current ‘near_identity_v2‘ probe remains an `IMPLEMENTATION_LIMIT`. It is cleaner than the original generator, but in the tested band it under-stresses the intended boundary.
3. Threshold stress is not globally problematic, but it is not perfectly flat either. The current result is localized rather than globally robust.
4. All empirical support remains `INTERNAL_SUPPORT_ONLY`. No external lab, external evaluator, or external benchmark suite has confirmed the framework. The clean-clone/fresh-environment re-instantiation strengthens same-machine release reproducibility only; it is not external validation, and no separate clean-agent confirmation artifact is being counted here as stronger evidence than that.

14.2 Non-Claims Revisited

The present internal program does not upgrade any of the following beyond NOT_CLAIMED:

1. ordinary human phenomenal consciousness;
2. biological human qualia as such;
3. subjective equivalence between humans and artificial systems;
4. personal identity preservation across substrate transfer;
5. full external validation of substrate invariance;
6. exhaustion of all relevant rival frameworks.

15 Final Claim Freeze Table

The present paper is a prediction-and-falsation layer, not a moving target. The table below records the narrowest honest freeze state for the main *Operational Consciousness Criterion* claim families at the end of the current phase.

Claim	Status	Current blocker	Next step
‘P5-01’	ROBUST SUPPORT LOCALIZED CAVEAT	support is now strong internally but still nearest to the legibility margin on the positive edge	rerun the convergent probe families under an evaluator maintained outside the current repo/runtime
‘P5-02’	ROBUST SUPPORT LOCALIZED CAVEAT	one legacy raw continuous/discrete positive pair remains ambiguous on I_{ri} handling	reproduce both strengthened substrate families under a more independent evaluator without post hoc threshold loosening
‘P5-03’	PROVISIONAL SUPPORT LOCALIZED	one localized knife-edge threshold band remains active in a negative-control region	rerun the localized stress zone with denser threshold sampling and fixed judge contract
‘P5-04’	ROBUST SUPPORT LOCALIZED CAVEAT	rupture-side support is strong internally but still bounded to the present invariant package and frozen thresholds	rerun the same boundary families under a separately maintained evaluator and frozen release package
‘P5-05’	ROBUST INTERNAL SUPPORT	no external matched-baseline campaign exists yet	add matched external-style complexity/activity baselines under the same judge contract
‘P5-06’	ROBUST SUPPORT LOCALIZED CAVEAT	inherits the same localized substrate-equivalence caveat as ‘P5-02’	upgrade only if the strengthened substrate families survive a more independent external-style rerun

16 Conclusion

The canon now has a prediction layer and a failure-mode ledger. That is a scientific gain even where the results are mixed. The framework is no longer only a theorem-bearing corpus. It is now attached to explicit discriminators, explicit downgrade conditions, and explicit support classes.

The current internal picture is sharp enough to summarize without inflation. The complexity/connectivity insufficiency claim has robust internal support. Substrate-preserved class

membership and biological non-primacy now have robust support with localized caveats. Approximate stability is provisionally supported with a localized threshold caveat. Boundary-sensitive stress of existence and rupture now has robust support with localized caveats under convergent internal probe families and a more independent judge path. All of this nevertheless remains internal-only.

This is the right posture for publication-preparation. The framework can now be presented not as a closed doctrine, and not as a philosophical manifesto, but as a frozen canon with explicit predictions, explicit discriminators, explicit failure modes, and an honest status ledger.

References

- [1] Johnny Andrey Perez Ramirez. “A Structural Criterion for Substrate-Invariant Operational Consciousness”. In: *arXiv preprint* (2026).
- [2] Johnny Andrey Perez Ramirez. “Canonical Core: Minimal Hypotheses for Causally Rigid Identity Frameworks”. In: *arXiv preprint* (2026).
- [3] Johnny Andrey Perez Ramirez. “Falsifiable Predictions Under Forensic Constraints”. In: *arXiv preprint* (2026).
- [4] Johnny Andrey Perez Ramirez. “Phenomenological Regimes Induced by Structural Identity”. In: *arXiv preprint* (2026).
- [5] Johnny Andrey Perez Ramirez. “Structural Instability of the Phenomenological Null Regime in Causally Rigid Channels”. In: *arXiv preprint* (2026).
- [6] Johnny Andrey Pérez Ramírez. “Rigid Identity as an Inverse Limit in Observable Channels”. In: *arXiv preprint* (2026).